

igraph Reference Manual

Gábor Csárdi, Department of Statistics, Harvard University

Tamás Nepusz, Department of Biological Physics, Eötvös Loránd University

Vincent Traag, Centre for Science and Technology Studies, Leiden University

Szabolcs Horvát, Center for Systems Biology Dres-

den, Max Planck Institute for Cell Biology and Genetics

Fabio Zanini, Lowy Cancer Research Centre, University of New South Wales

igraph Reference Manual

by Gábor Csárdi, Tamás Nepusz, Vincent Traag, Szabolcs Horvát, and Fabio Zanini

0.9.4

This manual is for igraph, version 0.9.4.

Copyright (C) 2005-2019 Gábor Csárdi and Tamás Nepusz. Copyright (C) 2020-2021 igraph development team. Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. A copy of the license is included in the section entitled “GNU Free Documentation License”.

Table of Contents

1. Introduction	1
igraph is free software	1
Citing igraph	2
2. Installation	3
Prerequisites	3
Installation	3
General build instructions	3
Specific instructions for Windows	4
Notable configuration options	5
Building the documentation	6
Notes for package maintainers	6
Auto-detection of dependencies	6
Shared and static builds	7
Cross-compiling	7
Additional notes	7
3. Tutorial	8
Compiling programs using igraph	8
Compiling with CMake	8
Compiling without CMake	9
Running the program	9
Creating your first graphs	9
Calculating various properties of graphs	11
4. About igraph graphs, the basic interface	13
The igraph data model	13
The basic interface	14
Graph constructors and destructors	14
Basic query operations	16
Adding and deleting vertices and edges	24
5. Error handling	27
Error handling basics	27
Error handlers	27
igraph_error_handler_t — The type of error handler functions.	27
igraph_error_handler_abort — Abort program in case of error.	28
igraph_error_handler_ignore — Ignore errors.	28
igraph_error_handler_printignore — Print and ignore errors.	28
Error codes	28
igraph_error_type_t — Error code type.	28
igraph_strerror — Textual description of an error.	31
Warning messages	32
igraph_warning_handler_t — The type of igraph warning handler func-	
tions.	32
igraph_set_warning_handler — Installs a warning handler.	32
IGRAPH_WARNING — Triggers a warning.	32
IGRAPH_WARNINGF — Triggers a warning, with printf-like syntax.	33
igraph_warning — Triggers a warning.	33
igraph_warningf — Triggers a warning, printf-like version.	33
igraph_warning_handler_ignore — Ignores all warnings.	34
igraph_warning_handler_print — Prints all warnings to the standard	
error.	34
Advanced topics	35
Writing error handlers	35
Error handling internals	35
Deallocating memory	37
Writing igraph functions with proper error handling	38
Fatal errors	39

Error handling and threads	41
6. Memory (de)allocation	42
igraph_malloc — Allocate memory that can be safely deallocated by igraph func-	
tions.	42
igraph_free — Deallocate memory that was allocated by igraph functions.	42
7. Data structure library: vector, matrix, other data types	44
About template types	44
Vectors	45
About igraph_vector_t objects	45
Constructors and Destructors	45
Initializing elements	47
Accessing elements	48
Vector views	50
Copying vectors	50
Exchanging elements	52
Vector operations	53
Vector comparisons	56
Finding minimum and maximum	59
Vector properties	61
Searching for elements	65
Resizing operations	67
Sorting	71
Set operations on sorted vectors	71
Pointer vectors (igraph_vector_ptr_t)	72
Matrices	78
About igraph_matrix_t objects	78
Matrix constructors and destructors	78
Initializing elements	79
Copying matrices	80
Accessing elements of a matrix	81
Operations on rows and columns	82
Matrix operations	86
Matrix comparisons	90
Combining matrices	92
Finding minimum and maximum	92
Matrix properties	95
Searching for elements	98
Resizing operations	99
Sparse matrices	101
About igraph_spmatrix_t objects	101
Sparse matrix constructors and destructors.	101
Accessing elements of a sparse matrix	102
Iterating over the non-zero elements of a sparse matrix	103
Matrix query operations	105
Matrix operations	107
Printing sparse matrices	109
Sparse matrices, another kind	109
About sparse matrices	109
Creating sparse matrix objects	110
Query properties of a sparse matrix	113
Operations on sparse matrices	119
Operations on sparse matrix iterators	124
Operations that change the internal representation	127
Decompositions and solving linear systems	128
Eigenvalues and eigenvectors	134
Conversion to other data types	135
Writing to a file, or to the screen	137
Stacks	137

igraph_stack_init — Initializes a stack.	137
igraph_stack_destroy — Destroys a stack object.	138
igraph_stack_reserve — Reserve memory.	138
igraph_stack_empty — Decides whether a stack object is empty.	138
igraph_stack_size — Returns the number of elements in a stack.	139
igraph_stack_clear — Removes all elements from a stack.	139
igraph_stack_push — Places an element on the top of a stack.	139
igraph_stack_pop — Removes and returns an element from the top of a stack.	140
igraph_stack_top — Query top element.	140
Double-ended queues	140
igraph_dqueue_init — Initialize a double ended queue (deque).	140
igraph_dqueue_destroy — Destroy a double ended queue.	141
igraph_dqueue_empty — Decide whether the queue is empty.	141
igraph_dqueue_full — Check whether the queue is full.	141
igraph_dqueue_clear — Remove all elements from the queue.	142
igraph_dqueue_size — Number of elements in the queue.	142
igraph_dqueue_head — Head of the queue.	142
igraph_dqueue_back — Tail of the queue.	143
igraph_dqueue_pop — Remove the head.	143
igraph_dqueue_pop_back — Remove the tail	143
igraph_dqueue_push — Appends an element.	143
Maximum and minimum heaps	144
igraph_heap_init — Initializes an empty heap object.	144
igraph_heap_init_array — Build a heap from an array.	144
igraph_heap_destroy — Destroys an initialized heap object.	145
igraph_heap_empty — Decides whether a heap object is empty.	145
igraph_heap_push — Add an element.	145
igraph_heap_top — Top element.	146
igraph_heap_delete_top — Return and removes the top element	146
igraph_heap_size — Number of elements	146
igraph_heap_reserve — Allocate more memory	147
String vectors	147
igraph_strvector_init — Initialize	147
igraph_strvector_copy — Initialization by copying.	147
igraph_strvector_destroy — Free allocated memory	148
STR — Indexing string vectors	148
igraph_strvector_get — Indexing	148
igraph_strvector_set — Set an element	149
igraph_strvector_set2 — Sets an element.	149
igraph_strvector_remove — Removes a single element from a string vector.	150
igraph_strvector_append — Concatenate two string vectors.	150
igraph_strvector_clear — Remove all elements	150
igraph_strvector_resize — Resize	151
igraph_strvector_size — Gives the size of a string vector.	151
igraph_strvector_add — Adds an element to the back of a string vector.	151
Adjacency lists	152
Adjacent vertices	152
Incident edges	155
Lazy adjacency list for vertices	157
Lazy incidence list for edges	159
Partial prefix sum trees	161
igraph_psumtree_init — Initializes a partial prefix sum tree.	161
igraph_psumtree_destroy — Destroys a partial prefix sum tree.	162
igraph_psumtree_size — Returns the size of the tree.	162
igraph_psumtree_get — Retrieves the value corresponding to an item in the tree.	162

igraph_psumtree_sum — Returns the sum of the values of the leaves in the tree.	163
igraph_psumtree_search — Finds an item in the tree, given a value.	163
igraph_psumtree_update — Updates the value associated to an item in the tree.	163
8. Random numbers	165
About random numbers in igraph, use cases	165
The default random number generator	165
igraph_rng_default — Query the default random number generator.	165
igraph_rng_set_default — Set the default igraph random number generator.	165
Creating random number generators	165
igraph_rng_init — Initialize a random number generator.	165
igraph_rng_destroy — Deallocate memory associated with a random number generator.	166
igraph_rng_seed — Set the seed of a random number generator.	166
igraph_rng_min — Query the minimum possible integer for a random number generator.	166
igraph_rng_max — Query the maximum possible integer for a random number generator.	167
igraph_rng_name — Query the type of a random number generator.	167
Generating random numbers	167
igraph_rng_get_integer — Generate an integer random number from an interval.	167
igraph_rng_get_unif — Generate real, uniform random numbers from an interval	168
igraph_rng_get_unif01 — Generate real, uniform random number from the unit interval	168
igraph_rng_get_normal — Normally distributed random numbers	169
igraph_rng_get_geom — Generate geometrically distributed random numbers	169
igraph_rng_get_binom — Generate binomially distributed random numbers	169
igraph_rng_get_gamma — Generate sample from a Gamma distribution	170
Supported random number generators	170
igraph_rngtype_mt19937 — The MT19937 random number generator.	170
igraph_rngtype_glibc2 — The random number generator introduced in GNU libc 2.	171
igraph_rngtype_rand — The old BSD rand/srand random number generator.	171
Use cases	172
Normal (default) use	172
Reproducible simulations	172
Changing the default generator	172
Using multiple generators	172
Example	173
9. Graph generators	174
Deterministic graph generators	174
igraph_create — Creates a graph with the specified edges.	174
igraph_small — Shorthand to create a small graph, giving the edges as arguments.	174
igraph_adjacency — Creates a graph from an adjacency matrix.	175
igraph_weighted_adjacency — Creates a graph from a weighted adjacency matrix.	176
igraph_adjlist — Creates a graph from an adjacency list.	177
igraph_star — Creates a <i>star</i> graph, every vertex connects only to the center. ..	178
igraph_lattice — Arbitrary dimensional square lattices.	178
igraph_ring — Creates a <i>ring</i> graph, a one dimensional lattice.	179

igraph_tree — Creates a tree in which almost all vertices have the same number of children.	180
igraph_full — Creates a full graph (directed or undirected, with or without loops).	181
igraph_full_citation — Creates a full citation graph	181
igraph_realize_degree_sequence — Generates a graph with the given degree sequence.	182
igraph_famous — Create a famous graph by simply providing its name.	184
igraph_lcf — Creates a graph from LCF notation.	186
igraph_lcf_vector — Creates a graph from LCF notation.	187
igraph_from_prufer — Generates a tree from a Prüfer sequence.	187
igraph_atlas — Create a small graph from the “Graph Atlas”.	188
igraph_de_bruijn — Generate a de Bruijn graph.	189
igraph_kautz — Generate a Kautz graph.	189
igraph_extended_chordal_ring — Create an extended chordal ring.	190
Games: randomized graph generators	191
igraph_grg_game — Generates a geometric random graph.	191
igraph_barabasi_game — Generates a graph based on the Barabási-Albert model.	191
igraph_erdos_renyi_game — Generates a random (Erdős-Rényi) graph.	193
igraph_watts_strogatz_game — The Watts-Strogatz small-world model.	194
igraph_rewire_edges — Rewires the edges of a graph with constant probability.	194
igraph_rewire_directed_edges — Rewires the chosen endpoint of directed edges.	195
igraph_degree_sequence_game — Generates a random graph with a given degree sequence.	196
igraph_k_regular_game — Generates a random graph where each vertex has the same degree.	197
igraph_static_fitness_game — Non-growing random graph with edge probabilities proportional to node fitness scores.	198
igraph_static_power_law_game — Generates a non-growing random graph with expected power-law degree distributions.	199
igraph_forest_fire_game — Generates a network according to the “forest fire game”.	200
igraph_rewire — Randomly rewires a graph while preserving the degree distribution.	201
igraph_growing_random_game — Generates a growing random graph.	202
igraph_callaway_traits_game — Simulates a growing network with vertex types.	202
igraph_establishment_game — Generates a graph with a simple growing model with vertex types.	203
igraph_preference_game — Generates a graph with vertex types and connection preferences.	204
igraph_asymmetric_preference_game — Generates a graph with asymmetric vertex types and connection preferences.	205
igraph_recent_degree_game — Stochastic graph generator based on the number of incident edges a node has gained recently.	206
igraph_barabasi_aging_game — Preferential attachment with aging of vertices.	207
igraph_recent_degree_aging_game — Preferential attachment based on the number of edges gained recently, with aging of vertices.	208
igraph_lastcit_game — Simulates a citation network, based on time passed since the last citation.	209
igraph_cited_type_game — Simulates a citation based on vertex types.	210
igraph_citing_cited_type_game — Simulates a citation network based on vertex types.	211

igraph_sbm_game — Sample from a stochastic block model.	212
igraph_hsbm_game — Hierarchical stochastic block model.	212
igraph_hsbm_list_game — Hierarchical stochastic block model, more general version.	213
igraph_dot_product_game — Generates a random dot product graph.	214
igraph_tree_game — Generates a random tree with the given number of nodes.	214
igraph_correlated_game — Generates a random graph correlated to an existing graph.	215
igraph_correlated_pair_game — Generates pairs of correlated random graphs.	216
igraph_simple_interconnected_islands_game — Generates a random graph made of several interconnected islands, each island being a random graph.	216
10. Games on graphs	218
Microscopic update rules	218
igraph_deterministic_optimal_imitation — Adopt a strategy via deterministic optimal imitation.	218
igraph_moran_process — The Moran process in a network setting.	219
igraph_roulette_wheel_imitation — Adopt a strategy via roulette wheel selection.	221
igraph_stochastic_imitation — Adopt a strategy via stochastic imitation with uniform selection.	222
Epidemic models	224
igraph_sir — Performs a number of SIR epidemics model runs on a graph.	224
igraph_sir_t — The result of one SIR model simulation.	224
igraph_sir_destroy — Deallocates memory associated with a SIR simulation run.	225
11. Vertex and edge selectors and sequences, iterators	226
About selectors, iterators	226
Vertex selector constructors	226
igraph_vs_all — Vertex set, all vertices of a graph.	226
igraph_vs_adj — Adjacent vertices of a vertex.	227
igraph_vs_nonadj — Non-adjacent vertices of a vertex.	227
igraph_vs_none — Empty vertex set.	228
igraph_vs_1 — Vertex set with a single vertex.	228
igraph_vs_vector — Vertex set based on a vector.	229
igraph_vs_vector_small — Create a vertex set by giving its elements.	229
igraph_vs_vector_copy — Vertex set based on a vector, with copying.	230
igraph_vs_seq — Vertex set, an interval of vertices.	230
Generic vertex selector operations	231
igraph_vs_copy — Creates a copy of a vertex selector.	231
igraph_vs_destroy — Destroy a vertex set.	231
igraph_vs_is_all — Check whether all vertices are included.	232
igraph_vs_size — Returns the size of the vertex selector.	232
igraph_vs_type — Returns the type of the vertex selector.	232
Immediate vertex selectors	232
igraph_vss_all — All vertices of a graph (immediate version).	232
igraph_vss_none — Empty vertex set (immediate version).	233
igraph_vss_1 — Vertex set with a single vertex (immediate version).	233
igraph_vss_vector — Vertex set based on a vector (immediate version).	234
igraph_vss_seq — An interval of vertices (immediate version).	234
Vertex iterators	234
igraph_vit_create — Creates a vertex iterator from a vertex selector.	234
igraph_vit_destroy — Destroys a vertex iterator.	235
Stepping over the vertices	235
IGRAPH_VIT_NEXT — Next vertex.	236
IGRAPH_VIT_END — Are we at the end?	236

IGRAPH_VIT_SIZE — Size of a vertex iterator.	236
IGRAPH_VIT_RESET — Reset a vertex iterator.	237
IGRAPH_VIT_GET — Query the current position.	237
Edge selector constructors	237
igraph_es_all — Edge set, all edges.	237
igraph_es_incident — Edges incident on a given vertex.	238
igraph_es_none — Empty edge selector.	238
igraph_es_1 — Edge selector containing a single edge.	239
igraph_es_vector — Handle a vector as an edge selector.	239
igraph_es_fromto — Edge selector, all edges between two vertex sets.	240
igraph_es_seq — Edge selector, a sequence of edge ids.	240
igraph_es_pairs — Edge selector, multiple edges defined by their endpoints in a vector.	241
igraph_es_pairs_small — Edge selector, multiple edges defined by their endpoints as arguments.	241
igraph_es_path — Edge selector, edge ids on a path.	242
igraph_es_vector_copy — Edge set, based on a vector, with copying.	242
Immediate edge selectors	243
igraph_ess_all — Edge set, all edges (immediate version)	243
igraph_ess_none — Immediate empty edge selector.	243
igraph_ess_1 — Immediate version of the single edge edge selector.	244
igraph_ess_vector — Immediate vector view edge selector.	244
igraph_ess_seq — Immediate version of the sequence edge selector.	244
Generic edge selector operations	245
igraph_es_as_vector — Transform edge selector into vector.	245
igraph_es_copy — Creates a copy of an edge selector.	245
igraph_es_destroy — Destroys an edge selector object.	246
igraph_es_is_all — Check whether an edge selector includes all edges.	246
igraph_es_size — Returns the size of the edge selector.	246
igraph_es_type — Returns the type of the edge selector.	247
Edge iterators	247
igraph_eit_create — Creates an edge iterator from an edge selector.	247
igraph_eit_destroy — Destroys an edge iterator.	247
Stepping over the edges	248
IGRAPH_EIT_NEXT — Next edge.	248
IGRAPH_EIT_END — Are we at the end?	248
IGRAPH_EIT_SIZE — Number of edges in the iterator.	248
IGRAPH_EIT_RESET — Reset an edge iterator.	249
IGRAPH_EIT_GET — Query an edge iterator.	249
12. Graph, vertex and edge attributes	250
The Attribute Handler Interface	250
igraph_attribute_table_t — Table of functions to perform operations on attributes	250
igraph_set_attribute_table — Attach an attribute table.	252
igraph_attribute_type_t — The possible types of the attributes.	253
Handling attribute combination lists	253
igraph_attribute_combination_init — Initialize attribute combina- tion list.	254
igraph_attribute_combination_add — Add combination record to at- tribute combination list.	254
igraph_attribute_combination_remove — Remove a record from an attribute combination list.	255
igraph_attribute_combination_destroy — Destroy attribute combi- nation list.	255
igraph_attribute_combination_type_t — The possible types of at- tribute combinations.	255
Accessing attributes from C	256
Query attributes	257

Set attributes	270
Remove attributes	282
13. Structural properties of graphs	286
Basic properties	286
igraph_are_connected — Decides whether two vertices are connected	286
(Shortest)-path related functions	286
igraph_shortest_paths — The length of the shortest paths between ver-	
tices.	286
igraph_shortest_paths_dijkstra — Weighted shortest path lengths	
between vertices.	287
igraph_shortest_paths_bellman_ford — Weighted shortest path	
lengths between vertices, allowing negative weights.	288
igraph_shortest_paths_johnson — Weighted shortest path lengths be-	
tween vertices, using Johnson's algorithm.	289
igraph_get_shortest_paths — Shortest paths from a vertex.	290
igraph_get_shortest_path — Shortest path from one vertex to another	
one.	291
igraph_get_shortest_paths_dijkstra — Weighted shortest paths	
from a vertex.	292
igraph_get_shortest_path_dijkstra — Weighted shortest path from	
one vertex to another one.	294
igraph_get_shortest_paths_bellman_ford — Weighted shortest	
paths from a vertex, allowing negative weights.	295
igraph_get_shortest_path_bellman_ford — Weighted shortest path	
from one vertex to another one.	296
igraph_get_all_shortest_paths — All shortest paths (geodesics) from	
a vertex.	297
igraph_get_all_shortest_paths_dijkstra — All weighted shortest	
paths (geodesics) from a vertex.	298
igraph_get_all_simple_paths — List all simple paths from one source. ..	299
igraph_average_path_length — Calculates the average unweighted	
shortest path length between all vertex pairs.	300
igraph_average_path_length_dijkstra — Calculates the average	
weighted shortest path length between all vertex pairs.	301
igraph_path_length_hist — Create a histogram of all shortest path	
lengths.	301
igraph_diameter — Calculates the diameter of a graph (longest geodesic).	302
igraph_diameter_dijkstra — Calculates the weighted diameter of a	
graph using Dijkstra's algorithm.	303
igraph_girth — The girth of a graph is the length of the shortest cycle in it.	304
igraph_eccentricity — Eccentricity of some vertices.	304
igraph_radius — Radius of a graph.	305
Efficiency measures	306
igraph_global_efficiency — Calculates the global efficiency of a net-	
work.	306
igraph_local_efficiency — Calculates the local efficiency around each	
vertex in a network.	306
igraph_average_local_efficiency — Calculates the average local ef-	
ficiency in a network.	307
Neighborhood of a vertex	308
igraph_neighborhood_size — Calculates the size of the neighborhood of	
a given vertex.	308
igraph_neighborhood — Calculate the neighborhood of vertices.	309
igraph_neighborhood_graphs — Create graphs from the neighbor-	
hood(s) of some vertex/vertices.	310
Local scan statistics	311
"Us" statistics	311
"Them" statistics	312

Pre-calculated neighborhoods	314
Graph components	315
igraph_subcomponent — The vertices in the same component as a given vertex.	315
igraph_clusters — Calculates the (weakly or strongly) connected components in a graph.	315
igraph_is_connected — Decides whether the graph is (weakly or strongly) connected.	316
igraph_decompose — Decompose a graph into connected components.	317
igraph_decompose_destroy — Free the memory allocated by igraph_decompose().	317
igraph_biconnected_components — Calculate biconnected components .	318
igraph_articulation_points — Find the articulation points in a graph. ...	319
igraph_bridges — Find all bridges in a graph.	319
Degree sequences	320
igraph_is_graphical — Is there a graph with the given degree sequence? ...	320
igraph_is_bigraphical — Is there a bipartite graph with the given bi-degree-sequence?	321
igraph_is_degree_sequence — Determines whether a degree sequence is valid.	322
igraph_is_graphical_degree_sequence — Determines whether a sequence of integers can be the degree sequence of some simple graph.	323
Centrality measures	323
igraph_closeness — Closeness centrality calculations for some vertices.	323
igraph_harmonic_centrality — Harmonic centrality for some vertices. ...	325
igraph_betweenness — Betweenness centrality of some vertices.	326
igraph_edge_betweenness — Betweenness centrality of the edges.	327
igraph_pagerank_algo_t — PageRank algorithm implementation	327
igraph_pagerank — Calculates the Google PageRank for the specified vertices.	328
igraph_personalized_pagerank — Calculates the personalized Google PageRank for the specified vertices.	329
igraph_personalized_pagerank_vs — Calculates the personalized Google PageRank for the specified vertices.	330
igraph_constraint — Burt's constraint scores.	332
igraph_maxdegree — The maximum degree in a graph (or set of vertices).	332
igraph_strength — Strength of the vertices, weighted vertex degree in other words.	333
igraph_eigenvector_centrality — Eigenvector centrality of the vertices	334
igraph_hub_score — Kleinberg's hub scores.	335
igraph_authority_score — Kleinberg's authority scores.	336
igraph_convergence_degree — Calculates the convergence degree of each edge in a graph.	336
Range-limited centrality measures	337
igraph_closeness_cutoff — Range limited closeness centrality.	337
igraph_harmonic_centrality_cutoff — Range limited harmonic centrality.	338
igraph_betweenness_cutoff — Range-limited betweenness centrality.	339
igraph_edge_betweenness_cutoff — Range-limited betweenness centrality of the edges.	340
Centralization	341
igraph_centralization — Calculate the centralization score from the node level scores	341
igraph_centralization_degree — Calculate vertex degree and graph centralization	342
igraph_centralization_betweenness — Calculate vertex betweenness and graph centralization	343

igraph_centralization_closeness — Calculate vertex closeness and graph centralization	343
igraph_centralization_eigenvector_centrality — Calculate eigenvector centrality scores and graph centralization	344
igraph_centralization_degree_tmax — Theoretical maximum for graph centralization based on degree	345
igraph_centralization_betweenness_tmax — Theoretical maximum for graph centralization based on betweenness	346
igraph_centralization_closeness_tmax — Theoretical maximum for graph centralization based on closeness	347
igraph_centralization_eigenvector_centrality_tmax — Theoretical maximum centralization for eigenvector centrality	348
Similarity measures	349
igraph_bibcoupling — Bibliographic coupling.	349
igraph_cocitation — Cocitation coupling.	349
igraph_similarity_jaccard — Jaccard similarity coefficient for the given vertices.	350
igraph_similarity_jaccard_pairs — Jaccard similarity coefficient for given vertex pairs.	351
igraph_similarity_jaccard_es — Jaccard similarity coefficient for a given edge selector.	352
igraph_similarity_dice — Dice similarity coefficient.	352
igraph_similarity_dice_pairs — Dice similarity coefficient for given vertex pairs.	353
igraph_similarity_dice_es — Dice similarity coefficient for a given edge selector.	354
igraph_similarity_inverse_log_weighted — Vertex similarity based on the inverse logarithm of vertex degrees.	355
Trees	356
igraph_minimum_spanning_tree — Calculates one minimum spanning tree of a graph.	356
igraph_minimum_spanning_tree_unweighted — Calculates one minimum spanning tree of an unweighted graph.	357
igraph_minimum_spanning_tree_prim — Calculates one minimum spanning tree of a weighted graph.	357
igraph_random_spanning_tree — Uniformly sample the spanning trees of a graph	358
igraph_is_tree — Decides whether the graph is a tree.	359
igraph_to_prufer — Converts a tree to its Prüfer sequence	360
Transitivity or clustering coefficient	360
igraph_transitivity_undirected — Calculates the transitivity (clustering coefficient) of a graph.	360
igraph_transitivity_local_undirected — Calculates the local transitivity (clustering coefficient) of a graph.	361
igraph_transitivity_avglocal_undirected — Average local transitivity (clustering coefficient).	362
igraph_transitivity_barrat — Weighted transitivity, as defined by A. Barrat.	363
Directedness conversion	363
igraph_to_directed — Convert an undirected graph to a directed one	363
igraph_to_undirected — Convert a directed graph to an undirected one.	364
Spectral properties	365
igraph_laplacian — Returns the Laplacian matrix of a graph	365
Non-simple graphs: Multiple and loop edges	366
igraph_is_simple — Decides whether the input graph is a simple graph.	366
igraph_is_loop — Find the loop edges in a graph.	366
igraph_is_multiple — Find the multiple edges in a graph.	367

igraph_has_multiple — Check whether the graph has at least one multiple edge.	367
igraph_count_multiple — Count the number of appearances of the edges in a graph.	368
Mixing patterns	368
igraph_assortativity_nominal — Assortativity of a graph based on vertex categories	368
igraph_assortativity — Assortativity based on numeric properties of vertices	369
igraph_assortativity_degree — Assortativity of a graph based on vertex degree	370
K-Cores	371
igraph_coreness — Finding the coreness of the vertices in a network.	371
Topological sorting, directed acyclic graphs	371
igraph_is_dag — Checks whether a graph is a directed acyclic graph (DAG) or not.	371
igraph_topological_sorting — Calculate a possible topological sorting of the graph.	372
igraph_feedback_arc_set — Calculates a feedback arc set of the graph using different	372
Maximum cardinality search and chordal graphs	373
igraph_maximum_cardinality_search — Maximum cardinality search.	373
igraph_is_chordal — Decides whether a graph is chordal.	374
Matchings	375
igraph_is_matching — Checks whether the given matching is valid for the given graph.	375
igraph_is_maximal_matching — Checks whether a matching in a graph is maximal.	376
igraph_maximum_bipartite_matching — Calculates a maximum matching in a bipartite graph.	376
Unfolding a graph into a tree	377
igraph_unfold_tree — Unfolding a graph into a tree, by possibly multiplying its vertices.	377
Other operations	378
igraph_density — Calculate the density of a graph.	378
igraph_reciprocity — Calculates the reciprocity of a directed graph.	379
igraph_diversity — Structural diversity index of the vertices	379
igraph_is_mutual — Check whether the edges of a directed graph are mutual.	380
igraph_avg_nearest_neighbor_degree — Average neighbor degree. ...	381
igraph_get_adjacency — Returns the adjacency matrix of a graph	382
igraph_get_stochastic — Stochastic adjacency matrix of a graph	382
igraph_get_stochastic_sparsemat — Stochastic adjacency matrix of a graph	383
igraph_get_edgelist — Returns the list of edges in a graph	384
14. Graph cycles	385
Eulerian cycles and paths	385
igraph_is_eulerian — Checks whether an Eulerian path or cycle exists	385
igraph_eulerian_cycle — Finds an Eulerian cycle	385
igraph_eulerian_path — Finds an Eulerian path	386
15. Graph visitors	387
Breadth-first search	387
igraph_bfs — Breadth-first search	387
igraph_bfs_simple — Breadth-first search, single-source version	388
igraph_bfshandler_t — Callback type for BFS function	389
Depth-first search	390
igraph_dfs — Depth-first search	390

igraph_dfshandler_t — Callback type for the DFS function	391
Random walks	391
igraph_random_walk — Perform a random walk on a graph	391
igraph_random_edge_walk — Perform a random walk on a graph and re- turn the traversed edges	392
16. Cliques and independent vertex sets	394
Cliques	394
igraph_cliques — Finds all or some cliques in a graph.	394
igraph_clique_size_hist — Counts cliques of each size in the graph.	394
igraph_cliques_callback — Calls a function for each clique in the graph.	395
igraph_clique_handler_t — Type of clique handler functions.	396
igraph_largest_cliques — Finds the largest clique(s) in a graph.	396
igraph_maximal_cliques — Finds all maximal cliques in a graph.	397
igraph_maximal_cliques_count — Count the number of maximal cliques in a graph	398
igraph_maximal_cliques_file — Find maximal cliques and write them to a file.	399
igraph_maximal_cliques_subset — Maximal cliques for a subset of initial vertices	399
igraph_maximal_cliques_hist — Counts the number of maximal cliques of each size in a graph.	400
igraph_maximal_cliques_callback — Finds maximal cliques in a graph and calls a function for each one.	401
igraph_clique_number — Finds the clique number of the graph.	402
Weighted cliques	402
igraph_weighted_cliques — Finds all cliques in a given weight range in a vertex weighted graph.	402
igraph_largest_weighted_cliques — Finds the largest weight clique(s) in a graph.	403
igraph_weighted_clique_number — Finds the weight of the largest weight clique in the graph.	404
Independent vertex sets	404
igraph_independent_vertex_sets — Finds all independent vertex sets in a graph.	404
igraph_largest_independent_vertex_sets — Finds the largest in- dependent vertex set(s) in a graph.	405
igraph_maximal_independent_vertex_sets — Finds all maximal in- dependent vertex sets of a graph.	406
igraph_independence_number — Finds the independence number of the graph.	406
17. Graph isomorphism	408
The simple interface	408
igraph_isomorphic — Decides whether two graphs are isomorphic	408
igraph_subisomorphic — Decide subgraph isomorphism.	409
The BLISS algorithm	409
igraph_bliss_sh_t — Splitting heuristics for Bliss.	410
igraph_bliss_info_t — Information about a BLISS run	410
igraph_canonical_permutation — Canonical permutation using Bliss	411
igraph_isomorphic_bliss — Graph isomorphism via Bliss	411
igraph_automorphisms — Number of automorphisms using Bliss	412
igraph_automorphism_group — Automorphism group generators using Bliss	413
The VF2 algorithm	413
igraph_isomorphic_vf2 — Isomorphism via VF2	414
igraph_count_isomorphisms_vf2 — Number of isomorphisms via VF2 ..	415
igraph_get_isomorphisms_vf2 — Collect all isomorphic mappings of two graphs.	416

igraph_isohandler_t — Callback type, called when an isomorphism was found	417
igraph_isocompat_t — Callback type, called to check whether two vertices or edges are compatible	417
igraph_isomorphic_function_vf2 — The generic VF2 interface	418
igraph_subisomorphic_vf2 — Decide subgraph isomorphism using VF2 ..	419
igraph_count_subisomorphisms_vf2 — Number of subgraph isomorphisms using VF2	420
igraph_get_subisomorphisms_vf2 — Return all subgraph isomorphic mappings.	421
igraph_subisomorphic_function_vf2 — Generic VF2 function for subgraph isomorphism problems	422
The LAD algorithm	424
igraph_subisomorphic_lad — Check subgraph isomorphism with the LAD algorithm	424
Functions for graphs with 3 or 4 vertices	425
igraph_isomorphic_34 — Graph isomorphism for 3-4 vertices	425
igraph_isoclass — Determine the isomorphism class of a graph with 3 or 4 vertices	425
igraph_isoclass_subgraph — The isomorphism class of a subgraph of a graph.	426
igraph_isoclass_create — Creates a graph from the given isomorphism class.	427
Utility functions	427
igraph_permute_vertices — Permute the vertices	427
igraph_simplify_and_colorize — Simplify the graph and compute self-loop and edge multiplicities.	428
18. Graph coloring	429
igraph_vertex_coloring_greedy — Computes a vertex coloring using a greedy algorithm.	429
igraph_coloring_greedy_t — Ordering heuristics for greedy graph coloring.	429
19. Graph motifs, dyad census and triad census	430
igraph_dyad_census — Calculating the dyad census as defined by Holland and Leinhardt.	430
igraph_triad_census — Triad census, as defined by Davis and Leinhardt	430
Finding triangles	432
igraph_adjacent_triangles — Count the number of triangles a vertex is part of.	432
igraph_list_triangles — Find all triangles in a graph.	432
Graph motifs	433
igraph_motifs_randesu — Count the number of motifs in a graph.	433
igraph_motifs_randesu_no — Count the total number of motifs in a graph.	434
igraph_motifs_randesu_estimate — Estimate the total number of motifs in a graph.	434
igraph_motifs_randesu_callback — Finds motifs in a graph and calls a function for each of them.	435
igraph_motifs_handler_t — Callback type for igraph_motifs_randesu_callback	436
20. Generating layouts for graph drawing	437
2D layout generators	437
igraph_layout_random — Places the vertices uniform randomly on a plane.	437
igraph_layout_circle — Places the vertices uniformly on a circle, in the order of vertex ids.	437
igraph_layout_star — Generates a star-like layout.	438
igraph_layout_grid — Places the vertices on a regular grid on the plane.	438

igraph_layout_graphopt — Optimizes vertex layout via the graphopt algorithm.	439
igraph_layout_bipartite — Simple layout for bipartite graphs.	440
The DrL layout generator	440
igraph_layout_fruchterman_reingold — Places the vertices on a plane according to the Fruchterman-Reingold algorithm.	445
igraph_layout_kamada_kawai — Places the vertices on a plane according to the Kamada-Kawai algorithm.	446
igraph_layout_gem — The GEM layout algorithm, as described in Arne Frick, Andreas Ludwig,	447
igraph_layout_davidson_harel — Davidson-Harel layout algorithm	447
igraph_layout_mds — Place the vertices on a plane using multidimensional scaling.	448
igraph_layout_lgl — Force based layout algorithm for large graphs.	449
igraph_layout_reingold_tilford — Reingold-Tilford layout for tree graphs	450
igraph_layout_reingold_tilford_circular — Circular Reingold-Tilford layout for trees	451
igraph_layout_sugiyama — Sugiyama layout algorithm for layered directed acyclic graphs.	452
3D layout generators	453
igraph_layout_random_3d — Places the vertices uniform randomly in a cube.	453
igraph_layout_sphere — Places vertices (more or less) uniformly on a sphere.	454
igraph_layout_grid_3d — Places the vertices on a regular grid in the 3D space.	454
igraph_layout_fruchterman_reingold_3d — 3D Fruchterman-Reingold algorithm.	455
igraph_layout_kamada_kawai_3d — 3D version of the Kamada-Kawai layout generator	456
Merging layouts	457
igraph_layout_merge_dla — Merge multiple layouts by using a DLA algorithm	457
21. Reading and writing graphs from and to files	458
Simple edge list and similar formats	458
igraph_read_graph_edgelist — Reads an edge list from a file and creates a graph.	458
igraph_write_graph_edgelist — Writes the edge list of a graph to a file.	458
igraph_read_graph_ncol — Reads a .ncol file used by LGL.	459
igraph_write_graph_ncol — Writes the graph to a file in .ncol format ..	460
igraph_read_graph_lgl — Reads a graph from an .lgl file	460
igraph_write_graph_lgl — Writes the graph to a file in .lgl format	461
igraph_read_graph_dimacs — Read a graph in DIMACS format.	462
igraph_write_graph_dimacs — Write a graph in DIMACS format.	463
Binary formats	464
igraph_read_graph_graphdb — Read a graph in the binary graph database format.	464
GraphML format	465
igraph_read_graph_graphml — Reads a graph from a GraphML file.	465
igraph_write_graph_graphml — Writes the graph to a file in GraphML format	465
GML format	466
igraph_read_graph_gml — Read a graph in GML format.	466
igraph_write_graph_gml — Write the graph to a stream in GML format	467
Pajek format	468
igraph_read_graph_pajek — Reads a file in Pajek format	468

igraph_write_graph_pajek — Writes a graph to a file in Pajek format.	469
UCINET's DL file format	470
igraph_read_graph_dl — Read a file in the DL format of UCINET	470
Graphviz format	470
igraph_write_graph_dot — Write the graph to a stream in DOT format	470
22. Maximum flows, minimum cuts and related measures	472
Maximum flows	472
igraph_maxflow — Maximum network flow between a pair of vertices	472
igraph_maxflow_value — Maximum flow in a network with the push/relabel algorithm	473
igraph_dominator_tree — Calculates the dominator tree of a flowgraph	474
igraph_maxflow_stats_t — A simple data type to return some statistics from the	475
Cuts and minimum cuts	475
igraph_st_mincut — Minimum cut between a source and a target vertex	475
igraph_st_mincut_value — The minimum s-t cut in a graph	476
igraph_all_st_cuts — List all edge-cuts between two vertices in a directed graph	477
igraph_all_st_mincuts — All minimum s-t cuts of a directed graph	478
igraph_mincut — Calculates the minimum cut in a graph.	479
igraph_mincut_value — The minimum edge cut in a graph	480
igraph_gomory_hu_tree — Gomory-Hu tree of a graph.	480
Connectivity	481
igraph_st_edge_connectivity — Edge connectivity of a pair of vertices	481
igraph_edge_connectivity — The minimum edge connectivity in a graph.	482
igraph_st_vertex_connectivity — The vertex connectivity of a pair of vertices	482
igraph_vertex_connectivity — The vertex connectivity of a graph	483
Edge- and vertex-disjoint paths	484
igraph_edge_disjoint_paths — The maximum number of edge-disjoint paths between two vertices.	484
igraph_vertex_disjoint_paths — Maximum number of vertex-disjoint paths between two vertices.	485
Graph adhesion and cohesion	485
igraph_adhesion — Graph adhesion, this is (almost) the same as edge connectivity.	485
igraph_cohesion — Graph cohesion, this is the same as vertex connectivity. ..	486
Cohesive blocks	487
igraph_cohesive_blocks — Identifies the hierarchical cohesive block structure of a graph	487
23. Vertex separators	488
igraph_is_separator — Would removing this set of vertices disconnect the graph?	488
igraph_is_minimal_separator — Decides whether a set of vertices is a minimal separator.	488
igraph_all_minimal_st_separators — List all vertex sets that are minimal (s,t) separators for some s and t.	489
igraph_minimum_size_separators — Find all minimum size separating vertex sets.	489
igraph_even_tarjan_reduction — Even-Tarjan reduction of a graph	490
24. Detecting community structure	492
Common functions related to community structure	492
igraph_modularity — Calculate the modularity of a graph with respect to some clusters or vertex types.	492
igraph_modularity_matrix — Calculate the modularity matrix	493

igraph_community_optimal_modularity — Calculate the community structure with the highest modularity value	494
igraph_community_to_membership — Create membership vector from community structure dendrogram	495
igraph_reindex_membership — Makes the IDs in a membership vector continuous	496
igraph_compare_communities — Compares community structures using various metrics	496
igraph_split_join_distance — Calculates the split-join distance of two community structures	497
Community structure based on statistical mechanics	498
igraph_community_spinglass — Community detection based on statistical mechanics	498
igraph_community_spinglass_single — Community of a single node based on statistical mechanics	500
Community structure based on eigenvectors of matrices	501
igraph_community_leading_eigenvector — Leading eigenvector community finding (proper version).	502
igraph_community_leading_eigenvector_callback_t — Callback for the leading eigenvector community finding method.	504
igraph_le_community_to_membership — Vertex membership from the leading eigenvector community structure	504
Walktrap: Community structure based on random walks	505
igraph_community_walktrap — This function is the implementation of the Walktrap community	505
Edge betweenness based community detection	506
igraph_community_edge_betweenness — Community finding based on edge betweenness.	506
igraph_community_eb_get_merges — Calculating the merges, i.e. the dendrogram for an edge betweenness community structure.	508
Community structure based on the optimization of modularity	509
igraph_community_fastgreedy — Finding community structure by greedy optimization of modularity.	509
igraph_community_multilevel — Finding community structure by multi-level optimization of modularity.	510
igraph_community_leiden — Finding community structure using the Leiden algorithm.	511
Fluid communities	512
igraph_community_fluid_communities — Community detection based on fluids interacting on the graph.	512
Label propagation	513
igraph_community_label_propagation — Community detection based on label propagation.	513
The InfoMAP algorithm	514
igraph_community_infomap — Find community structure that minimizes the expected	514
25. Graphlets	516
Introduction	516
Performing graphlet decomposition	516
igraph_graphlets — Calculate graphlets basis and project the graph on it	516
igraph_graphlets_candidate_basis — Calculate a candidate graphlets basis	517
igraph_graphlets_project — Project a graph on a graphlets basis	517
26. Hierarchical random graphs	519
Introduction	519
Representing HRGs	519
igraph_hrg_t — Data structure to store a hierarchical random graph	519
igraph_hrg_init — Allocate memory for a HRG.	520

igraph_hrg_destroy — Deallocate memory for an HRG.	520
igraph_hrg_size — Returns the size of the HRG, the number of leaf nodes. ...	520
igraph_hrg_resize — Resize a HRG.	521
Fitting HRGs	521
igraph_hrg_fit — Fit a hierarchical random graph model to a network	521
igraph_hrg_consensus — Calculate a consensus tree for a HRG.	521
HRG sampling	522
igraph_hrg_sample — Sample from a hierarchical random graph model	522
igraph_hrg_game — Generate a hierarchical random graph	523
Conversion to and from igraph graphs	523
igraph_hrg_dendrogram — Create a dendrogram from a hierarchical random graph.	523
igraph_hrg_create — Create a HRG from an igraph graph.	524
Predicting missing edges	524
igraph_hrg_predict — Predict missing edges in a graph, based on HRG models	524
27. Spectral coarse graining	526
Introduction	526
SCG in brief	526
Functions for performing SCG	527
References	527
SCG functions	527
igraph_scg_adjacency — Spectral coarse graining, symmetric case.	527
igraph_scg_stochastic — Spectral coarse graining, stochastic case.	529
igraph_scg_laplacian — Spectral coarse graining, Laplacian case.	531
igraph_scg_grouping — SCG problem solver.	533
igraph_scg_semiprojectors — Compute SCG semi-projectors for a given partition.	535
igraph_scg_norm_eps — Calculate SCG residuals.	536
28. Embedding of graphs	538
Spectral embedding	538
igraph_adjacency_spectral_embedding — Adjacency spectral embedding	538
igraph_laplacian_spectral_embedding — Spectral embedding of the Laplacian of a graph	539
igraph_dim_select — Dimensionality selection	540
29. Graph operators	542
Union and intersection	542
igraph_disjoint_union — Creates the union of two disjoint graphs	542
igraph_disjoint_union_many — The disjoint union of many graphs.	542
igraph_union — Calculates the union of two graphs.	543
igraph_union_many — Creates the union of many graphs.	544
igraph_intersection — Collect the common edges from two graphs.	544
igraph_intersection_many — The intersection of more than two graphs. ..	545
Other set-like operators	546
igraph_difference — Calculate the difference of two graphs	546
igraph_complementer — Create the complementer of a graph	546
igraph_compose — Calculates the composition of two graphs	547
Miscellaneous operators	548
igraph_connect_neighborhood — Connects every vertex to its neighborhood	548
igraph_contract_vertices — Replace multiple vertices with a single one.	548
igraph_induced_subgraph — Creates a subgraph induced by the specified vertices.	549
igraph_linegraph — Create the line graph of a graph.	550
igraph_simplify — Removes loop and/or multiple edges from the graph.	550

igraph_subgraph_edges — Creates a subgraph with the specified edges and their endpoints.	551
30. Using BLAS, LAPACK and ARPACK for igraph matrices and graphs	552
BLAS interface in igraph	552
igraph_blas_ddot — Dot product of two vectors.	552
igraph_blas_dgemv — Matrix-vector multiplication using BLAS, vector version.	552
igraph_blas_dgemv_array — Matrix-vector multiplication using BLAS, array version.	553
LAPACK interface in igraph	553
Matrix factorization, solving linear systems	554
Eigenvalues and eigenvectors of matrices	555
ARPACK interface in igraph	560
Data structures	560
ARPACK solvers	566
31. Bipartite, i.e. two-mode graphs	569
Bipartite networks in igraph	569
Create two-mode networks	569
igraph_create_bipartite — Create a bipartite graph.	569
igraph_full_bipartite — Create a full bipartite network.	570
igraph_bipartite_game — Generate a bipartite random graph (similar to Erdős-Rényi).	570
Incidence matrices	571
igraph_incidence — Creates a bipartite graph from an incidence matrix.	571
igraph_get_incidence — Convert a bipartite graph into an incidence matrix.	572
Project two-mode graphs	573
igraph_bipartite_projection_size — Calculate the number of vertices and edges in the bipartite projections.	573
igraph_bipartite_projection — Create one or both projections of a bipartite (two-mode) network.	574
Other operations on bipartite graphs	575
igraph_is_bipartite — Check whether a graph is bipartite.	575
32. Advanced igraph programming	576
Using igraph in multi-threaded programs	576
IGRAPH_THREAD_SAFE — Specifies whether igraph was built in thread-safe mode.	576
Thread-safe ARPACK library	576
Thread-safety of random number generators	576
Progress handlers	576
About progress handlers	576
Setting up progress handlers	577
Invoking the progress handler	578
Writing progress handlers	579
Writing igraph functions with progress reporting	580
Multi-threaded programs	580
Status handlers	580
Status reporting	580
Setting up status handlers	580
Invoking the status handler	581
33. Non-graph related functions	584
igraph version number	584
igraph_version — Return the version of the igraph C library	584
Running mean of a time series	584
igraph_running_mean — Calculates the running mean of a vector.	584
Random sampling from very long sequences	585
igraph_random_sample — Generates an increasing random sequence of integers.	585

Random sampling of spatial points	586
igraph_sample_sphere_surface — Sample points uniformly from the surface of a sphere	586
igraph_sample_sphere_volume — Sample points uniformly from the volume of a sphere	586
igraph_sample_dirichlet — Sample points from a Dirichlet distribution ..	587
Convex hull of a set of points on a plane	588
igraph_convex_hull — Determines the convex hull of a given set of points in the 2D plane	588
Fitting power-law distributions to empirical data	588
igraph_plfit_result_t — Result of fitting a power-law distribution to a vector	588
igraph_power_law_fit — Fits a power-law distribution to a vector of num- bers	589
34. Licenses for igraph and this manual	591
THE GNU GENERAL PUBLIC LICENSE	591
Preamble	591
GNU GENERAL PUBLIC LICENSE	591
How to Apply These Terms to Your New Programs	594
The GNU Free Documentation License	595
0. PREAMBLE	595
1. APPLICABILITY AND DEFINITIONS	596
2. VERBATIM COPYING	597
3. COPYING IN QUANTITY	597
4. MODIFICATIONS	597
5. COMBINING DOCUMENTS	599
6. COLLECTIONS OF DOCUMENTS	599
7. AGGREGATION WITH INDEPENDENT WORKS	599
8. TRANSLATION	599
9. TERMINATION	600
10. FUTURE REVISIONS OF THIS LICENSE	600
G.1.1 ADDENDUM: How to use this License for your documents	600
Index	601

List of Examples

4.1. File examples/simple/igraph_empty.c	14
4.2. File examples/simple/igraph_copy.c	15
4.3. File examples/simple/igraph_get_eid.c	19
4.4. File examples/simple/igraph_get_eids.c	20
4.5. File examples/simple/igraph_neighbors.c	22
4.6. File examples/simple/igraph_is_directed.c	23
4.7. File examples/simple/igraph_degree.c	24
4.8. File examples/simple/igraph_add_edges.c	25
4.9. File examples/simple/igraph_add_vertices.c	25
4.10. File examples/simple/igraph_delete_edges.c	26
4.11. File examples/simple/igraph_delete_vertices.c	26
6.1. File examples/simple/igraph_free.c	42
7.1. File examples/simple/igraph_fisher_yates_shuffle.c	53
7.2. File examples/simple/igraph_vector_ptr_sort.c	58
7.3. File examples/simple/igraph_vector_ptr_sort.c	59
7.4. File examples/simple/igraph_vector_ptr_sort.c	77
7.5. File examples/simple/igraph_sparsemat.c	110
7.6. File examples/simple/igraph_sparsemat3.c	110
7.7. File examples/simple/igraph_sparsemat4.c	110
7.8. File examples/simple/igraph_sparsemat6.c	110
7.9. File examples/simple/igraph_sparsemat7.c	110
7.10. File examples/simple/igraph_sparsemat8.c	110
7.11. File examples/simple/dqueue.c	140
7.12. File examples/simple/igraph_strvector.c	147
7.13. File examples/simple/adjlist.c	152
8.1. File examples/simple/random_seed.c	173
9.1. File examples/simple/igraph_create.c	174
9.2. File examples/simple/igraph_small.c	175
9.3. File examples/simple/igraph_adjacency.c	176
9.4. File examples/simple/igraph_weighted_adjacency.c	177
9.5. File examples/simple/igraph_star.c	178
9.6. File examples/simple/igraph_ring.c	180
9.7. File examples/simple/igraph_tree.c	180
9.8. File examples/simple/igraph_full.c	181
9.9. File examples/simple/igraph_lcf.c	187
9.10. File examples/simple/igraph_atlas.c	189
9.11. File examples/simple/igraph_grg_game.c	191
9.12. File examples/simple/igraph_barabasi_game.c	193
9.13. File examples/simple/igraph_barabasi_game2.c	193
9.14. File examples/simple/igraph_erdos_renyi_game.c	193
9.15. File examples/simple/igraph_degree_sequence_game.c	197
10.1. File examples/simple/igraph_deterministic_optimal_imitation.c	219
10.2. File examples/simple/igraph_roulette_wheel_imitation.c	222
10.3. File examples/simple/igraph_stochastic_imitation.c	224
11.1. File examples/simple/igraph_vs_nonadj.c	228
11.2. File examples/simple/igraph_vs_vector.c	229
11.3. File examples/simple/igraph_vs_seq.c	231
11.4. File examples/simple/igraph_es_pairs.c	241
12.1. File examples/simple/cattributes.c	257
12.2. File examples/simple/cattributes2.c	257
12.3. File examples/simple/cattributes3.c	257
12.4. File examples/simple/cattributes4.c	257
13.1. File examples/simple/dijkstra.c	288
13.2. File examples/simple/bellman_ford.c	289
13.3. File examples/simple/igraph_get_shortest_paths.c	291

13.4. File examples/simple/igraph_get_shortest_paths_dijkstra.c	294
13.5. File examples/simple/igraph_get_all_shortest_paths_dijkstra.c ..	299
13.6. File examples/simple/igraph_average_path_length.c	300
13.7. File examples/simple/igraph_grg_game.c	301
13.8. File examples/simple/igraph_diameter.c	303
13.9. File examples/simple/igraph_girth.c	304
13.10. File examples/simple/igraph_eccentricity.c	305
13.11. File examples/simple/igraph_radius.c	306
13.12. File examples/simple/igraph_decompose.c	317
13.13. File examples/simple/igraph_biconnected_components.c	319
13.14. File examples/simple/igraph_pagerank.c	329
13.15. File examples/simple/eigenvector_centrality.c	335
13.16. File examples/simple/centralization.c	342
13.17. File examples/simple/igraph_cocitation.c	349
13.18. File examples/simple/igraph_cocitation.c	350
13.19. File examples/simple/igraph_similarity.c	351
13.20. File examples/simple/igraph_similarity.c	351
13.21. File examples/simple/igraph_similarity.c	352
13.22. File examples/simple/igraph_similarity.c	353
13.23. File examples/simple/igraph_similarity.c	354
13.24. File examples/simple/igraph_similarity.c	355
13.25. File examples/simple/igraph_similarity.c	356
13.26. File examples/simple/igraph_minimum_spanning_tree.c	357
13.27. File examples/simple/igraph_minimum_spanning_tree.c	358
13.28. File examples/simple/igraph_tree.c	360
13.29. File examples/simple/igraph_transitivity.c	361
13.30. File examples/simple/igraph_to_undirected.c	365
13.31. File examples/simple/igraph_laplacian.c	366
13.32. File examples/simple/igraph_is_loop.c	367
13.33. File examples/simple/igraph_is_multiple.c	367
13.34. File examples/simple/igraph_has_multiple.c	368
13.35. File examples/simple/assortativity.c	369
13.36. File examples/simple/assortativity.c	370
13.37. File examples/simple/assortativity.c	371
13.38. File examples/simple/igraph_topological_sorting.c	372
13.39. File examples/simple/igraph_feedback_arc_set.c	373
13.40. File examples/simple/igraph_feedback_arc_set_ip.c	373
13.41. File examples/simple/igraph_maximum_bipartite_matching.c	376
13.42. File examples/simple/igraph_maximum_bipartite_matching.c	376
13.43. File examples/simple/igraph_maximum_bipartite_matching.c	377
13.44. File examples/simple/igraph_reciprocity.c	379
13.45. File examples/simple/igraph_knn.c	382
15.1. File examples/simple/igraph_bfs.c	388
15.2. File examples/simple/igraph_bfs_callback.c	388
15.3. File examples/simple/igraph_bfs_simple.c	389
16.1. File examples/simple/igraph_cliques.c	394
16.2. File examples/simple/igraph_maximal_cliques.c	398
16.3. File examples/simple/igraph_maximal_cliques.c	399
16.4. File examples/simple/igraph_independent_sets.c	405
17.1. File examples/simple/igraph_isomorphic_vf2.c	415
17.2. File examples/simple/igraph_subisomorphic_lad.c	425
18.1. File examples/simple/igraph_coloring.c	429
19.1. File examples/simple/igraph_motifs_randesu.c	434
19.2. File examples/simple/igraph_motifs_randesu.c	436
20.1. File examples/simple/igraph_layout_reingold_tilford.c	451
21.1. File examples/simple/igraph_read_graph_lgl.c	461
21.2. File examples/simple/igraph_write_graph_lgl.c	462
21.3. File examples/simple/igraph_read_graph_graphdb.c	464

21.4. File examples/simple/graphml.c	465
21.5. File examples/simple/graphml.c	466
21.6. File examples/simple/gml.c	467
21.7. File examples/simple/gml.c	468
21.8. File examples/simple/foreign.c	469
21.9. File examples/simple/igraph_write_graph_pajek.c	470
21.10. File examples/simple/igraph_read_graph_dl.c	470
21.11. File examples/simple/dot.c	471
22.1. File examples/simple/flow.c	473
22.2. File examples/simple/flow2.c	473
22.3. File examples/simple/dominator_tree.c	475
22.4. File examples/simple/igraph_all_st_mincuts.c	479
22.5. File examples/simple/igraph_mincut.c	480
22.6. File examples/simple/cohesive_blocks.c	487
23.1. File examples/simple/igraph_is_separator.c	488
23.2. File examples/simple/igraph_is_minimal_separator.c	489
23.3. File examples/simple/igraph_minimal_separators.c	489
23.4. File examples/simple/igraph_minimum_size_separators.c	490
23.5. File examples/simple/even_tarjan.c	491
24.1. File examples/simple/igraph_community_optimal_modularity.c	495
24.2. File examples/simple/igraph_community_leading_eigenvector.c	501
24.3. File examples/simple/walktrap.c	506
24.4. File examples/simple/igraph_community_edge_betweenness.c	507
24.5. File examples/simple/igraph_community_fastgreedy.c	510
24.6. File examples/simple/igraph_community_multilevel.c	511
24.7. File examples/simple/igraph_community_leiden.c	512
24.8. File examples/simple/igraph_community_fluid_communities.c	513
24.9. File examples/simple/igraph_community_label_propagation.c	514
27.1. File examples/simple/scg.c	529
27.2. File examples/simple/igraph_scg_grouping.c	535
27.3. File examples/simple/igraph_scg_grouping2.c	535
27.4. File examples/simple/igraph_scg_grouping3.c	535
27.5. File examples/simple/igraph_scg_grouping4.c	535
27.6. File examples/simple/igraph_scg_semiprojectors.c	536
27.7. File examples/simple/igraph_scg_semiprojectors2.c	536
27.8. File examples/simple/igraph_scg_semiprojectors3.c	536
29.1. File examples/simple/igraph_disjoint_union.c	542
29.2. File examples/simple/igraph_union.c	544
29.3. File examples/simple/igraph_union.c	544
29.4. File examples/simple/igraph_intersection.c	545
29.5. File examples/simple/igraph_difference.c	546
29.6. File examples/simple/igraph_complementer.c	547
29.7. File examples/simple/igraph_compose.c	548
29.8. File examples/simple/igraph_simplify.c	551
30.1. File examples/simple/blas.c	552
30.2. File examples/simple/blas.c	553
30.3. File examples/simple/igraph_lapack_dgesv.c	555
30.4. File examples/simple/igraph_lapack_dsyeivr.c	557
30.5. File examples/simple/igraph_lapack_dgeev.c	558
30.6. File examples/simple/igraph_lapack_dgeevx.c	560
31.1. File examples/simple/igraph_bipartite_create.c	569
31.2. File examples/simple/igraph_bipartite_projection.c	574
31.3. File examples/simple/igraph_bipartite_projection.c	575
33.1. File examples/simple/igraph_version.c	584
33.2. File examples/simple/igraph_random_sample.c	586
33.3. File examples/simple/igraph_convex_hull.c	588
33.4. File examples/simple/igraph_power_law_fit.c	590

Chapter 1. Introduction

igraph is a library for creating and manipulating graphs. You can look at it in two ways: first, igraph contains the implementation of quite a lot of graph algorithms. These include classic graph algorithms like graph isomorphism, graph girth and connectivity and also the new wave graph algorithms like transitivity, graph motifs and community structure detection. Skim through the table of contents or the index of this book to get an impression of what is available.

Second, igraph provides a platform for developing and/or implementing graph algorithms. It has an efficient data structure for representing graphs, and a number of other data structures like flexible vectors, stacks, heaps, queues, adjacency lists that are useful for implementing graph algorithms. In fact these data structures evolved along with the implementation of the classic and non-classic graph algorithms which make up the major part of the igraph library. This way, they were fine-tuned and checked for correctness several times.

Our main goal with developing igraph was to create a graph library which is efficient on large, but not extremely large graphs. More precisely, it is assumed that the graph(s) fit into the physical memory of the computer. Nowadays this means graphs with several million vertices and/or edges. Our definition of efficient is that it runs fast, both in theory and (more importantly) in practice.

We believe that one of the big strengths of igraph is that it can be embedded into a higher-level language or environment. Three such embeddings (or interfaces if you look at them another way) are currently being developed by us: an R package, a Python extension module, and a Mathematica (Wolfram Language) package. Others are likely to come. High level languages such as R or Python make it possible to use graph routines with much greater comfort, without actually writing a single line of C code. They have some, usually very small, speed penalty compared to the C version, but add ease of use and much flexibility. This manual, however, covers only the C library. If you want to use Python, R or the Wolfram Language, please see the documentation written specifically for these interfaces and come back here only if you are interested in some detail which is not covered in those documents.

We still consider igraph as a child project. It has much room for development and we are sure that it will improve a lot in the near future. Any feedback we can get from the users is very important for us, as most of the time these questions and comments guide us in what to add and what to improve.

igraph is open source and distributed under the terms of the GNU GPL. We strongly believe that all the algorithms used in science, let that be graph theory or not, should have an efficient open-source implementation allowing use and modification for anyone.

igraph is free software

igraph library

Copyright (C) 2003-2012 Gábor Csardi <csardi.gabor@gmail.com> 334 Harvard st, Cambridge MA, 02139, USA

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

Citing igraph

To cite igraph in publications, please use the following reference:

Gábor Csárdi, Tamás Nepusz: The igraph software package for complex network research. *InterJournal Complex Systems*, 1695, 2006.

The igraph C library is assigned the DOI 10.5281/zenodo.3630268 [<https://doi.org/10.5281/zenodo.3630268>] on Zenodo.

Chapter 2. Installation

Prerequisites

To build igraph from sources, you will need at least:

- CMake [<https://cmake.org>] 3.16 or later
- C and C++ compilers

Visual Studio 2015 and later are supported. Earlier Visual Studio versions may or may not work.

Certain features also require the following libraries:

- libxml2 [<http://www.xmlsoft.org/>], required for GraphML support

igraph bundles a number of libraries for convenience. However, it is preferable to use external versions of these libraries, which may improve performance. These are:

- GMP [<https://gmplib.org/>] (the bundled alternative is Mini-GMP)
- GLPK [<https://www.gnu.org/software/glpk/>]
- ARPACK [<https://www.caam.rice.edu/software/ARPACK/>]
- CXSparse from SuiteSparse [<https://people.engr.tamu.edu/davis/suitesparse.html>]
- A library providing a BLAS [<https://www.netlib.org/blas/>] API (available by default on macOS; OpenBLAS [<https://openblas.net>] is one option on other systems)
- A library providing a LAPACK [<https://www.netlib.org/lapack/>] API (available by default on macOS; OpenBLAS [<https://openblas.net>] is one option on other systems)

When building the development version of igraph, `bison`, `flex` and `git` are also required. Released versions do not require these tools.

To run the tests, `diff` is also required.

Installation

General build instructions

igraph uses a CMake-based build system [<https://cmake.org/cmake/help/latest/guide/user-interaction/index.html>]. To compile it,

- Enter the directory where the igraph sources are:

```
$ cd igraph
```

- Create a new directory. This is where igraph will be built:

```
$ mkdir build
$ cd build
```

- Run CMake, which will automatically configure igraph, and report the configuration:

```
$ cmake ..
```

To set a non-default installation location, such as `/opt/local`, use:

```
cmake .. -DCMAKE_INSTALL_PREFIX=/opt/local
```

- Check the output carefully, and ensure that all features you need are enabled. If CMake could not find certain libraries, some features such as GraphML support may have been automatically disabled.
- There are several ways to adjust the configuration:
 - Run `ccmake .` on Unix-like systems or `cmake-gui` on Windows for a convenient interface.
 - Simply edit the `CMakeCache.txt` file. Some of the relevant options are listed below.
- Once the configuration has been adjusted, run `cmake ..` again.
- Once `igraph` has been successfully configured, it can be built, tested and installed using:

```
$ cmake --build .  
$ cmake --build . --target check  
$ cmake --install .
```

Specific instructions for Windows

Microsoft Visual Studio

With Visual Studio, the steps to build `igraph` are generally the same as above. However, since the Visual Studio CMake generator is a multi-configuration one, we must specify the configuration (typically Release or Debug) with each build command using the `--config` option:

```
mkdir build  
cd build  
cmake ..  
cmake --build . --config Release  
cmake --build . --target check --config Release
```

When building the development version, `bison` and `flex` must be available on the system. `winflexbison` [<https://github.com/lexxmark/winflexbison>] for Bison version 3.x can be useful for this purpose—make sure that the executables are in the system `PATH`. The easiest installation option is probably by installing `winflexbison3` from the Chocolatey package manager [<https://chocolatey.org/packages/winflexbison3>].

vcpkg

Most external dependencies can be conveniently installed using `vcpkg` [<https://github.com/microsoft/vcpkg#quick-start-windows>]. Note that `igraph` bundles all dependencies except `libxml2`, which is needed for GraphML support.

In order to use `vcpkg` integrate it in the build environment by executing `vcpkg.exe integrate install` on the command line. When configuring `igraph`, point CMake to the correct `vcpkg.cmake` file using `-DCMAKE_TOOLCHAIN_FILE=...`, as instructed.

Additionally, it might be that you need to set the appropriate so-called triplet using `-DVCPKG_TARGET_TRIPLET` when running `cmake`, for example, setting it to `x64-windows` when using shared

builds of packages or `x64-windows-static` when using static builds. Similarly, you also need to specify this target triplet when installing packages. For example, to install `libxml2` as a shared library, use `vcpkg.exe install libxml2:x64-windows` and to install `libxml2` as a static library, use `vcpkg.exe install libxml2:x64-windows-static`. In addition, there is the possibility to use a static library with dynamic runtime linking using the `x64-windows-static-md` triplet.

There are some known issues with `igraph` when using certain external packages from `vcpkg`. When building against `OpenBLAS`, this results in a few differences in some unit tests, see issue #1491 [<https://github.com/igraph/igraph/issues/1491>].

MSYS2

MSYS2 can be installed from www.msys2.org. After installing MSYS2, ensure that it is up to date by opening a terminal and running `pacman -Syuu`.

The instructions below assume that you want to compile for a 64-bit target.

Install the following packages using `pacman -S`.

- Minimal requirements: `mingw-w64-x86_64-toolchain`, `mingw-w64-x86_64-cmake`.
- Optional dependencies that enable certain features: `mingw-w64-x86_64-gmp`, `mingw-w64-x86_64-libxml2`
- Optional external libraries for better performance: `mingw-w64-x86_64-openblas`, `mingw-w64-x86_64-suitesparse`, `mingw-w64-x86_64-arpack`, `mingw-w64-x86_64-glpk`
- Only needed for running the tests: `diffutils`
- Required only when building the development version: `git`, `bison`, `flex`

The following command will install of these at once:

```
pacman -S \
  mingw-w64-x86_64-toolchain mingw-w64-x86_64-cmake \
  mingw-w64-x86_64-gmp mingw-w64-x86_64-libxml2 \
  mingw-w64-x86_64-openblas mingw-w64-x86_64-suitesparse mingw-w64-x86_64-arpack \
  diffutils \
  git bison flex
```

In order to build `igraph`, follow the **General build instructions** above, paying attention to the following:

- When using MSYS2, start the “MSYS2 MinGW 64-bit” terminal, and *not* the “MSYS2 MSYS” one.
- Be sure to install the `mingw-w64-x86_64-cmake` package and not the `cmake` one. The latter will not work.
- When running `cmake`, pass the option `-G"MSYS Makefiles"`.
- Note that `ccmake` is not currently available. `cmake-gui` can be used only if the `mingw-w64-x86_64-qt5` package is installed.

Notable configuration options

The following options may be set to ON or OFF. Some of them have an AUTO setting, which chooses a reasonable default based on what libraries are available on the current system.

- `igraph` bundles some of its dependencies for convenience. The `IGRAPH_USE_INTERNAL_XXX` flags control whether these should be used instead of external versions. Set them to `ON` to use the bundled (“vendored”) versions. Generally, external versions are preferable as they may be newer and usually provide better performance.
- `IGRAPH_GLPK_SUPPORT`: whether to make use of the GLPK [<https://www.gnu.org/software/glpk/>] library. Some features, such as finding a minimum feedback arc set or finding communities through exact modularity optimization, require this.
- `IGRAPH_GRAPHML_SUPPORT`: whether to enable support for reading and writing GraphML [<http://graphml.graphdrawing.org/>] files. Requires the libxml2 [<http://xmlsoft.org/>] library.
- `IGRAPH_OPENMP_SUPPORT`: whether to use OpenMP parallelization to accelerate certain functions such as PageRank calculation. Compiler support is required.
- `IGRAPH_ENABLE_LTO`: whether to build `igraph` with link-time optimization, which improves performance. Not supported with all compilers.
- `IGRAPH_ENABLE_TLS`: whether to enable thread-local storage. Required when using `igraph` from multiple threads.
- `BUILD_SHARED_LIBS` [https://cmake.org/cmake/help/latest/variable/BUILD_SHARED_LIBS.html]: whether to build a shared library instead of a static one.
- `BLA_VENDOR`: controls which library to use for BLAS [<https://cmake.org/cmake/help/latest/module/FindBLAS.html>] and LAPACK [<https://cmake.org/cmake/help/latest/module/FindLAPACK.html>] functionality.
- `CMAKE_INSTALL_PREFIX` [https://cmake.org/cmake/help/latest/variable/CMAKE_INSTALL_PREFIX.html]: the location where `igraph` will be installed.

Building the documentation

Most users will not need to build the documentation, as the release tarball contains pre-built HTML documentation in the `doc` directory.

To build the documentation for the development version, simply build the `html` or `pdf` targets for the HTML and PDF versions of the documentation, respectively.

```
$ cmake --build . --target html
```

Building the HTML documentation requires Python 3, `xmlto` and `source-highlight`. Building the PDF documentation also requires `xsltproc`, `xmllint` and `fop`.

Notes for package maintainers

This section is for people who package `igraph` for Linux distros or other package managers. Please read it carefully before packaging `igraph`.

Auto-detection of dependencies

`igraph` bundles several of its dependencies (or simplified versions of its dependencies). During configuration time, it checks whether each dependency is present on the system. If yes, it uses it. Otherwise, it falls back to the bundled (“vendored”) version. In order to make configuration as deterministic as possible, you may want to disable this auto-detection. To do so, set each of the `IGRAPH_USE_INTERNAL_XXX` option described above. Additionally, set `BLA_VENDOR` to use the BLAS and LAPACK implementations of your choice. This should be the same BLAS and LAPACK library that `igraph`'s other dependencies, such as ARPACK and CXSparse are linked against.

For example, to force igraph to use external versions of all dependencies, and to use OpenBLAS for BLAS/LAPACK, use

```
$ cmake .. \  
-DIGRAPH_USE_INTERNAL_BLAS=OFF \  
-DIGRAPH_USE_INTERNAL_LAPACK=OFF \  
-DIGRAPH_USE_INTERNAL_ARPACK=OFF \  
-DIGRAPH_USE_INTERNAL_GLPK=OFF \  
-DIGRAPH_USE_INTERNAL_CXSPARSE=OFF \  
-DIGRAPH_USE_INTERNAL_GMP=OFF \  
-DBLA_VENDOR=OpenBLAS \  
-DIGRAPH_GRAPHML_SUPPORT=ON
```

Shared and static builds

On Windows, shared and static builds should not be installed in the same location. If you decide to do so anyway, keep in mind the following: Both builds contain an `igraph.lib` file. The static one should be renamed to avoid conflict. The headers from the static build are incompatible with the shared library. The headers from the shared build may be used with the static library, but `IGRAPH_STATIC` must be defined when compiling programs that will link to igraph statically.

These issues do not affect Unix-like systems.

Cross-compiling

When building igraph with an internal ARPACK, LAPACK or BLAS, it makes use of `f2c`, which compiles and runs the `arithchk` program at build time to detect the floating point characteristics of the current system. It writes the results into the `arith.h` header. Since running this program is not possible when cross-compiling, igraph's build system allows specifying a pre-generated version of this header file through the `F2C_EXTERNAL_ARITH_HEADER` CMake option. An example version of this header follows for the `x86_64` and `arm64` target architectures on macOS. Warning: Do not use this version of `arith.h` on other systems or architectures.

```
#define IEEE_8087  
#define Arith_Kind_ASF 1  
#define Long int  
#define Intcast (int)(long)  
#define Double_Align  
#define X64_bit_pointers  
#define NANCHECK  
#define QNaN0 0x0  
#define QNaN1 0x7ff80000
```

Additional notes

- As of igraph 0.9, there is no tangible benefit to using an external GMP, as igraph does not yet use GMP in any performance-critical way. The bundled Mini-GMP is sufficient.
- Link-time optimization noticeably improves the performance of some igraph functions. To enable it, use `-DIGRAPH_ENABLE_LTO=ON`. The `AUTO` setting is also supported, and will enable link-time optimization only if the current compiler supports it. Note that this is detected by CMake, and the detection is not always accurate.
- We saw occasional hangs on Windows when igraph was built for a 32-bit target with MinGW and linked to OpenBLAS. We believe this to be an issue with OpenBLAS, not igraph. On this platform, you may want to opt for a different BLAS/LAPACK or the bundled BLAS/LAPACK.

Chapter 3. Tutorial

Compiling programs using igraph

The following short example program demonstrates the basic usage of the **igraph** library.

```
#include <igraph.h>

int main() {
    igraph_real_t diameter;
    igraph_t graph;

    igraph_rng_seed(igraph_rng_default(), 42);

    igraph_erdos_renyi_game(&graph, IGRAPH_ERDOS_RENYI_GNM, 1000, 3000,
                           IGRAPH_UNDIRECTED, IGRAPH_NO_LOOPS);

    igraph_diameter(&graph, &diameter, 0, 0, 0, IGRAPH_UNDIRECTED, 1);
    printf("Diameter of a random graph with average degree %g: %g\n",
           2.0 * igraph_ecount(&graph) / igraph_vcount(&graph),
           (double) diameter);

    igraph_destroy(&graph);

    return 0;
}
```

This example illustrates a couple of points. First, programs using the **igraph** library should include the `igraph.h` header file. Second, **igraph** uses the `igraph_real_t` type for real numbers instead of `double`. Third, **igraph** graph objects are represented by the `igraph_t` data type. Fourth, the `igraph_erdos_renyi_game()` creates a graph and `igraph_destroy()` destroys it, i.e. deallocates the memory associated to it.

For compiling this program you need a C compiler. Optionally, CMake [<https://cmake.org>] can be used to automate the compilation.

Compiling with CMake

It is convenient to use CMake because it can automatically discover the necessary compilation flags on all operating systems. Many IDEs support CMake, and can work with CMake projects directly. To create a CMake project for this example program, create a file name `CMakeLists.txt` with the following contents:

```
cmake_minimum_required(VERSION 3.16)
project(igraph_test C)

find_package(igraph REQUIRED)

add_executable(igraph_test igraph_test.c)
target_link_libraries(igraph_test PUBLIC igraph::igraph)
```

To compile the project, create a new directory called `build`, and switch to it:

```
mkdir build
cd build
```


Run CMake to configure the project:

```
cmake ..
```

If **igraph** was installed at a non-standard location, specify its prefix using the `-DCMAKE_PREFIX_PATH=...` option. The prefix must be the same directory that was specified as the `CMAKE_INSTALL_PREFIX` when compiling **igraph**.

If configuration has succeeded, build the program using

```
cmake --build .
```

Compiling without CMake

On most Unix-like systems, the default C compiler is called **cc**. To compile the test program, you will need a command similar to the following:

```
cc igraph_test.c -I/usr/local/include/igraph -L/usr/local/lib -ligraph -o igraph_test
```

The exact form depends on where **igraph** was installed on your system, whether it was compiled as a shared or static library, and the external libraries it was linked to. The directory after the `-I` switch is the one containing the `igraph.h` file, while the one following `-L` should contain the library file itself, usually a file called `libigraph.a` (static library on macOS and Linux), `libigraph.so` (shared library on Linux), `libigraph.dylib` (shared library on macOS), `igraph.lib` (static library on Windows) or `igraph.dll` (shared library on Windows). If **igraph** was compiled as a static library, it is also necessary to manually link to all of its dependencies.

If your system has the **pkg-config** utility you are likely to get the necessary compile options by issuing the command

```
pkg-config --libs --cflags igraph
```

(if **igraph** was built as a shared library) or

```
pkg-config --static --libs --cflags igraph
```

(if **igraph** was built as a static library).

Running the program

On most systems, the executable can be run by simply typing its name like this:

```
./igraph_test
```

If you use dynamic linking and the **igraph** library is not in a standard place, you may need to add its location to the `LD_LIBRARY_PATH` (Linux), `DYLD_LIBRARY_PATH` (macOS) or `PATH` (Windows) environment variables.

Creating your first graphs

The functions generating graph objects are called graph generators. Stochastic (i.e. randomized) graph generators are called “games”.

igraph can handle directed and undirected graphs. Most graph generators are able to create both types of graphs and most other functions are usually also capable of handling both. E.g. `igraph_shortest_paths()` which (surprisingly) calculates shortest paths from a vertex to other vertices can calculate directed or undirected paths.

igraph has sophisticated ways for creating graphs. The simplest graphs are deterministic regular structures like star graphs (`igraph_star()`), ring graphs (`igraph_ring()`), lattices (`igraph_lattice()`) or trees (`igraph_tree()`).

The following example creates an undirected regular circular lattice, adds some random edges to it and calculates the average length of shortest paths between all pairs of vertices in the graph before and after adding the random edges. (The message is that some random edges can reduce path lengths a lot.)

```
#include <igraph.h>

int main() {
    igraph_t graph;
    igraph_vector_t dimvector;
    igraph_vector_t edges;
    igraph_real_t avg_path_len;
    int i;

    igraph_vector_init(&dimvector, 2);
    VECTOR(dimvector)[0]=30;
    VECTOR(dimvector)[1]=30;
    igraph_lattice(&graph, &dimvector, 0, IGRAPH_UNDIRECTED, 0, 1);

    igraph_average_path_length(&graph, &avg_path_len, NULL, IGRAPH_UNDIRECTED, 1)
    printf("Average path length (lattice): %g\n", (double) avg_path_len);

    igraph_rng_seed(igraph_rng_default(), 42);
    igraph_vector_init(&edges, 20);
    for (i=0; i < igraph_vector_size(&edges); i++) {
        VECTOR(edges)[i] = RNG_INTEGER(0, igraph_vcount(&graph) - 1);
    }

    igraph_add_edges(&graph, &edges, 0);
    igraph_average_path_length(&graph, &avg_path_len, NULL, IGRAPH_UNDIRECTED, 1)
    printf("Average path length (randomized lattice): %g\n", (double) avg_path_len);

    igraph_vector_destroy(&dimvector);
    igraph_vector_destroy(&edges);
    igraph_destroy(&graph);

    return 0;
}
```

This example illustrates some new points. **igraph** uses `igraph_vector_t` instead of plain C arrays. `igraph_vector_t` is superior to regular arrays in almost every sense. Vectors are created by the `igraph_vector_init()` function and, like graphs, they should be destroyed if not needed any more by calling `igraph_vector_destroy()` on them. A vector can be indexed by the `VECTOR()` function (right now it is a macro). Vectors can be resized, e.g. most **igraph** functions returning the result in a vector resize it to the size of the result.

`igraph_lattice()` takes a vector argument specifying the dimensions of the lattice. In this example we generate a 30x30 two dimensional lattice. See the documentation of `igraph_lattice()` in the reference manual for the other arguments.

The vertices in a graph are identified by an integer number between 0 and N-1, N is the number of vertices in the graph (this can be obtained by `igraph_vcount()`, as in the example).

The `igraph_add_edges()` function simply takes a graph and a vector of vertex ids defining the new edges. The first edge is between the first two vertex ids in the vector, the second edge is between the second two, etc. This way we add ten random edges to the lattice.

Note that in the example it is possible to add loop edges, edges pointing to the same vertex and multiple edges, more than one edge between the same pair of vertices. `igraph_t` can of course represent loops and multiple edges, although some routines expect simple graphs, i.e. graphs without loop and multiple edges, because for example some structural properties are ill-defined for non-simple graphs. Loop edges can be removed by calling `igraph_simplify()`.

Calculating various properties of graphs

In our next example we will calculate various centrality measures in a friendship graph. The friendship graph is from the famous Zachary karate club study. (Web search on 'Zachary karate' if you want to know more about this.) Centrality measures quantify how central is the position of individual vertices in the graph.

```
#include <igraph.h>

int main() {
    igraph_t graph;
    igraph_vector_t v;
    igraph_vector_t result;
    igraph_real_t edges[] = { 0, 1, 0, 2, 0, 3, 0, 4, 0, 5, 0, 6, 0, 7, 0, 8,
                              0,10, 0,11, 0,12, 0,13, 0,17, 0,19, 0,21, 0,31,
                              1, 2, 1, 3, 1, 7, 1,13, 1,17, 1,19, 1,21, 1,30,
                              2, 3, 2, 7, 2,27, 2,28, 2,32, 2, 9, 2, 8, 2,13,
                              3, 7, 3,12, 3,13, 4, 6, 4,10, 5, 6, 5,10, 5,16,
                              6,16, 8,30, 8,32, 8,33, 9,33,13,33,14,32,14,33,
                              15,32,15,33,18,32,18,33,19,33,20,32,20,33,
                              22,32,22,33,23,25,23,27,23,32,23,33,23,29,
                              24,25,24,27,24,31,25,31,26,29,26,33,27,33,
                              28,31,28,33,29,32,29,33,30,32,30,33,31,32,31,33,
                              32,33
    };

    igraph_vector_view(&v, edges, sizeof(edges) / sizeof(double));
    igraph_create(&graph, &v, 0, IGRAPH_UNDIRECTED);

    igraph_vector_init(&result, 0);

    igraph_degree(&graph, &result, igraph_vss_all(), IGRAPH_ALL, IGRAPH_LOOPS);
    printf("Maximum degree is      %10i, vertex %2i.\n",
           (int) igraph_vector_max(&result), (int) igraph_vector_which_max(&result));

    igraph_closeness(&graph, &result, NULL, NULL, igraph_vss_all(), IGRAPH_ALL,
                    /*weights=*/ NULL, /*normalized=*/ 0);
    printf("Maximum closeness is    %10g, vertex %2i.\n",
           (double) igraph_vector_max(&result), (int) igraph_vector_which_max(&result));

    igraph_betweenness(&graph, &result, igraph_vss_all(),
                      IGRAPH_UNDIRECTED, /*weights=*/ NULL);
    printf("Maximum betweenness is %10g, vertex %2i.\n",
           (double) igraph_vector_max(&result), (int) igraph_vector_which_max(&result));

    igraph_vector_destroy(&result);
    igraph_destroy(&graph);
}
```

```
    return 0;
}
```

This example reflects some new features. First of all, it shows a way to define a graph simply as defining a C array with its edges. Function `igraph_vector_view()` creates a *view* of a C array. It does not copy any data, this also means that you should not call `igraph_vector_destroy()` on a vector created this way. This vector is then used to create the undirected graph.

Then the degree, closeness and betweenness centrality of the vertices is calculated and the highest values are printed. Note that the vector (`result`) which returns the result from these functions has to be initialized first, and also that the functions resize it to be able to hold the result.

The `igraph_vss_all()` argument tells the functions to calculate the property for every vertex in the graph, it is shorthand for a *vertex selector* (`igraph_vs_t`). Vertex selectors help to perform operations on a subset of vertices, you can read more about them in one of the following chapters.

Chapter 4. About igraph graphs, the basic interface

The igraph data model

The igraph library can handle directed and undirected graphs. The igraph graphs are multisets of ordered (if directed) or unordered (if undirected) labeled pairs. The labels of the pairs plus the number of vertices always starts with zero and ends with the number of edges minus one. In addition to that a table of metadata is also attached to every graph, its most important entries are the number of vertices in the graph and whether the graph is directed or undirected.

Like the edges, the igraph vertices are also labeled by numbers between zero and the number of vertices minus one. So, to summarize, a directed graph can be imagined like this:

```
( vertices: 6,
  directed: yes,
  {
    (0,2),
    (2,2),
    (2,3),
    (3,3),
    (3,4),
    (3,4),
    (4,1)
  }
)
```

Here the edges are ordered pairs of vertex ids, and the graph is a multiset of edges plus some meta-data.

An undirected graph is like this:

```
( vertices: 6,
  directed: no,
  {
    {0,2},
    {2},
    {2,3},
    {3},
    {3,4},
    {3,4},
    {4,1}
  }
)
```

Here an edge is a set of one or two vertex ids, two for most of the time, except for loop edges. A graph is a multiset of edges plus meta data, just like in the directed case.

It is possible to convert a directed graph to an undirected one, see the `igraph_to_directed()` and `igraph_to_undirected()` functions.

Note that igraph has some limited support for graphs with multiple edges. The support means that multiple edges can be stored in igraph graphs, but for most functions (like `igraph_betweenness()`) it is not checked that they work well on graphs with multiple edges. To eliminate multiple edges from a graph, you can use `igraph_simplify()`.

The basic interface

This is the very minimal API in **igraph**. All the other functions use this minimal set for creating and manipulating graphs.

This is a very important principle since it makes possible to implement other data representations by implementing only this minimal set.

Graph constructors and destructors

igraph_empty — Creates an empty graph with some vertices and no edges.

```
int igraph_empty(igraph_t *graph, igraph_integer_t n, igraph_bool_t directed);
```

The most basic constructor, all the other constructors should call this to create a minimal graph object. Our use of the term "empty graph" in the above description should be distinguished from the mathematical definition of the empty or null graph. Strictly speaking, the empty or null graph in graph theory is the graph with no vertices and no edges. However by "empty graph" as used in **igraph** we mean a graph having zero or more vertices, but no edges.

Arguments:

graph: Pointer to a not-yet initialized graph object.

n: The number of vertices in the graph, a non-negative integer number is expected.

directed: Boolean; whether the graph is directed or not. Supported values are:

IGRAPH_DIRECTED	The graph will be <i>directed</i> .
IGRAPH_UNDIRECTED	The graph will be <i>undirected</i> .

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices.

Time complexity: $O(|V|)$ for a graph with $|V|$ vertices (and no edges).

Example 4.1. File `examples/simple/igraph_empty.c`

igraph_empty_attrs — Creates an empty graph with some vertices, no edges and some graph attributes.

```
int igraph_empty_attrs(igraph_t *graph, igraph_integer_t n, igraph_bool_t direc
```

Use this instead of `igraph_empty()` if you wish to add some graph attributes right after initialization. This function is currently not very interesting for the ordinary user. Just supply 0 here or use `igraph_empty()`.

Arguments:

graph: Pointer to a not-yet initialized graph object.

n: The number of vertices in the graph; a non-negative integer number is expected.

directed: Boolean; whether the graph is directed or not. Supported values are:

IGRAPH_DIRECTED Create a *directed* graph.

IGRAPH_UNDIRECTED Create an *undirected* graph.

attr: The attributes.

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices.

Time complexity: $O(|V|)$ for a graph with $|V|$ vertices (and no edges).

igraph_copy — Creates an exact (deep) copy of a graph.

```
int igraph_copy(igraph_t *to, const igraph_t *from);
```

This function deeply copies a graph object to create an exact replica of it. The new replica should be destroyed by calling `igraph_destroy()` on it when not needed any more.

You can also create a shallow copy of a graph by simply using the standard assignment operator, but be careful and do *not* destroy a shallow replica. To avoid this mistake, creating shallow copies is not recommended.

Arguments:

to: Pointer to an uninitialized graph object.

from: Pointer to the graph object to copy.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$ for a graph with $|V|$ vertices and $|E|$ edges.

Example 4.2. File `examples/simple/igraph_copy.c`

igraph_destroy — Frees the memory allocated for a graph object.

```
void igraph_destroy(igraph_t *graph);
```

This function should be called for every graph object exactly once.

This function invalidates all iterators (of course), but the iterators of a graph should be destroyed before the graph itself anyway.

Arguments:

graph: Pointer to the graph to free.

Time complexity: operating system specific.

Basic query operations

igraph_vcount — The number of vertices in a graph.

```
igraph_integer_t igraph_vcount(const igraph_t *graph);
```

Arguments:

graph: The graph.

Returns:

Number of vertices.

Time complexity: O(1)

igraph_ecount — The number of edges in a graph.

```
igraph_integer_t igraph_ecount(const igraph_t *graph);
```

Arguments:

graph: The graph.

Returns:

Number of edges.

Time complexity: O(1)

igraph_edge — Gives the head and tail vertices of an edge.

```
int igraph_edge(const igraph_t *graph, igraph_integer_t eid,  
               igraph_integer_t *from, igraph_integer_t *to);
```

Arguments:

graph: The graph object.

eid: The edge id.

from: Pointer to an `igraph_integer_t`. The tail (head) of the edge will be placed here for undirected (directed) graphs.

to: Pointer to an `igraph_integer_t`. The head (tail) of the edge will be placed here for undirected (directed) graphs.

Returns:

Error code. The current implementation always returns with success.

See also:

`igraph_get_eid()` for the opposite operation; `igraph_edges()` to get the endpoints of several edges; `IGRAPH_TO()`, `IGRAPH_FROM()` and `IGRAPH_OTHER()` for a faster but non-error-checked version.

Added in version 0.2.

Time complexity: $O(1)$.

`igraph_edges` — Gives the head and tail vertices of a series of edges.

```
int igraph_edges(const igraph_t *graph, igraph_es_t eids,
                 igraph_vector_t *edges);
```

Arguments:

graph: The graph object.

eids: Edge selector, the series of edges.

edges: Pointer to an initialized vector. The start and endpoints of each edge will be placed here.

Returns:

Error code.

See also:

`igraph_get_edgelist()` to get the endpoints of all edges; `igraph_get_eids()` and `igraph_get_eids_multi()` for the opposite operation; `igraph_edge()` for getting the endpoints of a single edge; `IGRAPH_TO()`, `IGRAPH_FROM()` and `IGRAPH_OTHER()` for a faster but non-error-checked method.

Time complexity: $O(k)$ where k is the number of edges in the selector.

`IGRAPH_FROM` — The source vertex of an edge.

```
#define IGRAPH_FROM(graph,eid)
```

Faster than `igraph_edge()`, but no error checking is done: *eid* is assumed to be valid.

Arguments:

graph: The graph.

eid: The edge ID.

Returns:

The source vertex of the edge.

See also:

`igraph_edge()` if error checking is desired.

IGRAPH_TO — The target vertex of an edge.

```
#define IGRAPH_TO(graph,eid)
```

Faster than `igraph_edge()`, but no error checking is done: *eid* is assumed to be valid.

Arguments:

graph: The graph object.

eid: The edge ID.

Returns:

The target vertex of the edge.

See also:

`igraph_edge()` if error checking is desired.

IGRAPH_OTHER — The other endpoint of an edge.

```
#define IGRAPH_OTHER(graph,eid,vid)
```

Typically used with undirected edges when one endpoint of the edge is known, and the other endpoint is needed. No error checking is done: *eid* and *vid* are assumed to be valid.

Arguments:

graph: The graph object.

eid: The edge ID.

vid: The vertex ID of one endpoint of an edge.

Returns:

The other endpoint of the edge.

See also:

`IGRAPH_TO()` and `IGRAPH_FROM()` to get the source and target of directed edges.

igraph_get_eid — Get the edge id from the end points of an edge.

```
int igraph_get_eid(const igraph_t *graph, igraph_integer_t *eid,
                  igraph_integer_t pfrom, igraph_integer_t pto,
                  igraph_bool_t directed, igraph_bool_t error);
```

For undirected graphs `pfrom` and `pto` are exchangeable.

Arguments:

- graph*: The graph object.
- eid*: Pointer to an integer, the edge id will be stored here.
- pfrom*: The starting point of the edge.
- pto*: The end point of the edge.
- directed*: Logical constant, whether to search for directed edges in a directed graph. Ignored for undirected graphs.
- error*: Logical scalar, whether to report an error if the edge was not found. If it is false, then -1 will be assigned to *eid*.

Returns:

Error code.

See also:

`igraph_edge()` for the opposite operation.

Time complexity: $O(\log(d))$, where d is smaller of the out-degree of `pfrom` and in-degree of `pto` if *directed* is true. If *directed* is false, then it is $O(\log(d)+\log(d2))$, where d is the same as before and $d2$ is the minimum of the out-degree of `pto` and the in-degree of `pfrom`.

Example 4.3. File `examples/simple/igraph_get_eid.c`

Added in version 0.2.

igraph_get_eids — Return edge ids based on the adjacent vertices.

```
int igraph_get_eids(const igraph_t *graph, igraph_vector_t *eids,
                  const igraph_vector_t *pairs,
                  const igraph_vector_t *path,
                  igraph_bool_t directed, igraph_bool_t error);
```

This function operates in two modes. If the `pairs` argument is not a null pointer, but the `path` argument is, then it searches for the edge ids of all pairs of vertices given in `pairs`. The pairs of vertex ids are taken consecutively from the vector, i.e. `VECTOR(pairs)[0]` and `VECTOR(pairs)[1]` give the first pair, `VECTOR(pairs)[2]` and `VECTOR(pairs)[3]` the second pair, etc.

If the `pairs` argument is a null pointer, and `path` is not a null pointer, then the `path` is interpreted as a path given by vertex ids and the edges along the path are returned.

If neither `pairs` nor `path` are null pointers, then both are considered (first `pairs` and then `path`), and the results are concatenated.

If the `error` argument is true, then it is an error to give pairs of vertices that are not connected. Otherwise -1 is reported for not connected vertices.

If there are multiple edges in the graph, then these are ignored; i.e. for a given pair of vertex ids, always the same edge id is returned, even if the pair is given multiple time in `pairs` or in `path`. See `igraph_get_eids_multi()` for a similar function that works differently in case of multiple edges.

Arguments:

- graph*: The input graph.
- eids*: Pointer to an initialized vector, the result is stored here. It will be resized as needed.
- pairs*: Vector giving pairs of vertices, or a null pointer.
- path*: Vector giving vertex ids along a path, or a null pointer.
- directed*: Logical scalar, whether to consider edge directions in directed graphs. This is ignored for undirected graphs.
- error*: Logical scalar, whether it is an error to supply non-connected vertices. If false, then -1 is returned for non-connected pairs.

Returns:

Error code.

Time complexity: $O(n \log(d))$, where n is the number of queried edges and d is the average degree of the vertices.

See also:

`igraph_get_eid()` for a single edge, `igraph_get_eids_multi()` for a version that handles multiple edges better (at a cost).

Example 4.4. File `examples/simple/igraph_get_eids.c`

`igraph_get_eids_multi` — Query edge ids based on their adjacent vertices, handle multiple edges.

```
int igraph_get_eids_multi(const igraph_t *graph, igraph_vector_t *eids,
                        const igraph_vector_t *pairs,
                        const igraph_vector_t *path,
                        igraph_bool_t directed, igraph_bool_t error);
```

This function operates in two modes. If the `pairs` argument is not a null pointer, but the `path` argument is, then it searches for the edge ids of all pairs of vertices given in `pairs`. The pairs of vertex ids are taken consecutively from the vector, i.e. `VECTOR(pairs)[0]` and `VECTOR(pairs)[1]` give the first pair, `VECTOR(pairs)[2]` and `VECTOR(pairs)[3]` the second pair, etc.

If the `pairs` argument is a null pointer, and `path` is not a null pointer, then the `path` is interpreted as a path given by vertex ids and the edges along the path are returned.

If the `error` argument is true, then it is an error to give pairs of vertices that are not connected. Otherwise -1 is returned for not connected vertex pairs.

An error is triggered if both `pairs` and `path` are non-null pointers.

This function handles multiple edges properly, i.e. if the same pair is given multiple times and they are indeed connected by multiple edges, then each time a different edge id is reported.

Arguments:

- graph*: The input graph.
- eids*: Pointer to an initialized vector, the result is stored here. It will be resized as needed.
- pairs*: Vector giving pairs of vertices, or a null pointer.
- path*: Vector giving vertex ids along a path, or a null pointer.
- directed*: Logical scalar, whether to consider edge directions in directed graphs. This is ignored for undirected graphs.
- error*: Logical scalar, whether to report an error if non-connected vertices are specified. If false, then -1 is returned for non-connected vertex pairs.

Returns:

Error code.

Time complexity: $O(|E| + n \log(d))$, where $|E|$ is the number of edges in the graph, n is the number of queried edges and d is the average degree of the vertices.

See also:

`igraph_get_eid()` for a single edge, `igraph_get_eids()` for a faster version that does not handle multiple edges.

igraph_neighbors — Adjacent vertices to a vertex.

```
int igraph_neighbors(const igraph_t *graph, igraph_vector_t *neis, igraph_integ
    igraph_neimode_t mode);
```

Arguments:

- graph*: The graph to work on.
- neis*: This vector will contain the result. The vector should be initialized beforehand and will be resized. Starting from igraph version 0.4 this vector is always sorted, the vertex ids are in increasing order.

pnode: The id of the node for which the adjacent vertices are to be searched.

mode: Defines the way adjacent vertices are searched in directed graphs. It can have the following values: IGRAPH_OUT, vertices reachable by an edge from the specified vertex are searched; IGRAPH_IN, vertices from which the specified vertex is reachable are searched; IGRAPH_ALL, both kinds of vertices are searched. This parameter is ignored for undirected graphs.

Returns:

Error code: IGRAPH_EINVAL: invalid vertex id. IGRAPH_EINVMODE: invalid mode argument. IGRAPH_ENOMEM: not enough memory.

Time complexity: $O(d)$, d is the number of adjacent vertices to the queried vertex.

Example 4.5. File `examples/simple/igraph_neighbors.c`

igraph_incident — Gives the incident edges of a vertex.

```
int igraph_incident(const igraph_t *graph, igraph_vector_t *eids, igraph_integer_t pnode, igraph_neimode_t mode);
```

Arguments:

graph: The graph object.

eids: An initialized vector_t object. It will be resized to hold the result.

pnode: A vertex id.

mode: Specifies what kind of edges to include for directed graphs. IGRAPH_OUT means only outgoing edges, IGRAPH_IN only incoming edges, IGRAPH_ALL both. This parameter is ignored for undirected graphs.

Returns:

Error code. IGRAPH_EINVAL: invalid *pnode* argument, IGRAPH_EINVMODE: invalid *mode* argument.

Added in version 0.2.

Time complexity: $O(d)$, the number of incident edges to *pnode*.

igraph_is_directed — Is this a directed graph?

```
igraph_bool_t igraph_is_directed(const igraph_t *graph);
```

Arguments:

graph: The graph.

Returns:

Logical value, TRUE if the graph is directed, FALSE otherwise.

Time complexity: $O(1)$

Example 4.6. File `examples/simple/igraph_is_directed.c`

igraph_is_same_graph — Are two graphs identical as labelled graphs?

```
int igraph_is_same_graph(const igraph_t *graph1, const igraph_t *graph2, igraph_t *res);
```

Two graphs are considered to be the same if they have the same vertex and edge sets. Graphs which are the same may have multiple different representations in igraph, hence the need for this function.

This function verifies that the two graphs have the same directedness, the same number of vertices, and that they contain precisely the same edges (regardless of their ordering) when written in terms of vertex indices. Graph attributes are not taken into account.

This concept is different from isomorphism. For example, the graphs 0-1, 2-1 and 1-2, 0-1 are considered the same because they only differ in the ordering of their edge lists and the ordering of vertices in an undirected edge. However, they are not the same as 0-2, 1-2, even though they are isomorphic to it. Note that this latter graph contains the edge 0-2 while the former two do not — thus their edge sets differ.

Arguments:

graph1: The first graph object.

graph2: The second graph object.

res: The result will be stored here.

Returns:

Error code.

Time complexity: $O(E)$, the number of edges in the graphs.

See also:

`igraph_isomorphic()` to test if two graphs are isomorphic.

igraph_degree — The degree of some vertices in a graph.

```
int igraph_degree(const igraph_t *graph, igraph_vector_t *res,  
                  const igraph_vs_t vids,  
                  igraph_neimode_t mode, igraph_bool_t loops);
```

This function calculates the in-, out- or total degree of the specified vertices.

Arguments:

graph: The graph.

res: Vector, this will contain the result. It should be initialized and will be resized to be the appropriate size.

vids: Vector, giving the vertex ids of which the degree will be calculated.

mode: Defines the type of the degree. Valid modes are: IGRAPH_OUT, out-degree; IGRAPH_IN, in-degree; IGRAPH_ALL, total degree (sum of the in- and out-degree). This parameter is ignored for undirected graphs.

loops: Boolean, gives whether the self-loops should be counted.

Returns:

Error code: IGRAPH_EINVAL: invalid vertex id. IGRAPH_EINVALMODE: invalid mode argument.

Time complexity: $O(v)$ if *loops* is TRUE, and $O(v*d)$ otherwise. *v* is the number of vertices for which the degree will be calculated, and *d* is their (average) degree.

See also:

`igraph_strength()` for the version that takes into account edge weights.

Example 4.7. File `examples/simple/igraph_degree.c`

Adding and deleting vertices and edges

`igraph_add_edge` — Adds a single edge to a graph.

```
int igraph_add_edge(igraph_t *graph, igraph_integer_t from, igraph_integer_t to
```

For directed graphs the edge points from *from* to *to*.

Note that if you want to add many edges to a big graph, then it is inefficient to add them one by one, it is better to collect them into a vector and add all of them via a single `igraph_add_edges()` call.

Arguments:

igraph: The graph.

from: The id of the first vertex of the edge.

to: The id of the second vertex of the edge.

Returns:

Error code.

See also:

`igraph_add_edges()` to add many edges, `igraph_delete_edges()` to remove edges and `igraph_add_vertices()` to add vertices.

Time complexity: $O(|V|+|E|)$, the number of edges plus the number of vertices.

igraph_add_edges — Adds edges to a graph object.

```
int igraph_add_edges(igraph_t *graph, const igraph_vector_t *edges,  
                    void *attr);
```

The edges are given in a vector, the first two elements define the first edge (the order is from, to for directed graphs). The vector should contain even number of integer numbers between zero and the number of vertices in the graph minus one (inclusive). If you also want to add new vertices, call `igraph_add_vertices()` first.

Arguments:

graph: The graph to which the edges will be added.

edges: The edges themselves.

attr: The attributes of the new edges, only used by high level interfaces currently, you can supply 0 here.

Returns:

Error code: `IGRAPH_EINVEVECTOR`: invalid (odd) edges vector length, `IGRAPH_EINVVID`: invalid vertex id in edges vector.

This function invalidates all iterators.

Time complexity: $O(|V|+|E|)$ where $|V|$ is the number of vertices and $|E|$ is the number of edges in the *new*, extended graph.

Example 4.8. File `examples/simple/igraph_add_edges.c`

igraph_add_vertices — Adds vertices to a graph.

```
int igraph_add_vertices(igraph_t *graph, igraph_integer_t nv, void *attr);
```

This function invalidates all iterators.

Arguments:

graph: The graph object to extend.

nv: Non-negative integer giving the number of vertices to add.

attr: The attributes of the new vertices, only used by high level interfaces, you can supply 0 here.

Returns:

Error code: `IGRAPH_EINVAL`: invalid number of new vertices.

Time complexity: $O(|V|)$ where $|V|$ is the number of vertices in the *new*, extended graph.

Example 4.9. File `examples/simple/igraph_add_vertices.c`

igraph_delete_edges — Removes edges from a graph.

```
int igraph_delete_edges(igraph_t *graph, igraph_es_t edges);
```

The edges to remove are given as an edge selector.

This function cannot remove vertices, they will be kept, even if they lose all their edges.

This function invalidates all iterators.

Arguments:

graph: The graph to work on.

edges: The edges to remove.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$ where $|V|$ and $|E|$ are the number of vertices and edges in the *original* graph, respectively.

Example 4.10. File `examples/simple/igraph_delete_edges.c`

igraph_delete_vertices — Removes vertices (with all their edges) from the graph.

```
int igraph_delete_vertices(igraph_t *graph, const igraph_vs_t vertices);
```

This function changes the ids of the vertices (except in some very special cases, but these should not be relied on anyway).

This function invalidates all iterators.

Arguments:

graph: The graph to work on.

vertices: The ids of the vertices to remove in a vector. The vector may contain the same id more than once.

Returns:

Error code: `IGRAPH_EINVVID`: invalid vertex id.

Time complexity: $O(|V|+|E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the original graph.

Example 4.11. File `examples/simple/igraph_delete_vertices.c`

Chapter 5. Error handling

Error handling basics

igraph functions can run into various problems preventing them from normal operation. The user might have supplied invalid arguments, e.g. a non-square matrix when a square-matrix was expected, or the program has run out of memory while some more memory allocation is required, etc.

By default **igraph** aborts the program when it runs into an error. While this behavior might be good enough for smaller programs, it is without doubt avoidable in larger projects. Please read further if your project requires more sophisticated error handling. You can safely skip the rest of this chapter otherwise.

Error handlers

If **igraph** runs into an error - an invalid argument was supplied to a function, or we've ran out of memory - the control is transferred to the *error handler* function.

The default error handler is `igraph_error_handler_abort` which prints an error message and aborts the program.

The `igraph_set_error_handler()` function can be used to set a new error handler function of type `igraph_error_handler_t`; see the documentation of this type for details.

There are two other predefined error handler functions, `igraph_error_handler_ignore` and `igraph_error_handler_printignore`. These deallocate the temporarily allocated memory (more about this later) and return with the error code. The latter also prints an error message. If you use these error handlers you need to take care about possible errors yourself by checking the return value of (almost) every non-void **igraph** function.

Independently of the error handler installed, all functions in the library do their best to leave their arguments *semantically* unchanged if an error happens. By semantically we mean that the implementation of an object supplied as an argument might change, but its “meaning” in most cases does not. The rare occasions when this rule is violated are documented in this manual.

`igraph_error_handler_t` — The type of error handler functions.

```
typedef void igraph_error_handler_t (const char * reason, const char * file,
                                     int line, int igraph_errno);
```

This is the type of the error handler functions.

Arguments:

<i>reason</i> :	Textual description of the error.
<i>file</i> :	The source file in which the error is noticed.
<i>line</i> :	The number of the line in the source file which triggered the error
<i>igraph_errno</i> :	The igraph error code.

igraph_error_handler_abort — Abort program in case of error.

```
IGRAPH_EXPORT igraph_error_handler_t igraph_error_handler_abort;
```

The default error handler, prints an error message and aborts the program.

igraph_error_handler_ignore — Ignore errors.

```
IGRAPH_EXPORT igraph_error_handler_t igraph_error_handler_ignore;
```

This error handler frees the temporarily allocated memory and returns with the error code.

igraph_error_handler_printignore — Print and ignore errors.

```
IGRAPH_EXPORT igraph_error_handler_t igraph_error_handler_printignore;
```

Frees temporarily allocated memory, prints an error message to the standard error and returns with the error code.

Error codes

Every **igraph** function which can fail return a single integer error code. Some functions are very simple and cannot run into any error, these may return other types, or void as well. The error codes are defined by the `igraph_error_type_t` enumeration.

igraph_error_type_t — Error code type.

```
typedef enum {  
    IGRAPH_SUCCESS           = 0,  
    IGRAPH_FAILURE          = 1,  
    IGRAPH_ENOMEM           = 2,  
    IGRAPH_PARSEERROR       = 3,  
    IGRAPH_EINVAL           = 4,  
    IGRAPH_EXISTS           = 5,  
    IGRAPH_EINVEVECTOR      = 6,  
    IGRAPH_EINVVID          = 7,  
    IGRAPH_NONSQUARE        = 8,  
    IGRAPH_EINVMODE         = 9,  
    IGRAPH_EFILE            = 10,  
    IGRAPH_UNIMPLEMENTED    = 12,  
    IGRAPH_INTERRUPTED      = 13,  
    IGRAPH_DIVERGED         = 14,  
    IGRAPH_ARPACK_PROD      = 15,  
    IGRAPH_ARPACK_NPOS      = 16,  
    IGRAPH_ARPACK_NEVNPOS   = 17,  
}
```

```
IGRAPH_ARPACK_NCVSMALL    = 18,
IGRAPH_ARPACK_NONPOSI     = 19,
IGRAPH_ARPACK_WHICHINV    = 20,
IGRAPH_ARPACK_BMATINV     = 21,
IGRAPH_ARPACK_WORKLSMALL  = 22,
IGRAPH_ARPACK_TRIDERR     = 23,
IGRAPH_ARPACK_ZEROSTART   = 24,
IGRAPH_ARPACK_MODEINV     = 25,
IGRAPH_ARPACK_MODEBMAT    = 26,
IGRAPH_ARPACK_ISHIFT      = 27,
IGRAPH_ARPACK_NEVBE       = 28,
IGRAPH_ARPACK_NOFACT      = 29,
IGRAPH_ARPACK_FAILED      = 30,
IGRAPH_ARPACK_HOWMNY      = 31,
IGRAPH_ARPACK_HOWMNYS     = 32,
IGRAPH_ARPACK_EVDIFF      = 33,
IGRAPH_ARPACK_SHUR        = 34,
IGRAPH_ARPACK_LAPACK      = 35,
IGRAPH_ARPACK_UNKNOWN     = 36,
IGRAPH_ENEGLOOP           = 37,
IGRAPH_EINTERNAL          = 38,
IGRAPH_ARPACK_MAXIT       = 39,
IGRAPH_ARPACK_NOSHIFT     = 40,
IGRAPH_ARPACK_REORDER     = 41,
IGRAPH_EDIVZERO           = 42,
IGRAPH_GLP_EBOUND         = 43,
IGRAPH_GLP_EROOT          = 44,
IGRAPH_GLP_ENOPFS         = 45,
IGRAPH_GLP_ENODFS         = 46,
IGRAPH_GLP_EFAIL          = 47,
IGRAPH_GLP_EMIPGAP        = 48,
IGRAPH_GLP_ETMLIM         = 49,
IGRAPH_GLP_ESTOP          = 50,
IGRAPH_EATTRIBUTES        = 51,
IGRAPH_EATTRCOMBINE       = 52,
IGRAPH_ELAPACK            = 53,
IGRAPH_EDRL               = 54,
IGRAPH_EOVERFLOW          = 55,
IGRAPH_EGLP               = 56,
IGRAPH_CPUTIME            = 57,
IGRAPH_EUNDERFLOW        = 58,
IGRAPH_ERWSTUCK           = 59,
IGRAPH_STOP               = 60 /* undocumented, used internally */
} igraph_error_type_t;
```

These are the possible values returned by **igraph** functions. Note that these are interesting only if you defined an error handler with `igraph_set_error_handler()`. Otherwise the program is aborted and the function causing the error never returns.

Values:

IGRAPH_SUCCESS:	The function successfully completed its task.
IGRAPH_FAILURE:	Something went wrong. You'll almost never meet this error as normally more specific error codes are used.
IGRAPH_ENOMEM:	There wasn't enough memory to allocate on the heap.
IGRAPH_PARSEERROR:	A parse error was found in a file.

IGRAPH_EINVAL:	A parameter's value is invalid. E.g. negative number was specified as the number of vertices.
IGRAPH_EXISTS:	A graph/vertex/edge attribute is already installed with the given name.
IGRAPH_EINVEVECTOR:	Invalid vector of vertex ids. A vertex id is either negative or bigger than the number of vertices minus one.
IGRAPH_EINVVID:	Invalid vertex id, negative or too big.
IGRAPH_NONSQUARE:	A non-square matrix was received while a square matrix was expected.
IGRAPH_EINVMODE:	Invalid mode parameter.
IGRAPH_EFILE:	A file operation failed. E.g. a file doesn't exist, or the user has no rights to open it.
IGRAPH_UNIMPLEMENTED:	Attempted to call an unimplemented or disabled (at compile-time) function.
IGRAPH_DIVERGED:	A numeric algorithm failed to converge.
IGRAPH_ARPACK_PROD:	Matrix-vector product failed.
IGRAPH_ARPACK_NPOS:	N must be positive.
IGRAPH_ARPACK_NEVNPOS:	NEV must be positive.
IGRAPH_ARPACK_NCVSMALL:	NCV must be bigger.
IGRAPH_ARPACK_NONPOSI:	Maximum number of iterations should be positive.
IGRAPH_ARPACK_WHICHINV:	Invalid WHICH parameter.
IGRAPH_ARPACK_BMATINV:	Invalid BMAT parameter.
IGRAPH_ARPACK_WORKLS-MALL:	WORKL is too small.
IGRAPH_ARPACK_TRIDERR:	LAPACK error in tridiagonal eigenvalue calculation.
IGRAPH_ARPACK_ZEROSTART:	Starting vector is zero.
IGRAPH_ARPACK_MODEINV:	MODE is invalid.
IGRAPH_ARPACK_MODEBMAT:	MODE and BMAT are not compatible.
IGRAPH_ARPACK_ISHIFT:	ISHIFT must be 0 or 1.
IGRAPH_ARPACK_NEVBE:	NEV and WHICH='BE' are incompatible.
IGRAPH_ARPACK_NOFACT:	Could not build an Arnoldi factorization.
IGRAPH_ARPACK_FAILED:	No eigenvalues to sufficient accuracy.
IGRAPH_ARPACK_HOWMNY:	HOWMNY is invalid.
IGRAPH_ARPACK_HOWMNYs:	HOWMNY='S' is not implemented.
IGRAPH_ARPACK_EVDIFF:	Different number of converged Ritz values.

IGRAPH_ARPACK_SHUR:	Error from calculation of a real Schur form.
IGRAPH_ARPACK_LAPACK:	LAPACK (dtrevc) error for calculating eigenvectors.
IGRAPH_ARPACK_UNKNOWN:	Unknown ARPACK error.
IGRAPH_ENEGLOOP:	Negative loop detected while calculating shortest paths.
IGRAPH_EINTERNAL:	Internal error, likely a bug in igraph.
IGRAPH_EDIVZERO:	Big integer division by zero.
IGRAPH_GLP_EBOUND:	GLPK error (GLP_EBOUND).
IGRAPH_GLP_EROOT:	GLPK error (GLP_EROOT).
IGRAPH_GLP_ENOPFS:	GLPK error (GLP_ENOPFS).
IGRAPH_GLP_ENODFS:	GLPK error (GLP_ENODFS).
IGRAPH_GLP_EFAIL:	GLPK error (GLP_EFAIL).
IGRAPH_GLP_EMIPGAP:	GLPK error (GLP_EMIPGAP).
IGRAPH_GLP_ETMLIM:	GLPK error (GLP_ETMLIM).
IGRAPH_GLP_ESTOP:	GLPK error (GLP_ESTOP).
IGRAPH_EATTRIBUTES:	Attribute handler error. The user is not expected to find this; it is signalled if some igraph function is not using the attribute handler interface properly.
IGRAPH_EATTRCOMBINE:	Unimplemented attribute combination method for the given attribute type.
IGRAPH_ELAPACK:	A LAPACK call resulted in an error.
IGRAPH_EDRL:	Internal error in the DrL layout generator.
IGRAPH_EOVERFLOW:	Integer or double overflow.
IGRAPH_EGLP:	Internal GLPK error.
IGRAPH_CPUTIME:	CPU time exceeded.
IGRAPH_EUNDERFLOW:	Integer or double underflow.
IGRAPH_ERWSTUCK:	Random walk got stuck.

igraph_strerror — Textual description of an error.

```
const char* igraph_strerror(const int igraph_errno);
```

This is a simple utility function, it gives a short general textual description for an **igraph** error code.

Arguments:

igraph_errno: The **igraph** error code.

Returns:

pointer to the textual description of the error code.

Warning messages

igraph also supports warning messages in addition to error messages. Warning messages typically do not terminate the program, but they are usually crucial to the user.

igraph warnings are handled similarly to errors. There is a separate warning handler function that is called whenever an **igraph** function triggers a warning. This handler can be set by the `igraph_set_warning_handler()` function. There are two predefined simple warning handlers, `igraph_warning_handler_ignore()` and `igraph_warning_handler_print()`, the latter being the default.

To trigger a warning, **igraph** functions typically use the `IGRAPH_WARNING()` macro, the `igraph_warning()` function, or if more flexibility is needed, `igraph_warningf()`.

igraph_warning_handler_t — The type of igraph warning handler functions.

```
typedef igraph_error_handler_t igraph_warning_handler_t;
```

Currently it is defined to have the same type as `igraph_error_handler_t`, although the last (error code) argument is not used.

igraph_set_warning_handler — Installs a warning handler.

```
igraph_warning_handler_t* igraph_set_warning_handler(igraph_warning_handler_t*
```

Install the supplied warning handler function.

Arguments:

new_handler: The new warning handler function to install. Supply a null pointer here to uninstall the current warning handler, without installing a new one.

Returns:

The current warning handler function.

IGRAPH_WARNING — Triggers a warning.

```
#define IGRAPH_WARNING(reason)
```

This is the usual way of triggering a warning from an **igraph** function. It calls `igraph_warning()`.

Arguments:

reason: The warning message.

IGRAPH_WARNINGF — Triggers a warning, with printf-like syntax.

```
#define IGRAPH_WARNINGF
```

igraph functions can use this macro when they notice a warning and want to pass on extra information to the user about what went wrong. It calls `igraph_warningf()` with the proper parameters and no error code.

Arguments:

reason: Textual description of the warning, a template string with the same syntax as the standard `printf` C library function.

...: The additional arguments to be substituted into the template string.

igraph_warning — Triggers a warning.

```
int igraph_warning(const char *reason, const char *file, int line,
                  int igraph_errno);
```

Call this function if you want to trigger a warning from within a function that uses **igraph**.

Arguments:

reason: Textual description of the warning.

file: The source file in which the warning was noticed.

line: The number of line in the source file which triggered the warning.

igraph_errno: Warnings could have potentially error codes as well, but this is currently not used in `igraph`.

Returns:

The supplied error code.

igraph_warningf — Triggers a warning, printf-like version.

```
int igraph_warningf(const char *reason, const char *file, int line,
                   int igraph_errno, ...);
```

This function is similar to `igraph_warning()`, but uses a `printf`-like syntax. It substitutes the additional arguments into the *reason* template string and calls `igraph_warning()`.

Arguments:

<i>reason</i> :	Textual description of the warning, a template string with the same syntax as the standard <code>printf</code> C library function.
<i>file</i> :	The source file in which the warning was noticed.
<i>line</i> :	The number of line in the source file which triggered the warning.
<i>igraph_errno</i> :	Warnings could have potentially error codes as well, but this is currently not used in <code>igraph</code> .
<i>...</i> :	The additional arguments to be substituted into the template string.

Returns:

The supplied error code.

`igraph_warning_handler_ignore` — Ignores all warnings.

```
void igraph_warning_handler_ignore(const char *reason, const char *file,  
                                  int line, int igraph_errno);
```

This warning handler function simply ignores all warnings.

Arguments:

<i>reason</i> :	Textual description of the warning.
<i>file</i> :	The source file in which the warning was noticed.
<i>line</i> :	The number of line in the source file which triggered the warning..
<i>igraph_errno</i> :	Warnings could have potentially error codes as well, but this is currently not used in <code>igraph</code> .

`igraph_warning_handler_print` — Prints all warnings to the standard error.

```
void igraph_warning_handler_print(const char *reason, const char *file,  
                                  int line, int igraph_errno);
```

This warning handler function simply prints all warnings to the standard error.

Arguments:

<i>reason</i> :	Textual description of the warning.
<i>file</i> :	The source file in which the warning was noticed.

<i>line</i> :	The number of line in the source file which triggered the warning..
<i>igraph_errno</i> :	Warnings could have potentially error codes as well, but this is currently not used in igraph.

Advanced topics

Writing error handlers

The contents of the rest of this chapter might be useful only for those who want to create an interface to **igraph** from another language. Most readers can safely skip to the next chapter.

You can write and install error handlers simply by defining a function of type `igraph_error_handler_t` and calling `igraph_set_error_handler()`. This feature is useful for interface writers, as **igraph** will have the chance to signal errors the appropriate way, e.g. the R interface defines an error handler which calls the `error()` function, as required by R, while the Python interface has an error handler which raises an exception according to the Python way.

If you want to write an error handler, your error handler should call `IGRAPH_FINALLY_FREE()` to deallocate all temporary memory to prevent memory leaks.

igraph_set_error_handler — Sets a new error handler.

```
igraph_error_handler_t* igraph_set_error_handler(igraph_error_handler_t* new_h
```

Installs a new error handler. If called with 0, it installs the default error handler (which is currently `igraph_error_handler_abort`).

Arguments:

new_handler: The error handler function to install.

Returns:

The old error handler function. This should be saved and restored if *new_handler* is not needed any more.

Error handling internals

If an error happens, the functions in the library call the `IGRAPH_ERROR()` macro with a textual description of the error and an **igraph** error code. This macro calls (through the `igraph_error()` function) the installed error handler. Another useful macro is `IGRAPH_CHECK()`. This checks the return value of its argument, which is normally a function call, and calls `IGRAPH_ERROR()` if it is not `IGRAPH_SUCCESS`.

IGRAPH_ERROR — Trigger an error.

```
#define IGRAPH_ERROR(reason, igraph_errno)
```

igraph functions usually use this macro when they notice an error. It calls `igraph_error()` with the proper parameters and if that returns the macro returns the "calling" function as well, with the error

code. If for some (suspicious) reason you want to call the error handler without returning from the current function, call `igraph_error()` directly.

Arguments:

reason: Textual description of the error. This should be something more descriptive than the text associated with the error code. E.g. if the error code is `IGRAPH_EINVAL`, its associated text (see `igraph_strerror()`) is "Invalid value" and this string should explain which parameter was invalid and maybe why.

igraph_errno: The **igraph** error code.

IGRAPH_ERRORF — Triggers an error, with printf-like syntax.

```
#define IGRAPH_ERRORF
```

igraph functions can use this macro when they notice an error and want to pass on extra information to the user about what went wrong. It calls `igraph_errorf()` with the proper parameters and if that returns the macro returns the "calling" function as well, with the error code. If for some (suspicious) reason you want to call the error handler without returning from the current function, call `igraph_errorf()` directly.

Arguments:

reason: Textual description of the error, a template string with the same syntax as the standard `printf` C library function. This should be something more descriptive than the text associated with the error code. E.g. if the error code is `IGRAPH_EINVAL`, its associated text (see `igraph_strerror()`) is "Invalid value" and this string should explain which parameter was invalid and maybe what was expected and what was recieved.

igraph_errno: The **igraph** error code.

...: The additional arguments to be substituted into the template string.

igraph_error — Triggers an error.

```
int igraph_error(const char *reason, const char *file, int line,  
                int igraph_errno);
```

igraph functions usually call this function (most often via the `IGRAPH_ERROR` macro) if they notice an error. It calls the currently installed error handler function with the supplied arguments.

Arguments:

reason: Textual description of the error.

file: The source file in which the error was noticed.

line: The number of line in the source file which triggered the error.

igraph_errno: The **igraph** error code.

Returns:

the error code (if it returns)

See also:

`igraph_errorf()`.

igraph_errorf — Triggers an error, printf-like version.

```
int igraph_errorf(const char *reason, const char *file, int line,
                  int igraph_errno, ...);
```

Arguments:

reason: Textual description of the error, interpreted as a `printf` format string.

file: The source file in which the error was noticed.

line: The line in the source file which triggered the error.

igraph_errno: The **igraph** error code.

...: Additional parameters, the values to substitute into the format string.

See also:

`igraph_error()`.

IGRAPH_CHECK — Checks the return value of a function call.

```
#define IGRAPH_CHECK(a)
```

Arguments:

a: An expression, usually a function call.

Executes the expression and checks its value. If this is not `IGRAPH_SUCCESS`, it calls `IGRAPH_ERROR` with the value as the error code. Here is an example usage:

```
IGRAPH_CHECK(vector_push_back(&v, 100));
```

There is only one reason to use this macro when writing **igraph** functions. If the user installs an error handler which returns to the auxiliary calling code (like `igraph_error_handler_ignore` and `igraph_error_handler_printignore`), and the **igraph** function signalling the error is called from another **igraph** function then we need to make sure that the error is propagated back to the auxiliary (i.e. non-igraph) calling function. This is achieved by using `IGRAPH_CHECK` on every **igraph** call which can return an error code.

Deallocating memory

If a function runs into an error (and the program is not aborted) the error handler should deallocate all temporary memory. This is done by storing the address and the destroy function of all temporary ob-

jects in a stack. The `IGRAPH_FINALLY` function declares an object as temporary by placing its address in the stack. If an **igraph** function returns with success it calls `IGRAPH_FINALLY_CLEAN()` with the number of objects to remove from the stack. If an error happens however, the error handler should call `IGRAPH_FINALLY_FREE()` to deallocate each object added to the stack. This means that the temporary objects allocated in the calling function (and etc.) will be freed as well.

IGRAPH_FINALLY — Registers an object for deallocation.

```
#define IGRAPH_FINALLY(func, ptr)
```

Arguments:

func: The address of the function which is normally called to destroy the object.

ptr: Pointer to the object itself.

This macro places the address of an object, together with the address of its destructor in a stack. This stack is used if an error happens to deallocate temporarily allocated objects to prevent memory leaks.

IGRAPH_FINALLY_CLEAN — Signals clean deallocation of objects.

```
void IGRAPH_FINALLY_CLEAN(int num);
```

Removes the specified number of objects from the stack of temporarily allocated objects. Most often this is called just before returning from a function.

Arguments:

num: The number of objects to remove from the bookkeeping stack.

IGRAPH_FINALLY_FREE — Deallocates all registered objects.

```
void IGRAPH_FINALLY_FREE(void);
```

Calls the destroy function for all objects in the stack of temporarily allocated objects. This is usually called only from an error handler. It is *not* appropriate to use it instead of destroying each unneeded object of a function, as it destroys the temporary objects of the caller function (and so on) as well.

Writing igraph functions with proper error handling

There are some simple rules to keep in order to have functions behaving well in erroneous situations. First, check the arguments of the functions and call `IGRAPH_ERROR()` if they are invalid. Second, call `IGRAPH_FINALLY` on each dynamically allocated object and call `IGRAPH_FINALLY_CLEAN()` with the proper argument before returning. Third, use `IGRAPH_CHECK` on all **igraph** function calls which can generate errors.

The size of the stack used for this bookkeeping is fixed, and small. If you want to allocate several objects, write a destroy function which can deallocate all of these. See the `adjlist.c` file in the **igraph** source for an example.

For some functions these mechanisms are simply not flexible enough. These functions should define their own error handlers and restore the error handler before they return.

Fatal errors

In some rare situations, **igraph** may encounter an internal error that cannot be fully handled. In this case, it will call the current fatal error handler. The default fatal error handler simply prints the error and aborts the program.

Fatal error handlers do not return. Typically, they might abort the the program immediately, or in the case of the high-level **igraph** interfaces, they might return to the top level using a `longjmp()`. The fatal error handler is only called when a serious error has occurred, and as a result **igraph** may be in an inconsistent state. The purpose of returning to the top level is to give the user a chance to save their work instead of aborting immediately. However, the program session should be restarted as soon as possible.

Most projects that use **igraph** will use the default fatal error handler.

igraph_fatal_handler_t — The type of **igraph** fatal error handler functions.

```
typedef void igraph_fatal_handler_t (const char *reason, const char *file, int line);
```

Functions of this type *must* not return. Typically they call `abort()` or do a `longjmp()`.

Arguments:

reason: Textual description of the error.

file: The source file in which the error is noticed.

line: The number of the line in the source file which triggered the error

igraph_fatal_handler_abort — Abort program in case of fatal error.

```
IGRAPH_EXPORT igraph_fatal_handler_t igraph_fatal_handler_abort;
```

The default fatal error handler, prints an error message and aborts the program.

IGRAPH_FATAL — Triggers a fatal error.

```
#define IGRAPH_FATAL(reason)
```

This is the usual way of triggering a fatal error from an **igraph** function. It calls `igraph_fatal()`.

Use this macro only in situations where the error cannot be handled. The normal way to handle errors is `IGRAPH_ERROR()`.

Arguments:

reason: The error message.

IGRAPH_FATALF — Triggers a fatal error, with printf-like syntax.

```
#define IGRAPH_FATALF
```

igraph functions can use this macro when a fatal error occurs and want to pass on extra information to the user about what went wrong. It calls `igraph_fatal()` with the proper parameters.

Arguments:

reason: Textual description of the error, a template string with the same syntax as the standard `printf` C library function.

...: The additional arguments to be substituted into the template string.

IGRAPH_ASSERT — igraph-specific replacement for `assert()`.

```
#define IGRAPH_ASSERT(condition)
```

This macro is like the standard `assert()`, but instead of calling `abort()`, it calls `igraph_fatal()`. This allows for returning the control to the calling program, e.g. returning to the top level in a high-level **igraph** interface.

Unlike `assert()`, `IGRAPH_ASSERT()` is not disabled when the `NDEBUG` macro is defined.

This macro is meant for internal use by **igraph**.

Since a typical fatal error handler does a `longjmp()`, avoid using this macro in C++ code. With most compilers, destructor will not be called when `longjmp()` leaves the current scope.

Arguments:

condition: The condition to be checked.

igraph_fatal — Triggers a fatal error.

```
IGRAPH_NORETURN void igraph_fatal(const char *reason, const char *file, int line)
```

This function triggers a fatal error. Typically it is called indirectly through `IGRAPH_FATAL()` or `IGRAPH_ASSERT()`.

Arguments:

reason: Textual description of the error.

file: The source file in which the error was noticed.

line: The number of line in the source file which triggered the error.

igraph_fatalf — Triggers a fatal error, printf-like syntax.


```
IGRAPH_NORETURN void igraph_fatal(const char *reason, const char *file, int line)
```

This function is similar to `igraph_fatal()`, but uses a `printf`-like syntax. It substitutes the additional arguments into the *reason* template string and calls `igraph_fatal()`.

Arguments:

- reason*: Textual description of the error.
- file*: The source file in which the error was noticed.
- line*: The number of line in the source file which triggered the error.
- ...*: The additional arguments to be substituted into the template string.

Error handling and threads

It is likely that the **igraph** error handling method is *not* thread-safe, mainly because of the static global stack which is used to store the address of the temporarily allocated objects. This issue might be addressed in a later version of **igraph**.

Chapter 6. Memory (de)allocation

igraph_malloc — Allocate memory that can be safely deallocated by igraph functions.

```
void *igraph_malloc(size_t n);
```

Some igraph functions, such as `igraph_vector_ptr_free_all()` and `igraph_vector_ptr_destroy_all()` can free memory that may have been allocated by the user. `igraph_malloc()` works exactly like `malloc()` from the C standard library, but it is guaranteed that it can be safely paired with the `free()` function used by igraph internally (which is also user-accessible through `igraph_free()`).

Arguments:

n: Number of bytes to be allocated.

Returns:

Pointer to the piece of allocated memory.

See also:

`igraph_free()`

igraph_free — Deallocate memory that was allocated by igraph functions.

```
void igraph_free(void *p);
```

Some igraph functions return a pointer vector (`igraph_vector_ptr_t`) containing pointers to other igraph or other data types. These data types are dynamically allocated and have to be deallocated manually when the user does not need them any more. This can be done by calling `igraph_free` on them.

Here is a complete example on how to use `igraph_free` properly.

Example 6.1. File `examples/simple/igraph_free.c`

Arguments:

p: Pointer to the piece of memory to be deallocated.

Returns:

Error code, currently always zero, meaning success.

Time complexity: platform dependent, ideally it should be $O(1)$.

See also:

`igraph_malloc()`

Chapter 7. Data structure library: vector, matrix, other data types

About template types

Some of the container types listed in this section are defined for many base types. This is similar to templates in C++ and generics in Ada, but it is implemented via preprocessor macros since the C language cannot handle it. Here is the list of template types and the all base types they currently support:

vector	Vector is currently defined for <code>igraph_real_t</code> , long int (long), char (char), <code>igraph_bool_t</code> (bool). The default is <code>igraph_real_t</code> .
matrix	Matrix is currently defined for <code>igraph_real_t</code> , long int (long), char (char), <code>igraph_bool_t</code> (bool). The default is <code>igraph_real_t</code> .
array3	Array3 is currently defined for <code>igraph_real_t</code> , long int (long), char (char), <code>igraph_bool_t</code> (bool). The default is <code>igraph_real_t</code> .
stack	Stack is currently defined for <code>igraph_real_t</code> , long int (long), char (char), <code>igraph_bool_t</code> (bool). The default is <code>igraph_real_t</code> .
double-ended queue	Dqueue is currently defined for <code>igraph_real_t</code> , long int (long), char (char), <code>igraph_bool_t</code> (bool). The default is <code>igraph_real_t</code> .
heap	Heap is currently defined for <code>igraph_real_t</code> , long int (long), char (char). In addition both maximum and minimum heaps are available. The default is the <code>igraph_real_t</code> maximum heap.

The name of the base element (in parentheses) is added to the function names, except for the default type.

Some examples:

- `igraph_vector_t` is a vector of `igraph_real_t` elements. Its functions are `igraph_vector_init`, `igraph_vector_destroy`, `igraph_vector_sort`, etc.
- `igraph_vector_bool_t` is a vector of `igraph_bool_t` elements, initialize it with `igraph_vector_bool_init`, destroy it with `igraph_vector_bool_destroy`, etc.
- `igraph_heap_t` is a maximum heap with `igraph_real_t` elements. The corresponding functions are `igraph_heap_init`, `igraph_heap_pop`, etc.
- `igraph_heap_min_t` is a minimum heap with `igraph_real_t` elements. The corresponding functions are called `igraph_heap_min_init`, `igraph_heap_min_pop`, etc.
- `igraph_heap_long_t` is a maximum heap with long int elements. Its function have the `igraph_heap_long_` prefix.
- `igraph_heap_min_long_t` is a minimum heap containing long int elements. Its functions have the `igraph_heap_min_long_` prefix.

Note that the `VECTOR` and the `MATRIX` macros can be used on *all* vector and matrix types.

Vectors

About `igraph_vector_t` objects

The `igraph_vector_t` data type is a simple and efficient interface to arrays containing numbers. It is something similar as (but much simpler than) the `vector` template in the C++ standard library.

Vectors are used extensively in **igraph**, all functions which expect or return a list of numbers use `igraph_vector_t` to achieve this.

The `igraph_vector_t` type usually uses $O(n)$ space to store n elements. Sometimes it uses more, this is because vectors can shrink, but even if they shrink, the current implementation does not free a single bit of memory.

The elements in an `igraph_vector_t` object are indexed from zero, we follow the usual C convention here.

The elements of a vector always occupy a single block of memory, the starting address of this memory block can be queried with the `VECTOR` macro. This way, vector objects can be used with standard mathematical libraries, like the GNU Scientific Library.

Constructors and Destructors

`igraph_vector_t` objects have to be initialized before using them, this is analogous to calling a constructor on them. There are a number of `igraph_vector_t` constructors, for your convenience. `igraph_vector_init()` is the basic constructor, it creates a vector of the given length, filled with zeros. `igraph_vector_copy()` creates a new identical copy of an already existing and initialized vector. `igraph_vector_init_copy()` creates a vector by copying a regular C array. `igraph_vector_init_seq()` creates a vector containing a regular sequence with increment one.

`igraph_vector_view()` is a special constructor, it allows you to handle a regular C array as a vector without copying its elements.

If a `igraph_vector_t` object is not needed any more, it should be destroyed to free its allocated memory by calling the `igraph_vector_t` destructor, `igraph_vector_destroy()`.

Note that vectors created by `igraph_vector_view()` are special, you mustn't call `igraph_vector_destroy()` on these.

`igraph_vector_init` — Initializes a vector object (constructor).

```
int igraph_vector_init(igraph_vector_t* v, int long size);
```

Every vector needs to be initialized before it can be used, and there are a number of initialization functions or otherwise called constructors. This function constructs a vector of the given size and initializes each entry to 0. Note that `igraph_vector_null()` can be used to set each element of a vector to zero. However, if you want a vector of zeros, it is much faster to use this function than to create a vector and then invoke `igraph_vector_null()`.

Every vector object initialized by this function should be destroyed (ie. the memory allocated for it should be freed) when it is not needed anymore, the `igraph_vector_destroy()` function is responsible for this.

Arguments:

v: Pointer to a not yet initialized vector object.

size: The size of the vector.

Returns:

error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, the amount of “time” required to allocate $O(n)$ elements, n is the number of elements.

igraph_vector_init_copy — Initializes a vector from an ordinary C array (constructor).

```
int igraph_vector_init_copy(igraph_vector_t *v,  
                           const igraph_real_t *data, long int length)
```

Arguments:

v: Pointer to an uninitialized vector object.

data: A regular C array.

length: The length of the C array.

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system specific, usually $O(\text{length})$.

igraph_vector_init_seq — Initializes a vector with a sequence.

```
int igraph_vector_init_seq(igraph_vector_t *v,  
                           igraph_real_t from, igraph_real_t to);
```

The vector will contain the numbers *from*, *from*+1, ..., *to*.

Arguments:

v: Pointer to an uninitialized vector object.

from: The lower limit in the sequence (inclusive).

to: The upper limit in the sequence (inclusive).

Returns:

Error code: IGRAPH_ENOMEM: out of memory.

Time complexity: $O(n)$, the number of elements in the vector.

igraph_vector_copy — Initializes a vector from another vector object (constructor).

```
int igraph_vector_copy(igraph_vector_t *to,  
                       const igraph_vector_t *from);
```

The contents of the existing vector object will be copied to the new one.

Arguments:

to: Pointer to a not yet initialized vector object.

from: The original vector object to copy.

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, usually $O(n)$, n is the size of the vector.

igraph_vector_destroy — Destroys a vector object.

```
void igraph_vector_destroy(igraph_vector_t* v);
```

All vectors initialized by `igraph_vector_init()` should be properly destroyed by this function. A destroyed vector needs to be reinitialized by `igraph_vector_init()`, `igraph_vector_init_copy()` or another constructor.

Arguments:

v: Pointer to the (previously initialized) vector object to destroy.

Time complexity: operating system dependent.

Initializing elements

igraph_vector_null — Sets each element in the vector to zero.

```
void igraph_vector_null(igraph_vector_t* v);
```

Note that `igraph_vector_init()` sets the elements to zero as well, so it makes no sense to call this function on a just initialized vector. Thus if you want to construct a vector of zeros, then you should use `igraph_vector_init()`.

Arguments:

v: The vector object.

Time complexity: $O(n)$, the size of the vector.

igraph_vector_fill — Fill a vector with a constant element

```
void igraph_vector_fill(igraph_vector_t* v, igraph_real_t e);
```

Sets each element of the vector to the supplied constant.

Arguments:

vector: The vector to work on.

e: The element to fill with.

Time complexity: $O(n)$, the size of the vector.

Accessing elements

The simplest way to access an element of a vector is to use the `VECTOR` macro. This macro can be used both for querying and setting `igraph_vector_t` elements. If you need a function, `igraph_vector_e()` queries and `igraph_vector_set()` sets an element of a vector. `igraph_vector_e_ptr()` returns the address of an element.

`igraph_vector_tail()` returns the last element of a non-empty vector. There is no `igraph_vector_head()` function however, as it is easy to write `VECTOR(v)[0]` instead.

VECTOR — Accessing an element of a vector.

```
#define VECTOR(v)
```

Usage:

```
VECTOR(v)[0]
```

to access the first element of the vector, you can also use this in assignments, like:

```
VECTOR(v)[10]=5;
```

Note that there are no range checks right now. This functionality might be redefined later as a real function instead of a `#define`.

Arguments:

v: The vector object.

Time complexity: $O(1)$.

igraph_vector_e — Access an element of a vector.

```
igraph_real_t igraph_vector_e(const igraph_vector_t* v, long int pos);
```

Arguments:

v: The `igraph_vector_t` object.

pos: The position of the element, the index of the first element is zero.

Returns:

The desired element.

See also:

`igraph_vector_e_ptr()` and the `VECTOR` macro.

Time complexity: $O(1)$.

`igraph_vector_e_ptr` — Get the address of an element of a vector

```
igraph_real_t* igraph_vector_e_ptr(const igraph_vector_t* v, long int pos);
```

Arguments:

v: The `igraph_vector_t` object.

pos: The position of the element, the position of the first element is zero.

Returns:

Pointer to the desired element.

See also:

`igraph_vector_e()` and the `VECTOR` macro.

Time complexity: $O(1)$.

`igraph_vector_set` — Assignment to an element of a vector.

```
void igraph_vector_set(igraph_vector_t* v,  
                      long int pos, igraph_real_t value);
```

Arguments:

v: The `igraph_vector_t` element.

pos: Position of the element to set.

value: New value of the element.

See also:

`igraph_vector_e()`.

igraph_vector_tail — Returns the last element in a vector.

```
igraph_real_t igraph_vector_tail(const igraph_vector_t *v);
```

It is an error to call this function on an empty vector, the result is undefined.

Arguments:

v: The vector object.

Returns:

The last element.

Time complexity: $O(1)$.

Vector views

igraph_vector_view — Handle a regular C array as a igraph_vector_t.

```
const igraph_vector_t*igraph_vector_view(const igraph_vector_t *v,  
                                         const igraph_real_t *data,  
                                         long int length);
```

This is a special `igraph_vector_t` constructor. It allows to handle a regular C array as a `igraph_vector_t` temporarily. Be sure that you *don't* ever call the destructor (`igraph_vector_destroy()`) on objects created by this constructor.

Arguments:

v: Pointer to an uninitialized `igraph_vector_t` object.

data: Pointer, the C array. It may not be `NULL`.

length: The length of the C array.

Returns:

Pointer to the vector object, the same as the *v* parameter, for convenience.

Time complexity: $O(1)$

Copying vectors

igraph_vector_copy_to — Copies the contents of a vector to a C array.

```
void igraph_vector_copy_to(const igraph_vector_t *v, igraph_real_t *to);
```

The C array should have sufficient length.

Arguments:

v: The vector object.

to: The C array.

Time complexity: $O(n)$, n is the size of the vector.

igraph_vector_update — Update a vector from another one.

```
int igraph_vector_update(igraph_vector_t *to,  
                        const igraph_vector_t *from);
```

After this operation the contents of *to* will be exactly the same as that of *from*. The vector *to* will be resized if it was originally shorter or longer than *from*.

Arguments:

to: The vector to update.

from: The vector to update from.

Returns:

Error code.

Time complexity: $O(n)$, the number of elements in *from*.

igraph_vector_append — Append a vector to another one.

```
int igraph_vector_append(igraph_vector_t *to,  
                        const igraph_vector_t *from);
```

The target vector will be resized (except when *from* is empty).

Arguments:

to: The vector to append to.

from: The vector to append, it is kept unchanged.

Returns:

Error code.

Time complexity: $O(n)$, the number of elements in the new vector.

igraph_vector_swap — Swap elements of two vectors.

```
int igraph_vector_swap(igraph_vector_t *v1, igraph_vector_t *v2);
```

The two vectors must have the same length, otherwise an error happens.

Arguments:

v1: The first vector.

v2: The second vector.

Returns:

Error code.

Time complexity: $O(n)$, the length of the vectors.

Exchanging elements

igraph_vector_swap_elements — Swap two elements in a vector.

```
int igraph_vector_swap_elements(igraph_vector_t *v,  
                                long int i, long int j);
```

Note that currently no range checking is performed.

Arguments:

v: The input vector.

i: Index of the first element.

j: Index of the second element (may be the same as the first one).

Returns:

Error code, currently always `IGRAPH_SUCCESS`.

Time complexity: $O(1)$.

igraph_vector_reverse — Reverse the elements of a vector.

```
int igraph_vector_reverse(igraph_vector_t *v);
```

The first element will be last, the last element will be first, etc.

Arguments:

v: The input vector.

Returns:

Error code, currently always `IGRAPH_SUCCESS`.

Time complexity: $O(n)$, the number of elements.

igraph_vector_shuffle — Shuffles a vector in-place using the Fisher-Yates method

```
int igraph_vector_shuffle(igraph_vector_t *v);
```

The Fisher-Yates shuffle ensures that every permutation is equally probable when using a proper randomness source. Of course this does not apply to pseudo-random generators as the cycle of these generators is less than the number of possible permutations of the vector if the vector is long enough.

Arguments:

`v`: The vector object.

Returns:

Error code, currently always `IGRAPH_SUCCESS`.

Time complexity: $O(n)$, n is the number of elements in the vector.

References:

(Fisher & Yates 1963) R. A. Fisher and F. Yates. *Statistical Tables for Biological, Agricultural and Medical Research*. Oliver and Boyd, 6th edition, 1963, page 37.

(Knuth 1998) D. E. Knuth. *Seminumerical Algorithms*, volume 2 of *The Art of Computer Programming*. Addison-Wesley, 3rd edition, 1998, page 145.

Example	7.1.	File	examples/simple/
igraph_fisher_yates_shuffle.c			

Vector operations

igraph_vector_add_constant — Add a constant to the vector.

```
void igraph_vector_add_constant(igraph_vector_t *v, igraph_real_t plus);
```

`plus` is added to every element of `v`. Note that overflow might happen.

Arguments:

`v`: The input vector.

`plus`: The constant to add.

Time complexity: $O(n)$, the number of elements.

igraph_vector_scale — Multiply all elements of a vector by a constant

```
void igraph_vector_scale(igraph_vector_t *v, igraph_real_t by);
```

Arguments:

v: The vector.

by: The constant.

Returns:

Error code. The current implementation always returns with success.

Added in version 0.2.

Time complexity: $O(n)$, the number of elements in a vector.

igraph_vector_add — Add two vectors.

```
int igraph_vector_add(igraph_vector_t *v1,  
                      const igraph_vector_t *v2);
```

Add the elements of *v2* to *v1*, the result is stored in *v1*. The two vectors must have the same length.

Arguments:

v1: The first vector, the result will be stored here.

v2: The second vector, its contents will be unchanged.

Returns:

Error code.

Time complexity: $O(n)$, the number of elements.

igraph_vector_sub — Subtract a vector from another one.

```
int igraph_vector_sub(igraph_vector_t *v1,  
                      const igraph_vector_t *v2);
```

Subtract the elements of *v2* from *v1*, the result is stored in *v1*. The two vectors must have the same length.

Arguments:

v1: The first vector, to subtract from. The result is stored here.

v2: The vector to subtract, it will be unchanged.

Returns:

Error code.

Time complexity: $O(n)$, the length of the vectors.

igraph_vector_mul — Multiply two vectors.

```
int igraph_vector_mul(igraph_vector_t *v1,
                      const igraph_vector_t *v2);
```

v1 will be multiplied by *v2*, elementwise. The two vectors must have the same length.

Arguments:

v1: The first vector, the result will be stored here.

v2: The second vector, it is left unchanged.

Returns:

Error code.

Time complexity: $O(n)$, the number of elements.

igraph_vector_div — Divide a vector by another one.

```
int igraph_vector_div(igraph_vector_t *v1,
                      const igraph_vector_t *v2);
```

v1 is divided by *v2*, elementwise. They must have the same length. If the base type of the vector can generate divide by zero errors then please make sure that *v2* contains no zero if you want to avoid trouble.

Arguments:

v1: The dividend. The result is also stored here.

v2: The divisor, it is left unchanged.

Returns:

Error code.

Time complexity: $O(n)$, the length of the vectors.

igraph_vector_floor — Transform a real vector to a long vector by flooring each element.

```
int igraph_vector_floor(const igraph_vector_t *from, igraph_vector_long_t *to);
```

Flooring means rounding down to the nearest integer.

Arguments:

from: The original real vector object.

to: Pointer to an initialized long vector. The result will be stored here.

Returns:

Error code: IGRAPH_ENOMEM: out of memory

Time complexity: $O(n)$, where n is the number of elements in the vector.

Vector comparisons

igraph_vector_all_e — Are all elements equal?

```
igraph_bool_t igraph_vector_all_e(const igraph_vector_t *lhs,
                                   const igraph_vector_t *rhs);
```

Arguments:

lhs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the *lhs* are all equal to the corresponding elements in *rhs*. Returns 0 (=false) if the lengths of the vectors don't match.

Time complexity: $O(n)$, the length of the vectors.

igraph_vector_all_l — Are all elements less?

```
igraph_bool_t igraph_vector_all_l(const igraph_vector_t *lhs,
                                   const igraph_vector_t *rhs);
```

Arguments:

lhs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the *lhs* are all less than the corresponding elements in *rhs*. Returns 0 (=false) if the lengths of the vectors don't match. If any element is NaN, it will return 0 (=false).

Time complexity: $O(n)$, the length of the vectors.

igraph_vector_all_g — Are all elements greater?


```
igraph_bool_t igraph_vector_all_g(const igraph_vector_t *lhs,  
    const igraph_vector_t *rhs);
```

Arguments:

lhs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the *lhs* are all greater than the corresponding elements in *rhs*. Returns 0 (=false) if the lengths of the vectors don't match. If any element is NaN, it will return 0 (=false).

Time complexity: $O(n)$, the length of the vectors.

igraph_vector_all_le — Are all elements less or equal?

```
igraph_bool_t  
igraph_vector_all_le(const igraph_vector_t *lhs,  
    const igraph_vector_t *rhs);
```

Arguments:

lhs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the *lhs* are all less than or equal to the corresponding elements in *rhs*. Returns 0 (=false) if the lengths of the vectors don't match. If any element is NaN, it will return 0 (=false).

Time complexity: $O(n)$, the length of the vectors.

igraph_vector_all_ge — Are all elements greater or equal?

```
igraph_bool_t  
igraph_vector_all_ge(const igraph_vector_t *lhs,  
    const igraph_vector_t *rhs);
```

Arguments:

lhs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the *lhs* are all greater than or equal to the corresponding elements in *rhs*. Returns 0 (=false) if the lengths of the vectors don't match. If any element is NaN, it will return 0 (=false).

Time complexity: $O(n)$, the length of the vectors.

igraph_vector_lex_cmp — Lexicographical comparison of two vectors.

```
int igraph_vector_lex_cmp(const void *lhs, const void *rhs);
```

If the elements of two vectors match but one is shorter, the shorter one comes first. Thus {1, 3} comes after {1, 2, 3}, but before {1, 3, 4}.

This function is typically used together with `igraph_vector_ptr_sort()`.

Arguments:

lhs: Pointer to a pointer to the first vector (interpreted as an `igraph_vector_t **`).

rhs: Pointer to a pointer to the second vector (interpreted as an `igraph_vector_t **`).

Returns:

-1 if *lhs* is lexicographically smaller, 0 if *lhs* and *rhs* are equal, else 1.

See also:

`igraph_vector_colex_cmp()` to compare vectors starting from the last element.

Time complexity: $O(n)$, the number of elements in the smaller vector.

Example 7.2. File `examples/simple/igraph_vector_ptr_sort.c`

igraph_vector_colex_cmp — Colexicographical comparison of two vectors.

```
int igraph_vector_colex_cmp(const void *lhs, const void *rhs);
```

This comparison starts from the last element of both vectors and moves backward. If the elements of two vectors match but one is shorter, the shorter one comes first. Thus {1, 2} comes after {3, 2, 1}, but before {0, 1, 2}.

This function is typically used together with `igraph_vector_ptr_sort()`.

Arguments:

lhs: Pointer to a pointer to the first vector (interpreted as an `igraph_vector_t **`).

rhs: Pointer to a pointer to the second vector (interpreted as an `igraph_vector_t **`).

Returns:

-1 if *lhs* in reverse order is lexicographically smaller than the reverse of *rhs*, 0 if *lhs* and *rhs* are equal, else 1.

See also:

`igraph_vector_lex_cmp()` to compare vectors starting from the first element.

Time complexity: $O(n)$, the number of elements in the smaller vector.

Example 7.3. File `examples/simple/igraph_vector_ptr_sort.c`

Finding minimum and maximum

`igraph_vector_min` — Smallest element of a vector.

```
igraph_real_t igraph_vector_min(const igraph_vector_t* v);
```

The vector must be non-empty.

Arguments:

v: The input vector.

Returns:

The smallest element of *v*, or NaN if any element is NaN.

Time complexity: $O(n)$, the number of elements.

`igraph_vector_max` — Largest element of a vector.

```
igraph_real_t igraph_vector_max(const igraph_vector_t* v);
```

If the size of the vector is zero, an arbitrary number is returned.

Arguments:

v: The vector object.

Returns:

The maximum element of *v*, or NaN if any element is NaN.

Time complexity: $O(n)$, the number of elements.

`igraph_vector_which_min` — Index of the smallest element.

```
long int igraph_vector_which_min(const igraph_vector_t* v);
```

The vector must be non-empty. If the smallest element is not unique, then the index of the first is returned. If the vector contains NaN values, the index of the first NaN value is returned.

Arguments:

v: The input vector.

Returns:

Index of the smallest element.

Time complexity: $O(n)$, the number of elements.

igraph_vector_which_max — Gives the index of the maximum element of the vector.

```
long int igraph_vector_which_max(const igraph_vector_t* v);
```

If the size of the vector is zero, -1 is returned. If the largest element is not unique, then the index of the first is returned. If the vector contains NaN values, the index of the first NaN value is returned.

Arguments:

v: The vector object.

Returns:

The index of the first maximum element.

Time complexity: $O(n)$, n is the size of the vector.

igraph_vector_minmax — Minimum and maximum elements of a vector.

```
int igraph_vector_minmax(const igraph_vector_t *v,  
                        igraph_real_t *min, igraph_real_t *max);
```

Handy if you want to have both the smallest and largest element of a vector. The vector is only traversed once. The vector must be non-empty. If a vector contains at least one NaN, both *min* and *max* will be NaN.

Arguments:

v: The input vector. It must contain at least one element.

min: Pointer to a base type variable, the minimum is stored here.

max: Pointer to a base type variable, the maximum is stored here.

Returns:

Error code.

Time complexity: $O(n)$, the number of elements.

igraph_vector_which_minmax — Index of the minimum and maximum elements

```
int igraph_vector_which_minmax(const igraph_vector_t *v,
                              long int *which_min, long int *which_max);
```

Handy if you need the indices of the smallest and largest elements. The vector is traversed only once. The vector must be non-empty. If the minimum or maximum is not unique, the index of the first minimum or the first maximum is returned, respectively. If a vector contains at least one NaN, both `which_min` and `which_max` will point to the first NaN value.

Arguments:

`v`: The input vector. It must contain at least one element.

`which_min`: The index of the minimum element will be stored here.

`which_max`: The index of the maximum element will be stored here.

Returns:

Error code.

Time complexity: $O(n)$, the number of elements.

Vector properties

igraph_vector_empty — Decides whether the size of the vector is zero.

```
igraph_bool_t igraph_vector_empty(const igraph_vector_t* v);
```

Arguments:

`v`: The vector object.

Returns:

Non-zero number (true) if the size of the vector is zero and zero (false) otherwise.

Time complexity: $O(1)$.

igraph_vector_size — Returns the size (=length) of the vector.

```
long int igraph_vector_size(const igraph_vector_t* v);
```

Arguments:

v: The vector object

Returns:

The size of the vector.

Time complexity: $O(1)$.

igraph_vector_capacity — Returns the allocated capacity of the vector

```
long int igraph_vector_capacity(const igraph_vector_t*v);
```

Note that this might be different from the size of the vector (as queried by `igraph_vector_size()`), and specifies how many elements the vector can hold, without reallocation.

Arguments:

v: Pointer to the (previously initialized) vector object to query.

Returns:

The allocated capacity.

See also:

`igraph_vector_size()`.

Time complexity: $O(1)$.

igraph_vector_sum — Calculates the sum of the elements in the vector.

```
igraph_real_t igraph_vector_sum(const igraph_vector_t *v);
```

For the empty vector 0.0 is returned.

Arguments:

v: The vector object.

Returns:

The sum of the elements.

Time complexity: $O(n)$, the size of the vector.

igraph_vector_prod — Calculates the product of the elements in the vector.

```
igraph_real_t igraph_vector_prod(const igraph_vector_t *v);
```

For the empty vector one (1) is returned.

Arguments:

v: The vector object.

Returns:

The product of the elements.

Time complexity: $O(n)$, the size of the vector.

igraph_vector_isininterval — Checks if all elements of a vector are in the given

```
igraph_bool_t igraph_vector_isininterval(const igraph_vector_t *v,  
                                         igraph_real_t low,  
                                         igraph_real_t high);
```

interval.

Arguments:

v: The vector object.

low: The lower limit of the interval (inclusive).

high: The higher limit of the interval (inclusive).

Returns:

True (positive integer) if all vector elements are in the interval, false (zero) otherwise. If any element is NaN, it will return 0 (=false).

Time complexity: $O(n)$, the number of elements in the vector.

igraph_vector_maxdifference — The maximum absolute difference of *m1* and *m2*

```
igraph_real_t igraph_vector_maxdifference(const igraph_vector_t *m1,  
                                           const igraph_vector_t *m2);
```

The element with the largest absolute value in $m1 - m2$ is returned. Both vectors must be non-empty, but they not need to have the same length, the extra elements in the longer vector are ignored. If any value is NaN in the shorter vector, the result will be NaN.

Arguments:

m1: The first vector.

m2: The second vector.

Returns:

The maximum absolute difference of $m1$ and $m2$.

Time complexity: $O(n)$, the number of elements in the shorter vector.

igraph_vector_order — Calculate the order of the elements in a vector.

```
int igraph_vector_order(const igraph_vector_t* v,  
                        const igraph_vector_t *v2,  
                        igraph_vector_t* res, igraph_real_t nodes);
```

The smallest element will have order zero, the second smallest order one, etc.

Arguments:

v: The original `igraph_vector_t` object.

v2: A secondary key, another `igraph_vector_t` object.

res: An initialized `igraph_vector_t` object, it will be resized to match the size of *v*. The result of the computation will be stored here.

nodes: Hint, the largest element in *v*.

Returns:

Error code: `IGRAPH_ENOMEM`: out of memory

Time complexity: $O()$

igraph_vector_is_nan — Check for each element if it is NaN.

```
int igraph_vector_is_nan(const igraph_vector_t *v, igraph_vector_bool_t *is_nan)
```

Arguments:

v: The `igraph_vector_t` object to check.

is_nan: The resulting boolean vector indicating for each element whether it is NaN or not.

Returns:

Error code, `IGRAPH_ENOMEM` if there is not enough memory. Note that this function *never* returns an error if the vector *is_nan* will already be large enough. Time complexity: $O(n)$, the number of elements.

igraph_vector_is_any_nan — Check if any element is NaN.


```
igraph_bool_t igraph_vector_is_any_nan(const igraph_vector_t *v);
```

Arguments:

v: The `igraph_vector_t` object to check.

Returns:

1 if any element is NaN, 0 otherwise.

Time complexity: $O(n)$, the number of elements.

Searching for elements

`igraph_vector_contains` — Linear search in a vector.

```
igraph_bool_t igraph_vector_contains(const igraph_vector_t *v,  
                                     igraph_real_t e);
```

Check whether the supplied element is included in the vector, by linear search.

Arguments:

v: The input vector.

e: The element to look for.

Returns:

TRUE if the element is found and FALSE otherwise.

Time complexity: $O(n)$, the length of the vector.

`igraph_vector_search` — Search from a given position

```
igraph_bool_t igraph_vector_search(const igraph_vector_t *v,  
                                   long int from, igraph_real_t what,  
                                   long int *pos);
```

The supplied element *what* is searched in vector *v*, starting from element index *from*. If found then the index of the first instance (after *from*) is stored in *pos*.

Arguments:

v: The input vector.

from: The index to start searching from. No range checking is performed.

what: The element to find.

pos: If not NULL then the index of the found element is stored here.

Returns:

Boolean, TRUE if the element was found, FALSE otherwise.

Time complexity: $O(m)$, the number of elements to search, the length of the vector minus the *from* argument.

igraph_vector_binsearch — Finds an element by binary searching a sorted vector.

```
igraph_bool_t igraph_vector_binsearch(const igraph_vector_t *v,  
                                     igraph_real_t what, long int *pos);
```

It is assumed that the vector is sorted. If the specified element (*what*) is not in the vector, then the position of where it should be inserted (to keep the vector sorted) is returned. If the vector contains any NaN values, the returned value is undefined and *pos* may point to any position.

Arguments:

v: The `igraph_vector_t` object.

what: The element to search for.

pos: Pointer to a long int. This is set to the position of an instance of *what* in the vector if it is present. If *v* does not contain *what* then *pos* is set to the position to which it should be inserted (to keep the the vector sorted of course).

Returns:

Positive integer (true) if *what* is found in the vector, zero (false) otherwise.

Time complexity: $O(\log(n))$, *n* is the number of elements in *v*.

igraph_vector_binsearch_slice — Finds an element by binary searching a sorted slice of a vector.

```
igraph_bool_t igraph_vector_binsearch_slice(const igraph_vector_t *v,  
                                             igraph_real_t what, long int *pos,  
                                             long int start, long int end);
```

It is assumed that the indicated slice of the vector, from *start* to *end*, is sorted. If the specified element (*what*) is not in the slice of the vector, then the position of where it should be inserted (to keep the vector sorted) is returned. If the indicated slice contains any NaN values, the returned value is undefined and *pos* may point to any position within the slice.

Arguments:

v: The `igraph_vector_t` object.

what: The element to search for.

pos: Pointer to a long int. This is set to the position of an instance of *what* in the slice of the vector if it is present. If *v* does not contain *what* then *pos* is set to the position to which it should be inserted (to keep the the vector sorted).

start: The start position of the slice to search (inclusive).

end: The end position of the slice to search (exclusive).

Returns:

Positive integer (true) if *what* is found in the vector, zero (false) otherwise.

Time complexity: $O(\log(n))$, n is the number of elements in the slice of v , i.e. $end - start$.

igraph_vector_binsearch2 — Binary search, without returning the index.

```
igraph_bool_t igraph_vector_binsearch2(const igraph_vector_t *v,
                                       igraph_real_t what);
```

It is assumed that the vector is sorted.

Arguments:

v: The `igraph_vector_t` object.

what: The element to search for.

Returns:

Positive integer (true) if *what* is found in the vector, zero (false) otherwise.

Time complexity: $O(\log(n))$, n is the number of elements in v .

Resizing operations

igraph_vector_clear — Removes all elements from a vector.

```
void igraph_vector_clear(igraph_vector_t* v);
```

This function simply sets the size of the vector to zero, it does not free any allocated memory. For that you have to call `igraph_vector_destroy()`.

Arguments:

v: The vector object.

Time complexity: $O(1)$.

igraph_vector_reserve — Reserves memory for a vector.

```
int igraph_vector_reserve(igraph_vector_t* v, long int size);
```

igraph vectors are flexible, they can grow and shrink. Growing however occasionally needs the data in the vector to be copied. In order to avoid this, you can call this function to reserve space for future growth of the vector.

Note that this function does *not* change the size of the vector. Let us see a small example to clarify things: if you reserve space for 100 elements and the size of your vector was (and still is) 60, then you can surely add additional 40 elements to your vector before it will be copied.

Arguments:

v: The vector object.

size: The new *allocated* size of the vector.

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, should be around $O(n)$, n is the new allocated size of the vector.

igraph_vector_resize — Resize the vector.

```
int igraph_vector_resize(igraph_vector_t* v, long int newsize);
```

Note that this function does not free any memory, just sets the size of the vector to the given one. It can on the other hand allocate more memory if the new size is larger than the previous one. In this case the newly appeared elements in the vector are *not* set to zero, they are uninitialized.

Arguments:

v: The vector object

newsize: The new size of the vector.

Returns:

Error code, IGRAPH_ENOMEM if there is not enough memory. Note that this function *never* returns an error if the vector is made smaller.

See also:

`igraph_vector_reserve()` for allocating memory for future extensions of a vector.
`igraph_vector_resize_min()` for deallocating the unnneded memory for a vector.

Time complexity: $O(1)$ if the new size is smaller, operating system dependent if it is larger. In the latter case it is usually around $O(n)$, n is the new size of the vector.

igraph_vector_resize_min — Deallocate the unused memory of a vector.

```
int igraph_vector_resize_min(igraph_vector_t*v);
```

Note that this function involves additional memory allocation and may result an out-of-memory error.

Arguments:

`v`: Pointer to an initialized vector.

Returns:

Error code.

See also:

`igraph_vector_resize()`, `igraph_vector_reserve()`.

Time complexity: operating system dependent.

igraph_vector_push_back — Appends one element to a vector.

```
int igraph_vector_push_back(igraph_vector_t* v, igraph_real_t e);
```

This function resizes the vector to be one element longer and sets the very last element in the vector to `e`.

Arguments:

`v`: The vector object.

`e`: The element to append to the vector.

Returns:

Error code: `IGRAPH_ENOMEM`: not enough memory.

Time complexity: operating system dependent. What is important is that a sequence of n subsequent calls to this function has time complexity $O(n)$, even if there hadn't been any space reserved for the new elements by `igraph_vector_reserve()`. This is implemented by a trick similar to the C++ vector class: each time more memory is allocated for a vector, the size of the additionally allocated memory is the same as the vector's current length. (We assume here that the time complexity of memory allocation is at most linear.)

igraph_vector_pop_back — Removes and returns the last element of a vector.

```
igraph_real_t igraph_vector_pop_back(igraph_vector_t* v);
```

It is an error to call this function with an empty vector.

Arguments:

`v`: The vector object.

Returns:

The removed last element.

Time complexity: $O(1)$.

igraph_vector_insert — Inserts a single element into a vector.

```
int igraph_vector_insert(igraph_vector_t *v, long int pos,
                        igraph_real_t value);
```

Note that this function does not do range checking. Insertion will shift the elements from the position given to the end of the vector one position to the right, and the new element will be inserted in the empty space created at the given position. The size of the vector will increase by one.

Arguments:

v: The vector object.

pos: The position where the new element is to be inserted.

value: The new element to be inserted.

igraph_vector_remove — Removes a single element from a vector.

```
void igraph_vector_remove(igraph_vector_t *v, long int elem);
```

Note that this function does not do range checking.

Arguments:

v: The vector object.

elem: The position of the element to remove.

Time complexity: $O(n - \text{elem})$, n is the number of elements in the vector.

igraph_vector_remove_section — Deletes a section from a vector.

```
void igraph_vector_remove_section(igraph_vector_t *v,
                                  long int from, long int to);
```

Note that this function does not do range checking. The result is undefined if you supply invalid limits.

Arguments:

v: The vector object.

from: The position of the first element to remove.

to: The position of the first element *not* to remove.

Time complexity: $O(n - \text{from})$, n is the number of elements in the vector.

Sorting

igraph_vector_sort — Sorts the elements of the vector into ascending order.

```
void igraph_vector_sort(igraph_vector_t *v);
```

If the vector contains any NaN values, the resulting ordering of NaN values is undefined and may appear anywhere in the vector.

Arguments:

v: Pointer to an initialized vector object.

Time complexity: $O(n \log n)$ for n elements.

igraph_vector_reverse_sort — Sorts the elements of the vector into descending order.

```
void igraph_vector_reverse_sort(igraph_vector_t *v);
```

If the vector contains any NaN values, the resulting ordering of NaN values is undefined and may appear anywhere in the vector.

Arguments:

v: Pointer to an initialized vector object.

Time complexity: $O(n \log n)$ for n elements.

Set operations on sorted vectors

igraph_vector_intersect_sorted — Calculates the intersection of two sorted vectors

```
int igraph_vector_intersect_sorted(const igraph_vector_t *v1,  
                                   const igraph_vector_t *v2, igraph_vector_t *result);
```

The elements that are contained in both vectors are stored in the result vector. All three vectors must be initialized.

Instead of the naive intersection which takes $O(n)$, this function uses the set intersection method of Ricardo Baeza-Yates, which is more efficient when one of the vectors is significantly smaller than the other, and gives similar performance on average when the two vectors are equal.

The algorithm keeps the multiplicities of the elements: if an element appears k_1 times in the first vector and k_2 times in the second, the result will include that element $\min(k_1, k_2)$ times.

Reference: Baeza-Yates R: A fast set intersection algorithm for sorted sequences. In: Lecture Notes in Computer Science, vol. 3109/2004, pp. 400–408, 2004. Springer Berlin/Heidelberg. ISBN: 978-3-540-22341-2.

Arguments:

v1: the first vector

v2: the second vector

result: the result vector, which will also be sorted.

Time complexity: $O(m \log(n))$ where m is the size of the smaller vector and n is the size of the larger one.

igraph_vector_difference_sorted — Calculates the difference between two sorted vectors (considered as sets)

```
int igraph_vector_difference_sorted(const igraph_vector_t *v1,
                                   const igraph_vector_t *v2, igraph_vector_t *result);
```

The elements that are contained in only the first vector but not the second are stored in the result vector. All three vectors must be initialized.

Arguments:

v1: the first vector

v2: the second vector

result: the result vector

Pointer vectors (igraph_vector_ptr_t)

The `igraph_vector_ptr_t` data type is very similar to the `igraph_vector_t` type, but it stores generic pointers instead of real numbers.

This type has the same space complexity as `igraph_vector_t`, and most implemented operations work the same way as for `igraph_vector_t`.

This type is mostly used to pass to or receive from a set of graphs to some **igraph** functions, such as `igraph_decompose()`, which decomposes a graph to connected components.

The same `VECTOR` macro used for ordinary vectors can be used for pointer vectors as well, please note that a typeless generic pointer will be provided by this macro and you may need to cast it to a specific pointer before starting to work with it.

Pointer vectors may have an associated item destructor function which takes a pointer and returns nothing. The item destructor will be called on each item in the pointer vector when it is destroyed by `igraph_vector_ptr_destroy()` or `igraph_vector_ptr_destroy_all()`, or when its elements are freed by `igraph_vector_ptr_free_all()`. Note that the semantics of an item destructor does not coincide with C++ destructors; for instance, when a pointer vector is resized to a smaller size, the extra items will *not* be destroyed automatically! Nevertheless, item destructors may become handy in many cases; for instance, a vector of graphs generated by `igraph_decompose()`

can be destroyed with a single call to `igraph_vector_ptr_destroy_all()` if the item destructor is set to `igraph_destroy()`.

igraph_vector_ptr_init — Initialize a pointer vector (constructor).

```
int igraph_vector_ptr_init(igraph_vector_ptr_t* v, int long size);
```

This is the constructor of the pointer vector data type. All pointer vectors constructed this way should be destroyed via calling `igraph_vector_ptr_destroy()`.

Arguments:

v: Pointer to an uninitialized `igraph_vector_ptr_t` object, to be created.

size: Integer, the size of the pointer vector.

Returns:

Error code: `IGRAPH_ENOMEM` if out of memory

Time complexity: operating system dependent, the amount of “time” required to allocate *size* elements.

igraph_vector_ptr_copy — Copy a pointer vector (constructor).

```
int igraph_vector_ptr_copy(igraph_vector_ptr_t *to, const igraph_vector_ptr_t *
```

This function creates a pointer vector by copying another one. This is shallow copy, only the pointers in the vector will be copied.

It is potentially dangerous to copy a pointer vector with an associated item destructor. The copied vector will inherit the item destructor, which may cause problems when both vectors are destroyed as the items might get destroyed twice. Make sure you know what you are doing when copying a pointer vector with an item destructor, or unset the item destructor on one of the vectors later.

Arguments:

to: Pointer to an uninitialized pointer vector object.

from: A pointer vector object.

Returns:

Error code: `IGRAPH_ENOMEM` if out of memory

Time complexity: $O(n)$ if allocating memory for n elements can be done in $O(n)$ time.

igraph_vector_ptr_destroy — Destroys a pointer vector.

```
void igraph_vector_ptr_destroy(igraph_vector_ptr_t* v);
```

The destructor for pointer vectors.

Arguments:

v: Pointer to the pointer vector to destroy.

Time complexity: operating system dependent, the “time” required to deallocate $O(n)$ bytes, n is the number of elements allocated for the pointer vector (not necessarily the number of elements in the vector).

igraph_vector_ptr_free_all — Frees all the elements of a pointer vector.

```
void igraph_vector_ptr_free_all(igraph_vector_ptr_t* v);
```

If an item destructor is set for this pointer vector, this function will first call the destructor on all elements of the vector and then free all the elements using `igraph_free()`. If an item destructor is not set, the elements will simply be freed.

Arguments:

v: Pointer to the pointer vector whose elements will be freed.

Time complexity: operating system dependent, the “time” required to call the destructor n times and then deallocate $O(n)$ pointers, each pointing to a memory area of arbitrary size. n is the number of elements in the pointer vector.

igraph_vector_ptr_destroy_all — Frees all the elements and destroys the pointer vector.

```
void igraph_vector_ptr_destroy_all(igraph_vector_ptr_t* v);
```

This function is equivalent to `igraph_vector_ptr_free_all()` followed by `igraph_vector_ptr_destroy()`.

Arguments:

v: Pointer to the pointer vector to destroy.

Time complexity: operating system dependent, the “time” required to deallocate $O(n)$ pointers, each pointing to a memory area of arbitrary size, plus the “time” required to deallocate $O(n)$ bytes, n being the number of elements allocated for the pointer vector (not necessarily the number of elements in the vector).

igraph_vector_ptr_size — Gives the number of elements in the pointer vector.

```
long int igraph_vector_ptr_size(const igraph_vector_ptr_t* v);
```

Arguments:

`v`: The pointer vector object.

Returns:

The size of the object, i.e. the number of pointers stored.

Time complexity: $O(1)$.

igraph_vector_ptr_clear — Removes all elements from a pointer vector.

```
void igraph_vector_ptr_clear(igraph_vector_ptr_t* v);
```

This function resizes a pointer to vector to zero length. Note that the pointed objects are *not* deallocated, you should call `igraph_free()` on them, or make sure that their allocated memory is freed in some other way, you'll get memory leaks otherwise. If you have set up an item destructor earlier, the destructor will be called on every element.

Note that the current implementation of this function does *not* deallocate the memory required for storing the pointers, so making a pointer vector smaller this way does not give back any memory. This behavior might change in the future.

Arguments:

`v`: The pointer vector to clear.

Time complexity: $O(1)$.

igraph_vector_ptr_push_back — Appends an element to the back of a pointer vector.

```
int igraph_vector_ptr_push_back(igraph_vector_ptr_t* v, void* e);
```

Arguments:

`v`: The pointer vector.

`e`: The new element to include in the pointer vector.

Returns:

Error code.

See also:

`igraph_vector_push_back()` for the corresponding operation of the ordinary vector type.

Time complexity: $O(1)$ or $O(n)$, n is the number of elements in the vector. The pointer vector implementation ensures that n subsequent `push_back` operations need $O(n)$ time to complete.

igraph_vector_ptr_insert — Inserts a single element into a pointer vector.

```
int igraph_vector_ptr_insert(igraph_vector_ptr_t* v, long int pos, void* e);
```

Note that this function does not do range checking. Insertion will shift the elements from the position given to the end of the vector one position to the right, and the new element will be inserted in the empty space created at the given position. The size of the vector will increase by one.

Arguments:

v: The pointer vector object.

pos: The position where the new element is inserted.

e: The inserted element

igraph_vector_ptr_e — Access an element of a pointer vector.

```
void *igraph_vector_ptr_e(const igraph_vector_ptr_t* v, long int pos);
```

Arguments:

v: Pointer to a pointer vector.

pos: The index of the pointer to return.

Returns:

The pointer at *pos* position.

Time complexity: $O(1)$.

igraph_vector_ptr_set — Assign to an element of a pointer vector.

```
void igraph_vector_ptr_set(igraph_vector_ptr_t* v, long int pos, void* value);
```

Arguments:

v: Pointer to a pointer vector.

pos: The index of the pointer to update.

value: The new pointer to set in the vector.

Time complexity: $O(1)$.

igraph_vector_ptr_resize — Resizes a pointer vector.

```
int igraph_vector_ptr_resize(igraph_vector_ptr_t* v, long int newsize);
```

Note that if a vector is made smaller the pointed object are not deallocated by this function and the item destructor is not called on the extra elements.

Arguments:

v: A pointer vector.

newsize: The new size of the pointer vector.

Returns:

Error code.

Time complexity: $O(1)$ if the vector is made smaller. Operating system dependent otherwise, the amount of “time” needed to allocate the memory for the vector elements.

igraph_vector_ptr_sort — Sorts the pointer vector based on an external comparison function.

```
void igraph_vector_ptr_sort(igraph_vector_ptr_t *v, int (*compar)(const void*, const void*))
```

Sometimes it is necessary to sort the pointers in the vector based on the property of the element being referenced by the pointer. This function allows us to sort the vector based on an arbitrary external comparison function which accepts two `void *` pointers `p1` and `p2` and returns an integer less than, equal to or greater than zero if the first argument is considered to be respectively less than, equal to, or greater than the second. `p1` and `p2` will point to the pointer in the vector, so they have to be double-dereferenced if one wants to get access to the underlying object the address of which is stored in `v`.

Arguments:

v: The pointer vector to be sorted.

compar: A `qsort`-compatible comparison function. It must take pointers to the elements of the pointer vector. For example, if the pointer vector contains `igraph_vector_t *` pointers, then the comparison function must interpret its arguments as `igraph_vector_t **`.

Example 7.4. File `examples/simple/igraph_vector_ptr_sort.c`

igraph_vector_ptr_get_item_destructor — Gets the current item destructor for this pointer vector.

```
igraph_finally_func_t* igraph_vector_ptr_get_item_destructor(const igraph_vector_ptr_t *v)
```

The item destructor is a function which will be called on every non-null pointer stored in this vector when `igraph_vector_ptr_destroy()`, `igraph_vector_ptr_destroy_all()` or `igraph_vector_ptr_free_all()` is called.

Returns:

The current item destructor.

Time complexity: $O(1)$.

igraph_vector_ptr_set_item_destructor — Sets the item destructor for this pointer vector.

```
igraph_finally_func_t* igraph_vector_ptr_set_item_destructor(  
    igraph_vector_ptr_t *v, igraph_finally_func_t *func);
```

The item destructor is a function which will be called on every non-null pointer stored in this vector when `igraph_vector_ptr_destroy()`, `igraph_vector_ptr_destroy_all()` or `igraph_vector_ptr_free_all()` is called.

Returns:

The old item destructor.

Time complexity: $O(1)$.

IGRAPH_VECTOR_PTR_SET_ITEM_DESTRUCTOR — Sets the item destructor for this pointer vector (macro version).

```
#define IGRAPH_VECTOR_PTR_SET_ITEM_DESTRUCTOR(v, func)
```

This macro is expanded to `igraph_vector_ptr_set_item_destructor()`, the only difference is that the second argument is automatically cast to an `igraph_finally_func_t*`. The cast is necessary in most cases as the destructor functions we use (such as `igraph_vector_destroy()`) take a pointer to some concrete igraph data type, while `igraph_finally_func_t` expects `void*`.

Matrices

About igraph_matrix_t objects

This type is just an interface to `igraph_vector_t`.

The `igraph_matrix_t` type usually stores n elements in $O(n)$ space, but not always. See the documentation of the vector type.

Matrix constructors and destructors

igraph_matrix_init — Initializes a matrix.

```
int igraph_matrix_init(igraph_matrix_t *m, long int nrow, long int ncol);
```

Every matrix needs to be initialized before using it. This is done by calling this function. A matrix has to be destroyed if it is not needed any more; see `igraph_matrix_destroy()`.

Arguments:

m: Pointer to a not yet initialized matrix object to be initialized.

nrow: The number of rows in the matrix.

ncol: The number of columns in the matrix.

Returns:

Error code.

Time complexity: usually $O(n)$, n is the number of elements in the matrix.

igraph_matrix_copy — Copies a matrix.

```
int igraph_matrix_copy(igraph_matrix_t *to, const igraph_matrix_t *from);
```

Creates a matrix object by copying from an existing matrix.

Arguments:

to: Pointer to an uninitialized matrix object.

from: The initialized matrix object to copy.

Returns:

Error code, `IGRAPH_ENOMEM` if there isn't enough memory to allocate the new matrix.

Time complexity: $O(n)$, the number of elements in the matrix.

igraph_matrix_destroy — Destroys a matrix object.

```
void igraph_matrix_destroy(igraph_matrix_t *m);
```

This function frees all the memory allocated for a matrix object. The destroyed object needs to be reinitialized before using it again.

Arguments:

m: The matrix to destroy.

Time complexity: operating system dependent.

Initializing elements

igraph_matrix_null — Sets all elements in a matrix to zero.

```
void igraph_matrix_null(igraph_matrix_t *m);
```

Arguments:

m: Pointer to an initialized matrix object.

Time complexity: $O(n)$, n is the number of elements in the matrix.

igraph_matrix_fill — Fill with an element.

```
void igraph_matrix_fill(igraph_matrix_t *m, igraph_real_t e);
```

Set the matrix to a constant matrix.

Arguments:

m: The input matrix.

e: The element to set.

Time complexity: $O(mn)$, the number of elements.

Copying matrices

igraph_matrix_copy_to — Copies a matrix to a regular C array.

```
void igraph_matrix_copy_to(const igraph_matrix_t *m, igraph_real_t *to);
```

The matrix is copied columnwise, as this is the format most programs and languages use. The C array should be of sufficient size; there are (of course) no range checks.

Arguments:

m: Pointer to an initialized matrix object.

to: Pointer to a C array; the place to copy the data to.

Returns:

Error code.

Time complexity: $O(n)$, n is the number of elements in the matrix.

igraph_matrix_update — Update from another matrix.

```
int igraph_matrix_update(igraph_matrix_t *to,  
                        const igraph_matrix_t *from);
```

This function replicates *from* in the matrix *to*. Note that *to* must be already initialized.

Arguments:

to: The result matrix.

from: The matrix to replicate; it is left unchanged.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_swap — Swap two matrices.

```
int igraph_matrix_swap(igraph_matrix_t *m1, igraph_matrix_t *m2);
```

The contents of the two matrices will be swapped. They must have the same dimensions.

Arguments:

m1: The first matrix.

m2: The second matrix.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements in the matrices.

Accessing elements of a matrix

MATRIX — Accessing an element of a matrix.

```
#define MATRIX(m,i,j)
```

Note that there are no range checks right now. This functionality might be redefined as a proper function later.

Arguments:

m: The matrix object.

i: The index of the row, starting with zero.

j: The index of the column, starting with zero.

Time complexity: $O(1)$.

igraph_matrix_e — Extract an element from a matrix.

```
igraph_real_t igraph_matrix_e(const igraph_matrix_t *m,  
                               long int row, long int col);
```

Use this if you need a function for some reason and cannot use the `MATRIX` macro. Note that no range checking is performed.

Arguments:

m: The input matrix.

row: The row index.

col: The column index.

Returns:

The element in the given row and column.

Time complexity: $O(1)$.

igraph_matrix_e_ptr — Pointer to an element of a matrix.

```
igraph_real_t* igraph_matrix_e_ptr(const igraph_matrix_t *m,
                                   long int row, long int col);
```

The function returns a pointer to an element. No range checking is performed.

Arguments:

m: The input matrix.

row: The row index.

col: The column index.

Returns:

Pointer to the element in the given row and column.

Time complexity: $O(1)$.

igraph_matrix_set — Set an element.

```
void igraph_matrix_set(igraph_matrix_t* m, long int row, long int col,
                       igraph_real_t value);
```

Set an element of a matrix. No range checking is performed.

Arguments:

m: The input matrix.

row: The row index.

col: The column index.

value: The new value of the element.

Time complexity: $O(1)$.

Operations on rows and columns

igraph_matrix_get_row — Extract a row.

```
int igraph_matrix_get_row(const igraph_matrix_t *m,
```

```
igraph_vector_t *res, long int index);
```

Extract a row from a matrix and return it as a vector.

Arguments:

m: The input matrix.
res: Pointer to an initialized vector; it will be resized if needed.
index: The index of the row to select.

Returns:

Error code.

Time complexity: $O(n)$, the number of columns in the matrix.

igraph_matrix_get_col — Select a column.

```
int igraph_matrix_get_col(const igraph_matrix_t *m,  
                          igraph_vector_t *res,  
                          long int index);
```

Extract a column of a matrix and return it as a vector.

Arguments:

m: The input matrix.
res: The result will be stored in this vector. It should be initialized and will be resized as needed.
index: The index of the column to select.

Returns:

Error code.

Time complexity: $O(n)$, the number of rows in the matrix.

igraph_matrix_set_row — Set a row from a vector.

```
int igraph_matrix_set_row(igraph_matrix_t *m,  
                          const igraph_vector_t *v, long int index);
```

Sets the elements of a row with the given vector. This has the effect of setting row *index* to have the elements in the vector *v*. The length of the vector and the number of columns in the matrix must match, otherwise an error is triggered.

Arguments:

m: The input matrix.
v: The vector containing the new elements of the row.
index: Index of the row to set.

Returns:

Error code.

Time complexity: $O(n)$, the number of columns in the matrix.

igraph_matrix_set_col — Set a column from a vector.

```
int igraph_matrix_set_col(igraph_matrix_t *m,
                          const igraph_vector_t *v, long int index);
```

Sets the elements of a column with the given vector. In effect, column `index` will be set with elements from the vector `v`. The length of the vector and the number of rows in the matrix must match, otherwise an error is triggered.

Arguments:

m: The input matrix.

v: The vector containing the new elements of the column.

index: Index of the column to set.

Returns:

Error code.

Time complexity: $O(m)$, the number of rows in the matrix.

igraph_matrix_swap_rows — Swap two rows.

```
int igraph_matrix_swap_rows(igraph_matrix_t *m,
                             long int i, long int j);
```

Swap two rows in the matrix.

Arguments:

m: The input matrix.

i: The index of the first row.

j: The index of the second row.

Returns:

Error code.

Time complexity: $O(n)$, the number of columns.

igraph_matrix_swap_cols — Swap two columns.

```
int igraph_matrix_swap_cols(igraph_matrix_t *m,
                             long int i, long int j);
```

Swap two columns in the matrix.

Arguments:

m: The input matrix.

i: The index of the first column.

j: The index of the second column.

Returns:

Error code.

Time complexity: $O(m)$, the number of rows.

igraph_matrix_select_rows — Select some rows of a matrix.

```
int igraph_matrix_select_rows(const igraph_matrix_t *m,
                             igraph_matrix_t *res,
                             const igraph_vector_t *rows);
```

This function selects some rows of a matrix and returns them in a new matrix. The result matrix should be initialized before calling the function.

Arguments:

m: The input matrix.

res: The result matrix. It should be initialized and will be resized as needed.

rows: Vector; it contains the row indices (starting with zero) to extract. Note that no range checking is performed.

Returns:

Error code.

Time complexity: $O(nm)$, n is the number of rows, m the number of columns of the result matrix.

igraph_matrix_select_cols — Select some columns of a matrix.

```
int igraph_matrix_select_cols(const igraph_matrix_t *m,
                              igraph_matrix_t *res,
                              const igraph_vector_t *cols);
```

This function selects some columns of a matrix and returns them in a new matrix. The result matrix should be initialized before calling the function.

Arguments:

m: The input matrix.

res: The result matrix. It should be initialized and will be resized as needed.

cols: Vector; it contains the column indices (starting with zero) to extract. Note that no range checking is performed.

Returns:

Error code.

Time complexity: $O(nm)$, n is the number of rows, m the number of columns of the result matrix.

igraph_matrix_select_rows_cols — Select some rows and columns of a matrix.

```
int igraph_matrix_select_rows_cols(const igraph_matrix_t *m,
    igraph_matrix_t *res,
    const igraph_vector_t *rows,
    const igraph_vector_t *cols);
```

This function selects some rows and columns of a matrix and returns them in a new matrix. The result matrix should be initialized before calling the function.

Arguments:

m: The input matrix.

res: The result matrix. It should be initialized and will be resized as needed.

rows: Vector; it contains the row indices (starting with zero) to extract. Note that no range checking is performed.

cols: Vector; it contains the column indices (starting with zero) to extract. Note that no range checking is performed.

Returns:

Error code.

Time complexity: $O(nm)$, n is the number of rows, m the number of columns of the result matrix.

Matrix operations

igraph_matrix_add_constant — Add a constant to every element.

```
void igraph_matrix_add_constant(igraph_matrix_t *m, igraph_real_t plus);
```

Arguments:

m: The input matrix.

plus: The constant to add.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_scale — Multiplies each element of the matrix by a constant.

```
void igraph_matrix_scale(igraph_matrix_t *m, igraph_real_t by);
```

Arguments:

m: The matrix.

by: The constant.

Added in version 0.2.

Time complexity: $O(n)$, the number of elements in the matrix.

igraph_matrix_add — Add two matrices.

```
int igraph_matrix_add(igraph_matrix_t *m1,  
                      const igraph_matrix_t *m2);
```

Add *m2* to *m1*, and store the result in *m1*. The dimensions of the matrices must match.

Arguments:

m1: The first matrix; the result will be stored here.

m2: The second matrix; it is left unchanged.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_sub — Difference of two matrices.

```
int igraph_matrix_sub(igraph_matrix_t *m1,  
                      const igraph_matrix_t *m2);
```

Subtract *m2* from *m1* and store the result in *m1*. The dimensions of the two matrices must match.

Arguments:

m1: The first matrix; the result is stored here.

m2: The second matrix; it is left unchanged.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_mul_elements — Elementwise multiplication.

```
int igraph_matrix_mul_elements(igraph_matrix_t *m1,  
                               const igraph_matrix_t *m2);
```

Multiply $m1$ by $m2$ elementwise and store the result in $m1$. The dimensions of the two matrices must match.

Arguments:

$m1$: The first matrix; the result is stored here.

$m2$: The second matrix; it is left unchanged.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_div_elements — Elementwise division.

```
int igraph_matrix_div_elements(igraph_matrix_t *m1,  
                               const igraph_matrix_t *m2);
```

Divide $m1$ by $m2$ elementwise and store the result in $m1$. The dimensions of the two matrices must match.

Arguments:

$m1$: The dividend. The result is store here.

$m2$: The divisor. It is left unchanged.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_sum — Sum of elements.

```
igraph_real_t igraph_matrix_sum(const igraph_matrix_t *m);
```

Returns the sum of the elements of a matrix.

Arguments:

m : The input matrix.

Returns:

The sum of the elements.

Time complexity: $O(mn)$, the number of elements in the matrix.

igraph_matrix_prod — Product of the elements.

```
igraph_real_t igraph_matrix_prod(const igraph_matrix_t *m);
```

Note this function can result in overflow easily, even for not too big matrices.

Arguments:

m: The input matrix.

Returns:

The product of the elements.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_rowsum — Rowwise sum.

```
int igraph_matrix_rowsum(const igraph_matrix_t *m,  
                        igraph_vector_t *res);
```

Calculate the sum of the elements in each row.

Arguments:

m: The input matrix.

res: Pointer to an initialized vector; the result is stored here. It will be resized if necessary.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements in the matrix.

igraph_matrix_colsum — Columnwise sum.

```
int igraph_matrix_colsum(const igraph_matrix_t *m,  
                        igraph_vector_t *res);
```

Calculate the sum of the elements in each column.

Arguments:

m: The input matrix.

res: Pointer to an initialized vector; the result is stored here. It will be resized if necessary.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements in the matrix.

igraph_matrix_transpose — Transpose a matrix.

```
int igraph_matrix_transpose(igraph_matrix_t *m);
```

Calculate the transpose of a matrix. Note that the function reallocates the memory used for the matrix.

Arguments:

m: The input (and output) matrix.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements in the matrix.

Matrix comparisons

igraph_matrix_all_e — Are all elements equal?

```
igraph_bool_t igraph_matrix_all_e(const igraph_matrix_t *lhs,  
                                  const igraph_matrix_t *rhs);
```

Arguments:

lhs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the *lhs* are all equal to the corresponding elements in *rhs*. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: $O(nm)$, the size of the matrices.

igraph_matrix_all_l — Are all elements less?

```
igraph_bool_t igraph_matrix_all_l(const igraph_matrix_t *lhs,  
                                  const igraph_matrix_t *rhs);
```

Arguments:

lhs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the *lhs* are all less than the corresponding elements in *rhs*. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: $O(nm)$, the size of the matrices.

igraph_matrix_all_g — Are all elements greater?

```
igraph_bool_t igraph_matrix_all_g(const igraph_matrix_t *lhs,  
                                  const igraph_matrix_t *rhs);
```

Arguments:

lhs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the *lhs* are all greater than the corresponding elements in *rhs*. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: $O(nm)$, the size of the matrices.

igraph_matrix_all_le — Are all elements less or equal?

```
igraph_bool_t  
igraph_matrix_all_le(const igraph_matrix_t *lhs,  
                    const igraph_matrix_t *rhs);
```

Arguments:

lhs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the *lhs* are all less than or equal to the corresponding elements in *rhs*. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: $O(nm)$, the size of the matrices.

igraph_matrix_all_ge — Are all elements greater or equal?

```
igraph_bool_t  
igraph_matrix_all_ge(const igraph_matrix_t *lhs,  
                    const igraph_matrix_t *rhs);
```

Arguments:

lhs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the *lhs* are all greater than or equal to the corresponding elements in *rhs*. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: $O(nm)$, the size of the matrices.

Combining matrices

igraph_matrix_rbind — Combine two matrices rowwise.

```
int igraph_matrix_rbind(igraph_matrix_t *to,
                        const igraph_matrix_t *from);
```

This function places the rows of *from* below the rows of *to* and stores the result in *to*. The number of columns in the two matrices must match.

Arguments:

to: The upper matrix; the result is also stored here.

from: The lower matrix. It is left unchanged.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements in the newly created matrix.

igraph_matrix_cbind — Combine matrices columnwise.

```
int igraph_matrix_cbind(igraph_matrix_t *to,
                        const igraph_matrix_t *from);
```

This function places the columns of *from* on the right of *to*, and stores the result in *to*.

Arguments:

to: The left matrix; the result is stored here too.

from: The right matrix. It is left unchanged.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements on the new matrix.

Finding minimum and maximum

igraph_matrix_min — Smallest element of a matrix.

```
igraph_real_t igraph_matrix_min(const igraph_matrix_t *m);
```

The matrix must be non-empty.

Arguments:

m: The input matrix.

Returns:

The smallest element of *m*, or NaN if any element is NaN.

Time complexity: $O(mn)$, the number of elements in the matrix.

igraph_matrix_max — Largest element of a matrix.

```
igraph_real_t igraph_matrix_max(const igraph_matrix_t *m);
```

If the matrix is empty, an arbitrary number is returned.

Arguments:

m: The matrix object.

Returns:

The maximum element of *m*, or NaN if any element is NaN.

Added in version 0.2.

Time complexity: $O(mn)$, the number of elements in the matrix.

igraph_matrix_which_min — Indices of the smallest element.

```
int igraph_matrix_which_min(const igraph_matrix_t *m,  
                           long int *i, long int *j);
```

The matrix must be non-empty. If the smallest element is not unique, then the indices of the first such element are returned. If the matrix contains NaN values, the indices of the first NaN value are returned.

Arguments:

m: The matrix.

i: Pointer to a long int. The row index of the minimum is stored here.

j: Pointer to a long int. The column index of the minimum is stored here.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_which_max — Indices of the largest element.

```
int igraph_matrix_which_max(const igraph_matrix_t *m,  
                           long int *i, long int *j);
```

The matrix must be non-empty. If the largest element is not unique, then the indices of the first such element are returned. If the matrix contains NaN values, the indices of the first NaN value are returned.

Arguments:

m: The matrix.

i: Pointer to a long int. The row index of the maximum is stored here.

j: Pointer to a long int. The column index of the maximum is stored here.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_minmax — Minimum and maximum elements of a matrix.

```
int igraph_matrix_minmax(const igraph_matrix_t *m,  
                        igraph_real_t *min, igraph_real_t *max);
```

Handy if you want to have both the smallest and largest element of a matrix. The matrix is only traversed once. The matrix must be non-empty. If a matrix contains at least one NaN, both *min* and *max* will be NaN.

Arguments:

m: The input matrix. It must contain at least one element.

min: Pointer to a base type variable. The minimum is stored here.

max: Pointer to a base type variable. The maximum is stored here.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_which_minmax — Indices of the minimum and maximum elements

```
int igraph_matrix_which_minmax(const igraph_matrix_t *m,  
                               long int *imin, long int *jmin,  
                               long int *imax, long int *jmax);
```

Handy if you need the indices of the smallest and largest elements. The matrix is traversed only once. The matrix must be non-empty. If the minimum or maximum is not unique, the index of the first minimum or the first maximum is returned, respectively. If a matrix contains at least one NaN, both `which_min` and `which_max` will point to the first NaN value.

Arguments:

m: The input matrix.

imin: Pointer to a long int, the row index of the minimum is stored here.

jmin: Pointer to a long int, the column index of the minimum is stored here.

imax: Pointer to a long int, the row index of the maximum is stored here.

jmax: Pointer to a long int, the column index of the maximum is stored here.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements.

Matrix properties

`igraph_matrix_empty` — Check for an empty matrix.

```
igraph_bool_t igraph_matrix_empty(const igraph_matrix_t *m);
```

It is possible to have a matrix with zero rows or zero columns, or even both. This functions checks for these.

Arguments:

m: The input matrix.

Returns:

Boolean, TRUE if the matrix contains zero elements, and FALSE otherwise.

Time complexity: $O(1)$.

`igraph_matrix_isnull` — Check for a null matrix.

```
igraph_bool_t igraph_matrix_isnull(const igraph_matrix_t *m);
```

Checks whether all elements are zero.

Arguments:

m: The input matrix.

Returns:

Boolean, TRUE if m contains only zeros and FALSE otherwise.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_size — The number of elements in a matrix.

```
long int igraph_matrix_size(const igraph_matrix_t *m);
```

Arguments:

m : Pointer to an initialized matrix object.

Returns:

The size of the matrix.

Time complexity: $O(1)$.

igraph_matrix_capacity — Returns the number of elements allocated for a matrix.

```
long int igraph_matrix_capacity(const igraph_matrix_t *m);
```

Note that this might be different from the size of the matrix (as queried by `igraph_matrix_size()`), and specifies how many elements the matrix can hold, without reallocation.

Arguments:

v : Pointer to the (previously initialized) matrix object to query.

Returns:

The allocated capacity.

See also:

`igraph_matrix_size()`, `igraph_matrix_nrow()`, `igraph_matrix_ncol()`.

Time complexity: $O(1)$.

igraph_matrix_nrow — The number of rows in a matrix.

```
long int igraph_matrix_nrow(const igraph_matrix_t *m);
```

Arguments:

m : Pointer to an initialized matrix object.

Returns:

The number of rows in the matrix.

Time complexity: $O(1)$.

igraph_matrix_ncol — The number of columns in a matrix.

```
long int igraph_matrix_ncol(const igraph_matrix_t *m);
```

Arguments:

m: Pointer to an initialized matrix object.

Returns:

The number of columns in the matrix.

Time complexity: $O(1)$.

igraph_matrix_is_symmetric — Check for symmetric matrix.

```
igraph_bool_t igraph_matrix_is_symmetric(const igraph_matrix_t *m);
```

A non-square matrix is not symmetric by definition.

Arguments:

m: The input matrix.

Returns:

Boolean, TRUE if the matrix is square and symmetric, FALSE otherwise.

Time complexity: $O(mn)$, the number of elements. $O(1)$ for non-square matrices.

igraph_matrix_maxdifference — Maximum absolute difference between two matrices.

```
igraph_real_t igraph_matrix_maxdifference(const igraph_matrix_t *m1,  
                                           const igraph_matrix_t *m2);
```

Calculate the maximum absolute difference of two matrices. Both matrices must be non-empty. If their dimensions differ then a warning is given and the comparison is performed by vectors columnwise from both matrices. The remaining elements in the larger vector are ignored.

Arguments:

m1: The first matrix.

m2: The second matrix.

Returns:

The element with the largest absolute value in $m_1 - m_2$.

Time complexity: $O(mn)$, the elements in the smaller matrix.

Searching for elements

igraph_matrix_contains — Search for an element.

```
igraph_bool_t igraph_matrix_contains(const igraph_matrix_t *m,  
                                     igraph_real_t e);
```

Search for the given element in the matrix.

Arguments:

m: The input matrix.

e: The element to search for.

Returns:

Boolean, TRUE if the matrix contains *e*, FALSE otherwise.

Time complexity: $O(mn)$, the number of elements.

igraph_matrix_search — Search from a given position.

```
igraph_bool_t igraph_matrix_search(const igraph_matrix_t *m,  
                                   long int from, igraph_real_t what,  
                                   long int *pos,  
                                   long int *row, long int *col);
```

Search for an element in a matrix and start the search from the given position. The search is performed columnwise.

Arguments:

m: The input matrix.

from: The position to search from, the positions are enumerated columnwise.

what: The element to search for.

pos: Pointer to a long int. If the element is found, then this is set to the position of its first appearance.

row: Pointer to a long int. If the element is found, then this is set to its row index.

col: Pointer to a long int. If the element is found, then this is set to its column index.

Returns:

Boolean, TRUE if the element is found, FALSE otherwise.

Time complexity: $O(mn)$, the number of elements.

Resizing operations

igraph_matrix_resize — Resizes a matrix.

```
int igraph_matrix_resize(igraph_matrix_t *m, long int nrow, long int ncol);
```

This function resizes a matrix by adding more elements to it. The matrix contains arbitrary data after resizing it. That is, after calling this function you cannot expect that element (i,j) in the matrix remains the same as before.

Arguments:

m: Pointer to an already initialized matrix object.

nrow: The number of rows in the resized matrix.

ncol: The number of columns in the resized matrix.

Returns:

Error code.

Time complexity: $O(1)$ if the matrix gets smaller, usually $O(n)$ if it gets larger, n is the number of elements in the resized matrix.

igraph_matrix_resize_min — Deallocates unused memory for a matrix.

```
int igraph_matrix_resize_min(igraph_matrix_t *m);
```

Note that this function might fail if there is not enough memory available.

Also note, that this function leaves the matrix intact, i.e. it does not destroy any of the elements. However, usually it involves copying the matrix in memory.

Arguments:

m: Pointer to an initialized matrix.

Returns:

Error code.

See also:

`igraph_matrix_resize()`.

Time complexity: operating system dependent.

igraph_matrix_add_rows — Adds rows to a matrix.

```
int igraph_matrix_add_rows(igraph_matrix_t *m, long int n);
```

Arguments:

m: The matrix object.

n: The number of rows to add.

Returns:

Error code, IGRAPH_ENOMEM if there isn't enough memory for the operation.

Time complexity: linear with the number of elements of the new, resized matrix.

igraph_matrix_add_cols — Adds columns to a matrix.

```
int igraph_matrix_add_cols(igraph_matrix_t *m, long int n);
```

Arguments:

m: The matrix object.

n: The number of columns to add.

Returns:

Error code, IGRAPH_ENOMEM if there is not enough memory to perform the operation.

Time complexity: linear with the number of elements of the new, resized matrix.

igraph_matrix_remove_row — Remove a row.

```
int igraph_matrix_remove_row(igraph_matrix_t *m, long int row);
```

A row is removed from the matrix.

Arguments:

m: The input matrix.

row: The index of the row to remove.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements in the matrix.

igraph_matrix_remove_col — Removes a column from a matrix.

```
int igraph_matrix_remove_col(igraph_matrix_t *m, long int col);
```

Arguments:

m: The matrix object.

col: The column to remove.

Returns:

Error code, always returns with success.

Time complexity: linear with the number of elements of the new, resized matrix.

Sparse matrices

About igraph_spmatrix_t objects

The `igraph_spmatrix_t` type stores a sparse matrix with the assumption that the number of nonzero elements in the matrix scales linearly with the row or column count of the matrix (so most of the elements are zero). Of course it can store an arbitrary real matrix, but if most of the elements are nonzero, one should use `igraph_matrix_t` instead.

The elements are stored in column compressed format, so the elements in the same column are stored adjacent in the computer's memory. The storage requirement for a sparse matrix is $O(n)$ where n is the number of nonzero elements. Actually it can be a bit larger, see the documentation of the vector type for an explanation.

Sparse matrix constructors and destructors.

`igraph_spmatrix_init` — Initializes a sparse matrix.

```
int igraph_spmatrix_init(igraph_spmatrix_t *m, long int nrow, long int ncol);
```

Every sparse matrix needs to be initialized before using it, this is done by calling this function. A matrix has to be destroyed if it is not needed any more, see `igraph_spmatrix_destroy()`.

Arguments:

m: Pointer to a not yet initialized sparse matrix object to be initialized.

nrow: The number of rows in the matrix.

ncol: The number of columns in the matrix.

Returns:

Error code.

Time complexity: operating system dependent.

`igraph_spmatrix_copy` — Copies a sparse matrix.

```
int igraph_spmatrix_copy(igraph_spmatrix_t *to, const igraph_spmatrix_t *from);
```

Creates a sparse matrix object by copying another one.

Arguments:

to: Pointer to an uninitialized sparse matrix object.

from: The initialized sparse matrix object to copy.

Returns:

Error code, IGRAPH_ENOMEM if there isn't enough memory to allocate the new sparse matrix.

Time complexity: $O(n)$, the number of elements in the matrix.

igraph_spmatrix_destroy — Destroys a sparse matrix object.

```
void igraph_spmatrix_destroy(igraph_spmatrix_t *m);
```

This function frees all the memory allocated for a sparse matrix object. The destroyed object needs to be reinitialized before using it again.

Arguments:

m: The matrix to destroy.

Time complexity: operating system dependent.

Accessing elements of a sparse matrix

igraph_spmatrix_e — Accessing an element of a sparse matrix.

```
igraph_real_t igraph_spmatrix_e(const igraph_spmatrix_t *m,  
                                long int row, long int col);
```

Note that there are no range checks right now.

Arguments:

m: The matrix object.

row: The index of the row, starting with zero.

col: The index of the column, starting with zero.

Time complexity: $O(\log n)$, where n is the number of nonzero elements in the requested column.

igraph_spmatrix_set — Setting an element of a sparse matrix.

```
int igraph_spmatrix_set(igraph_spmatrix_t *m, long int row, long int col,
                        igraph_real_t value);
```

Note that there are no range checks right now.

Arguments:

m: The matrix object.
row: The index of the row, starting with zero.
col: The index of the column, starting with zero.
value: The new value.

Time complexity: $O(\log n)$, where n is the number of nonzero elements in the requested column.

igraph_spmatrix_add_e — Adding a real value to an element of a sparse matrix.

```
int igraph_spmatrix_add_e(igraph_spmatrix_t *m, long int row, long int col,
                           igraph_real_t value);
```

Note that there are no range checks right now. This is implemented to avoid double lookup of a given element in the matrix by using `igraph_spmatrix_e()` and `igraph_spmatrix_set()` consecutively.

Arguments:

m: The matrix object.
row: The index of the row, starting with zero.
col: The index of the column, starting with zero.
value: The value to add.

Time complexity: $O(\log n)$, where n is the number of nonzero elements in the requested column.

Iterating over the non-zero elements of a sparse matrix

The `igraph_spmatrix_iter_t` type represents an iterator that can be used to step over the non-zero elements of a sparse matrix in columnwise order efficiently. In general, you shouldn't modify the elements of the matrix while iterating over it; doing so will probably invalidate the iterator, but there are no checks to prevent you from doing this.

To access the row index of the current element of the iterator, use its `ri` field. Similarly, the `ci` field stores the column index of the current element and the `value` field stores the value of the element.

igraph_spmatrix_iter_create — Creates a sparse matrix iterator corresponding to the given matrix.

```
int igraph_spmatrix_iter_create(igraph_spmatrix_iter_t *mit, const igraph_spmat
```

Arguments:

mit: pointer to the matrix iterator being initialized

m: pointer to the matrix we will be iterating over

Returns:

Error code. The current implementation is always successful.

Time complexity: $O(1)$.

igraph_spmatrix_iter_reset — Resets a sparse matrix iterator.

```
int igraph_spmatrix_iter_reset(igraph_spmatrix_iter_t *mit);
```

After resetting, the iterator will point to the first nonzero element (if any).

Arguments:

mit: pointer to the matrix iterator being reset

Returns:

Error code. The current implementation is always successful.

Time complexity: $O(1)$.

igraph_spmatrix_iter_next — Moves a sparse matrix iterator to the next nonzero element.

```
int igraph_spmatrix_iter_next(igraph_spmatrix_iter_t *mit);
```

You should call this function only if `igraph_spmatrix_iter_end()` returns FALSE (0).

Arguments:

mit: pointer to the matrix iterator being moved

Returns:

Error code. The current implementation is always successful.

Time complexity: $O(1)$.

igraph_spmatrix_iter_end — Checks whether there are more elements in the iterator.

```
igraph_bool_t igraph_spmatrix_iter_end(igraph_spmatrix_iter_t *mit);
```


You should call this function before calling `igraph_spmatrix_iter_next()` to make sure you have more elements in the iterator.

Arguments:

mit: pointer to the matrix iterator being checked

Returns:

TRUE (1) if there are more elements in the iterator, FALSE (0) otherwise.

Time complexity: $O(1)$.

igraph_spmatrix_iter_destroy — Frees the memory used by the iterator.

```
void igraph_spmatrix_iter_destroy(igraph_spmatrix_iter_t *mit);
```

The current implementation does not allocate any memory upon creation, so this function does nothing. However, since there is no guarantee that future implementations will not allocate any memory in `igraph_spmatrix_iter_create()`, you are still required to call this function whenever you are done with the iterator.

Arguments:

mit: pointer to the matrix iterator being destroyed

Time complexity: $O(1)$.

Matrix query operations

igraph_spmatrix_size — The number of elements in a sparse matrix.

```
long int igraph_spmatrix_size(const igraph_spmatrix_t *m);
```

Arguments:

m: Pointer to an initialized sparse matrix object.

Returns:

The size of the matrix.

Time complexity: $O(1)$.

igraph_spmatrix_nrow — The number of rows in a sparse matrix.

```
long int igraph_spmatrix_nrow(const igraph_spmatrix_t *m);
```

Arguments:

m: Pointer to an initialized sparse matrix object.

Returns:

The number of rows in the matrix.

Time complexity: $O(1)$.

igraph_spmatrix_ncol — The number of columns in a sparse matrix.

```
long int igraph_spmatrix_ncol(const igraph_spmatrix_t *m);
```

Arguments:

m: Pointer to an initialized sparse matrix object.

Returns:

The number of columns in the sparse matrix.

Time complexity: $O(1)$.

igraph_spmatrix_count_nonzero — The number of non-zero elements in a sparse matrix.

```
long int igraph_spmatrix_count_nonzero(const igraph_spmatrix_t *m);
```

Arguments:

m: Pointer to an initialized sparse matrix object.

Returns:

The size of the matrix.

Time complexity: $O(1)$.

igraph_spmatrix_max — Returns the maximum element of a matrix.

```
igraph_real_t igraph_spmatrix_max(const igraph_spmatrix_t *m,  
                                   igraph_real_t *ridx, igraph_real_t *cidx);
```

If the matrix is empty, zero is returned.

Arguments:

m: the matrix object.

ridx: the row index of the maximum element if not NULL.

cidx: the column index of the maximum element if not NULL.

Time complexity: $O(n)$, the number of nonzero elements in the matrix.

igraph_spmatrix_rowsums — Calculates the row sums of the matrix.

```
int igraph_spmatrix_rowsums(const igraph_spmatrix_t *m, igraph_vector_t *res);
```

Arguments:

m: The matrix.

res: An initialized `igraph_vector_t`, the result will be stored here. The vector will be resized as needed.

Time complexity: $O(n)$, the number of nonzero elements in the matrix.

igraph_spmatrix_colsums — Calculates the column sums of the matrix.

```
int igraph_spmatrix_colsums(const igraph_spmatrix_t *m, igraph_vector_t *res);
```

Arguments:

m: The matrix.

res: An initialized `igraph_vector_t`, the result will be stored here. The vector will be resized as needed.

Time complexity: $O(n)$, the number of nonzero elements in the matrix.

Matrix operations

igraph_spmatrix_scale — Multiplies each element of the sparse matrix by a constant.

```
void igraph_spmatrix_scale(igraph_spmatrix_t *m, igraph_real_t by);
```

Arguments:

m: The matrix.

by: The constant.

Time complexity: $O(n)$, the number of elements in the matrix.

igraph_spmatrix_add_rows — Adds rows to a sparse matrix.

```
int igraph_spmatrix_add_rows(igraph_spmatrix_t *m, long int n);
```

Arguments:

m: The sparse matrix object.

n: The number of rows to add.

Returns:

Error code.

Time complexity: $O(1)$.

igraph_spmatrix_add_cols — Adds columns to a sparse matrix.

```
int igraph_spmatrix_add_cols(igraph_spmatrix_t *m, long int n);
```

Arguments:

m: The sparse matrix object.

n: The number of columns to add.

Returns:

Error code.

Time complexity: $O(1)$.

igraph_spmatrix_resize — Resizes a sparse matrix.

```
int igraph_spmatrix_resize(igraph_spmatrix_t *m, long int nrow, long int ncol);
```

This function resizes a sparse matrix by adding more elements to it. The matrix retains its data even after resizing it, except for the data which lies outside the new boundaries (if the new size is smaller).

Arguments:

m: Pointer to an already initialized sparse matrix object.

nrow: The number of rows in the resized matrix.

ncol: The number of columns in the resized matrix.

Returns:

Error code.

Time complexity: $O(n)$. n is the number of elements in the old matrix.

Printing sparse matrices

igraph_sparsematrix_print — Prints a sparse matrix.

```
int igraph_sparsematrix_print(const igraph_sparsematrix_t* matrix);
```

Prints a sparse matrix to the standard output. Only the non-zero entries are printed.

Returns:

Error code.

Time complexity: $O(n)$, the number of non-zero elements.

igraph_sparsematrix_fprint — Prints a sparse matrix to the given file.

```
int igraph_sparsematrix_fprint(const igraph_sparsematrix_t* matrix, FILE *file);
```

Prints a sparse matrix to the given file. Only the non-zero entries are printed.

Returns:

Error code.

Time complexity: $O(n)$, the number of non-zero elements.

Sparse matrices, another kind

About sparse matrices

The `igraph_sparsemat_t` data type stores sparse matrices, i.e. matrices in which the majority of the elements are zero.

The data type is essentially a wrapper to some of the functions in the CXSparse library, by Tim Davis, see <http://faculty.cse.tamu.edu/davis/suitesparse.html>

Matrices can be stored in two formats: triplet and column-compressed. The triplet format is intended for sparse matrix initialization, as it is easy to add new (non-zero) elements to it. Most of the computations are done on sparse matrices in column-compressed format, after the user has converted the triplet matrix to column-compressed, via `igraph_sparsemat_compress()`.

Both formats are dynamic, in the sense that new elements can be added to them, possibly resulting the allocation of more memory.

Row and column indices follow the C convention and are zero-based.

Example 7.5. File `examples/simple/igraph_sparsemat.c`

Example 7.6. File `examples/simple/igraph_sparsemat3.c`

Example 7.7. File `examples/simple/igraph_sparsemat4.c`

Example 7.8. File `examples/simple/igraph_sparsemat6.c`

Example 7.9. File `examples/simple/igraph_sparsemat7.c`

Example 7.10. File `examples/simple/igraph_sparsemat8.c`

Creating sparse matrix objects

`igraph_sparsemat_init` — Initializes a sparse matrix, in triplet format.

```
int igraph_sparsemat_init(igraph_sparsemat_t *A, int rows, int cols, int nzmax)
```

This is the most common way to create a sparse matrix, together with the `igraph_sparsemat_entry()` function, which can be used to add the non-zero elements one by one. Once done, the user can call `igraph_sparsemat_compress()` to convert the matrix to column-compressed, to allow computations with it.

The user must call `igraph_sparsemat_destroy()` on the matrix to deallocate the memory, once the matrix is no more needed.

Arguments:

A: Pointer to a not yet initialized sparse matrix.

rows: The number of rows in the matrix.

cols: The number of columns.

nzmax: The maximum number of non-zero elements in the matrix. It is not compulsory to get this right, but it is useful for the allocation of the proper amount of memory.

Returns:

Error code.

Time complexity: TODO.

`igraph_sparsemat_copy` — Copies a sparse matrix.

```
int igraph_sparsemat_copy(igraph_sparsemat_t *to,  
                           const igraph_sparsemat_t *from);
```

Create a sparse matrix object, by copying another one. The source matrix can be either in triplet or column-compressed format.

Exactly the same amount of memory will be allocated to the copy matrix, as it is currently for the original one.

Arguments:

to: Pointer to an uninitialized sparse matrix, the copy will be created here.

from: The sparse matrix to copy.

Returns:

Error code.

Time complexity: $O(n+nz_{\max})$, the number of columns plus the maximum number of non-zero elements.

igraph_sparsemat_realloc — Allocates more (or less) memory for a sparse matrix.

```
int igraph_sparsemat_realloc(igraph_sparsemat_t *A, int nzmax);
```

Sparse matrices automatically allocate more memory, as needed. To control memory allocation, the user can call this function, to allocate memory for a given number of non-zero elements.

Arguments:

A: The sparse matrix, it can be in triplet or column-compressed format.

nzmax: The new maximum number of non-zero elements.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_destroy — Deallocates memory used by a sparse matrix.

```
void igraph_sparsemat_destroy(igraph_sparsemat_t *A);
```

One destroyed, the sparse matrix must be initialized again, before calling any other operation on it.

Arguments:

A: The sparse matrix to destroy.

Time complexity: $O(1)$.

igraph_sparsemat_eye — Creates a sparse identity matrix.

```
int igraph_sparsemat_eye(igraph_sparsemat_t *A, int n, int nzmax,  
                        igraph_real_t value,  
                        igraph_bool_t compress);
```

Arguments:

- A*: An uninitialized sparse matrix, the result is stored here.
- n*: The number of rows and number of columns in the matrix.
- nzmax*: The maximum number of non-zero elements, this essentially gives the amount of memory that will be allocated for matrix elements.
- value*: The value to store in the diagonal.
- compress*: Whether to create a column-compressed matrix. If false, then a triplet matrix is created.

Returns:

Error code.

Time complexity: $O(n)$.

igraph_sparsemat_diag — Creates a sparse diagonal matrix.

```
int igraph_sparsemat_diag(igraph_sparsemat_t *A, int nzmax,  
                          const igraph_vector_t *values,  
                          igraph_bool_t compress);
```

Arguments:

- A*: An uninitialized sparse matrix, the result is stored here.
- nzmax*: The maximum number of non-zero elements, this essentially gives the amount of memory that will be allocated for matrix elements.
- values*: The values to store in the diagonal, the size of the matrix defined by the length of this vector.
- compress*: Whether to create a column-compressed matrix. If false, then a triplet matrix is created.

Returns:

Error code.

Time complexity: $O(n)$, the length of the diagonal vector.

igraph_sparsemat_view — Initialize a sparse matrix and set all parameters.

```
int igraph_sparsemat_view(igraph_sparsemat_t *A, int nzmax, int m, int n,
```



```
int *p, int *i, double *x, int nz);
```

This function can be used to temporarily handle existing sparse matrix data, usually created by another software library, as an `igraph_sparsemat_t` object, and thus avoid unnecessary copying. It supports data stored in either the compressed sparse column format, or the (i, j, x) triplet format where i and j are the matrix indices of a non-zero element, and x is its value.

The compressed sparse column (or row) format is commonly used to represent sparse matrix data. It consists of three vectors, the p column pointers, the i row indices, and the x values. $p[k]$ is the number of non-zero entries in matrix columns $k-1$ and lower. $p[0]$ is always zero and $p[n]$ is always the total number of non-zero entries in the matrix. $i[1]$ is the row index of the 1-th stored element, while $x[1]$ is its value. If a matrix element with indices (j, k) is explicitly stored, it must be located between positions $p[k]$ and $p[k+1] - 1$ (inclusive) in the i and x vectors.

Do not call `igraph_sparsemat_destroy()` on a sparse matrix created with this function. Instead, `igraph_free()` must be called on the `cs` field of A to free the storage allocated by this function.

Warning: Matrices created with this function must not be used with functions that may reallocate the underlying storage, such as `igraph_sparsemat_entry()`.

Arguments:

- A:** The non-initialized sparse matrix.
- nzmax*:** The maximum number of entries, typically the actual number of entries.
- m*:** The number of matrix rows.
- n*:** The number of matrix columns.
- p*:** For a compressed matrix, this is the column pointer vector, and must be of size $n+1$. For a triplet format matrix, it is a vector of column indices and must be of size *nzmax*.
- i*:** The row vector. This should contain the row indices of the elements in x . It must be of size *nzmax*.
- x*:** The values of the non-zero elements of the sparse matrix. It must be of size *nzmax*.
- nz*:** For a compressed matrix, it must be -1. For a triplet format matrix, it must contain the number of entries.

Returns:

Error code.

Time complexity: $O(1)$.

Query properties of a sparse matrix

`igraph_sparsemat_index` — Extracts a submatrix or a single element.

```
int igraph_sparsemat_index(const igraph_sparsemat_t *A,
                           const igraph_vector_int_t *p,
                           const igraph_vector_int_t *q,
                           igraph_sparsemat_t *res,
```

```
igraph_real_t *constres);
```

This function indexes into a sparse matrix. It serves two purposes. First, it can extract submatrices from a sparse matrix. Second, as a special case, it can extract a single element from a sparse matrix.

Arguments:

- A*: The input matrix, it must be in column-compressed format.
- p*: An integer vector, or a null pointer. The selected row index or indices. A null pointer selects all rows.
- q*: An integer vector, or a null pointer. The selected column index or indices. A null pointer selects all columns.
- res*: Pointer to an uninitialized sparse matrix, or a null pointer. If not a null pointer, then the selected submatrix is stored here.
- constres*: Pointer to a real variable or a null pointer. If not a null pointer, then the first non-zero element in the selected submatrix is stored here, if there is one. Otherwise zero is stored here. This behavior is handy if one wants to select a single entry from the matrix.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_nrow — Number of rows.

```
long int igraph_sparsemat_nrow(const igraph_sparsemat_t *A);
```

Arguments:

- A*: The input matrix, in triplet or column-compressed format.

Returns:

The number of rows in the *A* matrix.

Time complexity: $O(1)$.

igraph_sparsemat_ncol — Number of columns.

```
long int igraph_sparsemat_ncol(const igraph_sparsemat_t *A);
```

Arguments:

- A*: The input matrix, in triplet or column-compressed format.

Returns:

The number of columns in the A matrix.

Time complexity: $O(1)$.

igraph_sparsemat_type — Type of a sparse matrix (triplet or column-compressed).

```
igraph_sparsemat_type_t igraph_sparsemat_type(const igraph_sparsemat_t *A);
```

Gives whether a sparse matrix is stored in the triplet format or in column-compressed format.

Arguments:

A: The input matrix.

Returns:

Either `IGRAPH_SPARSEMAT_CC` or `IGRAPH_SPARSEMAT_TRIPLET`.

Time complexity: $O(1)$.

igraph_sparsemat_is_triplet — Is this sparse matrix in triplet format?

```
igraph_bool_t igraph_sparsemat_is_triplet(const igraph_sparsemat_t *A);
```

Decides whether a sparse matrix is in triplet format.

Arguments:

A: The input matrix.

Returns:

One if the input matrix is in triplet format, zero otherwise.

Time complexity: $O(1)$.

igraph_sparsemat_is_cc — Is this sparse matrix in column-compressed format?

```
igraph_bool_t igraph_sparsemat_is_cc(const igraph_sparsemat_t *A);
```

Decides whether a sparse matrix is in column-compressed format.

Arguments:

A: The input matrix.

Returns:

One if the input matrix is in column-compressed format, zero otherwise.

Time complexity: $O(1)$.

igraph_sparsemat_getelements_sorted — Returns the sorted elements of a sparse matrix.

```
int igraph_sparsemat_getelements_sorted(const igraph_sparsemat_t *A,
                                       igraph_vector_int_t *i,
                                       igraph_vector_int_t *j,
                                       igraph_vector_t *x);
```

This function will sort a sparse matrix and return the elements in 3 vectors. Two vectors will indicate where the elements are located, and one will give the elements.

Arguments:

- A: A sparse matrix in either triplet or compressed form.
- i: An initialized int vector. This will store the rows of the returned elements.
- j: An initialized int vector. For a triplet matrix this will store the columns of the returned elements. For a compressed matrix, if the column index is k , then $j[k]$ is the index in x of the start of the k -th column, and the last element of j is the total number of elements. The total number of elements in the k -th column is therefore $j[k+1] - j[k]$. For example, if there is one element in the first column, and five in the second, j will be set to $\{0, 1, 6\}$.
- x: An initialized vector. The elements will be placed here.

Returns:

Error code.

Time complexity: $O(n)$, the number of stored elements in the sparse matrix.

igraph_sparsemat_min — Minimum of a sparse matrix.

```
igraph_real_t igraph_sparsemat_min(igraph_sparsemat_t *A);
```

Arguments:

- A: The input matrix, column-compressed.

Returns:

The minimum in the input matrix, or `IGRAPH_POSINFINITY` if the matrix has zero elements.

Time complexity: TODO.

igraph_sparsemat_max — Maximum of a sparse matrix.

```
igraph_real_t igraph_sparsemat_max(igraph_sparsemat_t *A);
```

Arguments:

A: The input matrix, column-compressed.

Returns:

The maximum in the input matrix, or IGRAPH_NEGINFINITY if the matrix has zero elements.

Time complexity: TODO.

igraph_sparsemat_minmax — Minimum and maximum of a sparse matrix.

```
int igraph_sparsemat_minmax(igraph_sparsemat_t *A,  
                             igraph_real_t *min, igraph_real_t *max);
```

Arguments:

A: The input matrix, column-compressed.

min: The minimum in the input matrix is stored here, or IGRAPH_POSINFINITY if the matrix has zero elements.

max: The maximum in the input matrix is stored here, or IGRAPH_NEGINFINITY if the matrix has zero elements.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_count_nonzero — Counts nonzero elements of a sparse matrix.

```
long int igraph_sparsemat_count_nonzero(igraph_sparsemat_t *A);
```

Arguments:

A: The input matrix, column-compressed.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_count_nonzerotol — Counts nonzero elements of a sparse matrix, ignoring elements close to zero.

```
long int igraph_sparsemat_count_nonzerotol(igraph_sparsemat_t *A,
                                           igraph_real_t tol);
```

Count the number of matrix entries that are closer to zero than *tol*.

Arguments:

The: input matrix, column-compressed.

Real: scalar, the tolerance.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_rowsums — Row-wise sums.

```
int igraph_sparsemat_rowsums(const igraph_sparsemat_t *A,
                             igraph_vector_t *res);
```

Arguments:

A: The input matrix, in triplet or column-compressed format.

res: An initialized vector, the result is stored here. It will be resized as needed.

Returns:

Error code.

Time complexity: $O(nz)$, the number of non-zero elements.

igraph_sparsemat_colsums — Column-wise sums.

```
int igraph_sparsemat_colsums(const igraph_sparsemat_t *A,
                              igraph_vector_t *res);
```

Arguments:

A: The input matrix, in triplet or column-compressed format.

res: An initialized vector, the result is stored here. It will be resized as needed.

Returns:

Error code.

Time complexity: $O(nz)$ for triplet matrices, $O(nz+n)$ for column-compressed ones, nz is the number of non-zero elements, n is the number of columns.

igraph_sparsemat_nonzero_storage — Returns number of stored entries of a sparse matrix.

```
int igraph_sparsemat_nonzero_storage(const igraph_sparsemat_t *A);
```

This function will return the number of stored entries of a sparse matrix. These entries can be zero, and multiple entries can be at the same position. Use `igraph_sparsemat_dupl()` to sum duplicate entries, and `igraph_sparsemat_dropzeros()` to remove zeros.

Arguments:

A: A sparse matrix in either triplet or compressed form.

Returns:

Number of stored entries.

Time complexity: $O(1)$.

Operations on sparse matrices

igraph_sparsemat_entry — Adds an element to a sparse matrix.

```
int igraph_sparsemat_entry(igraph_sparsemat_t *A, int row, int col,
                           igraph_real_t elem);
```

This function can be used to add the entries to a sparse matrix, after initializing it with `igraph_sparsemat_init()`. If you add multiple entries in the same position, they will all be saved, and the resulting value is the sum of all entries in that position.

Arguments:

A: The input matrix, it must be in triplet format.

row: The row index of the entry to add.

col: The column index of the entry to add.

elem: The value of the entry.

Returns:

Error code.

Time complexity: $O(1)$ on average.

igraph_sparsemat_fkeep — Filters the elements of a sparse matrix.

```
int igraph_sparsemat_fkeep(
    igraph_sparsemat_t *A,
    igraph_integer_t (*fkeep)(igraph_integer_t, igraph_integer_t, igraph_real_t,
    void *other
);
```

This function can be used to filter the (non-zero) elements of a sparse matrix. For all entries, it calls the supplied function and depending on the return values either keeps, or deleted the element from the matrix.

Arguments:

A: The input matrix, in column-compressed format.

fkeep: The filter function. It must take four arguments: the first is an `int`, the row index of the entry, the second is another `int`, the column index. The third is `igraph_real_t`, the value of the entry. The fourth element is a `void` pointer, the *other* argument is passed here. The function must return an `int`. If this is zero, then the entry is deleted, otherwise it is kept.

other: A `void` pointer that is passed to the filtering function.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_dropzeros — Drops the zero elements from a sparse matrix.

```
int igraph_sparsemat_dropzeros(igraph_sparsemat_t *A);
```

As a result of matrix operations, some of the entries in a sparse matrix might be zero. This function removes these entries.

Arguments:

A: The input matrix, it must be in column-compressed format.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_droptol — Drops the almost zero elements from a sparse matrix.

```
int igraph_sparsemat_droptol(igraph_sparsemat_t *A, igraph_real_t tol);
```

This function is similar to `igraph_sparsemat_dropzeros()`, but it also drops entries that are closer to zero than the given tolerance threshold.

Arguments:

A: The input matrix, it must be in column-compressed format.

tol: Real number, giving the tolerance threshold.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_scale — Scales a sparse matrix.

```
int igraph_sparsemat_scale(igraph_sparsemat_t *A, igraph_real_t by);
```

Multiplies all elements of a sparse matrix, by the given scalar.

Arguments:

A: The input matrix.

by: The scaling factor.

Returns:

Error code.

Time complexity: $O(nz)$, the number of non-zero elements in the matrix.

igraph_sparsemat_permute — Permutes the rows and columns of a sparse matrix.

```
int igraph_sparsemat_permute(const igraph_sparsemat_t *A,
                             const igraph_vector_int_t *p,
                             const igraph_vector_int_t *q,
                             igraph_sparsemat_t *res);
```

Arguments:

A: The input matrix, it must be in column-compressed format.

p: Integer vector, giving the permutation of the rows.

q: Integer vector, the permutation of the columns.

res: Pointer to an uninitialized sparse matrix, the result is stored here.

Returns:

Error code.

Time complexity: $O(m+n+nz)$, the number of rows plus the number of columns plus the number of non-zero elements in the matrix.

igraph_sparsemat_transpose — Transposes a sparse matrix.

```
int igraph_sparsemat_transpose(const igraph_sparsemat_t *A,
                               igraph_sparsemat_t *res,
                               int values);
```

Arguments:

- A*: The input matrix, column-compressed or triple format.
- res*: Pointer to an uninitialized sparse matrix, the result is stored here.
- values*: If this is non-zero, the matrix transpose is calculated the normal way. If it is zero, then only the pattern of the input matrix is stored in the result, the values are not.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_add — Sum of two sparse matrices.

```
int igraph_sparsemat_add(const igraph_sparsemat_t *A,
                         const igraph_sparsemat_t *B,
                         igraph_real_t alpha,
                         igraph_real_t beta,
                         igraph_sparsemat_t *res);
```

Arguments:

- A*: The first input matrix, in column-compressed format.
- B*: The second input matrix, in column-compressed format.
- alpha*: Real scalar, *A* is multiplied by *alpha* before the addition.
- beta*: Real scalar, *B* is multiplied by *beta* before the addition.
- res*: Pointer to an uninitialized sparse matrix, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_multiply — Matrix multiplication.

```
int igraph_sparsemat_multiply(const igraph_sparsemat_t *A,
                              const igraph_sparsemat_t *B,
                              igraph_sparsemat_t *res);
```

Multiplies two sparse matrices.

Arguments:

A: The first input matrix (left hand side), in column-compressed format.

B: The second input matrix (right hand side), in column-compressed format.

res: Pointer to an uninitialized sparse matrix, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_gaxpy — Matrix-vector product, added to another vector.

```
int igraph_sparsemat_gaxpy(const igraph_sparsemat_t *A,  
                           const igraph_vector_t *x,  
                           igraph_vector_t *res);
```

Arguments:

A: The input matrix, in column-compressed format.

x: The input vector, its size must match the number of columns in *A*.

res: This vector is added to the matrix-vector product and it is overwritten by the result.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_add_rows — Adds rows to a sparse matrix.

```
int igraph_sparsemat_add_rows(igraph_sparsemat_t *A, long int n);
```

The current matrix elements are retained and all elements in the new rows are zero.

Arguments:

A: The input matrix, in triplet or column-compressed format.

n: The number of rows to add.

Returns:

Error code.

Time complexity: $O(1)$.

igraph_sparsemat_add_cols — Adds columns to a sparse matrix.

```
int igraph_sparsemat_add_cols(igraph_sparsemat_t *A, long int n);
```

The current matrix elements are retained, and all elements in the new columns are zero.

Arguments:

A: The input matrix, in triplet or column-compressed format.

n: The number of columns to add.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_resize — Resizes a sparse matrix.

```
int igraph_sparsemat_resize(igraph_sparsemat_t *A, long int nrow,  
                             long int ncol, int nzmax);
```

This function resizes a sparse matrix. The resized sparse matrix will be empty.

Arguments:

A: The initialized sparse matrix to resize.

nrow: The new number of rows.

ncol: The new number of columns.

nzmax: The new maximum number of elements.

Returns:

Error code.

Time complexity: $O(nzmax)$, the maximum number of non-zero elements.

Operations on sparse matrix iterators

igraph_sparsemat_iterator_init — Initialize a sparse matrix iterator.

```
int igraph_sparsemat_iterator_init(igraph_sparsemat_iterator_t *it,  
                                   igraph_sparsemat_t *sparsemat);
```

Arguments:

it: A pointer to an uninitialized sparse matrix iterator.
sparsemat: Pointer to the sparse matrix.

Returns:

Error code. This will always return IGRAPH_SUCCESS

Time complexity: $O(n)$, the number of columns of the sparse matrix.

igraph_sparsemat_iterator_reset — Reset a sparse matrix iterator to the first element.

```
int igraph_sparsemat_iterator_reset(igraph_sparsemat_iterator_t *it);
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

Error code. This will always return IGRAPH_SUCCESS

Time complexity: $O(n)$, the number of columns of the sparse matrix.

igraph_sparsemat_iterator_end — Query if the iterator is past the last element.

```
igraph_bool_t  
igraph_sparsemat_iterator_end(const igraph_sparsemat_iterator_t *it);
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

true if the iterator is past the last element, false if it points to an element in a sparse matrix.

Time complexity: $O(1)$.

igraph_sparsemat_iterator_row — Return the row of the iterator.

```
int igraph_sparsemat_iterator_row(const igraph_sparsemat_iterator_t *it);
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The row of the element at the current iterator position.

Time complexity: $O(1)$.

igraph_sparsemat_iterator_col — Return the column of the iterator.

```
int igraph_sparsemat_iterator_col(const igraph_sparsemat_iterator_t *it);
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The column of the element at the current iterator position.

Time complexity: $O(1)$.

igraph_sparsemat_iterator_get — Return the element at the current iterator position.

```
igraph_real_t  
igraph_sparsemat_iterator_get(const igraph_sparsemat_iterator_t *it);
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The value of the element at the current iterator position.

Time complexity: $O(1)$.

igraph_sparsemat_iterator_next — Let a sparse matrix iterator go to the next element.

```
int igraph_sparsemat_iterator_next(igraph_sparsemat_iterator_t *it);
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The position of the iterator in the element vector.

Time complexity: $O(n)$, the number of columns of the sparse matrix.

igraph_sparsemat_iterator_idx — Returns the element vector index of a sparse matrix iterator.

```
int igraph_sparsemat_iterator_idx(const igraph_sparsemat_iterator_t *it);
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The position of the iterator in the element vector.

Time complexity: $O(1)$.

Operations that change the internal representation

igraph_sparsemat_compress — Converts a sparse matrix to column-compressed format.

```
int igraph_sparsemat_compress(const igraph_sparsemat_t *A,
                              igraph_sparsemat_t *res);
```

Converts a sparse matrix from triplet format to column-compressed format. Almost all sparse matrix operations require that the matrix is in column-compressed format.

Arguments:

A: The input matrix, it must be in triplet format.

res: Pointer to an uninitialized sparse matrix object, the compressed version of *A* is stored here.

Returns:

Error code.

Time complexity: $O(nz)$ where *nz* is the number of non-zero elements.

igraph_sparsemat_dupl — Removes duplicate elements from a sparse matrix.

```
int igraph_sparsemat_dupl(igraph_sparsemat_t *A);
```

It is possible that a column-compressed sparse matrix stores a single matrix entry in multiple pieces. The entry is then the sum of all its pieces. (Some functions create matrices like this.) This function eliminates the multiple pieces.

Arguments:

A: The input matrix, in column-compressed format.

Returns:

Error code.

Time complexity: TODO.

Decompositions and solving linear systems

igraph_sparsemat_symlu — Symbolic LU decomposition.

```
int igraph_sparsemat_symlu(long int order, const igraph_sparsemat_t *A,
                           igraph_sparsemat_symbolic_t *dis);
```

LU decomposition of sparse matrices involves two steps, the first is calling this function, and then `igraph_sparsemat_lu()`.

Arguments:

order: The ordering to use: 0 means natural ordering, 1 means minimum degree ordering of $A + A'$, 2 is minimum degree ordering of $A'A$ after removing the dense rows from A , and 3 is the minimum degree ordering of $A'A$.

A: The input matrix, in column-compressed format.

dis: The result of the symbolic analysis is stored here. Once not needed anymore, it must be destroyed by calling `igraph_sparsemat_symbolic_destroy()`.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_symbqr — Symbolic QR decomposition.

```
int igraph_sparsemat_symbqr(long int order, const igraph_sparsemat_t *A,
                             igraph_sparsemat_symbolic_t *dis);
```

QR decomposition of sparse matrices involves two steps, the first is calling this function, and then `igraph_sparsemat_qr()`.

Arguments:

order: The ordering to use: 0 means natural ordering, 1 means minimum degree ordering of $A + A'$, 2 is minimum degree ordering of $A'A$ after removing the dense rows from A , and 3 is the minimum degree ordering of $A'A$.

A: The input matrix, in column-compressed format.

dis: The result of the symbolic analysis is stored here. Once not needed anymore, it must be destroyed by calling `igraph_sparsemat_symbolic_destroy()`.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_lsolve — Solves a lower-triangular linear system.

```
int igraph_sparsemat_lsolve(const igraph_sparsemat_t *L,
                           const igraph_vector_t *b,
                           igraph_vector_t *res);
```

Solve the $Lx=b$ linear equation system, where the L coefficient matrix is square and lower-triangular, with a zero-free diagonal.

Arguments:

L: The input matrix, in column-compressed format.

b: The right hand side of the linear system.

res: An initialized vector, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_ltsolve — Solves an upper-triangular linear system.

```
int igraph_sparsemat_ltsolve(const igraph_sparsemat_t *L,
                             const igraph_vector_t *b,
                             igraph_vector_t *res);
```

Solve the $L'x=b$ linear equation system, where the L matrix is square and lower-triangular, with a zero-free diagonal.

Arguments:

L: The input matrix, in column-compressed format.

b: The right hand side of the linear system.

res: An initialized vector, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_usolve — Solves an upper-triangular linear system.

```
int igraph_sparsemat_usolve(const igraph_sparsemat_t *U,  
                           const igraph_vector_t *b,  
                           igraph_vector_t *res);
```

Solves the $Ux=b$ upper triangular system.

Arguments:

U: The input matrix, in column-compressed format.

b: The right hand side of the linear system.

res: An initialized vector, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_utsolve — Solves a lower-triangular linear system.

```
int igraph_sparsemat_utsolve(const igraph_sparsemat_t *U,  
                           const igraph_vector_t *b,  
                           igraph_vector_t *res);
```

This is the same as `igraph_sparsemat_usolve()`, but $U'x=b$ is solved, where the apostrophe denotes the transpose.

Arguments:

U: The input matrix, in column-compressed format.

b: The right hand side of the linear system.

res: An initialized vector, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_cholsol — Solves a symmetric linear system via Cholesky decomposition.

```
int igraph_sparsemat_cholsol(const igraph_sparsemat_t *A,  
                           const igraph_vector_t *b,
```

```
igraph_vector_t *res,  
int order);
```

Solve $Ax=b$, where A is a symmetric positive definite matrix.

Arguments:

A: The input matrix, in column-compressed format.

v: The right hand side.

res: An initialized vector, the result is stored here.

order: An integer giving the ordering method to use for the factorization. Zero is the natural ordering; if it is one, then the fill-reducing minimum-degree ordering of $A+A'$ is used.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_lusol — Solves a linear system via LU decomposition.

```
int igraph_sparsemat_lusol(const igraph_sparsemat_t *A,  
                           const igraph_vector_t *b,  
                           igraph_vector_t *res,  
                           int order,  
                           igraph_real_t tol);
```

Solve $Ax=b$, via LU factorization of A .

Arguments:

A: The input matrix, in column-compressed format.

b: The right hand side of the equation.

res: An initialized vector, the result is stored here.

order: The ordering method to use, zero means the natural ordering, one means the fill-reducing minimum-degree ordering of $A+A'$, two means the ordering of $A'*A$, after removing the dense rows from A . Three means the ordering of $A'*A$.

tol: Real number, the tolerance limit to use for the numeric LU factorization.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_lu — LU decomposition of a sparse matrix.

```
int igraph_sparsemat_lu(const igraph_sparsemat_t *A,
                       const igraph_sparsemat_symbolic_t *dis,
                       igraph_sparsemat_numeric_t *din, double tol);
```

Performs numeric sparse LU decomposition of a matrix.

Arguments:

A: The input matrix, in column-compressed format.

dis: The symbolic analysis for LU decomposition, coming from a call to the `igraph_sparsemat_symlu()` function.

din: The numeric decomposition, the result is stored here. It can be used to solve linear systems with changing right hand side vectors, by calling `igraph_sparsemat_luresol()`. Once not needed any more, it must be destroyed by calling `igraph_sparsemat_symbolic_destroy()` on it.

tol: The tolerance for the numeric LU decomposition.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_qr — QR decomposition of a sparse matrix.

```
int igraph_sparsemat_qr(const igraph_sparsemat_t *A,
                       const igraph_sparsemat_symbolic_t *dis,
                       igraph_sparsemat_numeric_t *din);
```

Numeric QR decomposition of a sparse matrix.

Arguments:

A: The input matrix, in column-compressed format.

dis: The result of the symbolic QR analysis, from the function `igraph_sparsemat_symbqr()`.

din: The result of the decomposition is stored here, it can be used to solve many linear systems with the same coefficient matrix and changing right hand sides, using the `igraph_sparsemat_qrresol()` function. Once not needed any more, one should call `igraph_sparsemat_numeric_destroy()` on it to free the allocated memory.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_luresol — Solves a linear system using a precomputed LU decomposition.

```
int igraph_sparsemat_luresol(const igraph_sparsemat_symbolic_t *dis,
                             const igraph_sparsemat_numeric_t *din,
                             const igraph_vector_t *b,
                             igraph_vector_t *res);
```

Uses the LU decomposition of a matrix to solve linear systems.

Arguments:

dis: The symbolic analysis of the coefficient matrix, the result of `igraph_sparsemat_symblu()`.

din: The LU decomposition, the result of a call to `igraph_sparsemat_lu()`.

b: A vector that defines the right hand side of the linear equation system.

res: An initialized vector, the solution of the linear system is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_qrresol — Solves a linear system using a precomputed QR decomposition.

```
int igraph_sparsemat_qrresol(const igraph_sparsemat_symbolic_t *dis,
                              const igraph_sparsemat_numeric_t *din,
                              const igraph_vector_t *b,
                              igraph_vector_t *res);
```

Solves a linear system using a QR decomposition of its coefficient matrix.

Arguments:

dis: Symbolic analysis of the coefficient matrix, the result of `igraph_sparsemat_symbqr()`.

din: The QR decomposition of the coefficient matrix, the result of `igraph_sparsemat_qr()`.

b: Vector, giving the right hand side of the linear equation system.

res: An initialized vector, the solution is stored here. It is resized as needed.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_symbolic_destroy — Deallocates memory after a symbolic decomposition.

```
void igraph_sparsemat_symbolic_destroy(igraph_sparsemat_symbolic_t *dis);
```

Frees the memory allocated by `igraph_sparsemat_symbqr()` or `igraph_sparsemat_symbilu()`.

Arguments:

dis: The symbolic analysis.

Time complexity: $O(1)$.

igraph_sparsemat_numeric_destroy — Deallocates memory after a numeric decomposition.

```
void igraph_sparsemat_numeric_destroy(igraph_sparsemat_numeric_t *din);
```

Frees the memory allocated by `igraph_sparsemat_qr()` or `igraph_sparsemat_lu()`.

Arguments:

din: The LU or QR decomposition.

Time complexity: $O(1)$.

Eigenvalues and eigenvectors

igraph_sparsemat_arpack_rssolve — Eigenvalues and eigenvectors of a symmetric sparse matrix via ARPACK.

```
int igraph_sparsemat_arpack_rssolve(const igraph_sparsemat_t *A,  
                                     igraph_arpack_options_t *options,  
                                     igraph_arpack_storage_t *storage,  
                                     igraph_vector_t *values,  
                                     igraph_matrix_t *vectors,  
                                     igraph_sparsemat_solve_t solvemethod);
```

Arguments:

A: input matrix, must be column-compressed.

options: It is passed to `igraph_arpack_rssolve()`. See `igraph_arpack_options_t` for the details. If `mode` is 1, then ARPACK uses regular mode, if `mode` is 3, then shift and invert mode is used and the `sigma` structure member defines the shift.

storage: Storage for ARPACK. See `igraph_arpack_rssolve()` and `igraph_arpack_storage_t` for details.

values: An initialized vector or a null pointer, the eigenvalues are stored here.

vectors: An initialised matrix, or a null pointer, the eigenvectors are stored here, in the columns.

solvemethod: The method to solve the linear system, if `mode` is 3, i.e. the shift and invert mode is used. Possible values:

IGRAPH_SPARSE-
MAT_SOLVE_LU

The linear system is solved using LU decomposition.

IGRAPH_SPARSE-
MAT_SOLVE_QR

The linear system is solved using QR decomposition.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_arpack_rnsolve — Eigenvalues and eigenvectors of a nonsymmetric sparse matrix via ARPACK.

```
int igraph_sparsemat_arpack_rnsolve(const igraph_sparsemat_t *A,  
                                     igraph_arpack_options_t *options,  
                                     igraph_arpack_storage_t *storage,  
                                     igraph_matrix_t *values,  
                                     igraph_matrix_t *vectors);
```

Eigenvalues and/or eigenvectors of a nonsymmetric sparse matrix.

Arguments:

A: The input matrix, in column-compressed mode.

options: ARPACK options, it is passed to `igraph_arpack_rnsolve()`. See also `igraph_arpack_options_t` for details.

storage: Storage for ARPACK, this is passed to `igraph_arpack_rnsolve()`. See `igraph_arpack_storage_t` for details.

values: An initialized matrix, or a null pointer. If not a null pointer, then the eigenvalues are stored here, the first column is the real part, the second column is the imaginary part.

vectors: An initialized matrix, or a null pointer. If not a null pointer, then the eigenvectors are stored here, please see `igraph_arpack_rnsolve()` for the format.

Returns:

Error code.

Time complexity: TODO.

Conversion to other data types

igraph_sparsemat — Creates an igraph graph from a sparse matrix.

```
int igraph_sparsemat(igraph_t *graph, const igraph_sparsemat_t *A,  
                     igraph_bool_t directed);
```

One edge is created for each non-zero entry in the matrix. If you have a symmetric matrix, and want to create an undirected graph, then delete the entries in the upper diagonal first, or call `igraph_simplify()` on the result graph to eliminate the multiple edges.

Arguments:

graph: Pointer to an uninitialized `igraph_t` object, the graph is stored here.

A: The input matrix, in triplet or column-compressed format.

directed: Boolean scalar, whether to create a directed graph.

Returns:

Error code.

Time complexity: TODO.

`igraph_get_sparsemat` — Converts an igraph graph to a sparse matrix.

```
int igraph_get_sparsemat(const igraph_t *graph, igraph_sparsemat_t *res);
```

If the graph is undirected, then a symmetric matrix is created.

Arguments:

graph: The input graph.

res: Pointer to an uninitialized sparse matrix. The result will be stored here.

Returns:

Error code.

Time complexity: TODO.

`igraph_matrix_as_sparsemat` — Converts a dense matrix to a sparse matrix.

```
int igraph_matrix_as_sparsemat(igraph_sparsemat_t *res,  
                               const igraph_matrix_t *mat,  
                               igraph_real_t tol);
```

Arguments:

res: An uninitialized sparse matrix, the result is stored here.

mat: The dense input matrix.

tol: Real scalar, the tolerance. Values closer than *tol* to zero are considered as zero, and will not be included in the sparse matrix.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements in the dense matrix.

igraph_sparsemat_as_matrix — Converts a sparse matrix to a dense matrix.

```
int igraph_sparsemat_as_matrix(igraph_matrix_t *res,  
                               const igraph_sparsemat_t *spmat);
```

Arguments:

res: Pointer to an initialized matrix, the result is stored here. It will be resized to the required size.

spmat: The input sparse matrix, in triplet or column-compressed format.

Returns:

Error code.

Time complexity: $O(mn)$, the number of elements in the dense matrix.

Writing to a file, or to the screen

igraph_sparsemat_print — Prints a sparse matrix to a file.

```
int igraph_sparsemat_print(const igraph_sparsemat_t *A,  
                           FILE *ostream);
```

Only the non-zero entries are printed. This function serves more as a debugging utility, as currently there is no function that could read back the printed matrix from the file.

Arguments:

A: The input matrix, triplet or column-compressed format.

ostream: The stream to print it to.

Returns:

Error code.

Time complexity: $O(nz)$ for triplet matrices, $O(n+nz)$ for column-compressed matrices. nz is the number of non-zero elements, n is the number columns in the matrix.

Stacks

igraph_stack_init — Initializes a stack.

```
int igraph_stack_init(igraph_stack_t* s, long int size);
```

The initialized stack is always empty.

Arguments:

s: Pointer to an uninitialized stack.

size: The number of elements to allocate memory for.

Returns:

Error code.

Time complexity: $O(size)$.

igraph_stack_destroy — Destroys a stack object.

```
void igraph_stack_destroy(igraph_stack_t* s);
```

Deallocate the memory used for a stack. It is possible to reinitialize a destroyed stack again by `igraph_stack_init()`.

Arguments:

s: The stack to destroy.

Time complexity: $O(1)$.

igraph_stack_reserve — Reserve memory.

```
int igraph_stack_reserve(igraph_stack_t* s, long int size);
```

Reserve memory for future use. The actual size of the stack is unchanged.

Arguments:

s: The stack object.

size: The number of elements to reserve memory for. If it is not bigger than the current size then nothing happens.

Returns:

Error code.

Time complexity: should be around $O(n)$, the new allocated size of the stack.

igraph_stack_empty — Decides whether a stack object is empty.

```
igraph_bool_t igraph_stack_empty(igraph_stack_t* s);
```

Arguments:

s: The stack object.

Returns:

Boolean, TRUE if the stack is empty, FALSE otherwise.

Time complexity: $O(1)$.

igraph_stack_size — Returns the number of elements in a stack.

```
long int igraph_stack_size(const igraph_stack_t* s);
```

Arguments:

s: The stack object.

Returns:

The number of elements in the stack.

Time complexity: $O(1)$.

igraph_stack_clear — Removes all elements from a stack.

```
void igraph_stack_clear(igraph_stack_t* s);
```

Arguments:

s: The stack object.

Time complexity: $O(1)$.

igraph_stack_push — Places an element on the top of a stack.

```
int igraph_stack_push(igraph_stack_t* s, igraph_real_t elem);
```

The capacity of the stack is increased, if needed.

Arguments:

s: The stack object.

elem: The element to push.

Returns:

Error code.

Time complexity: $O(1)$ if no reallocation is needed, $O(n)$ otherwise, but it is ensured that n push operations are performed in $O(n)$ time.

igraph_stack_pop — Removes and returns an element from the top of a stack.

```
igraph_real_t igraph_stack_pop(igraph_stack_t* s);
```

The stack must contain at least one element, call `igraph_stack_empty()` to make sure of this.

Arguments:

s: The stack object.

Returns:

The removed top element.

Time complexity: $O(1)$.

igraph_stack_top — Query top element.

```
igraph_real_t igraph_stack_top(const igraph_stack_t* s);
```

Returns the top element of the stack, without removing it. The stack must be non-empty.

Arguments:

s: The stack.

Returns:

The top element.

Time complexity: $O(1)$.

Double-ended queues

This is the classic data type of the double ended queue. Most of the time it is used if a First-In-First-Out (FIFO) behavior is needed. See the operations below.

Example 7.11. File `examples/simple/dqueue.c`

igraph_dqueue_init — Initialize a double ended queue (deque).

```
int igraph_dqueue_init(igraph_dqueue_t* q, long int size);
```

The queue will be always empty.

Arguments:

q: Pointer to an uninitialized deque.

size: How many elements to allocate memory for.

Returns:

Error code.

Time complexity: $O(size)$.

igraph_dqueue_destroy — Destroy a double ended queue.

```
void igraph_dqueue_destroy(igraph_dqueue_t* q);
```

Arguments:

q: The queue to destroy

Time complexity: $O(1)$.

igraph_dqueue_empty — Decide whether the queue is empty.

```
igraph_bool_t igraph_dqueue_empty(const igraph_dqueue_t* q);
```

Arguments:

q: The queue.

Returns:

Boolean, TRUE if *q* contains at least one element, FALSE otherwise.

Time complexity: $O(1)$.

igraph_dqueue_full — Check whether the queue is full.

```
igraph_bool_t igraph_dqueue_full(igraph_dqueue_t* q);
```

If a queue is full the next `igraph_dqueue_push()` operation will allocate more memory.

Arguments:

q : The queue.

Returns:

TRUE if q is full, FALSE otherwise.

Time complexity: $O(1)$.

igraph_dqueue_clear — Remove all elements from the queue.

```
void igraph_dqueue_clear(igraph_dqueue_t* q);
```

Arguments:

q : The queue

Time complexity: $O(1)$.

igraph_dqueue_size — Number of elements in the queue.

```
long int igraph_dqueue_size(const igraph_dqueue_t* q);
```

Arguments:

q : The queue.

Returns:

Integer, the number of elements currently in the queue.

Time complexity: $O(1)$.

igraph_dqueue_head — Head of the queue.

```
igraph_real_t igraph_dqueue_head(const igraph_dqueue_t* q);
```

The queue must contain at least one element.

Arguments:

q : The queue.

Returns:

The first element in the queue.

Time complexity: $O(1)$.

igraph_dqueue_back — Tail of the queue.

```
igraph_real_t igraph_dqueue_back(const igraph_dqueue_t* q);
```

The queue must contain at least one element.

Arguments:

q: The queue.

Returns:

The last element in the queue.

Time complexity: $O(1)$.

igraph_dqueue_pop — Remove the head.

```
igraph_real_t igraph_dqueue_pop(igraph_dqueue_t* q);
```

Removes and returns the first element in the queue. The queue must be non-empty.

Arguments:

q: The input queue.

Returns:

The first element in the queue.

Time complexity: $O(1)$.

igraph_dqueue_pop_back — Remove the tail

```
igraph_real_t igraph_dqueue_pop_back(igraph_dqueue_t* q);
```

Removes and returns the last element in the queue. The queue must be non-empty.

Arguments:

q: The queue.

Returns:

The last element in the queue.

Time complexity: $O(1)$.

igraph_dqueue_push — Appends an element.

```
int igraph_dqueue_push(igraph_dqueue_t* q, igraph_real_t elem);
```

Append an element to the end of the queue.

Arguments:

q: The queue.

elem: The element to append.

Returns:

Error code.

Time complexity: $O(1)$ if no memory allocation is needed, $O(n)$, the number of elements in the queue otherwise. But not that by allocating always twice as much memory as the current size of the queue we ensure that n push operations can always be done in at most $O(n)$ time. (Assuming memory allocation is at most linear.)

Maximum and minimum heaps

igraph_heap_init — Initializes an empty heap object.

```
int igraph_heap_init(igraph_heap_t* h, long int alloc_size);
```

Creates an empty heap, but allocates size for some elements.

Arguments:

h: Pointer to an uninitialized heap object.

alloc_size: Number of elements to allocate memory for.

Returns:

Error code.

Time complexity: $O(alloc_size)$, assuming memory allocation is a linear operation.

igraph_heap_init_array — Build a heap from an array.

```
int igraph_heap_init_array(igraph_heap_t *h, igraph_real_t* data, long int len)
```

Initializes a heap object from an array, the heap is also built of course (constructor).

Arguments:

h: Pointer to an uninitialized heap object.

data: Pointer to an array of base data type.

len: The length of the array at *data*.

Returns:

Error code.

Time complexity: $O(n)$, the number of elements in the heap.

igraph_heap_destroy — Destroys an initialized heap object.

```
void igraph_heap_destroy(igraph_heap_t* h);
```

Arguments:

h: The heap object.

Time complexity: $O(1)$.

igraph_heap_empty — Decides whether a heap object is empty.

```
igraph_bool_t igraph_heap_empty(igraph_heap_t* h);
```

Arguments:

h: The heap object.

Returns:

TRUE if the heap is empty, FALSE otherwise.

Time complexity: $O(1)$.

igraph_heap_push — Add an element.

```
int igraph_heap_push(igraph_heap_t* h, igraph_real_t elem);
```

Adds an element to the heap.

Arguments:

h: The heap object.

elem: The element to add.

Returns:

Error code.

Time complexity: $O(\log n)$, n is the number of elements in the heap if no reallocation is needed, $O(n)$ otherwise. It is ensured that n push operations are performed in $O(n \log n)$ time.

igraph_heap_top — Top element.

```
igraph_real_t igraph_heap_top(igraph_heap_t* h);
```

For maximum heaps this is the largest, for minimum heaps the smallest element of the heap.

Arguments:

h: The heap object.

Returns:

The top element.

Time complexity: $O(1)$.

igraph_heap_delete_top — Return and removes the top element

```
igraph_real_t igraph_heap_delete_top(igraph_heap_t* h);
```

Removes and returns the top element of the heap. For maximum heaps this is the largest, for minimum heaps the smallest element.

Arguments:

h: The heap object.

Returns:

The top element.

Time complexity: $O(\log n)$, n is the number of elements in the heap.

igraph_heap_size — Number of elements

```
long int igraph_heap_size(igraph_heap_t* h);
```

Gives the number of elements in a heap.

Arguments:

h: The heap object.

Returns:

The number of elements in the heap.

Time complexity: $O(1)$.

igraph_heap_reserve — Allocate more memory

```
int igraph_heap_reserve(igraph_heap_t* h, long int size);
```

Allocates memory for future use. The size of the heap is unchanged. If the heap is larger than the *size* parameter then nothing happens.

Arguments:

h: The heap object.

size: The number of elements to allocate memory for.

Returns:

Error code.

Time complexity: $O(\text{size})$ if *size* is larger than the current number of elements. $O(1)$ otherwise.

String vectors

The `igraph_strvector_t` type is a vector of strings. The current implementation is very simple and not too efficient. It works fine for not too many strings, e.g. the list of attribute names is returned in a string vector by `igraph_cattribute_list()`. Do not expect great performance from this type.

Example 7.12. File `examples/simple/igraph_strvector.c`

igraph_strvector_init — Initialize

```
int igraph_strvector_init(igraph_strvector_t *sv, long int len);
```

Reserves memory for the string vector, a string vector must be first initialized before calling other functions on it. All elements of the string vector are set to the empty string.

Arguments:

sv: Pointer to an initialized string vector.

len: The (initial) length of the string vector.

Returns:

Error code.

Time complexity: $O(\text{len})$.

igraph_strvector_copy — Initialization by copying.

```
int igraph_strvector_copy(igraph_strvector_t *to,
                          const igraph_strvector_t *from);
```

Initializes a string vector by copying another string vector.

Arguments:

to: Pointer to an uninitialized string vector.

from: The other string vector, to be copied.

Returns:

Error code.

Time complexity: $O(l)$, the total length of the strings in *from*.

igraph_strvector_destroy — Free allocated memory

```
void igraph_strvector_destroy(igraph_strvector_t *sv);
```

Destroy a string vector. It may be reinitialized with `igraph_strvector_init()` later.

Arguments:

sv: The string vector.

Time complexity: $O(l)$, the total length of the strings, maybe less depending on the memory manager.

STR — Indexing string vectors

```
#define STR(sv,i)
```

This is a macro which allows to query the elements of a string vector in simpler way than `igraph_strvector_get()`. Note this macro cannot be used to set an element, for that use `igraph_strvector_set()`.

Arguments:

sv: The string vector

i: The the index of the element.

Returns:

The element at position *i*.

Time complexity: $O(1)$.

igraph_strvector_get — Indexing

```
void igraph_strvector_get(const igraph_strvector_t *sv, long int idx,
                        char **value);
```

Query an element of a string vector. See also the STR macro for an easier way.

Arguments:

sv: The input string vector.

idx: The index of the element to query.

Pointer: to a char*, the address of the string is stored here.

Time complexity: O(1).

igraph_strvector_set — Set an element

```
int igraph_strvector_set(igraph_strvector_t *sv, long int idx,
                        const char *value);
```

The provided *value* is copied into the *idx* position in the string vector.

Arguments:

sv: The string vector.

idx: The position to set.

value: The new value.

Returns:

Error code.

Time complexity: O(l), the length of the new string. Maybe more, depending on the memory management, if reallocation is needed.

igraph_strvector_set2 — Sets an element.

```
int igraph_strvector_set2(igraph_strvector_t *sv, long int idx,
                        const char *value, int len);
```

This is almost the same as `igraph_strvector_set`, but the new value is not a zero terminated string, but its length is given.

Arguments:

sv: The string vector.

idx: The position to set.

value: The new value.

len: The length of the new value.

Returns:

Error code.

Time complexity: $O(l)$, the length of the new string. Maybe more, depending on the memory management, if reallocation is needed.

igraph_strvector_remove — Removes a single element from a string vector.

```
void igraph_strvector_remove(igraph_strvector_t *v, long int elem);
```

The string will be one shorter.

Arguments:

v: The string vector.

elem: The index of the element to remove.

Time complexity: $O(n)$, the length of the string.

igraph_strvector_append — Concatenate two string vectors.

```
int igraph_strvector_append(igraph_strvector_t *to,  
                           const igraph_strvector_t *from);
```

Arguments:

to: The first string vector, the result is stored here.

from: The second string vector, it is kept unchanged.

Returns:

Error code.

Time complexity: $O(n+l_2)$, n is the number of strings in the new string vector, l_2 is the total length of strings in the *from* string vector.

igraph_strvector_clear — Remove all elements

```
void igraph_strvector_clear(igraph_strvector_t *sv);
```

After this operation the string vector will be empty.

Arguments:

sv: The string vector.

Time complexity: $O(l)$, the total length of strings, maybe less, depending on the memory manager.

igraph_strvector_resize — Resize

```
int igraph_strvector_resize(igraph_strvector_t* v, long int newsize);
```

If the new size is bigger then empty strings are added, if it is smaller then the unneeded elements are removed.

Arguments:

v: The string vector.

newsize: The new size.

Returns:

Error code.

Time complexity: $O(n)$, the number of strings if the vector is made bigger, $O(1)$, the total length of the deleted strings if it is made smaller, maybe less, depending on memory management.

igraph_strvector_size — Gives the size of a string vector.

```
long int igraph_strvector_size(const igraph_strvector_t *sv);
```

Arguments:

sv: The string vector.

Returns:

The length of the string vector.

Time complexity: $O(1)$.

igraph_strvector_add — Adds an element to the back of a string vector.

```
int igraph_strvector_add(igraph_strvector_t *v, const char *value);
```

Arguments:

v: The string vector.

value: The string to add, it will be copied.

Returns:

Error code.

Time complexity: $O(n+l)$, n is the total number of strings, l is the length of the new string.

Adjacency lists

Sometimes it is easier to work with a graph which is in adjacency list format: a list of vectors; each vector contains the neighbor vertices or incident edges of a given vertex. Typically, this representation is good if we need to iterate over the neighbors of all vertices many times. E.g. when finding the shortest paths between all pairs of vertices or calculating closeness centrality for all the vertices.

The `igraph_adjlist_t` stores the adjacency lists of a graph. After creation it is independent of the original graph, it can be modified freely with the usual vector operations, the graph is not affected. E.g. the adjacency list can be used to rewire the edges of a graph efficiently. If one used the straightforward `igraph_delete_edges()` and `igraph_add_edges()` combination for this that needs $O(|V|+|E|)$ time for every single deletion and insertion operation, it is thus very slow if many edges are rewired. Extracting the graph into an adjacency list, do all the rewiring operations on the vectors of the adjacency list and then creating a new graph needs (depending on how exactly the rewiring is done) typically $O(|V|+|E|)$ time for the whole rewiring process.

Lazy adjacency lists are a bit different. When creating a lazy adjacency list, the neighbors of the vertices are not queried, only some memory is allocated for the vectors. When `igraph_lazy_adjlist_get()` is called for vertex v the first time, the neighbors of v are queried and stored in a vector of the adjacency list, so they don't need to be queried again. Lazy adjacency lists are handy if you have an at least linear operation (because initialization is generally linear in terms of the number of vertices), but you don't know how many vertices you will visit during the computation.

Example 7.13. File `examples/simple/adjlist.c`

Adjacent vertices

`igraph_adjlist_init` — Constructs an adjacency list of vertices from a given graph.

```
int igraph_adjlist_init(const igraph_t *graph, igraph_adjlist_t *al,
                       igraph_neimode_t mode, igraph_loops_t loops,
                       igraph_multiple_t multiple);
```

Creates a list of vectors containing the neighbors of all vertices in a graph. The adjacency list is independent of the graph after creation, e.g. the graph can be destroyed and modified, the adjacency list contains the state of the graph at the time of its initialization.

Arguments:

- | | |
|----------------|---|
| <i>graph</i> : | The input graph. |
| <i>al</i> : | Pointer to an uninitialized <code>igraph_adjlist_t</code> object. |
| <i>mode</i> : | Constant specifying whether outgoing (<code>IGRAPH_OUT</code>), incoming (<code>IGRAPH_IN</code>), or both (<code>IGRAPH_ALL</code>) types of neighbors to include in the adjacency list. It is ignored for undirected networks. |
| <i>loops</i> : | Specifies how to treat loop edges. <code>IGRAPH_NO_LOOPS</code> removes loop edges from the adjacency list. <code>IGRAPH_LOOPS_ONCE</code> makes each loop edge appear only once in the adjacency list of the corresponding vertex. <code>IGRAPH_LOOPS_TWICE</code> makes |

loop edges appear *twice* in the adjacency list of the corresponding vertex, but only if the graph is undirected or mode is set to IGRAPH_ALL.

multiple: Specifies how to treat multiple (parallel) edges. IGRAPH_NO_MULTIPLE collapses parallel edges into a single one; IGRAPH_MULTIPLE keeps the multiplicities of parallel edges so the same vertex will appear as many times in the adjacency list of another vertex as the number of parallel edges going between the two vertices.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

igraph_adjlist_init_empty — Initializes an empty adjacency list.

```
int igraph_adjlist_init_empty(igraph_adjlist_t *al, igraph_integer_t no_of_nodes)
```

Creates a list of vectors, one for each vertex. This is useful when you are *constructing* a graph using an adjacency list representation as it does not require your graph to exist yet.

Arguments:

no_of_nodes: The number of vertices

al: Pointer to an uninitialized igraph_adjlist_t object.

Returns:

Error code.

Time complexity: $O(|V|)$, linear in the number of vertices.

igraph_adjlist_init_complementer — Adjacency lists for the complementer graph.

```
int igraph_adjlist_init_complementer(const igraph_t *graph,  
                                     igraph_adjlist_t *al,  
                                     igraph_neimode_t mode,  
                                     igraph_bool_t loops);
```

This function creates adjacency lists for the complementer of the input graph. In the complementer graph all edges are present which are not present in the original graph. Multiple edges in the input graph are ignored.

Arguments:

graph: The input graph.

al: Pointer to a not yet initialized adjacency list.

mode: Constant specifying whether outgoing (IGRAPH_OUT), incoming (IGRAPH_IN), or both (IGRAPH_ALL) types of neighbors (in the complementer graph) to include in the adjacency list. It is ignored for undirected networks.

loops: Whether to consider loop edges.

Returns:

Error code.

Time complexity: $O(|V|^2+|E|)$, quadratic in the number of vertices.

igraph_adjlist_destroy — Deallocates an adjacency list.

```
void igraph_adjlist_destroy(igraph_adjlist_t *al);
```

Free all memory allocated for an adjacency list.

Arguments:

al: The adjacency list to destroy.

Time complexity: depends on memory management.

igraph_adjlist_get — Query a vector in an adjacency list.

```
#define igraph_adjlist_get(al,no)
```

Returns a pointer to an `igraph_vector_int_t` object from an adjacency list. The vector can be modified as desired.

Arguments:

al: The adjacency list object.

no: The vertex whose adjacent vertices will be returned.

Returns:

Pointer to the `igraph_vector_int_t` object.

Time complexity: $O(1)$.

igraph_adjlist_size — Returns the number of vertices in an adjacency list.

```
igraph_integer_t igraph_adjlist_size(const igraph_adjlist_t *al);
```

Arguments:

al: The adjacency list.

Returns:

The number of vertices in the adjacency list.

Time complexity: $O(1)$.

igraph_adjlist_clear — Removes all edges from an adjacency list.

```
void igraph_adjlist_clear(igraph_adjlist_t *al);
```

Arguments:

al: The adjacency list. Time complexity: depends on memory management, typically $O(n)$, where n is the total number of elements in the adjacency list.

igraph_adjlist_sort — Sorts each vector in an adjacency list.

```
void igraph_adjlist_sort(igraph_adjlist_t *al);
```

Sorts every vector of the adjacency list.

Arguments:

al: The adjacency list.

Time complexity: $O(n \log n)$, n is the total number of elements in the adjacency list.

igraph_adjlist_simplify — Simplifies an adjacency list.

```
int igraph_adjlist_simplify(igraph_adjlist_t *al);
```

Simplifies an adjacency list, i.e. removes loop and multiple edges.

Arguments:

al: The adjacency list.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of edges and vertices.

Incident edges

igraph_inclist_init — Initializes an incidence list.

```
int igraph_inclist_init(const igraph_t *graph,  
                        igraph_inclist_t *il,  
                        igraph_neimode_t mode,  
                        igraph_loops_t loops);
```

Creates a list of vectors containing the incident edges for all vertices. The incidence list is independent of the graph after creation, subsequent changes of the graph object do not update the incidence list, and changes to the incidence list do not update the graph.

When *mode* is `IGRAPH_IN` or `IGRAPH_OUT`, each edge ID will appear in the incidence list *once*. When *mode* is `IGRAPH_ALL`, each edge ID will appear in the incidence list *twice*, once for the source vertex and once for the target edge. It also means that the edge IDs of loop edges may potentially appear *twice* for the *same* vertex. Use the *loops* argument to control whether this will be the case (`IGRAPH_LOOPS_TWICE`) or not (`IGRAPH_LOOPS_ONCE` or `IGRAPH_NO_LOOPS`).

Arguments:

graph: The input graph.

il: Pointer to an uninitialized incidence list.

mode: Constant specifying whether incoming edges (`IGRAPH_IN`), outgoing edges (`IGRAPH_OUT`) or both (`IGRAPH_ALL`) to include in the incidence lists of directed graphs. It is ignored for undirected graphs.

loops: Specifies how to treat loop edges. `IGRAPH_NO_LOOPS` removes loop edges from the incidence list. `IGRAPH_LOOPS_ONCE` makes each loop edge appear only once in the incidence list of the corresponding vertex. `IGRAPH_LOOPS_TWICE` makes loop edges appear *twice* in the incidence list of the corresponding vertex, but only if the graph is undirected or *mode* is set to `IGRAPH_ALL`.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

igraph_inclist_destroy — Frees all memory allocated for an incidence list.

```
void igraph_inclist_destroy(igraph_inclist_t *il);
```

Arguments:

il: The incidence list to destroy.

Time complexity: depends on memory management.

igraph_inclist_get — Query a vector in an incidence list.

```
#define igraph_inclist_get(il,no)
```

Returns a pointer to an `igraph_vector_int_t` object from an incidence list containing edge ids. The vector can be modified, resized, etc. as desired.

Arguments:

il: Pointer to the incidence list.

no: The vertex for which the incident edges are returned.

Returns:

Pointer to an `igraph_vector_int_t` object.

Time complexity: $O(1)$.

`igraph_inclist_size` — Returns the number of vertices in an incidence list.

```
igraph_integer_t igraph_inclist_size(const igraph_inclist_t *il);
```

Arguments:

il: The incidence list.

Returns:

The number of vertices in the incidence list.

Time complexity: $O(1)$.

`igraph_inclist_clear` — Removes all edges from an incidence list.

```
void igraph_inclist_clear(igraph_inclist_t *il);
```

Arguments:

il: The incidence list.

Time complexity: depends on memory management, typically $O(n)$, where n is the total number of elements in the incidence list.

Lazy adjacency list for vertices

`igraph_lazy_adjlist_init` — Initialized a lazy adjacency list.

```
int igraph_lazy_adjlist_init(const igraph_t *graph,  
                             igraph_lazy_adjlist_t *al,  
                             igraph_neimode_t mode,  
                             igraph_loops_t loops,  
                             igraph_multiple_t multiple);
```

Create a lazy adjacency list for vertices. This function only allocates some memory for storing the vectors of an adjacency list, but the neighbor vertices are not queried, only at the `igraph_lazy_adjlist_get()` calls.

Arguments:

graph: The input graph.

al: Pointer to an uninitialized adjacency list object.

mode: Constant, it gives whether incoming edges (IGRAPH_IN), outgoing edges (IGRAPH_OUT) or both types of edges (IGRAPH_ALL) are considered. It is ignored for undirected graphs.

simplify: Constant, it gives whether to simplify the vectors in the adjacency list (IGRAPH_SIMPLIFY) or not (IGRAPH_DONT_SIMPLIFY).

Returns:

Error code.

Time complexity: $O(|V|)$, the number of vertices, possibly, but depends on the underlying memory management too.

igraph_lazy_adjlist_destroy — Deallocate a lazy adjacency list.

```
void igraph_lazy_adjlist_destroy(igraph_lazy_adjlist_t *al);
```

Free all allocated memory for a lazy adjacency list.

Arguments:

al: The adjacency list to deallocate.

Time complexity: depends on the memory management.

igraph_lazy_adjlist_get — Query neighbor vertices.

```
#define igraph_lazy_adjlist_get(al,no)
```

If the function is called for the first time for a vertex then the result is stored in the adjacency list and no further query operations are needed when the neighbors of the same vertex are queried again.

Arguments:

al: The lazy adjacency list.

no: The vertex ID to query.

Returns:

Pointer to a vector. It is allowed to modify it and modification does not affect the original graph.

Time complexity: $O(d)$, the number of neighbor vertices for the first time, $O(1)$ for subsequent calls.

igraph_lazy_adjlist_size — Returns the number of vertices in a lazy adjacency list.

```
igraph_integer_t igraph_lazy_adjlist_size(const igraph_lazy_adjlist_t *al);
```

Arguments:

al: The lazy adjacency list.

Returns:

The number of vertices in the lazy adjacency list.

Time complexity: $O(1)$.

igraph_lazy_adjlist_clear — Removes all edges from a lazy adjacency list.

```
void igraph_lazy_adjlist_clear(igraph_lazy_adjlist_t *al);
```

Arguments:

al: The lazy adjacency list. Time complexity: depends on memory management, typically $O(n)$, where n is the total number of elements in the adjacency list.

Lazy incidence list for edges

igraph_lazy_inclist_init — Initializes a lazy incidence list of edges.

```
int igraph_lazy_inclist_init(const igraph_t *graph,
                             igraph_lazy_inclist_t *il,
                             igraph_neimode_t mode,
                             igraph_loops_t loops);
```

Create a lazy incidence list for edges. This function only allocates some memory for storing the vectors of an incidence list, but the incident edges are not queried, only when `igraph_lazy_inclist_get()` is called.

When *mode* is `IGRAPH_IN` or `IGRAPH_OUT`, each edge ID will appear in the incidence list *once*. When *mode* is `IGRAPH_ALL`, each edge ID will appear in the incidence list *twice*, once for the source vertex and once for the target edge. It also means that the edge IDs of loop edges will appear *twice* for the *same* vertex.

Arguments:

graph: The input graph.

al: Pointer to an uninitialized incidence list.

mode: Constant, it gives whether incoming edges (`IGRAPH_IN`), outgoing edges (`IGRAPH_OUT`) or both types of edges (`IGRAPH_ALL`) are considered. It is ignored for undirected graphs.

loops: Specifies how to treat loop edges. `IGRAPH_NO_LOOPS` removes loop edges from the incidence list. `IGRAPH_LOOPS_ONCE` makes each loop edge appear only once in the incidence list of the corresponding vertex. `IGRAPH_LOOPS_TWICE` makes loop edges

appear *twice* in the incidence list of the corresponding vertex, but only if the graph is undirected or mode is set to IGRAPH_ALL.

Returns:

Error code.

Time complexity: $O(|V|)$, the number of vertices, possibly. But it also depends on the underlying memory management.

igraph_lazy_inclist_destroy — Deallocates a lazy incidence list.

```
void igraph_lazy_inclist_destroy(igraph_lazy_inclist_t *il);
```

Frees all allocated memory for a lazy incidence list.

Arguments:

al: The incidence list to deallocate.

Time complexity: depends on memory management.

igraph_lazy_inclist_get — Query incident edges.

```
#define igraph_lazy_inclist_get(al,no)
```

If the function is called for the first time for a vertex, then the result is stored in the incidence list and no further query operations are needed when the incident edges of the same vertex are queried again.

Arguments:

al: The lazy incidence list object.

no: The vertex id to query.

Returns:

Pointer to a vector. It is allowed to modify it and modification does not affect the original graph.

Time complexity: $O(d)$, the number of incident edges for the first time, $O(1)$ for subsequent calls with the same *no* argument.

igraph_lazy_inclist_size — Returns the number of vertices in a lazy incidence list.

```
igraph_integer_t igraph_lazy_inclist_size(const igraph_lazy_inclist_t *il);
```

Arguments:

il: The lazy incidence list.

Returns:

The number of vertices in the lazy incidence list.

Time complexity: $O(1)$.

igraph_lazy_inclist_clear — Removes all edges from a lazy incidence list.

```
void igraph_lazy_inclist_clear(igraph_lazy_inclist_t *il);
```

Arguments:

il: The lazy incidence list.

Time complexity: depends on memory management, typically $O(n)$, where n is the total number of elements in the incidence list.

Partial prefix sum trees

The `igraph_psumtree_t` data type represents a partial prefix sum tree. A partial prefix sum tree is a data structure that can be used to draw samples from a discrete probability distribution with dynamic probabilities that are updated frequently. This is achieved by creating a binary tree where the leaves are the items. Each leaf contains the probability corresponding to the items. Intermediate nodes of the tree always contain the sum of its two children. When the value of a leaf node is updated, the values of its ancestors are also updated accordingly.

Samples can be drawn from the probability distribution represented by the tree by generating a uniform random number between 0 (inclusive) and the value of the root of the tree (exclusive), and then following the branches of the tree as follows. In each step, the value in the current node is compared with the generated number. If the value in the node is larger, the left branch of the tree is taken; otherwise the generated number is decreased by the value in the node and the right branch of the tree is taken, until a leaf node is reached.

Note that the sampling process works only if all the values in the tree are non-negative. This is enforced by the object; in particular, trying to set a negative value for an item will produce an igraph error.

igraph_psumtree_init — Initializes a partial prefix sum tree.

```
int igraph_psumtree_init(igraph_psumtree_t *t, long int size);
```

The tree is initialized with a fixed number of elements. After initialization, the value corresponding to each element is zero.

Arguments:

t: The tree to initialize

size: The number of elements in the tree

Returns:

Error code, typically `IGRAPH_ENOMEM` if there is not enough memory

Time complexity: $O(n)$ for a tree containing n elements

`igraph_psumtree_destroy` — Destroys a partial prefix sum tree.

```
void igraph_psumtree_destroy(igraph_psumtree_t *t);
```

All partial prefix sum trees initialized by `igraph_psumtree_init()` should be properly destroyed by this function. A destroyed tree needs to be reinitialized by `igraph_psumtree_init()` if you want to use it again.

Arguments:

`t`: Pointer to the (previously initialized) tree to destroy.

Time complexity: operating system dependent.

`igraph_psumtree_size` — Returns the size of the tree.

```
long int igraph_psumtree_size(const igraph_psumtree_t *t);
```

Arguments:

`t`: The tree object

Returns:

The number of discrete items in the tree.

Time complexity: $O(1)$.

`igraph_psumtree_get` — Retrieves the value corresponding to an item in the tree.

```
igraph_real_t igraph_psumtree_get(const igraph_psumtree_t *t, long int idx);
```

Arguments:

`t`: The tree to query.

`idx`: The index of the item whose value is to be retrieved.

Returns:

The value corresponding to the item with the given index.

Time complexity: $O(1)$

igraph_psumtree_sum — Returns the sum of the values of the leaves in the tree.

```
igraph_real_t igraph_psumtree_sum(const igraph_psumtree_t *t);
```

Arguments:

t: The tree object

Returns:

The sum of the values of the leaves in the tree.

Time complexity: $O(1)$.

igraph_psumtree_search — Finds an item in the tree, given a value.

```
int igraph_psumtree_search(const igraph_psumtree_t *t, long int *idx,
                           igraph_real_t search);
```

This function finds the item with the lowest index where it holds that the sum of all the items with a *lower* index is less than or equal to the given value and that the sum of all the items with a lower index plus the item itself is larger than the given value.

If you think about the partial prefix sum tree as a tool to sample from a discrete probability distribution, then calling this function repeatedly with uniformly distributed random numbers in the range 0 (inclusive) to the sum of all values in the tree (exclusive) will sample the items in the tree with a probability that is proportional to their associated values.

Arguments:

t: The tree to query.

idx: The index of the item is returned here.

search: The value to use for the search.

Returns:

Error code; currently the search always succeeds.

Time complexity: $O(\log n)$, where n is the number of items in the tree.

igraph_psumtree_update — Updates the value associated to an item in the tree.

```
int igraph_psumtree_update(igraph_psumtree_t *t, long int idx,  
                           igraph_real_t new_value);
```

Arguments:

t: The tree to query.
idx: The index of the item to update.
new_value: The new value of the item.

Returns:

Error code, `IGRAPH_EINVAL` if the new value is negative or NaN, `IGRAPH_SUCCESS` if the operation was successful.

Time complexity: $O(\log n)$, where n is the number of items in the tree.

Chapter 8. Random numbers

About random numbers in igraph, use cases

Some algorithms in igraph, e.g. the generation of random graphs, require random number generators (RNGs). Prior to version 0.6 igraph did not have a sophisticated way to deal with random number generators at the C level, but this has changed. From version 0.6 different and multiple random number generators are supported.

The default random number generator

igraph_rng_default — Query the default random number generator.

```
igraph_rng_t *igraph_rng_default();
```

Returns:

A pointer to the default random number generator.

See also:

`igraph_rng_set_default()`

igraph_rng_set_default — Set the default igraph random number generator.

```
void igraph_rng_set_default(igraph_rng_t *rng);
```

Arguments:

rng: The random number generator to use as default from now on. Calling `igraph_rng_destroy()` on it, while it is still being used as the default will result in crashes and/or unpredictable results.

Time complexity: $O(1)$.

Creating random number generators

igraph_rng_init — Initialize a random number generator.

```
int igraph_rng_init(igraph_rng_t *rng, const igraph_rng_type_t *type);
```

This function allocates memory for a random number generator, with the given type, and sets its seed to the default.

Arguments:

rng: Pointer to an uninitialized RNG.

type: The type of the RNG, like `igraph_rngtype_glibc2`, `igraph_rngtype_mt19937` or `igraph_rngtype_rand`.

Returns:

Error code.

Time complexity: depends on the type of the generator, but usually it should be $O(1)$.

`igraph_rng_destroy` — Deallocate memory associated with a random number generator.

```
void igraph_rng_destroy(igraph_rng_t *rng);
```

Arguments:

rng: The RNG to destroy. Do not destroy an RNG that is used as the default igraph RNG.

Time complexity: $O(1)$.

`igraph_rng_seed` — Set the seed of a random number generator.

```
int igraph_rng_seed(igraph_rng_t *rng, unsigned long int seed);
```

Arguments:

rng: The RNG.

seed: The new seed.

Returns:

Error code.

Time complexity: usually $O(1)$, but may depend on the type of the RNG.

`igraph_rng_min` — Query the minimum possible integer for a random number generator.

```
unsigned long int igraph_rng_min(igraph_rng_t *rng);
```

Arguments:

rng: The RNG.

Returns:

The smallest possible integer that can be generated by calling `igraph_rng_get_integer()` on the RNG.

Time complexity: $O(1)$.

`igraph_rng_max` — Query the maximum possible integer for a random number generator.

```
unsigned long int igraph_rng_max(igraph_rng_t *rng);
```

Arguments:

rng: The RNG.

Returns:

The largest possible integer that can be generated by calling `igraph_rng_get_integer()` on the RNG.

Time complexity: $O(1)$.

`igraph_rng_name` — Query the type of a random number generator.

```
const char *igraph_rng_name(igraph_rng_t *rng);
```

Arguments:

rng: The RNG.

Returns:

The name of the type of the generator. Do not deallocate or change the returned string pointer.

Time complexity: $O(1)$.

Generating random numbers

`igraph_rng_get_integer` — Generate an integer random number from an interval.

```
long int igraph_rng_get_integer(igraph_rng_t *rng,  
                                long int l, long int h);
```

Arguments:

rng: Pointer to the RNG to use for the generation. Use `igraph_rng_default()` here to use the default igraph RNG.

l: Lower limit, inclusive, it can be negative as well.

h: Upper limit, inclusive, it can be negative as well, but it should be at least *l*.

Returns:

The generated random integer.

Time complexity: depends on the generator, but should be usually $O(1)$.

igraph_rng_get_unif — Generate real, uniform random numbers from an interval

```
igraph_real_t igraph_rng_get_unif(igraph_rng_t *rng,  
                                   igraph_real_t l, igraph_real_t h);
```

Arguments:

rng: Pointer to the RNG to use. Use `igraph_rng_default()` here to use the default igraph RNG.

l: The lower bound, it can be negative.

h: The upper bound, it can be negative, but it has to be larger than the lower bound.

Returns:

The generated uniformly distributed random number.

Time complexity: depends on the type of the RNG.

igraph_rng_get_unif01 — Generate real, uniform random number from the unit interval

```
igraph_real_t igraph_rng_get_unif01(igraph_rng_t *rng);
```

Arguments:

rng: Pointer to the RNG to use. Use `igraph_rng_default()` here to use the default igraph RNG.

Returns:

The generated uniformly distributed random number.

Time complexity: depends on the type of the RNG.

igraph_rng_get_normal — Normally distributed random numbers

```
igraph_real_t igraph_rng_get_normal(igraph_rng_t *rng,  
                                     igraph_real_t m, igraph_real_t s);
```

Arguments:

rng: Pointer to the RNG to use. Use `igraph_rng_default()` here to use the default igraph RNG.

m: The mean.

s: Standard deviation.

Returns:

The generated normally distributed random number.

Time complexity: depends on the type of the RNG.

igraph_rng_get_geom — Generate geometrically distributed random numbers

```
igraph_real_t igraph_rng_get_geom(igraph_rng_t *rng, igraph_real_t p);
```

Arguments:

rng: Pointer to the RNG to use. Use `igraph_rng_default()` here to use the default igraph RNG.

p: The probability of success in each trial. Must be larger than zero and smaller or equal to 1.

Returns:

The generated geometrically distributed random number.

Time complexity: depends on the type of the RNG.

igraph_rng_get_binom — Generate binomially distributed random numbers

```
igraph_real_t igraph_rng_get_binom(igraph_rng_t *rng, long int n,  
                                    igraph_real_t p);
```

Arguments:

rng: Pointer to the RNG to use. Use `igraph_rng_default()` here to use the default igraph RNG.

n: Number of observations.

p: Probability of an event.

Returns:

The generated binomially distributed random number.

Time complexity: depends on the type of the RNG.

igraph_rng_get_gamma — Generate sample from a Gamma distribution

```
igraph_real_t igraph_rng_get_gamma(igraph_rng_t *rng, igraph_real_t shape,
                                   igraph_real_t scale);
```

Arguments:

rng: Pointer to the RNG to use. Use `igraph_rng_default()` here to use the default igraph RNG.

shape: Shape parameter.

scale: Scale parameter.

Returns:

The generated sample

Time complexity: depends on RNG.

Supported random number generators

By default igraph uses the MT19937 generator. Prior to igraph version 0.6, the generator supplied by the standard C library was used. This means the GLIBC2 generator on GNU libc 2 systems, and maybe the RAND generator on others.

igraph_rngtype_mt19937 — The MT19937 random number generator.

```
const igraph_rng_type_t igraph_rngtype_mt19937 = {
    /* name= */      "MT19937",
    /* min= */       0,
    /* max= */       0xffffffffUL,
    /* init= */      igraph_rng_mt19937_init,
    /* destroy= */   igraph_rng_mt19937_destroy,
    /* seed= */      igraph_rng_mt19937_seed,
    /* get= */       igraph_rng_mt19937_get,
```

```
/* get_real= */  igraph_rng_mt19937_get_real,
/* get_norm= */  0,
/* get_geom= */  0,
/* get_binom= */ 0,
/* get_exp= */   0,
/* get_gamma= */ 0
};
```

The MT19937 generator of Makoto Matsumoto and Takuji Nishimura is a variant of the twisted generalized feedback shift-register algorithm, and is known as the “Mersenne Twister” generator. It has a Mersenne prime period of $2^{19937} - 1$ (about 10^{6000}) and is equi-distributed in 623 dimensions. It has passed the diehard statistical tests. It uses 624 words of state per generator and is comparable in speed to the other generators. The original generator used a default seed of 4357 and choosing `s` equal to zero in `gsl_rng_set` reproduces this. Later versions switched to 5489 as the default seed, you can choose this explicitly via `igraph_rng_seed()` instead if you require it.

For more information see, Makoto Matsumoto and Takuji Nishimura, “Mersenne Twister: A 623-dimensionally equidistributed uniform pseudorandom number generator”. *ACM Transactions on Modeling and Computer Simulation*, Vol. 8, No. 1 (Jan. 1998), Pages 3–30

The generator `igraph_rngtype_mt19937` uses the second revision of the seeding procedure published by the two authors above in 2002. The original seeding procedures could cause spurious artifacts for some seed values.

This generator was ported from the GNU Scientific Library.

`igraph_rngtype_glibc2` — The random number generator introduced in GNU libc 2.

```
const igraph_rng_type_t igraph_rngtype_glibc2 = {
/* name= */      "LIBC",
/* min= */      0,
/* max= */      0x7fffffffUL,
/* init= */      igraph_rng_glibc2_init,
/* destroy= */   igraph_rng_glibc2_destroy,
/* seed= */      igraph_rng_glibc2_seed,
/* get= */       igraph_rng_glibc2_get,
/* get_real= */  igraph_rng_glibc2_get_real,
/* get_norm= */  0,
/* get_geom= */  0,
/* get_binom= */ 0,
/* get_exp= */   0,
/* get_gamma= */ 0
};
```

This is a linear feedback shift register generator with a 128-byte buffer. This generator was the default prior to `igraph` version 0.6, at least on systems relying on GNU libc. This generator was ported from the GNU Scientific Library. It is a reimplementaion and does not call the system `glibc` generator.

`igraph_rngtype_rand` — The old BSD `rand/srand` random number generator.

```
const igraph_rng_type_t igraph_rngtype_rand = {
```

```
/* name= */      "RAND",
/* min=  */      0,
/* max=  */      0x7fffffffUL,
/* init= */      igraph_rng_rand_init,
/* destroy= */   igraph_rng_rand_destroy,
/* seed= */      igraph_rng_rand_seed,
/* get=  */      igraph_rng_rand_get,
/* get_real= */  igraph_rng_rand_get_real,
/* get_norm= */  0,
/* get_geom= */  0,
/* get_binom= */ 0,
/* get_exp= */   0,
/* get_gamma= */ 0
};
```

The sequence is $x_{n+1} = (a x_n + c) \bmod m$ with $a = 1103515245$, $c = 12345$ and $m = 2^{31} = 2147483648$. The seed specifies the initial value, x_1 .

The theoretical value of x_{10001} is 1910041713.

The period of this generator is 2^{31} .

This generator is not very good—the low bits of successive numbers are correlated.

This generator was ported from the GNU Scientific Library.

Use cases

Normal (default) use

If the user does not use any of the RNG functions explicitly, but calls some of the randomized igraph functions, then a default RNG is set up the first time an igraph function needs random numbers. The seed of this RNG is the output of the `time(0)` function call, using the `time` function from the standard C library. This ensures that igraph creates a different random graph, each time the C program is called.

The created default generator is stored internally and can be queried with the `igraph_rng_default()` function.

Reproducible simulations

If reproducible results are needed, then the user should set the seed of the default random number generator explicitly, using the `igraph_rng_seed()` function on the default generator, `igraph_rng_default()`. When setting the seed to the same number, igraph generates exactly the same random graph (or series of random graphs).

Changing the default generator

By default igraph uses the `igraph_rng_default()` random number generator. This can be changed any time by calling `igraph_rng_set_default()`, with an already initialized random number generator. Note that the old (replaced) generator is not destroyed, so no memory is deallocated.

Using multiple generators

igraph also provides functions to set up multiple random number generators, using the `igraph_rng_init()` function, and then generating random numbers from them, e.g. with `igraph_rng_get_integer()` and/or `igraph_rng_get_unif()` calls.

Note that initializing a new random number generator is independent of the generator that the igraph functions themselves use. If you want to replace that, then please use `igraph_rng_set_default()`.

Example

Example 8.1. File `examples/simple/random_seed.c`

Chapter 9. Graph generators

Graph generators create graphs.

Almost all functions which create graph objects are documented here. The exceptions are `igraph_induced_subgraph()` and alike, these create graphs based on another graph.

Deterministic graph generators

`igraph_create` — Creates a graph with the specified edges.

```
int igraph_create(igraph_t *graph, const igraph_vector_t *edges,
                  igraph_integer_t n, igraph_bool_t directed);
```

Arguments:

- graph*: An uninitialized graph object.
- edges*: The edges to add, the first two elements are the first edge, etc.
- n*: The number of vertices in the graph, if smaller or equal to the highest vertex id in the *edges* vector it will be increased automatically. So it is safe to give 0 here.
- directed*: Boolean, whether to create a directed graph or not. If yes, then the first edge points from the first vertex id in *edges* to the second, etc.

Returns:

Error code: `IGRAPH_EINVEVECTOR`: invalid edges vector (odd number of vertices).
`IGRAPH_EINVVID`: invalid (negative) vertex id.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of vertices, $|E|$ the number of edges in the graph.

Example 9.1. File `examples/simple/igraph_create.c`

`igraph_small` — Shorthand to create a small graph, giving the edges as arguments.

```
int igraph_small(igraph_t *graph, igraph_integer_t n, igraph_bool_t directed,
                 ...);
```

This function is handy when a relatively small graph needs to be created. Instead of giving the edges as a vector, they are given simply as arguments and a `-1` needs to be given after the last meaningful edge argument.

Note that only graphs which have vertices less than the highest value of the 'int' type can be created this way. If you give larger values then the result is undefined.

Arguments:

- graph*: Pointer to an uninitialized graph object. The result will be stored here.
- n*: The number of vertices in the graph; a nonnegative integer.
- directed*: Logical constant; gives whether the graph should be directed. Supported values are:
- | | |
|-------------------|---|
| IGRAPH_DIRECTED | The graph to be created will be <i>directed</i> . |
| IGRAPH_UNDIRECTED | The graph to be created will be <i>undirected</i> . |
- ...: The additional arguments giving the edges of the graph. Don't forget to supply an additional '-' after the last (meaningful) argument.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph to create.

Example 9.2. File `examples/simple/igraph_small.c`

igraph_adjacency — Creates a graph from an adjacency matrix.

```
int igraph_adjacency(igraph_t *graph, igraph_matrix_t *adjmatrix,
                    igraph_adjacency_t mode);
```

The order of the vertices in the matrix is preserved, i.e. the vertex corresponding to the first row/column will be vertex with id 0, the next row is for vertex 1, etc.

Arguments:

- graph*: Pointer to an uninitialized graph object.
- adjmatrix*: The adjacency matrix. How it is interpreted depends on the *mode* argument.
- mode*: Constant to specify how the given matrix is interpreted as an adjacency matrix. Possible values ($A(i,j)$ is the element in row i and column j in the adjacency matrix *adjmatrix*):
- | | |
|-----------------------|---|
| IGRAPH_ADJ_DIRECTED | the graph will be directed and an element gives the number of edges between two vertices. |
| IGRAPH_ADJ_UNDIRECTED | this is the same as IGRAPH_ADJ_MAX, for convenience. |
| IGRAPH_ADJ_MAX | undirected graph will be created and the number of edges between vertices i and j is $\max(A(i,j), A(j,i))$. |
| IGRAPH_ADJ_MIN | undirected graph will be created with $\min(A(i,j), A(j,i))$ edges between vertices i and j . |
| IGRAPH_ADJ_PLUS | undirected graph will be created with $A(i,j)+A(j,i)$ edges between vertices i and j . |

IGRAPH_ADJ_UPPER	undirected graph will be created, only the upper right triangle (including the diagonal) is used for the number of edges.
IGRAPH_ADJ_LOWER	undirected graph will be created, only the lower left triangle (including the diagonal) is used for creating the edges.

Returns:

Error code, IGRAPH_NONSQUARE: non-square matrix.

Time complexity: $O(|V||V|)$, $|V|$ is the number of vertices in the graph.

Example 9.3. File `examples/simple/igraph_adjacency.c`

igraph_weighted_adjacency — Creates a graph from a weighted adjacency matrix.

```
int igraph_weighted_adjacency(igraph_t *graph, igraph_matrix_t *adjmatrix,
                             igraph_adjacency_t mode, const char* attr,
                             igraph_bool_t loops);
```

The order of the vertices in the matrix is preserved, i.e. the vertex corresponding to the first row/column will be vertex with id 0, the next row is for vertex 1, etc.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.										
<i>adjmatrix</i> :	The weighted adjacency matrix. How it is interpreted depends on the <i>mode</i> argument. The common feature is that edges with zero weights are considered nonexistent (however, negative weights are permitted).										
<i>mode</i> :	Constant to specify how the given matrix is interpreted as an adjacency matrix. Possible values ($A(i,j)$ is the element in row i and column j in the adjacency matrix <i>adjmatrix</i>): <table><tr><td>IGRAPH_ADJ_DIRECTED</td><td>the graph will be directed and an element gives the weight of the edge between two vertices.</td></tr><tr><td>IGRAPH_ADJ_UNDIRECTED</td><td>this is the same as IGRAPH_ADJ_MAX, for convenience.</td></tr><tr><td>IGRAPH_ADJ_MAX</td><td>undirected graph will be created and the weight of the edge between vertices i and j is $\max(A(i,j), A(j,i))$.</td></tr><tr><td>IGRAPH_ADJ_MIN</td><td>undirected graph will be created with edge weight $\min(A(i,j), A(j,i))$ between vertices i and j.</td></tr><tr><td>IGRAPH_ADJ_PLUS</td><td>undirected graph will be created with edge weight $A(i,j)+A(j,i)$ between vertices i and j.</td></tr></table>	IGRAPH_ADJ_DIRECTED	the graph will be directed and an element gives the weight of the edge between two vertices.	IGRAPH_ADJ_UNDIRECTED	this is the same as IGRAPH_ADJ_MAX, for convenience.	IGRAPH_ADJ_MAX	undirected graph will be created and the weight of the edge between vertices i and j is $\max(A(i,j), A(j,i))$.	IGRAPH_ADJ_MIN	undirected graph will be created with edge weight $\min(A(i,j), A(j,i))$ between vertices i and j .	IGRAPH_ADJ_PLUS	undirected graph will be created with edge weight $A(i,j)+A(j,i)$ between vertices i and j .
IGRAPH_ADJ_DIRECTED	the graph will be directed and an element gives the weight of the edge between two vertices.										
IGRAPH_ADJ_UNDIRECTED	this is the same as IGRAPH_ADJ_MAX, for convenience.										
IGRAPH_ADJ_MAX	undirected graph will be created and the weight of the edge between vertices i and j is $\max(A(i,j), A(j,i))$.										
IGRAPH_ADJ_MIN	undirected graph will be created with edge weight $\min(A(i,j), A(j,i))$ between vertices i and j .										
IGRAPH_ADJ_PLUS	undirected graph will be created with edge weight $A(i,j)+A(j,i)$ between vertices i and j .										

	IGRAPH_ADJ_UPPER	undirected graph will be created, only the upper right triangle (including the diagonal) is used for the edge weights.
	IGRAPH_ADJ_LOWER	undirected graph will be created, only the lower left triangle (including the diagonal) is used for the edge weights.
<i>attr</i> :	the name of the attribute that will store the edge weights. If NULL , it will use weight as the attribute name.	
<i>loops</i> :	Logical scalar, whether to ignore the diagonal elements in the adjacency matrix.	

Returns:

Error code, IGRAPH_NONSQUARE: non-square matrix.

Time complexity: $O(|V||V|)$, $|V|$ is the number of vertices in the graph.

Example 9.4. File `examples/simple/igraph_weighted_adjacency.c`

igraph_adjlist — Creates a graph from an adjacency list.

```
int igraph_adjlist(igraph_t *graph, const igraph_adjlist_t *adjlist,
                  igraph_neimode_t mode, igraph_bool_t duplicate);
```

An adjacency list is a list of vectors, containing the neighbors of all vertices. For operations that involve many changes to the graph structure, it is recommended that you convert the graph into an adjacency list via `igraph_adjlist_init()`, perform the modifications (these are cheap for an adjacency list) and then recreate the igraph graph via this function.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.
<i>adjlist</i> :	The adjacency list.
<i>mode</i> :	Whether or not to create a directed graph. IGRAPH_ALL means an undirected graph, IGRAPH_OUT means a directed graph from an out-adjacency list (i.e. each list contains the successors of the corresponding vertices), IGRAPH_IN means a directed graph from an in-adjacency list
<i>duplicate</i> :	Logical, for undirected graphs this specified whether each edge is included twice, in the vectors of both adjacent vertices. If this is false (0), then it is assumed that every edge is included only once. This argument is ignored for directed graphs.

Returns:

Error code.

See also:

`igraph_adjlist_init()` for the opposite operation.

Time complexity: $O(|V|+|E|)$.

igraph_star — Creates a *star* graph, every vertex connects only to the center.

```
int igraph_star(igraph_t *graph, igraph_integer_t n, igraph_star_mode_t mode,
               igraph_integer_t center);
```

Arguments:

graph: Pointer to an uninitialized graph object, this will be the result.

n: Integer constant, the number of vertices in the graph.

mode: Constant, gives the type of the star graph to create. Possible values:

IGRAPH_STAR_OUT	directed star graph, edges point <i>from</i> the center to the other vertices.
IGRAPH_STAR_IN	directed star graph, edges point <i>to</i> the center from the other vertices.
IGRAPH_STAR_MUTUAL	directed star graph with mutual edges.
IGRAPH_STAR_UNDIRECTED	an undirected star graph is created.

center: Id of the vertex which will be the center of the graph.

Returns:

Error code:

IGRAPH_EINVVID invalid number of vertices.

IGRAPH_EINVAL invalid center vertex.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|V|)$, the number of vertices in the graph.

See also:

`igraph_lattice()`, `igraph_ring()`, `igraph_tree()` for creating other regular structures.

Example 9.5. File `examples/simple/igraph_star.c`

igraph_lattice — Arbitrary dimensional square lattices.

```
int igraph_lattice(igraph_t *graph, const igraph_vector_t *dimvector,
                  igraph_integer_t nei, igraph_bool_t directed, igraph_bool_t mutual,
                  igraph_bool_t circular);
```

Creates d-dimensional square lattices of the given size. Optionally, the lattice can be made periodic, and the neighbors within a given graph distance can be connected.

In the zero-dimensional case, the singleton graph is returned.

The vertices of the resulting graph are ordered such that the index of the vertex at position $(i_0, i_1, i_2, \dots, i_d)$ in a lattice of size (n_0, n_1, \dots, n_d) will be $i_0 + n_0 * i_1 + n_0 * n_1 * i_2 + \dots$.

Arguments:

- graph*: An uninitialized graph object.
- dimvector*: Vector giving the sizes of the lattice in each of its dimensions. The dimension of the lattice will be the same as the length of this vector.
- nei*: Integer value giving the distance (number of steps) within which two vertices will be connected.
- directed*: Boolean, whether to create a directed graph. If the *mutual* and *circular* arguments are not set to true, edges will be directed from lower-index vertices towards higher-index ones.
- mutual*: Boolean, if the graph is directed this gives whether to create all connections as mutual.
- circular*: Boolean, defines whether the generated lattice is periodic.

Returns:

Error code: IGRAPH_EINVAL: invalid (negative) dimension vector.

Time complexity: If *nei* is less than two then it is $O(|V|+|E|)$ (as far as I remember), $|V|$ and $|E|$ are the number of vertices and edges in the generated graph. Otherwise it is $O(|V|*d^k+|E|)$, d is the average degree of the graph, k is the *nei* argument.

igraph_ring — Creates a *ring* graph, a one dimensional lattice.

```
int igraph_ring(igraph_t *graph, igraph_integer_t n, igraph_bool_t directed,
                igraph_bool_t mutual, igraph_bool_t circular);
```

An undirected (circular) ring on n vertices is commonly known in graph theory as the cycle graph C_n .

Arguments:

- graph*: Pointer to an uninitialized graph object.
- n*: The number of vertices in the ring.
- directed*: Logical, whether to create a directed ring.
- mutual*: Logical, whether to create mutual edges in a directed ring. It is ignored for undirected graphs.

circular: Logical, if false, the ring will be open (this is not a real *ring* actually).

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices.

Time complexity: $O(|V|)$, the number of vertices in the graph.

See also:

`igraph_lattice()` for generating more general lattices.

Example 9.6. File `examples/simple/igraph_ring.c`

igraph_tree — Creates a tree in which almost all vertices have the same number of children.

```
int igraph_tree(igraph_t *graph, igraph_integer_t n, igraph_integer_t children,
                igraph_tree_mode_t type);
```

Arguments:

graph: Pointer to an uninitialized graph object.

n: Integer, the number of vertices in the graph.

children: Integer, the number of children of a vertex in the tree.

type: Constant, gives whether to create a directed tree, and if this is the case, also its orientation. Possible values:

IGRAPH_TREE_OUT	directed tree, the edges point from the parents to their children,
IGRAPH_TREE_IN	directed tree, the edges point from the children to their parents.
IGRAPH_TREE_UNDIRECTED	undirected tree.

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices. IGRAPH_INVMODE: invalid mode argument.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph.

See also:

`igraph_lattice()`, `igraph_star()` for creating other regular structures; `igraph_from_prufer()` for creating arbitrary trees; `igraph_tree_game()` for uniform random sampling of trees.

Example 9.7. File `examples/simple/igraph_tree.c`

igraph_full — Creates a full graph (directed or undirected, with or without loops).

```
int igraph_full(igraph_t *graph, igraph_integer_t n, igraph_bool_t directed,
               igraph_bool_t loops);
```

In a full graph every possible edge is present, every vertex is connected to every other vertex. A full graph in *igraph* should be distinguished from the concept of complete graphs as used in graph theory. If n is a positive integer, then the complete graph K_n on n vertices is the undirected simple graph with the following property. For any distinct pair (u,v) of vertices in K_n , uv (or equivalently vu) is an edge of K_n . In *igraph*, a full graph on n vertices can be K_n , a directed version of K_n , or K_n with at least one loop edge. In any case, if F is a full graph on n vertices as generated by *igraph*, then K_n is a subgraph of the undirected version of F .

Arguments:

graph: Pointer to an uninitialized graph object.

n: Integer, the number of vertices in the graph.

directed: Logical, whether to create a directed graph.

loops: Logical, whether to include self-edges (loops).

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of vertices, $|E|$ the number of edges in the graph. Of course this is the same as $O(|E|)=O(|V||V|)$ here.

See also:

`igraph_lattice()`, `igraph_star()`, `igraph_tree()` for creating other regular structures.

Example 9.8. File `examples/simple/igraph_full.c`

igraph_full_citation — Creates a full citation graph

```
int igraph_full_citation(igraph_t *graph, igraph_integer_t n,
                        igraph_bool_t directed);
```

This is a directed graph, where every $i \rightarrow j$ edge is present if and only if $j < i$. If the directed argument is zero then an undirected graph is created, and it is just a full graph.

Arguments:

graph: Pointer to an uninitialized graph object, the result is stored here.

n: The number of vertices.

directed: Whether to create a directed graph. If zero an undirected graph is created.

Returns:

Error code.

Time complexity: $O(|V|^2)$, as we have many edges.

igraph_realize_degree_sequence — Generates a graph with the given degree sequence.

```
int igraph_realize_degree_sequence(  
    igraph_t *graph,  
    const igraph_vector_t *outdeg, const igraph_vector_t *indeg,  
    igraph_edge_type_sw_t allowed_edge_types,  
    igraph_realize_degseq_t method);
```

This function generates an undirected graph that realizes a given degree sequence, or a directed graph that realized a given pair of out- and in-degree sequences.

Simple undirected graphs are constructed using the Havel-Hakimi algorithm (undirected case), or the analogous Kleitman-Wang algorithm (directed case). These algorithms work by choosing an arbitrary vertex and connecting all its stubs to other vertices of highest degree. In the directed case, the "highest" (in, out) degree pairs are determined based on lexicographic ordering. This step is repeated until all degrees have been connected up.

Loopless multigraphs are generated using an analogous algorithm: an arbitrary vertex is chosen, and it is connected with a single connection to a highest remaining degree vertex. If self-loops are also allowed, the same algorithm is used, but if a non-zero vertex remains at the end of the procedure, the graph is completed by adding self-loops to it. Thus, the result will contain at most one vertex with self-loops.

The *method* parameter controls the order in which the vertices to be connected are chosen.

References:

V. Havel, Poznámka o existenci konečných grafů (A remark on the existence of finite graphs), *Časopis pro pěstování matematiky* 80, 477-480 (1955). <http://eudml.org/doc/19050>

S. L. Hakimi, On Realizability of a Set of Integers as Degrees of the Vertices of a Linear Graph, *Journal of the SIAM* 10, 3 (1962). <https://www.jstor.org/stable/2098746>

D. J. Kleitman and D. L. Wang, Algorithms for Constructing Graphs and Digraphs with Given Valences and Factors, *Discrete Mathematics* 6, 1 (1973). [https://doi.org/10.1016/0012-365X\(73\)90037-X](https://doi.org/10.1016/0012-365X(73)90037-X)

Sz. Horvát and C. D. Modes, Connectivity matters: Construction and exact random sampling of connected graphs (2020). <https://arxiv.org/abs/2009.03747>

Arguments:

graph: Pointer to an uninitialized graph object.

outdeg: The degree sequence of an undirected graph (if *indeg* is NULL), or the out-degree sequence of a directed graph (if *indeg* is given).

<i>indeg</i> :	The in-degree sequence of a directed graph. Pass NULL to generate an undirected graph.	
<i>allowed_edge_types</i> :	The types of edges to allow in the graph. For directed graphs, only IGRAPH_SIMPLE_SW is implemented at this moment. For undirected graphs, the following values are valid:	
	IGRAPH_SIMPLE_SW	simple graphs (i.e. no self-loops or multi-edges allowed).
	IGRAPH_LOOPS_SW	single self-loops are allowed, but not multi-edges; currently not implemented.
	IGRAPH_MULTI_SW	multi-edges are allowed, but not self-loops.
	IGRAPH_LOOPS_SW IGRAPH_MULTI_SW	both self-loops and multi-edges are allowed.
<i>method</i> :	The method to generate the graph. Possible values:	
	IGRAPH_REALIZE_DEGSEQ_S- SMALLEST	The vertex with smallest remaining degree is selected first. The result is usually a graph with high negative degree assortativity. In the undirected case, this method is guaranteed to generate a connected graph, regardless of whether multi-edges are allowed, provided that a connected realization exists (see Horvát and Modes, 2020, as well as http://szhorvat.net/pelican/hh-connected-graphs.html). In the directed case it tends to generate weakly connected graphs, but this is not guaranteed.
	IGRAPH_REALIZE_DEGSEQ- Q_LARGEST	The vertex with the largest remaining degree is selected first. The result is usually a graph with high positive degree assortativity, and is often disconnected.
	IGRAPH_REALIZE_DEGSEQ- Q_INDEX	The vertices are selected in order of their index (i.e. their position in the degree vector). Note that sorting the degree vector and using the INDEX method is not equivalent to the SMALLEST method above, as SMALLEST uses the smallest <i>remaining</i> degree for selecting vertices, not the smallest <i>initial</i> degree.

Returns:

Error code:

IGRAPH_UNIMPLEMENTED	The requested method is not implemented.
IGRAPH_ENOMEM	There is not enough memory to perform the operation.
IGRAPH_EINVAL	Invalid method parameter, or invalid in- and/or out-degree vectors. The degree vectors should be non-negative, the length and sum of <i>outdeg</i> and <i>indeg</i> should match for directed graphs.

See also:

`igraph_is_graphical()` to test graphicality without generating a graph; `igraph_degree_sequence_game()` to generate random graphs with a given degree sequence; `igraph_k_regular_game()` to generate random regular graphs; `igraph_rewire()` to randomly rewire the edges of a graph while preserving its degree sequence.

igraph_famous — Create a famous graph by simply providing its name.

```
int igraph_famous(igraph_t *graph, const char *name);
```

The name of the graph can be simply supplied as a string. Note that this function creates graphs which don't take any parameters, there are separate functions for graphs with parameters, e.g. `igraph_full()` for creating a full graph.

The following graphs are supported:

Bull	The bull graph, 5 vertices, 5 edges, resembles the head of a bull if drawn properly.
Chvatal	This is the smallest triangle-free graph that is both 4-chromatic and 4-regular. According to the Grunbaum conjecture there exists an m -regular, m -chromatic graph with n vertices for every $m > 1$ and $n > 2$. The Chvatal graph is an example for $m=4$ and $n=12$. It has 24 edges.
Coxeter	A non-Hamiltonian cubic symmetric graph with 28 vertices and 42 edges.
Cubical	The Platonic graph of the cube. A convex regular polyhedron with 8 vertices and 12 edges.
Diamond	A graph with 4 vertices and 5 edges, resembles a schematic diamond if drawn properly.
Dodecahedral, Dodecahedron	Another Platonic solid with 20 vertices and 30 edges.
Folkman	The semisymmetric graph with minimum number of vertices, 20 and 40 edges. A semisymmetric graph is regular, edge transitive and not vertex transitive.

Franklin	This is a graph whose embedding to the Klein bottle can be colored with six colors, it is a counterexample to the necessity of the Heawood conjecture on a Klein bottle. It has 12 vertices and 18 edges.
Frucht	The Frucht Graph is the smallest cubical graph whose automorphism group consists only of the identity element. It has 12 vertices and 18 edges.
Grotzsch	The Grötzsch graph is a triangle-free graph with 11 vertices, 20 edges, and chromatic number 4. It is named after German mathematician Herbert Grötzsch, and its existence demonstrates that the assumption of planarity is necessary in Grötzsch's theorem that every triangle-free planar graph is 3-colorable.
Heawood	The Heawood graph is an undirected graph with 14 vertices and 21 edges. The graph is cubic, and all cycles in the graph have six or more edges. Every smaller cubic graph has shorter cycles, so this graph is the 6-cage, the smallest cubic graph of girth 6.
Herschel	The Herschel graph is the smallest nonhamiltonian polyhedral graph. It is the unique such graph on 11 nodes, and has 18 edges.
House	The house graph is a 5-vertex, 6-edge graph, the schematic draw of a house if drawn properly, basically a triangle on top of a square.
HouseX	The same as the house graph with an X in the square. 5 vertices and 8 edges.
Icosahedral, Icosahedron	A Platonic solid with 12 vertices and 30 edges.
Krackhardt_Kite	A social network with 10 vertices and 18 edges. Krackhardt, D. Assessing the Political Landscape: Structure, Cognition, and Power in Organizations. Admin. Sci. Quart. 35, 342-369, 1990.
Levi	The graph is a 4-arc transitive cubic graph, it has 30 vertices and 45 edges.
McGee	The McGee graph is the unique 3-regular 7-cage graph, it has 24 vertices and 36 edges.
Meredith	The Meredith graph is a quartic graph on 70 nodes and 140 edges that is a counterexample to the conjecture that every 4-regular 4-connected graph is Hamiltonian.
Noperfectmatching	A connected graph with 16 vertices and 27 edges containing no perfect matching. A matching in a graph is a set of pairwise non-incident edges; that is, no two edges share a common vertex. A perfect matching is a matching which covers all vertices of the graph.
Nonline	A graph whose connected components are the 9 graphs whose presence as a vertex-induced subgraph in a graph makes a non-line graph. It has 50 vertices and 72 edges.
Octahedral, Octahedron	Platonic solid with 6 vertices and 12 edges.
Petersen	A 3-regular graph with 10 vertices and 15 edges. It is the smallest hypohamiltonian graph, i.e. it is non-hamiltonian but removing any single vertex from it makes it Hamiltonian.

Robertson	The unique (4,5)-cage graph, i.e. a 4-regular graph of girth 5. It has 19 vertices and 38 edges.
Smallestcyclicgroup	A smallest nontrivial graph whose automorphism group is cyclic. It has 9 vertices and 15 edges.
Tetrahedral, Tetrahedron	Platonic solid with 4 vertices and 6 edges.
Thomassen	The smallest hypotraceable graph, on 34 vertices and 52 edges. A hypotraceable graph does not contain a Hamiltonian path but after removing any single vertex from it the remainder always contains a Hamiltonian path. A graph containing a Hamiltonian path is called traceable.
Tutte	Tait's Hamiltonian graph conjecture states that every 3-connected 3-regular planar graph is Hamiltonian. This graph is a counterexample. It has 46 vertices and 69 edges.
Uniquely3colorable	Returns a 12-vertex, triangle-free graph with chromatic number 3 that is uniquely 3-colorable.
Walther	An identity graph with 25 vertices and 31 edges. An identity graph has a single graph automorphism, the trivial one.
Zachary	Social network of friendships between 34 members of a karate club at a US university in the 1970s. See W. W. Zachary, An information flow model for conflict and fission in small groups, Journal of Anthropological Research 33, 452-473 (1977).

Arguments:

graph: Pointer to an uninitialized graph object.

name: Character constant, the name of the graph to be created, it is case insensitive.

Returns:

Error code, IGRAPH_EINVAL if there is no graph with the given name.

See also:

Other functions for creating graph structures: `igraph_ring()`, `igraph_tree()`, `igraph_lattice()`, `igraph_full()`.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph.

igraph_lcf — Creates a graph from LCF notation.

```
int igraph_lcf(igraph_t *graph, igraph_integer_t n, ...);
```

LCF is short for Lederberg-Coxeter-Frucht, it is a concise notation for 3-regular Hamiltonian graphs. It consists of three parameters: the number of vertices in the graph, a list of shifts giving additional edges to a cycle backbone, and another integer giving how many times the shifts should be performed. See <http://mathworld.wolfram.com/LCFNotation.html> for details.

Arguments:

graph: Pointer to an uninitialized graph object.

n: Integer, the number of vertices in the graph.

...: The shifts and the number of repeats for the shifts, plus an additional 0 to mark the end of the arguments.

Returns:

Error code.

See also:

See `igraph_lcf_vector()` for a similar function using a `vector_t` instead of the variable length argument list.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

Example 9.9. File `examples/simple/igraph_lcf.c`

igraph_lcf_vector — Creates a graph from LCF notation.

```
int igraph_lcf_vector(igraph_t *graph, igraph_integer_t n,
                     const igraph_vector_t *shifts,
                     igraph_integer_t repeats);
```

This function is essentially the same as `igraph_lcf()`, only the way for giving the arguments is different. See `igraph_lcf()` for details.

Arguments:

graph: Pointer to an uninitialized graph object.

n: Integer constant giving the number of vertices.

shifts: A vector giving the shifts.

repeats: An integer constant giving the number of repeats for the shifts.

Returns:

Error code.

See also:

`igraph_lcf()`, `igraph_extended_chordal_ring()`

Time complexity: $O(|V|+|E|)$, linear in the number of vertices plus the number of edges.

igraph_from_prufer — Generates a tree from a Prüfer sequence.

```
int igraph_from_prufer(igraph_t *graph, const igraph_vector_int_t *prufer);
```

A Prüfer sequence is a unique sequence of integers associated with a labelled tree. A tree on n vertices can be represented by a sequence of $n-2$ integers, each between 0 and $n-1$ (inclusive). The algorithm used by this function is based on Paulius Micikevičius, Saverio Caminiti, Narsingh Deo: Linear-time Algorithms for Encoding Trees as Sequences of Node Labels

Arguments:

graph: Pointer to an uninitialized graph object.

prufer: The Prüfer sequence

Returns:

Error code:

IGRAPH_ENOMEM there is not enough memory to perform the operation.

IGRAPH_EINVAL invalid Prüfer sequence given

See also:

`igraph_to_prufer()`, `igraph_tree()`, `igraph_tree_game()`

igraph_atlas — Create a small graph from the “Graph Atlas”.

```
int igraph_atlas(igraph_t *graph, int number);
```

The number of the graph is given as a parameter. The graphs are listed:

1. in increasing order of number of nodes;
2. for a fixed number of nodes, in increasing order of the number of edges;
3. for fixed numbers of nodes and edges, in increasing order of the degree sequence, for example 111223 < 112222;
4. for fixed degree sequence, in increasing number of automorphisms.

The data was converted from the NetworkX software package, see <http://networkx.github.io>.

See *An Atlas of Graphs* by Ronald C. Read and Robin J. Wilson, Oxford University Press, 1998.

Arguments:

graph: Pointer to an uninitialized graph object.

number: The number of the graph to generate.

Added in version 0.2.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

Example 9.10. File `examples/simple/igraph_atlas.c`

igraph_de_bruijn — Generate a de Bruijn graph.

```
int igraph_de_bruijn(igraph_t *graph, igraph_integer_t m, igraph_integer_t n);
```

A de Bruijn graph represents relationships between strings. An alphabet of m letters are used and strings of length n are considered. A vertex corresponds to every possible string and there is a directed edge from vertex v to vertex w if the string of v can be transformed into the string of w by removing its first letter and appending a letter to it.

Please note that the graph will have m to the power n vertices and even more edges, so probably you don't want to supply too big numbers for m and n .

De Bruijn graphs have some interesting properties, please see another source, e.g. Wikipedia for details.

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.

m: Integer, the number of letters in the alphabet.

n: Integer, the length of the strings.

Returns:

Error code.

See also:

`igraph_kautz()`.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

igraph_kautz — Generate a Kautz graph.

```
int igraph_kautz(igraph_t *graph, igraph_integer_t m, igraph_integer_t n);
```

A Kautz graph is a labeled graph, vertices are labeled by strings of length $n+1$ above an alphabet with $m+1$ letters, with the restriction that every two consecutive letters in the string must be different. There is a directed edge from a vertex v to another vertex w if it is possible to transform the string of v into the string of w by removing the first letter and appending a letter to it. For string length 1 the new letter cannot equal the old letter, so there are no loops.

Kautz graphs have some interesting properties, see e.g. Wikipedia for details.

Vincent Matossian wrote the first version of this function in R, thanks.

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.

m: Integer, $m+1$ is the number of letters in the alphabet.
n: Integer, $n+1$ is the length of the strings.

Returns:

Error code.

See also:

`igraph_de_bruijn()`.

Time complexity: $O(|V| * [(m+1)/m]^n + |E|)$, in practice it is more like $O(|V| + |E|)$. $|V|$ is the number of vertices, $|E|$ is the number of edges and m and n are the corresponding arguments.

igraph_extended_chordal_ring — Create an extended chordal ring.

```
int igraph_extended_chordal_ring(
    igraph_t *graph, igraph_integer_t nodes, const igraph_matrix_t *W,
    igraph_bool_t directed);
```

An extended chordal ring is a cycle graph with additional chords connecting its vertices. Each row L of the matrix W specifies a set of chords to be inserted, in the following way: vertex i will connect to a vertex $L[(i \bmod p)]$ steps ahead of it along the cycle, where p is the length of L . In other words, vertex i will be connected to vertex $(i + L[(i \bmod p)]) \bmod nodes$. If multiple edges are defined in this way, this will output a non-simple graph. The result can be simplified using `igraph_simplify()`.

See also Kotsis, G: Interconnection Topologies for Parallel Processing Systems, PARS Mitteilungen 11, 1-6, 1993. The igraph extended chordal rings are not identical to the ones in the paper. In igraph the matrix specifies which edges to add. In the paper, a condition is specified which should simultaneously hold between two endpoints and the reverse endpoints.

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.
nodes: Integer constant, the number of vertices in the graph. It must be at least 3.
W: The matrix specifying the extra edges. The number of columns should divide the number of total vertices. The elements are allowed to be negative.
directed: Whether the graph should be directed.

Returns:

Error code.

See also:

`igraph_ring()`, `igraph_lcf()`, `igraph_lcf_vector()`.

Time complexity: $O(|V| + |E|)$, the number of vertices plus the number of edges.

Games: randomized graph generators

Games are randomized graph generators. Randomization means that they generate a different graph every time you call them.

igraph_grg_game — Generates a geometric random graph.

```
int igraph_grg_game(igraph_t *graph, igraph_integer_t nodes,
                    igraph_real_t radius, igraph_bool_t torus,
                    igraph_vector_t *x, igraph_vector_t *y);
```

A geometric random graph is created by dropping points (i.e. vertices) randomly on the unit square and then connecting all those pairs which are less than `radius` apart in Euclidean distance.

Original code contributed by Keith Briggs, thanks Keith.

Arguments:

- graph*: Pointer to an uninitialized graph object.
- nodes*: The number of vertices in the graph.
- radius*: The radius within which the vertices will be connected.
- torus*: Logical constant. If true, periodic boundary conditions will be used, i.e. the vertices are assumed to be on a torus instead of a square.
- x*: An initialized vector or NULL. If not NULL, the points' x coordinates will be returned here.
- y*: An initialized vector or NULL. If not NULL, the points' y coordinates will be returned here.

Returns:

Error code.

Time complexity: TODO, less than $O(|V|^2 + |E|)$.

Example 9.11. File `examples/simple/igraph_grg_game.c`

igraph_barabasi_game — Generates a graph based on the Barabási-Albert model.

```
int igraph_barabasi_game(igraph_t *graph, igraph_integer_t n,
                        igraph_real_t power,
                        igraph_integer_t m,
                        const igraph_vector_t *outseq,
                        igraph_bool_t outpref,
                        igraph_real_t A,
                        igraph_bool_t directed,
                        igraph_barabasi_algorithm_t algo,
```

```
const igraph_t *start_from);
```

Arguments:

<i>graph</i> :	An uninitialized graph object.	
<i>n</i> :	The number of vertices in the graph.	
<i>power</i> :	Power of the preferential attachment. The probability that a vertex is cited is proportional to $d^{\text{power}+A}$, where d is its degree (see also the <i>outpref</i> argument), <i>power</i> and <i>A</i> are given by arguments. In the classic preferential attachment model <i>power</i> =1.	
<i>m</i> :	The number of outgoing edges generated for each vertex. (Only if <i>outseq</i> is NULL.)	
<i>outseq</i> :	Gives the (out-)degrees of the vertices. If this is constant, this can be a NULL pointer or an empty (but initialized!) vector, in this case <i>m</i> contains the constant out-degree. The very first vertex has by definition no outgoing edges, so the first number in this vector is ignored.	
<i>outpref</i> :	Boolean, if true not only the in- but also the out-degree of a vertex increases its citation probability. I.e., the citation probability is determined by the total degree of the vertices. Ignored and assumed to be true if the graph being generated is undirected.	
<i>A</i> :	The probability that a vertex is cited is proportional to $d^{\text{power}+A}$, where d is its degree (see also the <i>outpref</i> argument), <i>power</i> and <i>A</i> are given by arguments. In the previous versions of the function this parameter was implicitly set to one.	
<i>directed</i> :	Boolean, whether to generate a directed graph.	
<i>algo</i> :	The algorithm to use to generate the network. Possible values:	
	IGRAPH_BARABASI_BAG	This is the algorithm that was previously (before version 0.6) solely implemented in igraph. It works by putting the ids of the vertices into a bag (multiset, really), exactly as many times as their (in-)degree, plus once more. Then the required number of cited vertices are drawn from the bag, with replacement. This method might generate multiple edges. It only works if <i>power</i> =1 and <i>A</i> =1.
	IGRAPH_BARABASI_PSUMTREE	This algorithm uses a partial prefix-sum tree to generate the graph. It does not generate multiple edges and works for any <i>power</i> and <i>A</i> values.
	IGRAPH_BARABASI_PSUMTREE_MULTIPLE	This algorithm also uses a partial prefix-sum tree to generate the graph. The difference is, that now multiple edges are allowed. This method was implemented under the name <i>igraph_non-linear_barabasi_game</i> before version 0.6.
<i>start_from</i> :	Either a null pointer, or a graph. In the former case, the starting configuration is a clique of size <i>m</i> . In the latter case, the graph is a starting configuration. The graph	

must be non-empty, i.e. it must have at least one vertex. If a graph is supplied here and the *outseq* argument is also given, then *outseq* should only contain information on the vertices that are not in the *start_from* graph.

Returns:

Error code: IGRAPH_EINVAL: invalid *n*, *m* or *outseq* parameter.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

Example 9.12. File `examples/simple/igraph_barabasi_game.c`

Example 9.13. File `examples/simple/igraph_barabasi_game2.c`

igraph_erdos_renyi_game — Generates a random (Erd#s-Rényi) graph.

```
int igraph_erdos_renyi_game(igraph_t *graph, igraph_erdos_renyi_t type,
                             igraph_integer_t n, igraph_real_t p_or_m,
                             igraph_bool_t directed, igraph_bool_t loops);
```

Arguments:

graph: Pointer to an uninitialized graph object.

type: The type of the random graph, possible values:

IGRAPH_ERDOS_RENYI_GNM $G(n,m)$ graph, *m* edges are selected uniformly randomly in a graph with *n* vertices.

IGRAPH_ERDOS_RENYI_GNP $G(n,p)$ graph, every possible edge is included in the graph with probability *p*.

n: The number of vertices in the graph.

p_or_m: This is the *p* parameter for $G(n,p)$ graphs and the *m* parameter for $G(n,m)$ graphs.

directed: Logical, whether to generate a directed graph.

loops: Logical, whether to generate loops (self) edges.

Returns:

Error code: IGRAPH_EINVAL: invalid *type*, *n*, *p* or *m* parameter. IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph.

See also:

`igraph_barabasi_game()`, `igraph_growing_random_game()`

Example 9.14. File `examples/simple/igraph_erdos_renyi_game.c`

igraph_watts_strogatz_game — The Watts-Strogatz small-world model.

```
int igraph_watts_strogatz_game(igraph_t *graph, igraph_integer_t dim,
                               igraph_integer_t size, igraph_integer_t nei,
                               igraph_real_t p, igraph_bool_t loops,
                               igraph_bool_t multiple);
```

This function generates a graph according to the Watts-Strogatz model of small-world networks. The graph is obtained by creating a circular undirected lattice and then rewire the edges randomly with a constant probability.

See also: Duncan J Watts and Steven H Strogatz: Collective dynamics of “small world” networks, Nature 393, 440-442, 1998.

Arguments:

graph: The graph to initialize.

dim: The dimension of the lattice.

size: The size of the lattice along each dimension.

nei: The size of the neighborhood for each vertex. This is the same as the *nei* argument of `igraph_connect_neighborhood()`.

p: The rewiring probability. A real number between zero and one (inclusive).

loops: Logical, whether to generate loop edges.

multiple: Logical, whether to allow multiple edges in the generated graph.

Returns:

Error code.

See also:

`igraph_lattice()`, `igraph_connect_neighborhood()` and `igraph_rewire_edges()` can be used if more flexibility is needed, e.g. a different type of lattice.

Time complexity: $O(|V|*d^o+|E|)$, $|V|$ and $|E|$ are the number of vertices and edges, d is the average degree, o is the *nei* argument.

igraph_rewire_edges — Rewires the edges of a graph with constant probability.

```
int igraph_rewire_edges(igraph_t *graph, igraph_real_t prob,
                        igraph_bool_t loops, igraph_bool_t multiple);
```

This function rewires the edges of a graph with a constant probability. More precisely each end point of each edge is rewired to a uniformly randomly chosen vertex with constant probability *prob*.

Note that this function modifies the input *graph*, call `igraph_copy()` if you want to keep it.

Arguments:

graph: The input graph, this will be rewired, it can be directed or undirected.

prob: The rewiring probability a constant between zero and one (inclusive).

loops: Boolean, whether loop edges are allowed in the new graph, or not.

multiple: Boolean, whether multiple edges are allowed in the new graph.

Returns:

Error code.

See also:

`igraph_watts_strogatz_game()` uses this function for the rewiring.

Time complexity: $O(|V|+|E|)$.

`igraph_rewire_directed_edges` — Rewires the chosen endpoint of directed edges.

```
int igraph_rewire_directed_edges(igraph_t *graph, igraph_real_t prob,
                                igraph_bool_t loops, igraph_neimode_t mode);
```

This function rewires either the start or end of directed edges in a graph with a constant probability. Correspondingly, either the in-degree sequence or the out-degree sequence of the graph will be preserved.

Note that this function modifies the input *graph*, call `igraph_copy()` if you want to keep it.

Arguments:

graph: The input graph, this will be rewired, it can be directed or undirected. If it is directed, `igraph_rewire_edges()` will be called.

prob: The rewiring probability, a constant between zero and one (inclusive).

loops: Boolean, whether loop edges are allowed in the new graph, or not.

mode: The endpoints of directed edges to rewire. It is ignored for undirected graphs. Possible values:

`IGRAPH_OUT` rewire the end of each directed edge

`IGRAPH_IN` rewire the start of each directed edge

`IGRAPH_ALL` rewire both endpoints of each edge

Returns:

Error code.

See also:

`igraph_rewire_edges()`, `igraph_rewire()`

Time complexity: $O(|E|)$.

igraph_degree_sequence_game — Generates a random graph with a given degree sequence.

```
int igraph_degree_sequence_game(igraph_t *graph, const igraph_vector_t *out_deg,
                                const igraph_vector_t *in_deg,
                                igraph_degseq_t method);
```

Arguments:

graph: Pointer to an uninitialized graph object.

out_deg: The degree sequence for an undirected graph (if *in_deg* is NULL or of length zero), or the out-degree sequence of a directed graph (if *in_deg* is not of length zero).

in_deg: It is either a zero-length vector or NULL (if an undirected graph is generated), or the in-degree sequence.

method: The method to generate the graph. Possible values:

IGRAPH_DEGSEQ_SIMPLE

This method implements the configuration model. For undirected graphs, it puts all vertex IDs in a bag such that the multiplicity of a vertex in the bag is the same as its degree. Then it draws pairs from the bag until the bag becomes empty. This method may generate both loop (self) edges and multiple edges. For directed graphs, the algorithm is basically the same, but two separate bags are used for the in- and out-degrees. Undirected graphs are generated with probability proportional to $(\prod_{i < j} A_{ij} \prod_i A_{ii}!!)^{-1}$, where A denotes the adjacency matrix and $!!$ denotes the double factorial. Here A is assumed to have twice the number of self-loops on its diagonal. The corresponding expression for directed graphs is $(\prod_{i, j} A_{ij})^{-1}$. Thus the probability of all simple graphs (which only have 0s and 1s in the adjacency matrix) is the same, while that of non-simple ones depends on their edge and self-loop multiplicities.

IGRAPH_DEGSEQ_SIMPLE_NO_MULTIPLE

This method generates simple graphs. It is similar to IGRAPH_DEGSEQ_SIMPLE but tries to avoid multiple and loop edges and restarts the generation from scratch if it gets stuck. It can generate all simple realizations of a degree sequence, but it is not guaranteed to sample them uniformly. This method is relatively fast and it will eventually succeed if the provided de-

	gree sequence is graphical, but there is no upper bound on the number of iterations.
IGRAPH_DEGSEQ_SIMPLE_NO_MULTIPLE_UNIFORM	This method is identical to IGRAPH_DEGSEQ_SIMPLE, but if the generated graph is not simple, it rejects it and re-starts the generation. It generates all simple graphs with the same probability.
IGRAPH_DEGSEQ_VL	This method samples undirected <i>connected</i> graphs approximately uniformly. It is a Monte Carlo method based on degree-preserving edge swaps. This generator should be favoured if undirected and connected graphs are to be generated and execution time is not a concern. igraph uses the original implementation of Fabien Viger; for the algorithm, see https://www-complexnetworks.lip6.fr/~latapy/FV/generation.html and the paper https://arxiv.org/abs/cs/0502085

Returns:

Error code: IGRAPH_ENOMEM: there is not enough memory to perform the operation.
IGRAPH_EINVAL: invalid method parameter, or invalid in- and/or out-degree vectors. The degree vectors should be non-negative, *out_deg* should sum up to an even integer for undirected graphs; the length and sum of *out_deg* and *in_deg* should match for directed graphs.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges for IGRAPH_DEGSEQ_SIMPLE. The time complexity of the other modes is not known.

See also:

`igraph_barabasi_game()`, `igraph_erdos_renyi_game()`, `igraph_is_graphical()`

Example	9.15.	File	examples/simple/
igraph_degree_sequence_game.c			

igraph_k_regular_game — Generates a random graph where each vertex has the same degree.

```
int igraph_k_regular_game(igraph_t *graph,
                          igraph_integer_t no_of_nodes, igraph_integer_t k,
                          igraph_bool_t directed, igraph_bool_t multiple);
```

This game generates a directed or undirected random graph where the degrees of vertices are equal to a predefined constant *k*. For undirected graphs, at least one of *k* and the number of vertices must be even.

Currently, this game simply uses `igraph_degree_sequence_game` with the `SIMPLE_NO_MULTIPLE` method and appropriately constructed degree sequences. Therefore, it does not sample uniformly: while it can generate all *k*-regular graphs with the given number of vertices, it does not generate each one with the same probability.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.
<i>no_of_nodes</i> :	The number of nodes in the generated graph.
<i>k</i> :	The degree of each vertex in an undirected graph, or the out-degree and in-degree of each vertex in a directed graph.
<i>directed</i> :	Whether the generated graph will be directed.
<i>multiple</i> :	Whether to allow multiple edges in the generated graph.

Returns:

Error code: IGRAPH_EINVAL: invalid parameter; e.g., negative number of nodes, or odd number of nodes and odd *k* for undirected graphs. IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: $O(|V|+|E|)$ if *multiple* is true, otherwise not known.

igraph_static_fitness_game — Non-growing random graph with edge probabilities proportional to node fitness scores.

```
int igraph_static_fitness_game(igraph_t *graph, igraph_integer_t no_of_edges,
                               const igraph_vector_t *fitness_out, const igraph_t *graph_in,
                               igraph_bool_t loops, igraph_bool_t multiple);
```

This game generates a directed or undirected random graph where the probability of an edge between vertices *i* and *j* depends on the fitness scores of the two vertices involved. For undirected graphs, each vertex has a single fitness score. For directed graphs, each vertex has an out- and an in-fitness, and the probability of an edge from *i* to *j* depends on the out-fitness of vertex *i* and the in-fitness of vertex *j*.

The generation process goes as follows. We start from *N* disconnected nodes (where *N* is given by the length of the fitness vector). Then we randomly select two vertices *i* and *j*, with probabilities proportional to their fitnesses. (When the generated graph is directed, *i* is selected according to the out-fitnesses and *j* is selected according to the in-fitnesses). If the vertices are not connected yet (or if multiple edges are allowed), we connect them; otherwise we select a new pair. This is repeated until the desired number of links are created.

It can be shown that the *expected* degree of each vertex will be proportional to its fitness, although the actual, observed degree will not be. If you need to generate a graph with an exact degree sequence, consider `igraph_degree_sequence_game` instead.

This model is commonly used to generate static scale-free networks. To achieve this, you have to draw the fitness scores from the desired power-law distribution. Alternatively, you may use `igraph_static_power_law_game` which generates the fitnesses for you with a given exponent.

Reference: Goh K-I, Kahng B, Kim D: Universal behaviour of load distribution in scale-free networks. Phys Rev Lett 87(27):278701, 2001.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.
<i>fitness_out</i> :	A numeric vector containing the fitness of each vertex. For directed graphs, this specifies the out-fitness of each vertex.

fitness_in: If NULL, the generated graph will be undirected. If not NULL, this argument specifies the in-fitness of each vertex.

no_of_edges: The number of edges in the generated graph.

loops: Whether to allow loop edges in the generated graph.

multiple: Whether to allow multiple edges in the generated graph.

Returns:

Error code: IGRAPH_EINVAL: invalid parameter IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: $O(|V| + |E| \log |E|)$.

igraph_static_power_law_game — Generates a non-growing random graph with expected power-law degree distributions.

```
int igraph_static_power_law_game(igraph_t *graph,
                                igraph_integer_t no_of_nodes, igraph_integer_t
                                igraph_real_t exponent_out, igraph_real_t expo
                                igraph_bool_t loops, igraph_bool_t multiple,
                                igraph_bool_t finite_size_correction);
```

This game generates a directed or undirected random graph where the degrees of vertices follow power-law distributions with prescribed exponents. For directed graphs, the exponents of the in- and out-degree distributions may be specified separately.

The game simply uses `igraph_static_fitness_game` with appropriately constructed fitness vectors. In particular, the fitness of vertex i is $i^{-\alpha}$, where $\alpha = 1/(\gamma-1)$ and γ is the exponent given in the arguments.

To remove correlations between in- and out-degrees in case of directed graphs, the in-fitness vector will be shuffled after it has been set up and before `igraph_static_fitness_game` is called.

Note that significant finite size effects may be observed for exponents smaller than 3 in the original formulation of the game. This function provides an argument that lets you remove the finite size effects by assuming that the fitness of vertex i is $(i+i_0-1)^{-\alpha}$, where i_0 is a constant chosen appropriately to ensure that the maximum degree is less than the square root of the number of edges times the average degree; see the paper of Chung and Lu, and Cho et al for more details.

References:

Goh K-I, Kahng B, Kim D: Universal behaviour of load distribution in scale-free networks. Phys Rev Lett 87(27):278701, 2001.

Chung F and Lu L: Connected components in a random graph with given degree sequences. Annals of Combinatorics 6, 125-145, 2002.

Cho YS, Kim JS, Park J, Kahng B, Kim D: Percolation transitions in scale-free networks under the Achlioptas process. Phys Rev Lett 103:135702, 2009.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.
<i>no_of_nodes</i> :	The number of nodes in the generated graph.
<i>no_of_edges</i> :	The number of edges in the generated graph.
<i>exponent_out</i> :	The power law exponent of the degree distribution. For directed graphs, this specifies the exponent of the out-degree distribution. It must be greater than or equal to 2. If you pass <code>IGRAPH_INFINITY</code> here, you will get back an Erdos-Renyi random network.
<i>exponent_in</i> :	If negative, the generated graph will be undirected. If greater than or equal to 2, this argument specifies the exponent of the in-degree distribution. If non-negative but less than 2, an error will be generated.
<i>loops</i> :	Whether to allow loop edges in the generated graph.
<i>multiple</i> :	Whether to allow multiple edges in the generated graph.
<i>finite_size_correction</i> :	Whether to use the proposed finite size correction of Cho et al.

Returns:

Error code: `IGRAPH_EINVAL`: invalid parameter `IGRAPH_ENOMEM`: there is not enough memory for the operation.

Time complexity: $O(|V| + |E| \log |E|)$.

`igraph_forest_fire_game` — Generates a network according to the “forest fire game”.

```
int igraph_forest_fire_game(igraph_t *graph, igraph_integer_t nodes,
                           igraph_real_t fw_prob, igraph_real_t bw_factor,
                           igraph_integer_t pambs, igraph_bool_t directed);
```

The forest fire model intends to reproduce the following network characteristics, observed in real networks:

- Heavy-tailed in-degree distribution.
- Heavy-tailed out-degree distribution.
- Communities.
- Densification power-law. The network is densifying in time, according to a power-law rule.
- Shrinking diameter. The diameter of the network decreases in time.

The network is generated in the following way. One vertex is added at a time. This vertex connects to (*cites*) *ambs* vertices already present in the network, chosen uniformly random. Now, for each cited vertex *v* we do the following procedure:

1. We generate two random numbers, *x* and *y*, that are geometrically distributed with means $p/(1-p)$ and $rp/(1-rp)$. (*p* is *fw_prob*, *r* is *bw_factor*.) The new vertex cites *x* outgoing

neighbors and y incoming neighbors of v , from those which are not yet cited by the new vertex. If there are less than x or y such vertices available then we cite all of them.

2. The same procedure is applied to all the newly cited vertices.

See also: Jure Leskovec, Jon Kleinberg and Christos Faloutsos. Graphs over time: densification laws, shrinking diameters and possible explanations. *KDD '05: Proceeding of the eleventh ACM SIGKDD international conference on Knowledge discovery in data mining*, 177--187, 2005.

Note however, that the version of the model in the published paper is incorrect in the sense that it cannot generate the kind of graphs the authors claim. A corrected version is available from <http://cs.stanford.edu/people/jure/pubs/powergrowth-tkdd.pdf>, our implementation is based on this.

Arguments:

graph: Pointer to an uninitialized graph object.

nodes: The number of vertices in the graph.

fw_prob: The forward burning probability.

bw_factor: The backward burning ratio. The backward burning probability is calculated as `bw.factor*fw.prob`.

pambs: The number of ambassador vertices.

directed: Whether to create a directed graph.

Returns:

Error code.

Time complexity: TODO.

igraph_rewire — Randomly rewires a graph while preserving the degree distribution.

```
int igraph_rewire(igraph_t *graph, igraph_integer_t n, igraph_rewiring_t mode);
```

This function generates a new graph based on the original one by randomly rewiring edges while preserving the original graph's degree distribution. Please note that the rewiring is done "in place", so no new graph will be allocated. If you would like to keep the original graph intact, use `igraph_copy()` beforehand.

Arguments:

graph: The graph object to be rewired.

n: Number of rewiring trials to perform.

mode: The rewiring algorithm to be used. It can be one of the following flags:

<code>IGRAPH_REWIRING_SIMPLE</code>	Simple rewiring algorithm which chooses two arbitrary edges in each step (namely (a,b) and (c,d)) and substitutes them with (a,d) and (c,b) if they don't exist. The method will neither destroy nor create self-loops.
-------------------------------------	---

IGRAPH_REWIRING_SIMPLE_LOOPS

Same as IGRAPH_REWIRING_SIMPLE but allows the creation or destruction of self-loops.

Returns:

Error code:

IGRAPH_EINVMODE Invalid rewiring mode.

IGRAPH_EINVAL Graph unsuitable for rewiring (e.g. it has less than 4 nodes in case of IGRAPH_REWIRING_SIMPLE)

IGRAPH_ENOMEM Not enough memory for temporary data.

Time complexity: TODO.

igraph_growing_random_game — Generates a growing random graph.

```
int igraph_growing_random_game(igraph_t *graph, igraph_integer_t n,
                               igraph_integer_t m, igraph_bool_t directed,
                               igraph_bool_t citation);
```

This function simulates a growing random graph. We start out with one vertex. In each step a new vertex is added and a number of new edges are also added. These graphs are known to be different from standard (not growing) random graphs.

Arguments:

graph: Uninitialized graph object.

n: The number of vertices in the graph.

m: The number of edges to add in a time step (i.e. after adding a vertex).

directed: Boolean, whether to generate a directed graph.

citation: Boolean, if TRUE, the edges always originate from the most recently added vertex and are connected to a previous vertex.

Returns:

Error code: IGRAPH_EINVAL: invalid *n* or *m* parameter.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

igraph_callaway_traits_game — Simulates a growing network with vertex types.

```
int igraph_callaway_traits_game(igraph_t *graph, igraph_integer_t nodes,
                                igraph_integer_t types, igraph_integer_t edges_
```

```
const igraph_vector_t *type_dist,  
const igraph_matrix_t *pref_matrix,  
igraph_bool_t directed,  
igraph_vector_t *node_type_vec);
```

The different types of vertices prefer to connect other types of vertices with a given probability.

The simulation goes like this: in each discrete time step a new vertex is added to the graph. The type of this vertex is generated based on *type_dist*. Then two vertices are selected uniformly randomly from the graph. The probability that they will be connected depends on the types of these vertices and is taken from *pref_matrix*. Then another two vertices are selected and this is repeated *edges_per_step* times in each time step.

References:

D. S. Callaway, J. E. Hopcroft, J. M. Kleinberg, M. E. J. Newman, and S. H. Strogatz, Are randomly grown graphs really random? Phys. Rev. E 64, 041902 (2001). <https://doi.org/10.1103/PhysRevE.64.041902>

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph.
<i>nodes</i> :	The number of nodes in the graph.
<i>types</i> :	Number of node types.
<i>edges_per_step</i> :	The number of connections tried in each time step.
<i>type_dist</i> :	Vector giving the distribution of the vertex types. If NULL, the distribution is assumed to be uniform.
<i>pref_matrix</i> :	Matrix giving the connection probabilities for the vertex types.
<i>directed</i> :	Logical, whether to generate a directed graph.
<i>node_type_vec</i> :	An initialized vector or NULL. If not NULL, the type of each node will be stored here.

Returns:

Error code.

Added in version 0.2.

Time complexity: $O(|V|*k*\log(|V|))$, $|V|$ is the number of vertices, k is *edges_per_step*.

igraph_establishment_game — Generates a graph with a simple growing model with vertex types.

```
int igraph_establishment_game(igraph_t *graph, igraph_integer_t nodes,  
                              igraph_integer_t types, igraph_integer_t k,  
                              const igraph_vector_t *type_dist,  
                              const igraph_matrix_t *pref_matrix,  
                              igraph_bool_t directed,
```

```
igraph_vector_t *node_type_vec);
```

The simulation goes like this: a single vertex is added at each time step. This new vertex tries to connect to k vertices in the graph. The probability that such a connection is realized depends on the types of the vertices involved.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph.
<i>nodes</i> :	The number of vertices in the graph.
<i>types</i> :	The number of vertex types.
<i>k</i> :	The number of connections tried in each time step.
<i>type_dist</i> :	Vector giving the distribution of vertex types. If NULL, the distribution is assumed to be uniform.
<i>pref_matrix</i> :	Matrix giving the connection probabilities for different vertex types.
<i>directed</i> :	Logical, whether to generate a directed graph.
<i>node_type_vec</i> :	An initialized vector or NULL. If not NULL, the type of each node will be stored here.

Returns:

Error code.

Added in version 0.2.

Time complexity: $O(|V|*k*\log(|V|))$, $|V|$ is the number of vertices and k is the k parameter.

igraph_preference_game — Generates a graph with vertex types and connection preferences.

```
int igraph_preference_game(igraph_t *graph, igraph_integer_t nodes,
                           igraph_integer_t types,
                           const igraph_vector_t *type_dist,
                           igraph_bool_t fixed_sizes,
                           const igraph_matrix_t *pref_matrix,
                           igraph_vector_t *node_type_vec,
                           igraph_bool_t directed,
                           igraph_bool_t loops);
```

This is practically the nongrowing variant of `igraph_establishment_game()`. A given number of vertices are generated. Every vertex is assigned to a vertex type according to the given type probabilities. Finally, every vertex pair is evaluated and an edge is created between them with a probability depending on the types of the vertices involved.

In other words, this function generates a graph according to a block-model. Vertices are divided into groups (or blocks), and the probability the two vertices are connected depends on their groups only.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph.
<i>nodes</i> :	The number of vertices in the graph.
<i>types</i> :	The number of vertex types.
<i>type_dist</i> :	Vector giving the distribution of vertex types. If NULL, all vertex types will have equal probability. See also the <i>fixed_sizes</i> argument.
<i>fixed_sizes</i> :	Boolean. If true, then the number of vertices with a given vertex type is fixed and the <i>type_dist</i> argument gives these numbers for each vertex type. If true, and <i>type_dist</i> is NULL, then the function tries to make vertex groups of the same size. If this is not possible, then some groups will have an extra vertex.
<i>pref_matrix</i> :	Matrix giving the connection probabilities for different vertex types. This should be symmetric if the requested graph is undirected.
<i>node_type_vec</i> :	A vector where the individual generated vertex types will be stored. If NULL, the vertex types won't be saved.
<i>directed</i> :	Logical, whether to generate a directed graph. If undirected graphs are requested, only the lower left triangle of the preference matrix is considered.
<i>loops</i> :	Logical, whether loop edges are allowed.

Returns:

Error code.

Added in version 0.3.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph.

See also:

```
igraph_asymmetric_preference_game(),    igraph_establishment_game(),  
igraph_callaway_traits_game()
```

igraph_asymmetric_preference_game — Generates a graph with asymmetric vertex types and connection preferences.

```
int igraph_asymmetric_preference_game(igraph_t *graph, igraph_integer_t nodes,  
                                     igraph_integer_t out_types,  
                                     igraph_integer_t in_types,  
                                     const igraph_matrix_t *type_dist_matrix,  
                                     const igraph_matrix_t *pref_matrix,  
                                     igraph_vector_t *node_type_out_vec,  
                                     igraph_vector_t *node_type_in_vec,  
                                     igraph_bool_t loops);
```

This is the asymmetric variant of `igraph_preference_game()`. A given number of vertices are generated. Every vertex is assigned to an "outgoing" and an "incoming" vertex type according

to the given joint type probabilities. Finally, every vertex pair is evaluated and a directed edge is created between them with a probability depending on the "outgoing" type of the source vertex and the "incoming" type of the target vertex.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph.
<i>nodes</i> :	The number of vertices in the graph.
<i>out_types</i> :	The number of vertex out-types.
<i>in_types</i> :	The number of vertex in-types.
<i>type_dist_matrix</i> :	Matrix of size <i>out_types</i> * <i>in_types</i> , giving the joint distribution of vertex types. If NULL, incoming and outgoing vertex types are independent and uniformly distributed.
<i>pref_matrix</i> :	Matrix of size <i>out_types</i> * <i>in_types</i> , giving the connection probabilities for different vertex types.
<i>node_type_out_vec</i> :	A vector where the individual generated "outgoing" vertex types will be stored. If NULL, the vertex types won't be saved.
<i>node_type_in_vec</i> :	A vector where the individual generated "incoming" vertex types will be stored. If NULL, the vertex types won't be saved.
<i>loops</i> :	Logical, whether loop edges are allowed.

Returns:

Error code.

Added in version 0.3.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph.

See also:

`igraph_preference_game()`

igraph_recent_degree_game — Stochastic graph generator based on the number of incident edges a node has gained recently.

```
int igraph_recent_degree_game(igraph_t *graph, igraph_integer_t nodes,
                              igraph_real_t power,
                              igraph_integer_t time_window,
                              igraph_integer_t m,
                              const igraph_vector_t *outseq,
                              igraph_bool_t outpref,
                              igraph_real_t zero_appeal,
                              igraph_bool_t directed);
```

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.
<i>nodes</i> :	The number of vertices in the graph, this is the same as the number of time steps.
<i>power</i> :	The exponent, the probability that a node gains a new edge is proportional to the number of edges it has gained recently (in the last <i>window</i> time steps) to <i>power</i> .
<i>time_window</i> :	Integer constant, the size of the time window to use to count the number of recent edges.
<i>m</i> :	Integer constant, the number of edges to add per time step if the <i>outseq</i> parameter is a null pointer or a zero-length vector.
<i>outseq</i> :	The number of edges to add in each time step. This argument is ignored if it is a null pointer or a zero length vector. In this case the constant <i>m</i> parameter is used.
<i>outpref</i> :	Logical constant, if true the edges originated by a vertex also count as recent incident edges. For most applications it is reasonable to set it to false.
<i>zero_appeal</i> :	Constant giving the attractiveness of the vertices which haven't gained any edge recently.
<i>directed</i> :	Logical constant, whether to generate a directed graph.

Returns:

Error code.

Time complexity: $O(|V| \cdot \log(|V|) + |E|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges in the graph.

igraph_barabasi_aging_game — Preferential attachment with aging of vertices.

```
int igraph_barabasi_aging_game(igraph_t *graph,
                               igraph_integer_t nodes,
                               igraph_integer_t m,
                               const igraph_vector_t *outseq,
                               igraph_bool_t outpref,
                               igraph_real_t pa_exp,
                               igraph_real_t aging_exp,
                               igraph_integer_t aging_bins,
                               igraph_real_t zero_deg_appeal,
                               igraph_real_t zero_age_appeal,
                               igraph_real_t deg_coef,
                               igraph_real_t age_coef,
                               igraph_bool_t directed);
```

This game starts with one vertex (if *nodes* > 0). In each step a new node is added, and it is connected to *m* existing nodes. Existing nodes to connect to are chosen with probability dependent on their (in-)degree (*k*) and age (*l*). The degree-dependent part is *deg_coef* * *k*^{*pa_exp*} + *zero_deg_appeal*, while the age-dependent part is *age_coef* * *l*^{*aging_exp*} + *zero_age_appeal*, which are summed to obtain the final weight.

The age *l* is based on the number of vertices in the network and the *aging_bins* argument: the age of a node is incremented by 1 after each `floor(nodes / aging_bins) + 1` time steps.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.
<i>nodes</i> :	The number of vertices in the graph.
<i>m</i> :	The number of edges to add in each time step. Ignored if <i>outseq</i> is a non-zero length vector.
<i>outseq</i> :	The number of edges to add in each time step. If it is NULL or a zero-length vector then it is ignored and the <i>m</i> argument is used instead.
<i>outpref</i> :	Logical constant, whether the edges initiated by a vertex contribute to the probability to gain a new edge.
<i>pa_exp</i> :	The exponent of the preferential attachment, a small positive number usually, the value 1 yields the classic linear preferential attachment.
<i>aging_exp</i> :	The exponent of the aging, this is a negative number usually.
<i>aging_bins</i> :	Integer constant, the number of age bins to use.
<i>zero_deg_appeal</i> :	The degree dependent part of the attractiveness of the zero degree vertices.
<i>zero_age_appeal</i> :	The age dependent part of the attractiveness of the vertices of age zero. This parameter is usually zero.
<i>deg_coef</i> :	The coefficient for the degree.
<i>age_coef</i> :	The coefficient for the age.
<i>directed</i> :	Logical constant, whether to generate a directed graph.

Returns:

Error code.

Time complexity: $O((|V|+|V|/\text{aging_bins})*\log(|V|+|E|))$. $|V|$ is the number of vertices, $|E|$ the number of edges.

igraph_recent_degree_aging_game — Preferential attachment based on the number of edges gained recently, with aging of vertices.

```
int igraph_recent_degree_aging_game(igraph_t *graph,
                                     igraph_integer_t nodes,
                                     igraph_integer_t m,
                                     const igraph_vector_t *outseq,
                                     igraph_bool_t outpref,
                                     igraph_real_t pa_exp,
                                     igraph_real_t aging_exp,
                                     igraph_integer_t aging_bins,
                                     igraph_integer_t time_window,
                                     igraph_real_t zero_appeal,
                                     igraph_bool_t directed);
```


This game is very similar to `igraph_barabasi_aging_game()`, except that instead of the total number of incident edges the number of edges gained in the last `time_window` time steps are counted.

The degree dependent part of the attractiveness is given by k to the power of `pa_exp` plus `zero_appeal`; the age dependent part is l to the power to `aging_exp`. k is the number of edges gained in the last `time_window` time steps, l is the age of the vertex.

Arguments:

<code>graph</code> :	Pointer to an uninitialized graph object.
<code>nodes</code> :	The number of vertices in the graph.
<code>m</code> :	The number of edges to add in each time step. If the <code>outseq</code> argument is not a null vector or a zero-length vector then it is ignored.
<code>outseq</code> :	Vector giving the number of edges to add in each time step. If it is a null pointer or a zero-length vector then it is ignored and the <code>m</code> argument is used.
<code>outpref</code> :	Logical constant, if true the edges initiated by a vertex are also counted. Normally it is false.
<code>pa_exp</code> :	The exponent for the preferential attachment.
<code>aging_exp</code> :	The exponent for the aging, normally it is negative: old vertices gain edges with less probability.
<code>aging_bins</code> :	Integer constant, the number of age bins to use.
<code>time_window</code> :	The time window to use to count the number of incident edges for the vertices.
<code>zero_appeal</code> :	The degree dependent part of the attractiveness for zero degree vertices.
<code>directed</code> :	Logical constant, whether to create a directed graph.

Returns:

Error code.

Time complexity: $O((|V|+|V|/\text{aging_bins})\log(|V|+|E|))$. $|V|$ is the number of vertices, $|E|$ the number of edges.

`igraph_lastcit_game` — Simulates a citation network, based on time passed since the last citation.

```
int igraph_lastcit_game(igraph_t *graph,
                        igraph_integer_t nodes, igraph_integer_t edges_per_node,
                        igraph_integer_t agebins,
                        const igraph_vector_t *preference,
                        igraph_bool_t directed);
```

This is a quite special stochastic graph generator, it models an evolving graph. In each time step a single vertex is added to the network and it cites a number of other vertices (as specified by the `edges_per_step` argument). The cited vertices are selected based on the last time they were cited. Time is measured by the addition of vertices and it is binned into `agebins` bins. So if the current time step is t and the last citation to a given i vertex was made in time step t_0 , then $\lfloor (t-t_0)/\text{binwidth} \rfloor$

is calculated where $\text{binwidth} = \text{nodes}/\text{agebins}+1$, in the last expression $'/'$ denotes integer division, so the fraction part is omitted.

The *preference* argument specifies the preferences for the citation lags, i.e. its first elements contains the attractivity of the very recently cited vertices, etc. The last element is special, it contains the attractivity of the vertices which were never cited. This element should be bigger than zero.

Note that this function generates networks with multiple edges if *edges_per_step* is bigger than one, call `igraph_simplify()` on the result to get rid of these edges.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object, the result will be stored here.
<i>node</i> :	The number of vertices in the network.
<i>edges_per_node</i> :	The number of edges to add in each time step.
<i>agebins</i> :	The number of age bins to use.
<i>preference</i> :	Pointer to an initialized vector of length $\text{agebins}+1$. This contains the 'attractivity' of the various age bins, the last element is the attractivity of the vertices which were never cited, and it should be greater than zero. It is a good idea to have all positive values in this vector. Preferences cannot be negative.
<i>directed</i> :	Logical constant, whether to create directed networks.

Returns:

Error code.

See also:

`igraph_barabasi_aging_game()`.

Time complexity: $O(|V|*a+|E|*\log|V|)$, $|V|$ is the number of vertices, $|E|$ is the total number of edges, a is the *agebins* parameter.

igraph_cited_type_game — Simulates a citation based on vertex types.

```
int igraph_cited_type_game(igraph_t *graph, igraph_integer_t nodes,
                           const igraph_vector_t *types,
                           const igraph_vector_t *pref,
                           igraph_integer_t edges_per_step,
                           igraph_bool_t directed);
```

Function to create a network based on some vertex categories. This function creates a citation network: in each step a single vertex and *edges_per_step* citing edges are added. Nodes with different categories may have different probabilities to get cited, as given by the *pref* vector.

Note that this function might generate networks with multiple edges if *edges_per_step* is greater than one. You might want to call `igraph_simplify()` on the result to remove multiple edges.

Arguments:

graph: Pointer to an uninitialized graph object.

<i>nodes</i> :	The number of vertices in the network.
<i>types</i> :	Numeric vector giving the categories of the vertices, so it should contain <i>nodes</i> non-negative integer numbers. Types are numbered from zero.
<i>pref</i> :	The attractivity of the different vertex categories in a vector. Its length should be the maximum element in <i>types</i> plus one (types are numbered from zero).
<i>edges_per_step</i> :	Integer constant, the number of edges to add in each time step.
<i>directed</i> :	Logical constant, whether to create a directed network.

Returns:

Error code.

See also:

`igraph_citing_cited_type_game()` for a bit more general game.

Time complexity: $O((|V|+|E|)\log|V|)$, $|V|$ and $|E|$ are number of vertices and edges, respectively.

igraph_citing_cited_type_game — Simulates a citation network based on vertex types.

```
int igraph_citing_cited_type_game(igraph_t *graph, igraph_integer_t nodes,
                                   const igraph_vector_t *types,
                                   const igraph_matrix_t *pref,
                                   igraph_integer_t edges_per_step,
                                   igraph_bool_t directed);
```

This game is similar to `igraph_cited_type_game()` but here the category of the citing vertex is also considered.

An evolving citation network is modeled here, a single vertex and its *edges_per_step* citation are added in each time step. The odds the a given vertex is cited by the new vertex depends on the category of both the citing and the cited vertex and is given in the *pref* matrix. The categories of the citing vertex correspond to the rows, the categories of the cited vertex to the columns of this matrix. I.e. the element in row *i* and column *j* gives the probability that a *j* vertex is cited, if the category of the citing vertex is *i*.

Note that this function might generate networks with multiple edges if *edges_per_step* is greater than one. You might want to call `igraph_simplify()` on the result to remove multiple edges.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.
<i>nodes</i> :	The number of vertices in the network.
<i>types</i> :	A numeric matrix of length <i>nodes</i> , containing the categories of the vertices. The categories are numbered from zero.
<i>pref</i> :	The preference matrix, a square matrix is required, both the number of rows and columns should be the maximum element in <i>types</i> plus one (types are numbered from zero).

directed: Logical constant, whether to create a directed network.

Returns:

Error code.

Time complexity: $O((|V|+|E|)\log|V|)$, $|V|$ and $|E|$ are number of vertices and edges, respectively.

igraph_sbm_game — Sample from a stochastic block model.

```
int igraph_sbm_game(igraph_t *graph, igraph_integer_t n,
                    const igraph_matrix_t *pref_matrix,
                    const igraph_vector_int_t *block_sizes,
                    igraph_bool_t directed, igraph_bool_t loops);
```

This function samples graphs from a stochastic block model by (doing the equivalent of) Bernoulli trials for each potential edge with the probabilities given by the Bernoulli rate matrix, *pref_matrix*. See Faust, K., & Wasserman, S. (1992a). Blockmodels: Interpretation and evaluation. *Social Networks*, 14, 5–61.

The order of the vertex ids in the generated graph corresponds to the *block_sizes* argument.

Arguments:

graph: The output graph. This should be a pointer to an uninitialized graph.

n: Number of vertices.

pref_matrix: The matrix giving the Bernoulli rates. This is a $K \times K$ matrix, where K is the number of groups. The probability of creating an edge between vertices from groups i and j is given by element (i,j) .

block_sizes: An integer vector giving the number of vertices in each group.

directed: Boolean, whether to create a directed graph. If this argument is false, then *pref_matrix* must be symmetric.

loops: Boolean, whether to create self-loops.

Returns:

Error code.

Time complexity: $O(|V|+|E|+K^2)$, where $|V|$ is the number of vertices, $|E|$ is the number of edges, and K is the number of groups.

See also:

`igraph_erdos_renyi_game()` for a simple Bernoulli graph.

igraph_hsbm_game — Hierarchical stochastic block model.

```
int igraph_hsbm_game(igraph_t *graph, igraph_integer_t n,
                    igraph_integer_t m, const igraph_vector_t *rho,
                    const igraph_matrix_t *C, igraph_real_t p);
```

The function generates a random graph according to the hierarchical stochastic block model.

Arguments:

- graph*: The generated graph is stored here.
- n*: The number of vertices in the graph.
- m*: The number of vertices per block. n/m must be integer.
- rho*: The fraction of vertices per cluster, within a block. Must sum up to 1, and $\rho * m$ must be integer for all elements of ρ .
- C*: A square, symmetric numeric matrix, the Bernoulli rates for the clusters within a block. Its size must match the size of the `{rho}` vector.
- p*: The Bernoulli rate of connections between vertices in different blocks.

Returns:

Error code.

See also:

`igraph_sbm_game()` for the classic stochastic block model, `igraph_hsbm_list_game()` for a more general version.

`igraph_hsbm_list_game` — Hierarchical stochastic block model, more general version.

```
int igraph_hsbm_list_game(igraph_t *graph, igraph_integer_t n,
                        const igraph_vector_int_t *mlist,
                        const igraph_vector_ptr_t *rholist,
                        const igraph_vector_ptr_t *Clist,
                        igraph_real_t p);
```

The function generates a random graph according to the hierarchical stochastic block model.

Arguments:

- graph*: The generated graph is stored here.
- n*: The number of vertices in the graph.
- mlist*: An integer vector of block sizes.
- rholist*: A list of ρ vectors (`igraph_vector_t` objects), one for each block.
- Clist*: A list of square matrices (`igraph_matrix_t` objects), one for each block, giving the Bernoulli rates of connections within the block.

p: The Bernoulli rate of connections between vertices in different blocks.

Returns:

Error code.

See also:

`igraph_sbm_game()` for the classic stochastic block model, `igraph_hsbm_game()` for a simpler general version.

`igraph_dot_product_game` — Generates a random dot product graph.

```
int igraph_dot_product_game(igraph_t *graph, const igraph_matrix_t *vecs,
                           igraph_bool_t directed);
```

In this model, each vertex is represented by a latent position vector. Probability of an edge between two vertices are given by the dot product of their latent position vectors.

See also Christine Leigh Myers Nickel: Random dot product graphs, a model for social networks. Dissertation, Johns Hopkins University, Maryland, USA, 2006.

Arguments:

graph: The output graph is stored here.

vecs: A matrix in which each latent position vector is a column. The dot product of the latent position vectors should be in the [0,1] interval, otherwise a warning is given. For negative dot products, no edges are added; dot products that are larger than one always add an edge.

directed: Should the generated graph be directed?

Returns:

Error code.

Time complexity: $O(n*n*m)$, where n is the number of vertices, and m is the length of the latent vectors.

See also:

`igraph_sample_dirichlet()`, `igraph_sample_sphere_volume()`, `igraph_sample_sphere_surface()` for functions to generate the latent vectors.

`igraph_tree_game` — Generates a random tree with the given number of nodes.

```
int igraph_tree_game(igraph_t *graph, igraph_integer_t n, igraph_bool_t directed);
```

This function samples uniformly from the set of labelled trees, i.e. it generates each labelled tree with the same probability.

Arguments:

<i>graph</i> :	Pointer to an uninitialized graph object.	
<i>n</i> :	The number of nodes in the tree.	
<i>directed</i> :	Whether to create a directed tree. The edges are oriented away from the root.	
<i>method</i> :	The algorithm to use to generate the tree. Possible values:	
	IGRAPH_RANDOM_TREE_PRUFER	This algorithm samples Prüfer sequences uniformly, then converts them to trees. Directed trees are not currently supported.
	IGRAPH_RANDOM_LERW	This algorithm effectively performs a loop-erased random walk on the complete graph to uniformly sample its spanning trees (Wilson's algorithm).

Returns:

Error code: IGRAPH_ENOMEM: there is not enough memory to perform the operation.
IGRAPH_EINVAL: invalid tree size

See also:

`igraph_from_prufer()`

igraph_correlated_game — Generates a random graph correlated to an existing graph.

```
int igraph_correlated_game(const igraph_t *old_graph, igraph_t *new_graph,
                           igraph_real_t corr, igraph_real_t p,
                           const igraph_vector_t *permutation);
```

Sample a new graph by perturbing the adjacency matrix of a given graph and shuffling its vertices.

Arguments:

<i>old_graph</i> :	The original graph.
<i>new_graph</i> :	The new graph will be stored here.
<i>corr</i> :	A scalar in the unit interval, the target Pearson correlation between the adjacency matrices of the original the generated graph (the adjacency matrix being used as a vector).
<i>p</i> :	A numeric scalar, the probability of an edge between two vertices, it must be in the open (0,1) interval.
<i>permutation</i> :	A permutation to apply to the vertices of the generated graph. It can also be a null pointer, in which case the vertices will not be permuted.

Returns:

Error code

See also:

`igraph_correlated_pair_game()` for generating a pair of correlated random graphs in one go.

`igraph_correlated_pair_game` — Generates pairs of correlated random graphs.

```
int igraph_correlated_pair_game(igraph_t *graph1, igraph_t *graph2,
                                igraph_integer_t n, igraph_real_t corr, igraph_
                                igraph_bool_t directed,
                                const igraph_vector_t *permutation);
```

Sample two random graphs, with given correlation.

Arguments:

<i>graph1</i> :	The first graph will be stored here.
<i>graph2</i> :	The second graph will be stored here.
<i>n</i> :	The number of vertices in both graphs.
<i>corr</i> :	A scalar in the unit interval, the target Pearson correlation between the adjacency matrices of the original the generated graph (the adjacency matrix being used as a vector).
<i>p</i> :	A numeric scalar, the probability of an edge between two vertices, it must in the open (0,1) interval.
<i>directed</i> :	Whether to generate directed graphs.
<i>permutation</i> :	A permutation to apply to the vertices of the second graph. It can also be a null pointer, in which case the vertices will not be permuted.

Returns:

Error code

See also:

`igraph_correlated_game()` for generating a correlated pair to a given graph.

`igraph_simple_interconnected_islands_game` — Generates a random graph made of several interconnected islands, each island being a random graph.


```
int igraph_simple_interconnected_islands_game(  
    igraph_t *graph,  
    igraph_integer_t islands_n,  
    igraph_integer_t islands_size,  
    igraph_real_t islands_pin,  
    igraph_integer_t n_inter);
```

Arguments:

graph: Pointer to an uninitialized graph object.

islands_n: The number of islands in the graph.

islands_size: The size of islands in the graph.

islands_pin: The probability to create each possible edge into each island.

n_inter: The number of edges to create between two islands.

Returns:

Error code: IGRAPH_EINVAL: invalid parameter IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph.

Chapter 10. Games on graphs

Microscopic update rules

`igraph_deterministic_optimal_imitation` — Adopt a strategy via deterministic optimal imitation.

```
int igraph_deterministic_optimal_imitation(const igraph_t *graph,
                                           igraph_integer_t vid,
                                           igraph_optimal_t optimality,
                                           const igraph_vector_t *quantities,
                                           igraph_vector_t *strategies,
                                           igraph_neimode_t mode);
```

A simple deterministic imitation strategy where a vertex revises its strategy to that which yields a local optimal. Here "local" is with respect to the immediate neighbours of the vertex. The vertex retains its current strategy where this strategy yields a locally optimal quantity. The quantity in this case could be a measure such as fitness.

Arguments:

- graph*: The graph object representing the game network. This cannot be the empty or trivial graph, but must have at least two vertices and one edge. If *graph* has one vertex, then no strategy update would take place. Furthermore, if *graph* has at least two vertices but zero edges, then strategy update would also not take place.
- vid*: The vertex whose strategy is to be updated. It is assumed that *vid* represents a vertex in *graph*. No checking is performed and it is your responsibility to ensure that *vid* is indeed a vertex of *graph*. If an isolated vertex is provided, i.e. the input vertex has degree 0, then no strategy update would take place and *vid* would retain its current strategy. Strategy update would also not take place if the local neighbourhood of *vid* are its in-neighbours (respectively out-neighbours), but *vid* has zero in-neighbours (respectively out-neighbours). Loops are ignored in computing the degree (in, out, all) of *vid*.
- optimality*: Logical; controls the type of optimality to be used. Supported values are:
- | | |
|-----------------------------|---|
| <code>IGRAPH_MAXIMUM</code> | Use maximum deterministic imitation, where the strategy of the vertex with maximum quantity (e.g. fitness) would be adopted. We update the strategy of <i>vid</i> to that which yields a local maximum. |
| <code>IGRAPH_MINIMUM</code> | Use minimum deterministic imitation. That is, the strategy of the vertex with minimum quantity would be imitated. In other words, update to the strategy that yields a local minimum. |
- quantities*: A vector of quantities providing the quantity of each vertex in *graph*. Think of each entry of the vector as being generated by a function such as the fitness function for the game. So if the vector represents fitness quantities, then each vector entry is the fitness of some vertex. The length of this vector must be the same as the number of vertices in the vertex set of *graph*.
- strategies*: A vector of the current strategies for the vertex population. The updated strategy for *vid* would be stored here. Each strategy is identified with a nonnegative integer,

whose interpretation depends on the payoff matrix of the game. Generally we use the strategy ID as a row or column index of the payoff matrix. The length of this vector must be the same as the number of vertices in the vertex set of *graph*.

mode: Defines the sort of neighbourhood to consider for *vid*. If *graph* is undirected, then we use all the immediate neighbours of *vid*. Thus if you know that *graph* is undirected, then it is safe to pass the value *IGRAPH_ALL* here. Supported values are:

IGRAPH_OUT Use the out-neighbours of *vid*. This option is only relevant when *graph* is a directed graph.

IGRAPH_IN Use the in-neighbours of *vid*. Again this option is only relevant when *graph* is a directed graph.

IGRAPH_ALL Use both the in- and out-neighbours of *vid*. This option is only relevant if *graph* is a digraph. Also use this value if *graph* is undirected.

Returns:

The error code *IGRAPH_EINVAL* is returned in each of the following cases: (1) Any of the parameters *graph*, *quantities*, or *strategies* is a null pointer. (2) The vector *quantities* or *strategies* has a length different from the number of vertices in *graph*. (3) The parameter *graph* is the empty or null graph, i.e. the graph with zero vertices and edges.

Time complexity: $O(2d)$, where d is the degree of the vertex *vid*.

Example 10.1. File **examples/simple/igraph_deterministic_optimal_imitation.c**

igraph_moran_process — The Moran process in a network setting.

```
int igraph_moran_process(const igraph_t *graph,
                        const igraph_vector_t *weights,
                        igraph_vector_t *quantities,
                        igraph_vector_t *strategies,
                        igraph_neimode_t mode);
```

This is an extension of the classic Moran process to a network setting. The Moran process is a model of haploid (asexual) reproduction within a population having a fixed size. In the network setting, the Moran process operates on a weighted graph. At each time step a vertex *a* is chosen for reproduction and another vertex *b* is chosen for death. Vertex *a* gives birth to an identical clone *c*, which replaces *b*. Vertex *c* is a clone of *a* in that *c* inherits both the current quantity (e.g. fitness) and current strategy of *a*.

The graph *G* representing the game network is assumed to be simple, i.e. free of loops and without multiple edges. If, on the other hand, *G* has a loop incident on some vertex *v*, then it is possible that when *v* is chosen for reproduction it would forgo this opportunity. In particular, when *v* is chosen for reproduction and *v* is also chosen for death, the clone of *v* would be *v* itself with its current vertex ID. In effect *v* forgoes its chance for reproduction.

Arguments:

graph: The graph object representing the game network. This cannot be the empty or trivial graph, but must have at least two vertices and one edge. The Moran process will

not take place in each of the following cases: (1) If *graph* has one vertex. (2) If *graph* has at least two vertices but zero edges.

weights: A vector of all edge weights for *graph*. Thus *weights*[*i*] means the weight of the edge with edge ID *i*. For the purpose of the Moran process, each weight is assumed to be positive; it is your responsibility to ensure this condition holds. The length of this vector must be the same as the number of edges in *graph*.

quantities: A vector of quantities providing the quantity of each vertex in *graph*. The quantity of the new clone will be stored here. Think of each entry of the vector as being generated by a function such as the fitness function for the game. So if the vector represents fitness quantities, then each vector entry is the fitness of some vertex. The length of this vector must be the same as the number of vertices in the vertex set of *graph*. For the purpose of the Moran process, each vector entry is assumed to be nonnegative; no checks will be performed for this. It is your responsibility to ensure that at least one entry is positive. Furthermore, this vector cannot be a vector of zeros; this condition will be checked.

strategies: A vector of the current strategies for the vertex population. The strategy of the new clone will be stored here. Each strategy is identified with a nonnegative integer, whose interpretation depends on the payoff matrix of the game. Generally we use the strategy ID as a row or column index of the payoff matrix. The length of this vector must be the same as the number of vertices in the vertex set of *graph*.

mode: Defines the sort of neighbourhood to consider for the vertex a chosen for reproduction. This is only relevant if *graph* is directed. If *graph* is undirected, then it is safe to pass the value *IGRAPH_ALL* here. Supported values are:

IGRAPH_OUT Use the out-neighbours of *a*. This option is only relevant when *graph* is directed.

IGRAPH_IN Use the in-neighbours of *a*. Again this option is only relevant when *graph* is directed.

IGRAPH_ALL Use both the in- and out-neighbours of *a*. This option is only relevant if *graph* is directed. Also use this value if *graph* is undirected.

Returns:

The error code *IGRAPH_EINVAL* is returned in each of the following cases: (1) Any of the parameters *graph*, *weights*, *quantities* or *strategies* is a null pointer. (2) The vector *quantities* or *strategies* has a length different from the number of vertices in *graph*. (3) The vector *weights* has a length different from the number of edges in *graph*. (4) The parameter *graph* is the empty or null graph, i.e. the graph with zero vertices and edges. (5) The vector *weights*, or the combination of interest, sums to zero. (6) The vector *quantities*, or the combination of interest, sums to zero.

Time complexity: depends on the random number generator, but is usually $O(n)$ where *n* is the number of vertices in *graph*.

References:

- (Lieberman et al. 2005) E. Lieberman, C. Hauert, and M. A. Nowak. Evolutionary dynamics on graphs. *Nature*, 433(7023):312--316, 2005.
- (Moran 1958) P. A. P. Moran. Random processes in genetics. *Mathematical Proceedings of the Cambridge Philosophical Society*, 54(1):60--71, 1958.

igraph_roulette_wheel_imitation — Adopt a strategy via roulette wheel selection.

```
int igraph_roulette_wheel_imitation(const igraph_t *graph,
                                   igraph_integer_t vid,
                                   igraph_bool_t islocal,
                                   const igraph_vector_t *quantities,
                                   igraph_vector_t *strategies,
                                   igraph_neimode_t mode);
```

A simple stochastic imitation strategy where a vertex revises its strategy to that of a vertex *u* chosen proportionate to *u*'s quantity (e.g. fitness). This is a special case of stochastic imitation, where a candidate is not chosen uniformly at random but proportionate to its quantity.

Arguments:

- graph*: The graph object representing the game network. This cannot be the empty or trivial graph, but must have at least two vertices and one edge. If *graph* has one vertex, then no strategy update would take place. Furthermore, if *graph* has at least two vertices but zero edges, then strategy update would also not take place.
- vid*: The vertex whose strategy is to be updated. It is assumed that *vid* represents a vertex in *graph*. No checking is performed and it is your responsibility to ensure that *vid* is indeed a vertex of *graph*. If an isolated vertex is provided, i.e. the input vertex has degree 0, then no strategy update would take place and *vid* would retain its current strategy. Strategy update would also not take place if the local neighbourhood of *vid* are its in-neighbours (respectively out-neighbours), but *vid* has zero in-neighbours (respectively out-neighbours). Loops are ignored in computing the degree (in, out, all) of *vid*.
- islocal*: Boolean; this flag controls which perspective to use in computing the relative quantity. If true then we use the local perspective; otherwise we use the global perspective. The local perspective for *vid* is the set of all immediate neighbours of *vid*. In contrast, the global perspective for *vid* is the vertex set of *graph*.
- quantities*: A vector of quantities providing the quantity of each vertex in *graph*. Think of each entry of the vector as being generated by a function such as the fitness function for the game. So if the vector represents fitness quantities, then each vector entry is the fitness of some vertex. The length of this vector must be the same as the number of vertices in the vertex set of *graph*. For the purpose of roulette wheel selection, each vector entry is assumed to be nonnegative; no checks will be performed for this. It is your responsibility to ensure that at least one entry is nonzero. Furthermore, this vector cannot be a vector of zeros; this condition will be checked.
- strategies*: A vector of the current strategies for the vertex population. The updated strategy for *vid* would be stored here. Each strategy is identified with a nonnegative integer, whose interpretation depends on the payoff matrix of the game. Generally we use the strategy ID as a row or column index of the payoff matrix. The length of this vector must be the same as the number of vertices in the vertex set of *graph*.
- mode*: Defines the sort of neighbourhood to consider for *vid*. This is only relevant if we are considering the local perspective, i.e. if *islocal* is true. If we are considering the global perspective, then it is safe to pass the value `IGRAPH_ALL` here. If *graph* is undirected, then we use all the immediate neighbours of *vid*. Thus if you know that *graph* is undirected, then it is safe to pass the value `IGRAPH_ALL` here. Supported values are:

IGRAPH_OUT	Use the out-neighbours of <i>vid</i> . This option is only relevant when <i>graph</i> is a digraph and we are considering the local perspective.
IGRAPH_IN	Use the in-neighbours of <i>vid</i> . Again this option is only relevant when <i>graph</i> is a directed graph and we are considering the local perspective.
IGRAPH_ALL	Use both the in- and out-neighbours of <i>vid</i> . This option is only relevant if <i>graph</i> is a digraph. Also use this value if <i>graph</i> is undirected or we are considering the global perspective.

Returns:

The error code `IGRAPH_EINVAL` is returned in each of the following cases: (1) Any of the parameters *graph*, *quantities*, or *strategies* is a null pointer. (2) The vector *quantities* or *strategies* has a length different from the number of vertices in *graph*. (3) The parameter *graph* is the empty or null graph, i.e. the graph with zero vertices and edges. (4) The vector *quantities* sums to zero.

Time complexity: $O(n)$ where n is the number of vertices in the perspective to consider. If we consider the global perspective, then n is the number of vertices in the vertex set of *graph*. On the other hand, for the local perspective n is the degree of *vid*, excluding loops.

Reference:

(Yu & Gen 2010) X. Yu and M. Gen. *Introduction to Evolutionary Algorithms*. Springer, 2010, pages 18--20.

Example	10.2.	File	examples/simple/
		igraph_roulette_wheel_imitation.c	

igraph_stochastic_imitation — Adopt a strategy via stochastic imitation with uniform selection.

```
int igraph_stochastic_imitation(const igraph_t *graph,
                               igraph_integer_t vid,
                               igraph_imitate_algorithm_t algo,
                               const igraph_vector_t *quantities,
                               igraph_vector_t *strategies,
                               igraph_neimode_t mode);
```

A simple stochastic imitation strategy where a vertex revises its strategy to that of a vertex chosen uniformly at random from its local neighbourhood. This is called stochastic imitation via uniform selection, where the strategy to imitate is chosen via some random process. For the purposes of this function, we use uniform selection from a pool of candidates.

Arguments:

graph: The graph object representing the game network. This cannot be the empty or trivial graph, but must have at least two vertices and one edge. If *graph* has one vertex, then no strategy update would take place. Furthermore, if *graph* has at least two vertices but zero edges, then strategy update would also not take place.

vid: The vertex whose strategy is to be updated. It is assumed that *vid* represents a vertex in *graph*. No checking is performed and it is your responsibility to ensure

that *vid* is indeed a vertex of *graph*. If an isolated vertex is provided, i.e. the input vertex has degree 0, then no strategy update would take place and *vid* would retain its current strategy. Strategy update would also not take place if the local neighbourhood of *vid* are its in-neighbours (respectively out-neighbours), but *vid* has zero in-neighbours (respectively out-neighbours). Loops are ignored in computing the degree (in, out, all) of *vid*.

algo: This flag controls which algorithm to use in stochastic imitation. Supported values are:

IGRAPH_IMITATE_AUGMENTED Augmented imitation. Vertex *vid* imitates the strategy of the chosen vertex *u* provided that doing so would increase the quantity (e.g. fitness) of *vid*. Augmented imitation can be thought of as "imitate if better".

IGRAPH_IMITATE_BLIND Blind imitation. Vertex *vid* blindly imitates the strategy of the chosen vertex *u*, regardless of whether doing so would increase or decrease the quantity of *vid*.

IGRAPH_IMITATE_CONTRACT-
ED Contracted imitation. Here vertex *vid* imitates the strategy of the chosen vertex *u* if doing so would decrease the quantity of *vid*. Think of contracted imitation as "imitate if worse".

quantities: A vector of quantities providing the quantity of each vertex in *graph*. Think of each entry of the vector as being generated by a function such as the fitness function for the game. So if the vector represents fitness quantities, then each vector entry is the fitness of some vertex. The length of this vector must be the same as the number of vertices in the vertex set of *graph*.

strategies: A vector of the current strategies for the vertex population. The updated strategy for *vid* would be stored here. Each strategy is identified with a nonnegative integer, whose interpretation depends on the payoff matrix of the game. Generally we use the strategy ID as a row or column index of the payoff matrix. The length of this vector must be the same as the number of vertices in the vertex set of *graph*.

mode: Defines the sort of neighbourhood to consider for *vid*. If *graph* is undirected, then we use all the immediate neighbours of *vid*. Thus if you know that *graph* is undirected, then it is safe to pass the value IGRAPH_ALL here. Supported values are:

IGRAPH_OUT Use the out-neighbours of *vid*. This option is only relevant when *graph* is a directed graph.

IGRAPH_IN Use the in-neighbours of *vid*. Again this option is only relevant when *graph* is a directed graph.

IGRAPH_ALL Use both the in- and out-neighbours of *vid*. This option is only relevant if *graph* is a digraph. Also use this value if *graph* is undirected.

Returns:

The error code IGRAPH_EINVAL is returned in each of the following cases: (1) Any of the parameters *graph*, *quantities*, or *strategies* is a null pointer. (2) The vector *quantities* or *strategies* has a length different from the number of vertices in *graph*. (3) The parameter

graph is the empty or null graph, i.e. the graph with zero vertices and edges. (4) The parameter *algo* refers to an unsupported stochastic imitation algorithm.

Time complexity: depends on the uniform random number generator, but should usually be $O(1)$.

Example	10.3.	File	examples/simple/
igraph_stochastic_imitation.c			

Epidemic models

igraph_sir — Performs a number of SIR epidemics model runs on a graph.

```
int igraph_sir(const igraph_t *graph, igraph_real_t beta,
               igraph_real_t gamma, igraph_integer_t no_sim,
               igraph_vector_ptr_t *result);
```

The SIR model is a simple model from epidemiology. The individuals of the population might be in three states: susceptible, infected and recovered. Recovered people are assumed to be immune to the disease. Susceptibles become infected with a rate that depends on their number of infected neighbors. Infected people become recovered with a constant rate. See these parameters below.

This function runs multiple simulations, all starting with a single uniformly randomly chosen infected individual. A simulation is stopped when no infected individuals are left.

Arguments:

<i>graph</i> :	The graph to perform the model on. For directed graphs edge directions are ignored and a warning is given.
<i>beta</i> :	The rate of infection of an individual that is susceptible and has a single infected neighbor. The infection rate of a susceptible individual with n infected neighbors is n times β . Formally this is the rate parameter of an exponential distribution.
<i>gamma</i> :	The rate of recovery of an infected individual. Formally, this is the rate parameter of an exponential distribution.
<i>no_sim</i> :	The number of simulation runs to perform.
<i>result</i> :	The result of the simulation is stored here, in a list of <code>igraph_sir_t</code> objects. To deallocate memory, the user needs to call <code>igraph_sir_destroy</code> on each element, before destroying the pointer vector itself using <code>igraph_vector_ptr_destroy_all()</code> .

Returns:

Error code.

Time complexity: $O(\text{no_sim} * (|V| + |E| \log(|V|)))$.

igraph_sir_t — The result of one SIR model simulation.


```
typedef struct igraph_sir_t {  
    igraph_vector_t times;  
    igraph_vector_int_t no_s, no_i, no_r;  
} igraph_sir_t;
```

Data structure to store the results of one simulation of the SIR (susceptible-infected-recovered) model on a graph. It has the following members. They are all (real or integer) vectors, and they are of the same length.

Values:

`times`: A vector, the times of the events are stored here.

`no_s`: An integer vector, the number of susceptibles in each time step is stored here.

`no_i`: An integer vector, the number of infected individuals at each time step, is stored here.

`no_r`: An integer vector, the number of recovered individuals is stored here at each time step.

`igraph_sir_destroy` — Deallocates memory associated with a SIR simulation run.

```
void igraph_sir_destroy(igraph_sir_t *sir);
```

Arguments:

`sir`: The `igraph_sir_t` object storing the simulation.

Chapter 11. Vertex and edge selectors and sequences, iterators

About selectors, iterators

Everything about vertices and vertex selectors also applies to edges and edge selectors unless explicitly noted otherwise.

The vertex (and edge) selector notion was introduced in igraph 0.2. It is a way to reference a sequence of vertices or edges independently of the graph.

While this might sound quite mysterious, it is actually very simple. For example, all vertices of a graph can be selected by `igraph_vs_all()` and the graph independence means that `igraph_vs_all()` is not parametrized by a graph object. That is, `igraph_vs_all()` is the general *concept* of selecting all vertices of a graph. A vertex selector is then a way to specify the class of vertices to be visited. The selector might specify that all vertices of a graph or all the neighbours of a vertex are to be visited. A vertex selector is a way of saying that you want to visit a bunch of vertices, as opposed to a vertex iterator which is a concrete plan for visiting each of the chosen vertices of a specific graph.

To determine the actual vertex IDs implied by a vertex selector, you need to apply the concept of selecting vertices to a specific graph object. This can be accomplished by instantiating a vertex iterator using a specific vertex selection concept and a specific graph object. The notion of vertex iterators can be thought of in the following way. Given a specific graph object and the class of vertices to be visited, a vertex iterator is a road map, plan or route for how to visit the chosen vertices.

Some vertex selectors have *immediate* versions. These have the prefix `igraph_vss` instead of `igraph_vs`, e.g. `igraph_vss_all()` instead of `igraph_vs_all()`. The immediate versions are to be used in the parameter list of the igraph functions, such as `igraph_degree()`. These functions are not associated with any `igraph_vs_t` object, so they have no separate constructors and destructors (destroy functions).

Vertex selector constructors

Vertex selectors are created by vertex selector constructors, can be instantiated with `igraph_vit_create()`, and are destroyed with `igraph_vs_destroy()`.

`igraph_vs_all` — Vertex set, all vertices of a graph.

```
int igraph_vs_all(igraph_vs_t *vs);
```

Arguments:

`vs`: Pointer to an uninitialized `igraph_vs_t` object.

Returns:

Error code.

See also:

`igraph_vss_all()`, `igraph_vs_destroy()`

This selector includes all vertices of a given graph in increasing vertex id order.

Time complexity: $O(1)$.

igraph_vs_adj — Adjacent vertices of a vertex.

```
int igraph_vs_adj(igraph_vs_t *vs,  
                  igraph_integer_t vid, igraph_neimode_t mode);
```

All neighboring vertices of a given vertex are selected by this selector. The *mode* argument controls the type of the neighboring vertices to be selected. The vertices are visited in increasing vertex ID order, as of igraph version 0.4.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

vid: Vertex ID, the center of the neighborhood.

mode: Decides the type of the neighborhood for directed graphs. This parameter is ignored for undirected graphs. Possible values:

IGRAPH_OUT All vertices to which there is a directed edge from *vid*. That is, all the out-neighbors of *vid*.

IGRAPH_IN All vertices from which there is a directed edge to *vid*. In other words, all the in-neighbors of *vid*.

IGRAPH_ALL All vertices to which or from which there is a directed edge from/to *vid*. That is, all the neighbors of *vid* considered as if the graph is undirected.

Returns:

Error code.

See also:

`igraph_vs_destroy()`

Time complexity: $O(1)$.

igraph_vs_nonadj — Non-adjacent vertices of a vertex.

```
int igraph_vs_nonadj(igraph_vs_t *vs, igraph_integer_t vid,  
                     igraph_neimode_t mode);
```

All non-neighboring vertices of a given vertex. The *mode* argument controls the type of neighboring vertices *not* to select. Instead of selecting immediate neighbors of *vid* as is done by `igraph_vs_adj()`, the current function selects vertices that are *not* immediate neighbors of *vid*.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

vid: Vertex ID, the “center” of the non-neighborhood.

mode: The type of neighborhood not to select in directed graphs. Possible values:

IGRAPH_OUT	All vertices will be selected except those to which there is a directed edge from <i>vid</i> . That is, we select all vertices excluding the out-neighbors of <i>vid</i> .
IGRAPH_IN	All vertices will be selected except those from which there is a directed edge to <i>vid</i> . In other words, we select all vertices but the in-neighbors of <i>vid</i> .
IGRAPH_ALL	All vertices will be selected except those from or to which there is a directed edge to or from <i>vid</i> . That is, we select all vertices of <i>vid</i> except for its immediate neighbors.

Returns:

Error code.

See also:

`igraph_vs_destroy()`

Time complexity: $O(1)$.

Example 11.1. File `examples/simple/igraph_vs_nonadj.c`

igraph_vs_none — Empty vertex set.

```
int igraph_vs_none(igraph_vs_t *vs);
```

Creates an empty vertex selector.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

Returns:

Error code.

See also:

`igraph_vss_none()`, `igraph_vs_destroy()`

Time complexity: $O(1)$.

igraph_vs_1 — Vertex set with a single vertex.

```
int igraph_vs_1(igraph_vs_t *vs, igraph_integer_t vid);
```

This vertex selector selects a single vertex.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

vid: The vertex id to be selected.

Returns:

Error Code.

See also:

`igraph_vss_1()`, `igraph_vs_destroy()`

Time complexity: $O(1)$.

igraph_vs_vector — Vertex set based on a vector.

```
int igraph_vs_vector(igraph_vs_t *vs,  
                    const igraph_vector_t *v);
```

This function makes it possible to handle a `vector_t` temporarily as a vertex selector. The vertex selector should be thought of like a *view* to the vector. If you make changes to the vector that also affects the vertex selector. Destroying the vertex selector does not destroy the vector. (Of course.) Do not destroy the vector before destroying the vertex selector, or you might get strange behavior.

Arguments:

vs: Pointer to an uninitialized vertex selector.

v: Pointer to a `igraph_vector_t` object.

Returns:

Error code.

See also:

`igraph_vss_vector()`, `igraph_vs_destroy()`

Time complexity: $O(1)$.

Example 11.2. File `examples/simple/igraph_vs_vector.c`

igraph_vs_vector_small — Create a vertex set by giving its elements.

```
int igraph_vs_vector_small(igraph_vs_t *vs, ...);
```

This function can be used to create a vertex selector with a couple of vertices. Do not forget to include a `-1` after the last vertex id. The behavior of the function is undefined if you don't use a `-1` properly.

Note that the vertex ids supplied will be parsed as `int`'s so you cannot supply arbitrarily large (too large for `int`) vertex ids here.

Arguments:

`vs`: Pointer to an uninitialized vertex selector object.

`...`: Additional parameters, these will be the vertex ids to be included in the vertex selector. Supply a `-1` after the last vertex id.

Returns:

Error code.

See also:

`igraph_vs_destroy()`

Time complexity: $O(n)$, the number of vertex ids supplied.

`igraph_vs_vector_copy` — Vertex set based on a vector, with copying.

```
int igraph_vs_vector_copy(igraph_vs_t *vs,
                          const igraph_vector_t *v);
```

This function makes it possible to handle a `vector_t` permanently as a vertex selector. The vertex selector creates a copy of the original vector, so the vector can safely be destroyed after creating the vertex selector. Changing the original vector will not affect the vertex selector. The vertex selector is responsible for deleting the copy made by itself.

Arguments:

`vs`: Pointer to an uninitialized vertex selector.

`v`: Pointer to a `igraph_vector_t` object.

Returns:

Error code.

See also:

`igraph_vs_destroy()`

Time complexity: $O(1)$.

`igraph_vs_seq` — Vertex set, an interval of vertices.

```
int igraph_vs_seq(igraph_vs_t *vs,
```

```
igraph_integer_t from, igraph_integer_t to);
```

Creates a vertex selector containing all vertices with vertex id equal to or bigger than *from* and equal to or smaller than *to*.

Arguments:

vs: Pointer to an uninitialized vertex selector object.
from: The first vertex id to be included in the vertex selector.
to: The last vertex id to be included in the vertex selector.

Returns:

Error code.

See also:

```
igraph_vss_seq(), igraph_vs_destroy()
```

Time complexity: $O(1)$.

Example 11.3. File `examples/simple/igraph_vs_seq.c`

Generic vertex selector operations

igraph_vs_copy — Creates a copy of a vertex selector.

```
int igraph_vs_copy(igraph_vs_t* dest, const igraph_vs_t* src);
```

Arguments:

src: The selector being copied.
dest: An uninitialized selector that will contain the copy.

igraph_vs_destroy — Destroy a vertex set.

```
void igraph_vs_destroy(igraph_vs_t *vs);
```

This function should be called for all vertex selectors when they are not needed. The memory allocated for the vertex selector will be deallocated. Do not call this function on vertex selectors created with the immediate versions of the vertex selector constructors (starting with `igraph_vss`).

Arguments:

vs: Pointer to a vertex selector object.
Time complexity: operating system dependent, usually $O(1)$.

igraph_vs_is_all — Check whether all vertices are included.

```
igraph_bool_t igraph_vs_is_all(const igraph_vs_t *vs);
```

This function checks whether the vertex selector object was created by `igraph_vs_all()` or `igraph_vss_all()`. Note that the vertex selector might contain all vertices in a given graph but if it wasn't created by the two constructors mentioned here the return value will be `FALSE`.

Arguments:

vs: Pointer to a vertex selector object.

Returns:

TRUE (1) if the vertex selector contains all vertices and FALSE (0) otherwise.

Time complexity: $O(1)$.

igraph_vs_size — Returns the size of the vertex selector.

```
int igraph_vs_size(const igraph_t *graph, const igraph_vs_t *vs,  
                  igraph_integer_t *result);
```

The size of the vertex selector is the number of vertices it will yield when it is iterated over.

Arguments:

graph: The graph over which we will iterate.

result: The result will be returned here.

igraph_vs_type — Returns the type of the vertex selector.

```
int igraph_vs_type(const igraph_vs_t *vs);
```

Immediate vertex selectors

igraph_vss_all — All vertices of a graph (immediate version).


```
igraph_vs_t igraph_vss_all(void);
```

Immediate vertex selector for all vertices in a graph. It can be used conveniently when some vertex property (e.g. betweenness, degree, etc.) should be calculated for all vertices.

Returns:

A vertex selector for all vertices in a graph.

See also:

```
igraph_vs_all()
```

Time complexity: $O(1)$.

igraph_vss_none — Empty vertex set (immediate version).

```
igraph_vs_t igraph_vss_none(void);
```

The immediate version of the empty vertex selector.

Returns:

An empty vertex selector.

See also:

```
igraph_vs_none()
```

Time complexity: $O(1)$.

igraph_vss_1 — Vertex set with a single vertex (immediate version).

```
igraph_vs_t igraph_vss_1(igraph_integer_t vid);
```

The immediate version of the single-vertex selector.

Arguments:

vid: The vertex to be selected.

Returns:

A vertex selector containing a single vertex.

See also:

```
igraph_vs_1()
```

Time complexity: $O(1)$.

igraph_vss_vector — Vertex set based on a vector (immediate version).

```
igraph_vs_t igraph_vss_vector(const igraph_vector_t *v);
```

This is the immediate version of `igraph_vs_vector`.

Arguments:

v: Pointer to a `igraph_vector_t` object.

Returns:

A vertex selector object containing the vertices in the vector.

See also:

```
igraph_vs_vector()
```

Time complexity: $O(1)$.

igraph_vss_seq — An interval of vertices (immediate version).

```
igraph_vs_t igraph_vss_seq(igraph_integer_t from, igraph_integer_t to);
```

The immediate version of `igraph_vs_seq()`.

Arguments:

from: The first vertex id to be included in the vertex selector.

to: The last vertex id to be included in the vertex selector.

Returns:

Error code.

See also:

```
igraph_vs_seq()
```

Time complexity: $O(1)$.

Vertex iterators

igraph_vit_create — Creates a vertex iterator from a vertex selector.

```
int igraph_vit_create(const igraph_t *graph,
                     igraph_vs_t vs, igraph_vit_t *vit);
```

This function instantiates a vertex selector object with a given graph. This is the step when the actual vertex ids are created from the *logical* notion of the vertex selector based on the graph. E.g. a vertex selector created with `igraph_vs_all()` contains knowledge that *all* vertices are included in a (yet indefinite) graph. When instantiating it a vertex iterator object is created, this contains the actual vertex ids in the graph supplied as a parameter.

The same vertex selector object can be used to instantiate any number vertex iterators.

Arguments:

graph: An `igraph_t` object, a graph.

vs: A vertex selector object.

vit: Pointer to an uninitialized vertex iterator object.

Returns:

Error code.

See also:

`igraph_vit_destroy()`.

Time complexity: it depends on the vertex selector type. $O(1)$ for vertex selectors created with `igraph_vs_all()`, `igraph_vs_none()`, `igraph_vs_l`, `igraph_vs_vector`, `igraph_vs_seq()`, `igraph_vs_vector()`, `igraph_vs_vector_small()`. $O(d)$ for `igraph_vs_adj()`, d is the number of vertex ids to be included in the iterator. $O(|V|)$ for `igraph_vs_nonadj()`, $|V|$ is the number of vertices in the graph.

igraph_vit_destroy — Destroys a vertex iterator.

```
void igraph_vit_destroy(const igraph_vit_t *vit);
```

Deallocates memory allocated for a vertex iterator.

Arguments:

vit: Pointer to an initialized vertex iterator object.

See also:

`igraph_vit_create()`

Time complexity: operating system dependent, usually $O(1)$.

Stepping over the vertices

After creating an iterator with `igraph_vit_create()`, it points to the first vertex in the vertex determined by the vertex selector (if there is any). The `IGRAPH_VIT_NEXT()` macro steps to the next vertex, `IGRAPH_VIT_END()` checks whether there are more vertices to vis-

it, `IGRAPH_VIT_SIZE()` gives the total size of the vertices visited so far and to be visited. `IGRAPH_VIT_RESET()` resets the iterator, it will point to the first vertex again. Finally `IGRAPH_VIT_GET()` gives the current vertex pointed to by the iterator (call this only if `IGRAPH_VIT_END()` is false).

Here is an example on how to step over the neighbors of vertex 0:

```
igraph_vs_t vs;
igraph_vit_t vit;
...
igraph_vs_adj(&vs, 0, IGRAPH_ALL);
igraph_vit_create(&graph, vs, &vit);
while (!IGRAPH_VIT_END(vit)) {
    printf(" %li", (long int) IGRAPH_VIT_GET(vit));
    IGRAPH_VIT_NEXT(vit);
}
printf("\n");
...
igraph_vit_destroy(&vit);
igraph_vs_destroy(&vs);
```

IGRAPH_VIT_NEXT — Next vertex.

```
#define IGRAPH_VIT_NEXT(vit)
```

Steps the iterator to the next vertex. Only call this function if `IGRAPH_VIT_END()` returns false.

Arguments:

vit: The vertex iterator to step.

Time complexity: $O(1)$.

IGRAPH_VIT_END — Are we at the end?

```
#define IGRAPH_VIT_END(vit)
```

Checks whether there are more vertices to step to.

Arguments:

vit: The vertex iterator to check.

Returns:

Logical value, if true there are no more vertices to step to.

Time complexity: $O(1)$.

IGRAPH_VIT_SIZE — Size of a vertex iterator.

```
#define IGRAPH_VIT_SIZE(vit)
```

Gives the number of vertices in a vertex iterator.

Arguments:

vit: The vertex iterator.

Returns:

The number of vertices.

Time complexity: $O(1)$.

IGRAPH_VIT_RESET — Reset a vertex iterator.

```
#define IGRAPH_VIT_RESET(vit)
```

Resets a vertex iterator. After calling this macro the iterator will point to the first vertex.

Arguments:

vit: The vertex iterator.

Time complexity: $O(1)$.

IGRAPH_VIT_GET — Query the current position.

```
#define IGRAPH_VIT_GET(vit)
```

Gives the vertex id of the current vertex pointed to by the iterator.

Arguments:

vit: The vertex iterator.

Returns:

The vertex id of the current vertex.

Time complexity: $O(1)$.

Edge selector constructors

igraph_es_all — Edge set, all edges.

```
int igraph_es_all(igraph_es_t *es,  
                  igraph_edgeorder_type_t order);
```

Arguments:

es: Pointer to an uninitialized edge selector object.

order: Constant giving the order in which the edges will be included in the selector. Possible values: `IGRAPH_EDGEORDER_ID`, edge id order. `IGRAPH_EDGEORDER_FROM`, vertex id order, the id of the *source* vertex counts for directed graphs. The order of the incident edges of a given vertex is arbitrary. `IGRAPH_EDGEORDER_TO`, vertex id order, the id of the *target* vertex counts for directed graphs. The order of the incident edges of a given vertex is arbitrary. For undirected graph the latter two is the same.

Returns:

Error code.

See also:

`igraph_ess_all()`, `igraph_es_destroy()`

Time complexity: $O(1)$.

`igraph_es_incident` — Edges incident on a given vertex.

```
int igraph_es_incident(igraph_es_t *es,  
                      igraph_integer_t vid, igraph_neimode_t mode);
```

Arguments:

es: Pointer to an uninitialized edge selector object.

vid: Vertex id, of which the incident edges will be selected.

mode: Constant giving the type of the incident edges to select. This is ignored for undirected graphs. Possible values: `IGRAPH_OUT`, outgoing edges; `IGRAPH_IN`, incoming edges; `IGRAPH_ALL`, all edges.

Returns:

Error code.

See also:

`igraph_es_destroy()`

Time complexity: $O(1)$.

`igraph_es_none` — Empty edge selector.

```
int igraph_es_none(igraph_es_t *es);
```

Arguments:

es: Pointer to an uninitialized edge selector object to initialize.

Returns:

Error code.

See also:

`igraph_ess_none()`, `igraph_es_destroy()`

Time complexity: $O(1)$.

igraph_es_1 — Edge selector containing a single edge.

```
int igraph_es_1(igraph_es_t *es, igraph_integer_t eid);
```

Arguments:

es: Pointer to an uninitialized edge selector object.

eid: Edge id of the edge to select.

Returns:

Error code.

See also:

`igraph_ess_1()`, `igraph_es_destroy()`

Time complexity: $O(1)$.

igraph_es_vector — Handle a vector as an edge selector.

```
int igraph_es_vector(igraph_es_t *es,  
                    const igraph_vector_t *v);
```

Creates an edge selector which serves as a view to a vector containing edge ids. Do not destroy the vector before destroying the view. Many views can be created to the same vector.

Arguments:

es: Pointer to an uninitialized edge selector.

v: Vector containing edge ids.

Returns:

Error code.

See also:

`igraph_ess_vector()`, `igraph_es_destroy()`

Time complexity: $O(1)$.

`igraph_es_fromto` — Edge selector, all edges between two vertex sets.

```
int igraph_es_fromto(igraph_es_t *es,
                    igraph_vs_t from, igraph_vs_t to);
```

This function is not implemented yet.

Arguments:

es: Pointer to an uninitialized edge selector.

from: Vertex selector, their outgoing edges will be selected.

to: Vertex selector, their incoming edges will be selected from the previous selection.

Returns:

Error code.

See also:

`igraph_es_destroy()`

Time complexity: $O(1)$.

`igraph_es_seq` — Edge selector, a sequence of edge ids.

```
int igraph_es_seq(igraph_es_t *es,
                 igraph_integer_t from, igraph_integer_t to);
```

All edge ids between *from* and *to* will be included in the edge selection. This includes *from* and excludes *to*.

Arguments:

es: Pointer to an uninitialized edge selector object.

from: The first edge id to be included.

to: The last edge id to be included.

Returns:

Error code.

See also:

`igraph_ess_seq()`, `igraph_es_destroy()`

Time complexity: $O(1)$.

`igraph_es_pairs` — Edge selector, multiple edges defined by their endpoints in a vector.

```
int igraph_es_pairs(igraph_es_t *es, const igraph_vector_t *v,  
                    igraph_bool_t directed);
```

The edges between the given pairs of vertices will be included in the edge selection. The vertex pairs must be defined in the vector `v`, the first element of the vector is the first vertex of the first edge to be selected, the second element is the second vertex of the first edge, the third element is the first vertex of the second edge and so on.

Arguments:

es: Pointer to an uninitialized edge selector object.

v: The vector containing the endpoints of the edges.

directed: Whether the graph is directed or not.

Returns:

Error code.

See also:

`igraph_es_pairs_small()`, `igraph_es_destroy()`

Time complexity: $O(n)$, the number of edges being selected.

Example 11.4. File `examples/simple/igraph_es_pairs.c`

`igraph_es_pairs_small` — Edge selector, multiple edges defined by their endpoints as arguments.

```
int igraph_es_pairs_small(igraph_es_t *es, igraph_bool_t directed, ...);
```

The edges between the given pairs of vertices will be included in the edge selection. The vertex pairs must be given as the arguments of the function call, the third argument is the first vertex of the first edge, the fourth argument is the second vertex of the first edge, the fifth is the first vertex of the second edge and so on. The last element of the argument list must be -1 to denote the end of the argument list.

Arguments:

es: Pointer to an uninitialized edge selector object.

directed: Whether the graph is directed or not.

Returns:

Error code.

See also:

`igraph_es_pairs()`, `igraph_es_destroy()`

Time complexity: $O(n)$, the number of edges being selected.

igraph_es_path — Edge selector, edge ids on a path.

```
int igraph_es_path(igraph_es_t *es, const igraph_vector_t *v,  
                  igraph_bool_t directed);
```

This function takes a vector of vertices and creates a selector of edges between those vertices. Vector {0, 3, 4, 7} will select edges (0 -> 3), (3 -> 4), (4 -> 7). If these edges don't exist then trying to create an iterator using this selector will fail.

Arguments:

es: Pointer to an uninitialized edge selector object.

v: Pointer to a vector of vertex id's along the path.

directed: If edge directions should be taken into account. This will be ignored if the graph to select from is undirected.

Returns:

Error code.

See also:

`igraph_es_destroy()`

Time complexity: $O(n)$, the number of vertices.

igraph_es_vector_copy — Edge set, based on a vector, with copying.

```
int igraph_es_vector_copy(igraph_es_t *es, const igraph_vector_t *v);
```

This function makes it possible to handle a `vector_t` permanently as an edge selector. The edge selector creates a copy of the original vector, so the vector can safely be destroyed after creating the edge selector. Changing the original vector will not affect the edge selector. The edge selector is responsible for deleting the copy made by itself.

Arguments:

es: Pointer to an uninitialized edge selector.

v: Pointer to a `igraph_vector_t` object.

Returns:

Error code.

See also:

`igraph_es_destroy()`

Time complexity: $O(1)$.

Immediate edge selectors

`igraph_ess_all` — Edge set, all edges (immediate version)

```
igraph_es_t igraph_ess_all(igraph_edgeorder_type_t order);
```

The immediate version of the all-edges selector.

Arguments:

order: Constant giving the order of the edges in the edge selector. See `igraph_es_all()` for the possible values.

Returns:

The edge selector.

See also:

`igraph_es_all()`

Time complexity: $O(1)$.

`igraph_ess_none` — Immediate empty edge selector.

```
igraph_es_t igraph_ess_none(void);
```

Immediate version of the empty edge selector.

Returns:

Initialized empty edge selector.

See also:

`igraph_es_none()`

Time complexity: $O(1)$.

`igraph_ess_1` — Immediate version of the single edge edge selector.

```
igraph_es_t igraph_ess_1(igraph_integer_t eid);
```

Arguments:

eid: The id of the edge.

Returns:

The edge selector.

See also:

`igraph_es_1()`

Time complexity: $O(1)$.

`igraph_ess_vector` — Immediate vector view edge selector.

```
igraph_es_t igraph_ess_vector(const igraph_vector_t *v);
```

This is the immediate version of the vector of edge ids edge selector.

Arguments:

v: The vector of edge ids.

Returns:

Edge selector, initialized.

See also:

`igraph_es_vector()`

Time complexity: $O(1)$.

`igraph_ess_seq` — Immediate version of the sequence edge selector.

```
igraph_es_t igraph_ess_seq(igraph_integer_t from, igraph_integer_t to);
```

Arguments:

from: The first edge id to include.

to: The last edge id to include.

Returns:

The initialized edge selector.

See also:

`igraph_es_seq()`

Time complexity: $O(1)$.

Generic edge selector operations

`igraph_es_as_vector` — Transform edge selector into vector.

```
int igraph_es_as_vector(const igraph_t *graph, igraph_es_t es,  
                        igraph_vector_t *v);
```

Call this function on an edge selector to transform it into a vector. This is only implemented for sequence and vector selectors. If the edges do not exist in the graph, this will result in an error.

Arguments:

graph: Pointer to a graph to check if the edges in the selector exist.

es: An edge selector object.

v: Pointer to initialized vector. The result will be stored here.

Time complexity: $O(n)$, the number of edges in the selector.

`igraph_es_copy` — Creates a copy of an edge selector.

```
int igraph_es_copy(igraph_es_t* dest, const igraph_es_t* src);
```

Arguments:

src: The selector being copied.

dest: An uninitialized selector that will contain the copy.

See also:

```
igraph_es_destroy()
```

igraph_es_destroy — Destroys an edge selector object.

```
void igraph_es_destroy(igraph_es_t *es);
```

Call this function on an edge selector when it is not needed any more. Do *not* call this function on edge selectors created by immediate constructors, those don't need to be destroyed.

Arguments:

es: Pointer to an edge selector object.

Time complexity: operating system dependent, usually $O(1)$.

igraph_es_is_all — Check whether an edge selector includes all edges.

```
igraph_bool_t igraph_es_is_all(const igraph_es_t *es);
```

Arguments:

es: Pointer to an edge selector object.

Returns:

TRUE (1) if *es* was created with `igraph_es_all()` or `igraph_ess_all()`, and FALSE (0) otherwise.

Time complexity: $O(1)$.

igraph_es_size — Returns the size of the edge selector.

```
int igraph_es_size(const igraph_t *graph, const igraph_es_t *es,  
                  igraph_integer_t *result);
```

The size of the edge selector is the number of edges it will yield when it is iterated over.

Arguments:

graph: The graph over which we will iterate.

result: The result will be returned here.

igraph_es_type — Returns the type of the edge selector.

```
int igraph_es_type(const igraph_es_t *es);
```

Edge iterators

igraph_eit_create — Creates an edge iterator from an edge selector.

```
int igraph_eit_create(const igraph_t *graph,  
                     igraph_es_t es, igraph_eit_t *eit);
```

This function creates an edge iterator based on an edge selector and a graph.

The same edge selector can be used to create many edge iterators, also for different graphs.

Arguments:

graph: An `igraph_t` object for which the edge selector will be instantiated.

es: The edge selector to instantiate.

eit: Pointer to an uninitialized edge iterator.

Returns:

Error code.

See also:

```
igraph_eit_destroy()
```

Time complexity: depends on the type of the edge selector. For edge selectors created by `igraph_es_all()`, `igraph_es_none()`, `igraph_es_l()`, `igraph_es_vector()`, `igraph_es_seq()` it is $O(1)$. For `igraph_es_incident()` it is $O(d)$ where d is the number of incident edges of the vertex.

igraph_eit_destroy — Destroys an edge iterator.

```
void igraph_eit_destroy(const igraph_eit_t *eit);
```

Arguments:

eit: Pointer to an edge iterator to destroy.

See also:

`igraph_eit_create()`

Time complexity: operating system dependent, usually $O(1)$.

Stepping over the edges

Just like for vertex iterators, macros are provided for stepping over a sequence of edges: `IGRAPH_EIT_NEXT()` goes to the next edge, `IGRAPH_EIT_END()` checks whether there are more edges to visit, `IGRAPH_EIT_SIZE()` gives the number of edges in the edge sequence, `IGRAPH_EIT_RESET()` resets the iterator to the first edge and `IGRAPH_EIT_GET()` returns the id of the current edge.

IGRAPH_EIT_NEXT — Next edge.

```
#define IGRAPH_EIT_NEXT(eit)
```

Steps the iterator to the next edge. Call this function only if `IGRAPH_EIT_END()` returns false.

Arguments:

eit: The edge iterator to step.

Time complexity: $O(1)$.

IGRAPH_EIT_END — Are we at the end?

```
#define IGRAPH_EIT_END(eit)
```

Checks whether there are more edges to step to.

Arguments:

wit: The edge iterator to check.

Returns:

Logical value, if true there are no more edges to step to.

Time complexity: $O(1)$.

IGRAPH_EIT_SIZE — Number of edges in the iterator.

```
#define IGRAPH_EIT_SIZE(eit)
```

Gives the number of edges in an edge iterator.

Arguments:

eit: The edge iterator.

Returns:

The number of edges.

Time complexity: $O(1)$.

IGRAPH_EIT_RESET — Reset an edge iterator.

```
#define IGRAPH_EIT_RESET(eit)
```

Resets an edge iterator. After calling this macro the iterator will point to the first edge.

Arguments:

eit: The edge iterator.

Time complexity: $O(1)$.

IGRAPH_EIT_GET — Query an edge iterator.

```
#define IGRAPH_EIT_GET(eit)
```

Gives the edge id of the current edge pointed to by an iterator.

Arguments:

eit: The edge iterator.

Returns:

The id of the current edge.

Time complexity: $O(1)$.

Chapter 12. Graph, vertex and edge attributes

Attributes are numbers or strings (or basically any kind of data) associated with the vertices or edges of a graph, or with the graph itself. Eg. you may label vertices with symbolic names or attach numeric weights to the edges of a graph.

igraph attributes are designed to be flexible and extensible. In igraph attributes are implemented via an interface abstraction: any type implementing the functions in the interface, can be used for storing vertex, edge and graph attributes. This means that different attribute implementations can be used together with igraph. This is reasonable: if igraph is used from Python attributes can be of any Python type, from GNU R all R types are allowed. There is an experimental attribute implementation to be used when programming in C, but by default it is currently turned off.

First we briefly look over how attribute handlers can be implemented. This is not something a user does every day. It is rather typically the job of the high level interface writers. (But it is possible to write an interface without implementing attributes.) Then we show the experimental C attribute handler.

The Attribute Handler Interface

It is possible to attach an attribute handling interface to **igraph**. This is simply a table of functions, of type `igraph_attribute_table_t`. These functions are invoked to notify the attribute handling code about the structural changes in a graph. See the documentation of this type for details.

By default there is no attribute interface attached to **igraph**, to attach one, call `igraph_set_attribute_table` with your new table.

`igraph_attribute_table_t` — Table of functions to perform operations on attributes

```
typedef struct igraph_attribute_table_t {
    int (*init)(igraph_t *graph, igraph_vector_ptr_t *attr);
    void (*destroy)(igraph_t *graph);
    int (*copy)(igraph_t *to, const igraph_t *from, igraph_bool_t ga,
                igraph_bool_t va, igraph_bool_t ea);
    int (*add_vertices)(igraph_t *graph, long int nv, igraph_vector_ptr_t *attr);
    int (*permute_vertices)(const igraph_t *graph,
                            igraph_t *newgraph,
                            const igraph_vector_t *idx);
    int (*combine_vertices)(const igraph_t *graph,
                            igraph_t *newgraph,
                            const igraph_vector_ptr_t *merges,
                            const igraph_attribute_combination_t *comb);
    int (*add_edges)(igraph_t *graph, const igraph_vector_t *edges,
                    igraph_vector_ptr_t *attr);
    int (*permute_edges)(const igraph_t *graph,
                         igraph_t *newgraph, const igraph_vector_t *idx);
    int (*combine_edges)(const igraph_t *graph,
                         igraph_t *newgraph,
                         const igraph_vector_ptr_t *merges,
                         const igraph_attribute_combination_t *comb);
    int (*get_info)(const igraph_t *graph,
```

```
igraph_strvector_t *gnames, igraph_vector_t *gtypes,
igraph_strvector_t *vnames, igraph_vector_t *vtypes,
igraph_strvector_t *enames, igraph_vector_t *etypes);
igraph_bool_t (*has_attr)(const igraph_t *graph, igraph_attribute_elemtyp
                           const char *name);
int (*gettype)(const igraph_t *graph, igraph_attribute_type_t *type,
               igraph_attribute_elemtyp_t elemtype, const char *name);
int (*get_numeric_graph_attr)(const igraph_t *graph, const char *name,
                              igraph_vector_t *value);
int (*get_string_graph_attr)(const igraph_t *graph, const char *name,
                             igraph_strvector_t *value);
int (*get_bool_graph_attr)(const igraph_t *igraph, const char *name,
                           igraph_vector_bool_t *value);
int (*get_numeric_vertex_attr)(const igraph_t *graph, const char *name,
                               igraph_vs_t vs,
                               igraph_vector_t *value);
int (*get_string_vertex_attr)(const igraph_t *graph, const char *name,
                              igraph_vs_t vs,
                              igraph_strvector_t *value);
int (*get_bool_vertex_attr)(const igraph_t *graph, const char *name,
                            igraph_vs_t vs,
                            igraph_vector_bool_t *value);
int (*get_numeric_edge_attr)(const igraph_t *graph, const char *name,
                             igraph_es_t es,
                             igraph_vector_t *value);
int (*get_string_edge_attr)(const igraph_t *graph, const char *name,
                            igraph_es_t es,
                            igraph_strvector_t *value);
int (*get_bool_edge_attr)(const igraph_t *graph, const char *name,
                          igraph_es_t es,
                          igraph_vector_bool_t *value);
} igraph_attribute_table_t;
```

This type collects the functions defining an attribute handler. It has the following members:

Values:

<code>init:</code>	This function is called whenever a new graph object is created, right after it is created but before any vertices or edges are added. It is supposed to set the <code>attr</code> member of the <code>igraph_t</code> object. It is expected to return an error code.
<code>destroy:</code>	This function is called whenever the graph object is destroyed, right before freeing the allocated memory.
<code>copy:</code>	This function is called when copying a graph with <code>igraph_copy</code> , after the structure of the graph has been already copied. It is expected to return an error code.
<code>add_vertices:</code>	Called when vertices are added to a graph, before adding the vertices themselves. The number of vertices to add is supplied as an argument. Expected to return an error code.
<code>permute_vertices:</code>	Typically called when a new graph is created based on an existing one, e.g. if vertices are removed from a graph. The supplied index vector defines which old vertex a new vertex corresponds to. Its length must be the same as the number of vertices in the new graph.

<code>combine_vertices:</code>	This function is called when the creation of a new graph involves a merge (contraction, etc.) of vertices from another graph. The function is after the new graph was created. An argument specifies how several vertices from the old graph map to a single vertex in the new graph.
<code>add_edges:</code>	Called when new edges have been added. The number of new edges are supplied as well. It is expected to return an error code.
<code>permute_edges:</code>	Typically called when a new graph is created and some of the new edges should carry the attributes of some of the old edges. The <code>idx</code> vector shows the mapping between the old edges and the new ones. Its length is the same as the number of edges in the new graph, and for each edge it gives the id of the old edge (the edge in the old graph).
<code>combine_edges:</code>	This function is called when the creation of a new graph involves a merge (contraction, etc.) of edges from another graph. The function is after the new graph was created. An argument specifies how several edges from the old graph map to a single edge in the new graph.
<code>get_info:</code>	Query the attributes of a graph, the names and types should be returned.
<code>has_attr:</code>	Check whether a graph has the named graph/vertex/edge attribute.
<code>gettype:</code>	Query the type of a graph/vertex/edge attribute.
<code>get_numeric_graph_attr:</code>	Query a numeric graph attribute. The value should be placed as the first element of the <i>value</i> vector.
<code>get_string_graph_attr:</code>	Query a string graph attribute. The value should be placed as the first element of the <i>value</i> string vector.
<code>get_bool_graph_attr:</code>	Query a boolean graph attribute. The value should be placed as the first element of the <i>value</i> boolean vector.
<code>get_numeric_vertex_attr:</code>	Query a numeric vertex attribute, for the vertices included in <i>vs</i> .
<code>get_string_vertex_attr:</code>	Query a string vertex attribute, for the vertices included in <i>vs</i> .
<code>get_bool_vertex_attr:</code>	Query a boolean vertex attribute, for the vertices included in <i>vs</i> .
<code>get_numeric_edge_attr:</code>	Query a numeric edge attribute, for the edges included in <i>es</i> .
<code>get_string_edge_attr:</code>	Query a string edge attribute, for the edges included in <i>es</i> .
<code>get_bool_edge_attr:</code>	Query a boolean edge attribute, for the edges included in <i>es</i> .

Note that the `get_*_*_attr` are allowed to convert the attributes to numeric or string. E.g. if a vertex attribute is a GNU R complex data type, then `get_string_vertex_attribute` may serialize it into a string, but this probably makes sense only if `add_vertices` is able to deserialize it.

igraph_set_attribute_table — Attach an attribute table.

```
igraph_attribute_table_t *  
igraph_set_attribute_table(const igraph_attribute_table_t * table);
```

This function attaches attribute handling code to the igraph library. Note that the attribute handler table is *not* thread-local even if igraph is compiled in thread-local mode. In the vast majority of cases, this is not a significant restriction.

Arguments:

table: Pointer to an `igraph_attribute_table_t` object containing the functions for attribute manipulation. Supply NULL here if you don't want attributes.

Returns:

Pointer to the old attribute handling table.

Time complexity: $O(1)$.

`igraph_attribute_type_t` — The possible types of the attributes.

```
typedef enum { IGRAPH_ATTRIBUTE_DEFAULT = 0,  
               IGRAPH_ATTRIBUTE_NUMERIC = 1,  
               IGRAPH_ATTRIBUTE_BOOLEAN = 5,  
               IGRAPH_ATTRIBUTE_STRING = 2,  
               IGRAPH_ATTRIBUTE_R_OBJECT = 3,  
               IGRAPH_ATTRIBUTE_PY_OBJECT = 4  
             } igraph_attribute_type_t;
```

Note that this is only the type communicated by the attribute interface towards igraph functions. Eg. in the GNU R attribute handler, it is safe to say that all complex R object attributes are strings, as long as this interface is able to serialize them into strings. See also `igraph_attribute_table_t`.

Values:

IGRAPH_ATTRIBUTE_DEFAULT:	Currently not used for anything.
IGRAPH_ATTRIBUTE_NUMERIC:	Numeric attribute.
IGRAPH_ATTRIBUTE_BOOLEAN:	Logical values, true or false.
IGRAPH_ATTRIBUTE_STRING:	Attribute that can be converted to a string.
IGRAPH_ATTRIBUTE_R_OBJECT:	An R object. This is usually ignored by the igraph functions.
IGRAPH_ATTRIBUTE_PY_OBJECT:	A Python object. Usually ignored by the igraph functions.

Handling attribute combination lists

Several graph operations may collapse multiple vertices or edges into a single one. Attribute combination lists are used to indicate to the attribute handler how to combine the attributes of the original

vertices or edges and how to derive the final attribute value that is to be assigned to the collapsed vertex or edge. For example, `igraph_simplify()` removes loops and combines multiple edges into a single one; in case of a graph with an edge attribute named `weight` the attribute combination list can tell the attribute handler whether the weight of a collapsed edge should be the sum, the mean or some other function of the weights of the original edges that were collapsed into one.

One attribute combination list may contain several attribute combination records, one for each vertex or edge attribute that is to be handled during the operation.

`igraph_attribute_combination_init` — Initialize attribute combination list.

```
int igraph_attribute_combination_init(igraph_attribute_combination_t *comb);
```

Arguments:

comb: The uninitialized attribute combination list.

Returns:

Error code.

Time complexity: $O(1)$

`igraph_attribute_combination_add` — Add combination record to attribute combination list.

```
int igraph_attribute_combination_add(igraph_attribute_combination_t *comb,
                                     const char *name,
                                     igraph_attribute_combination_type_t type,
                                     igraph_function_pointer_t func);
```

Arguments:

comb: The attribute combination list.

name: The name of the attribute. If the name already exists the attribute combination record will be replaced. Use `NULL` to add a default combination record for all attributes not in the list.

type: The type of the attribute combination. See `igraph_attribute_combination_type_t` for the options.

func: Function to be used if *type* is `IGRAPH_ATTRIBUTE_COMBINE_FUNCTION`.

Returns:

Error code.

Time complexity: $O(n)$, where n is the number of current attribute combinations.

igraph_attribute_combination_remove — Remove a record from an attribute combination list.

```
int igraph_attribute_combination_remove(igraph_attribute_combination_t *comb,
                                       const char *name);
```

Arguments:

comb: The attribute combination list.

name: The attribute name of the attribute combination record to remove. It will be ignored if the named attribute does not exist. It can be NULL to remove the default combination record.

Returns:

Error code. This currently always returns IGRAPH_SUCCESS.

Time complexity: $O(n)$, where n is the number of records in the attribute combination list.

igraph_attribute_combination_destroy — Destroy attribute combination list.

```
void igraph_attribute_combination_destroy(igraph_attribute_combination_t *comb)
```

Arguments:

comb: The attribute combination list.

Time complexity: $O(n)$, where n is the number of records in the attribute combination list.

igraph_attribute_combination_type_t — The possible types of attribute combinations.

```
typedef enum {
    IGRAPH_ATTRIBUTE_COMBINE_IGNORE = 0,
    IGRAPH_ATTRIBUTE_COMBINE_DEFAULT = 1,
    IGRAPH_ATTRIBUTE_COMBINE_FUNCTION = 2,
    IGRAPH_ATTRIBUTE_COMBINE_SUM = 3,
    IGRAPH_ATTRIBUTE_COMBINE_PROD = 4,
    IGRAPH_ATTRIBUTE_COMBINE_MIN = 5,
    IGRAPH_ATTRIBUTE_COMBINE_MAX = 6,
    IGRAPH_ATTRIBUTE_COMBINE_RANDOM = 7,
    IGRAPH_ATTRIBUTE_COMBINE_FIRST = 8,
    IGRAPH_ATTRIBUTE_COMBINE_LAST = 9,
    IGRAPH_ATTRIBUTE_COMBINE_MEAN = 10,
    IGRAPH_ATTRIBUTE_COMBINE_MEDIAN = 11,
    IGRAPH_ATTRIBUTE_COMBINE_CONCAT = 12
} igraph_attribute_combination_type_t;
```

Values:

IGRAPH_ATTRIBUTE_COMBINE_IGNORE:	Ignore old attributes, use an empty value.
IGRAPH_ATTRIBUTE_COMBINE_DEFAULT:	Use the default way to combine attributes (decided by the attribute handler implementation).
IGRAPH_ATTRIBUTE_COMBINE_FUNCTION:	Supply your own function to combine attributes.
IGRAPH_ATTRIBUTE_COMBINE_SUM:	Take the sum of the attributes.
IGRAPH_ATTRIBUTE_COMBINE_PROD:	Take the product of the attributes.
IGRAPH_ATTRIBUTE_COMBINE_MIN:	Take the minimum attribute.
IGRAPH_ATTRIBUTE_COMBINE_MAX:	Take the maximum attribute.
IGRAPH_ATTRIBUTE_COMBINE_RANDOM:	Take a random attribute.
IGRAPH_ATTRIBUTE_COMBINE_FIRST:	Take the first attribute.
IGRAPH_ATTRIBUTE_COMBINE_LAST:	Take the last attribute.
IGRAPH_ATTRIBUTE_COMBINE_MEAN:	Take the mean of the attributes.
IGRAPH_ATTRIBUTE_COMBINE_MEDIAN:	Take the median of the attributes.
IGRAPH_ATTRIBUTE_COMBINE_CONCAT:	Concatenate the attributes.

Accessing attributes from C

There is an experimental attribute handler that can be used from C code. In this section we show how this works. This attribute handler is by default not attached (the default is no attribute handler), so we first need to attach it:

```
igraph_set_attribute_table(&igraph_cattribute_table);
```

Now the attribute functions are available. Please note that the attribute handler must be attached before you call any other igraph functions, otherwise you might end up with graphs without attributes and an active attribute handler, which might cause unexpected program behaviour. The rule is that you attach the attribute handler in the beginning of your `main()` and never touch it again. (Detaching the attribute handler might lead to memory leaks.)

It is not currently possible to have attribute handlers on a per-graph basis. All graphs in an application must be managed with the same attribute handler. (Including the default case when there is no attribute handler at all.)

The C attribute handler supports attaching real numbers and character strings as attributes. No vectors are allowed, i.e. every vertex might have an attribute called name, but it is not possible to have a coords graph (or other) attribute which is a vector of numbers.

Example 12.1. File `examples/simple/cattributes.c`

Example 12.2. File `examples/simple/cattributes2.c`

Example 12.3. File `examples/simple/cattributes3.c`

Example 12.4. File `examples/simple/cattributes4.c`

Query attributes

`igraph_cattribute_list` — List all attributes

```
int igraph_cattribute_list(const igraph_t *graph,
                          igraph_strvector_t *gnames, igraph_vector_t *gtypes,
                          igraph_strvector_t *vnames, igraph_vector_t *vtypes,
                          igraph_strvector_t *enames, igraph_vector_t *etypes)
```

See `igraph_attribute_type_t` for the various attribute types.

Arguments:

graph: The input graph.

gnames: String vector, the names of the graph attributes.

gtypes: Numeric vector, the types of the graph attributes.

vnames: String vector, the names of the vertex attributes.

vtypes: Numeric vector, the types of the vertex attributes.

enames: String vector, the names of the edge attributes.

etypes: Numeric vector, the types of the edge attributes.

Returns:

Error code.

Naturally, the string vector with the attribute names and the numeric vector with the attribute types are in the right order, i.e. the first name corresponds to the first type, etc. Time complexity: $O(A_g + A_v + A_e)$, the number of all attributes.

`igraph_cattribute_has_attr` — Checks whether a (graph, vertex or edge) attribute exists

```
igraph_bool_t igraph_cattribute_has_attr(const igraph_t *graph,
                                         igraph_attribute_element_t type,
                                         const char *name);
```

Arguments:

graph: The graph.

type: The type of the attribute, IGRAPH_ATTRIBUTE_GRAPH, IGRAPH_ATTRIBUTE_VERTEX or IGRAPH_ATTRIBUTE_EDGE.

name: Character constant, the name of the attribute.

Returns:

Logical value, TRUE if the attribute exists, FALSE otherwise.

Time complexity: $O(A)$, the number of (graph, vertex or edge) attributes, assuming attribute names are not too long.

igraph_cattribute_GAN — Query a numeric graph attribute.

```
igraph_real_t igraph_cattribute_GAN(const igraph_t *graph, const char *name);
```

Returns the value of the given numeric graph attribute. The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.

name: The name of the attribute to query.

Returns:

The value of the attribute.

See also:

GAN for a simpler interface.

Time complexity: $O(A_g)$, the number of graph attributes.

GAN — Query a numeric graph attribute.

```
#define GAN(graph,n)
```

This is shorthand for `igraph_cattribute_GAN()`.

Arguments:

graph: The graph.

n: The name of the attribute.

Returns:

The value of the attribute.

igraph_cattribute_GAB — Query a boolean graph attribute.

```
igraph_bool_t igraph_cattribute_GAB(const igraph_t *graph, const char *name);
```

Returns the value of the given numeric graph attribute. The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.

name: The name of the attribute to query.

Returns:

The value of the attribute.

See also:

GAB for a simpler interface.

Time complexity: $O(A_g)$, the number of graph attributes.

GAB — Query a boolean graph attribute.

```
#define GAB(graph,n)
```

This is shorthand for `igraph_cattribute_GAB()`.

Arguments:

graph: The graph.

n: The name of the attribute.

Returns:

The value of the attribute.

igraph_cattribute_GAS — Query a string graph attribute.

```
const char* igraph_cattribute_GAS(const igraph_t *graph, const char *name);
```

Returns a const pointer to the string graph attribute specified in *name*. The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.

name: The name of the attribute to query.

Returns:

The value of the attribute.

See also:

GAS for a simpler interface.

Time complexity: $O(A_g)$, the number of graph attributes.

GAS — Query a string graph attribute.

```
#define GAS(graph,n)
```

This is shorthand for `igraph_cattribute_GAS()`.

Arguments:

graph: The graph.

n: The name of the attribute.

Returns:

The value of the attribute.

igraph_cattribute_VAN — Query a numeric vertex attribute.

```
igraph_real_t igraph_cattribute_VAN(const igraph_t *graph, const char *name,  
                                     igraph_integer_t vid);
```

The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.

name: The name of the attribute.

vid: The id of the queried vertex.

Returns:

The value of the attribute.

See also:

VAN macro for a simpler interface.

Time complexity: $O(A_v)$, the number of vertex attributes.

VAN — Query a numeric vertex attribute.

```
#define VAN(graph,n,v)
```

This is shorthand for `igraph_cattribute_VAN()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: The id of the vertex.

Returns:

The value of the attribute.

igraph_cattribute_VANV — Query a numeric vertex attribute for many vertices

```
int igraph_cattribute_VANV(const igraph_t *graph, const char *name,  
                           igraph_vs_t vids, igraph_vector_t *result);
```

Arguments:

graph: The input graph.

name: The name of the attribute.

vids: The vertices to query.

result: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: $O(v)$, where v is the number of vertices in 'vids'.

VANV — Query a numeric vertex attribute for all vertices.

```
#define VANV(graph,n,vec)
```

This is a shorthand for `igraph_cattribute_VANV()`.

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_VAB — Query a boolean vertex attribute.

```
igraph_bool_t igraph_cattribute_VAB(const igraph_t *graph, const char *name,
                                     igraph_integer_t vid);
```

The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.

name: The name of the attribute.

vid: The id of the queried vertex.

Returns:

The value of the attribute.

See also:

VAB macro for a simpler interface.

Time complexity: $O(A_v)$, the number of vertex attributes.

VAB — Query a boolean vertex attribute.

```
#define VAB(graph,n,v)
```

This is shorthand for `igraph_cattribute_VAB()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: The id of the vertex.

Returns:

The value of the attribute.

igraph_cattribute_VABV — Query a boolean vertex attribute for many vertices

```
int igraph_cattribute_VABV(const igraph_t *graph, const char *name,
                           igraph_vs_t vids, igraph_vector_bool_t *result);
```

Arguments:

graph: The input graph.

name: The name of the attribute.

vids: The vertices to query.

result: Pointer to an initialized boolean vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: $O(v)$, where v is the number of vertices in 'vids'.

VABV — Query a boolean vertex attribute for all vertices.

```
#define VABV(graph,n,vec)
```

This is a shorthand for `igraph_cattribute_VABV()`.

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized boolean vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_VAS — Query a string vertex attribute.

```
const char* igraph_cattribute_VAS(const igraph_t *graph, const char *name,
                                   igraph_integer_t vid);
```

The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.

name: The name of the attribute.

vid: The id of the queried vertex.

Returns:

The value of the attribute.

See also:

The macro `VAS` for a simpler interface.

Time complexity: $O(A_v)$, the number of vertex attributes.

VAS — Query a string vertex attribute.

```
#define VAS(graph,n,v)
```

This is shorthand for `igraph_cattribute_VAS()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: The id of the vertex.

Returns:

The value of the attribute.

igraph_cattribute_VASV — Query a string vertex attribute for many vertices

```
int igraph_cattribute_VASV(const igraph_t *graph, const char *name,  
                           igraph_vs_t vids, igraph_strvector_t *result);
```

Arguments:

graph: The input graph.

name: The name of the attribute.

vids: The vertices to query.

result: Pointer to an initialized string vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: $O(v)$, where v is the number of vertices in 'vids'. (We assume that the string attributes have a bounded length.)

VASV — Query a string vertex attribute for all vertices.

```
#define VASV(graph,n,vec)
```

This is a shorthand for `igraph_cattribute_VASV()`.

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized string vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_EAN — Query a numeric edge attribute.

```
igraph_real_t igraph_cattribute_EAN(const igraph_t *graph, const char *name,  
                                     igraph_integer_t eid);
```

The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.

name: The name of the attribute.

eid: The id of the queried edge.

Returns:

The value of the attribute.

See also:

EAN for an easier interface.

Time complexity: $O(A_e)$, the number of edge attributes.

EAN — Query a numeric edge attribute.

```
#define EAN(graph,n,e)
```

This is shorthand for `igraph_cattribute_EAN()`.

Arguments:

graph: The graph.
n: The name of the attribute.
e: The id of the edge.

Returns:

The value of the attribute.

igraph_cattribute_EANV — Query a numeric edge attribute for many edges

```
int igraph_cattribute_EANV(const igraph_t *graph, const char *name,  
                           igraph_es_t eids, igraph_vector_t *result);
```

Arguments:

graph: The input graph.
name: The name of the attribute.
eids: The edges to query.
result: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: $O(e)$, where e is the number of edges in 'eids'.

EANV — Query a numeric edge attribute for all edges.

```
#define EANV(graph,n,vec)
```

This is a shorthand for `igraph_cattribute_EANV()`.

Arguments:

graph: The graph.
n: The name of the attribute.
vec: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_EAB — Query a boolean edge attribute.

```
igraph_bool_t igraph_cattribute_EAB(const igraph_t *graph, const char *name,
                                     igraph_integer_t eid);
```

The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.
name: The name of the attribute.
eid: The id of the queried edge.

Returns:

The value of the attribute.

See also:

EAB for an easier interface.

Time complexity: $O(A_e)$, the number of edge attributes.

EAB — Query a boolean edge attribute.

```
#define EAB(graph,n,e)
```

This is shorthand for `igraph_cattribute_EAB()`.

Arguments:

graph: The graph.
n: The name of the attribute.
e: The id of the edge.

Returns:

The value of the attribute.

igraph_cattribute_EABV — Query a boolean edge attribute for many edges

```
int igraph_cattribute_EABV(const igraph_t *graph, const char *name,
                           igraph_es_t eids, igraph_vector_bool_t *result);
```

Arguments:

graph: The input graph.

name: The name of the attribute.

eids: The edges to query.

result: Pointer to an initialized boolean vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: $O(e)$, where e is the number of edges in 'eids'.

EABV — Query a boolean edge attribute for all edges.

```
#define EABV(graph,n,vec)
```

This is a shorthand for `igraph_cattribute_EABV()`.

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_EAS — Query a string edge attribute.

```
const char* igraph_cattribute_EAS(const igraph_t *graph, const char *name,  
                                igraph_integer_t eid);
```

The attribute must exist, otherwise an error is triggered.

Arguments:

graph: The input graph.

name: The name of the attribute.

eid: The id of the queried edge.

Returns:

The value of the attribute.

\se EAS if you want to type less. Time complexity: $O(Ae)$, the number of edge attributes.

EAS — Query a string edge attribute.

```
#define EAS(graph,n,e)
```

This is shorthand for `igraph_cattribute_EAS()`.

Arguments:

graph: The graph.

n: The name of the attribute.

e: The id of the edge.

Returns:

The value of the attribute.

igraph_cattribute_EASV — Query a string edge attribute for many edges

```
int igraph_cattribute_EASV(const igraph_t *graph, const char *name,  
                           igraph_es_t eids, igraph_strvector_t *result);
```

Arguments:

graph: The input graph.

name: The name of the attribute.

vids: The edges to query.

result: Pointer to an initialized string vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: $O(e)$, where e is the number of edges in 'eids'. (We assume that the string attributes have a bounded length.)

EASV — Query a string edge attribute for all edges.

```
#define EASV(graph,n,vec)
```

This is a shorthand for `igraph_cattribute_EASV()`.

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized string vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Set attributes

igraph_cattribute_GAN_set — Set a numeric graph attribute

```
int igraph_cattribute_GAN_set(igraph_t *graph, const char *name,
                              igraph_real_t value);
```

Arguments:

graph: The graph.

name: Name of the graph attribute. If there is no such attribute yet, then it will be added.

value: The (new) value of the graph attribute.

Returns:

Error code.

\se SETGAN if you want to type less. Time complexity: O(1).

SETGAN — Set a numeric graph attribute

```
#define SETGAN(graph,n,value)
```

This is a shorthand for `igraph_cattribute_GAN_set()`.

Arguments:

graph: The graph.

n: The name of the attribute.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_GAB_set — Set a boolean graph attribute

```
int igraph_cattribute_GAB_set(igraph_t *graph, const char *name,
                              igraph_bool_t value);
```

Arguments:

graph: The graph.

name: Name of the graph attribute. If there is no such attribute yet, then it will be added.

value: The (new) value of the graph attribute.

Returns:

Error code.

\se SETGAN if you want to type less. Time complexity: O(1).

SETGAB — Set a boolean graph attribute

```
#define SETGAB(graph,n,value)
```

This is a shorthand for `igraph_cattribute_GAB_set()`.

Arguments:

graph: The graph.

n: The name of the attribute.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_GAS_set — Set a string graph attribute.

```
int igraph_cattribute_GAS_set(igraph_t *graph, const char *name,
                              const char *value);
```

Arguments:

graph: The graph.

name: Name of the graph attribute. If there is no such attribute yet, then it will be added.

value: The (new) value of the graph attribute. It will be copied.

Returns:

Error code.

\se SETGAS if you want to type less. Time complexity: O(1).

SETGAS — Set a string graph attribute

```
#define SETGAS(graph,n,value)
```

This is a shorthand for `igraph_cattribute_GAS_set()`.

Arguments:

graph: The graph.

n: The name of the attribute.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_VAN_set — Set a numeric vertex attribute

```
int igraph_cattribute_VAN_set(igraph_t *graph, const char *name,  
                              igraph_integer_t vid, igraph_real_t value);
```

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all vertices included in *vid*.

Arguments:

graph: The graph.

name: Name of the attribute.

vid: Vertices for which to set the attribute.

value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETVAN for a simpler way.

Time complexity: $O(n)$, the number of vertices if the attribute is new, $O(|vid|)$ otherwise.

SETVAN — Set a numeric vertex attribute

```
#define SETVAN(graph,n,vid,value)
```

This is a shorthand for `igraph_cattribute_VAN_set()`.

Arguments:

graph: The graph.
n: The name of the attribute.
vid: Ids of the vertices to set.
value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_VAB_set — Set a boolean vertex attribute

```
int igraph_cattribute_VAB_set(igraph_t *graph, const char *name,  
                              igraph_integer_t vid, igraph_bool_t value);
```

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all vertices included in *vid*.

Arguments:

graph: The graph.
name: Name of the attribute.
vid: Vertices for which to set the attribute.
value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETVAB for a simpler way.

Time complexity: $O(n)$, the number of vertices if the attribute is new, $O(|vid|)$ otherwise.

SETVAB — Set a boolean vertex attribute

```
#define SETVAB(graph,n,vid,value)
```

This is a shorthand for `igraph_cattribute_VAB_set()`.

Arguments:

graph: The graph.
n: The name of the attribute.
vid: Ids of the vertices to set.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_VAS_set — Set a string vertex attribute

```
int igraph_cattribute_VAS_set(igraph_t *graph, const char *name,
                              igraph_integer_t vid, const char *value);
```

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all vertices included in *vid*.

Arguments:

graph: The graph.

name: Name of the attribute.

vid: Vertices for which to set the attribute.

value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETVAS for a simpler way.

Time complexity: $O(n \cdot l)$, n is the number of vertices, l is the length of the string to set. If the attribute is not new then only $O(|vid| \cdot l)$.

SETVAS — Set a string vertex attribute

```
#define SETVAS(graph,n,vid,value)
```

This is a shorthand for `igraph_cattribute_VAS_set()`.

Arguments:

graph: The graph.

n: The name of the attribute.

vid: Ids of the vertices to set.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_EAN_set — Set a numeric edge attribute

```
int igraph_cattribute_EAN_set(igraph_t *graph, const char *name,
                              igraph_integer_t eid, igraph_real_t value);
```

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all edges included in *eid*.

Arguments:

graph: The graph.
name: Name of the attribute.
eid: Edges for which to set the attribute.
value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETEAN for a simpler way.

Time complexity: $O(e)$, the number of edges if the attribute is new, $O(|eid|)$ otherwise.

SETEAN — Set a numeric edge attribute

```
#define SETEAN(graph,n,eid,value)
```

This is a shorthand for `igraph_cattribute_EAN_set()`.

Arguments:

graph: The graph.
n: The name of the attribute.
eid: Ids of the edges to set.
value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_EAB_set — Set a boolean edge attribute

```
int igraph_cattribute_EAB_set(igraph_t *graph, const char *name,
```

```
igraph_integer_t eid, igraph_bool_t value);
```

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all edges included in *vid*.

Arguments:

graph: The graph.
name: Name of the attribute.
eid: Edges for which to set the attribute.
value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETEAB for a simpler way.

Time complexity: $O(e)$, the number of edges if the attribute is new, $O(|eid|)$ otherwise.

SETEAB — Set a boolean edge attribute

```
#define SETEAB(graph,n,eid,value)
```

This is a shorthand for `igraph_cattribute_EAB_set()`.

Arguments:

graph: The graph.
n: The name of the attribute.
eid: Ids of the edges to set.
value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_EAS_set — Set a string edge attribute

```
int igraph_cattribute_EAS_set(igraph_t *graph, const char *name,  
                              igraph_integer_t eid, const char *value);
```

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all edges included in *vid*.

Arguments:

graph: The graph.
name: Name of the attribute.
eid: Edges for which to set the attribute.
value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETEAS for a simpler way.

Time complexity: $O(e \cdot l)$, n is the number of edges, l is the length of the string to set. If the attribute is not new then only $O(|eid| \cdot l)$.

SETEAS — Set a string edge attribute

```
#define SETEAS(graph,n,eid,value)
```

This is a shorthand for `igraph_cattribute_EAS_set()`.

Arguments:

graph: The graph.
n: The name of the attribute.
eid: Ids of the edges to set.
value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_VAN_setv — Set a numeric vertex attribute for all vertices.

```
int igraph_cattribute_VAN_setv(igraph_t *graph, const char *name,  
                               const igraph_vector_t *v);
```

The attribute will be added if not present yet.

Arguments:

graph: The graph.
name: Name of the attribute.
v: The new attribute values. The length of this vector must match the number of vertices.

Returns:

Error code.

See also:

SETVANV for a simpler way.

Time complexity: $O(n)$, the number of vertices.

SETVANV — Set a numeric vertex attribute for all vertices

```
#define SETVANV(graph,n,v)
```

This is a shorthand for `igraph_cattribute_VAN_setv()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

Returns:

Error code.

igraph_cattribute_VAB_setv — Set a boolean vertex attribute for all vertices.

```
int igraph_cattribute_VAB_setv(igraph_t *graph, const char *name,  
                               const igraph_vector_bool_t *v);
```

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

v: The new attribute values. The length of this boolean vector must match the number of vertices.

Returns:

Error code.

See also:

SETVANV for a simpler way.

Time complexity: $O(n)$, the number of vertices.

SETVABV — Set a boolean vertex attribute for all vertices

```
#define SETVABV(graph,n,v)
```

This is a shorthand for `igraph_cattribute_VAB_setv()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

Returns:

Error code.

igraph_cattribute_VAS_setv — Set a string vertex attribute for all vertices.

```
int igraph_cattribute_VAS_setv(igraph_t *graph, const char *name,  
                               const igraph_strvector_t *sv);
```

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

sv: String vector, the new attribute values. The length of this vector must match the number of vertices.

Returns:

Error code.

See also:

SETVASV for a simpler way.

Time complexity: $O(n+l)$, n is the number of vertices, l is the total length of the strings.

SETVASV — Set a string vertex attribute for all vertices

```
#define SETVASV(graph,n,v)
```

This is a shorthand for `igraph_cattribute_VAS_setv()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

Returns:

Error code.

`igraph_cattribute_EAN_setv` — Set a numeric edge attribute for all edges.

```
int igraph_cattribute_EAN_setv(igraph_t *graph, const char *name,
                               const igraph_vector_t *v);
```

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

v: The new attribute values. The length of this vector must match the number of edges.

Returns:

Error code.

See also:

`SETEANV` for a simpler way.

Time complexity: $O(e)$, the number of edges.

`SETEANV` — Set a numeric edge attribute for all edges

```
#define SETEANV(graph,n,v)
```

This is a shorthand for `igraph_cattribute_EAN_setv()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

igraph_cattribute_EAB_setv — Set a boolean edge attribute for all edges.

```
int igraph_cattribute_EAB_setv(igraph_t *graph, const char *name,
                               const igraph_vector_bool_t *v);
```

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

v: The new attribute values. The length of this vector must match the number of edges.

Returns:

Error code.

See also:

SETEABV for a simpler way.

Time complexity: $O(e)$, the number of edges.

SETEABV — Set a boolean edge attribute for all edges

```
#define SETEABV(graph,n,v)
```

This is a shorthand for `igraph_cattribute_EAB_setv()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

igraph_cattribute_EAS_setv — Set a string edge attribute for all edges.

```
int igraph_cattribute_EAS_setv(igraph_t *graph, const char *name,
                               const igraph_strvector_t *sv);
```

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

sv: String vector, the new attribute values. The length of this vector must match the number of edges.

Returns:

Error code.

See also:

SETEASV for a simpler way.

Time complexity: $O(e+l)$, e is the number of edges, l is the total length of the strings.

SETEASV — Set a string edge attribute for all edges

```
#define SETEASV(graph,n,v)
```

This is a shorthand for `igraph_cattribute_EAS_setv()`.

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

Remove attributes

igraph_cattribute_remove_g — Remove a graph attribute

```
void igraph_cattribute_remove_g(igraph_t *graph, const char *name);
```

Arguments:

graph: The graph object.

name: Name of the graph attribute to remove.

See also:

DELGA for a simpler way.

DELGA — Remove a graph attribute.

```
#define DELGA(graph,n)
```

A shorthand for `igraph_cattribute_remove_g()`.

Arguments:

graph: The graph.

n: The name of the attribute to remove.

igraph_cattribute_remove_v — Remove a vertex attribute

```
void igraph_cattribute_remove_v(igraph_t *graph, const char *name);
```

Arguments:

graph: The graph object.

name: Name of the vertex attribute to remove.

See also:

DELVA for a simpler way.

DELVA — Remove a vertex attribute.

```
#define DELVA(graph,n)
```

A shorthand for `igraph_cattribute_remove_v()`.

Arguments:

graph: The graph.

n: The name of the attribute to remove.

igraph_cattribute_remove_e — Remove an edge attribute

```
void igraph_cattribute_remove_e(igraph_t *graph, const char *name);
```

Arguments:

graph: The graph object.

name: Name of the edge attribute to remove.

See also:

DELEA for a simpler way.

DELEA — Remove an edge attribute.

```
#define DELEA(graph,n)
```

A shorthand for `igraph_cattribute_remove_e()`.

Arguments:

graph: The graph.

n: The name of the attribute to remove.

igraph_cattribute_remove_all — Remove all graph/vertex/edge attributes

```
void igraph_cattribute_remove_all(igraph_t *graph, igraph_bool_t g,  
                                igraph_bool_t v, igraph_bool_t e);
```

Arguments:

graph: The graph object.

g: Boolean, whether to remove graph attributes.

v: Boolean, whether to remove vertex attributes.

e: Boolean, whether to remove edge attributes.

See also:

DELGAS, DELVAS, DELEAS, DELALL for simpler ways.

DELGAS — Remove all graph attributes.

```
#define DELGAS(graph)
```

Calls `igraph_cattribute_remove_all()`.

Arguments:

graph: The graph.

DELVAS — Remove all vertex attributes.

```
#define DELVAS(graph)
```

Calls `igraph_cattribute_remove_all()`.

Arguments:

graph: The graph.

DELEAS — Remove all edge attributes.

```
#define DELEAS(graph)
```

Calls `igraph_cattribute_remove_all()`.

Arguments:

graph: The graph.

DELALL — Remove all attributes.

```
#define DELALL(graph)
```

All graph, vertex and edges attributes will be removed. Calls `igraph_cattribute_remove_all()`.

Arguments:

graph: The graph.

Chapter 13. Structural properties of graphs

These functions usually calculate some structural property of a graph, like its diameter, the degree of the nodes, etc.

Basic properties

igraph_are_connected — Decides whether two vertices are connected

```
int igraph_are_connected(const igraph_t *graph,
                        igraph_integer_t v1, igraph_integer_t v2,
                        igraph_bool_t *res);
```

Arguments:

graph: The graph object.

v1: The first vertex.

v2: The second vertex.

res: Boolean, TRUE if there is an edge from *v1* to *v2*, FALSE otherwise.

Returns:

The error code IGRAPH_EINVVID is returned if an invalid vertex ID is given.

The function is of course symmetric for undirected graphs.

Time complexity: $O(\min(\log(d1), \log(d2)))$, *d1* is the (out-)degree of *v1* and *d2* is the (in-)degree of *v2*.

(Shortest)-path related functions

igraph_shortest_paths — The length of the shortest paths between vertices.

```
int igraph_shortest_paths(const igraph_t *graph, igraph_matrix_t *res,
                        const igraph_vs_t from, const igraph_vs_t to,
                        igraph_neimode_t mode);
```

Arguments:

graph: The graph object.

res: The result of the calculation, a matrix. A pointer to an initialized matrix, to be more precise. The matrix will be resized if needed. It will have the same number of rows as the length of the *from* argument, and its number of columns is the number of vertices in the *to* argument. One row of the matrix shows the distances from/to a given vertex to the ones in *to*. For the unreachable vertices `IGRAPH_INFINITY` is returned.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

mode: The type of shortest paths to be used for the calculation in directed graphs. Possible values:

<code>IGRAPH_OUT</code>	the lengths of the outgoing paths are calculated.
<code>IGRAPH_IN</code>	the lengths of the incoming paths are calculated.
<code>IGRAPH_ALL</code>	the directed graph is considered as an undirected one for the computation.

Returns:

Error code:

`IGRAPH_ENOMEM` not enough memory for temporary data.

`IGRAPH_EINVVID` invalid vertex id passed.

`IGRAPH_EINVMODE` invalid mode argument.

Time complexity: $O(n(|V|+|E|))$, n is the number of vertices to calculate, $|V|$ and $|E|$ are the number of vertices and edges in the graph.

See also:

`igraph_get_shortest_paths()` to get the paths themselves, `igraph_shortest_paths_dijkstra()` for the weighted version.

`igraph_shortest_paths_dijkstra` — Weighted shortest path lengths between vertices.

```
int igraph_shortest_paths_dijkstra(const igraph_t *graph,
                                   igraph_matrix_t *res,
                                   const igraph_vs_t from,
                                   const igraph_vs_t to,
                                   const igraph_vector_t *weights,
                                   igraph_neimode_t mode);
```

This function implements Dijkstra's algorithm to find the weighted shortest path lengths to all vertices from a single source. It is run independently for the given sources. It uses a binary heap for efficient implementation.

Arguments:

graph: The input graph, can be directed.

res: The result, a matrix. A pointer to an initialized matrix should be passed here. The matrix will be resized as needed. Each row contains the distances from a single source, to the

vertices given in the `to` argument. Unreachable vertices has distance `IGRAPH_INFINITY`.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, `igraph_shortest_paths()` is called.

mode: For directed graphs; whether to follow paths along edge directions (`IGRAPH_OUT`), or the opposite (`IGRAPH_IN`), or ignore edge directions completely (`IGRAPH_ALL`). It is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(s*|E|\log|E|+|V|)$, where $|V|$ is the number of vertices, $|E|$ the number of edges and s the number of sources.

See also:

`igraph_shortest_paths()` for a (slightly) faster unweighted version or `igraph_shortest_paths_bellman_ford()` for a weighted variant that works in the presence of negative edge weights (but no negative loops).

Example 13.1. File `examples/simple/dijkstra.c`

`igraph_shortest_paths_bellman_ford` — Weighted shortest path lengths between vertices, allowing negative weights.

```
int igraph_shortest_paths_bellman_ford(const igraph_t *graph,
                                       igraph_matrix_t *res,
                                       const igraph_vs_t from,
                                       const igraph_vs_t to,
                                       const igraph_vector_t *weights,
                                       igraph_neimode_t mode);
```

This function implements the Bellman-Ford algorithm to find the weighted shortest paths to all vertices from a single source, allowing negative weights. It is run independently for the given sources. If there are no negative weights, you are better off with `igraph_shortest_paths_dijkstra()`.

Arguments:

graph: The input graph, can be directed.

res: The result, a matrix. A pointer to an initialized matrix should be passed here, the matrix will be resized if needed. Each row contains the distances from a single source, to all vertices in the graph, in the order of vertex ids. For unreachable vertices the matrix contains `IGRAPH_INFINITY`.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

weights: The edge weights. There mustn't be any closed loop in the graph that has a negative total weight (since this would allow us to decrease the weight of any path containing at least a single vertex of this loop infinitely). Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, `igraph_shortest_paths()` is called.

mode: For directed graphs; whether to follow paths along edge directions (`IGRAPH_OUT`), or the opposite (`IGRAPH_IN`), or ignore edge directions completely (`IGRAPH_ALL`). It is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(s*|E|*|V|)$, where $|V|$ is the number of vertices, $|E|$ the number of edges and s the number of sources.

See also:

`igraph_shortest_paths()` for a faster unweighted version or `igraph_shortest_paths_dijkstra()` if you do not have negative edge weights.

Example 13.2. File `examples/simple/bellman_ford.c`

`igraph_shortest_paths_johnson` — Weighted shortest path lengths between vertices, using Johnson's algorithm.

```
int igraph_shortest_paths_johnson(const igraph_t *graph,
                                igraph_matrix_t *res,
                                const igraph_vs_t from,
                                const igraph_vs_t to,
                                const igraph_vector_t *weights);
```

See Wikipedia at http://en.wikipedia.org/wiki/Johnson's_algorithm for Johnson's algorithm. This algorithm works even if the graph contains negative edge weights, and it is worth using it if we calculate the shortest paths from many sources.

If no edge weights are supplied, then the unweighted version, `igraph_shortest_paths()` is called.

If all the supplied edge weights are non-negative, then Dijkstra's algorithm is used by calling `igraph_shortest_paths_dijkstra()`.

Arguments:

graph: The input graph. If negative weights are present, it should be directed.

res: Pointer to an initialized matrix, the result will be stored here, one line for each source vertex, one column for each target vertex.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

weights: Optional edge weights. If it is a null-pointer, then the unweighted breadth-first search based `igraph_shortest_paths()` will be called.

Returns:

Error code.

Time complexity: $O(s|V|\log|V|+|V||E|)$, $|V|$ and $|E|$ are the number of vertices and edges, s is the number of source vertices.

See also:

`igraph_shortest_paths()` for a faster unweighted version or `igraph_shortest_paths_dijkstra()` if you do not have negative edge weights, `igraph_shortest_paths_bellman_ford()` if you only need to calculate shortest paths from a couple of sources.

igraph_get_shortest_paths — Shortest paths from a vertex.

```
int igraph_get_shortest_paths(const igraph_t *graph,
                             igraph_vector_ptr_t *vertices,
                             igraph_vector_ptr_t *edges,
                             igraph_integer_t from, const igraph_vs_t to,
                             igraph_neimode_t mode,
                             igraph_vector_long_t *predecessors,
                             igraph_vector_long_t *inbound_edges);
```

If there is more than one geodesic between two vertices, this function gives only one of them.

Arguments:

<i>graph</i> :	The graph object.
<i>vertices</i> :	The result, the ids of the vertices along the paths. This is a pointer vector, each element points to a vector object. These should be initialized before passing them to the function, which will properly clear and/or resize them and fill the ids of the vertices along the geodesics from/to the vertices. Supply a null pointer here if you don't need these vectors.
<i>edges</i> :	The result, the ids of the edges along the paths. This is a pointer vector, each element points to a vector object. These should be initialized before passing them to the function, which will properly clear and/or resize them and fill the ids of the vertices along the geodesics from/to the vertices. Supply a null pointer here if you don't need these vectors.
<i>from</i> :	The id of the vertex from/to which the geodesics are calculated.
<i>to</i> :	Vertex sequence with the ids of the vertices to/from which the shortest paths will be calculated. A vertex might be given multiple times.
<i>mode</i> :	The type of shortest paths to be used for the calculation in directed graphs. Possible values: IGRAPH_OUT the outgoing paths are calculated.

IGRAPH_IN the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

predecessors: A pointer to an initialized igraph vector or null. If not null, a vector containing the predecessor of each vertex in the single source shortest path tree is returned here. The predecessor of vertex *i* in the tree is the vertex from which vertex *i* was reached. The predecessor of the start vertex (in the `from` argument) is itself by definition. If the predecessor is -1, it means that the given vertex was not reached from the source during the search. Note that the search terminates if all the vertices in `to` are reached.

inbound_edges: A pointer to an initialized igraph vector or null. If not null, a vector containing the inbound edge of each vertex in the single source shortest path tree is returned here. The inbound edge of vertex *i* in the tree is the edge via which vertex *i* was reached. The start vertex and vertices that were not reached during the search will have -1 in the corresponding entry of the vector. Note that the search terminates if all the vertices in `to` are reached.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVAL *from* is invalid vertex id, or the length of *to* is not the same as the length of *res*.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of vertices, $|E|$ the number of edges in the graph.

See also:

`igraph_shortest_paths()` if you only need the path length but not the paths themselves.

Example	13.3.	File	examples/simple/
<code>igraph_get_shortest_paths.c</code>			

`igraph_get_shortest_path` — Shortest path from one vertex to another one.

```
int igraph_get_shortest_path(const igraph_t *graph,
                             igraph_vector_t *vertices,
                             igraph_vector_t *edges,
                             igraph_integer_t from,
                             igraph_integer_t to,
                             igraph_neimode_t mode);
```

Calculates and returns a single unweighted shortest path from a given vertex to another one. If there are more than one shortest paths between the two vertices, then an arbitrary one is returned.

This function is a wrapper to `igraph_get_shortest_paths()`, for the special case when only one target vertex is considered.

Arguments:

<i>graph</i> :	The input graph, it can be directed or undirected. Directed paths are considered in directed graphs.
<i>vertices</i> :	Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex ids along the path are stored here, including the source and target vertices.
<i>edges</i> :	Pointer to an uninitialized vector or a null pointer. If not a null pointer, then the edge ids along the path are stored here.
<i>from</i> :	The id of the source vertex.
<i>to</i> :	The id of the target vertex.
<i>mode</i> :	A constant specifying how edge directions are considered in directed graphs. Valid modes are: <code>IGRAPH_OUT</code> , follows edge directions; <code>IGRAPH_IN</code> , follows the opposite directions; and <code>IGRAPH_ALL</code> , ignores edge directions. This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges in the graph.

See also:

`igraph_get_shortest_paths()` for the version with more target vertices.

`igraph_get_shortest_paths_dijkstra` — Weighted shortest paths from a vertex.

```
int igraph_get_shortest_paths_dijkstra(const igraph_t *graph,
                                      igraph_vector_ptr_t *vertices,
                                      igraph_vector_ptr_t *edges,
                                      igraph_integer_t from,
                                      igraph_vs_t to,
                                      const igraph_vector_t *weights,
                                      igraph_neimode_t mode,
                                      igraph_vector_long_t *predecessors,
                                      igraph_vector_long_t *inbound_edges);
```

If there is more than one path with the smallest weight between two vertices, this function gives only one of them.

Arguments:

<i>graph</i> :	The graph object.
<i>vertices</i> :	The result, the ids of the vertices along the paths. This is a pointer vector, each element points to a vector object. These should be initialized before passing them to the function, which will properly clear and/or resize them and fill the ids of the vertices along the geodesics from/to the vertices. Supply a null pointer here if you don't need these vectors. Normally, either this argument, or the

	edges should be non-null, but no error or warning is given if they are both null pointers.
<i>edges</i> :	The result, the ids of the edges along the paths. This is a pointer vector, each element points to a vector object. These should be initialized before passing them to the function, which will properly clear and/or resize them and fill the ids of the vertices along the geodesics from/to the vertices. Supply a null pointer here if you don't need these vectors. Normally, either this argument, or the <i>vertices</i> should be non-null, but no error or warning is given if they are both null pointers.
<i>from</i> :	The id of the vertex from/to which the geodesics are calculated.
<i>to</i> :	Vertex sequence with the ids of the vertices to/from which the shortest paths will be calculated. A vertex might be given multiple times. *
<i>weights</i> :	The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, <code>igraph_get_shortest_paths()</code> is called.
<i>mode</i> :	The type of shortest paths to be use for the calculation in directed graphs. Possible values: <div style="margin-left: 20px;"> <p><code>IGRAPH_OUT</code> the outgoing paths are calculated.</p> <p><code>IGRAPH_IN</code> the incoming paths are calculated.</p> <p><code>IGRAPH_ALL</code> the directed graph is considered as an undirected one for the computation.</p> </div>
<i>predecessors</i> :	A pointer to an initialized igraph vector or null. If not null, a vector containing the predecessor of each vertex in the single source shortest path tree is returned here. The predecessor of vertex <i>i</i> in the tree is the vertex from which vertex <i>i</i> was reached. The predecessor of the start vertex (in the <i>from</i> argument) is itself by definition. If the predecessor is -1, it means that the given vertex was not reached from the source during the search. Note that the search terminates if all the vertices in <i>to</i> are reached.
<i>inbound_edges</i> :	A pointer to an initialized igraph vector or null. If not null, a vector containing the inbound edge of each vertex in the single source shortest path tree is returned here. The inbound edge of vertex <i>i</i> in the tree is the edge via which vertex <i>i</i> was reached. The start vertex and vertices that were not reached during the search will have -1 in the corresponding entry of the vector. Note that the search terminates if all the vertices in <i>to</i> are reached.

Returns:

Error code:

<code>IGRAPH_ENOMEM</code>	not enough memory for temporary data.
<code>IGRAPH_EINVVID</code>	<i>from</i> is invalid vertex id, or the length of <i>to</i> is not the same as the length of <i>res</i> .
<code>IGRAPH_EINVMODE</code>	invalid mode argument.

Time complexity: $O(|E|\log|E|+|V|)$, where $|V|$ is the number of vertices and $|E|$ is the number of edges

See also:

`igraph_shortest_paths_dijkstra()` if you only need the path length but not the paths themselves, `igraph_get_shortest_paths()` if all edge weights are equal.

Example **13.4.** **File** **examples/simple/**
`igraph_get_shortest_paths_dijkstra.c`

`igraph_get_shortest_path_dijkstra` — Weighted shortest path from one vertex to another one.

```
int igraph_get_shortest_path_dijkstra(const igraph_t *graph,
                                     igraph_vector_t *vertices,
                                     igraph_vector_t *edges,
                                     igraph_integer_t from,
                                     igraph_integer_t to,
                                     const igraph_vector_t *weights,
                                     igraph_neimode_t mode);
```

Calculates a single (positively) weighted shortest path from a single vertex to another one, using Dijkstra's algorithm.

This function is a special case (and a wrapper) to `igraph_get_shortest_paths_dijkstra()`.

Arguments:

graph: The input graph, it can be directed or undirected.

vertices: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex ids along the path are stored here, including the source and target vertices.

edges: Pointer to an uninitialized vector or a null pointer. If not a null pointer, then the edge ids along the path are stored here.

from: The id of the source vertex.

to: The id of the target vertex.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, `igraph_get_shortest_paths()` is called.

mode: A constant specifying how edge directions are considered in directed graphs. `IGRAPH_OUT` follows edge directions, `IGRAPH_IN` follows the opposite directions, and `IGRAPH_ALL` ignores edge directions. This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(|E|\log|E|+|V|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges in the graph.

See also:

`igraph_get_shortest_paths_dijkstra()` for the version with more target vertices.

igraph_get_shortest_paths_bellman_ford — Weighted shortest paths from a vertex, allowing negative weights.

```
int igraph_get_shortest_paths_bellman_ford(const igraph_t *graph,
                                           igraph_vector_ptr_t *vertices,
                                           igraph_vector_ptr_t *edges,
                                           igraph_integer_t from,
                                           igraph_vs_t to,
                                           const igraph_vector_t *weights,
                                           igraph_neimode_t mode,
                                           igraph_vector_long_t *predecessors,
                                           igraph_vector_long_t *inbound_edges);
```

This function calculates weighted shortest paths from or to a single vertex, and allows negative weights. When there is more than one shortest path between two vertices, only one of them is returned. If there are no negative weights, you are better off with `igraph_get_shortest_paths_dijkstra()`.

Arguments:

<i>graph</i> :	The input graph, can be directed.
<i>vertices</i> :	The result, the ids of the vertices along the paths. This is a pointer vector, each element points to a vector object. These should be initialized before passing them to the function, which will properly clear and/or resize them and fill the ids of the vertices along the geodesics from/to the vertices. Supply a null pointer here if you don't need these vectors. Normally, either this argument, or the <i>edges</i> should be non-null, but no error or warning is given if they are both null pointers.
<i>edges</i> :	The result, the ids of the edges along the paths. This is a pointer vector, each element points to a vector object. These should be initialized before passing them to the function, which will properly clear and/or resize them and fill the ids of the vertices along the geodesics from/to the vertices. Supply a null pointer here if you don't need these vectors. Normally, either this argument, or the <i>vertices</i> should be non-null, but no error or warning is given if they are both null pointers.
<i>from</i> :	The id of the vertex from/to which the geodesics are calculated.
<i>to</i> :	Vertex sequence with the ids of the vertices to/from which the shortest paths will be calculated. A vertex might be given multiple times.
<i>weights</i> :	The edge weights. There mustn't be any closed loop in the graph that has a negative total weight (since this would allow us to decrease the weight of any path containing at least a single vertex of this loop infinitely). If this is a null pointer, then the unweighted version, <code>igraph_shortest_paths()</code> is called.
<i>mode</i> :	For directed graphs; whether to follow paths along edge directions (<code>IGRAPH_OUT</code>), or the opposite (<code>IGRAPH_IN</code>), or ignore edge directions completely (<code>IGRAPH_ALL</code>). It is ignored for undirected graphs.
<i>predecessors</i> :	A pointer to an initialized igraph vector or null. If not null, a vector containing the predecessor of each vertex in the single source shortest path tree is returned

inbound_edges: A pointer to an initialized igraph vector or null. If not null, a vector containing the inbound edge of each vertex in the single source shortest path tree is returned here. The inbound edge of vertex *i* in the tree is the edge via which vertex *i* was reached. The start vertex and vertices that were not reached during the search will have -1 in the corresponding entry of the vector. Note that the search terminates if all the vertices in *to* are reached.

Error code:

Time complexity: $O(|E| \cdot |V|)$, where $|V|$ is the number of vertices, $|E|$ the number of edges.

`igraph_shortest_paths()` for a faster unweighted version or `igraph_shortest_paths_dijkstra()` if you do not have negative edge weights.

Calculates a single (positively) weighted shortest path from a single vertex to another one, using Bellman-Ford algorithm.

Arguments:

vertices: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex ids along the path are stored here, including the source and target vertices.

edges: Pointer to an uninitialized vector or a null pointer. If not a null pointer, then the edge ids along the path are stored here.

from: The id of the source vertex.

to: The id of the target vertex.

weights: The edge weights. There mustn't be any closed loop in the graph that has a negative total weight (since this would allow us to decrease the weight of any path containing at least a single vertex of this loop infinitely). If this is a null pointer, then the unweighted version is called.

mode: A constant specifying how edge directions are considered in directed graphs. IGRAPH_OUT follows edge directions, IGRAPH_IN follows the opposite directions, and IGRAPH_ALL ignores edge directions. This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(|E|\log|E|+|V|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges in the graph.

See also:

`igraph_get_shortest_paths_bellman_ford()` for the version with more target vertices.

igraph_get_all_shortest_paths — All shortest paths (geodesics) from a vertex.

```
int igraph_get_all_shortest_paths(const igraph_t *graph,
                                igraph_vector_ptr_t *res,
                                igraph_vector_t *nrgeo,
                                igraph_integer_t from, const igraph_vs_t to,
                                igraph_neimode_t mode);
```

When there is more than one shortest path between two vertices, all of them will be returned.

Arguments:

graph: The graph object.

res: Pointer to an initialized pointer vector, the result will be stored here in `igraph_vector_t` objects. Each vector object contains the vertices along a shortest path from *from* to another vertex. The vectors are ordered according to their target vertex: first the shortest paths to vertex 0, then to vertex 1, etc. No data is included for unreachable vertices.

nrgeo: Pointer to an initialized `igraph_vector_t` object or NULL. If not NULL the number of shortest paths from *from* are stored here for every vertex in the graph. Note that the values will be accurate only for those vertices that are in the target vertex sequence (see *to*), since the search terminates as soon as all the target vertices have been found.

from: The id of the vertex from/to which the geodesics are calculated.

to: Vertex sequence with the ids of the vertices to/from which the shortest paths will be calculated. A vertex might be given multiple times.

mode: The type of shortest paths to be use for the calculation in directed graphs. Possible values:

IGRAPH_OUT	the lengths of the outgoing paths are calculated.
IGRAPH_IN	the lengths of the incoming paths are calculated.
IGRAPH_ALL	the directed graph is considered as an undirected one for the computation.

Returns:

Error code:

IGRAPH_ENOMEM	not enough memory for temporary data.
IGRAPH_EINVVID	<i>from</i> is invalid vertex id.
IGRAPH_EINVMODE	invalid mode argument.

Added in version 0.2.

Time complexity: $O(|V|+|E|)$ for most graphs, $O(|V|^2)$ in the worst case.

igraph_get_all_shortest_paths_dijkstra — All weighted shortest paths (geodesics) from a vertex.

```
int igraph_get_all_shortest_paths_dijkstra(const igraph_t *graph,
    igraph_vector_ptr_t *res,
    igraph_vector_t *nrgeo,
    igraph_integer_t from, igraph_vs_t to,
    const igraph_vector_t *weights,
    igraph_neimode_t mode);
```

Arguments:

graph: The graph object.

res: Pointer to an initialized pointer vector, the result will be stored here in `igraph_vector_t` objects. Each vector object contains the vertices along a shortest path from *from* to another vertex. The vectors are ordered according to their target vertex: first the shortest paths to vertex 0, then to vertex 1, etc. No data is included for unreachable vertices.

nrgeo: Pointer to an initialized `igraph_vector_t` object or NULL. If not NULL the number of shortest paths from *from* are stored here for every vertex in the graph. Note that the values will be accurate only for those vertices that are in the target vertex sequence (see *to*), since the search terminates as soon as all the target vertices have been found.

from: The id of the vertex from/to which the geodesics are calculated.

to: Vertex sequence with the ids of the vertices to/from which the shortest paths will be calculated. A vertex might be given multiple times.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error

is returned. If this is a null pointer, then the unweighted version, `igraph_get_all_shortest_paths()` is called.

mode: The type of shortest paths to be use for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing paths are calculated.

IGRAPH_IN the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVAL *from* is invalid vertex id, or the length of *to* is not the same as the length of *res*.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|E|\log|E|+|V|)$, where $|V|$ is the number of vertices and $|E|$ is the number of edges

See also:

`igraph_shortest_paths_dijkstra()` if you only need the path length but not the paths themselves, `igraph_get_all_shortest_paths()` if all edge weights are equal.

Example 13.5. File `examples/simple/igraph_get_all_shortest_paths_dijkstra.c`

igraph_get_all_simple_paths — List all simple paths from one source.

```
int igraph_get_all_simple_paths(const igraph_t *graph,
                              igraph_vector_int_t *res,
                              igraph_integer_t from,
                              const igraph_vs_t to,
                              igraph_integer_t cutoff,
                              igraph_neimode_t mode);
```

A path is simple if its vertices are unique, i.e. no vertex is visited more than once.

Note that potentially there are exponentially many paths between two vertices of a graph, and you may run out of memory when using this function, if your graph is lattice-like.

This function currently ignored multiple and loop edges.

Arguments:

graph: The input graph.

res: Initialized integer vector, all paths are returned here, separated by -1 markers. The paths are included in arbitrary order, as they are found.

from: The start vertex.

to: The target vertices.

cutoff: Maximum length of path that is considered. If negative, paths of all lengths are considered.

mode: The type of the paths to consider, it is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(n!)$ in the worst case, n is the number of vertices.

igraph_average_path_length — Calculates the average unweighted shortest path length between all vertex pairs.

```
int igraph_average_path_length(const igraph_t *graph,
                              igraph_real_t *res, igraph_real_t *unconn_pairs,
                              igraph_bool_t directed, igraph_bool_t unconn);
```

If no vertex pairs can be included in the calculation, for example because the graph has fewer than two vertices, or if the graph has no edges and `unconn` is set to `TRUE`, `NaN` is returned.

Arguments:

graph: The graph object.

res: Pointer to a real number, this will contain the result.

unconn_pairs: Pointer to a real number. If not a null pointer, the number of ordered vertex pairs where the second vertex is unreachable from the first one will be stored here.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

unconn: What to do if the graph is not connected. If `TRUE`, only those vertex pairs will be included in the calculation between which there is a path. If `FALSE`, `IGRAPH_INFINITY` is returned for disconnected graphs.

Returns:

Error code: `IGRAPH_ENOMEM`, not enough memory for data structures

Time complexity: $O(|V| |E|)$, the number of vertices times the number of edges.

See also:

`igraph_average_path_length_dijkstra()` for the weighted version.

Example	13.6.	File	examples/simple/
igraph_average_path_length.c			

igraph_average_path_length_dijkstra — Calculates the average weighted shortest path length between all vertex pairs.

```
int igraph_average_path_length_dijkstra(const igraph_t *graph,
                                       igraph_real_t *res, igraph_real_t *unconn,
                                       const igraph_vector_t *weights,
                                       igraph_bool_t directed, igraph_bool_t unconn_pairs);
```

If no vertex pairs can be included in the calculation, for example because the graph has fewer than two vertices, or if the graph has no edges and `unconn` is set to `TRUE`, NaN is returned.

All distinct ordered vertex pairs are taken into account.

Arguments:

<i>graph</i> :	The graph object.
<i>res</i> :	Pointer to a real number, this will contain the result.
<i>unconn_pairs</i> :	Pointer to a real number. If not a null pointer, the number of ordered vertex pairs where the second vertex is unreachable from the first one will be stored here.
<i>weights</i> :	The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, <code>igraph_average_path_length()</code> is called.
<i>directed</i> :	Boolean, whether to consider directed paths. Ignored for undirected graphs.
<i>unconn</i> :	If <code>TRUE</code> , only those pairs are considered for the calculation between which there is a path. If <code>FALSE</code> , <code>IGRAPH_INFINITY</code> is returned for disconnected graphs.

Returns:

Error code:

`IGRAPH_ENOMEM` not enough memory for data structures

`IGRAPH_EINVAL` invalid weight vector

Time complexity: $O(|V| |E| \log |E| + |V|)$, where $|V|$ is the number of vertices and $|E|$ is the number of edges.

See also:

`igraph_average_path_length()` for a slightly faster unweighted version.

Example 13.7. File `examples/simple/igraph_grg_game.c`

igraph_path_length_hist — Create a histogram of all shortest path lengths.

```
int igraph_path_length_hist(const igraph_t *graph, igraph_vector_t *res,
                           igraph_real_t *unconnected, igraph_bool_t directed)
```

This function calculates a histogram, by calculating the shortest path length between each pair of vertices. For directed graphs both directions might be considered and then every pair of vertices appears twice in the histogram.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the result is stored here. The first (i.e. zeroth) element contains the number of shortest paths of length 1, etc. The supplied vector is resized as needed.

unconnected: Pointer to a real number, the number of pairs for which the second vertex is not reachable from the first is stored here.

directed: Whether to consider directed paths in a directed graph (if not zero). This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(|V||E|)$, the number of vertices times the number of edges.

See also:

`igraph_average_path_length()` and `igraph_shortest_paths()`

igraph_diameter — Calculates the diameter of a graph (longest geodesic).

```
int igraph_diameter(const igraph_t *graph, igraph_real_t *pres,
                   igraph_integer_t *pfrom, igraph_integer_t *pto,
                   igraph_vector_t *path,
                   igraph_bool_t directed, igraph_bool_t unconn);
```

The diameter of a graph is the length of the longest shortest path it has. This function computes both the diameter, as well as the corresponding path. The diameter of the null graph is considered be infinity by convention. If the graph has no vertices, IGRAPH_NAN is returned.

Arguments:

graph: The graph object.

pres: Pointer to a real number, if not NULL then it will contain the diameter (the actual distance).

pfrom: Pointer to an integer, if not NULL it will be set to the source vertex of the diameter path. If the graph has no diameter path, it will be set to -1.

pto: Pointer to an integer, if not NULL it will be set to the target vertex of the diameter path. If the graph has no diameter path, it will be set to -1.

path: Pointer to an initialized vector. If not NULL the actual longest geodesic path will be stored here. The vector will be resized as needed.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

unconn: What to do if the graph is not connected. If TRUE the longest geodesic within a component will be returned, otherwise IGRAPH_INFINITY is returned.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: $O(|V||E|)$, the number of vertices times the number of edges.

See also:

`igraph_diameter_dijkstra()`

Example 13.8. File `examples/simple/igraph_diameter.c`

igraph_diameter_dijkstra — Calculates the weighted diameter of a graph using Dijkstra's algorithm.

```
int igraph_diameter_dijkstra(const igraph_t *graph,
                             const igraph_vector_t *weights,
                             igraph_real_t *pres,
                             igraph_integer_t *pfrom,
                             igraph_integer_t *pto,
                             igraph_vector_t *path,
                             igraph_bool_t directed,
                             igraph_bool_t unconn);
```

This function computes the weighted diameter of a graph. If the graph has no vertices, IGRAPH_NAN is returned.

Arguments:

graph: The input graph, can be directed or undirected.

pres: Pointer to a real number, if not NULL then it will contain the diameter (the actual distance).

pfrom: Pointer to an integer, if not NULL it will be set to the source vertex of the diameter path. If the graph has no diameter path, it will be set to -1.

pto: Pointer to an integer, if not NULL it will be set to the target vertex of the diameter path. If the graph has no diameter path, it will be set to -1.

path: Pointer to an initialized vector. If not NULL the actual longest geodesic path will be stored here. The vector will be resized as needed.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

unconn: What to do if the graph is not connected. If TRUE the longest geodesic within a component will be returned, otherwise IGRAPH_INFINITY is returned.

Returns:

Error code.

Time complexity: $O(|V||E|\log|E|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges.

See also:

`igraph_diameter()`

igraph_girth — The girth of a graph is the length of the shortest cycle in it.

```
int igraph_girth(const igraph_t *graph, igraph_integer_t *girth,
                igraph_vector_t *circle);
```

The current implementation works for undirected graphs only, directed graphs are treated as undirected graphs. Self-loops and multiple edges are ignored.

For graphs that contain no cycles, and only for such graphs, zero is returned. Note that in some applications, it is customary to define the girth of acyclic graphs to be infinity. However, infinity is not representable as an `igraph_integer_t`, therefore zero is used for this case.

This implementation is based on Alon Itai and Michael Rodeh: Finding a minimum circuit in a graph *Proceedings of the ninth annual ACM symposium on Theory of computing*, 1-10, 1977. The first implementation of this function was done by Keith Briggs, thanks Keith.

Arguments:

graph: The input graph.

girth: Pointer to an integer, if not NULL then the result will be stored here.

circle: Pointer to an initialized vector, the vertex ids in the shortest circle will be stored here. If NULL then it is ignored.

Returns:

Error code.

Time complexity: $O((|V|+|E|)^2)$, $|V|$ is the number of vertices, $|E|$ is the number of edges in the general case. If the graph has no cycles at all then the function needs $O(|V|+|E|)$ time to realize this and then it stops.

Example 13.9. File `examples/simple/igraph_girth.c`

igraph_eccentricity — Eccentricity of some vertices.

```
int igraph_eccentricity(const igraph_t *graph,
                      igraph_vector_t *res,
```



```
igraph_vs_t vids,  
igraph_neimode_t mode);
```

The eccentricity of a vertex is calculated by measuring the shortest distance from (or to) the vertex, to (or from) all vertices in the graph, and taking the maximum.

This implementation ignores vertex pairs that are in different components. Isolated vertices have eccentricity zero.

Arguments:

graph: The input graph, it can be directed or undirected.

res: Pointer to an initialized vector, the result is stored here.

vids: The vertices for which the eccentricity is calculated.

mode: What kind of paths to consider for the calculation: `IGRAPH_OUT`, paths that follow edge directions; `IGRAPH_IN`, paths that follow the opposite directions; and `IGRAPH_ALL`, paths that ignore edge directions. This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(v*(|V|+|E|))$, where $|V|$ is the number of vertices, $|E|$ is the number of edges and v is the number of vertices for which eccentricity is calculated.

See also:

`igraph_radius()`.

Example 13.10. File `examples/simple/igraph_eccentricity.c`

`igraph_radius` — Radius of a graph.

```
int igraph_radius(const igraph_t *graph, igraph_real_t *radius,  
igraph_neimode_t mode);
```

The radius of a graph is the defined as the minimum eccentricity of its vertices, see `igraph_eccentricity()`.

Arguments:

graph: The input graph, it can be directed or undirected.

radius: Pointer to a real variable, the result is stored here.

mode: What kind of paths to consider for the calculation: `IGRAPH_OUT`, paths that follow edge directions; `IGRAPH_IN`, paths that follow the opposite directions; and `IGRAPH_ALL`, paths that ignore edge directions. This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(|V|(|V|+|E|))$, where $|V|$ is the number of vertices and $|E|$ is the number of edges.

See also:

`igraph_eccentricity()`.

Example 13.11. File `examples/simple/igraph_radius.c`

Efficiency measures

`igraph_global_efficiency` — Calculates the global efficiency of a network.

```
int igraph_global_efficiency(const igraph_t *graph, igraph_real_t *res,
                             const igraph_vector_t *weights,
                             igraph_bool_t directed);
```

The global efficiency of a network is defined as the average of inverse distances between all pairs of vertices: $E_g = 1/(N*(N-1)) \sum_{i \neq j} 1/d_{ij}$, where N is the number of vertices. The inverse distance between pairs that are not reachable from each other is considered to be zero. For graphs with fewer than 2 vertices, NaN is returned.

Reference: V. Latora and M. Marchiori, Efficient Behavior of Small-World Networks, Phys. Rev. Lett. 87, 198701 (2001). <https://dx.doi.org/10.1103/PhysRevLett.87.198701>

Arguments:

graph: The graph object.

res: Pointer to a real number, this will contain the result.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, `igraph_average_path_length()` is used in calculating the global efficiency.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for data structures

IGRAPH_EINVAL invalid weight vector

Time complexity: $O(|V| |E| \log |E| + |V|)$ for weighted graphs and $O(|V| |E|)$ for unweighted ones. $|V|$ denotes the number of vertices and $|E|$ denotes the number of edges.

`igraph_local_efficiency` — Calculates the local efficiency around each vertex in a network.

```
int igraph_local_efficiency(const igraph_t *graph, igraph_vector_t *res,
```

```
const igraph_vs_t vids,  
const igraph_vector_t *weights,  
igraph_bool_t directed, igraph_neimode_t mode);
```

The local efficiency of a network around a vertex is defined as follows: We remove the vertex and compute the distances (shortest path lengths) between its neighbours through the rest of the network. The local efficiency around the removed vertex is the average of the inverse of these distances.

The inverse distance between two vertices which are not reachable from each other is considered to be zero. The local efficiency around a vertex with fewer than two neighbours is taken to be zero by convention.

Reference: I. Vragovi#, E. Louis, and A. Díaz-Guilera, Efficiency of informational transfer in regular and complex networks, Phys. Rev. E 71, 1 (2005). <http://dx.doi.org/10.1103/PhysRevE.71.036122>

Arguments:

- graph*: The graph object.
- res*: Pointer to an initialized vector, this will contain the result.
- vids*: The vertices around which the local efficiency will be calculated.
- weights*: The edge weights. All edge weights must be non-negative. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, `igraph_average_path_length()` is called.
- directed*: Boolean, whether to consider directed paths. Ignored for undirected graphs.
- mode*: How to determine the local neighborhood of each vertex in directed graphs. Ignored in undirected graphs.
- | | |
|-------------------------|---|
| <code>IGRAPH_ALL</code> | take both in- and out-neighbours; this is a reasonable default for high-level interfaces. |
| <code>IGRAPH_OUT</code> | take only out-neighbours |
| <code>IGRAPH_IN</code> | take only in-neighbours |

Returns:

Error code:

`IGRAPH_ENOMEM` not enough memory for data structures

`IGRAPH_EINVAL` invalid weight vector

Time complexity: $O(|E|^2 \log|E|)$ for weighted graphs and $O(|E|^2)$ for unweighted ones. $|E|$ denotes the number of edges.

See also:

`igraph_average_local_efficiency()`

igraph_average_local_efficiency — Calculates the average local efficiency in a network.

```
int igraph_average_local_efficiency(const igraph_t *graph, igraph_real_t *res,
                                   const igraph_vector_t *weights,
                                   igraph_bool_t directed, igraph_neimode_t mo
```

For the null graph, zero is returned by convention.

Arguments:

graph: The graph object.

res: Pointer to a real number, this will contain the result.

weights: The edge weights. They must be all non-negative. If a null pointer is given, all weights are assumed to be 1.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

mode: How to determine the local neighborhood of each vertex in directed graphs. Ignored in undirected graphs.

IGRAPH_ALL	take both in- and out-neighbours; this is a reasonable default for high-level interfaces.
IGRAPH_OUT	take only out-neighbours
IGRAPH_IN	take only in-neighbours

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for data structures

IGRAPH_EINVAL invalid weight vector

Time complexity: $O(|E|^2 \log|E|)$ for weighted graphs and $O(|E|^2)$ for unweighted ones. $|E|$ denotes the number of edges.

See also:

`igraph_local_efficiency()`

Neighborhood of a vertex

`igraph_neighborhood_size` — Calculates the size of the neighborhood of a given vertex.

```
int igraph_neighborhood_size(const igraph_t *graph, igraph_vector_t *res,
                             igraph_vs_t vids, igraph_integer_t order,
                             igraph_neimode_t mode,
                             igraph_integer_t mindist);
```

The neighborhood of a given order of a vertex includes all vertices which are closer to the vertex than the order. I.e., order 0 is always the vertex itself, order 1 is the vertex plus its immediate neighbors, order 2 is order 1 plus the immediate neighbors of the vertices in order 1, etc.

This function calculates the size of the neighborhood of the given order for the given vertices.

Arguments:

- graph*: The input graph.
- res*: Pointer to an initialized vector, the result will be stored here. It will be resized as needed.
- vids*: The vertices for which the calculation is performed.
- order*: Integer giving the order of the neighborhood.
- mode*: Specifies how to use the direction of the edges if a directed graph is analyzed. For `IGRAPH_OUT` only the outgoing edges are followed, so all vertices reachable from the source vertex in at most `order` steps are counted. For `IGRAPH_IN` all vertices from which the source vertex is reachable in at most `order` steps are counted. `IGRAPH_ALL` ignores the direction of the edges. This argument is ignored for undirected graphs.
- mindist*: The minimum distance to include a vertex in the counting. Vertices reachable with a path shorter than this value are excluded. If this is one, then the starting vertex is not counted. If this is two, then its neighbors are not counted either, etc.

Returns:

Error code.

See also:

`igraph_neighborhood()` for calculating the actual neighborhood, `igraph_neighborhood_graphs()` for creating separate graphs from the neighborhoods.

Time complexity: $O(n*d*o)$, where n is the number vertices for which the calculation is performed, d is the average degree, o is the order.

igraph_neighborhood — Calculate the neighborhood of vertices.

```
int igraph_neighborhood(const igraph_t *graph, igraph_vector_ptr_t *res,
                        igraph_vs_t vids, igraph_integer_t order,
                        igraph_neimode_t mode, igraph_integer_t mindist);
```

The neighborhood of a given order of a vertex includes all vertices which are closer to the vertex than the order. I.e., order 0 is always the vertex itself, order 1 is the vertex plus its immediate neighbors, order 2 is order 1 plus the immediate neighbors of the vertices in order 1, etc.

This function calculates the vertices within the neighborhood of the specified vertices.

Arguments:

- graph*: The input graph.
- res*: An initialized pointer vector. Note that the objects (pointers) in the vector will *not* be freed, but the pointer vector will be resized as needed. The result of the calculation will be stored here in `igraph_vector_t` objects.
- vids*: The vertices for which the calculation is performed.

- order*: Integer giving the order of the neighborhood.
- mode*: Specifies how to use the direction of the edges if a directed graph is analyzed. For IGRAPH_OUT only the outgoing edges are followed, so all vertices reachable from the source vertex in at most *order* steps are included. For IGRAPH_IN all vertices from which the source vertex is reachable in at most *order* steps are included. IGRAPH_ALL ignores the direction of the edges. This argument is ignored for undirected graphs.
- mindist*: The minimum distance to include a vertex in the counting. Vertices reachable with a path shorter than this value are excluded. If this is one, then the starting vertex is not counted. If this is two, then its neighbors are not counted either, etc.

Returns:

Error code.

See also:

`igraph_neighborhood_size()` to calculate the size of the neighborhood,
`igraph_neighborhood_graphs()` for creating graphs from the neighborhoods.

Time complexity: $O(n*d*o)$, n is the number of vertices for which the calculation is performed, d is the average degree, o is the order.

igraph_neighborhood_graphs — Create graphs from the neighborhood(s) of some vertex/vertices.

```
int igraph_neighborhood_graphs(const igraph_t *graph, igraph_vector_ptr_t *res,
                              igraph_vs_t vids, igraph_integer_t order,
                              igraph_neimode_t mode,
                              igraph_integer_t mindist);
```

The neighborhood of a given order of a vertex includes all vertices which are closer to the vertex than the order. Ie. order 0 is always the vertex itself, order 1 is the vertex plus its immediate neighbors, order 2 is order 1 plus the immediate neighbors of the vertices in order 1, etc.

This function finds every vertex in the neighborhood of a given parameter vertex and creates the induced subgraph from these vertices.

The first version of this function was written by Vincent Matossian, thanks Vincent.

Arguments:

- graph*: The input graph.
- res*: Pointer to a pointer vector, the result will be stored here, ie. *res* will contain pointers to `igraph_t` objects. It will be resized if needed but note that the objects in the pointer vector will not be freed.
- vids*: The vertices for which the calculation is performed.
- order*: Integer giving the order of the neighborhood.
- mode*: Specifies how to use the direction of the edges if a directed graph is analyzed. For IGRAPH_OUT only the outgoing edges are followed, so all vertices reachable from the source vertex in at most *order* steps are counted. For IGRAPH_IN all vertices from

which the source vertex is reachable in at most *order* steps are counted. IGRAPH_ALL ignores the direction of the edges. This argument is ignored for undirected graphs.

mindist: The minimum distance to include a vertex in the counting. Vertices reachable with a path shorter than this value are excluded. If this is one, then the starting vertex is not counted. If this is two, then its neighbors are not counted either, etc.

Returns:

Error code.

See also:

`igraph_neighborhood_size()` for calculating the neighborhood sizes only,
`igraph_neighborhood()` for calculating the neighborhoods (but not creating graphs).

Time complexity: $O(n*(|V|+|E|))$, where n is the number vertices for which the calculation is performed, $|V|$ and $|E|$ are the number of vertices and edges in the original input graph.

Local scan statistics

The scan statistic is a summary of the locality statistics that is computed from the local neighborhood of each vertex. For details, see Priebe, C. E., Conroy, J. M., Marchette, D. J., Park, Y. (2005). Scan Statistics on Enron Graphs. Computational and Mathematical Organization Theory.

"Us" statistics

`igraph_local_scan_0` — Local scan-statistics, $k=0$

```
int igraph_local_scan_0(const igraph_t *graph, igraph_vector_t *res,
                        const igraph_vector_t *weights,
                        igraph_neimode_t mode);
```

$K=0$ scan-statistics is arbitrarily defined as the vertex degree for unweighted, and the vertex strength for weighted graphs. See `igraph_degree()` and `igraph_strength()`.

Arguments:

graph: The input graph

res: An initialized vector, the results are stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN means incoming and IGRAPH_ALL means all edges.

Returns:

Error code.

`igraph_local_scan_1_ecount` — Local scan-statistics, $k=1$, edge count and sum of weights

```
int igraph_local_scan_1_ecount(const igraph_t *graph, igraph_vector_t *res,
                              const igraph_vector_t *weights,
                              igraph_neimode_t mode);
```

Count the number of edges or the sum the edge weights in the 1-neighborhood of vertices.

Arguments:

graph: The input graph

res: An initialized vector, the results are stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN means incoming and IGRAPH_ALL means all edges.

Returns:

Error code.

igraph_local_scan_k_ecount — Sum the number of edges or the weights in k-neighborhood of every vertex.

```
int igraph_local_scan_k_ecount(const igraph_t *graph, int k,
                              igraph_vector_t *res,
                              const igraph_vector_t *weights,
                              igraph_neimode_t mode);
```

Arguments:

graph: The input graph.

k: The size of the neighborhood, non-negative integer. The k=0 case is special, see `igraph_local_scan_0()`.

res: An initialized vector, the results are stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN means incoming and IGRAPH_ALL means all edges.

Returns:

Error code.

"Them" statistics

igraph_local_scan_0_them — Local THEM scan-statistics, k=0


```
int igraph_local_scan_0_them(const igraph_t *us, const igraph_t *them,
                             igraph_vector_t *res,
                             const igraph_vector_t *weights_them,
                             igraph_neimode_t mode);
```

K=0 scan-statistics is arbitrarily defined as the vertex degree for unweighted, and the vertex strength for weighted graphs. See `igraph_degree()` and `igraph_strength()`.

Arguments:

us: The input graph, to use to extract the neighborhoods.

them: The input graph to use for the actually counting.

res: An initialized vector, the results are stored here.

weights_them: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, `IGRAPH_OUT` means outgoing, `IGRAPH_IN` means incoming and `IGRAPH_ALL` means all edges.

Returns:

Error code.

igraph_local_scan_1_ecount_them — Local THEM scan-statistics, k=1, edge count and sum of weights

```
int igraph_local_scan_1_ecount_them(const igraph_t *us, const igraph_t *them,
                                     igraph_vector_t *res,
                                     const igraph_vector_t *weights_them,
                                     igraph_neimode_t mode);
```

Count the number of edges or the sum the edge weights in the 1-neighborhood of vertices.

Arguments:

us: The input graph to extract the neighborhoods.

them: The input graph to perform the counting.

weights_them: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, `IGRAPH_OUT` means outgoing, `IGRAPH_IN` means incoming and `IGRAPH_ALL` means all edges.

Returns:

Error code.

See also:

`igraph_local_scan_1_ecount()` for the US statistics.

igraph_local_scan_k_ecount_them — Local THEM scan-statistics, general function, edge count and sum of weights

```
int igraph_local_scan_k_ecount_them(const igraph_t *us, const igraph_t *them,
                                     int k, igraph_vector_t *res,
                                     const igraph_vector_t *weights_them,
                                     igraph_neimode_t mode);
```

Count the number of edges or the sum the edge weights in the k-neighborhood of vertices.

Arguments:

us: The input graph to extract the neighborhoods.

them: The input graph to perform the counting.

k: The size of the neighborhood, non-negative integer. The k=0 case is special, see `igraph_local_scan_0_them()`.

weights_them: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, `IGRAPH_OUT` means outgoing, `IGRAPH_IN` means incoming and `IGRAPH_ALL` means all edges.

Returns:

Error code.

See also:

`igraph_local_scan_1_ecount()` for the US statistics.

Pre-calculated neighborhoods

igraph_local_scan_neighborhood_ecount — Local scan-statistics with pre-calculated neighborhoods

```
int igraph_local_scan_neighborhood_ecount(const igraph_t *graph,
                                           igraph_vector_t *res,
                                           const igraph_vector_t *weights,
                                           const igraph_vector_ptr_t *neighborhoods);
```

Count the number of edges, or sum the edge weights in neighborhoods given as a parameter.

Arguments:

graph: The graph to perform the counting/summing in.

res: Initialized vector, the result is stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

neighborhoods: List of `igraph_vector_int_t` objects, the neighborhoods, one for each vertex in the graph.

Returns:

Error code.

Graph components

`igraph_subcomponent` — The vertices in the same component as a given vertex.

```
int igraph_subcomponent(const igraph_t *graph, igraph_vector_t *res, igraph_real_t vertex,
                        igraph_neimode_t mode);
```

Arguments:

graph: The graph object.

res: The result, vector with the ids of the vertices in the same component.

vertex: The id of the vertex of which the component is searched.

mode: Type of the component for directed graphs, possible values:

`IGRAPH_OUT` the set of vertices reachable *from* the *vertex*,

`IGRAPH_IN` the set of vertices from which the *vertex* is reachable.

`IGRAPH_ALL` the graph is considered as an undirected graph. Note that this is *not* the same as the union of the previous two.

Returns:

Error code:

`IGRAPH_ENOMEM` not enough memory for temporary data.

`IGRAPH_EINVVID` *vertex* is an invalid vertex id

`IGRAPH_EINVMODE` invalid mode argument passed.

Time complexity: $O(|V|+|E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the graph.

See also:

`igraph_induced_subgraph()` if you want a graph object consisting only a given set of vertices and the edges between them.

`igraph_clusters` — Calculates the (weakly or strongly) connected components in a graph.

```
int igraph_clusters(const igraph_t *graph, igraph_vector_t *membership,
                   igraph_vector_t *csize, igraph_integer_t *no,
                   igraph_connectedness_t mode);
```

Arguments:

- graph*: The graph object to analyze.
- membership*: First half of the result will be stored here. For every vertex the id of its component is given. The vector has to be preinitialized and will be resized. Alternatively this argument can be NULL, in which case it is ignored.
- csize*: The second half of the result. For every component it gives its size, the order is defined by the component ids. The vector has to be preinitialized and will be resized. Alternatively this argument can be NULL, in which case it is ignored.
- no*: Pointer to an integer, if not NULL then the number of clusters will be stored here.
- mode*: For directed graph this specifies whether to calculate weakly or strongly connected components. Possible values: IGRAPH_WEAK, IGRAPH_STRONG. This argument is ignored for undirected graphs.

Returns:

Error code: IGRAPH_EINVAL: invalid mode argument.

Time complexity: $O(|V|+|E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the graph.

igraph_is_connected — Decides whether the graph is (weakly or strongly) connected.

```
int igraph_is_connected(const igraph_t *graph, igraph_bool_t *res,
                       igraph_connectedness_t mode);
```

A graph with zero vertices (i.e. the null graph) is *not* connected by definition. This behaviour changed in igraph 0.9; earlier versions assumed that the null graph is connected. See the following issue on Github for the argument that led us to change the definition: <https://github.com/igraph/igraph/issues/1538>

Arguments:

- graph*: The graph object to analyze.
- res*: Pointer to a logical variable, the result will be stored here.
- mode*: For a directed graph this specifies whether to calculate weak or strong connectedness. Possible values: IGRAPH_WEAK, IGRAPH_STRONG. This argument is ignored for undirected graphs.

Returns:

Error code: IGRAPH_EINVAL: invalid mode argument.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph.

igraph_decompose — Decompose a graph into connected components.

```
int igraph_decompose(const igraph_t *graph, igraph_vector_ptr_t *components,
                    igraph_connectedness_t mode,
                    long int maxcompno, long int minelements);
```

Create separate graph for each component of a graph. Note that the vertex ids in the new graphs will be different than in the original graph. (Except if there is only one component in the original graph.)

Arguments:

- graph*: The original graph.
- components*: This pointer vector will contain pointers to the subcomponent graphs. It should be initialized before calling this function and will be resized to hold the graphs. Don't forget to call `igraph_destroy()` and `igraph_free()` on the elements of this pointer vector to free unneeded memory. Alternatively, you can simply call `igraph_decompose_destroy()` that does this for you.
- mode*: Either `IGRAPH_WEAK` or `IGRAPH_STRONG` for weakly and strongly connected components respectively.
- maxcompno*: The maximum number of components to return. The first *maxcompno* components will be returned (which hold at least *minelements* vertices, see the next parameter), the others will be ignored. Supply -1 here if you don't want to limit the number of components.
- minelements*: The minimum number of vertices a component should contain in order to place it in the *components* vector. Eg. supply 2 here to ignore isolated vertices.

Returns:

Error code, `IGRAPH_ENOMEM` if there is not enough memory to perform the operation.

Added in version 0.2.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

Example 13.12. File `examples/simple/igraph_decompose.c`

igraph_decompose_destroy — Free the memory allocated by `igraph_decompose()`.

```
void igraph_decompose_destroy(igraph_vector_ptr_t *complist);
```

Arguments:

complist: The list of graph components, as returned by `igraph_decompose()`.

Time complexity: $O(c)$, c is the number of components.

igraph_biconnected_components — Calculate biconnected components

```
int igraph_biconnected_components(const igraph_t *graph,
                                  igraph_integer_t *no,
                                  igraph_vector_ptr_t *tree_edges,
                                  igraph_vector_ptr_t *component_edges,
                                  igraph_vector_ptr_t *components,
                                  igraph_vector_t *articulation_points);
```

A graph is biconnected if the removal of any single vertex (and its incident edges) does not disconnect it.

A biconnected component of a graph is a maximal biconnected subgraph of it. The biconnected components of a graph can be given by the partition of its edges: every edge is a member of exactly one biconnected component. Note that this is not true for vertices: the same vertex can be part of many biconnected components.

Somewhat arbitrarily, igraph does not consider components containing a single vertex only as being biconnected. Isolated vertices will not be part of any of the biconnected components.

Arguments:

<i>graph</i> :	The input graph.
<i>no</i> :	The number of biconnected components will be stored here.
<i>tree_edges</i> :	If not a NULL pointer, then the found components are stored here, in a list of vectors. Every vector in the list is a biconnected component, represented by its edges. More precisely, a spanning tree of the biconnected component is returned. Note you'll have to destroy each vector first by calling <code>igraph_vector_destroy()</code> and then <code>igraph_free()</code> on it, plus you need to call <code>igraph_vector_ptr_destroy()</code> on the list to regain all allocated memory.
<i>component_edges</i> :	If not a NULL pointer, then the edges of the biconnected components are stored here, in the same form as for <i>tree_edges</i> .
<i>components</i> :	If not a NULL pointer, then the vertices of the biconnected components are stored here, in the same format as for the previous two arguments.
<i>articulation_points</i> :	If not a NULL pointer, then the articulation points of the graph are stored in this vector. A vertex is an articulation point if its removal increases the number of (weakly) connected components in the graph.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges, but only if you do not calculate *components* and *component_edges*. If you calculate *components*, then it is quadratic in the number of vertices. If you calculate *component_edges* as well, then it is cubic in the number of vertices.

See also:

```
igraph_articulation_points(), igraph_clusters().
```

Example	13.13.	File	examples/simple/
igraph_biconnected_components.c			

igraph_articulation_points — Find the articulation points in a graph.

```
int igraph_articulation_points(const igraph_t *graph,  
                              igraph_vector_t *res);
```

A vertex is an articulation point if its removal increases the number of connected components in the graph.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the articulation points will be stored here.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

See also:

```
igraph_biconnected_components(), igraph_clusters(), igraph_bridges()
```

igraph_bridges — Find all bridges in a graph.

```
int igraph_bridges(const igraph_t *graph, igraph_vector_t *bridges);
```

An edge is a bridge if its removal increases the number of (weakly) connected components in the graph.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the bridges will be stored here as edge indices.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

See also:

```
igraph_articulation_points(), igraph_biconnected_components(),  
igraph_clusters()
```

Degree sequences

`igraph_is_graphical` — Is there a graph with the given degree sequence?

```
int igraph_is_graphical(const igraph_vector_t *out_degrees,
                        const igraph_vector_t *in_degrees,
                        const igraph_edge_type_sw_t allowed_edge_types,
                        igraph_bool_t *res);
```

Determines whether a sequence of integers can be the degree sequence of some graph. The classical concept of graphicality assumes simple graphs. This function can perform the check also when either self-loops, multi-edge, or both are allowed in the graph.

For simple undirected graphs, the Erdős-Gallai conditions are checked using the linear-time algorithm of Cloteaux. If both self-loops and multi-edges are allowed, it is sufficient to check that the sum of degrees is even. If only multi-edges are allowed, but not self-loops, there is an additional condition that the sum of degrees be no smaller than twice the maximum degree. If at most one self-loop is allowed per vertex, but no multi-edges, a modified version of the Erdős-Gallai conditions are used (see Cairns & Mendan).

For simple directed graphs, the Fulkerson-Chen-Anstee theorem is used with the relaxation by Berger. If both self-loops and multi-edges are allowed, then it is sufficient to check that the sum of in- and out-degrees is the same. If only multi-edges are allowed, but not self-loops, there is an additional condition that the sum of out-degrees (or equivalently, in-degrees) is no smaller than the maximum total degree. If single self-loops are allowed, but not multi-edges, the problem is equivalent to realizability as a simple bipartite graph, thus the Gale-Ryser theorem can be used; see `igraph_is_bigraphical()` for more information.

References:

P. Erdős and T. Gallai, Gráfok elírt fokú pontokkal, Matematikai Lapok 11, pp. 264–274 (1960). https://users.renyi.hu/~p_erdos/1961-05.pdf

Z Király, Recognizing graphic degree sequences and generating all realizations. TR-2011-11, Egerváry Research Group, H-1117, Budapest, Hungary. ISSN 1587-4451 (2012). <http://bolyai.cs.elte.hu/egres/tr/egres-11-11.pdf>

B. Cloteaux, Is This for Real? Fast Graphicality Testing, Comput. Sci. Eng. 17, 91 (2015). <https://dx.doi.org/10.1109/MCSE.2015.125>

A. Berger, A note on the characterization of digraphic sequences, Discrete Math. 314, 38 (2014). <https://dx.doi.org/10.1016/j.disc.2013.09.010>

G. Cairns and S. Mendan, Degree Sequence for Graphs with Loops (2013). <https://arxiv.org/abs/1303.2145v1>

Arguments:

out_degrees: A vector of integers specifying the degree sequence for undirected graphs or the out-degree sequence for directed graphs.

in_degrees: A vector of integers specifying the in-degree sequence for directed graphs. For undirected graphs, it must be NULL.

allowed_edge_types: The types of edges to allow in the graph:

<code>IGRAPH_SIMPLE_SW</code>	simple graphs (i.e. no self-loops or multi-edges allowed).
<code>IGRAPH_LOOPS_SW</code>	single self-loops are allowed, but not multi-edges.
<code>IGRAPH_MULTI_SW</code>	multi-edges are allowed, but not self-loops.
<code>IGRAPH_LOOPS_SW IGRAPH_MULTI_SW</code>	both self-loops and multi-edges are allowed.
<code>res:</code>	Pointer to a Boolean. The result will be stored here.

Returns:

Error code.

See also:

`igraph_is_bigraphical()` to check if a bi-degree-sequence can be realized as a bipartite graph; `igraph_realize_degree_sequence()` to construct a graph with a given degree sequence.

Time complexity: $O(n^2)$ for simple directed graphs, $O(n \log n)$ for graphs with self-loops, and $O(n)$ for all other cases, where n is the length of the degree sequence(s).

igraph_is_bigraphical — Is there a bipartite graph with the given bi-degree-sequence?

```
int igraph_is_bigraphical(const igraph_vector_t *degrees1,
                        const igraph_vector_t *degrees2,
                        const igraph_edge_type_sw_t allowed_edge_types,
                        igraph_bool_t *res);
```

Determines whether two sequences of integers can be the degree sequences of a bipartite graph. Such a pair of degree sequence is called *bigraphical*.

When multi-edges are allowed, it is sufficient to check that the sum of degrees is the same in the two partitions. For simple graphs, the Gale-Ryser theorem is used with Berger's relaxation.

References:

H. J. Ryser, Combinatorial Properties of Matrices of Zeros and Ones, Can. J. Math. 9, 371 (1957). <https://dx.doi.org/10.4153/cjm-1957-044-3>

D. Gale, A theorem on flows in networks, Pacific J. Math. 7, 1073 (1957). <https://dx.doi.org/10.2140/pjm.1957.7.1073>

A. Berger, A note on the characterization of digraphic sequences, Discrete Math. 314, 38 (2014). <https://dx.doi.org/10.1016/j.disc.2013.09.010>

Arguments:

`degrees1:` A vector of integers specifying the degrees in the first partition

degrees2: A vector of integers specifying the degrees in the second partition

allowed_edge_types: The types of edges to allow in the graph:

IGRAPH_SIMPLE_SW simple graphs (i.e. no multi-edges allowed).

IGRAPH_MULTI_SW multi-edges are allowed.

res: Pointer to a Boolean. The result will be stored here.

Returns:

Error code.

See also:

`igraph_is_graphical()`

Time complexity: $O(n \log n)$ for simple graphs, $O(n)$ for multigraphs, where n is the length of the larger degree sequence.

igraph_is_degree_sequence — Determines whether a degree sequence is valid.

```
int igraph_is_degree_sequence(const igraph_vector_t *out_degrees,
                             const igraph_vector_t *in_degrees, igraph_bool_t
```

Warning

Deprecated since version 0.9. Please do not use this function in new code; use `igraph_is_graphical()` instead.

A sequence of n integers is a valid degree sequence if there exists some graph where the degree of the i -th vertex is equal to the i -th element of the sequence. Note that the graph may contain multiple or loop edges; if you are interested in whether the degrees of some *simple* graph may realize the given sequence, use `igraph_is_graphical_degree_sequence`.

In particular, the function checks whether all the degrees are non-negative. For undirected graphs, it also checks whether the sum of degrees is even. For directed graphs, the function checks whether the lengths of the two degree vectors are equal and whether their sums are also equal. These are known sufficient and necessary conditions for a degree sequence to be valid.

Arguments:

out_degrees: an integer vector specifying the degree sequence for undirected graphs or the out-degree sequence for directed graphs.

in_degrees: an integer vector specifying the in-degrees of the vertices for directed graphs. For undirected graphs, this must be null.

res: pointer to a boolean variable, the result will be stored here

Returns:

Error code.

Time complexity: $O(n)$, where n is the length of the degree sequence.

igraph_is_graphical_degree_sequence — Determines whether a sequence of integers can be the degree sequence of some simple graph.

```
int igraph_is_graphical_degree_sequence(const igraph_vector_t *out_degrees,  
                                       const igraph_vector_t *in_degrees, igraph_t graph)
```

Warning

Deprecated since version 0.9. Please do not use this function in new code; use `igraph_is_graphical()` instead.

References:

Hakimi SL: On the realizability of a set of integers as degrees of the vertices of a simple graph. J SIAM Appl Math 10:496-506, 1962.

PL Erdős, I Miklós and Z Toroczkai: A simple Havel-Hakimi type algorithm to realize graphical degree sequences of directed graphs. The Electronic Journal of Combinatorics 17(1):R66, 2010. <https://dx.doi.org/10.1017/S0963548317000499>

Z Kiraly: Recognizing graphic degree sequences and generating all realizations. TR-2011-11, Egervary Research Group, H-1117, Budapest, Hungary. ISSN 1587-4451, 2012. <https://www.c-s.elte.hu/egres/tr/egres-11-11.pdf>

Arguments:

out_degrees: an integer vector specifying the degree sequence for undirected graphs or the out-degree sequence for directed graphs.

in_degrees: an integer vector specifying the in-degrees of the vertices for directed graphs. For undirected graphs, this must be null.

res: pointer to a boolean variable, the result will be stored here

Returns:

Error code.

Time complexity: $O(n \log n)$ for undirected graphs, $O(n^2)$ for directed graphs, where n is the length of the degree sequence.

Centrality measures

igraph_closeness — Closeness centrality calculations for some vertices.

```
int igraph_closeness(const igraph_t *graph, igraph_vector_t *res,
                    igraph_vector_t *reachable_count, igraph_bool_t *all_reachable,
                    const igraph_vs_t vids, igraph_neimode_t mode,
                    const igraph_vector_t *weights,
                    igraph_bool_t normalized);
```

The closeness centrality of a vertex measures how easily other vertices can be reached from it (or the other way: how easily it can be reached from the other vertices). It is defined as the inverse of the mean distance to (or from) all other vertices.

Closeness centrality is meaningful only for connected graphs. If the graph is not connected, igraph computes the inverse of the mean distance to (or from) all *reachable* vertices. In undirected graphs, this is equivalent to computing the closeness separately in each connected component. The optional *all_reachable* output parameter is provided to help detect when the graph is disconnected.

While there is no universally adopted definition of closeness centrality for disconnected graphs, there have been some attempts for generalizing the concept to the disconnected case. One type of approach considers the mean distance only to reachable vertices, then re-scales the obtained centrality score by a factor that depends on the number of reachable vertices (i.e. the size of the component in the undirected case). To facilitate computing these generalizations of closeness centrality, the number of reachable vertices (not including the starting vertex) is returned in *reachable_count*.

In disconnected graphs, consider using the harmonic centrality, computable using `igraph_harmonic_centrality()`.

For isolated vertices, i.e. those having no associated paths, NaN is returned.

Arguments:

<i>graph</i> :	The graph object.
<i>res</i> :	The result of the computation, a vector containing the closeness centrality scores for the given vertices.
<i>reachable_count</i> :	If not NULL, this vector will contain the number of vertices reachable from each vertex for which the closeness is calculated (not including that vertex).
<i>all_reachable</i> :	Pointer to a Boolean. If not NULL, it indicates if all vertices of the graph were reachable from each vertex in <i>vids</i> . If false, the graph is non-connected. If true, and the graph is undirected, or if the graph is directed and <i>vids</i> contains all vertices, then the graph is connected.
<i>vids</i> :	The vertices for which the closeness centrality will be computed.
<i>mode</i> :	The type of shortest paths to be used for the calculation in directed graphs. Possible values: <ul style="list-style-type: none"> IGRAPH_OUT the lengths of the outgoing paths are calculated. IGRAPH_IN the lengths of the incoming paths are calculated. IGRAPH_ALL the directed graph is considered as an undirected one for the computation.
<i>weights</i> :	An optional vector containing edge weights for weighted closeness. No edge weight may be NaN. Supply a null pointer here for traditional, unweighted closeness.
<i>normalized</i> :	If true, the inverse of the mean distance to reachable vertices is returned. If false, the inverse of the sum of distances is returned.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVAL invalid vertex id passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(n|E|)$, n is the number of vertices for which the calculation is done and $|E|$ is the number of edges in the graph.

See also:

Other centrality types: `igraph_degree()`, `igraph_betweenness()`, `igraph_harmonic_centrality()`. See `igraph_closeness_cutoff()` for the range-limited closeness centrality.

`igraph_harmonic_centrality` — Harmonic centrality for some vertices.

```
int igraph_harmonic_centrality(const igraph_t *graph, igraph_vector_t *res,
                               const igraph_vs_t vids, igraph_neimode_t mode,
                               const igraph_vector_t *weights,
                               igraph_bool_t normalized);
```

The harmonic centrality of a vertex is the mean inverse distance to all other vertices. The inverse distance to an unreachable vertex is considered to be zero.

References:

M. Marchiori and V. Latora, Harmony in the small-world, *Physica A* 285, pp. 539-546 (2000). <https://doi.org/10.1016/S0378-4371%2800%2900311-3>

Y. Rochat, Closeness Centrality Extended to Unconnected Graphs: the Harmonic Centrality Index, ASNA 2009. <https://infoscience.epfl.ch/record/200525>

S. Vigna and P. Boldi, Axioms for Centrality, *Internet Mathematics* 10, (2014). <https://doi.org/10.1080/15427951.2013.865686>

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the harmonic centrality scores for the given vertices.

vids: The vertices for which the harmonic centrality will be computed.

mode: The type of shortest paths to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the lengths of the outgoing paths are calculated.

IGRAPH_IN the lengths of the incoming paths are calculated.

`IGRAPH_ALL` the directed graph is considered as an undirected one for the computation.

weights: An optional vector containing edge weights for weighted harmonic centrality. No edge weight may be NaN. If `NULL`, all weights are considered to be one.

normalized: Boolean, whether to normalize the result. If true, the result is the mean inverse path length to other vertices, i.e. it is normalized by the number of vertices minus one. If false, the result is the sum of inverse path lengths to other vertices.

Returns:

Error code:

`IGRAPH_ENOMEM` not enough memory for temporary data.

`IGRAPH_EINVVID` invalid vertex id passed.

`IGRAPH_EINVMODE` invalid mode argument.

Time complexity: $O(n|E|)$, where n is the number of vertices for which the calculation is done and $|E|$ is the number of edges in the graph.

See also:

Other centrality types: `igraph_closeness()`, `igraph_degree()`, `igraph_betweenness()`.

`igraph_betweenness` — Betweenness centrality of some vertices.

```
int igraph_betweenness(const igraph_t *graph, igraph_vector_t *res,
                      const igraph_vs_t vids, igraph_bool_t directed,
                      const igraph_vector_t* weights);
```

The betweenness centrality of a vertex is the number of geodesics going through it. If there are more than one geodesic between two vertices, the value of these geodesics are weighted by one over the number of geodesics.

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the betweenness scores for the specified vertices.

vids: The vertices of which the betweenness centrality scores will be calculated.

directed: Logical, if true directed paths will be considered for directed graphs. It is ignored for undirected graphs.

weights: An optional vector containing edge weights for calculating weighted betweenness. No edge weight may be NaN. Supply a null pointer here for unweighted betweenness.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex id passed in *vids*.

Time complexity: $O(|V||E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the graph. Note that the time complexity is independent of the number of vertices for which the score is calculated.

See also:

Other centrality types: `igraph_degree()`, `igraph_closeness()`. See `igraph_edge_betweenness()` for calculating the betweenness score of the edges in a graph. See `igraph_betweenness_cutoff()` to calculate the range-limited betweenness of the vertices in a graph.

igraph_edge_betweenness — Betweenness centrality of the edges.

```
int igraph_edge_betweenness(const igraph_t *graph, igraph_vector_t *result,
                           igraph_bool_t directed,
                           const igraph_vector_t *weights);
```

The betweenness centrality of an edge is the number of geodesics going through it. If there are more than one geodesics between two vertices, the value of these geodesics are weighted by one over the number of geodesics.

Arguments:

graph: The graph object.

result: The result of the computation, vector containing the betweenness scores for the edges.

directed: Logical, if true directed paths will be considered for directed graphs. It is ignored for undirected graphs.

weights: An optional weight vector for weighted edge betweenness. No edge weight may be NaN. Supply a null pointer here for the unweighted version.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: $O(|V||E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the graph.

See also:

Other centrality types: `igraph_degree()`, `igraph_closeness()`. See `igraph_edge_betweenness()` for calculating the betweenness score of the edges in a graph. See `igraph_edge_betweenness_cutoff()` to compute the range-limited betweenness score of the edges in a graph.

igraph_pagerank_algo_t — PageRank algorithm implementation

```
typedef enum {  
    IGRAPH_PAGERANK_ALGO_ARPACK = 1,  
    IGRAPH_PAGERANK_ALGO_PRPACK = 2  
} igraph_pagerank_algo_t;
```

Algorithms to calculate PageRank.

Values:

IGRAPH_PAGERANK_ALGO_ARPACK:	Use the ARPACK library, this was the PageRank implementation in igraph from version 0.5, until version 0.7.
IGRAPH_PAGERANK_ALGO_PRPACK:	Use the PRPACK library. Currently this implementation is recommended.

igraph_pagerank — Calculates the Google PageRank for the specified vertices.

```
int igraph_pagerank(const igraph_t *graph, igraph_pagerank_algo_t algo,  
                    igraph_vector_t *vector,  
                    igraph_real_t *value, const igraph_vs_t vids,  
                    igraph_bool_t directed, igraph_real_t damping,  
                    const igraph_vector_t *weights, igraph_arpack_options_t *options)
```

The PageRank centrality of a vertex is the fraction of time a random walker traversing the graph would spend on that vertex. The walker follows the out-edges with probabilities proportional to their weights. Additionally, in each step, it restarts the walk from a random vertex with probability $1 - \text{damping}$. If the random walker gets stuck in a sink vertex, it will also restart from a random vertex.

The PageRank centrality is mainly useful for directed graphs. In undirected graphs it converges to trivial values proportional to degrees as the damping factor approaches 1.

Starting from version 0.9, igraph has two PageRank implementations, and the user can choose between them. The first implementation is `IGRAPH_PAGERANK_ALGO_ARPACK`, based on the ARPACK library. This was the default before igraph version 0.7. The second and recommended implementation is `IGRAPH_PAGERANK_ALGO_PRPACK`. This is using the PRPACK package, see <https://github.com/dgleich/prpack>.

Note that the PageRank of a given vertex depends on the PageRank of all other vertices, so even if you want to calculate the PageRank for only some of the vertices, all of them must be calculated. Requesting the PageRank for only some of the vertices does not result in any performance increase at all.

References:

Sergey Brin and Larry Page: The Anatomy of a Large-Scale Hypertextual Web Search Engine. Proceedings of the 7th World-Wide Web Conference, Brisbane, Australia, April 1998.

Arguments:

<i>graph</i> :	The graph object.
<i>algo</i> :	The PageRank implementation to use. Possible values: <code>IGRAPH_PAGERANK_ALGO_ARPACK</code> , <code>IGRAPH_PAGERANK_ALGO_PRPACK</code> .
<i>vector</i> :	Pointer to an initialized vector, the result is stored here. It is resized as needed.

<i>value</i> :	Pointer to a real variable, the eigenvalue corresponding to the PageRank vector is stored here. It should be always exactly one.
<i>vids</i> :	The vertex ids for which the PageRank is returned.
<i>directed</i> :	Boolean, whether to consider the directedness of the edges. This is ignored for undirected graphs.
<i>damping</i> :	The damping factor ("d" in the original paper). Must be a probability in the range [0, 1]. A commonly used value is 0.85.
<i>weights</i> :	Optional edge weights. May be a NULL pointer, meaning unweighted edges, or a vector of non-negative values of the same length as the number of edges.
<i>options</i> :	Options for the ARPACK method. See <code>igraph_arpack_options_t</code> for details. Note that the function overwrites the <code>n</code> (number of vertices), <code>nev</code> (1), <code>ncv</code> (3) and <code>which</code> (LM) parameters and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code: `IGRAPH_ENOMEM`, not enough memory for temporary data. `IGRAPH_EINVVID`, invalid vertex id in *vids*.

Time complexity: depends on the input graph, usually it is $O(|E|)$, the number of edges.

See also:

`igraph_personalized_pagerank()` and `igraph_personalized_pagerank_vs()` for the personalized PageRank measure. See `igraph_arpack_rssolve()` and `igraph_arpack_rnsolve()` for the underlying machinery used by `IGRAPH_PAGERANK_ALGO_ARPACK`.

Example 13.14. File `examples/simple/igraph_pagerank.c`

`igraph_personalized_pagerank` — Calculates the personalized Google PageRank for the specified vertices.

```
int igraph_personalized_pagerank(const igraph_t *graph,
                                igraph_pagerank_algo_t algo, igraph_vector_t *
                                igraph_real_t *value, const igraph_vs_t vids,
                                igraph_bool_t directed, igraph_real_t damping,
                                const igraph_vector_t *reset,
                                const igraph_vector_t *weights,
                                igraph_arpack_options_t *options);
```

The personalized PageRank is similar to the original PageRank measure, but when the random walk is restarted, a new starting vertex is chosen non-uniformly, according to the distribution specified in *reset* (instead of the uniform distribution in the original PageRank measure). The *reset* distribution is used both when restarting randomly with probability $1 - \text{damping}$, and when the walker is forced to restart due to being stuck in a sink vertex (a vertex with no outgoing edges).

Note that the personalized PageRank of a given vertex depends on the personalized PageRank of all other vertices, so even if you want to calculate the personalized PageRank for only some of the vertices,

all of them must be calculated. Requesting the personalized PageRank for only some of the vertices does not result in any performance increase at all.

Arguments:

<i>graph</i> :	The graph object.
<i>algo</i> :	The PageRank implementation to use. Possible values: <code>IGRAPH_PAGERANK_ALGO_ARPACK</code> , <code>IGRAPH_PAGERANK_ALGO_PRPACK</code> .
<i>vector</i> :	Pointer to an initialized vector, the result is stored here. It is resized as needed.
<i>value</i> :	Pointer to a real variable, the eigenvalue corresponding to the PageRank vector is stored here. It should be always exactly one.
<i>vids</i> :	The vertex ids for which the PageRank is returned.
<i>directed</i> :	Boolean, whether to consider the directedness of the edges. This is ignored for undirected graphs.
<i>damping</i> :	The damping factor ("d" in the original paper). Must be a probability in the range [0, 1]. A commonly used value is 0.85.
<i>reset</i> :	The probability distribution over the vertices used when resetting the random walk. It is either a <code>NULL</code> pointer (denoting a uniform choice that results in the original PageRank measure) or a vector of the same length as the number of vertices.
<i>weights</i> :	Optional edge weights. May be a <code>NULL</code> pointer, meaning unweighted edges, or a vector of non-negative values of the same length as the number of edges.
<i>options</i> :	Options for the ARPACK method. See <code>igraph_arpack_options_t</code> for details. Note that the function overwrites the <code>n</code> (number of vertices), <code>nev</code> (1), <code>ncv</code> (3) and <code>which</code> (LM) parameters and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code: `IGRAPH_ENOMEM`, not enough memory for temporary data. `IGRAPH_EINVVID`, invalid vertex id in *vids* or an invalid reset vector in *reset*.

Time complexity: depends on the input graph, usually it is $O(|E|)$, the number of edges.

See also:

`igraph_pagerank()` for the non-personalized implementation, `igraph_personalized_pagerank_vs()` for a personalized implementation with resetting to specific vertices.

`igraph_personalized_pagerank_vs` — Calculates the personalized Google PageRank for the specified vertices.

```
int igraph_personalized_pagerank_vs(const igraph_t *graph,
                                   igraph_pagerank_algo_t algo, igraph_vector_t *vector,
                                   igraph_real_t *value, const igraph_vs_t vids,
                                   igraph_bool_t directed, igraph_real_t damping)
```

```
igraph_vs_t reset_vids,  
const igraph_vector_t *weights,  
igraph_arpack_options_t *options);
```

The personalized PageRank is similar to the original PageRank measure, but when the random walk is restarted, a new starting vertex is chosen according to a specified distribution. This distribution is used both when restarting randomly with probability $1 - \text{damping}$, and when the walker is forced to restart due to being stuck in a sink vertex (a vertex with no outgoing edges).

This simplified interface takes a vertex sequence and resets the random walk to one of the vertices in the specified vertex sequence, chosen uniformly. A typical application of personalized PageRank is when the random walk is reset to the same vertex every time - this can easily be achieved using `igraph_vss_1()` which generates a vertex sequence containing only a single vertex.

Note that the personalized PageRank of a given vertex depends on the personalized PageRank of all other vertices, so even if you want to calculate the personalized PageRank for only some of the vertices, all of them must be calculated. Requesting the personalized PageRank for only some of the vertices does not result in any performance increase at all.

Arguments:

<i>graph</i> :	The graph object.
<i>algo</i> :	The PageRank implementation to use. Possible values: <code>IGRAPH_PAGERANK_ALGO_ARPACK</code> , <code>IGRAPH_PAGERANK_ALGO_PRPACK</code> .
<i>vector</i> :	Pointer to an initialized vector, the result is stored here. It is resized as needed.
<i>value</i> :	Pointer to a real variable, the eigenvalue corresponding to the PageRank vector is stored here. It should be always exactly one.
<i>vids</i> :	The vertex ids for which the PageRank is returned.
<i>directed</i> :	Boolean, whether to consider the directedness of the edges. This is ignored for undirected graphs.
<i>damping</i> :	The damping factor ("d" in the original paper). Must be a probability in the range [0, 1]. A commonly used value is 0.85.
<i>reset_vids</i> :	IDs of the vertices used when resetting the random walk.
<i>weights</i> :	Optional edge weights, it is either a null pointer, then the edges are not weighted, or a vector of the same length as the number of edges.
<i>options</i> :	Options for the ARPACK method. See <code>igraph_arpack_options_t</code> for details. Note that the function overwrites the <code>n</code> (number of vertices), <code>nev</code> (1), <code>ncv</code> (3) and <code>which</code> (LM) parameters and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code: `IGRAPH_ENOMEM`, not enough memory for temporary data. `IGRAPH_EINVVID`, invalid vertex id in *vids* or an empty reset vertex sequence in *vids_reset*.

Time complexity: depends on the input graph, usually it is $O(|E|)$, the number of edges.

See also:

`igraph_pagerank()` for the non-personalized implementation.

igraph_constraint — Burt's constraint scores.

```
int igraph_constraint(const igraph_t *graph, igraph_vector_t *res,
                    igraph_vs_t vids, const igraph_vector_t *weights);
```

This function calculates Burt's constraint scores for the given vertices, also known as structural holes.

Burt's constraint is higher if ego has less, or mutually stronger related (i.e. more redundant) contacts. Burt's measure of constraint, $C[i]$, of vertex i 's ego network $V[i]$, is defined for directed and valued graphs,

$$C[i] = \text{sum}(\text{sum}((p[i,q] p[q,j])^2, q \in V[i], q \neq i, j \in V[], j \neq i))$$

for a graph of order (i.e. number of vertices) N , where proportional tie strengths are defined as

$$p[i,j] = (a[i,j] + a[j,i]) / \text{sum}(a[i,k] + a[k,i], k \in V[i], k \neq i),$$

$a[i,j]$ are elements of A and the latter being the graph adjacency matrix. For isolated vertices, constraint is undefined.

Burt, R.S. (2004). Structural holes and good ideas. *American Journal of Sociology* 110, 349-399.

The first R version of this function was contributed by Jeroen Bruggeman.

Arguments:

- graph*: A graph object.
- res*: Pointer to an initialized vector, the result will be stored here. The vector will be resized to have the appropriate size for holding the result.
- vids*: Vertex selector containing the vertices for which the constraint should be calculated.
- weights*: Vector giving the weights of the edges. If it is `NULL` then each edge is supposed to have the same weight.

Returns:

Error code.

Time complexity: $O(|V|+|E|+n*d^2)$, n is the number of vertices for which the constraint is calculated and d is the average degree, $|V|$ is the number of vertices, $|E|$ the number of edges in the graph. If the *weights* argument is `NULL` then the time complexity is $O(|V|+n*d^2)$.

igraph_maxdegree — The maximum degree in a graph (or set of vertices).

```
int igraph_maxdegree(const igraph_t *graph, igraph_integer_t *res,
                    igraph_vs_t vids, igraph_neimode_t mode,
                    igraph_bool_t loops);
```

The largest in-, out- or total degree of the specified vertices is calculated. If the graph has no vertices, or *vids* is empty, 0 is returned, as this is the smallest possible value for degrees.

Arguments:

- graph*: The input graph.
- res*: Pointer to an integer (`igraph_integer_t`), the result will be stored here.
- vids*: Vector giving the vertex IDs for which the maximum degree will be calculated.
- mode*: Defines the type of the degree. `IGRAPH_OUT`, out-degree, `IGRAPH_IN`, in-degree, `IGRAPH_ALL`, total degree (sum of the in- and out-degree). This parameter is ignored for undirected graphs.
- loops*: Boolean, gives whether the self-loops should be counted.

Returns:

Error code: `IGRAPH_EINVAL`: invalid vertex id. `IGRAPH_EINVALMODE`: invalid mode argument.

Time complexity: $O(v)$ if `loops` is `TRUE`, and $O(v*d)$ otherwise. v is the number of vertices for which the degree will be calculated, and d is their (average) degree.

`igraph_strength` — Strength of the vertices, weighted vertex degree in other words.

```
int igraph_strength(const igraph_t *graph, igraph_vector_t *res,
                   const igraph_vs_t vids, igraph_neimode_t mode,
                   igraph_bool_t loops, const igraph_vector_t *weights);
```

In a weighted network the strength of a vertex is the sum of the weights of all incident edges. In a non-weighted network this is exactly the vertex degree.

Arguments:

- graph*: The input graph.
- res*: Pointer to an initialized vector, the result is stored here. It will be resized as needed.
- vids*: The vertices for which the calculation is performed.
- mode*: Gives whether to count only outgoing (`IGRAPH_OUT`), incoming (`IGRAPH_IN`) edges or both (`IGRAPH_ALL`).
- loops*: A logical scalar, whether to count loop edges as well.
- weights*: A vector giving the edge weights. If this is a `NULL` pointer, then `igraph_degree()` is called to perform the calculation.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number vertices and edges.

See also:

`igraph_degree()` for the traditional, non-weighted version.

igraph_eigenvector_centrality — Eigenvector centrality of the vertices

```
int igraph_eigenvector_centrality(const igraph_t *graph,
                                igraph_vector_t *vector,
                                igraph_real_t *value,
                                igraph_bool_t directed, igraph_bool_t scale,
                                const igraph_vector_t *weights,
                                igraph_arnpack_options_t *options);
```

Eigenvector centrality is a measure of the importance of a node in a network. It assigns relative scores to all nodes in the network based on the principle that connections from high-scoring nodes contribute more to the score of the node in question than equal connections from low-scoring nodes. Specifically, the eigenvector centrality of each vertex is proportional to the sum of eigenvector centralities of its neighbors. In practice, the centralities are determined by calculating the eigenvector corresponding to the largest positive eigenvalue of the adjacency matrix. In the undirected case, this function considers the diagonal entries of the adjacency matrix to be *twice* the number of self-loops on the corresponding vertex.

The centrality scores returned by igraph can be normalized (using the *scale* parameter) such that the largest eigenvector centrality score is 1 (with one exception, see below).

In the directed case, the left eigenvector of the adjacency matrix is calculated. In other words, the centrality of a vertex is proportional to the sum of centralities of vertices pointing to it.

Eigenvector centrality is meaningful only for connected graphs. Graphs that are not connected should be decomposed into connected components, and the eigenvector centrality calculated for each separately. This function does not verify that the graph is connected. If it is not, in the undirected case the scores of all but one component will be zeros.

Also note that the adjacency matrix of a directed acyclic graph or the adjacency matrix of an empty graph does not possess positive eigenvalues, therefore the eigenvector centrality is not defined for these graphs. igraph will return an eigenvalue of zero in such cases. The eigenvector centralities will all be equal for an empty graph and will all be zeros for a directed acyclic graph. Such pathological cases can be detected by asking igraph to calculate the eigenvalue as well (using the *value* parameter, see below) and checking whether the eigenvalue is very close to zero.

Arguments:

- | | |
|-------------------|---|
| <i>graph</i> : | The input graph. It may be directed. |
| <i>vector</i> : | Pointer to an initialized vector, it will be resized as needed. The result of the computation is stored here. It can be a null pointer, then it is ignored. |
| <i>value</i> : | If not a null pointer, then the eigenvalue corresponding to the found eigenvector is stored here. |
| <i>directed</i> : | Boolean scalar, whether to consider edge directions in a directed graph. It is ignored for undirected graphs. |
| <i>scale</i> : | If not zero then the result will be scaled such that the absolute value of the maximum centrality is one. |
| <i>weights</i> : | A null pointer (=no edge weights), or a vector giving the weights of the edges. The algorithm might result complex numbers if some weights are negative. In this case only the real part is reported. |

options: Options to ARPACK. See `igraph_arnpack_options_t` for details. Note that the function overwrites the `n` (number of vertices) parameter and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code.

Time complexity: depends on the input graph, usually it is $O(|V|+|E|)$.

See also:

`igraph_pagerank` and `igraph_personalized_pagerank` for modifications of eigenvector centrality.

Example 13.15. File `examples/simple/eigenvector_centrality.c`

`igraph_hub_score` — Kleinberg's hub scores.

```
int igraph_hub_score(const igraph_t *graph, igraph_vector_t *vector,
                    igraph_real_t *value, igraph_bool_t scale,
                    const igraph_vector_t *weights,
                    igraph_arnpack_options_t *options);
```

The hub scores of the vertices are defined as the principal eigenvector of $A \cdot A^T$, where A is the adjacency matrix of the graph, A^T is its transposed.

See the following reference on the meaning of this score: J. Kleinberg. Authoritative sources in a hyperlinked environment. *Proc. 9th ACM-SIAM Symposium on Discrete Algorithms*, 1998. Extended version in *Journal of the ACM* 46(1999). Also appears as IBM Research Report RJ 10076, May 1997.

Arguments:

graph: The input graph. Can be directed and undirected.

vector: Pointer to an initialized vector, the result is stored here. If a null pointer then it is ignored.

value: If not a null pointer then the eigenvalue corresponding to the calculated eigenvector is stored here.

scale: If not zero then the result will be scaled such that the absolute value of the maximum centrality is one.

weights: A null pointer (=no edge weights), or a vector giving the weights of the edges.

options: Options to ARPACK. See `igraph_arnpack_options_t` for details. Note that the function overwrites the `n` (number of vertices) parameter and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code.

Time complexity: depends on the input graph, usually it is $O(|V|)$, the number of vertices.

See also:

`igraph_authority_score()` for the companion measure, `igraph_pagerank()`, `igraph_personalized_pagerank()`, `igraph_eigenvector_centrality()` for similar measures.

igraph_authority_score — Kleinerg's authority scores.

```
int igraph_authority_score(const igraph_t *graph, igraph_vector_t *vector,
                           igraph_real_t *value, igraph_bool_t scale,
                           const igraph_vector_t *weights,
                           igraph_arnpack_options_t *options);
```

The authority scores of the vertices are defined as the principal eigenvector of $A^T A$, where A is the adjacency matrix of the graph, A^T is its transposed.

See the following reference on the meaning of this score: J. Kleinberg. Authoritative sources in a hyperlinked environment. *Proc. 9th ACM-SIAM Symposium on Discrete Algorithms*, 1998. Extended version in *Journal of the ACM* 46(1999). Also appears as IBM Research Report RJ 10076, May 1997.

Arguments:

- graph*: The input graph. Can be directed and undirected.
- vector*: Pointer to an initialized vector, the result is stored here. If a null pointer then it is ignored.
- value*: If not a null pointer then the eigenvalue corresponding to the calculated eigenvector is stored here.
- scale*: If not zero then the result will be scaled such that the absolute value of the maximum centrality is one.
- weights*: A null pointer (=no edge weights), or a vector giving the weights of the edges.
- options*: Options to ARPACK. See `igraph_arnpack_options_t` for details. Note that the function overwrites the `n` (number of vertices) parameter and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code.

Time complexity: depends on the input graph, usually it is $O(|V|)$, the number of vertices.

See also:

`igraph_hub_score()` for the companion measure, `igraph_pagerank()`, `igraph_personalized_pagerank()`, `igraph_eigenvector_centrality()` for similar measures.

igraph_convergence_degree — Calculates the convergence degree of each edge in a graph.


```
int igraph_convergence_degree(const igraph_t *graph, igraph_vector_t *result,
                             igraph_vector_t *ins, igraph_vector_t *outs);
```

Let us define the input set of an edge (i, j) as the set of vertices where the shortest paths passing through (i, j) originate, and similarly, let us define the output set of an edge (i, j) as the set of vertices where the shortest paths passing through (i, j) terminate. The convergence degree of an edge is defined as the normalized value of the difference between the size of the input set and the output set, i.e. the difference of them divided by the sum of them. Convergence degrees are in the range (-1, 1); a positive value indicates that the edge is *convergent* since the shortest paths passing through it originate from a larger set and terminate in a smaller set, while a negative value indicates that the edge is *divergent* since the paths originate from a small set and terminate in a larger set.

Note that the convergence degree as defined above does not make sense in undirected graphs as there is no distinction between the input and output set. Therefore, for undirected graphs, the input and output sets of an edge are determined by orienting the edge arbitrarily while keeping the remaining edges undirected, and then taking the absolute value of the convergence degree.

Arguments:

- graph*: The input graph, it can be either directed or undirected.
- result*: Pointer to an initialized vector; the convergence degrees of each edge will be stored here. May be NULL if we are not interested in the exact convergence degrees.
- ins*: Pointer to an initialized vector; the size of the input set of each edge will be stored here. May be NULL if we are not interested in the sizes of the input sets.
- outs*: Pointer to an initialized vector; the size of the output set of each edge will be stored here. May be NULL if we are not interested in the sizes of the output sets.

Returns:

Error code.

Time complexity: $O(|V||E|)$, the number of vertices times the number of edges.

Range-limited centrality measures

igraph_closeness_cutoff — Range limited closeness centrality.

```
int igraph_closeness_cutoff(const igraph_t *graph, igraph_vector_t *res,
                           igraph_vector_t *reachable_count, igraph_bool_t *all_reachable,
                           const igraph_vs_t vids, igraph_neimode_t mode,
                           const igraph_vector_t *weights,
                           igraph_bool_t normalized,
                           igraph_real_t cutoff);
```

This function computes a range-limited version of closeness centrality by considering only those shortest paths whose length is no greater than the given cutoff value.

Arguments:

- graph*: The graph object.

<i>res</i> :	The result of the computation, a vector containing the range-limited closeness centrality scores for the given vertices.
<i>reachable_count</i> :	If not NULL, this vector will contain the number of vertices reachable within the cutoff distance from each vertex for which the range-limited closeness is calculated (not including that vertex).
<i>all_reachable</i> :	Pointer to a Boolean. If not NULL, it indicates if all vertices of the graph were reachable from each vertex in <i>vids</i> within the given cutoff distance.
<i>vids</i> :	The vertices for which the range limited closeness centrality will be computed.
<i>mode</i> :	The type of shortest paths to be used for the calculation in directed graphs. Possible values: <div> <div>IGRAPH_OUT</div> <div>the lengths of the outgoing paths are calculated.</div> </div> <div> <div>IGRAPH_IN</div> <div>the lengths of the incoming paths are calculated.</div> </div> <div> <div>IGRAPH_ALL</div> <div>the directed graph is considered as an undirected one for the computation.</div> </div>
<i>weights</i> :	An optional vector containing edge weights for weighted closeness. No edge weight may be NaN. Supply a null pointer here for traditional, unweighted closeness.
<i>normalized</i> :	If true, the inverse of the mean distance to vertices reachable within the cutoff is returned. If false, the inverse of the sum of distances is returned.
<i>cutoff</i> :	The maximal length of paths that will be considered. If negative, the exact closeness will be calculated (no upper limit on path lengths).

Returns:

Error code:

IGRAPH_ENOMEM	not enough memory for temporary data.
IGRAPH_EINVVID	invalid vertex id passed.
IGRAPH_EINVMODE	invalid mode argument.

Time complexity: $O(n|E|)$, n is the number of vertices for which the calculation is done and $|E|$ is the number of edges in the graph.

See also:

`igraph_closeness()` to calculate the exact closeness centrality.

igraph_harmonic_centrality_cutoff — Range limited harmonic centrality.

```
int igraph_harmonic_centrality_cutoff(const igraph_t *graph, igraph_vector_t *r,
                                     const igraph_vs_t vids, igraph_neimode_t mode,
                                     const igraph_vector_t *weights,
```

```
igraph_bool_t normalized,  
igraph_real_t cutoff);
```

This function computes the range limited version of harmonic centrality: only those shortest paths are considered whose length is not above the given cutoff. The inverse distance to vertices not reachable within the cutoff is considered to be zero.

Arguments:

- graph*: The graph object.
- res*: The result of the computation, a vector containing the range limited harmonic centrality scores for the given vertices.
- vids*: The vertices for which the harmonic centrality will be computed.
- mode*: The type of shortest paths to be used for the calculation in directed graphs. Possible values:
- IGRAPH_OUT the lengths of the outgoing paths are calculated.
 - IGRAPH_IN the lengths of the incoming paths are calculated.
 - IGRAPH_ALL the directed graph is considered as an undirected one for the computation.
- weights*: An optional vector containing edge weights for weighted harmonic centrality. No edge weight may be NaN. If NULL, all weights are considered to be one.
- normalized*: Boolean, whether to normalize the result. If true, the result is the mean inverse path length to other vertices. i.e. it is normalized by the number of vertices minus one. If false, the result is the sum of inverse path lengths to other vertices.
- cutoff*: The maximal length of paths that will be considered. The inverse distance to vertices that are not reachable within the cutoff path length is considered to be zero. Supply a negative value to compute the exact harmonic centrality, without any upper limit on the length of paths.

Returns:

Error code:

- IGRAPH_ENOMEM not enough memory for temporary data.
- IGRAPH_EINVVID invalid vertex id passed.
- IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(n|E|)$, where n is the number of vertices for which the calculation is done and $|E|$ is the number of edges in the graph.

See also:

Other centrality types: `igraph_closeness()`, `igraph_betweenness()`.

igraph_betweenness_cutoff — Range-limited betweenness centrality.

```
int igraph_betweenness_cutoff(const igraph_t *graph, igraph_vector_t *res,
                             const igraph_vs_t vids, igraph_bool_t directed,
                             const igraph_vector_t *weights, igraph_real_t cutoff);
```

This function computes a range-limited version of betweenness centrality by considering only those shortest paths whose length is no greater than the given cutoff value.

Arguments:

- graph*: The graph object.
- res*: The result of the computation, a vector containing the range-limited betweenness scores for the specified vertices.
- vids*: The vertices for which the range-limited betweenness centrality scores will be computed.
- directed*: Logical, if true directed paths will be considered for directed graphs. It is ignored for undirected graphs.
- weights*: An optional vector containing edge weights for calculating weighted betweenness. No edge weight may be NaN. Supply a null pointer here for unweighted betweenness.
- cutoff*: The maximal length of paths that will be considered. If negative, the exact betweenness will be calculated, and there will be no upper limit on path lengths.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex id passed in *vids*.

Time complexity: $O(|V||E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the graph. Note that the time complexity is independent of the number of vertices for which the score is calculated.

See also:

`igraph_betweenness()` to calculate the exact betweenness and `igraph_edge_betweenness_cutoff()` to calculate the range-limited edge betweenness.

igraph_edge_betweenness_cutoff — Range-limited betweenness centrality of the edges.

```
int igraph_edge_betweenness_cutoff(const igraph_t *graph, igraph_vector_t *result,
                                   igraph_bool_t directed,
                                   const igraph_vector_t *weights, igraph_real_t cutoff);
```

This function computes a range-limited version of edge betweenness centrality by considering only those shortest paths whose length is no greater than the given cutoff value.

Arguments:

- graph*: The graph object.
- result*: The result of the computation, vector containing the betweenness scores for the edges.

- directed*: Logical, if true directed paths will be considered for directed graphs. It is ignored for undirected graphs.
- weights*: An optional weight vector for weighted betweenness. No edge weight may be NaN. Supply a null pointer here for unweighted betweenness.
- cutoff*: The maximal length of paths that will be considered. If negative, the exact betweenness will be calculated (no upper limit on path lengths).

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: $O(|V||E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the graph.

See also:

`igraph_edge_betweenness()` to compute the exact edge betweenness and `igraph_betweenness_cutoff()` to compute the range-limited vertex betweenness.

Centralization

`igraph_centralization` — Calculate the centralization score from the node level scores

```
igraph_real_t igraph_centralization(const igraph_vector_t *scores,
                                   igraph_real_t theoretical_max,
                                   igraph_bool_t normalized);
```

For a centrality score defined on the vertices of a graph, it is possible to define a graph level centralization index, by calculating the sum of the deviation from the maximum centrality score. Consequently, the higher the centralization index of the graph, the more centralized the structure is.

In order to make graphs of different sizes comparable, the centralization index is usually normalized to a number between zero and one, by dividing the (unnormalized) centralization score of the most centralized structure with the same number of vertices.

For most centrality indices the most centralized structure is the star graph, a single center connected to all other nodes in the network. There are some variation depending on whether the graph is directed or not, whether loop edges are allowed, etc.

This function simply calculates the graph level index, if the node level scores and the theoretical maximum are given. It is called by all the measure-specific centralization functions.

Arguments:

- scores*: A vector containing the node-level centrality scores.
- theoretical_max*: The graph level centrality score of the most centralized graph with the same number of vertices. Only used if `normalized` set to true.
- normalized*: Boolean, whether to normalize the centralization by dividing the supplied theoretical maximum.

Returns:

The graph level index.

See also:

`igraph_centralization_degree()`, `igraph_centralization_betweenness()`, `igraph_centralization_closeness()`, and `igraph_centralization_eigenvector_centrality()` for specific centralization functions.

Time complexity: $O(n)$, the length of the score vector.

Example 13.16. File `examples/simple/centralization.c`

`igraph_centralization_degree` — Calculate vertex degree and graph centralization

```
int igraph_centralization_degree(const igraph_t *graph, igraph_vector_t *res,
                                igraph_neimode_t mode, igraph_bool_t loops,
                                igraph_real_t *centralization,
                                igraph_real_t *theoretical_max,
                                igraph_bool_t normalized);
```

This function calculates the degree of the vertices by passing its arguments to `igraph_degree()`; and it calculates the graph level centralization index based on the results by calling `igraph_centralization()`.

Arguments:

<i>graph</i> :	The input graph.
<i>res</i> :	A vector if you need the node-level degree scores, or a null pointer otherwise.
<i>mode</i> :	Constant that specifies the type of degree for directed graphs. Possible values: <code>IGRAPH_IN</code> , <code>IGRAPH_OUT</code> and <code>IGRAPH_ALL</code> . This argument is ignored for undirected graphs.
<i>loops</i> :	Boolean, whether to consider loop edges when calculating the degree (and the centralization).
<i>centralization</i> :	Pointer to a real number, the centralization score is placed here.
<i>theoretical_max</i> :	Pointer to real number or a null pointer. If not a null pointer, then the theoretical maximum graph centrality score for a graph with the same number of vertices is stored here.
<i>normalized</i> :	Boolean, whether to calculate a normalized centralization score. See <code>igraph_centralization()</code> for how the normalization is done.

Returns:

Error code.

See also:

`igraph_centralization()`, `igraph_degree()`.

Time complexity: the complexity of `igraph_degree()` plus $O(n)$, the number of vertices queried, for calculating the centralization score.

`igraph_centralization_betweenness` — Calculate vertex betweenness and graph centralization

```
int igraph_centralization_betweenness(const igraph_t *graph,
                                      igraph_vector_t *res,
                                      igraph_bool_t directed,
                                      igraph_real_t *centralization,
                                      igraph_real_t *theoretical_max,
                                      igraph_bool_t normalized);
```

This function calculates the betweenness centrality of the vertices by passing its arguments to `igraph_betweenness()`; and it calculates the graph level centralization index based on the results by calling `igraph_centralization()`.

Arguments:

<i>graph</i> :	The input graph.
<i>res</i> :	A vector if you need the node-level betweenness scores, or a null pointer otherwise.
<i>directed</i> :	Boolean, whether to consider directed paths when calculating betweenness.
<i>centralization</i> :	Pointer to a real number, the centralization score is placed here.
<i>theoretical_max</i> :	Pointer to real number or a null pointer. If not a null pointer, then the theoretical maximum graph centrality score for a graph with the same number vertices is stored here.
<i>normalized</i> :	Boolean, whether to calculate a normalized centralization score. See <code>igraph_centralization()</code> for how the normalization is done.

Returns:

Error code.

See also:

`igraph_centralization()`, `igraph_betweenness()`.

Time complexity: the complexity of `igraph_betweenness()` plus $O(n)$, the number of vertices queried, for calculating the centralization score.

`igraph_centralization_closeness` — Calculate vertex closeness and graph centralization

```
int igraph_centralization_closeness(const igraph_t *graph,
                                   igraph_vector_t *res,
                                   igraph_neimode_t mode,
                                   igraph_real_t *centralization,
                                   igraph_real_t *theoretical_max,
                                   igraph_bool_t normalized);
```

This function calculates the closeness centrality of the vertices by passing its arguments to `igraph_closeness()`; and it calculates the graph level centralization index based on the results by calling `igraph_centralization()`.

Arguments:

<i>graph</i> :	The input graph.
<i>res</i> :	A vector if you need the node-level closeness scores, or a null pointer otherwise.
<i>mode</i> :	Constant that specifies the type of closeness for directed graphs. Possible values: <code>IGRAPH_IN</code> , <code>IGRAPH_OUT</code> and <code>IGRAPH_ALL</code> . This argument is ignored for undirected graphs. See <code>igraph_closeness()</code> argument with the same name for more.
<i>centralization</i> :	Pointer to a real number, the centralization score is placed here.
<i>theoretical_max</i> :	Pointer to real number or a null pointer. If not a null pointer, then the theoretical maximum graph centrality score for a graph with the same number of vertices is stored here.
<i>normalized</i> :	Boolean, whether to calculate a normalized centralization score. See <code>igraph_centralization()</code> for how the normalization is done.

Returns:

Error code.

See also:

`igraph_centralization()`, `igraph_closeness()`.

Time complexity: the complexity of `igraph_closeness()` plus $O(n)$, the number of vertices queried, for calculating the centralization score.

`igraph_centralization_eigenvector_centrality` — Calculate eigenvector centrality scores and graph centralization

```
int igraph_centralization_eigenvector_centrality(
    const igraph_t *graph,
    igraph_vector_t *vector,
    igraph_real_t *value,
    igraph_bool_t directed,
    igraph_bool_t scale,
    igraph_arpack_options_t *options,
```



```
igraph_real_t *centralization,  
igraph_real_t *theoretical_max,  
igraph_bool_t normalized);
```

This function calculates the eigenvector centrality of the vertices by passing its arguments to `igraph_eigenvector_centrality()`; and it calculates the graph level centralization index based on the results by calling `igraph_centralization()`.

Arguments:

<i>graph</i> :	The input graph.
<i>vector</i> :	A vector if you need the node-level eigenvector centrality scores, or a null pointer otherwise.
<i>value</i> :	If not a null pointer, then the leading eigenvalue is stored here.
<i>scale</i> :	If not zero then the result will be scaled, such that the absolute value of the maximum centrality is one.
<i>options</i> :	Options to ARPACK. See <code>igraph_arpack_options_t</code> for details. Note that the function overwrites the <code>n</code> (number of vertices) parameter and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.
<i>centralization</i> :	Pointer to a real number, the centralization score is placed here.
<i>theoretical_max</i> :	Pointer to real number or a null pointer. If not a null pointer, then the theoretical maximum graph centrality score for a graph with the same number of vertices is stored here.
<i>normalized</i> :	Boolean, whether to calculate a normalized centralization score. See <code>igraph_centralization()</code> for how the normalization is done.

Returns:

Error code.

See also:

`igraph_centralization()`, `igraph_eigenvector_centrality()`.

Time complexity: the complexity of `igraph_eigenvector_centrality()` plus $O(|V|)$, the number of vertices for the calculating the centralization.

`igraph_centralization_degree_tmax` — Theoretical maximum for graph centralization based on degree

```
int igraph_centralization_degree_tmax(const igraph_t *graph,  
                                     igraph_integer_t nodes,  
                                     igraph_neimode_t mode,  
                                     igraph_bool_t loops,  
                                     igraph_real_t *res);
```

This function returns the theoretical maximum graph centrality based on vertex degree.

There are two ways to call this function, the first is to supply a graph as the `graph` argument, and then the number of vertices is taken from this object, and its directedness is considered as well. The `nodes` argument is ignored in this case. The `mode` argument is also ignored if the supplied graph is undirected.

The other way is to supply a null pointer as the `graph` argument. In this case the `nodes` and `mode` arguments are considered.

The most centralized structure is the star. More specifically, for undirected graphs it is the star, for directed graphs it is the in-star or the out-star.

Arguments:

graph: A graph object or a null pointer, see the description above.

nodes: The number of nodes. This is ignored if the `graph` argument is not a null pointer.

mode: Constant, whether the calculation is based on in-degree (`IGRAPH_IN`), out-degree (`IGRAPH_OUT`) or total degree (`IGRAPH_ALL`). This is ignored if the `graph` argument is not a null pointer and the given graph is undirected.

loops: Boolean scalar, whether to consider loop edges in the calculation.

res: Pointer to a real variable, the result is stored here.

Returns:

Error code.

Time complexity: $O(1)$.

See also:

`igraph_centralization_degree()` and `igraph_centralization()`.

`igraph_centralization_betweenness_tmax` — Theoretical maximum for graph centralization based on betweenness

```
int igraph_centralization_betweenness_tmax(const igraph_t *graph,
      igraph_integer_t nodes,
      igraph_bool_t directed,
      igraph_real_t *res);
```

This function returns the theoretical maximum graph centrality based on vertex betweenness.

There are two ways to call this function, the first is to supply a graph as the `graph` argument, and then the number of vertices is taken from this object, and its directedness is considered as well. The `nodes` argument is ignored in this case. The `directed` argument is also ignored if the supplied graph is undirected.

The other way is to supply a null pointer as the `graph` argument. In this case the `nodes` and `directed` arguments are considered.

The most centralized structure is the star.

Arguments:

- graph*: A graph object or a null pointer, see the description above.
- nodes*: The number of nodes. This is ignored if the *graph* argument is not a null pointer.
- directed*: Boolean scalar, whether to use directed paths in the betweenness calculation. This argument is ignored if *graph* is not a null pointer and it is undirected.
- res*: Pointer to a real variable, the result is stored here.

Returns:

Error code.

Time complexity: $O(1)$.

See also:

`igraph_centralization_betweenness()` and `igraph_centralization()`.

`igraph_centralization_closeness_tmax` — Theoretical maximum for graph centralization based on closeness

```
int igraph_centralization_closeness_tmax(const igraph_t *graph,
    igraph_integer_t nodes,
    igraph_neimode_t mode,
    igraph_real_t *res);
```

This function returns the theoretical maximum graph centrality based on vertex closeness.

There are two ways to call this function, the first is to supply a graph as the *graph* argument, and then the number of vertices is taken from this object, and its directedness is considered as well. The *nodes* argument is ignored in this case. The *mode* argument is also ignored if the supplied graph is undirected.

The other way is to supply a null pointer as the *graph* argument. In this case the *nodes* and *mode* arguments are considered.

The most centralized structure is the star.

Arguments:

- graph*: A graph object or a null pointer, see the description above.
- nodes*: The number of nodes. This is ignored if the *graph* argument is not a null pointer.
- mode*: Constant, specifies what kind of distances to consider to calculate closeness. See the *mode* argument of `igraph_closeness()` for details. This argument is ignored if *graph* is not a null pointer and it is undirected.
- res*: Pointer to a real variable, the result is stored here.

Returns:

Error code.

Time complexity: $O(1)$.

See also:

`igraph_centralization_closeness()` and `igraph_centralization()`.

`igraph_centralization_eigenvector_centrality_tmax` — Theoretical maximum centralization for eigenvector centrality

```
int igraph_centralization_eigenvector_centrality_tmax(
    const igraph_t *graph,
    igraph_integer_t nodes,
    igraph_bool_t directed,
    igraph_bool_t scale,
    igraph_real_t *res);
```

This function returns the theoretical maximum graph centrality based on vertex eigenvector centrality.

There are two ways to call this function, the first is to supply a graph as the `graph` argument, and then the number of vertices is taken from this object, and its directedness is considered as well. The `nodes` argument is ignored in this case. The `directed` argument is also ignored if the supplied graph is undirected.

The other way is to supply a null pointer as the `graph` argument. In this case the `nodes` and `directed` arguments are considered.

The most centralized directed structure is the in-star. The most centralized undirected structure is the graph with a single edge.

Arguments:

graph: A graph object or a null pointer, see the description above.

nodes: The number of nodes. This is ignored if the `graph` argument is not a null pointer.

directed: Boolean scalar, whether to consider edge directions. This argument is ignored if `graph` is not a null pointer and it is undirected.

scale: Whether to rescale the node-level centrality scores to have a maximum of one.

res: Pointer to a real variable, the result is stored here.

Returns:

Error code.

Time complexity: $O(1)$.

See also:

`igraph_centralization_closeness()` and `igraph_centralization()`.

Similarity measures

igraph_bibcoupling — Bibliographic coupling.

```
int igraph_bibcoupling(const igraph_t *graph, igraph_matrix_t *res,
                      const igraph_vs_t vids);
```

The bibliographic coupling of two vertices is the number of other vertices they both cite, `igraph_bibcoupling()` calculates this. The bibliographic coupling score for each given vertex and all other vertices in the graph will be calculated.

Arguments:

graph: The graph object to analyze.

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows is the same as the number of vertex ids in *vids*, the number of columns is the number of vertices in the graph.

vids: The vertex ids of the vertices for which the calculation will be done.

Returns:

Error code: `IGRAPH_EINVVID`: invalid vertex id.

Time complexity: $O(|V|d^2)$, $|V|$ is the number of vertices in the graph, d is the (maximum) degree of the vertices in the graph.

See also:

`igraph_cocitation()`

Example 13.17. File `examples/simple/igraph_cocitation.c`

igraph_cocitation — Cocitation coupling.

```
int igraph_cocitation(const igraph_t *graph, igraph_matrix_t *res,
                     const igraph_vs_t vids);
```

Two vertices are cocited if there is another vertex citing both of them. `igraph_cocitation()` simply counts how many times two vertices are cocited. The cocitation score for each given vertex and all other vertices in the graph will be calculated.

Arguments:

graph: The graph object to analyze.

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows is the same as the number of vertex ids in *vids*, the number of columns is the number of vertices in the graph.

vids: The vertex ids of the vertices for which the calculation will be done.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex id.

Time complexity: $O(|V|d^2)$, $|V|$ is the number of vertices in the graph, d is the (maximum) degree of the vertices in the graph.

See also:

`igraph_bibcoupling()`

Example 13.18. File `examples/simple/igraph_cocitation.c`

igraph_similarity_jaccard — Jaccard similarity coefficient for the given vertices.

```
int igraph_similarity_jaccard(const igraph_t *graph, igraph_matrix_t *res,
                             const igraph_vs_t vids, igraph_neimode_t mode, ig
```

The Jaccard similarity coefficient of two vertices is the number of common neighbors divided by the number of vertices that are neighbors of at least one of the two vertices being considered. This function calculates the pairwise Jaccard similarities for some (or all) of the vertices.

Arguments:

graph: The graph object to analyze

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows and columns is the same as the number of vertex ids in *vids*.

vids: The vertex ids of the vertices for which the calculation will be done.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH_IN the incoming edges will be considered for each node.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

loops: Whether to include the vertices themselves in the neighbor sets.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex id passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|V|^2 d)$, $|V|$ is the number of vertices in the vertex iterator given, d is the (maximum) degree of the vertices in the graph.

See also:

`igraph_similarity_dice()`, a measure very similar to the Jaccard coefficient

Example 13.19. File `examples/simple/igraph_similarity.c`

`igraph_similarity_jaccard_pairs` — Jaccard similarity coefficient for given vertex pairs.

```
int igraph_similarity_jaccard_pairs(const igraph_t *graph, igraph_vector_t *res,
                                   const igraph_vector_t *pairs, igraph_neimod_t mode,
```

The Jaccard similarity coefficient of two vertices is the number of common neighbors divided by the number of vertices that are neighbors of at least one of the two vertices being considered. This function calculates the pairwise Jaccard similarities for a list of vertex pairs.

Arguments:

graph: The graph object to analyze

res: Pointer to a vector, the result of the calculation will be stored here. The number of elements is the same as the number of pairs in *pairs*.

pairs: A vector that contains the pairs for which the similarity will be calculated. Each pair is defined by two consecutive elements, i.e. the first and second element of the vector specifies the first pair, the third and fourth element specifies the second pair and so on.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

`IGRAPH_OUT` the outgoing edges will be considered for each node.

`IGRAPH_IN` the incoming edges will be considered for each node.

`IGRAPH_ALL` the directed graph is considered as an undirected one for the computation.

loops: Whether to include the vertices themselves in the neighbor sets.

Returns:

Error code:

`IGRAPH_ENOMEM` not enough memory for temporary data.

`IGRAPH_EINVVID` invalid vertex id passed.

`IGRAPH_EINVMODE` invalid mode argument.

Time complexity: $O(nd)$, n is the number of pairs in the given vector, d is the (maximum) degree of the vertices in the graph.

See also:

`igraph_similarity_jaccard()` to calculate the Jaccard similarity between all pairs of a vertex set, or `igraph_similarity_dice()` and `igraph_similarity_dice_pairs()` for a measure very similar to the Jaccard coefficient

Example 13.20. File `examples/simple/igraph_similarity.c`

igraph_similarity_jaccard_es — Jaccard similarity coefficient for a given edge selector.

```
int igraph_similarity_jaccard_es(const igraph_t *graph, igraph_vector_t *res,
                                const igraph_es_t es, igraph_neimode_t mode, i
```

The Jaccard similarity coefficient of two vertices is the number of common neighbors divided by the number of vertices that are neighbors of at least one of the two vertices being considered. This function calculates the pairwise Jaccard similarities for the endpoints of edges in a given edge selector.

Arguments:

graph: The graph object to analyze

res: Pointer to a vector, the result of the calculation will be stored here. The number of elements is the same as the number of edges in *es*.

es: An edge selector that specifies the edges to be included in the result.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH_IN the incoming edges will be considered for each node.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

loops: Whether to include the vertices themselves in the neighbor sets.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex id passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(nd)$, n is the number of edges in the edge selector, d is the (maximum) degree of the vertices in the graph.

See also:

`igraph_similarity_jaccard()` and `igraph_similarity_jaccard_pairs()` to calculate the Jaccard similarity between all pairs of a vertex set or some selected vertex pairs, or `igraph_similarity_dice()`, `igraph_similarity_dice_pairs()` and `igraph_similarity_dice_es()` for a measure very similar to the Jaccard coefficient

Example 13.21. File `examples/simple/igraph_similarity.c`

igraph_similarity_dice — Dice similarity coefficient.


```
int igraph_similarity_dice(const igraph_t *graph, igraph_matrix_t *res,
                          const igraph_vs_t vids, igraph_neimode_t mode, igraph_t *g)
```

The Dice similarity coefficient of two vertices is twice the number of common neighbors divided by the sum of the degrees of the vertices. This function calculates the pairwise Dice similarities for some (or all) of the vertices.

Arguments:

graph: The graph object to analyze

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows and columns is the same as the number of vertex ids in *vids*.

vids: The vertex ids of the vertices for which the calculation will be done.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH_IN the incoming edges will be considered for each node.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

loops: Whether to include the vertices themselves as their own neighbors.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVAL invalid vertex id passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|V|^2 d)$, $|V|$ is the number of vertices in the vertex iterator given, d is the (maximum) degree of the vertices in the graph.

See also:

`igraph_similarity_jaccard()`, a measure very similar to the Dice coefficient

Example 13.22. File `examples/simple/igraph_similarity.c`

`igraph_similarity_dice_pairs` — Dice similarity coefficient for given vertex pairs.

```
int igraph_similarity_dice_pairs(const igraph_t *graph, igraph_vector_t *res,
                                const igraph_vector_t *pairs, igraph_neimode_t mode, igraph_t *g)
```

The Dice similarity coefficient of two vertices is twice the number of common neighbors divided by the sum of the degrees of the vertices. This function calculates the pairwise Dice similarities for a list of vertex pairs.

Arguments:

- graph*: The graph object to analyze
- res*: Pointer to a vector, the result of the calculation will be stored here. The number of elements is the same as the number of pairs in *pairs*.
- pairs*: A vector that contains the pairs for which the similarity will be calculated. Each pair is defined by two consecutive elements, i.e. the first and second element of the vector specifies the first pair, the third and fourth element specifies the second pair and so on.
- mode*: The type of neighbors to be used for the calculation in directed graphs. Possible values:
- IGRAPH_OUT the outgoing edges will be considered for each node.
 - IGRAPH_IN the incoming edges will be considered for each node.
 - IGRAPH_ALL the directed graph is considered as an undirected one for the computation.
- loops*: Whether to include the vertices themselves as their own neighbors.

Returns:

Error code:

- IGRAPH_ENOMEM not enough memory for temporary data.
- IGRAPH_EINVVID invalid vertex id passed.
- IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(nd)$, n is the number of pairs in the given vector, d is the (maximum) degree of the vertices in the graph.

See also:

`igraph_similarity_dice()` to calculate the Dice similarity between all pairs of a vertex set, or `igraph_similarity_jaccard()`, `igraph_similarity_jaccard_pairs()` and `igraph_similarity_jaccard_es()` for a measure very similar to the Dice coefficient

Example 13.23. File `examples/simple/igraph_similarity.c`

`igraph_similarity_dice_es` — Dice similarity coefficient for a given edge selector.

```
int igraph_similarity_dice_es(const igraph_t *graph, igraph_vector_t *res,
                             const igraph_es_t es, igraph_neimode_t mode, igraph_neimode_t mode2, igraph_neimode_t mode3)
```

The Dice similarity coefficient of two vertices is twice the number of common neighbors divided by the sum of the degrees of the vertices. This function calculates the pairwise Dice similarities for the endpoints of edges in a given edge selector.

Arguments:

- graph*: The graph object to analyze

res: Pointer to a vector, the result of the calculation will be stored here. The number of elements is the same as the number of edges in *es*.

es: An edge selector that specifies the edges to be included in the result.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

- `IGRAPH_OUT` the outgoing edges will be considered for each node.
- `IGRAPH_IN` the incoming edges will be considered for each node.
- `IGRAPH_ALL` the directed graph is considered as an undirected one for the computation.

loops: Whether to include the vertices themselves as their own neighbors.

Returns:

Error code:

- `IGRAPH_ENOMEM` not enough memory for temporary data.
- `IGRAPH_EINVAL` invalid vertex id passed.
- `IGRAPH_EINVMODE` invalid mode argument.

Time complexity: $O(nd)$, n is the number of pairs in the given vector, d is the (maximum) degree of the vertices in the graph.

See also:

`igraph_similarity_dice()` and `igraph_similarity_dice_pairs()` to calculate the Dice similarity between all pairs of a vertex set or some selected vertex pairs, or `igraph_similarity_jaccard()`, `igraph_similarity_jaccard_pairs()` and `igraph_similarity_jaccard_es()` for a measure very similar to the Dice coefficient

Example 13.24. File `examples/simple/igraph_similarity.c`

`igraph_similarity_inverse_log_weighted` — Vertex similarity based on the inverse logarithm of vertex degrees.

```
int igraph_similarity_inverse_log_weighted(const igraph_t *graph,
    igraph_matrix_t *res, const igraph_vs_t vids, igraph_neimode_t mode);
```

The inverse log-weighted similarity of two vertices is the number of their common neighbors, weighted by the inverse logarithm of their degrees. It is based on the assumption that two vertices should be considered more similar if they share a low-degree common neighbor, since high-degree common neighbors are more likely to appear even by pure chance.

Isolated vertices will have zero similarity to any other vertex. Self-similarities are not calculated.

See the following paper for more details: Lada A. Adamic and Eytan Adar: Friends and neighbors on the Web. *Social Networks*, 25(3):211-230, 2003.

Arguments:

graph: The graph object to analyze.

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows is the same as the number of vertex ids in *vids*, the number of columns is the number of vertices in the graph.

vids: The vertex ids of the vertices for which the calculation will be done.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT	the outgoing edges will be considered for each node. Nodes will be weighted according to their in-degree.
IGRAPH_IN	the incoming edges will be considered for each node. Nodes will be weighted according to their out-degree.
IGRAPH_ALL	the directed graph is considered as an undirected one for the computation. Every node is weighted according to its undirected degree.

Returns:

Error code: IGRAPH_EINVAL: invalid vertex id.

Time complexity: $O(|V|d^2)$, $|V|$ is the number of vertices in the graph, d is the (maximum) degree of the vertices in the graph.

Example 13.25. File `examples/simple/igraph_similarity.c`

Trees

`igraph_minimum_spanning_tree` — Calculates one minimum spanning tree of a graph.

```
int igraph_minimum_spanning_tree(const igraph_t* graph,
                                igraph_vector_t* res, const igraph_vector_t* w
```

If the graph has more minimum spanning trees (this is always the case, except if it is a forest) this implementation returns only the same one.

Directed graphs are considered as undirected for this computation.

If the graph is not connected then its minimum spanning forest is returned. This is the set of the minimum spanning trees of each component.

Arguments:

graph: The graph object.

res: An initialized vector, the IDs of the edges that constitute a spanning tree will be returned here. Use `igraph_subgraph_edges()` to extract the spanning tree as a separate graph object.

weights: A vector containing the weights of the edges in the same order as the simple edge iterator visits them (i.e. in increasing order of edge IDs).

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: $O(|V|+|E|)$ for the unweighted case, $O(|E| \log |V|)$ for the weighted case. $|V|$ is the number of vertices, $|E|$ the number of edges in the graph.

See also:

`igraph_minimum_spanning_tree_unweighted()` and `igraph_minimum_spanning_tree_prim()` if you only need the tree as a separate graph object.

Example	13.26.	File	examples/simple/
<code>igraph_minimum_spanning_tree.c</code>			

`igraph_minimum_spanning_tree_unweighted` — Calculates one minimum spanning tree of an unweighted graph.

```
int igraph_minimum_spanning_tree_unweighted(const igraph_t *graph,
                                             igraph_t *mst);
```

If the graph has more minimum spanning trees (this is always the case, except if it is a forest) this implementation returns only the same one.

Directed graphs are considered as undirected for this computation.

If the graph is not connected then its minimum spanning forest is returned. This is the set of the minimum spanning trees of each component.

Arguments:

graph: The graph object.

mst: The minimum spanning tree, another graph object. Do *not* initialize this object before passing it to this function, but be sure to call `igraph_destroy()` on it if you don't need it any more.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of vertices, $|E|$ the number of edges in the graph.

See also:

`igraph_minimum_spanning_tree_prim()` for weighted graphs, `igraph_minimum_spanning_tree()` if you need the IDs of the edges that constitute the spanning tree.

`igraph_minimum_spanning_tree_prim` — Calculates one minimum spanning tree of a weighted graph.

```
int igraph_minimum_spanning_tree_prim(const igraph_t *graph, igraph_t *mst,
                                     const igraph_vector_t *weights);
```

This function uses Prim's method for carrying out the computation, see Prim, R.C.: Shortest connection networks and some generalizations, Bell System Technical Journal, Vol. 36, 1957, 1389--1401.

If the graph has more than one minimum spanning tree, the current implementation returns always the same one.

Directed graphs are considered as undirected for this computation.

If the graph is not connected then its minimum spanning forest is returned. This is the set of the minimum spanning trees of each component.

Arguments:

graph: The graph object.

mst: The result of the computation, a graph object containing the minimum spanning tree of the graph. Do *not* initialize this object before passing it to this function, but be sure to call `igraph_destroy()` on it if you don't need it any more.

weights: A vector containing the weights of the edges in the same order as the simple edge iterator visits them (i.e. in increasing order of edge IDs).

Returns:

Error code: `IGRAPH_ENOMEM`, not enough memory. `IGRAPH_EINVAL`, length of weight vector does not match number of edges.

Time complexity: $O(|E| \log |V|)$, $|V|$ is the number of vertices, $|E|$ the number of edges in the graph.

See also:

`igraph_minimum_spanning_tree_unweighted()` for unweighted graphs, `igraph_minimum_spanning_tree()` if you need the IDs of the edges that constitute the spanning tree.

Example 13.27. File `examples/simple/igraph_minimum_spanning_tree.c`

`igraph_random_spanning_tree` — Uniformly sample the spanning trees of a graph

```
int igraph_random_spanning_tree(const igraph_t *graph, igraph_vector_t *res, ig
```

Performs a loop-erased random walk on the graph to uniformly sample its spanning trees. Edge directions are ignored.

Multi-graphs are supported, and edge multiplicities will affect the sampling frequency. For example, consider the 3-cycle graph $1=2-3-1$, with two edges between vertices 1 and 2. Due to these parallel edges, the trees $1-2-3$ and $3-1-2$ will be sampled with multiplicity 2, while the tree $2-3-1$ will be sampled with multiplicity 1.

Arguments:

- graph*: The input graph. Edge directions are ignored.
- res*: An initialized vector, the IDs of the edges that constitute a spanning tree will be returned here. Use `igraph_subgraph_edges()` to extract the spanning tree as a separate graph object.
- vid*: This parameter is relevant if the graph is not connected. If negative, a random spanning forest of all components will be generated. Otherwise, it should be the ID of a vertex. A random spanning tree of the component containing the vertex will be generated.

Returns:

Error code.

See also:

`igraph_minimum_spanning_tree()`, `igraph_random_walk()`

`igraph_is_tree` — Decides whether the graph is a tree.

```
int igraph_is_tree(const igraph_t *graph, igraph_bool_t *res, igraph_integer_t
```

An undirected graph is a tree if it is connected and has no cycles.

In the directed case, a possible additional requirement is that all edges are oriented away from a root (out-tree or arborescence) or all edges are oriented towards a root (in-tree or anti-arborescence). This test can be controlled using the *mode* parameter.

By convention, the null graph (i.e. the graph with no vertices) is considered not to be a tree.

Arguments:

- graph*: The graph object to analyze.
- res*: Pointer to a logical variable, the result will be stored here.
- root*: If not NULL, the root node will be stored here. When *mode* is `IGRAPH_ALL` or the graph is undirected, any vertex can be the root and *root* is set to 0 (the first vertex). When *mode* is `IGRAPH_OUT` or `IGRAPH_IN`, the root is set to the vertex with zero in- or out-degree, respectively.
- mode*: For a directed graph this specifies whether to test for an out-tree, an in-tree or ignore edge directions. The respective possible values are: `IGRAPH_OUT`, `IGRAPH_IN`, `IGRAPH_ALL`. This argument is ignored for undirected graphs.

Returns:

Error code: `IGRAPH_EINVAL`: invalid mode argument.

Time complexity: At most $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph.

See also:

```
igraph_is_weakly_connected()
```

Example 13.28. File `examples/simple/igraph_tree.c`

igraph_to_prufer — Converts a tree to its Prüfer sequence

```
int igraph_to_prufer(const igraph_t *graph, igraph_vector_int_t* prufer);
```

A Prüfer sequence is a unique sequence of integers associated with a labelled tree. A tree on $n \geq 2$ vertices can be represented by a sequence of $n-2$ integers, each between 0 and $n-1$ (inclusive).

Arguments:

graph: Pointer to an initialized graph object which must be a tree on $n \geq 2$ vertices.

prufer: A pointer to the integer vector that should hold the Prüfer sequence; the vector must be initialized and will be resized to $n - 2$.

Returns:

Error code:

IGRAPH_ENOMEM there is not enough memory to perform the operation.

IGRAPH_EINVAL the graph is not a tree or it has less than vertices

See also:

```
igraph_from_prufer()
```

Transitivity or clustering coefficient

igraph_transitivity_undirected — Calculates the transitivity (clustering coefficient) of a graph.

```
int igraph_transitivity_undirected(const igraph_t *graph,
                                   igraph_real_t *res,
                                   igraph_transitivity_mode_t mode);
```

The transitivity measures the probability that two neighbors of a vertex are connected. More precisely, this is the ratio of the triangles and connected triples in the graph, the result is a single real number. Directed graphs are considered as undirected ones and multi-edges are ignored.

Note that this measure is different from the local transitivity measure (see `igraph_transitivity_local_undirected()`) as it calculates a single value for the whole graph.

Clustering coefficient is an alternative name for transitivity.

References:

S. Wasserman and K. Faust: Social Network Analysis: Methods and Applications. Cambridge: Cambridge University Press, 1994.

Arguments:

graph: The graph object. Edge directions and multiplicities are ignored.

res: Pointer to a real variable, the result will be stored here.

mode: Defines how to treat graphs with no connected triples. IGRAPH_TRANSITIVITY_NAN returns NaN in this case, IGRAPH_TRANSITIVITY_ZERO returns zero.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory for temporary data.

See also:

`igraph_transitivity_local_undirected()`, `igraph_transitivity_avglocal_undirected()`.

Time complexity: $O(|V|*d^2)$, $|V|$ is the number of vertices in the graph, d is the average node degree.

Example 13.29. File `examples/simple/igraph_transitivity.c`

`igraph_transitivity_local_undirected` — Calculates the local transitivity (clustering coefficient) of a graph.

```
int igraph_transitivity_local_undirected(const igraph_t *graph,
    igraph_vector_t *res,
    const igraph_vs_t vids,
    igraph_transitivity_mode_t mode);
```

The transitivity measures the probability that two neighbors of a vertex are connected. In case of the local transitivity, this probability is calculated separately for each vertex.

Note that this measure is different from the global transitivity measure (see `igraph_transitivity_undirected()`) as it calculates a transitivity value for each vertex individually.

Clustering coefficient is an alternative name for transitivity.

References:

D. J. Watts and S. Strogatz: Collective dynamics of small-world networks. Nature 393(6684):440-442 (1998).

Arguments:

graph: The input graph. Edge directions and multiplicities are ignored.

res: Pointer to an initialized vector, the result will be stored here. It will be resized as needed.

vids: Vertex set, the vertices for which the local transitivity will be calculated.

mode: Defines how to treat vertices with degree less than two. `IGRAPH_TRANSITIVITY_NAN` returns NaN for these vertices, `IGRAPH_TRANSITIVITY_ZERO` returns zero.

Returns:

Error code.

See also:

`igraph_transitivity_undirected()`, `igraph_transitivity_avglocal_undirected()`.

Time complexity: $O(n*d^2)$, n is the number of vertices for which the transitivity is calculated, d is the average vertex degree.

`igraph_transitivity_avglocal_undirected` — Average local transitivity (clustering coefficient).

```
int igraph_transitivity_avglocal_undirected(const igraph_t *graph,
                                           igraph_real_t *res,
                                           igraph_transitivity_mode_t mode);
```

The transitivity measures the probability that two neighbors of a vertex are connected. In case of the average local transitivity, this probability is calculated for each vertex and then the average is taken. Vertices with less than two neighbors require special treatment, they will either be left out from the calculation or they will be considered as having zero transitivity, depending on the mode argument. Edge directions and edge multiplicities are ignored.

Note that this measure is different from the global transitivity measure (see `igraph_transitivity_undirected()`) as it simply takes the average local transitivity across the whole network.

Clustering coefficient is an alternative name for transitivity.

References:

D. J. Watts and S. Strogatz: Collective dynamics of small-world networks. *Nature* 393(6684):440-442 (1998).

Arguments:

graph: The input graph. Edge directions and multiplicities are ignored.

res: Pointer to a real variable, the result will be stored here.

mode: Defines how to treat vertices with degree less than two. `IGRAPH_TRANSITIVITY_NAN` leaves them out from averaging, `IGRAPH_TRANSITIVITY_ZERO` includes them with zero transitivity. The result will be NaN if the mode is `IGRAPH_TRANSITIVITY_NAN` and there are no vertices with more than one neighbor.

Returns:

Error code.

See also:

```
igraph_transitivity_undirected(), igraph_transitivity_local_undirected().
```

Time complexity: $O(|V|*d^2)$, $|V|$ is the number of vertices in the graph and d is the average degree.

igraph_transitivity_barrat — Weighted transitivity, as defined by A. Barrat.

```
int igraph_transitivity_barrat(const igraph_t *graph,
                              igraph_vector_t *res,
                              const igraph_vs_t vids,
                              const igraph_vector_t *weights,
                              igraph_transitivity_mode_t mode);
```

This is a local transitivity, i.e. a vertex-level index. For a given vertex i , from all triangles in which it participates we consider the weight of the edges incident on i . The transitivity is the sum of these weights divided by twice the strength of the vertex (see `igraph_strength()`) and the degree of the vertex minus one. See Alain Barrat, Marc Barthélemy, Romualdo Pastor-Satorras, Alessandro Vespignani: The architecture of complex weighted networks, Proc. Natl. Acad. Sci. USA 101, 3747 (2004) at <http://arxiv.org/abs/cond-mat/0311416> for the exact formula.

Arguments:

- graph*: The input graph. Edge directions are ignored for directed graphs. Note that the function does *not* work for non-simple graphs.
- res*: Pointer to an initialized vector, the result will be stored here. It will be resized as needed.
- vids*: The vertices for which the calculation is performed.
- weights*: Edge weights. If this is a null pointer, then a warning is given and `igraph_transitivity_local_undirected()` is called.
- mode*: Defines how to treat vertices with zero strength. `IGRAPH_TRANSITIVITY_NAN` says that the transitivity of these vertices is NaN, `IGRAPH_TRANSITIVITY_ZERO` says it is zero.

Returns:

Error code.

Time complexity: $O(|V|*d^2)$, $|V|$ is the number of vertices in the graph, d is the average node degree.

See also:

`igraph_transitivity_undirected()`, `igraph_transitivity_local_undirected()` and `igraph_transitivity_avglocal_undirected()` for other kinds of (non-weighted) transitivity.

Directedness conversion

igraph_to_directed — Convert an undirected graph to a directed one

```
int igraph_to_directed(igraph_t *graph,
                      igraph_to_directed_t mode);
```

If the supplied graph is directed, this function does nothing.

Arguments:

graph: The graph object to convert.

mode: Constant, specifies the details of how exactly the conversion is done. Possible values:

IGRAPH_TO_DIRECTED_ARBITRARY	The number of edges in the graph stays the same, an arbitrarily directed edge is created for each undirected edge.
IGRAPH_TO_DIRECTED_MUTUAL	Two directed edges are created for each undirected edge, one in each direction.
IGRAPH_TO_DIRECTED_RANDOM	Each undirected edge is converted to a randomly oriented directed one.
IGRAPH_TO_DIRECTED_ACYCLIC	Each undirected edge is converted to a directed edge oriented from a lower index vertex to a higher index one. If no self-loops were present, then the result is a directed acyclic graph.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

igraph_to_undirected — Convert a directed graph to an undirected one.

```
int igraph_to_undirected(igraph_t *graph,
                        igraph_to_undirected_t mode,
                        const igraph_attribute_combination_t *edge_comb);
```

If the supplied graph is undirected, this function does nothing.

Arguments:

graph: The graph object to convert.

mode: Constant, specifies the details of how exactly the conversion is done. Possible values: IGRAPH_TO_UNDIRECTED_EACH: the number of edges remains constant, an undirected edge is created for each directed one, this version might create graphs with multiple edges; IGRAPH_TO_UNDIRECTED_COLLAPSE: one undirected edge will be created for each pair of vertices which are connected with at least one directed edge, no multiple edges will be created. IGRAPH_TO_UNDIRECTED_MUTUAL creates an undirected edge for each pair of mutual edges in the directed graph. Non-mutual edges are lost. This mode might create multiple edges.

edge_comb: What to do with the edge attributes. See the igraph manual section about attributes for details.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

Example 13.30. File `examples/simple/igraph_to_undirected.c`

Spectral properties

`igraph_laplacian` — Returns the Laplacian matrix of a graph

```
int igraph_laplacian(const igraph_t *graph, igraph_matrix_t *res,
                    igraph_sparsemat_t *sparseres,
                    igraph_bool_t normalized,
                    const igraph_vector_t *weights);
```

The graph Laplacian matrix is similar to an adjacency matrix but contains -1's instead of 1's and the vertex degrees are included in the diagonal. So the result for edge $i-j$ is -1 if $i \neq j$ and is equal to the degree of vertex i if $i=j$. `igraph_laplacian` will work on a directed graph; in this case, the diagonal will contain the out-degrees. Loop edges will be ignored.

The normalized version of the Laplacian matrix has 1 in the diagonal and $-1/\sqrt{d[i]d[j]}$ if there is an edge from i to j .

The first version of this function was written by Vincent Matossian.

Arguments:

graph: Pointer to the graph to convert.

res: Pointer to an initialized matrix object, the result is stored here. It will be resized if needed. If it is a null pointer, then it is ignored. At least one of *res* and *sparseres* must be a non-null pointer.

sparseres: Pointer to an initialized sparse matrix object, the result is stored here, if it is not a null pointer. At least one of *res* and *sparseres* must be a non-null pointer.

normalized: Whether to create a normalized Laplacian matrix.

weights: An optional vector containing edge weights, to calculate the weighted Laplacian matrix. Set it to a null pointer to calculate the unweighted Laplacian.

Returns:

Error code.

Time complexity: $O(|V||V|)$, $|V|$ is the number of vertices in the graph.

Example 13.31. File `examples/simple/igraph_laplacian.c`

Non-simple graphs: Multiple and loop edges

`igraph_is_simple` — Decides whether the input graph is a simple graph.

```
int igraph_is_simple(const igraph_t *graph, igraph_bool_t *res);
```

A graph is a simple graph if it does not contain loop edges and multiple edges.

Arguments:

graph: The input graph.

res: Pointer to a boolean constant, the result is stored here.

Returns:

Error code.

See also:

`igraph_is_loop()` and `igraph_is_multiple()` to find the loops and multiple edges, `igraph_simplify()` to get rid of them, or `igraph_has_multiple()` to decide whether there is at least one multiple edge.

Time complexity: $O(|V|+|E|)$.

`igraph_is_loop` — Find the loop edges in a graph.

```
int igraph_is_loop(const igraph_t *graph, igraph_vector_bool_t *res,
                  igraph_es_t es);
```

A loop edge is an edge from a vertex to itself.

Arguments:

graph: The input graph.

res: Pointer to an initialized boolean vector for storing the result, it will be resized as needed.

es: The edges to check, for all edges supply `igraph_ess_all()` here.

Returns:

Error code.

See also:

`igraph_simplify()` to get rid of loop edges.

Time complexity: $O(e)$, the number of edges to check.

Example 13.32. File `examples/simple/igraph_is_loop.c`

`igraph_is_multiple` — Find the multiple edges in a graph.

```
int igraph_is_multiple(const igraph_t *graph, igraph_vector_bool_t *res,
                      igraph_es_t es);
```

An edge is a multiple edge if there is another edge with the same head and tail vertices in the graph.

Note that this function returns true only for the second or more appearances of the multiple edges.

Arguments:

graph: The input graph.

res: Pointer to a boolean vector, the result will be stored here. It will be resized as needed.

es: The edges to check. Supply `igraph_ess_all()` if you want to check all edges.

Returns:

Error code.

See also:

`igraph_count_multiple()`, `igraph_has_multiple()` and `igraph_simplify()`.

Time complexity: $O(e*d)$, e is the number of edges to check and d is the average degree (out-degree in directed graphs) of the vertices at the tail of the edges.

Example 13.33. File `examples/simple/igraph_is_multiple.c`

`igraph_has_multiple` — Check whether the graph has at least one multiple edge.

```
int igraph_has_multiple(const igraph_t *graph, igraph_bool_t *res);
```

An edge is a multiple edge if there is another edge with the same head and tail vertices in the graph.

Arguments:

graph: The input graph.

res: Pointer to a boolean variable, the result will be stored here.

Returns:

Error code.

See also:

`igraph_count_multiple()`, `igraph_is_multiple()` and `igraph_simplify()`.

Time complexity: $O(e*d)$, e is the number of edges to check and d is the average degree (out-degree in directed graphs) of the vertices at the tail of the edges.

Example 13.34. File `examples/simple/igraph_has_multiple.c`

`igraph_count_multiple` — Count the number of appearances of the edges in a graph.

```
int igraph_count_multiple(const igraph_t *graph, igraph_vector_t *res, igraph_edges_t es)
```

If the graph has no multiple edges then the result vector will be filled with ones. (An edge is a multiple edge if there is another edge with the same head and tail vertices in the graph.)

Arguments:

graph: The input graph.

res: Pointer to a vector, the result will be stored here. It will be resized as needed.

es: The edges to check. Supply `igraph_ess_all()` if you want to check all edges.

Returns:

Error code.

See also:

`igraph_is_multiple()` and `igraph_simplify()`.

Time complexity: $O(E d)$, E is the number of edges to check and d is the average degree (out-degree in directed graphs) of the vertices at the tail of the edges.

Mixing patterns

`igraph_assortativity_nominal` — Assortativity of a graph based on vertex categories

```
int igraph_assortativity_nominal(const igraph_t *graph,
                                const igraph_vector_t *types,
                                igraph_real_t *res,
                                igraph_bool_t directed);
```


Assuming the vertices of the input graph belong to different categories, this function calculates the assortativity coefficient of the graph. The assortativity coefficient is between minus one and one and it is one if all connections stay within categories, it is minus one, if the network is perfectly disassortative. For a randomly connected network it is (asymptotically) zero.

See equation (2) in M. E. J. Newman: Mixing patterns in networks, Phys. Rev. E 67, 026126 (2003) (<http://arxiv.org/abs/cond-mat/0209450>) for the proper definition.

Arguments:

- graph*: The input graph, it can be directed or undirected.
- types*: Vector giving the vertex types. They are assumed to be integer numbers, starting with zero.
- res*: Pointer to a real variable, the result is stored here.
- directed*: Boolean, it gives whether to consider edge directions in a directed graph. It is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(|E|+t)$, $|E|$ is the number of edges, t is the number of vertex types.

See also:

`igraph_assortativity` if the vertex types are defines by numeric values (e.g. vertex degree), instead of categories.

Example 13.35. File `examples/simple/assortativity.c`

`igraph_assortativity` — Assortativity based on numeric properties of vertices

```
int igraph_assortativity(const igraph_t *graph,
                        const igraph_vector_t *types1,
                        const igraph_vector_t *types2,
                        igraph_real_t *res,
                        igraph_bool_t directed);
```

This function calculates the assortativity coefficient of the input graph. This coefficient is basically the correlation between the actual connectivity patterns of the vertices and the pattern expected from the distribution of the vertex types.

See equation (21) in M. E. J. Newman: Mixing patterns in networks, Phys. Rev. E 67, 026126 (2003) (<http://arxiv.org/abs/cond-mat/0209450>) for the proper definition. The actual calculation is performed using equation (26) in the same paper for directed graphs, and equation (4) in M. E. J. Newman: Assortative mixing in networks, Phys. Rev. Lett. 89, 208701 (2002) (<http://arxiv.org/abs/cond-mat/0205405/>) for undirected graphs.

Arguments:

- graph*: The input graph, it can be directed or undirected.
- types1*: The vertex values, these can be arbitrary numeric values.

- types2*: A second value vector to be using for the incoming edges when calculating assortativity for a directed graph. Supply a null pointer here if you want to use the same values for outgoing and incoming edges. This argument is ignored (with a warning) if it is not a null pointer and undirected assortativity coefficient is being calculated.
- res*: Pointer to a real variable, the result is stored here.
- directed*: Boolean, whether to consider edge directions for directed graphs. It is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(|E|)$, linear in the number of edges of the graph.

See also:

`igraph_assortativity_nominal()` if you have discrete vertex categories instead of numeric labels, and `igraph_assortativity_degree()` for the special case of assortativity based on vertex degree.

Example 13.36. File `examples/simple/assortativity.c`

igraph_assortativity_degree — Assortativity of a graph based on vertex degree

```
int igraph_assortativity_degree(const igraph_t *graph,
                               igraph_real_t *res,
                               igraph_bool_t directed);
```

Assortativity based on vertex degree, please see the discussion at the documentation of `igraph_assortativity()` for details.

Arguments:

- graph*: The input graph, it can be directed or undirected.
- res*: Pointer to a real variable, the result is stored here.
- directed*: Boolean, whether to consider edge directions for directed graphs. This argument is ignored for undirected graphs. Supply 1 (=TRUE) here to do the natural thing, i.e. use directed version of the measure for directed graphs and the undirected version for undirected graphs.

Returns:

Error code.

Time complexity: $O(|E|+|V|)$, $|E|$ is the number of edges, $|V|$ is the number of vertices.

See also:

`igraph_assortativity()` for the general function calculating assortativity for any kind of numeric vertex values.

Example 13.37. File `examples/simple/assortativity.c`

K-Cores

`igraph_coreness` — Finding the coreness of the vertices in a network.

```
int igraph_coreness(const igraph_t *graph, igraph_vector_t *cores,
                    igraph_neimode_t mode);
```

The k -core of a graph is a maximal subgraph in which each vertex has at least degree k . (Degree here means the degree in the subgraph of course.). The coreness of a vertex is the highest order of a k -core containing the vertex.

This function implements the algorithm presented in Vladimir Batagelj, Matjaz Zaversnik: An $O(m)$ Algorithm for Cores Decomposition of Networks.

Arguments:

graph: The input graph.

cores: Pointer to an initialized vector, the result of the computation will be stored here. It will be resized as needed. For each vertex it contains the highest order of a core containing the vertex.

mode: For directed graph it specifies whether to calculate in-cores, out-cores or the undirected version. It is ignored for undirected graphs. Possible values: `IGRAPH_ALL` undirected version, `IGRAPH_IN` in-cores, `IGRAPH_OUT` out-cores.

Returns:

Error code.

Time complexity: $O(|E|)$, the number of edges.

Topological sorting, directed acyclic graphs

`igraph_is_dag` — Checks whether a graph is a directed acyclic graph (DAG) or not.

```
int igraph_is_dag(const igraph_t* graph, igraph_bool_t *res);
```

A directed acyclic graph (DAG) is a directed graph with no cycles.

Arguments:

graph: The input graph.

res: Pointer to a boolean constant, the result is stored here.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, where $|V|$ and $|E|$ are the number of vertices and edges in the original input graph.

See also:

`igraph_topological_sorting()` to get a possible topological sorting of a DAG.

igraph_topological_sorting — Calculate a possible topological sorting of the graph.

```
int igraph_topological_sorting(const igraph_t* graph, igraph_vector_t *res,
                              igraph_neimode_t mode);
```

A topological sorting of a directed acyclic graph (DAG) is a linear ordering of its vertices where each vertex comes before all nodes to which it has edges. Every DAG has at least one topological sort, and may have many. This function returns one possible topological sort among them. If the graph is not acyclic (it has at least one cycle), an error is raised.

Arguments:

graph: The input graph.

res: Pointer to a vector, the result will be stored here. It will be resized if needed.

mode: Specifies how to use the direction of the edges. For `IGRAPH_OUT`, the sorting order ensures that each vertex comes before all vertices to which it has edges, so vertices with no incoming edges go first. For `IGRAPH_IN`, it is quite the opposite: each vertex comes before all vertices from which it receives edges. Vertices with no outgoing edges go first.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, where $|V|$ and $|E|$ are the number of vertices and edges in the original input graph.

See also:

`igraph_is_dag()` if you are only interested in whether a given graph is a DAG or not, or `igraph_feedback_arc_set()` to find a set of edges whose removal makes the graph acyclic.

Example	13.38.	File	examples/simple/
igraph_topological_sorting.c			

igraph_feedback_arc_set — Calculates a feedback arc set of the graph using different

```
int igraph_feedback_arc_set(const igraph_t *graph, igraph_vector_t *result,
                           const igraph_vector_t *weights, igraph_fas_algorithm_t algo);
```

algorithms.

A feedback arc set is a set of edges whose removal makes the graph acyclic. We are usually interested in *minimum* feedback arc sets, i.e. sets of edges whose total weight is minimal among all the feedback arc sets.

For undirected graphs, the problem is simple: one has to find a maximum weight spanning tree and then remove all the edges not in the spanning tree. For directed graphs, this is an NP-hard problem, and various heuristics are usually used to find an approximate solution to the problem. This function implements a few of these heuristics.

Arguments:

graph: The graph object.

result: An initialized vector, the result will be returned here.

weights: Weight vector or NULL if no weights are specified.

algo: The algorithm to use to solve the problem if the graph is directed. Possible values:

IGRAPH_FAS_EXACT_IP	Finds a <i>minimum</i> feedback arc set using integer programming (IP). The complexity of this algorithm is exponential of course.
IGRAPH_FAS_APPROX_EADES	Finds a feedback arc set using the heuristic of Eades, Lin and Smyth (1993). This is guaranteed to be smaller than $ E /2 - V /6$, and it is linear in the number of edges (i.e. $O(E)$). For more details, see Eades P, Lin X and Smyth WF: A fast and effective heuristic for the feedback arc set problem. In: Proc Inf Process Lett 319-323, 1993.

Returns:

Error code: IGRAPH_EINVAL if an unknown method was specified or the weight vector is invalid.

Example 13.39. File `examples/simple/igraph_feedback_arc_set.c`

Example	13.40.	File	examples/simple/igraph_feedback_arc_set_ip.c
----------------	---------------	-------------	---

Time complexity: depends on *algo*, see the time complexities there.

Maximum cardinality search and chordal graphs

`igraph_maximum_cardinality_search` — Maximum cardinality search.

```
int igraph_maximum_cardinality_search(const igraph_t *graph,
                                     igraph_vector_t *alpha,
                                     igraph_vector_t *alpham1);
```

This function implements the maximum cardinality search algorithm. It computes a rank *alpha* for each vertex, such that visiting vertices in decreasing rank order corresponds to always choosing the vertex with the most already visited neighbors as the next one to visit.

Maximum cardinality search is useful in deciding the chordality of a graph. A graph is chordal if and only if any two neighbors of a vertex which are higher in rank than it are connected to each other.

References:

Robert E Tarjan and Mihalis Yannakakis: Simple linear-time algorithms to test chordality of graphs, test acyclicity of hypergraphs, and selectively reduce acyclic hypergraphs. SIAM Journal of Computation 13, 566--579, 1984. <https://doi.org/10.1137/0213035>

Arguments:

graph: The input graph. Edge directions will be ignored.

alpha: Pointer to an initialized vector, the result is stored here. It will be resized, as needed. Upon return it contains the rank of the each vertex in the range 0 to $n - 1$, where n is the number of vertices.

alpham1: Pointer to an initialized vector or a NULL pointer. If not NULL, then the inverse of *alpha* is stored here. In other words, the elements of *alpham1* are vertex IDs in reverse maximum cardinality search order.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in terms of the number of vertices and edges.

See also:

`igraph_is_chordal()`.

igraph_is_chordal — Decides whether a graph is chordal.

```
int igraph_is_chordal(const igraph_t *graph,
                     const igraph_vector_t *alpha,
                     const igraph_vector_t *alpham1,
                     igraph_bool_t *chordal,
                     igraph_vector_t *fill_in,
                     igraph_t *newgraph);
```

A graph is chordal if each of its cycles of four or more nodes has a chord, i.e. an edge joining two nodes that are not adjacent in the cycle. An equivalent definition is that any chordless cycles have at most three nodes. If either *alpha* or *alpham1* is given, then the other is calculated by taking simply the inverse. If neither are given, then `igraph_maximum_cardinality_search()` is called to calculate them.

Arguments:

graph: The input graph. Edge directions will be ignored.

alpha: Either an alpha vector coming from `igraph_maximum_cardinality_search()` (on the same graph), or a NULL pointer.

alpham1: Either an inverse alpha vector coming from `igraph_maximum_cardinality_search()` (on the same graph) or a NULL pointer.

chordal: Pointer to a boolean. If not NULL the result is stored here.

fill_in: Pointer to an initialized vector, or a NULL pointer. If not a NULL pointer, then the fill-in of the graph is stored here. The fill-in is the set of edges that are needed to make the graph chordal. The vector is resized as needed.

newgraph: Pointer to an uninitialized graph, or a NULL pointer. If not a null pointer, then a new triangulated graph is created here. This essentially means adding the fill-in edges to the original graph.

Returns:

Error code.

Time complexity: $O(n)$.

See also:

`igraph_maximum_cardinality_search()`.

Matchings

`igraph_is_matching` — Checks whether the given matching is valid for the given graph.

```
int igraph_is_matching(const igraph_t* graph,
                      const igraph_vector_bool_t* types, const igraph_vector_int_t* matching,
                      igraph_bool_t* result);
```

This function checks a matching vector and verifies whether its length matches the number of vertices in the given graph, its values are between -1 (inclusive) and the number of vertices (exclusive), and whether there exists a corresponding edge in the graph for every matched vertex pair. For bipartite graphs, it also verifies whether the matched vertices are in different parts of the graph.

Arguments:

graph: The input graph. It can be directed but the edge directions will be ignored.

types: If the graph is bipartite and you are interested in bipartite matchings only, pass the vertex types here. If the graph is non-bipartite, simply pass NULL.

matching: The matching itself. It must be a vector where element *i* contains the ID of the vertex that vertex *i* is matched to, or -1 if vertex *i* is unmatched.

result: Pointer to a boolean variable, the result will be returned here.

See also:

`igraph_is_maximal_matching()` if you are also interested in whether the matching is maximal (i.e. non-extendable).

Time complexity: $O(|V|+|E|)$ where $|V|$ is the number of vertices and $|E|$ is the number of edges.

Example **13.41.** **File** **examples/simple/**
igraph_maximum_bipartite_matching.c

igraph_is_maximal_matching — Checks whether a matching in a graph is maximal.

```
int igraph_is_maximal_matching(const igraph_t* graph,
                              const igraph_vector_bool_t* types, const igraph_
                              igraph_bool_t* result);
```

A matching is maximal if and only if there exists no unmatched vertex in a graph such that one of its neighbors is also unmatched.

Arguments:

graph: The input graph. It can be directed but the edge directions will be ignored.

types: If the graph is bipartite and you are interested in bipartite matchings only, pass the vertex types here. If the graph is non-bipartite, simply pass NULL.

matching: The matching itself. It must be a vector where element i contains the ID of the vertex that vertex i is matched to, or -1 if vertex i is unmatched.

result: Pointer to a boolean variable, the result will be returned here.

See also:

`igraph_is_matching()` if you are only interested in whether a matching vector is valid for a given graph.

Time complexity: $O(|V|+|E|)$ where $|V|$ is the number of vertices and $|E|$ is the number of edges.

Example **13.42.** **File** **examples/simple/**
igraph_maximum_bipartite_matching.c

igraph_maximum_bipartite_matching — Calculates a maximum matching in a bipartite graph.

```
int igraph_maximum_bipartite_matching(const igraph_t* graph,
                                      const igraph_vector_bool_t* types, igraph_
                                      igraph_real_t* matching_weight, igraph_
                                      const igraph_vector_t* weights, igraph_re
```

A matching in a bipartite graph is a partial assignment of vertices of the first kind to vertices of the second kind such that each vertex of the first kind is matched to at most one vertex of the second kind and vice versa, and matched vertices must be connected by an edge in the graph. The size (or

cardinality) of a matching is the number of edges. A matching is a maximum matching if there exists no other matching with larger cardinality. For weighted graphs, a maximum matching is a matching whose edges have the largest possible total weight among all possible matchings.

Maximum matchings in bipartite graphs are found by the push-relabel algorithm with greedy initialization and a global relabeling after every $n/2$ steps where n is the number of vertices in the graph.

References: Cherkassky BV, Goldberg AV, Martin P, Setubal JC and Stolfi J: Augment or push: A computational study of bipartite matching and unit-capacity flow algorithms. ACM Journal of Experimental Algorithmics 3, 1998.

Kaya K, Langguth J, Manne F and Ucar B: Experiments on push-relabel-based maximum cardinality matching algorithms for bipartite graphs. Technical Report TR/PA/11/33 of the Centre Europeen de Recherche et de Formation Avancee en Calcul Scientifique, 2011.

Arguments:

<i>graph</i> :	The input graph. It can be directed but the edge directions will be ignored.
<i>types</i> :	Boolean vector giving the vertex types of the graph.
<i>matching_size</i> :	The size of the matching (i.e. the number of matched vertex pairs will be returned here). It may be NULL if you don't need this.
<i>matching_weight</i> :	The weight of the matching if the edges are weighted, or the size of the matching again if the edges are unweighted. It may be NULL if you don't need this.
<i>matching</i> :	The matching itself. It must be a vector where element i contains the ID of the vertex that vertex i is matched to, or -1 if vertex i is unmatched.
<i>weights</i> :	A null pointer (=no edge weights), or a vector giving the weights of the edges. Note that the algorithm is stable only for integer weights.
<i>eps</i> :	A small real number used in equality tests in the weighted bipartite matching algorithm. Two real numbers are considered equal in the algorithm if their difference is smaller than <code>eps</code> . This is required to avoid the accumulation of numerical errors. It is advised to pass a value derived from the <code>DBL_EPSILON</code> constant in <code>float.h</code> here. If you are running the algorithm with no <code>weights</code> vector, this argument is ignored.

Returns:

Error code.

Time complexity: $O(\sqrt{|V|} |E|)$ for unweighted graphs (according to the technical report referenced above), $O(|V||E|)$ for weighted graphs.

Example	13.43.	File	examples/simple/
igraph_maximum_bipartite_matching.c			

Unfolding a graph into a tree

igraph_unfold_tree — Unfolding a graph into a tree, by possibly multiplying its vertices.

```
int igraph_unfold_tree(const igraph_t *graph, igraph_t *tree,
                      igraph_neimode_t mode, const igraph_vector_t *roots,
                      igraph_vector_t *vertex_index);
```

A graph is converted into a tree (or forest, if it is unconnected), by performing a breadth-first search on it, and replicating vertices that were found a second, third, etc. time.

Arguments:

graph: The input graph, it can be either directed or undirected.

tree: Pointer to an uninitialized graph object, the result is stored here.

mode: For directed graphs; whether to follow paths along edge directions (IGRAPH_OUT), or the opposite (IGRAPH_IN), or ignore edge directions completely (IGRAPH_ALL). It is ignored for undirected graphs.

roots: A numeric vector giving the root vertex, or vertices (if the graph is not connected), to start from.

vertex_index: Pointer to an initialized vector, or a null pointer. If not a null pointer, then a mapping from the vertices in the new graph to the ones in the original is created here.

Returns:

Error code.

Time complexity: $O(n+m)$, linear in the number vertices and edges.

Other operations

igraph_density — Calculate the density of a graph.

```
int igraph_density(const igraph_t *graph, igraph_real_t *res,
                  igraph_bool_t loops);
```

The density of a graph is simply the ratio number of edges and the number of possible edges. Note that density is ill-defined for graphs with multiple and/or loop edges, so consider calling `igraph_simplify()` on the graph if you know that it contains multiple or loop edges.

Arguments:

graph: The input graph object.

res: Pointer to a real number, the result will be stored here.

loops: Logical constant, whether to include loops in the calculation. If this constant is TRUE then loop edges are thought to be possible in the graph (this does not necessarily mean that the graph really contains any loops). If this is FALSE then the result is only correct if the graph does not contain loops.

Returns:

Error code.

Time complexity: $O(1)$.

igraph_reciprocity — Calculates the reciprocity of a directed graph.

```
int igraph_reciprocity(const igraph_t *graph, igraph_real_t *res,
                      igraph_bool_t ignore_loops,
                      igraph_reciprocity_t mode);
```

The measure of reciprocity defines the proportion of mutual connections, in a directed graph. It is most commonly defined as the probability that the opposite counterpart of a directed edge is also included in the graph. In adjacency matrix notation: $\text{sum}(i, j, (A \cdot A')_{ij}) / \text{sum}(i, j, A_{ij})$, where $A \cdot A'$ is the element-wise product of matrix A and its transpose. This measure is calculated if the *mode* argument is `IGRAPH_RECIPROCITY_DEFAULT`.

Prior to igraph version 0.6, another measure was implemented, defined as the probability of mutual connection between a vertex pair if we know that there is a (possibly non-mutual) connection between them. In other words, (unordered) vertex pairs are classified into three groups: (1) disconnected, (2) non-reciprocally connected, (3) reciprocally connected. The result is the size of group (3), divided by the sum of group sizes (2)+(3). This measure is calculated if *mode* is `IGRAPH_RECIPROCITY_RATIO`.

Arguments:

<i>graph</i> :	The graph object.
<i>res</i> :	Pointer to an <code>igraph_real_t</code> which will contain the result.
<i>ignore_loops</i> :	Whether to ignore loop edges.
<i>mode</i> :	Type of reciprocity to calculate, possible values are <code>IGRAPH_RECIPROCITY_DEFAULT</code> and <code>IGRAPH_RECIPROCITY_RATIO</code> , please see their description above.

Returns:

Error code: `IGRAPH_EINVAL`: graph has no edges `IGRAPH_ENOMEM`: not enough memory for temporary data.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges.

Example 13.44. File `examples/simple/igraph_reciprocity.c`

igraph_diversity — Structural diversity index of the vertices

```
int igraph_diversity(igraph_t *graph, const igraph_vector_t *weights,
                    igraph_vector_t *res, const igraph_vs_t vids);
```

This measure was defined in Nathan Eagle, Michael Macy and Rob Claxton: Network Diversity and Economic Development, Science 328, 1029--1031, 2010.

It is simply the (normalized) Shannon entropy of the incident edges' weights. $D(i) = H(i) / \log(k[i])$, and $H(i) = -\sum(p[i,j] \log(p[i,j])), j=1..k[i]$, where $p[i,j] = w[i,j] / \sum(w[i,l], l=1..k[i])$, $k[i]$ is the (total) degree of vertex i , and $w[i,j]$ is the weight of the edge(s) between vertex i and j . The diversity of isolated vertices will be NaN (not-a-number).

The measure works only if the graph is undirected and has no multiple edges. If the graph has multiple edges, simplify it first using `igraph_simplify()`. If the graph is directed, convert it into an undirected graph with `igraph_to_undirected()`.

Arguments:

graph: The undirected input graph.

weights: The edge weights, in the order of the edge ids, must have appropriate length.

res: An initialized vector, the results are stored here.

vids: Vertex selector that specifies the vertices which to calculate the measure.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear.

igraph_is_mutual — Check whether the edges of a directed graph are mutual.

```
int igraph_is_mutual(igraph_t *graph, igraph_vector_bool_t *res, igraph_es_t es)
```

An (A,B) edge is mutual if the graph contains the (B,A) edge, too.

An undirected graph only has mutual edges, by definition.

Edge multiplicity is not considered here, e.g. if there are two (A,B) edges and one (B,A) edge, then all three are considered to be mutual.

Loops are always mutual.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the result is stored here.

es: The sequence of edges to check. Supply `igraph_ess_all()` for all edges, see `igraph_ess_all()`.

Returns:

Error code.

Time complexity: $O(n \log(d))$, n is the number of edges supplied, d is the maximum in-degree of the vertices that are targets of the supplied edges. An upper limit of the time complexity is $O(n \log(|E|))$, $|E|$ is the number of edges in the graph.

igraph_avg_nearest_neighbor_degree — Average neighbor degree.

```
int igraph_avg_nearest_neighbor_degree(const igraph_t *graph,
                                       igraph_vs_t vids,
                                       igraph_neimode_t mode,
                                       igraph_neimode_t neighbor_degree_mode,
                                       igraph_vector_t *knn,
                                       igraph_vector_t *knnk,
                                       const igraph_vector_t *weights);
```

Calculates the average degree of the neighbors for each vertex (*knn*), and optionally, the same quantity as a function of the vertex degree (*knnk*).

For isolated vertices *knn* is set to NaN. The same is done in *knnk* for vertex degrees that don't appear in the graph.

The weighted version computes a weighted average of the neighbor degrees as $k_{nn,u} = 1/s_u \sum_v w_{uv} k_v$, where $s_u = \sum_v w_{uv}$ is the sum of the incident edge weights of vertex *u*, i.e. its strength. The sum runs over the neighbors *v* of vertex *u* as indicated by *mode*. *w_{uv}* denotes the weighted adjacency matrix and *k_v* is the neighbors' degree, specified by *neighbor_degree_mode*.

Reference: A. Barrat, M. Barthélemy, R. Pastor-Satorras, and A. Vespignani, The architecture of complex weighted networks, Proc. Natl. Acad. Sci. USA 101, 3747 (2004). <https://dx.doi.org/10.1073/pnas.0400087101>

Arguments:

<i>graph</i> :	The input graph. It may be directed.
<i>vids</i> :	The vertices for which the calculation is performed.
<i>mode</i> :	The type of neighbors to consider in directed graphs. IGRAPH_OUT considers out-neighbors, IGRAPH_IN in-neighbors and IGRAPH_ALL ignores edge directions.
<i>neighbor_degree_mode</i> :	The type of degree to average in directed graphs. IGRAPH_OUT averages out-degrees, IGRAPH_IN averages in-degrees and IGRAPH_ALL ignores edge directions for the degree calculation.
<i>vids</i> :	The vertices for which the calculation is performed.
<i>knn</i> :	Pointer to an initialized vector, the result will be stored here. It will be resized as needed. Supply a NULL pointer here, if you only want to calculate <i>knnk</i> .
<i>knnk</i> :	Pointer to an initialized vector, the average neighbor degree as a function of the vertex degree is stored here. The first (zeroth) element is for degree one vertices, etc. Supply a NULL pointer here if you don't want to calculate this.
<i>weights</i> :	Optional edge weights. Supply a null pointer here for the non-weighted version.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

Example 13.45. File `examples/simple/igraph_knn.c`

igraph_get_adjacency — Returns the adjacency matrix of a graph

```
int igraph_get_adjacency(const igraph_t *graph, igraph_matrix_t *res,
                        igraph_get_adjacency_t type, igraph_bool_t eids);
```

The result is an adjacency matrix. Entry i, j of the matrix contains the number of edges connecting vertex i to vertex j .

Arguments:

graph: Pointer to the graph to convert

res: Pointer to an initialized matrix object, it will be resized if needed.

type: Constant giving the type of the adjacency matrix to create for undirected graphs. It is ignored for directed graphs. Possible values:

IGRAPH_GET_ADJACENCY_UPPER — the upper right triangle of the matrix is used.

IGRAPH_GET_ADJACENCY_LOWER — the lower left triangle of the matrix is used.

IGRAPH_GET_ADJACENCY_BOTH — the whole matrix is used, a symmetric matrix is returned if the graph is undirected.

eids: Logical, if true, then the edges ids plus one are stored in the adjacency matrix, instead of the number of edges between the two vertices. (The plus one is needed, since edge ids start from zero, and zero means no edge in this case.)

Returns:

Error code: IGRAPH_EINVAL invalid type argument.

See also:

`igraph_get_adjacency_sparse` if you want a sparse matrix representation

Time complexity: $O(|V||V|)$, $|V|$ is the number of vertices in the graph.

igraph_get_stochastic — Stochastic adjacency matrix of a graph

```
int igraph_get_stochastic(const igraph_t *graph,
                        igraph_matrix_t *matrix,
                        igraph_bool_t column_wise);
```

Stochastic matrix of a graph. The stochastic matrix of a graph is its adjacency matrix, normalized row-wise or column-wise, such that the sum of each row (or column) is one.

Arguments:

graph: The input graph.

sparsemat: Pointer to an initialized matrix, the result is stored here.

column_wise: Whether to normalize column-wise. For undirected graphs this argument does not have any effect.

Returns:

Error code.

Time complexity: $O(|V||V|)$, quadratic in the number of vertices.

See also:

`igraph_get_stochastic_sparsemat()`, the sparse version of this function.

`igraph_get_stochastic_sparsemat` — Stochastic adjacency matrix of a graph

```
int igraph_get_stochastic_sparsemat(const igraph_t *graph,
                                   igraph_sparsemat_t *sparsemat,
                                   igraph_bool_t column_wise);
```

Stochastic matrix of a graph. The stochastic matrix of a graph is its adjacency matrix, normalized row-wise or column-wise, such that the sum of each row (or column) is one.

Arguments:

graph: The input graph.

sparsemat: Pointer to an uninitialized sparse matrix, the result is stored here.

column_wise: Whether to normalize column-wise. For undirected graphs this argument does not have any effect.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

See also:

`igraph_get_stochastic()`, the dense version of this function.

igraph_get_edgelist — Returns the list of edges in a graph

```
int igraph_get_edgelist(const igraph_t *graph, igraph_vector_t *res, igraph_bool_t bycol)
```

The order of the edges is given by the edge ids.

Arguments:

graph: Pointer to the graph object

res: Pointer to an initialized vector object, it will be resized.

bycol: Logical, if true, the edges will be returned columnwise, e.g. the first edge is `res[0] -> res[|E|]`, the second is `res[1] -> res[|E|+1]`, etc.

Returns:

Error code.

See also:

`igraph_edges()` to return the result only for some edge ids.

Time complexity: $O(|E|)$, the number of edges in the graph.

Chapter 14. Graph cycles

Eulerian cycles and paths

These functions calculate whether an Eulerian path or cycle exists and if so, can find them.

igraph_is_eulerian — Checks whether an Eulerian path or cycle exists

```
int igraph_is_eulerian(const igraph_t *graph, igraph_bool_t *has_path, igraph_bool_t *has_cycle)
```

An Eulerian path traverses each edge of the graph precisely once. A closed Eulerian path is referred to as an Eulerian cycle.

Arguments:

graph: The graph object.

has_path: Pointer to a Boolean, will be set to true if an Eulerian path exists.

has_cycle: Pointer to a Boolean, will be set to true if an Eulerian cycle exists.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

igraph_eulerian_cycle — Finds an Eulerian cycle

```
int igraph_eulerian_cycle(const igraph_t *graph, igraph_vector_t *edge_res, igraph_vector_t *vertex_res)
```

Finds an Eulerian cycle, if it exists. An Eulerian cycle is a closed path that traverses each edge precisely once.

This function uses Hierholzer's algorithm.

Arguments:

graph: The graph object.

edge_res: Pointer to an initialised vector. The indices of edges belonging to the cycle will be stored here. May be NULL if it is not needed by the caller.

vertex_res: Pointer to an initialised vector. The indices of vertices belonging to the cycle will be stored here. May be NULL if it is not needed by the caller.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID graph does not have an Eulerian cycle.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

igraph_eulerian_path — Finds an Eulerian path

```
int igraph_eulerian_path(const igraph_t *graph, igraph_vector_t *edge_res, igraph_vector_t *vertex_res)
```

Finds an Eulerian path, if it exists. An Eulerian path traverses each edge precisely once.

This function uses Hierholzer's algorithm.

Arguments:

graph: The graph object.

edge_res: Pointer to an initialised vector. The indices of edges belonging to the path will be stored here. May be NULL if it is not needed by the caller.

vertex_res: Pointer to an initialised vector. The indices of vertices belonging to the path will be stored here. May be NULL if it is not needed by the caller.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID graph does not have an Eulerian path.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

Chapter 15. Graph visitors

Breadth-first search

`igraph_bfs` — Breadth-first search

```
int igraph_bfs(const igraph_t *graph,
               igraph_integer_t root, const igraph_vector_t *roots,
               igraph_neimode_t mode, igraph_bool_t unreachable,
               const igraph_vector_t *restricted,
               igraph_vector_t *order, igraph_vector_t *rank,
               igraph_vector_t *father,
               igraph_vector_t *pred, igraph_vector_t *succ,
               igraph_vector_t *dist, igraph_bfshandler_t *callback,
               void *extra);
```

A simple breadth-first search, with a lot of different results and the possibility to call a callback whenever a vertex is visited. It is allowed to supply null pointers as the output arguments the user is not interested in, in this case they will be ignored.

If not all vertices can be reached from the supplied root vertex, then additional root vertices will be used, in the order of their vertex ids.

Consider using `igraph_bfs_simple` instead if you set most of the output arguments provided by this function to a null pointer.

Arguments:

<i>graph</i> :	The input graph.
<i>root</i> :	The id of the root vertex. It is ignored if the <i>roots</i> argument is not a null pointer.
<i>roots</i> :	Pointer to an initialized vector, or a null pointer. If not a null pointer, then it is a vector containing root vertices to start the BFS from. The vertices are considered in the order they appear. If a root vertex was already found while searching from another one, then no search is conducted from it.
<i>mode</i> :	For directed graphs, it defines which edges to follow. <code>IGRAPH_OUT</code> means following the direction of the edges, <code>IGRAPH_IN</code> means the opposite, and <code>IGRAPH_ALL</code> ignores the direction of the edges. This parameter is ignored for undirected graphs.
<i>unreachable</i> :	Logical scalar, whether the search should visit the vertices that are unreachable from the given root node(s). If true, then additional searches are performed until all vertices are visited.
<i>restricted</i> :	If not a null pointer, then it must be a pointer to a vector containing vertex ids. The BFS is carried out only on these vertices.
<i>order</i> :	If not null pointer, then the vertex ids of the graph are stored here, in the same order as they were visited.
<i>rank</i> :	If not a null pointer, then the rank of each vertex is stored here.
<i>father</i> :	If not a null pointer, then the id of the father of each vertex is stored here.

<i>pred:</i>	If not a null pointer, then the id of vertex that was visited before the current one is stored here. If there is no such vertex (the current vertex is the root of a search tree), then -1 is stored.
<i>succ:</i>	If not a null pointer, then the id of the vertex that was visited after the current one is stored here. If there is no such vertex (the current one is the last in a search tree), then -1 is stored.
<i>dist:</i>	If not a null pointer, then the distance from the root of the current search tree is stored here.
<i>callback:</i>	If not null, then it should be a pointer to a function of type <code>igraph_bfs_handler_t</code> . This function will be called, whenever a new vertex is visited.
<i>extra:</i>	Extra argument to pass to the callback function.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

Example 15.1. File `examples/simple/igraph_bfs.c`

Example 15.2. File `examples/simple/igraph_bfs_callback.c`

`igraph_bfs_simple` — Breadth-first search, single-source version

```
int igraph_bfs_simple(igraph_t *graph, igraph_integer_t vid, igraph_neimode_t mode,
                     igraph_vector_t *vids, igraph_vector_t *layers,
                     igraph_vector_t *parents);
```

An alternative breadth-first search implementation to cater for the simpler use-cases when only a single breadth-first search has to be conducted from a source node and most of the output arguments from `igraph_bfs` are not needed. It is allowed to supply null pointers as the output arguments the user is not interested in, in this case they will be ignored.

Arguments:

<i>graph:</i>	The input graph.
<i>vid:</i>	The id of the root vertex.
<i>mode:</i>	For directed graphs, it defines which edges to follow. <code>IGRAPH_OUT</code> means following the direction of the edges, <code>IGRAPH_IN</code> means the opposite, and <code>IGRAPH_ALL</code> ignores the direction of the edges. This parameter is ignored for undirected graphs.
<i>vids:</i>	If not a null pointer, then an initialized vector must be passed here. The ids of the vertices visited during the traversal will be stored here, in the same order as they were visited.
<i>layers:</i>	If not a null pointer, then an initialized vector must be passed here. The <i>i</i> -th element of the vector will contain the index into <i>vids</i> where the vertices that are at distance <i>i</i> from the root are stored. In other words, if you are interested in the vertices that are at distance <i>i</i> from the root, you need to look in the <i>vids</i> vector from <i>layers</i> [<i>i</i>] to <i>layers</i> [<i>i</i> +1].

parents: If not a null pointer, then an initialized vector must be passed here. The vector will be resized so its length is equal to the number of nodes, and it will contain the index of the parent node for each *visited* node. The values in the vector are undefined for vertices that were *not* visited.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

Example 15.3. File `examples/simple/igraph_bfs_simple.c`

igraph_bfshandler_t — Callback type for BFS function

```
typedef igraph_bool_t igraph_bfshandler_t(const igraph_t *graph,
    igraph_integer_t vid,
    igraph_integer_t pred,
    igraph_integer_t succ,
    igraph_integer_t rank,
    igraph_integer_t dist,
    void *extra);
```

`igraph_bfs()` is able to call a callback function, whenever a new vertex is found, while doing the breadth-first search. This callback function must be of type `igraph_bfshandler_t`. It has the following arguments:

Arguments:

graph: The graph that that algorithm is working on. Of course this must not be modified.

vid: The id of the vertex just found by the breadth-first search.

pred: The id of the previous vertex visited. It is -1 if there is no previous vertex, because the current vertex is the root of a search tree.

succ: The id of the next vertex that will be visited. It is -1 if there is no next vertex, because the current vertex is the last one in a search tree.

rank: The rank of the current vertex, it starts with zero.

dist: The distance (number of hops) of the current vertex from the root of the current search tree.

extra: The extra argument that was passed to `igraph_bfs()`.

Returns:

A logical value, if TRUE (=non-zero), that is interpreted as a request to stop the BFS and return to the caller. If a BFS is terminated like this, then all elements of the result vectors that were not yet calculated at the point of the termination contain NaN.

See also:

`igraph_bfs()`

Depth-first search

`igraph_dfs` — Depth-first search

```
int igraph_dfs(const igraph_t *graph, igraph_integer_t root,
               igraph_neimode_t mode, igraph_bool_t unreachable,
               igraph_vector_t *order,
               igraph_vector_t *order_out, igraph_vector_t *father,
               igraph_vector_t *dist, igraph_dfshandler_t *in_callback,
               igraph_dfshandler_t *out_callback,
               void *extra);
```

A simple depth-first search, with the possibility to call a callback whenever a vertex is discovered and/or whenever a subtree is finished. It is allowed to supply null pointers as the output arguments the user is not interested in, in this case they will be ignored.

If not all vertices can be reached from the supplied root vertex, then additional root vertices will be used, in the order of their vertex ids.

Arguments:

<i>graph</i> :	The input graph.
<i>root</i> :	The id of the root vertex.
<i>mode</i> :	For directed graphs, it defines which edges to follow. <code>IGRAPH_OUT</code> means following the direction of the edges, <code>IGRAPH_IN</code> means the opposite, and <code>IGRAPH_ALL</code> ignores the direction of the edges. This parameter is ignored for undirected graphs.
<i>unreachable</i> :	Logical scalar, whether the search should visit the vertices that are unreachable from the given root node(s). If true, then additional searches are performed until all vertices are visited.
<i>order</i> :	If not null pointer, then the vertex ids of the graph are stored here, in the same order as they were discovered.
<i>order_out</i> :	If not a null pointer, then the vertex ids of the graphs are stored here, in the order of the completion of their subtree.
<i>father</i> :	If not a null pointer, then the id of the father of each vertex is stored here.
<i>dist</i> :	If not a null pointer, then the distance from the root of the current search tree is stored here.
<i>in_callback</i> :	If not null, then it should be a pointer to a function of type <code>igraph_dfshandler_t</code> . This function will be called, whenever a new vertex is discovered.
<i>out_callback</i> :	If not null, then it should be a pointer to a function of type <code>igraph_dfshandler_t</code> . This function will be called, whenever the subtree of a vertex is completed.
<i>extra</i> :	Extra argument to pass to the callback function(s).

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

igraph_dfshandler_t — Callback type for the DFS function

```
typedef igraph_bool_t igraph_dfshandler_t(const igraph_t *graph,
    igraph_integer_t vid,
    igraph_integer_t dist,
    void *extra);
```

`igraph_dfs()` is able to call a callback function, whenever a new vertex is discovered, and/or whenever a subtree is completed. These callbacks must be of type `igraph_dfshandler_t`. They have the following arguments:

Arguments:

graph: The graph that that algorithm is working on. Of course this must not be modified.

vid: The id of the vertex just found by the depth-first search.

dist: The distance (number of hops) of the current vertex from the root of the current search tree.

extra: The extra argument that was passed to `igraph_dfs()`.

Returns:

A logical value, if TRUE (=non-zero), that is interpreted as a request to stop the DFS and return to the caller. If a DFS is terminated like this, then all elements of the result vectors that were not yet calculated at the point of the termination contain NaN.

See also:

`igraph_dfs()`

Random walks

igraph_random_walk — Perform a random walk on a graph

```
int igraph_random_walk(const igraph_t *graph, igraph_vector_t *walk,
    igraph_integer_t start, igraph_neimode_t mode,
    igraph_integer_t steps,
    igraph_random_walk_stuck_t stuck);
```

Performs a random walk with a given length on a graph, from the given start vertex. Edge directions are (potentially) considered, depending on the *mode* argument.

Arguments:

- graph*: The input graph, it can be directed or undirected. Multiple edges are respected, so are loop edges.
- walk*: An allocated vector, the result is stored here. It will be resized as needed.
- start*: The start vertex for the walk.
- steps*: The number of steps to take. If the random walk gets stuck, then the *stuck* argument specifies what happens.
- mode*: How to walk along the edges in directed graphs. `IGRAPH_OUT` means following edge directions, `IGRAPH_IN` means going opposite the edge directions, `IGRAPH_ALL` means ignoring edge directions. This argument is ignored for undirected graphs.
- stuck*: What to do if the random walk gets stuck. `IGRAPH_RANDOM_WALK_STUCK_RETURN` means that the function returns with a shorter walk; `IGRAPH_RANDOM_WALK_STUCK_ERROR` means that an error is reported. In both cases *walk* is truncated to contain the actual interrupted walk.

Returns:

Error code.

Time complexity: $O(l + d)$, where l is the length of the walk, and d is the total degree of the visited nodes.

`igraph_random_edge_walk` — Perform a random walk on a graph and return the traversed edges

```
int igraph_random_edge_walk(const igraph_t *graph,
                           const igraph_vector_t *weights,
                           igraph_vector_t *edgewalk,
                           igraph_integer_t start, igraph_neimode_t mode,
                           igraph_integer_t steps,
                           igraph_random_walk_stuck_t stuck);
```

Performs a random walk with a given length on a graph, from the given start vertex. Edge directions are (potentially) considered, depending on the *mode* argument.

Arguments:

- graph*: The input graph, it can be directed or undirected. Multiple edges are respected, so are loop edges.
- weights*: A vector of non-negative edge weights. It is assumed that at least one strictly positive weight is found among the outgoing edges of each vertex. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If it is a NULL pointer, all edges are considered to have equal weight.
- edgewalk*: An initialized vector; the indices of traversed edges are stored here. It will be resized as needed.
- start*: The start vertex for the walk.
- steps*: The number of steps to take. If the random walk gets stuck, then the *stuck* argument specifies what happens.

mode: How to walk along the edges in directed graphs. `IGRAPH_OUT` means following edge directions, `IGRAPH_IN` means going opposite the edge directions, `IGRAPH_ALL` means ignoring edge directions. This argument is ignored for undirected graphs.

stuck: What to do if the random walk gets stuck. `IGRAPH_RANDOM_WALK_STUCK_RETURN` means that the function returns with a shorter walk; `IGRAPH_RANDOM_WALK_STUCK_ERROR` means that an error is reported. In both cases, *edge-walk* is truncated to contain the actual interrupted walk.

Returns:

Error code.

Chapter 16. Cliques and independent vertex sets

These functions calculate various graph properties related to cliques and independent vertex sets.

Cliques

igraph_cliques — Finds all or some cliques in a graph.

```
int igraph_cliques(const igraph_t *graph, igraph_vector_ptr_t *res,
                  igraph_integer_t min_size, igraph_integer_t max_size);
```

Cliques are fully connected subgraphs of a graph.

If you are only interested in the size of the largest clique in the graph, use `igraph_clique_number()` instead.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, <http://users.aalto.fi/~pat/cliquer.html>

Arguments:

graph: The input graph.

res: Pointer to a pointer vector, the result will be stored here, i.e. *res* will contain pointers to `igraph_vector_t` objects which contain the indices of vertices involved in a clique. The pointer vector will be resized if needed but note that the objects in the pointer vector will not be freed.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_largest_cliques()` and `igraph_clique_number()`.

Time complexity: Exponential

Example 16.1. File `examples/simple/igraph_cliques.c`

igraph_clique_size_hist — Counts cliques of each size in the graph.

```
int igraph_clique_size_hist(const igraph_t *graph, igraph_vector_t *hist,
                           igraph_integer_t min_size, igraph_integer_t max_size)
```

Cliques are fully connected subgraphs of a graph.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, <http://users.aalto.fi/~pat/cliquer.html>

Arguments:

- graph*: The input graph.
- hist*: Pointer to an initialized vector. The result will be stored here. The first element will store the number of size-1 cliques, the second element the number of size-2 cliques, etc. For cliques smaller than *min_size*, zero counts will be returned.
- min_size*: Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.
- max_size*: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_cliques()` and `igraph_cliques_callback()`

Time complexity: Exponential

igraph_cliques_callback — Calls a function for each clique in the graph.

```
int igraph_cliques_callback(const igraph_t *graph,
                            igraph_integer_t min_size, igraph_integer_t max_size,
                            igraph_clique_handler_t *cliquehandler_fn, void *arg)
```

Cliques are fully connected subgraphs of a graph. This function enumerates all cliques within the given size range and calls *cliquehandler_fn* for each of them. The cliques are passed to the callback function as a pointer to an `igraph_vector_t`. Destroying and freeing this vector is left up to the user. Use `igraph_vector_destroy()` to destroy it first, then free it using `igraph_free()`.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, <http://users.aalto.fi/~pat/cliquer.html>

Arguments:

- graph*: The input graph.
- min_size*: Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

cliquehandler_fn: Callback function to be called for each clique. See also `igraph_clique_handler_t`.

arg: Extra argument to supply to *cliquehandler_fn*.

Returns:

Error code.

See also:

`igraph_cliques()`

Time complexity: Exponential

`igraph_clique_handler_t` — Type of clique handler functions.

```
typedef igraph_bool_t igraph_clique_handler_t(igraph_vector_t *clique, void *arg);
```

Callback type, called when a clique was found. See the details at the documentation of `igraph_cliques_callback()`.

Arguments:

clique: The current clique. Destroying and freeing this vector is left to the user. Use `igraph_vector_destroy()` and `igraph_free()` to do this.

arg: This extra argument was passed to `igraph_cliques_callback()` when it was called.

Returns:

Boolean, whether to continue with the clique search.

`igraph_largest_cliques` — Finds the largest clique(s) in a graph.

```
int igraph_largest_cliques(const igraph_t *graph, igraph_vector_ptr_t *res);
```

A clique is largest (quite intuitively) if there is no other clique in the graph which contains more vertices.

Note that this is not necessarily the same as a maximal clique, i.e. the largest cliques are always maximal but a maximal clique is not always largest.

The current implementation of this function searches for maximal cliques using `igraph_maximal_cliques()` and drops those that are not the largest.

The implementation of this function changed between igraph 0.5 and 0.6, so the order of the cliques and the order of vertices within the cliques will almost surely be different between these two versions.

Arguments:

graph: The input graph.

res: Pointer to an initialized pointer vector, the result will be stored here. It will be resized as needed. Note that vertices of a clique may be returned in arbitrary order.

Returns:

Error code.

See also:

`igraph_cliques()`, `igraph_maximal_cliques()`

Time complexity: $O(3^{(|V|/3)})$ worst case.

igraph_maximal_cliques — Finds all maximal cliques in a graph.

```
int igraph_maximal_cliques(const igraph_t *graph,
                           igraph_vector_ptr_t *res,
                           igraph_integer_t min_size,
                           igraph_integer_t max_size);
```

A maximal clique is a clique which can't be extended any more by adding a new vertex to it.

If you are only interested in the size of the largest clique in the graph, use `igraph_clique_number()` instead.

The current implementation uses a modified Bron-Kerbosch algorithm to find the maximal cliques, see: David Eppstein, Maarten Löffler, Darren Strash: Listing All Maximal Cliques in Sparse Graphs in Near-Optimal Time. Algorithms and Computation, Lecture Notes in Computer Science Volume 6506, 2010, pp 403-414.

The implementation of this function changed between igraph 0.5 and 0.6 and also between 0.6 and 0.7, so the order of the cliques and the order of vertices within the cliques will almost surely be different between these three versions.

Arguments:

graph: The input graph.

res: Pointer to a pointer vector, the result will be stored here, i.e. *res* will contain pointers to `igraph_vector_t` objects which contain the indices of vertices involved in a clique. The pointer vector will be resized if needed but note that the objects in the pointer vector will not be freed. Note that vertices of a clique may be returned in arbitrary order.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_maximal_independent_vertex_sets()`, `igraph_clique_number()`

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

Example 16.2. File `examples/simple/igraph_maximal_cliques.c`

`igraph_maximal_cliques_count` — Count the number of maximal cliques in a graph

```
int igraph_maximal_cliques_count(const igraph_t *graph,
                                igraph_integer_t *res,
                                igraph_integer_t min_size,
                                igraph_integer_t max_size);
```

The current implementation uses a modified Bron-Kerbosch algorithm to find the maximal cliques, see: David Eppstein, Maarten Löffler, Darren Strash: Listing All Maximal Cliques in Sparse Graphs in Near-Optimal Time. Algorithms and Computation, Lecture Notes in Computer Science Volume 6506, 2010, pp 403-414.

Arguments:

graph: The input graph.

res: Pointer to an `igraph_integer_t`; the number of maximal cliques will be stored here.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_maximal_cliques()`.

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

Example 16.3. File `examples/simple/igraph_maximal_cliques.c`

igraph_maximal_cliques_file — Find maximal cliques and write them to a file.

```
int igraph_maximal_cliques_file(const igraph_t *graph,
                                FILE *outfile,
                                igraph_integer_t min_size,
                                igraph_integer_t max_size);
```

This function enumerates all maximal cliques and writes them to file.

Edge directions are ignored.

Arguments:

- graph*: The input graph.
- outfile*: Pointer to the output file, it should be writable.
- min_size*: Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.
- max_size*: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_maximal_cliques()`.

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.*

igraph_maximal_cliques_subset — Maximal cliques for a subset of initial vertices

```
int igraph_maximal_cliques_subset(const igraph_t *graph,
                                   igraph_vector_int_t *subset,
                                   igraph_vector_ptr_t *res,
                                   igraph_integer_t *no,
                                   FILE *outfile,
                                   igraph_integer_t min_size,
                                   igraph_integer_t max_size);
```

This function enumerates all maximal cliques for a subset of initial vertices and writes them to file.

Edge directions are ignored.

Arguments:

graph: The input graph.

subset: Pointer to an `igraph_vector_int_t` containing the subset of initial vertices

res: Pointer to an `igraph_ptr_t`; the cliques will be stored here

no: Pointer to an `igraph_integer_t`; the number of maximal cliques will be stored here.

outfile: Pointer to the output file, it should be writable.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_maximal_cliques()`.

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

`igraph_maximal_cliques_hist` — Counts the number of maximal cliques of each size in a graph.

```
int igraph_maximal_cliques_hist(const igraph_t *graph,
                                igraph_vector_t *hist,
                                igraph_integer_t min_size,
                                igraph_integer_t max_size);
```

This function counts how many maximal cliques of each size are present in the graph. Size-1 maximal cliques are simply isolated vertices.

Edge directions are ignored.

Arguments:

graph: The input graph.

hist: Pointer to an initialized vector. The result will be stored here. The first element will store the number of size-1 maximal cliques, the second element the number of size-2 maximal cliques, etc. For cliques smaller than *min_size*, zero counts will be returned.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_maximal_cliques()`.

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

`igraph_maximal_cliques_callback` — Finds maximal cliques in a graph and calls a function for each one.

```
int igraph_maximal_cliques_callback(const igraph_t *graph,
                                   igraph_clique_handler_t *cliquehandler_fn,
                                   igraph_integer_t min_size, igraph_integer_t max_size);
```

This function enumerates all maximal cliques within the given size range and calls *cliquehandler_fn* for each of them. The cliques are passed to the callback function as a pointer to an `igraph_vector_t`. Destroying and freeing this vector is left up to the user. Use `igraph_vector_destroy()` to destroy it first, then free it using `igraph_free()`.

Edge directions are ignored.

Arguments:

<i>graph</i> :	The input graph.
<i>cliquehandler_fn</i> :	Callback function to be called for each clique. See also <code>igraph_clique_handler_t</code> .
<i>arg</i> :	Extra argument to supply to <i>cliquehandler_fn</i> .
<i>min_size</i> :	Integer giving the minimum size of the cliques to be returned. If negative or zero, no lower bound will be used.
<i>max_size</i> :	Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_maximal_cliques()`.

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

igraph_clique_number — Finds the clique number of the graph.

```
int igraph_clique_number(const igraph_t *graph, igraph_integer_t *no);
```

The clique number of a graph is the size of the largest clique.

Arguments:

graph: The input graph.

no: The clique number will be returned to the `igraph_integer_t` pointed by this variable.

Returns:

Error code.

See also:

`igraph_cliques()`, `igraph_largest_cliques()`.

Time complexity: $O(3^{(|V|/3)})$ worst case.

Weighted cliques

igraph_weighted_cliques — Finds all cliques in a given weight range in a vertex weighted graph.

```
int igraph_weighted_cliques(const igraph_t *graph,
                           const igraph_vector_t *vertex_weights, igraph_vector_t *res,
                           igraph_real_t min_weight, igraph_real_t max_weight,
```

Cliques are fully connected subgraphs of a graph. The weight of a clique is the sum of the weights of individual vertices within the clique.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, <http://users.aalto.fi/~pat/cliquer.html> Only positive integer vertex weights are supported.

Arguments:

graph: The input graph.

vertex_weights: A vector of vertex weights. The current implementation will truncate all weights to their integer parts.

res: Pointer to a pointer vector, the result will be stored here, i.e. *res* will contain pointers to `igraph_vector_t` objects which contain the indices of vertices involved in a clique. The pointer vector will be resized if needed but note that the objects in the pointer vector will not be freed.

min_weight: Integer giving the minimum weight of the cliques to be returned. If negative or zero, no lower bound will be used.

max_weight: Integer giving the maximum weight of the cliques to be returned. If negative or zero, no upper bound will be used.

maximal: If true, only maximal cliques will be returned

Returns:

Error code.

See also:

`igraph_cliques()`, `igraph_maximal_cliques()`

Time complexity: Exponential

`igraph_largest_weighted_cliques` — Finds the largest weight clique(s) in a graph.

```
int igraph_largest_weighted_cliques(const igraph_t *graph,
                                   const igraph_vector_t *vertex_weights, igraph_vector_t *res)
```

Finds the clique(s) having the largest weight in the graph.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, <http://users.aalto.fi/~pat/cliquer.html> Only positive integer vertex weights are supported.

Arguments:

graph: The input graph.

vertex_weights: A vector of vertex weights. The current implementation will truncate all weights to their integer parts.

res: Pointer to a pointer vector, the result will be stored here, i.e. *res* will contain pointers to `igraph_vector_t` objects which contain the indices of vertices involved in a clique. The pointer vector will be resized if needed but note that the objects in the pointer vector will not be freed.

Returns:

Error code.

See also:

`igraph_weighted_cliques()`, `igraph_weighted_clique_number()`,
`igraph_largest_cliques()`

Time complexity: TODO

igraph_weighted_clique_number — Finds the weight of the largest weight clique in the graph.

```
int igraph_weighted_clique_number(const igraph_t *graph,
                                  const igraph_vector_t *vertex_weights, igraph_real_t *res);
```

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, <http://users.aalto.fi/~pat/cliquer.html> Only positive integer vertex weights are supported.

Arguments:

graph: The input graph.

vertex_weights: A vector of vertex weights. The current implementation will truncate all weights to their integer parts.

res: The largest weight will be returned to the `igraph_real_t` pointed to by this variable.

Returns:

Error code.

See also:

`igraph_weighted_cliques()`, `igraph_largest_weighted_cliques()`,
`igraph_clique_number()`

Time complexity: TODO

Independent vertex sets

igraph_independent_vertex_sets — Finds all independent vertex sets in a graph.

```
int igraph_independent_vertex_sets(const igraph_t *graph,
                                   igraph_vector_ptr_t *res,
                                   igraph_integer_t min_size,
                                   igraph_integer_t max_size);
```

A vertex set is considered independent if there are no edges between them.

If you are interested in the size of the largest independent vertex set, use `igraph_independence_number()` instead.

The current implementation was ported to igraph from the Very Nauty Graph Library by Keith Briggs and uses the algorithm from the paper S. Tsukiyama, M. Ide, H. Ariyoshi and I. Shirawaka. A new algorithm for generating all the maximal independent sets. SIAM J Computing, 6:505--517, 1977.

Arguments:

graph: The input graph.

res: Pointer to a pointer vector, the result will be stored here, i.e. *res* will contain pointers to `igraph_vector_t` objects which contain the indices of vertices involved in an independent vertex set. The pointer vector will be resized if needed but note that the objects in the pointer vector will not be freed.

min_size: Integer giving the minimum size of the sets to be returned. If negative or zero, no lower bound will be used.

max_size: Integer giving the maximum size of the sets to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

`igraph_largest_independent_vertex_sets()`, `igraph_independence_number()`.

Time complexity: TODO

Example 16.4. File `examples/simple/igraph_independent_sets.c`

`igraph_largest_independent_vertex_sets` — Finds the largest independent vertex set(s) in a graph.

```
int igraph_largest_independent_vertex_sets(const igraph_t *graph,
                                           igraph_vector_ptr_t *res);
```

An independent vertex set is largest if there is no other independent vertex set with more vertices in the graph.

The current implementation was ported to igraph from the Very Nauty Graph Library by Keith Briggs and uses the algorithm from the paper S. Tsukiyama, M. Ide, H. Ariyoshi and I. Shirawaka. A new algorithm for generating all the maximal independent sets. *SIAM J Computing*, 6:505--517, 1977.

Arguments:

graph: The input graph.

res: Pointer to a pointer vector, the result will be stored here. It will be resized as needed.

Returns:

Error code.

See also:

`igraph_independent_vertex_sets()`, `igraph_maximal_independent_vertex_sets()`.

Time complexity: TODO

igraph_maximal_independent_vertex_sets — Finds all maximal independent vertex sets of a graph.

```
int igraph_maximal_independent_vertex_sets(const igraph_t *graph,
                                           igraph_vector_ptr_t *res);
```

A maximal independent vertex set is an independent vertex set which can't be extended any more by adding a new vertex to it.

The algorithm used here is based on the following paper: S. Tsukiyama, M. Ide, H. Ariyoshi and I. Shirawaka. A new algorithm for generating all the maximal independent sets. SIAM J Computing, 6:505--517, 1977.

The implementation was originally written by Kevin O'Neill and modified by K M Briggs in the Very Nauty Graph Library. I simply re-wrote it to use igraph's data structures.

If you are interested in the size of the largest independent vertex set, use `igraph_independence_number()` instead.

Arguments:

graph: The input graph.

res: Pointer to a pointer vector, the result will be stored here, i.e. *res* will contain pointers to `igraph_vector_t` objects which contain the indices of vertices involved in an independent vertex set. The pointer vector will be resized if needed but note that the objects in the pointer vector will not be freed.

Returns:

Error code.

See also:

`igraph_maximal_cliques()`, `igraph_independence_number()`

Time complexity: TODO.

igraph_independence_number — Finds the independence number of the graph.

```
int igraph_independence_number(const igraph_t *graph, igraph_integer_t *no);
```

The independence number of a graph is the cardinality of the largest independent vertex set.

The current implementation was ported to igraph from the Very Nauty Graph Library by Keith Briggs and uses the algorithm from the paper S. Tsukiyama, M. Ide, H. Ariyoshi and I. Shirawaka. A new algorithm for generating all the maximal independent sets. SIAM J Computing, 6:505--517, 1977.

Arguments:

graph: The input graph.

no: The independence number will be returned to the `igraph_integer_t` pointed by this variable.

Returns:

Error code.

See also:

`igraph_independent_vertex_sets()`.

Time complexity: TODO.

Chapter 17. Graph isomorphism

The simple interface

igraph provides four set of functions to deal with graph isomorphism problems.

The `igraph_isomorphic()` and `igraph_subisomorphic()` functions make up the first set (in addition with the `igraph_permute_vertices()` function). These functions choose the algorithm which is best for the supplied input graph. (The choice is not very sophisticated though, see their documentation for details.)

The VF2 graph (and subgraph) isomorphism algorithm is implemented in igraph, these functions are the second set. See `igraph_isomorphic_vf2()` and `igraph_subisomorphic_vf2()` for starters.

Functions for the Bliss algorithm constitute the third set, see `igraph_isomorphic_bliss()`.

Finally, the isomorphism classes of all graphs with three and four vertices are precomputed and stored in igraph, so for these small graphs there is a very simple fast way to decide isomorphism. See `igraph_isomorphic_34()`.

igraph_isomorphic — Decides whether two graphs are isomorphic

```
int igraph_isomorphic(const igraph_t *graph1, const igraph_t *graph2,
                     igraph_bool_t *iso);
```

In simple terms, two graphs are isomorphic if they become indistinguishable from each other once their vertex labels are removed (rendering the vertices within each graph indistinguishable). More precisely, two graphs are isomorphic if there is a one-to-one mapping from the vertices of the first one to the vertices of the second such that it transforms the edge set of the first graph into the edge set of the second. This mapping is called an *isomorphism*.

Currently, this function supports simple graphs and graphs with self-loops, but does not support multi-graphs.

This function decides which graph isomorphism algorithm to be used based on the input graphs. Right now it does the following:

1. If one graph is directed and the other undirected then an error is triggered.
2. If one of the graphs has multi-edges then an error is triggered.
3. If the two graphs does not have the same number of vertices and edges it returns with `FALSE`.
4. Otherwise, if the graphs have three or four vertices then an $O(1)$ algorithm is used with precomputed data.
5. Otherwise Bliss is used, see `igraph_isomorphic_bliss()`.

Please call the VF2 and Bliss functions directly if you need something more sophisticated, e.g. you need the isomorphic mapping.

Arguments:

graph1: The first graph.

graph2: The second graph.

iso: Pointer to a logical variable, will be set to TRUE (1) if the two graphs are isomorphic, and FALSE (0) otherwise.

Returns:

Error code.

See also:

`igraph_isoclass()`, `igraph_isoclass_subgraph()`, `igraph_isoclass_create()`.

Time complexity: exponential.

`igraph_subisomorphic` — Decide subgraph isomorphism.

```
int igraph_subisomorphic(const igraph_t *graph1, const igraph_t *graph2,
                        igraph_bool_t *iso);
```

Check whether *graph2* is isomorphic to a subgraph of *graph1*. Currently this function just calls `igraph_subisomorphic_vf2()` for all graphs.

Currently this function does not support non-simple graphs.

Arguments:

graph1: The first input graph, may be directed or undirected. This is supposed to be the bigger graph.

graph2: The second input graph, it must have the same directedness as *graph1*, or an error is triggered. This is supposed to be the smaller graph.

iso: Pointer to a boolean, the result is stored here.

Returns:

Error code.

Time complexity: exponential.

The BLISS algorithm

Bliss is a successor of the famous NAUTY algorithm and implementation. While using the same ideas in general, with better heuristics and data structures Bliss outperforms NAUTY on most graphs.

Bliss was developed and implemented by Tommi Junttila and Petteri Kaski at Helsinki University of Technology, Finland. For more information, see the Bliss homepage at <https://users.aalto.fi/~tjunttila/bliss/> and the following publication:

Tommi Junttila and Petteri Kaski: "Engineering an Efficient Canonical Labeling Tool for Large and Sparse Graphs" In ALENEX 2007, pages 135–149, 2007 <https://doi.org/10.1137/1.9781611972870.13>

Tommi Junttila and Petteri Kaski: "Conflict Propagation and Component Recursion for Canonical Labeling" in TAPAS 2011, pages 151–162, 2011. https://doi.org/10.1007/978-3-642-19754-3_16

Bliss works with both directed graphs and undirected graphs. It supports graphs with self-loops, but not graphs with multi-edges.

Bliss version 0.75 is included in igraph.

igraph_bliss_sh_t — Splitting heuristics for Bliss.

```
typedef enum { IGRAPH_BLISS_F = 0, IGRAPH_BLISS_FL,  
               IGRAPH_BLISS_FS, IGRAPH_BLISS_FM,  
               IGRAPH_BLISS_FLM, IGRAPH_BLISS_FSM  
             } igraph_bliss_sh_t;
```

IGRAPH_BLISS_FL provides good performance for many graphs, and is a reasonable default choice. IGRAPH_BLISS_FSM is recommended for graphs that have some combinatorial structure, and is the default of the Bliss library's command line tool.

Values:

IGRAPH_BLISS_F:	First non-singleton cell.
IGRAPH_BLISS_FL:	First largest non-singleton cell.
IGRAPH_BLISS_FS:	First smallest non-singleton cell.
IGRAPH_BLISS_FM:	First maximally non-trivially connected non-singleton cell.
IGRAPH_BLISS_FLM:	Largest maximally non-trivially connected non-singleton cell.
IGRAPH_BLISS_FSM:	Smallest maximally non-trivially connected non-singleton cell.

igraph_bliss_info_t — Information about a BLISS run

```
typedef struct igraph_bliss_info_t {  
    unsigned long nof_nodes;  
    unsigned long nof_leaf_nodes;  
    unsigned long nof_bad_nodes;  
    unsigned long nof_canupdates;  
    unsigned long nof_generators;  
    unsigned long max_level;  
    char *group_size;  
} igraph_bliss_info_t;
```

Some secondary information found by the BLISS algorithm is stored here. It is useful if you want to study the internal working of the algorithm.

Values:

<code>nof_nodes:</code>	The number of nodes in the search tree.
<code>nof_leaf_nodes:</code>	The number of leaf nodes in the search tree.
<code>nof_bad_nodes:</code>	Number of bad nodes.
<code>nof_canupdates:</code>	Number of canrep updates.
<code>nof_generators:</code>	Number of generators of the automorphism group.
<code>max_level:</code>	Maximum level.
<code>group_size:</code>	The size of the automorphism group of the graph, given as a string. It should be deallocated via <code>igraph_free()</code> if not needed any more.

See <http://www.tcs.hut.fi/Software/bliss/index.html> for details about the algorithm and these parameters.

`igraph_canonical_permutation` — Canonical permutation using Bliss

```
int igraph_canonical_permutation(const igraph_t *graph, const igraph_vector_int_t *
                                igraph_vector_t *labeling, igraph_bliss_sh_t sh)
```

This function computes the canonical permutation which transforms the graph into a canonical form by using the Bliss algorithm.

Arguments:

<i>graph</i> :	The input graph. Multiple edges between the same nodes are not supported and will cause an incorrect result to be returned.
<i>colors</i> :	An optional vertex color vector for the graph. Supply a null pointer if the graph is not colored.
<i>labeling</i> :	Pointer to a vector, the result is stored here. The permutation takes vertex 0 to the first element of the vector, vertex 1 to the second, etc. The vector will be resized as needed.
<i>sh</i> :	The splitting heuristics to be used in Bliss. See <code>igraph_bliss_sh_t</code> .
<i>info</i> :	If not NULL then information on Bliss internals is stored here. See <code>igraph_bliss_info_t</code> .

Returns:

Error code.

Time complexity: exponential, in practice it is fast for many graphs.

`igraph_isomorphic_bliss` — Graph isomorphism via Bliss

```
int igraph_isomorphic_bliss(const igraph_t *graph1, const igraph_t *graph2,
                           const igraph_vector_int_t *colors1, const igraph_vector_int_t *colors2,
                           igraph_bool_t *iso, igraph_vector_t *map12,
                           igraph_vector_t *map21, igraph_bliss_sh_t sh,
                           igraph_bliss_info_t *info1, igraph_bliss_info_t *info2);
```

This function uses the Bliss graph isomorphism algorithm, a successor of the famous NAUTY algorithm and implementation. Bliss is open source and licensed according to the GNU LGPL. See <https://users.aalto.fi/~tjunttil/bliss/> for details. Currently the 0.75 version of Bliss is included in igraph.

Arguments:

- graph1*: The first input graph. Multiple edges between the same nodes are not supported and will cause an incorrect result to be returned.
- graph2*: The second input graph. Multiple edges between the same nodes are not supported and will cause an incorrect result to be returned.
- colors1*: An optional vertex color vector for the first graph. Supply a null pointer if your graph is not colored.
- colors2*: An optional vertex color vector for the second graph. Supply a null pointer if your graph is not colored.
- iso*: Pointer to a boolean, the result is stored here.
- map12*: A vector or NULL pointer. If not NULL then an isomorphic mapping from *graph1* to *graph2* is stored here. If the input graphs are not isomorphic then this vector is cleared, i.e. it will have length zero.
- map21*: Similar to *map12*, but for the mapping from *graph2* to *graph1*.
- sh*: Splitting heuristics to be used for the graphs. See `igraph_bliss_sh_t`.
- info1*: If not NULL, information about the canonization of the first input graph is stored here. See `igraph_bliss_info_t` for details. Note that if the two graphs have different number of vertices or edges, then this is not filled.
- info2*: Same as *info1*, but for the second graph.

Returns:

Error code.

Time complexity: exponential, but in practice it is quite fast.

igraph_automorphisms — Number of automorphisms using Bliss

```
int igraph_automorphisms(const igraph_t *graph, const igraph_vector_int_t *colors,
                        igraph_bliss_sh_t sh, igraph_bliss_info_t *info);
```

The number of automorphisms of a graph is computed using Bliss. The result is returned as part of the *info* structure, in tag `group_size`. It is returned as a string, as it can be very high even for relatively small graphs. If the GNU MP library is used then this number is exact, otherwise a long double is used and it is only approximate. See also `igraph_bliss_info_t`.

Arguments:

- graph*: The input graph. Multiple edges between the same nodes are not supported and will cause an incorrect result to be returned.
- colors*: An optional vertex color vector for the graph. Supply a null pointer if the graph is not colored.
- sh*: The splitting heuristics to be used in Bliss. See `igraph_bliss_sh_t`.
- info*: The result is stored here, in particular in the `group_size` tag of *info*.

Returns:

Error code.

Time complexity: exponential, in practice it is fast for many graphs.

igraph_automorphism_group — Automorphism group generators using Bliss

```
int igraph_automorphism_group(  
    const igraph_t *graph, const igraph_vector_int_t *colors, igraph_vector_ptr_t  
    igraph_bliss_sh_t sh, igraph_bliss_info_t *info);
```

The generators of the automorphism group of a graph are computed using Bliss. The generator set may not be minimal and may depend on the splitting heuristics.

Arguments:

- graph*: The input graph. Multiple edges between the same nodes are not supported and will cause an incorrect result to be returned.
- colors*: An optional vertex color vector for the graph. Supply a null pointer if the graph is not colored.
- generators*: Must be an initialized pointer vector. It will contain pointers to `igraph_vector_t` objects representing generators of the automorphism group.
- sh*: The splitting heuristics to be used in Bliss. See `igraph_bliss_sh_t`.
- info*: If not `NULL` then information on Bliss internals is stored here. See `igraph_bliss_info_t`.

Returns:

Error code.

Time complexity: exponential, in practice it is fast for many graphs.

The VF2 algorithm

The VF2 algorithm can search for a subgraph in a larger graph, or check if two graphs are isomorphic. See P. Foggia, C. Sansone, M. Vento, An Improved algorithm for matching large graphs, Proc. of the 3rd IAPR-TC-15 International Workshop on Graph-based Representations, Italy, 2001.

VF2 supports both vertex and edge-colored graphs, as well as custom vertex or edge compatibility functions.

VF2 works with both directed and undirected graphs. Only simple graphs are supported. Self-loops or multi-edges must not be present in the graphs. Currently, the VF2 functions do not check that the input graph is simple: it is the responsibility of the user to pass in valid input.

igraph_isomorphic_vf2 — Isomorphism via VF2

```
int igraph_isomorphic_vf2(const igraph_t *graph1, const igraph_t *graph2,
                          const igraph_vector_int_t *vertex_color1,
                          const igraph_vector_int_t *vertex_color2,
                          const igraph_vector_int_t *edge_color1,
                          const igraph_vector_int_t *edge_color2,
                          igraph_bool_t *iso, igraph_vector_t *map12,
                          igraph_vector_t *map21,
                          igraph_isocompat_t *node_compat_fn,
                          igraph_isocompat_t *edge_compat_fn,
                          void *arg);
```

This function performs the VF2 algorithm via calling `igraph_isomorphic_function_vf2()`.

Note that this function cannot be used for deciding subgraph isomorphism, use `igraph_subisomorphic_vf2()` for that.

Arguments:

<i>graph1</i> :	The first graph, may be directed or undirected.
<i>graph2</i> :	The second graph. It must have the same directedness as <i>graph1</i> , otherwise an error is reported.
<i>vertex_color1</i> :	An optional color vector for the first graph. If color vectors are given for both graphs, then the isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here if your graphs are not colored.
<i>vertex_color2</i> :	An optional color vector for the second graph. See the previous argument for explanation.
<i>edge_color1</i> :	An optional edge color vector for the first graph. The matching edges in the two graphs must have matching colors as well. Supply a null pointer here if your graphs are not edge-colored.
<i>edge_color2</i> :	The edge color vector for the second graph.
<i>iso</i> :	Pointer to a logical constant, the result of the algorithm will be placed here.
<i>map12</i> :	Pointer to an initialized vector or a NULL pointer. If not a NULL pointer then the mapping from <i>graph1</i> to <i>graph2</i> is stored here. If the graphs are not isomorphic then the vector is cleared (i.e. has zero elements).
<i>map21</i> :	Pointer to an initialized vector or a NULL pointer. If not a NULL pointer then the mapping from <i>graph2</i> to <i>graph1</i> is stored here. If the graphs are not isomorphic then the vector is cleared (i.e. has zero elements).

node_compat_fn: A pointer to a function of type `igraph_isocompat_t`. This function will be called by the algorithm to determine whether two nodes are compatible.

edge_compat_fn: A pointer to a function of type `igraph_isocompat_t`. This function will be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions *node_compat_fn* and *edge_compat_fn*.

Returns:

Error code.

See also:

`igraph_subisomorphic_vf2()`, `igraph_count_isomorphisms_vf2()`,
`igraph_get_isomorphisms_vf2()`,

Time complexity: exponential, what did you expect?

Example 17.1. File `examples/simple/igraph_isomorphic_vf2.c`

`igraph_count_isomorphisms_vf2` — Number of isomorphisms via VF2

```
int igraph_count_isomorphisms_vf2(const igraph_t *graph1, const igraph_t *graph2,
                                  const igraph_vector_int_t *vertex_color1,
                                  const igraph_vector_int_t *vertex_color2,
                                  const igraph_vector_int_t *edge_color1,
                                  const igraph_vector_int_t *edge_color2,
                                  igraph_integer_t *count,
                                  igraph_isocompat_t *node_compat_fn,
                                  igraph_isocompat_t *edge_compat_fn,
                                  void *arg);
```

This function counts the number of isomorphic mappings between two graphs. It uses the generic `igraph_isomorphic_function_vf2()` function.

Arguments:

graph1: The first input graph, may be directed or undirected.

graph2: The second input graph, it must have the same directedness as *graph1*, or an error will be reported.

vertex_color1: An optional color vector for the first graph. If color vectors are given for both graphs, then the isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument for explanation.

edge_color1: An optional edge color vector for the first graph. The matching edges in the two graphs must have matching colors as well. Supply a null pointer here if your graphs are not edge-colored.

<i>edge_color2</i> :	The edge color vector for the second graph.
<i>count</i> :	Point to an integer, the result will be stored here.
<i>node_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two nodes are compatible.
<i>edge_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two edges are compatible.
<i>arg</i> :	Extra argument to supply to functions <i>node_compat_fn</i> and <i>edge_compat_fn</i> .

Returns:

Error code.

Time complexity: exponential.

igraph_get_isomorphisms_vf2 — Collect all isomorphic mappings of two graphs.

```
int igraph_get_isomorphisms_vf2(const igraph_t *graph1,
                                const igraph_t *graph2,
                                const igraph_vector_int_t *vertex_color1,
                                const igraph_vector_int_t *vertex_color2,
                                const igraph_vector_int_t *edge_color1,
                                const igraph_vector_int_t *edge_color2,
                                igraph_vector_ptr_t *maps,
                                igraph_isocompat_t *node_compat_fn,
                                igraph_isocompat_t *edge_compat_fn,
                                void *arg);
```

This function finds all the isomorphic mappings between two simple graphs. It uses the `igraph_isomorphic_function_vf2()` function. Call the function with the same graph as *graph1* and *graph2* to get automorphisms.

Arguments:

<i>graph1</i> :	The first input graph, may be directed or undirected.
<i>graph2</i> :	The second input graph, it must have the same directedness as <i>graph1</i> , or an error will be reported.
<i>vertex_color1</i> :	An optional color vector for the first graph. If color vectors are given for both graphs, then the isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here if your graphs are not colored.
<i>vertex_color2</i> :	An optional color vector for the second graph. See the previous argument for explanation.
<i>edge_color1</i> :	An optional edge color vector for the first graph. The matching edges in the two graphs must have matching colors as well. Supply a null pointer here if your graphs are not edge-colored.
<i>edge_color2</i> :	The edge color vector for the second graph.

<i>maps</i> :	Pointer vector. On return it is empty if the input graphs are not isomorphic. Otherwise it contains pointers to <code>igraph_vector_t</code> objects, each vector is an isomorphic mapping of <i>graph2</i> to <i>graph1</i> . Please note that you need to 1) Destroy the vectors via <code>igraph_vector_destroy()</code> , 2) free them via <code>igraph_free()</code> and then 3) call <code>igraph_vector_ptr_destroy()</code> on the pointer vector to deallocate all memory when <i>maps</i> is no longer needed.
<i>node_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two nodes are compatible.
<i>edge_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two edges are compatible.
<i>arg</i> :	Extra argument to supply to functions <i>node_compat_fn</i> and <i>edge_compat_fn</i> .

Returns:

Error code.

Time complexity: exponential.

`igraph_isohandler_t` — Callback type, called when an isomorphism was found

```
typedef igraph_bool_t igraph_isohandler_t(const igraph_vector_t *map12,  
                                           const igraph_vector_t *map21, void *arg);
```

See the details at the documentation of `igraph_isomorphic_function_vf2()`.

Arguments:

map12: The mapping from the first graph to the second.

map21: The mapping from the second graph to the first, the inverse of *map12* basically.

arg: This extra argument was passed to `igraph_isomorphic_function_vf2()` when it was called.

Returns:

Boolean, whether to continue with the isomorphism search.

`igraph_isocompat_t` — Callback type, called to check whether two vertices or edges are compatible

```
typedef igraph_bool_t igraph_isocompat_t(const igraph_t *graph1,  
                                           const igraph_t *graph2,  
                                           const igraph_integer_t g1_num,  
                                           const igraph_integer_t g2_num,
```

```
void *arg);
```

VF2 (subgraph) isomorphism functions can be restricted by defining relations on the vertices and/or edges of the graphs, and then checking whether the vertices (edges) match according to these relations.

This feature is implemented by two callbacks, one for vertices, one for edges. Every time igraph tries to match a vertex (edge) of the first (sub)graph to a vertex of the second graph, the vertex (edge) compatibility callback is called. The callback returns a logical value, giving whether the two vertices match.

Both callback functions are of type `igraph_isocompat_t`.

Arguments:

graph1: The first graph.

graph2: The second graph.

g1_num: The id of a vertex or edge in the first graph.

g2_num: The id of a vertex or edge in the second graph.

arg: Extra argument to pass to the callback functions.

Returns:

Logical scalar, whether vertex (or edge) *g1_num* in *graph1* is compatible with vertex (or edge) *g2_num* in *graph2*.

igraph_isomorphic_function_vf2 — The generic VF2 interface

```
int igraph_isomorphic_function_vf2(const igraph_t *graph1, const igraph_t *graph2,
                                   const igraph_vector_int_t *vertex_color1,
                                   const igraph_vector_int_t *vertex_color2,
                                   const igraph_vector_int_t *edge_color1,
                                   const igraph_vector_int_t *edge_color2,
                                   igraph_vector_t *map12,
                                   igraph_vector_t *map21,
                                   igraph_isohandler_t *isohandler_fn,
                                   igraph_isocompat_t *node_compat_fn,
                                   igraph_isocompat_t *edge_compat_fn,
                                   void *arg);
```

This function is an implementation of the VF2 isomorphism algorithm, see P. Foggia, C. Sansone, M. Vento, An Improved algorithm for matching large graphs, Proc. of the 3rd IAPR-TC-15 International Workshop on Graph-based Representations, Italy, 2001.

For using it you need to define a callback function of type `igraph_isohandler_t`. This function will be called whenever VF2 finds an isomorphism between the two graphs. The mapping between the two graphs will be also provided to this function. If the callback returns a nonzero value then the search is continued, otherwise it stops. The callback function must not destroy the mapping vectors that are passed to it.

Arguments:

<i>graph1</i> :	The first input graph.
<i>graph2</i> :	The second input graph.
<i>vertex_color1</i> :	An optional color vector for the first graph. If color vectors are given for both graphs, then the isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here if your graphs are not colored.
<i>vertex_color2</i> :	An optional color vector for the second graph. See the previous argument for explanation.
<i>edge_color1</i> :	An optional edge color vector for the first graph. The matching edges in the two graphs must have matching colors as well. Supply a null pointer here if your graphs are not edge-colored.
<i>edge_color2</i> :	The edge color vector for the second graph.
<i>map12</i> :	Pointer to an initialized vector or NULL. If not NULL and the supplied graphs are isomorphic then the permutation taking <i>graph1</i> to <i>graph2</i> is stored here. If not NULL and the graphs are not isomorphic then a zero-length vector is returned.
<i>map21</i> :	This is the same as <i>map12</i> , but for the permutation taking <i>graph2</i> to <i>graph1</i> .
<i>isohandler_fn</i> :	The callback function to be called if an isomorphism is found. See also <code>igraph_isohandler_t</code> .
<i>node_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two nodes are compatible.
<i>edge_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two edges are compatible.
<i>arg</i> :	Extra argument to supply to functions <i>isohandler_fn</i> , <i>node_compat_fn</i> and <i>edge_compat_fn</i> .

Returns:

Error code.

Time complexity: exponential.

igraph_subisomorphic_vf2 — Decide subgraph isomorphism using VF2

```
int igraph_subisomorphic_vf2(const igraph_t *graph1, const igraph_t *graph2,
                             const igraph_vector_int_t *vertex_color1,
                             const igraph_vector_int_t *vertex_color2,
                             const igraph_vector_int_t *edge_color1,
                             const igraph_vector_int_t *edge_color2,
                             igraph_bool_t *iso, igraph_vector_t *map12,
                             igraph_vector_t *map21,
                             igraph_isocompat_t *node_compat_fn,
                             igraph_isocompat_t *edge_compat_fn,
```

```
void *arg);
```

Decides whether a subgraph of *graph1* is isomorphic to *graph2*. It uses `igraph_subisomorphic_function_vf2()`.

Arguments:

<i>graph1</i> :	The first input graph, may be directed or undirected. This is supposed to be the larger graph.
<i>graph2</i> :	The second input graph, it must have the same directedness as <i>graph1</i> . This is supposed to be the smaller graph.
<i>vertex_color1</i> :	An optional color vector for the first graph. If color vectors are given for both graphs, then the subgraph isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here if your graphs are not colored.
<i>vertex_color2</i> :	An optional color vector for the second graph. See the previous argument for explanation.
<i>edge_color1</i> :	An optional edge color vector for the first graph. The matching edges in the two graphs must have matching colors as well. Supply a null pointer here if your graphs are not edge-colored.
<i>edge_color2</i> :	The edge color vector for the second graph.
<i>iso</i> :	Pointer to a boolean. The result of the decision problem is stored here.
<i>map12</i> :	Pointer to a vector or NULL. If not NULL, then an isomorphic mapping from <i>graph1</i> to <i>graph2</i> is stored here.
<i>map21</i> :	Pointer to a vector or NULL. If not NULL, then an isomorphic mapping from <i>graph2</i> to <i>graph1</i> is stored here.
<i>node_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two nodes are compatible.
<i>edge_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two edges are compatible.
<i>arg</i> :	Extra argument to supply to functions <i>node_compat_fn</i> and <i>edge_compat_fn</i> .

Returns:

Error code.

Time complexity: exponential.

igraph_count_subisomorphisms_vf2 — Number of subgraph isomorphisms using VF2

```
int igraph_count_subisomorphisms_vf2(const igraph_t *graph1, const igraph_t *graph2,
                                     const igraph_vector_int_t *vertex_color1,
                                     const igraph_vector_int_t *vertex_color2,
                                     const igraph_vector_int_t *edge_color1,
```

```
const igraph_vector_int_t *edge_color2,  
igraph_integer_t *count,  
igraph_isocompat_t *node_compat_fn,  
igraph_isocompat_t *edge_compat_fn,  
void *arg);
```

Count the number of isomorphisms between subgraphs of *graph1* and *graph2*. This function uses `igraph_subisomorphic_function_vf2()`.

Arguments:

- graph1*: The first input graph, may be directed or undirected. This is supposed to be the larger graph.
- graph2*: The second input graph, it must have the same directedness as *graph1*. This is supposed to be the smaller graph.
- vertex_color1*: An optional color vector for the first graph. If color vectors are given for both graphs, then the subgraph isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here if your graphs are not colored.
- vertex_color2*: An optional color vector for the second graph. See the previous argument for explanation.
- edge_color1*: An optional edge color vector for the first graph. The matching edges in the two graphs must have matching colors as well. Supply a null pointer here if your graphs are not edge-colored.
- edge_color2*: The edge color vector for the second graph.
- count*: Pointer to an integer. The number of subgraph isomorphisms is stored here.
- node_compat_fn*: A pointer to a function of type `igraph_isocompat_t`. This function will be called by the algorithm to determine whether two nodes are compatible.
- edge_compat_fn*: A pointer to a function of type `igraph_isocompat_t`. This function will be called by the algorithm to determine whether two edges are compatible.
- arg*: Extra argument to supply to functions *node_compat_fn* and *edge_compat_fn*.

Returns:

Error code.

Time complexity: exponential.

igraph_get_subisomorphisms_vf2 — Return all subgraph isomorphic mappings.

```
int igraph_get_subisomorphisms_vf2(const igraph_t *graph1,  
                                   const igraph_t *graph2,  
                                   const igraph_vector_int_t *vertex_color1,  
                                   const igraph_vector_int_t *vertex_color2,  
                                   const igraph_vector_int_t *edge_color1,
```

```
const igraph_vector_int_t *edge_color2,  
igraph_vector_ptr_t *maps,  
igraph_isocompat_t *node_compat_fn,  
igraph_isocompat_t *edge_compat_fn,  
void *arg);
```

This function collects all isomorphic mappings of *graph2* to a subgraph of *graph1*. It uses the `igraph_subisomorphic_function_vf2()` function. The graphs should be simple.

Arguments:

- graph1*: The first input graph, may be directed or undirected. This is supposed to be the larger graph.
- graph2*: The second input graph, it must have the same directedness as *graph1*. This is supposed to be the smaller graph.
- vertex_color1*: An optional color vector for the first graph. If color vectors are given for both graphs, then the subgraph isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here if your graphs are not colored.
- vertex_color2*: An optional color vector for the second graph. See the previous argument for explanation.
- edge_color1*: An optional edge color vector for the first graph. The matching edges in the two graphs must have matching colors as well. Supply a null pointer here if your graphs are not edge-colored.
- edge_color2*: The edge color vector for the second graph.
- maps*: Pointer vector. On return it contains pointers to `igraph_vector_t` objects, each vector is an isomorphic mapping of *graph2* to a subgraph of *graph1*. Please note that you need to 1) Destroy the vectors via `igraph_vector_destroy()`, 2) free them via `igraph_free()` and then 3) call `igraph_vector_ptr_destroy()` on the pointer vector to deallocate all memory when *maps* is no longer needed.
- node_compat_fn*: A pointer to a function of type `igraph_isocompat_t`. This function will be called by the algorithm to determine whether two nodes are compatible.
- edge_compat_fn*: A pointer to a function of type `igraph_isocompat_t`. This function will be called by the algorithm to determine whether two edges are compatible.
- arg*: Extra argument to supply to functions *node_compat_fn* and *edge_compat_fn*.

Returns:

Error code.

Time complexity: exponential.

igraph_subisomorphic_function_vf2 — Generic VF2 function for subgraph isomorphism problems

```
int igraph_subisomorphic_function_vf2(const igraph_t *graph1,
                                     const igraph_t *graph2,
                                     const igraph_vector_int_t *vertex_color1,
                                     const igraph_vector_int_t *vertex_color2,
                                     const igraph_vector_int_t *edge_color1,
                                     const igraph_vector_int_t *edge_color2,
                                     igraph_vector_t *map12,
                                     igraph_vector_t *map21,
                                     igraph_isohandler_t *isohandler_fn,
                                     igraph_isocompat_t *node_compat_fn,
                                     igraph_isocompat_t *edge_compat_fn,
                                     void *arg);
```

This function is the pair of `igraph_isomorphic_function_vf2()`, for subgraph isomorphism problems. It searches for subgraphs of *graph1* which are isomorphic to *graph2*. When it finds an isomorphic mapping it calls the supplied callback *isohandler_fn*. The mapping (and its inverse) and the additional *arg* argument are supplied to the callback.

Arguments:

<i>graph1</i> :	The first input graph, may be directed or undirected. This is supposed to be the larger graph.
<i>graph2</i> :	The second input graph, it must have the same directedness as <i>graph1</i> . This is supposed to be the smaller graph.
<i>vertex_color1</i> :	An optional color vector for the first graph. If color vectors are given for both graphs, then the subgraph isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here if your graphs are not colored.
<i>vertex_color2</i> :	An optional color vector for the second graph. See the previous argument for explanation.
<i>edge_color1</i> :	An optional edge color vector for the first graph. The matching edges in the two graphs must have matching colors as well. Supply a null pointer here if your graphs are not edge-colored.
<i>edge_color2</i> :	The edge color vector for the second graph.
<i>map12</i> :	Pointer to a vector or NULL. If not NULL, then an isomorphic mapping from <i>graph1</i> to <i>graph2</i> is stored here.
<i>map21</i> :	Pointer to a vector or NULL. If not NULL, then an isomorphic mapping from <i>graph2</i> to <i>graph1</i> is stored here.
<i>isohandler_fn</i> :	A pointer to a function of type <code>igraph_isohandler_t</code> . This will be called whenever a subgraph isomorphism is found. If the function returns with a non-zero value then the search is continued, otherwise it stops and the function returns.
<i>node_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two nodes are compatible.
<i>edge_compat_fn</i> :	A pointer to a function of type <code>igraph_isocompat_t</code> . This function will be called by the algorithm to determine whether two edges are compatible.
<i>arg</i> :	Extra argument to supply to functions <i>isohandler_fn</i> , <i>node_compat_fn</i> and <i>edge_compat_fn</i> .

Returns:

Error code.

Time complexity: exponential.

The LAD algorithm

The LAD algorithm can search for a subgraph in a larger graph, or check if two graphs are isomorphic. See Christine Solnon: AllDifferent-based Filtering for Subgraph Isomorphism. Artificial Intelligence, 174(12-13):850-864, 2010. <https://doi.org/10.1016/j.artint.2010.05.002> as well as the homepage of the LAD library at <http://liris.cnrs.fr/csolnon/LAD.html> The implementation in igraph is based on LADv1, but it is modified to use igraph's own memory allocation and error handling.

LAD uses the concept of domains to indicate vertex compatibility when matching the pattern graph. Domains can be used to implement matching of colored vertices.

LAD works with both directed and undirected graphs. Graphs with multi-edges are not supported.

igraph_subisomorphic_lad — Check subgraph isomorphism with the LAD algorithm

```
int igraph_subisomorphic_lad(const igraph_t *pattern, const igraph_t *target,
                             const igraph_vector_ptr_t *domains,
                             igraph_bool_t *iso, igraph_vector_t *map,
                             igraph_vector_ptr_t *maps,
                             igraph_bool_t induced, int time_limit);
```

Check whether *pattern* is isomorphic to a subgraph of *target*. The original LAD implementation by Christine Solnon was used as the basis of this code.

See more about LAD at <http://liris.cnrs.fr/csolnon/LAD.html> and in Christine Solnon: AllDifferent-based Filtering for Subgraph Isomorphism. Artificial Intelligence, 174(12-13):850-864, 2010. <https://doi.org/10.1016/j.artint.2010.05.002>

Arguments:

<i>pattern</i> :	The smaller graph, it can be directed or undirected.
<i>target</i> :	The bigger graph, it can be directed or undirected.
<i>domains</i> :	A pointer vector, or a null pointer. If a pointer vector, then it must contain pointers to <code>igraph_vector_t</code> objects and the length of the vector must match the number of vertices in the <i>pattern</i> graph. For each vertex, the ids of the compatible vertices in the target graph are listed.
<i>iso</i> :	Pointer to a boolean, or a null pointer. If not a null pointer, then the boolean is set to TRUE (1) if a subgraph isomorphism is found, and to FALSE (0) otherwise.
<i>map</i> :	Pointer to a vector or a null pointer. If not a null pointer and a subgraph isomorphism is found, the matching vertices from the target graph are listed here, for each vertex (in vertex id order) from the pattern graph.
<i>maps</i> :	Pointer vector or a null pointer. If not a null pointer, then all subgraph isomorphisms are stored in the pointer vector, in <code>igraph_vector_t</code> objects.
<i>induced</i> :	Boolean, whether to search for induced matching subgraphs.

time_limit: Processor time limit in seconds. Supply zero here for no limit. If the time limit is over, then the function signals an error.

Returns:

Error code

See also:

`igraph_subisomorphic_vf2()` for the VF2 algorithm.

Time complexity: exponential.

Example 17.2. File `examples/simple/igraph_subisomorphic_lad.c`

Functions for graphs with 3 or 4 vertices

`igraph_isomorphic_34` — Graph isomorphism for 3-4 vertices

```
int igraph_isomorphic_34(const igraph_t *graph1, const igraph_t *graph2,
                        igraph_bool_t *iso);
```

This function uses precomputed indices to decide isomorphism problems for graphs with only 3 or 4 vertices. Multi-edges and self-loops are ignored by this function.

Arguments:

graph1: The first input graph.

graph2: The second input graph. Must have the same directedness as *graph1*.

iso: Pointer to a boolean, the result is stored here.

Returns:

Error code.

Time complexity: $O(1)$.

`igraph_isoclass` — Determine the isomorphism class of a graph with 3 or 4 vertices

```
int igraph_isoclass(const igraph_t *graph, igraph_integer_t *isoclass);
```

All graphs with a given number of vertices belong to a number of isomorphism classes, with every graph in a given class being isomorphic to each other.

This function gives the isomorphism class (a number) of a graph. Two graphs have the same isomorphism class if and only if they are isomorphic.

The first isomorphism class is numbered zero and it is the empty graph, the last isomorphism class is the full graph. The number of isomorphism class for directed graphs with three vertices is 16 (between 0 and 15), for undirected graph it is only 4. For graphs with four vertices it is 218 (directed) and 11 (undirected).

Multi-edges and self-loops are ignored by this function.

Arguments:

graph: The graph object.

isoclass: Pointer to an integer, the isomorphism class will be stored here.

Returns:

Error code.

See also:

`igraph_isomorphic()`, `igraph_isoclass_subgraph()`, `igraph_isoclass_create()`, `igraph_motifs_randesu()`.

Because of some limitations this function works only for graphs with three or four vertices.

Time complexity: $O(|E|)$, the number of edges in the graph.

igraph_isoclass_subgraph — The isomorphism class of a subgraph of a graph.

```
int igraph_isoclass_subgraph(const igraph_t *graph, igraph_vector_t *vids,
                             igraph_integer_t *isoclass);
```

This function is only implemented for subgraphs with three or four vertices.

Arguments:

graph: The graph object.

vids: A vector containing the vertex ids to be considered as a subgraph. Each vertex id should be included at most once.

isoclass: Pointer to an integer, this will be set to the isomorphism class.

Returns:

Error code.

See also:

`igraph_isoclass()`, `igraph_isomorphic()`, `igraph_isoclass_create()`.

Time complexity: $O((d+n)*n)$, d is the average degree in the network, and n is the number of vertices in *vids*.

igraph_isoclass_create — Creates a graph from the given isomorphism class.

```
int igraph_isoclass_create(igraph_t *graph, igraph_integer_t size,
                           igraph_integer_t number, igraph_bool_t directed);
```

This function is implemented only for graphs with three or four vertices.

Arguments:

graph: Pointer to an uninitialized graph object.

size: The number of vertices to add to the graph.

number: The isomorphism class.

directed: Logical constant, whether to create a directed graph.

Returns:

Error code.

See also:

`igraph_isoclass()`, `igraph_isoclass_subgraph()`, `igraph_isomorphic()`.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the graph to create.

Utility functions

igraph_permute_vertices — Permute the vertices

```
int igraph_permute_vertices(const igraph_t *graph, igraph_t *res,
                           const igraph_vector_t *permutation);
```

This function creates a new graph from the input graph by permuting its vertices according to the specified mapping. Call this function with the output of `igraph_canonical_permutation()` to create the canonical form of a graph.

Arguments:

graph: The input graph.

res: Pointer to an uninitialized graph object. The new graph is created here.

permutation: The permutation to apply. Vertex 0 is mapped to the first element of the vector, vertex 1 to the second, etc. Note that it is not checked that the vector contains every element only once, and no range checking is performed either.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in terms of the number of vertices and edges.

igraph_simplify_and_colorize — Simplify the graph and compute self-loop and edge multiplicities.

```
int igraph_simplify_and_colorize(  
    const igraph_t *graph, igraph_t *res,  
    igraph_vector_int_t *vertex_color, igraph_vector_int_t *edge_color);
```

This function creates a vertex and edge colored simple graph from the input graph. The vertex colors are computed as the number of incident self-loops to each vertex in the input graph. The edge colors are computed as the number of parallel edges in the input graph that were merged to create each edge in the simple graph.

The resulting colored simple graph is suitable for use by isomorphism checking algorithms such as VF2, which only support simple graphs, but can consider vertex and edge colors.

Arguments:

<i>graph</i> :	The graph object, typically having self-loops or multi-edges.
<i>res</i> :	An uninitialized graph object. The result will be stored here
<i>vertex_color</i> :	Computed vertex colors corresponding to self-loop multiplicities.
<i>edge_color</i> :	Computed edge colors corresponding to edge multiplicities

Returns:

Error code.

See also:

`igraph_simplify()`, `igraph_isomorphic_vf2()`, `igraph_subisomorphic_vf2()`

Chapter 18. Graph coloring

igraph_vertex_coloring_greedy — Computes a vertex coloring using a greedy algorithm.

```
int igraph_vertex_coloring_greedy(const igraph_t *graph, igraph_vector_int_t *c
```

This function assigns a "color"—represented as a non-negative integer—to each vertex of the graph in such a way that neighboring vertices never have the same color. The obtained coloring is not necessarily minimal.

Vertices are colored one by one, choosing the smallest color index that differs from that of already colored neighbors. Colors are represented with non-negative integers 0, 1, 2, ...

Arguments:

graph: The input graph.

colors: Pointer to an initialized integer vector. The vertex colors will be stored here.

heuristic: The vertex ordering heuristic to use during greedy coloring. See `igraph_coloring_greedy_t`

Returns:

Error code.

Example 18.1. File `examples/simple/igraph_coloring.c`

igraph_coloring_greedy_t — Ordering heuristics for greedy graph coloring.

```
typedef enum {  
    IGRAPH_COLORING_GREEDY_COLORED_NEIGHBORS = 0  
} igraph_coloring_greedy_t;
```

Ordering heuristics for `igraph_vertex_coloring_greedy()`.

Values:

<code>IGRAPH_COLORING_GREEDY_COLORED_NEIGHBORS</code> :	Choose vertex with largest number of already colored neighbors.
---	---

Chapter 19. Graph motifs, dyad census and triad census

This section deals with functions which find small induced subgraphs in a graph. These were first defined for subgraphs of two and three vertices by Holland and Leinhardt, and named dyad census and triad census.

igraph_dyad_census — Calculating the dyad census as defined by Holland and Leinhardt.

```
int igraph_dyad_census(const igraph_t *graph, igraph_integer_t *mut,
                      igraph_integer_t *asym, igraph_integer_t *null);
```

Dyad census means classifying each pair of vertices of a directed graph into three categories: mutual (there is at least one edge from a to b and also from b to a); asymmetric (there is at least one edge either from a to b or from b to a, but not the other way) and null (no edges between a and b in either direction).

Holland, P.W. and Leinhardt, S. (1970). A Method for Detecting Structure in Sociometric Data. American Journal of Sociology, 70, 492-513.

Arguments:

- graph*: The input graph. For an undirected graph, there are no asymmetric connections.
- mut*: Pointer to an integer, the number of mutual dyads is stored here.
- asym*: Pointer to an integer, the number of asymmetric dyads is stored here.
- null*: Pointer to an integer, the number of null dyads is stored here. In case of an integer overflow (i.e. too many null dyads), -1 will be returned.

Returns:

Error code.

See also:

`igraph_reciprocity()`, `igraph_triad_census()`.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

igraph_triad_census — Triad census, as defined by Davis and Leinhardt

```
int igraph_triad_census(const igraph_t *graph, igraph_vector_t *res);
```

Calculating the triad census means classifying every triple of vertices in a directed graph. A triple can be in one of 16 states:

- 003 A, B, C, the empty graph.
- 012 A->B, C, a graph with a single directed edge.
- 102 A<->B, C, a graph with a mutual connection between two vertices.
- 021D A<-B->C, the binary out-tree.
- 021U A->B<-C, the binary in-tree.
- 021C A->B->C, the directed line.
- 111D A<->B<-C.
- 111U A<->B->C.
- 030T A->B<-C, A->C.
- 030C A<-B<-C, A->C.
- 201 A<->B<->C.
- 120D A<-B->C, A<->C.
- 120U A->B<-C, A<->C.
- 120C A->B->C, A<->C.
- 210 A->B<->C, A<->C.
- 300 A<->B<->C, A<->C, the complete graph.

See also Davis, J.A. and Leinhardt, S. (1972). The Structure of Positive Interpersonal Relations in Small Groups. In J. Berger (Ed.), Sociological Theories in Progress, Volume 2, 218-251. Boston: Houghton Mifflin.

This function calls `igraph_motifs_randesu()` which is an implementation of the FANMOD motif finder tool, see `igraph_motifs_randesu()` for details. Note that the order of the triads is not the same for `igraph_triad_census()` and `igraph_motifs_randesu()`.

Arguments:

- graph*: The input graph. A warning is given for undirected graphs, as the result is undefined for those.
- res*: Pointer to an initialized vector, the result is stored here in the same order as given in the list above. Note that this order is different than the one used by `igraph_motifs_randesu()`.

Returns:

Error code.

See also:

`igraph_motifs_randesu()`, `igraph_dyad_census()`.

Time complexity: TODO.

Finding triangles

igraph_adjacent_triangles — Count the number of triangles a vertex is part of.

```
int igraph_adjacent_triangles(const igraph_t *graph,
                              igraph_vector_t *res,
                              const igraph_vs_t vids);
```

Arguments:

graph: The input graph. Edge directions and multiplicities are ignored.

res: Initilialized vector, the results are stored here.

vids: The vertices to perform the calculation for.

Returns:

Error mode.

See also:

`igraph_list_triangles()` to list them.

Time complexity: $O(d^2 n)$, d is the average vertex degree of the queried vertices, n is their number.

igraph_list_triangles — Find all triangles in a graph.

```
int igraph_list_triangles(const igraph_t *graph,
                          igraph_vector_int_t *res);
```

Arguments:

graph: The input graph, edge directions are ignored. Multiple edges are ignored.

res: Pointer to an initialized integer vector, the result is stored here, in a long list of triples of vertex ids. Each triple is a triangle in the graph. Each triangle is listed exactly once.

Returns:

Error code.

See also:

`igraph_transitivity_undirected()` to count the triangles, `igraph_adjacent_triangles()` to count the triangles a vertex participates in.

Time complexity: $O(d^2 n)$, d is the average degree, n is the number of vertices.

Graph motifs

igraph_motifs_randesu — Count the number of motifs in a graph.

```
int igraph_motifs_randesu(const igraph_t *graph, igraph_vector_t *hist,
                          int size, const igraph_vector_t *cut_prob);
```

Motifs are small weakly connected induced subgraphs of a given structure in a graph. It is argued that the motif profile (i.e. the number of different motifs in the graph) is characteristic for different types of networks and network function is related to the motifs in the graph.

This function is able to find the different motifs of size three and four (i.e. the number of different subgraphs with three and four vertices) in the network.

In a big network the total number of motifs can be very large, so it takes a lot of time to find all of them, a sampling method can be used. This function is capable of doing sampling via the `cut_prob` argument. This argument gives the probability that a branch of the motif search tree will not be explored. See S. Wernicke and F. Rasche: FANMOD: a tool for fast network motif detection, Bioinformatics 22(9), 1152--1153, 2006 for details.

Set the `cut_prob` argument to a zero vector for finding all motifs.

Directed motifs will be counted in directed graphs and undirected motifs in undirected graphs.

Arguments:

- graph*: The graph to find the motifs in.
- hist*: The result of the computation, it gives the number of motifs found for each isomorphism class. See `igraph_isoclass()` for help about isomorphism classes. Note that this function does *not* count isomorphism classes that are not connected and will report NaN (more precisely `IGRAPH_NAN`) for them.
- size*: The size of the motifs to search for. Only three and four are implemented currently. The limitation is not in the motif finding code, but the graph isomorphism code.
- cut_prob*: Vector of probabilities for cutting the search tree at a given level. The first element is the first level, etc. Supply all zeros here (of length `size`) to find all motifs in a graph.

Returns:

Error code.

See also:

`igraph_motifs_randesu_estimate()` for estimating the number of motifs in a graph, this can help to set the `cut_prob` parameter; `igraph_motifs_randesu_no()` to calculate the total number of motifs of a given size in a graph; `igraph_motifs_randesu_callback()` for calling a callback function for every motif found; `igraph_subisomorphic_lad()` for finding subgraphs on more than 4 vertices.

Time complexity: TODO.

Example 19.1. File `examples/simple/igraph_motifs_randesu.c`

`igraph_motifs_randesu_no` — Count the total number of motifs in a graph.

```
int igraph_motifs_randesu_no(const igraph_t *graph, igraph_integer_t *no,
                             int size, const igraph_vector_t *cut_prob);
```

This function counts the total number of motifs in a graph, i.e. the number of (weakly) connected triplets or quadruplets, without assigning isomorphism classes to them.

Arguments:

graph: The graph object to study.

no: Pointer to an integer type, the result will be stored here.

size: The size of the motifs to count.

cut_prob: Vector giving the probabilities that a branch of the search tree will be cut at a given level.

Returns:

Error code.

See also:

`igraph_motifs_randesu()`, `igraph_motifs_randesu_estimate()`.

Time complexity: TODO.

`igraph_motifs_randesu_estimate` — Estimate the total number of motifs in a graph.

```
int igraph_motifs_randesu_estimate(const igraph_t *graph, igraph_integer_t *est,
                                   int size, const igraph_vector_t *cut_prob,
                                   igraph_integer_t sample_size,
                                   const igraph_vector_t *parsample);
```

This function estimates the total number of connected induced subgraphs, called motifs, of a fixed number of vertices. For example, an undirected complete graph on n vertices will have one motif of size n , and n motifs of size $n - 1$. As another example, one triangle and a separate vertex will have zero motifs of size four.

This function is useful for large graphs for which it is not feasible to count all the different motifs, because there are very many of them.

The total number of motifs is estimated by taking a sample of vertices and counts all motifs in which these vertices are included. (There is also a *cut_prob* parameter which gives the probabilities to cut a branch of the search tree.)

Directed motifs will be counted in directed graphs and undirected motifs in undirected graphs.

Arguments:

<i>graph</i> :	The graph object to study.
<i>est</i> :	Pointer to an integer type, the result will be stored here.
<i>size</i> :	The size of the motif to look for.
<i>cut_prob</i> :	Vector giving the probabilities to cut a branch of the search tree and omit counting the motifs in that branch. It contains a probability for each level. Supply size zeros here to count all the motifs in the sample.
<i>sample_size</i> :	The number of vertices to use as the sample. This parameter is only used if the <i>parsample</i> argument is a null pointer.
<i>parsample</i> :	Either pointer to an initialized vector or a null pointer. If a vector then the vertex ids in the vector are used as a sample. If a null pointer then the <i>sample_size</i> argument is used to create a sample of vertices drawn with uniform probability.

Returns:

Error code.

See also:

`igraph_motifs_randesu()`, `igraph_motifs_randesu_no()`.

Time complexity: TODO.

`igraph_motifs_randesu_callback` — Finds motifs in a graph and calls a function for each of them.

```
int igraph_motifs_randesu_callback(const igraph_t *graph, int size,
                                  const igraph_vector_t *cut_prob, igraph_motif_callback_t callback,
                                  void* extra);
```

Similarly to `igraph_motifs_randesu()`, this function is able to find the different motifs of size three and four (i.e. the number of different subgraphs with three and four vertices) in the network. However, instead of counting them, the function will call a callback function for each motif found to allow further tests or post-processing.

The *cut_prob* argument also allows sampling the motifs, just like for `igraph_motifs_randesu()`. Set the *cut_prob* argument to a zero vector for finding all motifs.

Arguments:

<i>graph</i> :	The graph to find the motifs in.
<i>size</i> :	The size of the motifs to search for. Only three and four are implemented currently. The limitation is not in the motif finding code, but the graph isomorphism code.
<i>cut_prob</i> :	Vector of probabilities for cutting the search tree at a given level. The first element is the first level, etc. Supply all zeros here (of length <i>size</i>) to find all motifs in a graph.

callback: A pointer to a function of type `igraph_motifs_handler_t`. This function will be called whenever a new motif is found.

extra: Extra argument to pass to the callback function.

Returns:

Error code.

Time complexity: TODO.

Example 19.2. File `examples/simple/igraph_motifs_randesu.c`

`igraph_motifs_handler_t` — Callback type for `igraph_motifs_randesu_callback`

```
typedef igraph_bool_t igraph_motifs_handler_t(const igraph_t *graph,
                                              igraph_vector_t *vids,
                                              int isoclass,
                                              void* extra);
```

`igraph_motifs_randesu_callback()` calls a specified callback function whenever a new motif is found during a motif search. This callback function must be of type `igraph_motifs_handler_t`. It has the following arguments:

Arguments:

graph: The graph that that algorithm is working on. Of course this must not be modified.

vids: The IDs of the vertices in the motif that has just been found. This vector is owned by the motif search algorithm, so do not modify or destroy it; make a copy of it if you need it later.

isoclass: The isomorphism class of the motif that has just been found. Use `igraph_isoclass` or `igraph_isoclass_subgraph` to find out which isomorphism class belongs to a given motif.

extra: The extra argument that was passed to `igraph_motifs_randesu_callback()`.

Returns:

A logical value, if TRUE (=non-zero), that is interpreted as a request to stop the motif search and return to the caller.

See also:

`igraph_motifs_randesu_callback()`

Chapter 20. Generating layouts for graph drawing

2D layout generators

Layout generator functions (or at least most of them) try to place the vertices and edges of a graph on a 2D plane or in 3D space in a way which visually pleases the human eye.

They take a graph object and a number of parameters as arguments and return an `igraph_matrix_t`, in which each row gives the coordinates of a vertex.

`igraph_layout_random` — Places the vertices uniformly randomly on a plane.

```
int igraph_layout_random(const igraph_t *graph, igraph_matrix_t *res);
```

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized as needed.

Returns:

Error code. The current implementation always returns with success.

Time complexity: $O(|V|)$, the number of vertices.

`igraph_layout_circle` — Places the vertices uniformly on a circle, in the order of vertex ids.

```
int igraph_layout_circle(const igraph_t *graph, igraph_matrix_t *res,
                        igraph_vs_t order);
```

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized as needed.

order: The order of the vertices on the circle. The vertices not included here, will be placed at (0,0). Supply `igraph_vss_all()` here for all vertices, in the order of their vertex ids.

Returns:

Error code.

Time complexity: $O(|V|)$, the number of vertices.

igraph_layout_star — Generates a star-like layout.

```
int igraph_layout_star(const igraph_t *graph, igraph_matrix_t *res,  
                      igraph_integer_t center, const igraph_vector_t *order);
```

Arguments:

- graph*: The input graph. Its edges are ignored by this function.
- res*: Pointer to an initialized matrix object. This will contain the result and will be resized as needed.
- center*: The id of the vertex to put in the center.
- order*: A numeric vector giving the order of the vertices (including the center vertex!). If a null pointer, then the vertices are placed in increasing vertex id order.

Returns:

Error code.

Time complexity: $O(|V|)$, linear in the number of vertices.

See also:

`igraph_layout_circle()` and other layout generators.

igraph_layout_grid — Places the vertices on a regular grid on the plane.

```
int igraph_layout_grid(const igraph_t *graph, igraph_matrix_t *res, long int width,
```

Arguments:

- graph*: Pointer to an initialized graph object.
- res*: Pointer to an initialized matrix object. This will contain the result and will be resized as needed.
- width*: The number of vertices in a single row of the grid. When zero or negative, the width of the grid will be the square root of the number of vertices, rounded up if needed.

Returns:

Error code. The current implementation always returns with success.

Time complexity: $O(|V|)$, the number of vertices.

igraph_layout_graphopt — Optimizes vertex layout via the graphopt algorithm.

```
int igraph_layout_graphopt(const igraph_t *graph, igraph_matrix_t *res,
                           igraph_integer_t niter,
                           igraph_real_t node_charge, igraph_real_t node_mass,
                           igraph_real_t spring_length,
                           igraph_real_t spring_constant,
                           igraph_real_t max_sa_movement,
                           igraph_bool_t use_seed);
```

This is a port of the graphopt layout algorithm by Michael Schmuhl. graphopt version 0.4.1 was rewritten in C and the support for layers was removed (might be added later) and a code was a bit reorganized to avoid some unnecessary steps is the node charge (see below) is zero.

Graphopt uses physical analogies for defining attracting and repelling forces among the vertices and then the physical system is simulated until it reaches an equilibrium. (There is no simulated annealing or anything like that, so a stable fixed point is not guaranteed.)

See also <http://www.schmuhl.org/graphopt/> for the original graphopt.

Arguments:

<i>graph</i> :	The input graph.
<i>res</i> :	Pointer to an initialized matrix, the result will be stored here and its initial contents are used as the starting point of the simulation if the <i>use_seed</i> argument is true. Note that in this case the matrix should have the proper size, otherwise a warning is issued and the supplied values are ignored. If no starting positions are given (or they are invalid) then a random starting position is used. The matrix will be resized if needed.
<i>niter</i> :	Integer constant, the number of iterations to perform. Should be a couple of hundred in general. If you have a large graph then you might want to only do a few iterations and then check the result. If it is not good enough you can feed it in again in the <i>res</i> argument. The original graphopt default is 500.
<i>node_charge</i> :	The charge of the vertices, used to calculate electric repulsion. The original graphopt default is 0.001.
<i>node_mass</i> :	The mass of the vertices, used for the spring forces. The original graphopt defaults to 30.
<i>spring_length</i> :	The length of the springs. The original graphopt defaults to zero.
<i>spring_constant</i> :	The spring constant, the original graphopt defaults to one.
<i>max_sa_movement</i> :	Real constant, it gives the maximum amount of movement allowed in a single step along a single axis. The original graphopt default is 5.
<i>use_seed</i> :	Logical scalar, whether to use the positions in <i>res</i> as a starting configuration. See also <i>res</i> above.

Returns:

Error code.

Time complexity: $O(n(|V|^2+|E|))$, n is the number of iterations, $|V|$ is the number of vertices, $|E|$ the number of edges. If *node_charge* is zero then it is only $O(n|E|)$.

igraph_layout_bipartite — Simple layout for bipartite graphs.

```
int igraph_layout_bipartite(const igraph_t *graph,
                           const igraph_vector_bool_t *types,
                           igraph_matrix_t *res, igraph_real_t hgap,
                           igraph_real_t vgap, long int maxiter);
```

The layout is created by first placing the vertices in two rows, according to their types. Then the positions within the rows are optimized to minimize edge crossings, by calling `igraph_layout_sugiyama()`.

Arguments:

- graph*: The input graph.
- types*: A boolean vector containing ones and zeros, the vertex types. Its length must match the number of vertices in the graph.
- res*: Pointer to an initialized matrix, the result, the x and y coordinates are stored here.
- hgap*: The preferred minimum horizontal gap between vertices in the same layer (i.e. vertices of the same type).
- vgap*: The distance between layers.
- maxiter*: Maximum number of iterations in the crossing minimization stage. 100 is a reasonable default; if you feel that you have too many edge crossings, increase this.

Returns:

Error code.

See also:

`igraph_layout_sugiyama()`.

The DrL layout generator

DrL is a sophisticated layout generator developed and implemented by Shawn Martin et al. As of October 2012 the original DrL homepage is unfortunately not available. You can read more about this algorithm in the following technical report: Martin, S., Brown, W.M., Klavans, R., Boyack, K.W., DrL: Distributed Recursive (Graph) Layout. SAND Reports, 2008. 2936: p. 1-10.

Only a subset of the complete DrL functionality is included in igraph, parallel runs and recursive, multi-level layouting is not supported.

The parameters of the layout are stored in an `igraph_layout_drl_options_t` structure, this can be initialized by calling the function `igraph_layout_drl_options_init()`. The fields

of this structure can then be adjusted by hand if needed. The layout is calculated by an `igraph_layout_drl()` call.

igraph_layout_drl_options_t — Parameters for the DrL layout generator

```
typedef struct igraph_layout_drl_options_t {
    igraph_real_t    edge_cut;
    igraph_integer_t init_iterations;
    igraph_real_t    init_temperature;
    igraph_real_t    init_attraction;
    igraph_real_t    init_damping_mult;
    igraph_integer_t liquid_iterations;
    igraph_real_t    liquid_temperature;
    igraph_real_t    liquid_attraction;
    igraph_real_t    liquid_damping_mult;
    igraph_integer_t expansion_iterations;
    igraph_real_t    expansion_temperature;
    igraph_real_t    expansion_attraction;
    igraph_real_t    expansion_damping_mult;
    igraph_integer_t cooldown_iterations;
    igraph_real_t    cooldown_temperature;
    igraph_real_t    cooldown_attraction;
    igraph_real_t    cooldown_damping_mult;
    igraph_integer_t crunch_iterations;
    igraph_real_t    crunch_temperature;
    igraph_real_t    crunch_attraction;
    igraph_real_t    crunch_damping_mult;
    igraph_integer_t simmer_iterations;
    igraph_real_t    simmer_temperature;
    igraph_real_t    simmer_attraction;
    igraph_real_t    simmer_damping_mult;
} igraph_layout_drl_options_t;
```

Values:

<code>edge_cut</code> :	The edge cutting parameter. Edge cutting is done in the late stages of the algorithm in order to achieve less dense layouts. Edges are cut if there is a lot of stress on them (a large value in the objective function sum). The edge cutting parameter is a value between 0 and 1 with 0 representing no edge cutting and 1 representing maximal edge cutting. The default value is 32/40.
<code>init_iterations</code> :	Number of iterations, initial phase.
<code>init_temperature</code> :	Start temperature, initial phase.
<code>init_attraction</code> :	Attraction, initial phase.
<code>init_damping_mult</code> :	Damping factor, initial phase.
<code>liquid_iterations</code> :	Number of iterations in the liquid phase.
<code>liquid_temperature</code> :	Start temperature in the liquid phase.
<code>liquid_attraction</code> :	Attraction in the liquid phase.

<code>liquid_damping_mult:</code>	Multiplicatie damping factor, liquid phase.
<code>expansion_iterations:</code>	Number of iterations in the expansion phase.
<code>expansion_temperature:</code>	Start temperature in the expansion phase.
<code>expansion_attraction:</code>	Attraction, expansion phase.
<code>expansion_damping_mult:</code>	Damping factor, expansion phase.
<code>cooldown_iterations:</code>	Number of iterations in the cooldown phase.
<code>cooldown_temperature:</code>	Start temperature in the cooldown phase.
<code>cooldown_attraction:</code>	Attraction in the cooldown phase.
<code>cooldown_damping_mult:</code>	Damping fact int the cooldown phase.
<code>crunch_iterations:</code>	Number of iterations in the crunch phase.
<code>crunch_temperature:</code>	Start temperature in the crunch phase.
<code>crunch_attraction:</code>	Attraction in the crunch phase.
<code>crunch_damping_mult:</code>	Damping factor in the crunch phase.
<code>simmer_iterations:</code>	Number of iterations in the simmer phase.
<code>simmer_temperature:</code>	Start temperature in te simmer phase.
<code>simmer_attraction:</code>	Attraction in the simmer phase.
<code>simmer_damping_mult:</code>	Multiplicative damping factor in the simmer phase.

igraph_layout_drl_default_t — Predefined parameter templates for the DrL layout generator

```
typedef enum { IGRAPH_LAYOUT_DRL_DEFAULT = 0,
               IGRAPH_LAYOUT_DRL_COARSEN,
               IGRAPH_LAYOUT_DRL_COARSEST,
               IGRAPH_LAYOUT_DRL_REFINE,
               IGRAPH_LAYOUT_DRL_FINAL
             } igraph_layout_drl_default_t;
```

These constants can be used to initialize a set of DrL parameters. These can then be modified according to the user's needs.

Values:

<code>IGRAPH_LAYOUT_DRL_DEFAULT:</code>	The deafult parameters.
<code>IGRAPH_LAYOUT_DRL_COARSEN:</code>	Slightly modified parameters to get a coarser layout.
<code>IGRAPH_LAYOUT_DRL_COARSEST:</code>	An even coarser layout.

IGRAPH_LAYOUT_DRL_RE- Refine an already calculated layout.
FINE:

IGRAPH_LAYOUT_DRL_FINAL: Finalize an already refined layout.

igraph_layout_drl_options_init — Initialize parameters for the DrL layout generator

```
int igraph_layout_drl_options_init(igraph_layout_drl_options_t *options,  
                                   igraph_layout_drl_default_t templ);
```

This function can be used to initialize the struct holding the parameters for the DrL layout generator. There are a number of predefined templates available, it is a good idea to start from one of these by modifying some parameters.

Arguments:

options: The struct to initialize.

templ: The template to use. Currently the following templates are supplied: IGRAPH_LAYOUT_DRL_DEFAULT, IGRAPH_LAYOUT_DRL_COARSEN, IGRAPH_LAYOUT_DRL_COARSEST, IGRAPH_LAYOUT_DRL_REFINE and IGRAPH_LAYOUT_DRL_FINAL.

Returns:

Error code.

Time complexity: $O(1)$.

igraph_layout_drl — The DrL layout generator

```
int igraph_layout_drl(const igraph_t *graph, igraph_matrix_t *res,  
                      igraph_bool_t use_seed,  
                      igraph_layout_drl_options_t *options,  
                      const igraph_vector_t *weights,  
                      const igraph_vector_bool_t *fixed);
```

This function implements the force-directed DrL layout generator. Please see more in the following technical report: Martin, S., Brown, W.M., Klavans, R., Boyack, K.W., DrL: Distributed Recursive (Graph) Layout. SAND Reports, 2008. 2936: p. 1-10.

Arguments:

graph: The input graph.

use_seed: Logical scalar, if true, then the coordinates supplied in the *res* argument are used as starting points.

res: Pointer to a matrix, the result layout is stored here. It will be resized as needed.

options: The parameters to pass to the layout generator.

weights: Edge weights, pointer to a vector. If this is a null pointer then every edge will have the same weight.

fixed: Pointer to a logical vector, or a null pointer. Originally, this argument was used in the DrL algorithm to keep the nodes marked with this argument as fixed; fixed nodes would then keep their positions in the initial stages of the algorithm. However, due to how the DrL code imported into igraph is organized, it seems that the argument does not do anything and we are not sure whether this is a bug or a feature in DrL. We are leaving the argument here in order not to break the API, but note that at the present stage it has no effect.

Returns:

Error code.

Time complexity: ???.

igraph_layout_drl_3d — The DrL layout generator, 3d version.

```
int igraph_layout_drl_3d(const igraph_t *graph, igraph_matrix_t *res,
                        igraph_bool_t use_seed,
                        igraph_layout_drl_options_t *options,
                        const igraph_vector_t *weights,
                        const igraph_vector_bool_t *fixed);
```

This function implements the force-directed DrL layout generator. Please see more in the technical report: Martin, S., Brown, W.M., Klavans, R., Boyack, K.W., DrL: Distributed Recursive (Graph) Layout. SAND Reports, 2008. 2936: p. 1-10.

This function uses a modified DrL generator that does the layout in three dimensions.

Arguments:

graph: The input graph.

use_seed: Logical scalar, if true, then the coordinates supplied in the *res* argument are used as starting points.

res: Pointer to a matrix, the result layout is stored here. It will be resized as needed.

options: The parameters to pass to the layout generator.

weights: Edge weights, pointer to a vector. If this is a null pointer then every edge will have the same weight.

fixed: Pointer to a logical vector, or a null pointer. This can be used to fix the position of some vertices. Vertices for which it is true will not be moved, but stay at the coordinates given in the *res* matrix. This argument is ignored if it is a null pointer or if *use_seed* is false.

Returns:

Error code.

Time complexity: ???.

See also:

`igraph_layout_drl()` for the standard 2d version.

igraph_layout_fruchterman_reingold — Places the vertices on a plane according to the Fruchterman-Reingold algorithm.

```
int igraph_layout_fruchterman_reingold(const igraph_t *graph,
                                       igraph_matrix_t *res,
                                       igraph_bool_t use_seed,
                                       igraph_integer_t niter,
                                       igraph_real_t start_temp,
                                       igraph_layout_grid_t grid,
                                       const igraph_vector_t *weight,
                                       const igraph_vector_t *minx,
                                       const igraph_vector_t *maxx,
                                       const igraph_vector_t *miny,
                                       const igraph_vector_t *maxy);
```

This is a force-directed layout, see Fruchterman, T.M.J. and Reingold, E.M.: Graph Drawing by Force-directed Placement. Software -- Practice and Experience, 21/11, 1129--1164, 1991.

Arguments:

- graph*: Pointer to an initialized graph object.
- res*: Pointer to an initialized matrix object. This will contain the result and will be resized as needed.
- use_seed*: Logical, if true the supplied values in the *res* argument are used as an initial layout, if false a random initial layout is used.
- niter*: The number of iterations to do. A reasonable default value is 500.
- start_temp*: Start temperature. This is the maximum amount of movement allowed along one axis, within one step, for a vertex. Currently it is decreased linearly to zero during the iteration.
- grid*: Whether to use the (fast but less accurate) grid based version of the algorithm. Possible values: `IGRAPH_LAYOUT_GRID`, `IGRAPH_LAYOUT_NOGRID`, `IGRAPH_LAYOUT_AUTOGRID`. The last one uses the grid based version only for large graphs, currently the ones with more than 1000 vertices.
- weight*: Pointer to a vector containing edge weights, the attraction along the edges will be multiplied by these. It will be ignored if it is a null-pointer.
- minx*: Pointer to a vector, or a `NULL` pointer. If not a `NULL` pointer then the vector gives the minimum “x” coordinate for every vertex.
- maxx*: Same as *minx*, but the maximum “x” coordinates.
- miny*: Pointer to a vector, or a `NULL` pointer. If not a `NULL` pointer then the vector gives the minimum “y” coordinate for every vertex.
- maxy*: Same as *miny*, but the maximum “y” coordinates.

Returns:

Error code.

Time complexity: $O(|V|^2)$ in each iteration, $|V|$ is the number of vertices in the graph.

igraph_layout_kamada_kawai — Places the vertices on a plane according the Kamada-Kawai algorithm.

```
int igraph_layout_kamada_kawai(const igraph_t *graph, igraph_matrix_t *res,
                               igraph_bool_t use_seed, igraph_integer_t maxiter,
                               igraph_real_t epsilon, igraph_real_t kkconst,
                               const igraph_vector_t *weights,
                               const igraph_vector_t *minx, const igraph_vector_t *maxx,
                               const igraph_vector_t *miny, const igraph_vector_t *maxy)
```

This is a force directed layout, see Kamada, T. and Kawai, S.: An Algorithm for Drawing General Undirected Graphs. Information Processing Letters, 31/1, 7--15, 1989.

Arguments:

- graph*: A graph object.
- res*: Pointer to an initialized matrix object. This will contain the result (x-positions in column zero and y-positions in column one) and will be resized if needed.
- use_seed*: Boolean, whether to use the values supplied in the *res* argument as the initial configuration. If zero and there are any limits on the X or Y coordinates, then a random initial configuration is used. Otherwise the vertices are placed on a circle of radius 1 as the initial configuration.
- maxiter*: The maximum number of iterations to perform. A reasonable default value is at least ten (or more) times the number of vertices.
- epsilon*: Stop the iteration, if the maximum delta value of the algorithm is smaller than still. It is safe to leave it at zero, and then *maxiter* iterations are performed.
- kkconst*: The Kamada-Kawai vertex attraction constant. Typical value: number of vertices.
- weights*: Edge weights, larger values will result longer edges.
- minx*: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum “x” coordinate for every vertex.
- maxx*: Same as *minx*, but the maximum “x” coordinates.
- miny*: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum “y” coordinate for every vertex.
- maxy*: Same as *miny*, but the maximum “y” coordinates.

Returns:

Error code.

Time complexity: $O(|V|)$ for each iteration, after an $O(|V|^2 \log|V|)$ initialization step. $|V|$ is the number of vertices in the graph.

igraph_layout_gem — The GEM layout algorithm, as described in Arne Frick, Andreas Ludwig,

```
int igraph_layout_gem(const igraph_t *graph, igraph_matrix_t *res,
                      igraph_bool_t use_seed, igraph_integer_t maxiter,
                      igraph_real_t temp_max, igraph_real_t temp_min,
                      igraph_real_t temp_init);
```

Heiko Mehldau: A Fast Adaptive Layout Algorithm for Undirected Graphs, Proc. Graph Drawing 1994, LNCS 894, pp. 388-403, 1995.

Arguments:

- graph*: The input graph. Edge directions are ignored in directed graphs.
- res*: The result is stored here. If the *use_seed* argument is true (non-zero), then this matrix is also used as the starting point of the algorithm.
- use_seed*: Boolean, whether to use the supplied coordinates in *res* as the starting point. If false (zero), then a uniform random starting point is used.
- maxiter*: The maximum number of iterations to perform. Updating a single vertex counts as an iteration. A reasonable default is $40 * n * n$, where n is the number of vertices. The original paper suggests $4 * n * n$, but this usually only works if the other parameters are set up carefully.
- temp_max*: The maximum allowed local temperature. A reasonable default is the number of vertices.
- temp_min*: The global temperature at which the algorithm terminates (even before reaching *maxiter* iterations). A reasonable default is $1/10$.
- temp_init*: Initial local temperature of all vertices. A reasonable default is the square root of the number of vertices.

Returns:

Error code.

Time complexity: $O(t * n * (n+e))$, where n is the number of vertices, e is the number of edges and t is the number of time steps performed.

igraph_layout_davidson_harel — Davidson-Harel layout algorithm

```
int igraph_layout_davidson_harel(const igraph_t *graph, igraph_matrix_t *res,
                                igraph_bool_t use_seed, igraph_integer_t maxiter,
                                igraph_integer_t fineiter, igraph_real_t cool_
                                igraph_real_t weight_node_dist, igraph_real_t w
                                igraph_real_t weight_edge_lengths,
                                igraph_real_t weight_edge_crossings,
                                igraph_real_t weight_node_edge_dist);
```

This function implements the algorithm by Davidson and Harel, see Ron Davidson, David Harel: Drawing Graphs Nicely Using Simulated Annealing. ACM Transactions on Graphics 15(4), pp. 301-331, 1996.

The algorithm uses simulated annealing and a sophisticated energy function, which is unfortunately hard to parameterize for different graphs. The original publication did not disclose any parameter values, and the ones below were determined by experimentation.

The algorithm consists of two phases, an annealing phase, and a fine-tuning phase. There is no simulated annealing in the second phase.

Our implementation tries to follow the original publication, as much as possible. The only major difference is that coordinates are explicitly kept within the bounds of the rectangle of the layout.

Arguments:

<i>graph</i> :	The input graph, edge directions are ignored.
<i>res</i> :	A matrix, the result is stored here. It can be used to supply start coordinates, see <i>use_seed</i> .
<i>use_seed</i> :	Boolean, whether to use the supplied <i>res</i> as start coordinates.
<i>maxiter</i> :	The maximum number of annealing iterations. A reasonable value for smaller graphs is 10.
<i>fineiter</i> :	The number of fine tuning iterations. A reasonable value is $\max(10, \log_2(n))$ where n is the number of vertices.
<i>cool_fact</i> :	Cooling factor. A reasonable value is 0.75.
<i>weight_node_dist</i> :	Weight for the node-node distances component of the energy function. Reasonable value: 1.0.
<i>weight_border</i> :	Weight for the distance from the border component of the energy function. It can be set to zero, if vertices are allowed to sit on the border.
<i>weight_edge_lengths</i> :	Weight for the edge length component of the energy function, a reasonable value is the density of the graph divided by 10.
<i>weight_edge_crossings</i> :	Weight for the edge crossing component of the energy function, a reasonable default is 1 minus the square root of the density of the graph.
<i>weight_node_edge_dist</i> :	Weight for the node-edge distance component of the energy function. A reasonable value is 1 minus the density, divided by 5.

Returns:

Error code.

Time complexity: one first phase iteration has time complexity $O(n^2+m^2)$, one fine tuning iteration has time complexity $O(mn)$. Time complexity might be smaller if some of the weights of the components of the energy function are set to zero.

igraph_layout_mds — Place the vertices on a plane using multidimensional scaling.


```
int igraph_layout_mds(const igraph_t* graph, igraph_matrix_t *res,
                     const igraph_matrix_t *dist, long int dim);
```

This layout requires a distance matrix, where the intersection of row *i* and column *j* specifies the desired distance between vertex *i* and vertex *j*. The algorithm will try to place the vertices in a space having a given number of dimensions in a way that approximates the distance relations prescribed in the distance matrix. igraph uses the classical multidimensional scaling by Torgerson; for more details, see Cox & Cox: Multidimensional Scaling (1994), Chapman and Hall, London.

If the input graph is disconnected, igraph will decompose it first into its subgraphs, lay out the subgraphs one by one using the appropriate submatrices of the distance matrix, and then merge the layouts using `igraph_layout_merge_dla`. Since `igraph_layout_merge_dla` works for 2D layouts only, you cannot run the MDS layout on disconnected graphs for more than two dimensions.

Warning: if the graph is symmetric to the exchange of two vertices (as is the case with leaves of a tree connecting to the same parent), classical multidimensional scaling may assign the same coordinates to these vertices.

Arguments:

graph: A graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized if needed.

dist: The distance matrix. It must be symmetric and this function does not check whether the matrix is indeed symmetric. Results are unspecified if you pass a non-symmetric matrix here. You can set this parameter to null; in this case, the shortest path lengths between vertices will be used as distances.

dim: The number of dimensions in the embedding space. For 2D layouts, supply 2 here.

Returns:

Error code.

Added in version 0.6.

Time complexity: usually around $O(|V|^2 \text{ dim})$.

igraph_layout_lgl — Force based layout algorithm for large graphs.

```
int igraph_layout_lgl(const igraph_t *graph, igraph_matrix_t *res,
                     igraph_integer_t maxit, igraph_real_t maxdelta,
                     igraph_real_t area, igraph_real_t coolexp,
                     igraph_real_t repulserad, igraph_real_t cellsize,
                     igraph_integer_t proot);
```

This is a layout generator similar to the Large Graph Layout algorithm and program (<http://lgl.sourceforge.net/>). But unlike LGL, this version uses a Fruchterman-Reingold style simulated annealing algorithm for placing the vertices. The speedup is achieved by placing the vertices on a grid and calculating the repulsion only for vertices which are closer to each other than a limit.

Arguments:

<i>graph</i> :	The (initialized) graph object to place.
<i>res</i> :	Pointer to an initialized matrix object to hold the result. It will be resized if needed.
<i>maxit</i> :	The maximum number of cooling iterations to perform for each layout step. A reasonable default is 150.
<i>maxdelta</i> :	The maximum length of the move allowed for a vertex in a single iteration. A reasonable default is the number of vertices.
<i>area</i> :	This parameter gives the area of the square on which the vertices will be placed. A reasonable default value is the number of vertices squared.
<i>coolexp</i> :	The cooling exponent. A reasonable default value is 1.5.
<i>repulserad</i> :	Determines the radius at which vertex-vertex repulsion cancels out attraction of adjacent vertices. A reasonable default value is <i>area</i> times the number of vertices.
<i>cellsize</i> :	The size of the grid cells, one side of the square. A reasonable default value is the fourth root of <i>area</i> (or the square root of the number of vertices if <i>area</i> is also left at its default value).
<i>proot</i> :	The root vertex, this is placed first, its neighbors in the first iteration, second neighbors in the second, etc. If negative then a random vertex is chosen.

Returns:

Error code.

Added in version 0.2.

Time complexity: ideally $O(\text{dia} \cdot \text{maxit} \cdot (|V| + |E|))$, $|V|$ is the number of vertices, dia is the diameter of the graph, worst case complexity is still $O(\text{dia} \cdot \text{maxit} \cdot (|V|^2 + |E|))$, this is the case when all vertices happen to be in the same grid cell.

igraph_layout_reingold_tilford — Reingold-Tilford layout for tree graphs

```
int igraph_layout_reingold_tilford(const igraph_t *graph,
                                   igraph_matrix_t *res,
                                   igraph_neimode_t mode,
                                   const igraph_vector_t *roots,
                                   const igraph_vector_t *rootlevel);
```

Arranges the nodes in a tree where the given node is used as the root. The tree is directed downwards and the parents are centered above its children. For the exact algorithm, see:

Reingold, E and Tilford, J: Tidier drawing of trees. IEEE Trans. Softw. Eng., SE-7(2):223--228, 1981

If the given graph is not a tree, a breadth-first search is executed first to obtain a possible spanning tree.

Arguments:

<i>graph</i> :	The graph object.
----------------	-------------------

res: The result, the coordinates in a matrix. The parameter should point to an initialized matrix object and will be resized.

mode: Specifies which edges to consider when building the tree. If it is `IGRAPH_OUT` then only the outgoing, if it is `IGRAPH_IN` then only the incoming edges of a parent are considered. If it is `IGRAPH_ALL` then all edges are used (this was the behavior in igraph 0.5 and before). This parameter also influences how the root vertices are calculated, if they are not given. See the *roots* parameter.

roots: The index of the root vertex or root vertices. If this is a non-empty vector then the supplied vertex ids are used as the roots of the trees (or a single tree if the graph is connected). If it is a null pointer of a pointer to an empty vector, then the root vertices are automatically calculated based on topological sorting, performed with the opposite mode than the *mode* argument. After the vertices have been sorted, one is selected from each component.

rootlevel: This argument can be useful when drawing forests which are not trees (i.e. they are unconnected and have tree components). It specifies the level of the root vertices for every tree in the forest. It is only considered if not a null pointer and the *roots* argument is also given (and it is not a null pointer of an empty vector).

Returns:

Error code.

Added in version 0.2.

See also:

`igraph_layout_reingold_tilford_circular()`.

Example	20.1.	File	examples/simple/
igraph_layout_reingold_tilford.c			

igraph_layout_reingold_tilford_circular — Circular Reingold-Tilford layout for trees

```
int igraph_layout_reingold_tilford_circular(const igraph_t *graph,
    igraph_matrix_t *res,
    igraph_neimode_t mode,
    const igraph_vector_t *roots,
    const igraph_vector_t *rootlevel);
```

This layout is almost the same as `igraph_layout_reingold_tilford()`, but the tree is drawn in a circular way, with the root vertex in the center.

Arguments:

graph: The graph object.

res: The result, the coordinates in a matrix. The parameter should point to an initialized matrix object and will be resized.

mode: Specifies which edges to consider when building the tree. If it is `IGRAPH_OUT` then only the outgoing, if it is `IGRAPH_IN` then only the incoming edges of a parent

are considered. If it is `IGRAPH_ALL` then all edges are used (this was the behavior in `igraph 0.5` and before). This parameter also influences how the root vertices are calculated, if they are not given. See the `roots` parameter.

roots: The index of the root vertex or root vertices. If this is a non-empty vector then the supplied vertex ids are used as the roots of the trees (or a single tree if the graph is connected). If it is a null pointer of a pointer to an empty vector, then the root vertices are automatically calculated based on topological sorting, performed with the opposite mode than the `mode` argument. After the vertices have been sorted, one is selected from each component.

rootlevel: This argument can be useful when drawing forests which are not trees (i.e. they are unconnected and have tree components). It specifies the level of the root vertices for every tree in the forest. It is only considered if not a null pointer and the `roots` argument is also given (and it is not a null pointer or an empty vector).

Returns:

Error code.

See also:

`igraph_layout_reingold_tilford()`.

`igraph_layout_sugiyama` — Sugiyama layout algorithm for layered directed acyclic graphs.

```
int igraph_layout_sugiyama(const igraph_t *graph, igraph_matrix_t *res,
                           igraph_t *extd_graph, igraph_vector_t *extd_to_orig,
                           const igraph_vector_t* layers, igraph_real_t hgap,
                           long int maxiter, const igraph_vector_t *weights);
```

This layout algorithm is designed for directed acyclic graphs where each vertex is assigned to a layer. Layers are indexed from zero, and vertices of the same layer will be placed on the same horizontal line. The X coordinates of vertices within each layer are decided by the heuristic proposed by Sugiyama et al to minimize edge crossings.

You can also try to lay out undirected graphs, graphs containing cycles, or graphs without an a priori layered assignment with this algorithm. `igraph` will try to eliminate cycles and assign vertices to layers, but there is no guarantee on the quality of the layout in such cases.

The Sugiyama layout may introduce "bends" on the edges in order to obtain a visually more pleasing layout. This is achieved by adding dummy nodes to edges spanning more than one layer. The resulting layout assigns coordinates not only to the nodes of the original graph but also to the dummy nodes. The layout algorithm will also return the extended graph with the dummy nodes. An edge in the original graph may either be mapped to a single edge in the extended graph or a *path* that starts and ends in the original source and target vertex and passes through multiple dummy vertices. In such cases, the user may also request the mapping of the edges of the extended graph back to the edges of the original graph.

For more details, see K. Sugiyama, S. Tagawa and M. Toda, "Methods for Visual Understanding of Hierarchical Systems". IEEE Transactions on Systems, Man and Cybernetics 11(2):109-125, 1981.

Arguments:

<i>graph</i> :	Pointer to an initialized graph object.
<i>res</i> :	Pointer to an initialized matrix object. This will contain the result and will be resized as needed. The first $ V $ rows of the layout will contain the coordinates of the original graph, the remaining rows contain the positions of the dummy nodes. Therefore, you can use the result both with <i>graph</i> or with <i>extended_graph</i> .
<i>extended_graph</i> :	Pointer to an uninitialized graph object or NULL. The extended graph with the added dummy nodes will be returned here. In this graph, each edge points downwards to lower layers, spans exactly one layer and the first $ V $ vertices coincide with the vertices of the original graph.
<i>extd_to_orig_eids</i> :	Pointer to a vector or NULL. If not NULL, the mapping from the edge IDs of the extended graph back to the edge IDs of the original graph will be stored here.
<i>layers</i> :	The layer index for each vertex or NULL if the layers should be determined automatically by <i>igraph</i> .
<i>hgap</i> :	The preferred minimum horizontal gap between vertices in the same layer.
<i>vgap</i> :	The distance between layers.
<i>maxiter</i> :	Maximum number of iterations in the crossing minimization stage. 100 is a reasonable default; if you feel that you have too many edge crossings, increase this.
<i>weights</i> :	Weights of the edges. These are used only if the graph contains cycles; <i>igraph</i> will tend to reverse edges with smaller weights when breaking the cycles.

3D layout generators

`igraph_layout_random_3d` — Places the vertices uniform randomly in a cube.

```
int igraph_layout_random_3d(const igraph_t *graph, igraph_matrix_t *res);
```

Vertex coordinates range from -1 to 1, and are placed in 3 columns of a matrix, with a row for each vertex.

Arguments:

<i>graph</i> :	The graph to place.
<i>res</i> :	Pointer to an initialized matrix object. It will be resized to hold the result.

Returns:

Error code. The current implementation always returns with success.

Added in version 0.2.

Time complexity: $O(|V|)$, the number of vertices.

igraph_layout_sphere — Places vertices (more or less) uniformly on a sphere.

```
int igraph_layout_sphere(const igraph_t *graph, igraph_matrix_t *res);
```

The algorithm was described in the following paper: Distributing many points on a sphere by E.B. Saff and A.B.J. Kuijlaars, *Mathematical Intelligencer* 19.1 (1997) 5--11.

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized as needed.

Returns:

Error code. The current implementation always returns with success.

Added in version 0.2.

Time complexity: $O(|V|)$, the number of vertices in the graph.

igraph_layout_grid_3d — Places the vertices on a regular grid in the 3D space.

```
int igraph_layout_grid_3d(const igraph_t *graph, igraph_matrix_t *res,  
                          long int width, long int height);
```

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized as needed.

width: The number of vertices in a single row of the grid. When zero or negative, the width is determined automatically.

height: The number of vertices in a single column of the grid. When zero or negative, the height is determined automatically.

Returns:

Error code. The current implementation always returns with success.

Time complexity: $O(|V|)$, the number of vertices.

igraph_layout_fruchterman_reingold_3d — 3D Fruchterman-Reingold algorithm.

```
int igraph_layout_fruchterman_reingold_3d(const igraph_t *graph,
    igraph_matrix_t *res,
    igraph_bool_t use_seed,
    igraph_integer_t niter,
    igraph_real_t start_temp,
    const igraph_vector_t *weight,
    const igraph_vector_t *minx,
    const igraph_vector_t *maxx,
    const igraph_vector_t *miny,
    const igraph_vector_t *maxy,
    const igraph_vector_t *minz,
    const igraph_vector_t *maxz);
```

This is the 3D version of the force based Fruchterman-Reingold layout (see `igraph_layout_fruchterman_reingold` for the 2D version)

Arguments:

- graph*: Pointer to an initialized graph object.
- res*: Pointer to an initialized matrix object. This will contain the result and will be resized as needed.
- use_seed*: Logical, if true the supplied values in the *res* argument are used as an initial layout, if false a random initial layout is used.
- niter*: The number of iterations to do. A reasonable default value is 500.
- start_temp*: Start temperature. This is the maximum amount of movement allowed along one axis, within one step, for a vertex. Currently it is decreased linearly to zero during the iteration.
- weight*: Pointer to a vector containing edge weights, the attraction along the edges will be multiplied by these. It will be ignored if it is a null-pointer.
- minx*: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum “x” coordinate for every vertex.
- maxx*: Same as *minx*, but the maximum “x” coordinates.
- miny*: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum “y” coordinate for every vertex.
- maxy*: Same as *miny*, but the maximum “y” coordinates.
- minz*: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum “z” coordinate for every vertex.
- maxz*: Same as *minz*, but the maximum “z” coordinates.

Returns:

Error code.

Added in version 0.2.

Time complexity: $O(|V|^2)$ in each iteration, $|V|$ is the number of vertices in the graph.

igraph_layout_kamada_kawai_3d — 3D version of the Kamada-Kawai layout generator

```
int igraph_layout_kamada_kawai_3d(const igraph_t *graph, igraph_matrix_t *res,
                                   igraph_bool_t use_seed, igraph_integer_t maxiter,
                                   igraph_real_t epsilon, igraph_real_t kkconst,
                                   const igraph_vector_t *weights,
                                   const igraph_vector_t *minx, const igraph_vector_t *maxx,
                                   const igraph_vector_t *miny, const igraph_vector_t *maxy,
                                   const igraph_vector_t *minz, const igraph_vector_t *maxz)
```

This is a force directed layout, see Kamada, T. and Kawai, S.: An Algorithm for Drawing General Undirected Graphs. Information Processing Letters, 31/1, 7--15, 1989.

Arguments:

- graph*: A graph object.
- res*: Pointer to an initialized matrix object. This will contain the result (x-positions in column zero and y-positions in column one) and will be resized if needed.
- use_seed*: Boolean, whether to use the values supplied in the *res* argument as the initial configuration. If zero and there are any limits on the X, Y or Z coordinates, then a random initial configuration is used. Otherwise the vertices are placed uniformly on a sphere of radius 1 as the initial configuration.
- maxiter*: The maximum number of iterations to perform. A reasonable default value is at least ten (or more) times the number of vertices.
- epsilon*: Stop the iteration, if the maximum delta value of the algorithm is smaller than still. It is safe to leave it at zero, and then *maxiter* iterations are performed.
- kkconst*: The Kamada-Kawai vertex attraction constant. Typical value: number of vertices.
- weights*: Edge weights, larger values will result longer edges.
- minx*: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum “x” coordinate for every vertex.
- maxx*: Same as *minx*, but the maximum “x” coordinates.
- miny*: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum “y” coordinate for every vertex.
- maxy*: Same as *miny*, but the maximum “y” coordinates.
- minz*: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum “z” coordinate for every vertex.
- maxz*: Same as *minz*, but the maximum “z” coordinates.

Returns:

Error code.

Time complexity: $O(|V|)$ for each iteration, after an $O(|V|^2 \log|V|)$ initialization step. $|V|$ is the number of vertices in the graph.

Merging layouts

igraph_layout_merge_dla — Merge multiple layouts by using a DLA algorithm

```
int igraph_layout_merge_dla(igraph_vector_ptr_t *thegraphs,  
                             igraph_vector_ptr_t *coords,  
                             igraph_matrix_t *res);
```

First each layout is covered by a circle. Then the layout of the largest graph is placed at the origin. Then the other layouts are placed by the DLA algorithm, larger ones first and smaller ones last.

Arguments:

thegraphs: Pointer vector containing the graph object of which the layouts will be merged.

coords: Pointer vector containing matrix objects with the 2d layouts of the graphs in *thegraphs*.

res: Pointer to an initialized matrix object, the result will be stored here. It will be resized if needed.

Returns:

Error code.

Added in version 0.2. This function is experimental.

Time complexity: TODO.

Chapter 21. Reading and writing graphs from and to files

These functions can write a graph to a file, or read a graph from a file.

Note that as **igraph** uses the traditional C streams, it is possible to read/write files from/to memory, at least on GNU operating systems supporting “non-standard” streams.

Simple edge list and similar formats

igraph_read_graph_edgelist — Reads an edge list from a file and creates a graph.

```
int igraph_read_graph_edgelist(igraph_t *graph, FILE *instream,
                               igraph_integer_t n, igraph_bool_t directed);
```

This format is simply a series of an even number of non-negative integers separated by whitespace. The integers represent vertex IDs. Placing each edge (i.e. pair of integers) on a separate line is not required, but it is recommended for readability. Edges of directed graphs are assumed to be in "from, to" order.

Arguments:

- graph*: Pointer to an uninitialized graph object.
- instream*: Pointer to a stream, it should be readable.
- n*: The number of vertices in the graph. If smaller than the largest integer in the file it will be ignored. It is thus safe to supply zero here.
- directed*: Logical, if true the graph is directed, if false it will be undirected.

Returns:

Error code: `IGRAPH_PARSEERROR`: if there is a problem reading the file, or the file is syntactically incorrect.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges. It is assumed that reading an integer requires $O(1)$ time.

igraph_write_graph_edgelist — Writes the edge list of a graph to a file.

```
int igraph_write_graph_edgelist(const igraph_t *graph, FILE *outstream);
```

One edge is written per line, separated by a single space. For directed graphs edges are written in from, to order.

Arguments:

graph: The graph object to write.

ostream: Pointer to a stream, it should be writable.

Returns:

Error code: `IGRAPH_EFILE` if there is an error writing the file.

Time complexity: $O(|E|)$, the number of edges in the graph. It is assumed that writing an integer to the file requires $O(1)$ time.

igraph_read_graph_ncol — Reads a .ncol file used by LGL.

```
int igraph_read_graph_ncol(igraph_t *graph, FILE *instream,
                           igraph_strvector_t *predefnames,
                           igraph_bool_t names,
                           igraph_add_weights_t weights,
                           igraph_bool_t directed);
```

Also useful for creating graphs from “named” (and optionally weighted) edge lists.

This format is used by the Large Graph Layout program (<http://lgl.sourceforge.net>), and it is simply a symbolic weighted edge list. It is a simple text file with one edge per line. An edge is defined by two symbolic vertex names separated by whitespace. (The symbolic vertex names themselves cannot contain whitespace. They might follow by an optional number, this will be the weight of the edge; the number can be negative and can be in scientific notation. If there is no weight specified to an edge it is assumed to be zero.

The resulting graph is always undirected. LGL cannot deal with files which contain multiple or loop edges, this is however not checked here, as **igraph** is happy with these.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: Pointer to a stream, it should be readable.

predefnames: Pointer to the symbolic names of the vertices in the file. If `NULL` is given here then vertex ids will be assigned to vertex names in the order of their appearance in the `\c .ncol` file. If it is not `NULL` and some unknown vertex names are found in the `\c .ncol` file then new vertex ids will be assigned to them.

names: Logical value, if `TRUE` the symbolic names of the vertices will be added to the graph as a vertex attribute called “name”.

weights: Whether to add the weights of the edges to the graph as an edge attribute called “weight”. `IGRAPH_ADD_WEIGHTS_YES` adds the weights (even if they are not present in the file, in this case they are assumed to be zero). `IGRAPH_ADD_WEIGHTS_NO` does not add any edge attribute. `IGRAPH_ADD_WEIGHTS_IF_PRESENT` adds the attribute if and only if there is at least one explicit edge weight in the input file.

directed: Whether to create a directed graph. As this format was originally used only for undirected graphs there is no information in the file about the directedness of the graph. Set this parameter to `IGRAPH_DIRECTED` or `IGRAPH_UNDIRECTED` to create a directed or undirected graph.

Returns:

Error code: `IGRAPH_PARSEERROR`: if there is a problem reading the file, or the file is syntactically incorrect.

Time complexity: $O(|V|+|E|\log(|V|))$ if we neglect the time required by the parsing. As usual $|V|$ is the number of vertices, while $|E|$ is the number of edges.

See also:

`igraph_read_graph_lgl()`, `igraph_write_graph_ncol()`

igraph_write_graph_ncol — Writes the graph to a file in .ncol format

```
int igraph_write_graph_ncol(const igraph_t *graph, FILE *outstream,
                           const char *names, const char *weights);
```

`.ncol` is a format used by LGL, see `igraph_read_graph_ncol()` for details.

Note that having multiple or loop edges in an `.ncol` file breaks the LGL software but **igraph** does not check for this condition.

Arguments:

graph: The graph to write.

outstream: The stream object to write to, it should be writable.

names: The name of the vertex attribute, if symbolic names are written to the file. If not, supply 0 here.

weights: The name of the edge attribute, if they are also written to the file. If you don't want weights, supply 0 here.

Returns:

Error code: `IGRAPH_EFILE` if there is an error writing the file.

Time complexity: $O(|E|)$, the number of edges. All file operations are expected to have time complexity $O(1)$.

See also:

`igraph_read_graph_ncol()`, `igraph_write_graph_lgl()`

igraph_read_graph_lgl — Reads a graph from an .lgl file

```
int igraph_read_graph_lgl(igraph_t *graph, FILE *instream,
                          igraph_bool_t names,
```

```
igraph_add_weights_t weights,  
igraph_bool_t directed);
```

The `.lgl` format is used by the Large Graph Layout visualization software (<http://lgl.sourceforge.net>), it can describe undirected optionally weighted graphs. From the LGL manual:

The second format is the LGL file format (`.lgl` file suffix). This is yet another graph file format that tries to be as stingy as possible with space, yet keeping the edge file in a human readable (not binary) format. The format itself is like the following:

```
# vertex1name  
vertex2name [optionalWeight]  
vertex3name [optionalWeight]
```

Here, the first vertex of an edge is preceded with a pound sign '#'. Then each vertex that shares an edge with that vertex is listed one per line on subsequent lines.

LGL cannot handle loop and multiple edges or directed graphs, but in **igraph** it is not an error to have multiple and loop edges.

Arguments:

- graph*: Pointer to an uninitialized graph object.
- istream*: A stream, it should be readable.
- names*: Logical value, if TRUE the symbolic names of the vertices will be added to the graph as a vertex attribute called "name".
- weights*: Whether to add the weights of the edges to the graph as an edge attribute called "weight". `IGRAPH_ADD_WEIGHTS_YES` adds the weights (even if they are not present in the file, in this case they are assumed to be zero). `IGRAPH_ADD_WEIGHTS_NO` does not add any edge attribute. `IGRAPH_ADD_WEIGHTS_IF_PRESENT` adds the attribute if and only if there is at least one explicit edge weight in the input file.
- directed*: Whether to create a directed graph. As this format was originally used only for undirected graphs there is no information in the file about the directedness of the graph. Set this parameter to `IGRAPH_DIRECTED` or `IGRAPH_UNDIRECTED` to create a directed or undirected graph.

Returns:

Error code: `IGRAPH_PARSEERROR`: if there is a problem reading the file, or the file is syntactically incorrect.

Time complexity: $O(|V|+|E|\log(|V|))$ if we neglect the time required by the parsing. As usual $|V|$ is the number of vertices, while $|E|$ is the number of edges.

See also:

```
igraph_read_graph_ncol(), igraph_write_graph_lgl()
```

Example 21.1. File `examples/simple/igraph_read_graph_lgl.c`

igraph_write_graph_lgl — Writes the graph to a file in `.lgl` format

```
int igraph_write_graph_lgl(const igraph_t *graph, FILE *outstream,
                           const char *names, const char *weights,
                           igraph_bool_t isolates);
```

`.lgl` is a format used by LGL, see `igraph_read_graph_lgl()` for details.

Note that having multiple or loop edges in an `.lgl` file breaks the LGL software but **igraph** does not check for this condition.

Arguments:

graph: The graph to write.

outstream: The stream object to write to, it should be writable.

names: The name of the vertex attribute, if symbolic names are written to the file. If not supply 0 here.

weights: The name of the edge attribute, if they are also written to the file. If you don't want weights supply 0 here.

isolates: Logical, if TRUE isolated vertices are also written to the file. If FALSE they will be omitted.

Returns:

Error code: `IGRAPH_EFILE` if there is an error writing the file.

Time complexity: $O(|E|)$, the number of edges if *isolates* is FALSE, $O(|V|+|E|)$ otherwise. All file operations are expected to have time complexity $O(1)$.

See also:

`igraph_read_graph_lgl()`, `igraph_write_graph_ncol()`

Example 21.2. File `examples/simple/igraph_write_graph_lgl.c`

igraph_read_graph_dimacs — Read a graph in DIMACS format.

```
int igraph_read_graph_dimacs(igraph_t *graph, FILE *instream,
                             igraph_strvector_t *problem,
                             igraph_vector_t *label,
                             igraph_integer_t *source,
                             igraph_integer_t *target,
                             igraph_vector_t *capacity,
                             igraph_bool_t directed);
```

This function reads the DIMACS file format, more specifically the version for network flow problems, see the files at <ftp://dimacs.rutgers.edu/pub/netflow/general-info/>

This is a line-oriented text file (ASCII) format. The first character of each line defines the type of the line. If the first character is `c` the line is a comment line and it is ignored. There is one problem line (`p` in the file, it must appear before any node and arc descriptor lines. The problem line has three

fields separated by spaces: the problem type (min, max or asn), the number of vertices and number of edges in the graph. Exactly two node identification lines are expected (n), one for the source, one for the target vertex. These have two fields: the id of the vertex and the type of the vertex, either s (=source) or t (=target). Arc lines start with a and have three fields: the source vertex, the target vertex and the edge capacity.

Vertex ids are numbered from 1.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: The file to read from.

source: Pointer to an integer, the id of the source node will be stored here. (The igraph vertex id, which is one less than the actual number in the file.) It is ignored if NULL.

target: Pointer to an integer, the (igraph) id of the target node will be stored here. It is ignored if NULL.

capacity: Pointer to an initialized vector, the capacity of the edges will be stored here if not NULL.

directed: Boolean, whether to create a directed graph.

Returns:

Error code.

Time complexity: $O(|V|+|E|+c)$, the number of vertices plus the number of edges, plus the size of the file in characters.

See also:

`igraph_write_graph_dimacs()`

igraph_write_graph_dimacs — Write a graph in DIMACS format.

```
int igraph_write_graph_dimacs(const igraph_t *graph, FILE *outstream,
                             long int source, long int target,
                             const igraph_vector_t *capacity);
```

This function writes a graph to an output stream in DIMACS format, describing a maximum flow problem. See <ftp://dimacs.rutgers.edu/pub/netflow/general-info/>

This file format is discussed in the documentation of `igraph_read_graph_dimacs()`, see that for more information.

Arguments:

graph: The graph to write to the stream.

outstream: The stream.

source: Integer, the id of the source vertex for the maximum flow.

target: Integer, the id of the target vertex.

capacity: Pointer to an initialized vector containing the edge capacity values.

Returns:

Error code.

Time complexity: $O(|E|)$, the number of edges in the graph.

See also:

`igraph_read_graph_dimacs()`

Binary formats

`igraph_read_graph_graphdb` — Read a graph in the binary graph database format.

```
int igraph_read_graph_graphdb(igraph_t *graph, FILE *instream,
                              igraph_bool_t directed);
```

This is a binary format, used in the graph database for isomorphism testing. From the (now defunct) graph database homepage:

The graphs are stored in a compact binary format, one graph per file. The file is composed of 16 bit words, which are represented using the so-called little-endian convention, i.e. the least significant byte of the word is stored first.

Then, for each node, the file contains the list of edges coming out of the node itself. The list is represented by a word encoding its length, followed by a word for each edge, representing the destination node of the edge. Node numeration is 0-based, so the first node of the graph has index 0.

Only unlabelled graphs are implemented.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: The stream to read from.

directed: Logical scalar, whether to create a directed graph.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

Example	21.3.	File	<code>examples/simple/</code>
<code>igraph_read_graph_graphdb.c</code>			

GraphML format

igraph_read_graph_graphml — Reads a graph from a GraphML file.

```
int igraph_read_graph_graphml(igraph_t *graph, FILE *instream,
                              int index);
```

GraphML is an XML-based file format for representing various types of graphs. Currently only the most basic import functionality is implemented in igraph: it can read GraphML files without nested graphs and hyperedges. Attributes of the graph are loaded only if an attribute interface is attached, i.e. if you use igraph from R or Python.

Graph attribute names are taken from the `attr.name` attributes of the key tags in the GraphML file. Since `attr.name` is not mandatory, igraph will fall back to the `id` attribute of the key tag if `attr.name` is missing.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: A stream, it should be readable.

index: If the GraphML file contains more than one graph, the one specified by this index will be loaded. Indices start from zero, so supply zero here if your GraphML file contains only a single graph.

Returns:

Error code: `IGRAPH_PARSEERROR`: if there is a problem reading the file, or the file is syntactically incorrect. `IGRAPH_UNIMPLEMENTED`: the GraphML functionality was disabled at compile-time

Example 21.4. File `examples/simple/graphml.c`

igraph_write_graph_graphml — Writes the graph to a file in GraphML format

```
int igraph_write_graph_graphml(const igraph_t *graph, FILE *outstream,
                              igraph_bool_t prefixattr);
```

GraphML is an XML-based file format for representing various types of graphs. See the GraphML Primer (<http://graphml.graphdrawing.org/primer/graphml-primer.html>) for detailed format description.

Arguments:

graph: The graph to write.

outstream: The stream object to write to, it should be writable.

prefixattr: Logical value, whether to put a prefix in front of the attribute names to ensure uniqueness if the graph has vertex and edge (or graph) attributes with the same name.

Returns:

Error code: `IGRAPH_EFILE` if there is an error writing the file.

Time complexity: $O(|V|+|E|)$ otherwise. All file operations are expected to have time complexity $O(1)$.

Example 21.5. File `examples/simple/graphml.c`

GML format

`igraph_read_graph_gml` — Read a graph in GML format.

```
int igraph_read_graph_gml(igraph_t *graph, FILE *instream);
```

GML is a simple textual format, see <http://www.fim.uni-passau.de/en/fim/faculty/chairs/theoretische-informatik/projects.html> for details.

Although all syntactically correct GML can be parsed, we implement only a subset of this format, some attributes might be ignored. Here is a list of all the differences:

1. Only node and edge attributes are used, and only if they have a simple type: integer, real or string. So if an attribute is an array or a record, then it is ignored. This is also true if only some values of the attribute are complex.
2. Top level attributes except for `Version` and the first graph attribute are completely ignored.
3. Graph attributes except for node and edge are completely ignored.
4. There is no maximum line length.
5. There is no maximum keyword length.
6. Character entities in strings are not interpreted.
7. We allow `inf` (infinity) and `nan` (not a number) as a real number. This is case insensitive, so `nan`, `NaN` and `NAN` are equal.

Please contact us if you cannot live with these limitations of the GML parser.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: The stream to read the GML file from.

Returns:

Error code.

Time complexity: should be proportional to the length of the file.

See also:

`igraph_read_graph_graphml()` for a more modern format,
`igraph_write_graph_gml()` for writing GML files.

Example 21.6. File `examples/simple/gml.c`

`igraph_write_graph_gml` — Write the graph to a stream in GML format

```
int igraph_write_graph_gml(const igraph_t *graph, FILE *outstream,  
                           const igraph_vector_t *id, const char *creator);
```

GML is a quite general textual format, see <http://www.fim.uni-passau.de/en/fim/faculty/chairs/theoretische-informatik/projects.html> for details.

The graph, vertex and edges attributes are written to the file as well, if they are numeric or string.

As igraph is more forgiving about attribute names, it might be necessary to simplify the them before writing to the GML file. This way we'll have a syntactically correct GML file. The following simple procedure is performed on each attribute name: first the alphanumeric characters are extracted, the others are ignored. Then if the first character is not a letter then the attribute name is prefixed with "igraph". Note that this might result identical names for two attributes, igraph does not check this.

The "id" vertex attribute is treated specially. If the *id* argument is not 0 then it should be a numeric vector with the vertex ids and the "id" vertex attribute is ignored (if there is one). If *id* is 0 and there is a numeric "id" vertex attribute that is used instead. If ids are not specified in either way then the regular igraph vertex ids are used.

Note that whichever way vertex ids are specified, their uniqueness is not checked.

If the graph has edge attributes named "source" or "target" they're silently ignored. GML uses these attributes to specify the edges, so we cannot write them to the file. Rename them before calling this function if you want to preserve them.

Arguments:

<i>graph</i> :	The graph to write to the stream.
<i>outstream</i> :	The stream to write the file to.
<i>id</i> :	Either NULL or a numeric vector with the vertex ids. See details above.
<i>creator</i> :	An optional string to write to the stream in the creator line. If this is 0 then the current date and time is added.

Returns:

Error code.

Time complexity: should be proportional to the number of characters written to the file.

See also:

`igraph_read_graph_gml()` for reading GML files, `igraph_read_graph_graphml()` for a more modern format.

Example 21.7. File `examples/simple/gml.c`

Pajek format

`igraph_read_graph_pajek` — Reads a file in Pajek format

```
int igraph_read_graph_pajek(igraph_t *graph, FILE *instream);
```

Arguments:

graph: Pointer to an uninitialized graph object.

file: An already opened file handler.

Returns:

Error code.

Only a subset of the Pajek format is implemented. This is partially because this format is not very well documented, but also because **igraph** does not support some Pajek features, like multigraphs.

Starting from version 0.6.1 **igraph** reads bipartite (two-mode) graphs from Pajek files and add the `type` vertex attribute for them. Warnings are given for invalid edges, i.e. edges connecting vertices of the same type.

The list of the current limitations:

1. Only `.net` files are supported, Pajek project files (`.paj`) are not. These might be supported in the future if there is need for it.
2. Time events networks are not supported.
3. Hypergraphs (i.e. graphs with non-binary edges) are not supported.
4. Graphs with both directed and non-directed edges are not supported, are they cannot be represented in **igraph**.
5. Only Pajek networks are supported, permutations, hierarchies, clusters and vectors are not.
6. Graphs with multiple edge sets are not supported.

If there are attribute handlers installed, **igraph** also reads the vertex and edge attributes from the file. Most attributes are renamed to be more informative: `color` instead of `c`, `xfact` instead of `x_fact`, `yfact` instead of `y_fact`, `labeldist` instead of `lr`, `labeldegree2` instead of `lphi`, `framewidth` instead of `bw`, `fontsize` instead of `fos`, `rotation` instead of `phi`, `radius` instead of `r`, `diamondratio` instead of `q`, `labeldegree` instead of `la`, `vertexsize` instead of `size`, `color` instead of `ic`, `framecolor` instead of `bc`, `labelcolor` instead of `lc`, these belong to vertices.

Edge attributes are also renamed, `s` to `arrowsize`, `w` to `edgewidth`, `h1` to `hook1`, `h2` to `hook2`, `a1` to `angle1`, `a2` to `angle2`, `k1` to `velocity1`, `k2` to `velocity2`, `ap` to `arrowpos`, `lp` to

labelpos, lr to labelangle, lphi to labelangle2, la to labeldegree, fos to font-size, a to arrowtype, p to linepattern, l to label, lc to labelcolor, c to color.

In addition the following vertex attributes might be added: id if there are vertex ids in the file, x and y or x and y and z if there are vertex coordinates in the file.

The weight edge attribute might be added if there are edge weights present.

See the pajek homepage: <http://vlado.fmf.uni-lj.si/pub/networks/pajek/> for more info on Pajek and the Pajek manual: <http://vlado.fmf.uni-lj.si/pub/networks/pajek/doc/pajekman.pdf> for information on the Pajek file format.

Time complexity: $O(|V|+|E|+|A|)$, $|V|$ is the number of vertices, $|E|$ the number of edges, $|A|$ the number of attributes (vertex + edge) in the graph if there are attribute handlers installed.

See also:

`igraph_write_graph_pajek()` for writing Pajek files,
`igraph_read_graph_graphml()` for reading GraphML files.

Example 21.8. File `examples/simple/foreign.c`

`igraph_write_graph_pajek` — Writes a graph to a file in Pajek format.

```
int igraph_write_graph_pajek(const igraph_t *graph, FILE *outstream);
```

The Pajek vertex and edge parameters (like color) are determined by the attributes of the vertices and edges, of course this requires an attribute handler to be installed. The names of the corresponding vertex and edge attributes are listed at `igraph_read_graph_pajek()`, e.g. the `color` vertex attributes determines the color (`c` in Pajek) parameter.

As of version 0.6.1 `igraph` writes bipartite graphs into Pajek files correctly, i.e. they will be also bipartite when read into Pajek. As Pajek is less flexible for bipartite graphs (the numeric IDs of the vertices must be sorted according to vertex type), `igraph` might need to reorder the vertices when writing a bipartite Pajek file. This effectively means that numeric vertex IDs usually change when a bipartite graph is written to a Pajek file, and then read back into `igraph`.

Early versions of Pajek supported only Windows-style line endings in Pajek files, but recent versions support both Windows and Unix line endings. `igraph` therefore uses the platform-native line endings when the input file is opened in text mode, and uses Unix-style line endings when the input file is opened in binary mode. If you are using an old version of Pajek, you are on Unix and you are having problems reading files written by `igraph` on a Windows machine, convert the line endings manually with a text editor or with `unix2dos` or `iconv` from the command line).

Arguments:

graph: The graph object to write.

outstream: The file to write to. It should be opened and writable. Make sure that you open the file in binary format if you use MS Windows, otherwise end of line characters will be messed up. (`igraph` will be able to read back these messed up files, but Pajek won't.)

Returns:

Error code.

Time complexity: $O(|V|+|E|+|A|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges, $|A|$ the number of attributes (vertex + edge) in the graph if there are attribute handlers installed.

See also:

`igraph_read_graph_pajek()` for reading Pajek graphs,
`igraph_write_graph_graphml()` for writing a graph in GraphML format, this suites
igraph graphs better.

Example 21.9. File `examples/simple/igraph_write_graph_pajek.c`

UCINET's DL file format

`igraph_read_graph_dl` — Read a file in the DL format of UCINET

```
int igraph_read_graph_dl(igraph_t *graph, FILE *instream,
                        igraph_bool_t directed);
```

This is a simple textual file format used by UCINET. See <http://www.analytictech.com/networks/dataentry.htm> for examples. All the forms described here are supported by igraph. Vertex names and edge weights are also supported and they are added as attributes. (If an attribute handler is attached.)

Note the specification does not mention whether the format is case sensitive or not. For igraph DL files are case sensitive, i.e. `Larry` and `larry` are not the same.

Arguments:

graph: Pointer to an uninitialized graph object.
instream: The stream to read the DL file from.
directed: Logical scalar, whether to create a directed file.

Returns:

Error code.

Time complexity: linear in terms of the number of edges and vertices, except for the matrix format, which is quadratic in the number of vertices.

Example 21.10. File `examples/simple/igraph_read_graph_dl.c`

Graphviz format

`igraph_write_graph_dot` — Write the graph to a stream in DOT format

```
int igraph_write_graph_dot(const igraph_t *graph, FILE* ostream);
```

DOT is the format used by the widely known GraphViz software, see <http://www.graphviz.org> for details. The grammar of the DOT format can be found here: <http://www.graphviz.org/doc/info/lang.html>

This is only a preliminary implementation, only the vertices and the edges are written but not the attributes or any visualization information.

Arguments:

graph: The graph to write to the stream.

ostream: The stream to write the file to.

Time complexity: should be proportional to the number of characters written to the file.

See also:

`igraph_write_graph_graphml()` for a more modern format.

Example 21.11. File `examples/simple/dot.c`

Chapter 22. Maximum flows, minimum cuts and related measures

Maximum flows

`igraph_maxflow` — Maximum network flow between a pair of vertices

```
int igraph_maxflow(const igraph_t *graph, igraph_real_t *value,
                  igraph_vector_t *flow, igraph_vector_t *cut,
                  igraph_vector_t *partition, igraph_vector_t *partition2,
                  igraph_integer_t source, igraph_integer_t target,
                  const igraph_vector_t *capacity,
                  igraph_maxflow_stats_t *stats);
```

This function implements the Goldberg-Tarjan algorithm for calculating value of the maximum flow in a directed or undirected graph. The algorithm was given in Andrew V. Goldberg, Robert E. Tarjan: A New Approach to the Maximum-Flow Problem, Journal of the ACM, 35(4), 921-940, 1988.

The input of the function is a graph, a vector of real numbers giving the capacity of the edges and two vertices of the graph, the source and the target. A flow is a function assigning positive real numbers to the edges and satisfying two requirements: (1) the flow value is less than the capacity of the edge and (2) at each vertex except the source and the target, the incoming flow (i.e. the sum of the flow on the incoming edges) is the same as the outgoing flow (i.e. the sum of the flow on the outgoing edges). The value of the flow is the incoming flow at the target vertex. The maximum flow is the flow with the maximum value.

Arguments:

<i>graph</i> :	The input graph, either directed or undirected.
<i>value</i> :	Pointer to a real number, the value of the maximum will be placed here, unless it is a null pointer.
<i>flow</i> :	If not a null pointer, then it must be a pointer to an initialized vector. The vector will be resized, and the flow on each edge will be placed in it, in the order of the edge ids. For undirected graphs this argument is bit trickier, since for these the flow direction is not predetermined by the edge direction. For these graphs the elements of the <i>flow</i> vector can be negative, this means that the flow goes from the bigger vertex id to the smaller one. Positive values mean that the flow goes from the smaller vertex id to the bigger one.
<i>cut</i> :	A null pointer or a pointer to an initialized vector. If not a null pointer, then the minimum cut corresponding to the maximum flow is stored here, i.e. all edge ids that are part of the minimum cut are stored in the vector.
<i>partition</i> :	A null pointer or a pointer to an initialized vector. If not a null pointer, then the first partition of the minimum cut that corresponds to the maximum flow will be placed here. The first partition is always the one that contains the source vertex.
<i>partition2</i> :	A null pointer or a pointer to an initialized vector. If not a null pointer, then the second partition of the minimum cut that corresponds to the maximum flow will be placed here. The second partition is always the one that contains the target vertex.

source: The id of the source vertex.

target: The id of the target vertex.

capacity: Vector containing the capacity of the edges. If NULL, then every edge is considered to have capacity 1.0.

stats: Counts of the number of different operations preformed by the algorithm are stored here.

Returns:

Error code.

Time complexity: $O(|V|^3)$. In practice it is much faster, but i cannot prove a better lower bound for the data structure i've used. In fact, this implementation runs much faster than the `hi_pr` implementation discussed in B. V. Cherkassky and A. V. Goldberg: On implementing the push-relabel method for the maximum flow problem, (Algorithmica, 19:390--410, 1997) on all the graph classes i've tried.

See also:

`igraph_mincut_value()`, `igraph_edge_connectivity()`, `igraph_vertex_connectivity()` for properties based on the maximum flow.

Example 22.1. File `examples/simple/flow.c`

Example 22.2. File `examples/simple/flow2.c`

`igraph_maxflow_value` — Maximum flow in a network with the push/relabel algorithm

```
int igraph_maxflow_value(const igraph_t *graph, igraph_real_t *value,
                        igraph_integer_t source, igraph_integer_t target,
                        const igraph_vector_t *capacity,
                        igraph_maxflow_stats_t *stats);
```

This function implements the Goldberg-Tarjan algorithm for calculating value of the maximum flow in a directed or undirected graph. The algorithm was given in Andrew V. Goldberg, Robert E. Tarjan: A New Approach to the Maximum-Flow Problem, Journal of the ACM, 35(4), 921-940, 1988.

The input of the function is a graph, a vector of real numbers giving the capacity of the edges and two vertices of the graph, the source and the target. A flow is a function assigning positive real numbers to the edges and satisfying two requirements: (1) the flow value is less than the capacity of the edge and (2) at each vertex except the source and the target, the incoming flow (i.e. the sum of the flow on the incoming edges) is the same as the outgoing flow (i.e. the sum of the flow on the outgoing edges). The value of the flow is the incoming flow at the target vertex. The maximum flow is the flow with the maximum value.

According to a theorem by Ford and Fulkerson (L. R. Ford Jr. and D. R. Fulkerson. Maximal flow through a network. Canadian J. Math., 8:399-404, 1956.) the maximum flow between two vertices is the same as the minimum cut between them (also called the minimum s-t cut). So `igraph_st_mincut_value()` gives the same result in all cases as `igraph_maxflow_value()`.

Note that the value of the maximum flow is the same as the minimum cut in the graph.

Arguments:

- graph*: The input graph, either directed or undirected.
- value*: Pointer to a real number, the result will be placed here.
- source*: The id of the source vertex.
- target*: The id of the target vertex.
- capacity*: Vector containing the capacity of the edges. If NULL, then every edge is considered to have capacity 1.0.
- stats*: Counts of the number of different operations preformed by the algorithm are stored here.

Returns:

Error code.

Time complexity: $O(|V|^3)$.

See also:

`igraph_maxflow()` to calculate the actual flow. `igraph_mincut_value()`, `igraph_edge_connectivity()`, `igraph_vertex_connectivity()` for properties based on the maximum flow.

igraph_dominator_tree — Calculates the dominator tree of a flowgraph

```
int igraph_dominator_tree(const igraph_t *graph,
                          igraph_integer_t root,
                          igraph_vector_t *dom,
                          igraph_t *domtree,
                          igraph_vector_t *leftout,
                          igraph_neimode_t mode);
```

A flowgraph is a directed graph with a distinguished start (or root) vertex r , such that for any vertex v , there is a path from r to v . A vertex v dominates another vertex w (not equal to v), if every path from r to w contains v . Vertex v is the immediate dominator of w , $v = \text{idom}(w)$, if v dominates w and every other dominator of w dominates v . The edges $\{(\text{idom}(w), w) \mid w \text{ is not } r\}$ form a directed tree, rooted at r , called the dominator tree of the graph. Vertex v dominates vertex w if and only if v is an ancestor of w in the dominator tree.

This function implements the Lengauer-Tarjan algorithm to construct the dominator tree of a directed graph. For details please see Thomas Lengauer, Robert Endre Tarjan: A fast algorithm for finding dominators in a flowgraph, ACM Transactions on Programming Languages and Systems (TOPLAS) I/1, 121--141, 1979.

Arguments:

- graph*: A directed graph. If it is not a flowgraph, and it contains some vertices not reachable from the root vertex, then these vertices will be collected in the `leftout` vector.
- root*: The id of the root (or source) vertex, this will be the root of the tree.

<i>dom</i> :	Pointer to an initialized vector or a null pointer. If not a null pointer, then the immediate dominator of each vertex will be stored here. For vertices that are not reachable from the root, NaN is stored here. For the root vertex itself, -1 is added.
<i>domtree</i> :	Pointer to an uninitialized <code>igraph_t</code> , or NULL. If not a null pointer, then the dominator tree is returned here. The graph contains the vertices that are unreachable from the root (if any), these will be isolates.
<i>leftout</i> :	Pointer to an initialized vector object, or NULL. If not NULL, then the ids of the vertices that are unreachable from the root vertex (and thus not part of the dominator tree) are stored here.
<i>mode</i> :	Constant, must be <code>IGRAPH_IN</code> or <code>IGRAPH_OUT</code> . If it is <code>IGRAPH_IN</code> , then all directions are considered as opposite to the original one in the input graph.

Returns:

Error code.

Time complexity: very close to $O(|E|+|V|)$, linear in the number of edges and vertices. More precisely, it is $O(|V|+|E|\alpha(|E|,|V|))$, where $\alpha(|E|,|V|)$ is a functional inverse of Ackermann's function.

Example 22.3. File `examples/simple/dominator_tree.c`

`igraph_maxflow_stats_t` — A simple data type to return some statistics from the

```
typedef struct {
    int nopush, norelabel, nogap, nogapnodes, nobfs;
```

push-relabel maximum flow solver.

Arguments:

<i>nopush</i> :	The number of push operations performed.
<i>norelabel</i> :	The number of relabel operations performed.
<i>nogap</i> :	The number of times the gap heuristics was used.
<i>nogapnodes</i> :	The total number of vertices that were omitted from further calculations because of the gap heuristics.
<i>nobfs</i> :	The number of times the reverse BFS was run to assign good values to the height function. This includes an initial run before the whole algorithm, so it is always at least one.

Cuts and minimum cuts

`igraph_st_mincut` — Minimum cut between a source and a target vertex

```
int igraph_st_mincut(const igraph_t *graph, igraph_real_t *value,
                    igraph_vector_t *cut, igraph_vector_t *partition,
                    igraph_vector_t *partition2,
                    igraph_integer_t source, igraph_integer_t target,
                    const igraph_vector_t *capacity);
```

Finds the edge set that has the smallest total capacity among all edge sets that disconnect the source and target vertices.

The calculation is performed using maximum flow techniques, by calling `igraph_maxflow()`.

Arguments:

<i>graph</i> :	The input graph.
<i>value</i> :	Pointer to a real variable, the value of the cut is stored here.
<i>cut</i> :	Pointer to a real vector, the edge ids that are included in the cut are stored here. This argument is ignored if it is a null pointer.
<i>partition</i> :	Pointer to a real vector, the vertex ids of the vertices in the first partition of the cut are stored here. The first partition is always the one that contains the source vertex. This argument is ignored if it is a null pointer.
<i>partition2</i> :	Pointer to a real vector, the vertex ids of the vertices in the second partition of the cut are stored here. The second partition is always the one that contains the target vertex. This argument is ignored if it is a null pointer.
<i>source</i> :	Integer, the id of the source vertex.
<i>target</i> :	Integer, the id of the target vertex.
<i>capacity</i> :	Vector containing the capacity of the edges. If a null pointer, then every edge is considered to have capacity 1.0.

Returns:

Error code.

See also:

`igraph_maxflow()`.

Time complexity: see `igraph_maxflow()`.

igraph_st_mincut_value — The minimum s-t cut in a graph

```
int igraph_st_mincut_value(const igraph_t *graph, igraph_real_t *value,
                           igraph_integer_t source, igraph_integer_t target,
                           const igraph_vector_t *capacity);
```

The minimum s-t cut in a weighted (=valued) graph is the total minimum edge weight needed to remove from the graph to eliminate all paths from a given vertex (*source*) to another vertex (*target*). Directed paths are considered in directed graphs, and undirected paths in undirected graphs.

The minimum s-t cut between two vertices is known to be same as the maximum flow between these two vertices. So this function calls `igraph_maxflow_value()` to do the calculation.

Arguments:

- graph*: The input graph.
- value*: Pointer to a real variable, the result will be stored here.
- source*: The id of the source vertex.
- target*: The id of the target vertex.
- capacity*: Pointer to the capacity vector, it should contain non-negative numbers and its length should be the same the the number of edges in the graph. It can be a null pointer, then every edge has unit capacity.

Returns:

Error code.

Time complexity: $O(|V|^3)$, see also the discussion for `igraph_maxflow_value()`, $|V|$ is the number of vertices.

igraph_all_st_cuts — List all edge-cuts between two vertices in a directed graph

```
int igraph_all_st_cuts(const igraph_t *graph,
                      igraph_vector_ptr_t *cuts,
                      igraph_vector_ptr_t *partitionls,
                      igraph_integer_t source,
                      igraph_integer_t target);
```

This function lists all edge-cuts between a source and a target vertex. Every cut is listed exactly once. The implemented algorithm is described in JS Provan and DR Shier: A Paradigm for listing (s,t)-cuts in graphs, *Algorithmica* 15, 351--372, 1996.

Arguments:

- graph*: The input graph, is must be directed.
- cuts*: An initialized pointer vector, the cuts are stored here. It is a list of pointers to `igraph_vector_t` objects. Each vector will contain the ids of the edges in the cut. This argument is ignored if it is a null pointer. To free all memory allocated for *cuts*, you need call `igraph_vector_destroy()` and then `igraph_free()` on each element, before destroying the pointer vector itself.
- partitionls*: An initialized pointer vector, the list of vertex sets, generating the actual edge cuts, are stored here. Each vector contains a set of vertex ids. If *X* is such a set, then all edges going from *X* to the complement of *X* form an (s,t) edge-cut in the graph. This argument is ignored if it is a null pointer. To free all memory allocated for *partitionls*, you need call `igraph_vector_destroy()` and then `igraph_free()` on each element, before destroying the pointer vector itself.
- source*: The id of the source vertex.
- target*: The id of the target vertex.

Returns:

Error code.

Time complexity: $O(n(|V|+|E|))$, where $|V|$ is the number of vertices, $|E|$ is the number of edges, and n is the number of cuts.

igraph_all_st_mincuts — All minimum s-t cuts of a directed graph

```
int igraph_all_st_mincuts(const igraph_t *graph, igraph_real_t *value,
                          igraph_vector_ptr_t *cuts,
                          igraph_vector_ptr_t *partitionls,
                          igraph_integer_t source,
                          igraph_integer_t target,
                          const igraph_vector_t *capacity);
```

This function lists all edge cuts between two vertices, in a directed graph, with minimum total capacity. Possibly, multiple cuts may have the same total capacity, although there is often only one minimum cut in weighted graphs. It is recommended to supply integer-values capacities. Otherwise, not all minimum cuts may be detected because of numerical roundoff errors. The implemented algorithm is described in JS Provan and DR Shier: A Paradigm for listing (s,t)-cuts in graphs, *Algorithmica* 15, 351--372, 1996.

Arguments:

- | | |
|----------------------|--|
| <i>graph</i> : | The input graph, it must be directed. |
| <i>value</i> : | Pointer to a real number, the value of the minimum cut is stored here, unless it is a null pointer. |
| <i>cuts</i> : | An initialized pointer vector, the cuts are stored here. It is a list of pointers to <code>igraph_vector_t</code> objects. Each vector will contain the ids of the edges in the cut. This argument is ignored if it is a null pointer. To free all memory allocated for <i>cuts</i> , you need call <code>igraph_vector_destroy()</code> and then <code>igraph_free()</code> on each element, before destroying the pointer vector itself. |
| <i>partitionls</i> : | An initialized pointer vector, the list of vertex sets, generating the actual edge cuts, are stored here. Each vector contains a set of vertex ids. If X is such a set, then all edges going from X to the complement of X form an (s,t) edge-cut in the graph. This argument is ignored if it is a null pointer. |
| <i>source</i> : | The id of the source vertex. |
| <i>target</i> : | The id of the target vertex. |
| <i>capacity</i> : | Vector of edge capacities. All capacities must be strictly positive. If this is a null pointer, then all edges are assumed to have capacity one. |

Returns:

Error code.

Time complexity: $O(n(|V|+|E|))+O(F)$, where $|V|$ is the number of vertices, $|E|$ is the number of edges, and n is the number of cuts; $O(F)$ is the time complexity of the maximum flow algorithm, see `igraph_maxflow()`.

Example 22.4. File `examples/simple/igraph_all_st_mincuts.c`

igraph_mincut — Calculates the minimum cut in a graph.

```
int igraph_mincut(const igraph_t *graph,
                  igraph_real_t *value,
                  igraph_vector_t *partition,
                  igraph_vector_t *partition2,
                  igraph_vector_t *cut,
                  const igraph_vector_t *capacity);
```

This function calculates the minimum cut in a graph. The minimum cut is the minimum set of edges which needs to be removed to disconnect the graph. The minimum is calculated using the weights (*capacity*) of the edges, so the cut with the minimum total capacity is calculated.

For directed graphs an implementation based on calculating $2|V|-2$ maximum flows is used. For undirected graphs we use the Stoer-Wagner algorithm, as described in M. Stoer and F. Wagner: A simple min-cut algorithm, Journal of the ACM, 44 585-591, 1997.

The first implementation of the actual cut calculation for undirected graphs was made by Gregory Benison, thanks Greg.

Arguments:

- graph*: The input graph.
- value*: Pointer to a float, the value of the cut will be stored here.
- partition*: Pointer to an initialized vector, the ids of the vertices in the first partition after separating the graph will be stored here. The vector will be resized as needed. This argument is ignored if it is a NULL pointer.
- partition2*: Pointer to an initialized vector the ids of the vertices in the second partition will be stored here. The vector will be resized as needed. This argument is ignored if it is a NULL pointer.
- cut*: Pointer to an initialized vector, the ids of the edges in the cut will be stored here. This argument is ignored if it is a NULL pointer.
- capacity*: A numeric vector giving the capacities of the edges. If a null pointer then all edges have unit capacity.

Returns:

Error code.

See also:

`igraph_mincut_value()`, a simpler interface for calculating the value of the cut only.

Time complexity: for directed graphs it is $O(|V|^4)$, but see the remarks at `igraph_maxflow()`. For undirected graphs it is $O(|V||E|+|V|^2 \log|V|)$. $|V|$ and $|E|$ are the number of vertices and edges respectively.

Example 22.5. File `examples/simple/igraph_mincut.c`

igraph_mincut_value — The minimum edge cut in a graph

```
int igraph_mincut_value(const igraph_t *graph, igraph_real_t *res,  
                        const igraph_vector_t *capacity);
```

The minimum edge cut in a graph is the total minimum weight of the edges needed to remove from the graph to make the graph *not* strongly connected. (If the original graph is not strongly connected then this is zero.) Note that in undirected graphs strong connectedness is the same as weak connectedness.

The minimum cut can be calculated with maximum flow techniques, although the current implementation does this only for directed graphs and a separate non-flow based implementation is used for undirected graphs. See Mechthild Stoer and Frank Wagner: A simple min-cut algorithm, Journal of the ACM 44 585--591, 1997. For directed graphs the maximum flow is calculated between a fixed vertex and all the other vertices in the graph and this is done in both directions. Then the minimum is taken to get the minimum cut.

Arguments:

graph: The input graph.

res: Pointer to a real variable, the result will be stored here.

capacity: Pointer to the capacity vector, it should contain the same number of non-negative numbers as the number of edges in the graph. If a null pointer then all edges will have unit capacity.

Returns:

Error code.

See also:

`igraph_mincut()`, `igraph_maxflow_value()`, `igraph_st_mincut_value()`.

Time complexity: $O(\log(|V|) \cdot |V|^2)$ for undirected graphs and $O(|V|^4)$ for directed graphs, but see also the discussion at the documentation of `igraph_maxflow_value()`.

igraph_gomory_hu_tree — Gomory-Hu tree of a graph.

```
int igraph_gomory_hu_tree(const igraph_t *graph, igraph_t *tree,  
                          igraph_vector_t *flows, const igraph_vector_t *capaci
```

The Gomory-Hu tree is a concise representation of the value of all the maximum flows (or minimum cuts) in a graph. The vertices of the tree correspond exactly to the vertices of the original graph in the same order. Edges of the Gomory-Hu tree are annotated by flow values. The value of the maximum flow (or minimum cut) between an arbitrary (u,v) vertex pair in the original graph is then given by the

minimum flow value (i.e. edge annotation) along the shortest path between u and v in the Gomory-Hu tree.

This implementation uses Gusfield's algorithm to construct the Gomory-Hu tree. See the following paper for more details:

Gusfield D: Very simple methods for all pairs network flow analysis. SIAM J Comput 19(1):143-155, 1990.

Arguments:

graph: The input graph.

tree: Pointer to an uninitialized graph; the result will be stored here.

flows: Pointer to an uninitialized vector; the flow values corresponding to each edge in the Gomory-Hu tree will be returned here. You may pass a NULL pointer here if you are not interested in the flow values.

capacity: Vector containing the capacity of the edges. If NULL, then every edge is considered to have capacity 1.0.

Returns:

Error code.

Time complexity: $O(|V|^4)$ since it performs a max-flow calculation between vertex zero and every other vertex and max-flow is $O(|V|^3)$.

See also:

`igraph_maxflow()`

Connectivity

`igraph_st_edge_connectivity` — Edge connectivity of a pair of vertices

```
int igraph_st_edge_connectivity(const igraph_t *graph, igraph_integer_t *res,
                               igraph_integer_t source,
                               igraph_integer_t target);
```

The edge connectivity of two vertices (*source* and *target*) in a graph is the minimum number of edges that have to be deleted from the graph to eliminate all paths from *source* to *target*.

This function uses the maximum flow algorithm to calculate the edge connectivity.

Arguments:

graph: The input graph, it has to be directed.

res: Pointer to an integer, the result will be stored here.

source: The id of the source vertex.

target: The id of the target vertex.

Returns:

Error code.

Time complexity: $O(|V|^3)$.

See also:

`igraph_maxflow_value()`, `igraph_edge_connectivity()`, `igraph_st_vertex_connectivity()`, `igraph_vertex_connectivity()`.

`igraph_edge_connectivity` — The minimum edge connectivity in a graph.

```
int igraph_edge_connectivity(const igraph_t *graph, igraph_integer_t *res,
                             igraph_bool_t checks);
```

This is the minimum of the edge connectivity over all pairs of vertices in the graph.

The edge connectivity of a graph is the same as group adhesion as defined in Douglas R. White and Frank Harary: The cohesiveness of blocks in social networks: node connectivity and conditional density, Sociological Methodology 31:305--359, 2001.

Arguments:

graph: The input graph.

res: Pointer to an integer, the result will be stored here.

checks: Logical constant. Whether to check that the graph is connected and also the degree of the vertices. If the graph is not (strongly) connected then the connectivity is obviously zero. Otherwise if the minimum degree is one then the edge connectivity is also one. It is a good idea to perform these checks, as they can be done quickly compared to the connectivity calculation itself. They were suggested by Peter McMahan, thanks Peter.

Returns:

Error code.

Time complexity: $O(\log(|V|)*|V|^2)$ for undirected graphs and $O(|V|^4)$ for directed graphs, but see also the discussion at the documentation of `igraph_maxflow_value()`.

See also:

`igraph_st_edge_connectivity()`, `igraph_maxflow_value()`, `igraph_vertex_connectivity()`.

`igraph_st_vertex_connectivity` — The vertex connectivity of a pair of vertices

```
int igraph_st_vertex_connectivity(const igraph_t *graph,
```

```
igraph_integer_t *res,  
igraph_integer_t source,  
igraph_integer_t target,  
igraph_vconn_nei_t neighbors);
```

The vertex connectivity of two vertices (*source* and *target*) is the minimum number of vertices that have to be deleted to eliminate all paths from *source* to *target*. Directed paths are considered in directed graphs.

The vertex connectivity of a pair is the same as the number of different (i.e. node-independent) paths from *source* to *target*.

The current implementation uses maximum flow calculations to obtain the result.

Arguments:

graph: The input graph.

res: Pointer to an integer, the result will be stored here.

source: The id of the source vertex.

target: The id of the target vertex.

neighbors: A constant giving what to do if the two vertices are connected. Possible values: `IGRAPH_VCONN_NEI_ERROR`, stop with an error message, `IGRAPH_VCONN_NEGATIVE`, return -1. `IGRAPH_VCONN_NUMBER_OF_NODES`, return the number of nodes. `IGRAPH_VCONN_IGNORE`, ignore the fact that the two vertices are connected and calculate the number of vertices needed to eliminate all paths except for the trivial (direct) paths between *source* and *vertex*. TODO: what about neighbors?

Returns:

Error code.

Time complexity: $O(|V|^3)$, but see the discussion at `igraph_maxflow_value()`.

See also:

```
igraph_vertex_connectivity(),  
igraph_maxflow_value(),  
igraph_edge_connectivity(),
```

igraph_vertex_connectivity — The vertex connectivity of a graph

```
int igraph_vertex_connectivity(const igraph_t *graph, igraph_integer_t *res,  
                              igraph_bool_t checks);
```

The vertex connectivity of a graph is the minimum vertex connectivity along each pairs of vertices in the graph.

The vertex connectivity of a graph is the same as group cohesion as defined in Douglas R. White and Frank Harary: The cohesiveness of blocks in social networks: node connectivity and conditional density, Sociological Methodology 31:305--359, 2001.

Arguments:

- graph*: The input graph.
- res*: Pointer to an integer, the result will be stored here.
- checks*: Logical constant. Whether to check that the graph is connected and also the degree of the vertices. If the graph is not (strongly) connected then the connectivity is obviously zero. Otherwise if the minimum degree is one then the vertex connectivity is also one. It is a good idea to perform these checks, as they can be done quickly compared to the connectivity calculation itself. They were suggested by Peter McMahan, thanks Peter.

Returns:

Error code.

Time complexity: $O(|V|^5)$.

See also:

`igraph_st_vertex_connectivity()`, `igraph_maxflow_value()`, and
`igraph_edge_connectivity()`.

Edge- and vertex-disjoint paths

`igraph_edge_disjoint_paths` — The maximum number of edge-disjoint paths between two vertices.

```
int igraph_edge_disjoint_paths(const igraph_t *graph, igraph_integer_t *res,
                              igraph_integer_t source,
                              igraph_integer_t target);
```

A set of paths between two vertices is called edge-disjoint if they do not share any edges. The maximum number of edge-disjoint paths are calculated by this function using maximum flow techniques. Directed paths are considered in directed graphs.

Note that the number of disjoint paths is the same as the edge connectivity of the two vertices using uniform edge weights.

Arguments:

- graph*: The input graph, can be directed or undirected.
- res*: Pointer to an integer variable, the result will be stored here.
- source*: The id of the source vertex.
- target*: The id of the target vertex.

Returns:

Error code.

Time complexity: $O(|V|^3)$, but see the discussion at `igraph_maxflow_value()`.

See also:

`igraph_vertex_disjoint_paths()`, `igraph_st_edge_connectivity()`,
`igraph_maxflow_value()`.

`igraph_vertex_disjoint_paths` — Maximum number of vertex-disjoint paths between two vertices.

```
int igraph_vertex_disjoint_paths(const igraph_t *graph, igraph_integer_t *res,
                                igraph_integer_t source,
                                igraph_integer_t target);
```

A set of paths between two vertices is called vertex-disjoint if they share no vertices. The calculation is performed by using maximum flow techniques.

Note that the number of vertex-disjoint paths is the same as the vertex connectivity of the two vertices in most cases (if the two vertices are not connected by an edge).

Arguments:

graph: The input graph.

res: Pointer to an integer variable, the result will be stored here.

source: The id of the source vertex.

target: The id of the target vertex.

Returns:

Error code.

Time complexity: $O(|V|^3)$.

See also:

`igraph_edge_disjoint_paths()`, `igraph_vertex_connectivity()`,
`igraph_maxflow_value()`.

Graph adhesion and cohesion

`igraph_adhesion` — Graph adhesion, this is (almost) the same as edge connectivity.

```
int igraph_adhesion(const igraph_t *graph, igraph_integer_t *res,
                    igraph_bool_t checks);
```

This quantity is defined by White and Harary in The cohesiveness of blocks in social networks: node connectivity and conditional density, (Sociological Methodology 31:305--359, 2001) and basically it is the edge connectivity of the graph with uniform edge weights.

Arguments:

- graph*: The input graph, either directed or undirected.
- res*: Pointer to an integer, the result will be stored here.
- checks*: Logical constant. Whether to check that the graph is connected and also the degree of the vertices. If the graph is not (strongly) connected then the cohesion is obviously zero. Otherwise if the minimum degree is one then the cohesion is also one. It is a good idea to perform these checks, as they can be done quickly compared to the edge connectivity calculation itself. They were suggested by Peter McMahan, thanks Peter. *

Returns:

Error code.

Time complexity: $O(\log(|V|)*|V|^2)$ for undirected graphs and $O(|V|^4)$ for directed graphs, but see also the discussion at the documentation of `igraph_maxflow_value()`.

See also:

`igraph_cohesion()`, `igraph_maxflow_value()`, `igraph_edge_connectivity()`, `igraph_mincut_value()`.

igraph_cohesion — Graph cohesion, this is the same as vertex connectivity.

```
int igraph_cohesion(const igraph_t *graph, igraph_integer_t *res,
                    igraph_bool_t checks);
```

This quantity was defined by White and Harary in “The cohesiveness of blocks in social networks: node connectivity and conditional density”, (Sociological Methodology 31:305--359, 2001) and it is the same as the vertex connectivity of a graph.

Arguments:

- graph*: The input graph.
- res*: Pointer to an integer variable, the result will be stored here.
- checks*: Logical constant. Whether to check that the graph is connected and also the degree of the vertices. If the graph is not (strongly) connected then the cohesion is obviously zero. Otherwise if the minimum degree is one then the cohesion is also one. It is a good idea to perform these checks, as they can be done quickly compared to the vertex connectivity calculation itself. They were suggested by Peter McMahan, thanks Peter.

Returns:

Error code.

Time complexity: $O(|V|^4)$, $|V|$ is the number of vertices. In practice it is more like $O(|V|^2)$, see `igraph_maxflow_value()`.

See also:

```
igraph_vertex_connectivity(),igraph_adhesion(),igraph_maxflow_val-  
ue().
```

Cohesive blocks

igraph_cohesive_blocks — Identifies the hierarchical cohesive block structure of a graph

```
int igraph_cohesive_blocks(const igraph_t *graph,  
                           igraph_vector_ptr_t *blocks,  
                           igraph_vector_t *cohesion,  
                           igraph_vector_t *parent,  
                           igraph_t *block_tree);
```

Cohesive blocking is a method of determining hierarchical subsets of graph vertices based on their structural cohesion (or vertex connectivity). For a given graph G , a subset of its vertices S is said to be maximally k -cohesive if there is no superset of S with vertex connectivity greater than or equal to k . Cohesive blocking is a process through which, given a k -cohesive set of vertices, maximally l -cohesive subsets are recursively identified with $l > k$. Thus a hierarchy of vertex subsets is found, with the entire graph G at its root. See the following reference for details: J. Moody and D. R. White. Structural cohesion and embeddedness: A hierarchical concept of social groups. *American Sociological Review*, 68(1):103--127, Feb 2003.

This function implements cohesive blocking and calculates the complete cohesive block hierarchy of a graph.

Arguments:

- graph*: The input graph. It must be undirected and simple. See `igraph_is_simple()`.
- blocks*: If not a null pointer, then it must be an initialized vector of pointers and the cohesive blocks are stored here. Each block is encoded with a numeric vector, that contains the vertex ids of the block.
- cohesion*: If not a null pointer, then it must be an initialized vector and the cohesion of the blocks is stored here, in the same order as the blocks in the *blocks* pointer vector.
- parent*: If not a null pointer, then it must be an initialized vector and the block hierarchy is stored here. For each block, the id (i.e. the position in the *blocks* pointer vector) of its parent block is stored. For the top block in the hierarchy, -1 is stored.
- block_tree*: If not a null pointer, then it must be a pointer to an uninitialized graph, and the block hierarchy is stored here as an igraph graph. The vertex ids correspond to the order of the blocks in the *blocks* vector.

Returns:

Error code.

Time complexity: TODO.

Example 22.6. File `examples/simple/cohesive_blocks.c`

Chapter 23. Vertex separators

igraph_is_separator — Would removing this set of vertices disconnect the graph?

```
int igraph_is_separator(const igraph_t *graph,
                       const igraph_vs_t candidate,
                       igraph_bool_t *res);
```

Arguments:

graph: The input graph. It may be directed, but edge directions are ignored.

candidate: The candidate separator. It must not contain all vertices.

res: Pointer to a Boolean variable, the result is stored here.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number vertices and edges.

Example 23.1. File `examples/simple/igraph_is_separator.c`

igraph_is_minimal_separator — Decides whether a set of vertices is a minimal separator.

```
int igraph_is_minimal_separator(const igraph_t *graph,
                               const igraph_vs_t candidate,
                               igraph_bool_t *res);
```

A set of vertices is a minimal separator, if the removal of the vertices disconnects the graph, and this is not true for any subset of the set.

This implementation first checks that the given candidate is a separator, by calling `igraph_is_separator()`. If it is a separator, then it checks that each subset of size $n-1$, where n is the size of the candidate, is not a separator.

Arguments:

graph: The input graph. It may be directed, but edge directions are ignored.

candidate: Pointer to a vector of long integers, the candidate minimal separator.

res: Pointer to a boolean variable, the result is stored here.

Returns:

Error code.

Time complexity: $O(n(|V|+|E|))$, $|V|$ is the number of vertices, $|E|$ is the number of edges, n is the number vertices in the candidate separator.

Example	23.2.	File	examples/simple/
igraph_is_minimal_separator.c			

igraph_all_minimal_st_separators — List all vertex sets that are minimal (s,t) separators for some s and t.

```
int igraph_all_minimal_st_separators(const igraph_t *graph,
                                     igraph_vector_ptr_t *separators);
```

This function lists all vertex sets that are minimal (s,t) separators for some (s,t) vertex pair.

See more about the implemented algorithm in Anne Berry, Jean-Paul Bordat and Olivier Cogis: Generating All the Minimal Separators of a Graph, In: Peter Widmayer, Gabriele Neyer and Stephan Eidenbenz (editors): Graph-theoretic concepts in computer science, 1665, 167--172, 1999. Springer.

Arguments:

<i>graph</i> :	The input graph. It may be directed, but edge directions are ignored.
<i>separators</i> :	An initialized pointer vector, the separators are stored here. It is a list of pointers to <code>igraph_vector_t</code> objects. Each vector will contain the ids of the vertices in the separator. To free all memory allocated for separators, you need call <code>igraph_vector_destroy()</code> and then <code>igraph_free()</code> on each element, before destroying the pointer vector itself.

Returns:

Error code.

Time complexity: $O(n|V|^3)$, $|V|$ is the number of vertices, n is the number of separators.

Example	23.3.	File	examples/simple/
igraph_minimum_size_separators.c			

igraph_minimum_size_separators — Find all minimum size separating vertex sets.

```
int igraph_minimum_size_separators(const igraph_t *graph,
                                    igraph_vector_ptr_t *separators);
```

This function lists all separator vertex sets of minimum size. A vertex set is a separator if its removal disconnects the graph.

The implementation is based on the following paper: Arkady Kanevsky: Finding all minimum-size separating vertex sets in a graph, Networks 23, 533--541, 1993.

Arguments:

graph: The input graph, which must be undirected.

separators: An initialized pointer vector, the separators are stored here. It is a list of pointers to `igraph_vector_t` objects. Each vector will contain the ids of the vertices in the separator. To free all memory allocated for separators, you need call `igraph_vector_destroy()` and then `igraph_free()` on each element, before destroying the pointer vector itself.

Returns:

Error code.

Time complexity: TODO.

Example **23.4.** **File** **examples/simple/**
igraph_minimum_size_separators.c

`igraph_even_tarjan_reduction` — Even-Tarjan reduction of a graph

```
int igraph_even_tarjan_reduction(const igraph_t *graph, igraph_t *graphbar,  
                                igraph_vector_t *capacity);
```

A digraph is created with twice as many vertices and edges. For each original vertex i , two vertices $i' = i$ and $i'' = i' + n$ are created, with a directed edge from i' to i'' . For each original directed edge from i to j , two new edges are created, from i' to j'' and from i'' to j' .

This reduction is used in the paper (observation 2): Arkady Kanevsky: Finding all minimum-size separating vertex sets in a graph, Networks 23, 533--541, 1993.

The original paper where this reduction was conceived is Shimon Even and R. Endre Tarjan: Network Flow and Testing Graph Connectivity, SIAM J. Comput., 4(4), 507–518.

Arguments:

graph: A graph. Although directness is not checked, this function is commonly used only on directed graphs.

graphbar: Pointer to a new directed graph that will contain the reduction, with twice as many vertices and edges.

capacity: Pointer to an initialized vector or a null pointer. If not a null pointer, then it will be filled the capacity from the reduction: the first $|E|$ elements are 1, the remaining $|E|$ are equal to $|V|$ (which is used to mean infinity).

Returns:

Error code.

Time complexity: $O(|E|+|V|)$.

Example 23.5. File `examples/simple/even_tarjan.c`

Chapter 24. Detecting community structure

Common functions related to community structure

igraph_modularity — Calculate the modularity of a graph with respect to some clusters or vertex types.

```
int igraph_modularity(const igraph_t *graph,
                     const igraph_vector_t *membership,
                     const igraph_vector_t *weights,
                     const igraph_real_t resolution,
                     const igraph_bool_t directed,
                     igraph_real_t *modularity);
```

The modularity of a graph with respect to some clustering of the vertices (or assignment of vertex types) measures how strongly separated the different clusters are from each other compared to a random null model. It is defined as

$$Q = 1/(2m) \sum_{i,j} (A_{ij} - \gamma * k_i * k_j / (2m)) * d(c_i, c_j),$$

where m is the number of edges, A_{ij} is the adjacency matrix, k_i is the degree of vertex i , c_i is the cluster that vertex i belongs to (or its vertex type), $d(i, j) = 1$ if $i = j$ and 0 otherwise, and the sum goes over all i, j pairs of vertices.

The resolution parameter γ allows weighting the random null model, which might be useful when finding partitions with a high modularity. Maximizing modularity with higher values of the resolution parameter typically results in more, smaller clusters when finding partitions with a high modularity. Lower values typically results in fewer, larger clusters. The original definition of modularity is retrieved when setting $\gamma = 1$.

Modularity can also be calculated on directed graphs. This only requires a relatively modest change

$$Q = 1/(m) \sum_{i,j} (A_{ij} - \gamma * k^{\text{out}}_i * k^{\text{in}}_j / m) * d(c_i, c_j),$$

where k^{out}_i is the out-degree of node i and k^{in}_j is the in-degree of node j .

Modularity on weighted graphs is also meaningful. When taking edge weights into account, A_{ij} equals the weight of the corresponding edge (or 0 if there is no edge), k_i is the strength (i.e. the weighted degree) of vertex i , with similar counterparts for a directed graph, and m is the total weight of all edges.

Note that the modularity is not well-defined for graphs with no edges. igraph returns NaN for graphs with no edges; see <https://github.com/igraph/igraph/issues/1539> for a detailed discussion.

For the original definition of modularity, see Newman, M. E. J., and Girvan, M. (2004). Finding and evaluating community structure in networks. *Physical Review E* 69, 026113. <https://doi.org/10.1103/PhysRevE.69.026113>

For the directed definition of modularity, see Leicht, E. A., and Newman, M. E. J. (2008). Community Structure in Directed Networks. *Physical Review Letters* 100, 118703. <https://doi.org/10.1103/PhysRevLett.100.118703>

For the introduction of the resolution parameter, see Reichardt, J., and Bornholdt, S. (2006). Statistical mechanics of community detection. *Physical Review E* 74, 016110. <https://doi.org/10.1103/PhysRevE.74.016110>

Arguments:

<i>graph</i> :	The input graph.
<i>membership</i> :	Numeric vector of integer values which gives the type of each vertex, i.e. the cluster to which it belongs. It does not have to be consecutive, i.e. empty communities are allowed.
<i>weights</i> :	Weight vector or NULL if no weights are specified.
<i>resolution</i> :	Resolution parameter. Must be greater than or equal to 0. Set it to 1 to use the classical definition of modularity.
<i>directed</i> :	Whether to use the directed or undirected version of modularity. Ignored for undirected graphs.
<i>modularity</i> :	Pointer to a real number, the result will be stored here.

Returns:

Error code.

See also:

`igraph_modularity_matrix()`

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges.

igraph_modularity_matrix — Calculate the modularity matrix

```
int igraph_modularity_matrix(const igraph_t *graph,
                             const igraph_vector_t *weights,
                             const igraph_real_t resolution,
                             igraph_matrix_t *modmat,
                             igraph_bool_t directed);
```

This function returns the modularity matrix defined as

$$B_{ij} = A_{ij} - \gamma * k_i * k_j / (2m)$$

for undirected graphs, where A_{ij} is the adjacency matrix, γ is the resolution parameter, k_i is the degree of vertex i , and m is the number of edges in the graph. When there are no edges, or the weights add up to zero, the result is undefined.

For directed graphs the modularity matrix is changed to

$$B_{ij} = A_{ij} - \gamma * k^{\text{out}}_i * k^{\text{in}}_j / m \text{ where } k^{\text{out}}_i \text{ is the out-degree of node } i \text{ and } k^{\text{in}}_j \text{ is the in-degree of node } j.$$

Note that self-loops in undirected graphs are multiplied by 2 in this implementation. If weights are specified, the weighted counterparts are used.

Arguments:

<i>graph</i> :	The input graph.
<i>weights</i> :	Edge weights, pointer to a vector. If this is a null pointer then every edge is assumed to have a weight of 1.
<i>resolution</i> :	Resolution parameter. Must be greater than or equal to 0. Default is 1. Lower values favor fewer, larger communities; higher values favor more, smaller communities.
<i>modmat</i> :	Pointer to an initialized matrix in which the modularity matrix is stored.
<i>directed</i> :	For directed graphs: if the edges should be treated as undirected. For undirected graphs this is ignored.

See also:

`igraph_modularity()`

`igraph_community_optimal_modularity` — Calculate the community structure with the highest modularity value

```
int igraph_community_optimal_modularity(const igraph_t *graph,
                                       igraph_real_t *modularity,
                                       igraph_vector_t *membership,
                                       const igraph_vector_t *weights);
```

This function calculates the optimal community structure for a graph, in terms of maximal modularity score.

The calculation is done by transforming the modularity maximization into an integer programming problem, and then calling the GLPK library to solve that. Please see Ulrik Brandes et al.: On Modularity Clustering, IEEE Transactions on Knowledge and Data Engineering 20(2):172-188, 2008.

Note that modularity optimization is an NP-complete problem, and all known algorithms for it have exponential time complexity. This means that you probably don't want to run this function on larger graphs. Graphs with up to fifty vertices should be fine, graphs with a couple of hundred vertices might be possible.

Arguments:

<i>graph</i> :	The input graph. It is always treated as undirected.
<i>modularity</i> :	Pointer to a real number, or a null pointer. If it is not a null pointer, then a optimal modularity value is returned here.
<i>membership</i> :	Pointer to a vector, or a null pointer. If not a null pointer, then the membership vector of the optimal community structure is stored here.
<i>weights</i> :	Vector giving the weights of the edges. If it is NULL then each edge is supposed to have the same weight.

Returns:

Error code.

See also:

`igraph_modularity()`, `igraph_community_fastgreedy()` for an algorithm that finds a local optimum in a greedy way.

Time complexity: exponential in the number of vertices.

Example **24.1.** **File** **examples/simple/**
igraph_community_optimal_modularity.c

igraph_community_to_membership — Create membership vector from community structure dendrogram

```
int igraph_community_to_membership(const igraph_matrix_t *merges,
                                   igraph_integer_t nodes,
                                   igraph_integer_t steps,
                                   igraph_vector_t *membership,
                                   igraph_vector_t *csize);
```

This function creates a membership vector from a community structure dendrogram. A membership vector contains for each vertex the id of its graph component, the graph components are numbered from zero, see the same argument of `igraph_clusters()` for an example of a membership vector.

Many community detection algorithms return with a *merges* matrix, `igraph_community_walktrap()` and `igraph_community_edge_betweenness()` are two examples. The matrix contains the merge operations performed while mapping the hierarchical structure of a network. If the matrix has $n-1$ rows, where n is the number of vertices in the graph, then it contains the hierarchical structure of the whole network and it is called a dendrogram.

This function performs *steps* merge operations as prescribed by the *merges* matrix and returns the current state of the network.

If *merges* is not a complete dendrogram, it is possible to take *steps* steps if *steps* is not bigger than the number lines in *merges*.

Arguments:

<i>merges</i> :	The two-column matrix containing the merge operations. See <code>igraph_community_walktrap()</code> for the detailed syntax.
<i>nodes</i> :	The number of leaf nodes in the dendrogram.
<i>steps</i> :	Integer constant, the number of steps to take.
<i>membership</i> :	Pointer to an initialized vector, the membership results will be stored here, if not NULL. The vector will be resized as needed.
<i>csize</i> :	Pointer to an initialized vector, or NULL. If not NULL then the sizes of the components will be stored here, the vector will be resized as needed.

See also:

`igraph_community_walktrap()`, `igraph_community_edge_betweenness()`, `igraph_community_fastgreedy()` for community structure detection algorithms.

Time complexity: $O(|V|)$, the number of vertices in the graph.

igraph_reindex_membership — Makes the IDs in a membership vector continuous

```
int igraph_reindex_membership(igraph_vector_t *membership,
                             igraph_vector_t *new_to_old,
                             igraph_integer_t *nb_clusters);
```

This function reindexes component IDs in a membership vector in a way that the new IDs start from zero and go up to $C-1$, where C is the number of unique component IDs in the original vector. The supplied membership is expected to fall in the range 0, ..., $n-1$.

Arguments:

- membership*: Numeric vector which gives the type of each vertex, i.e. the component to which it belongs. The vector will be altered in-place.
- new_to_old*: Pointer to a vector which will contain the old component ID for each new one, or NULL, in which case it is not returned. The vector will be resized as needed.
- nb_clusters*: Pointer to an integer for the number of distinct clusters. If not NULL, this will be updated to reflect the number of distinct clusters found in membership.

Time complexity: should be $O(n)$ for n elements.

igraph_compare_communities — Compares community structures using various metrics

```
int igraph_compare_communities(const igraph_vector_t *comm1,
                              const igraph_vector_t *comm2, igraph_real_t* res,
                              igraph_community_comparison_t method);
```

This function assesses the distance between two community structures using the variation of information (VI) metric of Meila (2003), the normalized mutual information (NMI) of Danon et al (2005), the split-join distance of van Dongen (2000), the Rand index of Rand (1971) or the adjusted Rand index of Hubert and Arabie (1985).

References:

Meila M: Comparing clusterings by the variation of information. In: Schölkopf B, Warmuth MK (eds.). Learning Theory and Kernel Machines: 16th Annual Conference on Computational Learning Theory and 7th Kernel Workshop, COLT/Kernel 2003, Washington, DC, USA. Lecture Notes in Computer Science, vol. 2777, Springer, 2003. ISBN: 978-3-540-40720-1.

Danon L, Diaz-Guilera A, Duch J, Arenas A: Comparing community structure identification. J Stat Mech P09008, 2005.

van Dongen S: Performance criteria for graph clustering and Markov cluster experiments. Technical Report INS-R0012, National Research Institute for Mathematics and Computer Science in the Netherlands, Amsterdam, May 2000.

Rand WM: Objective criteria for the evaluation of clustering methods. J Am Stat Assoc 66(336):846-850, 1971.

Hubert L and Arabie P: Comparing partitions. *Journal of Classification* 2:193-218, 1985.

Arguments:

comm1: the membership vector of the first community structure

comm2: the membership vector of the second community structure

result: the result is stored here.

method: the comparison method to use. IGRAPH_COMMCMP_VI selects the variation of information (VI) metric of Meila (2003), IGRAPH_COMMCMP_NMI selects the normalized mutual information measure proposed by Danon et al (2005), IGRAPH_COMMCMP_SPLIT_JOIN selects the split-join distance of van Dongen (2000), IGRAPH_COMMCMP_P_RAND selects the unadjusted Rand index (1971) and IGRAPH_COMMCMP_ADJUSTED_RAND selects the adjusted Rand index.

Returns:

Error code.

Time complexity: $O(n \log(n))$.

igraph_split_join_distance — Calculates the split-join distance of two community structures

```
int igraph_split_join_distance(const igraph_vector_t *comm1,
                              const igraph_vector_t *comm2, igraph_integer_t *distance12,
                              igraph_integer_t *distance21);
```

The split-join distance between partitions A and B is the sum of the projection distance of A from B and the projection distance of B from A. The projection distance is an asymmetric measure and it is defined as follows:

First, each set in partition A is evaluated against all sets in partition B. For each set in partition A, the best matching set in partition B is found and the overlap size is calculated. (Matching is quantified by the size of the overlap between the two sets). Then, the maximal overlap sizes for each set in A are summed together and subtracted from the number of elements in A.

The split-join distance will be returned in two arguments, *distance12* will contain the projection distance of the first partition from the second, while *distance21* will be the projection distance of the second partition from the first. This makes it easier to detect whether a partition is a subpartition of the other, since in this case, the corresponding distance will be zero.

Reference:

van Dongen S: Performance criteria for graph clustering and Markov cluster experiments. Technical Report INS-R0012, National Research Institute for Mathematics and Computer Science in the Netherlands, Amsterdam, May 2000.

Arguments:

comm1: the membership vector of the first community structure

comm2: the membership vector of the second community structure

distance12: pointer to an *igraph_integer_t*, the projection distance of the first community structure from the second one will be returned here.

distance21: pointer to an `igraph_integer_t`, the projection distance of the second community structure from the first one will be returned here.

Returns:

Error code.

\see `igraph_compare_communities()` with the `IGRAPH_COMMCMP_SPLIT_JOIN` method if you are not interested in the individual distances but only the sum of them. Time complexity: $O(n \log(n))$.

Community structure based on statistical mechanics

`igraph_community_spinglass` — Community detection based on statistical mechanics

```
int igraph_community_spinglass(const igraph_t *graph,
                              const igraph_vector_t *weights,
                              igraph_real_t *modularity,
                              igraph_real_t *temperature,
                              igraph_vector_t *membership,
                              igraph_vector_t *csize,
                              igraph_integer_t spins,
                              igraph_bool_t parupdate,
                              igraph_real_t starttemp,
                              igraph_real_t stoptemp,
                              igraph_real_t coolfact,
                              igraph_spincomm_update_t update_rule,
                              igraph_real_t gamma,
                              /* the rest is for the NegSpin implementation */
                              igraph_spinglass_implementation_t implementation,
                              /* igraph_matrix_t *adhesion, */
                              /* igraph_matrix_t *normalised_a */
                              /* igraph_real_t *polarization, */
                              igraph_real_t gamma_minus);
```

This function implements the community structure detection algorithm proposed by Joerg Reichardt and Stefan Bornholdt. The algorithm is described in their paper: Statistical Mechanics of Community Detection, <http://arxiv.org/abs/cond-mat/0603718>.

From version 0.6, `igraph` also supports an extension to the algorithm that allows negative edge weights. This is described in V. A. Traag and Jeroen Bruggeman: Community detection in networks with positive and negative links, <http://arxiv.org/abs/0811.2329>.

Arguments:

<i>graph</i> :	The input graph, it may be directed but the direction of the edges is not used in the algorithm.
<i>weights</i> :	The vector giving the edge weights, it may be <code>NULL</code> , in which case all edges are weighted equally. The edge weights must be positive unless using the <code>IGRAPH_SPINCOMM_IMP_NEG</code> implementation. This condition is not verified by the function.

<i>modularity:</i>	Pointer to a real number, if not NULL then the modularity score of the solution will be stored here. This is the generalized modularity that simplifies to the one defined in M. E. J. Newman and M. Girvan, Phys. Rev. E 69, 026113 (2004), if the gamma parameter is one.
<i>temperature:</i>	Pointer to a real number, if not NULL then the temperature at the end of the algorithm will be stored here.
<i>membership:</i>	Pointer to an initialized vector or NULL. If not NULL then the result of the clustering will be stored here. For each vertex, the number of its cluster is given, with the first cluster numbered zero. The vector will be resized as needed.
<i>csize:</i>	Pointer to an initialized vector or NULL. If not NULL then the sizes of the clusters will stored here in cluster number order. The vector will be resized as needed.
<i>spins:</i>	Integer giving the number of spins, i.e. the maximum number of clusters. Usually it is not a program to give a high number here, the default was 25 in the original code. Even if the number of spins is high the number of clusters in the result might be small.
<i>parupdate:</i>	A logical constant, whether to update all spins in parallel. The default for this argument was FALSE (i.e. 0) in the original code. It is not implemented in the IGRAPH_SPINCOMM_INP_NEG implementation.
<i>starttemp:</i>	Real number, the temperature at the start. The value of this argument was 1.0 in the original code.
<i>stoptemp:</i>	Real number, the algorithm stops at this temperature. The default was 0.01 in the original code.
<i>coolfact:</i>	Real number, the cooling factor for the simulated annealing. The default was 0.99 in the original code.
<i>update_rule:</i>	The type of the update rule. Possible values: IGRAPH_SPINCOMM_UPDATE_SIMPLE and IGRAPH_SPINCOMM_UPDATE_CONFIG. Basically this parameter defines the null model based on which the actual clustering is done. If this is IGRAPH_SPINCOMM_UPDATE_SIMPLE then the random graph (i.e. $G(n,p)$), if it is IGRAPH_SPINCOMM_UPDATE then the configuration model is used. The configuration means that the baseline for the clustering is a random graph with the same degree distribution as the input graph.
<i>gamma:</i>	Real number. The gamma parameter of the algorithm. This defines the weight of the missing and existing links in the quality function for the clustering. The default value in the original code was 1.0, which is equal weight to missing and existing edges. Smaller values make the existing links contribute more to the energy function which is minimized in the algorithm. Bigger values make the missing links more important. (If my understanding is correct.)
<i>implementation:</i>	Constant, chooses between the two implementations of the spin-glass algorithm that are included in igraph. IGRAPH_SPINCOMM_IMP_ORIG selects the original implementation, this is faster, IGRAPH_SPINCOMM_INP_NEG selects a new implementation by Vincent Traag that allows negative edge weights.
<i>gamma_minus:</i>	Real number. Parameter for the IGRAPH_SPINCOMM_IMP_NEG implementation. This specifies the balance between the importance of present and non-present negative weighted edges in a community. Smaller values of <i>gamma_minus</i> lead to communities with lesser negative intra-connectivi-

ty. If this argument is set to zero, the algorithm reduces to a graph coloring algorithm, using the number of spins as the number of colors.

Returns:

Error code.

See also:

`igraph_community_spinglass_single()` for calculating the community of a single vertex.

Time complexity: TODO.

igraph_community_spinglass_single — Community of a single node based on statistical mechanics

```
int igraph_community_spinglass_single(const igraph_t *graph,
                                     const igraph_vector_t *weights,
                                     igraph_integer_t vertex,
                                     igraph_vector_t *community,
                                     igraph_real_t *cohesion,
                                     igraph_real_t *adhesion,
                                     igraph_integer_t *inner_links,
                                     igraph_integer_t *outer_links,
                                     igraph_integer_t spins,
                                     igraph_spincomm_update_t update_rule,
                                     igraph_real_t gamma);
```

This function implements the community structure detection algorithm proposed by Joerg Reichardt and Stefan Bornholdt. It is described in their paper: Statistical Mechanics of Community Detection, <http://arxiv.org/abs/cond-mat/0603718>.

This function calculates the community of a single vertex without calculating all the communities in the graph.

Arguments:

<i>graph</i> :	The input graph, it may be directed but the direction of the edges is not used in the algorithm.
<i>weights</i> :	Pointer to a vector with the weights of the edges. Alternatively NULL can be supplied to have the same weight for every edge.
<i>vertex</i> :	The vertex id of the vertex of which ths community is calculated.
<i>community</i> :	Pointer to an initialized vector, the result, the ids of the vertices in the community of the input vertex will be stored here. The vector will be resized as needed.
<i>cohesion</i> :	Pointer to a real variable, if not NULL the cohesion index of the community will be stored here.
<i>adhesion</i> :	Pointer to a real variable, if not NULL the adhesion index of the community will be stored here.
<i>inner_links</i> :	Pointer to an integer, if not NULL the number of edges within the community is stored here.

outer_links: Pointer to an integer, if not NULL the number of edges between the community and the rest of the graph will be stored here.

spins: The number of spins to use, this can be higher than the actual number of clusters in the network, in which case some clusters will contain zero vertices.

update_rule: The type of the update rule. Possible values: `IGRAPH_SPINCOMM_UPDATE_SIMPLE` and `IGRAPH_SPINCOMM_UPDATE_CONFIG`. Basically this parameter defined the null model based on which the actual clustering is done. If this is `IGRAPH_SPINCOMM_UPDATE_SIMPLE` then the random graph (ie. $G(n,p)$), if it is `IGRAPH_SPINCOMM_UPDATE_CONFIG` then the configuration model is used. The configuration means that the baseline for the clustering is a random graph with the same degree distribution as the input graph.

gamma: Real number. The gamma parameter of the algorithm. This defined the weight of the missing and existing links in the quality function for the clustering. The default value in the original code was 1.0, which is equal weight to missing and existing edges. Smaller values make the existing links contribute more to the energy function which is minimized in the algorithm. Bigger values make the missing links more important. (If my understanding is correct.)

Returns:

Error code.

See also:

`igraph_community_spinglass()` for the traditional version of the algorithm.

Time complexity: TODO.

Community structure based on eigenvectors of matrices

The function documented in these section implements the “leading eigenvector” method developed by Mark Newman and published in MEJ Newman: Finding community structure using the eigenvectors of matrices, Phys Rev E 74:036104 (2006).

The heart of the method is the definition of the modularity matrix, B , which is $B=A-P$, A being the adjacency matrix of the (undirected) network, and P contains the probability that certain edges are present according to the “configuration model” In other words, a P_{ij} element of P is the probability that there is an edge between vertices i and j in a random network in which the degrees of all vertices are the same as in the input graph.

The leading eigenvector method works by calculating the eigenvector of the modularity matrix for the largest positive eigenvalue and then separating vertices into two community based on the sign of the corresponding element in the eigenvector. If all elements in the eigenvector are of the same sign that means that the network has no underlying community structure. Check Newman's paper to understand why this is a good method for detecting community structure.

The leading eigenvector community structure detection method is implemented in `igraph_community_leading_eigenvector()`. After the initial split, the following splits are done in a way to optimize modularity regarding to the original network. Note that any further refinement, for example using Kernighan-Lin, as proposed in Section V.A of Newman (2006), is not implemented here.

Example	24.2.	File	examples/simple/
<code>igraph_community_leading_eigenvector.c</code>			

igraph_community_leading_eigenvector — Leading eigenvector community finding (proper version).

```
int igraph_community_leading_eigenvector(const igraph_t *graph,
    const igraph_vector_t *weights,
    igraph_matrix_t *merges,
    igraph_vector_t *membership,
    igraph_integer_t steps,
    igraph_arnpack_options_t *options,
    igraph_real_t *modularity,
    igraph_bool_t start,
    igraph_vector_t *eigenvalues,
    igraph_vector_ptr_t *eigenvectors,
    igraph_vector_t *history,
    igraph_community_leading_eigenvector_callback_t *callback,
    void *callback_extra);
```

Newman's leading eigenvector method for detecting community structure. This is the proper implementation of the recursive, divisive algorithm: each split is done by maximizing the modularity regarding the original network, see MEJ Newman: Finding community structure in networks using the eigenvectors of matrices, Phys Rev E 74:036104 (2006).

Arguments:

<i>graph</i> :	The undirected input graph.
<i>weights</i> :	The weights of the edges, or a null pointer for unweighted graphs.
<i>merges</i> :	The result of the algorithm, a matrix containing the information about the splits performed. The matrix is built in the opposite way however, it is like the result of an agglomerative algorithm. If at the end of the algorithm (after <i>steps</i> steps was done) there are “p” communities, then these are numbered from zero to “p-1”. The first line of the matrix contains the first “merge” (which is in reality the last split) of two communities into community “p”, the merge in the second line forms community “p+1”, etc. The matrix should be initialized before calling and will be resized as needed. This argument is ignored if it is NULL.
<i>membership</i> :	The membership of the vertices after all the splits were performed will be stored here. The vector must be initialized before calling and will be resized as needed. This argument is ignored if it is NULL. This argument can also be used to supply a starting configuration for the community finding, in the format of a membership vector. In this case the <i>start</i> argument must be set to 1.
<i>steps</i> :	The maximum number of steps to perform. It might happen that some component (or the whole network) has no underlying community structure and no further steps can be done. If you want as many steps as possible then supply the number of vertices in the network here.
<i>options</i> :	The options for ARPACK. <i>n</i> is always overwritten. <i>ncv</i> is set to at least 4.
<i>modularity</i> :	If not a null pointer, then it must be a pointer to a real number and the modularity score of the final division is stored here.
<i>start</i> :	Boolean, whether to use the community structure given in the <i>membership</i> argument as a starting point.

<i>eigenvalues:</i>	Pointer to an initialized vector or a null pointer. If not a null pointer, then the eigenvalues calculated along the community structure detection are stored here. The non-positive eigenvalues, that do not result a split, are stored as well.	
<i>eigenvectors:</i>	If not a null pointer, then the eigenvectors that are calculated in each step of the algorithm, are stored here, in a pointer vector. Each eigenvector is stored in an <code>igraph_vector_t</code> object. The user is responsible of deallocating the memory that belongs to the individual vectors, by calling first <code>igraph_vector_destroy()</code> , and then <code>igraph_free()</code> on them.	
<i>history:</i>	Pointer to an initialized vector or a null pointer. If not a null pointer, then a trace of the algorithm is stored here, encoded numerically. The various operations:	
	<code>IGRAPH_LEVC_HIST_START_FULL</code>	Start the algorithm from an initial state where each connected component is a separate community.
	<code>IGRAPH_LEVC_HIST_START_GIVEN</code>	Start the algorithm from a given community structure. The next value in the vector contains the initial number of communities.
	<code>IGRAPH_LEVC_HIST_SPLIT</code>	Split a community into two communities. The id of the splitted community is given in the next element of the history vector. The id of the first new community is the same as the id of the splitted community. The id of the second community equals to the number of communities before the split.
	<code>IGRAPH_LEVC_HIST_FAILED</code>	Tried to split a community, but it was not worth it, as it does not result in a bigger modularity value. The id of the community is given in the next element of the vector.
<i>callback:</i>	A null pointer or a function of type <code>igraph_community_leading_eigenvector_callback_t</code> . If given, this callback function is called after each eigenvector/eigenvalue calculation. If the callback returns a non-zero value, then the community finding algorithm stops. See the arguments passed to the callback at the documentation of <code>igraph_community_leading_eigenvector_callback_t</code> .	
<i>callback_extra:</i>	Extra argument to pass to the callback function.	

Returns:

Error code.

See also:

`igraph_community_walktrap()` and `igraph_community_spinglass()` for other community structure detection methods.

Time complexity: $O(|E|+|V|^2*\text{steps})$, $|V|$ is the number of vertices, $|E|$ the number of edges, “steps” the number of splits performed.

igraph_community_leading_eigenvector_callback_t — Callback for the leading eigenvector community finding method.

```
typedef int igraph_community_leading_eigenvector_callback_t(  
    const igraph_vector_t *membership,  
    long int comm,  
    igraph_real_t eigenvalue,  
    const igraph_vector_t *eigenvector,  
    igraph_arnpack_function_t *arnpack_multiplier,  
    void *arnpack_extra,  
    void *extra);
```

The leading eigenvector community finding implementation in igraph is able to call a callback function, after each eigenvalue calculation. This callback function must be of `igraph_community_leading_eigenvector_callback_t` type. The following arguments are passed to the callback:

Arguments:

<i>membership:</i>	The actual membership vector, before recording the potential change implied by the newly found eigenvalue.
<i>comm:</i>	The id of the community that the algorithm tried to split in the last iteration. The community ids are indexed from zero here!
<i>eigenvalue:</i>	The eigenvalue the algorithm has just found.
<i>eigenvector:</i>	The eigenvector corresponding to the eigenvalue the algorithm just found.
<i>arnpack_multiplier:</i>	A function that was passed to <code>igraph_arnpack_rssolve()</code> to solve the last eigenproblem.
<i>arnpack_extra:</i>	The extra argument that was passed to the ARPACK solver.
<i>extra:</i>	Extra argument that as passed to <code>igraph_community_leading_eigenvector()</code> .

See also:

```
igraph_community_leading_eigenvector(),    igraph_arnpack_function_t,  
igraph_arnpack_rssolve().
```

igraph_le_community_to_membership — Vertex membership from the leading eigenvector community structure

```
int igraph_le_community_to_membership(const igraph_matrix_t *merges,  
                                     igraph_integer_t steps,
```



```
igraph_vector_t *membership,
igraph_vector_t *csize);
```

This function creates a membership vector from the result of `igraph_community_leading_eigenvector()`. It takes `membership` and performs `steps` merges, according to the supplied merges matrix.

Arguments:

merges: The two-column matrix containing the merge operations. See `igraph_community_walktrap()` for the detailed syntax. This is usually from the output of the leading eigenvector community structure detection routines.

steps: The number of steps to make according to merges.

membership: Initially the starting membership vector, on output the resulting membership vector, after performing `steps` merges.

csize: Optionally the sizes of the communities is stored here, if this is not a null pointer, but an initialized vector.

Returns:

Error code.

Time complexity: $O(|V|)$, the number of vertices.

Walktrap: Community structure based on random walks

`igraph_community_walktrap` — This function is the implementation of the Walktrap community

```
int igraph_community_walktrap(const igraph_t *graph,
                             const igraph_vector_t *weights,
                             int steps,
                             igraph_matrix_t *merges,
                             igraph_vector_t *modularity,
                             igraph_vector_t *membership);
```

finding algorithm, see Pascal Pons, Matthieu Latapy: Computing communities in large networks using random walks, <https://arxiv.org/abs/physics/0512106>

Currently the original C++ implementation is used in igraph, see <https://www-complexnetworks.lip6.fr/~latapy/PP/walktrap.html> We are grateful to Matthieu Latapy and Pascal Pons for providing this source code.

In contrast to the original implementation, isolated vertices are allowed in the graph and they are assumed to have a single incident loop edge with weight 1.

Arguments:

graph: The input graph, edge directions are ignored.

weights: Numeric vector giving the weights of the edges. If it is a NULL pointer then all edges will have equal weights. The weights are expected to be positive.

steps: Integer constant, the length of the random walks.

merges: Pointer to a matrix, the merges performed by the algorithm will be stored here (if not NULL). Each merge is a row in a two-column matrix and contains the ids of the merged clusters. Clusters are numbered from zero and cluster numbers smaller than the number of nodes in the network belong to the individual vertices as singleton clusters. In each step a new cluster is created from two other clusters and its id will be one larger than the largest cluster id so far. This means that before the first merge we have n clusters (the number of vertices in the graph) numbered from zero to $n-1$. The first merge creates cluster n , the second cluster $n+1$, etc.

modularity: Pointer to a vector. If not NULL then the modularity score of the current clustering is stored here after each merge operation.

membership: Pointer to a vector. If not a NULL pointer, then the membership vector corresponding to the maximal modularity score is stored here. If it is not a NULL pointer, then neither *modularity* nor *merges* may be NULL.

Returns:

Error code.

See also:

`igraph_community_spinglass()`, `igraph_community_edge_betweenness()`.

Time complexity: $O(|E||V|^2)$ in the worst case, $O(|V|^2 \log|V|)$ typically, $|V|$ is the number of vertices, $|E|$ is the number of edges.

Example 24.3. File `examples/simple/walktrap.c`

Edge betweenness based community detection

`igraph_community_edge_betweenness` — Community finding based on edge betweenness.

```
int igraph_community_edge_betweenness(const igraph_t *graph,
                                     igraph_vector_t *result,
                                     igraph_vector_t *edge_betweenness,
                                     igraph_matrix_t *merges,
                                     igraph_vector_t *bridges,
                                     igraph_vector_t *modularity,
                                     igraph_vector_t *membership,
                                     igraph_bool_t directed,
                                     const igraph_vector_t *weights);
```

Community structure detection based on the betweenness of the edges in the network. The algorithm was invented by M. Girvan and M. Newman, see: M. Girvan and M. E. J. Newman: Community structure in social and biological networks, Proc. Nat. Acad. Sci. USA 99, 7821-7826 (2002).

The idea is that the betweenness of the edges connecting two communities is typically high, as many of the shortest paths between nodes in separate communities go through them. So we gradually remove the edge with highest betweenness from the network, and recalculate edge betweenness after every removal. This way sooner or later the network falls off to two components, then after a while one of these components falls off to two smaller components, etc. until all edges are removed. This is a divisive hierarchical approach, the result is a dendrogram.

Arguments:

<i>graph</i> :	The input graph.
<i>result</i> :	Pointer to an initialized vector, the result will be stored here, the ids of the removed edges in the order of their removal. It will be resized as needed. It may be NULL if the edge IDs are not needed by the caller.
<i>edge_betweenness</i> :	Pointer to an initialized vector or NULL. In the former case the edge betweenness of the removed edge is stored here. The vector will be resized as needed.
<i>merges</i> :	Pointer to an initialized matrix or NULL. If not NULL then merges performed by the algorithm are stored here. Even if this is a divisive algorithm, we can replay it backwards and note which two clusters were merged. Clusters are numbered from zero, see the <i>merges</i> argument of <code>igraph_community_walktrap()</code> for details. The matrix will be resized as needed.
<i>bridges</i> :	Pointer to an initialized vector of NULL. If not NULL then all edge removals which separated the network into more components are marked here.
<i>modularity</i> :	If not a null pointer, then the modularity values of the different divisions are stored here, in the order corresponding to the merge matrix. The modularity values will take weights into account if <i>weights</i> is not null.
<i>membership</i> :	If not a null pointer, then the membership vector, corresponding to the highest modularity value, is stored here.
<i>directed</i> :	Logical constant, whether to calculate directed betweenness (i.e. directed paths) for directed graphs. It is ignored for undirected graphs.
<i>weights</i> :	An optional vector containing edge weights. If null, the unweighted edge betweenness scores will be calculated and used. If not null, the weighted edge betweenness scores will be calculated and used.

Returns:

Error code.

See also:

`igraph_community_eb_get_merges()`, `igraph_community_spinglass()`,
`igraph_community_walktrap()`.

Time complexity: $O(|V||E|^2)$, as the betweenness calculation requires $O(|V||E|)$ and we do it $|E|-1$ times.

Example **24.4.** **File** **examples/simple/**
igraph_community_edge_betweenness.c

igraph_community_eb_get_merges — Calculating the merges, i.e. the dendrogram for an edge betweenness community structure.

```
int igraph_community_eb_get_merges(const igraph_t *graph,
                                   const igraph_bool_t directed,
                                   const igraph_vector_t *edges,
                                   const igraph_vector_t *weights,
                                   igraph_matrix_t *res,
                                   igraph_vector_t *bridges,
                                   igraph_vector_t *modularity,
                                   igraph_vector_t *membership);
```

This function is handy if you have a sequence of edges which are gradually removed from the network and you would like to know how the network falls apart into separate components. The edge sequence may come from the `igraph_community_edge_betweenness()` function, but this is not necessary. Note that `igraph_community_edge_betweenness()` can also calculate the dendrogram, via its *merges* argument.

Arguments:

- | | |
|---------------------|---|
| <i>graph</i> : | The input graph. |
| <i>edges</i> : | Vector containing the edges to be removed from the network, all edges are expected to appear exactly once in the vector. |
| <i>directed</i> : | Whether to use the directed or undirected version of modularity. Will be ignored for undirected graphs. |
| <i>weights</i> : | An optional vector containing edge weights. If null, the unweighted modularity scores will be calculated. If not null, the weighted modularity scores will be calculated. Ignored if both <i>modularity</i> and <i>membership</i> are NULL pointers. |
| <i>res</i> : | Pointer to an initialized matrix, if not NULL then the dendrogram will be stored here, in the same form as for the <code>igraph_community_walktrap()</code> function: the matrix has two columns and each line is a merge given by the ids of the merged components. The component ids are numbered from zero and component ids smaller than the number of vertices in the graph belong to individual vertices. The non-trivial components containing at least two vertices are numbered from <i>n</i> , where <i>n</i> is the number of vertices in the graph. So if the first line contains <i>a</i> and <i>b</i> that means that components <i>a</i> and <i>b</i> are merged into component <i>n</i> , the second line creates component <i>n</i> +1, etc. The matrix will be resized as needed. |
| <i>bridges</i> : | Pointer to an initialized vector or NULL. If not null then the index of the edge removals which split the network will be stored here. The vector will be resized as needed. |
| <i>modularity</i> : | If not a null pointer, then the modularity values for the different divisions, corresponding to the merges matrix, will be stored here. |
| <i>membership</i> : | If not a null pointer, then the membership vector for the best division (in terms of modularity) will be stored here. |

Returns:

Error code.

See also:

`igraph_community_edge_betweenness()`.

Time complexity: $O(|E|+|V|\log|V|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges.

Community structure based on the optimization of modularity

`igraph_community_fastgreedy` — Finding community structure by greedy optimization of modularity.

```
int igraph_community_fastgreedy(const igraph_t *graph,
                               const igraph_vector_t *weights,
                               igraph_matrix_t *merges,
                               igraph_vector_t *modularity,
                               igraph_vector_t *membership);
```

This function implements the fast greedy modularity optimization algorithm for finding community structure, see A Clauset, MEJ Newman, C Moore: Finding community structure in very large networks, <http://www.arxiv.org/abs/cond-mat/0408187> for the details.

Some improvements proposed in K Wakita, T Tsurumi: Finding community structure in mega-scale social networks, <http://www.arxiv.org/abs/cs.CY/0702048v1> have also been implemented.

Arguments:

- | | |
|---------------------|--|
| <i>graph</i> : | The input graph. It must be a graph without multiple edges. This is checked and an error message is given for graphs with multiple edges. |
| <i>weights</i> : | Potentially a numeric vector containing edge weights. Supply a null pointer here for unweighted graphs. The weights are expected to be non-negative. |
| <i>merges</i> : | Pointer to an initialized matrix or NULL, the result of the computation is stored here. The matrix has two columns and each merge corresponds to one merge, the ids of the two merged components are stored. The component ids are numbered from zero and the first n components are the individual vertices, n is the number of vertices in the graph. Component n is created in the first merge, component $n+1$ in the second merge, etc. The matrix will be resized as needed. If this argument is NULL then it is ignored completely. |
| <i>modularity</i> : | Pointer to an initialized vector or NULL pointer, in the former case the modularity scores along the stages of the computation are recorded here. The vector will be resized as needed. |
| <i>membership</i> : | Pointer to a vector. If not a null pointer, then the membership vector corresponding to the best split (in terms of modularity) is stored here. |

Returns:

Error code.

See also:

`igraph_community_walktrap()`, `igraph_community_edge_betweenness()` for other community detection algorithms, `igraph_community_to_membership()` to convert the dendrogram to a membership vector.

Time complexity: $O(|E||V|\log|V|)$ in the worst case, $O(|E|+|V|\log^2|V|)$ typically, $|V|$ is the number of vertices, $|E|$ is the number of edges.

Example	24.5.	File	examples/simple/
<code>igraph_community_fastgreedy.c</code>			

`igraph_community_multilevel` — Finding community structure by multi-level optimization of modularity.

```
int igraph_community_multilevel(const igraph_t *graph,
                              const igraph_vector_t *weights,
                              const igraph_real_t resolution,
                              igraph_vector_t *membership,
                              igraph_matrix_t *memberships, igraph_vector_t *
```

This function implements the multi-level modularity optimization algorithm for finding community structure, see Blondel, V. D., Guillaume, J.-L., Lambiotte, R., & Lefebvre, E. (2008). Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment*, 10008(10), 6. <https://doi.org/10.1088/1742-5468/2008/10/P10008> for the details (preprint: <http://arxiv.org/abs/0803.0476>). The algorithm is sometimes known as the "Louvain" algorithm.

The algorithm is based on the modularity measure and a hierarchical approach. Initially, each vertex is assigned to a community on its own. In every step, vertices are re-assigned to communities in a local, greedy way: in a random order, each vertex is moved to the community with which it achieves the highest contribution to modularity. When no vertices can be reassigned, each community is considered a vertex on its own, and the process starts again with the merged communities. The process stops when there is only a single vertex left or when the modularity cannot be increased any more in a step.

The resolution parameter `gamma` allows finding communities at different resolutions. Higher values of the resolution parameter typically result in more, smaller communities. Lower values typically result in fewer, larger communities. The original definition of modularity is retrieved when setting `gamma=1`. Note that the returned modularity value is calculated using the indicated resolution parameter. See `igraph_modularity()` for more details. This function was contributed by Tom Gregorovic.

Arguments:

<i>graph</i> :	The input graph. It must be an undirected graph.
<i>weights</i> :	Numeric vector containing edge weights. If <code>NULL</code> , every edge has equal weight. The weights are expected to be non-negative.
<i>resolution</i> :	Resolution parameter. Must be greater than or equal to 0. Lower values favor fewer, larger communities; higher values favor more, smaller communities. Set it to 1 to use the classical definition of modularity.
<i>membership</i> :	The membership vector, the result is returned here. For each vertex it gives the ID of its community. The vector must be initialized and it will be resized accordingly.
<i>memberships</i> :	Numeric matrix that will contain the membership vector after each level, if not <code>NULL</code> . It must be initialized and it will be resized accordingly.

modularity: Numeric vector that will contain the modularity score after each level, if not NULL. It must be initialized and it will be resized accordingly.

Returns:

Error code.

Time complexity: in average near linear on sparse graphs.

Example **24.6.** **File** **examples/simple/igraph_community_multilevel.c**

igraph_community_leiden — Finding community structure using the Leiden algorithm.

```
int igraph_community_leiden(const igraph_t *graph,
                           const igraph_vector_t *edge_weights, const igraph_v
                           const igraph_real_t resolution_parameter, const igr
                           igraph_vector_t *membership, igraph_integer_t *nb_c
```

This function implements the Leiden algorithm for finding community structure, see Traag, V. A., Waltman, L., & van Eck, N. J. (2019). From Louvain to Leiden: guaranteeing well-connected communities. *Scientific reports*, 9(1), 5233. <http://dx.doi.org/10.1038/s41598-019-41695-z>

It is similar to the multilevel algorithm, often called the Louvain algorithm, but it is faster and yields higher quality solutions. It can optimize both modularity and the Constant Potts Model, which does not suffer from the resolution-limit (see preprint <http://arxiv.org/abs/1104.3083>).

The Leiden algorithm consists of three phases: (1) local moving of nodes, (2) refinement of the partition and (3) aggregation of the network based on the refined partition, using the non-refined partition to create an initial partition for the aggregate network. In the local move procedure in the Leiden algorithm, only nodes whose neighborhood has changed are visited. The refinement is done by restarting from a singleton partition within each cluster and gradually merging the subclusters. When aggregating, a single cluster may then be represented by several nodes (which are the subclusters identified in the refinement).

The Leiden algorithm provides several guarantees. The Leiden algorithm is typically iterated: the output of one iteration is used as the input for the next iteration. At each iteration all clusters are guaranteed to be connected and well-separated. After an iteration in which nothing has changed, all nodes and some parts are guaranteed to be locally optimally assigned. Finally, asymptotically, all subsets of all clusters are guaranteed to be locally optimally assigned. For more details, please see Traag, Waltman & van Eck (2019).

The objective function being optimized is

$$1 / 2m \sum_{ij} (A_{ij} - \gamma n_i n_j) d(s_i, s_j)$$

where m is the total edge weight, A_{ij} is the weight of edge (i, j) , γ is the so-called resolution parameter, n_i is the node weight of node i , s_i is the cluster of node i and $d(x, y) = 1$ if and only if $x = y$ and 0 otherwise. By setting $n_i = k_i$, the degree of node i , and dividing γ by $2m$, you effectively obtain an expression for modularity. Hence, the standard modularity will be optimized when you supply the degrees as `node_weights` and by supplying as a resolution parameter $1.0 / (2 * m)$, with m the number of edges.

Arguments:

<i>graph:</i>	The input graph. It must be an undirected graph.
<i>edge_weights:</i>	Numeric vector containing edge weights. If NULL, every edge has equal weight of 1. The weights need not be non-negative.
<i>node_weights:</i>	Numeric vector containing node weights.
<i>resolution_parameter:</i>	The resolution parameter used, which is represented by gamma in the objective function mentioned in the documentation.
<i>beta:</i>	The randomness used in the refinement step when merging. A small amount of randomness (beta = 0.01) typically works well.
<i>start:</i>	Start from membership vector. If this is true, the optimization will start from the provided membership vector. If this is false, the optimization will start from a singleton partition.
<i>membership:</i>	The membership vector. This is both used as the initial membership from which optimisation starts and is updated in place. It must hence be properly initialized. When finding clusters from scratch it is typically started using a singleton clustering. This can be achieved using <code>igraph_vector_init_seq</code> .
<i>nb_clusters:</i>	The number of clusters contained in membership. Must not be a NULL pointer.
<i>quality:</i>	The quality of the partition, in terms of the objective function as included in the documentation. If NULL the quality will not be calculated.

Returns:

Error code.

Time complexity: near linear on sparse graphs.

Example 24.7. File `examples/simple/igraph_community_leiden.c`

Fluid communities

`igraph_community_fluid_communities` — Community detection based on fluids interacting on the graph.

```
int igraph_community_fluid_communities(const igraph_t *graph,
                                     igraph_integer_t no_of_communities,
                                     igraph_vector_t *membership,
                                     igraph_real_t *modularity);
```

The algorithm is based on the simple idea of several fluids interacting in a non-homogeneous environment (the graph topology), expanding and contracting based on their interaction and density. This function implements the community detection method described in: Parés F, Gasulla DG, et. al. (2018) Fluid Communities: A Competitive, Scalable and Diverse Community Detection Algorithm. In: Complex Networks & Their Applications VI: Proceedings of Complex Networks 2017 (The Sixth International Conference on Complex Networks and Their Applications), Springer, vol 689, p 229.

Arguments:

<i>graph</i> :	The input graph. The graph must be simple and connected. Empty graphs are not supported as well as single vertex graphs. Edge directions are ignored. Weights are not considered.
<i>no_of_communities</i> :	The number of communities to be found. Must be greater than 0 and fewer than number of vertices in the graph.
<i>membership</i> :	The result vector mapping vertices to the communities they are assigned to.
<i>modularity</i> :	If not a null pointer, then it must be a pointer to a real number. The modularity score of the detected community structure is stored here.

Returns:

Error code.

Time complexity: $O(|E|)$

Example **24.8.** **File** **examples/simple/igraph_community_fluid_communities.c**

Label propagation

igraph_community_label_propagation — Community detection based on label propagation.

```
int igraph_community_label_propagation(const igraph_t *graph,
                                      igraph_vector_t *membership,
                                      const igraph_vector_t *weights,
                                      const igraph_vector_t *initial,
                                      igraph_vector_bool_t *fixed,
                                      igraph_real_t *modularity);
```

This function implements the community detection method described in: Raghavan, U.N. and Albert, R. and Kumara, S.: Near linear time algorithm to detect community structures in large-scale networks. Phys Rev E 76, 036106. (2007). This version extends the original method by the ability to take edge weights into consideration and also by allowing some labels to be fixed.

Weights are taken into account as follows: when the new label of node *i* is determined, the algorithm iterates over all edges incident on node *i* and calculate the total weight of edges leading to other nodes with label 0, 1, 2, ..., *k* - 1 (where *k* is the number of possible labels). The new label of node *i* will then be the label whose edges (among the ones incident on node *i*) have the highest total weight.

Arguments:

<i>graph</i> :	The input graph, should be undirected to make sense.
<i>membership</i> :	The membership vector, the result is returned here. For each vertex it gives the ID of its community (label).
<i>weights</i> :	The weight vector, it should contain a positive weight for all the edges.

initial: The initial state. If NULL, every vertex will have a different label at the beginning. Otherwise it must be a vector with an entry for each vertex. Non-negative values denote different labels, negative entries denote vertices without labels.

fixed: Boolean vector denoting which labels are fixed. Of course this makes sense only if you provided an initial state, otherwise this element will be ignored. Also note that vertices without labels cannot be fixed. If they are, this vector will be modified to make it consistent with *initial*.

modularity: If not a null pointer, then it must be a pointer to a real number. The modularity score of the detected community structure is stored here.

Returns:

Error code.

Time complexity: $O(m+n)$

Example **24.9.** **File** **examples/simple/**
igraph_community_label_propagation.c

The InfoMAP algorithm

igraph_community_infomap — Find community structure that minimizes the expected

```
int igraph_community_infomap(const igraph_t * graph,
                             const igraph_vector_t *e_weights,
                             const igraph_vector_t *v_weights,
                             int nb_trials,
                             igraph_vector_t *membership,
                             igraph_real_t *codelength);
```

description length of a random walker trajectory. Implementation of the InfoMap community detection algorithm. of Martin Rosvall and Carl T. Bergstrom. See : Visualization of the math and the map generator: www.mapequation.org [2] The original paper: M. Rosvall and C. T. Bergstrom, Maps of information flow reveal community structure in complex networks, PNAS 105, 1118 (2008) [<http://dx.doi.org/10.1073/pnas.0706851105> , <http://arxiv.org/abs/0707.0609>] [3] A more detailed paper: M. Rosvall, D. Axelsson, and C. T. Bergstrom, The map equation, Eur. Phys. J. Special Topics 178, 13 (2009). [<http://dx.doi.org/10.1140/epjst/e2010-01179-1> , <http://arxiv.org/abs/0906.1405>]

The original C++ implementation of Martin Rosvall is used, see http://www.tp.umu.se/~rosvall/downloads/infomap_undir.tgz . Intergration in igraph has be done by Emmanuel Navarro (who is grateful to * Martin Rosvall and Carl T. Bergstrom for providing this source code.)

Note that the graph must not contain isolated vertices.

If you want to specify a random seed (as in original implementation) you can use `igraph_rng_seed()`.

Arguments:

graph: The input graph.

e_weights: Numeric vector giving the weights of the edges. If it is a NULL pointer then all edges will have equal weights. The weights are expected to be positive.

v_weights: Numeric vector giving the weights of the vertices. If it is a NULL pointer then all vertices will have equal weights. The weights are expected to be positive.

nb_trials: The number of attempts to partition the network (can be any integer value equal or larger than 1).

membership: Pointer to a vector. The membership vector is stored here.

code_length: Pointer to a real. If not NULL the code length of the partition is stored here.

Returns:

Error code.

See also:

`igraph_community_spinglass()`, `igraph_community_edge_betweenness()`,
`igraph_community_walktrap()`.

Time complexity: TODO.

Chapter 25. Graphlets

Introduction

Graphlet decomposition models a weighted undirected graph via the union of potentially overlapping dense social groups. This is done by a two-step algorithm. In the first step, a candidate set of groups (a candidate basis) is created by finding cliques in the thresholded input graph. In the second step, the graph is projected onto the candidate basis, resulting in a weight coefficient for each clique in the candidate basis.

For more information on graphlet decomposition, see Hossein Azari Soufiani and Edoardo M Airolti: "Graphlet decomposition of a weighted network", <https://arxiv.org/abs/1203.2821> and <http://proceedings.mlr.press/v22/azari12/azari12.pdf>

`igraph` contains three functions for performing the graphlet decomposition of a graph. The first is `igraph_graphlets()`, which performs both steps of the method and returns a list of subgraphs with their corresponding weights. The other two functions correspond to the first and second steps of the algorithm, and they are useful if the user wishes to perform them individually: `igraph_graphlets_candidate_basis()` and `igraph_graphlets_project()`.

Note: The term "graphlet" is used for several unrelated concepts in the literature. If you are looking to count induced subgraphs, see `igraph_motifs_randesu()` and `igraph_subisomorphic_lad()`.

Performing graphlet decomposition

`igraph_graphlets` — Calculate graphlets basis and project the graph on it

```
int igraph_graphlets(const igraph_t *graph,
                    const igraph_vector_t *weights,
                    igraph_vector_ptr_t *cliques,
                    igraph_vector_t *Mu, int niter);
```

This function simply calls `igraph_graphlets_candidate_basis()` and `igraph_graphlets_project()`, and then orders the graphlets according to decreasing weights.

Arguments:

- graph*: The input graph, it must be a simple graph, edge directions are ignored.
- weights*: Weights of the edges, a vector.
- cliques*: An initialized vector of pointers. The graphlet basis is stored here. Each element of the pointer vector will be a vector of vertex ids.
- Mu*: An initialized vector, the weights of the graphlets will be stored here.
- niter*: Integer scalar, the number of iterations to perform for the projection step.

Returns:

Error code.

See also: `igraph_graphlets_candidate_basis()` and `igraph_graphlets_project()`.

igraph_graphlets_candidate_basis — Calculate a candidate graphlets basis

```
int igraph_graphlets_candidate_basis(const igraph_t *graph,
                                     const igraph_vector_t *weights,
                                     igraph_vector_ptr_t *cliques,
                                     igraph_vector_t *thresholds);
```

Arguments:

- graph*: The input graph, it must be a simple graph, edge directions are ignored.
- weights*: Weights of the edges, a vector.
- cliques*: An initialized vector of pointers. The graphlet basis is stored here. Each element of the pointer vector will be a vector of vertex ids. Each elements must be destroyed using `igraph_vector_destroy()` and `igraph_free()`.
- thresholds*: An initialized vector, the (highest possible) weight thresholds for finding the basis subgraphs are stored here.

Returns:

Error code.

See also: `igraph_graphlets()` and `igraph_graphlets_project()`.

igraph_graphlets_project — Project a graph on a graphlets basis

```
int igraph_graphlets_project(const igraph_t *graph,
                             const igraph_vector_t *weights,
                             const igraph_vector_ptr_t *cliques,
                             igraph_vector_t *Mu, igraph_bool_t startMu,
                             int niter);
```

Note that the graph projected does not have to be the same that was used to calculate the graphlet basis, but it is assumed that it has the same number of vertices, and the vertex ids of the two graphs match.

Arguments:

- graph*: The input graph, it must be a simple graph, edge directions are ignored.
- weights*: Weights of the edges in the input graph, a vector.
- cliques*: The graphlet basis, a pointer vector, in which each element is a vector of vertex ids.
- Mu*: An initialized vector, the weights of the graphlets will be stored here. This vector is also used to initialize the the weight vector for the iterative algorithm, if the `startMu` argument is true (non-zero).

startMu: If true (non-zero), then the supplied Mu vector is used as the starting point of the iteration. Otherwise a constant 1 vector is used.

niter: Integer scalar, the number of iterations to perform.

Returns:

Error code.

See also: `igraph_graphlets()` and `igraph_graphlets_candidate_basis()`.

Chapter 26. Hierarchical random graphs

Introduction

A hierarchical random graph is an ensemble of undirected graphs with n vertices. It is defined via a binary tree with n leaf and $n-1$ internal vertices, where the internal vertices are labeled with probabilities. The probability that two vertices are connected in the random graph is given by the probability label at their closest common ancestor.

Please read the following two articles for more about hierarchical random graphs: A. Clauset, C. Moore, and M.E.J. Newman. Hierarchical structure and the prediction of missing links in networks. *Nature* 453, 98 - 101 (2008); and A. Clauset, C. Moore, and M.E.J. Newman. Structural Inference of Hierarchies in Networks. In E. M. Airolidi et al. (Eds.): *ICML 2006 Ws, Lecture Notes in Computer Science* 4503, 1-13. Springer-Verlag, Berlin Heidelberg (2007).

igraph contains functions for fitting HRG models to a given network (`igraph_hrg_fit`), for generating networks from a given HRG ensemble (`igraph_hrg_game`, `igraph_hrg_sample`), converting an igraph graph to a HRG and back (`igraph_hrg_create`, `igraph_hrg_dendrogram`), for calculating a consensus tree from a set of sampled HRGs (`igraph_hrg_consensus`) and for predicting missing edges in a network based on its HRG models (`igraph_hrg_predict`).

The igraph HRG implementation is heavily based on the code published by Aaron Clauset, at his website, <http://tuvalu.santafe.edu/~aaronc/hierarchy/>

Representing HRGs

`igraph_hrg_t` — Data structure to store a hierarchical random graph

```
typedef struct igraph_hrg_t {
    igraph_vector_t left, right, prob, edges, vertices;
} igraph_hrg_t;
```

A hierarchical random graph (HRG) can be given as a binary tree, where the internal vertices are labeled with real numbers.

Note that you don't necessarily have to know this internal representation for using the HRG functions, just pass the HRG objects created by one igraph function, to another igraph function.

It has the following members:

Values:

- | | |
|----------------------|---|
| <code>left</code> : | Vector that contains the left children of the internal tree vertices. The first vertex is always the root vertex, so the first element of the vector is the left child of the root vertex. Internal vertices are denoted with negative numbers, starting from -1 and going down, i.e. the root vertex is -1. Leaf vertices are denoted by non-negative number, starting from zero and up. |
| <code>right</code> : | Vector that contains the right children of the vertices, with the same encoding as the <code>left</code> vector. |

- prob:** The connection probabilities attached to the internal vertices, the first number belongs to the root vertex (i.e. internal vertex -1), the second to internal vertex -2, etc.
- edges:** The number of edges in the subtree below the given internal vertex.
- vertices:** The number of vertices in the subtree below the given internal vertex, including itself.

igraph_hrg_init — Allocate memory for a HRG.

```
int igraph_hrg_init(igraph_hrg_t *hrg, int n);
```

This function must be called before passing an `igraph_hrg_t` to an `igraph` function.

Arguments:

- hrg:** Pointer to the HRG data structure to initialize.
- n:** The number of vertices in the graph that is modeled by this HRG. It can be zero, if this is not yet known.

Returns:

Error code.

Time complexity: $O(n)$, the number of vertices in the graph.

igraph_hrg_destroy — Deallocate memory for an HRG.

```
void igraph_hrg_destroy(igraph_hrg_t *hrg);
```

The HRG data structure can be reinitialized again with an `igraph_hrg_destroy` call.

Arguments:

- hrg:** Pointer to the HRG data structure to deallocate.

Time complexity: operating system dependent.

igraph_hrg_size — Returns the size of the HRG, the number of leaf nodes.

```
int igraph_hrg_size(const igraph_hrg_t *hrg);
```

Arguments:

- hrg:** Pointer to the HRG.

Returns:

The number of leaf nodes in the HRG.

Time complexity: $O(1)$.

igraph_hrg_resize — Resize a HRG.

```
int igraph_hrg_resize(igraph_hrg_t *hrg, int newsize);
```

Arguments:

hrg: Pointer to an initialized (see `igraph_hrg_init`) HRG.

newsize: The new size, i.e. the number of leaf nodes.

Returns:

Error code.

Time complexity: $O(n)$, n is the new size.

Fitting HRGs

igraph_hrg_fit — Fit a hierarchical random graph model to a network

```
int igraph_hrg_fit(const igraph_t *graph,
                  igraph_hrg_t *hrg,
                  igraph_bool_t start,
                  int steps);
```

Arguments:

graph: The igraph graph to fit the model to. Edge directions are ignored in directed graphs.

hrg: Pointer to an initialized HRG, the result of the fitting is stored here. It can also be used to pass a HRG to the function, that can be used as the starting point of the Markov Chain Monte Carlo fitting, if the *start* argument is true.

start: Logical, whether to start the fitting from the given HRG.

steps: Integer, the number of MCMC steps to take in the fitting procedure. If this is zero, then the fitting stop is a convergence criteria is fulfilled.

Returns:

Error code.

Time complexity: TODO.

igraph_hrg_consensus — Calculate a consensus tree for a HRG.

```
int igraph_hrg_consensus(const igraph_t *graph,
                        igraph_vector_t *parents,
                        igraph_vector_t *weights,
                        igraph_hrg_t *hrg,
                        igraph_bool_t start,
                        int num_samples);
```

The calculation can be started from the given HRG (*hrg*), or (if *start* is false), a HRG is first fitted to the given graph.

Arguments:

graph: The input graph.

parents: An initialized vector, the results are stored here. For each vertex, the id of its parent vertex is stored, or -1, if the vertex is the root vertex in the tree. The first *n* vertex ids (from 0) refer to the original vertices of the graph, the other ids refer to vertex groups.

weights: Numeric vector, counts the number of times a given tree split occurred in the generated network samples, for each internal vertices. The order is the same as in *parents*.

hrg: A hierarchical random graph. It is used as a starting point for the sampling, if the *start* argument is true. It is modified along the MCMC.

start: Logical, whether to use the supplied HRG (in *hrg*) as a starting point for the MCMC.

num_samples: The number of samples to generate for creating the consensus tree.

Returns:

Error code.

Time complexity: TODO.

HRG sampling

igraph_hrg_sample — Sample from a hierarchical random graph model

```
int igraph_hrg_sample(const igraph_t *input_graph,
                    igraph_t *sample,
                    igraph_vector_ptr_t *samples,
                    igraph_integer_t no_samples,
                    igraph_hrg_t *hrg,
                    igraph_bool_t start);
```

Sample from a hierarchical random graph ensemble. The ensemble can be given as a graph (*input_graph*), or as a HRG object (*hrg*). If a graph is given, then first an MCMC optimization is performed to find the optimal fitting model; then the MCMC is used to sample the graph(s).

Arguments:

input_graph: An igraph graph, or a null pointer. If not a null pointer, then a HRG is first fitted to the graph, possibly starting from the given HRG, if the `start` argument is true. If it is a null pointer, then the given HRG is used as a starting point, to find the optimum of the Markov chain, before the sampling.

sample: Pointer to an uninitialized graph, or a null pointer. If only one sample is requested, and it is not a null pointer, then the sample is stored here.

samples: An initialized vector of pointers. If more than one samples are requested, then they are stored here. Note that to free this data structure, you need to call `igraph_destroy()` on each graph first, then `igraph_free()` on all pointers, and finally `igraph_vector_ptr_destroy`.

no_samples: The number of samples to generate.

hrg: A HRG. It is modified during the sampling.

start: Logical, whether to start the MCMC from the given HRG.

Returns:

Error code.

Time complexity: TODO.

igraph_hrg_game — Generate a hierarchical random graph

```
int igraph_hrg_game(igraph_t *graph,
                    const igraph_hrg_t *hrg);
```

This function is a simple shortcut to `igraph_hrg_sample`. It creates a single graph, from the given HRG.

Arguments:

graph: Pointer to an uninitialized graph, the new graph is created here.

hrg: The hierarchical random graph model to sample from. It is modified during the MCMC process.

Returns:

Error code.

Time complexity: TODO.

Conversion to and from igraph graphs

igraph_hrg_dendrogram — Create a dendrogram from a hierarchical random graph.

```
int igraph_hrg_dendrogram(igraph_t *graph,
                          const igraph_hrg_t *hrg);
```

Creates the igraph graph equivalent of an `igraph_hrg_t` data structure.

Arguments:

graph: Pointer to an uninitialized graph, the result is stored here.

hrg: The hierarchical random graph to convert.

Returns:

Error code.

Time complexity: $O(n)$, the number of vertices in the graph.

igraph_hrg_create — Create a HRG from an igraph graph.

```
int igraph_hrg_create(igraph_hrg_t *hrg,
                      const igraph_t *graph,
                      const igraph_vector_t *prob);
```

Arguments:

hrg: Pointer to an initialized `igraph_hrg_t`. The result is stored here.

graph: The igraph graph to convert. It must be a directed binary tree, with $n-1$ internal and n leaf vertices. The root vertex must have in-degree zero.

prob: The vector of probabilities, this is used to label the internal nodes of the hierarchical random graph. The values corresponding to the leaves are ignored.

Returns:

Error code.

Time complexity: $O(n)$, the number of vertices in the tree.

Predicting missing edges

igraph_hrg_predict — Predict missing edges in a graph, based on HRG models

```
int igraph_hrg_predict(const igraph_t *graph,
                      igraph_vector_t *edges,
                      igraph_vector_t *prob,
                      igraph_hrg_t *hrg,
                      igraph_bool_t start,
                      int num_samples,
```

```
int num_bins);
```

Samples HRG models for a network, and estimated the probability that an edge was falsely observed as non-existent in the network.

Arguments:

<i>graph</i> :	The input graph.
<i>edges</i> :	The list of missing edges is stored here, the first two elements are the first edge, the next two the second edge, etc.
<i>prob</i> :	Vector of probabilities for the existence of missing edges, in the order corresponding to <i>edges</i> .
<i>hrg</i> :	A HRG, it is used as a starting point if <i>start</i> is true. It is also modified during the MCMC sampling.
<i>start</i> :	Logical, whether to start the MCMC from the given HRG.
<i>num_samples</i> :	The number of samples to generate.
<i>num_bins</i> :	Controls the resolution of the edge probabilities. Higher numbers result higher resolution.

Returns:

Error code.

Time complexity: TODO.

Chapter 27. Spectral coarse graining

Introduction

The SCG functions provide a framework, called Spectral Coarse Graining (SCG), for reducing large graphs while preserving their *spectral-related features*, that is features closely related with the eigenvalues and eigenvectors of a graph matrix (which for now can be the adjacency, the stochastic, or the Laplacian matrix).

Common examples of such features comprise the first-passage-time of random walkers on Markovian graphs, thermodynamic properties of lattice models in statistical physics (e.g. Ising model), and the epidemic threshold of epidemic network models (SIR and SIS models).

SCG differs from traditional clustering schemes by producing a *coarse-grained graph* (not just a partition of the vertices), representative of the original one. As shown in [1], Principal Component Analysis can be viewed as a particular SCG, called *exact SCG*, where the matrix to be coarse-grained is the covariance matrix of some data set.

SCG should be of interest to practitioners of various fields dealing with problems where matrix eigenpairs play an important role, as for instance is the case of dynamical processes on networks.

SCG in brief

The main idea of SCG is to operate on a matrix a shrinkage operation specifically designed to preserve some of the matrix eigenpairs while not altering other important matrix features (such as its structure). Mathematically, this idea was expressed as follows. Consider a (complex) $n \times n$ matrix M and form the product

$$M' = LMR^*,$$

where $n' < n$ and L, R are from $C[n' \times n]$ and are such that $LR^* = I[n']$ (R^* denotes the conjugate transpose of R). Under these assumptions, it can be shown that $P = R^*L$ is an n' -rank projector and that, if (λ, v) is a (right) eigenpair of M (i.e. $Mv = \lambda v$) and P is orthogonal, there exists an eigenvalue λ' of M' such that

$$|\lambda - \lambda'| \leq \text{const } \|e[P](v)\| [1 + O(\|e[P](v)\|^2)],$$

where $\|e[P](v)\| = \|v - Pv\|$. Hence, if P (or equivalently L, R) is chosen so as to make $\|e[P](v)\|$ as small as possible, one can preserve to any desired level the original eigenvalue λ in the coarse-grained matrix M' ; under extra assumptions on M , this result can be generalized to eigenvectors [1]. This leads to the following generic definition of a SCG problem.

Given $M (C[n \times n])$ and (λ, v) , a (right) eigenpair of M to be preserved by the coarse graining, the problem is to find a projector P solving

$$\min(\|e[P](v)\|, p \text{ in } \Omega),$$

where Ω is a set of projectors in $C[n \times n]$ described by some ad hoc constraints $c[1], \dots, c[r]$ (e.g. $c[1]: P \text{ in } R[n \times n], c[2]: P = t(P), c[3]: P[i, j] \geq 0$, etc).

Choosing pertinent constraints to solve the SCG problem is of great importance in applications. For instance, in the absence of constraints the SCG problem is solved trivially by $P = vv^*$ (v is assumed normalized). We have designed a particular constraint, called *homogeneous mixing*, which ensures that vertices belonging to the same group are merged consistently from a physical point of view (see [1] for details). Under this constraint the SCG problem reduces to finding the partition of $1, \dots, n$ (labeling the original vertices) minimizing

$$\|e[P](v)\|^2 = \sum ([v(i) - (Pv)(i)]^2; \alpha = 1, \dots, n', i \text{ in } \alpha),$$

where α denotes a group (i.e. a block) in a partition of $\{1, \dots, n\}$, and $|\alpha|$ is the number of elements in α .

If M is symmetric or stochastic, for instance, then it may be desirable (or mandatory) to choose L, R so that M' is symmetric or stochastic as well. This *structural constraint* has led to the construction of particular semi-projectors for symmetric [1], stochastic [3] and Laplacian [2] matrices, that are made available.

In short, the coarse graining of matrices and graphs involves:

1. Retrieving a matrix or a graph matrix M from the problem.
2. Computing the eigenpairs of M to be preserved in the coarse-grained graph or matrix.
3. Setting some problem-specific constraints (e.g. dimension of the coarse-grained object).
4. Solving the constrained SCG problem, that is finding P' .
5. Computing from P' two semi-projectors L' and R' (e.g. following the method proposed in [1]).
6. Working out the product $M'=L'MR'^*$ and, if needed, defining from M' a coarse-grained graph.

Functions for performing SCG

The main functions are `igraph_scg_adjacency()`, `igraph_scg_laplacian()` and `igraph_scg_stochastic()`. These functions handle all the steps involved in the Spectral Coarse Graining (SCG) of some particular matrices and graphs as described above and in reference [1]. In more details, they compute some prescribed eigenpairs of a matrix or a graph matrix, (for now adjacency, Laplacian and stochastic matrices are available), work out an optimal partition to preserve the eigenpairs, and finally output a coarse-grained matrix or graph along with other useful information.

These steps can also be carried out independently: (1) Use `igraph_get_adjacency()`, `igraph_get_sparsemat()`, `igraph_laplacian()`, `igraph_get_stochastic()` or `igraph_get_stochastic_sparsemat()` to compute a matrix M . (2) Work out some prescribed eigenpairs of M e.g. by means of `igraph_arnold_solve()` or `igraph_arnold_rnsolve()`. (3) Invoke one of the four algorithms of the function `igraph_scg_grouping()` to get a partition that will preserve the eigenpairs in the coarse-grained matrix. (4) Compute the semi-projectors L and R using `igraph_scg_semiprojectors()` and from there the coarse-grained matrix $M'=LMR^*$. If necessary, construct a coarse-grained graph from M' (e.g. as in [1]).

References

- [1] D. Morton de Lachapelle, D. Gfeller, and P. De Los Rios, Shrinking Matrices while Preserving their Eigenpairs with Application to the Spectral Coarse Graining of Graphs. Submitted to *SIAM Journal on Matrix Analysis and Applications*, 2008. <http://people.epfl.ch/david.morton>
- [2] D. Gfeller, and P. De Los Rios, Spectral Coarse Graining and Synchronization in Oscillator Networks. *Physical Review Letters*, **100**(17), 2008. <http://arxiv.org/abs/0708.2055>
- [3] D. Gfeller, and P. De Los Rios, Spectral Coarse Graining of Complex Networks, *Physical Review Letters*, **99**(3), 2007. <http://arxiv.org/abs/0706.0812>

SCG functions

`igraph_scg_adjacency` — Spectral coarse graining, symmetric case.

```
int igraph_scg_adjacency(const igraph_t *graph,
                        const igraph_matrix_t *matrix,
                        const igraph_sparsemat_t *sparsemat,
                        const igraph_vector_t *ev,
                        igraph_integer_t nt,
                        const igraph_vector_t *nt_vec,
                        igraph_scg_algorithm_t algo,
                        igraph_vector_t *values,
                        igraph_matrix_t *vectors,
                        igraph_vector_t *groups,
                        igraph_bool_t use_arpack,
                        igraph_integer_t maxiter,
                        igraph_t *scg_graph,
                        igraph_matrix_t *scg_matrix,
                        igraph_sparsemat_t *scg_sparsemat,
                        igraph_matrix_t *L,
                        igraph_matrix_t *R,
                        igraph_sparsemat_t *Lsparse,
                        igraph_sparsemat_t *Rsparse);
```

This function handles all the steps involved in the Spectral Coarse Graining (SCG) of some matrices and graphs as described in the reference below.

Arguments:

<i>graph</i> :	The input graph. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>matrix</i> :	The input matrix. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>sparsemat</i> :	The input sparse matrix. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>ev</i> :	A vector of positive integers giving the indexes of the eigenpairs to be preserved. 1 designates the eigenvalue with largest algebraic value, 2 the one with second largest algebraic value, etc.
<i>nt</i> :	Positive integer. When <i>algo</i> is IGRAPH_SCG_OPTIMUM, it gives the number of groups to partition each eigenvector separately. When <i>algo</i> is IGRAPH_SCG_INTERV or IGRAPH_SCG_INTERV_KM, it gives the number of intervals to partition each eigenvector. This is ignored when <i>algo</i> is IGRAPH_SCG_EXACT.
<i>nt_vec</i> :	A numeric vector of length one or the length must match the number of eigenvectors given in <i>V</i> , or a NULL pointer. If not NULL, then this argument gives the number of groups or intervals, and <i>nt</i> is ignored. Different number of groups or intervals can be specified for each eigenvector.
<i>algo</i> :	The algorithm to solve the SCG problem. Possible values: IGRAPH_SCG_OPTIMUM, IGRAPH_SCG_INTERV_KM, IGRAPH_SCG_INTERV and IGRAPH_SCG_EXACT. Please see the details about them above.
<i>values</i> :	If this is not NULL and the eigenvectors are re-calculated, then the eigenvalues are stored here.
<i>vectors</i> :	If this is not NULL, and not a zero-length matrix, then it is interpreted as the eigenvectors to use for the coarse-graining. Otherwise the eigenvectors are re-calculated, and they are stored here. (If this is not NULL.)

<i>groups</i> :	If this is not NULL, and not a zero-length vector, then it is interpreted as the vector of group labels. (Group labels are integers from zero and are sequential.) Otherwise group labels are re-calculated and stored here, if this argument is not a null pointer.
<i>use_arpak</i> :	Whether to use ARPACK for solving the eigenproblem. Currently ARPACK is not implemented.
<i>maxiter</i> :	A positive integer giving the number of iterations of the k-means algorithm when <i>algo</i> is IGRAPH_SCG_INTERV_KM. It is ignored in other cases. A reasonable (initial) value for this argument is 100.
<i>scg_graph</i> :	If not a NULL pointer, then the coarse-grained graph is returned here.
<i>scg_matrix</i> :	If not a NULL pointer, then it must be an initialised matrix, and the coarse-grained matrix is returned here.
<i>scg_sparsemat</i> :	If not a NULL pointer, then the coarse grained matrix is returned here, in sparse matrix form.
<i>L</i> :	If not a NULL pointer, then it must be an initialized matrix and the left semi-projector is returned here.
<i>R</i> :	If not a NULL pointer, then it must be an initialized matrix and the right semi-projector is returned here.
<i>Lsparse</i> :	If not a NULL pointer, then the left semi-projector is returned here.
<i>Rsparse</i> :	If not a NULL pointer, then the right semi-projector is returned here.

Returns:

Error code.

Time complexity: TODO.

See also:

`igraph_scg_grouping()`, `igraph_scg_semiprojectors()`, `igraph_scg_stochastic()` and `igraph_scg_laplacian()`.

Example 27.1. File `examples/simple/scg.c`

`igraph_scg_stochastic` — Spectral coarse graining, stochastic case.

```
int igraph_scg_stochastic(const igraph_t *graph,
                          const igraph_matrix_t *matrix,
                          const igraph_sparsemat_t *sparsemat,
                          const igraph_vector_t *ev,
                          igraph_integer_t nt,
                          const igraph_vector_t *nt_vec,
                          igraph_scg_algorithm_t algo,
                          igraph_scg_norm_t norm,
                          igraph_vector_complex_t *values,
                          igraph_matrix_complex_t *vectors,
```

```
igraph_vector_t *groups,  
igraph_vector_t *p,  
igraph_bool_t use_arnpack,  
igraph_integer_t maxiter,  
igraph_t *scg_graph,  
igraph_matrix_t *scg_matrix,  
igraph_sparsemat_t *scg_sparsemat,  
igraph_matrix_t *L,  
igraph_matrix_t *R,  
igraph_sparsemat_t *Lsparse,  
igraph_sparsemat_t *Rsparse);
```

This function handles all the steps involved in the Spectral Coarse Graining (SCG) of some matrices and graphs as described in the reference below.

Arguments:

<i>graph</i> :	The input graph. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>matrix</i> :	The input matrix. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>sparsemat</i> :	The input sparse matrix. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>ev</i> :	A vector of positive integers giving the indexes of the eigenpairs to be preserved. 1 designates the eigenvalue with largest magnitude, 2 the one with second largest magnitude, etc.
<i>nt</i> :	Positive integer. When <i>algo</i> is IGRAPH_SCG_OPTIMUM, it gives the number of groups to partition each eigenvector separately. When <i>algo</i> is IGRAPH_SCG_INTERV or IGRAPH_SCG_INTERV_KM, it gives the number of intervals to partition each eigenvector. This is ignored when <i>algo</i> is IGRAPH_SCG_EXACT.
<i>nt_vec</i> :	A numeric vector of length one or the length must match the number of eigenvectors given in <i>V</i> , or a NULL pointer. If not NULL, then this argument gives the number of groups or intervals, and <i>nt</i> is ignored. Different number of groups or intervals can be specified for each eigenvector.
<i>algo</i> :	The algorithm to solve the SCG problem. Possible values: IGRAPH_SCG_OPTIMUM, IGRAPH_SCG_INTERV_KM, IGRAPH_SCG_INTERV and IGRAPH_SCG_EXACT. Please see the details about them above.
<i>norm</i> :	Either IGRAPH_SCG_NORM_ROW or IGRAPH_SCG_NORM_COL. Specifies whether the rows or the columns of the stochastic matrix sum up to one.
<i>values</i> :	If this is not NULL and the eigenvectors are re-calculated, then the eigenvalues are stored here.
<i>vectors</i> :	If this is not NULL, and not a zero-length matrix, then it is interpreted as the eigenvectors to use for the coarse-graining. Otherwise the eigenvectors are re-calculated, and they are stored here. (If this is not NULL.)
<i>groups</i> :	If this is not NULL, and not a zero-length vector, then it is interpreted as the vector of group labels. (Group labels are integers from zero and are sequential.) Otherwise group labels are re-calculated and stored here, if this argument is not a null pointer.

<i>p</i> :	If this is not NULL, and not zero length, then it is interpreted as the stationary probability distribution of the Markov chain corresponding to the input matrix/graph. Its length must match the number of vertices in the input graph (or number of rows in the input matrix). If not given, then the stationary distribution is calculated and stored here. (Unless this argument is a NULL pointer, in which case it is not stored.)
<i>use_arpack</i> :	Whether to use ARPACK for solving the eigenproblem. Currently ARPACK is not implemented.
<i>maxiter</i> :	A positive integer giving the number of iterations of the k-means algorithm when <i>algo</i> is IGRAPH_SCG_INTERV_KM. It is ignored in other cases. A reasonable (initial) value for this argument is 100.
<i>scg_graph</i> :	If not a NULL pointer, then the coarse-grained graph is returned here.
<i>scg_matrix</i> :	If not a NULL pointer, then it must be an initialised matrix, and the coarse-grained matrix is returned here.
<i>scg_sparsemat</i> :	If not a NULL pointer, then the coarse grained matrix is returned here, in sparse matrix form.
<i>L</i> :	If not a NULL pointer, then it must be an initialized matrix and the left semi-projector is returned here.
<i>R</i> :	If not a NULL pointer, then it must be an initialized matrix and the right semi-projector is returned here.
<i>Lsparse</i> :	If not a NULL pointer, then the left semi-projector is returned here.
<i>Rsparse</i> :	If not a NULL pointer, then the right semi-projector is returned here.

Returns:

Error code.

Time complexity: TODO.

See also:

`igraph_scg_grouping()`, `igraph_scg_semiprojectors()`, `igraph_scg_adjacency()` and `igraph_scg_laplacian()`.

igraph_scg_laplacian — Spectral coarse graining, Laplacian case.

```
int igraph_scg_laplacian(const igraph_t *graph,
                        const igraph_matrix_t *matrix,
                        const igraph_sparsemat_t *sparsemat,
                        const igraph_vector_t *ev,
                        igraph_integer_t nt,
                        const igraph_vector_t *nt_vec,
                        igraph_scg_algorithm_t algo,
                        igraph_scg_norm_t norm,
                        igraph_scg_direction_t direction,
                        igraph_vector_complex_t *values,
```

```
igraph_matrix_complex_t *vectors,  
igraph_vector_t *groups,  
igraph_bool_t use_arpack,  
igraph_integer_t maxiter,  
igraph_t *scg_graph,  
igraph_matrix_t *scg_matrix,  
igraph_sparsemat_t *scg_sparsemat,  
igraph_matrix_t *L,  
igraph_matrix_t *R,  
igraph_sparsemat_t *Lsparse,  
igraph_sparsemat_t *Rsparse);
```

This function handles all the steps involved in the Spectral Coarse Graining (SCG) of some matrices and graphs as described in the reference below.

Arguments:

<i>graph</i> :	The input graph. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>matrix</i> :	The input matrix. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>sparsemat</i> :	The input sparse matrix. Exactly one of <i>graph</i> , <i>matrix</i> and <i>sparsemat</i> must be given, the other two must be NULL pointers.
<i>ev</i> :	A vector of positive integers giving the indexes of the eigenpairs to be preserved. 1 designates the eigenvalue with largest magnitude, 2 the one with second largest magnitude, etc.
<i>nt</i> :	Positive integer. When <i>algo</i> is IGRAPH_SCG_OPTIMUM, it gives the number of groups to partition each eigenvector separately. When <i>algo</i> is IGRAPH_SCG_INTERV or IGRAPH_SCG_INTERV_KM, it gives the number of intervals to partition each eigenvector. This is ignored when <i>algo</i> is IGRAPH_SCG_EXACT.
<i>nt_vec</i> :	A numeric vector of length one or the length must match the number of eigenvectors given in <i>V</i> , or a NULL pointer. If not NULL, then this argument gives the number of groups or intervals, and <i>nt</i> is ignored. Different number of groups or intervals can be specified for each eigenvector.
<i>algo</i> :	The algorithm to solve the SCG problem. Possible values: IGRAPH_SCG_OPTIMUM, IGRAPH_SCG_INTERV_KM, IGRAPH_SCG_INTERV and IGRAPH_SCG_EXACT. Please see the details about them above.
<i>norm</i> :	Either IGRAPH_SCG_NORM_ROW or IGRAPH_SCG_NORM_COL. Specifies whether the rows or the columns of the Laplacian matrix sum up to zero.
<i>direction</i> :	Whether to work with left or right eigenvectors. Possible values: IGRAPH_SCG_DIRECTION_DEFAULT, IGRAPH_SCG_DIRECTION_LEFT, IGRAPH_SCG_DIRECTION_RIGHT. This argument is currently ignored and right eigenvectors are always used.
<i>values</i> :	If this is not NULL and the eigenvectors are re-calculated, then the eigenvalues are stored here.
<i>vectors</i> :	If this is not NULL, and not a zero-length matrix, then it is interpreted as the eigenvectors to use for the coarse-graining. Otherwise the eigenvectors are re-calculated, and they are stored here. (If this is not NULL.)

<i>groups</i> :	If this is not NULL, and not a zero-length vector, then it is interpreted as the vector of group labels. (Group labels are integers from zero and are sequential.) Otherwise group labels are re-calculated and stored here, if this argument is not a null pointer.
<i>use_arpack</i> :	Whether to use ARPACK for solving the eigenproblem. Currently ARPACK is not implemented.
<i>maxiter</i> :	A positive integer giving the number of iterations of the k-means algorithm when <i>algo</i> is IGRAPH_SCG_INTERV_KM. It is ignored in other cases. A reasonable (initial) value for this argument is 100.
<i>scg_graph</i> :	If not a NULL pointer, then the coarse-grained graph is returned here.
<i>scg_matrix</i> :	If not a NULL pointer, then it must be an initialised matrix, and the coarse-grained matrix is returned here.
<i>scg_sparsemat</i> :	If not a NULL pointer, then the coarse grained matrix is returned here, in sparse matrix form.
<i>L</i> :	If not a NULL pointer, then it must be an initialized matrix and the left semi-projector is returned here.
<i>R</i> :	If not a NULL pointer, then it must be an initialized matrix and the right semi-projector is returned here.
<i>Lsparse</i> :	If not a NULL pointer, then the left semi-projector is returned here.
<i>Rsparse</i> :	If not a NULL pointer, then the right semi-projector is returned here.

Returns:

Error code.

Time complexity: TODO.

See also:

`igraph_scg_grouping()`, `igraph_scg_semiprojectors()`, `igraph_scg_stochastic()` and `igraph_scg_adjacency()`.

`igraph_scg_grouping` — SCG problem solver.

```
int igraph_scg_grouping(const igraph_matrix_t *V,
                        igraph_vector_t *groups,
                        igraph_integer_t nt,
                        const igraph_vector_t *nt_vec,
                        igraph_scg_matrix_t mtype,
                        igraph_scg_algorithm_t algo,
                        const igraph_vector_t *p,
                        igraph_integer_t maxiter);
```

This function solves the Spectral Coarse Graining (SCG) problem; either exactly, or approximately but faster.

The algorithm IGRAPH_SCG_OPTIMUM solves the SCG problem exactly for each eigenvector in *V*. The running time of this algorithm is $O(\max(nt) m^2)$ for the symmetric and Laplacian matrix

problems. It is $O(m^3)$ for the stochastic problem. Here m is the number of rows in V . In all three cases, the memory usage is $O(m^2)$.

The algorithms `IGRAPH_SCG_INTERV` and `IGRAPH_SCG_INTERV_KM` solve the SCG problem approximately by performing a (for now) constant binning of the components of the eigenvectors, that is `nt_vec[i]` constant-size bins are used to partition the i th eigenvector in V . When `algo` is `IGRAPH_SCG_INTERV_KM`, the (Lloyd) k-means algorithm is run on each partition obtained by `IGRAPH_SCG_INTERV` to improve accuracy.

Once a minimizing partition (either exact or approximate) has been found for each eigenvector, the final grouping is worked out as follows: two vertices are grouped together in the final partition if they are grouped together in each minimizing partition. In general, the size of the final partition is not known in advance when the number of columns in V is larger than one.

Finally, the algorithm `IGRAPH_SCG_EXACT` groups the vertices with equal components in each eigenvector. The last three algorithms essentially have linear running time and memory load.

Arguments:

- `V`: The matrix of eigenvectors to be preserved by coarse graining, each column is an eigenvector.
- `groups`: Pointer to an initialized vector; the result of the SCG is stored here.
- `nt`: Positive integer. When `algo` is `IGRAPH_SCG_OPTIMUM`, it gives the number of groups to partition each eigenvector separately. When `algo` is `IGRAPH_SCG_INTERV` or `IGRAPH_SCG_INTERV_KM`, it gives the number of intervals to partition each eigenvector. This is ignored when `algo` is `IGRAPH_SCG_EXACT`.
- `nt_vec`: May be (1) a numeric vector of length one, or (2) a vector of the same length as the number of eigenvectors given in V , or (3) a `NULL` pointer. If not `NULL`, then this argument gives the number of groups or intervals, and `nt` is ignored. Different number of groups or intervals can be specified for each eigenvector.
- `mttype`: The type of semi-projectors used in the SCG. Possible values are `IGRAPH_SCG_SYMMETRIC`, `IGRAPH_SCG_STOCHASTIC` and `IGRAPH_SCG_LAPLACIAN`.
- `algo`: The algorithm to solve the SCG problem. Possible values: `IGRAPH_SCG_OPTIMUM`, `IGRAPH_SCG_INTERV_KM`, `IGRAPH_SCG_INTERV` and `IGRAPH_SCG_EXACT`. Please see the details about them above.
- `p`: A probability vector, or `NULL`. This argument must be given if `mttype` is `IGRAPH_SCG_STOCHASTIC`, but it is ignored otherwise. For the stochastic case it gives the stationary probability distribution of a Markov chain, the one specified by the graph/matrix under study.
- `maxiter`: A positive integer giving the number of iterations of the k-means algorithm when `algo` is `IGRAPH_SCG_INTERV_KM`. It is ignored in other cases. A reasonable (initial) value for this argument is 100.

Returns:

Error code.

Time complexity: see description above.

See also:

`igraph_scg_adjacency()`, `igraph_scg_laplacian()`, `igraph_scg_stochastic()`.

Example 27.2. File `examples/simple/igraph_scg_grouping.c`

Example 27.3. File `examples/simple/igraph_scg_grouping2.c`

Example 27.4. File `examples/simple/igraph_scg_grouping3.c`

Example 27.5. File `examples/simple/igraph_scg_grouping4.c`

igraph_scg_semiprojectors — Compute SCG semi-projectors for a given partition.

```
int igraph_scg_semiprojectors(const igraph_vector_t *groups,
                             igraph_scg_matrix_t mtype,
                             igraph_matrix_t *L,
                             igraph_matrix_t *R,
                             igraph_sparsemat_t *Lsparse,
                             igraph_sparsemat_t *Rsparse,
                             const igraph_vector_t *p,
                             igraph_scg_norm_t norm);
```

The three types of semi-projectors are defined as follows. Let $\gamma(j)$ label the group of vertex j in a partition of all the vertices.

The symmetric semi-projectors are defined as

$$L[\alpha, j] = R[\alpha, j] = 1/\sqrt{|\alpha|} \delta[\alpha, \gamma(j)],$$

the (row) Laplacian semi-projectors as

$$L[\alpha, j] = 1/|\alpha| \delta[\alpha, \gamma(j)]$$

and

$$R[\alpha, j] = \delta[\alpha, \gamma(j)],$$

and the (row) stochastic semi-projectors as

$$L[\alpha, j] = p[1][j] / \sum(p[1][k]; k \text{ in } \gamma(j)) \delta[\alpha, \gamma(j)]$$

and

$$R[\alpha, j] = \delta[\alpha, \gamma(j)],$$

where $p[1]$ is the (left) eigenvector associated with the one-eigenvalue of the stochastic matrix. L and R are defined in a symmetric way when `norm` is `IGRAPH_SCG_NORM_COL`. All these semi-projectors verify various properties described in the reference.

Arguments:

groups: A vector of integers, giving the group label of every vertex in the partition. Group labels should start at zero and should be sequential.

mtype: The type of semi-projectors. For now `IGRAPH_SCG_SYMMETRIC`, `IGRAPH_SCG_STOCHASTIC` and `IGRAPH_SCG_LAPLACIAN` are supported.

- L*: If not a NULL pointer, then it must be a pointer to an initialized matrix. The left semi-projector is stored here.
- R*: If not a NULL pointer, then it must be a pointer to an initialized matrix. The right semi-projector is stored here.
- Lsparse*: If not a NULL pointer, then it must be a pointer to an uninitialized sparse matrix. The left semi-projector is stored here.
- Rsparse*: If not a NULL pointer, then it must be a pointer to an uninitialized sparse matrix. The right semi-projector is stored here.
- p*: NULL, or a probability vector of the same length as *groups*. *p* is the stationary probability distribution of a Markov chain when *mtype* is IGRAPH_SCG_STOCHASTIC. This argument is ignored in all other cases.
- norm*: Either IGRAPH_SCG_NORM_ROW or IGRAPH_SCG_NORM_COL. Specifies whether the rows or the columns of the Laplacian matrix sum up to zero, or whether the rows or the columns of the stochastic matrix sum up to one.

Returns:

Error code.

Time complexity: TODO.

See also:

`igraph_scg_adjacency()`, `igraph_scg_stochastic()` and `igraph_scg_laplacian()`, `igraph_scg_grouping()`.

Example 27.6. **File** `examples/simple/igraph_scg_semiprojectors.c`

Example 27.7. **File** `examples/simple/igraph_scg_semiprojectors2.c`

Example 27.8. **File** `examples/simple/igraph_scg_semiprojectors3.c`

igraph_scg_norm_eps — Calculate SCG residuals.

```
int igraph_scg_norm_eps(const igraph_matrix_t *V,
                        const igraph_vector_t *groups,
                        igraph_vector_t *eps,
                        igraph_scg_matrix_t mtype,
                        const igraph_vector_t *p,
                        igraph_scg_norm_t norm);
```

Computes $|v[i] - Pv[i]|$, where $v[i]$ is the i -th eigenvector in V and P is the projector corresponding to the *mtype* argument.

Arguments:

- V*: The matrix of eigenvectors to be preserved by coarse graining, each column is an eigenvector.
- groups*: A vector of integers, giving the group label of every vertex in the partition. Group labels should start at zero and should be sequential.
- eps*: Pointer to a real value, the result is stored here.
- mtype*: The type of semi-projectors. For now `IGRAPH_SCG_SYMMETRIC`, `IGRAPH_SCG_STOCHASTIC` and `IGRAPH_SCG_LAPLACIAN` are supported.
- p*: NULL, or a probability vector of the same length as *groups*. *p* is the stationary probability distribution of a Markov chain when *mtype* is `IGRAPH_SCG_STOCHASTIC`. This argument is ignored in all other cases.
- norm*: Either `IGRAPH_SCG_NORM_ROW` or `IGRAPH_SCG_NORM_COL`. Specifies whether the rows or the columns of the Laplacian matrix sum up to zero, or whether the rows or the columns of the stochastic matrix sum up to one.

Returns:

Error code.

Time complexity: TODO.

See also:

`igraph_scg_adjacency()`, `igraph_scg_stochastic()` and
`igraph_scg_laplacian()`, `igraph_scg_grouping()`, `igraph_scg_semiprojectors()`.

Chapter 28. Embedding of graphs

Spectral embedding

`igraph_adjacency_spectral_embedding` — Adjacency spectral embedding

```
int igraph_adjacency_spectral_embedding(const igraph_t *graph,
                                       igraph_integer_t n,
                                       const igraph_vector_t *weights,
                                       igraph_eigen_which_position_t which,
                                       igraph_bool_t scaled,
                                       igraph_matrix_t *X,
                                       igraph_matrix_t *Y,
                                       igraph_vector_t *D,
                                       const igraph_vector_t *cvec,
                                       igraph_arnpack_options_t *options);
```

Spectral decomposition of the adjacency matrices of graphs. This function computes an n -dimensional Euclidean representation of the graph based on its adjacency matrix, A . This representation is computed via the singular value decomposition of the adjacency matrix, $A=U D V^T$. In the case, where the graph is a random dot product graph generated using latent position vectors in R^n for each vertex, the embedding will provide an estimate of these latent vectors.

For undirected graphs, the latent positions are calculated as $X = U^n D^{(1/2)}$ where U^n equals to the first n columns of U , and $D^{(1/2)}$ is a diagonal matrix containing the square root of the selected singular values on the diagonal.

For directed graphs, the embedding is defined as the pair $X = U^n D^{(1/2)}$, $Y = V^n D^{(1/2)}$. (For undirected graphs $U=V$, so it is sufficient to keep one of them.)

Arguments:

- graph*: The input graph, can be directed or undirected.
- n*: An integer scalar. This value is the embedding dimension of the spectral embedding. Should be smaller than the number of vertices. The largest n -dimensional non-zero singular values are used for the spectral embedding.
- weights*: Optional edge weights. Supply a null pointer for unweighted graphs.
- which*: Which eigenvalues (or singular values, for directed graphs) to use, possible values:
- `IGRAPH_EIGEN_LM` the ones with the largest magnitude
 - `IGRAPH_EIGEN_LA` the (algebraic) largest ones
 - `IGRAPH_EIGEN_SA` the (algebraic) smallest ones.
- For directed graphs, `IGRAPH_EIGEN_LM` and `IGRAPH_EIGEN_LA` are the same because singular values are used for the ordering instead of eigenvalues.
- scaled*: Whether to return X and Y (if *scaled* is true), or U and V .
- X*: Initialized matrix, the estimated latent positions are stored here.

- Y*: Initialized matrix or a null pointer. If not a null pointer, then the second half of the latent positions are stored here. (For undirected graphs, this always equals *X*.)
- D*: Initialized vector or a null pointer. If not a null pointer, then the eigenvalues (for undirected graphs) or the singular values (for directed graphs) are stored here.
- cvec*: A numeric vector, its length is the number vertices in the graph. This vector is added to the diagonal of the adjacency matrix, before performing the SVD.
- options*: Options to ARPACK. See `igraph_arnpack_options_t` for details. Note that the function overwrites the *n* (number of vertices), *nev* and *which* parameters and it always starts the calculation from a random start vector.

Returns:

Error code.

igraph_laplacian_spectral_embedding — Spectral embedding of the Laplacian of a graph

```
int igraph_laplacian_spectral_embedding(const igraph_t *graph,
                                       igraph_integer_t n,
                                       const igraph_vector_t *weights,
                                       igraph_eigen_which_position_t which,
                                       igraph_laplacian_spectral_embedding_type_t type,
                                       igraph_bool_t scaled,
                                       igraph_matrix_t *X,
                                       igraph_matrix_t *Y,
                                       igraph_vector_t *D,
                                       igraph_arnpack_options_t *options);
```

This function essentially does the same as `igraph_adjacency_spectral_embedding`, but works on the Laplacian of the graph, instead of the adjacency matrix.

Arguments:

- graph*: The input graph.
- n*: The number of eigenvectors (or singular vectors if the graph is directed) to use for the embedding.
- weights*: Optional edge weights. Supply a null pointer for unweighted graphs.
- which*: Which eigenvalues (or singular values, for directed graphs) to use, possible values:
- `IGRAPH_EIGEN_LM` the ones with the largest magnitude
 - `IGRAPH_EIGEN_LA` the (algebraic) largest ones
 - `IGRAPH_EIGEN_SA` the (algebraic) smallest ones.
- For directed graphs, `IGRAPH_EIGEN_LM` and `IGRAPH_EIGEN_LA` are the same because singular values are used for the ordering instead of eigenvalues.
- type*: The type of the Laplacian to use. Various definitions exist for the Laplacian of a graph, and one can choose between them with this argument. Possible values:

<code>IGRAPH_EMBEDDING_D_A</code>	means $D - A$ where D is the degree matrix and A is the adjacency matrix
<code>IGRAPH_EMBEDDING_DAD</code>	means D_i times A times D_i , where D_i is the inverse of the square root of the degree matrix;
<code>IGRAPH_EMBEDDING_I_DAD</code>	means $I - D_i A D_i$, where I is the identity matrix.
<i>scaled</i> :	Whether to return X and Y (if <i>scaled</i> is true), or U and V .
<i>X</i> :	Initialized matrix, the estimated latent positions are stored here.
<i>Y</i> :	Initialized matrix or a null pointer. If not a null pointer, then the second half of the latent positions are stored here. (For undirected graphs, this always equals X .)
<i>D</i> :	Initialized vector or a null pointer. If not a null pointer, then the eigenvalues (for undirected graphs) or the singular values (for directed graphs) are stored here.
<i>options</i> :	Options to ARPACK. See <code>igraph_arnpack_options_t</code> for details. Note that the function overwrites the <i>n</i> (number of vertices), <i>nev</i> and <i>which</i> parameters and it always starts the calculation from a random start vector.

Returns:

Error code.

See also:

`igraph_adjacency_spectral_embedding` to embed the adjacency matrix.

igraph_dim_select — Dimensionality selection

```
int igraph_dim_select(const igraph_vector_t *sv, igraph_integer_t *dim);
```

Dimensionality selection for singular values using profile likelihood.

The input of the function is a numeric vector which contains the measure of "importance" for each dimension.

For spectral embedding, these are the singular values of the adjacency matrix. The singular values are assumed to be generated from a Gaussian mixture distribution with two components that have different means and same variance. The dimensionality d is chosen to maximize the likelihood when the d largest singular values are assigned to one component of the mixture and the rest of the singular values assigned to the other component.

This function can also be used for the general separation problem, where we assume that the left and the right of the vector are coming from two normal distributions, with different means, and we want to know their border.

Arguments:

sv: A numeric vector, the ordered singular values.

dim: The result is stored here.

Returns:

Error code.

Time complexity: $O(n)$, n is the number of values in `sv`.

See also:

`igraph_adjacency_spectral_embedding()`.

Chapter 29. Graph operators

Union and intersection

igraph_disjoint_union — Creates the union of two disjoint graphs

```
int igraph_disjoint_union(igraph_t *res, const igraph_t *left,
                          const igraph_t *right);
```

First the vertices of the second graph will be relabeled with new vertex ids to have two disjoint sets of vertex ids, then the union of the two graphs will be formed. If the two graphs have $|V1|$ and $|V2|$ vertices and $|E1|$ and $|E2|$ edges respectively then the new graph will have $|V1|+|V2|$ vertices and $|E1|+|E2|$ edges.

Both graphs need to have the same directedness, i.e. either both directed or both undirected.

The current version of this function cannot handle graph, vertex and edge attributes, they will be lost.

Arguments:

res: Pointer to an uninitialized graph object, the result will stored here.

left: The first graph.

right: The second graph.

Returns:

Error code.

See also:

`igraph_disjoint_union_many()` for creating the disjoint union of more than two graphs,
`igraph_union()` for non-disjoint union.

Time complexity: $O(|V1|+|V2|+|E1|+|E2|)$.

Example 29.1. File `examples/simple/igraph_disjoint_union.c`

igraph_disjoint_union_many — The disjoint union of many graphs.

```
int igraph_disjoint_union_many(igraph_t *res,
                              const igraph_vector_ptr_t *graphs);
```

First the vertices in the graphs will be relabeled with new vertex ids to have pairwise disjoint vertex id sets and then the union of the graphs is formed. The number of vertices and edges in the result is the total number of vertices and edges in the graphs.

All graphs need to have the same directedness, i.e. either all directed or all undirected. If the graph list has length zero, the result will be a *directed* graph with no vertices.

The current version of this function cannot handle graph, vertex and edge attributes, they will be lost.

Arguments:

res: Pointer to an uninitialized graph object, the result of the operation will be stored here.

graphs: Pointer vector, contains pointers to initialized graph objects.

Returns:

Error code.

See also:

`igraph_disjoint_union()` for an easier syntax if you have only two graphs, `igraph_union_many()` for non-disjoint union.

Time complexity: $O(|V|+|E|)$, the number of vertices plus the number of edges in the result.

igraph_union — Calculates the union of two graphs.

```
int igraph_union(igraph_t *res,
                 const igraph_t *left, const igraph_t *right,
                 igraph_vector_t *edge_map1, igraph_vector_t *edge_map2);
```

The number of vertices in the result is that of the larger graph from the two arguments. The result graph contains edges which are present in at least one of the operand graphs.

Arguments:

res: Pointer to an uninitialized graph object, the result will be stored here.

left: The first graph.

right: The second graph.

edge_map1: Pointer to an initialized vector or a null pointer. If not a null pointer, it will contain a mapping from the edges of the first argument graph (*left*) to the edges of the result graph.

edge_map2: The same as *edge_map1*, but for the second graph, *right*.

Returns:

Error code.

See also:

`igraph_union_many()` for the union of many graphs, `igraph_intersection()` and `igraph_difference()` for other operators.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of vertices, $|E|$ the number of edges in the result graph.

Example 29.2. File `examples/simple/igraph_union.c`

igraph_union_many — Creates the union of many graphs.

```
int igraph_union_many(igraph_t *res, const igraph_vector_ptr_t *graphs,
                      igraph_vector_ptr_t *edgemaps);
```

The result graph will contain as many vertices as the largest graph among the arguments does, and an edge will be included in it if it is part of at least one operand graph.

The directedness of the operand graphs must be the same. If the graph list has length zero, the result will be a *directed* graph with no vertices.

Arguments:

- res*: Pointer to an uninitialized graph object, this will contain the result.
- graphs*: Pointer vector, contains pointers to the operands of the union operator, graph objects of course.
- edgemaps*: If not a null pointer, then it must be an initialized pointer vector and the mappings of edges from the graphs to the result graph will be stored here, in the same order as *graphs*. Each mapping is stored in a separate `igraph_vector_t` object.

Returns:

Error code.

See also:

`igraph_union()` for the union of two graphs, `igraph_intersection_many()`, `igraph_intersection()` and `igraph_difference` for other operators.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of vertices in largest graph and $|E|$ is the number of edges in the result graph.

Example 29.3. File `examples/simple/igraph_union.c`

igraph_intersection — Collect the common edges from two graphs.

```
int igraph_intersection(igraph_t *res,
                       const igraph_t *left, const igraph_t *right,
                       igraph_vector_t *edge_map1,
                       igraph_vector_t *edge_map2);
```

The result graph contains only edges present both in the first and the second graph. The number of vertices in the result graph is the same as the larger from the two arguments.

Arguments:

- res*: Pointer to an uninitialized graph object. This will contain the result of the operation.
- left*: The first operand, a graph object.
- right*: The second operand, a graph object.
- edge_map1*: Null pointer, or an initialized `igraph_vector_t`. If the latter, then a mapping from the edges of the result graph, to the edges of the *left* input graph is stored here.
- edge_map2*: Null pointer, or an `igraph_vector_t`. The same as *edge_map1*, but for the *right* input graph.

Returns:

Error code.

See also:

`igraph_intersection_many()` to calculate the intersection of many graphs at once, `igraph_union()`, `igraph_difference()` for other operators.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of nodes, $|E|$ is the number of edges in the smaller graph of the two. (The one containing less vertices is considered smaller.)

Example 29.4. File `examples/simple/igraph_intersection.c`

`igraph_intersection_many` — The intersection of more than two graphs.

```
int igraph_intersection_many(igraph_t *res,
                             const igraph_vector_ptr_t *graphs,
                             igraph_vector_ptr_t *edgemaps);
```

This function calculates the intersection of the graphs stored in the *graphs* argument. Only those edges will be included in the result graph which are part of every graph in *graphs*.

The number of vertices in the result graph will be the maximum number of vertices in the argument graphs.

Arguments:

- res*: Pointer to an uninitialized graph object, the result of the operation will be stored here.
- graphs*: Pointer vector, contains pointers to graphs objects, the operands of the intersection operator.
- edgemaps*: If not a null pointer, then it must be an initialized pointer vector and the mappings of edges from the graphs to the result graph will be stored here, in the same order as *graphs*. Each mapping is stored in a separate `igraph_vector_t` object. For the edges that are not in the intersection, -1 is stored.

Returns:

Error code.

See also:

`igraph_intersection()` for the intersection of two graphs, `igraph_union_many()`, `igraph_union()` and `igraph_difference()` for other operators.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges in the smallest graph (i.e. the graph having the less vertices).

Other set-like operators

`igraph_difference` — Calculate the difference of two graphs

```
int igraph_difference(igraph_t *res,
                     const igraph_t *orig, const igraph_t *sub);
```

The number of vertices in the result is the number of vertices in the original graph, i.e. the left, first operand. In the results graph only edges will be included from *orig* which are not present in *sub*.

Arguments:

res: Pointer to an uninitialized graph object, the result will be stored here.

orig: The left operand of the operator, a graph object.

sub: The right operand of the operator, a graph object.

Returns:

Error code.

See also:

`igraph_intersection()` and `igraph_union()` for other operators.

Time complexity: $O(|V|+|E|)$, $|V|$ is the number vertices in the smaller graph, $|E|$ is the number of edges in the result graph.

Example 29.5. File `examples/simple/igraph_difference.c`

`igraph_complementer` — Create the complementer of a graph

```
int igraph_complementer(igraph_t *res, const igraph_t *graph,
                       igraph_bool_t loops);
```

The complementer graph means that all edges which are not part of the original graph will be included in the result.

Arguments:

res: Pointer to an uninitialized graph object.

graph: The original graph.

loops: Whether to add loop edges to the complementer graph.

Returns:

Error code.

See also:

`igraph_union()`, `igraph_intersection()` and `igraph_difference()`.

Time complexity: $O(|V|+|E1|+|E2|)$, $|V|$ is the number of vertices in the graph, $|E1|$ is the number of edges in the original and $|E2|$ in the complementer graph.

Example 29.6. File `examples/simple/igraph_complementer.c`

igraph_compose — Calculates the composition of two graphs

```
int igraph_compose(igraph_t *res, const igraph_t *g1, const igraph_t *g2,
                  igraph_vector_t *edge_map1, igraph_vector_t *edge_map2);
```

The composition of graphs contains the same number of vertices as the bigger graph of the two operands. It contains an (i,j) edge if and only if there is a k vertex, such that the first graph contains an (i,k) edge and the second graph a (k,j) edge.

This is of course exactly the composition of two binary relations.

Two graphs must have the same directedness, otherwise the function returns with an error message. Note that for undirected graphs the two relations are by definition symmetric.

Arguments:

res: Pointer to an uninitialized graph object, the result will be stored here.

g1: The first operand, a graph object.

g2: The second operand, another graph object.

edge_map1: If not a null pointer, then it must be a pointer to an initialized vector, and a mapping from the edges of the result graph to the edges of the first graph is stored here.

edge_map2: If not a null pointer, then it must be a pointer to an initialized vector, and a mapping from the edges of the result graph to the edges of the second graph is stored here.

Returns:

Error code.

Time complexity: $O(|V|*d1*d2)$, $|V|$ is the number of vertices in the first graph, $d1$ and $d2$ the average degree in the first and second graphs.

Example 29.7. File `examples/simple/igraph_compose.c`

Miscellaneous operators

igraph_connect_neighborhood — Connects every vertex to its neighborhood

```
int igraph_connect_neighborhood(igraph_t *graph, igraph_integer_t order,
                                igraph_neimode_t mode);
```

This function adds new edges to the input graph. Each vertex is connected to all vertices reachable by at most *order* steps from it (unless a connection already existed). In other words, the *order* power of the graph is computed.

Note that the input graph is modified in place, no new graph is created. Call `igraph_copy()` if you want to keep the original graph as well.

For undirected graphs reachability is always symmetric: if vertex A can be reached from vertex B in at most *order* steps, then the opposite is also true. Only one undirected (A,B) edge will be added in this case.

Arguments:

graph: The input graph, this is the output graph as well.

order: Integer constant, it gives the distance within which the vertices will be connected to the source vertex.

mode: Constant, it specifies how the neighborhood search is performed for directed graphs. If `IGRAPH_OUT` then vertices reachable from the source vertex will be connected, `IGRAPH_IN` is the opposite. If `IGRAPH_ALL` then the directed graph is considered as an undirected one.

Returns:

Error code.

See also:

`igraph_lattice()` uses this function to connect the neighborhood of the vertices.

Time complexity: $O(|V|*d^k)$, $|V|$ is the number of vertices in the graph, d is the average degree and k is the *order* argument.

igraph_contract_vertices — Replace multiple vertices with a single one.

```
int igraph_contract_vertices(igraph_t *graph,
```

```
const igraph_vector_t *mapping,  
const igraph_attribute_combination_t *vertex_comb)
```

This function creates a new graph, by merging several vertices into one. The vertices in the new graph correspond to sets of vertices in the input graph.

Arguments:

graph: The input graph, it can be directed or undirected.

mapping: A vector giving the mapping. For each vertex in the original graph, it should contain its id in the new graph.

vertex_comb: What to do with the vertex attributes. See the igraph manual section about attributes for details.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices plus edges.

igraph_induced_subgraph — Creates a subgraph induced by the specified vertices.

```
int igraph_induced_subgraph(const igraph_t *graph, igraph_t *res,  
                           const igraph_vs_t vids, igraph_subgraph_implementation_t impl)
```

This function collects the specified vertices and all edges between them to a new graph. As the vertex ids in a graph always start with zero, this function very likely needs to reassign ids to the vertices.

Arguments:

graph: The graph object.

res: The subgraph, another graph object will be stored here, do *not* initialize this object before calling this function, and call `igraph_destroy()` on it if you don't need it any more.

vids: A vertex selector describing which vertices to keep.

impl: This parameter selects which implementation should we use when constructing the new graph. Basically there are two possibilities: `IGRAPH_SUBGRAPH_COPY_AND_DELETE` copies the existing graph and deletes the vertices that are not needed in the new graph, while `IGRAPH_SUBGRAPH_CREATE_FROM_SCRATCH` constructs the new graph from scratch without copying the old one. The latter is more efficient if you are extracting a relatively small subpart of a very large graph, while the former is better if you want to extract a subgraph whose size is comparable to the size of the whole graph. There is a third possibility: `IGRAPH_SUBGRAPH_AUTO` will select one of the two methods automatically based on the ratio of the number of vertices in the new and the old graph.

Returns:

Error code: `IGRAPH_ENOMEM`, not enough memory for temporary data. `IGRAPH_EINVVID`, invalid vertex id in *vids*.

Time complexity: $O(|V|+|E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the original graph.

See also:

`igraph_delete_vertices()` to delete the specified set of vertices from a graph, the opposite of this function.

igraph_linegraph — Create the line graph of a graph.

```
int igraph_linegraph(const igraph_t *graph, igraph_t *linegraph);
```

The line graph $L(G)$ of a G undirected graph is defined as follows. $L(G)$ has one vertex for each edge in G and two different vertices in $L(G)$ are connected by an edge if their corresponding edges share an end point. In a multigraph, if two end points are shared, two edges are created. The vertex of a loop is counted as two end points.

The line graph $L(G)$ of a G directed graph is slightly different, $L(G)$ has one vertex for each edge in G and two vertices in $L(G)$ are connected by a directed edge if the target of the first vertex's corresponding edge is the same as the source of the second vertex's corresponding edge.

Edge i in the original graph will correspond to vertex i in the line graph.

The first version of this function was contributed by Vincent Matossian, thanks.

Arguments:

graph: The input graph, may be directed or undirected.

linegraph: Pointer to an uninitialized graph object, the result is stored here.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, the number of edges plus the number of vertices.

igraph_simplify — Removes loop and/or multiple edges from the graph.

```
int igraph_simplify(igraph_t *graph, igraph_bool_t multiple,  
                   igraph_bool_t loops,  
                   const igraph_attribute_combination_t *edge_comb);
```

Arguments:

graph: The graph object.

multiple: Logical, if true, multiple edges will be removed.

loops: Logical, if true, loops (self edges) will be removed.

edge_comb: What to do with the edge attributes. See the igraph manual section about attributes for details.

Returns:

Error code: IGRAPH_ENOMEM if we are out of memory.

Time complexity: $O(|V|+|E|)$.

Example 29.8. File `examples/simple/igraph_simplify.c`

igraph_subgraph_edges — Creates a subgraph with the specified edges and their endpoints.

```
int igraph_subgraph_edges(const igraph_t *graph, igraph_t *res,
                        const igraph_es_t eids, igraph_bool_t delete_vertices)
```

This function collects the specified edges and their endpoints to a new graph. As the vertex ids in a graph always start with zero, this function very likely needs to reassign ids to the vertices.

Arguments:

<i>graph</i> :	The graph object.
<i>res</i> :	The subgraph, another graph object will be stored here, do <i>not</i> initialize this object before calling this function, and call <code>igraph_destroy()</code> on it if you don't need it any more.
<i>eids</i> :	An edge selector describing which edges to keep.
<i>delete_vertices</i> :	Whether to delete the vertices not incident on any of the specified edges as well. If <code>FALSE</code> , the number of vertices in the result graph will always be equal to the number of vertices in the input graph.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVEID, invalid edge id in *eids*.

Time complexity: $O(|V|+|E|)$, $|V|$ and $|E|$ are the number of vertices and edges in the original graph.

See also:

`igraph_delete_edges()` to delete the specified set of edges from a graph, the opposite of this function.

Chapter 30. Using BLAS, LAPACK and ARPACK for igraph matrices and graphs

BLAS interface in igraph

BLAS is a highly optimized library for basic linear algebra operations such as vector-vector, matrix-vector and matrix-matrix product. Please see <http://www.netlib.org/blas/> for details and a reference implementation in Fortran. igraph contains some wrapper functions that can be used to call BLAS routines in a somewhat more user-friendly way. Not all BLAS routines are included in igraph, and even those which are included might not have wrappers; the extension of the set of wrapped functions will probably be driven by igraph's internal requirements. The wrapper functions usually substitute double-precision floating point arrays used by BLAS with `igraph_vector_t` and `igraph_matrix_t` instances and also remove those parameters (such as the number of rows/columns) that can be inferred from the passed arguments directly.

`igraph_blas_ddot` — Dot product of two vectors.

```
int igraph_blas_ddot(const igraph_vector_t *v1, const igraph_vector_t *v2,
                    igraph_real_t *res);
```

Arguments:

v1: The first vector.

v2: The second vector.

res: Pointer to a real, the result will be stored here.

Time complexity: $O(n)$ where n is the length of the vectors.

Example 30.1. File `examples/simple/blas.c`

`igraph_blas_dgemv` — Matrix-vector multiplication using BLAS, vector version.

```
void igraph_blas_dgemv(igraph_bool_t transpose, igraph_real_t alpha,
                      const igraph_matrix_t* a, const igraph_vector_t* x,
                      igraph_real_t beta, igraph_vector_t* y);
```

This function is a somewhat more user-friendly interface to the `dgemv` function in BLAS. `dgemv` performs the operation $y = \alpha * A * x + \beta * y$, where x and y are vectors and A is an appropriately sized matrix (symmetric or non-symmetric).

Arguments:

transpose: whether to transpose the matrix A

alpha: the constant *alpha*

a: the matrix *A*

x: the vector *x*

beta: the constant *beta*

y: the vector *y* (which will be modified in-place)

Time complexity: $O(nk)$ if the matrix is of size $n \times k$

See also:

`igraph_blas_dgemv_array` if you have arrays instead of vectors.

Example 30.2. File `examples/simple/blas.c`

`igraph_blas_dgemv_array` — Matrix-vector multiplication using BLAS, array version.

```
void igraph_blas_dgemv_array(igraph_bool_t transpose, igraph_real_t alpha,
                             const igraph_matrix_t* a, const igraph_real_t* x,
                             igraph_real_t beta, igraph_real_t* y);
```

This function is a somewhat more user-friendly interface to the `dgemv` function in BLAS. `dgemv` performs the operation $y = \alpha A x + \beta y$, where x and y are vectors and A is an appropriately sized matrix (symmetric or non-symmetric).

Arguments:

transpose: whether to transpose the matrix *A*

alpha: the constant *alpha*

a: the matrix *A*

x: the vector *x* as a regular C array

beta: the constant *beta*

y: the vector *y* as a regular C array (which will be modified in-place)

Time complexity: $O(nk)$ if the matrix is of size $n \times k$

See also:

`igraph_blas_dgemv` if you have vectors instead of arrays.

LAPACK interface in igraph

LAPACK is written in Fortran90 and provides routines for solving systems of simultaneous linear equations, least-squares solutions of linear systems of equations, eigenvalue problems, and singular value problems. The associated matrix factorizations (LU, Cholesky, QR, SVD, Schur, generalized Schur) are also provided, as are related computations such as reordering of the Schur factorizations and estimating condition numbers. Dense and banded matrices are handled, but not general sparse

matrices. In all areas, similar functionality is provided for real and complex matrices, in both single and double precision.

igraph provides an interface to a very limited set of LAPACK functions, using the regular igraph data structures.

See more about LAPACK at <http://www.netlib.org/lapack/>

Matrix factorization, solving linear systems

igraph_lapack_dgetrf — LU factorization of a general M-by-N matrix.

```
int igraph_lapack_dgetrf(igraph_matrix_t *a, igraph_vector_int_t *ipiv,  
                        int *info);
```

The factorization has the form $A = P * L * U$ where P is a permutation matrix, L is lower triangular with unit diagonal elements (lower trapezoidal if $m > n$), and U is upper triangular (upper trapezoidal if $m < n$).

Arguments:

- a*: The input/output matrix. On entry, the M-by-N matrix to be factored. On exit, the factors L and U from the factorization $A = P * L * U$; the unit diagonal elements of L are not stored.
- ipiv*: An integer vector, the pivot indices are stored here, unless it is a null pointer. Row i of the matrix was interchanged with row $ipiv[i]$.
- info*: LAPACK error code. Zero on successful exit. If its value is a positive number i , it indicates that $U(i,i)$ is exactly zero. The factorization has been completed, but the factor U is exactly singular, and division by zero will occur if it is used to solve a system of equations. If LAPACK returns an error, i.e. a negative *info* value, then an igraph error is generated as well.

Returns:

Error code.

Time complexity: TODO.

igraph_lapack_dgetrs — Solve general system of linear equations using LU factorization.

```
int igraph_lapack_dgetrs(igraph_bool_t transpose, const igraph_matrix_t *a,  
                        const igraph_vector_int_t *ipiv, igraph_matrix_t *b);
```

This function calls LAPACK to solve a system of linear equations $A * X = B$ or $A' * X = B$ with a general N-by-N matrix A using the LU factorization computed by `igraph_lapack_dgetrf`.

Arguments:

- transpose*: Logical scalar, whether to transpose the input matrix.
- a*: A matrix containing the L and U factors from the factorization $A = P * L * U$. L is expected to be unitriangular, diagonal entries are those of U . If A is singular, no warning or error will be given and random output will be returned.

ipiv: An integer vector, the pivot indices from `igraph_lapack_dgetrf()` must be given here. Row *i* of *A* was interchanged with row `ipiv[i]`.

b: The right hand side matrix must be given here. The solution will also be placed here.

Returns:

Error code.

Time complexity: TODO.

igraph_lapack_dgesv — Solve system of linear equations with LU factorization

```
int igraph_lapack_dgesv(igraph_matrix_t *a, igraph_vector_int_t *ipiv,
                        igraph_matrix_t *b, int *info);
```

This function computes the solution to a real system of linear equations $A * X = B$, where *A* is an *N*-by-*N* matrix and *X* and *B* are *N*-by-*NRHS* matrices.

The LU decomposition with partial pivoting and row interchanges is used to factor *A* as $A = P * L * U$, where *P* is a permutation matrix, *L* is unit lower triangular, and *U* is upper triangular. The factored form of *A* is then used to solve the system of equations $A * X = B$.

Arguments:

a: Matrix. On entry the *N*-by-*N* coefficient matrix, on exit, the factors *L* and *U* from the factorization $A = P * L * U$; the unit diagonal elements of *L* are not stored.

ipiv: An integer vector or a null pointer. If not a null pointer, then the pivot indices that define the permutation matrix *P*, are stored here. Row *i* of the matrix was interchanged with row `IPIV(i)`.

b: Matrix, on entry the right hand side matrix should be stored here. On exit, if there was no error, and the *info* argument is zero, then it contains the solution matrix *X*.

info: The LAPACK info code. If it is positive, then `U(info,info)` is exactly zero. In this case the factorization has been completed, but the factor *U* is exactly singular, so the solution could not be computed.

Returns:

Error code.

Time complexity: TODO.

Example 30.3. File `examples/simple/igraph_lapack_dgesv.c`

Eigenvalues and eigenvectors of matrices

igraph_lapack_dsyevr — Selected eigenvalues and optionally eigenvectors of a symmetric matrix

```
int igraph_lapack_dsyevr(const igraph_matrix_t *A,
                        igraph_lapack_dsyev_which_t which,
                        igraph_real_t vl, igraph_real_t vu, int vestimate,
                        int il, int iu, igraph_real_t abstol,
                        igraph_vector_t *values, igraph_matrix_t *vectors,
                        igraph_vector_int_t *support);
```

Calls the DSYEVR LAPACK function to compute selected eigenvalues and, optionally, eigenvectors of a real symmetric matrix A. Eigenvalues and eigenvectors can be selected by specifying either a range of values or a range of indices for the desired eigenvalues.

See more in the LAPACK documentation.

Arguments:

- A:** Matrix, on entry it contains the symmetric input matrix. Only the leading N-by-N upper triangular part is used for the computation.
- which:** Constant that gives which eigenvalues (and possibly the corresponding eigenvectors) to calculate. Possible values are IGRAPH_LAPACK_DSYEV_ALL, all eigenvalues; IGRAPH_LAPACK_DSYEV_INTERVAL, all eigenvalues in the half-open interval (vl,vu]; IGRAPH_LAPACK_DSYEV_SELECT, the il-th through iu-th eigenvalues.
- vl:** If *which* is IGRAPH_LAPACK_DSYEV_INTERVAL, then this is the lower bound of the interval to be searched for eigenvalues. See also the *vestimate* argument.
- vu:** If *which* is IGRAPH_LAPACK_DSYEV_INTERVAL, then this is the upper bound of the interval to be searched for eigenvalues. See also the *vestimate* argument.
- vestimate:** An upper bound for the number of eigenvalues in the (vl,vu] interval, if *which* is IGRAPH_LAPACK_DSYEV_INTERVAL. Memory is allocated only for the given number of eigenvalues (and eigenvectors), so this upper bound must be correct.
- il:** The index of the smallest eigenvalue to return, if *which* is IGRAPH_LAPACK_DSYEV_SELECT.
- iu:** The index of the targets eigenvalue to return, if *which* is IGRAPH_LAPACK_DSYEV_SELECT.
- abstol:** The absolute error tolerance for the eigenvalues. An approximate eigenvalue is accepted as converged when it is determined to lie in an interval [a,b] of width less than or equal to $abstol + EPS * \max(|a|,|b|)$, where EPS is the machine precision.
- values:** An initialized vector, the eigenvalues are stored here, unless it is a null pointer. It will be resized as needed.
- vectors:** An initialized matrix, the eigenvectors are stored in its columns, unless it is a null pointer. It will be resized as needed.
- support:** An integer vector. If not a null pointer, then it will be resized to $(2 * \max(1, M))$ (M is the total number of eigenvalues found). Then the support of the eigenvectors in *vectors* is stored here, i.e., the indices indicating the nonzero elements in *vectors*. The i-th eigenvector is nonzero only in elements support(2*i-1) through support(2*i).

Returns:

Error code.

Time complexity: TODO.

Example 30.4. File `examples/simple/igraph_lapack_dsyevr.c`

igraph_lapack_dgeev — Eigenvalues and optionally eigenvectors of a non-symmetric matrix

```
int igraph_lapack_dgeev(const igraph_matrix_t *A,
                        igraph_vector_t *valuesreal,
                        igraph_vector_t *valuesimag,
                        igraph_matrix_t *vectorsleft,
                        igraph_matrix_t *vectorsright,
                        int *info);
```

This function calls LAPACK to compute, for an N-by-N real nonsymmetric matrix A, the eigenvalues and, optionally, the left and/or right eigenvectors.

The right eigenvector $v(j)$ of A satisfies $A * v(j) = \lambda(j) * v(j)$ where $\lambda(j)$ is its eigenvalue. The left eigenvector $u(j)$ of A satisfies $u(j)^* * H * A = \lambda(j) * u(j)^* * H$ where $u(j)^* * H$ denotes the conjugate transpose of $u(j)$.

The computed eigenvectors are normalized to have Euclidean norm equal to 1 and largest component real.

Arguments:

- | | |
|-----------------------|--|
| <i>A</i> : | matrix. On entry it contains the N-by-N input matrix. |
| <i>valuesreal</i> : | Pointer to an initialized vector, or a null pointer. If not a null pointer, then the real parts of the eigenvalues are stored here. The vector will be resized as needed. |
| <i>valuesimag</i> : | Pointer to an initialized vector, or a null pointer. If not a null pointer, then the imaginary parts of the eigenvalues are stored here. The vector will be resized as needed. |
| <i>vectorsleft</i> : | Pointer to an initialized matrix, or a null pointer. If not a null pointer, then the left eigenvectors are stored in the columns of the matrix. The matrix will be resized as needed. |
| <i>vectorsright</i> : | Pointer to an initialized matrix, or a null pointer. If not a null pointer, then the right eigenvectors are stored in the columns of the matrix. The matrix will be resized as needed. |
| <i>info</i> : | This argument is used for two purposes. As an input argument it gives whether an igraph error should be generated if the QR algorithm fails to compute all eigenvalues. If <i>info</i> is non-zero, then an error is generated, otherwise only a warning is given. On exit it contains the LAPACK error code. Zero means successful exit. A negative values means that some of the arguments had an illegal value, this always triggers an igraph error. An i positive value means that the QR algorithm failed to compute all the eigenvalues, and no eigenvectors have been computed; element $i+1:N$ of <i>valuesreal</i> and <i>valuesimag</i> contain eigenvalues which have converged. This case only generates an igraph error, if <i>info</i> was non-zero on entry. |

Returns:

Error code.

Time complexity: TODO.

Example 30.5. File `examples/simple/igraph_lapack_dgeev.c`

igraph_lapack_dgeevx — Eigenvalues/vectors of nonsymmetric matrices, expert mode

```
int igraph_lapack_dgeevx(igraph_lapack_dgeevx_balance_t balance,
                        const igraph_matrix_t *A,
                        igraph_vector_t *valuesreal,
                        igraph_vector_t *valuesimag,
                        igraph_matrix_t *vectorsleft,
                        igraph_matrix_t *vectorsright,
                        int *ilo, int *ihi, igraph_vector_t *scale,
                        igraph_real_t *abnrm,
                        igraph_vector_t *rconde,
                        igraph_vector_t *rcondv,
                        int *info);
```

This function calculates the eigenvalues and optionally the left and/or right eigenvectors of a nonsymmetric N -by- N real matrix.

Optionally also, it computes a balancing transformation to improve the conditioning of the eigenvalues and eigenvectors (*ilo*, *ihi*, *scale*, and *abnrm*), reciprocal condition numbers for the eigenvalues (*rconde*), and reciprocal condition numbers for the right eigenvectors (*rcondv*).

The right eigenvector $v(j)$ of A satisfies $A * v(j) = \text{lambda}(j) * v(j)$ where $\text{lambda}(j)$ is its eigenvalue. The left eigenvector $u(j)$ of A satisfies $u(j)^H * A = \text{lambda}(j) * u(j)^H$ where $u(j)^H$ denotes the conjugate transpose of $u(j)$.

The computed eigenvectors are normalized to have Euclidean norm equal to 1 and largest component real.

Balancing a matrix means permuting the rows and columns to make it more nearly upper triangular, and applying a diagonal similarity transformation $D * A * D^{-1}$, where D is a diagonal matrix, to make its rows and columns closer in norm and the condition numbers of its eigenvalues and eigenvectors smaller. The computed reciprocal condition numbers correspond to the balanced matrix. Permuting rows and columns will not change the condition numbers (in exact arithmetic) but diagonal scaling will. For further explanation of balancing, see section 4.10.2 of the LAPACK Users' Guide.

Arguments:

<i>balance</i> :	Scalar that indicated, whether the input matrix should be balanced. Possible values:
IGRAPH_LAPACK_DGEEVX_BALANCE_NONE	no not diagonally scale or permute.
IGRAPH_LAPACK_DGEEVX_BALANCE_PERM	perform permutations to make the matrix more nearly upper triangular. Do not diagonally scale.
IGRAPH_LAPACK_DGEEVX_BALANCE_SCALE	diagonally scale the matrix, i.e. replace A by $D * A * D^{-1}$, where D is a diagonal matrix, chosen to make the rows and columns of A more equal in norm. Do not permute.

	IGRAPH_LA- PACK_DGEEVX_BALANCE_BOTH	both diagonally scale and permute A.
<i>A</i> :	The input matrix, must be square.	
<i>valuesreal</i> :	An initialized vector, or a NULL pointer. If not a NULL pointer, then the real parts of the eigenvalues are stored here. The vector will be resized, as needed.	
<i>valuesimag</i> :	An initialized vector, or a NULL pointer. If not a NULL pointer, then the imaginary parts of the eigenvalues are stored here. The vector will be resized, as needed.	
<i>vectorsleft</i> :	An initialized matrix or a NULL pointer. If not a null pointer, then the left eigenvectors are stored here. The order corresponds to the eigenvalues and the eigenvectors are stored in a compressed form. If the <i>j</i> -th eigenvalue is real then column <i>j</i> contains the corresponding eigenvector. If the <i>j</i> -th and (<i>j</i> +1)-th eigenvalues form a complex conjugate pair, then the <i>j</i> -th and (<i>j</i> +1)-th columns contain their corresponding eigenvectors.	
<i>vectorsright</i> :	An initialized matrix or a NULL pointer. If not a null pointer, then the right eigenvectors are stored here. The format is the same, as for the <i>vectorsleft</i> argument.	
<i>ilo</i> :		
<i>ihi</i> :	<i>ilo</i> and <i>ihi</i> are integer values determined when <i>A</i> was balanced. The balanced $A(i,j) = 0$ if $I > J$ and $J = 1, \dots, ilo-1$ or $I = ihi+1, \dots, N$.	
<i>scale</i> :	<p>Pointer to an initialized vector or a NULL pointer. If not a NULL pointer, then details of the permutations and scaling factors applied when balancing <i>A</i>, are stored here. If <i>P</i>(<i>j</i>) is the index of the row and column interchanged with row and column <i>j</i>, and <i>D</i>(<i>j</i>) is the scaling factor applied to row and column <i>j</i>, then</p> <p>$scale(J) = P(J), \text{ for } J = 1, \dots, ilo-1$</p> <p>$scale(J) = D(J), \text{ for } J = ilo, \dots, ihi$</p> <p>$scale(J) = P(J) \text{ for } J = ihi+1, \dots, N.$</p> <p>The order in which the interchanges are made is <i>N</i> to <i>ihi</i>+1, then 1 to <i>ilo</i>-1.</p>	
<i>abnorm</i> :	Pointer to a real variable, the one-norm of the balanced matrix is stored here. (The one-norm is the maximum of the sum of absolute values of elements in any column.)	
<i>rconde</i> :	An initialized vector or a NULL pointer. If not a null pointer, then the reciprocal condition numbers of the eigenvalues are stored here.	
<i>rcondv</i> :	An initialized vector or a NULL pointer. If not a null pointer, then the reciprocal condition numbers of the right eigenvectors are stored here.	
<i>info</i> :	This argument is used for two purposes. As an input argument it gives whether an igraph error should be generated if the QR algorithm fails to compute all eigenvalues. If <i>info</i> is non-zero, then an error is generated, otherwise only a warning is given. On exit it contains the LAPACK error code. Zero means successful exit. A negative values means that some of the arguments had an illegal value, this always triggers an igraph error. An i positive value means that the QR algorithm failed to compute all the eigenvalues, and no eigenvectors have	

been computed; element $i+1:N$ of *valuesreal* and *valuesimag* contain eigenvalues which have converged. This case only generated an igraph error, if *info* was non-zero on entry.

Returns:

Error code.

Time complexity: TODO

Example 30.6. File `examples/simple/igraph_lapack_dgeevx.c`

ARPACK interface in igraph

ARPACK is a library for solving large scale eigenvalue problems. The package is designed to compute a few eigenvalues and corresponding eigenvectors of a general n by n matrix A . It is most appropriate for large sparse or structured matrices A where structured means that a matrix-vector product $w \leftarrow Av$ requires order n rather than the usual order n^2 floating point operations. Please see <http://www.caam.rice.edu/software/ARPACK/> for details.

The eigenvalue calculation in ARPACK (in the simplest case) involves the calculation of the Av product where A is the matrix we work with and v is an arbitrary vector. A user-defined function of type `igraph_arpack_function_t` is expected to perform this product. If the product can be done efficiently, e.g. if the matrix is sparse, then ARPACK is usually able to calculate the eigenvalues very quickly.

In igraph, eigenvalue/eigenvector calculations usually involve the following steps:

1. Initialization of an `igraph_arpack_options_t` data structure using `igraph_arpack_options_init`.
2. Setting some options in the initialized `igraph_arpack_options_t` object.
3. Defining a function of type `igraph_arpack_function_t`. The input of this function is a vector, and the output should be the output matrix multiplied by the input vector.
4. Calling `igraph_arpack_rssolve()` (is the matrix is symmetric), or `igraph_arpack_rnsolve()`.

The `igraph_arpack_options_t` object can be used multiple times.

If we have many eigenvalue problems to solve, then it might worth to create an `igraph_arpack_storage_t` object, and initialize it via `igraph_arpack_storage_init()`. This structure contains all memory needed for ARPACK (with the given upper limit regarding to the size of the eigenvalue problem). Then many problems can be solved using the same `igraph_arpack_storage_t` object, without always reallocating the required memory. The `igraph_arpack_storage_t` object needs to be destroyed by calling `igraph_arpack_storage_destroy()` on it, when it is not needed any more.

igraph does not contain all ARPACK routines, only the ones dealing with symmetric and non-symmetric eigenvalue problems using double precision real numbers.

Data structures

`igraph_arpack_options_t` — Options for ARPACK


```
typedef struct igraph_arpack_options_t {
    /* INPUT */
    char bmat[1];          /* I-standard problem, G-generalized */
    int n;                 /* Dimension of the eigenproblem */
    char which[2];         /* LA, SA, LM, SM, BE */
    int nev;               /* Number of eigenvalues to be computed */
    igraph_real_t tol;     /* Stopping criterion */
    int ncv;               /* Number of columns in V */
    int ldv;               /* Leading dimension of V */
    int ishift;            /* 0-reverse comm., 1-exact with tridiagonal */
    int mxiter;            /* Maximum number of update iterations to take */
    int nb;                /* Block size on the recurrence, only 1 works */
    int mode;              /* The kind of problem to be solved (1-5)
                           1: A*x=l*x, A symmetric
                           2: A*x=l*M*x, A symm. M pos. def.
                           3: K*x = l*M*x, K symm., M pos. semidef.
                           4: K*x = l*KG*x, K s. pos. semidef. KG s. indef.
                           5: A*x = l*M*x, A symm., M symm. pos. semidef. */
    int start;             /* 0: random, 1: use the supplied vector */
    int lworkl;            /* Size of temporary storage, default is fine */
    igraph_real_t sigma;   /* The shift for modes 3,4,5 */
    igraph_real_t sigmai;  /* The imaginary part of shift for rnsolve */
    /* OUTPUT */
    int info;              /* What happened, see docs */
    int ierr;              /* What happened in the dseupd call */
    int noiter;            /* The number of iterations taken */
    int nconv;
    int numop;             /* Number of OP*x operations */
    int numopb;            /* Number of B*x operations if BMAT='G' */
    int numreo;            /* Number of steps of re-orthogonalizations */
    /* INTERNAL */
    int iparam[11];
    int ipntr[14];
} igraph_arpack_options_t;
```

This data structure contains the options of the ARPACK eigenvalue solver routines. It must be initialized by calling `igraph_arpack_options_init()` on it. Then it can be used for multiple ARPACK calls, as the ARPACK solvers do not modify it. Input options:

Values:

bmat: Character. Whether to solve a standard ('I') or a generalized problem ('G').

n: Dimension of the eigenproblem.

which: Specifies which eigenvalues/vectors to compute. Possible values for symmetric matrices:

- LA Compute `nev` largest (algebraic) eigenvalues.
- SA Compute `nev` smallest (algebraic) eigenvalues.
- LM Compute `nev` largest (in magnitude) eigenvalues.
- SM Compute `nev` smallest (in magnitude) eigenvalues.
- BE Compute `nev` eigenvalues, half from each end of the spectrum. When `nev` is odd, compute one more from the high end than from the low end.

Possible values for non-symmetric matrices:

LM	Compute <code>nev</code> largest (in magnitude) eigenvalues.
SM	Compute <code>nev</code> smallest (in magnitude) eigenvalues.
LR	Compute <code>nev</code> eigenvalues of largest real part.
SR	Compute <code>nev</code> eigenvalues of smallest real part.
LI	Compute <code>nev</code> eigenvalues of largest imaginary part.
SI	Compute <code>nev</code> eigenvalues of smallest imaginary part.
<code>nev</code> :	The number of eigenvalues to be computed.
<code>tol</code> :	Stopping criterion: the relative accuracy of the Ritz value is considered acceptable if its error is less than <code>tol</code> times its estimated value. If this is set to zero then machine precision is used.
<code>ncv</code> :	Number of Lanczos vectors to be generated. Setting this to zero means that <code>igraph_arpack_rssolve</code> and <code>igraph_arpack_rnsolve</code> will determine a suitable value for <code>ncv</code> automatically.
<code>ldv</code> :	Numeric scalar. It should be set to zero in the current igraph implementation.
<code>ishift</code> :	Either zero or one. If zero then the shifts are provided by the user via reverse communication. If one then exact shifts with respect to the reduced tridiagonal matrix <code>T</code> . Please always set this to one.
<code>mxiter</code> :	Maximum number of Arnoldi update iterations allowed.
<code>nb</code> :	Blocksize to be used in the recurrence. Please always leave this on the default value, one.
<code>mode</code> :	The type of the eigenproblem to be solved. Possible values if the input matrix is symmetric: <ol style="list-style-type: none"> 1. $A*x = \lambda*x$, A is symmetric. 2. $A*x = \lambda*M*x$, A is symmetric, M is symmetric positive definite. 3. $K*x = \lambda*M*x$, K is symmetric, M is symmetric positive semi-definite. 4. $K*x = \lambda*KG*x$, K is symmetric positive semi-definite, KG is symmetric indefinite. 5. $A*x = \lambda*M*x$, A is symmetric, M is symmetric positive semi-definite. (Cayley transformed mode.) Please note that only <code>mode == 1</code> was tested and other values might not work properly. Possible values if the input matrix is not symmetric: <ol style="list-style-type: none"> 1. $A*x = \lambda*x$. 2. $A*x = \lambda*M*x$, M is symmetric positive definite. 3. $A*x = \lambda*M*x$, M is symmetric semi-definite. 4. $A*x = \lambda*M*x$, M is symmetric semi-definite. Please note that only <code>mode == 1</code> was tested and other values might not work properly.
<code>start</code> :	Whether to use the supplied starting vector (1), or use a random starting vector (0). The starting vector must be supplied in the first column of the <code>vectors</code> argument of the <code>igraph_arpack_rssolve()</code> or <code>igraph_arpack_rnsolve()</code> call.

Output options:

Values:

info: Error flag of ARPACK. Possible values:

- 0 Normal exit.
- 1 Maximum number of iterations taken.
- 3 No shifts could be applied during a cycle of the Implicitly restarted Arnoldi iteration. One possibility is to increase the size of `ncv` relative to `nev`. ARPACK can return other error flags as well, but these are converted to igraph errors, see `igraph_error_type_t`.

ierr: Error flag of the second ARPACK call (one eigenvalue computation usually involves two calls to ARPACK). This is always zero, as other error codes are converted to igraph errors.

noiter: Number of Arnoldi iterations taken.

nconv: Number of converged Ritz values. This represents the number of Ritz values that satisfy the convergence criterion.

numop: Total number of matrix-vector multiplications.

numopb: Not used currently.

numreo: Total number of steps of re-orthogonalization.

Internal options:

Values:

lworkl: Do not modify this option.

sigma: The shift for the shift-invert mode.

sigmai: The imaginary part of the shift, for the non-symmetric or complex shift-invert mode.

iparam: Do not modify this option.

ipntr: Do not modify this option.

igraph_arnpack_storage_t — Storage for ARPACK

```
typedef struct igraph_arnpack_storage_t {
    int maxn, maxncv, maxldv;
    igraph_real_t *v;
    igraph_real_t *workl;
    igraph_real_t *workd;
    igraph_real_t *d;
    igraph_real_t *resid;
    igraph_real_t *ax;
    int *select;
    igraph_real_t *di;           /* These two only for non-symmetric problems */
    igraph_real_t *workev;
```

```
} igraph_arnpack_storage_t;
```

Public members, do not modify them directly, these are considered to be read-only.

Values:

`maxn`: Maximum rank of matrix.

`maxncv`: Maximum NCV.

`maxldv`: Maximum LDV.

These members are considered to be private:

Values:

`workl`: Working memory.

`workd`: Working memory.

`d`: Memory for eigenvalues.

`resid`: Memory for residuals.

`ax`: Working memory.

`select`: Working memory.

`di`: Memory for eigenvalues, non-symmetric case only.

`workev`: Working memory, non-symmetric case only.

igraph_arnpack_function_t — Type of the ARPACK callback function

```
typedef int igraph_arnpack_function_t(igraph_real_t *to, const igraph_real_t *from,
                                      int n, void *extra);
```

Arguments:

to: Pointer to an `igraph_real_t`, the result of the matrix-vector product is expected to be stored here.

from: Pointer to an `igraph_real_t`, the input matrix should be multiplied by the vector stored here.

n: The length of the vector (which is the same as the order of the input matrix).

extra: Extra argument to the matrix-vector calculation function. This is coming from the `igraph_arnpack_rssolve()` or `igraph_arnpack_rnsolve()` function.

Returns:

Error code, if not zero, then the ARPACK solver considers this as an error, stops and calls the igraph error handler.

igraph_arpack_options_init — Initialize ARPACK options

```
void igraph_arpack_options_init(igraph_arpack_options_t *o);
```

Initializes ARPACK options, set them to default values. You can always pass the initialized `igraph_arpack_options_t` object to built-in igraph functions without any modification. The built-in igraph functions modify the options to perform their calculation, e.g. `igraph_pagerank()` always searches for the eigenvalue with the largest magnitude, regardless of the supplied value.

If you want to implement your own function involving eigenvalue calculation using ARPACK, however, you will likely need to set up the fields for yourself.

Arguments:

o: The `igraph_arpack_options_t` object to initialize.

Time complexity: $O(1)$.

igraph_arpack_storage_init — Initialize ARPACK storage

```
int igraph_arpack_storage_init(igraph_arpack_storage_t *s, long int maxn,  
                              long int maxncv, long int maxldv,  
                              igraph_bool_t symm);
```

You only need this function if you want to run multiple eigenvalue calculations using ARPACK, and want to spare the memory allocation/deallocation between each two runs. Otherwise it is safe to supply a null pointer as the `storage` argument of both `igraph_arpack_rssolve()` and `igraph_arpack_rnsolve()` to make memory allocated and deallocated automatically.

Don't forget to call the `igraph_arpack_storage_destroy()` function on the storage object if you don't need it any more.

Arguments:

s: The `igraph_arpack_storage_t` object to initialize.

maxn: The maximum order of the matrices.

maxncv: The maximum NCV parameter intended to use.

maxldv: The maximum LDV parameter intended to use.

symm: Whether symmetric or non-symmetric problems will be solved using this `igraph_arpack_storage_t`. (You cannot use the same storage both with symmetric and non-symmetric solvers.)

Returns:

Error code.

Time complexity: $O(\text{maxncv} * (\text{maxldv} + \text{maxn}))$.

igraph_arpack_storage_destroy — Deallocate ARPACK storage

```
void igraph_arpack_storage_destroy(igraph_arpack_storage_t *s);
```

Arguments:

s: The `igraph_arpack_storage_t` object for which the memory will be deallocated.

Time complexity: operating system dependent.

ARPACK solvers

`igraph_arpack_rssolve` — ARPACK solver for symmetric matrices

```
int igraph_arpack_rssolve(igraph_arpack_function_t *fun, void *extra,  
                           igraph_arpack_options_t *options,  
                           igraph_arpack_storage_t *storage,  
                           igraph_vector_t *values, igraph_matrix_t *vectors);
```

This is the ARPACK solver for symmetric matrices. Please use `igraph_arpack_rnsolve()` for non-symmetric matrices.

Arguments:

fun: Pointer to an `igraph_arpack_function_t` object, the function that performs the matrix-vector multiplication.

extra: An extra argument to be passed to *fun*.

options: An `igraph_arpack_options_t` object.

storage: An `igraph_arpack_storage_t` object, or a null pointer. In the latter case memory allocation and deallocation is performed automatically. Either this or the *vectors* argument must be non-null if the ARPACK iteration is started from a given starting vector. If both are given *vectors* take precedence.

values: If not a null pointer, then it should be a pointer to an initialized vector. The eigenvalues will be stored here. The vector will be resized as needed.

vectors: If not a null pointer, then it must be a pointer to an initialized matrix. The eigenvectors will be stored in the columns of the matrix. The matrix will be resized as needed. Either this or the *vectors* argument must be non-null if the ARPACK iteration is started from a given starting vector. If both are given *vectors* take precedence.

Returns:

Error code.

Time complexity: depends on the matrix-vector multiplication. Usually a small number of iterations is enough, so if the matrix is sparse and the matrix-vector multiplication can be done in $O(n)$ time (the number of vertices), then the eigenvalues are found in $O(n)$ time as well.

`igraph_arpack_rnsolve` — ARPACK solver for non-symmetric matrices

```
int igraph_arpack_rnsolve(igraph_arpack_function_t *fun, void *extra,
                          igraph_arpack_options_t *options,
                          igraph_arpack_storage_t *storage,
                          igraph_matrix_t *values, igraph_matrix_t *vectors);
```

Please always consider calling `igraph_arpack_rssolve()` if your matrix is symmetric, it is much faster. `igraph_arpack_rnsolve()` for non-symmetric matrices.

Note that ARPACK is not called for 2x2 matrices as an exact algebraic solution exists in these cases.

Arguments:

- fun*: Pointer to an `igraph_arpack_function_t` object, the function that performs the matrix-vector multiplication.
- extra*: An extra argument to be passed to *fun*.
- options*: An `igraph_arpack_options_t` object.
- storage*: An `igraph_arpack_storage_t` object, or a null pointer. In the latter case memory allocation and deallocation is performed automatically.
- values*: If not a null pointer, then it should be a pointer to an initialized matrix. The (possibly complex) eigenvalues will be stored here. The matrix will have two columns, the first column contains the real, the second the imaginary parts of the eigenvalues. The matrix will be resized as needed.
- vectors*: If not a null pointer, then it must be a pointer to an initialized matrix. The eigenvectors will be stored in the columns of the matrix. The matrix will be resized as needed. Note that real eigenvalues will have real eigenvectors in a single column in this matrix; however, complex eigenvalues come in conjugate pairs and the result matrix will store the eigenvector corresponding to the eigenvalue with *positive* imaginary part only. Since in this case the eigenvector is also complex, it will occupy *two* columns in the eigenvector matrix (the real and the imaginary parts, in this order). Caveat: if the eigenvalue vector returns only the eigenvalue with the *negative* imaginary part for a complex conjugate eigenvalue pair, the result vector will *still* store the eigenvector corresponding to the eigenvalue with the positive imaginary part (since this is how ARPACK works).

Returns:

Error code.

Time complexity: depends on the matrix-vector multiplication. Usually a small number of iterations is enough, so if the matrix is sparse and the matrix-vector multiplication can be done in $O(n)$ time (the number of vertices), then the eigenvalues are found in $O(n)$ time as well.

igraph_arpack_unpack_complex — Make the result of the non-symmetric ARPACK solver more readable

```
int igraph_arpack_unpack_complex(igraph_matrix_t *vectors, igraph_matrix_t *values,
                                long int nev);
```

This function works on the output of `igraph_arpack_rnsolve` and brushes it up a bit: it only keeps *nev* eigenvalues/vectors and every eigenvector is stored in two columns of the *vectors* matrix.

The output of the non-symmetric ARPACK solver is somewhat hard to parse, as real eigenvectors occupy only one column in the matrix, and the complex conjugate eigenvectors are not stored at all (usually). The other problem is that the solver might return more eigenvalues than requested. The common use of this function is to call it directly after `igraph_arpack_rnsolve` with its `vectors` and `values` argument and `options->nev` as `nev`. This will add the vectors for eigenvalues with a negative imaginary part and return all vectors as 2 columns, a real and imaginary part.

Arguments:

- vectors*: The eigenvector matrix, as returned by `igraph_arpack_rnsolve`. It will be resized, typically it will be larger.
- values*: The eigenvalue matrix, as returned by `igraph_arpack_rnsolve`. It will be resized, typically extra, unneeded rows (=eigenvalues) will be removed.
- nev*: The number of eigenvalues/vectors to keep. Can be less or equal than the number originally requested from ARPACK.

Returns:

Error code.

Time complexity: linear in the number of elements in the *vectors* matrix.

Chapter 31. Bipartite, i.e. two-mode graphs

Bipartite networks in igraph

A bipartite network contains two kinds of vertices and connections are only possible between two vertices of different kinds. There are many natural examples, e.g. movies and actors as vertices and a movie is connected to all participating actors, etc.

igraph does not have direct support for bipartite networks, at least not at the C language level. In other words the `igraph_t` structure does not contain information about the vertex types. The C functions for bipartite networks usually have an additional input argument to graph, called `types`, a boolean vector giving the vertex types.

Most functions creating bipartite networks are able to create this extra vector, you just need to supply an initialized boolean vector to them.

Create two-mode networks

`igraph_create_bipartite` — Create a bipartite graph.

```
int igraph_create_bipartite(igraph_t *graph, const igraph_vector_bool_t *types,
                           const igraph_vector_t *edges,
                           igraph_bool_t directed);
```

This is a simple wrapper function to create a bipartite graph. It does a little more than `igraph_create()`, e.g. it checks that the graph is indeed bipartite with respect to the given `types` vector. If there is an edge connecting two vertices of the same kind, then an error is reported.

Arguments:

- graph*: Pointer to an uninitialized graph object, the result is created here.
- types*: Boolean vector giving the vertex types. The length of the vector defines the number of vertices in the graph.
- edges*: Vector giving the edges of the graph. The highest vertex id in this vector must be smaller than the length of the *types* vector.
- directed*: Boolean scalar, whether to create a directed graph.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

Example 31.1. File `examples/simple/igraph_bipartite_create.c`

igraph_full_bipartite — Create a full bipartite network.

```
int igraph_full_bipartite(igraph_t *graph,
                          igraph_vector_bool_t *types,
                          igraph_integer_t n1, igraph_integer_t n2,
                          igraph_bool_t directed,
                          igraph_neimode_t mode);
```

A bipartite network contains two kinds of vertices and connections are only possible between two vertices of different kind. There are many natural examples, e.g. movies and actors as vertices and a movie is connected to all participating actors, etc.

igraph does not have direct support for bipartite networks, at least not at the C language level. In other words the `igraph_t` structure does not contain information about the vertex types. The C functions for bipartite networks usually have an additional input argument to `graph`, called `types`, a boolean vector giving the vertex types.

Most functions creating bipartite networks are able to create this extra vector, you just need to supply an initialized boolean vector to them.

Arguments:

graph: Pointer to an `igraph_t` object, the graph will be created here.

types: Pointer to a boolean vector. If not a null pointer, then the vertex types will be stored here.

n1: Integer, the number of vertices of the first kind.

n2: Integer, the number of vertices of the second kind.

directed: Boolean, whether to create a directed graph.

mode: A constant that gives the type of connections for directed graphs. If `IGRAPH_OUT`, then edges point from vertices of the first kind to vertices of the second kind; if `IGRAPH_IN`, then the opposite direction is realized; if `IGRAPH_ALL`, then mutual edges will be created.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

See also:

`igraph_full()` for non-bipartite full graphs.

igraph_bipartite_game — Generate a bipartite random graph (similar to Erdős-Rényi).

```
int igraph_bipartite_game(igraph_t *graph, igraph_vector_bool_t *types,
```

```
igraph_erdos_renyi_t type,  
igraph_integer_t n1, igraph_integer_t n2,  
igraph_real_t p, igraph_integer_t m,  
igraph_bool_t directed, igraph_neimode_t mode);
```

Similarly to unipartite (one-mode) networks, we can define the $G(n,p)$, and $G(n,m)$ graph classes for bipartite graphs, via their generating process. In $G(n,p)$ every possible edge between top and bottom vertices is realized with probability p , independently of the rest of the edges. In $G(n,m)$, we uniformly choose m edges to realize.

Arguments:

- graph*: Pointer to an uninitialized igraph graph, the result is stored here.
- types*: Pointer to an initialized boolean vector, or a null pointer. If not a null pointer, then the vertex types are stored here. Bottom vertices come first, $n1$ of them, then $n2$ top vertices.
- type*: The type of the random graph, possible values:
- | | |
|------------------------|---|
| IGRAPH_ERDOS_RENYI_GNM | $G(n,m)$ graph, m edges are selected uniformly randomly in a graph with n vertices. |
| IGRAPH_ERDOS_RENYI_GNP | $G(n,p)$ graph, every possible edge is included in the graph with probability p . |
- n1*: The number of bottom vertices.
- n2*: The number of top vertices.
- p*: The connection probability for $G(n,p)$ graphs. It is ignored for $G(n,m)$ graphs.
- m*: The number of edges for $G(n,m)$ graphs. It is ignored for $G(n,p)$ graphs.
- directed*: Boolean, whether to generate a directed graph. See also the *mode* argument.
- mode*: Specifies how to direct the edges in directed graphs. If it is `IGRAPH_OUT`, then directed edges point from bottom vertices to top vertices. If it is `IGRAPH_IN`, edges point from top vertices to bottom vertices. `IGRAPH_OUT` and `IGRAPH_IN` do not generate mutual edges. If this argument is `IGRAPH_ALL`, then each edge direction is considered independently and mutual edges might be generated. This argument is ignored for undirected graphs.

Returns:

Error code.

See also:

`igraph_erdos_renyi_game`.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

Incidence matrices

`igraph_incidence` — Creates a bipartite graph from an incidence matrix.

```
int igraph_incidence(igraph_t *graph, igraph_vector_bool_t *types,
                    const igraph_matrix_t *incidence,
                    igraph_bool_t directed,
                    igraph_neimode_t mode, igraph_bool_t multiple);
```

A bipartite (or two-mode) graph contains two types of vertices and edges always connect vertices of different types. An incidence matrix is an $n \times m$ matrix, n and m are the number of vertices of the two types, respectively. Nonzero elements in the matrix denote edges between the two corresponding vertices.

Note that this function can operate in two modes, depending on the *multiple* argument. If it is `FALSE` (i.e. 0), then a single edge is created for every non-zero element in the incidence matrix. If *multiple* is `TRUE` (i.e. 1), then the matrix elements are rounded up to the closest non-negative integer to get the number of edges to create between a pair of vertices.

This function does not create multiple edges if *multiple* is `FALSE`, but might create some if it is `TRUE`.

Arguments:

- graph*: Pointer to an uninitialized graph object.
- types*: Pointer to an initialized boolean vector, or a null pointer. If not a null pointer, then the vertex types are stored here. It is resized as needed.
- incidence*: The incidence matrix.
- directed*: Gives whether to create an undirected or a directed graph.
- mode*: Specifies the direction of the edges in a directed graph. If `IGRAPH_OUT`, then edges point from vertices of the first kind (corresponding to rows) to vertices of the second kind (corresponding to columns); if `IGRAPH_IN`, then the opposite direction is realized; if `IGRAPH_ALL`, then mutual edges will be created.
- multiple*: How to interpret the incidence matrix elements. See details below.

Returns:

Error code.

Time complexity: $O(n \cdot m)$, the size of the incidence matrix.

igraph_get_incidence — Convert a bipartite graph into an incidence matrix.

```
int igraph_get_incidence(const igraph_t *graph,
                        const igraph_vector_bool_t *types,
                        igraph_matrix_t *res,
                        igraph_vector_t *row_ids,
                        igraph_vector_t *col_ids);
```

Arguments:

- graph*: The input graph, edge directions are ignored.

types: Boolean vector containing the vertex types. All vertices in one part of the graph should have type 0, the others type 1.

res: Pointer to an initialized matrix, the result is stored here. An element of the matrix gives the number of edges (irrespectively of their direction) between the two corresponding vertices. The rows will correspond to vertices with type 0, the columns correspond to vertices with type 1.

row_ids: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex ids (in the graph) corresponding to the rows of the result matrix are stored here.

col_ids: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex ids corresponding to the columns of the result matrix are stored here.

Returns:

Error code.

Time complexity: $O(n*m)$, n and m are number of vertices of the two different kind.

See also:

`igraph_incidence()` for the opposite operation.

Project two-mode graphs

`igraph_bipartite_projection_size` — Calculate the number of vertices and edges in the bipartite projections.

```
int igraph_bipartite_projection_size(const igraph_t *graph,
                                     const igraph_vector_bool_t *types,
                                     igraph_integer_t *vcount1,
                                     igraph_integer_t *ecount1,
                                     igraph_integer_t *vcount2,
                                     igraph_integer_t *ecount2);
```

This function calculates the number of vertices and edges in the two projections of a bipartite network. This is useful if you have a big bipartite network and you want to estimate the amount of memory you would need to calculate the projections themselves.

Arguments:

graph: The input graph.

types: Boolean vector giving the vertex types of the graph.

vcount1: Pointer to an `igraph_integer_t`, the number of vertices in the first projection is stored here.

ecount1: Pointer to an `igraph_integer_t`, the number of edges in the first projection is stored here.

vcount2: Pointer to an `igraph_integer_t`, the number of vertices in the second projection is stored here.

ecount2: Pointer to an `igraph_integer_t`, the number of edges in the second projection is stored here.

Returns:

Error code.

See also:

`igraph_bipartite_projection()` to calculate the actual projection.

Time complexity: $O(|V|*d^2+|E|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges, d is the average (total) degree of the graphs.

Example **31.2.** **File** **examples/simple/
igraph_bipartite_projection.c**

igraph_bipartite_projection — Create one or both projections of a bipartite (two-mode) network.

```
int igraph_bipartite_projection(const igraph_t *graph,
                               const igraph_vector_bool_t *types,
                               igraph_t *proj1,
                               igraph_t *proj2,
                               igraph_vector_t *multiplicity1,
                               igraph_vector_t *multiplicity2,
                               igraph_integer_t probe1);
```

Creates one or both projections of a bipartite graph.

Arguments:

graph: The bipartite input graph. Directedness of the edges is ignored.

types: Boolean vector giving the vertex types of the graph.

proj1: Pointer to an uninitialized graph object, the first projection will be created here. If a null pointer, then it is ignored, see also the *probe1* argument.

proj2: Pointer to an uninitialized graph object, the second projection is created here, if it is not a null pointer. See also the *probe1* argument.

multiplicity1: Pointer to a vector, or a null pointer. If not the latter, then the multiplicity of the edges is stored here. E.g. if there is an A-C-B and also an A-D-B triple in the bipartite graph (but no more X, such that A-X-B is also in the graph), then the multiplicity of the A-B edge in the projection will be 2.

multiplicity2: The same as *multiplicity1*, but for the other projection.

probe1: This argument can be used to specify the order of the projections in the resulting list. When it is non-negative, then it is considered as a vertex ID and the projection containing this vertex will be the first one in the result. Setting this argument to a non-negative value implies that *proj1* must be a non-null pointer. If you don't care about the ordering of the projections, pass -1 here.

Returns:

Error code.

See also:

`igraph_bipartite_projection_size()` to calculate the number of vertices and edges in the projections, without creating the projection graphs themselves.

Time complexity: $O(|V|*d^2+|E|)$, $|V|$ is the number of vertices, $|E|$ is the number of edges, d is the average (total) degree of the graphs.

Example	31.3.	File	examples/simple/
igraph_bipartite_projection.c			

Other operations on bipartite graphs

igraph_is_bipartite — Check whether a graph is bipartite.

```
int igraph_is_bipartite(const igraph_t *graph,
                        igraph_bool_t *res,
                        igraph_vector_bool_t *types);
```

This function checks whether a graph is bipartite. It tries to find a mapping that gives a possible division of the vertices into two classes, such that no two vertices of the same class are connected by an edge.

The existence of such a mapping is equivalent of having no circuits of odd length in the graph. A graph with loop edges cannot be bipartite.

Note that the mapping is not necessarily unique, e.g. if the graph has at least two components, then the vertices in the separate components can be mapped independently.

Arguments:

graph: The input graph.

res: Pointer to a boolean, the result is stored here.

types: Pointer to an initialized boolean vector, or a null pointer. If not a null pointer and a mapping was found, then it is stored here. If not a null pointer, but no mapping was found, the contents of this vector is invalid.

Returns:

Error code.

Time complexity: $O(|V|+|E|)$, linear in the number of vertices and edges.

Chapter 32. Advanced igraph programming

Using igraph in multi-threaded programs

The igraph library is considered thread-safe if it has been compiled with thread-local storage enabled, i.e. the `IGRAPH_ENABLE_TLS` setting was toggled to ON and the current platform supports this feature. To check whether an igraph build is thread-safe, use the `IGRAPH_THREAD_SAFE` macro. When linking to external versions of igraph's dependencies, it is the responsibility of the user to check that these dependencies were also compiled to be thread-safe.

IGRAPH_THREAD_SAFE — Specifies whether igraph was built in thread-safe mode.

```
#define IGRAPH_THREAD_SAFE
```

This macro is defined to 1 if the current build of the igraph library is built in thread-safe mode, and 0 if it is not. A thread-safe igraph library attempts to use thread-local data structures instead of global ones, but note that this is not (and can not) be guaranteed for third-party libraries that igraph links to.

Thread-safe ARPACK library

Note that igraph is only thread-safe if it was built with the internal ARPACK library, i.e. the one that comes with igraph. The standard ARPACK library is not thread-safe.

Thread-safety of random number generators

The default random number generator that igraph uses is *not* guaranteed to be thread-safe. You need to set a different random number generator instance for every thread that you want to use igraph from. This is especially important if you set the seed of the random number generator to ensure reproducibility; sharing a random number generator between threads would break reproducibility as the order in which the various threads are scheduled is random, and therefore they would still receive random numbers in an unpredictable order from the shared random number generator.

Progress handlers

About progress handlers

It is often useful to report the progress of some long calculation, to allow the user to follow the computation and guess the total running time. A couple of igraph functions support this at the time of writing, hopefully more will support it in the future.

To see the progress of a computation, the user has to install a progress handler, as there is none installed by default. If an igraph function supports progress reporting, then it calls the installed progress handler periodically, and passes a percentage value to it, the percentage of computation already performed. To install a progress handler, you need to call `igraph_set_progress_handler()`. Currently there is a single pre-defined progress handler, called `igraph_progress_handler_stderr()`.

Setting up progress handlers

igraph_progress_handler_t — Type of progress handler functions

```
typedef int igraph_progress_handler_t(const char *message, igraph_real_t percent, void *data);
```

This is the type of the igraph progress handler functions. There is currently one such predefined function, `igraph_progress_handler_stderr()`, but the user can write and set up more sophisticated ones.

Arguments:

message: A string describing the function or algorithm that is reporting the progress. Current igraph functions always use the name *message* argument if reporting from the same function.

percent: Numeric, the percentage that was completed by the algorithm or function.

data: User-defined data. Current igraph functions that report progress pass a null pointer here. Users can write their own progress handlers and functions with progress reporting, and then pass some meaningful context here.

Returns:

If the return value of the progress handler is not `IGRAPH_SUCCESS`, then `igraph_progress()` returns the error code `IGRAPH_INTERRUPTED`. The `IGRAPH_PROGRESS()` macro frees all memory and finishes the igraph function with error code `IGRAPH_INTERRUPTED` in this case.

igraph_set_progress_handler — .

```
igraph_progress_handler_t *  
igraph_set_progress_handler(igraph_progress_handler_t new_handler);
```

Install a progress handler, or remove the current handler. There is a single simple predefined progress handler: `igraph_progress_handler_stderr()`.

Arguments:

new_handler: Pointer to a function of type `igraph_progress_handler_t`, the progress handler function to install. To uninstall the current progress handler, this argument can be a null pointer.

Returns:

Pointer to the previously installed progress handler function.

Time complexity: $O(1)$.

igraph_progress_handler_stderr — A simple predefined progress handler.

```
int igraph_progress_handler_stderr(const char *message, igraph_real_t percent,
                                   void* data);
```

This simple progress handler first prints *message*, and then the percentage complete value in a short message to standard error.

Arguments:

message: A string describing the function or algorithm that is reporting the progress. Current igraph functions always use the same *message* argument if reporting from the same function.

percent: Numeric, the percentage that was completed by the algorithm or function.

data: User-defined data. Current igraph functions that report progress pass a null pointer here. Users can write their own progress handlers and functions with progress reporting, and then pass some meaningful context here.

Returns:

This function always returns with `IGRAPH_SUCCESS`.

Time complexity: $O(1)$.

Invoking the progress handler

IGRAPH_PROGRESS — Report progress.

```
#define IGRAPH_PROGRESS(message, percent, data)
```

The standard way to report progress from an igraph function

Arguments:

message: A string, a textual message that references the calculation under progress.

percent: Numeric scalar, the percentage that is complete.

data: User-defined data, this can be used in user-defined progress handler functions, from user-written igraph functions.

Returns:

If the progress handler returns with `IGRAPH_INTERRUPTED`, then this macro frees up the igraph allocated memory for temporary data and returns to the caller with `IGRAPH_INTERRUPTED`.

igraph_progress — Report progress

```
int igraph_progress(const char *message, igraph_real_t percent, void *data);
```

Note that the usual way to report progress is the `IGRAPH_PROGRESS` macro, as that takes care of the return value of the progress handler.

Arguments:

message: A string describing the function or algorithm that is reporting the progress. Current igraph functions always use the name *message* argument if reporting from the same function.

percent: Numeric, the percentage that was completed by the algorithm or function.

data: User-defined data. Current igraph functions that report progress pass a null pointer here. Users can write their own progress handlers and functions with progress reporting, and then pass some meaningful context here.

Returns:

If there is a progress handler installed and it does not return `IGRAPH_SUCCESS`, then `IGRAPH_INTERRUPTED` is returned.

Time complexity: $O(1)$.

igraph_progressf — Report progress, printf-like version

```
int igraph_progressf(const char *message, igraph_real_t percent, void *data,
...);
```

This is a more flexible version of `igraph_progress()`, with a printf-like template string. First the template string is filled with the additional arguments and then `igraph_progress()` is called.

Note that there is an upper limit for the length of the *message* string, currently 1000 characters.

Arguments:

message: A string describing the function or algorithm that is reporting the progress. For this function this is a template string, using the same syntax as the standard `libc printf` function.

percent: Numeric, the percentage that was completed by the algorithm or function.

data: User-defined data. Current igraph functions that report progress pass a null pointer here. Users can write their own progress handlers and functions with progress reporting, and then pass some meaningful context here.

...: Additional argument that were specified in the *message* argument.

Returns:

If there is a progress handler installed and it does not return `IGRAPH_SUCCESS`, then `IGRAPH_INTERRUPTED` is returned. \return

Writing progress handlers

To write a new progress handler, one needs to create a function of type `igraph_progress_handler_t`. The new progress handler can then be installed with the `igraph_set_progress_handler()` function.

One can assume that the first progress handler call from a calculation will be call with zero as the *percentage* argument, and the last call from a function will have 100 as the *percentage* argu-

ment. Note, however, that if an error happens in the middle of a computation, then the 100 percent call might be omitted.

Writing igraph functions with progress reporting

If you want to write a function that uses igraph and supports progress reporting, you need to include `igraph_progress()` calls in your function, usually via the `IGRAPH_PROGRESS()` macro.

It is good practice to always include a call to `igraph_progress()` with a zero *percentage* argument, before the computation; and another call with 100 *percentage* value after the computation is completed.

It is also good practice *not* to call `igraph_progress()` too often, as this would slow down the computation. It might not be worth to support progress reporting in functions with linear or log-linear time complexity, as these are fast, even with a large amount of data. For functions with quadratic or higher time complexity make sure that the time complexity of the progress reporting is constant or at least linear. In practice this means having at most $O(n)$ progress checks and at most 100 `igraph_progress()` calls.

Multi-threaded programs

In multi-threaded programs, each thread has its own progress handler, if thread-local storage is supported and igraph is thread-safe. See the `IGRAPH_THREAD_SAFE` macro for checking whether an igraph build is thread-safe.

Status handlers

Status reporting

In addition to the possibility of reporting the progress of an igraph computation via `igraph_progress()`, it is also possible to report simple status messages from within igraph functions, without having to judge how much of the computation was performed already. For this one needs to install a status handler function.

Status handler functions must be of type `igraph_status_handler_t` and they can be installed by a call to `igraph_set_status_handler()`. Currently there is a simple predefined status handler function, called `igraph_status_handler_stderr()`, but the user can define new ones.

igraph functions report their status via a call to the `IGRAPH_STATUS()` or the `IGRAPH_STATUSF()` macro.

Setting up status handlers

`igraph_status_handler_t` — The type of the igraph status handler functions

```
typedef int igraph_status_handler_t(const char *message, void *data);
```

Arguments:

message: The status message.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null pointer here.

igraph_set_status_handler — Install or uninstall a status handler function.

```
igraph_status_handler_t *  
igraph_set_status_handler(igraph_status_handler_t new_handler);
```

To uninstall the currently installed status handler, call this function with a null pointer.

Arguments:

new_handler: The status handler function to install.

Returns:

The previously installed status handler function.

Time complexity: $O(1)$.

igraph_status_handler_stderr — A simple predefined status handler function.

```
int igraph_status_handler_stderr(const char *message, void *data);
```

A simple status handler function, that writes the status message to the standard error.

Arguments:

message: The status message.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null pointer here.

Returns:

Error code.

Time complexity: $O(1)$.

Invoking the status handler

IGRAPH_STATUS — Report the status of an igraph function.

```
#define IGRAPH_STATUS(message, data)
```

Typically this function is called only a handful of times from an igraph function. E.g. if an algorithm has three major steps, then it is logical to call it three times, to signal the three major steps.

Arguments:

message: The status message.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null pointer here.

Returns:

If the status handler returns with a value other than `IGRAPH_SUCCESS`, then the function that called this macro returns as well, with error code `IGRAPH_INTERRUPTED`.

IGRAPH_STATUSF — Report the status from an igraph function

```
#define IGRAPH_STATUSF(args)
```

This is the more flexible version of `IGRAPH_STATUS()`, having a printf-like syntax. As this macro takes variable number of arguments, they must be all supplied as a single argument, enclosed in parentheses. Then `igraph_statusf()` is called with the given arguments.

Arguments:

args: The arguments to pass to `igraph_statusf()`.

Returns:

If the status handler returns with a value other than `IGRAPH_SUCCESS`, then the function that called this macro returns as well, with error code `IGRAPH_INTERRUPTED`.

igraph_status — Report status from an igraph function.

```
int igraph_status(const char *message, void *data);
```

It calls the installed status handler function, if there is one. Otherwise it does nothing. Note that the standard way to report the status from an igraph function is the `IGRAPH_STATUS` or `IGRAPH_STATUSF` macro, as these take care of the termination of the calling function if the status handler returns with `IGRAPH_INTERRUPTED`.

Arguments:

message: The status message.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null pointer here.

Returns:

Error code. If a status handler function was called and it did not return with `IGRAPH_SUCCESS`, then `IGRAPH_INTERRUPTED` is returned by `igraph_status()`.

Time complexity: $O(1)$.

igraph_statusf — Report status, more flexible printf-like version.

```
int igraph_statusf(const char *message, void *data, ...);
```

This is the more flexible version of `igraph_status()`, that has a syntax similar to the `printf` standard C library function. It substitutes the values of the additional arguments into the *message* template string and calls `igraph_status()`.

Arguments:

message: Status message template string, the syntax is the same as for the `printf` function.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null pointer here.

...: The additional arguments to fill the template given in the *message* argument.

Returns:

Error code. If a status handler function was called and it did not return with `IGRAPH_SUCCESS`, then `IGRAPH_INTERRUPTED` is returned by `igraph_status()`.

Chapter 33. Non-graph related functions

igraph version number

igraph_version — Return the version of the igraph C library

```
int igraph_version(const char **version_string,
                  int *major,
                  int *minor,
                  int *subminor);
```

Arguments:

version_string: Pointer to a string pointer. If not null, it is set to the igraph version string, e.g. "0.6" or "0.5.3". This string should not be modified or deallocated.

major: If not a null pointer, then it is set to the major igraph version. E.g. for version "0.5.3" this is 0.

minor: If not a null pointer, then it is set to the minor igraph version. E.g. for version "0.5.3" this is 5.

subminor: If not a null pointer, then it is set to the subminor igraph version. E.g. for version "0.5.3" this is 3.

Returns:

Error code.

Time complexity: $O(1)$.

Example 33.1. File `examples/simple/igraph_version.c`

Running mean of a time series

igraph_running_mean — Calculates the running mean of a vector.

```
int igraph_running_mean(const igraph_vector_t *data, igraph_vector_t *res,
                       igraph_integer_t binwidth);
```

The running mean is defined by the mean of the previous *binwidth* values.

Arguments:

data: The vector containing the data.

res: The vector containing the result. This should be initialized before calling this function and will be resized.

binwidth: Integer giving the width of the bin for the running mean calculation.

Returns:

Error code.

Time complexity: $O(n)$, n is the length of the data vector.

Random sampling from very long sequences

`igraph_random_sample` — Generates an increasing random sequence of integers.

```
int igraph_random_sample(igraph_vector_t *res, igraph_real_t l, igraph_real_t h,
                          igraph_integer_t length);
```

This function generates an increasing sequence of random integer numbers from a given interval. The algorithm is taken literally from (Vitter 1987). This method can be used for generating numbers from a *very* large interval. It is primarily created for randomly selecting some edges from the sometimes huge set of possible edges in a large graph.

Note that the type of the lower and the upper limit is `igraph_real_t`, not `igraph_integer_t`. This does not mean that you can pass fractional numbers there; these values must still be integral, but we need the longer range of `igraph_real_t` in several places in the library (for instance, when generating Erdos-Renyi graphs).

Arguments:

res: Pointer to an initialized vector. This will hold the result. It will be resized to the proper size.

l: The lower limit of the generation interval (inclusive). This must be less than or equal to the upper limit, and it must be integral. Passing a fractional number here results in undefined behaviour.

h: The upper limit of the generation interval (inclusive). This must be greater than or equal to the lower limit, and it must be integral. Passing a fractional number here results in undefined behaviour.

length: The number of random integers to generate.

Returns:

The error code `IGRAPH_EINVAL` is returned in each of the following cases: (1) The given lower limit is greater than the given upper limit, i.e. $l > h$. (2) Assuming that $l < h$ and N is the sample size, the above error code is returned if $N > |h - l|$, i.e. the sample size exceeds the size of the candidate pool.

Time complexity: according to (Vitter 1987), the expected running time is $O(\text{length})$.

Reference:

(Vitter 1987) J. S. Vitter. An efficient algorithm for sequential random sampling. *ACM Transactions on Mathematical Software*, 13(1):58--67, 1987.

Example 33.2. File `examples/simple/igraph_random_sample.c`

Random sampling of spatial points

`igraph_sample_sphere_surface` — Sample points uniformly from the surface of a sphere

```
int igraph_sample_sphere_surface(igraph_integer_t dim, igraph_integer_t n,
                                igraph_real_t radius,
                                igraph_bool_t positive,
                                igraph_matrix_t *res);
```

The center of the sphere is at the origin.

Arguments:

dim: The dimension of the random vectors.

n: The number of vectors to sample.

radius: Radius of the sphere, it must be positive.

positive: Whether to restrict sampling to the positive orthant.

res: Pointer to an initialized matrix, the result is stored here, each column will be a sampled vector. The matrix is resized, as needed.

Returns:

Error code.

Time complexity: $O(n \cdot \text{dim} \cdot g)$, where g is the time complexity of generating a standard normal random number.

See also:

`igraph_sample_sphere_volume()`, `igraph_sample_dirichlet()` for other similar samplers.

`igraph_sample_sphere_volume` — Sample points uniformly from the volume of a sphere

```
int igraph_sample_sphere_volume(igraph_integer_t dim, igraph_integer_t n,
                                igraph_real_t radius,
                                igraph_bool_t positive,
                                igraph_matrix_t *res);
```

The center of the sphere is at the origin.

Arguments:

dim: The dimension of the random vectors.

n: The number of vectors to sample.

radius: Radius of the sphere, it must be positive.

positive: Whether to restrict sampling to the positive orthant.

res: Pointer to an initialized matrix, the result is stored here, each column will be a sampled vector. The matrix is resized, as needed.

Returns:

Error code.

Time complexity: $O(n \cdot \text{dim} \cdot g)$, where g is the time complexity of generating a standard normal random number.

See also:

`igraph_sample_sphere_surface()`, `igraph_sample_dirichlet()` for other similar samplers.

`igraph_sample_dirichlet` — Sample points from a Dirichlet distribution

```
int igraph_sample_dirichlet(igraph_integer_t n, const igraph_vector_t *alpha,
                             igraph_matrix_t *res);
```

Arguments:

n: The number of vectors to sample.

alpha: The parameters of the Dirichlet distribution. They must be positive. The length of this vector gives the dimension of the generated samples.

res: Pointer to an initialized matrix, the result is stored here, one sample in each column. It will be resized, as needed.

Returns:

Error code.

Time complexity: $O(n \cdot \text{dim} \cdot g)$, where dim is the dimension of the sample vectors, set by the length of α , and g is the time complexity of sampling from a Gamma distribution.

See also:

`igraph_sample_sphere_surface()` and `igraph_sample_sphere_volume()` for other methods to sample latent vectors.

Convex hull of a set of points on a plane

igraph_convex_hull — Determines the convex hull of a given set of points in the 2D plane

```
int igraph_convex_hull(const igraph_matrix_t *data, igraph_vector_t *resverts,
                      igraph_matrix_t *rescoords);
```

The convex hull is determined by the Graham scan algorithm. See the following reference for details:

Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. Introduction to Algorithms, Second Edition. MIT Press and McGraw-Hill, 2001. ISBN 0262032937. Pages 949-955 of section 33.3: Finding the convex hull.

Arguments:

data: vector containing the coordinates. The length of the vector must be even, since it contains X-Y coordinate pairs.

resverts: the vector containing the result, e.g. the vector of vertex indices used as the corners of the convex hull. Supply NULL here if you are only interested in the coordinates of the convex hull corners.

rescoords: the matrix containing the coordinates of the selected corner vertices. Supply NULL here if you are only interested in the vertex indices.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory

Time complexity: $O(n \log(n))$ where n is the number of vertices

Example 33.3. File `examples/simple/igraph_convex_hull.c`

Fitting power-law distributions to empirical data

igraph_plfit_result_t — Result of fitting a power-law distribution to a vector

```
typedef struct igraph_plfit_result_t {
    igraph_bool_t continuous;
    double alpha;
    double xmin;
    double L;
    double D;
    double p;
} igraph_plfit_result_t;
```

This data structure contains the result of `igraph_power_law_fit()`, which tries to fit a power-law distribution to a vector of numbers. The structure contains the following members:

Values:

<code>continuous</code> :	Whether the fitted power-law distribution was continuous or discrete.
<code>alpha</code> :	The exponent of the fitted power-law distribution.
<code>xmin</code> :	The minimum value from which the power-law distribution was fitted. In other words, only the values larger than <code>xmin</code> were used from the input vector.
<code>L</code> :	The log-likelihood of the fitted parameters; in other words, the probability of observing the input vector given the parameters.
<code>D</code> :	The test statistic of a Kolmogorov-Smirnov test that compares the fitted distribution with the input vector. Smaller scores denote better fit.
<code>p</code> :	The p-value of the Kolmogorov-Smirnov test. Small p-values (less than 0.05) indicate that the test rejected the hypothesis that the original data could have been drawn from the fitted power-law distribution.

`igraph_power_law_fit` — Fits a power-law distribution to a vector of numbers

```
int igraph_power_law_fit(const igraph_vector_t* data, igraph_plfit_result_t* result,
                        igraph_real_t xmin, igraph_bool_t force_continuous);
```

This function fits a power-law distribution to a vector containing samples from a distribution (that is assumed to follow a power-law of course). In a power-law distribution, it is generally assumed that $P(X=x)$ is proportional to $x^{-\alpha}$, where x is a positive number and α is greater than 1. In many real-world cases, the power-law behaviour kicks in only above a threshold value *xmin*. The goal of this functions is to determine *alpha* if *xmin* is given, or to determine *xmin* and the corresponding value of *alpha*.

The function uses the maximum likelihood principle to determine *alpha* for a given *xmin*; in other words, the function will return the *alpha* value for which the probability of drawing the given sample is the highest. When *xmin* is not given in advance, the algorithm will attempt to find the optimal *xmin* value for which the p-value of a Kolmogorov-Smirnov test between the fitted distribution and the original sample is the largest. The function uses the method of Clauset, Shalizi and Newman to calculate the parameters of the fitted distribution. See the following reference for details:

Aaron Clauset, Cosma R. Shalizi and Mark E.J. Newman: Power-law distributions in empirical data. SIAM Review 51(4):661-703, 2009.

Arguments:

<code>data</code> :	vector containing the samples for which a power-law distribution is to be fitted. Note that you have to provide the <i>samples</i> , not the probability density function or the cumulative distribution function. For example, if you wish to fit a power-law to the degrees of a graph, you can use the output of <code>igraph_degree</code> directly as an input argument to <code>igraph_power_law_fit</code>
<code>result</code> :	the result of the fitting algorithm. See <code>igraph_plfit_result_t</code> for more details.

xmin: the minimum value in the sample vector where the power-law behaviour is expected to kick in. Samples smaller than *xmin* will be ignored by the algorithm. Pass zero here if you want to include all the samples. If *xmin* is negative, the algorithm will attempt to determine its best value automatically.

force_continuous: assume that the samples in the *data* argument come from a continuous distribution even if the sample vector contains integer values only (by chance). If this argument is false, *igraph* will assume a continuous distribution if at least one sample is non-integer and assume a discrete distribution otherwise.

Returns:

Error code: `IGRAPH_ENOMEM`: not enough memory `IGRAPH_EINVAL`: one of the arguments is invalid `IGRAPH_EOVERFLOW`: overflow during the fitting process `IGRAPH_EUNDERFLOW`: underflow during the fitting process `IGRAPH_FAILURE`: the underlying algorithm signaled a failure without returning a more specific error code

Time complexity: in the continuous case, $O(n \log(n))$ if *xmin* is given. In the discrete case, the time complexity is dominated by the complexity of the underlying L-BFGS algorithm that is used to optimize α . If *xmin* is not given, the time complexity is multiplied by the number of unique samples in the input vector (although it should be faster in practice).

Example 33.4. File `examples/simple/igraph_power_law_fit.c`

Chapter 34. Licenses for igraph and this manual

THE GNU GENERAL PUBLIC LICENSE

Copyright © 1989, 1991 Free Software Foundation, Inc. 51 Franklin St, Fifth Floor, Boston, MA 02110-1301 USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software--to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation's software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Library General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps: (1) copyright the software, and (2) offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author's protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors' reputations.

Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone's free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

GNU GENERAL PUBLIC LICENSE

TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

0. This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The "Program", below, refers to any such program or work, and a "work based on the Program" means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term "modification".) Each licensee is addressed as "you".

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

1. You may copy and distribute verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

2. You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of Section 1 above, provided that you also meet all of these conditions:

- a. You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.
- b. You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.
- c. If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you provide a warranty) and that users may redistribute the program under these conditions, and telling the user how to view a copy of this License. (Exception: if the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

3. You may copy and distribute the Program (or a work based on it, under Section 2) in object code or executable form under the terms of Sections 1 and 2 above provided that you also do one of the following:

- a. Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- b. Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,

- c. Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.

4. You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

5. You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Program or works based on it.

6. Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.

7. If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

8. If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may

add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.

9. The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and "any later version", you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

10. If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

NO WARRANTY

11. BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

12. IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the "copyright" line and a pointer to where the full notice is found.

<one line to give the program's name and a brief idea of what it does.>
Copyright (C) <year> <name of author>

This program is free software; you can redistribute it and/or modify
it under the terms of the GNU General Public License as published by

the Free Software Foundation; either version 2 of the License, or
(at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
GNU General Public License for more details.

You should have received a copy of the GNU General Public License
along with this program; if not, write to the Free Software
Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like this when it starts in an interactive mode:

```
Gnomovision version 69, Copyright (C) year name of author
Gnomovision comes with ABSOLUTELY NO WARRANTY; for details type `show w'.
This is free software, and you are welcome to redistribute it
under certain conditions; type `show c' for details.
```

The hypothetical commands ``show w'` and ``show c'` should show the appropriate parts of the General Public License. Of course, the commands you use may be called something other than ``show w'` and ``show c'`; they could even be mouse-clicks or menu items--whatever suits your program.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a "copyright disclaimer" for the program, if necessary. Here is a sample; alter the names:

```
Yoyodyne, Inc., hereby disclaims all copyright interest in the program
`Gnomovision' (which makes passes at compilers) written by James Hacker.
```

```
<signature of Ty Coon>, 1 April 1989
Ty Coon, President of Vice
```

This General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Library General Public License instead of this License.

The GNU Free Documentation License

Copyright © 2000, 2001, 2002 Free Software Foundation, Inc. 51 Franklin St, Fifth Floor, Boston, MA 02110-1301 USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document "free" in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or noncommercially. Secondly, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of "copyleft", which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The "Document", below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as "you". You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A "Modified Version" of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A "Secondary Section" is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document's overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The "Invariant Sections" are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The "Cover Texts" are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A "Transparent" copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not "Transparent" is called "Opaque".

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The "Title Page" means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, "Title Page" means the text near the most prominent appearance of the work's title, preceding the beginning of the body of the text.

A section "Entitled XYZ" means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ

stands for a specific section name mentioned below, such as "Acknowledgements", "Dedications", "Endorsements", or "History".) To "Preserve the Title" of such a section when you modify the Document means that it remains a section "Entitled XYZ" according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

- A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section

of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.

- B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
- C. State on the Title page the name of the publisher of the Modified Version, as the publisher.
- D. Preserve all the copyright notices of the Document.
- E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
- F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
- G. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.
- H. Include an unaltered copy of this License.
- I. Preserve the section Entitled "History", Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
- J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
- K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.
- L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
- M. Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.
- N. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.
- O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties--for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by)

any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled "History" in the various original documents, forming one section Entitled "History"; likewise combine any sections Entitled "Acknowledgements", and any sections Entitled "Dedications". You must delete all sections Entitled "Endorsements".

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an "aggregate" if the copyright resulting from the compilation is not used to limit the legal rights of the compilation's users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document's Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in

addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled "Acknowledgements", "Dedications", or "History", the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided for under this License. Any other attempt to copy, modify, sublicense or distribute the Document is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <http://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License "or any later version" applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation.

G.1.1 ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

Copyright (c) YEAR YOUR NAME.
Permission is granted to copy, distribute and/or modify this document
under the terms of the GNU Free Documentation License, Version 1.2
or any later version published by the Free Software Foundation;
with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts.
A copy of the license is included in the section entitled "GNU
Free Documentation License".

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the "with...Texts." line with this:

with the Invariant Sections being LIST THEIR TITLES, with the
Front-Cover Texts being LIST, and with the Back-Cover Texts being LIST.

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

Index

A

add_edge, 24
add_edges, 25
add_vertices, 25
adhesion, 485
adjacency, 175
adjacency_spectral_embedding, 538
adjacent_triangles, 432
adjlist, 177
adjlist_clear, 155
adjlist_destroy, 154
adjlist_get, 154
adjlist_init, 152
adjlist_init_complementer, 153
adjlist_init_empty, 153
adjlist_simplify, 155
adjlist_size, 154
adjlist_sort, 155
all_minimal_st_separators, 489
all_st_cuts, 477
all_st_mincuts, 478
are_connected, 286
arpack_function_t, 564
arpack_options_init, 565
arpack_options_t, 560
arpack_rnsolve, 566
arpack_rssolve, 566
arpack_storage_destroy, 565
arpack_storage_init, 565
arpack_storage_t, 563
arpack_unpack_complex, 567
articulation_points, 319
ASSERT, 40
assortativity, 369
assortativity_degree, 370
assortativity_nominal, 368
asymmetric_preference_game, 205
atlas, 188
attribute_combination_add, 254
attribute_combination_destroy, 255
attribute_combination_init, 254
attribute_combination_remove, 255
attribute_combination_type_t, 255
attribute_table_t, 250
attribute_type_t, 253
authority_score, 336
automorphisms, 412
automorphism_group, 413
average_local_efficiency, 307
average_path_length, 300
average_path_length_dijkstra, 301
avg_nearest_neighbor_degree, 381

B

barabasi_aging_game, 207
barabasi_game, 191
betweenness, 326
betweenness_cutoff, 339
bfs, 387
bfshandler_t, 389
bfs_simple, 388
bibcoupling, 349
biconnected_components, 318
bipartite_game, 570
bipartite_projection, 574
bipartite_projection_size, 573
blas_ddot, 552
blas_dgemv, 552
blas_dgemv_array, 553
bliss_info_t, 410
bliss_sh_t, 410
bridges, 319

C

callaway_traits_game, 202
canonical_permutation, 411
cattribute_EAB, 267
cattribute_EABV, 267
cattribute_EAB_set, 275
cattribute_EAB_setv, 281
cattribute_EAN, 265
cattribute_EANV, 266
cattribute_EAN_set, 275
cattribute_EAN_setv, 280
cattribute_EAS, 268
cattribute_EASV, 269
cattribute_EAS_set, 276
cattribute_EAS_setv, 281
cattribute_GAB, 259
cattribute_GAB_set, 270
cattribute_GAN, 258
cattribute_GAN_set, 270
cattribute_GAS, 259
cattribute_GAS_set, 271
cattribute_has_attr, 257
cattribute_list, 257
cattribute_remove_all, 284
cattribute_remove_e, 283
cattribute_remove_g, 282
cattribute_remove_v, 283
cattribute_VAB, 262
cattribute_VABV, 263
cattribute_VAB_set, 273
cattribute_VAB_setv, 278
cattribute_VAN, 260
cattribute_VANV, 261
cattribute_VAN_set, 272
cattribute_VAN_setv, 277
cattribute_VAS, 263
cattribute_VASV, 264

cattribute_VAS_set, 274
cattribute_VAS_setv, 279
centralization, 341
centralization_betweenness, 343
centralization_betweenness_tmax, 346
centralization_closeness, 343
centralization_closeness_tmax, 347
centralization_degree, 342
centralization_degree_tmax, 345
centralization_eigenvector centrality, 344
centralization_eigenvector centrality_tmax, 348
CHECK, 37
cited_type_game, 210
citing_cited_type_game, 211
cliques, 394
cliques_callback, 395
clique_handler_t, 396
clique_number, 402
clique_size_hist, 394
closeness, 323
closeness_cutoff, 337
clusters, 315
cocitation, 349
cohesion, 486
cohesive_blocks, 487
coloring_greedy_t, 429
community_eb_get_merges, 508
community_edge_betweenness, 506
community_fastgreedy, 509
community_fluid_communities, 512
community_infomap, 514
community_label_propagation, 513
community_leading_eigenvector, 502
community_leading_eigenvector_callback_t, 504
community_leiden, 511
community_multilevel, 510
community_optimal_modularity, 494
community_springlass, 498
community_springlass_single, 500
community_to_membership, 495
community_walktrap, 505
compare_communities, 496
complementer, 546
compose, 547
connect_neighborhood, 548
constraint, 332
contract_vertices, 548
convergence_degree, 336
convex_hull, 588
copy, 15
coreness, 371
correlated_game, 215
correlated_pair_game, 216
count_isomorphisms_vf2, 415
count_multiple, 368
count_subisomorphisms_vf2, 420
create, 174
create_bipartite, 569

D

decompose, 317
decompose_destroy, 317
degree, 23
degree_sequence_game, 196
DELALL, 285
DELEA, 283
DELEAS, 285
delete_edges, 26
delete_vertices, 26
DELGA, 282
DELGAS, 284
DELVA, 283
DELVAS, 284
density, 378
destroy, 15
deterministic_optimal_imitation, 218
de_bruijn, 189
dfs, 390
dfshandler_t, 391
diameter, 302
diameter_dijkstra, 303
difference, 546
dim_select, 540
disjoint_union, 542
disjoint_union_many, 542
diversity, 379
dominator_tree, 474
dot_product_game, 214
dqueue_back, 143
dqueue_clear, 142
dqueue_destroy, 141
dqueue_empty, 141
dqueue_full, 141
dqueue_head, 142
dqueue_init, 140
dqueue_pop, 143
dqueue_pop_back, 143
dqueue_push, 143
dqueue_size, 142
dyad_census, 430

E

EAB, 267
EABV, 268
EAN, 265
EANV, 266
EAS, 269
EASV, 269
eccentricity, 304
ecount, 16
edge, 16
edges, 17
edge_betweenness, 327
edge_betweenness_cutoff, 340
edge_connectivity, 482
edge_disjoint_paths, 484

eigenvector centrality, 334
eit_create, 247
eit_destroy, 247
EIT_END, 248
EIT_GET, 249
EIT_NEXT, 248
EIT_RESET, 249
EIT_SIZE, 248
empty, 14
empty_attrs, 14
erdos_renyi_game, 193
ERROR, 35
error, 36
ERRORF, 36
errorf, 37
error_handler_abort, 28
error_handler_ignore, 28
error_handler_printignore, 28
error_handler_t, 27
error_type_t, 28
ess_1, 244
ess_all, 243
ess_none, 243
ess_seq, 244
ess_vector, 244
establishment_game, 203
es_1, 239
es_all, 237
es_as_vector, 245
es_copy, 245
es_destroy, 246
es_fromto, 240
es_incident, 238
es_is_all, 246
es_none, 238
es_pairs, 241
es_pairs_small, 241
es_path, 242
es_seq, 240
es_size, 246
es_type, 247
es_vector, 239
es_vector_copy, 242
eulerian_cycle, 385
eulerian_path, 386
even_tarjan_reduction, 490
extended_chordal_ring, 190

F

famous, 184
FATAL, 39
fatal, 40
FATALF, 40
fatalf, 40
fatal_handler_abort, 39
fatal_handler_t, 39
feedback_arc_set, 372
FINALLY, 38

FINALLY_CLEAN, 38
FINALLY_FREE, 38
forest_fire_game, 200
free, 42
FROM, 17
from_prufer, 187
full, 181
full_bipartite, 570
full_citation, 181

G

GAB, 259
GAN, 258
GAS, 260
get_adjacency, 382
get_all_shortest_paths, 297
get_all_shortest_paths_dijkstra, 298
get_all_simple_paths, 299
get_edgelist, 384
get_eid, 19
get_eids, 19
get_eids_multi, 20
get_incidence, 572
get_isomorphisms_vf2, 416
get_shortest_path, 291
get_shortest_paths, 290
get_shortest_paths_bellman_ford, 295
get_shortest_paths_dijkstra, 292
get_shortest_path_bellman_ford, 296
get_shortest_path_dijkstra, 294
get_sparsemat, 136
get_stochastic, 382
get_stochastic_sparsemat, 383
get_subisomorphisms_vf2, 421
girth, 304
global_efficiency, 306
gomory_hu_tree, 480
graphlets, 516
graphlets_candidate_basis, 517
graphlets_project, 517
grg_game, 191
growing_random_game, 202

H

harmonic centrality, 325
harmonic centrality_cutoff, 338
has_multiple, 367
heap_delete_top, 146
heap_destroy, 145
heap_empty, 145
heap_init, 144
heap_init_array, 144
heap_push, 145
heap_reserve, 147
heap_size, 146
heap_top, 146
hrg_consensus, 521

hrg_create, 524
hrg_dendrogram, 523
hrg_destroy, 520
hrg_fit, 521
hrg_game, 523
hrg_init, 520
hrg_predict, 524
hrg_resize, 521
hrg_sample, 522
hrg_size, 520
hrg_t, 519
hsbm_game, 212
hsbm_list_game, 213
hub_score, 335

I

incidence, 571
incident, 22
inclist_clear, 157
inclist_destroy, 156
inclist_get, 156
inclist_init, 155
inclist_size, 157
independence_number, 406
independent_vertex_sets, 404
induced_subgraph, 549
intersection, 544
intersection_many, 545
isoclass, 425
isoclass_create, 427
isoclass_subgraph, 426
isocompat_t, 417
isohandler_t, 417
isomorphic, 408
isomorphic_3d, 425
isomorphic_bliss, 411
isomorphic_function_vf2, 418
isomorphic_vf2, 414
is_bigraphical, 321
is_bipartite, 575
is_chordal, 374
is_connected, 316
is_dag, 371
is_degree_sequence, 322
is_directed, 22
is_eulerian, 385
is_graphical, 320
is_graphical_degree_sequence, 323
is_loop, 366
is_matching, 375
is_maximal_matching, 376
is_minimal_separator, 488
is_multiple, 367
is_mutual, 380
is_same_graph, 23
is_separator, 488
is_simple, 366
is_tree, 359

K

kautz, 189
k_regular_game, 197

L

lapack_dgeev, 557
lapack_dgeevx, 558
lapack_dgesv, 555
lapack_dgetrf, 554
lapack_dgetrs, 554
lapack_dsyevr, 555
laplacian, 365
laplacian_spectral_embedding, 539
largest_cliques, 396
largest_independent_vertex_sets, 405
largest_weighted_cliques, 403
lastcit_game, 209
lattice, 178
layout_bipartite, 440
layout_circle, 437
layout_davidson_harel, 447
layout_drl, 443
layout_drl_3d, 444
layout_drl_default_t, 442
layout_drl_options_init, 443
layout_drl_options_t, 441
layout_fruchterman_reingold, 445
layout_fruchterman_reingold_3d, 455
layout_gem, 447
layout_graphopt, 439
layout_grid, 438
layout_grid_3d, 454
layout_kamada_kawai, 446
layout_kamada_kawai_3d, 456
layout_lgl, 449
layout_mds, 448
layout_merge_dla, 457
layout_random, 437
layout_random_3d, 453
layout_reingold_tilford, 450
layout_reingold_tilford_circular, 451
layout_sphere, 454
layout_star, 438
layout_sugiyama, 452
lazy_adjlist_clear, 159
lazy_adjlist_destroy, 158
lazy_adjlist_get, 158
lazy_adjlist_init, 157
lazy_adjlist_size, 158
lazy_inclist_clear, 161
lazy_inclist_destroy, 160
lazy_inclist_get, 160
lazy_inclist_init, 159
lazy_inclist_size, 160
lcf, 186
lcf_vector, 187
le_community_to_membership, 504

linegraph, 550
list_triangles, 432
local_efficiency, 306
local_scan_0, 311
local_scan_0_them, 312
local_scan_1_ecount, 311
local_scan_1_ecount_them, 313
local_scan_k_ecount, 312
local_scan_k_ecount_them, 314
local_scan_neighborhood_ecount, 314

M

malloc, 42
MATRIX, 81
matrix_add, 87
matrix_add_cols, 100
matrix_add_constant, 86
matrix_add_rows, 99
matrix_all_e, 90
matrix_all_g, 91
matrix_all_ge, 91
matrix_all_l, 90
matrix_all_le, 91
matrix_as_sparsemat, 136
matrix_capacity, 96
matrix_cbind, 92
matrix_colsum, 89
matrix_contains, 98
matrix_copy, 79
matrix_copy_to, 80
matrix_destroy, 79
matrix_div_elements, 88
matrix_e, 81
matrix_empty, 95
matrix_e_ptr, 82
matrix_fill, 80
matrix_get_col, 83
matrix_get_row, 82
matrix_init, 78
matrix_isnull, 95
matrix_is_symmetric, 97
matrix_max, 93
matrix_maxdifference, 97
matrix_min, 92
matrix_minmax, 94
matrix_mul_elements, 88
matrix_ncol, 97
matrix_nrow, 96
matrix_null, 79
matrix_prod, 89
matrix_rbind, 92
matrix_remove_col, 100
matrix_remove_row, 100
matrix_resize, 99
matrix_resize_min, 99
matrix_rowsum, 89
matrix_scale, 87
matrix_search, 98
matrix_select_cols, 85
matrix_select_rows, 85
matrix_select_rows_cols, 86
matrix_set, 82
matrix_set_col, 84
matrix_set_row, 83
matrix_size, 96
matrix_sub, 87
matrix_sum, 88
matrix_swap, 81
matrix_swap_cols, 84
matrix_swap_rows, 84
matrix_transpose, 90
matrix_update, 80
matrix_which_max, 94
matrix_which_min, 93
matrix_which_minmax, 94
maxdegree, 332
maxflow, 472
maxflow_stats_t, 475
maxflow_value, 473
maximal_cliques, 397
maximal_cliques_callback, 401
maximal_cliques_count, 398
maximal_cliques_file, 399
maximal_cliques_hist, 400
maximal_cliques_subset, 399
maximal_independent_vertex_sets, 406
maximum_bipartite_matching, 376
maximum_cardinality_search, 373
mincut, 479
mincut_value, 480
minimum_size_separators, 489
minimum_spanning_tree, 356
minimum_spanning_tree_prim, 357
minimum_spanning_tree_unweighted, 357
modularity, 492
modularity_matrix, 493
moran_process, 219
motifs_handler_t, 436
motifs_randesu, 433
motifs_randesu_callback, 435
motifs_randesu_estimate, 434
motifs_randesu_no, 434

N

neighborhood, 309
neighborhood_graphs, 310
neighborhood_size, 308
neighbors, 21

O

OTHER, 18

P

pagerank, 328
pagerank_algo_t, 327

path_length_hist, 301
permute_vertices, 427
personalized_pagerank, 329
personalized_pagerank_vs, 330
plfit_result_t, 588
power_law_fit, 589
preference_game, 204
PROGRESS, 578
progress, 578
progressf, 579
progress_handler_stderr, 577
progress_handler_t, 577
psumtree_destroy, 162
psumtree_get, 162
psumtree_init, 161
psumtree_search, 163
psumtree_size, 162
psumtree_sum, 163
psumtree_update, 163

R

radius, 305
random_edge_walk, 392
random_sample, 585
random_spanning_tree, 358
random_walk, 391
read_graph_dimacs, 462
read_graph_dl, 470
read_graph_edgelist, 458
read_graph_gml, 466
read_graph_graphdb, 464
read_graph_graphml, 465
read_graph_lgl, 460
read_graph_ncol, 459
read_graph_pajek, 468
realize_degree_sequence, 182
recent_degree_aging_game, 208
recent_degree_game, 206
reciprocity, 379
reindex_membership, 496
rewire, 201
rewire_directed_edges, 195
rewire_edges, 194
ring, 179
rngtype_glibc2, 171
rngtype_mt19937, 170
rngtype_rand, 171
rng_default, 165
rng_destroy, 166
rng_get_binom, 169
rng_get_gamma, 170
rng_get_geom, 169
rng_get_integer, 167
rng_get_normal, 169
rng_get_unif, 168
rng_get_unif01, 168
rng_init, 165
rng_max, 167
rng_min, 166
rng_name, 167
rng_seed, 166
rng_set_default, 165
roulette_wheel_imitation, 221
running_mean, 584

S

sample_dirichlet, 587
sample_sphere_surface, 586
sample_sphere_volume, 586
sbm_game, 212
scg_adjacency, 527
scg_grouping, 533
scg_laplacian, 531
scg_norm_eps, 536
scg_semiprojectors, 535
scg_stochastic, 529
SETEAB, 276
SETEABV, 281
SETEAN, 275
SETEANV, 280
SETEAS, 277
SETEASV, 282
SETGAB, 271
SETGAN, 270
SETGAS, 272
SETVAB, 273
SETVABV, 279
SETVAN, 272
SETVANV, 278
SETVAS, 274
SETVASV, 279
set_attribute_table, 252
set_error_handler, 35
set_progress_handler, 577
set_status_handler, 581
set_warning_handler, 32
shortest_paths, 286
shortest_paths_bellman_ford, 288
shortest_paths_dijkstra, 287
shortest_paths_johnson, 289
similarity_dice, 352
similarity_dice_es, 354
similarity_dice_pairs, 353
similarity_inverse_log_weighted, 355
similarity_jaccard, 350
similarity_jaccard_es, 352
similarity_jaccard_pairs, 351
simple_interconnected_islands_game, 216
simplify, 550
simplify_and_colorize, 428
sir, 224
sir_destroy, 225
sir_t, 224
small, 174
sparsemat, 135
sparsemat_add, 122

sparsemat_add_cols, 124
sparsemat_add_rows, 123
sparsemat_arpack_rnsolve, 135
sparsemat_arpack_rssolve, 134
sparsemat_as_matrix, 137
sparsemat_cholsol, 130
sparsemat_colsums, 118
sparsemat_compress, 127
sparsemat_copy, 110
sparsemat_count_nonzero, 117
sparsemat_count_nonzerotol, 117
sparsemat_destroy, 111
sparsemat_diag, 112
sparsemat_droptol, 120
sparsemat_dropzeros, 120
sparsemat_dupl, 127
sparsemat_entry, 119
sparsemat_eye, 111
sparsemat_fkeep, 119
sparsemat_gaxpy, 123
sparsemat_getelements_sorted, 116
sparsemat_index, 113
sparsemat_init, 110
sparsemat_is_cc, 115
sparsemat_is_triplet, 115
sparsemat_iterator_col, 126
sparsemat_iterator_end, 125
sparsemat_iterator_get, 126
sparsemat_iterator_idx, 127
sparsemat_iterator_init, 124
sparsemat_iterator_next, 126
sparsemat_iterator_reset, 125
sparsemat_iterator_row, 125
sparsemat_lsolve, 129
sparsemat_ltsolve, 129
sparsemat_lu, 131
sparsemat_luresol, 132
sparsemat_lusol, 131
sparsemat_max, 116
sparsemat_min, 116
sparsemat_minmax, 117
sparsemat_multiply, 122
sparsemat_ncol, 114
sparsemat_nonzero_storage, 119
sparsemat_nrow, 114
sparsemat_numeric_destroy, 134
sparsemat_permute, 121
sparsemat_print, 137
sparsemat_qr, 132
sparsemat_qresol, 133
sparsemat_realloc, 111
sparsemat_resize, 124
sparsemat_rowsums, 118
sparsemat_scale, 121
sparsemat_symblu, 128
sparsemat_symbolic_destroy, 133
sparsemat_symbqr, 128
sparsemat_transpose, 122
sparsemat_type, 115
sparsemat_utsolve, 130
sparsemat_utsolve, 130
sparsemat_view, 112
split_join_distance, 497
spmatrix_add_cols, 108
spmatrix_add_e, 103
spmatrix_add_rows, 108
spmatrix_colsums, 107
spmatrix_copy, 101
spmatrix_count_nonzero, 106
spmatrix_destroy, 102
spmatrix_e, 102
spmatrix_fprint, 109
spmatrix_init, 101
spmatrix_iter_create, 103
spmatrix_iter_destroy, 105
spmatrix_iter_end, 104
spmatrix_iter_next, 104
spmatrix_iter_reset, 104
spmatrix_max, 106
spmatrix_ncol, 106
spmatrix_nrow, 105
spmatrix_print, 109
spmatrix_resize, 108
spmatrix_rowsums, 107
spmatrix_scale, 107
spmatrix_set, 102
spmatrix_size, 105
stack_clear, 139
stack_destroy, 138
stack_empty, 138
stack_init, 137
stack_pop, 140
stack_push, 139
stack_reserve, 138
stack_size, 139
stack_top, 140
star, 178
static_fitness_game, 198
static_power_law_game, 199
STATUS, 581
status, 582
STATUSF, 582
statusf, 582
status_handler_stderr, 581
status_handler_t, 580
stochastic_imitation, 222
STR, 148
strength, 333
strerror, 31
strvector_add, 151
strvector_append, 150
strvector_clear, 150
strvector_copy, 147
strvector_destroy, 148
strvector_get, 148
strvector_init, 147

strvector_remove, 150
strvector_resize, 151
strvector_set, 149
strvector_set2, 149
strvector_size, 151
st_edge_connectivity, 481
st_mincut, 475
st_mincut_value, 476
st_vertex_connectivity, 482
subcomponent, 315
subgraph_edges, 551
subisomorphic, 409
subisomorphic_function_vf2, 422
subisomorphic_lad, 424
subisomorphic_vf2, 419

T

THREAD_SAFE, 576
TO, 18
topological_sorting, 372
to_directed, 363
to_prufer, 360
to_undirected, 364
transitivity_avglocal_undirected, 362
transitivity_barrat, 363
transitivity_local_undirected, 361
transitivity_undirected, 360
tree, 180
tree_game, 214
triad_census, 430

U

unfold_tree, 377
union, 543
union_many, 544

V

VAB, 262
VABV, 263
VAN, 261
VANV, 261
VAS, 264
VASV, 265
vcount, 16
VECTOR, 48
vector_add, 54
vector_add_constant, 53
vector_all_e, 56
vector_all_g, 56
vector_all_ge, 57
vector_all_l, 56
vector_all_le, 57
vector_append, 51
vector_binsearch, 66
vector_binsearch2, 67
vector_binsearch_slice, 66
vector_capacity, 62

vector_clear, 67
vector_colex_cmp, 58
vector_contains, 65
vector_copy, 47
vector_copy_to, 50
vector_destroy, 47
vector_difference_sorted, 72
vector_div, 55
vector_e, 48
vector_empty, 61
vector_e_ptr, 49
vector_fill, 48
vector_floor, 55
vector_init, 45
vector_init_copy, 46
vector_init_seq, 46
vector_insert, 70
vector_intersect_sorted, 71
vector_isininterval, 63
vector_is_any_nan, 64
vector_is_nan, 64
vector_lex_cmp, 58
vector_max, 59
vector_maxdifference, 63
vector_min, 59
vector_minmax, 60
vector_mul, 55
vector_null, 47
vector_order, 64
vector_pop_back, 69
vector_prod, 62
vector_ptr_clear, 75
vector_ptr_copy, 73
vector_ptr_destroy, 73
vector_ptr_destroy_all, 74
vector_ptr_e, 76
vector_ptr_free_all, 74
vector_ptr_get_item_destructor, 77
vector_ptr_init, 73
vector_ptr_insert, 76
vector_ptr_push_back, 75
vector_ptr_resize, 76
vector_ptr_set, 76
vector_ptr_set_item_destructor, 78
VECTOR_PTR_SET_ITEM_DESTRUCTOR, 78
vector_ptr_size, 74
vector_ptr_sort, 77
vector_push_back, 69
vector_remove, 70
vector_remove_section, 70
vector_reserve, 67
vector_resize, 68
vector_resize_min, 68
vector_reverse, 52
vector_reverse_sort, 71
vector_scale, 53
vector_search, 65
vector_set, 49

vector_shuffle, 53
vector_size, 61
vector_sort, 71
vector_sub, 54
vector_sum, 62
vector_swap, 51
vector_swap_elements, 52
vector_tail, 50
vector_update, 51
vector_view, 50
vector_which_max, 60
vector_which_min, 59
vector_which_minmax, 61
version, 584
vertex_coloring_greedy, 429
vertex_connectivity, 483
vertex_disjoint_paths, 485
vit_create, 234
vit_destroy, 235
VIT_END, 236
VIT_GET, 237
VIT_NEXT, 236
VIT_RESET, 237
VIT_SIZE, 236
vss_1, 233
vss_all, 232
vss_none, 233
vss_seq, 234
vss_vector, 234
vs_1, 228
vs_adj, 227
vs_all, 226
vs_copy, 231
vs_destroy, 231
vs_is_all, 232
vs_nonadj, 227
vs_none, 228
vs_seq, 230
vs_size, 232
vs_type, 232
vs_vector, 229
vs_vector_copy, 230
vs_vector_small, 229
write_graph_edgelist, 458
write_graph_gml, 467
write_graph_graphml, 465
write_graph_lgl, 461
write_graph_ncol, 460
write_graph_pajek, 469

W

WARNING, 32
warning, 33
WARNINGF, 33
warningf, 33
warning_handler_ignore, 34
warning_handler_print, 34
warning_handler_t, 32
watts_strogatz_game, 194
weighted_adjacency, 176
weighted_cliques, 402
weighted_clique_number, 404
write_graph_dimacs, 463
write_graph_dot, 470