igraph Reference Manual

Gábor Csárdi, Department of Statistics, Harvard University
Tamás Nepusz, Department of Biological Physics, Eötvös Loránd University
Vincent Traag, Centre for Science and Technology Studies, Leiden University
Szabolcs Horvát, Center for Systems Biology Dresden, Max Planck Institute for Cell Biology and Genetics
Fabio Zanini, Lowy Cancer Research Centre, University of New South Wales
Daniel Noom, jitjit software development

igraph Reference Manual

by Gábor Csárdi, Tamás Nepusz, Vincent Traag, Szabolcs Horvát, Fabio Zanini, and Daniel Noom 0.10.0

This manual is for igraph, version 0.10.0.

Copyright (C) 2005-2019 Gábor Csárdi and Tamás Nepusz. Copyright (C) 2020-2022 igraph development team. Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. A copy of the license is included in the section entitled "GNU Free Documentation License".

Table of Contents

1.	Introduction	1
	igraph is free software	1
	Citing igraph	2
2.	Installation	3
	Prerequisites	3
	Installation	
	General build instructions	
	Specific instructions for Windows	
	Notable configuration options	
	Building the documentation	
	Notes for package maintainers	
	Auto-detection of dependencies	
	Shared and static builds	
	Cross-compiling	
2	Additional notes	
3.	Tutorial	
	Compiling programs using igraph	
	Compiling with CMake	
	Compiling without CMake	
	Running the program	
	Creating your first graphs	
	Calculating various properties of graphs	
4.	Basic data types and interface	
	The igraph data model	
	General conventions of igraph functions	
	Atomic data types	
	The basic interface	
	Graph constructors and destructors	
	Basic query operations	
	Adding and deleting vertices and edges	
	Miscellaneous macros and helper functions	28
	IGRAPH_VCOUNT_MAX — The maximum number of vertices supported in	
	igraph graphs.	28
	IGRAPH_ECOUNT_MAX — The maximum number of edges supported in igraph	
	graphs.	29
	igraph_expand_path_to_pairs — Helper function to convert a sequence	
	of vertex IDs describing a path into a "pairs" vector.	29
	igraph_invalidate_cache — Invalidates the internal cache of an igraph	
	graph.	29
	igraph_is_same_graph — Are two graphs identical as labelled graphs?	
5.	Error handling	31
	Error handling basics	
	Error handlers	. 31
	igraph_error_handler_t — The type of error handler functions	
	igraph_error_handler_abort — Abort program in case of error	
	igraph_error_handler_ignore — Ignore errors	
	igraph_error_handler_printignore — Print and ignore errors	
	Error codes	
	igraph_error_t — Return type for functions returning an error code	
	igraph_error_type_t — Error code type	
	igraph_strerror — Textual description of an error	
	Warning messages	
	igraph_warning_handler_t — The type of igraph warning handler func-	. 50
	tions	36
	igraph_set_warning_handler — Installs a warning handler	
	Tarapii_bee_wariitiig_nanatet — mbanb a wannig nanatei	50

IGRAPH_WARNING — Triggers a warning.	. 37
IGRAPH_WARNINGF — Triggers a warning, with printf-like syntax	. 37
igraph_warning — Reports a warning	. 37
igraph_warningf — Reports a warning, printf-like version	. 38
igraph_warning_handler_ignore — Ignores all warnings	
igraph_warning_handler_print — Prints all warnings to the standard	
error.	. 38
Advanced topics	
Writing error handlers	
Error handling internals	
Deallocating memory	
Writing igraph functions with proper error handling	
Fatal errors	
Error handling and threads	
6. Memory (de)allocation	. 47
igraph_malloc — Allocate memory that can be safely deallocated by igraph func-	
tions.	. 47
igraph_calloc — Allocate memory that can be safely deallocated by igraph func-	
tions	. 47
igraph_realloc — Reallocate memory that can be safely deallocated by igraph	
functions.	. 48
igraph_free — Deallocate memory that was allocated by igraph functions	. 48
7. Data structure library: vector, matrix, other data types	
About template types	
Vectors	
About igraph_vector_t objects	
Constructors and destructors	
Initializing elements	
Accessing elements	
e e e e e e e e e e e e e e e e e e e	
Vector views	
Copying vectors	
Exchanging elements	
Vector operations	
Vector comparisons	
Finding minimum and maximum	
Vector properties	
Searching for elements	
Resizing operations	. 75
Complex vector operations	. 78
Sorting	. 81
Set operations on sorted vectors	. 82
Pointer vectors (igraph_vector_ptr_t)	83
Deprecated functions	. 90
Matrices	. 92
About igraph_matrix_t objects	92
Matrix constructors and destructors	
Initializing elements	
Accessing elements of a matrix	
Matrix views	
Copying matrices	
Operations on rows and columns	
Matrix operations	
Matrix comparisons	
Combining matrices	
Finding minimum and maximum	
Matrix properties	
Searching for elements	
Resizing operations	115

Complex matrix operations	
Deprecated functions	
Sparse matrices	121
About sparse matrices	121
Creating sparse matrix objects	121
Query properties of a sparse matrix	125
Operations on sparse matrices	
Operations on sparse matrix iterators	
Operations that change the internal representation	
Decompositions and solving linear systems	
Eigenvalues and eigenvectors	
Conversion to other data types	
Writing to a file, or to the screen	
Deprecated functions	
Stacks	
igraph_stack_init — Initializes a stack.	
igraph_stack_destroy — Destroys a stack object	
igraph_stack_reserve — Reserve memory.	
igraph_stack_empty — Decides whether a stack object is empty	
igraph_stack_size — Returns the number of elements in a stack	
igraph_stack_clear — Removes all elements from a stack	
igraph_stack_push — Places an element on the top of a stack	153
igraph_stack_pop — Removes and returns an element from the top of a	
stack.	
igraph_stack_top — Query top element	
Double-ended queues	154
igraph_dqueue_init — Initialize a double ended queue (deque)	
igraph_dqueue_destroy — Destroy a double ended queue	155
igraph_dqueue_empty — Decide whether the queue is empty	155
igraph_dqueue_full — Check whether the queue is full	155
igraph_dqueue_clear — Remove all elements from the queue	
igraph_dqueue_size — Number of elements in the queue	
igraph_dqueue_head — Head of the queue	156
igraph_dqueue_back — Tail of the queue.	157
igraph_dqueue_pop — Remove the head.	
igraph_dqueue_pop_back — Removes the tail	
igraph_dqueue_push — Appends an element	
Maximum and minimum heaps	
igraph_heap_init — Initializes an empty heap object.	
igraph_heap_init_array — Build a heap from an array	
igraph_heap_destroy — Destroys an initialized heap object	
igraph_heap_clear — Removes all elements from a heap.	
igraph_heap_empty — Decides whether a heap object is empty	
igraph_heap_push — Add an element.	
igraph_heap_top — Top element.	
igraph_heap_delete_top — Removes and returns the top element	
igraph_heap_size — Number of elements in the heap	
igraph_heap_reserve — Reserves memory for a heap	
String vectors	
igraph_strvector_init — Initializes a string vector	
igraph_strvector_init_copy — Initialization by copying	162
igraph_strvector_destroy — Frees the memory allocated for the string	
vector.	162
STR — Indexing string vectors.	162
igraph_strvector_get — Retrieves an element of the string vector	163
igraph_strvector_set — Takes elements at given positions from a string	
vector.	163

191apii_strvector_set_1eii — Sets an element of the string vector given
a buffer and its size.
igraph_strvector_push_back — Adds an element to the back of a string
vector.
igraph_strvector_push_back_len — Adds a string of the given length
to the back of a string vector.
igraph_strvector_remove — Removes a single element from a string
vector.
igraph_strvector_remove_section — Removes a section from a string
vector.
igraph_strvector_append — Concatenates two string vectors
igraph_strvector_merge — Moves the contents of a string vector to the
end of another.
igraph_strvector_clear — Removes all elements from a string vector
igraph_strvector_resize — Resizes a string vector.
igraph_strvector_reserve — Reserves memory for a string vector
igraph_strvector_resize_min — Deallocates the unused memory of a
string vector.
igraph_structor_size — Returns the size of a string vector.
igraph_strvector_capacity — Returns the capacity of a string vector
Deprecated functions
Lists of vectors, matrices and graphs
About igraph_vector_list_t objects
Constructors and destructors
Accessing elements
Vector properties
Resizing operations
Sorting and reordering
Adjacency lists
Adjacent vertices
Incident edges
Lazy adjacency list for vertices
Lazy incidence list for edges
Partial prefix sum trees
igraph_psumtree_init — Initializes a partial prefix sum tree.
igraph_psumtree_destroy — Destroys a partial prefix sum tree
igraph_psumtree_size — Returns the size of the tree.
igraph_psumtree_get — Retrieves the value corresponding to an item in
the tree.
igraph_psumtree_sum — Returns the sum of the values of the leaves in the
treeigraph_psumtree_search — Finds an item in the tree, given a value
igraph_psumtree_search — Finds an item in the tree, given a value
tree
8. Random numbers
About random numbers in igraph
The default random number generator
igraph_rng_default — Query the default random number generator
igraph_rng_default — Query the default random number generator
ator.
Creating random number generators
igraph_rng_init — Initializes a random number generator
igraph_rng_destroy — Deallocates memory associated with a random
number generator.
igraph_rng_seed — Seeds a random number generator.
igraph_rng_bits — The number of random bits that a random number gen-
erator can produces in a single round.

erator	.98 .99 .99 .99 .00 201 201 202 202
Generating random numbers	.99 .99 .99 .00 .00 .01 .01 .02 .02 .02 .03
igraph_rng_get_integer — Generate an integer random number from an interval	.99 .99 .99 .00 .00 .01 .01 .02 .02 .03
interval	.99 .99 200 201 201 202 202 203
igraph_rng_get_unif01 — Samples uniformly from the unit interval	.99 .99 200 201 201 202 202 203
igraph_rng_get_unif — Samples real numbers from a given interval	.99 200 201 201 202 202 203
igraph_rng_get_normal — Samples from a normal distribution. 2 igraph_rng_get_exp — Samples from an exponential distribution. 2 igraph_rng_get_gamma — Samples from a gamma distribution. 2 igraph_rng_get_binom — Samples from a binomial distribution. 2 igraph_rng_get_geom — Samples from a geometric distribution. 2 igraph_rng_get_pois — Samples from a Poisson distribution. 2 Supported random number generators	200 201 201 202 202 203
igraph_rng_get_exp — Samples from an exponential distribution	200 201 201 202 202 203
igraph_rng_get_gamma — Samples from a gamma distribution	201 201 202 202 203
igraph_rng_get_binom — Samples from a binomial distribution.	201 202 202 203
igraph_rng_get_geom — Samples from a geometric distribution	202 202 203
igraph_rng_get_pois — Samples from a Poisson distribution. 2 Supported random number generators	202 203
Supported random number generators	203
igraph_rngtype_mt19937 — The MT19937 random number generator 20 igraph_rngtype_glibc2 — The random number generator introduced in GNU libc 2	
igraph_rngtype_glibc2 — The random number generator introduced in GNU libc 2	203
GNU libc 2	
	:03
igraph_rngtype_pcg32 — The PCG random number generator (32-bit ver-	
sion)	:04
igraph_rngtype_pcg64 — The PCG random number generator (64-bit ver-	
sion)	
Use cases	
Normal (default) use	
Reproducible simulations	:05
Changing the default generator	
Using multiple generators	
Example 2	
9. Graph generators	:07
Deterministic graph generators	
igraph_create — Creates a graph with the specified edges	:07
igraph_small — Shorthand to create a small graph, giving the edges as argu-	
ments	
igraph_adjacency — Creates a graph from an adjacency matrix 20	.08
igraph_weighted_adjacency — Creates a graph from a weighted adja-	
cency matrix	:09
igraph_sparse_adjacency — Creates a graph from a sparse adjacency	
	210
igraph_sparse_weighted_adjacency — Creates a graph from a weight-	
ed sparse adjacency matrix	
igraph_adjlist — Creates a graph from an adjacency list	
igraph_star — Creates a <i>star</i> graph, every vertex connects only to the center 2	
igraph_wheel — Creates a <i>wheel</i> graph, a union of a star and a cycle graph 2	
igraph_square_lattice — Arbitrary dimensional square lattices 2	
igraph_ring — Creates a cycle graph or a path graph	.14
igraph_kary_tree — Creates a k-ary tree in which almost all vertices have	
k children	:15
igraph_symmetric_tree — Creates a symmetric tree with the specified	
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	216
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	216 217
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	216 217 218
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	216 217 218
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	216 217 218 218
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	216 217 218 218 219
igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level	216 217 218 218 219

1	igraph_lcf — Creates a graph from LCF notation	224
j	igraph_lcf_vector — Creates a graph from LCF notation	225
j	igraph_from_prufer — Generates a tree from a Prüfer sequence	225
j	igraph_atlas — Create a small graph from the "Graph Atlas"	226
i	igraph_de_bruijn — Generate a de Bruijn graph	226
j	igraph_kautz — Generate a Kautz graph	227
j	igraph_circulant — Creates a circulant graph	227
j	igraph_generalized_petersen — Creates a Generalized Petersen	
٤	graph	228
j	igraph_extended_chordal_ring — Create an extended chordal ring	229
Games	s: Randomized graph generators	229
j	igraph_grg_game — Generates a geometric random graph	230
j	igraph_barabasi_game — Generates a graph based on the Barabási-Albert	
r	model	230
j	igraph_erdos_renyi_game — Generates a random (Erd#s-Rényi) graph	232
j	igraph_watts_strogatz_game — The Watts-Strogatz small-world mod-	
e	el	232
	igraph_rewire_edges — Rewires the edges of a graph with constant proba-	
	pility.	233
	igraph_rewire_directed_edges — Rewires the chosen endpoint of di-	
	rected edges.	234
	igraph_degree_sequence_game — Generates a random graph with a giv-	
	en degree sequence.	235
	igraph_k_regular_game — Generates a random graph where each vertex	
	has the same degree.	236
	igraph_static_fitness_game — Non-growing random graph with edge	
	probabilities proportional to node fitness scores.	237
	igraph_static_power_law_game — Generates a non-growing random	
	graph with expected power-law degree distributions.	238
	igraph_forest_fire_game — Generates a network according to the "for-	
	est fire game".	239
	igraph_rewire — Randomly rewires a graph while preserving its degree se-	
	quence.	240
	igraph_growing_random_game — Generates a growing random graph	
	igraph_callaway_traits_game — Simulates a growing network with	
	vertex types.	241
	igraph_establishment_game — Generates a graph with a simple growing	
		242
	igraph_preference_game — Generates a graph with vertex types and con-	
	nection preferences.	243
	igraph_asymmetric_preference_game — Generates a graph with	
	asymmetric vertex types and connection preferences.	244
	igraph_recent_degree_game — Stochastic graph generator based on the	
r	number of incident edges a node has gained recently.	245
	igraph_barabasi_aging_game — Preferential attachment with aging of	
		246
	igraph_recent_degree_aging_game — Preferential attachment based	
C	on the number of edges gained recently, with aging of vertices	247
	igraph_lastcit_game — Simulates a citation network, based on time	
ŗ	passed since the last citation.	248
	igraph_cited_type_game — Simulates a citation based on vertex types	249
	igraph_citing_cited_type_game — Simulates a citation network based	
	on vertex types.	250
	igraph_sbm_game — Sample from a stochastic block model	
	igraph_hsbm_game — Hierarchical stochastic block model	
	igraph_hsbm_list_game — Hierarchical stochastic block model, more gen-	
	eral version.	252
	igraph dot product game — Generates a random dot product graph	

igraph_tree_game — Generates a random tree with the given number of	
nodes.	. 25
igraph_correlated_game — Generates a random graph correlated to an	
existing graph.	25
igraph_correlated_pair_game — Generates pairs of correlated random	
graphs.	25
igraph_simple_interconnected_islands_game — Generates a ran-	
dom graph made of several interconnected islands, each island being a random	
graph.	
Deprecated functions	
igraph_lattice — Arbitrary dimensional square lattices (deprecated)	25
igraph_tree — Creates a k-ary tree in which almost all vertices have k chil-	
dren (deprecated alias).	
10. Games on graphs	25
Microscopic update rules	25
igraph_deterministic_optimal_imitation — Adopt a strategy via	
deterministic optimal imitation.	25
igraph_moran_process — The Moran process in a network setting	25
igraph_roulette_wheel_imitation — Adopt a strategy via roulette	
wheel selection.	26
igraph_stochastic_imitation — Adopt a strategy via stochastic imita-	
tion with uniform selection.	26
Epidemic models	
igraph_sir — Performs a number of SIR epidemics model runs on a graph	
igraph_sir_t — The result of one SIR model simulation.	
igraph_sir_destroy — Deallocates memory associated with a SIR simula-	
tion run.	26
11. Vertex and edge selectors and sequences, iterators	
About selectors, iterators	
Vertex selector constructors	
igraph_vs_all — Vertex set, all vertices of a graph	
igraph_vs_adj — Adjacent vertices of a vertex	
igraph_vs_nonadj — Non-adjacent vertices of a vertex	
igraph_vs_none — Empty vertex set	
igraph_vs_1 — Vertex set with a single vertex.	
igraph_vs_vector — Vertex set based on a vector	
igraph_vs_vector_small — Create a vertex set by giving its elements	
igraph_vs_vector_copy — Vertex set based on a vector, with copying	
igraph_vs_range — Vertex set, an interval of vertices	
Generic vertex selector operations	
igraph_vs_copy — Creates a copy of a vertex selector	
igraph_vs_destroy — Destroy a vertex set	
igraph_vs_is_all — Check whether all vertices are included	
igraph_vs_size — Returns the size of the vertex selector	
igraph_vs_type — Returns the type of the vertex selector	
Immediate vertex selectors	
igraph_vss_all — All vertices of a graph (immediate version).	
igraph_vss_none — Empty vertex set (immediate version).	
igraph_vss_1 — Vertex set with a single vertex (immediate version)	
igraph_vss_vector — Vertex set based on a vector (immediate version)	
igraph_vss_range — An interval of vertices (immediate version)	
Vertex iterators	
igraph_vit_create — Creates a vertex iterator from a vertex selector	
igraph_vit_destroy — Destroys a vertex iterator	
Stepping over the vertices	
IGRAPH_VIT_NEXT — Next vertex	
IGRAPH_VIT_END — Are we at the end?	
IGRAPH_VIT_SIZE — Size of a vertex iterator	27

IGRAPH_VIT_RESET — Reset a vertex iterator	
IGRAPH_VIT_GET — Query the current position	277
Edge selector constructors	277
igraph_es_all — Edge set, all edges	277
igraph_es_incident — Edges incident on a given vertex	278
igraph_es_none — Empty edge selector	
igraph_es_1 — Edge selector containing a single edge	
igraph_es_vector — Handle a vector as an edge selector	
igraph_es_range — Edge selector, a sequence of edge IDs	
igraph_es_pairs — Edge selector, multiple edges defined by their endpoints	
in a vector.	280
igraph_es_pairs_small — Edge selector, multiple edges defined by their	200
endpoints as arguments.	281
igraph_es_path — Edge selector, edge IDs on a path.	
igraph_es_vector_copy — Edge set, based on a vector, with copying	
Immediate edge selectors	
igraph_ess_all — Edge set, all edges (immediate version)	
igraph_ess_none — Immediate empty edge selector.	
igraph_ess_1 — Immediate version of the single edge edge selector	
igraph_ess_vector — Immediate vector view edge selector	
igraph_ess_range — Immediate version of the sequence edge selector	
Generic edge selector operations	
igraph_es_as_vector — Transform edge selector into vector	
igraph_es_copy — Creates a copy of an edge selector	
igraph_es_destroy — Destroys an edge selector object	
igraph_es_is_all — Check whether an edge selector includes all edges	
igraph_es_size — Returns the size of the edge selector	
igraph_es_type — Returns the type of the edge selector	286
Edge iterators	286
igraph_eit_create — Creates an edge iterator from an edge selector	286
igraph_eit_destroy — Destroys an edge iterator	287
Stepping over the edges	287
IGRAPH_EIT_NEXT — Next edge.	287
IGRAPH_EIT_END — Are we at the end?	288
IGRAPH_EIT_SIZE — Number of edges in the iterator.	
IGRAPH_EIT_RESET — Reset an edge iterator.	
IGRAPH_EIT_GET — Query an edge iterator.	
Deprecated functions	
igraph_es_seq — Edge selector, a sequence of edge IDs, with inclusive end-	
	289
igraph_ess_seq — Immediate version of the sequence edge selector, with in-	
	289
igraph_vs_seq — Vertex set, an interval of vertices with inclusive endpoints	207
•	290
	290
igraph_vss_seq — An interval of vertices with inclusive endpoints (immediate vertices at a vertice)	200
······································	290
12. Graph, vertex and edge attributes	
The Attribute Handler Interface	292
igraph_attribute_table_t — Table of functions to perform operations	202
	292
igraph_set_attribute_table — Attach an attribute table	
igraph_attribute_type_t — The possible types of the attributes	
	295
igraph_attribute_combination_init — Initialize attribute combina-	
tion list.	296
	270
igraph_attribute_combination_add — Add combination record to at-	270
	296

igraph_attribute_combination_remove — Remove a record from an	
attribute combination list.	297
igraph_attribute_combination_destroy — Destroy attribute combi-	
nation list.	297
igraph_attribute_combination_type_t — The possible types of at-	
tribute combinations.	297
igraph_attribute_combination — Initialize attribute combination list	
and add records.	298
Accessing attributes from C	299
Query attributes	300
Set attributes	
Remove attributes	
Custom attribute combination functions	
13. Structural properties of graphs	
Basic properties	
igraph_are_connected — Decides whether two vertices are connected	
Sparsifiers	
igraph_spanner — Calculates a spanner of a graph with a given stretch fac-	343
tor	220
(Shortest)-path related functions	
igraph_distances — Length of the shortest paths between vertices	330
igraph_distances_dijkstra — Weighted shortest path lengths between	221
vertices.	331
igraph_distances_bellman_ford — Weighted shortest path lengths be-	
tween vertices, allowing negative weights.	332
igraph_distances_johnson — Weighted shortest path lengths between	
vertices, using Johnson's algorithm.	
igraph_get_shortest_paths — Shortest paths from a vertex	334
igraph_get_shortest_path — Shortest path from one vertex to another	
one.	335
igraph_get_shortest_paths_dijkstra — Weighted shortest paths	
from a vertex.	336
igraph_get_shortest_path_dijkstra — Weighted shortest path from	
one vertex to another one.	337
<pre>igraph_get_shortest_paths_bellman_ford — Weighted shortest</pre>	
paths from a vertex, allowing negative weights.	338
igraph_get_shortest_path_bellman_ford — Weighted shortest path	
from one vertex to another one.	339
igraph_get_all_shortest_paths — All shortest paths (geodesics) from	
a vertex.	340
igraph_get_all_shortest_paths_dijkstra — All weighted shortest	
paths (geodesics) from a vertex.	341
igraph_get_k_shortest_paths — k shortest paths between two ver-	
tices.	343
igraph_get_all_simple_paths — List all simple paths from one source	
igraph_average_path_length — Calculates the average unweighted	٥.,
shortest path length between all vertex pairs.	344
igraph_average_path_length_dijkstra — Calculates the average	5-1-1
weighted shortest path length between all vertex pairs.	3/15
	343
igraph_path_length_hist — Create a histogram of all shortest path	2/4
lengths.	346
igraph_diameter — Calculates the diameter of a graph (longest geodesic)	<i>541</i>
igraph_diameter_dijkstra — Calculates the weighted diameter of a	246
graph using Dijkstra's algorithm.	348
igraph_girth — The girth of a graph is the length of the shortest cycle in it	
igraph_eccentricity — Eccentricity of some vertices	349
igraph_eccentricity_dijkstra — Eccentricity of some vertices, using	_
weighted edges.	350

igraph_graph_center — Central vertices of a graph	
igraph_radius — Radius of a graph	351
igraph_pseudo_diameter — Approximation and lower bound of diame-	
ter.	352
igraph_vertex_path_from_edge_path — Converts a path of edge IDs	
to the traversed vertex IDs.	353
Widest-path related functions	
igraph_get_widest_path — Widest path from one vertex to another one	
igraph_get_widest_paths — Widest paths from a single vertex	
igraph_widest_path_widths_dijkstra — Widths of widest paths be-	334
•	255
tween vertices.	333
igraph_widest_path_widths_floyd_warshall — Widths of widest	
paths between vertices.	
Efficiency measures	357
igraph_global_efficiency — Calculates the global efficiency of a net-	
work.	357
igraph_local_efficiency — Calculates the local efficiency around each	
vertex in a network.	358
igraph_average_local_efficiency — Calculates the average local ef-	
ficiency in a network.	359
Neighborhood of a vertex	
igraph_neighborhood_size — Calculates the size of the neighborhood of	
a given vertex.	360
igraph_neighborhood — Calculate the neighborhood of vertices	
igraph_neighborhood_graphs — Create graphs from the neighbor-	501
hood(s) of some vertex/vertices.	261
Local scan statistics	
"Us" statistics	
"Them" statistics	
Pre-calculated subsets	
Graph components	367
igraph_subcomponent — The vertices in the same component as a given	
vertex.	367
<pre>igraph_connected_components — Calculates the (weakly or strongly)</pre>	
connected components in a graph.	367
igraph_clusters — Calculates the (weakly or strongly) connected compo-	
nents in a graph (deprecated alias).	368
igraph_is_connected — Decides whether the graph is (weakly or strongly)	
connected.	368
igraph_decompose — Decomposes a graph into connected components	
igraph_decompose_destroy — Frees the contents of a pointer vector hold-	
ing graphs.	370
igraph_biconnected_components — Calculates biconnected compo-	570
nents.	370
igraph_articulation_points — Finds the articulation points in a graph	
igraph_bridges — Finds all bridges in a graph.	
Degree sequences	
igraph_is_graphical — Is there a graph with the given degree sequence?	3/2
igraph_is_bigraphical — Is there a bipartite graph with the given bi-de-	
gree-sequence?	
Centrality measures	
igraph_closeness — Closeness centrality calculations for some vertices	
igraph_harmonic_centrality — Harmonic centrality for some vertices	
igraph_betweenness — Betweenness centrality of some vertices	377
igraph_edge_betweenness — Betweenness centrality of the edges	378
igraph_pagerank_algo_t — PageRank algorithm implementation	
igraph_pagerank — Calculates the Google PageRank for the specified ver-	
tices.	379

igraph_personalized_pagerank — Calculates the personalized Google	
PageRank for the specified vertices.	380
igraph_personalized_pagerank_vs — Calculates the personalized	
Google PageRank for the specified vertices.	381
igraph_constraint — Burt's constraint scores	383
igraph_maxdegree — The maximum degree in a graph (or set of vertices)	383
igraph_strength — Strength of the vertices, also called weighted vertex de-	20.4
gree.	384
igraph_eigenvector_centrality — Eigenvector centrality of the vertices.	385
igraph_hub_score — Kleinberg's hub scores	386
igraph_authority_score — Kleinerg's authority scores	
igraph_convergence_degree — Calculates the convergence degree of	
each edge in a graph.	388
Range-limited centrality measures	
igraph_closeness_cutoff — Range limited closeness centrality	
igraph_harmonic_centrality_cutoff — Range limited harmonic cen-	
trality.	
igraph_betweenness_cutoff — Range-limited betweenness centrality	391
igraph_edge_betweenness_cutoff — Range-limited betweenness cen-	
trality of the edges.	391
Subset-limited Centrality Measures	392
igraph_betweenness_subset — Betweenness centrality for a subset of	
source and target vertices.	392
igraph_edge_betweenness_subset — Edge betweenness centrality for a	
subset of source and target vertices.	393
Centralization	394
igraph_centralization — Calculate the centralization score from the	
node level scores.	394
igraph_centralization_degree — Calculate vertex degree and graph centralization.	395
igraph_centralization_betweenness — Calculate vertex between-	
ness and graph centralization.	396
igraph_centralization_closeness — Calculate vertex closeness and	
graph centralization.	396
igraph_centralization_eigenvector_centrality — Calculate	
eigenvector centrality scores and graph centralization.	397
igraph_centralization_degree_tmax — Theoretical maximum for	
graph centralization based on degree.	398
igraph_centralization_betweenness_tmax — Theoretical maxi-	
mum for graph centralization based on betweenness.	399
igraph_centralization_closeness_tmax — Theoretical maximum	
for graph centralization based on closeness.	400
igraph_centralization_eigenvector_centrality_tmax — The-	
oretical maximum centralization for eigenvector centrality.	401
Similarity measures	
igraph_bibcoupling — Bibliographic coupling	
igraph_cocitation — Cocitation coupling	
igraph_similarity_jaccard — Jaccard similarity coefficient for the giv-	
en vertices.	403
igraph_similarity_jaccard_pairs — Jaccard similarity coefficient for	. 50
given vertex pairs.	404
igraph_similarity_jaccard_es — Jaccard similarity coefficient for a	
given edge selector.	405
igraph_similarity_dice — Dice similarity coefficient	
igraph_similarity_dice_pairs — Dice similarity coefficient for given	
vertex pairs.	406

igraph_similarity_dice_es — Dice similarity coefficient for a given	
edge selector.	. 40
igraph_similarity_inverse_log_weighted — Vertex similarity	
based on the inverse logarithm of vertex degrees.	. 40
Trees and forests	40
igraph_minimum_spanning_tree — Calculates one minimum spanning	
tree of a graph.	. 40
igraph_minimum_spanning_tree_unweighted — Calculates one min-	
imum spanning tree of an unweighted graph.	. 41
igraph_minimum_spanning_tree_prim — Calculates one minimum	
spanning tree of a weighted graph.	. 41
igraph_random_spanning_tree — Uniformly samples the spanning trees	
of a graph.	. 41
igraph_is_tree — Decides whether the graph is a tree	41
igraph_is_forest — Decides whether the graph is a forest	
igraph_to_prufer — Converts a tree to its Prüfer sequence	
Transitivity or clustering coefficient	
igraph_transitivity_undirected — Calculates the transitivity (clus-	
tering coefficient) of a graph.	. 41
igraph_transitivity_local_undirected — Calculates the local tran-	, F1
sitivity (clustering coefficient) of a graph.	. 41
igraph_transitivity_avglocal_undirected — Average local tran-	
sitivity (clustering coefficient).	. 41
igraph_transitivity_barrat — Weighted transitivity, as defined by A.	. +1
Barrat.	11
Directedness conversion	
igraph_to_directed — Convert an undirected graph to a directed one	
igraph_to_undirected — Convert an undirected graph to an undirected one	
Spectral properties	
1 1 1	
igraph_get_laplacian — Returns the Laplacian matrix of a graph.	. 41
igraph_get_laplacian_sparse — Returns the Laplacian of a graph in a	41
sparse matrix format.	. 41
igraph_laplacian_normalization_t — Normalization methods for a	40
Laplacian matrix.	
Non-simple graphs: Multiple and loop edges	
igraph_is_simple — Decides whether the input graph is a simple graph	
igraph_is_loop — Find the loop edges in a graph	
igraph_is_multiple — Find the multiple edges in a graph	. 42
igraph_has_multiple — Check whether the graph has at least one multiple	
edge.	. 42
igraph_count_multiple — Count the number of appearances of the edges	
in a graph.	
Mixing patterns	42
igraph_assortativity_nominal — Assortativity of a graph based on	
vertex categories.	. 42
igraph_assortativity — Assortativity based on numeric properties of	
vertices.	. 42
igraph_assortativity_degree — Assortativity of a graph based on ver-	
tex degree.	
K-Cores and K-Trusses	
igraph_coreness — Finding the coreness of the vertices in a network	. 42
igraph_trussness — Finding the "trussness" of the edges in a network	. 42
Topological sorting, directed acyclic graphs	
igraph_is_dag — Checks whether a graph is a directed acyclic graph	
(DAG)	. 42
igraph_topological_sorting — Calculate a possible topological sorting	
of the graph.	42

igraph_feedback_arc_set — Feedback arc set of a graph using exact or	
heuristic methods.	
Maximum cardinality search and chordal graphs	430
$igraph_maximum_cardinality_search-Maximum$ cardinality	
search.	
igraph_is_chordal — Decides whether a graph is chordal	
Matchings	432
igraph_is_matching — Checks whether the given matching is valid for the	
given graph.	432
igraph_is_maximal_matching — Checks whether a matching in a graph	
is maximal.	432
igraph_maximum_bipartite_matching — Calculates a maximum	
matching in a bipartite graph.	433
Unfolding a graph into a tree	434
igraph_unfold_tree — Unfolding a graph into a tree, by possibly multipli-	
cating its vertices.	434
Other operations	
igraph_density — Calculate the density of a graph	
igraph_reciprocity — Calculates the reciprocity of a directed graph	
igraph_diversity — Structural diversity index of the vertices	
igraph_is_mutual — Check whether some edges of a directed graph are mu-	
tual.	437
igraph_avg_nearest_neighbor_degree — Average neighbor degree	
igraph_get_adjacency — The adjacency matrix of a graph	
igraph_get_adjacency_sparse — Returns the adjacency matrix of a	750
graph in a sparse matrix format.	<i>11</i> 0
igraph_get_stochastic — Stochastic adjacency matrix of a graph	
igraph_get_stochastic_sparse — The stochastic adjacency matrix of a	440
graph graph. graph.	441
igraph_get_edgelist — The list of edges in a graph.	
igraph_is_acyclic — Checks whether a graph is acyclic or not	
Deprecated functions	
igraph_shortest_paths — Length of the shortest paths between vertices	442
igraph_shortest_paths_dijkstra — Weighted shortest path lengths	4.40
between vertices (deprecated).	443
igraph_shortest_paths_bellman_ford — Weighted shortest path	4.40
lengths between vertices, allowing negative weights (deprecated)	443
igraph_shortest_paths_johnson — Weighted shortest path lengths be-	4.40
tween vertices, using Johnson's algorithm (deprecated).	443
igraph_get_stochastic_sparsemat — Stochastic adjacency matrix of a	
graph (deprecated).	444
igraph_get_sparsemat — Converts an igraph graph to a sparse matrix	
` 1	444
igraph_laplacian — Returns the Laplacian matrix of a graph (deprecated)	
14. Graph cycles	
Eulerian cycles and paths	
igraph_is_eulerian — Checks whether an Eulerian path or cycle exists	
igraph_eulerian_cycle — Finds an Eulerian cycle	
igraph_eulerian_path — Finds an Eulerian path	
15. Graph visitors	
Breadth-first search	
igraph_bfs — Breadth-first search	
igraph_bfs_simple — Breadth-first search, single-source version	449
igraph_bfshandler_t — Callback type for BFS function	450
Depth-first search	
igraph_dfs — Depth-first search	451
igraph_dfshandler_t — Callback type for the DFS function	
Random walks	453

igraph_random_walk — Performs a random walk on a graph	
Deprecated functions	454
igraph_random_edge_walk — Performs a random walk on a graph and re-	
turns the traversed edges.	
16. Cliques and independent vertex sets	
Cliques	
igraph_cliques — Finds all or some cliques in a graph	455
igraph_clique_size_hist — Counts cliques of each size in the graph	455
igraph_cliques_callback — Calls a function for each clique in the	
graph.	456
igraph_clique_handler_t — Type of clique handler functions	457
igraph_largest_cliques — Finds the largest clique(s) in a graph	
igraph_maximal_cliques — Finds all maximal cliques in a graph	
igraph_maximal_cliques_count — Count the number of maximal	
cliques in a graph	459
igraph_maximal_cliques_file — Find maximal cliques and write them	
to a file.	460
igraph_maximal_cliques_subset — Maximal cliques for a subset of	700
initial vertices	460
igraph_maximal_cliques_hist — Counts the number of maximal	- U(
cliques of each size in a graph.	46
igraph_maximal_cliques_callback — Finds maximal cliques in a	+0.
graph and calls a function for each one.	16
• •	
igraph_clique_number — Finds the clique number of the graph	
Weighted cliques	40.
igraph_weighted_cliques — Finds all cliques in a given weight range in	10
a vertex weighted graph.	46.
igraph_largest_weighted_cliques — Finds the largest weight	a -
clique(s) in a graph.	464
igraph_weighted_clique_number — Finds the weight of the largest	
weight clique in the graph.	
Independent vertex sets	465
igraph_independent_vertex_sets — Finds all independent vertex sets	
in a graph.	465
igraph_largest_independent_vertex_sets — Finds the largest in-	
dependent vertex set(s) in a graph.	466
igraph_maximal_independent_vertex_sets — Finds all maximal in-	
dependent vertex sets of a graph.	46
igraph_independence_number — Finds the independence number of the	
graph.	46
17. Graph isomorphism	
The simple interface	
igraph_isomorphic — Are two graphs isomorphic?	
igraph_subisomorphic — Decide subgraph isomorphism	
The BLISS algorithm	
igraph_bliss_sh_t — Splitting heuristics for Bliss	
igraph_bliss_info_t — Information about a BLISS run	
igraph_canonical_permutation — Canonical permutation using Bliss	
igraph_isomorphic_bliss — Graph isomorphism via Bliss	
igraph_automorphisms — Number of automorphisms using Bliss	
igraph_automorphism_group — Automorphism group generators using	+/.
	17
Bliss.	
The VF2 algorithm	
igraph_isomorphic_vf2 — Isomorphism via VF2	47.
igraph_count_isomorphisms_vf2 — Number of isomorphisms via	. –
VF2.	47
igraph_get_isomorphisms_vf2 — Collect all isomorphic mappings of	
two graphs.	47

igrapn_get_isomorphisms_vi2_callback — The generic VF2 inter-
face
igraph_isohandler_t — Callback type, called when an isomorphism was
found
igraph_isocompat_t — Callback type, called to check whether two vertices
or edges are compatible
$\verb igraph_subisomorphic_vf2Decide subgraph isomorphism using VF2 .$
igraph_count_subisomorphisms_vf2 — Number of subgraph isomor-
phisms using VF2
igraph_get_subisomorphisms_vf2 — Return all subgraph isomorphic
mappings.
igraph_get_subisomorphisms_vf2_callback—Generic VF2 func-
tion for subgraph isomorphism problems.
Deprecated aliases
The LAD algorithm
igraph_subisomorphic_lad — Check subgraph isomorphism with the
LAD algorithm
Functions for small graphs
igraph_isoclass — Determine the isomorphism class of small graphs
igraph_isoclass_subgraph — The isomorphism class of a subgraph of a
graph
igraph_isoclass_create — Creates a graph from the given isomorphism
class.
igraph_graph_count — The number of unlabelled graphs on the given
number of vertices.
Utility functions
igraph_permute_vertices — Permute the vertices
igraph_simplify_and_colorize — Simplify the graph and compute
self-loop and edge multiplicities.
Deprecated functions
igraph_isomorphic_34 — Graph isomorphism for 3-4 vertices (deprecated
alias).
18. Graph coloring
igraph_vertex_coloring_greedy — Computes a vertex coloring using a
greedy algorithm.
igraph_coloring_greedy_t — Ordering heuristics for greedy graph coloring
igraph_is_perfect — Checks if the graph is perfect
19. Graph motifs, dyad census and triad census
igraph_dyad_census — Dyad census, as defined by Holland and Leinhardt
igraph_triad_census — Triad census, as defined by Davis and Leinhardt
Finding triangles
igraph_adjacent_triangles — Count the number of triangles a vertex is
part of
igraph_list_triangles — Find all triangles in a graph
Graph motifs
igraph_motifs_randesu — Count the number of motifs in a graph
igraph_motifs_randesu_no — Count the total number of motifs in a
graph.
igraph_motifs_randesu_estimate — Estimate the total number of mo-
tifs in a graph.
igraph_motifs_randesu_callback — Finds motifs in a graph and calls
a function for each of them.
igraph_motifs_handler_t — Callback type for igraph_motif-
s_randesu_callback
20. Generating layouts for graph drawing
2D layout generators
igraph_layout_random — Places the vertices uniform randomly on a
plane

igraph_layout_circle — Places the vertices uniformly on a circle in arbi-	
trary order.	
· · · · · · · · · · · · · · ·	503
igraph_layout_grid — Places the vertices on a regular grid on the plane	503
igraph_layout_graphopt — Optimizes vertex layout via the graphopt al-	504
gorithm.	
igraph_layout_bipartite — Simple layout for bipartite graphs	
The DrL layout generator	505
igraph_layout_fruchterman_reingold — Places the vertices on a	500
	509
igraph_layout_kamada_kawai — Places the vertices on a plane according	<i>5</i> 1 1
ϵ	511
igraph_layout_gem — Layout graph according to GEM algorithm.	
igraph_layout_davidson_harel — Davidson-Harel layout algorithm	312
igraph_layout_mds — Place the vertices on a plane using multidimensional	51 4
scaling.	
igraph_layout_lgl — Force based layout algorithm for large graphs	
Layouts for trees and acyclic graphs	515
igraph_layout_reingold_tilford — Reingold-Tilford layout for tree	~1~
graphs.	515
igraph_layout_reingold_tilford_circular — Circular Rein-	~1.
gold-Tilford layout for trees.	
igraph_roots_for_tree_layout — Roots suitable for a nice tree layout	51/
igraph_layout_sugiyama — Sugiyama layout algorithm for layered direct-	<i>E</i> 1 0
<i>y U</i> 1	518
igraph_layout_umap — Layout using Uniform Manifold Approximation	<i>5</i> 10
3	519
, e	521
igraph_layout_random_3d — Places the vertices uniform randomly in a	521
cube	321
	521
igraph_layout_grid_3d — Places the vertices on a regular grid in the 3D	<i>J</i> 21
space.	522
igraph_layout_fruchterman_reingold_3d — 3D Fruchterman-Rein-	322
	522
igraph_layout_kamada_kawai_3d — 3D version of the Kamada-Kawai	322
layout generator.	523
igraph_layout_umap_3d — 3D layout using UMAP.	
Merging layouts	
igraph_layout_merge_dla — Merges multiple layouts by using a DLA al-	323
gorithm.	525
21. Reading and writing graphs from and to files	
Simple edge list and similar formats	
igraph_read_graph_edgelist — Reads an edge list from a file and cre-	321
ates a graph.	527
igraph_write_graph_edgelist — Writes the edge list of a graph to a	321
file.	527
igraph_read_graph_ncol — Reads an .ncol file used by LGL	
igraph_write_graph_ncol — Writes the graph to a file in .ncol format	
igraph_read_graph_lgl — Reads a graph from an .lgl file	
igraph_write_graph_lgl — Writes the graph to a file in .lgl format	
igraph_read_graph_dimacs_flow — Read a graph in DIMACS format	
igraph_write_graph_dimacs_flow — Write a graph in DIMACS for-	
mat.	532
Binary formats	
igraph_read_graph_graphdb — Read a graph in the binary graph data-	
base format.	533

GraphML format	
igraph_read_graph_graphml — Reads a graph from a GraphML file	534
igraph_write_graph_graphml — Writes the graph to a file in GraphML	
format.	534
GML format	
igraph_read_graph_gml — Read a graph in GML format	
igraph_write_graph_gml — Write the graph to a stream in GML format	
Pajek format	
igraph_read_graph_pajek — Reads a file in Pajek format	
igraph_write_graph_pajek — Writes a graph to a file in Pajek format	
UCINET's DL file format	
igraph_read_graph_dl — Reads a file in the DL format of UCINET	
Graphviz format	
igraph_write_graph_dot — Write the graph to a stream in DOT format	
Convenience functions for locale change	540
igraph_enter_safelocale — Temporarily set the C locale	540
igraph_exit_safelocale — Temporarily set the C locale	
Deprecated functions	
igraph_read_graph_dimacs — Read a graph in DIMACS format (depre-	٠
cated alias).	5/12
igraph_write_graph_dimacs — Write a graph in DIMACS format (dep-	372
· · · · · · · · · · · · · · · ·	<i>5</i> 42
recated alias).	
22. Maximum flows, minimum cuts and related measures	
Maximum flows	
igraph_maxflow — Maximum network flow between a pair of vertices	543
igraph_maxflow_value — Maximum flow in a network with the push/rela-	
bel algorithm.	544
igraph_dominator_tree — Calculates the dominator tree of a flowgraph	545
igraph_maxflow_stats_t — A simple data type to return some statistics	
	546
Cuts and minimum cuts	
igraph_st_mincut — Minimum cut between a source and a target vertex	
igraph_st_mincut_value — The minimum s-t cut in a graph	
igraph_all_st_cuts — List all edge-cuts between two vertices in a directed	517
graphgraph are the state of the state	510
igraph_all_st_mincuts — All minimum s-t cuts of a directed graph	
igraph_mincut — Calculates the minimum cut in a graph	
igraph_mincut_value — The minimum edge cut in a graph	
igraph_gomory_hu_tree — Gomory-Hu tree of a graph	
Connectivity	552
igraph_st_edge_connectivity — Edge connectivity of a pair of ver-	
tices.	552
igraph_edge_connectivity — The minimum edge connectivity in a	
graph.	553
igraph_st_vertex_connectivity — The vertex connectivity of a pair of	
vertices.	553
igraph_vertex_connectivity — The vertex connectivity of a graph	
Edge- and vertex-disjoint paths	
igraph_edge_disjoint_paths — The maximum number of edge-disjoint	333
	555
paths between two vertices.	223
igraph_vertex_disjoint_paths — Maximum number of vertex-disjoint	
paths between two vertices.	
Graph adhesion and cohesion	556
igraph_adhesion — Graph adhesion, this is (almost) the same as edge con-	
nectivity.	
igraph_cohesion — Graph cohesion, this is the same as vertex connectivity	557
Cohesive blocks	558

igraph_cohesive_blocks — Identifies the hierarchical cohesive block	
structure of a graph	
3. Vertex separators	55
igraph_is_separator — Would removing this set of vertices disconnect the graph?	55
igraph_is_minimal_separator — Decides whether a set of vertices is a mini-	
mal separator	55
igraph_all_minimal_st_separators — List all vertex sets that are minimal	
(s,t) separators for some s and t.	56
igraph_minimum_size_separators — Find all minimum size separating ver-	
tex sets.	56
igraph_even_tarjan_reduction — Even-Tarjan reduction of a graph	56
4. Detecting community structure	
Common functions related to community structure	
igraph_modularity — Calculates the modularity of a graph with respect to	
some clusters or vertex types.	56
igraph_modularity_matrix — Calculates the modularity matrix	
igraph_community_optimal_modularity — Calculate the community	50
	<i>5 (</i>
structure with the highest modularity value	30
igraph_community_to_membership — Creates a membership vector	
from a community structure dendrogram.	56
igraph_reindex_membership — Makes the IDs in a membership vector	_
contiguous.	56
igraph_compare_communities — Compares community structures using	
various metrics.	56
igraph_split_join_distance — Calculates the split-join distance of two	
community structures.	
Community structure based on statistical mechanics	57
igraph_community_spinglass — Community detection based on statisti-	
cal mechanics.	57
igraph_community_spinglass_single — Community of a single node	
based on statistical mechanics.	57
Community structure based on eigenvectors of matrices	57
igraph_community_leading_eigenvector — Leading eigenvector	
community finding (proper version).	57
igraph_community_leading_eigenvector_callback_t — Call-	
back for the leading eigenvector community finding method.	57
igraph_le_community_to_membership — Vertex membership from the	
leading eigenvector community structure	
Walktrap: Community structure based on random walks	
igraph_community_walktrap — This function is the implementation of	5
the Walktrap community	57
Edge betweenness based community detection	
	31
igraph_community_edge_betweenness — Community finding based on	
edge betweenness.	5
igraph_community_eb_get_merges — Calculating the merges, i.e. the	
dendrogram for an edge betweenness community structure.	
Community structure based on the optimization of modularity	58
igraph_community_fastgreedy — Finding community structure by	
greedy optimization of modularity.	58
igraph_community_multilevel — Finding community structure by mul-	
ti-level optimization of modularity.	58
igraph_community_leiden — Finding community structure using the Lei-	
den algorithm.	58
Fluid communities	
igraph_community_fluid_communities — Community detection based	
on fluids interacting on the graph.	. 58
Label propagation	

igraph_community_label_propagation — Community detection based	
on label propagation.	. 585
The InfoMAP algorithm	. 587
igraph_community_infomap — Find community structure that minimizes	
the expected description length of a random walker trajectory	
25. Graphlets	
Introduction	589
Performing graphlet decomposition	
igraph_graphlets — Calculate graphlets basis and project the graph on it	. 589
igraph_graphlets_candidate_basis — Calculate a candidate	
graphlets basis	. 590
igraph_graphlets_project — Project a graph on a graphlets basis	590
26. Hierarchical random graphs	. 592
Introduction	592
Representing HRGs	. 592
igraph_hrg_t — Data structure to store a hierarchical random graph	. 592
igraph_hrg_init — Allocate memory for a HRG	. 593
igraph_hrg_destroy — Deallocate memory for an HRG	. 593
igraph_hrg_size — Returns the size of the HRG, the number of leaf nodes	. 593
igraph_hrg_resize — Resize a HRG	. 594
Fitting HRGs	. 594
igraph_hrg_fit — Fit a hierarchical random graph model to a network	. 594
igraph_hrg_consensus — Calculate a consensus tree for a HRG	
HRG sampling	595
igraph_hrg_sample — Sample from a hierarchical random graph model	. 595
igraph_hrg_game — Generate a hierarchical random graph	. 596
Conversion to and from igraph graphs	. 596
igraph_hrg_dendrogram — Create a dendrogram from a hierarchical ran-	
dom graph.	596
igraph_hrg_create — Create a HRG from an igraph graph	597
Predicting missing edges	. 597
igraph_hrg_predict — Predict missing edges in a graph, based on HRG	
models.	597
27. Embedding of graphs	599
Spectral embedding	
igraph_adjacency_spectral_embedding — Adjacency spectral em-	
bedding	. 599
igraph_laplacian_spectral_embedding — Spectral embedding of the	
Laplacian of a graph	
igraph_dim_select — Dimensionality selection	601
28. Graph operators	. 603
Union and intersection	. 603
igraph_disjoint_union — Creates the union of two disjoint graphs	. 603
igraph_disjoint_union_many — The disjint union of many graphs	
igraph_union — Calculates the union of two graphs	. 604
igraph_union_many — Creates the union of many graphs	. 605
igraph_intersection — Collect the common edges from two graphs	
igraph_intersection_many — The intersection of more than two graphs	
Other set-like operators	
igraph_difference — Calculates the difference of two graphs	
igraph_complementer — Creates the complementer of a graph	
igraph_compose — Calculates the composition of two graphs	
Miscellaneous operators	
igraph_connect_neighborhood — Graph power: connect each vertex to	
its neighborhood.	609
igraph_contract_vertices — Replace multiple vertices with a single	
one.	609

igraph_induced_subgraph — Creates a subgraph induced by the specified	
vertices.	
igraph_linegraph — Create the line graph of a graph	611
igraph_simplify — Removes loop and/or multiple edges from the graph	
igraph_subgraph_edges — Creates a subgraph with the specified edges	
and their endpoints.	612
igraph_reverse_edges — Reverses some edges of a directed graph	
29. Using BLAS, LAPACK and ARPACK for igraph matrices and graphs	
BLAS interface in igraph	
igraph_blas_ddot — Dot product of two vectors	
igraph_blas_dnrm2 — Euclidean norm of a vector.	
igraph_blas_dgemv — Matrix-vector multiplication using BLAS, vector	011
version.	615
igraph_blas_dgemm — Matrix-matrix multiplication using BLAS	
	013
igraph_blas_dgemv_array — Matrix-vector multiplication using BLAS,	C1 C
array version.	
LAPACK interface in igraph	
Matrix factorization, solving linear systems	
Eigenvalues and eigenvectors of matrices	
ARPACK interface in igraph	
Data structures	
ARPACK solvers	
30. Bipartite, i.e. two-mode graphs	632
Bipartite networks in igraph	632
Create two-mode networks	632
igraph_create_bipartite — Create a bipartite graph	632
igraph_full_bipartite — Create a full bipartite network	633
igraph_bipartite_game — Generate a bipartite random graph (similar to	
Erd#s-Rényi).	633
Incidence matrices	. 634
igraph_incidence — Creates a bipartite graph from an incidence matrix	
igraph_get_incidence — Convert a bipartite graph into an incidence ma-	
trix.	635
Project two-mode graphs	
igraph_bipartite_projection_size — Calculate the number of ver-	. 050
tices and edges in the bipartite projections.	636
igraph_bipartite_projection — Create one or both projections of a bi-	. 050
· ·	. 637
partite (two-mode) network.	
Other operations on bipartite graphs	
igraph_is_bipartite — Check whether a graph is bipartite	
31. Advanced igraph programming	
Using igraph in multi-threaded programs	639
IGRAPH_THREAD_SAFE — Specifies whether igraph was built in thread-safe	
mode	
Thread-safe ARPACK library	
Thread-safety of random number generators	
Progress handlers	
About progress handlers	639
Setting up progress handlers	640
Invoking the progress handler	641
Writing progress handlers	642
Writing igraph functions with progress reporting	
Multi-threaded programs	
Status handlers	
Status reporting	
Setting up status handlers	
Invoking the status handler	
32. Non-graph related functions	

	igraph version number	647
	igraph_version — Return the version of the igraph C library	647
	Running mean of a time series	647
	igraph_running_mean — Calculates the running mean of a vector	647
	Random sampling from very long sequences	648
	igraph_random_sample — Generates an increasing random sequence of in-	
	tegers.	648
	Random sampling of spatial points	649
	igraph_sample_sphere_surface — Sample points uniformly from the	
	surface of a sphere.	649
	igraph_sample_sphere_volume — Sample points uniformly from the	
	volume of a sphere.	649
	igraph_sample_dirichlet — Sample points from a Dirichlet distribution	
	Convex hull of a set of points on a plane	
	igraph_convex_hull — Determines the convex hull of a given set of points	
	in the 2D plane.	650
	Fitting power-law distributions to empirical data	
	igraph_plfit_result_t — Result of fitting a power-law distribution to a	001
	vector	651
	igraph_power_law_fit — Fits a power-law distribution to a vector of num-	001
	bers.	652
	igraph_plfit_result_calculate_p_value — Calculates the p-value	032
	of a fitted power-law model.	653
	Comparing floats with a tolerance	
	igraph_cmp_epsilon — Compare two double-precision floats with a toler-	054
	ance	654
	igraph_almost_equals — Compare two double-precision floats with a tol-	054
	erance.	654
	igraph_complex_almost_equals — Compare two complex numbers	054
	with a tolerance.	654
33 I	icenses for igraph and this manual	
JJ. L	THE GNU GENERAL PUBLIC LICENSE	
	Preamble	
	GNU GENERAL PUBLIC LICENSE	
	How to Apply These Terms to Your New Programs	
	The GNU Free Documentation License	
	0. PREAMBLE	
	1. APPLICABILITY AND DEFINITIONS	
	2. VERBATIM COPYING	
	3. COPYING IN QUANTITY	
	4. MODIFICATIONS	
	5. COMBINING DOCUMENTS	
	6. COLLECTIONS OF DOCUMENTS	
	7. AGGREGATION WITH INDEPENDENT WORKS	
	8. TRANSLATION	
	9. TERMINATION 10. FUTURE REVISIONS OF THIS LICENSE	
[m.d	G.1.1 ADDENDUM: How to use this License for your documents	
mue>	K	000

List of Examples

4.1. File examples/simple/creation.c	. 16
4.2. File examples/simple/igraph_copy.c	. 17
4.3. File examples/simple/igraph_is_directed.c	. 19
4.4. File examples/simple/igraph_get_eid.c	
4.5. File examples/simple/igraph_get_eids.c	. 23
4.6. File examples/simple/igraph_neighbors.c	. 24
4.7. File examples/simple/igraph_degree.c	. 25
4.8. File examples/simple/creation.c	. 26
4.9. File examples/simple/creation.c	. 27
4.10. File examples/simple/igraph_delete_edges.c	27
4.11. File examples/simple/igraph_delete_vertices.c	28
4.12. File examples/simple/igraph_delete_vertices.c	28
7.1. File examples/simple/igraph_fisher_yates_shuffle.c	59
7.2. File examples/simple/igraph_vector_int_list_sort.c	65
7.3. File examples/simple/igraph_vector_int_list_sort.c	66
7.4. File examples/simple/igraph_sparsemat.c	121
7.5. File examples/simple/igraph_sparsemat3.c	
7.6. File examples/simple/igraph_sparsemat4.c	121
7.7. File examples/simple/igraph_sparsemat6.c	
7.8. File examples/simple/igraph_sparsemat7.c	
7.9. File examples/simple/igraph_sparsemat8.c	121
7.10. File examples/simple/dqueue.c	154
7.11. File examples/simple/igraph_strvector.c	161
7.12. File examples/simple/adjlist.c	
8.1. File examples/simple/random_seed.c	206
9.1. File examples/simple/igraph_create.c	
9.2. File examples/simple/igraph_small.c	
9.3. File examples/simple/igraph_adjacency.c	
9.4. File examples/simple/igraph_weighted_adjacency.c	
9.5. File examples/simple/igraph_star.c	
9.6. File examples/simple/igraph_ring.c	
9.7. File examples/simple/igraph_kary_tree.c	
9.8. File examples/simple/igraph_symmetric_tree.c	
9.9. File examples/simple/igraph_regular_tree.c	
9.10. File examples/simple/igraph_full.c	
9.11. File examples/simple/igraph_realize_degree_sequence.c	
9.12. File examples/simple/igraph_lcf.c	
9.13. File examples/simple/igraph_atlas.c	
9.14. File examples/simple/igraph_grg_game.c	
9.15. File examples/simple/igraph_barabasi_game.c	
9.16. File examples/simple/igraph_barabasi_game2.c	
9.17. File examples/simple/igraph_erdos_renyi_game.c	
9.18. File examples/simple/igraph_degree_sequence_game.c	
10.1. File examples/simple/igraph_deterministic_optimal_imitation.c.	
10.2. File examples/simple/igraph_roulette_wheel_imitation.c	
10.3. File examples/simple/igraph_stochastic_imitation.c	
11.1. File examples/simple/igraph_vs_nonadj.c	
11.2. File examples/simple/igraph_vs_vector.c	
11.3. File examples/simple/igraph_vs_seq.c	
11.4. File examples/simple/igraph_es_pairs.c	
11.5. File examples/simple/igraph_vs_seq.c	
12.1. File examples/simple/igraph_attribute_combination.c	
12.2. File examples/simple/cattributes.c	
12.3. File examples/simple/cattributes2.c	299 299
17.4 PDE EXAMOTES/STUDIE/CALL CLOUL ASS C	/ 44

12.5. File examples/simple/cattributes4.c	
13.1. File examples/simple/dijkstra.c	
13.2. File examples/simple/bellman_ford.c	
13.3. File examples/simple/igraph_get_shortest_paths.c	
13.4. File examples/simple/igraph_get_shortest_paths_dijkstra.c	
13.5. File examples/simple/igraph_get_all_shortest_paths_dijkstra.c.	343
13.6. File examples/simple/igraph_average_path_length.c	345
13.7. File examples/simple/igraph_grg_game.c	346
13.8. File examples/simple/igraph_diameter.c	348
13.9. File examples/simple/igraph_girth.c	
13.10. File examples/simple/igraph_eccentricity.c	350
13.11. File examples/simple/igraph_radius.c	352
13.12. File examples/simple/igraph_decompose.c	370
13.13. File examples/simple/igraph_biconnected_components.c	
13.14. File examples/simple/igraph_pagerank.c	380
13.15. File examples/simple/eigenvector_centrality.c	386
13.16. File examples/simple/centralization.c	395
13.17. File examples/simple/igraph_cocitation.c	402
13.18. File examples/simple/igraph_cocitation.c	403
13.19. File examples/simple/igraph_similarity.c	404
13.20. File examples/simple/igraph_similarity.c	404
13.21. File examples/simple/igraph_similarity.c	405
13.22. File examples/simple/igraph_similarity.c	406
13.23. File examples/simple/igraph_similarity.c	407
13.24. File examples/simple/igraph_similarity.c	408
13.25. File examples/simple/igraph_similarity.c	409
13.26. File examples/simple/igraph_minimum_spanning_tree.c	410
13.27. File examples/simple/igraph_minimum_spanning_tree.c	411
13.28. File examples/simple/igraph_kary_tree.c	413
13.29. File examples/simple/igraph_transitivity.c	415
13.30. File examples/simple/igraph_to_undirected.c	418
13.31. File examples/simple/igraph_get_laplacian.c	419
13.32. File examples/simple/igraph_get_laplacian_sparse.c	420
13.33. File examples/simple/igraph_is_loop.c	422
13.34. File examples/simple/igraph_is_multiple.c	422
13.35. File examples/simple/igraph_has_multiple.c	423
13.36. File examples/simple/igraph_assortativity_nominal.c	425
13.37. File examples/simple/igraph_assortativity_degree.c	427
13.38. File examples/simple/igraph_topological_sorting.c	429
13.39. File examples/simple/igraph_feedback_arc_set.c	430
13.40. File examples/simple/igraph_feedback_arc_set_ip.c	430
13.41. File examples/simple/igraph_maximum_bipartite_matching.c	432
13.42. File examples/simple/igraph_maximum_bipartite_matching.c	433
13.43. File examples/simple/igraph_maximum_bipartite_matching.c	434
13.44. File examples/simple/igraph_reciprocity.c	436
13.45. File examples/simple/igraph_knn.c	438
15.1. File examples/simple/igraph_bfs.c	449
15.2. File examples/simple/igraph_bfs_callback.c	449
15.3. File examples/simple/igraph_bfs_simple.c	450
16.1. File examples/simple/igraph_cliques.c	455
16.2. File examples/simple/igraph_maximal_cliques.c	
16.3. File examples/simple/igraph_maximal_cliques.c	
16.4. File examples/simple/igraph_independent_sets.c	
17.1. File examples/simple/igraph_isomorphic_vf2.c	
17.2. File examples/simple/igraph_subisomorphic_lad.c	
18.1. File examples/simple/igraph_coloring.c	
19.1. File examples/simple/igraph_list_triangles.c	
19.2. File examples/simple/igraph_motifs_randesu.c	

19.3. File examples/simple/igrapn_motils_randesu.c :	500
20.1. File examples/simple/igraph_layout_reingold_tilford.c :	516
21.1. File examples/simple/igraph_read_graph_lgl.c	531
21.2. File examples/simple/igraph_write_graph_lgl.c	531
21.3. File examples/simple/igraph_read_graph_graphdb.c	534
21.4. File examples/simple/graphml.c	534
21.5. File examples/simple/graphml.c	535
21.6. File examples/simple/gml.c	536
21.7. File examples/simple/gml.c	537
21.8. File examples/simple/foreign.c	538
21.9. File examples/simple/igraph_write_graph_pajek.c	539
21.10. File examples/simple/igraph_read_graph_dl.c	
21.11. File examples/simple/dot.c	
21.12. File examples/simple/safelocale.c	
22.1. File examples/simple/flow.c	
22.2. File examples/simple/flow2.c	
22.3. File examples/simple/dominator_tree.c	
22.4. File examples/simple/igraph_all_st_mincuts.c	
22.5. File examples/simple/igraph_mincut.c	
22.6. File examples/simple/cohesive_blocks.c	
23.1. File examples/simple/igraph_is_separator.c	
23.2. File examples/simple/igraph_is_minimal_separator.c	
23.3. File examples/simple/igraph_minimal_separators.c	
23.4. File examples/simple/igraph_minimum_size_separators.c	
23.5. File examples/simple/even_tarjan.c	
24.1. File examples/simple/igraph_community_optimal_modularity.c	
24.2. File examples/simple/igraph_community_leading_eigenvector.c	
24.3. File examples/simple/walktrap.c	
24.4. File examples/simple/igraph_community_edge_betweenness.c	
24.5. File examples/simple/igraph_community_fastgreedy.c	
24.6. File examples/simple/igraph_community_multilevel.c	
24.7. File examples/simple/igraph_community_leiden.c	
24.8. File examples/simple/igraph_community_label_propagation.c	
28.1. File examples/simple/igraph_disjoint_union.c	
28.2. File examples/simple/igraph_union.c	
28.3. File examples/simple/igraph_union.c	
28.4. File examples/simple/igraph_intersection.c	
	607
28.6. File examples/simple/igraph_complementer.c	608
28.7. File examples/simple/igraph_compose.c	
28.8. File examples/simple/igraph_simplify.c	
29.1. File examples/simple/blas.c	
29.2. File examples/simple/blas.c	
29.3. File examples/simple/blas_dgemm.c	
29.4. File examples/simple/igraph_lapack_dgesv.c	
29.5. File examples/simple/igraph_lapack_dsyevr.c	
29.6. File examples/simple/igraph_lapack_dgeev.c	
29.7. File examples/simple/igraph_lapack_dgeevx.c	
30.1. File examples/simple/igraph_bipartite_create.c	
30.2. File examples/simple/igraph_bipartite_projection.c	
30.3. File examples/simple/igraph_bipartite_projection.c	
32.1. File examples/simple/igraph_version.c	
32.2. File examples/simple/igraph_random_sample.c	
32.3. File examples/simple/igraph power law fit.c	

Chapter 1. Introduction

igraph is a library for creating and manipulating graphs. You can look at it in two ways: first, igraph contains the implementation of quite a lot of graph algorithms. These include classic graph algorithms like graph isomorphism, graph girth and connectivity and also the new wave graph algorithms like transitivity, graph motifs and community structure detection. Skim through the table of contents or the index of this book to get an impression of what is available.

Second, igraph provides a platform for developing and/or implementing graph algorithms. It has an efficient data structure for representing graphs, and a number of other data structures like flexible vectors, stacks, heaps, queues, adjacency lists that are useful for implementing graph algorithms. In fact these data structures evolved along with the implementation of the classic and non-classic graph algorithms which make up the major part of the igraph library. This way, they were fine-tuned and checked for correctness several times.

Our main goal with developing igraph was to create a graph library which is efficient on large, but not extremely large graphs. More precisely, it is assumed that the graph(s) fit into the physical memory of the computer. Nowadays this means graphs with several million vertices and/or edges. Our definition of efficient is that it runs fast, both in theory and (more importantly) in practice.

We believe that one of the big strengths of igraph is that it can be embedded into a higher-level language or environment. Three such embeddings (or interfaces if you look at them another way) are currently being developed by us: an R package, a Python extension module, and a Mathematica (Wolfram Language) package. Others are likely to come. High level languages such as R or Python make it possible to use graph routines with much greater comfort, without actually writing a single line of C code. They have some, usually very small, speed penalty compared to the C version, but add ease of use and much flexibility. This manual, however, covers only the C library. If you want to use Python, R or the Wolfram Language, please see the documentation written specifically for these interfaces and come back here only if you are interested in some detail which is not covered in those documents.

We still consider igraph as a child project. It has much room for development and we are sure that it will improve a lot in the near future. Any feedback we can get from the users is very important for us, as most of the time these questions and comments guide us in what to add and what to improve.

igraph is open source and distributed under the terms of the GNU GPL. We strongly believe that all the algorithms used in science, let that be graph theory or not, should have an efficient open-source implementation allowing use and modification for anyone.

igraph is free software

igraph library

Copyright (C) 2003-2004 Gábor Csárdi <csardi.gabor@gmail.com>

Copyright (C) 2005-2019 Gábor Csárdi <csardi.gabor@gmail.com> and Tamás Nepusz <ntamas@gmail.com>

Copyright (C) 2020-2022 The igraph development team

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc.

Citing igraph

To cite igraph in publications, please use the following reference:

Gábor Csárdi, Tamás Nepusz: The igraph software package for complex network research. InterJournal Complex Systems, 1695, 2006.

The igraph C library is assigned the DOI 10.5281/zenodo.3630268 [https://doi.org/10.5281/zenodo.3630268] on Zenodo.

Chapter 2. Installation

This chapter describes building igraph from source code and installing it. The source archive of the latest stable release is always available from the igraph website [https://igraph.org/c/#downloads]. igraph is also included in many Linux distributions, as well as several package managers such as vcp-kg [https://vcpkg.io/] (convenient on Windows), MacPorts [https://www.macports.org/] (macOS) and Homebrew [https://brew.sh/] (macOS), which provide an easier means of installation. If you decide to use them, please consult their documentation on how to install packages.

Prerequisites

To build igraph from sources, you will need at least:

- CMake [https://cmake.org] 3.18 or later
- C and C++ compilers

Visual Studio 2015 and later are supported. Earlier Visual Studio versions may or may not work.

Certain features also require the following libraries:

• libxml2 [http://www.xmlsoft.org/], required for GraphML support

igraph bundles a number of libraries for convenience. However, it is preferable to use external versions of these libraries, which may improve performance. These are:

- GMP [https://gmplib.org/] (the bundled alternative is Mini-GMP)
- GLPK [https://www.gnu.org/software/glpk/] (version 4.57 or later)
- ARPACK [https://www.caam.rice.edu/software/ARPACK/]
- A library providing a BLAS [https://www.netlib.org/blas/] API (available by default on macOS; OpenBLAS [https://www.openblas.net] is one option on other systems)
- A library providing a LAPACK [https://www.netlib.org/lapack/] API (available by default on macOS; OpenBLAS [https://openblas.net] is one option on other systems)

When building the development version of igraph, bison, flex and git are also required. Released versions do not require these tools.

To run the tests, diff is also required.

Installation

General build instructions

igraph uses a CMake-based build system [https://cmake.org/cmake/help/latest/guide/user-interaction/index.html]. To compile it,

- Enter the directory where the igraph sources are:
 - \$ cd igraph
- Create a new directory. This is where igraph will be built:
 - \$ mkdir build

```
$ cd build
```

• Run CMake, which will automatically configure igraph, and report the configuration:

```
$ cmake ..
To set a non-default installation location, such as /opt/local, use:
```

```
cmake .. -DCMAKE_INSTALL_PREFIX=/opt/local
```

- Check the output carefully, and ensure that all features you need are enabled. If CMake could not
 find certain libraries, some features such as GraphML support may have been automatically disabled.
- There are several ways to adjust the configuration:
 - Run ccmake . on Unix-like systems or cmake-gui on Windows for a convenient interface.
 - Simply edit the CMakeCache.txt file. Some of the relevant options are listed below.
- Once the configuration has been adjusted, run cmake . . again.
- Once igraph has been successfully configured, it can be built, tested and installed using:

```
$ cmake --build .
$ cmake --build . --target check
$ cmake --install .
```

Specific instructions for Windows

Microsoft Visual Studio

With Visual Studio, the steps to build igraph are generally the same as above. However, since the Visual Studio CMake generator is a multi-configuration one, we must specify the configuration (typically Release or Debug) with each build command using the --config option:

```
mkdir build
cd build
cmake ..
cmake --build . --config Release
cmake --build . --target check --config Release
```

When building the development version, bison and flex must be available on the system. winflexbison [https://github.com/lexxmark/winflexbison] for Bison version 3.x can be useful for this purpose—make sure that the executables are in the system PATH. The easiest installation option is probably by installing winflexbison3 from the Chocolatey package manager [https://chocolatey.org/packages/winflexbison3].

vcpkg

Most external dependencies can be conveniently installed using vcpkg [https://github.com/microsoft/vcpkg#quick-start-windows]. Note that igraph bundles all dependencies except libxml2, which is needed for GraphML support.

In order to use vcpkg integrate it in the build environment by executing vcpkg.exe integrate install on the command line. When configuring igraph, point CMake to the correct vcpkg.c-make file using -DCMAKE TOOLCHAIN FILE=..., as instructed.

Additionally, it might be that you need to set the appropriate so-called triplet using -DVCPKG_TAR-GET_TRIPLET when running cmake, for exampling, setting it to x64-windows when using shared builds of packages or x64-windows-static when using static builds. Similarly, you also need to specify this target triplet when installing packages. For example, to install libxml2 as a shared library, use vcpkg.exe install libxml2:x64-windows and to install libxml2 as a static library, use vcpkg.exe install libxml2:x64-windows-static. In addition, there is the possibility to use a static library with dynamic runtime linking using the x64-windows-static-md triplet.

MSYS2

MSYS2 can be installed from msys2.org [https://www.msys2.org/]. After installing MSYS2, ensure that it is up to date by opening a terminal and running pacman -Syuu.

The instructions below assume that you want to compile for a 64-bit target.

Install the following packages using pacman -S.

- Minimal requirements: mingw-w64-x86_64-toolchain, mingw-w64-x86_64-cmake.
- Optional dependencies that enable certain features: mingw-w64-x86_64-gmp, mingw-w64-x86_64-libxml2
- Optional external libraries for better performance: mingw-w64-x86_64-openblas, mingw-w64-x86_64-suitesparse, mingw-w64-x86_64-arpack, mingw-w64-x86_64-glpk
- Only needed for running the tests: diffutils
- Required only when building the development version: git, bison, flex

The following command will install of these at once:

```
pacman -S \
  mingw-w64-x86_64-toolchain mingw-w64-x86_64-cmake \
  mingw-w64-x86_64-gmp mingw-w64-x86_64-libxml2 \
  mingw-w64-x86_64-openblas mingw-w64-x86_64-suitesparse mingw-w64-x86_64-arpac
  mingw-w64-x86_64-glpk diffutils git bison flex
```

In order to build igraph, follow the **General build instructions** above, paying attention to the following:

- When using MSYS2, start the "MSYS2 MinGW 64-bit" terminal, and not the "MSYS2 MSYS" one.
- Be sure to install the mingw-w64-x86_64-cmake package and not the cmake one. The latter will not work.
- When running cmake, pass the option -G"MSYS Makefiles".
- Note that ccmake is not currently available. cmake-gui can be used only if the mingw-w64-x86_64-qt5 package is installed.

Notable configuration options

The following options may be set to ON or OFF. Some of them have an AUTO setting, which chooses a reasonable default based on what libraries are available on the current system.

• igraph bundles some of its dependencies for convenience. The IGRAPH_USE_INTERNAL_XXX flags control whether these should be used instead of external versions. Set them to ON to use the

bundled ("vendored") versions. Generally, external versions are preferable as they may be newer and usually provide better performance.

- IGRAPH_GLPK_SUPPORT: whether to make use of the GLPK [https://www.gnu.org/soft-ware/glpk/] library. Some features, such as finding a minimum feedback arc set or finding communities through exact modularity optimization, require this.
- IGRAPH_GRAPHML_SUPPORT: whether to enable support for reading and writing GraphML [http://graphml.graphdrawing.org/] files. Requires the libxml2 [http://xmlsoft.org/] library.
- IGRAPH_OPENMP_SUPPORT: whether to use OpenMP parallelization to accelerate certain functions such as PageRank calculation. Compiler support is required.
- IGRAPH_ENABLE_LTO: whether to build igraph with link-time optimization, which improves performance. Not supported with all compilers.
- IGRAPH_ENABLE_TLS: whether to enable thread-local storage. Required when using igraph from multiple threads.
- IGRAPH_WARNINGS_AS_ERRORS: whether to treat compiler warnings as errors. We strive to
 eliminate all compiler warnings during development so this switch is turned on by default. If your
 compiler prints warnings for some parts of the code that we did not anticipate, you can turn off this
 option to prevent the warnings from stopping the compilation.
- BUILD_SHARED_LIBS [https://cmake.org/cmake/help/latest/variable/BUILD_SHARED_LIBS.html]: whether to build a shared library instead of a static one.
- BLA_VENDOR: controls which library to use for BLAS [https://cmake.org/cmake/help/latest/module/FindBLAS.html] and LAPACK [https://cmake.org/cmake/help/latest/module/FindLA-PACK.html] functionality.
- CMAKE_INSTALL_PREFIX [https://cmake.org/cmake/help/latest/variable/CMAKE_INSTAL-L_PREFIX.html]: the location where igraph will be installed.

Building the documentation

Most users will not need to build the documentation, as the release tarball contains pre-built HTML documentation in the doc directory.

To build the documentation for the development version, simply build the html or pdf targets for the HTML and PDF versions of the documentation, respectively.

```
$ cmake --build . --target html
```

Building the HTML documentation requires Python 3, xmlto and source-highlight. Building the PDF documentation also requires xsltproc, xmllint and fop.

Notes for package maintainers

This section is for people who package igraph for Linux distros or other package managers. Please read it carefully before packaging igraph.

Auto-detection of dependencies

igraph bundles several of its dependencies (or simplified versions of its dependencies). During configuration time, it checks whether each dependency is present on the system. If yes, it uses it. Otherwise, it falls back to the bundled ("vendored") version. In order to make configuration as deterministic as

possible, you may want to disable this auto-detection. To do so, set each of the IGRAPH_USE_INTERNAL_XXX option described above. Additionally, set BLA_VENDOR to use the BLAS and LAPACK implementations of your choice. This should be the same BLAS and LAPACK library that igraph's other dependencies (e.g., ARPACK) are linked against.

For example, to force igraph to use external versions of all dependencies, and to use OpenBLAS for BLAS/LAPACK, use

```
$ cmake .. \
    -DIGRAPH_USE_INTERNAL_BLAS=OFF \
    -DIGRAPH_USE_INTERNAL_LAPACK=OFF \
    -DIGRAPH_USE_INTERNAL_ARPACK=OFF \
    -DIGRAPH_USE_INTERNAL_GLPK=OFF \
    -DIGRAPH_USE_INTERNAL_GMP=OFF \
    -DBLA_VENDOR=OpenBLAS \
    -DIGRAPH_GRAPHML_SUPPORT=ON
```

Shared and static builds

On Windows, shared and static builds should not be installed in the same location. If you decide to do so anyway, keep in mind the following: Both builds contain an igraph.lib file. The static one should be renamed to avoid conflict. The headers from the static build are incompatible with the shared library. The headers from the shared build may be used with the static library, but IGRAPH_STATIC must be defined when compiling programs that will link to igraph statically.

These issues do not affect Unix-like systems.

Cross-compiling

When building igraph with an internal ARPACK, LAPACK or BLAS, it makes use of f2c, which compiles and runs the arithchk program at build time to detect the floating point characteristics of the current system. It writes the results into the arith.h header. However, running this program is not possible when cross-compiling without providing a userspace emulator that can run executables of the target platform on the host system. Therefore, when cross-compiling, you either need to provide such an emulator with the CMAKE_CROSSCOMPILING_EMULATOR option, or you need to specify a pre-generated version of the arith.h header file through the F2C_EXTERNAL_ARITH_HEADER CMake option. An example version of this header follows for the x86_64 and arm64 target architectures on macOS. Warning: Do not use this version of arith.h on other systems or architectures.

```
#define IEEE_8087
#define Arith_Kind_ASL 1
#define Long int
#define Intcast (int)(long)
#define Double_Align
#define X64_bit_pointers
#define NANCHECK
#define QNaN0 0x0
#define QNaN1 0x7ff80000
```

igraph also checks whether the endianness of uint64_t matches the endianness of double on the platform being compiled. This is needed to ensure that certain functions in igraph's random number generator work properly. However, it is not possible to execute this check when cross-compiling without an emulator, so in this case igraph simply assumes that the endianness matches (which is the case for the vast majority of platforms anyway). The only case where you might run into problems is when you cross-compile for Apple Silicon (arm64) from an Intel-based Mac, in which case CMake might not realize that you are cross-compiling and will try to execute the check anyway. You can

work around this by setting IEEE754_DOUBLE_ENDIANNESS_MATCHES to ON explicitly before invoking CMake.

Providing an emulator in CMAKE_CROSSCOMPILING_EMULATOR has the added benefit that you can run the compiled unit tests on the host platform. We have experimented with cross-compiling to 64-bit ARM CPUs (aarch64) on 64-bit Intel CPUs (amd64), and we can confirm that using qemu-aarch64 works as a cross-compiling emulator in this setup.

Additional notes

- As of igraph 0.10, there is no tangible benefit to using an external GMP, as igraph does not yet use GMP in any performance-critical way. The bundled Mini-GMP is sufficient.
- Link-time optimization noticeably improves the performance of some igraph functions. To enable it, use <code>-DIGRAPH_ENABLE_LTO=ON</code>. The AUTO setting is also supported, and will enable link-time optimization only if the current compiler supports it. Note that this is detected by CMake, and the detection is not always accurate.
- We saw occasional hangs on Windows when igraph was built for a 32-bit target with MinGW and linked to OpenBLAS. We believe this to be an issue with OpenBLAS, not igraph. On this platform, you may want to opt for a different BLAS/LAPACK or the bundled BLAS/LAPACK.

Chapter 3. Tutorial

Compiling programs using igraph

The following short example program demonstrates the basic usage of the **igraph** library.

```
#include <igraph.h>
int main() {
  igraph_integer_t num_vertices = 1000;
  igraph_integer_t num_edges = 1000;
  igraph_real_t diameter;
  igraph_t graph;
  igraph rng seed(igraph rng default(), 42);
  igraph_erdos_renyi_game(&graph, IGRAPH_ERDOS_RENYI_GNM,
                          num vertices, num edges,
                          IGRAPH UNDIRECTED, IGRAPH NO LOOPS);
  igraph_diameter(&graph, &diameter, NULL, NULL, NULL, NULL, IGRAPH_UNDIRECTED,
  printf("Diameter of a random graph with average degree %g: %g\n",
          2.0 * igraph_ecount(&graph) / igraph_vcount(&graph),
          (double) diameter);
  igraph destroy(&graph);
  return 0;
```

This example illustrates a couple of points:

- First, programs using the **igraph** library should include the igraph. h header file.
- Second, **igraph** uses the igraph_integer_t type for integers instead of int or long int, and it also uses the igraph_real_t type for real numbers instead of double. Depending on how **igraph** was compiled, and whether you are using a 32-bit or 64-bit system, igraph_integer_t may be a 32-bit or 64-bit integer.
- Third, **igraph** graph objects are represented by the igraph_t data type.
- Fourth, the igraph_erdos_renyi_game() creates a graph and igraph_destroy() destroys it, i.e. deallocates the memory associated to it.

For compiling this program you need a C compiler. Optionally, CMake [https://cmake.org] can be used to automate the compilation.

Compiling with CMake

It is convenient to use CMake because it can automatically discover the necessary compilation flags on all operating systems. Many IDEs support CMake, and can work with CMake projects directly. To create a CMake project for this example program, create a file name CMakeLists.txt with the following contents:

```
cmake_minimum_required(VERSION 3.18)
project(igraph test)
```

```
find_package(igraph REQUIRED)
add_executable(igraph_test igraph_test.c)
target_link_libraries(igraph_test PUBLIC igraph::igraph)
```

To compile the project, create a new directory called build in the root of the **igraph** source tree, and switch to it:

```
mkdir build cd build
```

Run CMake to configure the project:

```
cmake ..
```

If **igraph** was installed at a non-standard location, specify its prefix using the -DCMAKE_PRE-FIX_PATH=... option. The prefix must be the same directory that was specified as the CMAKE_INSTALL_PREFIX when compiling igraph.

If configuration has succeeded, build the program using

```
cmake --build .
```

C++ must be enabled in igraph projects

Parts of **igraph** are implemented in C++; therefore, any CMake target that depends on **igraph** should use the C++ linker. Furthermore, OpenMP support in igraph works correctly only if C++ is enabled in the CMake project. The script that finds **igraph** on the host machine will throw an error if C++ support is not enabled in the CMake project.

C++ support is enabled by default when no languages are explicitly specified in CMake's project command. If you do specify some languages explicitly, make sure to also include CXX.

Compiling without CMake

On most Unix-like systems, the default C compiler is called **cc**. To compile the test program, you will need a command similar to the following:

```
cc igraph_test.c -I/usr/local/include/igraph -L/usr/local/lib -ligraph -o igraph
```

The exact form depends on where **igraph** was installed on your system, whether it was compiled as a shared or static library, and the external libraries it was linked to. The directory after the -I switch is the one containing the igraph.h file, while the one following -L should contain the library file itself, usually a file called libigraph.a (static library on macOS and Linux), libigraph.so (shared library on Linux), libigraph.dylib (shared library on macOS), igraph.lib (static library on Windows) or igraph.dll (shared library on Windows). If **igraph** was compiled as a static library, it is also necessary to manually link to all of its dependencies.

If your system has the **pkg-config** utility you are likely to get the necessary compile options by issuing the command

```
pkg-config --libs --cflags igraph
```

(if igraph was built as a shared library) or

```
pkg-config --static --libs --cflags igraph (if igraph was built as a static library).
```

Running the program

On most systems, the executable can be run by simply typing its name like this:

```
./igraph_test
```

If you use dynamic linking and the **igraph** library is not in a standard place, you may need to add its location to the LD_LIBRARY_PATH (Linux), DYLD_LIBRARY_PATH (macOS) or PATH (Windows) environment variables.

Creating your first graphs

The functions generating graph objects are called graph generators. Stochastic (i.e. randomized) graph generators are called "games".

igraph can handle directed and undirected graphs. Most graph generators are able to create both types of graphs and most other functions are usually also capable of handling both. E.g., igraph_get_shortest_paths(), which calculates shortest paths from a vertex to other vertices, can calculate directed or undirected paths.

igraph has sophisticated ways for creating graphs. The simplest graphs are deterministic regular structures like star graphs (igraph_star()), ring graphs (igraph_ring()), lattices (igraph_square_lattice()) or trees (igraph_kary_tree()).

The following example creates an undirected regular circular lattice, adds some random edges to it and calculates the average length of shortest paths between all pairs of vertices in the graph before and after adding the random edges. (The message is that some random edges can reduce path lengths a lot.)

```
#include <igraph.h>
```

```
int main() {
  igraph_t graph;
  igraph_vector_int_t dimvector;
  igraph_vector_int_t edges;
  igraph_vector_bool_t periodic;
  igraph_real_t avg_path_len;
  igraph_vector_int_init(&dimvector, 2);
  VECTOR(dimvector)[0]=30;
  VECTOR(dimvector)[1]=30;
  igraph_vector_bool_init(&periodic, 2);
  igraph_vector_bool_fill(&periodic, true);
  igraph_square_lattice(&graph, &dimvector, 0, IGRAPH_UNDIRECTED, /* mutual= */
  igraph_average_path_length(&graph, &avg_path_len, NULL, IGRAPH_UNDIRECTED, /*
  printf("Average path length (lattice):
                                                     %g\n", (double) avg_path_le
  igraph_rng_seed(igraph_rng_default(), 42); /* seed RNG before first use */
```

igraph vector int init(&edges, 20);

```
for (igraph_integer_t i=0; i < igraph_vector_int_size(&edges); i++) {
    VECTOR(edges)[i] = RNG_INTEGER(0, igraph_vcount(&graph) - 1);
}

igraph_add_edges(&graph, &edges, NULL);
igraph_average_path_length(&graph, &avg_path_len, NULL, IGRAPH_UNDIRECTED, /*
printf("Average path length (randomized lattice): %g\n", (double) avg_path_length_vector_bool_destroy(&periodic);
igraph_vector_int_destroy(&dimvector);
igraph_vector_int_destroy(&edges);
igraph_destroy(&graph);

return 0;
}</pre>
```

This example illustrates some new points. <code>igraph</code> uses <code>igraph_vector_t</code> and its related types (<code>igraph_vector_int_t</code>, <code>igraph_vector_bool_t</code> and so on) instead of plain C arrays. <code>igraph_vector_t</code> is superior to regular arrays in almost every sense. Vectors are created by the <code>igraph_vector_init()</code> function and, like graphs, they should be destroyed if not needed any more by calling <code>igraph_vector_tor_destroy()</code> on them. A vector can be indexed by the <code>VECTOR()</code> function (right now it is a macro). The elements of a vector are of type <code>igraph_real_t</code> for <code>igraph_vector_t</code>, and of type <code>igraph_integer_t</code> for <code>igraph_vector_int_t</code>. As you might expect, <code>igraph_vector_bool_t</code> holds <code>igraph_bool_t</code> values. Vectors can be resized and most <code>igraph</code> functions returning the result in a vector automatically resize it to the size they need.

igraph_square_lattice() takes an integer vector argument specifying the dimensions of the lattice. In this example we generate a 30x30 two dimensional periodic lattice. See the documentation of igraph_square_lattice() in the reference manual for the other arguments.

The vertices in a graph are identified by a *vertex ID*, an integer between 0 and N-1, where N is the number of vertices in the graph. The vertex count can be retrieved using igraph_vcount(), as in the example.

The <code>igraph_add_edges()</code> function simply takes a graph and a vector of vertex IDs defining the new edges. The first edge is between the first two vertex IDs in the vector, the second edge is between the second two, etc. This way we add ten random edges to the lattice.

Note that this example program may add *loop edges*, edges pointing a vertex to itself, or *multiple edges*, more than one edge between the same pair of vertices. igraph_t can of course represent loops and multiple edges, although some routines expect simple graphs, i.e. graphs which contain neither of these. This is because some structural properties are ill-defined for non-simple graphs. Loop and multi-edges can be removed by calling igraph_simplify().

Calculating various properties of graphs

In our next example we will calculate various centrality measures in a friendship graph. The friendship graph is from the famous Zachary karate club study. (Do a web search on "Zachary karate" if you want to know more about this.) Centrality measures quantify how central is the position of individual vertices in the graph.

```
#include <igraph.h>
int main() {
  igraph_t graph;
  igraph_vector_int_t v;
  igraph_vector_int_t result;
  igraph vector t result real;
```

```
igraph_integer_t edges[] = { 0,1, 0,2, 0,3, 0,4, 0,5, 0,6, 0,7, 0,8,
                               0,10, 0,11, 0,12, 0,13, 0,17, 0,19, 0,21, 0,31,
                               1, 2, 1, 3, 1, 7, 1,13, 1,17, 1,19, 1,21, 1,30,
                               2, 3, 2, 7, 2,27, 2,28, 2,32, 2, 9, 2, 8, 2,13,
                               3, 7, 3,12, 3,13, 4, 6, 4,10, 5, 6, 5,10, 5,16,
                               6,16, 8,30, 8,32, 8,33, 9,33, 13,33, 14,32, 14,3
                               15,32, 15,33, 18,32, 18,33, 19,33, 20,32, 20,33,
                               22,32, 22,33, 23,25, 23,27, 23,32, 23,33, 23,29,
                               24,25, 24,27, 24,31, 25,31, 26,29, 26,33, 27,33,
                               28,31, 28,33, 29,32, 29,33, 30,32, 30,33, 31,32,
                               31,33, 32,33 };
 igraph_vector_int_view(&v, edges, sizeof(edges) / sizeof(edges[0]));
 igraph_create(&graph, &v, 0, IGRAPH_UNDIRECTED);
 igraph_vector_int_init(&result, 0);
 igraph_vector_init(&result_real, 0);
 igraph_degree(&graph, &result, igraph_vss_all(), IGRAPH_ALL, IGRAPH_LOOPS);
 printf("Maximum degree is
                                 %10" IGRAPH_PRId ", vertex %2" IGRAPH_PRId ".\
         igraph_vector_int_max(&result),
         igraph_vector_int_which_max(&result));
 igraph_closeness(&graph, &result_real, NULL, NULL, igraph_vss_all(), IGRAPH_A
                   /* weights= */ NULL, /* normalized= */ false);
 printf("Maximum closeness is
                                 %10g, vertex %2" IGRAPH_PRId ".\n",
         (double) igraph_vector_max(&result_real),
         igraph_vector_which_max(&result_real));
 igraph_betweenness(&graph, &result_real, igraph_vss_all(),
                     IGRAPH_UNDIRECTED, /* weights= */ NULL);
 printf("Maximum betweenness is %10g, vertex %2" IGRAPH_PRId ".\n",
         (double) igraph_vector_max(&result_real),
         igraph_vector_which_max(&result_real));
 igraph_vector_int_destroy(&result);
 igraph_vector_destroy(&result_real);
 igraph_destroy(&graph);
 return 0;
}
```

This example demonstrates some new operations. First of all, it shows a way to create a graph a list of edges stored in a plain C array. Function <code>igraph_vector_view()</code> creates a *view* of a C array. It does not copy any data, which means that you must not call <code>igraph_vector_destroy()</code> on a vector created this way. This vector is then used to create the undirected graph.

Then the degree, closeness and betweenness centrality of the vertices is calculated and the highest values are printed. Note that the vector result, into which these functions will write their result, must be initialized first, and also that the functions resize it to be able to hold the result.

Notice that in order to print values of type igraph_integer_t, we used the IGRAPH_PRId format macro constant. This macro is similar to the standard PRI constants defined in stdint.h, and expands to the correct printf format specifier on each platform that **igraph** supports.

The igraph_vss_all() argument tells the functions to calculate the property for every vertex in the graph. It is shorthand for a *vertex selector*, represented by type igraph_vs_t. Vertex selectors help perform operations on a subset of vertices. You can read more about them in one of the following chapters.

Chapter 4. Basic data types and interface

The igraph data model

The igraph library can handle directed and undirected graphs. The igraph graphs are multisets of ordered (if directed) or unordered (if undirected) labeled pairs. The labels of the pairs plus the number of vertices always starts with zero and ends with the number of edges minus one. In addition to that, a table of metadata is also attached to every graph, its most important entries being the number of vertices in the graph and whether the graph is directed or undirected.

Like the edges, the igraph vertices are also labeled by numbers between zero and the number of vertices minus one. So, to summarize, a directed graph can be imagined like this:

```
( vertices: 6,
    directed: yes,
    {
      (0,2),
      (2,2),
      (3,2),
      (3,3),
      (3,4),
      (4,3),
      (4,1)
    }
)
```

Here the edges are ordered pairs or vertex ids, and the graph is a multiset of edges plus some metadata.

An undirected graph is like this:

```
( vertices: 6,
   directed: no,
   {
     (0,2),
     (2,2),
     (2,3),
     (3,3),
     (3,4),
     (3,4),
     (3,4),
     (1,4)
   }
)
```

Here, an edge is an unordered pair of two vertex IDs. A graph is a multiset of edges plus metadata, just like in the directed case.

It is possible to convert between directed and undirected graphs, see the <code>igraph_to_directed()</code> and <code>igraph_to_undirected()</code> functions.

igraph aims to robustly support multigraphs, i.e. graphs which have more than one edge between some pairs of vertices, as well as graphs with self-loops. Most functions which do not support such graphs

will check their input and issue an error if it is not valid. Those rare functions which do not perform this check clearly indicate this in their documentation. To eliminate multiple edges from a graph, you can use igraph simplify().

General conventions of igraph functions

igraph has a simple and consistent interface. Most functions check their input for validity and display an informative error message when something goes wrong. In order to support this, the majority of functions return an error code. In basic usage, this code can be ignored, as the default behaviour is to abort the program immediately upon error. See the section on error handling for more information on this topic.

Results are typically returned through *output arguments*, i.e. pointers to a data structure into which the result will be written. In almost all cases, this data structure is expected to be pre-initialized. A few simple functions communicate their result directly through their return value—these functions can never encounter an error.

Atomic data types

igraph introduces a few aliases to standard C data types that are then used throughout the library. The most important of these types is igraph_integer_t, which is an alias to either a 32-bit or a 64-bit *signed* integer, depending on whether igraph was compiled in 32-bit or 64-bit mode. The size of igraph_integer_t also influences the maximum number of vertices that an igraph graph can represent as the number of vertices is stored in a variable of type igraph_integer_t.

Since the size of a variable of type igraph_integer_t may change depending on how igraph is compiled, you cannot simply use %d or %ld as a placeholder for igraph integers in printf format strings. igraph provides the IGRAPH_PRId macro, which maps to d, ld or lld depending on the size of igraph_integer_t, and you must use this macro in printf format strings to avoid compiler warnings.

Similarly to how igraph_integer_t maps to the standard size signed integer in the library, igraph_uint_t maps to a 32-bit or a 64-bit *unsigned* integer. It is guaranteed that the size of igraph_integer_t is the same as the size of igraph_uint_t. igraph provides IGRAPH_PRIu as a format string placeholder for variables of type igraph_uint_t.

Real numbers (i.e. quantities that can potentially be fractional or infinite) are represented with a type named igraph_real_t. Currently igraph_real_t is always aliased to double, but it is still good practice to use igraph_real_t in your own code for sake of consistency.

Boolean values are represented with a type named igraph_bool_t. It tries to be as small as possible since it only needs to represent a truth value. For printing purposes, you can treat it as an integer and use %d in format strings as a placeholder for an igraph_bool_t.

Upper and lower limits of igraph_integer_t and igraph_uint_t are provided by the constants named IGRAPH_INTEGER_MIN, IGRAPH_INTEGER_MAX, IGRAPH_UINT_MIN and IGRAPH_UINT_MAX.

The basic interface

This is the very minimal API in **igraph**. All the other functions use this minimal set for creating and manipulating graphs.

This is a very important principle since it makes possible to implement other data representations by implementing only this minimal set.

This section lists all the functions and macros that are considered as part of the core API from the point of view of the *users* of igraph. Some of these functions and macros have sensible default implementations that simply call some other core function (e.g., igraph_empty() calls igraph_empty_attrs() with a null attribute table pointer). If you wish to experiment with implementing an alternative data type, the actual number of functions that you need to replace is lower as you can rely on the same default implementations in most cases.

Graph constructors and destructors

igraph_empty — Creates an empty graph with some vertices and no edges.

```
igraph_error_t igraph_empty(igraph_t *graph, igraph_integer_t n, igraph_bool_t
```

The most basic constructor, all the other constructors should call this to create a minimal graph object. Our use of the term "empty graph" in the above description should be distinguished from the mathematical definition of the empty or null graph. Strictly speaking, the empty or null graph in graph theory is the graph with no vertices and no edges. However by "empty graph" as used in igraph we mean a graph having zero or more vertices, but no edges.

Arguments:

graph: Pointer to a not-yet initialized graph object.

n: The number of vertices in the graph, a non-negative integer number is expected.

directed: Boolean; whether the graph is directed or not. Supported values are:

IGRAPH_DIRECTED The graph will be *directed*.

IGRAPH UNDIRECTED The graph will be undirected.

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices.

Time complexity: O(|V|) for a graph with |V| vertices (and no edges).

Example 4.1. File examples/simple/creation.c

igraph_empty_attrs — Creates an empty graph with some vertices, no edges and some graph attributes.

```
igraph_error_t igraph_empty_attrs(igraph_t *graph, igraph_integer_t n, igraph_b
```

Use this instead of igraph_empty() if you wish to add some graph attributes right after initialization. This function is currently not very interesting for the ordinary user. Just supply 0 here or use igraph_empty().

This function does not set any vertex attributes. To create a graph which has vertex attributes, call this function specifying 0 vertices, then use <code>igraph_add_vertices()</code> to add vertices and their attributes.

Arguments:

graph: Pointer to a not-yet initialized graph object.

n: The number of vertices in the graph; a non-negative integer number is expected.

directed: Boolean; whether the graph is directed or not. Supported values are:

IGRAPH DIRECTED Create a directed graph.

IGRAPH_UNDIRECTED Create an undirected graph.

attr: The graph attributes. Supply NULL if not graph attributes are to be set.

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices.

See also:

igraph_empty() to create an empty graph without attributes; igraph_add_vertices()
and igraph_add_edges() to add vertices and edges, possibly with associated attributes.

Time complexity: O(|V|) for a graph with |V| vertices (and no edges).

igraph_copy — Creates an exact (deep) copy of a graph.

```
igraph_error_t igraph_copy(igraph_t *to, const igraph_t *from);
```

This function deeply copies a graph object to create an exact replica of it. The new replica should be destroyed by calling <code>igraph_destroy()</code> on it when not needed any more.

You can also create a shallow copy of a graph by simply using the standard assignment operator, but be careful and do *not* destroy a shallow replica. To avoid this mistake, creating shallow copies is not recommended.

Arguments:

to: Pointer to an uninitialized graph object.

from: Pointer to the graph object to copy.

Returns:

Error code.

Time complexity: O(|V|+|E|) for a graph with |V| vertices and |E| edges.

Example 4.2. File examples/simple/igraph_copy.c

igraph_destroy — Frees the memory allocated for a graph object.

```
void igraph_destroy(igraph_t *graph);
```

This function should be called for every graph object exactly once.

This function invalidates all iterators (of course), but the iterators of a graph should be destroyed before the graph itself anyway.

Arguments:

graph: Pointer to the graph to free.

Time complexity: operating system specific.

Basic query operations

igraph_vcount — The number of vertices in a graph.

```
igraph_integer_t igraph_vcount(const igraph_t *graph);
```

Arguments:

graph: The graph.

Returns:

Number of vertices.

Time complexity: O(1)

igraph_ecount — The number of edges in a graph.

```
igraph_integer_t igraph_ecount(const igraph_t *graph);
```

Arguments:

graph: The graph.

Returns:

Number of edges.

Time complexity: O(1)

igraph_is_directed — Is this a directed graph?

```
igraph_bool_t igraph_is_directed(const igraph_t *graph);
```

Arguments:

graph: The graph.

Returns:

Logical value, true if the graph is directed, false otherwise.

Time complexity: O(1)

Example 4.3. File examples/simple/igraph_is_directed.c

igraph_edge — Returns the head and tail vertices of an edge.

```
igraph_error_t igraph_edge(
    const igraph_t *graph, igraph_integer_t eid,
    igraph_integer_t *from, igraph_integer_t *to
);
```

Arguments:

graph: The graph object.

eid: The edge ID.

from: Pointer to an igraph_integer_t. The tail (source) of the edge will be placed here.

to: Pointer to an igraph_integer_t. The head (target) of the edge will be placed here.

Returns:

Error code. The current implementation always returns with success.

See also:

igraph_get_eid() for the opposite operation; igraph_edges() to get the endpoints of several edges; IGRAPH_TO(), IGRAPH_FROM() and IGRAPH_OTHER() for a faster but nonerror-checked version.

Added in version 0.2.

Time complexity: O(1).

igraph_edges — Gives the head and tail vertices of a series of edges.

igraph_error_t igraph_edges(const igraph_t *graph, igraph_es_t eids, igraph_vec

Arguments:

graph: The graph object.

eids: Edge selector, the series of edges.

edges: Pointer to an initialized vector. The start and endpoints of each edge will be placed here.

Returns:

Error code.

See also:

igraph_get_edgelist() to get the endpoints of all edges; igraph_get_eids() for the
opposite operation; igraph_edge() for getting the endpoints of a single edge; IGRAPH_TO(),
IGRAPH_FROM() and IGRAPH_OTHER() for a faster but non-error-checked method.

Time complexity: O(k) where k is the number of edges in the selector.

IGRAPH_FROM — The source vertex of an edge.

```
#define IGRAPH_FROM(graph,eid)
```

Faster than igraph_edge(), but no error checking is done: eid is assumed to be valid.

Arguments:

graph: The graph.

eid: The edge ID.

Returns:

The source vertex of the edge.

See also:

igraph_edge() if error checking is desired.

IGRAPH_TO — The target vertex of an edge.

```
#define IGRAPH_TO(graph,eid)
```

Faster than igraph_edge(), but no error checking is done: eid is assumed to be valid.

Arguments:

graph: The graph object.

eid: The edge ID.

Returns:

The target vertex of the edge.

See also:

igraph_edge() if error checking is desired.

IGRAPH_OTHER — The other endpoint of an edge.

```
#define IGRAPH_OTHER(graph,eid,vid)
```

Typically used with undirected edges when one endpoint of the edge is known, and the other endpoint is needed. No error checking is done: eid and vid are assumed to be valid.

Arguments:

graph: The graph object.

eid: The edge ID.

vid: The vertex ID of one endpoint of an edge.

Returns:

The other endpoint of the edge.

See also:

IGRAPH_TO() and IGRAPH_FROM() to get the source and target of directed edges.

igraph_get_eid — Get the edge ID from the end points of an edge.

For undirected graphs pfrom and pto are exchangeable.

Arguments:

graph: The graph object.

eid: Pointer to an integer, the edge ID will be stored here.

pfrom: The starting point of the edge.

pto: The end point of the edge.

directed: Logical constant, whether to search for directed edges in a directed graph. Ignored for

undirected graphs.

error: Logical scalar, whether to report an error if the edge was not found. If it is false, then

-1 will be assigned to eid. Note that invalid vertex IDs in input arguments (pfrom

or pto) always return an error code.

Returns:

Error code.

See also:

igraph_edge() for the opposite operation, igraph_get_all_eids_between() to retrieve all edge IDs between a pair of vertices.

Time complexity: O(log (d)), where d is smaller of the out-degree of pfrom and in-degree of pto if directed is true. If directed is false, then it is O(log(d)+log(d2)), where d is the same as before and d2 is the minimum of the out-degree of pto and the in-degree of pfrom.

Example 4.4. File examples/simple/igraph_get_eid.c

Added in version 0.2.

igraph_get_eids — Return edge IDs based on the adjacent vertices.

The pairs of vertex IDs for which the edges are looked up are taken consecutively from the pairs vector, i.e. VECTOR(pairs)[0] and VECTOR(pairs)[1] specify the first pair, VECTOR(pairs)[2] and VECTOR(pairs)[3] the second pair, etc.

If you have a sequence of vertex IDs that describe a *path* on the graph, use igraph_ex-pand_path_to_pairs() to convert them to a list of vertex pairs along the path.

If the error argument is true, then it is an error to specify pairs of vertices that are not connected. Otherwise -1 is reported for vertex pairs without at least one edge between them.

If there are multiple edges in the graph, then these are ignored; i.e. for a given pair of vertex IDs, igraph always returns the same edge ID, even if the pair appears multiple times in pairs.

Arguments:

graph: The input graph.

eids: Pointer to an initialized vector, the result is stored here. It will be resized as needed.

pairs: Vector giving pairs of vertices to fetch the edges for.

directed: Logical scalar, whether to consider edge directions in directed graphs. This is ignored

for undirected graphs.

error: Logical scalar, whether it is an error to supply non-connected vertices. If false, then

-1 is returned for non-connected pairs.

Returns:

Error code.

Time complexity: O(n log(d)), where n is the number of queried edges and d is the average degree of the vertices.

See also:

igraph_get_eid() for a single edge.

Example 4.5. File examples/simple/igraph_get_eids.c

igraph_get_all_eids_between — Returns all edge IDs between a pair of vertices.

```
igraph_error_t igraph_get_all_eids_between(
    const igraph_t *graph, igraph_vector_int_t *eids,
    igraph_integer_t source, igraph_integer_t target, igraph_bool_t directed
);
```

For undirected graphs source and target are exchangeable.

Arguments:

graph: The input graph.

eids: Pointer to an initialized vector, the result is stored here. It will be resized as needed.

source: The ID of the source vertex

target: The ID of the target vertex

directed: Logical scalar, whether to consider edge directions in directed graphs. This is ignored

for undirected graphs.

Returns:

Error code.

Time complexity: TODO

See also:

igraph_get_eid() for a single edge.

igraph_neighbors — Adjacent vertices to a vertex.

Arguments:

graph: The graph to work on.

neis: This vector will contain the result. The vector should be initialized beforehand and will be

resized. Starting from igraph version 0.4 this vector is always sorted, the vertex IDs are in increasing order. If one neighbor is connected with multiple edges, the neighbor will

be returned multiple times.

pnode: The id of the node for which the adjacent vertices are to be searched.

mode: Defines the way adjacent vertices are searched in directed graphs. It can have the follow-

ing values: IGRAPH_OUT, vertices reachable by an edge from the specified vertex are searched; IGRAPH_IN, vertices from which the specified vertex is reachable are searched; IGRAPH_ALL, both kinds of vertices are searched. This parameter is ignored for undi-

rected graphs.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex ID. IGRAPH_EINVMODE: invalid mode argument. IGRAPH_ENOMEM: not enough memory.

Time complexity: O(d), d is the number of adjacent vertices to the queried vertex.

Example 4.6. File examples/simple/igraph_neighbors.c

igraph_incident — Gives the incident edges of a vertex.

Arguments:

graph: The graph object.

eids: An initialized vector. It will be resized to hold the result.

pnode: A vertex ID.

mode: Specifies what kind of edges to include for directed graphs. IGRAPH OUT means only

outgoing edges, IGRAPH_IN only incoming edges, IGRAPH_ALL both. This parameter

is ignored for undirected graphs.

Returns:

Error code. IGRAPH_EINVVID: invalid pnode argument, IGRAPH_EINVMODE: invalid mode argument.

Added in version 0.2.

Time complexity: O(d), the number of incident edges to pnode.

igraph_degree — The degree of some vertices in a graph.

```
igraph_error_t igraph_degree(const igraph_t *graph, igraph_vector_int_t *res,
```

```
const igraph_vs_t vids,
igraph_neimode_t mode, igraph_bool_t loops);
```

This function calculates the in-, out- or total degree of the specified vertices.

This function returns the result as a vector of igraph_integer_t values. In applications where igraph_real_t is desired, use igraph_strength() with NULL weights.

Arguments:

graph: The graph.

res: Integer vector, this will contain the result. It should be initialized and will be resized to

be the appropriate size.

vids: Vertex selector, giving the vertex IDs of which the degree will be calculated.

mode: Defines the type of the degree for directed graphs. Valid modes are: IGRAPH_OUT, out-

degree; IGRAPH_IN, in-degree; IGRAPH_ALL, total degree (sum of the in- and out-de-

gree). This parameter is ignored for undirected graphs.

loops: Boolean, gives whether the self-loops should be counted.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex ID. IGRAPH_EINVMODE: invalid mode argument.

Time complexity: O(v) if loops is true, and O(v*d) otherwise. v is the number of vertices for which the degree will be calculated, and d is their (average) degree.

See also:

igraph_strength() for the version that takes into account edge weights.

Example 4.7. File examples/simple/igraph degree.c

Adding and deleting vertices and edges

igraph_add_edge — Adds a single edge to a graph.

```
igraph_error_t igraph_add_edge(igraph_t *graph, igraph_integer_t from, igraph_integer_t from igraph_integ
```

For directed graphs the edge points from $from\ to\ to$.

Note that if you want to add many edges to a big graph, then it is inefficient to add them one by one, it is better to collect them into a vector and add all of them via a single igraph_add_edges() call.

Arguments:

igraph: The graph.

from: The id of the first vertex of the edge.

to: The id of the second vertex of the edge.

Returns:

Error code.

See also:

igraph_add_edges() to add many edges, igraph_delete_edges() to remove edges
and igraph_add_vertices() to add vertices.

Time complexity: O(|V|+|E|), the number of edges plus the number of vertices.

igraph_add_edges — Adds edges to a graph object.

The edges are given in a vector, the first two elements define the first edge (the order is from, to for directed graphs). The vector should contain even number of integer numbers between zero and the number of vertices in the graph minus one (inclusive). If you also want to add new vertices, call igraph_add_vertices() first.

Arguments:

graph: The graph to which the edges will be added.

edges: The edges themselves.

attr: The attributes of the new edges. You can supply a null pointer here if you do not need

edge attributes.

Returns:

Error code: IGRAPH_EINVEVECTOR: invalid (odd) edges vector length, IGRAPH_EINVVID: invalid vertex ID in edges vector.

This function invalidates all iterators.

Time complexity: O(|V|+|E|) where |V| is the number of vertices and |E| is the number of edges in the *new*, extended graph.

Example 4.8. File examples/simple/creation.c

igraph_add_vertices — Adds vertices to a graph.

```
igraph_error_t igraph_add_vertices(igraph_t *graph, igraph_integer_t nv, void *
```

This function invalidates all iterators.

Arguments:

graph: The graph object to extend.

nv: Non-negative integer specifying the number of vertices to add.

attr: The attributes of the new vertices. You can supply a null pointer here if you do not need vertex attributes.

Returns:

Error code: IGRAPH_EINVAL: invalid number of new vertices.

Time complexity: O(|V|) where |V| is the number of vertices in the *new*, extended graph.

Example 4.9. File examples/simple/creation.c

igraph_delete_edges — Removes edges from a graph.

```
igraph_error_t igraph_delete_edges(igraph_t *graph, igraph_es_t edges);
```

The edges to remove are specified as an edge selector.

This function cannot remove vertices; vertices will be kept even if they lose all their edges.

This function invalidates all iterators.

Arguments:

graph: The graph to work on.

edges: The edges to remove.

Returns:

Error code.

Time complexity: O(|V|+|E|) where |V| and |E| are the number of vertices and edges in the *original* graph, respectively.

Example 4.10. File examples/simple/igraph_delete_edges.c

igraph_delete_vertices — Removes some vertices (with all their edges) from the graph.

```
igraph_error_t igraph_delete_vertices(igraph_t *graph, const igraph_vs_t vertic
```

This function changes the IDs of the vertices (except in some very special cases, but these should not be relied on anyway).

This function invalidates all iterators.

Arguments:

graph: The graph to work on.

vertices: The IDs of the vertices to remove, in a vector. The vector may contain the same ID

more than once.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex ID.

Time complexity: O(|V|+|E|), |V| and |E| are the number of vertices and edges in the original graph.

Example 4.11. File examples/simple/igraph_delete_vertices.c

igraph_delete_vertices_idx — Removes some vertices (with all their edges) from the graph.

```
igraph_error_t igraph_delete_vertices_idx(
    igraph_t *graph, const igraph_vs_t vertices, igraph_vector_int_t *idx,
    igraph_vector_int_t *invidx
);
```

This function changes the IDs of the vertices (except in some very special cases, but these should not be relied on anyway). You can use the idx argument to obtain the mapping from old vertex IDs to the new ones, and the newidx argument to obtain the reverse mapping.

This function invalidates all iterators.

Arguments:

graph: The graph to work on.

vertices: The IDs of the vertices to remove, in a vector. The vector may contain the same ID

more than once.

idx: An optional pointer to a vector that provides the mapping from the vertex IDs before

the removal to the vertex IDs after the removal. You can supply NULL here if you

are not interested.

invidx: An optional pointer to a vector that provides the mapping from the vertex IDs after

the removal to the vertex IDs before the removal. You can supply NULL here if you

are not interested.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex ID.

Time complexity: O(|V|+|E|), |V| and |E| are the number of vertices and edges in the original graph.

Example 4.12. File examples/simple/igraph_delete_vertices.c

Miscellaneous macros and helper functions

IGRAPH_VCOUNT_MAX — The maximum number of vertices supported in igraph graphs.

```
#define IGRAPH_VCOUNT_MAX
```

The value of this constant is one less than <code>IGRAPH_INTEGER_MAX</code>. When igraph is compiled in 32-bit mode, this means that you are limited to $2^{31}-2$ (about 2.1 billion) vertices. In 64-bit mode, the limit is $2^{63}-2$ so you are much more likely to hit out-of-memory issues due to other reasons before reaching this limit.

IGRAPH_ECOUNT_MAX — The maximum number of edges supported in igraph graphs.

```
#define IGRAPH_ECOUNT_MAX
```

The value of this constant is half of IGRAPH_INTEGER_MAX. When igraph is compiled in 32-bit mode, this means that you are limited to approximately 2^{30} (about 1.07 billion) vertices. In 64-bit mode, the limit is approximately 2^{62} so you are much more likely to hit out-of-memory issues due to other reasons before reaching this limit.

igraph_expand_path_to_pairs — Helper function to convert a sequence of vertex IDs describing a path into a "pairs" vector.

```
igraph_error_t igraph_expand_path_to_pairs(igraph_vector_int_t* path);
```

This function is useful when you have a sequence of vertex IDs in a graph and you would like to retrieve the IDs of the edges between them. The function duplicates all but the first and the last elements in the vector, effectively converting the path into a vector of vertex IDs that can be passed to igraph_get_eids().

Arguments:

path: the input vector. It will be modified in-place and it will be resized as needed. When the vector contains less than two vertex IDs, it will be cleared.

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory to expand the vector.

igraph_invalidate_cache — Invalidates the internal cache of an igraph graph.

```
void igraph_invalidate_cache(const igraph_t* graph);
```

igraph graphs cache some basic properties about themselves in an internal data structure. This function invalidates the contents of the cache and forces a recalculation of the cached properties the next time they are needed.

You should not need to call this function during normal usage; however, we might ask you to call this function explicitly if we suspect that you are running into a bug in igraph's cache handling. A tell-tale sign of an invalid cache entry is that the result of a cached igraph function (such as igraph_is_dag() or igraph_is_simple()) is different before and after a cache invalidation.

Arguments:

graph: The graph whose cache is to be invalidated.

Time complexity: O(1).

igraph_is_same_graph — Are two graphs identical as labelled graphs?

```
igraph_error_t igraph_is_same_graph(const igraph_t *graph1, const igraph_t *graph
```

Two graphs are considered to be the same if they have the same vertex and edge sets. Graphs which are the same may have multiple different representations in igraph, hence the need for this function.

This function verifies that the two graphs have the same directedness, the same number of vertices, and that they contain precisely the same edges (regardless of their ordering) when written in terms of vertex indices. Graph attributes are not taken into account.

This concept is different from isomorphism. For example, the graphs 0-1, 2-1 and 1-2, 0-1 are considered the same because they only differ in the ordering of their edge lists and the ordering of vertices in an undirected edge. However, they are not the same as 0-2, 1-2, even though they are isomorphic to it. Note that this latter graph contains the edge 0-2 while the former two do not — thus their edge sets differ.

Arguments:

graph1: The first graph object.

graph2: The second graph object.

res: The result will be stored here.

Returns:

Error code.

Time complexity: O(E), the number of edges in the graphs.

See also:

igraph_isomorphic() to test if two graphs are isomorphic.

Chapter 5. Error handling

Error handling basics

igraph functions can run into various problems preventing them from normal operation. The user might have supplied invalid arguments, e.g. a non-square matrix when a square-matrix was expected, or the program has run out of memory while some more memory allocation is required, etc.

By default **igraph** aborts the program when it runs into an error. While this behavior might be good enough for smaller programs, it is without doubt avoidable in larger projects. Please read further if your project requires more sophisticated error handling. You can safely skip the rest of this chapter otherwise.

Error handlers

If **igraph** runs into an error - an invalid argument was supplied to a function, or we've ran out of memory - the control is transferred to the *error handler* function.

The default error handler is igraph_error_handler_abort which prints an error message and aborts the program.

The igraph_set_error_handler() function can be used to set a new error handler function of type igraph_error_handler_t; see the documentation of this type for details.

There are two other predefined error handler functions, <code>igraph_error_handler_ignore</code> and <code>igraph_error_handler_printignore</code>. These deallocate the temporarily allocated memory (more about this later) and return with the error code. The latter also prints an error message. If you use these error handlers you need to take care about possible errors yourself by checking the return value of (almost) every non-void <code>igraph</code> function.

Independently of the error handler installed, all functions in the library do their best to leave their arguments *semantically* unchanged if an error happens. By semantically we mean that the implementation of an object supplied as an argument might change, but its "meaning" in most cases does not. The rare occasions when this rule is violated are documented in this manual.

igraph_error_handler_t — The type of error handler functions.

This is the type of the error handler functions.

Arguments:

reason: Textual description of the error.

file: The source file in which the error is noticed.

1 ine: The number of the line in the source file which triggered the error

igraph_errno: The **igraph** error code.

igraph_error_handler_abort — Abort program in case of error.

```
IGRAPH_EXPORT igraph_error_handler_t igraph_error_handler_abort;
```

The default error handler, prints an error message and aborts the program.

igraph_error_handler_ignore — Ignore errors.

```
IGRAPH_EXPORT igraph_error_handler_t igraph_error_handler_ignore;
```

This error handler frees the temporarily allocated memory and returns with the error code.

igraph_error_handler_printignore — Print and ignore errors.

```
IGRAPH_EXPORT igraph_error_handler_t igraph_error_handler_printignore;
```

Frees temporarily allocated memory, prints an error message to the standard error and returns with the error code.

Error codes

Every **igraph** function which can fail return a single integer error code. Some functions are very simple and cannot run into any error, these may return other types, or void as well. The error codes are defined by the <code>igraph_error_type_t</code> enumeration.

igraph_error_t — Return type for functions returning an error code.

```
typedef igraph_error_type_t igraph_error_t;
```

This type is used as the return type of igraph functions that return an error code. It is a type alias because igraph_error_t used to be an int, and was used slightly differenly than igraph_error_type_t.

igraph_error_type_t — Error code type.

```
typedef enum {
    IGRAPH_SUCCESS = 0,
    IGRAPH_FAILURE = 1,
    IGRAPH_ENOMEM = 2,
    IGRAPH_PARSEERROR = 3,
    IGRAPH_EINVAL = 4,
    IGRAPH_EXISTS = 5,
    IGRAPH_EINVEVECTOR = 6,
    IGRAPH EINVVID = 7,
```

```
IGRAPH_NONSQUARE
                             = 8,
    IGRAPH_EINVMODE
                             = 9,
    IGRAPH EFILE
                             = 10,
   IGRAPH_UNIMPLEMENTED
                             = 12,
   IGRAPH_INTERRUPTED
                             = 13,
    IGRAPH_DIVERGED
                             = 14,
                            = 15,
                                      /* unused, reserved */
   IGRAPH ARPACK PROD
                             = 16,
    IGRAPH_ARPACK_NPOS
                             = 17,
    IGRAPH_ARPACK_NEVNPOS
    IGRAPH_ARPACK_NCVSMALL
                             = 18,
    IGRAPH_ARPACK_NONPOSI
                             = 19,
    IGRAPH_ARPACK_WHICHINV
                             = 20,
                             = 21,
    IGRAPH_ARPACK_BMATINV
   IGRAPH_ARPACK_WORKLSMALL = 22,
    IGRAPH_ARPACK_TRIDERR
                             = 23,
   IGRAPH_ARPACK_ZEROSTART = 24,
   IGRAPH_ARPACK_MODEINV
                             = 25.
                             = 26.
   IGRAPH_ARPACK_MODEBMAT
                             = 27,
    IGRAPH_ARPACK_ISHIFT
    IGRAPH_ARPACK_NEVBE
                             = 28,
    IGRAPH_ARPACK_NOFACT
                             = 29,
                             = 30,
    IGRAPH_ARPACK_FAILED
                             = 31,
    IGRAPH_ARPACK_HOWMNY
   IGRAPH_ARPACK_HOWMNYS
                             = 32,
                             = 33,
   IGRAPH_ARPACK_EVDIFF
    IGRAPH_ARPACK_SHUR
                             = 34,
   IGRAPH_ARPACK_LAPACK
                            = 35,
   IGRAPH ARPACK UNKNOWN
                            = 36,
                             = 37,
    IGRAPH_ENEGLOOP
                             = 38,
    IGRAPH_EINTERNAL
   IGRAPH_ARPACK_MAXIT
                             = 39,
                             = 40,
   IGRAPH_ARPACK_NOSHIFT
                             = 41,
   IGRAPH_ARPACK_REORDER
    IGRAPH_EDIVZERO
                             = 42,
                             = 43,
    IGRAPH_GLP_EBOUND
                            = 44,
   IGRAPH_GLP_EROOT
                            = 45,
    IGRAPH_GLP_ENOPFS
    IGRAPH_GLP_ENODFS
                            = 46,
   IGRAPH_GLP_EFAIL
                             = 47,
                             = 48.
    IGRAPH_GLP_EMIPGAP
    IGRAPH_GLP_ETMLIM
                             = 49,
                             = 50,
   IGRAPH_GLP_ESTOP
                            = 51,
    IGRAPH_EATTRIBUTES
                            = 52,
    IGRAPH_EATTRCOMBINE
   IGRAPH_ELAPACK
                             = 53,
                             = 54,
   IGRAPH_EDRL
   IGRAPH_EOVERFLOW
                             = 55,
                             = 56,
   IGRAPH_EGLP
   IGRAPH_CPUTIME
                             = 57,
                            = 58,
    IGRAPH_EUNDERFLOW
                            = 59,
    IGRAPH_ERWSTUCK
                             = 60,
    IGRAPH_STOP
                             = 61
    IGRAPH_ERANGE
} igraph_error_type_t;
```

These are the possible values returned by **igraph** functions. Note that these are interesting only if you defined an error handler with <code>igraph_set_error_handler()</code>. Otherwise the program is aborted and the function causing the error never returns.

Values:

IGRAPH_SUCCESS: The function successfully completed its task.

IGRAPH_FAILURE: Something went wrong. You'll almost never meet this error as

normally more specific error codes are used.

IGRAPH_ENOMEM: There wasn't enough memory to allocate on the heap.

IGRAPH_PARSEERROR: A parse error was found in a file.

IGRAPH_EINVAL: A parameter's value is invalid. E.g. negative number was spec-

ified as the number of vertices.

IGRAPH_EXISTS: A graph/vertex/edge attribute is already installed with the given

name.

IGRAPH EINVEVECTOR: Invalid vector of vertex IDs. A vertex ID is either negative or

bigger than the number of vertices minus one.

IGRAPH_EINVVID: Invalid vertex ID, negative or too big.

IGRAPH_NONSQUARE: A non-square matrix was received while a square matrix was

expected.

IGRAPH_EINVMODE: Invalid mode parameter.

IGRAPH_EFILE: A file operation failed. E.g. a file doesn't exist, or the user has

no rights to open it.

IGRAPH_UNIMPLEMENTED: Attempted to call an unimplemented or disabled (at com-

pile-time) function.

IGRAPH_DIVERGED: A numeric algorithm failed to converge.

IGRAPH_ARPACK_PROD: Matrix-vector product failed (not used any more).

IGRAPH_ARPACK_NPOS: N must be positive.

IGRAPH_ARPACK_NEVNPOS: NEV must be positive.

IGRAPH_ARPACK_NCVSMALL: NCV must be bigger.

IGRAPH_ARPACK_NONPOSI: Maximum number of iterations should be positive.

IGRAPH_ARPACK_WHICHINV: Invalid WHICH parameter.

IGRAPH_ARPACK_BMATINV: Invalid BMAT parameter.

IGRAPH_ARPACK_WORKLS-

MALL:

WORKL is too small.

IGRAPH_ARPACK_TRIDERR: LAPACK error in tridiagonal eigenvalue calculation.

IGRAPH_ARPACK_ZEROSTART: Starting vector is zero.

IGRAPH_ARPACK_MODEINV: MODE is invalid.

IGRAPH ARPACK MODEBMAT: MODE and BMAT are not compatible.

IGRAPH_ARPACK_ISHIFT: ISHIFT must be 0 or 1.

IGRAPH_ARPACK_NEVBE: NEV and WHICH='BE' are incompatible.

IGRAPH_ARPACK_NOFACT: Could not build an Arnoldi factorization.

IGRAPH_ARPACK_FAILED: No eigenvalues to sufficient accuracy.

IGRAPH_ARPACK_HOWMNY: HOWMNY is invalid.

IGRAPH_ARPACK_HOWMNYS: HOWMNY='S' is not implemented.

IGRAPH_ARPACK_EVDIFF: Different number of converged Ritz values.

IGRAPH_ARPACK_SHUR: Error from calculation of a real Schur form.

IGRAPH_ARPACK_LAPACK: LAPACK (dtrevc) error for calculating eigenvectors.

IGRAPH_ARPACK_UNKNOWN: Unknown ARPACK error.

IGRAPH_ENEGLOOP: Negative loop detected while calculating shortest paths.

IGRAPH_EINTERNAL: Internal error, likely a bug in igraph.

IGRAPH_EDIVZERO: Big integer division by zero.

IGRAPH_GLP_EBOUND: GLPK error (GLP_EBOUND).

IGRAPH_GLP_EROOT: GLPK error (GLP_EROOT).

IGRAPH_GLP_ENOPFS: GLPK error (GLP_ENOPFS).

IGRAPH_GLP_ENODFS: GLPK error (GLP_ENODFS).

IGRAPH_GLP_EFAIL: GLPK error (GLP_EFAIL).

IGRAPH_GLP_EMIPGAP: GLPK error (GLP_EMIPGAP).

IGRAPH_GLP_ETMLIM: GLPK error (GLP_ETMLIM).

IGRAPH_GLP_ESTOP: GLPK error (GLP_ESTOP).

IGRAPH EATTRIBUTES: Attribute handler error. The user is not expected to find this;

it is signalled if some igraph function is not using the attribute

handler interface properly.

IGRAPH_EATTRCOMBINE: Unimplemented attribute combination method for the given at-

tribute type.

IGRAPH_ELAPACK: A LAPACK call resulted in an error.

IGRAPH_EDRL: Internal error in the DrL layout generator.

IGRAPH_EOVERFLOW: Integer or double overflow.

IGRAPH_EGLP: Internal GLPK error.

IGRAPH_CPUTIME: CPU time exceeded.

IGRAPH_EUNDERFLOW: Integer or double underflow.

IGRAPH_ERWSTUCK: Random walk got stuck.

igraph_strerror — Textual description of an error.

```
const char* igraph_strerror(const igraph_error_t igraph_errno);
```

This is a simple utility function, it gives a short general textual description for an **igraph** error code.

Arguments:

igraph_errno: The **igraph** error code.

Returns:

pointer to the textual description of the error code.

Warning messages

igraph also supports warning messages in addition to error messages. Warning messages typically do not terminate the program, but they are usually crucial to the user.

igraph warnings are handled similarly to errors. There is a separate warning handler function that is called whenever an **igraph** function triggers a warning. This handler can be set by the <code>igraph_set_warning_handler()</code> function. There are two predefined simple warning handlers, <code>igraph_warning_handler_ignore()</code> and <code>igraph_warning_handler_print()</code>, the latter being the default.

To trigger a warning, **igraph** functions typically use the IGRAPH_WARNING() macro, the igraph_warning() function, or if more flexibility is needed, igraph_warningf().

igraph_warning_handler_t — The type of igraph warning handler functions.

```
typedef void igraph_warning_handler_t (const char *reason, const char *file, in
```

Currently it is defined to have the same type as igraph_error_handler_t, although the last (error code) argument is not used.

igraph_set_warning_handler — Installs a warning handler.

igraph_warning_handler_t* igraph_set_warning_handler(igraph_warning_handler_t*

Install the supplied warning handler function.

Arguments:

new_handler: The new warning handler function to install. Supply a null pointer here to uninstall

the current warning handler, without installing a new one.

Returns:

The current warning handler function.

IGRAPH_WARNING — Triggers a warning.

#define IGRAPH_WARNING(reason)

This is the usual way of triggering a warning from an igraph function. It calls igraph_warning().

Arguments:

reason: The warning message.

IGRAPH_WARNINGF — Triggers a warning, with printf-like syntax.

#define IGRAPH_WARNINGF

igraph functions can use this macro when they notice a warning and want to pass on extra information to the user about what went wrong. It calls <code>igraph_warningf()</code> with the proper parameters and no error code.

Arguments:

reason: Textual description of the warning, a template string with the same syntax as the standard printf C library function.

. . .: The additional arguments to be substituted into the template string.

igraph_warning — Reports a warning.

void igraph_warning(const char *reason, const char *file, int line);

Call this function if you want to trigger a warning from within a function that uses **igraph**.

Arguments:

reason: Textual description of the warning.

file: The source file in which the warning was noticed.

1ine: The number of line in the source file which triggered the warning.

igraph_errno: Warnings could have potentially error codes as well, but this is currently not

used in igraph.

Returns:

The supplied error code.

igraph_warningf — Reports a warning, printf-like version.

IGRAPH_FUNCATTR_PRINTFLIKE(1,4)
IGRAPH_EXPORT void igraph_warningf(const char *reason, const char *file, int li

This function is similar to <code>igraph_warning()</code>, but uses a printf-like syntax. It substitutes the additional arguments into the <code>reason</code> template string and calls <code>igraph_warning()</code>.

Arguments:

reason: Textual description of the warning, a template string with the same syntax as

the standard printf C library function.

file: The source file in which the warning was noticed.

1ine: The number of line in the source file which triggered the warning.

igraph_errno: Warnings could have potentially error codes as well, but this is currently not

used in igraph.

...: The additional arguments to be substituted into the template string.

igraph_warning_handler_ignore — Ignores all warnings.

void igraph_warning_handler_ignore(const char *reason, const char *file, int li

This warning handler function simply ignores all warnings.

Arguments:

reason: Textual description of the warning.

file: The source file in which the warning was noticed.

1ine: The number of line in the source file which triggered the warning...

igraph_errno: Warnings could have potentially error codes as well, but this is currently not

used in igraph.

igraph_warning_handler_print — Prints all warnings to the standard error.

void igraph_warning_handler_print(const char *reason, const char *file, int lin

This warning handler function simply prints all warnings to the standard error.

Arguments:

reason: Textual description of the warning.

file: The source file in which the warning was noticed.

1 ine: The number of line in the source file which triggered the warning...

igraph_errno: Warnings could have potentially error codes as well, but this is currently not

used in igraph.

Advanced topics

Writing error handlers

The contents of the rest of this chapter might be useful only for those who want to create an interface to **igraph** from another language, or use igraph from a GUI application. Most readers can safely skip to the next chapter.

You can write and install error handlers simply by defining a function of type <code>igraph_error_handler_t</code> and calling <code>igraph_set_error_handler()</code>. This feature is useful for interface writers, as <code>igraph</code> will have the chance to signal errors the appropriate way. For example, the R interface uses R's native printing facilities to communicate errors, while the Python interface converts them into Python exceptions.

The two main tasks of the error handler are to report the error (i.e. print the error message) and ensure proper resource cleanup. This is ensured by calling IGRAPH_FINALLY_FREE(), which deallocates some of the temporary memory to avoid memory leaks. Note that this may invalidate the error message buffer reason passed to the error handler. Do not access it after having called IGRAPH_FINALLY FREE().

As of **igraph** 0.10, temporary memory is dellocated in stages, through multiple calls to the error handler (and indirectly to IGRAPH_FINALLY_FREE()). Therefore, error handlers that do not abort the program immediately are expected to return. The error handler should not perform a longjmp, as this may lead to some of the memory not getting freed.

igraph_set_error_handler — Sets a new error handler.

```
igraph_error_handler_t* igraph_set_error_handler(igraph_error_handler_t* new_h
```

Installs a new error handler. If called with NULL, it installs the default error handler (which is currently igraph_error_handler_abort).

Arguments:

new handler: The error handler function to install.

Returns:

The old error handler function. This should be saved and restored if <code>new_handler</code> is not needed any more.

Error handling internals

If an error happens, the functions in the library call the IGRAPH_ERROR() macro with a textual description of the error and an **igraph** error code. This macro calls (through the igraph error()

function) the installed error handler. Another useful macro is <code>IGRAPH_CHECK()</code>. This checks the return value of its argument, which is normally a function call, and calls <code>IGRAPH_ERROR()</code> if it is not <code>IGRAPH_SUCCESS</code>.

IGRAPH_ERROR — Triggers an error.

```
#define IGRAPH_ERROR(reason, igraph_errno)
```

igraph functions usually use this macro when they notice an error. It calls <code>igraph_error()</code> with the proper parameters and if that returns the macro returns the "calling" function as well, with the error code. If for some (suspicious) reason you want to call the error handler without returning from the current function, call <code>igraph_error()</code> directly.

Arguments:

reason: Textual description of the error. This should be something more descriptive than

the text associated with the error code. E.g. if the error code is IGRAPH_EIN-VAL, its associated text (see igraph_strerror()) is "Invalid value" and this string should explain which parameter was invalid and maybe why.

igraph_errno: The **igraph** error code.

IGRAPH_ERRORF — Triggers an error, with printf-like syntax.

#define IGRAPH_ERRORF

igraph functions can use this macro when they notice an error and want to pass on extra information to the user about what went wrong. It calls <code>igraph_errorf()</code> with the proper parameters and if that returns the macro returns the "calling" function as well, with the error code. If for some (suspicious) reason you want to call the error handler without returning from the current function, call <code>igraph_errorf()</code> directly.

Arguments:

reason: Textual description of the error, a template string with the same syntax as

the standard printf C library function. This should be something more descriptive than the text associated with the error code. E.g. if the error code is ${\tt IGRAPH_EINVAL}$, its associated text (see ${\tt igraph_strerror}()$) is "Invalid value" and this string should explain which parameter was invalid and

maybe what was expected and what was recieved.

igraph_errno: The igraph error code.

. . .: The additional arguments to be substituted into the template string.

igraph_error — Reports an error.

igraph functions usually call this function (most often via the IGRAPH_ERROR macro) if they notice an error. It calls the currently installed error handler function with the supplied arguments.

Arguments:

reason: Textual description of the error.

file: The source file in which the error was noticed.

1 ine: The number of line in the source file which triggered the error.

igraph_errno: The **igraph** error code.

Returns:

the error code (if it returns)

See also:

igraph_errorf().

igraph_errorf — Reports an error, printf-like version.

Arguments:

reason: Textual description of the error, interpreted as a printf format string.

file: The source file in which the error was noticed.

line: The line in the source file which triggered the error.

igraph_errno: The **igraph** error code.

. . .: Additional parameters, the values to substitute into the format string.

See also:

igraph_error().

IGRAPH CHECK — Checks the return value of a function call.

```
#define IGRAPH_CHECK(expr)
```

Arguments:

expr: An expression, usually a function call. It is guaranteed to be evaluated only once.

Executes the expression and checks its value. If this is not IGRAPH_SUCCESS, it calls IGRAPH_ERROR with the value as the error code. Here is an example usage:

```
IGRAPH_CHECK(vector_push_back(&v, 100));
```

There is only one reason to use this macro when writing <code>igraph</code> functions. If the user installs an error handler which returns to the auxiliary calling code (like <code>igraph_error_handler_ignore</code> and <code>igraph_error_handler_printignore</code>), and the <code>igraph</code> function signalling the error is called from another <code>igraph</code> function then we need to make sure that the error is propagated back to the auxiliary (i.e. non-igraph) calling function. This is achieved by using <code>IGRAPH_CHECK</code> on every <code>igraph</code> call which can return an error code.

IGRAPH_CHECK_CALLBACK — Checks the return value of a callback.

```
#define IGRAPH_CHECK_CALLBACK(expr, code)
```

Identical to IGRAPH_CHECK, but treats IGRAPH_STOP as a normal (non-erroneous) return code. This macro is used in some igraph functions that allow the user to hook into a long-running calculation with a callback function. When the user-defined callback function returns IGRAPH_SUCCESS, the calculation will proceed normally. Returning IGRAPH_STOP from the callback will terminate the calculation without reporting an error. Returning any other value from the callback is treated as an error code, and igraph will trigger the necessary cleanup functions before exiting the function.

Note that IGRAPH_CHECK_CALLBACK does not handle IGRAPH_STOP by any means except returning it in the variable pointed to by code. It is the responsibility of the caller to handle IGRAPH_STOP accordingly.

Arguments:

expr: An expression, usually a call to a user-defined callback function. It is guaranteed to be evaluated only once.

code: Pointer to an optional variable of type igraph_error_t; the value of this variable will be set to the error code if it is not a null pointer.

Deallocating memory

If a function runs into an error (and the program is not aborted) the error handler should deallocate all temporary memory. This is done by storing the address and the destroy function of all temporary objects in a stack. The <code>IGRAPH_FINALLY</code> function declares an object as temporary by placing its address in the stack. If an <code>igraph</code> function returns with success it calls <code>IGRAPH_FINALLY_CLEAN()</code> with the number of objects to remove from the stack. If an error happens however, the error handler should call <code>IGRAPH_FINALLY_FREE()</code> to deallocate each object added to the stack. This means that the temporary objects allocated in the calling function (and etc.) will be freed as well.

IGRAPH_FINALLY — Registers an object for deallocation.

```
#define IGRAPH_FINALLY(func, ptr)
```

This macro places the address of an object, together with the address of its destructor in a stack. This stack is used if an error happens to deallocate temporarily allocated objects to prevent memory leaks. After manual deallocation, objects are removed from the stack using IGRAPH_FINALLY_CLEAN().

Arguments:

func: The function which is normally called to destroy the object.

ptr: Pointer to the object itself.

IGRAPH_FINALLY_CLEAN — Signals clean deallocation of objects.

```
void IGRAPH_FINALLY_CLEAN(int num);
```

Removes the specified number of objects from the stack of temporarily allocated objects. It is typically called immediately after manually destroying the objects:

```
igraph_vector_t vector;
igraph_vector_init(&vector, 10);
IGRAPH_FINALLY(igraph_vector_destroy, &vector);
// use vector
igraph_vector_destroy(&vector);
IGRAPH_FINALLY_CLEAN(1);
```

Arguments:

num: The number of objects to remove from the bookkeeping stack.

IGRAPH_FINALLY_FREE — Deallocates objects registered at the current level.

```
void IGRAPH FINALLY FREE(void);
```

Calls the destroy function for all objects in the current level of the stack of temporarily allocated objects, i.e. up to the nearest mark set by IGRAPH_FINALLY_ENTER(). This function must only be called from an error handler. It is *not* appropriate to use it instead of destroying each unneeded object of a function, as it destroys the temporary objects of the caller function (and so on) as well.

Writing igraph functions with proper error handling

There are some simple rules to keep in order to have functions behaving well in erroneous situations. First, check the arguments of the functions and call <code>IGRAPH_ERROR()</code> if they are invalid. Second, call <code>IGRAPH_FINALLY</code> on each dynamically allocated object and call <code>IGRAPH_FINALLY_CLEAN()</code> with the proper argument before returning. Third, use <code>IGRAPH_CHECK</code> on all <code>igraph</code> function calls which can generate errors.

The size of the stack used for this bookkeeping is fixed, and small. If you want to allocate several objects, write a destroy function which can deallocate all of these. See the adjlist.c file in the **igraph** source for an example.

For some functions these mechanisms are simply not flexible enough. These functions should define their own error handlers and restore the error handler before they return.

Fatal errors

In some rare situations, **igraph** may encounter an internal error that cannot be fully handled. In this case, it will call the current fatal error handler. The default fatal error handler simply prints the error and aborts the program.

Fatal error handlers do not return. Typically, they might abort the the program immediately, or in the case of the high-level **igraph** interfaces, they might return to the top level using a <code>longjmp()</code>. The fatal error handler is only called when a serious error has occurred, and as a result igraph may be in an inconsistent state. The purpose of returning to the top level is to give the user a chance to save their work instead of aborting immediately. However, the program session should be restarted as soon as possible.

Most projects that use **igraph** will use the default fatal error handler.

igraph_fatal_handler_t — The type of igraph fatal error handler functions.

typedef void igraph_fatal_handler_t (const char *reason, const char *file, int Functions of this type *must* not return. Typically they call abort() or do a longjmp().

Arguments:

reason: Textual description of the error.

file: The source file in which the error is noticed.

line: The number of the line in the source file which triggered the error

igraph_fatal_handler_abort — Abort program in case of fatal error.

```
IGRAPH_EXPORT igraph_fatal_handler_t igraph_fatal_handler_abort;
```

The default fatal error handler, prints an error message and aborts the program.

IGRAPH_FATAL — Triggers a fatal error.

```
#define IGRAPH_FATAL(reason)
```

This is the usual way of triggering a fatal error from an igraph function. It calls igraph_fatal().

Use this macro only in situations where the error cannot be handled. The normal way to handle errors is <code>IGRAPH_ERROR()</code>.

Arguments:

reason: The error message.

IGRAPH_FATALF — Triggers a fatal error, with printf-like syntax.

```
#define IGRAPH_FATALF
```

igraph functions can use this macro when a fatal error occurs and want to pass on extra information to the user about what went wrong. It calls igraph_fatalf() with the proper parameters.

Arguments:

reason: Textual description of the error, a template string with the same syntax as the standard

printf C library function.

. . .: The additional arguments to be substituted into the template string.

IGRAPH_ASSERT — igraph-specific replacement for assert().

#define IGRAPH ASSERT(condition)

This macro is like the standard assert(), but instead of calling abort(), it calls igraph_fa-tal(). This allows for returning the control to the calling program, e.g. returning to the top level in a high-level **igraph** interface.

Unlike assert(), IGRAPH_ASSERT() is not disabled when the NDEBUG macro is defined.

This macro is meant for internal use by **igraph**.

Since a typical fatal error handler does a longjmp(), avoid using this macro in C++ code. With most compilers, destructor will not be called when longjmp() leaves the current scope.

Arguments:

condition: The condition to be checked.

igraph_fatal — Triggers a fatal error.

IGRAPH_FUNCATTR_NORETURN void igraph_fatal(const char *reason, const char *fil

This function triggers a fatal error. Typically it is called indirectly through IGRAPH_FATAL() or IGRAPH_ASSERT().

Arguments:

reason: Textual description of the error.

file: The source file in which the error was noticed.

line: The number of line in the source file which triggered the error.

igraph_fatalf — Triggers a fatal error, printf-like syntax.

```
IGRAPH_FUNCATTR_PRINTFLIKE(1,4)
IGRAPH_EXPORT IGRAPH_FUNCATTR_NORETURN void igraph_fatalf(const char *reason, c
```

This function is similar to <code>igraph_fatal()</code>, but uses a printf-like syntax. It substitutes the additional arguments into the <code>reason</code> template string and calls <code>igraph_fatal()</code>.

Arguments:

reason: Textual description of the error.

file: The source file in which the error was noticed.

1ine: The number of line in the source file which triggered the error.

. . .: The additional arguments to be substituted into the template string.

Error handling and threads

It is likely that the **igraph** error handling method is *not* thread-safe, mainly because of the static global stack which is used to store the address of the temporarily allocated objects. This issue might be addressed in a later version of **igraph**.

Chapter 6. Memory (de)allocation

igraph_malloc — Allocate memory that can be safely deallocated by igraph functions.

```
void *igraph_malloc(size_t size);
```

This function behaves like malloc(), but it ensures that at least one byte is allocated even when the caller asks for zero bytes.

Arguments:

size: Number of bytes to be allocated. Zero is treated as one byte.

Returns:

Pointer to the piece of allocated memory; NULL if the allocation failed.

See also:

```
igraph_calloc(), igraph_realloc(), igraph_free()
```

igraph_calloc — Allocate memory that can be safely deallocated by igraph functions.

```
void *igraph_calloc(size_t count, size_t size);
```

This function behaves like calloc(), but it ensures that at least one byte is allocated even when the caller asks for zero bytes.

Arguments:

count: Number of items to be allocated.

size: Size of a single item to be allocated.

Returns:

Pointer to the piece of allocated memory; NULL if the allocation failed.

See also:

```
igraph_malloc(), igraph_realloc(), igraph_free()
```

igraph_realloc — Reallocate memory that can be safely deallocated by igraph functions.

```
void *igraph_realloc(void* ptr, size_t size);
```

This function behaves like realloc(), but it ensures that at least one byte is allocated even when the caller asks for zero bytes.

Arguments:

ptr: The pointer to reallocate.

size: Number of bytes to be allocated.

Returns:

Pointer to the piece of allocated memory; NULL if the allocation failed.

See also:

```
igraph_free(),igraph_malloc()
```

igraph_free — Deallocate memory that was allocated by igraph functions.

```
void igraph_free(void *ptr);
```

This function exposes the free() function used internally by igraph.

Arguments:

ptr: Pointer to the piece of memory to be deallocated.

Time complexity: platform dependent, ideally it should be O(1).

See also:

```
igraph_calloc(), igraph_malloc(), igraph_realloc()
```

Chapter 7. Data structure library: vector, matrix, other data types

About template types

Some of the container types listed in this section are defined for many base types. This is similar to templates in C++ and generics in Ada, but it is implemented via preprocessor macros since the C language cannot handle it. Here is the list of template types and the all base types they currently support:

vector Vector is currently defined for igraph_real_t, igraph_integer_t

(int), char (char), igraph_bool_t (bool) and igraph_complex_t

(complex). The default is igraph_real_t.

matrix Matrix is currently defined for igraph_real_t, igraph_integer_t

(int), char (char), igraph_bool_t (bool) and igraph_complex_t

(complex). The default is igraph_real_t.

array3 Array3 is currently defined for igraph_real_t, igraph_integer_t

(int), char (char) and igraph_bool_t (bool). The default is

igraph_real_t.

stack Stack is currently defined for igraph_real_t, igraph_integer_t

(int), char (char), igraph_bool_t (bool) and void* (ptr). The de-

fault is igraph_real_t.

double-ended queue Dqueue is currently defined for igraph_real_t, igraph_integer_t

(int), char (char) and igraph bool t (bool). The default is

igraph_real_t.

heap Heap is currently defined for igraph real t, igraph integer t (int),

char (char). In addition both maximum and minimum heaps are

available. The default is the igraph_real_t maximum heap.

igraph_real_t and igraph_integer_t (int). The default is igraph_re-

al_t.

igraph_real_t only.

The name of the base element (in parentheses) is added to the function names, except for the default type.

Some examples:

- igraph_vector_t is a vector of igraph_real_t elements. Its functions are igraph_vector_init, igraph_vector_destroy, igraph_vector_sort, etc.
- igraph_vector_bool_t is a vector of igraph_bool_t elements; initialize it with igraph_vector_bool_init, destroy it with igraph_vector_bool_destroy, etc.
- igraph_heap_t is a maximum heap with igraph_real_t elements. The corresponding functions are igraph_heap_init, igraph_heap_pop, etc.
- igraph_heap_min_t is a minimum heap with igraph_real_t elements. The corresponding functions are called igraph_heap_min_init, igraph_heap_min_pop, etc.

- igraph_heap_int_t is a maximum heap with igraph_integer_t elements. Its functions have the igraph_heap_int_ prefix.
- igraph_heap_min_int_t is a minimum heap containing igraph_integer_t elements. Its functions have the igraph_heap_min_int_ prefix.
- igraph_vector_list_t is a list of (floating-point) vectors; each element in this data structure is an igraph_vector_t. Similarly, igraph_matrix_list_t is a list of (floating-point) matrices; each element in this data structure is an igraph_matrix_t.
- igraph_vector_int_list_t is a list of integer vectors; each element in this data structure is an igraph_vector_int_t.

Note that the VECTOR and the MATRIX macros can be used on *all* vector and matrix types. VECTOR cannot be used on *lists* of vectors, though, only on the individual vectors in the list.

Vectors

About igraph_vector_t objects

The igraph_vector_t data type is a simple and efficient interface to arrays containing numbers. It is something similar to (but much simpler than) the vector template in the C++ standard library.

There are multiple variants of igraph_vector_t; the basic variant stores doubles, but there is also igraph_vector_int_t for integers (of type igraph_integer_t), igraph_vector_bool_t for booleans (of type igraph_bool_t) and so on. Vectors are used extensively in **igraph**; all functions that expect or return a list of numbers use igraph_vector_t or igraph_vector_int_t to achieve this. Integer vectors are typically used when the vector is supposed to hold vertex or edge identifiers, while igraph_vector_t is used when the vector is expected to hold fractional numbers or infinities.

The igraph_vector_t type and its variants usually use O(n) space to store n elements. Sometimes they use more, this is because vectors can shrink, but even if they shrink, the current implementation does not free a single bit of memory.

The elements in an igraph_vector_t object and its variants are indexed from zero, we follow the usual C convention here.

The elements of a vector always occupy a single block of memory, the starting address of this memory block can be queried with the VECTOR macro. This way, vector objects can be used with standard mathematical libraries, like the GNU Scientific Library.

Almost all of the functions described below for igraph_vector_t also exist for all the other vector type variants. These variants are not documented separately; you can simply replace vector with vector_int, vector_bool or something similar if you need a function for another variant. For instance, to initialize a vector of type igraph_vector_int_t, you need to use igraph_vector_int_t_init() and not igraph_vector_init().

Constructors and destructors

igraph_vector_t objects have to be initialized before using them, this is analogous to calling a constructor on them. There are a number of igraph_vector_t constructors, for your convenience. igraph_vector_init() is the basic constructor, it creates a vector of the given length, filled with zeros. igraph_vector_init_copy() creates a new identical copy of an already existing and initialized vector. igraph_vector_init_array() creates a vector by copying a regular C array. igraph_vector_init_range() creates a vector containing a regular sequence with increment one.

igraph_vector_view() is a special constructor, it allows you to handle a regular C array as a vector without copying its elements.

If a igraph_vector_t object is not needed any more, it should be destroyed to free its allocated memory by calling the igraph_vector_t destructor, igraph_vector_destroy().

Note that vectors created by igraph_vector_view() are special, you must not call igraph_vector_destroy() on these.

igraph_vector_init — Initializes a vector object (constructor).

```
igraph_error_t igraph_vector_init(igraph_vector_t* v, igraph_integer_t size);
```

Every vector needs to be initialized before it can be used, and there are a number of initialization functions or otherwise called constructors. This function constructs a vector of the given size and initializes each entry to 0. Note that <code>igraph_vector_null()</code> can be used to set each element of a vector to zero. However, if you want a vector of zeros, it is much faster to use this function than to create a vector and then invoke <code>igraph_vector_null()</code>.

Every vector object initialized by this function should be destroyed (ie. the memory allocated for it should be freed) when it is not needed anymore, the <code>igraph_vector_destroy()</code> function is responsible for this.

Arguments:

v: Pointer to a not yet initialized vector object.

size: The size of the vector.

Returns:

error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, the amount of "time" required to allocate O(n) elements, n is the number of elements.

igraph_vector_init_array — Initializes a vector from an ordinary C array (constructor).

Arguments:

v: Pointer to an uninitialized vector object.

data: A regular C array.

length: The length of the C array.

Returns:

Error code: IGRAPH ENOMEM if there is not enough memory.

Time complexity: operating system specific, usually O(length).

igraph_vector_init_copy — Initializes a vector from another vector object (constructor).

```
igraph_error_t igraph_vector_init_copy(
    igraph_vector_t *to, const igraph_vector_t *from
);
```

The contents of the existing vector object will be copied to the new one.

Arguments:

to: Pointer to a not yet initialized vector object.

from: The original vector object to copy.

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, usually O(n), n is the size of the vector.

igraph_vector_init_range — Initializes a vector with a range.

```
igraph_error_t igraph_vector_init_range(igraph_vector_t *v, igraph_real_t from,
```

The vector will contain the numbers start, start+1, ..., end-1. Note that the range is closed from the left and open from the right, according to C conventions.

Arguments:

v: Pointer to an uninitialized vector object.

start: The lower limit in the range (inclusive).

end: The upper limit in the range (exclusive).

Returns:

Error code: IGRAPH_ENOMEM: out of memory.

Time complexity: O(n), the number of elements in the vector.

igraph_vector_destroy — Destroys a vector object.

```
void igraph_vector_destroy(igraph_vector_t* v);
```

All vectors initialized by igraph_vector_init() should be properly destroyed by this function. A destroyed vector needs to be reinitialized by igraph_vector_init(), igraph_vector init array() or another constructor.

Arguments:

v: Pointer to the (previously initialized) vector object to destroy.

Time complexity: operating system dependent.

Initializing elements

igraph_vector_null — Sets each element in the vector to zero.

```
void igraph_vector_null(igraph_vector_t* v);
```

Note that igraph_vector_init() sets the elements to zero as well, so it makes no sense to call this function on a just initialized vector. Thus if you want to construct a vector of zeros, then you should use igraph_vector_init().

Arguments:

v: The vector object.

Time complexity: O(n), the size of the vector.

igraph_vector_fill — Fill a vector with a constant element.

```
void igraph_vector_fill(igraph_vector_t* v, igraph_real_t e);
```

Sets each element of the vector to the supplied constant.

Arguments:

vector: The vector to work on.

e: The element to fill with.

Time complexity: O(n), the size of the vector.

igraph_vector_range — Updates a vector to store a range.

```
igraph_error_t igraph_vector_range(igraph_vector_t *v, igraph_real_t from, igraph_
```

Sets the elements of the vector to contain the numbers start, start+1, ..., end-1. Note that the range is closed from the left and open from the right, according to C conventions.

Arguments:

v: The vector to update.

start: The lower limit in the range (inclusive).

end: The upper limit in the range (exclusive).

Returns:

Error code: IGRAPH_ENOMEM: out of memory.

Time complexity: O(n), the number of elements in the vector.

Accessing elements

The simplest way to access an element of a vector is to use the VECTOR macro. This macro can be used both for querying and setting igraph_vector_t elements. If you need a function, igraph_vector_get() queries and igraph_vector_set() sets an element of a vector. igraph_vector_get_ptr() returns the address of an element.

 $igraph_vector_tail()$ returns the last element of a non-empty vector. There is no $igraph_vector_head()$ function however, as it is easy to write VECTOR(v)[0] instead.

VECTOR — Accessing an element of a vector.

```
#define VECTOR(v)
```

Usage:

VECTOR(v)[0]

to access the first element of the vector, you can also use this in assignments, like:

```
VECTOR(v)[10]=5;
```

Note that there are no range checks right now. This functionality might be redefined later as a real function instead of a #define.

Arguments:

v: The vector object.

Time complexity: O(1).

igraph_vector_get — Access an element of a vector.

```
igraph_real_t igraph_vector_get(const igraph_vector_t* v, igraph_integer_t pos)
```

Arguments:

v: The igraph_vector_t object.

pos: The position of the element, the index of the first element is zero.

Returns:

The desired element.

See also:

igraph_vector_get_ptr() and the VECTOR macro.

Time complexity: O(1).

igraph_vector_get_ptr — Get the address of an element of a vector.

```
igraph_real_t* igraph_vector_get_ptr(const igraph_vector_t* v, igraph_integer_t
```

Arguments:

v: The igraph_vector_t object.

pos: The position of the element, the position of the first element is zero.

Returns:

Pointer to the desired element.

See also:

```
igraph_vector_get() and the VECTOR macro.
```

Time complexity: O(1).

igraph_vector_set — Assignment to an element of a vector.

```
void igraph_vector_set(igraph_vector_t* v, igraph_integer_t pos, igraph_real_t
```

Arguments:

v: The igraph_vector_t element.

pos: Position of the element to set.

value: New value of the element.

See also:

igraph_vector_get().

igraph_vector_tail — Returns the last element in a vector.

```
igraph_real_t igraph_vector_tail(const igraph_vector_t *v);
```

It is an error to call this function on an empty vector, the result is undefined.

Arguments:

v: The vector object.

Returns:

The last element.

Time complexity: O(1).

Vector views

igraph_vector_view — Handle a regular C array as a igraph_vector_t.

This is a special igraph_vector_t constructor. It allows to handle a regular C array as a igraph_vector_t temporarily. Be sure that you *don't* ever call the destructor (igraph_vector_destroy()) on objects created by this constructor.

Arguments:

v: Pointer to an uninitialized igraph_vector_t object.

data: Pointer, the C array. It may not be NULL, except when length is zero.

length: The length of the C array.

Returns:

Pointer to the vector object, the same as the v parameter, for convenience.

Time complexity: O(1)

Copying vectors

igraph_vector_copy_to — Copies the contents of a vector to a C array.

```
void igraph_vector_copy_to(const igraph_vector_t *v, igraph_real_t *to);
```

The C array should have sufficient length.

Arguments:

v: The vector object.

to: The C array.

Time complexity: O(n), n is the size of the vector.

 $\verb|igraph_vector_update -- Update a vector from another one.|\\$

After this operation the contents of to will be exactly the same as that of from. The vector to will be resized if it was originally shorter or longer than from.

Arguments:

to: The vector to update.

from: The vector to update from.

Returns:

Error code.

Time complexity: O(n), the number of elements in *from*.

igraph_vector_append — Append a vector to another one.

The target vector will be resized (except when from is empty).

Arguments:

to: The vector to append to.

from: The vector to append, it is kept unchanged.

Returns:

Error code.

Time complexity: O(n), the number of elements in the new vector.

igraph_vector_swap — Swap all elements of two vectors.

```
igraph_error_t igraph_vector_swap(igraph_vector_t *v1, igraph_vector_t *v2);
```

Arguments:

v1: The first vector.

v2: The second vector.

Returns:

Error code.

Time complexity: O(1).

Exchanging elements

igraph_vector_swap_elements — Swap two elements in a vector.

Note that currently no range checking is performed.

Arguments:

- v: The input vector.
- i: Index of the first element.
- *j*: Index of the second element (may be the same as the first one).

Returns:

Error code, currently always IGRAPH_SUCCESS.

Time complexity: O(1).

igraph_vector_reverse — Reverse the elements of a vector.

```
igraph_error_t igraph_vector_reverse(igraph_vector_t *v);
```

The first element will be last, the last element will be first, etc.

Arguments:

v: The input vector.

Returns:

Error code, currently always IGRAPH_SUCCESS.

Time complexity: O(n), the number of elements.

igraph_vector_shuffle — Shuffles a vector in-place using the Fisher-Yates method.

```
igraph_error_t igraph_vector_shuffle(igraph_vector_t *v);
```

The Fisher-Yates shuffle ensures that every permutation is equally probable when using a proper randomness source. Of course this does not apply to pseudo-random generators as the cycle of these generators is less than the number of possible permutations of the vector if the vector is long enough.

Arguments:

v: The vector object.

Returns:

Error code, currently always IGRAPH_SUCCESS.

Time complexity: O(n), n is the number of elements in the vector.

References:

(Fisher & Yates 1963) R. A. Fisher and F. Yates. Statistical Tables for Biological, Agri-

cultural and Medical Research. Oliver and Boyd, 6th edition, 1963,

page 37.

(Knuth 1998) D. E. Knuth. Seminumerical Algorithms, volume 2 of The Art of

Computer Programming. Addison-Wesley, 3rd edition, 1998, page

145.

Example 7.1. File examples/simple/igraph_fisher_yates_shuffle.c

igraph_vector_permute — Permutes the elements of a vector in place according to an index vector.

```
igraph_error_t igraph_vector_permute(igraph_vector_t* v, const igraph_vector_in
```

This function takes a vector v and a corresponding index vector ind, and permutes the elements of v such that v[ind[i]] is moved to become v[i] after the function is executed.

It is an error to call this function with an index vector that does not represent a valid permutation. Each element in the index vector must be between 0 and the length of the vector minus one (inclusive), and each such element must appear only once. The function does not attempt to validate the index vector.

The index vector that this function takes is compatible with the index vector returned from igraph_vector_qsort_ind(); passing in the index vector from igraph_vector_q-sort_ind() will sort the original vector.

As a special case, this function allows the index vector to be *shorter* than the vector being permuted, in which case the elements whose indices do not occur in the index vector will be removed from the vector.

Arguments:

v: the vector to permute

ind: the index vector

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: O(n), the size of the vector.

Vector operations

igraph_vector_add_constant — Add a constant to the vector.

void igraph_vector_add_constant(igraph_vector_t *v, igraph_real_t plus);
plus is added to every element of v. Note that overflow might happen.

Arguments:

v: The input vector.

plus: The constant to add.

Time complexity: O(n), the number of elements.

igraph_vector_scale — Multiplies all elements of a vector by a constant.

```
void igraph_vector_scale(igraph_vector_t *v, igraph_real_t by);
```

Arguments:

v: The vector.

by: The constant.

Returns:

Error code. The current implementation always returns with success.

Added in version 0.2.

Time complexity: O(n), the number of elements in a vector.

igraph_vector_add — Add two vectors.

Add the elements of v2 to v1, the result is stored in v1. The two vectors must have the same length.

Arguments:

v1: The first vector, the result will be stored here.

v2: The second vector, its contents will be unchanged.

Returns:

Error code.

Time complexity: O(n), the number of elements.

igraph_vector_sub — Subtract a vector from another one.

Subtract the elements of v2 from v1, the result is stored in v1. The two vectors must have the same length.

Arguments:

- v1: The first vector, to subtract from. The result is stored here.
- v2: The vector to subtract, it will be unchanged.

Returns:

Error code.

Time complexity: O(n), the length of the vectors.

igraph_vector_mul — Multiply two vectors.

v1 will be multiplied by v2, elementwise. The two vectors must have the same length.

Arguments:

- v1: The first vector, the result will be stored here.
- v2: The second vector, it is left unchanged.

Returns:

Error code.

Time complexity: O(n), the number of elements.

igraph_vector_div — Divide a vector by another one.

v1 is divided by v2, elementwise. They must have the same length. If the base type of the vector can generate divide by zero errors then please make sure that v2 contains no zero if you want to avoid trouble.

Arguments:

- v1: The dividend. The result is also stored here.
- v2: The divisor, it is left unchanged.

Returns:

Error code.

Time complexity: O(n), the length of the vectors.

igraph_vector_floor — Transform a real vector to an integer vector by flooring each element.

```
igraph_error_t igraph_vector_floor(const igraph_vector_t *from, igraph_vector_i
```

Flooring means rounding down to the nearest integer.

Arguments:

from: The original real vector object.

to: Pointer to an initialized integer vector. The result will be stored here.

Returns:

Error code: IGRAPH_ENOMEM: out of memory

Time complexity: O(n), where n is the number of elements in the vector.

Vector comparisons

igraph_vector_all_e — Are all elements equal?

Checks element-wise equality of two vectors. For vectors containing floating point values, consider using igraph_matrix_all_almost_e().

Arguments:

1hs: The first vector.

rhs: The second vector.

Returns:

True if the elements in the 1hs are all equal to the corresponding elements in rhs. Returns false if the lengths of the vectors don't match.

Time complexity: O(n), the length of the vectors.

igraph_vector_all_almost_e — Are all elements almost equal?

Checks if the elements of two vectors are equal within a relative tolerance.

Arguments:

1hs: The first vector.

rhs: The second vector.

eps: Relative tolerance, see igraph_almost_equals() for details.

Returns:

True if the two vectors are almost equal, false if there is at least one differing element or if the vectors are not of the same size.

igraph_vector_all 1 — Are all elements less?

Arguments:

1hs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the 1hs are all less than the corresponding elements in rhs. Returns 0 (=false) if the lengths of the vectors don't match. If any element is NaN, it will return 0 (=false).

Time complexity: O(n), the length of the vectors.

igraph_vector_all_g — Are all elements greater?

Arguments:

1hs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the 1hs are all greater than the corresponding elements in rhs. Returns 0 (=false) if the lengths of the vectors don't match. If any element is NaN, it will return 0 (=false).

Time complexity: O(n), the length of the vectors.

igraph_vector_all_le — Are all elements less or equal?

Arguments:

1hs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the 1hs are all less than or equal to the corresponding elements in rhs. Returns 0 (=false) if the lengths of the vectors don't match. If any element is NaN, it will return 0 (=false).

Time complexity: O(n), the length of the vectors.

igraph_vector_all_ge — Are all elements greater or equal?

Arguments:

1hs: The first vector.

rhs: The second vector.

Returns:

Positive integer (=true) if the elements in the 1hs are all greater than or equal to the corresponding elements in rhs. Returns 0 (=false) if the lengths of the vectors don't match. If any element is NaN, it will return 0 (=false).

Time complexity: O(n), the length of the vectors.

igraph_vector_zapsmall — Replaces small elements of a vector by exact zeros.

```
igraph_error_t igraph_vector_zapsmall(igraph_vector_t *v, igraph_real_t tol);
```

Vector elements which are smaller in magnitude than the given absolute tolerance will be replaced by exact zeros. The default tolerance corresponds to two-thirds of the representable digits of igraph_real_t, i.e. DBL_EPSILON^(2/3) which is approximately 10^-10.

Arguments:

- v: The vector to process, it will be changed in-place.
- to1: Tolerance value. Numbers smaller than this in magnitude will be replaced by zeros. Pass in zero to use the default tolerance. Must not be negative.

Returns:

Error code.

See also:

igraph_vector_all_almost_e() and igraph_almost_equals() to perform comparisons with relative tolerances.

igraph_vector_lex_cmp — Lexicographical comparison of two vectors (type-safe variant).

```
int igraph_vector_lex_cmp(
    const igraph_vector_t *lhs, const igraph_vector_t *rhs
);
```

If the elements of two vectors match but one is shorter, the shorter one comes first. Thus $\{1, 3\}$ comes after $\{1, 2, 3\}$, but before $\{1, 3, 4\}$.

This function is typically used together with igraph_vector_list_sort().

Arguments:

1hs: Pointer to the first vector.

rhs: Pointer to the second vector.

Returns:

-1 if 1hs is lexicographically smaller, 0 if 1hs and rhs are equal, else 1.

See also:

igraph_vector_lex_cmp_untyped() for an untyped variant of this function, or igraph_vector_colex_cmp() to compare vectors starting from the last element.

Time complexity: O(n), the number of elements in the smaller vector.

```
Example 7.2. File examples/simple/igraph_vector_int_list_sort.c
```

igraph_vector_lex_cmp_untyped — Lexicographical comparison of two vectors (non-type-safe).

```
int igraph_vector_lex_cmp_untyped(const void *lhs, const void *rhs);
```

If the elements of two vectors match but one is shorter, the shorter one comes first. Thus $\{1, 3\}$ comes after $\{1, 2, 3\}$, but before $\{1, 3, 4\}$.

This function is typically used together with igraph_vector_ptr_sort().

Arguments:

1hs: Pointer to a pointer to the first vector (interpreted as an igraph_vector_t **).

rhs: Pointer to a pointer to the second vector (interpreted as an igraph_vector_t **).

Returns:

-1 if 1hs is lexicographically smaller, 0 if 1hs and rhs are equal, else 1.

See also:

```
igraph_vector_lex_cmp() for a type-safe variant of this function, or igraph_vec-
tor_colex_cmp_untyped() to compare vectors starting from the last element.
```

Time complexity: O(n), the number of elements in the smaller vector.

igraph_vector_colex_cmp — Colexicographical comparison of two vectors.

```
int igraph_vector_colex_cmp(
     const igraph_vector_t *lhs, const igraph_vector_t *rhs
);
```

This comparison starts from the last element of both vectors and moves backward. If the elements of two vectors match but one is shorter, the shorter one comes first. Thus $\{1, 2\}$ comes after $\{3, 2, 1\}$, but before $\{0, 1, 2\}$.

This function is typically used together with igraph_vector_list_sort().

Arguments:

1hs: Pointer to a pointer to the first vector.

rhs: Pointer to a pointer to the second vector.

Returns:

-1 if 1hs in reverse order is lexicographically smaller than the reverse of rhs, 0 if 1hs and rhs are equal, else 1.

See also:

igraph_vector_colex_cmp_untyped() for an untyped variant of this function, or igraph_vector_lex_cmp() to compare vectors starting from the first element.

Time complexity: O(n), the number of elements in the smaller vector.

```
Example 7.3. File examples/simple/igraph_vector_int_list_sort.c
```

igraph_vector_colex_cmp_untyped — Colexicographical comparison of two vectors.

```
int igraph_vector_colex_cmp_untyped(const void *lhs, const void *rhs);
```

This comparison starts from the last element of both vectors and moves backward. If the elements of two vectors match but one is shorter, the shorter one comes first. Thus $\{1, 2\}$ comes after $\{3, 2, 1\}$, but before $\{0, 1, 2\}$.

This function is typically used together with igraph_vector_ptr_sort().

Arguments:

1hs: Pointer to a pointer to the first vector (interpreted as an igraph_vector_t **).

rhs: Pointer to a pointer to the second vector (interpreted as an igraph_vector_t **).

Returns:

-1 if 1hs in reverse order is lexicographically smaller than the reverse of rhs, 0 if 1hs and rhs are equal, else 1.

See also:

igraph_vector_colex_cmp() for a type-safe variant of this function, igraph_vector_lex_cmp_untyped() to compare vectors starting from the first element.

Time complexity: O(n), the number of elements in the smaller vector.

Finding minimum and maximum

igraph_vector_min — Smallest element of a vector.

```
igraph_real_t igraph_vector_min(const igraph_vector_t* v);
```

The vector must be non-empty.

Arguments:

v: The input vector.

Returns:

The smallest element of v, or NaN if any element is NaN.

Time complexity: O(n), the number of elements.

igraph_vector_max — Largest element of a vector.

```
igraph_real_t igraph_vector_max(const igraph_vector_t* v);
```

If the size of the vector is zero, an arbitrary number is returned.

Arguments:

v: The vector object.

Returns:

The maximum element of v, or NaN if any element is NaN.

Time complexity: O(n), the number of elements.

igraph_vector_which_min — Index of the smallest element.

```
igraph_integer_t igraph_vector_which_min(const igraph_vector_t* v);
```

The vector must be non-empty. If the smallest element is not unique, then the index of the first is returned. If the vector contains NaN values, the index of the first NaN value is returned.

Arguments:

v: The input vector.

Returns:

Index of the smallest element.

Time complexity: O(n), the number of elements.

igraph_vector_which_max — Gives the index of the maximum element of the vector.

```
igraph_integer_t igraph_vector_which_max(const igraph_vector_t* v);
```

If the size of the vector is zero, -1 is returned. If the largest element is not unique, then the index of the first is returned. If the vector contains NaN values, the index of the first NaN value is returned.

Arguments:

v: The vector object.

Returns:

The index of the first maximum element.

Time complexity: O(n), n is the size of the vector.

igraph_vector_minmax — Minimum and maximum elements of a vector.

Handy if you want to have both the smallest and largest element of a vector. The vector is only traversed once. The vector must be non-empty. If a vector contains at least one NaN, both min and max will be NaN.

Arguments:

v: The input vector. It must contain at least one element.

min: Pointer to a base type variable, the minimum is stored here.

max: Pointer to a base type variable, the maximum is stored here.

Time complexity: O(n), the number of elements.

igraph_vector_which_minmax — Index of the minimum and maximum elements.

Handy if you need the indices of the smallest and largest elements. The vector is traversed only once. The vector must be non-empty. If the minimum or maximum is not unique, the index of the first minimum or the first maximum is returned, respectively. If a vector contains at least one NaN, both which_min and which_max will point to the first NaN value.

Arguments:

v: The input vector. It must contain at least one element.

which_min: The index of the minimum element will be stored here.

which max: The index of the maximum element will be stored here.

Time complexity: O(n), the number of elements.

Vector properties

igraph_vector_empty — Decides whether the size of the vector is zero.

```
igraph_bool_t igraph_vector_empty(const igraph_vector_t* v);
```

Arguments:

v: The vector object.

Returns:

Non-zero number (true) if the size of the vector is zero and zero (false) otherwise.

Time complexity: O(1).

igraph_vector_size — Returns the size (=length) of the vector.

igraph_integer_t igraph_vector_size(const igraph_vector_t* v);
Arguments:
v: The vector object

Returns:

The size of the vector.

Time complexity: O(1).

igraph_vector_capacity — Returns the allocated capacity of the vector.

```
igraph_integer_t igraph_vector_capacity(const igraph_vector_t*v);
```

Note that this might be different from the size of the vector (as queried by igraph_vector_size()), and specifies how many elements the vector can hold, without reallocation.

Arguments:

v: Pointer to the (previously initialized) vector object to query.

Returns:

The allocated capacity.

See also:

```
igraph\_vector\_size().
```

Time complexity: O(1).

igraph_vector_sum — Calculates the sum of the elements in the vector.

```
igraph_real_t igraph_vector_sum(const igraph_vector_t *v);
```

For the empty vector 0.0 is returned.

Arguments:

v: The vector object.

Returns:

The sum of the elements.

Time complexity: O(n), the size of the vector.

igraph_vector_prod — Calculates the product of the elements in the vector.

```
igraph_real_t igraph_vector_prod(const igraph_vector_t *v);
```

For the empty vector one (1) is returned.

Arguments:

v: The vector object.

Returns:

The product of the elements.

Time complexity: O(n), the size of the vector.

igraph_vector_isininterval — Checks if all elements of a vector are in the given interval.

Arguments:

v: The vector object.

low: The lower limit of the interval (inclusive).

high: The higher limit of the interval (inclusive).

Returns:

True (positive integer) if the vector is empty or all vector elements are in the interval, false (zero) otherwise. If any element is NaN, it will return 0 (=false).

Time complexity: O(n), the number of elements in the vector.

igraph_vector_maxdifference — The maximum absolute difference of m1 and m2.

The element with the largest absolute value in m1 - m2 is returned. Both vectors must be non-empty, but they not need to have the same length, the extra elements in the longer vector are ignored. If any value is NaN in the shorter vector, the result will be NaN.

Arguments:

m1: The first vector.

m2: The second vector.

Returns:

The maximum absolute difference of m1 and m2.

Time complexity: O(n), the number of elements in the shorter vector.

igraph_vector_is_nan — Check for each element if it is NaN.

igraph_error_t igraph_vector_is_nan(const igraph_vector_t *v, igraph_vector_boo

Arguments:

v: The igraph_vector_t object to check.

is_nan: The resulting boolean vector indicating for each element whether it is NaN or not.

Returns:

Error code, IGRAPH_ENOMEM if there is not enough memory. Note that this function *never* returns an error if the vector *is_nan* will already be large enough.

Time complexity: O(n), the number of elements.

igraph_vector_is_any_nan — Check if any element is NaN.

```
igraph_bool_t igraph_vector_is_any_nan(const igraph_vector_t *v);
```

Arguments:

v: The igraph_vector_t object to check.

Returns:

1 if any element is NaN, 0 otherwise.

Time complexity: O(n), the number of elements.

Searching for elements

igraph_vector_contains — Linear search in a vector.

Check whether the supplied element is included in the vector, by linear search.

Arguments:

v: The input vector.

e: The element to look for.

Returns:

true if the element is found and false otherwise.

Time complexity: O(n), the length of the vector.

igraph_vector_search — Searches in a vector from a given position.

The supplied element what is searched in vector v, starting from element index from. If found then the index of the first instance (after from) is stored in pos.

Arguments:

v: The input vector.

from: The index to start searching from. No range checking is performed.

what: The element to find.

pos: If not NULL then the index of the found element is stored here.

Returns:

Boolean, true if the element was found, false otherwise.

Time complexity: O(m), the number of elements to search, the length of the vector minus the from argument.

igraph_vector_binsearch — Finds an element by binary searching a sorted vector.

It is assumed that the vector is sorted. If the specified element (*what*) is not in the vector, then the position of where it should be inserted (to keep the vector sorted) is returned. If the vector contains any NaN values, the returned value is undefined and *pos* may point to any position.

Arguments:

v: The igraph_vector_t object.

what: The element to search for.

pos:

Pointer to an igraph_integer_t. This is set to the position of an instance of what in the vector if it is present. If v does not contain what then pos is set to the position to which it should be inserted (to keep the the vector sorted of course).

Returns:

Positive integer (true) if what is found in the vector, zero (false) otherwise.

Time complexity: $O(\log(n))$, n is the number of elements in v.

igraph_vector_binsearch_slice — Finds an element by binary searching a sorted slice of a vector.

It is assumed that the indicated slice of the vector, from start to end, is sorted. If the specified element (what) is not in the slice of the vector, then the position of where it should be inserted (to keep the slice sorted) is returned. Note that this means that the returned index will point inside the slice (including its endpoints), but will not evaluate values outside the slice. If the indicated slice contains any NaN values, the returned value is undefined and pos may point to any position within the slice.

Arguments:

v: The igraph_vector_t object.

what: The element to search for.

pos: Pointer to an igraph_integer_t. This is set to the position of an instance of what in the slice

of the vector if it is present. If v does not contain what then pos is set to the position to

which it should be inserted (to keep the the vector sorted).

start: The start position of the slice to search (inclusive).

end: The end position of the slice to search (exclusive).

Returns:

Positive integer (true) if what is found in the vector, zero (false) otherwise.

Time complexity: O(log(n)), n is the number of elements in the slice of v, i.e. end - start.

igraph_vector_binsearch2 — Binary search, without returning the index.

It is assumed that the vector is sorted.

Arguments:

v: The igraph_vector_t object.

what: The element to search for.

Returns:

Positive integer (true) if what is found in the vector, zero (false) otherwise.

Time complexity: $O(\log(n))$, n is the number of elements in v.

Resizing operations

igraph_vector_clear — Removes all elements from a vector.

```
void igraph_vector_clear(igraph_vector_t* v);
```

This function simply sets the size of the vector to zero, it does not free any allocated memory. For that you have to call <code>igraph_vector_destroy()</code>.

Arguments:

v: The vector object.

Time complexity: O(1).

igraph_vector_reserve — Reserves memory for a vector.

```
igraph_error_t igraph_vector_reserve(igraph_vector_t* v, igraph_integer_t capac
```

igraph vectors are flexible, they can grow and shrink. Growing however occasionally needs the data in the vector to be copied. In order to avoid this, you can call this function to reserve space for future growth of the vector.

Note that this function does *not* change the size of the vector. Let us see a small example to clarify things: if you reserve space for 100 elements and the size of your vector was (and still is) 60, then you can surely add additional 40 elements to your vector before it will be copied.

Arguments:

v: The vector object.

capacity: The new allocated size of the vector.

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, should be around O(n), n is the new allocated size of the vector.

igraph vector resize — Resize the vector.

```
igraph_error_t igraph_vector_resize(igraph_vector_t* v, igraph_integer_t new_si
```

Note that this function does not free any memory, just sets the size of the vector to the given one. It can on the other hand allocate more memory if the new size is larger than the previous one. In this case the newly appeared elements in the vector are *not* set to zero, they are uninitialized.

Arguments:

v: The vector object

new_size: The new size of the vector.

Returns:

Error code, IGRAPH_ENOMEM if there is not enough memory. Note that this function *never* returns an error if the vector is made smaller.

See also:

igraph_vector_reserve() for allocating memory for future extensions of a vector.
igraph_vector_resize_min() for deallocating the unnneded memory for a vector.

Time complexity: O(1) if the new size is smaller, operating system dependent if it is larger. In the latter case it is usually around O(n), n is the new size of the vector.

igraph_vector_resize_min — Deallocate the unused memory of a vector.

```
void igraph_vector_resize_min(igraph_vector_t*v);
```

This function attempts to deallocate the unused reserved storage of a vector. If it succeeds, igraph_vector_size() and igraph_vector_capacity() will be the same. The data in the vector is always preserved, even if deallocation is not successful.

Arguments:

v: Pointer to an initialized vector.

See also:

```
igraph_vector_resize(), igraph_vector_reserve().
```

Time complexity: operating system dependent, O(n) at worst.

igraph_vector_push_back — Appends one element to a vector.

```
igraph_error_t igraph_vector_push_back(igraph_vector_t* v, igraph_real_t e);
```

This function resizes the vector to be one element longer and sets the very last element in the vector to e.

Arguments:

- v: The vector object.
- e: The element to append to the vector.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory.

Time complexity: operating system dependent. What is important is that a sequence of n subsequent calls to this function has time complexity O(n), even if there hadn't been any space reserved for the new elements by <code>igraph_vector_reserve()</code>. This is implemented by a trick similar to the C ++ vector class: each time more memory is allocated for a vector, the size of the additionally allocated memory is the same as the vector's current length. (We assume here that the time complexity of memory allocation is at most linear.)

igraph_vector_pop_back — Removes and returns the last element of a vector.

```
igraph_real_t igraph_vector_pop_back(igraph_vector_t* v);
```

It is an error to call this function with an empty vector.

Arguments:

v: The vector object.

Returns:

The removed last element.

Time complexity: O(1).

igraph_vector_insert — Inserts a single element into a vector.

Note that this function does not do range checking. Insertion will shift the elements from the position given to the end of the vector one position to the right, and the new element will be inserted in the empty space created at the given position. The size of the vector will increase by one.

Arguments:

v: The vector object.

pos: The position where the new element is to be inserted.

value: The new element to be inserted.

igraph_vector_remove — Removes a single element from a vector.

```
void igraph_vector_remove(igraph_vector_t *v, igraph_integer_t elem);
```

Note that this function does not do range checking.

Arguments:

v: The vector object.

elem: The position of the element to remove.

Time complexity: O(n-elem), n is the number of elements in the vector.

igraph_vector_remove_section — Deletes a section from a vector.

Arguments:

v: The vector object.

from: The position of the first element to remove.

to: The position of the first element *not* to remove.

Time complexity: O(n-from), n is the number of elements in the vector.

Complex vector operations

igraph_vector_complex_real — Gives the real part of a complex vector.

Arguments:

v: Pointer to a complex vector.

real: Pointer to an initialized vector. The result will be stored here.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the vector.

igraph_vector_complex_imag — Gives the imaginary part of a complex vector.

Arguments:

v: Pointer to a complex vector.

real: Pointer to an initialized vector. The result will be stored here.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the vector.

igraph_vector_complex_realimag — Gives the real and imaginary parts of a complex vector.

Arguments:

v: Pointer to a complex vector.

real: Pointer to an initialized vector. The real part will be stored here.

imag: Pointer to an initialized vector. The imaginary part will be stored here.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the vector.

igraph_vector_complex_create — Creates a complex vector from a real and imaginary part.

Arguments:

v: Pointer to an uninitialized complex vector.

real: Pointer to the real part of the complex vector.

imag: Pointer to the imaginary part of the complex vector.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the vector.

igraph_vector_complex_create_polar — Creates a complex matrix from a magnitude and an angle.

Arguments:

m: Pointer to an uninitialized complex vector.

r: Pointer to a real vector containing magnitudes.

theta: Pointer to a real vector containing arguments (phase angles).

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the vector.

igraph_vector_complex_all_almost_e — Are all elements almost equal?

Checks if the elements of two complex vectors are equal within a relative tolerance.

Arguments:

1hs: The first vector.

rhs: The second vector.

eps: Relative tolerance, see igraph_complex_almost_equals() for details.

Returns:

True if the two vectors are almost equal, false if there is at least one differing element or if the vectors are not of the same size.

igraph_vector_complex_zapsmall — Replaces small elements of a complex vector by exact zeros.

igraph_error_t igraph_vector_complex_zapsmall(igraph_vector_complex_t *v, igraph_

Similarly to igraph_vector_zapsmall(), small elements will be replaced by zeros. The operation is performed separately on the real and imaginary parts of the numbers. This way, complex numbers with a large real part and tiny imaginary part will effectively be transformed to real numbers. The default tolerance corresponds to two-thirds of the representable digits of igraph_real_t, i.e. DBL EPSILON^(2/3) which is approximately 10^-10.

Arguments:

- v: The vector to process, it will be changed in-place.
- to1: Tolerance value. Real and imaginary parts smaller than this in magnitude will be replaced by zeros. Pass in zero to use the default tolerance. Must not be negative.

Returns:

Error code.

See also:

igraph_vector_complex_all_almost_e() and igraph_complex_almost_equals() to perform comparisons with relative tolerances.

Sorting

igraph_vector_sort — Sorts the elements of the vector into ascending order.

```
void igraph_vector_sort(igraph_vector_t *v);
```

If the vector contains any NaN values, the resulting ordering of NaN values is undefined and may appear anywhere in the vector.

Arguments:

v: Pointer to an initialized vector object.

Time complexity: $O(n \log n)$ for n elements.

igraph_vector_reverse_sort — Sorts the elements of the vector into descending order.

```
void igraph_vector_reverse_sort(igraph_vector_t *v);
```

If the vector contains any NaN values, the resulting ordering of NaN values is undefined and may appear anywhere in the vector.

Arguments:

v: Pointer to an initialized vector object.

Time complexity: O(n log n) for n elements.

igraph_vector_qsort_ind — Returns a permutation of indices that sorts a vector.

Takes an unsorted array v as input and computes an array of indices inds such that v[inds[i]], with i increasing from 0, is an ordered array (either ascending or descending, depending on v order). The order of indices for identical elements is not defined. If the vector contains any NaN values, the ordering of NaN values is undefined.

Arguments:

v: the array to be sorted

inds: the output array of indices. This must be initialized, but will be resized

order: whether the output array should be sorted in ascending or descending order. Use

IGRAPH_ASCENDING for ascending and IGRAPH_DESCENDING for descending order.

Returns:

Error code.

This routine uses igraph's built-in qsort routine. Algorithm: 1) create an array of pointers to the elements of v. 2) Pass this array to qsort. 3) after sorting the difference between the pointer value and the first pointer value gives its original position in the array. Use this to set the values of inds.

Set operations on sorted vectors

igraph_vector_intersect_sorted — Calculates the intersection of two sorted vectors.

The elements that are contained in both vectors are stored in the result vector. All three vectors must be initialized.

Instead of the naive intersection which takes O(n), this function uses the set intersection method of Ricardo Baeza-Yates, which is more efficient when one of the vectors is significantly smaller than the other, and gives similar performance on average when the two vectors are equal.

The algorithm keeps the multiplicities of the elements: if an element appears k1 times in the first vector and k2 times in the second, the result will include that element min(k1, k2) times.

Reference: Baeza-Yates R: A fast set intersection algorithm for sorted sequences. In: Lecture Notes in Computer Science, vol. 3109/2004, pp. 400--408, 2004. Springer Berlin/Heidelberg. ISBN: 978-3-540-22341-2.

v1: the first vector

v2: the second vector

result: the result vector, which will also be sorted.

Time complexity: $O(m \log(n))$ where m is the size of the smaller vector and n is the size of the larger one.

igraph_vector_difference_sorted — Calculates the difference between two sorted vectors (considered as sets).

The elements that are contained in only the first vector but not the second are stored in the result vector. All three vectors must be initialized.

Arguments:

v1: the first vector

v2: the second vector

result: the result vector

Pointer vectors (igraph_vector_ptr_t)

The igraph_vector_ptr_t data type is very similar to the igraph_vector_t type, but it stores generic pointers instead of real numbers.

This type has the same space complexity as igraph_vector_t, and most implemented operations work the same way as for igraph_vector_t.

The same VECTOR macro used for ordinary vectors can be used for pointer vectors as well, please note that a typeless generic pointer will be provided by this macro and you may need to cast it to a specific pointer before starting to work with it.

Pointer vectors may have an associated item destructor function which takes a pointer and returns nothing. The item destructor will be called on each item in the pointer vector when it is destroyed by igraph_vector_ptr_destroy() or igraph_vector_ptr_destroy_all(), or when its elements are freed by igraph_vector_ptr_free_all(). Note that the semantics of an item destructor does not coincide with C++ destructors; for instance, when a pointer vector is resized to a smaller size, the extra items will *not* be destroyed automatically! Nevertheless, item destructors may become handy in many cases; for instance, a vector of graphs generated by some function can be destroyed with a single call to igraph_vector_ptr_destroy_all() if the item destructor is set to igraph_destroy().

igraph_vector_ptr_init — Initialize a pointer vector (constructor).

```
igraph_error_t igraph_vector_ptr_init(igraph_vector_ptr_t* v, igraph_integer_t
```

This is the constructor of the pointer vector data type. All pointer vectors constructed this way should be destroyed via calling igraph_vector_ptr_destroy().

Arguments:

v: Pointer to an uninitialized igraph_vector_ptr_t object, to be created.

size: Integer, the size of the pointer vector.

Returns:

Error code: IGRAPH_ENOMEM if out of memory

Time complexity: operating system dependent, the amount of "time" required to allocate size elements.

igraph_vector_ptr_init_copy — Initializes a pointer vector from another one (constructor).

```
igraph_error_t igraph_vector_ptr_init_copy(igraph_vector_ptr_t *to, const igraph_
```

This function creates a pointer vector by copying another one. This is shallow copy, only the pointers in the vector will be copied.

It is potentially dangerous to copy a pointer vector with an associated item destructor. The copied vector will inherit the item destructor, which may cause problems when both vectors are destroyed as the items might get destroyed twice. Make sure you know what you are doing when copying a pointer vector with an item destructor, or unset the item destructor on one of the vectors later.

Arguments:

to: Pointer to an uninitialized pointer vector object.

from: A pointer vector object.

Returns:

Error code: IGRAPH_ENOMEM if out of memory

Time complexity: O(n) if allocating memory for n elements can be done in O(n) time.

igraph_vector_ptr_destroy — Destroys a pointer vector.

```
void igraph_vector_ptr_destroy(igraph_vector_ptr_t* v);
```

The destructor for pointer vectors.

Arguments:

v: Pointer to the pointer vector to destroy.

Time complexity: operating system dependent, the "time" required to deallocate O(n) bytes, n is the number of elements allocated for the pointer vector (not necessarily the number of elements in the vector).

igraph_vector_ptr_free_all — Frees all the elements of a pointer vector.

```
void igraph_vector_ptr_free_all(igraph_vector_ptr_t* v);
```

If an item destructor is set for this pointer vector, this function will first call the destructor on all elements of the vector and then free all the elements using <code>igraph_free()</code>. If an item destructor is not set, the elements will simply be freed.

Arguments:

v: Pointer to the pointer vector whose elements will be freed.

Time complexity: operating system dependent, the "time" required to call the destructor n times and then deallocate O(n) pointers, each pointing to a memory area of arbitrary size. n is the number of elements in the pointer vector.

igraph_vector_ptr_destroy_all — Frees all the elements and destroys the pointer vector.

```
void igraph_vector_ptr_destroy_all(igraph_vector_ptr_t* v);
```

This function is equivalent to igraph_vector_ptr_free_all() followed by igraph_vector_ptr_destroy().

Arguments:

v: Pointer to the pointer vector to destroy.

Time complexity: operating system dependent, the "time" required to deallocate O(n) pointers, each pointing to a memory area of arbitrary size, plus the "time" required to deallocate O(n) bytes, n being the number of elements allocated for the pointer vector (not necessarily the number of elements in the vector).

igraph_vector_ptr_size — Gives the number of elements in the pointer vector.

```
igraph_integer_t igraph_vector_ptr_size(const igraph_vector_ptr_t* v);
```

Arguments:

v: The pointer vector object.

Returns:

The size of the object, i.e. the number of pointers stored.

Time complexity: O(1).

igraph_vector_ptr_clear — Removes all elements from a pointer vector.

```
void igraph_vector_ptr_clear(igraph_vector_ptr_t* v);
```

This function resizes a pointer to vector to zero length. Note that the pointed objects are *not* deallocated, you should call <code>igraph_free()</code> on them, or make sure that their allocated memory is freed in some other way, you'll get memory leaks otherwise. If you have set up an item destructor earlier, the destructor will be called on every element.

Note that the current implementation of this function does *not* deallocate the memory required for storing the pointers, so making a pointer vector smaller this way does not give back any memory. This behavior might change in the future.

Arguments:

v: The pointer vector to clear.

Time complexity: O(1).

igraph_vector_ptr_push_back — Appends an element to the back of a pointer vector.

```
igraph_error_t igraph_vector_ptr_push_back(igraph_vector_ptr_t* v, void* e);
```

Arguments:

- v: The pointer vector.
- e: The new element to include in the pointer vector.

Returns:

Error code.

See also:

igraph_vector_push_back() for the corresponding operation of the ordinary vector type.

Time complexity: O(1) or O(n), n is the number of elements in the vector. The pointer vector implementation ensures that n subsequent push_back operations need O(n) time to complete.

igraph_vector_ptr_pop_back — Removes and returns the last element of a pointer vector.

```
void *igraph_vector_ptr_pop_back(igraph_vector_ptr_t *v);
```

It is an error to call this function with an empty vector.

v: The pointer vector.

Returns:

The removed last element.

Time complexity: O(1).

igraph_vector_ptr_insert — Inserts a single element into a pointer vector.

```
igraph_error_t igraph_vector_ptr_insert(igraph_vector_ptr_t* v, igraph_integer_
```

Note that this function does not do range checking. Insertion will shift the elements from the position given to the end of the vector one position to the right, and the new element will be inserted in the empty space created at the given position. The size of the vector will increase by one.

Arguments:

v: The pointer vector object.

pos: The position where the new element is inserted.

e: The inserted element

igraph_vector_ptr_get — Access an element of a pointer vector.

```
void *igraph_vector_ptr_get(const igraph_vector_ptr_t* v, igraph_integer_t pos)
```

Arguments:

v: Pointer to a pointer vector.

pos: The index of the pointer to return.

Returns:

The pointer at pos position.

Time complexity: O(1).

igraph_vector_ptr_set — Assign to an element of a pointer vector.

```
void igraph_vector_ptr_set(igraph_vector_ptr_t* v, igraph_integer_t pos, void*
```

Arguments:

v: Pointer to a pointer vector.

pos: The index of the pointer to update.

value: The new pointer to set in the vector.

Time complexity: O(1).

igraph_vector_ptr_resize — Resizes a pointer vector.

igraph_error_t igraph_vector_ptr_resize(igraph_vector_ptr_t* v, igraph_integer_

Note that if a vector is made smaller the pointed object are not deallocated by this function and the item destructor is not called on the extra elements.

Arguments:

v: A pointer vector.

newsize: The new size of the pointer vector.

Returns:

Error code.

Time complexity: O(1) if the vector if made smaller. Operating system dependent otherwise, the amount of "time" needed to allocate the memory for the vector elements.

igraph_vector_ptr_sort — Sorts the pointer vector based on an external comparison function.

```
void igraph_vector_ptr_sort(igraph_vector_ptr_t *v, int (*compar)(const void*,
```

Sometimes it is necessary to sort the pointers in the vector based on the property of the element being referenced by the pointer. This function allows us to sort the vector based on an arbitrary external comparison function which accepts two void * pointers p1 and p2 and returns an integer less than, equal to or greater than zero if the first argument is considered to be respectively less than, equal to, or greater than the second. p1 and p2 will point to the pointer in the vector, so they have to be double-dereferenced if one wants to get access to the underlying object the address of which is stored in v.

Arguments:

v: The pointer vector to be sorted.

compar: A quot-compatible comparison function. It must take pointers to the elements of the pointer vector. For example, if the pointer vector contains igraph_vector_t * pointers, then the comparison function must interpret its arguments as igraph_vector_t **.

igraph_vector_ptr_sort_ind — Returns a permutation of indices that sorts a vector of pointers.

Takes an unsorted array v as input and computes an array of indices inds such that v[inds[i]], with iincreasing from 0, is an ordered array (either ascending or descending, depending on v order). The order of indices for identical elements is not defined.

Arguments:

v: the array to be sorted

inds: the output array of indices. This must be initialized, but will be resized

cmp: a comparator function that takes two elements of the pointer vector being sorted (these are

constant pointers on their own) and returns a negative value if the item "pointed to" by the first pointer is smaller than the item "pointed to" by the second pointer, a positive value if

it is larger, or zero if the two items are equal

Returns:

Error code.

This routine uses the C library quort routine. Algorithm: 1) create an array of pointers to the elements of v. 2) Pass this array to quort. 3) after sorting the difference between the pointer value and the first pointer value gives its original position in the array. Use this to set the values of inds.

igraph_vector_ptr_permute — Permutes the elements of a pointer vector in place according to an index vector.

```
igraph_error_t igraph_vector_ptr_permute(igraph_vector_ptr_t* v, const igraph_v
```

This function takes a vector \mathbf{v} and a corresponding index vector ind, and permutes the elements of \mathbf{v} such that $\mathbf{v}[\inf[i]]$ is moved to become $\mathbf{v}[i]$ after the function is executed.

It is an error to call this function with an index vector that does not represent a valid permutation. Each element in the index vector must be between 0 and the length of the vector minus one (inclusive), and each such element must appear only once. The function does not attempt to validate the index vector.

The index vector that this function takes is compatible with the index vector returned from igraph_vector_ptr_sort_ind(); passing in the index vector from igraph_vector_p-tr_sort_ind() will sort the original vector.

As a special case, this function allows the index vector to be *shorter* than the vector being permuted, in which case the elements whose indices do not occur in the index vector will be removed from the vector.

Arguments:

v: the vector to permute

ind: the index vector

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: O(n), the size of the vector.

igraph_vector_ptr_get_item_destructor — Gets the current item destructor for this pointer vector.

```
igraph_finally_func_t* igraph_vector_ptr_get_item_destructor(const igraph_vector)
```

The item destructor is a function which will be called on every non-null pointer stored in this vector when igraph_vector_ptr_destroy(), igraph_vector_ptr_destroy_all() or igraph_vector_ptr_free_all() is called.

Returns:

The current item destructor.

Time complexity: O(1).

igraph_vector_ptr_set_item_destructor — Sets the item destructor for this pointer vector.

```
igraph_finally_func_t* igraph_vector_ptr_set_item_destructor(
    igraph_vector_ptr_t *v, igraph_finally_func_t *func);
```

The item destructor is a function which will be called on every non-null pointer stored in this vector when igraph_vector_ptr_destroy(), igraph_vector_ptr_destroy_all() or igraph_vector_ptr_free_all() is called.

Returns:

The old item destructor.

Time complexity: O(1).

IGRAPH_VECTOR_PTR_SET_ITEM_DESTRUCTOR — Sets the item destructor for this pointer vector (macro version).

```
#define IGRAPH_VECTOR_PTR_SET_ITEM_DESTRUCTOR(v, func)
```

This macro is expanded to <code>igraph_vector_ptr_set_item_destructor()</code>, the only difference is that the second argument is automatically cast to an <code>igraph_finally_func_t*</code>. The cast is necessary in most cases as the destructor functions we use (such as <code>igraph_vector_destroy())</code> take a pointer to some concrete <code>igraph</code> data type, while <code>igraph_finally_func_t</code> expects <code>void*</code>

Deprecated functions

igraph_vector_copy — Initializes a vector from another vector object (deprecated alias).

Warning

Deprecated since version 0.10. Please do not use this function in new code; use igraph_vector_init_copy() instead.

igraph_vector_e — Access an element of a vector (deprecated alias).

igraph_real_t igraph_vector_e(const igraph_vector_t* v, igraph_integer_t pos);

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_vector_get() instead.

igraph_vector_e_ptr — Get the address of an element of a vector.

igraph_real_t* igraph_vector_e_ptr(const igraph_vector_t* v, igraph_integer_t p

Arguments:

v: The igraph_vector_t object.

pos: The position of the element, the position of the first element is zero.

Returns:

Pointer to the desired element.

See also:

igraph_vector_get() and the VECTOR macro.

Time complexity: O(1).

igraph_vector_init_seq — Initializes a vector with a sequence, inclusive endpoints (deprecated).

```
igraph_error_t igraph_vector_init_seq(igraph_vector_t *v, igraph_real_t from, id
```

The vector will contain the numbers from, from+1, ..., to. Note that both endpoints are *inclusive*, contrary to typical usage of ranges in C.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_vector_init_range() instead.

Arguments:

v: Pointer to an uninitialized vector object.

from: The lower limit in the sequence (inclusive).

to: The upper limit in the sequence (inclusive).

Returns:

Error code: IGRAPH_ENOMEM: out of memory.

Time complexity: O(n), the number of elements in the vector.

igraph_vector_ptr_copy — Initializes a pointer vector from another one (deprecated alias).

igraph_error_t igraph_vector_ptr_copy(igraph_vector_ptr_t *to, const igraph_vec

Warning

Deprecated since version 0.10. Please do not use this function in new code; use igraph_vector_ptr_init_copy() instead.

igraph_vector_ptr_e — Access an element of a pointer vector (deprecated alias).

void *igraph_vector_ptr_e(const igraph_vector_ptr_t* v, igraph_integer_t pos);

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_vector_ptr_get() instead.

Matrices

About igraph_matrix_t objects

This type is just an interface to igraph_vector_t.

The igraph_matrix_t type usually stores n elements in O(n) space, but not always. See the documentation of the vector type.

Matrix constructors and destructors

igraph_matrix_init — Initializes a matrix.

Every matrix needs to be initialized before using it. This is done by calling this function. A matrix has to be destroyed if it is not needed any more; see <code>igraph_matrix_destroy()</code>.

Arguments:

m: Pointer to a not yet initialized matrix object to be initialized.

nrow: The number of rows in the matrix.

ncol: The number of columns in the matrix.

Returns:

Error code.

Time complexity: usually O(n), n is the number of elements in the matrix.

igraph_matrix_init_copy — Copies a matrix.

```
igraph_error_t igraph_matrix_init_copy(igraph_matrix_t *to, const igraph_matrix_
```

Creates a matrix object by copying from an existing matrix.

Arguments:

to: Pointer to an uninitialized matrix object.

from: The initialized matrix object to copy.

Returns:

Error code, IGRAPH_ENOMEM if there isn't enough memory to allocate the new matrix.

Time complexity: O(n), the number of elements in the matrix.

igraph_matrix_destroy — Destroys a matrix object.

```
void igraph_matrix_destroy(igraph_matrix_t *m);
```

This function frees all the memory allocated for a matrix object. The destroyed object needs to be reinitialized before using it again.

Arguments:

m: The matrix to destroy.

Time complexity: operating system dependent.

Initializing elements

igraph_matrix_null — Sets all elements in a matrix to zero.

```
void igraph_matrix_null(igraph_matrix_t *m);
```

Arguments:

m: Pointer to an initialized matrix object.

Time complexity: O(n), n is the number of elements in the matrix.

igraph_matrix_fill — Fill with an element.

```
void igraph_matrix_fill(igraph_matrix_t *m, igraph_real_t e);
```

Set the matrix to a constant matrix.

Arguments:

- *m*: The input matrix.
- e: The element to set.

Time complexity: O(mn), the number of elements.

Accessing elements of a matrix

MATRIX — Accessing an element of a matrix.

```
#define MATRIX(m,i,j)
```

Note that there are no range checks right now. This functionality might be redefined as a proper function later.

Arguments:

- m: The matrix object.
- i: The index of the row, starting with zero.
- *j*: The index of the column, starting with zero.

Time complexity: O(1).

igraph_matrix_get — Extract an element from a matrix.

Use this if you need a function for some reason and cannot use the MATRIX macro. Note that no range checking is performed.

Arguments:

m: The input matrix.

row: The row index.

col: The column index.

Returns:

The element in the given row and column.

Time complexity: O(1).

igraph_matrix_get_ptr — Pointer to an element of a matrix.

The function returns a pointer to an element. No range checking is performed.

Arguments:

m: The input matrix.

row: The row index.

col: The column index.

Returns:

Pointer to the element in the given row and column.

Time complexity: O(1).

igraph_matrix_set — Set an element.

Set an element of a matrix. No range checking is performed.

Arguments:

m: The input matrix.

row: The row index.

col: The column index.

value: The new value of the element.

Time complexity: O(1).

Matrix views

igraph_matrix_view — Creates a matrix view into an existing array.

This function lets you treat an existing C array as a matrix. The elements of the matrix are assumed to be stored in column-major order in the array, i.e. the elements of the first column are stored first, followed by the second column and so on.

Since this function creates a view into an existing array, you must *not* destroy the <code>igraph_ma-trix_t</code> object when you are done with it. Similarly, you must *not* call any function on it that may attempt to modify the size of the matrix. Modifying an element in the matrix will modify the underlying array as the two share the same memory block.

Arguments:

m: Pointer to a not yet initialized matrix object where the view will be created.

data: The array that the matrix provides a view into.

nrow: The number of rows in the matrix.

ncol: The number of columns in the matrix.

Returns:

Pointer to the matrix object, the same as the *m* parameter, for convenience.

Time complexity: O(1).

igraph_matrix_view_from_vector — Creates a matrix view that treats an existing vector as a matrix.

```
const igraph_matrix_t *igraph_matrix_view_from_vector(
    const igraph_matrix_t *m, const igraph_vector_t *v,
    igraph_integer_t nrow
);
```

This function lets you treat an existing igraph vector as a matrix. The elements of the matrix are assumed to be stored in column-major order in the vector, i.e. the elements of the first column are stored first, followed by the second column and so on.

Since this function creates a view into an existing vector, you must *not* destroy the igraph_matrix t object when you are done with it. Similarly, you must *not* call any function on it that may

attempt to modify the size of the vector. Modifying an element in the matrix will modify the underlying vector as the two share the same memory block.

Additionally, you must *not* attempt to grow the underlying vector by any vector operation as that may result in a re-allocation of the backing memory block of the vector, and the matrix view will not be informed about the re-allocation so it will point to an invalid memory area afterwards.

Arguments:

m: Pointer to a not yet initialized matrix object where the view will be created.

v: The vector that the matrix will provide a view into.

nrow: The number of rows in the matrix. The number of columns will be derived implicitly from the size of the vector. If the number of items in the vector is not divisible by the number of rows, the last few elements of the vector will not be covered by the view.

Returns:

Error code.

Time complexity: O(1).

Copying matrices

igraph_matrix_copy_to — Copies a matrix to a regular C array.

```
void igraph_matrix_copy_to(const igraph_matrix_t *m, igraph_real_t *to);
```

The matrix is copied columnwise, as this is the format most programs and languages use. The C array should be of sufficient size; there are (of course) no range checks.

Arguments:

m: Pointer to an initialized matrix object.

to: Pointer to a C array; the place to copy the data to.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the matrix.

igraph_matrix_update — Update from another matrix.

This function replicates from in the matrix to. Note that to must be already initialized.

Arguments:

to: The result matrix.

from: The matrix to replicate; it is left unchanged.

Returns:

Error code.

Time complexity: O(mn), the number of elements.

igraph_matrix_swap — Swap two matrices.

```
igraph_error_t igraph_matrix_swap(igraph_matrix_t *m1, igraph_matrix_t *m2);
```

The contents of the two matrices will be swapped.

Arguments:

m1: The first matrix.

m2: The second matrix.

Returns:

Error code.

Time complexity: O(1).

Operations on rows and columns

igraph_matrix_get_row — Extract a row.

Extract a row from a matrix and return it as a vector.

Arguments:

m: The input matrix.

res: Pointer to an initialized vector; it will be resized if needed.

index: The index of the row to select.

Returns:

Error code.

Time complexity: O(n), the number of columns in the matrix.

igraph_matrix_get_col — Select a column.

```
igraph_error_t igraph_matrix_get_col(const igraph_matrix_t *m,
```

```
igraph_vector_t *res,
igraph_integer_t index);
```

Extract a column of a matrix and return it as a vector.

Arguments:

m: The input matrix.

res: The result will we stored in this vector. It should be initialized and will be resized as needed.

index: The index of the column to select.

Returns:

Error code.

Time complexity: O(n), the number of rows in the matrix.

igraph_matrix_set_row — Set a row from a vector.

Sets the elements of a row with the given vector. This has the effect of setting row index to have the elements in the vector v. The length of the vector and the number of columns in the matrix must match, otherwise an error is triggered.

Arguments:

m: The input matrix.

v: The vector containing the new elements of the row.

index: Index of the row to set.

Returns:

Error code.

Time complexity: O(n), the number of columns in the matrix.

igraph_matrix_set_col — Set a column from a vector.

Sets the elements of a column with the given vector. In effect, column index will be set with elements from the vector v. The length of the vector and the number of rows in the matrix must match, otherwise an error is triggered.

Arguments:

m: The input matrix.

v: The vector containing the new elements of the column.

index: Index of the column to set.

Returns:

Error code.

Time complexity: O(m), the number of rows in the matrix.

igraph_matrix_swap_rows — Swap two rows.

Swap two rows in the matrix.

Arguments:

- m: The input matrix.
- i: The index of the first row.
- 7: The index of the second row.

Returns:

Error code.

Time complexity: O(n), the number of columns.

igraph_matrix_swap_cols — Swap two columns.

Swap two columns in the matrix.

Arguments:

- *m*: The input matrix.
- *i*: The index of the first column.
- *j*: The index of the second column.

Returns:

Error code.

Time complexity: O(m), the number of rows.

igraph_matrix_select_rows — Select some rows of a matrix.

Data structure library: vector, matrix, other data types

This function selects some rows of a matrix and returns them in a new matrix. The result matrix should be initialized before calling the function.

Arguments:

m: The input matrix.

res: The result matrix. It should be initialized and will be resized as needed.

rows: Vector; it contains the row indices (starting with zero) to extract. Note that no range checking is performed.

Returns:

Error code.

Time complexity: O(nm), n is the number of rows, m the number of columns of the result matrix.

igraph_matrix_select_cols — Select some columns of a matrix.

This function selects some columns of a matrix and returns them in a new matrix. The result matrix should be initialized before calling the function.

Arguments:

m: The input matrix.

res: The result matrix. It should be initialized and will be resized as needed.

cols: Vector; it contains the column indices (starting with zero) to extract. Note that no range checking is performed.

Returns:

Error code.

Time complexity: O(nm), n is the number of rows, m the number of columns of the result matrix.

igraph_matrix_select_rows_cols — Select some rows and columns of a matrix.

This function selects some rows and columns of a matrix and returns them in a new matrix. The result matrix should be initialized before calling the function.

Arguments:

m: The input matrix.

res: The result matrix. It should be initialized and will be resized as needed.

rows: Vector; it contains the row indices (starting with zero) to extract. Note that no range checking

is performed.

cols: Vector; it contains the column indices (starting with zero) to extract. Note that no range

checking is performed.

Returns:

Error code.

Time complexity: O(nm), n is the number of rows, m the number of columns of the result matrix.

Matrix operations

igraph_matrix_add_constant — Add a constant to every element.

```
void igraph_matrix_add_constant(igraph_matrix_t *m, igraph_real_t plus);
```

Arguments:

m: The input matrix.

plud: The constant to add.

Time complexity: O(mn), the number of elements.

igraph_matrix_scale — Multiplies each element of the matrix by a constant.

```
void igraph_matrix_scale(igraph_matrix_t *m, igraph_real_t by);
```

Arguments:

m: The matrix.

by: The constant.

Added in version 0.2.

Time complexity: O(n), the number of elements in the matrix.

igraph_matrix_add — Add two matrices.

Add *m2* to *m1*, and store the result in *m1*. The dimensions of the matrices must match.

Arguments:

- *n*1: The first matrix; the result will be stored here.
- *m2*: The second matrix; it is left unchanged.

Returns:

Error code.

Time complexity: O(mn), the number of elements.

igraph_matrix_sub — Difference of two matrices.

Subtract *m2* from *m1* and store the result in *m1*. The dimensions of the two matrices must match.

Arguments:

- *m*1: The first matrix; the result is stored here.
- *m2*: The second matrix; it is left unchanged.

Returns:

Error code.

Time complexity: O(mn), the number of elements.

igraph_matrix_mul_elements — Elementwise multiplication.

Multiply m1 by m2 elementwise and store the result in m1. The dimensions of the two matrices must match.

Arguments:

- *m*1: The first matrix; the result is stored here.
- *m2*: The second matrix; it is left unchanged.

Returns:

Error code.

Time complexity: O(mn), the number of elements.

igraph_matrix_div_elements — Elementwise division.

Divide m1 by m2 elementwise and store the result in m1. The dimensions of the two matrices must match.

Arguments:

m1: The dividend. The result is store here.

m2: The divisor. It is left unchanged.

Returns:

Error code.

Time complexity: O(mn), the number of elements.

igraph_matrix_sum — Sum of elements.

```
igraph_real_t igraph_matrix_sum(const igraph_matrix_t *m);
```

Returns the sum of the elements of a matrix.

Arguments:

m: The input matrix.

Returns:

The sum of the elements.

Time complexity: O(mn), the number of elements in the matrix.

igraph_matrix_prod — Product of the elements.

```
igraph_real_t igraph_matrix_prod(const igraph_matrix_t *m);
```

Note this function can result in overflow easily, even for not too big matrices.

Arguments:

m: The input matrix.

Returns:

The product of the elements.

Time complexity: O(mn), the number of elements.

igraph matrix rowsum — Rowwise sum.

Calculate the sum of the elements in each row.

Arguments:

m: The input matrix.

res: Pointer to an initialized vector; the result is stored here. It will be resized if necessary.

Returns:

Error code.

Time complexity: O(mn), the number of elements in the matrix.

igraph_matrix_colsum — Columnwise sum.

Calculate the sum of the elements in each column.

Arguments:

m: The input matrix.

res: Pointer to an initialized vector; the result is stored here. It will be resized if necessary.

Returns:

Error code.

Time complexity: O(mn), the number of elements in the matrix.

igraph_matrix_transpose — Transpose a matrix.

```
igraph_error_t igraph_matrix_transpose(igraph_matrix_t *m);
```

Calculate the transpose of a matrix. Note that the function reallocates the memory used for the matrix.

Arguments:

n: The input (and output) matrix.

Returns:

Error code.

Time complexity: O(mn), the number of elements in the matrix.

Matrix comparisons

igraph_matrix_all_e — Are all elements equal?

Checks element-wise equality of two matrices. For matrices containing floating point values, consider using igraph_matrix_all_almost_e().

Arguments:

1hs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the 1hs are all equal to the corresponding elements in rhs. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: O(nm), the size of the matrices.

igraph_matrix_all_almost_e — Are all elements almost equal?

Checks if the elements of two matrices are equal within a relative tolerance.

Arguments:

1hs: The first matrix.

rhs: The second matrix.

eps: Relative tolerance, see igraph_almost_equals() for details.

Returns:

True if the two matrices are almost equal, false if there is at least one differing element or if the matrices are not of the same dimensions.

igraph_matrix_all_1 — Are all elements less?

1hs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the 1hs are all less than the corresponding elements in rhs. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: O(nm), the size of the matrices.

igraph_matrix_all_g — Are all elements greater?

Arguments:

1hs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the 1hs are all greater than the corresponding elements in rhs. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: O(nm), the size of the matrices.

igraph_matrix_all_le — Are all elements less or equal?

Arguments:

1hs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the 1hs are all less than or equal to the corresponding elements in rhs. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: O(nm), the size of the matrices.

igraph_matrix_all_ge — Are all elements greater or equal?

1hs: The first matrix.

rhs: The second matrix.

Returns:

Positive integer (=true) if the elements in the 1hs are all greater than or equal to the corresponding elements in rhs. Returns 0 (=false) if the dimensions of the matrices don't match.

Time complexity: O(nm), the size of the matrices.

igraph_matrix_zapsmall — Replaces small elements of a matrix by exact zeros.

```
igraph_error_t igraph_matrix_zapsmall(igraph_matrix_t *m, igraph_real_t tol);
```

Matrix elements which are smaller in magnitude than the given absolute tolerance will be replaced by exact zeros. The default tolerance corresponds to two-thirds of the representable digits of igraph_real_t, i.e. DBL_EPSILON^(2/3) which is approximately 10^-10.

Arguments:

m: The matrix to process, it will be changed in-place.

to1: Tolerance value. Numbers smaller than this in magnitude will be replaced by zeros. Pass in zero to use the default tolerance. Must not be negative.

Returns:

Error code.

See also:

 $igraph_matrix_all_almost_e()$ and $igraph_almost_equals()$ to perform comparisons with relative tolerances.

Combining matrices

igraph matrix rbind — Combine two matrices rowwise.

This function places the rows of from below the rows of to and stores the result in to. The number of columns in the two matrices must match.

to: The upper matrix; the result is also stored here.

from: The lower matrix. It is left unchanged.

Returns:

Error code.

Time complexity: O(mn), the number of elements in the newly created matrix.

igraph_matrix_cbind — Combine matrices columnwise.

This function places the columns of from on the right of to, and stores the result in to.

Arguments:

to: The left matrix; the result is stored here too.

from: The right matrix. It is left unchanged.

Returns:

Error code.

Time complexity: O(mn), the number of elements on the new matrix.

Finding minimum and maximum

igraph_matrix_min — Smallest element of a matrix.

```
igraph_real_t igraph_matrix_min(const igraph_matrix_t *m);
```

The matrix must be non-empty.

Arguments:

m: The input matrix.

Returns:

The smallest element of *m*, or NaN if any element is NaN.

Time complexity: O(mn), the number of elements in the matrix.

igraph_matrix_max — Largest element of a matrix.

```
igraph_real_t igraph_matrix_max(const igraph_matrix_t *m);
```

If the matrix is empty, an arbitrary number is returned.

Arguments:

m: The matrix object.

Returns:

The maximum element of m, or NaN if any element is NaN.

Added in version 0.2.

Time complexity: O(mn), the number of elements in the matrix.

igraph_matrix_which_min — Indices of the smallest element.

The matrix must be non-empty. If the smallest element is not unique, then the indices of the first such element are returned. If the matrix contains NaN values, the indices of the first NaN value are returned.

Arguments:

- m: The matrix.
- i: Pointer to an igraph_integer_t. The row index of the minimum is stored here.
- 7: Pointer to an igraph integer t. The column index of the minimum is stored here.

Time complexity: O(mn), the number of elements.

igraph_matrix_which_max — Indices of the largest element.

The matrix must be non-empty. If the largest element is not unique, then the indices of the first such element are returned. If the matrix contains NaN values, the indices of the first NaN value are returned.

Arguments:

- *m*: The matrix.
- i: Pointer to an igraph_integer_t. The row index of the maximum is stored here.
- *j*: Pointer to an igraph_integer_t. The column index of the maximum is stored here.

Time complexity: O(mn), the number of elements.

igraph_matrix_minmax — Minimum and maximum elements of a matrix.

Handy if you want to have both the smallest and largest element of a matrix. The matrix is only traversed once. The matrix must be non-empty. If a matrix contains at least one NaN, both min and max will be NaN.

Arguments:

m: The input matrix. It must contain at least one element.

min: Pointer to a base type variable. The minimum is stored here.

max: Pointer to a base type variable. The maximum is stored here.

Time complexity: O(mn), the number of elements.

igraph_matrix_which_minmax — Indices of the minimum and maximum elements.

Handy if you need the indices of the smallest and largest elements. The matrix is traversed only once. The matrix must be non-empty. If the minimum or maximum is not unique, the index of the first minimum or the first maximum is returned, respectively. If a matrix contains at least one NaN, both which_min and which_max will point to the first NaN value.

Arguments:

m: The input matrix.

imin: Pointer to an igraph_integer_t, the row index of the minimum is stored here.

jmin: Pointer to an igraph_integer_t, the column index of the minimum is stored here.

imax: Pointer to an igraph_integer_t, the row index of the maximum is stored here.

jmax: Pointer to an igraph_integer_t, the column index of the maximum is stored here.

Time complexity: O(mn), the number of elements.

Matrix properties

igraph_matrix_empty — Check for an empty matrix.

```
igraph_bool_t igraph_matrix_empty(const igraph_matrix_t *m);
```

It is possible to have a matrix with zero rows or zero columns, or even both. This functions checks for these.

Arguments:

m: The input matrix.

Returns:

Boolean, true if the matrix contains zero elements, and false otherwise.

Time complexity: O(1).

igraph_matrix_isnull — Check for a null matrix.

```
igraph_bool_t igraph_matrix_isnull(const igraph_matrix_t *m);
```

Checks whether all elements are zero.

Arguments:

m: The input matrix.

Returns:

Boolean, true is m contains only zeros and false otherwise.

Time complexity: O(mn), the number of elements.

igraph_matrix_size — The number of elements in a matrix.

```
igraph_integer_t igraph_matrix_size(const igraph_matrix_t *m);
```

Arguments:

m: Pointer to an initialized matrix object.

Returns:

The size of the matrix.

Time complexity: O(1).

igraph_matrix_capacity — Returns the number of elements allocated for a matrix.

```
igraph_integer_t igraph_matrix_capacity(const igraph_matrix_t *m);
```

Note that this might be different from the size of the matrix (as queried by igraph_ma-trix_size(), and specifies how many elements the matrix can hold, without reallocation.

Arguments:

v: Pointer to the (previously initialized) matrix object to query.

Returns:

The allocated capacity.

See also:

```
igraph_matrix_size(), igraph_matrix_nrow(), igraph_matrix_ncol().
Time complexity: O(1).
```

igraph_matrix_nrow — The number of rows in a matrix.

```
igraph_integer_t igraph_matrix_nrow(const igraph_matrix_t *m);
```

Arguments:

m: Pointer to an initialized matrix object.

Returns:

The number of rows in the matrix.

Time complexity: O(1).

igraph_matrix_ncol — The number of columns in a matrix.

```
igraph_integer_t igraph_matrix_ncol(const igraph_matrix_t *m);
```

Arguments:

m: Pointer to an initialized matrix object.

Returns:

The number of columns in the matrix.

Time complexity: O(1).

igraph_matrix_is_symmetric — Check for symmetric matrix.

```
igraph_bool_t igraph_matrix_is_symmetric(const igraph_matrix_t *m);
```

A non-square matrix is not symmetric by definition.

Arguments:

m: The input matrix.

Returns:

Boolean, true if the matrix is square and symmetric, false otherwise.

Time complexity: O(mn), the number of elements. O(1) for non-square matrices.

igraph_matrix_maxdifference — Maximum absolute difference between two matrices.

Calculate the maximum absolute difference of two matrices. Both matrices must be non-empty. If their dimensions differ then a warning is given and the comparison is performed by vectors columnwise from both matrices. The remaining elements in the larger vector are ignored.

Arguments:

m1: The first matrix.

m2: The second matrix.

Returns:

The element with the largest absolute value in m1 - m2.

Time complexity: O(mn), the elements in the smaller matrix.

Searching for elements

igraph_matrix_contains — Search for an element.

Search for the given element in the matrix.

Arguments:

m: The input matrix.

e: The element to search for.

Returns:

Boolean, true if the matrix contains e, false otherwise.

Time complexity: O(mn), the number of elements.

igraph_matrix_search — Search from a given position.

Search for an element in a matrix and start the search from the given position. The search is performed columnwise.

m: The input matrix.

from: The position to search from, the positions are enumerated columnwise.

what: The element to search for.

pos: Pointer to an igraph_integer_t. If the element is found, then this is set to the position of its

first appearance.

row: Pointer to an igraph_integer_t. If the element is found, then this is set to its row index.

col: Pointer to an igraph_integer_t. If the element is found, then this is set to its column index.

Returns:

Boolean, true if the element is found, false otherwise.

Time complexity: O(mn), the number of elements.

Resizing operations

igraph_matrix_resize — Resizes a matrix.

```
igraph_error_t igraph_matrix_resize(igraph_matrix_t *m, igraph_integer_t nrow,
```

This function resizes a matrix by adding more elements to it. The matrix contains arbitrary data after resizing it. That is, after calling this function you cannot expect that element (i,j) in the matrix remains the same as before.

Arguments:

m: Pointer to an already initialized matrix object.

nrow: The number of rows in the resized matrix.

ncol: The number of columns in the resized matrix.

Returns:

Error code.

Time complexity: O(1) if the matrix gets smaller, usually O(n) if it gets larger, n is the number of elements in the resized matrix.

igraph_matrix_resize_min — Deallocates unused memory for a matrix.

```
void igraph_matrix_resize_min(igraph_matrix_t *m);
```

This function attempts to deallocate the unused reserved storage of a matrix.

Arguments:

m: Pointer to an initialized matrix.

See also:

```
igraph_matrix_resize().
```

Time complexity: operating system dependent, O(n) at worst.

igraph_matrix_add_rows — Adds rows to a matrix.

```
igraph_error_t igraph_matrix_add_rows(igraph_matrix_t *m, igraph_integer_t n);
```

Arguments:

- m: The matrix object.
- n: The number of rows to add.

Returns:

Error code, IGRAPH_ENOMEM if there isn't enough memory for the operation.

Time complexity: linear with the number of elements of the new, resized matrix.

igraph matrix add cols — Adds columns to a matrix.

```
igraph_error_t igraph_matrix_add_cols(igraph_matrix_t *m, igraph_integer_t n);
```

Arguments:

- m: The matrix object.
- n: The number of columns to add.

Returns:

Error code, IGRAPH_ENOMEM if there is not enough memory to perform the operation.

Time complexity: linear with the number of elements of the new, resized matrix.

igraph_matrix_remove_row — Remove a row.

```
igraph_error_t igraph_matrix_remove_row(igraph_matrix_t *m, igraph_integer_t ro
A row is removed from the matrix.
```

Arguments:

m: The input matrix.

row: The index of the row to remove.

Returns:

Error code.

Time complexity: O(mn), the number of elements in the matrix.

igraph_matrix_remove_col — Removes a column from a matrix.

```
igraph_error_t igraph_matrix_remove_col(igraph_matrix_t *m, igraph_integer_t co
```

Arguments:

m: The matrix object.

col: The column to remove.

Returns:

Error code, always returns with success.

Time complexity: linear with the number of elements of the new, resized matrix.

Complex matrix operations

igraph_matrix_complex_real — Gives the real part of a complex matrix.

Arguments:

m: Pointer to a complex matrix.

real: Pointer to an initialized matrix. The result will be stored here.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the matrix.

igraph_matrix_complex_imag — Gives the imaginary part of a complex matrix.

m: Pointer to a complex matrix.

imag: Pointer to an initialized matrix. The result will be stored here.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the matrix.

igraph_matrix_complex_realimag — Gives the real and imaginary parts of a complex matrix.

Arguments:

m: Pointer to a complex matrix.

real: Pointer to an initialized matrix. The real part will be stored here.

imag: Pointer to an initialized matrix. The imaginary part will be stored here.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the matrix.

igraph_matrix_complex_create — Creates a complex matrix from a real and imaginary part.

Arguments:

m: Pointer to an uninitialized complex matrix.

real: Pointer to the real part of the complex matrix.

imag: Pointer to the imaginary part of the complex matrix.

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the matrix.

igraph_matrix_complex_create_polar — Creates a complex matrix from a magnitude and an angle.

Arguments:

m: Pointer to an uninitialized complex matrix.

r: Pointer to a real matrix containing magnitudes.

theta: Pointer to a real matrix containing arguments (phase angles).

Returns:

Error code.

Time complexity: O(n), n is the number of elements in the matrix.

igraph_matrix_complex_all_almost_e — Are all elements almost equal?

Checks if the elements of two complex matrices are equal within a relative tolerance.

Arguments:

1hs: The first matrix.

rhs: The second matrix.

eps: Relative tolerance, see igraph complex almost equals() for details.

Returns:

True if the two matrices are almost equal, false if there is at least one differing element or if the matrices are not of the same dimensions.

igraph_matrix_complex_zapsmall — Replaces small elements of a complex matrix by exact zeros.

```
igraph_error_t igraph_matrix_complex_zapsmall(igraph_matrix_complex_t *m, igraph_error_t
```

Similarly to igraph_matrix_zapsmall(), small elements will be replaced by zeros. The operation is performed separately on the real and imaginary parts of the numbers. This way, complex

numbers with a large real part and tiny imaginary part will effectively be transformed to real numbers. The default tolerance corresponds to two-thirds of the representable digits of igraph_real_t, i.e. DBL_EPSILON^(2/3) which is approximately 10^-10.

Arguments:

m: The matrix to process, it will be changed in-place.

to1: Tolerance value. Real and imaginary parts smaller than this in magnitude will be replaced by zeros. Pass in zero to use the default tolerance. Must not be negative.

Returns:

Error code.

See also:

igraph_matrix_complex_all_almost_e() and igraph_complex_almost_e-quals() to perform comparisons with relative tolerances.

Deprecated functions

igraph_matrix_copy — Copies a matrix (deprecated alias).

```
igraph_error_t igraph_matrix_copy(igraph_matrix_t *to, const igraph_matrix_t *f
```

Warning

Deprecated since version 0.10. Please do not use this function in new code; use igraph_matrix_init_copy() instead.

igraph_matrix_e — Extract an element from a matrix (deprecated alias).

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_matrix_get() instead.

igraph_matrix_e_ptr — Pointer to an element of a matrix.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_matrix_get_ptr() instead.

Sparse matrices

About sparse matrices

The igraph_sparsemat_t data type stores sparse matrices, i.e. matrices in which the majority of the elements are zero.

The data type is essentially a wrapper to some of the functions in the CXSparse library, by Tim Davis, see http://faculty.cse.tamu.edu/davis/suitesparse.html

Matrices can be stored in two formats: triplet and column-compressed. The triplet format is intended for sparse matrix initialization, as it is easy to add new (non-zero) elements to it. Most of the computations are done on sparse matrices in column-compressed format, after the user has converted the triplet matrix to column-compressed, via igraph_sparsemat_compress().

Both formats are dynamic, in the sense that new elements can be added to them, possibly resulting the allocation of more memory.

Row and column indices follow the C convention and are zero-based.

Example 7.4. File examples/simple/igraph_sparsemat.c

Example 7.5. File examples/simple/igraph_sparsemat3.c

Example 7.6. File examples/simple/igraph_sparsemat4.c

Example 7.7. File examples/simple/igraph_sparsemat6.c

Example 7.8. File examples/simple/igraph_sparsemat7.c

Example 7.9. File examples/simple/igraph sparsemat8.c

Creating sparse matrix objects

igraph_sparsemat_init — Initializes a sparse matrix, in triplet format.

This is the most common way to create a sparse matrix, together with the igraph_sparse-mat_entry() function, which can be used to add the non-zero elements one by one. Once done, the

user can call $igraph_sparsemat_compress()$ to convert the matrix to column-compressed, to allow computations with it.

The user must call igraph_sparsemat_destroy() on the matrix to deallocate the memory, once the matrix is no more needed.

Arguments:

A: Pointer to a not yet initialized sparse matrix.

rows: The number of rows in the matrix.

cols: The number of columns.

nzmax: The maximum number of non-zero elements in the matrix. It is not compulsory to get this

right, but it is useful for the allocation of the proper amount of memory.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_init_copy — Copies a sparse matrix.

```
igraph_error_t igraph_sparsemat_init_copy(
    igraph_sparsemat_t *to, const igraph_sparsemat_t *from
);
```

Create a sparse matrix object, by copying another one. The source matrix can be either in triplet or column-compressed format.

Exactly the same amount of memory will be allocated to the copy matrix, as it is currently for the original one.

Arguments:

to: Pointer to an uninitialized sparse matrix, the copy will be created here.

from: The sparse matrix to copy.

Returns:

Error code.

Time complexity: O(n+nzmax), the number of columns plus the maximum number of non-zero elements.

igraph_sparsemat_init_diag — Creates a sparse diagonal matrix.

```
igraph_error_t igraph_sparsemat_init_diag(
    igraph_sparsemat_t *A, igraph_integer_t nzmax, const igraph_vector_t *value
    igraph_bool_t compress
);
```

A: An uninitialized sparse matrix, the result is stored here.

nzmax: The maximum number of non-zero elements, this essentially gives the amount of

memory that will be allocated for matrix elements.

values: The values to store in the diagonal, the size of the matrix defined by the length of

this vector.

compress: Whether to create a column-compressed matrix. If false, then a triplet matrix is created.

Returns:

Error code.

Time complexity: O(n), the length of the diagonal vector.

igraph_sparsemat_init_eye — Creates a sparse identity matrix.

```
igraph_error_t igraph_sparsemat_init_eye(
    igraph_sparsemat_t *A, igraph_integer_t n, igraph_integer_t nzmax,
    igraph_real_t value, igraph_bool_t compress
);
```

Arguments:

A: An uninitialized sparse matrix, the result is stored here.

n: The number of rows and number of columns in the matrix.

nzmax: The maximum number of non-zero elements, this essentially gives the amount of

memory that will be allocated for matrix elements.

value: The value to store in the diagonal.

compress: Whether to create a column-compressed matrix. If false, then a triplet matrix is created.

Returns:

Error code.

Time complexity: O(n).

igraph_sparsemat_realloc — Allocates more (or less) memory for a sparse matrix.

```
igraph_error_t igraph_sparsemat_realloc(igraph_sparsemat_t *A, igraph_integer_t
```

Sparse matrices automatically allocate more memory, as needed. To control memory allocation, the user can call this function, to allocate memory for a given number of non-zero elements.

A: The sparse matrix, it can be in triplet or column-compressed format.

nzmax: The new maximum number of non-zero elements.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_destroy — Deallocates memory used by a sparse matrix.

```
void igraph_sparsemat_destroy(igraph_sparsemat_t *A);
```

One destroyed, the sparse matrix must be initialized again, before calling any other operation on it.

Arguments:

A: The sparse matrix to destroy.

Time complexity: O(1).

igraph_sparsemat_view — Initialize a sparse matrix and set all parameters.

This function can be used to temporarily handle existing sparse matrix data, usually created by another software library, as an $igraph_sparsemat_t$ object, and thus avoid unnecessary copying. It supports data stored in either the compressed sparse column format, or the (i, j, x) triplet format where i and j are the matrix indices of a non-zero element, and x is its value.

The compressed sparse column (or row) format is commonly used to represent sparse matrix data. It consists of three vectors, the p column pointers, the i row indices, and the x values. p[k] is the number of non-zero entires in matrix columns k-1 and lower. p[0] is always zero and p[n] is always the total number of non-zero entires in the matrix. i[1] is the row index of the 1-th stored element, while x[1] is its value. If a matrix element with indices (j, k) is explicitly stored, it must be located between positions p[k] and p[k+1] - 1 (inclusive) in the i and x vectors.

Do not call igraph_sparsemat_destroy() on a sparse matrix created with this function. Instead, igraph_free() must be called on the cs field of A to free the storage allocated by this function.

Warning: Matrices created with this function must not be used with functions that may reallocate the underlying storage, such as igraph_sparsemat_entry().

Arguments:

A: The non-initialized sparse matrix.

nzmax: The maximum number of entries, typically the actual number of entries.

m: The number of matrix rows.

n: The number of matrix columns.

p: For a compressed matrix, this is the column pointer vector, and must be of size n+1. For a triplet format matrix, it is a vector of column indices and must be of size nzmax.

i: The row vector. This should contain the row indices of the elements in x. It must be of size nzmax.

x: The values of the non-zero elements of the sparse matrix. It must be of size nzmax.

nz: For a compressed matrix, is must be -1. For a triplet format matrix, is must contain the number of entries.

Returns:

Error code.

Time complexity: O(1).

Query properties of a sparse matrix

igraph_sparsemat_index — Extracts a submatrix or a single element.

This function indexes into a sparse matrix. It serves two purposes. First, it can extract submatrices from a sparse matrix. Second, as a special case, it can extract a single element from a sparse matrix.

Arguments:

A: The input matrix, it must be in column-compressed format.

p: An integer vector, or a null pointer. The selected row index or indices. A null pointer

selects all rows.

q: An integer vector, or a null pointer. The selected column index or indices. A null

pointer selects all columns.

res: Pointer to an uninitialized sparse matrix, or a null pointer. If not a null pointer, then

the selected submatrix is stored here.

constres: Pointer to a real variable or a null pointer. If not a null pointer, then the first non-

zero element in the selected submatrix is stored here, if there is one. Otherwise zero is stored here. This behavior is handy if one wants to select a single entry from the

matrix.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_nrow — Number of rows.

igraph_integer_t igraph_sparsemat_nrow(const igraph_sparsemat_t *A);

Arguments:

A: The input matrix, in triplet or column-compressed format.

Returns:

The number of rows in the A matrix.

Time complexity: O(1).

igraph_sparsemat_ncol — Number of columns.

igraph_integer_t igraph_sparsemat_ncol(const igraph_sparsemat_t *A);

Arguments:

A: The input matrix, in triplet or column-compressed format.

Returns:

The number of columns in the A matrix.

Time complexity: O(1).

igraph_sparsemat_type — Type of a sparse matrix (triplet or column-compressed).

igraph_sparsemat_type_t igraph_sparsemat_type(const igraph_sparsemat_t *A);

Gives whether a sparse matrix is stored in the triplet format or in column-compressed format.

Arguments:

A: The input matrix.

Returns:

Either IGRAPH SPARSEMAT CC or IGRAPH SPARSEMAT TRIPLET.

Time complexity: O(1).

igraph_sparsemat_is_triplet — Is this sparse matrix in triplet format?

igraph_bool_t igraph_sparsemat_is_triplet(const igraph_sparsemat_t *A);

Decides whether a sparse matrix is in triplet format.

Arguments:

A: The input matrix.

Returns:

One if the input matrix is in triplet format, zero otherwise.

Time complexity: O(1).

igraph_sparsemat_is_cc — Is this sparse matrix in column-compressed format?

```
igraph_bool_t igraph_sparsemat_is_cc(const igraph_sparsemat_t *A);
```

Decides whether a sparse matrix is in column-compressed format.

Arguments:

A: The input matrix.

Returns:

One if the input matrix is in column-compressed format, zero otherwise.

Time complexity: O(1).

igraph_sparsemat_is_symmetric — Returns whether a sparse matrix is symmetric.

```
igraph_error_t igraph_sparsemat_is_symmetric(const igraph_sparsemat_t *A, igraph_sparsemat_t *A)
```

Arguments:

A: The input matrix

result: Pointer to an igraph_bool_t; the result is provided here.

Returns:

Error code.

igraph_sparsemat_get — Return the value of a single element from a sparse matrix.

```
igraph_real_t igraph_sparsemat_get(
     const igraph_sparsemat_t *A, igraph_integer_t row, igraph_integer_t col
);
```

A: The input matrix, in triplet or column-compressed format.

row: The row index

col: The column index

Returns:

The value of the cell with the given row and column indices in the matrix; zero if the indices are out of bounds.

Time complexity: TODO.

igraph_sparsemat_getelements — Returns all elements of a sparse matrix.

This function will return the elements of a sparse matrix in three vectors. Two vectors will indicate where the elements are located, and one will specify the elements themselves.

Arguments:

- A: A sparse matrix in either triplet or compressed form.
- *i*: An initialized integer vector. This will store the rows of the returned elements.
- j: An initialized integer vector. For a triplet matrix this will store the columns of the returned elements. For a compressed matrix, if the column index is k, then j[k] is the index in x of the start of the k-th column, and the last element of j is the total number of elements. The total number of elements in the k-th column is therefore j[k+1] j[k]. For example, if there is one element in the first column, and five in the second, j will be set to {0, 1, 6}.
- x: An initialized vector. The elements will be placed here.

Returns:

Error code.

Time complexity: O(n), the number of stored elements in the sparse matrix.

igraph_sparsemat_getelements_sorted — Returns all elements of a sparse matrix, sorted by row and column indices.

This function will sort a sparse matrix and return the elements in three vectors. Two vectors will indicate where the elements are located, and one will specify the elements themselves.

Sorting is done based on the *indices* of the elements, not their numeric values. The returned entries will be sorted by column indices; entries in the same column are then sorted by row indices.

Arguments:

- A: A sparse matrix in either triplet or compressed form.
- i: An initialized integer vector. This will store the rows of the returned elements.
- j: An initialized integer vector. For a triplet matrix this will store the columns of the returned elements. For a compressed matrix, if the column index is k, then j[k] is the index in x of the start of the k-th column, and the last element of j is the total number of elements. The total number of elements in the k-th column is therefore j[k+1] j[k]. For example, if there is one element in the first column, and five in the second, j will be set to {0, 1, 6}.
- x: An initialized vector. The elements will be placed here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_min — Minimum of a sparse matrix.

```
igraph_real_t igraph_sparsemat_min(igraph_sparsemat_t *A);
```

Arguments:

A: The input matrix, column-compressed.

Returns:

The minimum in the input matrix, or IGRAPH_POSINFINITY if the matrix has zero elements.

Time complexity: TODO.

igraph_sparsemat_max — Maximum of a sparse matrix.

```
igraph_real_t igraph_sparsemat_max(igraph_sparsemat_t *A);
```

Arguments:

A: The input matrix, column-compressed.

Returns:

The maximum in the input matrix, or IGRAPH_NEGINFINITY if the matrix has zero elements.

Time complexity: TODO.

igraph_sparsemat_minmax — Minimum and maximum of a sparse matrix.

Arguments:

A: The input matrix, column-compressed.

min: The minimum in the input matrix is stored here, or IGRAPH_POSINFINITY if the matrix has zero elements.

max: The maximum in the input matrix is stored here, or IGRAPH_NEGINFINITY if the matrix has zero elements.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_count_nonzero — Counts nonzero elements of a sparse matrix.

```
igraph_integer_t igraph_sparsemat_count_nonzero(igraph_sparsemat_t *A);
```

Arguments:

A: The input matrix, column-compressed.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_count_nonzerotol — Counts nonzero elements of a sparse matrix, ignoring elements close to zero.

Count the number of matrix entries that are closer to zero than to1.

The: input matrix, column-compressed.

Real: scalar, the tolerance.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_rowsums — Row-wise sums.

Arguments:

A: The input matrix, in triplet or column-compressed format.

res: An initialized vector, the result is stored here. It will be resized as needed.

Returns:

Error code.

Time complexity: O(nz), the number of non-zero elements.

igraph_sparsemat_colsums — Column-wise sums.

Arguments:

A: The input matrix, in triplet or column-compressed format.

res: An initialized vector, the result is stored here. It will be resized as needed.

Returns:

Error code.

Time complexity: O(nz) for triplet matrices, O(nz+n) for column-compressed ones, nz is the number of non-zero elements, n is the number of columns.

igraph_sparsemat_nonzero_storage — Returns number of stored entries of a sparse matrix.

```
igraph_integer_t igraph_sparsemat_nonzero_storage(const igraph_sparsemat_t *A);
```

This function will return the number of stored entries of a sparse matrix. These entries can be zero, and multiple entries can be at the same position. Use <code>igraph_sparsemat_dupl()</code> to sum duplicate entries, and <code>igraph_sparsemat_dropzeros()</code> to remove zeros.

Arguments:

A: A sparse matrix in either triplet or compressed form.

Returns:

Number of stored entries.

Time complexity: O(1).

Operations on sparse matrices

igraph_sparsemat_entry — Adds an element to a sparse matrix.

This function can be used to add the entries to a sparse matrix, after initializing it with <code>igraph_s-parsemat_init()</code>. If you add multiple entries in the same position, they will all be saved, and the resulting value is the sum of all entries in that position.

Arguments:

A: The input matrix, it must be in triplet format.

row: The row index of the entry to add.

col: The column index of the entry to add.

elem: The value of the entry.

Returns:

Error code.

Time complexity: O(1) on average.

igraph_sparsemat_fkeep — Filters the elements of a sparse matrix.

```
igraph_error_t igraph_sparsemat_fkeep(
   igraph_sparsemat_t *A,
   igraph_integer_t (*fkeep)(igraph_integer_t, igraph_integer_t, igraph_real_t
   void *other
);
```

This function can be used to filter the (non-zero) elements of a sparse matrix. For all entries, it calls the supplied function and depending on the return values either keeps, or deleted the element from the matrix.

A: The input matrix, in column-compressed format.

fkeep: The filter function. It must take four arguments: the first is an igraph_integer_t, the

row index of the entry, the second is another <code>igraph_integer_t</code>, the column index. The third is <code>igraph_real_t</code>, the value of the entry. The fourth element is a <code>void</code> pointer, the <code>other</code> argument is passed here. The function must return an <code>int</code>. If this is

zero, then the entry is deleted, otherwise it is kept.

other: A void pointer that is passed to the filtering function.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_dropzeros — Drops the zero elements from a sparse matrix.

```
igraph_error_t igraph_sparsemat_dropzeros(igraph_sparsemat_t *A);
```

As a result of matrix operations, some of the entries in a sparse matrix might be zero. This function removes these entries.

Arguments:

A: The input matrix, it must be in column-compressed format.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_droptol — Drops the almost zero elements from a sparse matrix.

```
igraph_error_t igraph_sparsemat_droptol(igraph_sparsemat_t *A, igraph_real_t to
```

This function is similar to $igraph_sparsemat_dropzeros()$, but it also drops entries that are closer to zero than the given tolerance threshold.

Arguments:

A: The input matrix, it must be in column-compressed format.

to1: Real number, giving the tolerance threshold.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_scale — Scales a sparse matrix.

```
igraph_error_t igraph_sparsemat_scale(igraph_sparsemat_t *A, igraph_real_t by);
```

Multiplies all elements of a sparse matrix, by the given scalar.

Arguments:

A: The input matrix.

by: The scaling factor.

Returns:

Error code.

Time complexity: O(nz), the number of non-zero elements in the matrix.

igraph_sparsemat_permute — Permutes the rows and columns of a sparse matrix.

Arguments:

- A: The input matrix, it must be in column-compressed format.
- p: Integer vector, giving the permutation of the rows.
- *q*: Integer vector, the permutation of the columns.
- res: Pointer to an uninitialized sparse matrix, the result is stored here.

Returns:

Error code.

Time complexity: O(m+n+nz), the number of rows plus the number of columns plus the number of non-zero elements in the matrix.

igraph_sparsemat_transpose — Transposes a sparse matrix.

```
igraph_error_t igraph_sparsemat_transpose(
    const igraph_sparsemat_t *A, igraph_sparsemat_t *res
);
```

A: The input matrix, column-compressed or triple format.

res: Pointer to an uninitialized sparse matrix, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_add — Sum of two sparse matrices.

Arguments:

A: The first input matrix, in column-compressed format.

B: The second input matrix, in column-compressed format.

alpha: Real scalar, A is multiplied by alpha before the addition.

beta: Real scalar, B is multiplied by beta before the addition.

res: Pointer to an uninitialized sparse matrix, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_multiply — Matrix multiplication.

Multiplies two sparse matrices.

Arguments:

A: The first input matrix (left hand side), in column-compressed format.

B: The second input matrix (right hand side), in column-compressed format.

res: Pointer to an uninitialized sparse matrix, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_gaxpy — Matrix-vector product, added to another vector.

Arguments:

A: The input matrix, in column-compressed format.

x: The input vector, its size must match the number of columns in A.

res: This vector is added to the matrix-vector product and it is overwritten by the result.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_add_rows — Adds rows to a sparse matrix.

```
igraph_error_t igraph_sparsemat_add_rows(igraph_sparsemat_t *A, igraph_integer_
```

The current matrix elements are retained and all elements in the new rows are zero.

Arguments:

- A: The input matrix, in triplet or column-compressed format.
- n: The number of rows to add.

Returns:

Error code.

Time complexity: O(1).

igraph_sparsemat_add_cols — Adds columns to a sparse matrix.

```
igraph_error_t igraph_sparsemat_add_cols(igraph_sparsemat_t *A, igraph_integer_
```

The current matrix elements are retained, and all elements in the new columns are zero.

- A: The input matrix, in triplet or column-compressed format.
- n: The number of columns to add.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_resize — Resizes a sparse matrix and clears all the elements.

This function resizes a sparse matrix. The resized sparse matrix will become empty, even if it contained nonzero entries.

Arguments:

A: The initialized sparse matrix to resize.

nrow: The new number of rows.

ncol: The new number of columns.

nzmax: The new maximum number of elements.

Returns:

Error code.

Time complexity: O(nzmax), the maximum number of non-zero elements.

igraph_sparsemat_sort — Sorts all elements of a sparse matrix by row and column indices.

This function will sort the elements of a sparse matrix such that iterating over the entries will return them sorted by column indices; elements in the same column are then sorted by row indices.

Arguments:

A: A sparse matrix in either triplet or compressed form.

sorted: An uninitialized sparse matrix; the result will be returned here. The result will be in triplet form if the input was in triplet form, otherwise it will be in compressed form. Note that sorting is more efficient when the matrix is already in compressed form.

Returns:

Error code.

Time complexity: TODO

Operations on sparse matrix iterators

igraph_sparsemat_iterator_init — Initialize a sparse matrix iterator.

```
igraph_error_t igraph_sparsemat_iterator_init(
    igraph_sparsemat_iterator_t *it, const igraph_sparsemat_t *sparsemat
);
```

Arguments:

it: A pointer to an uninitialized sparse matrix iterator.

sparsemat: Pointer to the sparse matrix.

Returns:

Error code. This will always return IGRAPH_SUCCESS

Time complexity: O(n), the number of columns of the sparse matrix.

igraph_sparsemat_iterator_reset — Reset a sparse matrix iterator to the first element.

```
igraph_error_t igraph_sparsemat_iterator_reset(igraph_sparsemat_iterator_t *it)
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

Error code. This will always return IGRAPH_SUCCESS

Time complexity: O(n), the number of columns of the sparse matrix.

igraph_sparsemat_iterator_end — Query if the iterator is past the last element.

```
igraph_bool_t
igraph_sparsemat_iterator_end(const igraph_sparsemat_iterator_t *it);
```

it: A pointer to the sparse matrix iterator.

Returns:

true if the iterator is past the last element, false if it points to an element in a sparse matrix.

Time complexity: O(1).

igraph_sparsemat_iterator_row — Return the row of the iterator.

igraph_integer_t igraph_sparsemat_iterator_row(const igraph_sparsemat_iterator_

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The row of the element at the current iterator position.

Time complexity: O(1).

igraph_sparsemat_iterator_col — Return the column of the iterator.

igraph_integer_t igraph_sparsemat_iterator_col(const igraph_sparsemat_iterator_

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The column of the element at the current iterator position.

Time complexity: O(1).

igraph_sparsemat_iterator_get — Return the element at the current iterator position.

```
igraph_real_t
igraph_sparsemat_iterator_get(const igraph_sparsemat_iterator_t *it);
```

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The value of the element at the current iterator position.

Time complexity: O(1).

igraph_sparsemat_iterator_next — Let a sparse matrix iterator go to the next element.

igraph_integer_t igraph_sparsemat_iterator_next(igraph_sparsemat_iterator_t *it

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The position of the iterator in the element vector.

Time complexity: O(n), the number of columns of the sparse matrix.

igraph_sparsemat_iterator_idx — Returns the element vector index of a sparse matrix iterator.

igraph_integer_t igraph_sparsemat_iterator_idx(const igraph_sparsemat_iterator_

Arguments:

it: A pointer to the sparse matrix iterator.

Returns:

The position of the iterator in the element vector.

Time complexity: O(1).

Operations that change the internal representation

igraph_sparsemat_compress — Converts a sparse matrix to column-compressed format.

Converts a sparse matrix from triplet format to column-compressed format. Almost all sparse matrix operations require that the matrix is in column-compressed format.

Arguments:

A: The input matrix, it must be in triplet format.

res: Pointer to an uninitialized sparse matrix object, the compressed version of A is stored here.

Returns:

Error code.

Time complexity: O(nz) where nz is the number of non-zero elements.

igraph_sparsemat_dupl — Removes duplicate elements from a sparse matrix.

```
igraph_error_t igraph_sparsemat_dupl(igraph_sparsemat_t *A);
```

It is possible that a column-compressed sparse matrix stores a single matrix entry in multiple pieces. The entry is then the sum of all its pieces. (Some functions create matrices like this.) This function eliminates the multiple pieces.

Arguments:

A: The input matrix, in column-compressed format.

Returns:

Error code.

Time complexity: TODO.

Decompositions and solving linear systems

igraph_sparsemat_symblu — Symbolic LU decomposition.

LU decomposition of sparse matrices involves two steps, the first is calling this function, and then igraph_sparsemat_lu().

Arguments:

order: The ordering to use: 0 means natural ordering, 1 means minimum degree ordering of A +A', 2 is minimum degree ordering of A'A after removing the dense rows from A, and 3 is the minimum degree ordering of A'A.

A: The input matrix, in column-compressed format.

dis: The result of the symbolic analysis is stored here. Once not needed anymore, it must be destroyed by calling igraph_sparsemat_symbolic_destroy().

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_symbqr — Symbolic QR decomposition.

QR decomposition of sparse matrices involves two steps, the first is calling this function, and then $igraph_sparsemat_qr()$.

Arguments:

order: The ordering to use: 0 means natural ordering, 1 means minimum degree ordering of A +A', 2 is minimum degree ordering of A'A after removing the dense rows from A, and 3

is the minimum degree ordering of A'A.

A: The input matrix, in column-compressed format.

dis: The result of the symbolic analysis is stored here. Once not needed anymore, it must be destroyed by calling igraph_sparsemat_symbolic_destroy().

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_lsolve — Solves a lower-triangular linear system.

Solve the Lx=b linear equation system, where the L coefficient matrix is square and lower-triangular, with a zero-free diagonal.

Arguments:

L: The input matrix, in column-compressed format.

b: The right hand side of the linear system.

res: An initialized vector, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_ltsolve — Solves an upper-triangular linear system.

Solve the L'x=b linear equation system, where the L matrix is square and lower-triangular, with a zero-free diagonal.

Arguments:

L: The input matrix, in column-compressed format.

b: The right hand side of the linear system.

res: An initialized vector, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_usolve — Solves an upper-triangular linear system.

Solves the Ux=b upper triangular system.

Arguments:

U: The input matrix, in column-compressed format.

b: The right hand side of the linear system.

res: An initialized vector, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_utsolve — Solves a lower-triangular linear system.

This is the same as igraph_sparsemat_usolve(), but U'x=b is solved, where the apostrophe denotes the transpose.

U: The input matrix, in column-compressed format.

b: The right hand side of the linear system.

res: An initialized vector, the result is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_cholsol — Solves a symmetric linear system via Cholesky decomposition.

Solve Ax=b, where A is a symmetric positive definite matrix.

Arguments:

A: The input matrix, in column-compressed format.

v: The right hand side.

res: An initialized vector, the result is stored here.

order: An integer giving the ordering method to use for the factorization. Zero is the natural ordering; if it is one, then the fill-reducing minimum-degree ordering of A+A' is used.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_lusol — Solves a linear system via LU decomposition.

Solve Ax=b, via LU factorization of A.

Arguments:

A: The input matrix, in column-compressed format.

b: The right hand side of the equation.

res: An initialized vector, the result is stored here.

order: The ordering method to use, zero means the natural ordering, one means the fill-reducing

minimum-degree ordering of A+A', two means the ordering of A'*A, after removing the

dense rows from A. Three means the ordering of A'*A.

to1: Real number, the tolerance limit to use for the numeric LU factorization.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_lu — LU decomposition of a sparse matrix.

Performs numeric sparse LU decomposition of a matrix.

Arguments:

A: The input matrix, in column-compressed format.

dis: The symbolic analysis for LU decomposition, coming from a call to the igraph_sparse-mat_symblu() function.

din: The numeric decomposition, the result is stored here. It can be used to solve linear systems with changing right hand side vectors, by calling igraph_sparsemat_luresol(). Once not needed any more, it must be destroyed by calling igraph_sparsemat_symbolic_destroy() on it.

to1: The tolerance for the numeric LU decomposition.

Returns:

Error code.

Time complexity: TODO.

$igraph_sparsemat_qr$ — QR decomposition of a sparse matrix.

Numeric QR decomposition of a sparse matrix.

Arguments:

A: The input matrix, in column-compressed format.

- dis: The result of the symbolic QR analysis, from the function igraph_sparsemat_symbqr().
- din: The result of the decomposition is stored here, it can be used to solve many linear systems with the same coefficient matrix and changing right hand sides, using the igraph_sparse-mat_qrresol() function. Once not needed any more, one should call igraph_sparsemat_numeric_destroy() on it to free the allocated memory.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_luresol — Solves a linear system using a precomputed LU decomposition.

Uses the LU decomposition of a matrix to solve linear systems.

Arguments:

- dis: The symbolic analysis of the coefficient matrix, the result of igraph_sparsemat_symblu().
- din: The LU decomposition, the result of a call to igraph_sparsemat_lu().
- b: A vector that defines the right hand side of the linear equation system.
- res: An initialized vector, the solution of the linear system is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_qrresol — Solves a linear system using a precomputed QR decomposition.

Solves a linear system using a QR decomposition of its coefficient matrix.

Arguments:

dis: Symbolic analysis of the coefficient matrix, the result of igraph_sparsemat_symbol.

din: The QR decomposition of the coefficient matrix, the result of igraph_sparsemat_qr().

b: Vector, giving the right hand side of the linear equation system.

res: An initialized vector, the solution is stored here. It is resized as needed.

Returns:

Error code.

Time complexity: TODO.

Time complexity: O(1).

igraph_sparsemat_symbolic_destroy — Deallocates memory after a symbolic decomposition.

```
void igraph_sparsemat_symbolic_destroy(igraph_sparsemat_symbolic_t *dis);
Frees the memory allocated by igraph_sparsemat_symbqr() or igraph_sparse-
mat_symblu().

Arguments:
dis: The symbolic analysis.
```

igraph_sparsemat_numeric_destroy — Deallocates memory after a numeric decomposition.

```
void igraph_sparsemat_numeric_destroy(igraph_sparsemat_numeric_t *din);
Frees the memoty allocated by igraph_sparsemat_qr() or igraph_sparsemat_lu().
Arguments:
din: The LU or QR decomposition.
Time complexity: O(1).
```

Eigenvalues and eigenvectors

igraph_sparsemat_arpack_rssolve — Eigenvalues and eigenvectors of a symmetric sparse matrix via ARPACK.

The: input matrix, must be column-compressed.

options: It is passed to igraph_arpack_rssolve(). Supply NULL here to use the

defaults. See igraph_arpack_options_t for the details. If mode is 1, then ARPACK uses regular mode, if mode is 3, then shift and invert mode is used and

the sigma structure member defines the shift.

storage: Storage for ARPACK. See igraph_arpack_rssolve() and

igraph_arpack_storage_t for details.

values: An initialized vector or a null pointer, the eigenvalues are stored here.

vectors: An initialised matrix, or a null pointer, the eigenvectors are stored here, in the

columns.

solvemethod: The method to solve the linear system, if mode is 3, i.e. the shift and invert mode

is used. Possible values:

IGRAPH_SPARSE- The linear system is solved using LU de-

MAT_SOLVE_LU composition.

IGRAPH_SPARSE- The linear system is solved using QR de-

MAT_SOLVE_QR composition.

Returns:

Error code.

Time complexity: TODO.

igraph_sparsemat_arpack_rnsolve — Eigenvalues and eigenvectors of a nonsymmetric sparse matrix via ARPACK.

Eigenvalues and/or eigenvectors of a nonsymmetric sparse matrix.

Arguments:

A: The input matrix, in column-compressed mode.

 $options: \quad ARPACK\ options, it\ is\ passed\ to\ \verb"igraph_arpack_rnsolve" (\).\ Supply\ \verb"NULL" here$

to use the defaults. See also igraph_arpack_options_t for details.

storage: Storage for ARPACK, this is passed to igraph_arpack_rnsolve(). See

igraph_arpack_storage_t for details.

values: An initialized matrix, or a null pointer. If not a null pointer, then the eigenvalues are

stored here, the first column is the real part, the second column is the imaginary part.

vectors: An initialized matrix, or a null pointer. If not a null pointer, then the eigenvectors are

stored here, please see igraph_arpack_rnsolve() for the format.

Returns:

Error code.

Time complexity: TODO.

Conversion to other data types

igraph_sparsemat — Creates an igraph graph from a sparse matrix.

One edge is created for each non-zero entry in the matrix. If you have a symmetric matrix, and want to create an undirected graph, then delete the entries in the upper diagonal first, or call <code>igraph_simplify()</code> on the result graph to eliminate the multiple edges.

Arguments:

graph: Pointer to an uninitialized igraph_t object, the graphs is stored here.

A: The input matrix, in triplet or column-compressed format.

directed: Boolean scalar, whether to create a directed graph.

Returns:

Error code.

Time complexity: TODO.

igraph_matrix_as_sparsemat — Converts a dense matrix to a sparse matrix.

Arguments:

res: An uninitialized sparse matrix, the result is stored here.

mat: The dense input matrix.

to1: Real scalar, the tolerance. Values closer than to1 to zero are considered as zero, and will not be included in the sparse matrix.

Returns:

Error code.

Time complexity: O(mn), the number of elements in the dense matrix.

igraph_sparsemat_as_matrix — Converts a sparse matrix to a dense matrix.

Arguments:

res: Pointer to an initialized matrix, the result is stored here. It will be resized to the required

size.

spmat: The input sparse matrix, in triplet or column-compressed format.

Returns:

Error code.

Time complexity: O(mn), the number of elements in the dense matrix.

Writing to a file, or to the screen

igraph_sparsemat_print — Prints a sparse matrix to a file.

Only the non-zero entries are printed. This function serves more as a debugging utility, as currently there is no function that could read back the printed matrix from the file.

Arguments:

A: The input matrix, triplet or column-compressed format.

outstream: The stream to print it to.

Returns:

Error code.

Time complexity: O(nz) for triplet matrices, O(n+nz) for column-compressed matrices. nz is the number of non-zero elements, n is the number columns in the matrix.

Deprecated functions

igraph_sparsemat_copy — Copies a sparse matrix (deprecated alias).

```
igraph_error_t igraph_sparsemat_copy(
    igraph_sparsemat_t *to, const igraph_sparsemat_t *from
);
```

Warning

Deprecated since version 0.10. Please do not use this function in new code; use $igraph_s-parsemat_init_copy()$ instead.

igraph_sparsemat_diag — Creates a sparse diagonal matrix (deprecated alias).

```
igraph_error_t igraph_sparsemat_diag(
    igraph_sparsemat_t *A, igraph_integer_t nzmax, const igraph_vector_t *value
    igraph_bool_t compress
);
```

Warning

Deprecated since version 0.10. Please do not use this function in new code; use igraph_s-parsemat_init_diag() instead.

igraph_sparsemat_eye — Creates a sparse identity matrix (deprecated alias).

```
igraph_error_t igraph_sparsemat_eye(
    igraph_sparsemat_t *A, igraph_integer_t n, igraph_integer_t nzmax,
    igraph_real_t value, igraph_bool_t compress
);
```

Warning

Deprecated since version 0.10. Please do not use this function in new code; use igraph_s-parsemat_init_eye() instead.

Stacks

igraph_stack_init — Initializes a stack.

```
igraph_error_t igraph_stack_init(igraph_stack_t* s, igraph_integer_t size);
```

The initialized stack is always empty.

Arguments:

s: Pointer to an uninitialized stack.

size: The number of elements to allocate memory for.

Returns:

Error code.

Time complexity: O(size).

igraph_stack_destroy — Destroys a stack object.

```
void igraph_stack_destroy(igraph_stack_t* s);
```

Deallocate the memory used for a stack. It is possible to reinitialize a destroyed stack again by igraph_stack_init().

Arguments:

s: The stack to destroy.

Time complexity: O(1).

igraph_stack_reserve — Reserve memory.

```
igraph_error_t igraph_stack_reserve(igraph_stack_t* s, igraph_integer_t capacit
```

Reserve memory for future use. The actual size of the stack is unchanged.

Arguments:

s: The stack object.

size: The number of elements to reserve memory for. If it is not bigger than the current size then nothing happens.

Returns:

Error code.

Time complexity: should be around O(n), the new allocated size of the stack.

igraph_stack_empty — Decides whether a stack object is empty.

```
igraph_bool_t igraph_stack_empty(igraph_stack_t* s);
```

Arguments:

s: The stack object.

Returns:

Boolean, true if the stack is empty, false otherwise.

Time complexity: O(1).

igraph_stack_size — Returns the number of elements in a stack.

```
igraph_integer_t igraph_stack_size(const igraph_stack_t* s);
```

Arguments:

s: The stack object.

Returns:

The number of elements in the stack.

Time complexity: O(1).

igraph_stack_clear — Removes all elements from a stack.

```
void igraph_stack_clear(igraph_stack_t* s);
```

Arguments:

s: The stack object.

Time complexity: O(1).

igraph_stack_push — Places an element on the top of a stack.

```
igraph_error_t igraph_stack_push(igraph_stack_t* s, igraph_real_t elem);
```

The capacity of the stack is increased, if needed.

Arguments:

s: The stack object.

elem: The element to push.

Returns:

Error code.

Time complexity: O(1) is no reallocation is needed, O(n) otherwise, but it is ensured that n push operations are performed in O(n) time.

igraph_stack_pop — Removes and returns an element from the top of a stack.

```
igraph_real_t igraph_stack_pop(igraph_stack_t* s);
```

The stack must contain at least one element, call igraph_stack_empty() to make sure of this.

Arguments:

s: The stack object.

Returns:

The removed top element.

Time complexity: O(1).

igraph_stack_top — Query top element.

```
igraph_real_t igraph_stack_top(const igraph_stack_t* s);
```

Returns the top element of the stack, without removing it. The stack must be non-empty.

Arguments:

s: The stack.

Returns:

The top element.

Time complexity: O(1).

Double-ended queues

This is the classic data type of the double ended queue. Most of the time it is used if a First-In-First-Out (FIFO) behavior is needed. See the operations below.

Example 7.10. File examples/simple/dqueue.c

igraph_dqueue_init — Initialize a double ended queue (deque).

igraph_error_t igraph_dqueue_init(igraph_dqueue_t* q, igraph_integer_t capacity

The queue will be always empty.

Arguments:

q: Pointer to an uninitialized deque.

capacity: How many elements to allocate memory for.

Returns:

Error code.

Time complexity: O(capacity).

igraph_dqueue_destroy — Destroy a double ended queue.

```
void igraph_dqueue_destroy(igraph_dqueue_t* q);
```

Arguments:

q: The queue to destroy

Time complexity: O(1).

igraph_dqueue_empty — Decide whether the queue is empty.

```
igraph_bool_t igraph_dqueue_empty(const igraph_dqueue_t* q);
```

Arguments:

q: The queue.

Returns:

Boolean, true if q contains at least one element, false otherwise.

Time complexity: O(1).

igraph_dqueue_full — Check whether the queue is full.

```
igraph_bool_t igraph_dqueue_full(igraph_dqueue_t* q);
```

If a queue is full the next igraph_dqueue_push() operation will allocate more memory.

Arguments:

q: The queue.

Returns:

```
true if q is full, false otherwise.
```

Time complecity: O(1).

igraph_dqueue_clear — Remove all elements from the queue.

```
void igraph_dqueue_clear(igraph_dqueue_t* q);
```

Arguments:

q: The queue.

Time complexity: O(1).

igraph_dqueue_size — Number of elements in the queue.

```
igraph_integer_t igraph_dqueue_size(const igraph_dqueue_t* q);
```

Arguments:

q: The queue.

Returns:

Integer, the number of elements currently in the queue.

Time complexity: O(1).

igraph_dqueue_head — Head of the queue.

```
igraph_real_t igraph_dqueue_head(const igraph_dqueue_t* q);
```

The queue must contain at least one element.

Arguments:

g: The queue.

Returns:

The first element in the queue.

Time complexity: O(1).

igraph_dqueue_back — Tail of the queue.

igraph_real_t igraph_dqueue_back(const igraph_dqueue_t* q);

The queue must contain at least one element.

Arguments:

q: The queue.

Returns:

The last element in the queue.

Time complexity: O(1).

igraph_dqueue_pop — Remove the head.

```
igraph_real_t igraph_dqueue_pop(igraph_dqueue_t* q);
```

Removes and returns the first element in the queue. The queue must be non-empty.

Arguments:

q: The input queue.

Returns:

The first element in the queue.

Time complexity: O(1).

igraph_dqueue_pop_back — Removes the tail.

```
igraph_real_t igraph_dqueue_pop_back(igraph_dqueue_t* q);
```

Removes and returns the last element in the queue. The queue must be non-empty.

Arguments:

g: The queue.

Returns:

The last element in the queue.

Time complexity: O(1).

igraph_dqueue_push — Appends an element.

igraph_error_t igraph_dqueue_push(igraph_dqueue_t* q, igraph_real_t elem);

Append an element to the end of the queue.

Arguments:

q: The queue.

elem: The element to append.

Returns:

Error code.

Time complexity: O(1) if no memory allocation is needed, O(n), the number of elements in the queue otherwise. But note that by allocating always twice as much memory as the current size of the queue we ensure that n push operations can always be done in at most O(n) time. (Assuming memory allocation is at most linear.)

Maximum and minimum heaps

igraph_heap_init — Initializes an empty heap object.

igraph_error_t igraph_heap_init(igraph_heap_t* h, igraph_integer_t capacity);

Creates an *empty* heap, and also allocates memory for some elements.

Arguments:

h: Pointer to an uninitialized heap object.

capacity: Number of elements to allocate memory for.

Returns:

Error code.

Time complexity: O(alloc_size), assuming memory allocation is a linear operation.

igraph_heap_init_array — Build a heap from an array.

Initializes a heap object from an array. The heap is also built of course (constructor).

igraph_error_t igraph_heap_init_array(igraph_heap_t *h, igraph_real_t* data, ig

Arguments:

h: Pointer to an uninitialized heap object.

data: Pointer to an array of base data type.

len: The length of the array at data.

Returns:

Error code.

Time complexity: O(n), the number of elements in the heap.

igraph_heap_destroy — Destroys an initialized heap object.

```
void igraph_heap_destroy(igraph_heap_t* h);
```

Arguments:

h: The heap object.

Time complexity: O(1).

igraph_heap_clear — Removes all elements from a heap.

```
void igraph_heap_clear(igraph_heap_t* h);
```

This function simply sets the size of the heap to zero, it does not free any allocated memory. For that you have to call <code>igraph_heap_destroy()</code>.

Arguments:

h: The heap object.

Time complexity: O(1).

igraph_heap_empty — Decides whether a heap object is empty.

```
igraph_bool_t igraph_heap_empty(const igraph_heap_t* h);
```

Arguments:

h: The heap object.

Returns:

true if the heap is empty, false otherwise.

TIme complexity: O(1).

igraph_heap_push — Add an element.

igraph_error_t igraph_heap_push(igraph_heap_t* h, igraph_real_t elem);

Adds an element to the heap.

Arguments:

h: The heap object.

elem: The element to add.

Returns:

Error code.

Time complexity: $O(\log n)$, n is the number of elements in the heap if no reallocation is needed, O(n) otherwise. It is ensured that n push operations are performed in $O(n \log n)$ time.

igraph_heap_top — Top element.

```
igraph_real_t igraph_heap_top(const igraph_heap_t* h);
```

For maximum heaps this is the largest, for minimum heaps the smallest element of the heap.

Arguments:

h: The heap object.

Returns:

The top element.

Time complexity: O(1).

igraph_heap_delete_top — Removes and returns the top element.

```
igraph_real_t igraph_heap_delete_top(igraph_heap_t* h);
```

Removes and returns the top element of the heap. For maximum heaps this is the largest, for minimum heaps the smallest element.

Arguments:

h: The heap object.

Returns:

The top element.

Time complexity: O(log n), n is the number of elements in the heap.

igraph_heap_size — Number of elements in the heap.

igraph_integer_t igraph_heap_size(const igraph_heap_t* h);

Gives the number of elements in a heap.

Arguments:

h: The heap object.

Returns:

The number of elements in the heap.

Time complexity: O(1).

igraph_heap_reserve — Reserves memory for a heap.

```
igraph_error_t igraph_heap_reserve(igraph_heap_t* h, igraph_integer_t capacity)
```

Allocates memory for future use. The size of the heap is unchanged. If the heap is larger than the capacity parameter then nothing happens.

Arguments:

h: The heap object.

capacity: The number of elements to allocate memory for.

Returns:

Error code.

Time complexity: O(capacity) if capacity is larger than the current number of elements. O(1) otherwise.

String vectors

The igraph_strvector_t type is a vector of null-terminated strings. It is used internally for storing graph attribute names as well as string attributes in the C attribute handler.

This container automatically manages the memory of its elements. The strings within an igraph_strvector_t should be considered constant, and not modified directly. Functions that add new elements always make copies of the string passed to them.

Example 7.11. File examples/simple/igraph_strvector.c

igraph_strvector_init — Initializes a string vector.

igraph_error_t igraph_strvector_init(igraph_strvector_t *sv, igraph_integer_t s

Reserves memory for the string vector, a string vector must be first initialized before calling other functions on it. All elements of the string vector are set to the empty string.

Arguments:

sv: Pointer to an initialized string vector.

len: The (initial) length of the string vector.

Returns:

Error code.

Time complexity: O(1en).

igraph_strvector_init_copy — Initialization by copying.

Initializes a string vector by copying another string vector.

Arguments:

to: Pointer to an uninitialized string vector.

from: The other string vector, to be copied.

Returns:

Error code.

Time complexity: O(l), the total length of the strings in from.

igraph_strvector_destroy — Frees the memory allocated for the string vector.

```
void igraph_strvector_destroy(igraph_strvector_t *sv);
```

Destroy a string vector. It may be reinitialized with igraph_strvector_init() later.

Arguments:

sv: The string vector.

Time complexity: O(l), the total length of the strings, maybe less depending on the memory manager.

STR — Indexing string vectors.

```
#define STR(sv,i)
```

This is a macro that allows to query the elements of a string vector, just like <code>igraph_strvector_get()</code>, but without the overhead of a function call. Note this macro cannot be used to set an element. Use <code>igraph_strvector_set()</code> to set an element instead.

Arguments:

sv: The string vector

i: The the index of the element.

Returns:

The element at position i.

Time complexity: O(1).

igraph_strvector_get — Retrieves an element of the string vector.

```
const char* igraph_strvector_get(const igraph_strvector_t *sv, igraph_integer_t
```

Query an element of a string vector. See also the STR macro for an easier way.

Arguments:

sv: The input string vector.

idx: The index of the element to query.

Time complexity: O(1).

igraph_strvector_set — Takes elements at given positions from a string vector.

Arguments:

sv: The string vector.

newv: An initialized string vector, it will be resized as needed.

idx: An integer vector of indices to take from sv.

Returns:

Error code.

igraph_strvector_set_len — Sets an element of the string vector given a buffer and its size.

This is almost the same as igraph_strvector_set, but the new value is not a zero terminated string, but its length is given.

Arguments:

sv: The string vector.

idx: The position to set.

value: The new value.

len: The length of the new value.

Returns:

Error code.

Time complexity: O(l), the length of the new string. Maybe more, depending on the memory management, if reallocation is needed.

igraph_strvector_push_back — Adds an element to the back of a string vector.

igraph_error_t igraph_strvector_push_back(igraph_strvector_t *sv, const char *v.

Arguments:

sv: The string vector.

value: The string to add; it will be copied.

Returns:

Error code.

Time complexity: O(n+1), n is the total number of strings, l is the length of the new string.

igraph_strvector_push_back_len — Adds a string of the given length to the back of a string vector.

igraph_error_t igraph_strvector_push_back_len(

```
igraph_strvector_t *sv,
const char *value, igraph_integer_t len);
```

Arguments:

sv: The string vector.

value: The start of the string to add. At most len characters will be copied.

len: The length of the string.

Returns:

Error code.

Time complexity: O(n+l), n is the total number of strings, l is the length of the new string.

igraph_strvector_remove — Removes a single element from a string vector.

```
void igraph_strvector_remove(igraph_strvector_t *sv, igraph_integer_t elem);
```

The string will be one shorter.

Arguments:

sv: The string vector.

elem: The index of the element to remove.

Time complexity: O(n), the length of the string.

igraph_strvector_remove_section — Removes a section from a string vector.

This function removes the range [from, to) from the string vector.

Arguments:

sv: The string vector.

from: The position of the first element to remove.

to: The position of the first element *not* to remove.

igraph_strvector_append — Concatenates two string vectors.

Appends the contents of the *from* vector to the *to* vector. If the *from* vector is no longer needed after this operation, use <code>igraph_strvector_merge()</code> for better performance.

Arguments:

to: The first string vector, the result is stored here.

from: The second string vector, it is kept unchanged.

Returns:

Error code.

See also:

```
igraph_strvector_merge()
```

Time complexity: O(n+12), n is the number of strings in the new string vector, 12 is the total length of strings in the from string vector.

igraph_strvector_merge — Moves the contents of a string vector to the end of another.

```
igraph_error_t igraph_strvector_merge(igraph_strvector_t *to, igraph_strvector_
```

Transfers the contents of the *from* vector to the end of *to*, clearing *from* in the process. If this operation fails, both vectors are left intact. This function does not copy or reallocate individual strings, therefore it performs better than igraph_strvector_append().

Arguments:

to: The target vector. The contents of from will be appended to it.

from: The source vector. It will be cleared.

Returns:

Error code.

See also:

```
igraph_strvector_append()
```

Time complexity: O(12) if to has sufficient capacity, O(2*11+12) otherwise, where 11 and 12 are the lengths of to and \from respectively.

igraph_strvector_clear — Removes all elements from a string vector.

```
void igraph_strvector_clear(igraph_strvector_t *sv);
```

After this operation the string vector will be empty.

Arguments:

sv: The string vector.

Time complexity: O(1), the total length of strings, maybe less, depending on the memory manager.

igraph_strvector_resize — Resizes a string vector.

```
igraph_error_t igraph_strvector_resize(igraph_strvector_t *sv, igraph_integer_t
```

If the new size is bigger then empty strings are added, if it is smaller then the unneeded elements are removed.

Arguments:

sv: The string vector.

newsize: The new size.

Returns:

Error code.

Time complexity: O(n), the number of strings if the vector is made bigger, O(l), the total length of the deleted strings if it is made smaller, maybe less, depending on memory management.

igraph_strvector_reserve — Reserves memory for a string vector.

```
igraph_error_t igraph_strvector_reserve(igraph_strvector_t *sv, igraph_integer_
```

igraph string vectors are flexible, they can grow and shrink. Growing however occasionally needs the data in the vector to be copied. In order to avoid this, you can call this function to reserve space for future growth of the vector.

Note that this function does *not* change the size of the string vector. Let us see a small example to clarify things: if you reserve space for 100 strings and the size of your vector was (and still is) 60, then you can surely add additional 40 strings to your vector before it will be copied.

Arguments:

sv: The string vector object.

capacity: The new allocated size of the string vector.

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, should be around O(n), n is the new allocated size of the vector.

igraph_strvector_resize_min — Deallocates the unused memory of a string vector.

```
void igraph_strvector_resize_min(igraph_strvector_t *sv);
```

This function attempts to deallocate the unused reserved storage of a string vector. If it succeeds, igraph_strvector_size() and igraph_strvector_capacity() will be the same. The data in the string vector is always preserved, even if deallocation is not successful.

Arguments:

sv: The string vector.

Time complexity: Operating system dependent, at most O(n).

igraph_strvector_size — Returns the size of a string vector.

```
igraph_integer_t igraph_strvector_size(const igraph_strvector_t *sv);
```

Arguments:

sv: The string vector.

Returns:

The length of the string vector.

Time complexity: O(1).

igraph_strvector_capacity — Returns the capacity of a string vector.

```
igraph_integer_t igraph_strvector_capacity(const igraph_strvector_t *sv);
```

Arguments:

sv: The string vector.

Returns:

The capacity of the string vector.

Time complexity: O(1).

Deprecated functions

igraph_strvector_copy — Initialization by copying (deprecated alias).

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_strvector_init_copy() instead.

igraph_strvector_add — Adds an element to the back of a string vector (deprecated alias).

```
igraph_error_t igraph_strvector_add(igraph_strvector_t *sv, const char *value);
```

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use $igraph_strvector_push_back()$ instead.

igraph_strvector_set2 — Sets an element of the string vector given a buffer and its size (deprecated alias).

```
igraph_error_t igraph_strvector_set2(
    igraph_strvector_t *sv, igraph_integer_t idx, const char *value, size_t len
);
```

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_strvector_set_len() instead.

Lists of vectors, matrices and graphs

About igraph_vector_list_t objects

The igraph_vector_list_t data type is essentially a list of igraph_vector_t objects with automatic memory management. It is something similar to (but much simpler than) the vector template in the C++ standard library where the elements are vectors themselves.

There are multiple variants of igraph_vector_list_t; the basic variant stores vectors of doubles (i.e. each item is an igraph_vector_t), but there is also igraph_vector_int_list_t for integers (where each item is an igraph_vector_int_t), igraph_matrix_list_t for matrices of doubles and so on. The following list summarizes the variants that are currently available in the library:

- igraph_vector_list_t for lists of vectors of floating-point numbers (igraph_vector_t)
- igraph_vector_int_list_t for lists of integer vectors (igraph_vector_int_t)
- igraph_matrix_list_t for lists of matrices of floating-point numbers (igraph_matrix_t)
- igraph_graph_list_t for lists of graphs (igraph_t)

Lists of vectors are used in **igraph** in many cases, e.g., when returning lists of paths, cliques or vertex sets. Functions that expect or return a list of numeric vectors typically use igraph_vector_list_t or igraph_vector_int_list_t to achieve this. Lists of integer vectors are used when the vectors in the list are supposed to hold vertex or edge identifiers, while lists of floating-point vectors are used when the vectors are expected to hold fractional numbers or infinities.

The elements in an igraph_vector_list_t object and its variants are indexed from zero, we follow the usual C convention here.

Almost all of the functions described below for igraph_vector_list_t also exist for all the other vector list variants. These variants are not documented separately; you can simply replace vector_list with, say, vector_int_list if you need a function for another variant. For instance, to initialize a list of integer vectors, you need to use igraph_vector_int_list_init() and not igraph_vector_list_init().

Before diving into a detailed description of the functions related to lists of vectors, we must also talk about the *ownership* rules of these objects. The most important rule is that the vectors in the list are owned by the list itself, meaning that the user is *not* responsible for allocating memory for the vectors or for freeing the memory associated to the vectors. It is the responsibility of the list to allocate and initialize the vectors when new items are created in the list, and it is also the responsibility of the list to destroy the items when they are removed from the list without passing on their ownership to the user. As a consequence, the list may not contain "uninitialized" or "null" items; each item is initialized when it comes to existence. If you create a list containing one million vectors, you are not only allocating memory for one million igraph_vector_t object but you are also initializing one million vectors. Also, if you have a list containing one million vectors and you clear the list by calling igraph_vector_list_clear(), the list will implicitly destroy these lists, and any pointers that you may hold to the items become invalid at once.

Speaking about pointers, the typical way of working with vectors in a list is to obtain a pointer to one of the items via the <code>igraph_vector_list_get_ptr()</code> method and then passing this pointer onwards to functions that manipulate <code>igraph_vector_t</code> objects. However, note that the pointers are <code>ephemeral</code> in the sense that they may be invalidated any time when the list is modified because a modification may involve the re-allocation of the internal storage of the list if more space is needed, and the pointers that you obtained will not follow the reallocation. This limitation does not appear often in real-world usage of <code>igraph_vector_list_t</code> and in general, the advantages of the automatic memory management outweigh this limitation.

Constructors and destructors

igraph_vector_list_t objects have to be initialized before using them, this is analogous to calling a constructor on them. igraph_vector_list_init() is the basic constructor; it creates a list of the given length and also initializes each vector in the newly created list to zero length.

If an igraph_vector_list_t object is not needed any more, it should be destroyed to free its allocated memory by calling the igraph_vector_list_t destructor, igraph_vector_list_destroy(). Calling the destructor also destroys all the vectors inside the vector list due to the own-

ership rules. If you want to keep a few of the vectors in the vector list, you need to copy them with igraph_vector_init_copy() or igraph_vector_update(), or you need to remove them from the list and take ownership by calling igraph_vector_list_pop_back(), igraph_vector_list_remove() or igraph_vector_list_remove_fast().

igraph_vector_list_init — Initializes a list of vectors (constructor).

```
igraph_error_t igraph_vector_list_init(igraph_vector_list_t* v, igraph_integer_
```

This function constructs a list of vectors of the given size, and initializes each vector in the newly created list to become an empty vector.

Vector objects initialized by this function are *owned* by the list, and they will be destroyed automatically when the list is destroyed with <code>igraph_vector_list_destroy()</code>.

Arguments:

v: Pointer to a not yet initialized list of vectors.

size: The size of the list.

Returns:

error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, the amount of "time" required to allocate O(n) elements and initialize the corresponding vectors; n is the number of elements.

igraph_vector_list_destroy — Destroys a list of vectors object.

```
void igraph_vector_list_destroy(igraph_vector_list_t* v);
```

All lists initialized by igraph_vector_list_init() should be properly destroyed by this function. A destroyed list of vectors needs to be reinitialized by igraph_vector_list_init() if you want to use it again.

Vectors that are in the list when it is destroyed are also destroyed implicitly.

Arguments:

v: Pointer to the (previously initialized) list object to destroy.

Time complexity: operating system dependent.

Accessing elements

Elements of a vector list may be accessed with the <code>igraph_vector_list_get_ptr()</code> function. The function returns a *pointer* to the vector with a given index inside the list, and you may then pass this pointer onwards to other functions that can query or manipulate vectors. The pointer itself is guaranteed to stay valid as long as the list itself is not modified; however, *any* modification to the

list will invalidate the pointer, even modifications that are seemingly unrelated to the vector that the pointer points to (such as adding a new vector at the end of the list). This is because the list data structure may be forced to re-allocate its internal storage if a new element does not fit into the already allocated space, and there are no guarantees that the re-allocated block remains at the same memory location (typically it gets moved elsewhere).

Note that the standard VECTOR macro that works for ordinary vectors does not work for lists of vectors to access the i-th element (but of course you can use it to index into an existing vector that you retrieved from the vector list with <code>igraph_vector_list_get_ptr()</code>). This is because the macro notation would allow one to overwrite the vector in the list with another one without the list knowing about it, so the list would not be able to destroy the vector that was overwritten by a new one.

igraph_vector_list_tail_ptr() returns a pointer to the last vector in the list, or NULL if the list is empty. There is no igraph_vector_list_head_ptr(), however, as it is easy to write igraph_vector_list_get_ptr(v, 0) instead.

igraph_vector_list_get_ptr — Retrieve the address of a vector in the vector list.

```
igraph_vector_t* igraph_vector_list_get_ptr(const igraph_vector_list_t* v, igraph_vector_li
```

Arguments:

v: The list object.

pos: The position of the vector in the list. The position of the first vector is zero.

Returns:

A pointer to the vector. It remains valid as long as the underlying list of vectors is not modified.

Time complexity: O(1).

igraph_vector_list_tail_ptr — Retrieve the address of the last vector in the vector list.

```
igraph_vector_t* igraph_vector_list_tail_ptr(const igraph_vector_list_t *v);
```

Arguments:

v: The list object.

Returns:

A pointer to the last vector in the list, or NULL if the list is empty.

Time complexity: O(1).

igraph_vector_list_set — Sets the vector at the given index in the list.

```
void igraph_vector_list_set(igraph_vector_list_t* v, igraph_integer_t pos, igraph_integer_t
```

This function destroys the vector that is already at the given index *pos* in the list, and replaces it with the vector pointed to by *e*. The ownership of the vector pointed to by *e* is taken by the list so the user is not responsible for destroying *e* any more; it will be destroyed when the list itself is destroyed or if *e* gets removed from the list without passing on the ownership to somewhere else.

Arguments:

v: The list object.

pos: The index to modify in the list.

e: The vector to set in the list.

Time complexity: O(1).

igraph_vector_list_replace — Replaces the vector at the given index in the list with another one.

```
void igraph_vector_list_replace(igraph_vector_list_t* v, igraph_integer_t pos,
```

This function replaces the vector that is already at the given index *pos* in the list with the vector pointed to by *e*. The ownership of the vector pointed to by *e* is taken by the list so the user is not responsible for destroying *e* any more. At the same time, the ownership of the vector that *was* in the list at position *pos* will be transferred to the caller and *e* will be updated to point to it, so the caller becomes responsible for destroying it when it does not need the vector any more.

Arguments:

v: The list object.

pos: The index to modify in the list.

e: The vector to swap with the one already in the list.

Time complexity: O(1).

Vector properties

igraph_vector_list_empty — Decides whether the size of the list is zero.

```
igraph_bool_t igraph_vector_list_empty(const igraph_vector_list_t* v);
```

Arguments:

v: The list object.

Returns:

Non-zero number (true) if the size of the list is zero and zero (false) otherwise.

Time complexity: O(1).

igraph_vector_list_size — Returns the size (=length) of the vector.

```
igraph_integer_t igraph_vector_list_size(const igraph_vector_list_t* v);
```

Arguments:

v: The list object

Returns:

The size of the list.

Time complexity: O(1).

igraph_vector_list_capacity — Returns the allocated capacity of the list.

```
igraph_integer_t igraph_vector_list_capacity(const igraph_vector_list_t* v);
```

Note that this might be different from the size of the list (as queried by igraph_vector_list_size()), and specifies how many vectors the list can hold, without reallocation.

Arguments:

v: Pointer to the (previously initialized) list object to query.

Returns:

The allocated capacity.

See also:

```
igraph_vector_list_size().
```

Time complexity: O(1).

Resizing operations

igraph_vector_list_clear — Removes all elements from a
list of vectors.

```
void igraph_vector_list_clear(igraph_vector_list_t* v);
```

This function sets the size of the list to zero, and it also destroys all the vectors that were placed in the list before clearing it.

Arguments:

v: The list object.

Time complexity: O(n), n is the number of items being deleted.

igraph_vector_list_reserve — Reserves memory for a list.

igraph_error_t igraph_vector_list_reserve(igraph_vector_list_t* v, igraph_integ

igraph lists are flexible, they can grow and shrink. Growing however occasionally needs the data in the list to be copied. In order to avoid this, you can call this function to reserve space for future growth of the list.

Note that this function does *not* change the size of the list, neither does it initialize any new vectors. Let us see a small example to clarify things: if you reserve space for 100 elements and the size of your list was (and still is) 60, then you can surely add additional 40 new vectors to your list before it will be copied.

Arguments:

v: The list object.

capacity: The new allocated size of the list.

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory.

Time complexity: operating system dependent, should be around O(n), n is the new allocated size of the list.

igraph_vector_list_resize — Resize the list of vectors.

 $igraph_error_t \ igraph_vector_list_resize(igraph_vector_list_t* \ v, \ igraph_intege$

Note that this function does not free any memory, just sets the size of the list to the given one. It can on the other hand allocate more memory if the new size is larger than the previous one.

When the new size is larger than the current size, the newly added vectors in the list are initialized to empty vectors. When the new size is smaller than the current size, the vectors that were removed from the end of the list are destroyed automatically.

Arguments:

v: The list object

new size: The new size of the list.

Returns:

Error code, IGRAPH_ENOMEM if there is not enough memory. Note that this function *never* returns an error if the list is made smaller.

See also:

igraph_vector_list_reserve() for allocating memory for future extensions of a list.

Time complexity: O(m) if the new size is smaller (m is the number of items that were removed from the list), operating system dependent if the new size is larger. In the latter case it is usually around O(n), where n is the new size of the vector.

igraph_vector_list_push_back — Append an existing vector to the list, transferring ownership.

```
igraph_error_t igraph_vector_list_push_back(igraph_vector_list_t* v, igraph_vec
```

This function resizes the list to be one element longer, and sets the very last element in the list to the specified vector e. The list takes ownership of the vector so the user is not responsible for freeing e any more; the vector will be destroyed when the list itself is destroyed or if e gets removed from the list without passing on the ownership to somewhere else.

Arguments:

- v: The list object.
- e: Pointer to the vector to append to the list.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory.

Time complexity: operating system dependent. What is important is that a sequence of n subsequent calls to this function has time complexity O(n), even if there hadn't been any space reserved for the new elements by <code>igraph_vector_list_reserve()</code>. This is implemented by a trick similar to the C++ vector class: each time more memory is allocated for a vector, the size of the additionally allocated memory is the same as the vector's current length. (We assume here that the time complexity of memory allocation is at most linear).

igraph_vector_list_push_back_copy — Append the copy of a vector to the list.

```
igraph_error_t igraph_vector_list_push_back_copy(igraph_vector_list_t* v, const
```

This function resizes the list to be one element longer, and copies the specified vector given as an argument to the last element. The newly added element is owned by the list, but the ownership of the original vector is retained at the caller.

Arguments:

- v: The list object.
- e: Pointer to the vector to copy to the end of the list.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory.

Time complexity: same as igraph_vector_list_push_back() plus the time needed to copy the vector (which is O(n) for n elements in the vector).

igraph_vector_list_push_back_new — Append a new vector to the list.

```
igraph_error_t igraph_vector_list_push_back_new(igraph_vector_list_t* v, igraph_
```

This function resizes the list to be one element longer. The newly added element will be an empty vector that is owned by the list. A pointer to the newly added element is returned in the last argument if it is not NULL.

Arguments:

v: The list object.

result: Pointer to a vector pointer; this will be updated to point to the newly added vector. May be NULL if you do not need a pointer to the newly added vector.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory.

Time complexity: same as igraph_vector_list_push_back().

igraph_vector_list_pop_back — Remove the last item from the vector list and transfer ownership to the caller.

```
igraph_vector_t igraph_vector_list_pop_back(igraph_vector_list_t* v);
```

This function removes the last vector from the list. The vector that was removed from the list is returned and its ownership is passed back to the caller; in other words, the caller becomes responsible for destroying the vector when it is not needed any more.

It is an error to call this function with an empty vector.

Arguments:

v: The list object.

result: Pointer to an igraph_vector_t object; it will be updated to the item that was removed from the list. Ownership of this vector is passed on to the caller.

Time complexity: O(1).

igraph_vector_list_insert — Insert an existing vector into the list, transferring ownership.

igraph_error_t igraph_vector_list_insert(igraph_vector_list_t* v, igraph_intege

This function inserts e into the list at the given index, moving other items towards the end of the list as needed. The list takes ownership of the vector so the user is not responsible for freeing e any more; the vector will be destroyed when the list itself is destroyed or if e gets removed from the list without passing on the ownership to somewhere else.

Arguments:

v: The list object.

pos: The position where the new element is to be inserted.

e: Pointer to the vector to insert into the list.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory.

Time complexity: O(n).

igraph_vector_list_insert_copy — Insert the copy of a vector to the list.

```
igraph_error_t igraph_vector_list_insert_copy(igraph_vector_list_t* v, igraph_i
```

This function inserts a copy of *e* into the list at the given index, moving other items towards the end of the list as needed. The newly added element is owned by the list, but the ownership of the original vector is retained at the caller.

Arguments:

v: The list object.

pos: The position where the new element is to be inserted.

e: Pointer to the vector to copy to the end of the list.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory.

Time complexity: same as igraph_vector_list_insert() plus the time needed to copy the vector (which is O(n) for n elements in the vector).

igraph_vector_list_insert_new — Insert a new vector into the list.

igraph_error_t igraph_vector_list_insert_new(igraph_vector_list_t* v, igraph_in

This function inserts a newly created empty vector into the list at the given index, moving other items towards the end of the list as needed. The newly added vector is owned by the list. A pointer to the new element is returned in the last argument if it is not NULL.

Arguments:

v: The list object.

pos: The position where the new element is to be inserted.

result: Pointer to a vector pointer; this will be updated to point to the newly added vector. May

be NULL if you do not need a pointer to the newly added vector.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory.

Time complexity: same as igraph_vector_list_push_back().

igraph_vector_list_remove — Remove the item at the given index from the vector list and transfer ownership to the caller.

```
igraph_error_t igraph_vector_list_remove(igraph_vector_list_t* v, igraph_intege
```

This function removes the vector at the given index from the list, and moves all subsequent items in the list by one slot to the left to fill the gap. The vector that was removed from the list is returned in e and its ownership is passed back to the caller; in other words, the caller becomes responsible for destroying the vector when it is not needed any more.

Arguments:

v: The list object.

index: Index of the item to be removed.

result: Pointer to an igraph_vector_t object; it will be updated to the item that was re-

moved from the list. Ownership of this vector is passed on to the caller. It is an error to

supply a null pointer here.

See also:

igraph_vector_list_discard() if you are not interested in the item that was removed, igraph_vector_list_remove_fast() if you do not care about the order of the items in the list.

Time complexity: O(n), where n is the number of items in the list.

igraph_vector_list_remove_fast — Remove the item at the given index in the vector list, move the last item to its place and transfer ownership to the caller.

```
igraph_error_t igraph_vector_list_remove_fast(igraph_vector_list_t* v, igraph_i;
```

This function removes the vector at the given index from the list, moves the last item in the list to *index* to fill the gap, and then transfers ownership of the removed vector back to the caller; in other words, the caller becomes responsible for destroying the vector when it is not needed any more.

Arguments:

v: The list object.

index: Index of the item to be removed.

result: Pointer to an igraph_vector_t object; it will be updated to the item that was re-

moved from the list. Ownership of this vector is passed on to the caller. It is an error to

supply a null pointer here.

See also:

igraph_vector_list_remove() if you want to preserve the order of the items in the list, igraph_vector_list_discard_fast() if you are not interested in the item that was removed.

Time complexity: O(1).

igraph_vector_list_discard — Discard the item at the given index in the vector list.

```
void igraph_vector_list_discard(igraph_vector_list_t* v, igraph_integer_t index
```

This function removes the vector at the given index from the list, and moves all subsequent items in the list by one slot to the left to fill the gap. The vector that was removed from the list is destroyed automatically.

Arguments:

v: The list object.

index: Index of the item to be discarded and destroyed.

See also:

igraph_vector_list_discard_fast() if you do not care about the order of the items in the list, igraph_vector_list_remove() if you want to gain ownership of the item that was removed instead of destroying it.

Time complexity: O(n), where n is the number of items in the list.

igraph_vector_list_discard_back — Discard the last item in the vector list.

```
void igraph_vector_list_discard_back(igraph_vector_list_t* v);
```

This function removes the last vector from the list and destroys it.

Arguments:

v: The list object.

Time complexity: O(1).

igraph_vector_list_discard_fast — Discard the item at the given index in the vector list and move the last item to its place.

```
void igraph_vector_list_discard_fast(igraph_vector_list_t* v, igraph_integer_t
```

This function removes the vector at the given index from the list, and moves the last item in the list to index to fill the gap. The vector that was removed from the list is destroyed automatically.

Arguments:

v: The list object.

index: Index of the item to be discarded and destroyed.

See also:

igraph_vector_list_discard() if you want to preserve the order of the items in the list, igraph_vector_list_remove_fast() if you want to gain ownership of the item that was removed instead of destroying it.

Time complexity: O(1).

Sorting and reordering

igraph_vector_list_permute — Permutes the elements of a list in place according to an index vector.

```
igraph_error_t igraph_vector_list_permute(igraph_vector_list_t* v, const igraph_
```

This function takes a list v and a corresponding index vector index, and permutes the elements of v such that v[index[i]] is moved to become v[i] after the function is executed.

It is an error to call this function with an index vector that does not represent a valid permutation. Each element in the index vector must be between 0 and the length of the list minus one (inclusive), and each such element must appear only once. The function does not attempt to validate the index vector. Memory may be leaked if the index vector does not satisfy these conditions.

The index vector that this function takes is compatible with the index vector returned from igraph_vector_list_sort_ind(); passing in the index vector from igraph_vector_list_sort_ind() will sort the original vector.

Arguments:

v: the list to permute

index: the index vector

Time complexity: O(n), the number of items in the list.

igraph_vector_list_sort — Sorts the elements of the list into ascending order.

void igraph_vector_list_sort(igraph_vector_list_t *v, int (*cmp)(const igraph_v

Arguments:

v: Pointer to an initialized list object.

cmp: A comparison function that takes pointers to two vectors and returns zero if the two vectors are considered equal, any negative number if the first vector is smaller and any positive number if the second vector is smaller.

Returns:

Error code.

Time complexity: O(n log n) for n elements.

igraph_vector_list_sort_ind — Returns a permutation of indices that sorts the list.

```
igraph_error_t igraph_vector_list_sort_ind(
    igraph_vector_list_t *v, igraph_vector_int_t *inds,
    int (*cmp)(const igraph_vector_t*, const igraph_vector_t*)
);
```

Takes an unsorted list v as input and computes an array of indices inds such that v[inds[i]], with i increasing from 0, is an ordered array according to the comparison function cmp. The order of indices for identical elements is not defined.

Arguments:

v: the list to be sorted

inds: the output array of indices. This must be initialized, but will be resized

cmp: A comparison function that takes pointers to two vectors and returns zero if the two vectors are considered equal, any negative number if the first vector is smaller and any positive number if the second vector is smaller.

Returns:

Error code.

Time complexity: O(n log n) for n elements.

igraph_vector_list_swap — Swaps all elements of two vector lists.

```
igraph_error_t igraph_vector_list_swap(igraph_vector_list_t *v1, igraph_vector_
```

Arguments:

v1: The first list.

v2: The second list.

Returns:

Error code.

Time complexity: O(1).

igraph_vector_list_swap_elements — Swap two elements in a vector list.

igraph_error_t igraph_vector_list_swap_elements(igraph_vector_list_t *v1, igraph_error_t

Note that currently no range checking is performed.

Arguments:

v: The input list.

i: Index of the first element.

j: Index of the second element (may be the same as the first one).

Returns:

Error code, currently always IGRAPH_SUCCESS.

Time complexity: O(1).

Adjacency lists

Sometimes it is easier to work with a graph which is in adjacency list format: a list of vectors; each vector contains the neighbor vertices or incident edges of a given vertex. Typically, this representation is good if we need to iterate over the neighbors of all vertices many times. E.g. when finding the shortest paths between all pairs of vertices or calculating closeness centrality for all the vertices.

The igraph_adjlist_t stores the adjacency lists of a graph. After creation it is independent of the original graph, it can be modified freely with the usual vector operations, the graph is not affected. E.g. the adjacency list can be used to rewire the edges of a graph efficiently. If one used the straightforward $igraph_delete_edges()$ and $igraph_add_edges()$ combination for this that needs O(|V|+|E|) time for every single deletion and insertion operation, it is thus very slow if many edges are rewired. Extracting the graph into an adjacency list, do all the rewiring operations on the vectors of the adjacency list and then creating a new graph needs (depending on how exactly the rewiring is done) typically O(|V|+|E|) time for the whole rewiring process.

Lazy adjacency lists are a bit different. When creating a lazy adjacency list, the neighbors of the vertices are not queried, only some memory is allocated for the vectors. When <code>igraph_lazy_ad-jlist_get()</code> is called for vertex v the first time, the neighbors of v are queried and stored in a vector of the adjacency list, so they don't need to be queried again. Lazy adjacency lists are handy if you have an at least linear operation (because initialization is generally linear in terms of the number of vertices), but you don't know how many vertices you will visit during the computation.

Example 7.12. File examples/simple/adjlist.c

Adjacent vertices

igraph_adjlist_init — Constructs an adjacency list of vertices from an incidence list.

```
igraph_error_t igraph_adjlist_init_from_inclist(
    const igraph_t *graph, igraph_adjlist_t *al, const igraph_inclist_t *il);
```

In some algorithms it is useful to have an adjacency list *and* an incidence list representation of the same graph, and in many cases it is the most useful if they are consistent with each other, i.e. if can be guaranteed that the vertex ID in the i-th entry of the adjacency list of vertex v is the *other* endpoint of the edge in the i-th entry of the incidence list of the same vertex. This function creates such an adjacency list from the corresponding incidence list by looking up the endpoints of each edge in the incidence list and constructing the corresponding adjacency vectors.

The adjacency list is independent of the graph or the incidence list after creation; in other words, modifications that are made to the graph or the incidence list are not reflected in the adjacency list.

Arguments:

graph: The input graph.

a1: Pointer to an uninitialized igraph_adjlist_t object.

i1: Pointer to an *initialized* igraph_inclist_t object that will be converted into an adjacency list.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

igraph_adjlist_init_empty — Initializes an empty adjacency list.

```
igraph_error_t igraph_adjlist_init_empty(igraph_adjlist_t *al, igraph_integer_t
```

Creates a list of vectors, one for each vertex. This is useful when you are *constructing* a graph using an adjacency list representation as it does not require your graph to exist yet.

Arguments:

no_of_nodes: The number of vertices

a1: Pointer to an uninitialized igraph_adjlist_t object.

Returns:

Error code.

Time complexity: O(|V|), linear in the number of vertices.

igraph_adjlist_init_complementer — Adjacency lists for the complementer graph.

This function creates adjacency lists for the complementer of the input graph. In the complementer graph all edges are present which are not present in the original graph. Multiple edges in the input graph are ignored.

Arguments:

graph: The input graph.

a1: Pointer to a not yet initialized adjacency list.

mode: Constant specifying whether outgoing (IGRAPH_OUT), incoming (IGRAPH_IN), or both

(IGRAPH_ALL) types of neighbors (in the complementer graph) to include in the adja-

cency list. It is ignored for undirected networks.

loops: Whether to consider loop edges.

Returns:

Error code.

Time complexity: $O(|V|^2+|E|)$, quadratic in the number of vertices.

igraph_adjlist_destroy — Deallocates an adjacency list.

```
void igraph_adjlist_destroy(igraph_adjlist_t *al);
```

Free all memory allocated for an adjacency list.

Arguments:

a1: The adjacency list to destroy.

Time complexity: depends on memory management.

igraph_adjlist_get — Query a vector in an adjacency list.

```
#define igraph_adjlist_get(al,no)
```

Returns a pointer to an igraph_vector_int_t object from an adjacency list. The vector can be modified as desired.

Arguments:

a1: The adjacency list object.

no: The vertex whose adjacent vertices will be returned.

Returns:

Pointer to the igraph_vector_int_t object.

Time complexity: O(1).

igraph_adjlist_size — Returns the number of vertices in an adjacency list.

```
igraph_integer_t igraph_adjlist_size(const igraph_adjlist_t *al);
```

Arguments:

a1: The adjacency list.

Returns:

The number of vertices in the adjacency list.

Time complexity: O(1).

igraph_adjlist_clear — Removes all edges from an adjacency list.

```
void igraph_adjlist_clear(igraph_adjlist_t *al);
```

Arguments:

a1: The adjacency list. Time complexity: depends on memory management, typically O(n), where n is the total number of elements in the adjacency list.

igraph_adjlist_sort — Sorts each vector in an adjacency list.

```
void igraph_adjlist_sort(igraph_adjlist_t *al);
```

Sorts every vector of the adjacency list.

Arguments:

a1: The adjacency list.

Time complexity: O(n log n), n is the total number of elements in the adjacency list.

igraph_adjlist_simplify — Simplifies an adjacency list.

```
igraph_error_t igraph_adjlist_simplify(igraph_adjlist_t *al);
```

Simplifies an adjacency list, i.e. removes loop and multiple edges.

Arguments:

a1: The adjacency list.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of edges and vertices.

Incident edges

igraph_inclist_init — Initializes an incidence list.

Creates a list of vectors containing the incident edges for all vertices. The incidence list is independent of the graph after creation, subsequent changes of the graph object do not update the incidence list, and changes to the incidence list do not update the graph.

When mode is IGRAPH_IN or IGRAPH_OUT, each edge ID will appear in the incidence list once. When mode is IGRAPH_ALL, each edge ID will appear in the incidence list twice, once for the source vertex and once for the target edge. It also means that the edge IDs of loop edges may potentially appear twice for the same vertex. Use the loops argument to control whether this will be the case (IGRAPH_LOOPS_TWICE) or not (IGRAPH_LOOPS_ONCE or IGRAPH_NO_LOOPS).

Arguments:

graph: The input graph.

i 1: Pointer to an uninitialized incidence list.

mode: Constant specifying whether incoming edges (IGRAPH_IN), outgoing edges (IGRAPH_OUT) or both (IGRAPH_ALL) to include in the incidence lists of directed

graphs. It is ignored for undirected graphs.

100ps: Specifies how to treat loop edges. IGRAPH_NO_LOOPS removes loop edges from the

incidence list. IGRAPH_LOOPS_ONCE makes each loop edge appear only once in the incidence list of the corresponding vertex. IGRAPH_LOOPS_TWICE makes loop edges appear *twice* in the incidence list of the corresponding vertex, but only if the graph is

undirected or mode is set to IGRAPH ALL.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

igraph_inclist_destroy — Frees all memory allocated for an incidence list.

void igraph_inclist_destroy(igraph_inclist_t *il);

Arguments:

eal: The incidence list to destroy.

Time complexity: depends on memory management.

igraph_inclist_get — Query a vector in an incidence list.

```
#define igraph_inclist_get(il,no)
```

Returns a pointer to an igraph_vector_int_t object from an incidence list containing edge IDs. The vector can be modified, resized, etc. as desired.

Arguments:

i 1: Pointer to the incidence list.

no: The vertex for which the incident edges are returned.

Returns:

Pointer to an igraph_vector_int_t object.

Time complexity: O(1).

igraph_inclist_size — Returns the number of vertices in an incidence list.

```
igraph_integer_t igraph_inclist_size(const igraph_inclist_t *il);
```

Arguments:

i 1: The incidence list.

Returns:

The number of vertices in the incidence list.

Time complexity: O(1).

igraph_inclist_clear — Removes all edges from an incidence list.

```
void igraph_inclist_clear(igraph_inclist_t *il);
```

Arguments:

i1: The incidence list.

Time complexity: depends on memory management, typically O(n), where n is the total number of elements in the incidence list.

Lazy adjacency list for vertices

igraph_lazy_adjlist_init — Initialized a lazy adjacency list.

Create a lazy adjacency list for vertices. This function only allocates some memory for storing the vectors of an adjacency list, but the neighbor vertices are not queried, only at the <code>igraph_lazy_ad-jlist_get()</code> calls.

Arguments:

graph: The input graph.

a1: Pointer to an uninitialized adjacency list object.

mode: Constant, it gives whether incoming edges (IGRAPH_IN), outgoing edges (IGR-

PAH_OUT) or both types of edges (IGRAPH_ALL) are considered. It is ignored for

undirected graphs.

simplify: Constant, it gives whether to simplify the vectors in the adjacency list

(IGRAPH_SIMPLIFY) or not (IGRAPH_DONT_SIMPLIFY).

Returns:

Error code.

Time complexity: O(|V|), the number of vertices, possibly, but depends on the underlying memory management too.

igraph_lazy_adjlist_destroy — Deallocate a lazt adjacency list.

```
void igraph_lazy_adjlist_destroy(igraph_lazy_adjlist_t *al);
```

Free all allocated memory for a lazy adjacency list.

Arguments:

a1: The adjacency list to deallocate.

Time complexity: depends on the memory management.

igraph_lazy_adjlist_get — Query neighbor vertices.

```
#define igraph_lazy_adjlist_get(al,no)
```

If the function is called for the first time for a vertex then the result is stored in the adjacency list and no further query operations are needed when the neighbors of the same vertex are queried again.

Arguments:

a1: The lazy adjacency list.

no: The vertex ID to query.

Returns:

Pointer to a vector, or NULL upon error. It is safe to modify this vector, modification does not affect the original graph.

Time complexity: O(d), the number of neighbor vertices for the first time, O(1) for subsequent calls.

igraph_lazy_adjlist_size — Returns the number of vertices in a lazy adjacency list.

```
igraph_integer_t igraph_lazy_adjlist_size(const igraph_lazy_adjlist_t *al);
```

Arguments:

a1: The lazy adjacency list.

Returns:

The number of vertices in the lazy adjacency list.

Time complexity: O(1).

igraph_lazy_adjlist_clear — Removes all edges from a lazy adjacency list.

```
void igraph_lazy_adjlist_clear(igraph_lazy_adjlist_t *al);
```

Arguments:

a1: The lazy adjacency list. Time complexity: depends on memory management, typically O(n), where n is the total number of elements in the adjacency list.

Lazy incidence list for edges

igraph_lazy_inclist_init — Initializes a lazy incidence list of edges.

```
igraph_error_t igraph_lazy_inclist_init(const igraph_t *graph,
```

Data structure library: vector, matrix, other data types

```
igraph_lazy_inclist_t *il,
igraph_neimode_t mode,
igraph loops t loops);
```

Create a lazy incidence list for edges. This function only allocates some memory for storing the vectors of an incidence list, but the incident edges are not queried, only when <code>igraph_lazy_in-clist_get()</code> is called.

When mode is IGRAPH_IN or IGRAPH_OUT, each edge ID will appear in the incidence list once. When mode is IGRAPH_ALL, each edge ID will appear in the incidence list twice, once for the source vertex and once for the target edge. It also means that the edge IDs of loop edges will appear twice for the same vertex.

Arguments:

graph: The input graph.

a1: Pointer to an uninitialized incidence list.

mode: Constant, it gives whether incoming edges (IGRAPH_IN), outgoing edges (IGRAPH_OUT) or both types of edges (IGRAPH_ALL) are considered. It is ignored for

undirected graphs.

loops: Specifies how to treat loop edges. IGRAPH_NO_LOOPS removes loop edges from the

incidence list. <code>IGRAPH_LOOPS_ONCE</code> makes each loop edge appear only once in the incidence list of the corresponding vertex. <code>IGRAPH_LOOPS_TWICE</code> makes loop edges appear <code>twice</code> in the incidence list of the corresponding vertex, but only if the graph is

undirected or mode is set to IGRAPH_ALL.

Returns:

Error code.

Time complexity: O(|V|), the number of vertices, possibly. But it also depends on the underlying memory management.

igraph_lazy_inclist_destroy — Deallocates a lazy incidence list.

```
void igraph_lazy_inclist_destroy(igraph_lazy_inclist_t *il);
```

Frees all allocated memory for a lazy incidence list.

Arguments:

a1: The incidence list to deallocate.

Time complexity: depends on memory management.

igraph_lazy_inclist_get — Query incident edges.

```
#define igraph_lazy_inclist_get(al,no)
```

If the function is called for the first time for a vertex, then the result is stored in the incidence list and no further query operations are needed when the incident edges of the same vertex are queried again.

Arguments:

a1: The lazy incidence list object.

no: The vertex ID to query.

Returns:

Pointer to a vector, or NULL upon error. It is safe to modify this vector, modification does not affect the original graph.

Time complexity: O(d), the number of incident edges for the first time, O(1) for subsequent calls with the same *no* argument.

igraph_lazy_inclist_size — Returns the number of vertices in a lazy incidence list.

```
igraph_integer_t igraph_lazy_inclist_size(const igraph_lazy_inclist_t *il);
```

Arguments:

i1: The lazy incidence list.

Returns:

The number of vertices in the lazy incidence list.

Time complexity: O(1).

igraph_lazy_inclist_clear — Removes all edges from a lazy incidence list.

```
void igraph_lazy_inclist_clear(igraph_lazy_inclist_t *il);
```

Arguments:

i1: The lazy incidence list.

Time complexity: depends on memory management, typically O(n), where n is the total number of elements in the incidence list.

Partial prefix sum trees

The igraph_psumtree_t data type represents a partial prefix sum tree. A partial prefix sum tree is a data structure that can be used to draw samples from a discrete probability distribution with dynamic probabilities that are updated frequently. This is achieved by creating a binary tree where the leaves are the items. Each leaf contains the probability corresponding to the items. Intermediate nodes of the tree always contain the sum of its two children. When the value of a leaf node is updated, the values of its ancestors are also updated accordingly.

Samples can be drawn from the probability distribution represented by the tree by generating a uniform random number between 0 (inclusive) and the value of the root of the tree (exclusive), and then

following the branches of the tree as follows. In each step, the value in the current node is compared with the generated number. If the value in the node is larger, the left branch of the tree is taken; otherwise the generated number is decreased by the value in the node and the right branch of the tree is taken, until a leaf node is reached.

Note that the sampling process works only if all the values in the tree are non-negative. This is enforced by the object; in particular, trying to set a negative value for an item will produce an igraph error.

igraph_psumtree_init — Initializes a partial prefix sum tree.

```
igraph_error_t igraph_psumtree_init(igraph_psumtree_t *t, igraph_integer_t size
```

The tree is initialized with a fixed number of elements. After initialization, the value corresponding to each element is zero.

Arguments:

t: The tree to initialize.

size: The number of elements in the tree. It must be at least one.

Returns:

Error code, typically IGRAPH_ENOMEM if there is not enough memory.

Time complexity: O(n) for a tree containing n elements

igraph_psumtree_destroy — Destroys a partial prefix sum tree.

```
void igraph_psumtree_destroy(igraph_psumtree_t *t);
```

All partial prefix sum trees initialized by <code>igraph_psumtree_init()</code> should be properly destroyed by this function. A destroyed tree needs to be reinitialized by <code>igraph_psumtree_init()</code> if you want to use it again.

Arguments:

t: Pointer to the (previously initialized) tree to destroy.

Time complexity: operating system dependent.

igraph_psumtree_size — Returns the size of the tree.

```
igraph_integer_t igraph_psumtree_size(const igraph_psumtree_t *t);
```

Arguments:

t: The tree object

Returns:

The number of discrete items in the tree.

Time complexity: O(1).

igraph_psumtree_get — Retrieves the value corresponding to an item in the tree.

```
igraph_real_t igraph_psumtree_get(const igraph_psumtree_t *t, igraph_integer_t
```

Arguments:

t: The tree to query.

idx: The index of the item whose value is to be retrieved.

Returns:

The value corresponding to the item with the given index.

Time complexity: O(1)

igraph_psumtree_sum — Returns the sum of the values of the leaves in the tree.

```
igraph_real_t igraph_psumtree_sum(const igraph_psumtree_t *t);
```

Arguments:

t: The tree object

Returns:

The sum of the values of the leaves in the tree.

Time complexity: O(1).

igraph_psumtree_search — Finds an item in the tree, given a value.

igraph_error_t igraph_psumtree_search(const igraph_psumtree_t *t, igraph_intege:

```
igraph_real_t search);
```

This function finds the item with the lowest index where it holds that the sum of all the items with a *lower* index is less than or equal to the given value and that the sum of all the items with a lower index plus the item itself is larger than the given value.

If you think about the partial prefix sum tree as a tool to sample from a discrete probability distribution, then calling this function repeatedly with uniformly distributed random numbers in the range 0 (inclusive) to the sum of all values in the tree (exclusive) will sample the items in the tree with a probability that is proportional to their associated values.

Arguments:

t: The tree to query.

idx: The index of the item is returned here.

search: The value to use for the search. Must be in the interval [0, sum), where sum is the

sum of all elements (leaves) in the tree.

Returns:

Error code; currently the search always succeeds.

Time complexity: O(log n), where n is the number of items in the tree.

igraph_psumtree_update — Updates the value associated to an item in the tree.

Arguments:

t: The tree to query.

idx: The index of the item to update.

new_value: The new value of the item.

Returns:

Error code, IGRAPH_EINVAL if the new value is negative or NaN, IGRAPH_SUCCESS if the operation was successful.

Time complexity: O(log n), where n is the number of items in the tree.

Chapter 8. Random numbers

About random numbers in igraph

Some algorithms in igraph, such as sampling from random graph models, require random number generators (RNGs). igraph includes a flexible RNG framework that allows hooking up arbitrary random number generators, and comes with several ready-to-use generators. This framework is used in igraph's high-level interfaces to integrate with the host language's own RNG.

The default random number generator

igraph_rng_default — Query the default random number generator.

```
igraph_rng_t *igraph_rng_default(void);
```

Returns:

A pointer to the default random number generator.

See also:

igraph_rng_set_default()

igraph_rng_set_default — Set the default igraph random number generator.

```
void igraph_rng_set_default(igraph_rng_t *rng);
```

This function *copies* the internal structure of the given igraph_rng_t object to igraph's internal default RNG structure. The structure itself contains two pointers only, one to the "methods" of the RNG and one to the memory buffer holding the internal state of the RNG. This means that if you keep on generating random numbers from the RNG after setting it as the default, it will affect the state of the default RNG as well because the two share the same state pointer. However, do *not* expect igraph_rng_default() to return the same pointer as the one you passed in here - the state is shared, but the entire structure is not.

Arguments:

rng: The random number generator to use as default from now on. Calling igraph_rng_de-stroy() on it, while it is still being used as the default will result in crashes and/or unpredictable results.

Time complexity: O(1).

Creating random number generators

igraph_rng_init — Initializes a random number generator.

```
igraph_error_t igraph_rng_init(igraph_rng_t *rng, const igraph_rng_type_t *type
```

This function allocates memory for a random number generator, with the given type, and sets its seed to the default.

Arguments:

rng: Pointer to an uninitialized RNG.

type: The type of the RNG, such as igraph_rngtype_mt19937, igraph_rng-type_glibc2, igraph_rngtype_pcg32 or igraph_rngtype_pcg64.

Returns:

Error code.

igraph_rng_destroy — Deallocates memory associated with a random number generator.

```
void igraph_rng_destroy(igraph_rng_t *rng);
```

Arguments:

rng: The RNG to destroy. Do not destroy an RNG that is used as the default igraph RNG.

Time complexity: O(1).

igraph_rng_seed — Seeds a random number generator.

```
igraph_error_t igraph_rng_seed(igraph_rng_t *rng, igraph_uint_t seed);
```

Arguments:

rng: The RNG.

seed: The new seed.

Returns:

Error code.

Time complexity: usually O(1), but may depend on the type of the RNG.

igraph_rng_bits — The number of random bits that a random number generator can produces in a single round.

```
igraph_integer_t igraph_rng_bits(const igraph_rng_t* rng);
```

Arguments:

rng: The RNG.

Returns:

The number of random bits that can be generated in a single round with the RNG.

Time complexity: O(1).

igraph_rng_max — The maximum possible integer for a random number generator.

```
igraph_uint_t igraph_rng_max(const igraph_rng_t *rng);
```

Note that this number is only for informational purposes; it returns the maximum possible integer that can be generated with the RNG with a single call to its internals. It is derived directly from the number of random *bits* that the RNG can generate in a single round. When this is smaller than what would be needed by other RNG functions like <code>igraph_rng_get_integer()</code>, igraph will call the RNG multiple times to generate more random bits.

Arguments:

rng: The RNG.

Returns:

The largest possible integer that can be generated in a single round with the RNG.

Time complexity: O(1).

igraph_rng_name — The type of a random number generator.

```
const char *igraph_rng_name(const igraph_rng_t *rng);
```

Arguments:

rng: The RNG.

Returns:

The name of the type of the generator. Do not deallocate or change the returned string.

Time complexity: O(1).

Generating random numbers

igraph_rng_get_integer — Generate an integer random number from an interval.

```
igraph_integer_t igraph_rng_get_integer(
    igraph_rng_t *rng, igraph_integer_t l, igraph_integer_t h
);
```

Arguments:

rng: Pointer to the RNG to use for the generation. Use igraph_rng_default() here to use the default igraph RNG.

1: Lower limit, inclusive, it can be negative as well.

h: Upper limit, inclusive, it can be negative as well, but it should be at least 1.

Returns:

The generated random integer.

Time complexity: O(log2(h-l) / bits) where bits is the value of igraph_rng_bits(rng).

igraph_rng_get_unif01 — Samples uniformly from the unit interval.

```
igraph_real_t igraph_rng_get_unif01(igraph_rng_t *rng);
```

Generates uniformly distributed real numbers from the [0, 1) half-open interval.

Arguments:

rng: Pointer to the RNG to use. Use igraph_rng_default() here to use the default igraph
RNG.

Returns:

The generated uniformly distributed random number.

Time complexity: depends on the type of the RNG.

igraph_rng_get_unif — Samples real numbers from a given interval.

Generates uniformly distributed real numbers from the [1, h) half-open interval.

Arguments:

- rng: Pointer to the RNG to use. Use igraph_rng_default() here to use the default igraph
 RNG.
- 1: The lower bound, it can be negative.
- h: The upper bound, it can be negative, but it has to be larger than the lower bound.

Returns:

The generated uniformly distributed random number.

Time complexity: depends on the type of the RNG.

igraph_rng_get_normal — Samples from a normal distribution.

Generates random variates from a normal distribution with probability density

```
\exp(-(x - m)^2 / (2 s^2)).
```

Arguments:

- *m*: The mean.
- s: The standard deviation.

Returns:

The generated normally distributed random number.

Time complexity: depends on the type of the RNG.

igraph_rng_get_exp — Samples from an exponential distribution.

```
igraph_real_t igraph_rng_get_exp(igraph_rng_t *rng, igraph_real_t rate);
```

Generates random variates from an exponential distribution with probability density proportional to

```
exp(-rate x).
```

Arguments:

rng: Pointer to the RNG to use. Use igraph_rng_default() here to use the default igraph

RNG.

rate: Rate parameter.

Returns:

The generated sample.

Time complexity: depends on the RNG.

igraph_rng_get_gamma — Samples from a gamma distribution.

Generates random variates from a gamma distribution with probability density proportional to

```
x^{(shape-1)} exp(-x / scale).
```

Arguments:

rng: Pointer to the RNG to use. Use igraph_rng_default() here to use the default igraph

RNG.

shape: Shape parameter.

scale: Scale parameter.

Returns:

The generated sample.

Time complexity: depends on the RNG.

igraph_rng_get_binom — Samples from a binomial distribution.

```
igraph_real_t igraph_rng_get_binom(igraph_rng_t *rng, igraph_integer_t n, igraph_
```

Generates random variates from a binomial distribution. The number k is generated with probability

```
(n \land choose k) p^k (1-p)^n(n-k), k = 0, 1, ..., n.
```

Arguments:

rng: Pointer to the RNG to use. Use igraph_rng_default() here to use the default igraph RNG.

- n: Number of observations.
- p: Probability of an event.

Returns:

The generated binomially distributed random number.

Time complexity: depends on the RNG.

igraph_rng_get_geom — Samples from a geometric distribution.

```
igraph_real_t igraph_rng_get_geom(igraph_rng_t *rng, igraph_real_t p); Generates random variates from a geometric distribution. The number k is generated with probability (1 - p)^k p, k = 0, 1, 2, \ldots
```

Arguments:

rng: Pointer to the RNG to use. Use igraph_rng_default() here to use the default igraph RNG

p: The probability of success in each trial. Must be larger than zero and smaller or equal to 1.

Returns:

The generated geometrically distributed random number.

Time complexity: depends on the RNG.

igraph_rng_get_pois — Samples from a Poisson distribution.

Arguments:

rng: Pointer to the RNG to use. Use igraph_rng_default() here to use the default igraph RNG.

rate: The rate parameter of the Poisson distribution. Must not be negative.

Returns:

The generated geometrically distributed random number.

Time complexity: depends on the RNG.

Supported random number generators

By default igraph uses the MT19937 generator. Prior to igraph version 0.6, the generator supplied by the standard C library was used. This means the GLIBC2 generator on GNU libc 2 systems, and maybe the BSD RAND generator on others. The RAND generator was removed due to poor statistical properties in version 0.10. The PCG32 generator was added in version 0.10.

igraph_rngtype_mt19937 — The MT19937 random number generator.

```
const igraph_rng_type_t igraph_rngtype_mt19937 = {
    /* name= */ "MT19937",
    /* Hame-
/* bits= */ 32,
/* init= */ igraph_rng_mt19937_init,
--- m+19937_destr
    /* destroy= */ igraph_rng_mt19937_destroy,
    /* seed= */
                    igraph_rng_mt19937_seed,
    /* get= */
                     igraph_rng_mt19937_get,
    /* get int= */ 0,
    /* get_real= */ 0,
    /* get_norm= */ 0,
    /* get_geom= */ 0,
    /* get_binom= */ 0,
    /* get_exp= */ 0,
    /* get_gamma= */ 0,
    /* get pois= */ 0
};
```

The MT19937 generator of Makoto Matsumoto and Takuji Nishimura is a variant of the twisted generalized feedback shift-register algorithm, and is known as the "Mersenne Twister" generator. It has a Mersenne prime period of 2^19937 - 1 (about 10^6000) and is equi-distributed in 623 dimensions. It has passed the diehard statistical tests. It uses 624 words of state per generator and is comparable in speed to the other generators. The original generator used a default seed of 4357 and choosing s equal to zero in <code>igraph_rng_mt19937_seed()</code> reproduces this. Later versions switched to 5489 as the default seed, you can choose this explicitly via <code>igraph_rng_seed()</code> instead if you require it.

For more information see, Makoto Matsumoto and Takuji Nishimura, "Mersenne Twister: A 623-dimensionally equidistributed uniform pseudorandom number generator". ACM Transactions on Modeling and Computer Simulation, Vol. 8, No. 1 (Jan. 1998), Pages 3–30

The generator igraph_rngtype_mt19937 uses the second revision of the seeding procedure published by the two authors above in 2002. The original seeding procedures could cause spurious artifacts for some seed values.

This generator was ported from the GNU Scientific Library.

igraph_rngtype_glibc2 — The random number generator introduced in GNU libc 2.

```
const igraph_rng_type_t igraph_rngtype_glibc2 = {
   /* name= */    "LIBC",
   /* bits= */    31,
   /* init= */    igraph rng glibc2 init,
```

```
/* destroy= */ igraph_rng_glibc2_destroy,
/* seed= */ igraph_rng_glibc2_seed,
/* get= */ igraph_rng_glibc2_get,
/* get_int= */ 0,
/* get_real= */ 0,
/* get_norm= */ 0,
/* get_geom= */ 0,
/* get_binom= */ 0,
/* get_exp= */ 0,
/* get_exp= */ 0,
/* get_pois= */ 0
};
```

This is a linear feedback shift register generator with a 128-byte buffer. This generator was the default prior to igraph version 0.6, at least on systems relying on GNU libc. This generator was ported from the GNU Scientific Library. It is a reimplementation and does not call the system glibc generator.

igraph_rngtype_pcg32 — The PCG random number generator (32-bit version).

```
const igraph_rng_type_t igraph_rngtype_pcg32 = {
    /* name= */ "PCG32",
    /* bits= */ 32,
/* init= */ igraph_rng_pcg32_init,
    /* destroy= */ igraph_rng_pcg32_destroy,
/* seed= */ igraph_rng_pcg32_seed,
/* get= */ igraph_rng_pcg32_get
     /* get= */
                        igraph_rng_pcg32_get,
     /* get_int= */
                        Ο,
     /* get_real= */ 0,
     /* get_norm= */ 0,
     /* get_geom= */ 0,
     /* get_binom= */ 0,
     /* get_exp= */ 0,
     /* get_gamma= */ 0,
     /* get_pois= */ 0
};
```

This is an implementation of the PCG random number generator; see https://www.pcg-random.org for more details. This implementation returns 32 random bits in a single iteration.

The generator was ported from the original source code published by the authors at https://github.com/imneme/pcg-c.

igraph_rngtype_pcg64 — The PCG random number generator (64-bit version).

This is an implementation of the PCG random number generator; see https://www.pcg-random.org for more details. This implementation returns 64 random bits in a single iteration. It is only available on 64-bit plaforms with compilers that provide the __uint128_t type.

PCG64 typically provides better performance than PCG32 when sampling floating point numbers or very large integers, as it can provide twice as many random bits in a single generation round.

The generator was ported from the original source code published by the authors at https://github.com/imneme/pcg-c.

Use cases

Normal (default) use

If the user does not use any of the RNG functions explicitly, but calls some of the randomized igraph functions, then a default RNG is set up the first time an igraph function needs random numbers. The seed of this RNG is the output of the time(0) function call, using the time function from the standard C library. This ensures that igraph creates a different random graph, each time the C program is called.

The created default generator is stored internally and can be queried with the igraph_rng_default() function.

Reproducible simulations

If reproducible results are needed, then the user should set the seed of the default random number generator explicitly, using the <code>igraph_rng_seed()</code> function on the default generator, <code>igraph_rng_default()</code>. When setting the seed to the same number, igraph generates exactly the same random graph (or series of random graphs).

Changing the default generator

By default igraph uses the <code>igraph_rng_default()</code> random number generator. This can be changed any time by calling <code>igraph_rng_set_default()</code>, with an already initialized random number generator. Note that the old (replaced) generator is not destroyed, so no memory is deallocated.

Using multiple generators

igraph also provides functions to set up multiple random number generators, using the igraph_rng_init() function, and then generating random numbers from them, e.g. with igraph_rng_get_integer() and/or igraph_rng_get_unif() calls.

Note that initializing a new random number generator is independent of the generator that the igraph functions themselves use. If you want to replace that, then please use <code>igraph_rng_set_de-fault()</code>.

Example

 $Example \ 8.1. \ File \ {\tt examples/simple/random_seed.c}$

Chapter 9. Graph generators

Graph generators create graphs.

Almost all functions which create graph objects are documented here. The exceptions are igraph_induced_subgraph() and alike, these create graphs based on another graph.

Deterministic graph generators

igraph_create — Creates a graph with the specified edges.

Arguments:

graph: An uninitialized graph object.

edges: The edges to add, the first two elements are the first edge, etc.

n: The number of vertices in the graph, if smaller or equal to the highest vertex ID in the

edges vector it will be increased automatically. So it is safe to give 0 here.

directed: Boolean, whether to create a directed graph or not. If yes, then the first edge points

from the first vertex ID in edges to the second, etc.

Returns:

Error code: IGRAPH_EINVEVECTOR: invalid edges vector (odd number of vertices). IGRAPH_EINVVID: invalid (negative) vertex ID.

Time complexity: O(|V|+|E|), |V| is the number of vertices, |E| the number of edges in the graph.

Example 9.1. File examples/simple/igraph create.c

igraph_small — Shorthand to create a small graph, giving the edges as arguments.

```
igraph_error_t igraph_small(igraph_t *graph, igraph_integer_t n, igraph_bool_t int first, ...);
```

This function is handy when a relatively small graph needs to be created. Instead of giving the edges as a vector, they are given simply as arguments and a '-1' needs to be given after the last meaningful edge argument.

Note that only graphs which have vertices less than the highest value of the 'int' type can be created this way. If you give larger values then the result is undefined.

Arguments:

graph: Pointer to an uninitialized graph object. The result will be stored here.

n: The number of vertices in the graph; a nonnegative integer.

directed: Logical constant; gives whether the graph should be directed. Supported values are:

IGRAPH_DIRECTED The graph to be created will be *directed*.

IGRAPH_UNDIRECTED The graph to be created will be *undirected*.

...: The additional arguments giving the edges of the graph. Don't forget to supply an

additional '-1' after the last (meaningful) argument. The first parameter is present

for technical reasons and represents the first variadic argument.

Returns:

Error code.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph to create.

Example 9.2. File examples/simple/igraph_small.c

igraph_adjacency — Creates a graph from an adjacency matrix.

```
igraph_error_t igraph_adjacency(
    igraph_t *graph, const igraph_matrix_t *adjmatrix, igraph_adjacency_t mode,
    igraph_loops_t loops
);
```

The order of the vertices in the matrix is preserved, i.e. the vertex corresponding to the first row/column will be vertex with id 0, the next row is for vertex 1, etc.

Arguments:

graph: Pointer to an uninitialized graph object.

adjmatrix: The adjacency matrix. How it is interpreted depends on the mode argument.

mode: Constant to specify how the given matrix is interpreted as an adjacency matrix. Pos-

sible values (A(i,j) is the element in row i and column j in the adjacency matrix ad-

jmatrix):

IGRAPH_ADJ_DIRECTED the graph will be directed and an element gives the

number of edges between two vertices.

IGRAPH_ADJ_UNDIRECTED this is the same as IGRAPH_ADJ_MAX, for con-

venience.

IGRAPH_ADJ_MAX undirected graph will be created and the number

of edges between vertices i and j is max(A(i,j),

A(j,i)).

 ${\tt IGRAPH_ADJ_MIN} \qquad \qquad {\tt undirected \ graph \ will \ be \ created \ with \ min} (A(i,j),$

A(j,i)) edges between vertices i and j.

IGRAPH_ADJ_PLUS undirected graph will be created with A(i,j)+A(j,i)

edges between vertices i and j.

IGRAPH_ADJ_UPPER undirected graph will be created, only the upper

right triangle (including the diagonal) is used for

the number of edges.

IGRAPH_ADJ_LOWER undirected graph will be created, only the lower

left triangle (including the diagonal) is used for

creating the edges.

loops: Constant to specify how the diagonal of the matrix should be treated when creating

loop edges.

IGRAPH_NO_LOOPS Ignore the diagonal of the input matrix and do not cre-

ate loops.

incident on the corresponding vertex.

IGRAPH_LOOPS_TWICE Treat the diagonal entries as *twice* the number of loop

edges incident on the corresponding vertex. Odd numbers in the diagonal will return an error code.

Returns:

Error code, IGRAPH_NONSQUARE: non-square matrix. IGRAPH_EINVAL: Negative entry was found in adjacency matrix, or an odd number was found in the diagonal with IGRAPH_LOOPS_TWICE

Time complexity: O(|V||V|), |V| is the number of vertices in the graph.

Example 9.3. File examples/simple/igraph_adjacency.c

igraph_weighted_adjacency — Creates a graph from a weighted adjacency matrix.

```
igraph_error_t igraph_weighted_adjacency(
    igraph_t *graph, const igraph_matrix_t *adjmatrix, igraph_adjacency_t mode,
    igraph_vector_t *weights, igraph_loops_t loops
);
```

The order of the vertices in the matrix is preserved, i.e. the vertex corresponding to the first row/column will be vertex with id 0, the next row is for vertex 1, etc.

Arguments:

graph: Pointer to an uninitialized graph object.

adjmatrix: The weighted adjacency matrix. How it is interpreted depends on the mode argu-

ment. The common feature is that edges with zero weights are considered nonexis-

tent (however, negative weights are permitted).

mode: Constant to specify how the given matrix is interpreted as an adjacency matrix. Pos-

sible values (A(i,j) is the element in row i and column j in the adjacency matrix ad-

jmatrix):

IGRAPH_ADJ_DIRECTED	the graph will be directed	d and an element gives the
---------------------	----------------------------	----------------------------

weight of the edge between two vertices.

IGRAPH_ADJ_UNDIRECTED this is the same as IGRAPH_ADJ_MAX, for con-

venience.

IGRAPH_ADJ_MAX undirected graph will be created and the weight

of the edge between vertices i and j is max(A(i,j),

A(j,i)).

IGRAPH_ADJ_MIN undirected graph will be created with edge weight

min(A(i,j), A(j,i)) between vertices i and j.

IGRAPH_ADJ_PLUS undirected graph will be created with edge weight

A(i,j)+A(j,i) between vertices i and j.

IGRAPH_ADJ_UPPER undirected graph will be created, only the upper

right triangle (including the diagonal) is used for

the edge weights.

IGRAPH_ADJ_LOWER undirected graph will be created, only the lower

left triangle (including the diagonal) is used for

the edge weights.

weights: Pointer to an initialized vector, the weights will be stored here.

loops: Constant to specify how the diagonal of the matrix should be treated when creating

loop edges.

IGRAPH_NO_LOOPS Ignore the diagonal of the input matrix and do not cre-

ate loops.

IGRAPH_LOOPS_ONCE Treat the diagonal entries as the weight of the loop

edge incident on the corresponding vertex.

IGRAPH_LOOPS_TWICE Treat the diagonal entries as twice the weight of the

loop edge incident on the corresponding vertex.

Returns:

Error code, IGRAPH_NONSQUARE: non-square matrix.

Time complexity: O(|V||V|), |V| is the number of vertices in the graph.

Example 9.4. File examples/simple/igraph_weighted_adjacency.c

igraph_sparse_adjacency — Creates a graph from a sparse adjacency matrix.

```
igraph_error_t igraph_sparse_adjacency(igraph_t *graph, igraph_sparsemat_t *adjacency_t mode, igraph_loops_t loops);
```

This has the same functionality as $igraph_adjacency()$, but uses a column-compressed adjacency matrix. Time complexity: O(|E|), where |E| is the number of edges in the graph.

igraph_sparse_weighted_adjacency — Creates a graph from a weighted sparse adjacency matrix.

```
igraph_error_t igraph_sparse_weighted_adjacency(
    igraph_t *graph, igraph_sparsemat_t *adjmatrix, igraph_adjacency_t mode,
    igraph_vector_t *weights, igraph_loops_t loops
);
```

This has the same functionality as $igraph_weighted_adjacency()$, but uses a column-compressed adjacency matrix. Time complexity: O(|E|), where |E| is the number of edges in the graph.

igraph_adjlist — Creates a graph from an adjacency list.

An adjacency list is a list of vectors, containing the neighbors of all vertices. For operations that involve many changes to the graph structure, it is recommended that you convert the graph into an adjacency list via igraph_adjlist_init(), perform the modifications (these are cheap for an adjacency list) and then recreate the igraph graph via this function.

Arguments:

graph: Pointer to an uninitialized graph object.

adjlist: The adjacency list.

mode: Whether or not to create a directed graph. IGRAPH_ALL means an undirected graph,

IGRAPH_OUT means a directed graph from an out-adjacency list (i.e. each list contains the successors of the corresponding vertices), IGRAPH_IN means a directed

graph from an in-adjacency list

duplicate: Logical, for undirected graphs this specified whether each edge is included twice, in

the vectors of both adjacent vertices. If this is false (0), then it is assumed that every

edge is included only once. This argument is ignored for directed graphs.

Returns:

Error code.

See also:

igraph_adjlist_init() for the opposite operation.

Time complexity: O(|V|+|E|).

igraph_star — Creates a star graph, every vertex connects only to the center.

Arguments:

graph: Pointer to an uninitialized graph object, this will be the result.

n: Integer constant, the number of vertices in the graph.

mode: Constant, gives the type of the star graph to create. Possible values:

IGRAPH_STAR_OUT directed star graph, edges point from the center to the

other vertices.

IGRAPH_STAR_IN directed star graph, edges point to the center from the

other vertices.

IGRAPH_STAR_MUTUAL directed star graph with mutual edges.

IGRAPH_STAR_UNDIRECTED an undirected star graph is created.

center: Id of the vertex which will be the center of the graph.

Returns:

Error code:

IGRAPH_EINVVID invalid number of vertices.

IGRAPH_EINVAL invalid center vertex.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(|V|), the number of vertices in the graph.

See also:

igraph_square_lattice(), igraph_ring(), igraph_kary_tree() for creating
other regular structures.

Example 9.5. File examples/simple/igraph_star.c

igraph_wheel — Creates a wheel graph, a union of a star and a cycle graph.

A wheel graph on n vertices can be thought of as a wheel with n-1 spokes. The cycle graph part makes up the rim, while the star graph part adds the spokes.

Note that the two and three-vertex wheel graphs are non-simple: The two-vertex wheel graph contains a self-loop, while the three-vertex wheel graph contains parallel edges (a 1-cycle and a 2-cycle, respectively).

Arguments:

graph: Pointer to an uninitialized graph object, this will be the result.

n: Integer constant, the number of vertices in the graph.

mode: Constant, gives the type of the star graph to create. Possible values:

IGRAPH_WHEEL_OUT directed wheel graph, edges point from the center

to the other vertices.

IGRAPH_WHEEL_IN directed wheel graph, edges point to the center from

the other vertices.

IGRAPH_WHEEL_MUTUAL directed wheel graph with mutual edges.

IGRAPH_WHEEL_UNDIRECTED an undirected wheel graph is created.

center: Id of the vertex which will be the center of the graph.

Returns:

Error code:

IGRAPH_EINVVID invalid number of vertices.

IGRAPH_EINVAL invalid center vertex.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(|V|), the number of vertices in the graph.

See also:

igraph_square_lattice — Arbitrary dimensional square lattices.

```
igraph_error_t igraph_square_lattice(
    igraph_t *graph, const igraph_vector_int_t *dimvector, igraph_integer_t nei
    igraph_bool_t directed, igraph_bool_t mutual, const igraph_vector_bool_t *p
);
```

Creates d-dimensional square lattices of the given size. Optionally, the lattice can be made periodic, and the neighbors within a given graph distance can be connected.

In the zero-dimensional case, the singleton graph is returned.

The vertices of the resulting graph are ordered such that the index of the vertex at position (i_0 , i_1 , i_2 , ..., i_d) in a lattice of size (n_0 , n_1 , ..., n_d) will be $i_0 + n_0$ * $i_1 + n_0 * n_1 * i_2 + ...$

Arguments:

graph: An uninitialized graph object.

dimvector: Vector giving the sizes of the lattice in each of its dimensions. The dimension of the

lattice will be the same as the length of this vector.

nei: Integer value giving the distance (number of steps) within which two vertices will

be connected.

directed: Boolean, whether to create a directed graph. If the mutual and circular argu-

ments are not set to true, edges will be directed from lower-index vertices towards

higher-index ones.

mutual: Boolean, if the graph is directed this gives whether to create all connections as mu-

tual.

periodic: Boolean vector, defines whether the generated lattice is periodic along each dimen-

sion. The length of this vector must match the length of dimvector. This parame-

ter may also be NULL, which implies that the lattice will not be periodic.

Returns:

Error code: IGRAPH_EINVAL: invalid (negative) dimension vector or mismatch between the length of the dimension vector and the periodicity vector.

Time complexity: If nei is less than two then it is O(|V|+|E|) (as far as I remember), |V| and |E| are the number of vertices and edges in the generated graph. Otherwise it is $O(|V|*d^k+|E|)$, d is the average degree of the graph, k is the nei argument.

igraph_ring — Creates a cycle graph or a path graph.

A circular ring on n vertices is commonly known in graph theory as the cycle graph, and often denoted by C_n. Removing a single edge from the cycle graph C_n results in the path graph P_n. This function can generate both.

When n is 1 or 2, the result may not be a simple graph: the one-cycle contains a self-loop and the undirected or reciprocally connected directed two-cycle contains parallel edges.

Arguments:

graph: Pointer to an uninitialized graph object.

n: The number of vertices in the graph.

directed: Logical, whether to create a directed graph. All edges will be oriented in the same

direction along the cycle or path.

mutual: Logical, whether to create mutual edges in directed graphs. It is ignored for undirected

graphs.

circular: Logical, whether to create a closed ring (a cycle) or an open path.

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices.

Time complexity: O(|V|), the number of vertices in the graph.

See also:

igraph_lattice() for generating more general lattices.

Example 9.6. File examples/simple/igraph_ring.c

igraph_kary_tree — Creates a k-ary tree in which almost all vertices have k children.

To obtain a completely symmetric tree with 1 layers, where each vertex has precisely children descendants, use $n = (children^(1+1) - 1) / (children - 1)$. Such trees are often called k-ary trees, where k refers to the number of children.

Note that for n=0, the null graph is returned, which is not considered to be a tree by $igraph_is_tree()$.

Arguments:

graph: Pointer to an uninitialized graph object.

n: Integer, the number of vertices in the graph.

children: Integer, the number of children of a vertex in the tree.

type: Constant, gives whether to create a directed tree, and if this is the case, also its orien-

tation. Possible values:

IGRAPH_TREE_OUT directed tree, the edges point from the parents to

their children,

IGRAPH_TREE_IN directed tree, the edges point from the children to

their parents.

IGRAPH_TREE_UNDIRECTED undirected tree.

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices. IGRAPH_INVMODE: invalid mode argument.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

See also:

igraph_lattice(), igraph_star() for creating other regular structures;
igraph_from_prufer() for creating arbitrary trees; igraph_tree_game() for uniform
random sampling of trees.

Example 9.7. File examples/simple/igraph_kary_tree.c

igraph_symmetric_tree — Creates a symmetric tree with the specified number of branches at each level.

This function creates a tree in which all vertices at distance d from the root have branch-ing_counts[d] children.

Arguments:

graph: Pointer to an uninitialized graph object.

branches: Vector detailing the number of branches at each level.

type: Constant, gives whether to create a directed tree, and if this is the case, also its orien-

tation. Possible values:

IGRAPH_TREE_OUT directed tree, the edges point from the parents to

their children,

IGRAPH_TREE_IN directed tree, the edges point from the children to

their parents.

IGRAPH_TREE_UNDIRECTED undirected tree.

Returns:

Error code: IGRAPH_INVMODE: invalid mode argument. IGRAPH_EINVAL: invalid number of children.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

See also:

igraph_kary_tree(), igraph_regular_tree() and igraph_star() for creating other regular tree structures; igraph_from_prufer() for creating arbitrary trees; igraph_tree_game() for uniform random sampling of trees.

Example 9.8. File examples/simple/igraph_symmetric_tree.c

igraph_regular_tree — Creates a regular tree.

```
igraph_error_t igraph_regular_tree(igraph_t *graph, igraph_integer_t h, igraph_
```

All vertices of a regular tree, except its leaves, have the same total degree k. This is different from a k-ary tree (igraph_kary_tree()), where all vertices have the same number of children, thus the degree of the root is one less than the degree of the other internal vertices. Regular trees are also referred to as Bethe lattices.

Arguments:

graph: Pointer to an uninitialized graph object.

h: The height of the tree, i.e. the distance between the root and the leaves.

k: The degree of the regular tree.

type: Constant, gives whether to create a directed tree, and if this is the case, also its orientation.

Possible values:

IGRAPH_TREE_OUT directed tree, the edges point from the parents to their

children,

IGRAPH_TREE_IN directed tree, the edges point from the children to their

parents.

IGRAPH_TREE_UNDIRECTED undirected tree.

Returns:

Error code.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

See also:

igraph_kary_tree() to create k-ary tree where each vertex has the same number of children, i.e. out-degree, instead of the same total degree. igraph_symmetric_tree() to use a different number of children at each level.

Example 9.9. File examples/simple/igraph_regular_tree.c

igraph_full — Creates a full graph (directed or undirected, with or without loops).

In a full graph every possible edge is present, every vertex is connected to every other vertex. A full graph in igraph should be distinguished from the concept of complete graphs as used in graph theory. If n is a positive integer, then the complete graph K_n on n vertices is the undirected simple graph with the following property. For any distinct pair (u,v) of vertices in K_n, uv (or equivalently vu) is an edge of K_n. In igraph, a full graph on n vertices can be K_n, a directed version of K_n, or K_n with at least one loop edge. In any case, if F is a full graph on n vertices as generated by igraph, then K_n is a subgraph of the undirected version of F.

Arguments:

graph: Pointer to an uninitialized graph object.

n: Integer, the number of vertices in the graph.

directed: Logical, whether to create a directed graph.

100ps: Logical, whether to include self-edges (loops).

Returns:

Error code: IGRAPH_EINVAL: invalid number of vertices.

Time complexity: O(|V|+|E|), |V| is the number of vertices, |E| the number of edges in the graph. Of course this is the same as O(|E|)=O(|V||V|) here.

See also:

igraph_square_lattice(), igraph_star(), igraph_kary_tree() for creating
other regular structures.

Example 9.10. File examples/simple/igraph_full.c

igraph_full_citation — Creates a full citation graph.

This is a directed graph, where every i->j edge is present if and only if j<i. If the directed argument is zero then an undirected graph is created, and it is just a full graph.

Arguments:

graph: Pointer to an uninitialized graph object, the result is stored here.

n: The number of vertices.

directed: Whether to created a directed graph. If zero an undirected graph is created.

Returns:

Error code.

Time complexity: $O(|V|^2)$, as we have many edges.

igraph_full_multipartite — Create a full multipartite graph.

A multipartite graph contains two or more types of vertices and connections are only possible between two vertices of different types. This function creates a complete multipartite graph.

Arguments:

graph: Pointer to an igraph_t object, the graph will be created here.

types: Pointer to an integer vector. If not a null pointer, the type of each vertex will be stored

here.

n: Pointer to an integer vector, the number of vertices of each type.

directed: Boolean, whether to create a directed graph.

mode: A constant that gives the type of connections for directed graphs. If IGRAPH_OUT,

then edges point from vertices of low-index vertices to high-index vertices; if

IGRAPH_IN, then the opposite direction is realized; if IGRAPH_ALL, then mutual edges will be created.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

See also:

igraph_full_bipartite() for full bipartite graphs.

igraph turan — Create a Turán graph.

```
igraph_error_t igraph_turan(igraph_t *graph,
                             igraph_vector_int_t *types,
                             igraph_integer_t n,
                            igraph_integer_t r);
```

Turán graphs are complete multipartite graphs with the property that the sizes of the partitions are as close to equal as possible. This function generates undirected graphs. The null graph is returned when the number of vertices is zero. A complete graph is returned if the number of partitions is greater than the number of vertices.

Arguments:

graph: Pointer to an igraph_t object, the graph will be created here.

Pointer to an integer vector. If not a null pointer, the type (partition index) of each vertex types: will be stored here.

Integer, the number of vertices in the graph. n:

Integer, the number of partitions of the graph, must be positive. r:

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

See also:

igraph_full_multipartite() for full multipartite graphs.

igraph_realize_degree_sequence — Generates a graph with the given degree sequence.

```
igraph_error_t igraph_realize_degree_sequence(
        igraph_t *graph,
        const igraph_vector_int_t *outdeg, const igraph_vector_int_t *indeg,
```

```
igraph_edge_type_sw_t allowed_edge_types,
igraph_realize_degseq_t method);
```

This function generates an undirected graph that realizes a given degree sequence, or a directed graph that realized a given pair of out- and in-degree sequences.

Simple undirected graphs are constructed using the Havel-Hakimi algorithm (undirected case), or the analogous Kleitman-Wang algorithm (directed case). These algorithms work by choosing an arbitrary vertex and connecting all its stubs to other vertices of highest degree. In the directed case, the "highest" (in, out) degree pairs are determined based on lexicographic ordering. This step is repeated until all degrees have been connected up.

Loopless multigraphs are generated using an analogous algorithm: an arbitrary vertex is chosen, and it is connected with a single connection to a highest remaining degee vertex. If self-loops are also allowed, the same algorithm is used, but if a non-zero vertex remains at the end of the procedure, the graph is completed by adding self-loops to it. Thus, the result will contain at most one vertex with self-loops.

The method parameter controls the order in which the vertices to be connected are chosen.

References:

- V. Havel, Poznámka o existenci kone#ných graf# (A remark on the existence of finite graphs), #asopis pro p#stování matematiky 80, 477-480 (1955). http://eudml.org/doc/19050
- S. L. Hakimi, On Realizability of a Set of Integers as Degrees of the Vertices of a Linear Graph, Journal of the SIAM 10, 3 (1962). https://www.jstor.org/stable/2098770
- D. J. Kleitman and D. L. Wang, Algorithms for Constructing Graphs and Digraphs with Given Valences and Factors, Discrete Mathematics 6, 1 (1973). https://doi.org/10.1016/0012-365X %2873%2990037-X
- Sz. Horvát and C. D. Modes, Connectivity matters: Construction and exact random sampling of connected graphs (2020). https://arxiv.org/abs/2009.03747

Arguments:

graph:	Pointer to an uninitialized graph object.		
outdeg:	The degree sequence of an undirected graph (if <i>indeg</i> is NULL), or the out-degree sequence of a directed graph (if <i>indeg</i> is given).		
indeg:	The in-degree sequence of a directed graph. Pass NULL to generate an undirected graph.		
allowed_edge_types:	The types of edges to allow in the graph. For directed graphs, only IGRAPH_SIMPLE_SW is implemented at this moment. For undirected graphs, the following values are valid:		
	IGRAPH_SIMPLE_SW	simple graphs (i.e. no self-loops or multi-edges allowed).	
	IGRAPH_LOOPS_SW	single self-loops are allowed, but not multi-edges; currently not implemented.	
	IGRAPH_MULTI_SW	multi-edges are allowed, but not self-loops.	
	IGRAPH_LOOPS_SW IGRAPH_MULTI_SW	both self-loops and multi-edges are allowed.	

method:

The method to generate the graph. Possible values:

IGRAPH_REALIZE_DEGSEQ_S-

MALLEST

The vertex with smallest remaining degree is selected first. The result is usually a graph with high negative degree assortativity. In the undirected case, this method is guaranteed to generate a connected graph, regardless of whether multi-edges are allowed, provided that a connected realization exists (see Horvát and Modes, 2020, as well as http://szhorvat.net/pelican/hh-

connected-graphs.html). In the directed case it tends to generate weakly connected graphs, but this is not guaranteed.

IGRAPH REALIZE DEGSE-Q_LARGEST

IGRAPH_REALIZE_DEGSE-

The vertex with the largest remaining degree is selected first. The result is usually a graph with high positive degree assortativity, and is often disconnected.

Q_INDEX

The vertices are selected in order of their index (i.e. their position in the degree vector). Note that sorting the degree vector and using the INDEX method is not equivalent to the SMALLEST method above, as SMALLEST uses the smallest remaining degree for selecting vertices, not the smallest initial degree.

Returns:

Error code:

IGRAPH_UNIMPLEMENTED The requested method is not implemented.

There is not enough memory to perform the operation. IGRAPH_ENOMEM

Invalid method parameter, or invalid in- and/or out-degree vectors. IGRAPH_EINVAL

The degree vectors should be non-negative, the length and sum of

outdeg and indeg should match for directed graphs.

See also:

igraph_is_graphical() to test graphicality without generating a graph; igraph_degree_sequence_game() to generate random graphs with a given degree sequence; igraph_k_regular_game() to generate random regular graphs; igraph_rewire() to randomly rewire the edges of a graph while preserving its degree sequence.

Example 9.11. File examples/simple/igraph_realize_degree_sequence.c

igraph_famous — Create a famous graph by simply providing its name.

igraph_error_t igraph_famous(igraph_t *graph, const char *name);

The name of the graph can be simply supplied as a string. Note that this function creates graphs which don't take any parameters, there are separate functions for graphs with parameters, e.g. igraph_full() for creating a full graph.

The following graphs are supported:

Bull The bull graph, 5 vertices, 5 edges, resembles the head of a bull

if drawn properly.

Chvatal This is the smallest triangle-free graph that is both 4-chromatic

and 4-regular. According to the Grunbaum conjecture there exists an m-regular, m-chromatic graph with n vertices for every m>1 and n>2. The Chvatal graph is an example for m=4 and

n=12. It has 24 edges.

Coxeter A non-Hamiltonian cubic symmetric graph with 28 vertices and

42 edges.

Cubical The Platonic graph of the cube. A convex regular polyhedron

with 8 vertices and 12 edges.

Diamond A graph with 4 vertices and 5 edges, resembles a schematic

diamond if drawn properly.

Dodecahedral, Dodecahe-

dron

Another Platonic solid with 20 vertices and 30 edges.

Folkman The semisymmetric graph with minimum number of vertices,

20 and 40 edges. A semisymmetric graph is regular, edge tran-

sitive and not vertex transitive.

Franklin This is a graph whose embedding to the Klein bottle can be

colored with six colors, it is a counterexample to the necessity of the Heawood conjecture on a Klein bottle. It has 12 vertices

and 18 edges.

Frucht The Frucht Graph is the smallest cubical graph whose automor-

phism group consists only of the identity element. It has 12 ver-

tices and 18 edges.

Grotzsch The Grötzsch graph is a triangle-free graph with 11 vertices, 20

edges, and chromatic number 4. It is named after German mathematician Herbert Grötzsch, and its existence demonstrates that the assumption of planarity is necessary in Grötzsch's theorem

that every triangle-free planar graph is 3-colorable.

Heawood The Heawood graph is an undirected graph with 14 vertices and

21 edges. The graph is cubic, and all cycles in the graph have

six or more edges. Every smaller cubic graph has shorter cycles, so this graph is the 6-cage, the smallest cubic graph of girth 6.

Herschel The Herschel graph is the smallest nonhamiltonian polyhedral

graph. It is the unique such graph on 11 nodes, and has 18 edges.

House The house graph is a 5-vertex, 6-edge graph, the schematic

draw of a house if drawn properly, basically a triangle on top

of a square.

HouseX The same as the house graph with an X in the square. 5 vertices

and 8 edges.

Icosahedral, Icosahedron A Platonic solid with 12 vertices and 30 edges.

Krackhardt_Kite A social network with 10 vertices and 18 edges. Krackhardt,

D. Assessing the Political Landscape: Structure, Cognition, and Power in Organizations. Admin. Sci. Quart. 35, 342-369, 1990.

Levi The graph is a 4-arc transitive cubic graph, it has 30 vertices

and 45 edges.

McGee The McGee graph is the unique 3-regular 7-cage graph, it has

24 vertices and 36 edges.

Meredith The Meredith graph is a quartic graph on 70 nodes and 140

edges that is a counterexample to the conjecture that every 4-

regular 4-connected graph is Hamiltonian.

Noperfectmatching A connected graph with 16 vertices and 27 edges containing

no perfect matching. A matching in a graph is a set of pairwise non-incident edges; that is, no two edges share a common vertex. A perfect matching is a matching which covers all vertices

of the graph.

Nonline A graph whose connected components are the 9 graphs whose

presence as a vertex-induced subgraph in a graph makes a non-

line graph. It has 50 vertices and 72 edges.

Octahedral, Octahedron Platonic solid with 6 vertices and 12 edges.

Petersen A 3-regular graph with 10 vertices and 15 edges. It is the small-

est hypohamiltonian graph, i.e. it is non-hamiltonian but re-

moving any single vertex from it makes it Hamiltonian.

Robertson The unique (4,5)-cage graph, i.e. a 4-regular graph of girth 5.

It has 19 vertices and 38 edges.

Smallestcyclicgroup A smallest nontrivial graph whose automorphism group is

cyclic. It has 9 vertices and 15 edges.

Tetrahedral, Tetrahedron Platonic solid with 4 vertices and 6 edges.

Thomassen The smallest hypotraceable graph, on 34 vertices and 52 edges.

A hypotracable graph does not contain a Hamiltonian path but after removing any single vertex from it the remainder always contains a Hamiltonian path. A graph containing a Hamiltonian

path is called traceable.

Tutte Tait's Hamiltonian graph conjecture states that every 3-con-

nected 3-regular planar graph is Hamiltonian. This graph is a

counterexample. It has 46 vertices and 69 edges.

Uniquely3colorable Returns a 12-vertex, triangle-free graph with chromatic number

3 that is uniquely 3-colorable.

Walther An identity graph with 25 vertices and 31 edges. An identity

graph has a single graph automorphism, the trivial one.

Zachary Social network of friendships between 34 members of a karate

club at a US university in the 1970s. See W. W. Zachary, An information flow model for conflict and fission in small groups, Journal of Anthropological Research 33, 452-473 (1977).

Arguments:

graph: Pointer to an uninitialized graph object.

name: Character constant, the name of the graph to be created, it is case insensitive.

Returns:

Error code, IGRAPH_EINVAL if there is no graph with the given name.

See also:

```
Other functions for creating graph structures: igraph_ring(), igraph_kary_tree(), igraph_square_lattice(), igraph_full().
```

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

igraph_lcf — Creates a graph from LCF notation.

```
igraph_error_t igraph_lcf(igraph_t *graph, igraph_integer_t n, ...);
```

LCF is short for Lederberg-Coxeter-Frucht, it is a concise notation for 3-regular Hamiltonian graphs. It consists of three parameters: the number of vertices in the graph, a list of shifts giving additional edges to a cycle backbone, and another integer giving how many times the shifts should be performed. See http://mathworld.wolfram.com/LCFNotation.html for details.

Arguments:

graph: Pointer to an uninitialized graph object.

n: Integer, the number of vertices in the graph.

...: The shifts and the number of repeats for the shifts, plus an additional 0 to mark the end of the arguments.

Returns:

Error code.

See also:

See igraph_lcf_vector() for a similar function using a vector_t instead of the variable length argument list.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

Example 9.12. File examples/simple/igraph_lcf.c

igraph_lcf_vector — Creates a graph from LCF notation.

This function is essentially the same as igraph_lcf(), only the way for giving the arguments is different. See igraph_lcf() for details.

Arguments:

graph: Pointer to an uninitialized graph object.

n: Integer constant giving the number of vertices.

shifts: A vector giving the shifts.

repeats: An integer constant giving the number of repeats for the shifts.

Returns:

Error code.

See also:

```
igraph_lcf(), igraph_extended_chordal_ring()
```

Time complexity: O(|V|+|E|), linear in the number of vertices plus the number of edges.

igraph_from_prufer — Generates a tree from a Prüfer sequence.

```
igraph_error_t igraph_from_prufer(igraph_t *graph, const igraph_vector_int_t *p.
```

A Prüfer sequence is a unique sequence of integers associated with a labelled tree. A tree on n vertices can be represented by a sequence of n-2 integers, each between 0 and n-1 (inclusive). The algorithm used by this function is based on Paulius Micikevi#ius, Saverio Caminiti, Narsingh Deo: Linear-time Algorithms for Encoding Trees as Sequences of Node Labels

Arguments:

graph: Pointer to an uninitialized graph object.

prufer: The Prüfer sequence

Returns:

Error code:

IGRAPH_ENOMEM there is not enough memory to perform the operation.

IGRAPH_EINVAL invalid Prüfer sequence given

See also:

```
igraph_to_prufer(), igraph_kary_tree(), igraph_tree_game()
```

igraph_atlas — Create a small graph from the "Graph Atlas".

```
igraph_error_t igraph_atlas(igraph_t *graph, igraph_integer_t number);
```

The number of the graph is given as a parameter. The graphs are listed:

- 1. in increasing order of number of nodes;
- 2. for a fixed number of nodes, in increasing order of the number of edges;
- 3. for fixed numbers of nodes and edges, in increasing order of the degree sequence, for example 111223 < 112222;
- 4. for fixed degree sequence, in increasing number of automorphisms.

The data was converted from the NetworkX software package, see http://networkx.github.io .

See An Atlas of Graphs by Ronald C. Read and Robin J. Wilson, Oxford University Press, 1998.

Arguments:

graph: Pointer to an uninitialized graph object.

number: The number of the graph to generate.

Added in version 0.2.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

Example 9.13. File examples/simple/igraph_atlas.c

igraph_de_bruijn — Generate a de Bruijn graph.

```
igraph_error_t igraph_de_bruijn(igraph_t *graph, igraph_integer_t m, igraph_int
```

A de Bruijn graph represents relationships between strings. An alphabet of m letters are used and strings of length n are considered. A vertex corresponds to every possible string and there is a directed edge from vertex v to vertex w if the string of v can be transformed into the string of w by removing its first letter and appending a letter to it.

Please note that the graph will have m to the power n vertices and even more edges, so probably you don't want to supply too big numbers for m and n.

De Bruijn graphs have some interesting properties, please see another source, e.g. Wikipedia for details.

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.

m: Integer, the number of letters in the alphabet.

n: Integer, the length of the strings.

Returns:

Error code.

See also:

```
igraph_kautz().
```

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

igraph_kautz — Generate a Kautz graph.

```
igraph_error_t igraph_kautz(igraph_t *graph, igraph_integer_t m, igraph_integer_
```

A Kautz graph is a labeled graph, vertices are labeled by strings of length n+1 above an alphabet with m+1 letters, with the restriction that every two consecutive letters in the string must be different. There is a directed edge from a vertex v to another vertex w if it is possible to transform the string of v into the string of w by removing the first letter and appending a letter to it. For string length 1 the new letter cannot equal the old letter, so there are no loops.

Kautz graphs have some interesting properties, see e.g. Wikipedia for details.

Vincent Matossian wrote the first version of this function in R, thanks.

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.

m: Integer, m+1 is the number of letters in the alphabet.

n: Integer, n+1 is the length of the strings.

Returns:

Error code.

See also:

```
igraph_de_bruijn().
```

Time complexity: $O(|V|^* [(m+1)/m]^n + |E|)$, in practice it is more like O(|V| + |E|). |V| is the number of vertices, |E| is the number of edges and m and n are the corresponding arguments.

igraph_circulant — Creates a circulant graph.

igraph_error_t igraph_circulant(igraph_t *graph, igraph_integer_t n, const igraph_

A circulant graph G(n, shifts) consists of n vertices $v_0, ..., v_{n-1}$ such that for each s_i in the list of offsets shifts, v_j is connected to $v_i(j + s_i)$ mod n for all j.

The function can generate either directed or undirected graphs. It does not generate multi-edges or self-loops.

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.

n: Integer, the number of vertices in the circulant graph.

shifts: Integer vector, a list of the offsets within the circulant graph.

directed: Boolean, whether to create a directed graph.

Returns:

Error code.

See also:

```
igraph_ring(), igraph_generalized_petersen(), igraph_extend-
ed chordal ring()
```

Time complexity: O(|V||shifts|), the number of vertices in the graph times the number of shifts.

igraph_generalized_petersen — Creates a Generalized Petersen graph.

```
igraph_error_t igraph_generalized_petersen(igraph_t *graph, igraph_integer_t n,
```

The generalized Petersen graph G(n, k) consists of n vertices $v_0, ..., v_n$ forming an "outer" cycle graph, and n additional vertices $u_0, ..., u_n$ forming an "inner" circulant graph where u_i is connected to $u_i + k \mod n$. Additionally, all v_i are connected to u_i .

G(n, k) has 2n vertices and 3n edges. The Petersen graph itself is G(5, 2).

Reference:

M. E. Watkins, A Theorem on Tait Colorings with an Application to the Generalized Petersen Graphs, Journal of Combinatorial Theory 6, 152-164 (1969). https://doi.org/10.1016%2FS0021-9800%2869%2980116-X

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.

n: Integer, n is the number of vertices in the inner and outer cycle/circulant graphs. It must be at least 3.

k: Integer, k is the shift of the circulant graph. It must be positive and less than n/2.

Returns:

Error code.

See also:

igraph_famous() for the original Petersen graph.

Time complexity: O(|V|), the number of vertices in the graph.

igraph_extended_chordal_ring — Create an extended chordal ring.

```
igraph_error_t igraph_extended_chordal_ring(
    igraph_t *graph, igraph_integer_t nodes, const igraph_matrix_int_t *W,
    igraph_bool_t directed);
```

An extended chordal ring is a cycle graph with additional chords connecting its vertices. Each row L of the matrix W specifies a set of chords to be inserted, in the following way: vertex i will connect to a vertex L[(i mod p)] steps ahead of it along the cycle, where p is the length of L. In other words, vertex i will be connected to vertex (i + L[(i mod p)]) mod nodes. If multiple edges are defined in this way, this will output a non-simple graph. The result can be simplified using $igraph_simplify()$.

See also Kotsis, G: Interconnection Topologies for Parallel Processing Systems, PARS Mitteilungen 11, 1-6, 1993. The igraph extended chordal rings are not identical to the ones in the paper. In igraph the matrix specifies which edges to add. In the paper, a condition is specified which should simultaneously hold between two endpoints and the reverse endpoints.

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.

nodes: Integer constant, the number of vertices in the graph. It must be at least 3.

W: The matrix specifying the extra edges. The number of columns should divide the num-

ber of total vertices. The elements are allowed to be negative.

directed: Whether the graph should be directed.

Returns:

Error code.

See also:

```
igraph_ring(), igraph_lcf(), igraph_lcf_vector().
```

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

Games: Randomized graph generators

Games are randomized graph generators. Randomization means that they generate a different graph every time you call them.

igraph_grg_game — Generates a geometric random graph.

A geometric random graph is created by dropping points (i.e. vertices) randomly on the unit square and then connecting all those pairs which are strictly less than radius apart in Euclidean distance.

Original code contributed by Keith Briggs, thanks Keith.

Arguments:

graph: Pointer to an uninitialized graph object.

nodes: The number of vertices in the graph.

radius: The radius within which the vertices will be connected.

torus: Logical constant. If true, periodic boundary conditions will be used, i.e. the vertices are

assumed to be on a torus instead of a square.

x: An initialized vector or NULL. If not NULL, the points' x coordinates will be returned here.

y: An initialized vector or NULL. If not NULL, the points' y coordinates will be returned here.

Returns:

Error code.

Time complexity: TODO, less than $O(|V|^2+|E|)$.

Example 9.14. File examples/simple/igraph_grg_game.c

igraph_barabasi_game — Generates a graph based on the Barabási-Albert model.

Arguments:

graph: An uninitialized graph object.

n: The number of vertices in the graph.

power: Power of the preferential attachment. The probability that a vertex is cited is pro-

portional to d^power+A, where d is its degree (see also the outpref argument), power and A are given by arguments. In the classic preferential attachment model

power=1.

m: The number of outgoing edges generated for each vertex. (Only if outseq is

NULL.)

outseq: Gives the (out-)degrees of the vertices. If this is constant, this can be a NULL

pointer or an empty (but initialized!) vector, in this case *m* contains the constant out-degree. The very first vertex has by definition no outgoing edges, so the first

number in this vector is ignored.

outpref: Boolean, if true not only the in- but also the out-degree of a vertex increases its

citation probability. I.e., the citation probability is determined by the total degree of the vertices. Ignored and assumed to be true if the graph being generated is

undirected.

A: The probability that a vertex is cited is proportional to d^power+A, where d is its

degree (see also the *outpref* argument), power and A are given by arguments. In the previous versions of the function this parameter was implicitly set to one.

directed: Boolean, whether to generate a directed graph.

algo: The algorithm to use to generate the network. Possible values:

IGRAPH_BARABASI_BAG This is the algorithm that was previously

(before version 0.6) solely implemented in igraph. It works by putting the IDs of the vertices into a bag (multiset, really), exactly as many times as their (in-)degree, plus once more. Then the required number of cited vertices are drawn from the bag, with replacement. This method might generate multiple edges. It only works if power=1

and A=1.

IGRAPH_BARABASI_PSUMTREE This algorithm uses a partial prefix-sum tree

to generate the graph. It does not generate multiple edges and works for any power and

A values.

IGRAPH_BARABASI_PSUMTREE_MITTISTIAtgorithm also uses a partial pre-

PLE fix-sum tree to generate the graph. The difference is, that now multiple edges

are allowed. This method was implemented under the name igraph_non-linear_barabasi_game before ver-

sion 0.6.

start_from: Either a null pointer, or a graph. In the former case, the starting configuration is a

clique of size m. In the latter case, the graph is a starting configuration. The graph must be non-empty, i.e. it must have at least one vertex. If a graph is supplied here and the outseq argument is also given, then outseq should only contain

information on the vertices that are not in the start_from graph.

Returns:

Error code: IGRAPH_EINVAL: invalid n, m or outseq parameter.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

Example 9.15. File examples/simple/igraph_barabasi_game.c

Example 9.16. File examples/simple/igraph_barabasi_game2.c

igraph_erdos_renyi_game — Generates a random (Erd#s-Rényi) graph.

Arguments:

graph: Pointer to an uninitialized graph object.

type: The type of the random graph, possible values:

IGRAPH_ERDOS_RENYI_GNM G(n,m) graph, m edges are selected uniformly ran-

domly in a graph with n vertices.

IGRAPH_ERDOS_RENYI_GNP G(n,p) graph, every possible edge is included in

the graph with probability p.

n: The number of vertices in the graph.

 p_or_m : This is the p parameter for G(n,p) graphs and the m parameter for G(n,m) graphs.

directed: Logical, whether to generate a directed graph.

loops: Logical, whether to generate loops (self) edges.

Returns:

Error code: IGRAPH_EINVAL: invalid type, n, p or m parameter. IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

See also:

```
igraph_barabasi_game(),igraph_growing_random_game()
```

Example 9.17. File examples/simple/igraph_erdos_renyi_game.c

igraph_watts_strogatz_game — The Watts-Strogatz small-world model.

This function generates a graph according to the Watts-Strogatz model of small-world networks. The graph is obtained by creating a circular undirected lattice and then rewire the edges randomly with a constant probability.

See also: Duncan J Watts and Steven H Strogatz: Collective dynamics of "small world" networks, Nature 393, 440-442, 1998.

Arguments:

graph: The graph to initialize.

dim: The dimension of the lattice.

size: The size of the lattice along each dimension.

nei: The size of the neighborhood for each vertex. This is the same as the nei argument

of igraph_connect_neighborhood().

p: The rewiring probability. A real number between zero and one (inclusive).

loops: Logical, whether to generate loop edges.

multiple: Logical, whether to allow multiple edges in the generated graph.

Returns:

Error code.

See also:

Time complexity: $O(|V|*d^o+|E|)$, |V| and |E| are the number of vertices and edges, d is the average degree, o is the nei argument.

igraph_rewire_edges — Rewires the edges of a graph with constant probability.

This function rewires the edges of a graph with a constant probability. More precisely each end point of each edge is rewired to a uniformly randomly chosen vertex with constant probability *prob*.

Note that this function modifies the input graph, call igraph_copy() if you want to keep it.

Arguments:

graph: The input graph, this will be rewired, it can be directed or undirected.

prob: The rewiring probability a constant between zero and one (inclusive).

100ps: Boolean, whether loop edges are allowed in the new graph, or not.

multiple: Boolean, whether multiple edges are allowed in the new graph.

Returns:

Error code.

See also:

igraph_watts_strogatz_game() uses this function for the rewiring.

Time complexity: O(|V|+|E|).

igraph_rewire_directed_edges — Rewires the chosen endpoint of directed edges.

This function rewires either the start or end of directed edges in a graph with a constant probability. Correspondingly, either the in-degree sequence or the out-degree sequence of the graph will be preserved.

Note that this function modifies the input graph, call igraph_copy() if you want to keep it.

This function can produce multiple edges between two vertices.

Arguments:

graph: The input graph, this will be rewired, it can be directed or undirected. If it is undirected or

mode is set to IGRAPH_ALL, igraph_rewire_edges() will be called.

prob: The rewiring probability, a constant between zero and one (inclusive).

100ps: Boolean, whether loop edges are allowed in the new graph, or not.

mode: The endpoints of directed edges to rewire. It is ignored for undirected graphs. Possible

values:

IGRAPH_OUT rewire the end of each directed edge

IGRAPH_IN rewire the start of each directed edge

IGRAPH_ALL rewire both endpoints of each edge

Returns:

Error code.

See also:

```
igraph_rewire_edges(), igraph_rewire()
```

Time complexity: O(|E|).

igraph_degree_sequence_game — Generates a random graph with a given degree sequence.

Arguments:

graph: Pointer to an uninitialized graph object.

out_deg: The degree sequence for an undirected graph (if in_seq is NULL or of length zero), or the out-degree sequence of a directed graph (if in_deq is not of length zero).

in_deg: It is either a zero-length vector or NULL (if an undirected graph is generated), or the

in-degree sequence.

method: The method to generate the graph. Possible values:

IGRAPH_DEGSEQ_CONFIGU-RATION

el. For undirected graphs, it puts all vertex IDs in a bag such that the multiplicity of a vertex in the bag is the same as its degree. Then it draws pairs from the bag until the bag becomes empty. This method may generate both loop (self) edges and multiple edges. For directed graphs, the algorithm is basically the same, but two separate bags are used for the in- and out-degrees. Undirected graphs are generated with probability proportional to (\prod_{i<j}</pre> A_{ij} $\prod_i A_{ii} !!)^{-1}$, where A denotes the adjacency matrix and !! denotes the double factorial. Here A is assumed to have twice the number of self-loops on its diagonal. The corresponding expression for directed graphs is $(\prod_{i,j} A_{ij}!)^{-1}$. Thus the probability of all simple graphs (which only have 0s and 1s in the adjacency matrix) is the same, while that of non-simple ones depends on their edge and self-loop multiplicities.

This method implements the configuration mod-

iGRAPH_DEGSEQ_CONFIGU-RATION SIMPLE This method is identical to IGRAPH_DEGSE-Q_CONFIGURATION, but if the generated graph is not simple, it rejects it and re-starts the generation. It generates all simple graphs with the same probability.

IGRAPH_DEGSEQ_FAST_HEUR_SIMPLE

This method generates simple graphs. It is similar to IGRAPH_DEGSEQ_CONFIGURATION but tries to avoid multiple and loop edges and restarts the generation from scratch if it gets stuck. It can generate all simple realizations of a degree sequence, but it is not guaranteed to sample them uniformly. This method is relative-

ly fast and it will eventually succeed if the provided degree sequence is graphical, but there is no upper bound on the number of iterations.

IGRAPH_DEGSEQ_EDGE_SWITCHING_SIMPLE

This is an MCMC sampler based on degree-preserving edge switches. It generates simple undirected or directed graphs. It uses igraph_re-alize_degree_sequence() to construct an initial graph, then rewires it using igraph_rewire().

IGRAPH_DEGSEQ_VL

This method samples undirected *connected* graphs approximately uniformly. It is a Monte Carlo method based on degree-preserving edge switches. This generator should be favoured if undirected and connected graphs are to be generated and execution time is not a concern. igraph uses the original implementation of Fabien Viger; for the algorithm, see https://www-complexnetworks.lip6.fr/~latapy/FV/generation.html and the paper https://arxiv.org/abs/cs/0502085

Returns:

Error code: IGRAPH_ENOMEM: there is not enough memory to perform the operation. IGRAPH_EINVAL: invalid method parameter, or invalid in- and/or out-degree vectors. The degree vectors should be non-negative, out_deg should sum up to an even integer for undirected graphs; the length and sum of out_deg and in_deg should match for directed graphs.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges for IGRAPH_DEGSEQ_SIMPLE. The time complexity of the other modes is not known.

See also:

```
igraph_barabasi_game(),igraph_erdos_renyi_game(),igraph_is_graphi-
cal()
```

Example 9.18. File examples/simple/igraph_degree_sequence_game.c

igraph_k_regular_game — Generates a random graph where each vertex has the same degree.

This game generates a directed or undirected random graph where the degrees of vertices are equal to a predefined constant k. For undirected graphs, at least one of k and the number of vertices must be even.

Currently, this game simply uses igraph_degree_sequence_game with the SIM-PLE_NO_MULTIPLE method and appropriately constructed degree sequences. Thefore, it does not sample uniformly: while it can generate all k-regular graphs with the given number of vertices, it does not generate each one with the same probability.

Arguments:

graph: Pointer to an uninitialized graph object.

no_of_nodes: The number of nodes in the generated graph.

k: The degree of each vertex in an undirected graph, or the out-degree and in-degree

of each vertex in a directed graph.

directed: Whether the generated graph will be directed.

multiple: Whether to allow multiple edges in the generated graph.

Returns:

Error code: IGRAPH_EINVAL: invalid parameter; e.g., negative number of nodes, or odd number of nodes and odd k for undirected graphs. IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: O(|V|+|E|) if multiple is true, otherwise not known.

igraph_static_fitness_game — Non-growing random graph with edge probabilities proportional to node fitness scores.

This game generates a directed or undirected random graph where the probability of an edge between vertices i and j depends on the fitness scores of the two vertices involved. For undirected graphs, each vertex has a single fitness score. For directed graphs, each vertex has an out- and an in-fitness, and the probability of an edge from i to j depends on the out-fitness of vertex i and the in-fitness of vertex j.

The generation process goes as follows. We start from N disconnected nodes (where N is given by the length of the fitness vector). Then we randomly select two vertices i and j, with probabilities proportional to their fitnesses. (When the generated graph is directed, i is selected according to the out-fitnesses and j is selected according to the in-fitnesses). If the vertices are not connected yet (or if multiple edges are allowed), we connect them; otherwise we select a new pair. This is repeated until the desired number of links are created.

It can be shown that the *expected* degree of each vertex will be proportional to its fitness, although the actual, observed degree will not be. If you need to generate a graph with an exact degree sequence, consider igraph_degree_sequence_game instead.

This model is commonly used to generate static scale-free networks. To achieve this, you have to draw the fitness scores from the desired power-law distribution. Alternatively, you may use igraph_static_power_law_game which generates the fitnesses for you with a given exponent.

Reference: Goh K-I, Kahng B, Kim D: Universal behaviour of load distribution in scale-free networks. Phys Rev Lett 87(27):278701, 2001.

Arguments:

graph: Pointer to an uninitialized graph object.

fitness_out: A numeric vector containing the fitness of each vertex. For directed graphs, this

specifies the out-fitness of each vertex.

fitness_in: If NULL, the generated graph will be undirected. If not NULL, this argument spec-

ifies the in-fitness of each vertex.

no_of_edges: The number of edges in the generated graph.

100ps: Whether to allow loop edges in the generated graph.

multiple: Whether to allow multiple edges in the generated graph.

Returns:

Error code: IGRAPH_EINVAL: invalid parameter IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: $O(|V| + |E| \log |E|)$.

igraph_static_power_law_game — Generates a non-growing random graph with expected power-law degree distributions.

This game generates a directed or undirected random graph where the degrees of vertices follow power-law distributions with prescribed exponents. For directed graphs, the exponents of the in- and out-degree distributions may be specified separately.

The game simply uses $igraph_static_fitness_game$ with appropriately constructed fitness vectors. In particular, the fitness of vertex i is i^{-alpha} , where alpha = 1/(gamma-1) and gamma is the exponent given in the arguments.

To remove correlations between in- and out-degrees in case of directed graphs, the in-fitness vector will be shuffled after it has been set up and before igraph_static_fitness_game is called.

Note that significant finite size effects may be observed for exponents smaller than 3 in the original formulation of the game. This function provides an argument that lets you remove the finite size effects by assuming that the fitness of vertex i is $(i+i0-1)^{-alpha}$, where i0 is a constant chosen appropriately to ensure that the maximum degree is less than the square root of the number of edges times the average degree; see the paper of Chung and Lu, and Cho et al for more details.

References:

Goh K-I, Kahng B, Kim D: Universal behaviour of load distribution in scale-free networks. Phys Rev Lett 87(27):278701, 2001.

Chung F and Lu L: Connected components in a random graph with given degree sequences. Annals of Combinatorics 6, 125-145, 2002.

Cho YS, Kim JS, Park J, Kahng B, Kim D: Percolation transitions in scale-free networks under the Achlioptas process. Phys Rev Lett 103:135702, 2009.

Arguments:

graph: Pointer to an uninitialized graph object.

no_of_nodes: The number of nodes in the generated graph.

no_of_edges: The number of edges in the generated graph.

exponent_out: The power law exponent of the degree distribution. For directed

graphs, this specifies the exponent of the out-degree distribution. It must be greater than or equal to 2. If you pass IGRAPH_IN-FINITY here, you will get back an Erdos-Renyi random net-

work.

exponent_in: If negative, the generated graph will be undirected. If greater than

or equal to 2, this argument specifies the exponent of the in-degree distribution. If non-negative but less than 2, an error will be

generated.

100ps: Whether to allow loop edges in the generated graph.

multiple: Whether to allow multiple edges in the generated graph.

finite_size_correction: Whether to use the proposed finite size correction of Cho et al.

Returns:

Error code: IGRAPH_EINVAL: invalid parameter IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: $O(|V| + |E| \log |E|)$.

igraph_forest_fire_game — Generates a network according to the "forest fire game".

The forest fire model intends to reproduce the following network characteristics, observed in real networks:

- Heavy-tailed in-degree distribution.
- · Heavy-tailed out-degree distribution.
- · Communities.
- Densification power-law. The network is densifying in time, according to a power-law rule.
- Shrinking diameter. The diameter of the network decreases in time.

The network is generated in the following way. One vertex is added at a time. This vertex connects to (cites) ambs vertices already present in the network, chosen uniformly random. Now, for each cited vertex v we do the following procedure:

1. We generate two random numbers, x and y, that are geometrically distributed with means p/(1-p) and rp(1-rp). (p is fw_prob , r is bw_factor .) The new vertex cites x outgoing

neighbors and y incoming neighbors of v, from those which are not yet cited by the new vertex. If there are less than x or y such vertices available then we cite all of them.

2. The same procedure is applied to all the newly cited vertices.

See also: Jure Leskovec, Jon Kleinberg and Christos Faloutsos. Graphs over time: densification laws, shrinking diameters and possible explanations. *KDD '05: Proceeding of the eleventh ACM SIGKDD international conference on Knowledge discovery in data mining*, 177--187, 2005.

Note however, that the version of the model in the published paper is incorrect in the sense that it cannot generate the kind of graphs the authors claim. A corrected version is available from http://cs.stanford.edu/people/jure/pubs/powergrowth-tkdd.pdf, our implementation is based on this.

Arguments:

graph: Pointer to an uninitialized graph object.

nodes: The number of vertices in the graph.

fw_prob: The forward burning probability.

bw_factor: The backward burning ratio. The backward burning probability is calculated as

bw_factor * fw_prob.

pambs: The number of ambassador vertices.

directed: Whether to create a directed graph.

Returns:

Error code.

Time complexity: TODO.

igraph_rewire — Randomly rewires a graph while preserving its degree sequence.

```
igraph_error_t igraph_rewire(igraph_t *graph, igraph_integer_t n, igraph_rewiri)
```

This function generates a new graph based on the original one by randomly rewiring edges while preserving the original graph's degree sequence. The rewiring is done "in place", so no new graph will be allocated. If you would like to keep the original graph intact, use <code>igraph_copy()</code> beforehand. All graph attributes will be lost.

Arguments:

graph: The graph object to be rewired.

n: Number of rewiring trials to perform.

mode: The rewiring algorithm to be used. It can be one of the following flags:

IGRAPH_REWIRING_SIMPLE Simple rewiring algorithm which chooses two arbi-

trary edges in each step (namely (a,b) and (c,d)) and substitutes them with (a,d) and (c,b) if they don't exist. The method will neither destroy nor create self-

loops.

IGRAPH_REWIRING_SIM-PLE_LOOPS Same as IGRAPH_REWIRING_SIMPLE but allows the creation or destruction of self-loops.

Returns:

Error code:

IGRAPH_EINVMODE Invalid rewiring mode.

IGRAPH_EINVAL Graph unsuitable for rewiring (e.g. it has less than 4 nodes in case of

IGRAPH_REWIRING_SIMPLE)

IGRAPH_ENOMEM Not enough memory for temporary data.

Time complexity: TODO.

igraph_growing_random_game — Generates a growing random graph.

This function simulates a growing random graph. We start out with one vertex. In each step a new vertex is added and a number of new edges are also added. These graphs are known to be different from standard (not growing) random graphs.

Arguments:

graph: Uninitialized graph object.

n: The number of vertices in the graph.

m: The number of edges to add in a time step (i.e. after adding a vertex).

directed: Boolean, whether to generate a directed graph.

citation: Boolean, if true, the edges always originate from the most recently added vertex and

are connected to a previous vertex.

Returns:

Error code: IGRAPH_EINVAL: invalid *n* or *m* parameter.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

igraph_callaway_traits_game — Simulates a growing network with vertex types.

```
const igraph_vector_t *type_dist,
const igraph_matrix_t *pref_matrix,
igraph_bool_t directed,
igraph_vector_int_t *node_type_vec);
```

The different types of vertices prefer to connect other types of vertices with a given probability.

The simulation goes like this: in each discrete time step a new vertex is added to the graph. The type of this vertex is generated based on $type_dist$. Then two vertices are selected uniformly randomly from the graph. The probability that they will be connected depends on the types of these vertices and is taken from $pref_matrix$. Then another two vertices are selected and this is repeated $edges_per_step$ times in each time step.

References:

D. S. Callaway, J. E. Hopcroft, J. M. Kleinberg, M. E. J. Newman, and S. H. Strogatz, Are randomly grown graphs really random? Phys. Rev. E 64, 041902 (2001). https://doi.org/10.1103/Phys-RevE.64.041902

Arguments:

graph: Pointer to an uninitialized graph.

nodes: The number of nodes in the graph.

types: Number of node types.

edges_per_step: The number of connections tried in each time step.

type_dist: Vector giving the distribution of the vertex types. If NULL, the distribution

is assumed to be uniform.

pref_matrix:
Matrix giving the connection probabilities for the vertex types.

directed: Logical, whether to generate a directed graph.

node_type_vec: An initialized vector or NULL. If not NULL, the type of each node will be

stored here.

Returns:

Error code.

Added in version 0.2.

Time complexity: O(|V|*k*log(|V|)), |V| is the number of vertices, k is $edges_per_step$.

igraph_establishment_game — Generates a graph with a simple growing model with vertex types.

```
igraph_vector_int_t *node_type_vec);
```

The simulation goes like this: a single vertex is added at each time step. This new vertex tries to connect to *k* vertices in the graph. The probability that such a connection is realized depends on the types of the vertices involved.

Arguments:

graph: Pointer to an uninitialized graph.

nodes: The number of vertices in the graph.

types: The number of vertex types.

k: The number of connections tried in each time step.

type_dist: Vector giving the distribution of vertex types. If NULL, the distribution is as-

sumed to be uniform.

pref_matrix: Matrix giving the connection probabilities for different vertex types.

directed: Logical, whether to generate a directed graph.

node type vec: An initialized vector or NULL. If not NULL, the type of each node will be

stored here.

Returns:

Error code.

Added in version 0.2.

Time complexity: O(|V|*k*log(|V|)), |V| is the number of vertices and k is the k parameter.

igraph_preference_game — Generates a graph with vertex types and connection preferences.

igraph_bool_t loops);

This is practically the nongrowing variant of <code>igraph_establishment_game()</code>. A given number of vertices are generated. Every vertex is assigned to a vertex type according to the given type probabilities. Finally, every vertex pair is evaluated and an edge is created between them with a probability depending on the types of the vertices involved.

In other words, this function generates a graph according to a block-model. Vertices are divided into groups (or blocks), and the probability the two vertices are connected depends on their groups only.

Arguments:

graph: Pointer to an uninitialized graph.

nodes: The number of vertices in the graph.

types: The number of vertex types.

type_dist: Vector giving the distribution of vertex types. If NULL, all vertex types will

have equal probability. See also the fixed_sizes argument.

fixed_sizes: Boolean. If true, then the number of vertices with a given vertex type is fixed

and the type_dist argument gives these numbers for each vertex type. If true, and type_dist is NULL, then the function tries to make vertex groups of the same size. If this is not possible, then some groups will have an extra

vertex.

pref_matrix: Matrix giving the connection probabilities for different vertex types. This

should be symmetric if the requested graph is undirected.

node_type_vec: A vector where the individual generated vertex types will be stored. If NULL,

the vertex types won't be saved.

directed: Logical, whether to generate a directed graph. If undirected graphs are request-

ed, only the lower left triangle of the preference matrix is considered.

loops: Logical, whether loop edges are allowed.

Returns:

Error code.

Added in version 0.3.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

See also:

igraph_asymmetric_preference_game — Generates a graph with asymmetric vertex types and connection preferences.

igraph_bool_t loops);

This is the asymmetric variant of igraph_preference_game(). A given number of vertices are generated. Every vertex is assigned to an "outgoing" and an "incoming" vertex type according

to the given joint type probabilities. Finally, every vertex pair is evaluated and a directed edge is created between them with a probability depending on the "outgoing" type of the source vertex and the "incoming" type of the target vertex.

Arguments:

graph: Pointer to an uninitialized graph.

nodes: The number of vertices in the graph.

no_out_types: The number of vertex out-types.

no_in_types: The number of vertex in-types.

type_dist_matrix: Matrix of size out_types * in_types, giving the joint distribu-

tion of vertex types. If NULL, incoming and outgoing vertex types are

independent and uniformly distributed.

pref_matrix: Matrix of size out_types * in_types, giving the connection prob-

abilities for different vertex types.

node_type_out_vec: A vector where the individual generated "outgoing" vertex types will be

stored. If NULL, the vertex types won't be saved.

node_type_in_vec: A vector where the individual generated "incoming" vertex types will be

stored. If NULL, the vertex types won't be saved.

Logical, whether loop edges are allowed.

Returns:

Error code.

Added in version 0.3.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

See also:

```
igraph_preference_game()
```

igraph_recent_degree_game — Stochastic graph generator based on the number of incident edges a node has gained recently.

igraph_bool_t directed);

Arguments:

graph: Pointer to an uninitialized graph object.

nodes: The number of vertices in the graph, this is the same as the number of time steps.

power: The exponent, the probability that a node gains a new edge is proportional to the

number of edges it has gained recently (in the last window time steps) to power.

time_window: Integer constant, the size of the time window to use to count the number of recent

edges.

m: Integer constant, the number of edges to add per time step if the outseq para-

meter is a null pointer or a zero-length vector.

outseq: The number of edges to add in each time step. This argument is ignored if it is a

null pointer or a zero length vector. In this case the constant *m* parameter is used.

outpref: Logical constant, if true the edges originated by a vertex also count as recent

incident edges. For most applications it is reasonable to set it to false.

zero_appeal: Constant giving the attractiveness of the vertices which haven't gained any edge

recently.

directed: Logical constant, whether to generate a directed graph.

Returns:

Error code.

Time complexity: O(|V|*log(|V|)+|E|), |V| is the number of vertices, |E| is the number of edges in the graph.

igraph_barabasi_aging_game — Preferential attachment with aging of vertices.

This game starts with one vertex (if nodes > 0). In each step a new node is added, and it is connected to m existing nodes. Existing nodes to connect to are chosen with probability dependent on their (in-)degree (k) and age (1). The degree-dependent part is $deg_coef * k^pa_exp + zero_deg_ap-peal$, while the age-dependent part is $age_coef * l^aging_exp + zero_age_appeal$, which are summed to obtain the final weight.

The age 1 is based on the number of vertices in the network and the aging_bins argument: the age of a node is incremented by 1 after each floor(nodes / aging_bins) + 1 time steps.

Arguments:

graph: Pointer to an uninitialized graph object.

nodes: The number of vertices in the graph.

m: The number of edges to add in each time step. Ignored if outseq is a non-

zero length vector.

outseq: The number of edges to add in each time step. If it is NULL or a zero-length

vector then it is ignored and the *m* argument is used instead.

outpref: Logical constant, whether the edges initiated by a vertex contribute to the

probability to gain a new edge.

pa_exp: The exponent of the preferential attachment, a small positive number usu-

ally, the value 1 yields the classic linear preferential attachment.

aging_exp: The exponent of the aging, this is a negative number usually.

aging_bins: Integer constant, the number of age bins to use.

zero_deg_appea1: The degree dependent part of the attractiveness of the zero degree vertices.

zero_age_appea1: The age dependent part of the attractiveness of the vertices of age zero. This

parameter is usually zero.

deg_coef: The coefficient for the degree.

age_coef: The coefficient for the age.

directed: Logical constant, whether to generate a directed graph.

Returns:

Error code.

Time complexity: $O((|V|+|V|/aging_bins)*log(|V|)+|E|)$. |V| is the number of vertices, |E| the number of edges.

igraph_recent_degree_aging_game — Preferential attachment based on the number of edges gained recently, with aging of vertices.

This game is very similar to <code>igraph_barabasi_aging_game()</code>, except that instead of the total number of incident edges the number of edges gained in the last <code>time_window</code> time steps are counted.

The degree dependent part of the attractiveness is given by k to the power of pa_exp plus zero_ap-peal; the age dependent part is 1 to the power to aging_exp. k is the number of edges gained in the last time_window time steps, 1 is the age of the vertex.

Arguments:

graph: Pointer to an uninitialized graph object.

nodes: The number of vertices in the graph.

m: The number of edges to add in each time step. If the *outseq* argument is not a

null vector or a zero-length vector then it is ignored.

outseq: Vector giving the number of edges to add in each time step. If it is a null pointer

or a zero-length vector then it is ignored and the *m* argument is used.

outpref: Logical constant, if true the edges initiated by a vertex are also counted. Normally

it is false.

pa_exp: The exponent for the preferential attachment.

aging_exp: The exponent for the aging, normally it is negative: old vertices gain edges with

less probability.

aging_bins: Integer constant, the number of age bins to use.

time_window: The time window to use to count the number of incident edges for the vertices.

zero_appeal: The degree dependent part of the attractiveness for zero degree vertices.

directed: Logical constant, whether to create a directed graph.

Returns:

Error code.

Time complexity: $O((|V|+|V|/aging_bins)*log(|V|)+|E|)$. |V| is the number of vertices, |E| the number of edges.

igraph_lastcit_game — Simulates a citation network, based on time passed since the last citation.

This is a quite special stochastic graph generator, it models an evolving graph. In each time step a single vertex is added to the network and it cites a number of other vertices (as specified by the <code>edges_per_step</code> argument). The cited vertices are selected based on the last time they were cited. Time is measured by the addition of vertices and it is binned into <code>agebins</code> bins. So if the current time step is t and the last citation to a given i vertex was made in time step t0, then \c (t-t0)/binwidth

is calculated where binwidth is nodes/agebins+1, in the last expression '/' denotes integer division, so the fraction part is omitted.

The *preference* argument specifies the preferences for the citation lags, i.e. its first elements contains the attractivity of the very recently cited vertices, etc. The last element is special, it contains the attractivity of the vertices which were never cited. This element should be bigger than zero.

Note that this function generates networks with multiple edges if edges_per_step is bigger than one, call igraph_simplify() on the result to get rid of these edges.

Arguments:

graph: Pointer to an uninitialized graph object, the result will be stored here.

node: The number of vertices in the network.

edges_per_node: The number of edges to add in each time step.

agebins: The number of age bins to use.

preference: Pointer to an initialized vector of length agebins+1. This contains the `at-

tractivity' of the various age bins, the last element is the attractivity of the vertices which were never cited, and it should be greater than zero. It is a good idea to have all positive values in this vector. Preferences cannot be negative.

directed: Logical constant, whether to create directed networks.

Returns:

Error code.

See also:

```
igraph_barabasi_aging_game().
```

Time complexity: O(|V|*a+|E|*log|V|), |V| is the number of vertices, |E| is the total number of edges, a is the *agebins* parameter.

igraph_cited_type_game — Simulates a citation based on vertex types.

Function to create a network based on some vertex categories. This function creates a citation network: in each step a single vertex and <code>edges_per_step</code> citing edges are added. Nodes with different categories may have different probabilities to get cited, as given by the <code>pref</code> vector.

Note that this function might generate networks with multiple edges if <code>edges_per_step</code> is greater than one. You might want to call <code>igraph_simplify()</code> on the result to remove multiple edges.

Arguments:

graph: Pointer to an uninitialized graph object.

nodes: The number of vertices in the network.

types: Numeric vector giving the categories of the vertices, so it should contain

nodes non-negative integer numbers. Types are numbered from zero.

pref: The attractivity of the different vertex categories in a vector. Its length should

be the maximum element in types plus one (types are numbered from zero).

edges_per_step: Integer constant, the number of edges to add in each time step.

directed: Logical constant, whether to create a directed network.

Returns:

Error code.

See also:

 $\verb|igraph_citing_cited_type_game|() | for a bit more general game.$

Time complexity: $O((|V|+|E|)\log|V|)$, |V| and |E| are number of vertices and edges, respectively.

igraph_citing_cited_type_game — Simulates a citation network based on vertex types.

This game is similar to igraph_cited_type_game() but here the category of the citing vertex is also considered.

An evolving citation network is modeled here, a single vertex and its <code>edges_per_step</code> citation are added in each time step. The odds the a given vertex is cited by the new vertex depends on the category of both the citing and the cited vertex and is given in the <code>pref</code> matrix. The categories of the citing vertex correspond to the rows, the categories of the cited vertex to the columns of this matrix. I.e. the element in row <code>i</code> and column <code>j</code> gives the probability that a <code>j</code> vertex is cited, if the category of the citing vertex is <code>i</code>.

Note that this function might generate networks with multiple edges if <code>edges_per_step</code> is greater than one. You might want to call <code>igraph_simplify()</code> on the result to remove multiple edges.

Arguments:

graph: Pointer to an uninitialized graph object.

nodes: The number of vertices in the network.

types: A numeric vector of length nodes, containing the categories of the vertices.

The categories are numbered from zero.

pref: The preference matrix, a square matrix is required, both the number of rows

and columns should be the maximum element in types plus one (types are

numbered from zero).

edges_per_step: Integer constant, the number of edges to add in each time step.

directed: Logical constant, whether to create a directed network.

Returns:

Error code.

Time complexity: $O((|V|+|E|)\log|V|)$, |V| and |E| are number of vertices and edges, respectively.

igraph_sbm_game — Sample from a stochastic block model.

This function samples graphs from a stochastic block model by (doing the equivalent of) Bernoulli trials for each potential edge with the probabilities given by the Bernoulli rate matrix, pref_ma-trix. See Faust, K., & Wasserman, S. (1992a). Blockmodels: Interpretation and evaluation. Social Networks, 14, 5—61.

The order of the vertex IDs in the generated graph corresponds to the block_sizes argument.

Arguments:

graph: The output graph. This should be a pointer to an uninitialized graph.

n: Number of vertices.

pref_matrix: The matrix giving the Bernoulli rates. This is a KxK matrix, where K is the num-

ber of groups. The probability of creating an edge between vertices from groups

i and j is given by element (i,j).

block sizes: An integer vector giving the number of vertices in each group.

directed: Boolean, whether to create a directed graph. If this argument is false, then pre-

f matrix must be symmetric.

loops: Boolean, whether to create self-loops.

Returns:

Error code.

Time complexity: $O(|V|+|E|+K^2)$, where |V| is the number of vertices, |E| is the number of edges, and K is the number of groups.

See also:

igraph_erdos_renyi_game() for a simple Bernoulli graph.

igraph_hsbm_game — Hierarchical stochastic block model.

The function generates a random graph according to the hierarchical stochastic block model.

Arguments:

graph: The generated graph is stored here.

n: The number of vertices in the graph.

m: The number of vertices per block. n/m must be integer.

rho: The fraction of vertices per cluster, within a block. Must sum up to 1, and rho * m must

be integer for all elements of rho.

C: A square, symmetric numeric matrix, the Bernoulli rates for the clusters within a block.

Its size must mach the size of the rho vector.

p: The Bernoulli rate of connections between vertices in different blocks.

Returns:

Error code.

See also:

igraph_sbm_game() for the classic stochastic block model, igraph_hsbm_list_game()
for a more general version.

igraph_hsbm_list_game — Hierarchical stochastic block model, more general version.

The function generates a random graph according to the hierarchical stochastic block model.

Arguments:

graph: The generated graph is stored here.

n: The number of vertices in the graph.

mlist: An integer vector of block sizes.

rholist: A list of rho vectors (igraph_vector_t objects), one for each block.

Clist: A list of square matrices (igraph_matrix_t objects), one for each block, specifying

the Bernoulli rates of connections within the block.

p: The Bernoulli rate of connections between vertices in different blocks.

Returns:

Error code.

See also:

igraph_sbm_game() for the classic stochastic block model, igraph_hsbm_game() for a simpler general version.

igraph_dot_product_game — Generates a random dot product graph.

In this model, each vertex is represented by a latent position vector. Probability of an edge between two vertices are given by the dot product of their latent position vectors.

See also Christine Leigh Myers Nickel: Random dot product graphs, a model for social networks. Dissertation, Johns Hopkins University, Maryland, USA, 2006.

Arguments:

graph: The output graph is stored here.

vecs: A matrix in which each latent position vector is a column. The dot product of the

latent position vectors should be in the [0,1] interval, otherwise a warning is given. For negative dot products, no edges are added; dot products that are larger than one

always add an edge.

directed: Should the generated graph be directed?

Returns:

Error code.

Time complexity: O(n*n*m), where n is the number of vertices, and m is the length of the latent vectors.

See also:

igraph_tree_game — Generates a random tree with the given number of nodes.

```
igraph_error_t igraph_tree_game(igraph_t *graph, igraph_integer_t n, igraph_boo
```

This function samples uniformly from the set of labelled trees, i.e. it generates each labelled tree with the same probability.

Note that for n=0, the null graph is returned, which is not considered to be a tree by $igraph_is_tree()$.

Arguments:

graph: Pointer to an uninitialized graph object.

n: The number of nodes in the tree.

directed: Whether to create a directed tree. The edges are oriented away from the root.

method: The algorithm to use to generate the tree. Possible values:

IGRAPH_RAN- This algorithm samples Prüfer sequences uni-DOM_TREE_PRUFER formly, then converts them to trees. Directed

trees are not currently supported.

IGRAPH_RANDOM_LERW This algorithm effectively performs a loop-

erased random walk on the complete graph to uniformly sample its spanning trees (Wilson's

algorithm).

Returns:

Error code: IGRAPH_ENOMEM: there is not enough memory to perform the operation. IGRAPH_EINVAL: invalid tree size

See also:

igraph_from_prufer()

igraph_correlated_game — Generates a random graph correlated to an existing graph.

Sample a new graph by perturbing the adjacency matrix of a given simple graph and shuffling its vertices.

Arguments:

old_graph: The original graph, it must be simple.

new_graph: The new graph will be stored here.

corr: A scalar in the unit interval [0,1], the target Pearson correlation between the ad-

jacency matrices of the original and the generated graph (the adjacency matrix

being used as a vector).

p: A numeric scalar, the probability of an edge between two vertices, it must in the

open (0,1) interval. Typically, the density of old_graph.

permutation: A permutation to apply to the vertices of the generated graph. It can also be a null

pointer, in which case the vertices will not be permuted.

Returns:

Error code

See also:

igraph_correlated_pair_game() for generating a pair of correlated random graphs in one go.

igraph_correlated_pair_game — Generates pairs of correlated random graphs.

Sample two random graphs, with given correlation.

Arguments:

graph1: The first graph will be stored here.

graph2: The second graph will be stored here.

n: The number of vertices in both graphs.

corr: A scalar in the unit interval, the target Pearson correlation between the adjacency

matrices of the original the generated graph (the adjacency matrix being used as

a vector).

p: A numeric scalar, the probability of an edge between two vertices, it must in the

open (0,1) interval.

directed: Whether to generate directed graphs.

permutation: A permutation to apply to the vertices of the second graph. It can also be a null

pointer, in which case the vertices will not be permuted.

Returns:

Error code

See also:

igraph_correlated_game() for generating a correlated pair to a given graph.

igraph_simple_interconnected_islands_game — Generates a random graph made of several interconnected islands, each island being a random graph.

```
igraph_error_t igraph_simple_interconnected_islands_game(
    igraph_t *graph,
    igraph_integer_t islands_n,
    igraph_integer_t islands_size,
    igraph_real_t islands_pin,
    igraph_integer_t n_inter);
```

All islands are of the same size. Within an island, each edge is generated with the same probability. A fixed number of additional edges are then generated for each unordered pair of islands to connect them. The generated graph is guaranteed to be simple.

Arguments:

graph: Pointer to an uninitialized graph object.

islands_n: The number of islands in the graph.

islands_size: The size of islands in the graph.

islands_pin: The probability to create each possible edge within islands.

n_inter: The number of edges to create between two islands. It may be larger than is-

lands_size squared, but in this case it is assumed to be islands_size

squared.

Returns:

Error code: IGRAPH_EINVAL: invalid parameter IGRAPH_ENOMEM: there is not enough memory for the operation.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

Deprecated functions

igraph_lattice — Arbitrary dimensional square lattices (deprecated).

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_square_lattice() instead.

igraph_tree — Creates a k-ary tree in which almost all vertices have k children (deprecated alias).

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use $igraph_kary_tree()$ instead.

Chapter 10. Games on graphs

Microscopic update rules

igraph_deterministic_optimal_imitation — Adopt a strategy via deterministic optimal imitation.

```
igraph_error_t igraph_deterministic_optimal_imitation(const igraph_t *graph,
        igraph_integer_t vid,
        igraph_optimal_t optimality,
        const igraph_vector_t *quantities,
        igraph_vector_int_t *strategies,
        igraph_neimode_t mode);
```

A simple deterministic imitation strategy where a vertex revises its strategy to that which yields a local optimum. Here "local" is with respect to the immediate neighbours of the vertex. The vertex retains its current strategy where this strategy yields a locally optimal quantity. The quantity in this case could be a measure such as fitness.

Arguments:

graph:

The graph object representing the game network. This cannot be the empty or trivial graph, but must have at least two vertices and one edge. If graph has one vertex, then no strategy update would take place. Furthermore, if graph has at least two vertices but zero edges, then strategy update would also not take place.

vid:

The vertex whose strategy is to be updated. It is assumed that vid represents a vertex in graph. No checking is performed and it is your responsibility to ensure that vid is indeed a vertex of graph. If an isolated vertex is provided, i.e. the input vertex has degree 0, then no strategy update would take place and vid would retain its current strategy. Strategy update would also not take place if the local neighbourhood of vid are its in-neighbours (respectively out-neighbours), but vid has zero in-neighbours (respectively out-neighbours). Loops are ignored in computing the degree (in, out, all) of vid.

optimality:

Logical; controls the type of optimality to be used. Supported values are:

IGRAPH_MAXIMUM Use maximum deterministic imitation, where the strategy of the vertex with maximum quantity (e.g. fitness) would be adopted. We update the strategy of vid to that which yields a local maximum.

IGRAPH_MINIMUM

Use minimum deterministic imitation. That is, the strategy of the vertex with minimum quantity would be imitated. In other words, update to the strategy that yields a local minimum.

quantities:

A vector of quantities providing the quantity of each vertex in graph. Think of each entry of the vector as being generated by a function such as the fitness function for the game. So if the vector represents fitness quantities, then each vector entry is the fitness of some vertex. The length of this vector must be the same as the number of vertices in the vertex set of graph.

strategies:

A vector of the current strategies for the vertex population. The updated strategy for vid would be stored here. Each strategy is identified with a nonnegative integer,

whose interpretation depends on the payoff matrix of the game. Generally we use the strategy ID as a row or column index of the payoff matrix. The length of this vector must be the same as the number of vertices in the vertex set of *graph*.

mode:

Defines the sort of neighbourhood to consider for *vid*. If *graph* is undirected, then we use all the immediate neighbours of *vid*. Thus if you know that *graph* is undirected, then it is safe to pass the value *IGRAPH_ALL* here. Supported values are:

IGRAPH_OUT Use the out-neighbours of *vid*. This option is only relevant when *graph* is a directed graph.

IGRAPH_IN Use the in-neighbours of vid. Again this option is only relevant when graph is a directed graph.

IGRAPH_ALL Use both the in- and out-neighbours of vid. This option is only

relevant if *graph* is a digraph. Also use this value if *graph* is undirected.

Returns:

The error code *IGRAPH_EINVAL* is returned in each of the following cases: (1) Any of the parameters *graph*, *quantities*, or *strategies* is a null pointer. (2) The vector *quantities* or *strategies* has a length different from the number of vertices in *graph*. (3) The parameter *graph* is the empty or null graph, i.e. the graph with zero vertices and edges.

Time complexity: O(2d), where d is the degree of the vertex vid.

Example 10.1. File examples/simple/igraph_deterministic_optimal_imitation.c

igraph_moran_process — The Moran process in a network setting.

This is an extension of the classic Moran process to a network setting. The Moran process is a model of haploid (asexual) reproduction within a population having a fixed size. In the network setting, the Moran process operates on a weighted graph. At each time step a vertex a is chosen for reproduction and another vertex b is chosen for death. Vertex a gives birth to an identical clone c, which replaces b. Vertex c is a clone of a in that c inherits both the current quantity (e.g. fitness) and current strategy of a.

The graph G representing the game network is assumed to be simple, i.e. free of loops and without multiple edges. If, on the other hand, G has a loop incident on some vertex v, then it is possible that when v is chosen for reproduction it would forgo this opportunity. In particular, when v is chosen for reproduction and v is also chosen for death, the clone of v would be v itself with its current vertex ID. In effect v forgoes its chance for reproduction.

Arguments:

graph:

The graph object representing the game network. This cannot be the empty or trivial graph, but must have at least two vertices and one edge. The Moran process will

not take place in each of the following cases: (1) If *graph* has one vertex. (2) If *graph* has at least two vertices but zero edges.

weights:

A vector of all edge weights for *graph*. Thus weights[i] means the weight of the edge with edge ID i. For the purpose of the Moran process, each weight is assumed to be positive; it is your responsibility to ensure this condition holds. The length of this vector must be the same as the number of edges in *graph*.

quantities:

A vector of quantities providing the quantity of each vertex in <code>graph</code>. The quantity of the new clone will be stored here. Think of each entry of the vector as being generated by a function such as the fitness function for the game. So if the vector represents fitness quantities, then each vector entry is the fitness of some vertex. The length of this vector must be the same as the number of vertices in the vertex set of <code>graph</code>. For the purpose of the Moran process, each vector entry is assumed to be nonnegative; no checks will be performed for this. It is your responsibility to ensure that at least one entry is positive. Furthermore, this vector cannot be a vector of zeros; this condition will be checked.

strategies:

A vector of the current strategies for the vertex population. The strategy of the new clone will be stored here. Each strategy is identified with a nonnegative integer, whose interpretation depends on the payoff matrix of the game. Generally we use the strategy ID as a row or column index of the payoff matrix. The length of this vector must be the same as the number of vertices in the vertex set of graph.

mode:

Defines the sort of neighbourhood to consider for the vertex a chosen for reproduction. This is only relevant if <code>graph</code> is directed. If <code>graph</code> is undirected, then it is safe to pass the value <code>IGRAPH_ALL</code> here. Supported values are:

IGRAPH_OUT Use the out-neighbours of a. This option is only relevant when graph is directed.

IGRAPH_IN Use the in-neighbours of a. Again this option is only relevant when graph is directed.

Use both the in- and out-neighbours of a. This option is only relevant if *graph* is directed. Also use this value if *graph* is undirected.

Returns:

The error code *IGRAPH_EINVAL* is returned in each of the following cases: (1) Any of the parameters *graph*, *weights*, *quantities* or *strategies* is a null pointer. (2) The vector *quantities* or *strategies* has a length different from the number of vertices in *graph*. (3) The vector *weights* has a length different from the number of edges in *graph*. (4) The parameter *graph* is the empty or null graph, i.e. the graph with zero vertices and edges. (5) The vector *weights*, or the combination of interest, sums to zero. (6) The vector *quantities*, or the combination of interest, sums to zero.

Time complexity: depends on the random number generator, but is usually O(n) where n is the number of vertices in graph.

References:

(Lieberman et al. 2005) E. Lieberman, C. Hauert, and M. A. Nowak. Evolutionary dynamics on graphs. *Nature*, 433(7023):312--316, 2005.

(Moran 1958) P. A. P. Moran. Random processes in genetics. *Mathematical Proceedings of the Cambridge Philosophical Society*, 54(1):60--71, 1958.

igraph_roulette_wheel_imitation — Adopt a strategy via roulette wheel selection.

A simple stochastic imitation strategy where a vertex revises its strategy to that of a vertex u chosen proportionate to u's quantity (e.g. fitness). This is a special case of stochastic imitation, where a candidate is not chosen uniformly at random but proportionate to its quantity.

Arguments:

graph: The graph object representing the game network. This cannot be the empty or trivial

graph, but must have at least two vertices and one edge. If graph has one vertex, then no strategy update would take place. Furthermore, if graph has at least two

vertices but zero edges, then strategy update would also not take place.

vid: The vertex whose strategy is to be updated. It is assumed that vid represents a

vertex in graph. No checking is performed and it is your responsibility to ensure that vid is indeed a vertex of graph. If an isolated vertex is provided, i.e. the input vertex has degree 0, then no strategy update would take place and vid would retain its current strategy. Strategy update would also not take place if the local neighbourhood of vid are its in-neighbours (respectively out-neighbours), but vid has zero in-neighbours (respectively out-neighbours). Loops are ignored in computing

the degree (in, out, all) of vid.

islocal: Boolean; this flag controls which perspective to use in computing the relative quan-

tity. If true then we use the local perspective; otherwise we use the global perspective. The local perspective for vid is the set of all immediate neighbours of vid.

In contrast, the global perspective for vid is the vertex set of graph.

quantities: A vector of quantities providing the quantity of each vertex in graph. Think of

each entry of the vector as being generated by a function such as the fitness function for the game. So if the vector represents fitness quantities, then each vector entry is the fitness of some vertex. The length of this vector must be the same as the number of vertices in the vertex set of graph. For the purpose of roulette wheel selection, each vector entry is assumed to be nonnegative; no checks will be performed for this. It is your responsibility to ensure that at least one entry is nonzero. Further-

more, this vector cannot be a vector of zeros; this condition will be checked.

strategies: A vector of the current strategies for the vertex population. The updated strategy for

vid would be stored here. Each strategy is identified with a nonnegative integer, whose interpretation depends on the payoff matrix of the game. Generally we use the strategy ID as a row or column index of the payoff matrix. The length of this

vector must be the same as the number of vertices in the vertex set of graph.

mode: Defines the sort of neighbourhood to consider for vid. This is only relevant if we

are considering the local perspective, i.e. if <code>islocal</code> is true. If we are considering the global perspective, then it is safe to pass the value <code>IGRAPH_ALL</code> here. If <code>graph</code> is undirected, then we use all the immediate neighbours of <code>vid</code>. Thus if you know that <code>graph</code> is undirected, then it is safe to pass the value <code>IGRAPH_ALL</code>

here. Supported values are:

Use the out-neighbours of vid. This option is only relevant when IGRAPH_OUT

graph is a digraph and we are considering the local perspective.

Use the in-neighbours of vid. Again this option is only relevant IGRAPH_IN

when graph is a directed graph and we are considering the local

perspective.

Use both the in- and out-neighbours of vid. This option is only IGRAPH_ALL

relevant if graph is a digraph. Also use this value if graph is

undirected or we are considering the global perspective.

Returns:

The error code IGRAPH_EINVAL is returned in each of the following cases: (1) Any of the parameters graph, quantities, or strategies is a null pointer. (2) The vector quantities or strategies has a length different from the number of vertices in graph. (3) The parameter graph is the empty or null graph, i.e. the graph with zero vertices and edges. (4) The vector quantities sums to zero.

Time complexity: O(n) where n is the number of vertices in the perspective to consider. If we consider the global perspective, then n is the number of vertices in the vertex set of graph. On the other hand, for the local perspective n is the degree of vid, excluding loops.

Reference:

(Yu & Gen 2010) X. Yu and M. Gen. Introduction to Evolutionary Algorithms. Springer, 2010, pages 18--20.

Example 10.2. File igraph_roulette_wheel_imitation.c

examples/simple/

igraph stochastic imitation — Adopt a strategy via stochastic imitation with uniform selection.

```
igraph_error_t igraph_stochastic_imitation(const igraph_t *graph,
                                igraph_integer_t vid,
                                igraph_imitate_algorithm_t algo,
                                const igraph_vector_t *quantities,
                                igraph_vector_int_t *strategies,
                                igraph_neimode_t mode);
```

A simple stochastic imitation strategy where a vertex revises its strategy to that of a vertex chosen uniformly at random from its local neighbourhood. This is called stochastic imitation via uniform selection, where the strategy to imitate is chosen via some random process. For the purposes of this function, we use uniform selection from a pool of candidates.

Arguments:

graph: The graph object representing the game network. This cannot be the empty or trivial

graph, but must have at least two vertices and one edge. If graph has one vertex, then no strategy update would take place. Furthermore, if graph has at least two

vertices but zero edges, then strategy update would also not take place.

The vertex whose strategy is to be updated. It is assumed that vid represents a vid:

vertex in graph. No checking is performed and it is your responsibility to ensure

that vid is indeed a vertex of graph. If an isolated vertex is provided, i.e. the input vertex has degree 0, then no strategy update would take place and vid would retain its current strategy. Strategy update would also not take place if the local neighbourhood of vid are its in-neighbours (respectively out-neighbours), but vid has zero in-neighbours (respectively out-neighbours). Loops are ignored in computing the degree (in, out, all) of vid.

algo:

This flag controls which algorithm to use in stochastic imitation. Supported values are:

IGRAPH_IMITATE_AUGMENTED Augmented imitation. Vertex vid imitates

the strategy of the chosen vertex u provided that doing so would increase the quantity (e.g. fitness) of *vid*. Augmented imitation can be thought of as "imitate if better".

IGRAPH_IMITATE_BLIND Blind imitation. Vertex vid blindly imi-

tates the strategy of the chosen vertex u, regardless of whether doing so would increase

or decrease the quantity of vid.

IGRAPH_IMITATE_CONTRACT-

ED

Contracted imitation. Here vertex *vid* imitates the strategy of the chosen vertex u if doing so would decrease the quantity of *vid*. Think of contracted imitation as "imitation if warms."

itate if worse".

quantities:

A vector of quantities providing the quantity of each vertex in *graph*. Think of each entry of the vector as being generated by a function such as the fitness function for the game. So if the vector represents fitness quantities, then each vector entry is the fitness of some vertex. The length of this vector must be the same as the number of vertices in the vertex set of *graph*.

strategies:

A vector of the current strategies for the vertex population. The updated strategy for vid would be stored here. Each strategy is identified with a nonnegative integer, whose interpretation depends on the payoff matrix of the game. Generally we use the strategy ID as a row or column index of the payoff matrix. The length of this vector must be the same as the number of vertices in the vertex set of graph.

mode:

Defines the sort of neighbourhood to consider for *vid*. If *graph* is undirected, then we use all the immediate neighbours of *vid*. Thus if you know that *graph* is undirected, then it is safe to pass the value *IGRAPH_ALL* here. Supported values are:

IGRAPH_OUT Use the out-neighbours of vid. This option is only relevant when

graph is a directed graph.

IGRAPH_IN Use the in-neighbours of vid. Again this option is only relevant

when graph is a directed graph.

IGRAPH_ALL Use both the in- and out-neighbours of vid. This option is only

relevant if graph is a digraph. Also use this value if graph is

undirected.

Returns:

The error code *IGRAPH_EINVAL* is returned in each of the following cases: (1) Any of the parameters *graph*, *quantities*, or *strategies* is a null pointer. (2) The vector *quantities* or *strategies* has a length different from the number of vertices in *graph*. (3) The parameter

graph is the empty or null graph, i.e. the graph with zero vertices and edges. (4) The parameter algo refers to an unsupported stochastic imitation algorithm.

Time complexity: depends on the uniform random number generator, but should usually be O(1).

File **Example** 10.3. examples/simple/ igraph_stochastic_imitation.c

Epidemic models

igraph sir — Performs a number of SIR epidemics model runs on a graph.

```
igraph_error_t igraph_sir(const igraph_t *graph, igraph_real_t beta,
               igraph_real_t gamma, igraph_integer_t no_sim,
               igraph_vector_ptr_t *result);
```

The SIR model is a simple model from epidemiology. The individuals of the population might be in three states: susceptible, infected and recovered. Recovered people are assumed to be immune to the disease. Susceptibles become infected with a rate that depends on their number of infected neighbors. Infected people become recovered with a constant rate. See these parameters below.

This function runs multiple simulations, all starting with a single uniformly randomly chosen infected individual. A simulation is stopped when no infected individuals are left.

Arguments:

The graph to perform the model on. For directed graphs edge directions are ignored and graph:

a warning is given.

beta: The rate of infection of an individual that is susceptible and has a single infected neighbor.

The infection rate of a susceptible individual with n infected neighbors is n times beta.

Formally this is the rate parameter of an exponential distribution.

The rate of recovery of an infected individual. Formally, this is the rate parameter of an gamma:

exponential distribution.

no sim: The number of simulation runs to perform.

result: The result of the simulation is stored here, in a list of igraph sir t objects. To deal-

> locate memory, the user needs to call igraph_sir_destroy on each element, before destroying the pointer vector itself using igraph_vector_ptr_destroy_all().

Returns:

Error code.

Time complexity: $O(no_sim * (|V| + |E| log(|V|)))$.

igraph sir t — The result of one SIR model simulation.

```
typedef struct igraph_sir_t {
    igraph_vector_t times;
    igraph_vector_int_t no_s, no_i, no_r;
} igraph_sir_t;
```

Data structure to store the results of one simulation of the SIR (susceptible-infected-recovered) model on a graph. It has the following members. They are all (real or integer) vectors, and they are of the same length.

Values:

times: A vector, the times of the events are stored here.

no_s: An integer vector, the number of susceptibles in each time step is stored here.

no_i: An integer vector, the number of infected individuals at each time step, is stored here.

no_r: An integer vector, the number of recovered individuals is stored here at each time step.

igraph_sir_destroy — Deallocates memory associated with a SIR simulation run.

```
void igraph_sir_destroy(igraph_sir_t *sir);
```

Arguments:

sir: The igraph_sir_t object storing the simulation.

Chapter 11. Vertex and edge selectors and sequences, iterators

About selectors, iterators

Everything about vertices and vertex selectors also applies to edges and edge selectors unless explicitly noted otherwise.

The vertex (and edge) selector notion was introduced in igraph 0.2. It is a way to reference a sequence of vertices or edges independently of the graph.

While this might sound quite mysterious, it is actually very simple. For example, all vertices of a graph can be selected by <code>igraph_vs_all()</code> and the graph independence means that <code>igraph_vs_al-l()</code> is not parametrized by a graph object. That is, <code>igraph_vs_all()</code> is the general <code>concept</code> of selecting all vertices of a graph. A vertex selector is then a way to specify the class of vertices to be visited. The selector might specify that all vertices of a graph or all the neighbours of a vertex are to be visited. A vertex selector is a way of saying that you want to visit a bunch of vertices, as opposed to a vertex iterator which is a concrete plan for visiting each of the chosen vertices of a specific graph.

To determine the actual vertex IDs implied by a vertex selector, you need to apply the concept of selecting vertices to a specific graph object. This can be accomplished by instantiating a vertex iterator using a specific vertex selection concept and a specific graph object. The notion of vertex iterators can be thought of in the following way. Given a specific graph object and the class of vertices to be visited, a vertex iterator is a road map, plan or route for how to visit the chosen vertices.

Some vertex selectors have *immediate* versions. These have the prefix <code>igraph_vss</code> instead of <code>igraph_vs</code>, e.g. <code>igraph_vss_all()</code> instead of <code>igraph_vs_all()</code>. The immediate versions are to be used in the parameter list of the igraph functions, such as <code>igraph_degree()</code>. These functions are not associated with any <code>igraph_vs_t</code> object, so they have no separate constructors and destructors (destroy functions).

Vertex selector constructors

Vertex selectors are created by vertex selector constructors, can be instantiated with igraph_vit_create(), and are destroyed with igraph_vs_destroy().

igraph_vs_all — Vertex set, all vertices of a graph.

```
igraph_error_t igraph_vs_all(igraph_vs_t *vs);

Arguments:

vs: Pointer to an uninitialized igraph_vs_t object.

Returns:
    Error code.

See also:
```

igraph vss all(), igraph vs destroy()

This selector includes all vertices of a given graph in increasing vertex ID order.

Time complexity: O(1).

igraph_vs_adj — Adjacent vertices of a vertex.

All neighboring vertices of a given vertex are selected by this selector. The mode argument controls the type of the neighboring vertices to be selected. The vertices are visited in increasing vertex ID order, as of igraph version 0.4.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

vid: Vertex ID, the center of the neighborhood.

mode: Decides the type of the neighborhood for directed graphs. This parameter is ignored for undirected graphs. Possible values:

IGRAPH_OUT All vertices to which there is a directed edge from vid. That is, all the out-neighbors of vid.

 ${\tt IGRAPH_IN} \qquad {\tt All \ vertices \ from \ which \ there \ is \ a \ directed \ edge \ to \ vid. \ In \ other \ words,}$

all the in-neighbors of vid.

IGRAPH_ALL All vertices to which or from which there is a directed edge from/to vid.

That is, all the neighbors of vid considered as if the graph is undirected.

Returns:

Error code.

See also:

```
igraph_vs_destroy()
```

Time complexity: O(1).

igraph_vs_nonadj — Non-adjacent vertices of a vertex.

All non-neighboring vertices of a given vertex. The *mode* argument controls the type of neighboring vertices *not* to select. Instead of selecting immediate neighbors of vid as is done by igraph_vs_adj(), the current function selects vertices that are *not* immediate neighbors of vid.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

vid: Vertex ID, the "center" of the non-neighborhood.

mode: The type of neighborhood not to select in directed graphs. Possible values:

IGRAPH_OUT All vertices will be selected except those to which there is a directed edge

from vid. That is, we select all vertices excluding the out-neighbors of

vid.

IGRAPH_IN All vertices will be selected except those from which there is a directed

edge to vid. In other words, we select all vertices but the in-neighbors

of vid.

IGRAPH_ALL All vertices will be selected except those from or to which there is a di-

rected edge to or from vid. That is, we select all vertices of vid except

for its immediate neighbors.

Returns:

Error code.

See also:

```
igraph_vs_destroy()
```

Time complexity: O(1).

Example 11.1. File examples/simple/igraph_vs_nonadj.c

igraph_vs_none — Empty vertex set.

```
igraph_error_t igraph_vs_none(igraph_vs_t *vs);
```

Creates an empty vertex selector.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

Returns:

Error code.

See also:

```
igraph_vss_none(), igraph_vs_destroy()
```

Time complexity: O(1).

igraph_vs_1 — Vertex set with a single vertex.

```
igraph_error_t igraph_vs_1(igraph_vs_t *vs, igraph_integer_t vid);
```

This vertex selector selects a single vertex.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

vid: The vertex ID to be selected.

Returns:

Error Code.

See also:

```
igraph_vss_1(), igraph_vs_destroy()
```

Time complexity: O(1).

igraph_vs_vector — Vertex set based on a vector.

This function makes it possible to handle an igraph_vector_int_t temporarily as a vertex selector. The vertex selector should be thought of as a *view* into the vector. If you make changes to the vector that also affects the vertex selector. Destroying the vertex selector does not destroy the vector. Do not destroy the vector before destroying the vertex selector, or you might get strange behavior. Since selectors are not tied to any specific graph, this function does not check whether the vertex IDs in the vector are valid.

Arguments:

vs: Pointer to an uninitialized vertex selector.

v: Pointer to a igraph_vector_int_t object.

Returns:

Error code.

See also:

```
igraph_vss_vector(), igraph_vs_destroy()
```

Time complexity: O(1).

Example 11.2. File examples/simple/igraph_vs_vector.c

igraph_vs_vector_small — Create a vertex set by giving its elements.

```
igraph_error_t igraph_vs_vector_small(igraph_vs_t *vs, ...);
```

This function can be used to create a vertex selector with a few of vertices. Do not forget to include a -1 after the last vertex ID. The behavior of the function is undefined if you don't use a -1 properly.

Note that the vertex IDs supplied will be parsed as value of type int so you cannot supply arbitrarily large (too large for int) vertex IDs here.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

...: Additional parameters, these will be the vertex IDs to be included in the vertex selector. Supply a -1 after the last vertex ID.

Returns:

Error code.

See also:

```
igraph_vs_destroy()
```

Time complexity: O(n), the number of vertex IDs supplied.

igraph_vs_vector_copy — Vertex set based on a vector, with copying.

```
igraph_error_t igraph_vs_vector_copy(igraph_vs_t *vs, const igraph_vector_int_t
```

This function makes it possible to handle an igraph_vector_int_t permanently as a vertex selector. The vertex selector creates a copy of the original vector, so the vector can safely be destroyed after creating the vertex selector. Changing the original vector will not affect the vertex selector. The vertex selector is responsible for deleting the copy made by itself. Since selectors are not tied to any specific graph, this function does not check whether the vertex IDs in the vector are valid.

Arguments:

vs: Pointer to an uninitialized vertex selector.

v: Pointer to a igraph_vector_int_t object.

Returns:

Error code.

See also:

```
igraph_vs_destroy()
```

Time complexity: O(1).

igraph_vs_range — Vertex set, an interval of vertices.

```
igraph_error_t igraph_vs_range(igraph_vs_t *vs, igraph_integer_t start, igraph_
```

Creates a vertex selector containing all vertices with vertex ID equal to or bigger than from and smaller than to. Note that the interval is closed from the left and open from the right, following C conventions.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

start: The first vertex ID to be included in the vertex selector.

end: The first vertex ID not to be included in the vertex selector.

Returns:

Error code.

See also:

```
igraph_vss_range(), igraph_vs_destroy()
Time complexity: O(1).
```

Example 11.3. File examples/simple/igraph_vs_seq.c

Generic vertex selector operations

igraph_vs_copy — Creates a copy of a vertex selector.

```
igraph_error_t igraph_vs_copy(igraph_vs_t* dest, const igraph_vs_t* src);
```

Arguments:

src: The selector being copied.

dest: An uninitialized selector that will contain the copy.

igraph_vs_destroy — Destroy a vertex set.

```
void igraph_vs_destroy(igraph_vs_t *vs);
```

This function should be called for all vertex selectors when they are not needed. The memory allocated for the vertex selector will be deallocated. Do not call this function on vertex selectors created with the immediate versions of the vertex selector constructors (starting with igraph_vss).

Arguments:

vs: Pointer to a vertex selector object.

Time complexity: operating system dependent, usually O(1).

igraph_vs_is_all — Check whether all vertices are included.

```
igraph_bool_t igraph_vs_is_all(const igraph_vs_t *vs);
```

This function checks whether the vertex selector object was created by <code>igraph_vs_all()</code> or <code>igraph_vss_all()</code>. Note that the vertex selector might contain all vertices in a given graph but if it wasn't created by the two constructors mentioned here the return value will be false.

Arguments:

vs: Pointer to a vertex selector object.

Returns:

true if the vertex selector contains all vertices and false otherwise.

Time complexity: O(1).

igraph_vs_size — Returns the size of the vertex selector.

The size of the vertex selector is the number of vertices it will yield when it is iterated over.

Arguments:

graph: The graph over which we will iterate.

result: The result will be returned here.

igraph_vs_type — Returns the type of the vertex selector.

```
igraph_vs_type_t igraph_vs_type(const igraph_vs_t *vs);
```

Immediate vertex selectors

igraph_vss_all — All vertices of a graph (immediate version).

```
igraph_vs_t igraph_vss_all(void);
```

Immediate vertex selector for all vertices in a graph. It can be used conveniently when some vertex property (e.g. betweenness, degree, etc.) should be calculated for all vertices.

Returns:

A vertex selector for all vertices in a graph.

See also:

```
igraph_vs_all()
Time complexity: O(1).
```

igraph_vss_none — Empty vertex set (immediate version).

```
igraph_vs_t igraph_vss_none(void);
```

The immediate version of the empty vertex selector.

Returns:

An empty vertex selector.

See also:

```
igraph_vs_none()
```

Time complexity: O(1).

igraph_vss_1 — Vertex set with a single vertex (immediate version).

```
igraph_vs_t igraph_vss_1(igraph_integer_t vid);
```

The immediate version of the single-vertex selector.

Arguments:

vid: The vertex to be selected.

Returns:

A vertex selector containing a single vertex.

See also:

```
igraph_vs_1()
```

Time complexity: O(1).

igraph_vss_vector — Vertex set based on a vector (immediate version).

```
igraph_vs_t igraph_vss_vector(const igraph_vector_int_t *v);
```

This is the immediate version of igraph_vs_vector.

Arguments:

v: Pointer to a igraph_vector_int_t object.

Returns:

A vertex selector object containing the vertices in the vector.

See also:

```
igraph_vs_vector()
```

Time complexity: O(1).

igraph_vss_range — An interval of vertices (immediate version).

```
igraph_vs_t igraph_vss_range(igraph_integer_t start, igraph_integer_t end);
The immediate version of igraph_vs_range().
```

Arguments:

start: The first vertex ID to be included in the vertex selector.

end: The first vertex ID not to be included in the vertex selector.

Returns:

Error code.

See also:

```
igraph_vs_range()
```

Time complexity: O(1).

Vertex iterators

igraph_vit_create — Creates a vertex iterator from a vertex selector.

This function instantiates a vertex selector object with a given graph. This is the step when the actual vertex IDs are created from the *logical* notion of the vertex selector based on the graph. E.g. a vertex selector created with <code>igraph_vs_all()</code> contains knowledge that *all* vertices are included in a (yet indefinite) graph. When instantiating it a vertex iterator object is created, this contains the actual vertex IDs in the graph supplied as a parameter.

The same vertex selector object can be used to instantiate any number vertex iterators.

Arguments:

graph: An igraph_t object, a graph.

vs: A vertex selector object.

vit: Pointer to an uninitialized vertex iterator object.

Returns:

Error code.

See also:

```
igraph_vit_destroy().
```

Time complexity: it depends on the vertex selector type. O(1) for vertex selectors created with $igraph_vs_all()$, $igraph_vs_none()$, $igraph_vs_1$, $igraph_vs_vector$, $igraph_vs_range()$, $igraph_vs_vector()$, $igraph_vs_vector_small()$. O(d) for $igraph_vs_adj()$, d is the number of vertex IDs to be included in the iterator. O(|V|) for $igraph_vs_nonadj()$, |V| is the number of vertices in the graph.

igraph_vit_destroy — Destroys a vertex iterator.

```
void igraph_vit_destroy(const igraph_vit_t *vit);
```

Deallocates memory allocated for a vertex iterator.

Arguments:

vit: Pointer to an initialized vertex iterator object.

See also:

```
igraph_vit_create()
```

Time complexity: operating system dependent, usually O(1).

Stepping over the vertices

After creating an iterator with <code>igraph_vit_create()</code>, it points to the first vertex in the vertex determined by the vertex selector (if there is any). The <code>IGRAPH_VIT_NEXT()</code> macro steps to the next vertex, <code>IGRAPH_VIT_END()</code> checks whether there are more vertices to vis-

it, IGRAPH_VIT_SIZE() gives the total size of the vertices visited so far and to be visited. IGRAPH_VIT_RESET() resets the iterator, it will point to the first vertex again. Finally IGRAPH_VIT_GET() gives the current vertex pointed to by the iterator (call this only if IGRAPH_VIT_END() is false).

Here is an example on how to step over the neighbors of vertex 0:

```
igraph_vs_t vs;
igraph_vit_t vit;
...
igraph_vs_adj(&vs, 0, IGRAPH_ALL);
igraph_vit_create(&graph, vs, &vit);
while (!IGRAPH_VIT_END(vit)) {
   printf(" %" IGRAPH_PRId, IGRAPH_VIT_GET(vit));
   IGRAPH_VIT_NEXT(vit);
}
printf("\n");
...
igraph_vit_destroy(&vit);
igraph_vs_destroy(&vs);
```

IGRAPH_VIT_NEXT — Next vertex.

```
#define IGRAPH_VIT_NEXT(vit)
```

Steps the iterator to the next vertex. Only call this function if IGRAPH_VIT_END() returns false.

Arguments:

vit: The vertex iterator to step.

Time complexity: O(1).

IGRAPH_VIT_END — Are we at the end?

```
#define IGRAPH_VIT_END(vit)
```

Checks whether there are more vertices to step to.

Arguments:

vit: The vertex iterator to check.

Returns:

Logical value, if true there are no more vertices to step to.

Time complexity: O(1).

IGRAPH_VIT_SIZE — Size of a vertex iterator.

```
#define IGRAPH_VIT_SIZE(vit)
```

Gives the number of vertices in a vertex iterator.

Arguments:

vit: The vertex iterator.

Returns:

The number of vertices.

Time complexity: O(1).

IGRAPH VIT RESET — Reset a vertex iterator.

```
#define IGRAPH_VIT_RESET(vit)
```

Resets a vertex iterator. After calling this macro the iterator will point to the first vertex.

Arguments:

vit: The vertex iterator.

Time complexity: O(1).

IGRAPH_VIT_GET — Query the current position.

```
#define IGRAPH_VIT_GET(vit)
```

Gives the vertex ID of the current vertex pointed to by the iterator.

Arguments:

vit: The vertex iterator.

Returns:

The vertex ID of the current vertex.

Time complexity: O(1).

Edge selector constructors

igraph_es_all — Edge set, all edges.

Arguments:

es: Pointer to an uninitialized edge selector object.

order:

Constant giving the order in which the edges will be included in the selector. Possible values: IGRAPH_EDGEORDER_ID, edge ID order. IGRAPH_EDGEORDER_FROM, vertex ID order, the id of the *source* vertex counts for directed graphs. The order of the incident edges of a given vertex is arbitrary. IGRAPH_EDGEORDER_TO, vertex ID order, the ID of the *target* vertex counts for directed graphs. The order of the incident edges of a given vertex is arbitrary. For undirected graph the latter two is the same.

Returns:

Error code.

See also:

```
igraph_ess_all(),igraph_es_destroy()
```

Time complexity: O(1).

igraph_es_incident — Edges incident on a given vertex.

Arguments:

es: Pointer to an uninitialized edge selector object.

vid: Vertex ID, of which the incident edges will be selected.

mode:

Constant giving the type of the incident edges to select. This is ignored for undirected graphs. Possible values: IGRAPH_OUT, outgoing edges; IGRAPH_IN, incoming edges; IGRAPH_ALL, all edges.

Returns:

Error code.

See also:

```
igraph_es_destroy()
```

Time complexity: O(1).

igraph_es_none — Empty edge selector.

```
igraph_error_t igraph_es_none(igraph_es_t *es);
```

Arguments:

es: Pointer to an uninitialized edge selector object to initialize.

Returns:

Error code.

See also:

```
igraph_ess_none(), igraph_es_destroy()
```

Time complexity: O(1).

igraph_es_1 — Edge selector containing a single edge.

```
igraph_error_t igraph_es_1(igraph_es_t *es, igraph_integer_t eid);
```

Arguments:

es: Pointer to an uninitialized edge selector object.

eid: Edge ID of the edge to select.

Returns:

Error code.

See also:

```
igraph_ess_1(), igraph_es_destroy()
```

Time complexity: O(1).

igraph_es_vector — Handle a vector as an edge selector.

```
igraph_error_t igraph_es_vector(igraph_es_t *es, const igraph_vector_int_t *v);
```

Creates an edge selector which serves as a view into a vector containing edge IDs. Do not destroy the vector before destroying the edge selector. Since selectors are not tied to any specific graph, this function does not check whether the edge IDs in the vector are valid.

Arguments:

es: Pointer to an uninitialized edge selector.

v: Vector containing edge IDs.

Returns:

Error code.

See also:

```
igraph_ess_vector(), igraph_es_destroy()
Time complexity: O(1).
```

igraph_es_range — Edge selector, a sequence of edge IDs.

```
igraph_error_t igraph_es_range(igraph_es_t *es, igraph_integer_t start, igraph_
```

Creates an edge selector containing all edges with edge ID equal to or bigger than from and smaller than to. Note that the interval is closed from the left and open from the right, following C conventions.

Arguments:

vs: Pointer to an uninitialized edge selector object.

start: The first edge ID to be included in the edge selector.

end: The first edge ID not to be included in the edge selector.

Returns:

Error code.

See also:

```
igraph_ess_range(), igraph_es_destroy()
```

Time complexity: O(1).

igraph_es_pairs — Edge selector, multiple edges defined by their endpoints in a vector.

The edges between the given pairs of vertices will be included in the edge selection. The vertex pairs must be defined in the vector v, the first element of the vector is the first vertex of the first edge to be selected, the second element is the second vertex of the first edge, the third element is the first vertex of the second edge and so on.

Arguments:

es: Pointer to an uninitialized edge selector object.

v: The vector containing the endpoints of the edges.

directed: Whether the graph is directed or not.

Returns:

Error code.

See also:

```
igraph_es_pairs_small(), igraph_es_destroy()
```

Time complexity: O(n), the number of edges being selected.

Example 11.4. File examples/simple/igraph_es_pairs.c

igraph_es_pairs_small — Edge selector, multiple edges defined by their endpoints as arguments.

```
igraph_error_t igraph_es_pairs_small(igraph_es_t *es, igraph_bool_t directed, if
```

The edges between the given pairs of vertices will be included in the edge selection. The vertex pairs must be given as the arguments of the function call, the third argument is the first vertex of the first edge, the fourth argument is the second vertex of the first edge, the fifth is the first vertex of the second edge and so on. The last element of the argument list must be -1 to denote the end of the argument list.

Note that the vertex IDs supplied will be parsed as int's so you cannot supply arbitrarily large (too large for int) vertex IDs here.

Arguments:

es: Pointer to an uninitialized edge selector object.

directed: Whether the graph is directed or not.

. . .: The additional arguments give the edges to be included in the selector, as pairs of ver-

tex IDs. The last argument must be -1. The first parameter is present for technical

reasons and represents the first variadic argument.

Returns:

Error code.

See also:

```
igraph_es_pairs(), igraph_es_destroy()
```

Time complexity: O(n), the number of edges being selected.

igraph_es_path — Edge selector, edge IDs on a path.

This function takes a vector of vertices and creates a selector of edges between those vertices. Vector $\{0, 3, 4, 7\}$ will select edges (0 -> 3), (3 -> 4), (4 -> 7). If these edges don't exist then trying to create an iterator using this selector will fail.

Arguments:

Vertex and edge selectors and sequences, iterators

es: Pointer to an uninitialized edge selector object.

v: Pointer to a vector of vertex IDs along the path.

directed: If edge directions should be taken into account. This will be ignored if the graph to

select from is undirected.

Returns:

Error code.

See also:

```
igraph_es_destroy()
```

Time complexity: O(n), the number of vertices.

igraph_es_vector_copy — Edge set, based on a vector, with copying.

```
igraph_error_t igraph_es_vector_copy(igraph_es_t *es, const igraph_vector_int_t
```

This function makes it possible to handle an igraph_vector_int_t permanently as an edge selector. The edge selector creates a copy of the original vector, so the vector can safely be destroyed after creating the edge selector. Changing the original vector will not affect the edge selector. The edge selector is responsible for deleting the copy made by itself. Since selectors are not tied to any specific graph, this function does not check whether the edge IDs in the vector are valid.

Arguments:

es: Pointer to an uninitialized edge selector.

v: Pointer to a igraph_vector_int_t object.

Returns:

Error code.

See also:

```
igraph_es_destroy()
```

Time complexity: O(1).

Immediate edge selectors

igraph_ess_all — Edge set, all edges (immediate version).

```
igraph_es_t igraph_ess_all(igraph_edgeorder_type_t order);
```

The immediate version of the all-edges selector.

Arguments:

order: Constant giving the order of the edges in the edge selector. See igraph_es_all() for the possible values.

Returns:

The edge selector.

See also:

```
igraph_es_all()
```

Time complexity: O(1).

igraph_ess_none — Immediate empty edge selector.

```
igraph_es_t igraph_ess_none(void);
```

Immediate version of the empty edge selector.

Returns:

Initialized empty edge selector.

See also:

```
igraph_es_none()
```

Time complexity: O(1).

igraph_ess_1 — Immediate version of the single edge edge selector.

```
igraph_es_t igraph_ess_1(igraph_integer_t eid);
```

Arguments:

eid: The ID of the edge.

Returns:

The edge selector.

See also:

```
igraph_es_1()
```

Time complexity: O(1).

igraph_ess_vector — Immediate vector view edge selector.

```
igraph_es_t igraph_ess_vector(const igraph_vector_int_t *v);
```

This is the immediate version of the vector of edge IDs edge selector.

Arguments:

v: The vector of edge IDs.

Returns:

Edge selector, initialized.

See also:

```
igraph_es_vector()
```

Time complexity: O(1).

igraph_ess_range — Immediate version of the sequence edge selector.

```
igraph_es_t igraph_ess_range(igraph_integer_t start, igraph_integer_t end);
```

Arguments:

start: The first edge ID to be included in the edge selector.

end: The first edge ID not to be included in the edge selector.

Returns:

The initialized edge selector.

See also:

```
igraph_es_range()
```

Time complexity: O(1).

Generic edge selector operations

igraph_es_as_vector — Transform edge selector into vector.

Call this function on an edge selector to transform it into a vector. This is only implemented for sequence and vector selectors. If the edges do not exist in the graph, this will result in an error.

Arguments:

graph: Pointer to a graph to check if the edges in the selector exist.

es: An edge selector object.

v: Pointer to initialized vector. The result will be stored here.

Time complexity: O(n), the number of edges in the selector.

igraph_es_copy — Creates a copy of an edge selector.

```
igraph_error_t igraph_es_copy(igraph_es_t* dest, const igraph_es_t* src);
```

Arguments:

src: The selector being copied.

dest: An uninitialized selector that will contain the copy.

See also:

igraph_es_destroy()

igraph_es_destroy — Destroys an edge selector object.

```
void igraph_es_destroy(igraph_es_t *es);
```

Call this function on an edge selector when it is not needed any more. Do *not* call this function on edge selectors created by immediate constructors, those don't need to be destroyed.

Arguments:

es: Pointer to an edge selector object.

Time complexity: operating system dependent, usually O(1).

igraph_es_is_all — Check whether an edge selector includes all edges.

```
igraph_bool_t igraph_es_is_all(const igraph_es_t *es);
```

Arguments:

es: Pointer to an edge selector object.

Returns:

true if es was created with igraph_es_all() or igraph_ess_all(), and false otherwise.

Time complexity: O(1).

igraph_es_size — Returns the size of the edge selector.

The size of the edge selector is the number of edges it will yield when it is iterated over.

Arguments:

graph: The graph over which we will iterate.

result: The result will be returned here.

igraph_es_type — Returns the type of the edge selector.

```
igraph_es_type_t igraph_es_type(const igraph_es_t *es);
```

Edge iterators

igraph_eit_create — Creates an edge iterator from an edge selector.

This function creates an edge iterator based on an edge selector and a graph.

The same edge selector can be used to create many edge iterators, also for different graphs.

Arguments:

graph: An igraph_t object for which the edge selector will be instantiated.

es: The edge selector to instantiate.

eit: Pointer to an uninitialized edge iterator.

Returns:

Error code.

See also:

```
igraph_eit_destroy()
```

Time complexity: depends on the type of the edge selector. For edge selectors created by $igraph_es_all()$, $igraph_es_none()$, $igraph_es_1()$, $igraph_es_vector()$, $igraph_es_seq()$ it is O(1). For $igraph_es_incident()$ it is O(d) where d is the number of incident edges of the vertex.

igraph_eit_destroy — Destroys an edge iterator.

```
void igraph_eit_destroy(const igraph_eit_t *eit);
```

Arguments:

eit: Pointer to an edge iterator to destroy.

See also:

```
igraph_eit_create()
```

Time complexity: operating system dependent, usually O(1).

Stepping over the edges

Just like for vertex iterators, macros are provided for stepping over a sequence of edges: IGRAPH_EIT_NEXT() goes to the next edge, IGRAPH_EIT_END() checks whether there are more edges to visit, IGRAPH_EIT_SIZE() gives the number of edges in the edge sequence, IGRAPH_EIT_RESET() resets the iterator to the first edge and IGRAPH_EIT_GET() returns the id of the current edge.

IGRAPH_EIT_NEXT — Next edge.

```
#define IGRAPH EIT NEXT(eit)
```

Steps the iterator to the next edge. Call this function only if IGRAPH_EIT_END() returns false.

Arguments:

eit: The edge iterator to step.

Time complexity: O(1).

IGRAPH_EIT_END — Are we at the end?

```
#define IGRAPH_EIT_END(eit)
```

Checks whether there are more edges to step to.

Arguments:

wit: The edge iterator to check.

Returns:

Logical value, if true there are no more edges to step to.

Time complexity: O(1).

IGRAPH_EIT_SIZE — Number of edges in the iterator.

```
#define IGRAPH_EIT_SIZE(eit)
```

Gives the number of edges in an edge iterator.

Arguments:

eit: The edge iterator.

Returns:

The number of edges.

Time complexity: O(1).

IGRAPH_EIT_RESET — Reset an edge iterator.

```
#define IGRAPH_EIT_RESET(eit)
```

Resets an edge iterator. After calling this macro the iterator will point to the first edge.

Arguments:

eit: The edge iterator.

Time complexity: O(1).

IGRAPH_EIT_GET — Query an edge iterator.

```
#define IGRAPH_EIT_GET(eit)
```

Gives the edge ID of the current edge pointed to by an iterator.

Arguments:

eit: The edge iterator.

Returns:

The id of the current edge.

Time complexity: O(1).

Deprecated functions

igraph_es_seq — Edge selector, a sequence of edge IDs, with inclusive endpoints (deprecated).

igraph_error_t igraph_es_seq(igraph_es_t *es, igraph_integer_t from, igraph_int All edge IDs between from and to (inclusive) will be included in the edge selection.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_es_range() instead.

Arguments:

es: Pointer to an uninitialized edge selector object.

from: The first edge ID to be included.

to: The last edge ID to be included.

Returns:

Error code.

See also:

```
igraph_ess_seq(), igraph_es_destroy()
```

Time complexity: O(1).

igraph_ess_seq — Immediate version of the sequence edge selector, with inclusive endpoints.

```
igraph_es_t igraph_ess_seq(igraph_integer_t from, igraph_integer_t to);
```

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use $igraph_ess_range()$ instead.

Arguments:

from: The first edge ID to include.

to: The last edge ID to include.

Returns:

The initialized edge selector.

See also:

```
igraph_es_seq()
```

Time complexity: O(1).

igraph_vs_seq — Vertex set, an interval of vertices with inclusive endpoints (deprecated).

```
igraph_error_t igraph_vs_seq(igraph_vs_t *vs, igraph_integer_t from, igraph_int
```

Creates a vertex selector containing all vertices with vertex ID equal to or bigger than from and equal to or smaller than to. Note that both endpoints are inclusive, contrary to C conventions.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_vs_range() instead.

Arguments:

vs: Pointer to an uninitialized vertex selector object.

from: The first vertex ID to be included in the vertex selector.

to: The last vertex ID to be included in the vertex selector.

Returns:

Error code.

See also:

```
igraph_vs_range(), igraph_vss_seq(), igraph_vs_destroy()
```

Time complexity: O(1).

Example 11.5. File examples/simple/igraph_vs_seq.c

igraph_vss_seq — An interval of vertices with inclusive endpoints (immediate version, deprecated).

Vertex and edge selectors and sequences, iterators

igraph_vs_t igraph_vss_seq(igraph_integer_t from, igraph_integer_t to);
The immediate version of igraph_vs_seq().

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use $igraph_vss_range()$ instead.

Arguments:

from: The first vertex ID to be included in the vertex selector.

to: The last vertex ID to be included in the vertex selector.

Returns:

Error code.

See also:

```
igraph_vss_range(), igraph_vs_seq()
```

Time complexity: O(1).

Chapter 12. Graph, vertex and edge attributes

Attributes are numbers or strings (or basically any kind of data) associated with the vertices or edges of a graph, or with the graph itself. E.g. you may label vertices with symbolic names or attach numeric weights to the edges of a graph.

igraph attributes are designed to be flexible and extensible. In igraph attributes are implemented via an interface abstraction: any type implementing the functions in the interface, can be used for storing vertex, edge and graph attributes. This means that different attribute implementations can be used together with igraph. This is reasonable: if igraph is used from Python attributes can be of any Python type, from GNU R all R types are allowed. There is an experimental attribute implementation to be used when programming in C, but by default it is currently turned off.

First we briefly look over how attribute handlers can be implemented. This is not something a user does every day. It is rather typically the job of the high level interface writers. (But it is possible to write an interface without implementing attributes.) Then we show the experimental C attribute handler.

The Attribute Handler Interface

It is possible to attach an attribute handling interface to **igraph**. This is simply a table of functions, of type igraph_attribute_table_t. These functions are invoked to notify the attribute handling code about the structural changes in a graph. See the documentation of this type for details.

By default there is no attribute interface attached to **igraph**, to attach one, call igraph_set_at-tribute_table with your new table.

igraph_attribute_table_t — Table of functions to perform operations on attributes

```
typedef struct igraph_attribute_table_t {
            igraph_error_t (*init)(igraph_t *graph, igraph_vector_ptr_t *attr);
                                                          (*destroy)(igraph_t *graph);
            igraph_error_t (*copy)(igraph_t *to, const igraph_t *from, igraph_bool_t ga
                                                                                  igraph_bool_t va, igraph_bool_t ea);
            igraph_error_t (*add_vertices)(igraph_t *graph, igraph_integer_t nv, igraph_
            igraph_error_t (*permute_vertices)(const igraph_t *graph,
                                                                                                                      igraph_t *newgraph,
                                                                                                                      const igraph_vector_int_t *idx);
            igraph_error_t (*combine_vertices)(const igraph_t *graph,
                                                                                                                      igraph t *newgraph,
                                                                                                                      const igraph_vector_int_list_t *merges,
                                                                                                                      const igraph_attribute_combination_t *compared to the combination of the combination
            igraph_error_t (*add_edges)(igraph_t *graph, const igraph_vector_int_t *edg
                                                                                                 igraph_vector_ptr_t *attr);
            igraph_error_t (*permute_edges)(const igraph_t *graph,
                                                                                                             igraph_t *newgraph, const igraph_vector_int
            igraph_error_t (*combine_edges)(const igraph_t *graph,
                                                                                                             igraph_t *newgraph,
                                                                                                             const igraph_vector_int_list_t *merges,
                                                                                                             const igraph_attribute_combination_t *comb)
            igraph_error_t (*get_info)(const igraph_t *graph,
```

```
igraph_strvector_t *gnames, igraph_vector_int_t
                               igraph_strvector_t *vnames, igraph_vector_int_t
                               igraph_strvector_t *enames, igraph_vector_int_t
   igraph_bool_t (*has_attr)(const igraph_t *graph, igraph_attribute_elemtype_
                              const char *name);
   igraph_error_t (*gettype)(const igraph_t *graph, igraph_attribute_type_t *t
                              igraph_attribute_elemtype_t elemtype, const char
   igraph_error_t (*get_numeric_graph_attr)(const igraph_t *graph, const char
                                             igraph_vector_t *value);
   igraph_error_t (*get_string_graph_attr)(const igraph_t *graph, const char *;
                                            igraph_strvector_t *value);
   igraph_error_t (*get_bool_graph_attr)(const igraph_t *igraph, const char *n
                                          igraph_vector_bool_t *value);
   igraph_error_t (*get_numeric_vertex_attr)(const igraph_t *graph, const char
                                              igraph_vs_t vs,
                                              igraph_vector_t *value);
   igraph_error_t (*get_string_vertex_attr)(const igraph_t *graph, const char
                                             igraph_vs_t vs,
                                             igraph_strvector_t *value);
   igraph_error_t (*get_bool_vertex_attr)(const igraph_t *graph, const char *n
                                           igraph_vs_t vs,
                                           igraph_vector_bool_t *value);
   igraph_error_t (*get_numeric_edge_attr)(const igraph_t *graph, const char *:
                                            igraph_es_t es,
                                            igraph_vector_t *value);
   igraph_error_t (*get_string_edge_attr)(const igraph_t *graph, const char *n
                                           igraph_es_t es,
                                           igraph_strvector_t *value);
   igraph_error_t (*get_bool_edge_attr)(const igraph_t *graph, const char *nam
                                         igraph_es_t es,
                                         igraph_vector_bool_t *value);
} igraph_attribute_table_t;
```

This type collects the functions defining an attribute handler. It has the following members:

Values:

init: This function is called whenever a new graph object is cre-

ated, right after it is created but before any vertices or edges are added. It is supposed to set the attr member of the igraph_t object. It is expected to return an error code.

destroy: This function is called whenever the graph object is destroyed,

right before freeing the allocated memory.

copy: This function is called when copying a graph with

igraph_copy, after the structure of the graph has been al-

ready copied. It is expected to return an error code.

add_vertices: Called when vertices are added to a graph, before adding the

vertices themselves. The number of vertices to add is supplied

as an argument. Expected to return an error code.

permute_vertices: Typically called when a new graph is created based on an ex-

isting one, e.g. if vertices are removed from a graph. The supplied index vector defines which old vertex a new vertex corresponds to. Its length must be the same as the number of vertices

in the new graph.

combine_vertices: This function is called when the creation of a new graph in-

volves a merge (contraction, etc.) of vertices from another graph. The function is after the new graph was created. An argument specifies how several vertices from the old graph map

to a single vertex in the new graph.

add_edges: Called when new edges have been added. The number of new

edges are supplied as well. It is expected to return an error code.

permute_edges: Typically called when a new graph is created and some of the

new edges should carry the attributes of some of the old edges. The idx vector shows the mapping between the old edges and the new ones. Its length is the same as the number of edges in the new graph, and for each edge it gives the id of the old edge

(the edge in the old graph).

combine_edges: This function is called when the creation of a new graph in-

volves a merge (contraction, etc.) of edges from another graph. The function is after the new graph was created. An argument specifies how several edges from the old graph map to a single

edge in the new graph.

get_info: Query the attributes of a graph, the names and types should be

returned.

has_attr: Check whether a graph has the named graph/vertex/edge at-

tribute.

gettype: Query the type of a graph/vertex/edge attribute.

get_numeric_graph_attr: Query a numeric graph attribute. The value should be placed as

the first element of the value vector.

get_string_graph_attr: Query a string graph attribute. The value should be placed as

the first element of the value string vector.

get_bool_graph_attr: Query a boolean graph attribute. The value should be placed as

the first element of the value boolean vector.

get_numeric_vertex_attr: Query a numeric vertex attribute, for the vertices included in

vs.

get_string_vertex_attr: Query a string vertex attribute, for the vertices included in vs.

get_bool_vertex_attr: Query a boolean vertex attribute, for the vertices included in

VS.

get_numeric_edge_attr: Query a numeric edge attribute, for the edges included in es.

get_string_edge_attr: Query a string edge attribute, for the edges included in es.

get_bool_edge_attr: Query a boolean edge attribute, for the edges included in es.

Note that the get_*_*_attr are allowed to convert the attributes to numeric or string. E.g. if a vertex attribute is a GNU R complex data type, then get_string_vertex_attribute may serialize it into a string, but this probably makes sense only if add_vertices is able to describing it.

igraph_set_attribute_table — Attach an attribute table.

```
igraph_attribute_table_t *
igraph_set_attribute_table(const igraph_attribute_table_t * table);
```

This function attaches attribute handling code to the igraph library. Note that the attribute handler table is *not* thread-local even if igraph is compiled in thread-local mode. In the vast majority of cases, this is not a significant restriction.

Arguments:

table: Pointer to an igraph_attribute_table_t object containing the functions for attribute manipulation. Supply NULL here if you don't want attributes.

Returns:

Pointer to the old attribute handling table.

Time complexity: O(1).

igraph_attribute_type_t — The possible types of the attributes.

Note that this is only the type communicated by the attribute interface towards igraph functions. E.g. in the R attribute handler, it is safe to say that all complex R object attributes are strings, as long as this interface is able to serialize them into strings. See also igraph_attribute_table_t.

Values:

IGRAPH_AT- Currently used internally as a "null value" or "placeholder val-TRIBUTE_UNSPECIFIED: ue" in some algorithms. Attribute records with this type must

not be passed to igraph functions.

IGRAPH ATTRIBUTE NU- Numeric attribute.

MERIC:

IGRAPH_AT- Logical values, true or false.

TRIBUTE_BOOLEAN:

IGRAPH ATTRIBUTE STRING: Attribute that can be converted to a string.

IGRAPH_ATTRIBUTE_OBJECT: Custom attribute type, to be used for special data types by client

applications. The R and Python interfaces use this for attributes that hold R or Python objects. Usually ignored by igraph func-

tions.

Handling attribute combination lists

Several graph operations may collapse multiple vertices or edges into a single one. Attribute combination lists are used to indicate to the attribute handler how to combine the attributes of the original

vertices or edges and how to derive the final attribute value that is to be assigned to the collapsed vertex or edge. For example, <code>igraph_simplify()</code> removes loops and combines multiple edges into a single one; in case of a graph with an edge attribute named weight the attribute combination list can tell the attribute handler whether the weight of a collapsed edge should be the sum, the mean or some other function of the weights of the original edges that were collapsed into one.

One attribute combination list may contain several attribute combination records, one for each vertex or edge attribute that is to be handled during the operation.

igraph_attribute_combination_init — Initialize attribute combination list.

igraph_error_t igraph_attribute_combination_init(igraph_attribute_combination_t

Arguments:

comb: The uninitialized attribute combination list.

Returns:

Error code.

Time complexity: O(1)

igraph_attribute_combination_add — Add combination record to attribute combination list.

Arguments:

comb: The attribute combination list.

name: The name of the attribute. If the name already exists the attribute combination record will be replaced. Use NULL to add a default combination record for all atributes not in the list.

type: The type of the attribute combination. See igraph_attribute_combination_type_t for the options.

func: Function to be used if type is IGRAPH_ATTRIBUTE_COMBINE_FUNCTION. This function is called by the concrete attribute handler attached to igraph, and its calling signature depends completely on the attribute handler. For instance, if you are using attributes from C and you have attached the C attribute handler, you need to follow the documentation of the C attribute handler for more details.

Returns:

Error code.

Time complexity: O(n), where n is the number of current attribute combinations.

igraph_attribute_combination_remove — Remove a record from an attribute combination list.

Arguments:

comb: The attribute combination list.

name: The attribute name of the attribute combination record to remove. It will be ignored if the named attribute does not exist. It can be NULL to remove the default combination record.

Returns:

Error code. This currently always returns IGRAPH_SUCCESS.

Time complexity: O(n), where n is the number of records in the attribute combination list.

igraph_attribute_combination_destroy — Destroy attribute combination list.

```
void igraph_attribute_combination_destroy(igraph_attribute_combination_t *comb)
```

Arguments:

comb: The attribute combination list.

Time complexity: O(n), where n is the number of records in the attribute combination list.

igraph_attribute_combination_type_t — The possible types of attribute combinations.

```
typedef enum {
    IGRAPH_ATTRIBUTE_COMBINE_IGNORE = 0,
    IGRAPH_ATTRIBUTE_COMBINE_DEFAULT = 1,
    IGRAPH_ATTRIBUTE_COMBINE_FUNCTION = 2,
    IGRAPH_ATTRIBUTE_COMBINE_SUM = 3,
    IGRAPH_ATTRIBUTE_COMBINE_PROD = 4,
    IGRAPH_ATTRIBUTE_COMBINE_MIN = 5,
    IGRAPH_ATTRIBUTE_COMBINE_MAX = 6,
    IGRAPH_ATTRIBUTE_COMBINE_RANDOM = 7,
    IGRAPH_ATTRIBUTE_COMBINE_FIRST = 8,
    IGRAPH_ATTRIBUTE_COMBINE_LAST = 9,
    IGRAPH_ATTRIBUTE_COMBINE_MEAN = 10,
    IGRAPH_ATTRIBUTE_COMBINE_MEAN = 11,
```

```
IGRAPH_ATTRIBUTE_COMBINE_CONCAT = 12
} igraph_attribute_combination_type_t;
```

Values:

IGRAPH_ATTRIBUTE_COM-Ignore old attributes, use an empty value. BINE_IGNORE: Use the default way to combine attributes (decided by the at-IGRAPH_ATTRIBUTE_COMtribute handler implementation). BINE_DEFAULT: IGRAPH_ATTRIBUTE_COM-Supply your own function to combine attributes. BINE_FUNCTION: Take the sum of the attributes. IGRAPH_ATTRIBUTE_COM-BINE SUM: IGRAPH ATTRIBUTE COM-Take the product of the attributes. BINE_PROD: IGRAPH ATTRIBUTE COM-Take the minimum attribute. BINE_MIN: Take the maximum attribute. IGRAPH_ATTRIBUTE_COM-BINE_MAX: Take a random attribute. IGRAPH_ATTRIBUTE_COM-BINE_RANDOM: Take the first attribute. IGRAPH_ATTRIBUTE_COM-BINE FIRST: IGRAPH_ATTRIBUTE_COM-Take the last attribute. BINE_LAST: Take the mean of the attributes. IGRAPH_ATTRIBUTE_COM-BINE_MEAN: Take the median of the attributes. IGRAPH_ATTRIBUTE_COM-BINE MEDIAN: Concatenate the attributes.

igraph_attribute_combination — Initialize attribute combination list and add records.

```
igraph_error_t igraph_attribute_combination(
        igraph_attribute_combination_t *comb, ...);
```

Arguments:

The uninitialized attribute combination list.

A list of 'name, type[, func]', where: . . .:

IGRAPH_ATTRIBUTE_COM-

BINE_CONCAT:

name: The name of the attribute. If the name already exists the attribute combination record will

be replaced. Use NULL to add a default combination record for all atributes not in the list.

type: The type of the attribute combination. See igraph_attribute_combina-

tion_type_t for the options.

func: Function to be used if type is IGRAPH_ATTRIBUTE_COMBINE_FUNCTION. The list

is closed by setting the name to IGRAPH_NO_MORE_ATTRIBUTES.

Returns:

Error code.

Time complexity: $O(n^2)$, where n is the number attribute combinations records to add.

Example 12.1. File examples/simple/igraph_attribute_combination.c

Accessing attributes from C

There is an experimental attribute handler that can be used from C code. In this section we show how this works. This attribute handler is by default not attached (the default is no attribute handler), so we first need to attach it:

```
igraph_set_attribute_table(&igraph_cattribute_table);
```

Now the attribute functions are available. Please note that the attribute handler must be attached before you call any other igraph functions, otherwise you might end up with graphs without attributes and an active attribute handler, which might cause unexpected program behaviour. The rule is that you attach the attribute handler in the beginning of your main() and never touch it again. Detaching the attribute handler might lead to memory leaks.

It is not currently possible to have attribute handlers on a per-graph basis. All graphs in an application must be managed with the same attribute handler. This also applies to the default case when there is no attribute handler at all.

The C attribute handler supports attaching real numbers, boolean values and character strings as attributes. No vector values are allowed. For example, vertices have a name attribute holding a single string value for each vertex, but it is not possible to have a coords attribute which is a vector of numbers per vertex.

The functions documented in this section are specific to the C attribute handler. Code using these functions will not function when a different attribute handler is attached.

Example 12.2. File examples/simple/cattributes.c

Example 12.3. File examples/simple/cattributes2.c

Example 12.4. File examples/simple/cattributes3.c

Example 12.5. File examples/simple/cattributes4.c

Query attributes

igraph_cattribute_list — List all attributes.

See igraph_attribute_type_t for the various attribute types.

Arguments:

graph: The input graph.

gnames: String vector, the names of the graph attributes.

gtypes: Numeric vector, the types of the graph attributes.

vnames: String vector, the names of the vertex attributes.

vtypes: Numeric vector, the types of the vertex attributes.

enames: String vector, the names of the edge attributes.

etypes: Numeric vector, the types of the edge attributes.

Returns:

Error code.

Naturally, the string vector with the attribute names and the numeric vector with the attribute types are in the right order, i.e. the first name corresponds to the first type, etc. Time complexity: O(Ag+Av+Ae), the number of all attributes.

igraph_cattribute_has_attr — Checks whether a (graph, vertex or edge) attribute exists.

Arguments:

graph: The graph.

type: The type of the attribute, IGRAPH_ATTRIBUTE_GRAPH, IGRAPH_ATTRIBUTE_VERTEX or IGRAPH_ATTRIBUTE_EDGE.

_ _ _

name: Character constant, the name of the attribute.

Returns:

Logical value, true if the attribute exists, false otherwise.

Time complexity: O(A), the number of (graph, vertex or edge) attributes, assuming attribute names are not too long.

igraph_cattribute_GAN — Query a numeric graph attribute.

```
igraph_real_t igraph_cattribute_GAN(const igraph_t *graph, const char *name);
```

Returns the value of the given numeric graph attribute. If the attribute does not exist, a warning is issued and NaN is returned.

Arguments:

graph: The input graph.

name: The name of the attribute to query.

Returns:

The value of the attribute.

See also:

GAN for a simpler interface.

Time complexity: O(Ag), the number of graph attributes.

GAN — Query a numeric graph attribute.

```
#define GAN(graph,n)
```

This is shorthand for igraph_cattribute_GAN().

Arguments:

graph: The graph.

n: The name of the attribute.

Returns:

The value of the attribute.

igraph_cattribute_GAB — Query a boolean graph attribute.

```
igraph_bool_t igraph_cattribute_GAB(const igraph_t *graph, const char *name);
```

Returns the value of the given boolean graph attribute. If the attribute does not exist, a warning is issued and false is returned.

Arguments:

graph: The input graph.

name: The name of the attribute to query.

Returns:

The value of the attribute.

See also:

GAB for a simpler interface.

Time complexity: O(Ag), the number of graph attributes.

GAB — Query a boolean graph attribute.

```
#define GAB(graph,n)
```

This is shorthand for igraph_cattribute_GAB().

Arguments:

graph: The graph.

n: The name of the attribute.

Returns:

The value of the attribute.

igraph_cattribute_GAS — Query a string graph attribute.

```
const char *igraph_cattribute_GAS(const igraph_t *graph, const char *name);
```

Returns a const pointer to the string graph attribute specified in *name*. The value must not be modified. If the attribute does not exist, a warning is issued and an empty string is returned.

Arguments:

graph: The input graph.

name: The name of the attribute to query.

Returns:

The value of the attribute.

See also:

GAS for a simpler interface.

Time complexity: O(Ag), the number of graph attributes.

GAS — Query a string graph attribute.

```
#define GAS(graph,n)
```

This is shorthand for igraph_cattribute_GAS().

Arguments:

graph: The graph.

n: The name of the attribute.

Returns:

The value of the attribute.

igraph_cattribute_VAN — Query a numeric vertex attribute.

If the attribute does not exist, a warning is issued and NaN is returned. See <code>igraph_cattribut-e_VANV()</code> for an error-checked version.

Arguments:

graph: The input graph.

name: The name of the attribute.

vid: The id of the queried vertex.

Returns:

The value of the attribute.

See also:

VAN macro for a simpler interface.

Time complexity: O(Av), the number of vertex attributes.

VAN — Query a numeric vertex attribute.

```
#define VAN(graph,n,v)
```

This is shorthand for igraph_cattribute_VAN().

Arguments:

graph: The graph.

n: The name of the attribute.

v: The id of the vertex.

Returns:

The value of the attribute.

igraph_cattribute_VANV — Query a numeric vertex attribute for many vertices.

Arguments:

graph: The input graph.

name: The name of the attribute.

vids: The vertices to query.

result: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: O(v), where v is the number of vertices in 'vids'.

VANV — Query a numeric vertex attribute for all vertices.

```
#define VANV(graph,n,vec)
```

This is a shorthand for igraph_cattribute_VANV().

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_VAB — Query a boolean vertex attribute.

If the vertex attribute does not exist, a warning is issued and false is returned. See igraph_cattribute VABV() for an error-checked version.

Arguments:

graph: The input graph.

name: The name of the attribute.

vid: The id of the queried vertex.

Returns:

The value of the attribute.

See also:

VAB macro for a simpler interface.

Time complexity: O(Av), the number of vertex attributes.

VAB — Query a boolean vertex attribute.

```
#define VAB(graph,n,v)
```

This is shorthand for igraph_cattribute_VAB().

Arguments:

graph: The graph.

n: The name of the attribute.

v: The id of the vertex.

Returns:

The value of the attribute.

igraph_cattribute_VABV — Query a boolean vertex attribute for many vertices.

Arguments:

graph: The input graph.

name: The name of the attribute.

vids: The vertices to query.

result: Pointer to an initialized boolean vector, the result is stored here. It will be resized, if

needed.

Returns:

Error code.

Time complexity: O(v), where v is the number of vertices in 'vids'.

VABV — Query a boolean vertex attribute for all vertices.

```
#define VABV(graph,n,vec)
```

This is a shorthand for igraph_cattribute_VABV().

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized boolean vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_VAS — Query a string vertex attribute.

Returns a const pointer to the string vertex attribute specified in *name*. The value must not be modified. If the vertex attribute does not exist, a warning is issued and an empty string is returned. See <code>igraph_cattribute_VASV()</code> for an error-checked version.

Arguments:

graph: The input graph.

name: The name of the attribute.

vid: The id of the queried vertex.

Returns:

The value of the attribute.

See also:

The macro VAS for a simpler interface.

Time complexity: O(Av), the number of vertex attributes.

VAS — Query a string vertex attribute.

```
#define VAS(graph,n,v)
```

This is shorthand for igraph_cattribute_VAS().

Arguments:

graph: The graph.

n: The name of the attribute.

v: The id of the vertex.

Returns:

The value of the attribute.

igraph_cattribute_VASV — Query a string vertex attribute for many vertices.

Arguments:

graph: The input graph.

name: The name of the attribute.

vids: The vertices to query.

result: Pointer to an initialized string vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: O(v), where v is the number of vertices in 'vids'. (We assume that the string attributes have a bounded length.)

VASV — Query a string vertex attribute for all vertices.

```
#define VASV(graph,n,vec)
```

This is a shorthand for igraph_cattribute_VASV().

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized string vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_EAN — Query a numeric edge attribute.

If the attribute does not exist, a warning is issued and NaN is returned. See <code>igraph_cattribut-e_EANV()</code> for an error-checked version.

Arguments:

graph: The input graph.

name: The name of the attribute.

eid: The id of the queried edge.

Returns:

The value of the attribute.

See also:

EAN for an easier interface.

Time complexity: O(Ae), the number of edge attributes.

EAN — Query a numeric edge attribute.

```
#define EAN(graph,n,e)
```

This is shorthand for igraph_cattribute_EAN().

Arguments:

graph: The graph.

n: The name of the attribute.

e: The id of the edge.

Returns:

The value of the attribute.

igraph_cattribute_EANV — Query a numeric edge attribute for many edges.

Arguments:

graph: The input graph.

name: The name of the attribute.

eids: The edges to query.

result: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: O(e), where e is the number of edges in 'eids'.

EANV — Query a numeric edge attribute for all edges.

```
#define EANV(graph,n,vec)
```

This is a shorthand for igraph_cattribute_EANV().

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_EAB — Query a boolean edge attribute.

If the edge attribute does not exist, a warning is issued and false is returned. See igraph_cattribute_EABV() for an error-checked version.

Arguments:

graph: The input graph.

name: The name of the attribute.

eid: The id of the queried edge.

Returns:

The value of the attribute.

See also:

EAB for an easier interface.

Time complexity: O(Ae), the number of edge attributes.

EAB — Query a boolean edge attribute.

```
#define EAB(graph,n,e)
```

This is shorthand for igraph_cattribute_EAB().

Arguments:

graph: The graph.

n: The name of the attribute.

e: The id of the edge.

Returns:

The value of the attribute.

igraph_cattribute_EABV — Query a boolean edge attribute for many edges.

Arguments:

graph: The input graph.

name: The name of the attribute.

eids: The edges to query.

result: Pointer to an initialized boolean vector, the result is stored here. It will be resized, if

needed.

Returns:

Error code.

Time complexity: O(e), where e is the number of edges in 'eids'.

EABV — Query a boolean edge attribute for all edges.

```
#define EABV(graph,n,vec)
```

This is a shorthand for igraph_cattribute_EABV().

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

igraph_cattribute_EAS — Query a string edge attribute.

Returns a const pointer to the string edge attribute specified in *name*. The value must not be modified. If the edge attribute does not exist, a warning is issued and an empty string is returned. See <code>igraph_cattribute_EASV()</code> for an error-checked version.

Arguments:

graph: The input graph.

name: The name of the attribute.

eid: The id of the queried edge.

Returns:

The value of the attribute.

 $\$ EAS if you want to type less. Time complexity: O(Ae), the number of edge attributes.

EAS — Query a string edge attribute.

```
#define EAS(graph,n,e)
```

This is shorthand for igraph_cattribute_EAS().

Arguments:

graph: The graph.

n: The name of the attribute.

e: The id of the edge.

Returns:

The value of the attribute.

igraph_cattribute_EASV — Query a string edge attribute for many edges.

Arguments:

graph: The input graph.

name: The name of the attribute.

vids: The edges to query.

result: Pointer to an initialized string vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Time complexity: O(e), where e is the number of edges in 'eids'. (We assume that the string attributes have a bounded length.)

EASV — Query a string edge attribute for all edges.

```
#define EASV(graph,n,vec)
```

This is a shorthand for igraph_cattribute_EASV().

Arguments:

graph: The graph.

n: The name of the attribute.

vec: Pointer to an initialized string vector, the result is stored here. It will be resized, if needed.

Returns:

Error code.

Set attributes

igraph_cattribute_GAN_set — Set a numeric graph attribute.

Arguments:

graph: The graph.

name: Name of the graph attribute. If there is no such attribute yet, then it will be added.

value: The (new) value of the graph attribute.

Returns:

Error code.

\se SETGAN if you want to type less. Time complexity: O(1).

SETGAN — Set a numeric graph attribute

```
#define SETGAN(graph,n,value)
```

This is a shorthand for igraph_cattribute_GAN_set().

Arguments:

graph: The graph.

n: The name of the attribute.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_GAB_set — Set a boolean graph attribute.

Arguments:

graph: The graph.

name: Name of the graph attribute. If there is no such attribute yet, then it will be added.

value: The (new) value of the graph attribute.

Returns:

Error code.

\se SETGAN if you want to type less. Time complexity: O(1).

SETGAB — Set a boolean graph attribute

```
#define SETGAB(graph,n,value)
```

This is a shorthand for igraph_cattribute_GAB_set().

Arguments:

graph: The graph.

n: The name of the attribute.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_GAS_set — Set a string graph attribute.

Arguments:

graph: The graph.

name: Name of the graph attribute. If there is no such attribute yet, then it will be added.

value: The (new) value of the graph attribute. It will be copied.

Returns:

Error code.

\se SETGAS if you want to type less. Time complexity: O(1).

SETGAS — Set a string graph attribute

```
#define SETGAS(graph,n,value)
```

This is a shorthand for igraph_cattribute_GAS_set().

Arguments:

graph: The graph.

n: The name of the attribute.

value: The new value of the attribute.

Returns:

Error code.

igraph cattribute VAN set — Set a numeric vertex attribute.

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all vertices included in *vid*.

Arguments:

graph: The graph.

name: Name of the attribute.

vid: Vertices for which to set the attribute.

value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETVAN for a simpler way.

Time complexity: O(n), the number of vertices if the attribute is new, O(|vid|) otherwise.

SETVAN — Set a numeric vertex attribute

```
#define SETVAN(graph,n,vid,value)
```

This is a shorthand for igraph_cattribute_VAN_set().

Arguments:

graph: The graph.

n: The name of the attribute.

vid: Ids of the vertices to set.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_VAB_set — Set a boolean vertex attribute.

```
igraph_error_t igraph_cattribute_VAB_set(igraph_t *graph, const char *name,
```

```
igraph_integer_t vid, igraph_bool_t value);
```

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all vertices included in *vid*.

Arguments:

graph: The graph.

name: Name of the attribute.

vid: Vertices for which to set the attribute.

value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETVAB for a simpler way.

Time complexity: O(n), the number of vertices if the attribute is new, O(|vid|) otherwise.

SETVAB — Set a boolean vertex attribute

```
#define SETVAB(graph,n,vid,value)
```

This is a shorthand for igraph_cattribute_VAB_set().

Arguments:

graph: The graph.

n: The name of the attribute.

vid: Ids of the vertices to set.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_VAS_set — Set a string vertex attribute.

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all vertices included in *vid*.

Arguments:

graph: The graph.

name: Name of the attribute.

vid: Vertices for which to set the attribute.

value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETVAS for a simpler way.

Time complexity: O(n*l), n is the number of vertices, l is the length of the string to set. If the attribute if not new then only O(|vid|*l).

SETVAS — Set a string vertex attribute

```
#define SETVAS(graph,n,vid,value)
```

This is a shorthand for igraph_cattribute_VAS_set().

Arguments:

graph: The graph.

n: The name of the attribute.

vid: Ids of the vertices to set.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_EAN_set — Set a numeric edge attribute.

The attribute will be added if not present already. If present it will be overwritten. The same value is set for all edges included in vid.

Arguments:

graph: The graph.

name: Name of the attribute.

eid: Edges for which to set the attribute.

value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETEAN for a simpler way.

Time complexity: O(e), the number of edges if the attribute is new, O(|eid|) otherwise.

SETEAN — Set a numeric edge attribute

```
#define SETEAN(graph,n,eid,value)
```

This is a shorthand for igraph_cattribute_EAN_set().

Arguments:

graph: The graph.

n: The name of the attribute.

eid: Ids of the edges to set.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_EAB_set — Set a boolean edge attribute.

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all edges included in *vid*.

Arguments:

graph: The graph.

name: Name of the attribute.

eid: Edges for which to set the attribute.

value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETEAB for a simpler way.

Time complexity: O(e), the number of edges if the attribute is new, O(|eid|) otherwise.

SETEAB — Set a boolean edge attribute

```
#define SETEAB(graph,n,eid,value)
```

This is a shorthand for igraph_cattribute_EAB_set().

Arguments:

graph: The graph.

n: The name of the attribute.

eid: Ids of the edges to set.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_EAS_set — Set a string edge attribute.

The attribute will be added if not present already. If present it will be overwritten. The same *value* is set for all edges included in *vid*.

Arguments:

graph: The graph.

name: Name of the attribute.

eid: Edges for which to set the attribute.

value: The (new) value of the attribute.

Returns:

Error code.

See also:

SETEAS for a simpler way.

Time complexity: $O(e^*l)$, n is the number of edges, l is the length of the string to set. If the attribute if not new then only $O(|eid|^*l)$.

SETEAS — Set a string edge attribute

```
#define SETEAS(graph,n,eid,value)
```

This is a shorthand for igraph_cattribute_EAS_set().

Arguments:

graph: The graph.

n: The name of the attribute.

eid: Ids of the edges to set.

value: The new value of the attribute.

Returns:

Error code.

igraph_cattribute_VAN_setv — Set a numeric vertex attribute for all vertices.

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

v: The new attribute values. The length of this vector must match the number of vertices.

Returns:

Error code.

See also:

SETVANV for a simpler way.

Time complexity: O(n), the number of vertices.

SETVANV — Set a numeric vertex attribute for all vertices

```
#define SETVANV(graph,n,v)
```

This is a shorthand for igraph_cattribute_VAN_setv().

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

Returns:

Error code.

igraph_cattribute_VAB_setv — Set a boolean vertex attribute for all vertices.

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

v: The new attribute values. The length of this boolean vector must match the number of

vertices.

Returns:

Error code.

See also:

SETVANV for a simpler way.

Time complexity: O(n), the number of vertices.

SETVABV — Set a boolean vertex attribute for all vertices

```
#define SETVABV(graph,n,v)
```

This is a shorthand for igraph_cattribute_VAB_setv().

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

Returns:

Error code.

igraph_cattribute_VAS_setv — Set a string vertex attribute for all vertices.

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

sv: String vector, the new attribute values. The length of this vector must match the number

of vertices.

Returns:

Error code.

See also:

SETVASV for a simpler way.

Time complexity: O(n+l), n is the number of vertices, l is the total length of the strings.

SETVASV — Set a string vertex attribute for all vertices

```
#define SETVASV(graph,n,v)
```

This is a shorthand for igraph_cattribute_VAS_setv().

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

Returns:

Error code.

igraph_cattribute_EAN_setv — Set a numeric edge attribute for all edges.

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

v: The new attribute values. The length of this vector must match the number of edges.

Returns:

Error code.

See also:

SETEANV for a simpler way.

Time complexity: O(e), the number of edges.

SETEANV — Set a numeric edge attribute for all edges

```
#define SETEANV(graph,n,v)
```

This is a shorthand for igraph_cattribute_EAN_setv().

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

igraph_cattribute_EAB_setv — Set a boolean edge attribute for all edges.

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

v: The new attribute values. The length of this vector must match the number of edges.

Returns:

Error code.

See also:

SETEABV for a simpler way.

Time complexity: O(e), the number of edges.

SETEABV — Set a boolean edge attribute for all edges

```
#define SETEABV(graph,n,v)
```

This is a shorthand for igraph_cattribute_EAB_setv().

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

igraph_cattribute_EAS_setv — Set a string edge attribute for all edges.

The attribute will be added if not present yet.

Arguments:

graph: The graph.

name: Name of the attribute.

sv: String vector, the new attribute values. The length of this vector must match the number

of edges.

Returns:

Error code.

See also:

SETEASV for a simpler way.

Time complexity: O(e+l), e is the number of edges, l is the total length of the strings.

SETEASV — Set a string edge attribute for all edges

```
#define SETEASV(graph,n,v)
```

This is a shorthand for igraph_cattribute_EAS_setv().

Arguments:

graph: The graph.

n: The name of the attribute.

v: Vector containing the new values of the attributes.

Remove attributes

igraph_cattribute_remove_g — Remove a graph attribute.

```
void igraph_cattribute_remove_g(igraph_t *graph, const char *name);
```

Arguments:

graph: The graph object.

name: Name of the graph attribute to remove.

See also:

DELGA for a simpler way.

DELGA — Remove a graph attribute.

```
#define DELGA(graph,n)
```

A shorthand for igraph_cattribute_remove_g().

Arguments:

graph: The graph.

n: The name of the attribute to remove.

${\tt igraph_cattribute_remove_v-Remove\ a\ vertex\ attribute}.$

```
void igraph_cattribute_remove_v(igraph_t *graph, const char *name);
```

Arguments:

graph: The graph object.

name: Name of the vertex attribute to remove.

See also:

DELVA for a simpler way.

DELVA — Remove a vertex attribute.

```
#define DELVA(graph,n)
A shorthand for igraph_cattribute_remove_v().
Arguments:
graph: The graph.
```

n: The name of the attribute to remove.

igraph_cattribute_remove_e — Remove an edge attribute.

```
void igraph_cattribute_remove_e(igraph_t *graph, const char *name);

Arguments:
graph: The graph object.
name: Name of the edge attribute to remove.

See also:
```

DELEA — Remove an edge attribute.

DELEA for a simpler way.

```
#define DELEA(graph,n)
A shorthand for igraph_cattribute_remove_e().
Arguments:
graph: The graph.
n: The name of the attribute to remove.
```

igraph_cattribute_remove_all — Remove all graph/vertex/edge attributes.

```
void igraph_cattribute_remove_all(igraph_t *graph, igraph_bool_t g,
```

```
igraph_bool_t v, igraph_bool_t e);
```

Arguments:

graph: The graph object.

g: Boolean, whether to remove graph attributes.

v: Boolean, whether to remove vertex attributes.

e: Boolean, whether to remove edge attributes.

See also:

DELGAS, DELVAS, DELEAS, DELALL for simpler ways.

DELGAS — Remove all graph attributes.

```
#define DELGAS(graph)
Calls igraph_cattribute_remove_all().
Arguments:
graph: The graph.
```

DELVAS — Remove all vertex attributes.

```
#define DELVAS(graph)
Calls igraph_cattribute_remove_all().
Arguments:
graph: The graph.
```

DELEAS — Remove all edge attributes.

```
#define DELEAS(graph)
Calls igraph_cattribute_remove_all().
Arguments:
graph: The graph.
```

DELALL — Remove all attributes.

```
#define DELALL(graph)
```

All graph, vertex and edges attributes will be removed. Calls igraph_cattribute_remove_all().

Arguments:

graph: The graph.

Custom attribute combination functions

The C attribute handler supports combining the attributes of multiple vertices of edges into a single attribute during a vertex or edge contraction operation via a user-defined function. This is achieved by setting the type of the attribute combination to IGRAPH_ATTRIBUTE_COMBINE_FUNCTION and passing in a pointer to the custom combination function when specifying attribute combinations in igraph_attribute_combination() or igraph_attribute_combination_add(). For the C attribute handler, the signature of the function depends on the type of the underlying attribute. For numeric attributes, use:

```
igraph_error_t function(const igraph_vector_t *input, igraph_real_t *output);
```

where *input* will receive a vector containing the value of the attribute for all the vertices or edges being combined, and *output* must be filled by the function to the combined value. Similarly, for Boolean attributes, the function takes a boolean vector in *input* and must return the combined Boolean value in *output*:

```
igraph_error_t function(const igraph_vector_bool_t *input, igraph_bool_t *outp
```

For string attributes, the signature is slightly different:

```
igraph_error_t function(const igraph_strvector_t *input, char **output);
```

In case of strings, all strings in the input vector are *owned* by igraph and must not be modified or freed in the combination handler. The string returned to the caller in *output* remains owned by the caller; igraph will make a copy it and store the copy in the appropriate part of the data structure holding the vertex or edge attributes.

Chapter 13. Structural properties of graphs

These functions usually calculate some structural property of a graph, like its diameter, the degree of the nodes, etc.

Basic properties

igraph_are_connected — Decides whether two vertices are connected.

Arguments:

graph: The graph object.

v1: The first vertex.

v2: The second vertex.

res: Boolean, true if there is an edge from v1 to v2, false otherwise.

Returns:

The error code IGRAPH_EINVVID is returned if an invalid vertex ID is given.

The function is of course symmetric for undirected graphs.

Time complexity: O(min(log(d1), log(d2))), d1 is the (out-)degree of v1 and d2 is the (in-)degree of v2

Sparsifiers

igraph_spanner — Calculates a spanner of a graph with a given stretch factor.

A spanner of a graph G = (V,E) with a stretch t is a subgraph H = (V,Es) such that Es is a subset of E and the distance between any pair of nodes in H is at most t times the distance in G. The returned graph is always a spanner of the given graph with the specified stretch. For weighted graphs the number of edges in the spanner is $O(k * n^{(1 + 1/k)})$, where k is k = (stretch + 1)/2, m is the number of edges and n is the number of nodes in G. For unweighted graphs the number of edges is $O(n^{(1 + 1/k)} + kn)$.

This function is based on the algorithm of Baswana and Sen: "A Simple and Linear Time Randomized Algorithm for Computing Sparse Spanners in Weighted Graphs"

Arguments:

graph: An undirected connected graph object. If the graph is directed, the directions of the

edges will be ignored.

spanenr: An initialized vector, the IDs of the edges that constitute the calculated spanner will be

returned here. Use igraph_subgraph_edges() to extract the spanner as a sepa-

rate graph object.

stretch: The stretch factor of the spanner.

weights: The edge weights or NULL.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

Time complexity: The algorithm is a randomized Las Vegas algorithm. The expected running time is O(km) where k is the value mentioned above.

(Shortest)-path related functions

igraph_distances — Length of the shortest paths between vertices.

Arguments:

graph: The graph object.

res: The result of the calculation, a matrix. A pointer to an initialized matrix, to be more precise.

The matrix will be resized if needed. It will have the same number of rows as the length of the from argument, and its number of columns is the number of vertices in the to argument. One row of the matrix shows the distances from/to a given vertex to the ones in

to. For the unreachable vertices IGRAPH_INFINITY is returned.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

mode: The type of shortest paths to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the lengths of the outgoing paths are calculated.

IGRAPH_IN the lengths of the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(n(|V|+|E|)), n is the number of vertices to calculate, |V| and |E| are the number of vertices and edges in the graph.

See also:

igraph_get_shortest_paths() to get the paths themselves, igraph_distances_dijkstra() for the weighted version with non-negative weights, igraph_distances_bellman_ford() if you also have negative weights.

igraph_distances_dijkstra — Weighted shortest path lengths between vertices.

This function implements Dijkstra's algorithm, which can find the weighted shortest path lengths from a source vertex to all other vertices. This function allows specifying a set of source and target vertices. The algorithm is run independently for each source and the results are retained only for the specified targets. This implementation uses a binary heap for efficiency.

Arguments:

graph: The input graph, can be directed.

res: The result, a matrix. A pointer to an initialized matrix should be passed here. The matrix

will be resized as needed. Each row contains the distances from a single source, to the vertices given in the to argument. Unreachable vertices has distance IGRAPH_IN-

FINITY.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to

work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, igraph_dis-

tances() is called.

mode: For directed graphs; whether to follow paths along edge directions (IGRAPH_OUT), or

the opposite (IGRAPH_IN), or ignore edge directions completely (IGRAPH_ALL). It

is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(s^*|E|\log|E|+|V|)$, where |V| is the number of vertices, |E| the number of edges and s the number of sources.

See also:

igraph_distances() for a (slightly) faster unweighted version or igraph_distances_bellman_ford() for a weighted variant that works in the presence of negative edge weights (but no negative loops)

Example 13.1. File examples/simple/dijkstra.c

igraph_distances_bellman_ford — Weighted shortest path lengths between vertices, allowing negative weights.

This function implements the Bellman-Ford algorithm to find the weighted shortest paths to all vertices from a single source, allowing negative weights. It is run independently for the given sources. If there are no negative weights, you are better off with igraph_distances_dijkstra().

Arguments:

graph: The input graph, can be directed.

res: The result, a matrix. A pointer to an initialized matrix should be passed here, the matrix

will be resized if needed. Each row contains the distances from a single source, to all vertices in the graph, in the order of vertex IDs. For unreachable vertices the matrix

contains IGRAPH INFINITY.

from: The source vertices.

to: The target vertices.

weights: The edge weights. There must not be any closed loop in the graph that has a negative

total weight (since this would allow us to decrease the weight of any path containing at least a single vertex of this loop infinitely). Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the

unweighted version, igraph_distances() is called.

mode: For directed graphs; whether to follow paths along edge directions (IGRAPH_OUT), or

the opposite (IGRAPH_IN), or ignore edge directions completely (IGRAPH_ALL). It

is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(s^*|E|^*|V|)$, where |V| is the number of vertices, |E| the number of edges and s the number of sources.

See also:

igraph_distances() for a faster unweighted version or igraph_distances_dijkstra() if you do not have negative edge weights.

Example 13.2. File examples/simple/bellman_ford.c

igraph_distances_johnson — Weighted shortest path lengths between vertices, using Johnson's algorithm.

See Wikipedia at http://en.wikipedia.org/wiki/Johnson's_algorithm for Johnson's algorithm. This algorithm works even if the graph contains negative edge weights, and it is worth using it if we calculate the shortest paths from many sources.

If no edge weights are supplied, then the unweighted version, igraph_distances() is called.

If all the supplied edge weights are non-negative, then Dijkstra's algorithm is used by calling igraph_distances_dijkstra().

Arguments:

graph: The input graph. If negative weights are present, it should be directed.

res: Pointer to an initialized matrix, the result will be stored here, one line for each source

vertex, one column for each target vertex.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

weights: Optional edge weights. If it is a null-pointer, then the unweighted breadth-first search

based igraph_distances() will be called.

Returns:

Error code.

Time complexity: O(s|V|log|V|+|V||E|), |V| and |E| are the number of vertices and edges, s is the number of source vertices.

See also:

igraph_distances() for a faster unweighted version, igraph_distances_dijkstra() if you do not have negative edge weights, igraph_distances_bellman_ford() if you only need to calculate shortest paths from a couple of sources.

igraph_get_shortest_paths — Shortest paths from a vertex.

If there is more than one geodesic between two vertices, this function gives only one of them.

Arguments:

graph: The graph object.

vertices: The result, the IDs of the vertices along the paths. This is a list of integer

vectors where each element is an igraph_vector_int_t object. The list will be resized as needed. Supply a null pointer here if you don't need these

vectors.

edges: The result, the IDs of the edges along the paths. This is a list of integer vectors

where each element is an igraph_vector_int_t object. The list will be resized as needed. Supply a null pointer here if you don't need these vectors.

from: The id of the vertex from/to which the geodesics are calculated.

to: Vertex sequence with the IDs of the vertices to/from which the shortest paths

will be calculated. A vertex might be given multiple times.

mode: The type of shortest paths to be used for the calculation in directed graphs.

Possible values:

IGRAPH_OUT the outgoing paths are calculated.

IGRAPH_IN the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the

computation.

parents: A pointer to an initialized igraph vector or null. If not null, a vector containing

the parent of each vertex in the single source shortest path tree is returned here. The parent of vertex i in the tree is the vertex from which vertex i was reached. The parent of the start vertex (in the from argument) is -1. If the parent is -2, it means that the given vertex was not reached from the source during the search. Note that the search terminates if all the vertices in to are reached.

inbound_edges: A pointer to an initialized igraph vector or null. If not null, a vector contain-

ing the inbound edge of each vertex in the single source shortest path tree is returned here. The inbound edge of vertex i in the tree is the edge via which vertex i was reached. The start vertex and vertices that were not reached during the search will have -1 in the corresponding entry of the vector. Note that the

search terminates if all the vertices in to are reached.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID from is invalid vertex ID

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(|V|+|E|), |V| is the number of vertices, |E| the number of edges in the graph.

See also:

igraph_distances() if you only need the path lengths but not the paths themselves.

Example 13.3. File examples/simple/igraph_get_shortest_paths.c

igraph_get_shortest_path — Shortest path from one vertex to another one.

Calculates and returns a single unweighted shortest path from a given vertex to another one. If there are more than one shortest paths between the two vertices, then an arbitrary one is returned.

This function is a wrapper to <code>igraph_get_shortest_paths()</code>, for the special case when only one target vertex is considered.

Arguments:

graph: The input graph, it can be directed or undirected. Directed paths are considered in

directed graphs.

vertices: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex

IDs along the path are stored here, including the source and target vertices.

edges: Pointer to an initialized vector or a null pointer. If not a null pointer, then the edge

IDs along the path are stored here.

from: The id of the source vertex.

to: The id of the target vertex.

mode: A constant specifying how edge directions are considered in directed graphs. Valid

modes are: IGRAPH_OUT, follows edge directions; IGRAPH_IN, follows the opposite directions; and IGRAPH_ALL, ignores edge directions. This argument is ignored

for undirected graphs.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges in the graph.

See also:

igraph_get_shortest_paths() for the version with more target vertices.

igraph_get_shortest_paths_dijkstra — Weighted shortest paths from a vertex.

If there is more than one path with the smallest weight between two vertices, this function gives only one of them.

Arguments:

graph: The graph object.

vertices: The result, the IDs of the vertices along the paths. This is a list of integer

vectors where each element is an igraph_vector_int_t object. The list will be resized as needed. Supply a null pointer here if you don't need these

vectors.

edges: The result, the IDs of the edges along the paths. This is a list of integer vectors

where each element is an igraph_vector_int_t object. The list will be resized as needed. Supply a null pointer here if you don't need these vectors.

from: The id of the vertex from/to which the geodesics are calculated.

to: Vertex sequence with the IDs of the vertices to/from which the shortest paths

will be calculated. A vertex might be given multiple times. *

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algo-

rithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted ver-

sion, igraph_get_shortest_paths() is called.

mode: The type of shortest paths to be use for the calculation in directed graphs.

Possible values:

IGRAPH_OUT the outgoing paths are calculated.

IGRAPH_IN the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the

computation.

parents: A pointer to an initialized igraph vector or null. If not null, a vector containing

the parent of each vertex in the single source shortest path tree is returned here.

The parent of vertex i in the tree is the vertex from which vertex i was reached. The parent of the start vertex (in the from argument) is -1. If the parent is -2, it means that the given vertex was not reached from the source during the search. Note that the search terminates if all the vertices in to are reached.

inbound_edges:

A pointer to an initialized igraph vector or null. If not null, a vector containing the inbound edge of each vertex in the single source shortest path tree is returned here. The inbound edge of vertex i in the tree is the edge via which vertex i was reached. The start vertex and vertices that were not reached during the search will have -1 in the corresponding entry of the vector. Note that the search terminates if all the vertices in to are reached.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID from is invalid vertex ID

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|E|\log|E|+|V|)$, where |V| is the number of vertices and |E| is the number of edges

See also:

igraph_distances_dijkstra() if you only need the path length but not the paths themselves, igraph_get_shortest_paths() if all edge weights are equal.

Example 13.4. File examples/simple/igraph_get_shortest_paths_dijkstra.c

igraph_get_shortest_path_dijkstra — Weighted shortest path from one vertex to another one.

Calculates a single (positively) weighted shortest path from a single vertex to another one, using Dijk-stra's algorithm.

This function is a special case (and a wrapper) to igraph_get_shortest_paths_dijk-stra().

Arguments:

graph: The input graph, it can be directed or undirected.

vertices: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex

IDs along the path are stored here, including the source and target vertices.

edges: Pointer to an initialized vector or a null pointer. If not a null pointer, then the edge

IDs along the path are stored here.

from: The id of the source vertex.

to: The id of the target vertex.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algo-

rithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version,

igraph_get_shortest_paths() is called.

mode: A constant specifying how edge directions are considered in directed graphs.

IGRAPH_OUT follows edge directions, IGRAPH_IN follows the opposite directions, and IGRAPH_ALL ignores edge directions. This argument is ignored for undirected

graphs.

Returns:

Error code.

Time complexity: $O(|E|\log|E|+|V|)$, |V| is the number of vertices, |E| is the number of edges in the graph.

See also:

igraph_get_shortest_paths_dijkstra() for the version with more target vertices.

igraph_get_shortest_paths_bellman_ford — Weighted shortest paths from a vertex, allowing negative weights.

This function calculates weighted shortest paths from or to a single vertex, and allows negative weights. When there is more than one shortest path between two vertices, only one of them is returned. If there are no negative weights, you are better off with <code>igraph_get_shortest_paths_dijk-stra()</code>.

Arguments:

graph: The input graph, can be directed.

vertices: The result, the IDs of the vertices along the paths. This is a list of integer

vectors where each element is an igraph_vector_int_t object. The list will be resized as needed. Supply a null pointer here if you don't need these

vectors.

edges: The result, the IDs of the edges along the paths. This is a list of integer vectors

where each element is an igraph_vector_int_t object. The list will be resized as needed. Supply a null pointer here if you don't need these vectors.

from: The id of the vertex from/to which the geodesics are calculated.

to: Vertex sequence with the IDs of the vertices to/from which the shortest paths

will be calculated. A vertex might be given multiple times.

weights: The edge weights. There must not be any closed loop in the graph that has a

negative total weight (since this would allow us to decrease the weight of any path containing at least a single vertex of this loop infinitely). If this is a null pointer, then the unweighted version, igraph_get_shortest_paths()

is called.

mode: For directed graphs; whether to follow paths along edge directions

(IGRAPH_OUT), or the opposite (IGRAPH_IN), or ignore edge directions

completely (IGRAPH_ALL). It is ignored for undirected graphs.

parents: A pointer to an initialized igraph vector or null. If not null, a vector containing

the parent of each vertex in the single source shortest path tree is returned here. The parent of vertex i in the tree is the vertex from which vertex i was reached. The parent of the start vertex (in the from argument) is -1. If the parent is -2, it means that the given vertex was not reached from the source during the search. Note that the search terminates if all the vertices in to are reached.

inbound_edges: A pointer to an initialized igraph vector or null. If not null, a vector contain-

ing the inbound edge of each vertex in the single source shortest path tree is returned here. The inbound edge of vertex i in the tree is the edge via which vertex i was reached. The start vertex and vertices that were not reached during the search will have -1 in the corresponding entry of the vector. Note that the

search terminates if all the vertices in to are reached.

Returns:

Error code:

IGRAPH_ENOMEM Not enough memory for temporary data.

IGRAPH_EINVAL The weight vector doesn't math the number of edges.

IGRAPH_EINVVID from is invalid vertex ID

IGRAPH_ENEGLOOP Bellman-ford algorithm encounted a negative loop.

Time complexity: $O(|E|^*|V|)$, where |V| is the number of vertices, |E| the number of edges.

See also:

igraph_get_shortest_paths() for a faster unweighted version or igraph_get_shortest_paths_dijkstra() if you do not have negative edge weights.

igraph_get_shortest_path_bellman_ford — Weighted shortest path from one vertex to another one.

Calculates a single (positively) weighted shortest path from a single vertex to another one, using Bellman-Ford algorithm.

This function is a special case (and a wrapper) to igraph_get_shortest_paths_bell-man ford().

Arguments:

graph: The input graph, it can be directed or undirected.

vertices: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex

IDs along the path are stored here, including the source and target vertices.

edges: Pointer to an initialized vector or a null pointer. If not a null pointer, then the edge

IDs along the path are stored here.

from: The id of the source vertex.

to: The id of the target vertex.

weights: The edge weights. There must not be any closed loop in the graph that has a negative

total weight (since this would allow us to decrease the weight of any path containing at least a single vertex of this loop infinitely). If this is a null pointer, then the unweighted

version is called.

mode: A constant specifying how edge directions are considered in directed graphs.

IGRAPH_OUT follows edge directions, IGRAPH_IN follows the opposite directions, and IGRAPH_ALL ignores edge directions. This argument is ignored for undirected

graphs.

Returns:

Error code.

Time complexity: $O(|E|\log|E|+|V|)$, |V| is the number of vertices, |E| is the number of edges in the graph.

See also:

igraph_get_shortest_paths_bellman_ford() for the version with more target vertices

igraph_get_all_shortest_paths — All shortest paths (geodesics) from a vertex.

igraph_neimode_t mode);

When there is more than one shortest path between two vertices, all of them will be returned.

Arguments:

graph: The graph object.

vertices: The result, the IDs of the vertices along the paths. This is a list of integer vectors where

each element is an <code>igraph_vector_int_t</code> object. Each vector object contains the vertices along a shortest path from <code>from</code> to another vertex. The vectors are ordered according to their target vertex: first the shortest paths to vertex 0, then to vertex 1, etc. No data is included for unreachable vertices. The list will be resized as needed.

Supply a null pointer here if you don't need these vectors.

edges: The result, the IDs of the edges along the paths. This is a list of integer vectors where

each element is an <code>igraph_vector_int_t</code> object. Each vector object contains the edges along a shortest path from <code>from</code> to another vertex. The vectors are ordered according to their target vertex: first the shortest paths to vertex 0, then to vertex 1, etc. No data is included for unreachable vertices. The list will be resized as needed.

Supply a null pointer here if you don't need these vectors.

nrgeo: Pointer to an initialized igraph_vector_int_t object or NULL. If not NULL the

number of shortest paths from *from* are stored here for every vertex in the graph. Note that the values will be accurate only for those vertices that are in the target vertex sequence (see *to*), since the search terminates as soon as all the target vertices have

been found.

from: The id of the vertex from/to which the geodesics are calculated.

to: Vertex sequence with the IDs of the vertices to/from which the shortest paths will be

calculated. A vertex might be given multiple times.

mode: The type of shortest paths to be use for the calculation in directed graphs. Possible

values:

IGRAPH_OUT the lengths of the outgoing paths are calculated.

IGRAPH_IN the lengths of the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the com-

putation.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

 ${\tt IGRAPH_EINVVID} \qquad \textit{from} \ is \ invalid \ vertex \ ID.$

IGRAPH_EINVMODE invalid mode argument.

Added in version 0.2.

Time complexity: O(|V|+|E|) for most graphs, $O(|V|^2)$ in the worst case.

igraph_get_all_shortest_paths_dijkstra — All weighted shortest paths (geodesics) from a vertex.

Arguments:

graph: The graph object.

vertices: Pointer to an initialized integer vector list or NULL. If not NULL, then each vector

object contains the vertices along a shortest path from from to another vertex. The vectors are ordered according to their target vertex: first the shortest paths to vertex

0, then to vertex 1, etc. No data is included for unreachable vertices.

edges: Pointer to an initialized integer vector list or NULL. If not NULL, then each vector

object contains the edges along a shortest path from from to another vertex. The vectors are ordered according to their target vertex: first the shortest paths to vertex

0, then to vertex 1, etc. No data is included for unreachable vertices.

nrgeo: Pointer to an initialized igraph_vector_int_t object or NULL. If not NULL the number

of shortest paths from from are stored here for every vertex in the graph. Note that the values will be accurate only for those vertices that are in the target vertex sequence (see to), since the search terminates as soon as all the target vertices have been found.

from: The id of the vertex from/to which the geodesics are calculated.

to: Vertex sequence with the IDs of the vertices to/from which the shortest paths will be

calculated. A vertex might be given multiple times.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to

work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, <code>igraph_get_al-</code>

l_shortest_paths() is called.

mode: The type of shortest paths to be use for the calculation in directed graphs. Possible

values:

IGRAPH_OUT the outgoing paths are calculated.

IGRAPH IN the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the com-

putation.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID from is an invalid vertex ID

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|E|\log|E|+|V|)$, where |V| is the number of vertices and |E| is the number of edges

See also:

igraph_distances_dijkstra() if you only need the path length but not the paths themselves, igraph_get_all_shortest_paths() if all edge weights are equal.

Example 13.5. File examples/simple/igraph_get_all_shortest_paths_dijkstra.c

igraph_get_k_shortest_paths — k shortest paths between two vertices.

```
igraph_error_t igraph_get_k_shortest_paths(
    const igraph_t *graph, const igraph_vector_t *weights,
    igraph_vector_int_list_t *vertex_paths,
    igraph_vector_int_list_t *edge_paths,
    igraph_integer_t k, igraph_integer_t from, igraph_integer_t to,
    igraph_neimode_t mode
);
```

Arguments:

graph: The graph object.

weights: The edge weights of the graph. Can be NULL for an unweighted graph. Infinite

weights will be treated as missing edges.

vertex_paths: Pointer to an initialized list of integer vectors, the result will be stored here in

igraph_vector_int_t objects. Each vector object contains the vertex IDs along the kth shortest path between *from* and *to*, where k is the vector list

index. May be NULL if the vertex paths are not needed.

edge_paths: Pointer to an initialized list of integer vectors, the result will be stored here in

igraph_vector_int_t objects. Each vector object contains the edge IDs along the kth shortest path between from and to, where k is the vector list

index. May be NULL if the edge paths are not needed.

k: The number of paths.

from: The ID of the vertex from which the paths are calculated.

to: The ID of the vertex to which the paths are calculated.

mode: The type of paths to be used for the calculation in directed graphs. Possible

values:

IGRAPH_OUT The outgoing paths of from are calculated.

 $IGRAPH_IN$ The incoming paths of from are calculated.

IGRAPH_ALL The directed graph is considered as an undirected one for the

computation.

Returns:

Error code:

```
IGRAPH_ENOMEM Not enough memory for temporary data.

IGRAPH_EINVVID from or to is an invalid vertex id.

IGRAPH_EINVMODE Invalid mode argument.

IGRAPH_EINVAL Invalid argument.
```

Time complexity: $k |V| (|V| \log |V| + |E|)$, where |V| is the number of vertices, and |E| is the number of edges.

igraph_get_all_simple_paths — List all simple paths from one source.

A path is simple if its vertices are unique, i.e. no vertex is visited more than once.

Note that potentially there are exponentially many paths between two vertices of a graph, and you may run out of memory when using this function, if your graph is lattice-like.

This function currently ignored multiple and loop edges.

Arguments:

graph: The input graph.

res: Initialized integer vector, all paths are returned here, separated by -1 markers. The paths

are included in arbitrary order, as they are found.

from: The start vertex.

to: The target vertices.

cutoff: Maximum length of path that is considered. If negative, paths of all lengths are consid-

ered.

mode: The type of the paths to consider, it is ignored for undirected graphs.

Returns:

Error code.

Time complexity: O(n!) in the worst case, n is the number of vertices.

igraph_average_path_length — Calculates the average unweighted shortest path length between all vertex pairs.

If no vertex pairs can be included in the calculation, for example because the graph has fewer than two vertices, or if the graph has no edges and unconn is set to true, NaN is returned.

Arguments:

graph: The graph object.

res: Pointer to a real number, this will contain the result.

unconn_pairs: Pointer to a real number. If not a null pointer, the number of ordered vertex pairs

where the second vertex is unreachable from the first one will be stored here.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

unconn: What to do if the graph is not connected. If true, only those vertex pairs

will be included in the calculation between which there is a path. If false,

IGRAPH_INFINITY is returned for disconnected graphs.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for data structures

Time complexity: O(|V||E|), the number of vertices times the number of edges.

See also:

igraph_average_path_length_dijkstra() for the weighted version.

Example 13.6. File examples/simple/igraph_average_path_length.c

igraph_average_path_length_dijkstra — Calculates the average weighted shortest path length between all vertex pairs.

If no vertex pairs can be included in the calculation, for example because the graph has fewer than two vertices, or if the graph has no edges and unconn is set to true, NaN is returned.

All distinct ordered vertex pairs are taken into account.

Arguments:

graph: The graph object.

res: Pointer to a real number, this will contain the result.

unconn_pairs: Pointer to a real number. If not a null pointer, the number of ordered vertex pairs

where the second vertex is unreachable from the first one will be stored here.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algo-

rithm to work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version,

igraph_average_path_length() is called.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

unconn: If true, only those pairs are considered for the calculation between which there

is a path. If false, IGRAPH_INFINITY is returned for disconnected graphs.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for data structures

IGRAPH_EINVAL invalid weight vector

Time complexity: $O(|V| |E| \log |E| + |V|)$, where |V| is the number of vertices and |E| is the number of edges.

See also:

igraph_average_path_length() for a slightly faster unweighted version.

Example 13.7. File examples/simple/igraph_grg_game.c

igraph_path_length_hist — Create a histogram of all shortest path lengths.

This function calculates a histogram, by calculating the shortest path length between each pair of vertices. For directed graphs both directions might be considered and then every pair of vertices appears twice in the histogram.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the result is stored here. The first (i.e. zeroth)

element contains the number of shortest paths of length 1, etc. The supplied vector

is resized as needed.

unconnected: Pointer to a real number, the number of pairs for which the second vertex is not

reachable from the first is stored here.

directed: Whether to consider directed paths in a directed graph (if not zero). This argument

is ignored for undirected graphs.

Returns:

Error code.

Time complexity: O(|V||E|), the number of vertices times the number of edges.

See also:

igraph_average_path_length() and igraph_distances()

igraph_diameter — Calculates the diameter of a graph (longest geodesic).

The diameter of a graph is the length of the longest shortest path it has. This function computes both the diameter, as well as the corresponding path. The diameter of the null graph is considered be infinity by convention. If the graph has no vertices, IGRAPH NAN is returned.

Arguments:

graph: The graph object.

res: Pointer to a real number, if not NULL then it will contain the diameter (the actual

distance).

from: Pointer to an integer, if not NULL it will be set to the source vertex of the diameter

path. If the graph has no diameter path, it will be set to -1.

to: Pointer to an integer, if not NULL it will be set to the target vertex of the diameter

path. If the graph has no diameter path, it will be set to -1.

vertex_path: Pointer to an initialized vector. If not NULL the actual longest geodesic path in

terms of vertices will be stored here. The vector will be resized as needed.

edge_path: Pointer to an initialized vector. If not NULL the actual longest geodesic path in

terms of edges will be stored here. The vector will be resized as needed.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

unconn: What to do if the graph is not connected. If true the longest geodesic within a

component will be returned, otherwise IGRAPH_INFINITY is returned.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: O(|V||E|), the number of vertices times the number of edges.

See also:

```
igraph_diameter_dijkstra()
```

Example 13.8. File examples/simple/igraph_diameter.c

igraph_diameter_dijkstra — Calculates the weighted diameter of a graph using Dijkstra's algorithm.

This function computes the weighted diameter of a graph. If the graph has no vertices, IGRAPH_NAN is returned.

Arguments:

graph: The input graph, can be directed or undirected.

weights: The edge weights of the graph. Can be NULL for an unweighted graph.

res: Pointer to a real number, if not NULL then it will contain the diameter (the actual

distance).

from: Pointer to an integer, if not NULL it will be set to the source vertex of the diameter

path. If the graph has no diameter path, it will be set to -1.

to: Pointer to an integer, if not NULL it will be set to the target vertex of the diameter

path. If the graph has no diameter path, it will be set to -1.

vertex_path: Pointer to an initialized vector. If not NULL the actual longest geodesic path in

terms of vertices will be stored here. The vector will be resized as needed.

edge_path: Pointer to an initialized vector. If not NULL the actual longest geodesic path in

terms of edges will be stored here. The vector will be resized as needed.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

unconn: What to do if the graph is not connected. If true the longest geodesic within a

component will be returned, otherwise IGRAPH INFINITY is returned.

Returns:

Error code.

Time complexity: O(|V||E|*log|E|), |V| is the number of vertices, |E| is the number of edges.

See also:

```
igraph_diameter()
```

igraph_girth — The girth of a graph is the length of the shortest cycle in it.

The current implementation works for undirected graphs only, directed graphs are treated as undirected graphs. Self-loops and multiple edges are ignored.

For graphs that contain no cycles, and only for such graphs, infinity is returned.

This implementation is based on Alon Itai and Michael Rodeh: Finding a minimum circuit in a graph *Proceedings of the ninth annual ACM symposium on Theory of computing*, 1-10, 1977. The first implementation of this function was done by Keith Briggs, thanks Keith.

Arguments:

graph: The input graph.

girth: Pointer to an igraph_real_t, if not NULL then the result will be stored here.

circle: Pointer to an initialized vector, the vertex IDs in the shortest circle will be stored here.

If NULL then it is ignored.

Returns:

Error code.

Time complexity: $O((|V|+|E|)^2)$, |V| is the number of vertices, |E| is the number of edges in the general case. If the graph has no cycles at all then the function needs O(|V|+|E|) time to realize this and then it stops.

Example 13.9. File examples/simple/igraph_girth.c

igraph_eccentricity — Eccentricity of some vertices.

The eccentricity of a vertex is calculated by measuring the shortest distance from (or to) the vertex, to (or from) all vertices in the graph, and taking the maximum.

This implementation ignores vertex pairs that are in different components. Isolated vertices have eccentricity zero.

Arguments:

graph: The input graph, it can be directed or undirected.

res: Pointer to an initialized vector, the result is stored here.

vids: The vertices for which the eccentricity is calculated.

mode: What kind of paths to consider for the calculation: IGRAPH_OUT, paths that follow edge

directions; IGRAPH_IN, paths that follow the opposite directions; and IGRAPH_ALL,

paths that ignore edge directions. This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(v^*(|V|+|E|))$, where |V| is the number of vertices, |E| is the number of edges and v is the number of vertices for which eccentricity is calculated.

See also:

```
igraph_radius().
```

Example 13.10. File examples/simple/igraph_eccentricity.c

igraph_eccentricity_dijkstra — Eccentricity of some vertices, using weighted edges.

The eccentricity of a vertex is calculated by measuring the shortest distance from (or to) the vertex, to (or from) all vertices in the graph, and taking the maximum.

This implementation ignores vertex pairs that are in different components. Isolated vertices have eccentricity zero.

Arguments:

graph: The input graph, it can be directed or undirected.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to

work. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, igraph_eccen-

tricity() is called.

res: Pointer to an initialized vector, the result is stored here.

vids: The vertices for which the eccentricity is calculated.

mode: What kind of paths to consider for the calculation: IGRAPH_OUT, paths that fol-

low edge directions; IGRAPH_IN, paths that follow the opposite directions; and IGRAPH_ALL, paths that ignore edge directions. This argument is ignored for undi-

rected graphs.

Returns:

Error code.

igraph_graph_center — Central vertices of a graph.

```
igraph_error_t igraph_graph_center(
    const igraph_t *graph, igraph_vector_int_t *res, igraph_neimode_t mode
);
```

The central vertices of a graph are calculated by finding the vertices with the minimum eccentricity. This concept is typically applied to connected graphs. In undirected disconnected graphs, the calculation is effectively done per connected component.

Arguments:

graph: The input graph, it can be directed or undirected.

res: Pointer to an initialized vector, the result is stored here.

mode: What kind of paths to consider for the calculation: IGRAPH_OUT, paths that follow edge

directions; IGRAPH_IN, paths that follow the opposite directions; and IGRAPH_ALL, paths that ignore edge directions. This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: O(|V|(|V|+|E|)), where |V| is the number of vertices and |E| is the number of edges.

See also:

igraph_eccentricity().

igraph_radius — Radius of a graph.

The radius of a graph is the defined as the minimum eccentricity of its vertices, see igraph_eccentricity().

Arguments:

graph: The input graph, it can be directed or undirected.

radius: Pointer to a real variable, the result is stored here.

mode: What kind of paths to consider for the calculation: IGRAPH_OUT, paths that follow edge

directions; IGRAPH_IN, paths that follow the opposite directions; and IGRAPH_ALL, paths that ignore edge directions. This argument is ignored for undirected graphs.

Returns:

Error code.

Time complexity: O(|V|(|V|+|E|)), where |V| is the number of vertices and |E| is the number of edges.

See also:

```
igraph_eccentricity().
```

Example 13.11. File examples/simple/igraph_radius.c

igraph_pseudo_diameter — Approximation and lower bound of diameter.

This algorithm finds a pseudo-peripheral vertex and returns its eccentricity. This value can be used as an approximation and lower bound of the diameter of a graph.

A pseudo-peripheral vertex is a vertex v, such that for every vertex u which is as far away from v as possible, v is also as far away from u as possible. The process of finding one depends on where the search starts, and for a disconnected graph the maximum diameter found will be that of the component vid_start is in.

Arguments:

graph: The input graph, if it is directed, its edge directions are ignored.

diameter: Pointer to a real variable, the result is stored here.

vid_start: Id of the starting vertex. If this is negative, a random starting vertex is chosen.

from: Pointer to an integer, if not NULL it will be set to the source vertex of the diameter

path. If unconn is false, and a disconnected graph is detected, this is set to -1.

to: Pointer to an integer, if not NULL it will be set to the target vertex of the diameter

path. If unconn is false, and a disconnected graph is detected, this is set to -1.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

unconn: What to do if the graph is not connected. If true the longest geodesic within a

component will be returned, otherwise IGRAPH_INFINITY is returned.

Returns:

Error code.

Time complexity: O(|V||E|)), where |V| is the number of vertices and |E| is the number of edges.

See also:

```
igraph_eccentricity(), igraph_diameter().
```

igraph_vertex_path_from_edge_path — Converts a path of edge IDs to the traversed vertex IDs.

```
igraph_error_t igraph_vertex_path_from_edge_path(
   const igraph_t *graph, igraph_integer_t start,
   const igraph_vector_int_t *edge_path, igraph_vector_int_t *vertex_path,
   igraph_neimode_t mode
);
```

This function is useful when you have a sequence of edge IDs representing a continuous path in a graph and you would like to obtain the vertex IDs that the path traverses. The function is used implicitly by several shortest path related functions to convert a path of edge IDs to the corresponding representation that describes the path in terms of vertex IDs instead.

Arguments:

graph: the graph that the edge IDs refer to

start: the start vertex of the path

edge_path: the sequence of edge IDs that describe the path

vertex_path: the sequence of vertex IDs traversed will be returned here

Returns:

Error code: IGRAPH_ENOMEM if there is not enough memory, IGRAPH_EINVAL if the edge path does not start at the given vertex or if there is at least one edge whose start vertex does not match the end vertex of the previous edge

Widest-path related functions

igraph_get_widest_path — Widest path from one vertex to another one.

Calculates a single widest path from a single vertex to another one, using Dijkstra's algorithm.

This function is a special case (and a wrapper) to igraph_get_widest_paths().

Arguments:

graph: The input graph, it can be directed or undirected.

vertices: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex

IDs along the path are stored here, including the source and target vertices.

edges: Pointer to an initialized vector or a null pointer. If not a null pointer, then the edge

IDs along the path are stored here.

from: The id of the source vertex.

to: The id of the target vertex.

weights: The edge weights. Edge weights can be negative. If this is a null pointer or if any edge

weight is NaN, then an error is returned.

mode: A constant specifying how edge directions are considered in directed graphs.

IGRAPH_OUT follows edge directions, IGRAPH_IN follows the opposite directions, and IGRAPH_ALL ignores edge directions. This argument is ignored for undirected

graphs.

Returns:

Error code.

 $Time\ complexity:\ O(|E|log|E|+|V|),\ |V|\ is\ the\ number\ of\ vertices,\ |E|\ is\ the\ number\ of\ edges\ in\ the\ graph.$

See also:

igraph_get_widest_paths() for the version with more target vertices.

igraph_get_widest_paths — Widest paths from a single vertex.

Calculates the widest paths from a single node to all other specified nodes, using a modified Dijkstra's algorithm. If there is more than one path with the largest width between two vertices, this function gives only one of them.

Arguments:

graph: The graph object.

vertices: The result, the IDs of the vertices along the paths. This is a list of integer

vectors where each element is an igraph_vector_int_t object. The list will be resized as needed. Supply a null pointer here if you don't need these

vectors.

edges: The result, the IDs of the edges along the paths. This is a list of integer vectors

where each element is an igraph_vector_int_t object. The list will be resized as needed. Supply a null pointer here if you don't need these vectors.

from: The id of the vertex from/to which the widest paths are calculated.

to: Vertex sequence with the IDs of the vertices to/from which the widest paths

will be calculated. A vertex might be given multiple times.

weights: The edge weights. Edge weights can be negative. If this is a null pointer or if

any edge weight is NaN, then an error is returned.

mode: The type of widest paths to be used for the calculation in directed graphs.

Possible values:

IGRAPH_OUT the outgoing paths are calculated.

IGRAPH_IN the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the

computation.

parents: A pointer to an initialized igraph vector or null. If not null, a vector containing

the parent of each vertex in the single source widest path tree is returned here. The parent of vertex i in the tree is the vertex from which vertex i was reached. The parent of the start vertex (in the from argument) is -1. If the parent is -2, it means that the given vertex was not reached from the source during the search. Note that the search terminates if all the vertices in to are reached.

inbound_edges: A pointer to an initialized igraph vector or null. If not null, a vector containing

the inbound edge of each vertex in the single source widest path tree is returned here. The inbound edge of vertex i in the tree is the edge via which vertex i was reached. The start vertex and vertices that were not reached during the search will have -1 in the corresponding entry of the vector. Note that the search

terminates if all the vertices in to are reached.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID from is invalid vertex ID

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|E|\log|E|+|V|)$, where |V| is the number of vertices in the graph and |E| is the number of edges

See also:

igraph_widest_path_widths_dijkstra() or igraph_widest_path_widths_floyd_warshall() if you only need the widths of the paths but not the paths themselves.

igraph_widest_path_widths_dijkstra — Widths of widest paths between vertices.

```
const igraph_vs_t to,
const igraph_vector_t *weights,
igraph neimode t mode);
```

This function implements a modified Dijkstra's algorithm, which can find the widest path widths from a source vertex to all other vertices. This function allows specifying a set of source and target vertices. The algorithm is run independently for each source and the results are retained only for the specified targets. This implementation uses a binary heap for efficiency.

Arguments:

graph: The input graph, can be directed.

res: The result, a matrix. A pointer to an initialized matrix should be passed here. The matrix

will be resized as needed. Each row contains the widths from a single source, to the vertices given in the to argument. Unreachable vertices have width IGRAPH_NEGIN-FINITY, and vertices have a width of IGRAPH_POSINFINITY to themselves.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

weights: The edge weights. Edge weights can be negative. If this is a null pointer or if any edge

weight is NaN, then an error is returned.

mode: For directed graphs; whether to follow paths along edge directions (IGRAPH_OUT), or

the opposite (IGRAPH_IN), or ignore edge directions completely (IGRAPH_ALL). It

is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(s^*(|E|log|E|+|V|))$, where |V| is the number of vertices in the graph, |E| the number of edges and s the number of sources.

See also:

 $igraph_widest_path_widths_floyd_warshall()$ for a variant that runs faster on dense graphs.

igraph_widest_path_widths_floyd_warshall — Widths of widest paths between vertices.

This function implements a modified Floyd Warshalls algorithm, to find the widest path widths from a set of source vertices to all other target vertices.

Arguments:

graph: The input graph, can be directed.

res: The result, a matrix. A pointer to an initialized matrix should be passed here. The matrix

will be resized as needed. Each row contains the widths from a single source, to the vertices given in the to argument. Unreachable vertices have width IGRAPH_NEGIN-FINITY, and vertices have a width of IGRAPH_POSINFINITY to themselves.

from: The source vertices.

to: The target vertices. It is not allowed to include a vertex twice or more.

weights: The edge weights. Edge weights can be negative. If this is a null pointer or if any edge

weight is NaN, then an error is returned.

mode: For directed graphs; whether to follow paths along edge directions (IGRAPH_OUT), or

the opposite (IGRAPH_IN), or ignore edge directions completely (IGRAPH_ALL). It

is ignored for undirected graphs.

Returns:

Error code.

Time complexity: $O(|V|^{\Lambda}3)$, where |V| is the number of vertices in the graph.

See also:

igraph_widest_path_widths_dijkstra() for a variant that runs faster on sparse graphs.

Efficiency measures

igraph_global_efficiency — Calculates the global efficiency of a network.

The global efficiency of a network is defined as the average of inverse distances between all pairs of vertices: $E_g = 1/(N*(N-1))$ sum_{i.l. = j} 1/d_i, where N is the number of vertices. The inverse distance between pairs that are not reachable from each other is considered to be zero. For graphs with fewer than 2 vertices, NaN is returned.

Reference: V. Latora and M. Marchiori, Efficient Behavior of Small-World Networks, Phys. Rev. Lett. 87, 198701 (2001). https://dx.doi.org/10.1103/PhysRevLett.87.198701

Arguments:

graph: The graph object.

res: Pointer to a real number, this will contain the result.

weights: The edge weights. All edge weights must be non-negative for Dijkstra's algorithm to

work. Additionally, no edge weight may be NaN. If either case does not hold, an error

is returned. If this is a null pointer, then the unweighted version, igraph_average_path_length() is used in calculating the global efficiency.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for data structures

IGRAPH_EINVAL invalid weight vector

Time complexity: $O(|V| |E| \log |E| + |V|)$ for weighted graphs and O(|V| |E|) for unweighted ones. |V| denotes the number of vertices and |E| denotes the number of edges.

igraph_local_efficiency — Calculates the local efficiency around each vertex in a network.

The local efficiency of a network around a vertex is defined as follows: We remove the vertex and compute the distances (shortest path lengths) between its neighbours through the rest of the network. The local efficiency around the removed vertex is the average of the inverse of these distances.

The inverse distance between two vertices which are not reachable from each other is considered to be zero. The local efficiency around a vertex with fewer than two neighbours is taken to be zero by convention.

Reference: I. Vragovi#, E. Louis, and A. Díaz-Guilera, Efficiency of informational transfer in regular and complex networks, Phys. Rev. E 71, 1 (2005). http://dx.doi.org/10.1103/PhysRevE.71.036122

Arguments:

graph: The graph object.

res: Pointer to an initialized vector, this will contain the result.

vids: The vertices around which the local efficiency will be calculated.

weights: The edge weights. All edge weights must be non-negative. Additionally, no edge

weight may be NaN. If either case does not hold, an error is returned. If this is a null pointer, then the unweighted version, igraph_average_path_length()

is called.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

mode: How to determine the local neighborhood of each vertex in directed graphs. Ignored

in undirected graphs.

IGRAPH_ALL take both in- and out-neighbours; this is a reasonable default for

high-level interfaces.

IGRAPH_OUT take only out-neighbours

IGRAPH_IN take only in-neighbours

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for data structures

IGRAPH_EINVAL invalid weight vector

Time complexity: $O(|E|^2 \log |E|)$ for weighted graphs and $O(|E|^2)$ for unweighted ones. |E| denotes the number of edges.

See also:

igraph_average_local_efficiency()

igraph_average_local_efficiency — Calculates the average local efficiency in a network.

For the null graph, zero is returned by convention.

Arguments:

graph: The graph object.

res: Pointer to a real number, this will contain the result.

weights: The edge weights. They must be all non-negative. If a null pointer is given, all weights

are assumed to be 1.

directed: Boolean, whether to consider directed paths. Ignored for undirected graphs.

mode: How to determine the local neighborhood of each vertex in directed graphs. Ignored

in undirected graphs.

IGRAPH_ALL take both in- and out-neighbours; this is a reasonable default for

high-level interfaces.

IGRAPH_OUT take only out-neighbours

IGRAPH_IN take only in-neighbours

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for data structures

IGRAPH_EINVAL invalid weight vector

Time complexity: $O(|E|^2 \log |E|)$ for weighted graphs and $O(|E|^2)$ for unweighted ones. |E| denotes the number of edges.

See also:

```
igraph_local_efficiency()
```

Neighborhood of a vertex

igraph_neighborhood_size — Calculates the size of the neighborhood of a given vertex.

The neighborhood of a given order of a vertex includes all vertices which are closer to the vertex than the order. I.e., order 0 is always the vertex itself, order 1 is the vertex plus its immediate neighbors, order 2 is order 1 plus the immediate neighbors of the vertices in order 1, etc.

This function calculates the size of the neighborhood of the given order for the given vertices.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the result will be stored here. It will be resized as needed.

vids: The vertices for which the calculation is performed.

order: Integer giving the order of the neighborhood.

mode: Specifies how to use the direction of the edges if a directed graph is analyzed. For

IGRAPH_OUT only the outgoing edges are followed, so all vertices reachable from the source vertex in at most order steps are counted. For IGRAPH_IN all vertices from which the source vertex is reachable in at most order steps are counted. IGRAPH_ALL ignores the direction of the edges. This argument is ignored for undirected graphs.

mindist: The minimum distance to include a vertex in the counting. Vertices reachable with a

path shorter than this value are excluded. If this is one, then the starting vertex is not counted. If this is two, then its neighbors are not counted either, etc.

Returns:

Error code.

See also:

igraph_neighborhood() for calculating the actual neighborhood, igraph_neighborhood_graphs() for creating separate graphs from the neighborhoods.

Time complexity: O(n*d*o), where n is the number vertices for which the calculation is performed, d is the average degree, o is the order.

igraph_neighborhood — Calculate the neighborhood of vertices.

The neighborhood of a given order of a vertex includes all vertices which are closer to the vertex than the order. I.e., order 0 is always the vertex itself, order 1 is the vertex plus its immediate neighbors, order 2 is order 1 plus the immediate neighbors of the vertices in order 1, etc.

This function calculates the vertices within the neighborhood of the specified vertices.

Arguments:

graph: The input graph.

res: An initialized list of integer vectors. The result of the calculation will be stored here.

The list will be resized as needed.

vids: The vertices for which the calculation is performed.

order: Integer giving the order of the neighborhood.

mode: Specifies how to use the direction of the edges if a directed graph is analyzed. For

IGRAPH_OUT only the outgoing edges are followed, so all vertices reachable from the source vertex in at most order steps are included. For IGRAPH_IN all vertices from which the source vertex is reachable in at most order steps are included. IGRAPH_ALL ignores the direction of the edges. This argument is ignored for undi-

rected graphs.

mindist: The minimum distance to include a vertex in the counting. Vertices reachable with a

path shorter than this value are excluded. If this is one, then the starting vertex is not

counted. If this is two, then its neighbors are not counted either, etc.

Returns:

Error code.

See also:

igraph_neighborhood_size() to calculate the size of the neighborhood, igraph_neighborhood_graphs() for creating graphs from the neighborhoods.

Time complexity: O(n*d*o), n is the number of vertices for which the calculation is performed, d is the average degree, o is the order.

igraph_neighborhood_graphs — Create graphs from the neighborhood(s) of some vertex/vertices.

```
igraph_neimode_t mode,
igraph_integer_t mindist);
```

The neighborhood of a given order of a vertex includes all vertices which are closer to the vertex than the order. Ie. order 0 is always the vertex itself, order 1 is the vertex plus its immediate neighbors, order 2 is order 1 plus the immediate neighbors of the vertices in order 1, etc.

This function finds every vertex in the neighborhood of a given parameter vertex and creates the induced subgraph from these vertices.

The first version of this function was written by Vincent Matossian, thanks Vincent.

Arguments:

graph: The input graph.

res: Pointer to a list of graphs, the result will be stored here. Each item in the list is an

igraph_t object. The list will be resized as needed.

vids: The vertices for which the calculation is performed.

order: Integer giving the order of the neighborhood.

mode: Specifies how to use the direction of the edges if a directed graph is analyzed. For

IGRAPH_OUT only the outgoing edges are followed, so all vertices reachable from the source vertex in at most order steps are counted. For IGRAPH_IN all vertices from which the source vertex is reachable in at most order steps are counted. IGRAPH_ALL ignores the direction of the edges. This argument is ignored for undirected graphs.

mindist: The minimum distance to include a vertex in the counting. Vertices reachable with a

path shorter than this value are excluded. If this is one, then the starting vertex is not

counted. If this is two, then its neighbors are not counted either, etc.

Returns:

Error code.

See also:

igraph_neighborhood_size() for calculating the neighborhood sizes only igraph_neighborhood() for calculating the neighborhoods (but not creating graphs).

Time complexity: O(n*(|V|+|E|)), where n is the number vertices for which the calculation is performed, |V| and |E| are the number of vertices and edges in the original input graph.

Local scan statistics

The scan statistic is a summary of the locality statistics that is computed from the local neighborhood of each vertex. For details, see Priebe, C. E., Conroy, J. M., Marchette, D. J., Park, Y. (2005). Scan Statistics on Enron Graphs. Computational and Mathematical Organization Theory.

"Us" statistics

igraph_local_scan_0 — Local scan-statistics, k=0

```
igraph_error_t igraph_local_scan_0(const igraph_t *graph, igraph_vector_t *res,
```

```
const igraph_vector_t *weights,
igraph_neimode_t mode);
```

K=0 scan-statistics is arbitrarily defined as the vertex degree for unweighted, and the vertex strength for weighted graphs. See <code>igraph_degree()</code> and <code>igraph_strength()</code>.

Arguments:

graph: The input graph

res: An initialized vector, the results are stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN means in-

coming and IGRAPH_ALL means all edges.

Returns:

Error code.

igraph_local_scan_1_ecount — Local scan-statistics, k=1, edge count and sum of weights

Count the number of edges or the sum the edge weights in the 1-neighborhood of vertices.

Arguments:

graph: The input graph

res: An initialized vector, the results are stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN means in-

coming and IGRAPH_ALL means all edges.

Returns:

Error code.

igraph_local_scan_k_ecount — Sum the number of edges or the weights in k-neighborhood of every vertex.

Arguments:

graph: The input graph.

k: The size of the neighborhood, non-negative integer. The k=0 case is special, see

igraph_local_scan_0().

res: An initialized vector, the results are stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN means in-

coming and IGRAPH_ALL means all edges.

Returns:

Error code.

"Them" statistics

igraph_local_scan_0_them — Local THEM scan-statistics, k=0

K=0 scan-statistics is arbitrarily defined as the vertex degree for unweighted, and the vertex strength for weighted graphs. See <code>igraph_degree()</code> and <code>igraph_strength()</code>.

Arguments:

us: The input graph, to use to extract the neighborhoods.

them: The input graph to use for the actually counting.

res: An initialized vector, the results are stored here.

weights them: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN

means incoming and IGRAPH_ALL means all edges.

Returns:

Error code.

igraph_local_scan_1_ecount_them — Local THEM scan-statistics, k=1, edge count and sum of weights

```
igraph_neimode_t mode);
```

Count the number of edges or the sum the edge weights in the 1-neighborhood of vertices.

Arguments:

us: The input graph to extract the neighborhoods.

them: The input graph to perform the counting.

weights_them: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN

means incoming and IGRAPH_ALL means all edges.

Returns:

Error code.

See also:

igraph_local_scan_1_ecount() for the US statistics.

igraph_local_scan_k_ecount_them — Local THEM scan-statistics, edge count or sum of weights.

Count the number of edges or the sum the edge weights in the k-neighborhood of vertices.

Arguments:

us: The input graph to extract the neighborhoods.

them: The input graph to perform the counting.

k: The size of the neighborhood, non-negative integer. The k=0 case is special, see

igraph_local_scan_0_them().

weights_them: Weight vector for weighted graphs, null pointer for unweighted graphs.

mode: Type of the neighborhood, IGRAPH_OUT means outgoing, IGRAPH_IN

means incoming and IGRAPH_ALL means all edges.

Returns:

Error code.

See also:

igraph_local_scan_1_ecount() for the US statistics.

Pre-calculated subsets

igraph_local_scan_neighborhood_ecount — Local scan-statistics with pre-calculated neighborhoods

Count the number of edges, or sum the edge weights in neighborhoods given as a parameter.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_local_scan_subset_ecount() instead.

Arguments:

graph: The graph to perform the counting/summing in.

res: Initialized vector, the result is stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

neighborhoods: List of igraph_vector_int_t objects, the neighborhoods, one for each vertex in

the graph.

Returns:

Error code.

igraph_local_scan_subset_ecount — Local scan-statistics of subgraphs induced by subsets of vertices.

Count the number of edges, or sum the edge weights in induced subgraphs from vertices given as a parameter.

Arguments:

graph: The graph to perform the counting/summing in.

res: Initialized vector, the result is stored here.

weights: Weight vector for weighted graphs, null pointer for unweighted graphs.

subsets: List of igraph_vector_int_t objects, the vertex subsets.

Returns:

Error code.

Graph components

igraph_subcomponent — The vertices in the same component as a given vertex.

```
igraph_error_t igraph_subcomponent(
    const igraph_t *graph, igraph_vector_int_t *res, igraph_integer_t vertex,
    igraph_neimode_t mode
);
```

Arguments:

graph: The graph object.

res: The result, vector with the IDs of the vertices in the same component.

vertex: The id of the vertex of which the component is searched.

mode: Type of the component for directed graphs, possible values:

IGRAPH_OUT the set of vertices reachable from the vertex,

IGRAPH IN the set of vertices from which the *vertex* is reachable.

IGRAPH_ALL the graph is considered as an undirected graph. Note that this is not the

same as the union of the previous two.

Returns:

Error code:

```
IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID vertex is an invalid vertex ID
```

IGRAPH_EINVMODE invalid mode argument passed.

Time complexity: O(|V|+|E|), |V| and |E| are the number of vertices and edges in the graph.

See also:

igraph_induced_subgraph() if you want a graph object consisting only a given set of vertices and the edges between them.

igraph_connected_components — Calculates the (weakly or strongly) connected components in a graph.

```
igraph_error_t igraph_connected_components(
    const igraph_t *graph, igraph_vector_int_t *membership,
    igraph_vector_int_t *csize, igraph_integer_t *no, igraph_connectedness_t mode);
```

Arguments:

graph: The graph object to analyze.

membership: First half of the result will be stored here. For every vertex the id of its component

is given. The vector has to be preinitialized and will be resized. Alternatively this

argument can be NULL, in which case it is ignored.

csize: The second half of the result. For every component it gives its size, the order is de-

fined by the component ids. The vector has to be preinitialized and will be resized.

Alternatively this argument can be NULL, in which case it is ignored.

no: Pointer to an integer, if not NULL then the number of clusters will be stored here.

mode: For directed graph this specifies whether to calculate weakly or strongly connected

components. Possible values: IGRAPH_WEAK, IGRAPH_STRONG. This argument

is ignored for undirected graphs.

Returns:

Error code: IGRAPH_EINVAL: invalid mode argument.

Time complexity: O(|V|+|E|), |V| and |E| are the number of vertices and edges in the graph.

igraph_clusters — Calculates the (weakly or strongly) connected components in a graph (deprecated alias).

Warning

Deprecated since version 0.10. Please do not use this function in new code; use igraph_connected_components() instead.

igraph_is_connected — Decides whether the graph is (weakly or strongly) connected.

A graph is considered connected when any of its vertices is reachable from any other. A directed graph with this property is called *strongly* connected. A directed graph that would be connected when ignoring the directions of its edges is called *weakly* connected.

A graph with zero vertices (i.e. the null graph) is *not* connected by definition. This behaviour changed in igraph 0.9; earlier versions assumed that the null graph is connected. See the following issue on Github for the argument that led us to change the definition: https://github.com/igraph/igraph/issues/1539

The return value of this function is cached in the graph itself, separately for weak and strong connectivity. Calling the function multiple times with no modifications to the graph in between will return a cached value in O(1) time.

Arguments:

graph: The graph object to analyze.

res: Pointer to a logical variable, the result will be stored here.

mode: For a directed graph this specifies whether to calculate weak or strong connectedness. Pos-

sible values: IGRAPH_WEAK, IGRAPH_STRONG. This argument is ignored for undirect-

ed graphs.

Returns:

Error code: IGRAPH_EINVAL: invalid mode argument.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph.

igraph_decompose — Decomposes a graph into connected components.

Creates a separate graph for each component of a graph. Note that the vertex IDs in the new graphs will be different than in the original graph, except when there is only a single component in the original graph.

Arguments:

graph: The original graph.

components: This list of graphs will contain the individual components. It should be initialized

before calling this function and will be resized to hold the graphs.

mode: Either IGRAPH_WEAK or IGRAPH_STRONG for weakly and strongly connected

components respectively.

maxcompno: The maximum number of components to return. The first maxcompno compo-

nents will be returned (which hold at least *minelements* vertices, see the next parameter), the others will be ignored. Supply -1 here if you don't want to limit

the number of components.

minelements: The minimum number of vertices a component should contain in order to place

it in the *components* vector. Eg. supply 2 here to ignore isolated vertices.

Returns:

Error code, IGRAPH_ENOMEM if there is not enough memory to perform the operation.

Added in version 0.2.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

Example 13.12. File examples/simple/igraph_decompose.c

igraph_decompose_destroy — Frees the contents of a pointer vector holding graphs.

```
void igraph_decompose_destroy(igraph_vector_ptr_t *complist);
```

This function destroys and frees all igraph_t objects held in <code>complist</code>. However, it does not destroy <code>complist</code> itself. Use <code>igraph_vector_ptr_destroy()</code> to destroy <code>complist</code>.

Arguments:

complist: The list of graphs to destroy.

Time complexity: O(n), n is the number of items.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code.

igraph_biconnected_components — Calculates biconnected components.

A graph is biconnected if the removal of any single vertex (and its incident edges) does not disconnect it.

A biconnected component of a graph is a maximal biconnected subgraph of it. The biconnected components of a graph can be given by the partition of its edges: every edge is a member of exactly one biconnected component. Note that this is not true for vertices: the same vertex can be part of many biconnected components.

Note that some authors do not consider the graph consisting of two connected vertices as biconnected, however, igraph does.

Somewhat arbitrarily, igraph does not consider components containing a single vertex only as being biconnected. Isolated vertices will not be part of any of the biconnected components.

Arguments:

graph: The input graph.

no: If not a NULL pointer, the number of biconnected components will

be stored here.

tree_edges: If not a NULL pointer, then the found components are stored here,

in a list of vectors. Every vector in the list is a biconnected component, represented by its edges. More precisely, a spanning tree of the

biconnected component is returned.

component_edges: If not a NULL pointer, then the edges of the biconnected components

are stored here, in the same form as for tree_edges.

components: If not a NULL pointer, then the vertices of the biconnected compo-

nents are stored here, in the same format as for the previous two ar-

guments.

articulation_points: If not a NULL pointer, then the articulation points of the graph are

stored in this vector. A vertex is an articulation point if its removal increases the number of (weakly) connected components in the graph.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges, but only if you do not calculate components and component_edges. If you calculate components, then it is quadratic in the number of vertices. If you calculate component_edges as well, then it is cubic in the number of vertices.

See also:

igraph_articulation_points(), igraph_clusters().

Example 13.13. File examples/simple/igraph_biconnected_components.c

igraph_articulation_points — Finds the articulation points in a graph.

igraph_error_t igraph_articulation_points(const igraph_t *graph, igraph_vector_

A vertex is an articulation point if its removal increases the number of connected components in the graph.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the articulation points will be stored here.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

See also:

```
igraph_biconnected_components(), igraph_clusters(), igraph_bridges()
```

igraph_bridges — Finds all bridges in a graph.

```
igraph_error_t igraph_bridges(const igraph_t *graph, igraph_vector_int_t *bridg
```

An edge is a bridge if its removal increases the number of (weakly) connected components in the graph.

Arguments:

graph: The input graph. It will be treated as undirected.

res: Pointer to an initialized vector, the bridges will be stored here as edge indices.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

See also:

Degree sequences

igraph_is_graphical — Is there a graph with the given degree sequence?

Determines whether a sequence of integers can be the degree sequence of some graph. The classical concept of graphicality assumes simple graphs. This function can perform the check also when either self-loops, multi-edge, or both are allowed in the graph.

For simple undirected graphs, the Erd#s-Gallai conditions are checked using the linear-time algorithm of Cloteaux. If both self-loops and multi-edges are allowed, it is sufficient to chek that that sum of degrees is even. If only multi-edges are allowed, but not self-loops, there is an additional condition that the sum of degrees be no smaller than twice the maximum degree. If at most one self-loop is allowed per vertex, but no multi-edges, a modified version of the Erd#s-Gallai conditions are used (see Cairns & Mendan).

For simple directed graphs, the Fulkerson-Chen-Anstee theorem is used with the relaxation by Berger. If both self-loops and multi-edges are allowed, then it is sufficient to check that the sum of in- and

out-degrees is the same. If only multi-edges are allowed, but not self loops, there is an additional condition that the sum of out-degrees (or equivalently, in-degrees) is no smaller than the maximum total degree. If single self-loops are allowed, but not multi-edges, the problem is equivalent to realizability as a simple bipartite graph, thus the Gale-Ryser theorem can be used; see <code>igraph_is_bigraph-ical()</code> for more information.

References:

- P. Erd#s and T. Gallai, Gráfok el#írt fokú pontokkal, Matematikai Lapok 11, pp. 264–274 (1960). https://users.renyi.hu/~p_erdos/1961-05.pdf
- Z Király, Recognizing graphic degree sequences and generating all realizations. TR-2011-11, Egerváry Research Group, H-1117, Budapest, Hungary. ISSN 1587-4451 (2012). http://bolyai.cs.elte.hu/egres/tr/egres-11-11.pdf
- B. Cloteaux, Is This for Real? Fast Graphicality Testing, Comput. Sci. Eng. 17, 91 (2015). https://dx.doi.org/10.1109/MCSE.2015.125
- A. Berger, A note on the characterization of digraphic sequences, Discrete Math. 314, 38 (2014). https://dx.doi.org/10.1016/j.disc.2013.09.010
- G. Cairns and S. Mendan, Degree Sequence for Graphs with Loops (2013). https://arxiv.org/abs/1303.2145v1

Arguments:

out_degrees: A vector of integers specifying the degree sequence for undirected

graphs or the out-degree sequence for directed graphs.

in_degrees: A vector of integers specifying the in-degree sequence for directed

graphs. For undirected graphs, it must be NULL.

allowed_edge_types: The types of edges to allow in the graph:

IGRAPH_SIMPLE_SW simple graphs (i.e. no self-loops

or multi-edges allowed).

IGRAPH_LOOPS_SW single self-loops are allowed,

but not multi-edges.

IGRAPH_MULTI_SW multi-edges are allowed, but

not self-loops.

IGRAPH_LOOPS_SW | both self-loops and multi-edges

IGRAPH_MULTI_SW are allowed.

res: Pointer to a Boolean. The result will be stored here.

Returns:

Error code.

See also:

igraph_is_bigraphical() to check if a bi-degree-sequence can be realized as a bipartite
graph; igraph_realize_degree_sequence() to construct a graph with a given degree
sequence.

Time complexity: $O(n^2)$ for simple directed graphs, $O(n \log n)$ for graphs with self-loops, and O(n) for all other cases, where n is the length of the degree sequence(s).

igraph_is_bigraphical — Is there a bipartite graph with the given bi-degree-sequence?

Determines whether two sequences of integers can be the degree sequences of a bipartite graph. Such a pair of degree sequence is called *bigraphical*.

When multi-edges are allowed, it is sufficient to check that the sum of degrees is the same in the two partitions. For simple graphs, the Gale-Ryser theorem is used with Berger's relaxation.

References:

H. J. Ryser, Combinatorial Properties of Matrices of Zeros and Ones, Can. J. Math. 9, 371 (1957). https://dx.doi.org/10.4153/cjm-1957-044-3

D. Gale, A theorem on flows in networks, Pacific J. Math. 7, 1073 (1957). https://dx.doi.org/10.2140/pjm.1957.7.1073

A. Berger, A note on the characterization of digraphic sequences, Discrete Math. 314, 38 (2014). https://dx.doi.org/10.1016/j.disc.2013.09.010

Arguments:

degrees 1: A vector of integers specifying the degrees in the first partition

degrees 2: A vector of integers specifying the degrees in the second partition

allowed_edge_types: The types of edges to allow in the graph:

IGRAPH_SIMPLE_SW simple graphs (i.e. no multi-edges allowed).

IGRAPH_MULTI_SW multi-edges are allowed.

res: Pointer to a Boolean. The result will be stored here.

Returns:

Error code.

See also:

```
igraph_is_graphical()
```

Time complexity: $O(n \log n)$ for simple graphs, O(n) for multigraphs, where n is the length of the larger degree sequence.

Centrality measures

igraph_closeness — Closeness centrality calculations for some vertices.

The closeness centrality of a vertex measures how easily other vertices can be reached from it (or the other way: how easily it can be reached from the other vertices). It is defined as the inverse of the mean distance to (or from) all other vertices.

Closeness centrality is meaningful only for connected graphs. If the graph is not connected, igraph computes the inverse of the mean distance to (or from) all *reachable* vertices. In undirected graphs, this is equivalent to computing the closeness separately in each connected component. The optional <code>all_reachable</code> output parameter is provided to help detect when the graph is disconnected.

While there is no universally adopted definition of closeness centrality for disconnected graphs, there have been some attempts for generalizing the concept to the disconnected case. One type of approach considers the mean distance only to reachable vertices, then re-scales the obtained certrality score by a factor that depends on the number of reachable vertices (i.e. the size of the component in the undirected case). To facilitate computing these generalizations of closeness centrality, the number of reachable vertices (not including the starting vertex) is returned in reachable_count.

In disconnected graphs, consider using the harmonic centrality, computable using igraph_harmonic_centrality().

For isolated vertices, i.e. those having no associated paths, NaN is returned.

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the closeness centrality

scores for the given vertices.

reachable_count: If not NULL, this vector will contain the number of vertices reachable from

each vertex for which the closeness is calculated (not including that vertex).

all_reachable: Pointer to a Boolean. If not NULL, it indicates if all vertices of the graph

were reachable from each vertex in *vids*. If false, the graph is non-connected. If true, and the graph is undirected, or if the graph is directed and

vids contains all vertices, then the graph is connected.

vids: The vertices for which the closeness centrality will be computed.

mode: The type of shortest paths to be used for the calculation in directed graphs.

Possible values:

IGRAPH_OUT the lengths of the outgoing paths are calculated.

IGRAPH_IN the lengths of the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for

the computation.

weights: An optional vector containing edge weights for weighted closeness. No

edge weight may be NaN. Supply a null pointer here for traditional, un-

weighted closeness.

normalized: If true, the inverse of the mean distance to reachable vetices is returned. If

false, the inverse of the sum of distances is returned.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(n|E|), n is the number of vertices for which the calculation is done and |E| is the number of edges in the graph.

See also:

Other centrality types: igraph_degree(), igraph_betweenness(), igraph_harmon-ic_centrality(). See igraph_closeness_cutoff() for the range-limited closeness centrality.

igraph_harmonic_centrality — Harmonic centrality for some vertices.

The harmonic centrality of a vertex is the mean inverse distance to all other vertices. The inverse distance to an unreachable vertex is considered to be zero.

References:

M. Marchiori and V. Latora, Harmony in the small-world, Physica A 285, pp. 539-546 (2000). https://doi.org/10.1016/S0378-4371%2800%2900311-3

Y. Rochat, Closeness Centrality Extended to Unconnected Graphs: the Harmonic Centrality Index, ASNA 2009. https://infoscience.epfl.ch/record/200525

S. Vigna and P. Boldi, Axioms for Centrality, Internet Mathematics 10, (2014). https://doi.org/10.1080/15427951.2013.865686

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the harmonic centrality scores

for the given vertices.

vids: The vertices for which the harmonic centrality will be computed.

mode: The type of shortest paths to be used for the calculation in directed graphs. Possible

values:

IGRAPH_OUT the lengths of the outgoing paths are calculated.

IGRAPH_IN the lengths of the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the com-

putation.

weights: An optional vector containing edge weights for weighted harmonic centrality. No

edge weight may be NaN. If NULL, all weights are considered to be one.

normalized: Boolean, whether to normalize the result. If true, the result is the mean inverse path

length to other vertices, i.e. it is normalized by the number of vertices minus one.

If false, the result is the sum of inverse path lengths to other vertices.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(n|E|), where n is the number of vertices for which the calculation is done and |E| is the number of edges in the graph.

See also:

Other centrality types: igraph_closeness(), igraph_degree(), igraph_betweenness().

igraph_betweenness — Betweenness centrality of some vertices.

The betweenness centrality of a vertex is the number of geodesics going through it. If there are more than one geodesic between two vertices, the value of these geodesics are weighted by one over the number of geodesics.

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the betweenness scores for the spec-

ified vertices.

vids: The vertices of which the betweenness centrality scores will be calculated.

directed: Logical, if true directed paths will be considered for directed graphs. It is ignored for

undirected graphs.

weights: An optional vector containing edge weights for calculating weighted betweenness. No

edge weight may be NaN. Supply a null pointer here for unweighted betweenness.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex ID passed in *vids*.

Time complexity: O(|V||E|), |V| and |E| are the number of vertices and edges in the graph. Note that the time complexity is independent of the number of vertices for which the score is calculated.

See also:

Other centrality types: igraph_degree(), igraph_closeness(). See igraph_edge_betweenness() for calculating the betweenness score of the edges in a graph. See igraph_betweenness_cutoff() to calculate the range-limited betweenness of the vertices in a graph.

igraph_edge_betweenness — Betweenness centrality of the edges.

The betweenness centrality of an edge is the number of geodesics going through it. If there are more than one geodesics between two vertices, the value of these geodesics are weighted by one over the number of geodesics.

Arguments:

graph: The graph object.

result: The result of the computation, vector containing the betweenness scores for the edges.

directed: Logical, if true directed paths will be considered for directed graphs. It is ignored for

undirected graphs.

weights: An optional weight vector for weighted edge betweenness. No edge weight may be

NaN. Supply a null pointer here for the unweighted version.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: O(|V||E|), |V| and |E| are the number of vertices and edges in the graph.

See also:

Other centrality types: igraph_degree(), igraph_closeness(). See igraph_edge_betweenness() for calculating the betweenness score of the edges in a graph. See igraph_edge_betweenness_cutoff() to compute the range-limited betweenness score of the edges in a graph.

igraph_pagerank_algo_t — PageRank algorithm implementation

```
typedef enum {
    IGRAPH_PAGERANK_ALGO_ARPACK = 1,
    IGRAPH_PAGERANK_ALGO_PRPACK = 2
} igraph_pagerank_algo_t;
```

Algorithms to calculate PageRank.

Values:

IGRAPH_PAGERANK_ALUse the ARPACK library, this was the PageRank implementation in igraph from version 0.5, until version 0.7.

IGRAPH_PAGERANK_ALUse the PRPACK library. Currently this implementation is rec-

GO_PRPACK: ommended.

igraph_pagerank — Calculates the Google PageRank for the specified vertices.

The PageRank centrality of a vertex is the fraction of time a random walker traversing the graph would spend on that vertex. The walker follows the out-edges with probabilities proportional to their weights. Additionally, in each step, it restarts the walk from a random vertex with probability 1 - damping. If the random walker gets stuck in a sink vertex, it will also restart from a random vertex.

The PageRank centrality is mainly useful for directed graphs. In undirected graphs it converges to trivial values proportional to degrees as the damping factor approaches 1.

Starting from version 0.9, igraph has two PageRank implementations, and the user can choose between them. The first implementation is ${\tt IGRAPH_PAGERANK_ALGO_ARPACK}$, based on the ARPACK library. This was the default before igraph version 0.7. The second and recommended implementation is ${\tt IGRAPH_PAGERANK_ALGO_PRPACK}$. This is using the PRPACK package, see https://github.com/dgleich/prpack.

Note that the PageRank of a given vertex depends on the PageRank of all other vertices, so even if you want to calculate the PageRank for only some of the vertices, all of them must be calculated. Requesting the PageRank for only some of the vertices does not result in any performance increase at all.

References:

Sergey Brin and Larry Page: The Anatomy of a Large-Scale Hypertextual Web Search Engine. Proceedings of the 7th World-Wide Web Conference, Brisbane, Australia, April 1998.

Arguments:

graph: The graph object.

algo: The PageRank implementation to use. Possible values: IGRAPH_PAGERANK_AL-

GO ARPACK, IGRAPH PAGERANK ALGO PRPACK.

vector: Pointer to an initialized vector, the result is stored here. It is resized as needed.

value: Pointer to a real variable, the eigenvalue corresponding to the PageRank vector is

stored here. It should be always exactly one.

vids: The vertex IDs for which the PageRank is returned.

directed: Boolean, whether to consider the directedness of the edges. This is ignored for undi-

rected graphs.

damping: The damping factor ("d" in the original paper). Must be a probability in the range [0,

1]. A commonly used value is 0.85.

weights: Optional edge weights. May be a NULL pointer, meaning unweighted edges, or a vec-

tor of non-negative values of the same length as the number of edges.

options: Options for the ARPACK method. See igraph_arpack_options_t for details.

Supply NULL here to use the defaults. Note that the function overwrites the n (number of vertices), nev (1), ncv (3) and which (LM) parameters and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex ID in *vids*.

Time complexity: depends on the input graph, usually it is O(|E|), the number of edges.

See also:

igraph_personalized_pagerank() and igraph_personalized_pagerank_vs() for the personalized PageRank measure. See igraph_arpack_rssolve() and
igraph_arpack_rnsolve() for the underlying machinery used by IGRAPH_PAGERANK_ALGO_ARPACK.

Example 13.14. File examples/simple/igraph_pagerank.c

igraph_personalized_pagerank — Calculates the personalized Google PageRank for the specified vertices.

The personalized PageRank is similar to the original PageRank measure, but when the random walk is restarted, a new starting vertex is chosen non-uniformly, according to the distribution specified in reset (instead of the uniform distribution in the original PageRank measure). The reset distribution is used both when restarting randomly with probability 1 - damping, and when the walker is forced to restart due to being stuck in a sink vertex (a vertex with no outgoing edges).

Note that the personalized PageRank of a given vertex depends on the personalized PageRank of all other vertices, so even if you want to calculate the personalized PageRank for only some of the vertices,

all of them must be calculated. Requesting the personalized PageRank for only some of the vertices does not result in any performance increase at all.

Arguments:

graph: The graph object.

algo: The PageRank implementation to use. Possible values: IGRAPH_PAGERANK_AL-

GO_ARPACK, IGRAPH_PAGERANK_ALGO_PRPACK.

vector: Pointer to an initialized vector, the result is stored here. It is resized as needed.

value: Pointer to a real variable, the eigenvalue corresponding to the PageRank vector is

stored here. It should be always exactly one.

vids: The vertex IDs for which the PageRank is returned.

directed: Boolean, whether to consider the directedness of the edges. This is ignored for undi-

rected graphs.

damping: The damping factor ("d" in the original paper). Must be a probability in the range [0,

1]. A commonly used value is 0.85.

reset: The probability distribution over the vertices used when resetting the random walk. It

is either a NULL pointer (denoting a uniform choice that results in the original PageR-

ank measure) or a vector of the same length as the number of vertices.

weights: Optional edge weights. May be a NULL pointer, meaning unweighted edges, or a vec-

tor of non-negative values of the same length as the number of edges.

options: Options for the ARPACK method. See igraph_arpack_options_t for details.

Supply NULL here to use the defaults. Note that the function overwrites the n (number of vertices), nev (1), ncv (3) and which (LM) parameters and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex ID in *vids* or an invalid reset vector in *reset*.

Time complexity: depends on the input graph, usually it is O(|E|), the number of edges.

See also:

igraph_pagerank() for the non-personalized implementation, igraph_personalized_pagerank_vs() for a personalized implementation with resetting to specific vertices.

igraph_personalized_pagerank_vs — Calculates the personalized Google PageRank for the specified vertices.

```
igraph_bool_t directed, igraph_real_t dampi
igraph_vs_t reset_vids,
const igraph_vector_t *weights,
igraph_arpack_options_t *options);
```

The personalized PageRank is similar to the original PageRank measure, but when the random walk is restarted, a new starting vertex is chosen according to a specified distribution. This distribution is used both when restarting randomly with probability 1 - damping, and when the walker is forced to restart due to being stuck in a sink vertex (a vertex with no outgoing edges).

This simplified interface takes a vertex sequence and resets the random walk to one of the vertices in the specified vertex sequence, chosen uniformly. A typical application of personalized PageRank is when the random walk is reset to the same vertex every time - this can easily be achieved using igraph_vss_1() which generates a vertex sequence containing only a single vertex.

Note that the personalized PageRank of a given vertex depends on the personalized PageRank of all other vertices, so even if you want to calculate the personalized PageRank for only some of the vertices, all of them must be calculated. Requesting the personalized PageRank for only some of the vertices does not result in any performance increase at all.

Arguments:

graph: The graph object.

algo: The PageRank implementation to use. Possible values: IGRAPH_PAGER-

ANK_ALGO_ARPACK, IGRAPH_PAGERANK_ALGO_PRPACK.

vector: Pointer to an initialized vector, the result is stored here. It is resized as needed.

value: Pointer to a real variable, the eigenvalue corresponding to the PageRank vector is

stored here. It should be always exactly one.

vids: The vertex IDs for which the PageRank is returned.

directed: Boolean, whether to consider the directedness of the edges. This is ignored for

undirected graphs.

damping: The damping factor ("d" in the original paper). Must be a probability in the range

[0, 1]. A commonly used value is 0.85.

reset_vids: IDs of the vertices used when resetting the random walk.

weights: Optional edge weights, it is either a null pointer, then the edges are not weighted,

or a vector of the same length as the number of edges.

options: Options for the ARPACK method. See igraph_arpack_options_t for de-

tails. Supply NULL here to use the defaults. Note that the function overwrites the n (number of vertices), nev (1), ncv (3) and which (LM) parameters and it always starts the calculation from a non-random vector calculated based on the degree of

the vertices.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex ID in *vids* or an empty reset vertex sequence in *vids_reset*.

Time complexity: depends on the input graph, usually it is O(|E|), the number of edges.

See also:

igraph_pagerank() for the non-personalized implementation.

igraph_constraint — Burt's constraint scores.

This function calculates Burt's constraint scores for the given vertices, also known as structural holes.

Burt's constraint is higher if ego has less, or mutually stronger related (i.e. more redundant) contacts. Burt's measure of constraint, C[i], of vertex i's ego network V[i], is defined for directed and valued graphs,

```
C[i] = sum( sum( (p[i,q] p[q,j])^2, q in V[i], q!=i,j), j in V[], j!=i)
```

for a graph of order (i.e. number of vertices) N, where proportional tie strengths are defined as

```
p[i,j]=(a[i,j]+a[j,i]) / sum(a[i,k]+a[k,i], k in V[i], k != i),
```

a[i,j] are elements of A and the latter being the graph adjacency matrix. For isolated vertices, constraint is undefined.

Burt, R.S. (2004). Structural holes and good ideas. American Journal of Sociology 110, 349-399.

The first R version of this function was contributed by Jeroen Bruggeman.

Arguments:

graph: A graph object.

res: Pointer to an initialized vector, the result will be stored here. The vector will be resized

to have the appropriate size for holding the result.

vids: Vertex selector containing the vertices for which the constraint should be calculated.

weights: Vector giving the weights of the edges. If it is NULL then each edge is supposed to have

the same weight.

Returns:

Error code.

Time complexity: $O(|V|+E|+n*d^2)$, n is the number of vertices for which the constraint is calculated and d is the average degree, |V| is the number of vertices, |E| the number of edges in the graph. If the weights argument is NULL then the time complexity is $O(|V|+n*d^2)$.

igraph_maxdegree — The maximum degree in a graph (or set of vertices).

```
igraph_bool_t loops);
```

The largest in-, out- or total degree of the specified vertices is calculated. If the graph has no vertices, or *vids* is empty, 0 is returned, as this is the smallest possible value for degrees.

Arguments:

graph: The input graph.

res: Pointer to an integer (igraph_integer_t), the result will be stored here.

vids: Vector giving the vertex IDs for which the maximum degree will be calculated.

mode: Defines the type of the degree. IGRAPH_OUT, out-degree, IGRAPH_IN, in-degree,

IGRAPH_ALL, total degree (sum of the in- and out-degree). This parameter is ignored for

undirected graphs.

loops: Boolean, gives whether the self-loops should be counted.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex ID. IGRAPH_EINVMODE: invalid mode argument.

Time complexity: O(v) if loops is true, and O(v*d) otherwise. v is the number of vertices for which the degree will be calculated, and d is their (average) degree.

igraph_strength — Strength of the vertices, also called weighted vertex degree.

In a weighted network the strength of a vertex is the sum of the weights of all incident edges. In a non-weighted network this is exactly the vertex degree.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the result is stored here. It will be resized as needed.

vids: The vertices for which the calculation is performed.

mode: Gives whether to count only outgoing (IGRAPH_OUT), incoming (IGRAPH_IN) edges

or both (IGRAPH_ALL).

loops: A logical scalar, whether to count loop edges as well.

weights: A vector giving the edge weights. If this is a NULL pointer, then igraph_degree()

is called to perform the calculation.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number vertices and edges.

See also:

igraph_degree() for the traditional, non-weighted version.

igraph_eigenvector_centrality — Eigenvector centrality of the vertices.

Eigenvector centrality is a measure of the importance of a node in a network. It assigns relative scores to all nodes in the network based on the principle that connections from high-scoring nodes contribute more to the score of the node in question than equal connections from low-scoring nodes. Specifically, the eigenvector centrality of each vertex is proportional to the sum of eigenvector centralities of its neighbors. In practice, the centralities are determined by calculating the eigenvector corresponding to the largest positive eigenvalue of the adjacency matrix. In the undirected case, this function considers the diagonal entries of the adjacency matrix to be *twice* the number of self-loops on the corresponding vertex.

In the weighted case, the eigenvector centrality of a vertex is proportional to the weighted sum of centralities of its neighbours, i.e. c_i = sum_j w_ij c_j, where w_ij is the weight of the edge connecting vertices i and j. The weights of parallel edges are added up.

The centrality scores returned by igraph can be normalized (using the scale parameter) such that the largest eigenvector centrality score is 1 (with one exception, see below).

In the directed case, the left eigenvector of the adjacency matrix is calculated. In other words, the centrality of a vertex is proportional to the sum of centralities of vertices pointing to it.

Eigenvector centrality is meaningful only for (strongly) connected graphs. Undirected graphs that are not connected should be decomposed into connected components, and the eigenvector centrality calculated for each separately. This function does not verify that the graph is connected. If it is not, in the undirected case the scores of all but one component will be zeros.

Also note that the adjacency matrix of a directed acyclic graph or the adjacency matrix of an empty graph does not possess positive eigenvalues, therefore the eigenvector centrality is not defined for these graphs. igraph will return an eigenvalue of zero in such cases. The eigenvector centralities will all be equal for an empty graph and will all be zeros for a directed acyclic graph. Such pathological cases can be detected by asking igraph to calculate the eigenvalue as well (using the value parameter, see below) and checking whether the eigenvalue is very close to zero.

Arguments:

graph: The input graph. It may be directed.

vector: Pointer to an initialized vector, it will be resized as needed. The result of the compu-

tation is stored here. It can be a null pointer, then it is ignored.

value: If not a null pointer, then the eigenvalue corresponding to the found eigenvector is

stored here.

directed: Boolean scalar, whether to consider edge directions in a directed graph. It is ignored

for undirected graphs.

scale: If not zero then the result will be scaled such that the absolute value of the maximum

centrality is one.

weights: A null pointer (indicating no edge weights), or a vector giving the weights of the edges.

Weights should be positive to guarantee a meaningful result. The algorithm might produce complex numbers when some weights are negative and the graph is directed.

In this case only the real part is reported.

options: Options to ARPACK. See igraph_arpack_options_t for details. Supply

NULL here to use the defaults. Note that the function overwrites the n (number of vertices) parameter and it always starts the calculation from a non-random vector cal-

culated based on the degree of the vertices.

Returns:

Error code.

Time complexity: depends on the input graph, usually it is O(|V|+|E|).

See also:

igraph_pagerank and igraph_personalized_pagerank for modifications of eigenvector centrality.

Example 13.15. File examples/simple/eigenvector_centrality.c

igraph_hub_score — Kleinberg's hub scores.

The hub scores of the vertices are defined as the principal eigenvector of A*A^T, where A is the adjacency matrix of the graph, A^T is its transposed.

See the following reference on the meaning of this score: J. Kleinberg. Authoritative sources in a hyperlinked environment. *Proc. 9th ACM-SIAM Symposium on Discrete Algorithms*, 1998. Extended version in *Journal of the ACM* 46(1999). Also appears as IBM Research Report RJ 10076, May 1997.

Arguments:

graph: The input graph. Can be directed and undirected.

vector: Pointer to an initialized vector, the result is stored here. If a null pointer then it is ignored.

value: If not a null pointer then the eigenvalue corresponding to the calculated eigenvector is

stored here.

scale: If not zero then the result will be scaled such that the absolute value of the maximum

centrality is one.

weights: A null pointer (=no edge weights), or a vector giving the weights of the edges.

options: Options to ARPACK. See igraph_arpack_options_t for details. Note that the

function overwrites the n (number of vertices) parameter and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code.

Time complexity: depends on the input graph, usually it is O(|V|), the number of vertices.

See also:

igraph_authority_score() for the companion measure, igraph_pagerank(), igraph_personalized_pagerank(), igraph_eigenvector_centrality() for similar measures.

igraph_authority_score — Kleinerg's authority scores.

The authority scores of the vertices are defined as the principal eigenvector of A^T*A, where A is the adjacency matrix of the graph, A^T is its transposed.

See the following reference on the meaning of this score: J. Kleinberg. Authoritative sources in a hyperlinked environment. *Proc. 9th ACM-SIAM Symposium on Discrete Algorithms*, 1998. Extended version in *Journal of the ACM* 46(1999). Also appears as IBM Research Report RJ 10076, May 1997.

Arguments:

graph: The input graph. Can be directed and undirected.

vector: Pointer to an initialized vector, the result is stored here. If a null pointer then it is ignored.

value: If not a null pointer then the eigenvalue corresponding to the calculated eigenvector is

stored here.

scale: If not zero then the result will be scaled such that the absolute value of the maximum

centrality is one.

weights: A null pointer (=no edge weights), or a vector giving the weights of the edges.

options: Options to ARPACK. See igraph arpack options t for details. Note that the

function overwrites the n (number of vertices) parameter and it always starts the calculation from a non-random vector calculated based on the degree of the vertices.

Returns:

Error code.

Time complexity: depends on the input graph, usually it is O(|V|), the number of vertices.

See also:

igraph_hub_score() for the companion measure, igraph_pagerank(), igraph_personalized_pagerank(), igraph_eigenvector_centrality() for similar measures.

igraph_convergence_degree — Calculates the convergence degree of each edge in a graph.

Let us define the input set of an edge (i, j) as the set of vertices where the shortest paths passing through (i, j) originate, and similarly, let us defined the output set of an edge (i, j) as the set of vertices where the shortest paths passing through (i, j) terminate. The convergence degree of an edge is defined as the normalized value of the difference between the size of the input set and the output set, i.e. the difference of them divided by the sum of them. Convergence degrees are in the range (-1, 1); a positive value indicates that the edge is *convergent* since the shortest paths passing through it originate from a larger set and terminate in a smaller set, while a negative value indicates that the edge is *divergent* since the paths originate from a small set and terminate in a larger set.

Note that the convergence degree as defined above does not make sense in undirected graphs as there is no distinction between the input and output set. Therefore, for undirected graphs, the input and output sets of an edge are determined by orienting the edge arbitrarily while keeping the remaining edges undirected, and then taking the absolute value of the convergence degree.

Arguments:

graph: The input graph, it can be either directed or undirected.

result: Pointer to an initialized vector; the convergence degrees of each edge will be stored here.

May be NULL if we are not interested in the exact convergence degrees.

ins: Pointer to an initialized vector; the size of the input set of each edge will be stored here.

May be NULL if we are not interested in the sizes of the input sets.

outs: Pointer to an initialized vector; the size of the output set of each edge will be stored here.

May be NULL if we are not interested in the sizes of the output sets.

Returns:

Error code.

Time complexity: O(|V||E|), the number of vertices times the number of edges.

Range-limited centrality measures

igraph_closeness_cutoff — Range limited closeness centrality.

```
const igraph_vector_t *weights,
igraph_bool_t normalized,
igraph real t cutoff);
```

This function computes a range-limited version of closeness centrality by considering only those shortest paths whose length is no greater then the given cutoff value.

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the range-limited close-

ness centrality scores for the given vertices.

reachable_count: If not NULL, this vector will contain the number of vertices reachable within

the cutoff distance from each vertex for which the range-limited closeness

is calculated (not including that vertex).

all_reachable: Pointer to a Boolean. If not NULL, it indicates if all vertices of the graph

were reachable from each vertex in vids within the given cutoff distance.

vids: The vertices for which the range limited closeness centrality will be com-

puted.

mode: The type of shortest paths to be used for the calculation in directed graphs.

Possible values:

IGRAPH_OUT the lengths of the outgoing paths are calculated.

IGRAPH_IN the lengths of the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for

the computation.

weights: An optional vector containing edge weights for weighted closeness. No

edge weight may be NaN. Supply a null pointer here for traditional, un-

weighted closeness.

normalized: If true, the inverse of the mean distance to vertices reachable within the

cutoff is returned. If false, the inverse of the sum of distances is returned.

cutoff: The maximal length of paths that will be considered. If negative, the exact

closeness will be calculated (no upper limit on path lengths).

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(n|E|), n is the number of vertices for which the calculation is done and |E| is the number of edges in the graph.

See also:

igraph_closeness() to calculate the exact closeness centrality.

igraph_harmonic_centrality_cutoff — Range limited harmonic centrality.

```
igraph_error_t igraph_harmonic_centrality_cutoff(const igraph_t *graph, igraph_reimode_t const igraph_vs_t vids, igraph_neimode_t const igraph_vector_t *weights, igraph_bool_t normalized, igraph_real_t cutoff);
```

This function computes the range limited version of harmonic centrality: only those shortest paths are considered whose length is not above the given cutoff. The inverse distance to vertices not reachable within the cutoff is considered to be zero.

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the range limited harmonic cen-

trality scores for the given vertices.

vids: The vertices for which the harmonic centrality will be computed.

mode: The type of shortest paths to be used for the calculation in directed graphs. Possible

values:

IGRAPH_OUT the lengths of the outgoing paths are calculated.

IGRAPH_IN the lengths of the incoming paths are calculated.

IGRAPH_ALL the directed graph is considered as an undirected one for the com-

putation.

weights: An optional vector containing edge weights for weighted harmonic centrality. No

edge weight may be NaN. If NULL, all weights are considered to be one.

normalized: Boolean, whether to normalize the result. If true, the result is the mean inverse path

length to other vertices. i.e. it is normalized by the number of vertices minus one.

If false, the result is the sum of inverse path lengths to other vertices.

cutoff: The maximal length of paths that will be considered. The inverse distance to ver-

tices that are not reachable within the cutoff path length is considered to be zero. Supply a negative value to compute the exact harmonic centrality, without any up-

per limit on the length of paths.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(n|E|), where n is the number of vertices for which the calculation is done and |E| is the number of edges in the graph.

See also:

Other centrality types: igraph_closeness(), igraph_betweenness().

igraph_betweenness_cutoff — Range-limited betweenness centrality.

This function computes a range-limited version of betweenness centrality by considering only those shortest paths whose length is no greater then the given cutoff value.

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the range-limited betweenness

scores for the specified vertices.

vids: The vertices for which the range-limited betweenness centrality scores will be com-

puted.

directed: Logical, if true directed paths will be considered for directed graphs. It is ignored for

undirected graphs.

weights: An optional vector containing edge weights for calculating weighted betweenness. No

edge weight may be NaN. Supply a null pointer here for unweighted betweenness.

cutoff: The maximal length of paths that will be considered. If negative, the exact between-

ness will be calculated, and there will be no upper limit on path lengths.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex ID passed in *vids*.

Time complexity: O(|V||E|), |V| and |E| are the number of vertices and edges in the graph. Note that the time complexity is independent of the number of vertices for which the score is calculated.

See also:

 $igraph_betweenness() \ to \ calculate \ the \ exact \ betweenness \ and \ igraph_edge_betweenness.$

igraph_edge_betweenness_cutoff — Range-limited betweenness centrality of the edges.

```
const igraph_vector_t *weights, igraph_real_
```

This function computes a range-limited version of edge betweenness centrality by considering only those shortest paths whose length is no greater then the given cutoff value.

Arguments:

graph: The graph object.

result: The result of the computation, vector containing the betweenness scores for the edges.

directed: Logical, if true directed paths will be considered for directed graphs. It is ignored for

undirected graphs.

weights: An optional weight vector for weighted betweenness. No edge weight may be NaN.

Supply a null pointer here for unweighted betweenness.

cutoff: The maximal length of paths that will be considered. If negative, the exact between-

ness will be calculated (no upper limit on path lengths).

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: O(|V||E|), |V| and |E| are the number of vertices and edges in the graph.

See also:

igraph_edge_betweenness() to compute the exact edge betweenness and igraph_betweenness_cutoff() to compute the range-limited vertex betweenness.

Subset-limited Centrality Measures

igraph_betweenness_subset — Betweenness centrality for a subset of source and target vertices.

This function computes the subset-limited version of betweenness centrality by considering only those shortest paths that lie between vertices in a given source and target subset.

Arguments:

graph: The graph object.

res: The result of the computation, a vector containing the betweenness score for the subset

of vertices.

vids: The vertices for which the subset-limited betweenness centrality scores will be com-

puted.

directed: Logical, if true directed paths will be considered for directed graphs. It is ignored for

undirected graphs.

weights: An optional vector containing edge weights for calculating weighted betweenness. No

edge weight may be NaN. Supply a null pointer here for unweighted betweenness.

sources: A vertex selector for the sources of the shortest paths taken into considuration in the

betweenness calculation.

targets: A vertex selector for the targets of the shortest paths taken into considuration in the

betweenness calculation.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex ID passed in *vids*, *sources* or *targets*

Time complexity: O(|S||E|), |S| The number of vertices in the subset |E| The number of edges in the graph

See also:

igraph_betweenness() to calculate the exact vertex betweenness and igraph_betweenness_cutoff() to calculate the range-limited vertex betweenness.

igraph_edge_betweenness_subset — Edge betweenness centrality for a subset of source and target vertices.

This function computes the subset-limited version of edge betweenness centrality by considering only those shortest paths that lie between vertices in a given source and target subset.

Arguments:

graph: The graph object.

res: The result of the computation, vector containing the betweenness scores for the edges.

eids: The edges for which the subset-limited betweenness centrality scores will be comput-

ed.

directed: Logical, if true directed paths will be considered for directed graphs. It is ignored for

undirected graphs.

weights: An optional weight vector for weighted betweenness. No edge weight may be NaN.

Supply a null pointer here for unweighted betweenness.

sources: A vertex selector for the sources of the shortest paths taken into considuration in the

betweenness calculation.

targets:

A vertex selector for the targets of the shortest paths taken into considuration in the betweenness calculation.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex ID passed in sources or targets

Time complexity: O(|S||E|), |S| The number of vertices in the subset |E| The number of edges in the graph

See also:

igraph_edge_betweenness() to compute the exact edge betweenness and igraph_edge_betweenness_cutoff() to compute the range-limited edge betweenness.

Centralization

igraph_centralization — Calculate the centralization score from the node level scores.

For a centrality score defined on the vertices of a graph, it is possible to define a graph level centralization index, by calculating the sum of the deviation from the maximum centrality score. Consequently, the higher the centralization index of the graph, the more centralized the structure is.

In order to make graphs of different sizes comparable, the centralization index is usually normalized to a number between zero and one, by dividing the (unnormalized) centralization score of the most centralized structure with the same number of vertices.

For most centrality indices the most centralized structure is the star graph, a single center connected to all other nodes in the network. There are some variation depending on whether the graph is directed or not, whether loop edges are allowed, etc.

This function simply calculates the graph level index, if the node level scores and the theoretical maximum are given. It is called by all the measure-specific centralization functions.

Arguments:

scores: A vector containing the node-level centrality scores.

theoretical_max: The graph level centrality score of the most centralized graph with the same

number of vertices. Only used if normalized set to true.

normalized: Boolean, whether to normalize the centralization by dividing the supplied

theoretical maximum.

Returns:

The graph level index.

See also:

Time complexity: O(n), the length of the score vector.

Example 13.16. File examples/simple/centralization.c

igraph_centralization_degree — Calculate vertex degree and graph centralization.

This function calculates the degree of the vertices by passing its arguments to <code>igraph_degree()</code>; and it calculates the graph level centralization index based on the results by calling <code>igraph_centralization()</code>.

Arguments:

graph: The input graph.

res: A vector if you need the node-level degree scores, or a null pointer other-

wise.

mode: Constant the specifies the type of degree for directed graphs. Possible val-

ues: IGRAPH_IN, IGRAPH_OUT and IGRAPH_ALL. This argument is

ignored for undirected graphs.

100ps: Boolean, whether to consider loop edges when calculating the degree (and

the centralization).

centralization: Pointer to a real number, the centralization score is placed here.

theoretical_max: Pointer to real number or a null pointer. If not a null pointer, then the theo-

retical maximum graph centrality score for a graph with the same number

vertices is stored here.

normalized: Boolean, whether to calculate a normalized centralization score. See

igraph_centralization() for how the normalization is done.

Returns:

Error code.

See also:

igraph_centralization(), igraph_degree().

Time complexity: the complexity of $igraph_degree()$ plus O(n), the number of vertices queried, for calculating the centralization score.

igraph_centralization_betweenness — Calculate vertex betweenness and graph centralization.

This function calculates the betweenness centrality of the vertices by passing its arguments to igraph_betweenness(); and it calculates the graph level centralization index based on the results by calling igraph_centralization().

Arguments:

graph: The input graph.

res: A vector if you need the node-level betweenness scores, or a null pointer

otherwise.

directed: Boolean, whether to consider directed paths when calculating betweenness.

centralization: Pointer to a real number, the centralization score is placed here.

theoretical_max: Pointer to real number or a null pointer. If not a null pointer, then the theo-

retical maximum graph centrality score for a graph with the same number

vertices is stored here.

normalized: Boolean, whether to calculate a normalized centralization score. See

igraph_centralization() for how the normalization is done.

Returns:

Error code.

See also:

```
igraph_centralization(), igraph_betweenness().
```

Time complexity: the complexity of $igraph_betweenness()$ plus O(n), the number of vertices queried, for calculating the centralization score.

igraph_centralization_closeness — Calculate vertex closeness and graph centralization.

```
igraph_real_t *centralization,
igraph_real_t *theoretical_max,
igraph bool t normalized);
```

This function calculates the closeness centrality of the vertices by passing its arguments to igraph_closeness(); and it calculates the graph level centralization index based on the results by calling igraph_centralization().

Arguments:

graph: The input graph.

res: A vector if you need the node-level closeness scores, or a null pointer oth-

erwise.

mode: Constant the specifies the type of closeness for directed graphs. Possible

values: IGRAPH_IN, IGRAPH_OUT and IGRAPH_ALL. This argument is ignored for undirected graphs. See igraph_closeness() argument

with the same name for more.

centralization: Pointer to a real number, the centralization score is placed here.

theoretical_max: Pointer to real number or a null pointer. If not a null pointer, then the theo-

retical maximum graph centrality score for a graph with the same number

vertices is stored here.

normalized: Boolean, whether to calculate a normalized centralization score. See

igraph_centralization() for how the normalization is done.

Returns:

Error code.

See also:

```
igraph_centralization(), igraph_closeness().
```

Time complexity: the complexity of igraph_closeness() plus O(n), the number of vertices queried, for calculating the centralization score.

igraph_centralization_eigenvector_centrality — Calculate eigenvector centrality scores and graph centralization.

```
igraph_error_t igraph_centralization_eigenvector_centrality(
   const igraph_t *graph,
   igraph_vector_t *vector,
   igraph_real_t *value,
   igraph_bool_t directed,
   igraph_bool_t scale,
   igraph_arpack_options_t *options,
   igraph_real_t *centralization,
   igraph_real_t *theoretical_max,
   igraph bool t normalized);
```

This function calculates the eigenvector centrality of the vertices by passing its arguments to igraph_eigenvector_centrality); and it calculates the graph level centralization index based on the results by calling igraph_centralization().

Arguments:

graph: The input graph.

vector: A vector if you need the node-level eigenvector centrality scores, or a null

pointer otherwise.

value: If not a null pointer, then the leading eigenvalue is stored here.

scale: If not zero then the result will be scaled, such that the absolute value of the

maximum centrality is one.

options: Options to ARPACK. See igraph_arpack_options_t for details.

Note that the function overwrites the n (number of vertices) parameter and it always starts the calculation from a non-random vector calculated based

on the degree of the vertices.

centralization: Pointer to a real number, the centralization score is placed here.

theoretical_max: Pointer to real number or a null pointer. If not a null pointer, then the theo-

retical maximum graph centrality score for a graph with the same number

vertices is stored here.

normalized: Boolean, whether to calculate a normalized centralization score. See

igraph_centralization() for how the normalization is done.

Returns:

Error code.

See also:

```
igraph_centralization(), igraph_eigenvector_centrality().
```

Time complexity: the complexity of $igraph_eigenvector_centrality()$ plus O(|V|), the number of vertices for the calculating the centralization.

igraph_centralization_degree_tmax — Theoretical maximum for graph centralization based on degree.

This function returns the theoretical maximum graph centrality based on vertex degree.

There are two ways to call this function, the first is to supply a graph as the graph argument, and then the number of vertices is taken from this object, and its directedness is considered as well. The

nodes argument is ignored in this case. The mode argument is also ignored if the supplied graph is undirected.

The other way is to supply a null pointer as the graph argument. In this case the nodes and mode arguments are considered.

The most centralized structure is the star. More specifically, for undirected graphs it is the star, for directed graphs it is the in-star or the out-star.

Arguments:

graph: A graph object or a null pointer, see the description above.

nodes: The number of nodes. This is ignored if the graph argument is not a null pointer.

mode: Constant, whether the calculation is based on in-degree (IGRAPH IN), out-degree

(IGRAPH_OUT) or total degree (IGRAPH_ALL). This is ignored if the graph argument

is not a null pointer and the given graph is undirected.

100ps: Boolean scalar, whether to consider loop edges in the calculation.

res: Pointer to a real variable, the result is stored here.

Returns:

Error code.

Time complexity: O(1).

See also:

igraph_centralization_degree() and igraph_centralization().

igraph_centralization_betweenness_tmax — Theoretical maximum for graph centralization based on betweenness.

This function returns the theoretical maximum graph centrality based on vertex betweenness.

There are two ways to call this function, the first is to supply a graph as the graph argument, and then the number of vertices is taken from this object, and its directedness is considered as well. The nodes argument is ignored in this case. The directed argument is also ignored if the supplied graph is undirected.

The other way is to supply a null pointer as the graph argument. In this case the nodes and directed arguments are considered.

The most centralized structure is the star.

Arguments:

graph: A graph object or a null pointer, see the description above.

nodes: The number of nodes. This is ignored if the graph argument is not a null pointer.

directed: Boolean scalar, whether to use directed paths in the betweenness calculation. This

argument is ignored if graph is not a null pointer and it is undirected.

res: Pointer to a real variable, the result is stored here.

Returns:

Error code.

Time complexity: O(1).

See also:

igraph_centralization_betweenness() and igraph_centralization().

igraph_centralization_closeness_tmax — Theoretical maximum for graph centralization based on closeness.

This function returns the theoretical maximum graph centrality based on vertex closeness.

There are two ways to call this function, the first is to supply a graph as the graph argument, and then the number of vertices is taken from this object, and its directedness is considered as well. The nodes argument is ignored in this case. The mode argument is also ignored if the supplied graph is undirected.

The other way is to supply a null pointer as the graph argument. In this case the nodes and mode arguments are considered.

The most centralized structure is the star.

Arguments:

graph: A graph object or a null pointer, see the description above.

nodes: The number of nodes. This is ignored if the graph argument is not a null pointer.

mode: Constant, specifies what kinf of distances to consider to calculate closeness. See the mode

argument of igraph_closeness() for details. This argument is ignored if graph is

not a null pointer and it is undirected.

res: Pointer to a real variable, the result is stored here.

Returns:

Error code.

Time complexity: O(1).

See also:

igraph_centralization_closeness() and igraph_centralization().

igraph_centralization_eigenvector_centrality_tmax — Theoretical maximum centralization for eigenvector centrality.

```
igraph_error_t igraph_centralization_eigenvector_centrality_tmax(
   const igraph_t *graph,
   igraph_integer_t nodes,
   igraph_bool_t directed,
   igraph_bool_t scale,
   igraph_real_t *res);
```

This function returns the theoretical maximum graph centrality based on vertex eigenvector centrality.

There are two ways to call this function, the first is to supply a graph as the graph argument, and then the number of vertices is taken from this object, and its directedness is considered as well. The nodes argument is ignored in this case. The directed argument is also ignored if the supplied graph is undirected.

The other way is to supply a null pointer as the graph argument. In this case the nodes and directed arguments are considered.

The most centralized directed structure is the in-star. The most centralized undirected structure is the graph with a single edge.

Arguments:

graph: A graph object or a null pointer, see the description above.

nodes: The number of nodes. This is ignored if the graph argument is not a null pointer.

directed: Boolean scalar, whether to consider edge directions. This argument is ignored if

graph is not a null pointer and it is undirected.

scale: Whether to rescale the node-level centrality scores to have a maximum of one.

res: Pointer to a real variable, the result is stored here.

Returns:

Error code.

Time complexity: O(1).

See also:

igraph_centralization_closeness() and igraph_centralization().

Similarity measures

igraph_bibcoupling — Bibliographic coupling.

The bibliographic coupling of two vertices is the number of other vertices they both cite, igraph_bibcoupling() calculates this. The bibliographic coupling score for each given vertex and all other vertices in the graph will be calculated.

Arguments:

graph: The graph object to analyze.

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows

is the same as the number of vertex IDs in vids, the number of columns is the number

of vertices in the graph.

vids: The vertex IDs of the vertices for which the calculation will be done.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex ID.

Time complexity: $O(|V|d^2)$, |V| is the number of vertices in the graph, d is the (maximum) degree of the vertices in the graph.

See also:

```
igraph_cocitation()
```

Example 13.17. File examples/simple/igraph_cocitation.c

igraph_cocitation — Cocitation coupling.

Two vertices are cocited if there is another vertex citing both of them. igraph_cocitation() simply counts how many times two vertices are cocited. The cocitation score for each given vertex and all other vertices in the graph will be calculated.

Arguments:

graph: The graph object to analyze.

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows

is the same as the number of vertex IDs in vids, the number of columns is the number

of vertices in the graph.

vids: The vertex IDs of the vertices for which the calculation will be done.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex ID.

Time complexity: $O(|V|d^2)$, |V| is the number of vertices in the graph, d is the (maximum) degree of the vertices in the graph.

See also:

igraph_bibcoupling()

Example 13.18. File examples/simple/igraph_cocitation.c

igraph_similarity_jaccard — Jaccard similarity coefficient for the given vertices.

The Jaccard similarity coefficient of two vertices is the number of common neighbors divided by the number of vertices that are neighbors of at least one of the two vertices being considered. This function calculates the pairwise Jaccard similarities for some (or all) of the vertices.

Arguments:

graph: The graph object to analyze

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows

and columns is the same as the number of vertex IDs in vids.

vids: The vertex IDs of the vertices for which the calculation will be done.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH_IN the incoming edges will be considered for each node.

 ${\tt IGRAPH_ALL} \quad \text{the directed graph is considered as an undirected one for the computation}.$

100ps: Whether to include the vertices themselves in the neighbor sets.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|V|^2 d)$, |V| is the number of vertices in the vertex iterator given, d is the (maximum) degree of the vertices in the graph.

See also:

igraph_similarity_dice(), a measure very similar to the Jaccard coefficient

Example 13.19. File examples/simple/igraph_similarity.c

igraph_similarity_jaccard_pairs — Jaccard similarity coefficient for given vertex pairs.

The Jaccard similarity coefficient of two vertices is the number of common neighbors divided by the number of vertices that are neighbors of at least one of the two vertices being considered. This function calculates the pairwise Jaccard similarities for a list of vertex pairs.

Arguments:

graph: The graph object to analyze

res: Pointer to a vector, the result of the calculation will be stored here. The number of elements

is the same as the number of pairs in pairs.

pairs: A vector that contains the pairs for which the similarity will be calculated. Each pair is

defined by two consecutive elements, i.e. the first and second element of the vector specific at the first and second element of the vector specific at the second arise and second element.

ifies the first pair, the third and fourth element specifies the second pair and so on.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH_IN the incoming edges will be considered for each node.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

100ps: Whether to include the vertices themselves in the neighbor sets.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(nd), n is the number of pairs in the given vector, d is the (maximum) degree of the vertices in the graph.

See also:

igraph_similarity_jaccard() to calculate the Jaccard similarity between all
pairs of a vertex set, or igraph_similarity_dice() and igraph_similarity_dice_pairs() for a measure very similar to the Jaccard coefficient

Example 13.20. File examples/simple/igraph_similarity.c

igraph_similarity_jaccard_es — Jaccard similarity coefficient for a given edge selector.

The Jaccard similarity coefficient of two vertices is the number of common neighbors divided by the number of vertices that are neighbors of at least one of the two vertices being considered. This function calculates the pairwise Jaccard similarities for the endpoints of edges in a given edge selector.

Arguments:

graph: The graph object to analyze

res: Pointer to a vector, the result of the calculation will be stored here. The number of elements

is the same as the number of edges in es.

es: An edge selector that specifies the edges to be included in the result.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH_IN the incoming edges will be considered for each node.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

100ps: Whether to include the vertices themselves in the neighbor sets.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(nd), n is the number of edges in the edge selector, d is the (maximum) degree of the vertices in the graph.

See also:

igraph_similarity_jaccard() and igraph_similarity_jaccard_pairs() to calculate the Jaccard similarity between all pairs of a vertex set or some selected vertex pairs, or igraph_similarity_dice(), igraph_similarity_dice_pairs() and igraph_similarity_dice_es() for a measure very similar to the Jaccard coefficient

Example 13.21. File examples/simple/igraph_similarity.c

igraph_similarity_dice — Dice similarity coefficient.

The Dice similarity coefficient of two vertices is twice the number of common neighbors divided by the sum of the degrees of the vertices. This function calculates the pairwise Dice similarities for some (or all) of the vertices.

Arguments:

graph: The graph object to analyze.

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows

and columns is the same as the number of vertex IDs in vids.

vids: The vertex IDs of the vertices for which the calculation will be done.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH_IN the incoming edges will be considered for each node.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

100ps: Whether to include the vertices themselves as their own neighbors.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: $O(|V|^2 d)$, where |V| is the number of vertices in the vertex iterator given, and d is the (maximum) degree of the vertices in the graph.

See also:

igraph_similarity_jaccard(), a measure very similar to the Dice coefficient

Example 13.22. File examples/simple/igraph_similarity.c

igraph_similarity_dice_pairs — Dice similarity coefficient for given vertex pairs.

The Dice similarity coefficient of two vertices is twice the number of common neighbors divided by the sum of the degrees of the vertices. This function calculates the pairwise Dice similarities for a list of vertex pairs.

Arguments:

graph: The graph object to analyze

res: Pointer to a vector, the result of the calculation will be stored here. The number of elements

is the same as the number of pairs in pairs.

pairs: A vector that contains the pairs for which the similarity will be calculated. Each pair is

defined by two consecutive elements, i.e. the first and second element of the vector specifies the first pair, the third and fourth element specifies the second pair and so on.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH IN the incoming edges will be considered for each node.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

loops: Whether to include the vertices themselves as their own neighbors.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(nd), n is the number of pairs in the given vector, d is the (maximum) degree of the vertices in the graph.

See also:

igraph_similarity_dice() to calculate the Dice similarity between all pairs of a vertex
set, or igraph_similarity_jaccard(), igraph_similarity_jaccard_pairs()
and igraph_similarity_jaccard_es() for a measure very similar to the Dice coefficient

Example 13.23. File examples/simple/igraph_similarity.c

igraph_similarity_dice_es — Dice similarity coefficient for a given edge selector.

The Dice similarity coefficient of two vertices is twice the number of common neighbors divided by the sum of the degrees of the vertices. This function calculates the pairwise Dice similarities for the endpoints of edges in a given edge selector.

Arguments:

graph: The graph object to analyze

res: Pointer to a vector, the result of the calculation will be stored here. The number of elements

is the same as the number of edges in es.

es: An edge selector that specifies the edges to be included in the result.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node.

IGRAPH_IN the incoming edges will be considered for each node.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

100ps: Whether to include the vertices themselves as their own neighbors.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID invalid vertex ID passed.

IGRAPH_EINVMODE invalid mode argument.

Time complexity: O(nd), n is the number of pairs in the given vector, d is the (maximum) degree of the vertices in the graph.

See also:

igraph_similarity_dice() and igraph_similarity_dice_pairs() to calculate the Dice similarity between all pairs of a vertex set or some selected vertex pairs, or igraph_similarity_jaccard(), igraph_similarity_jaccard_pairs() and igraph_similarity_jaccard_es() for a measure very similar to the Dice coefficient

Example 13.24. File examples/simple/igraph_similarity.c

igraph_similarity_inverse_log_weighted — Vertex similarity based on the inverse logarithm of vertex degrees.

The inverse log-weighted similarity of two vertices is the number of their common neighbors, weighted by the inverse logarithm of their degrees. It is based on the assumption that two vertices should be considered more similar if they share a low-degree common neighbor, since high-degree common neighbors are more likely to appear even by pure chance.

Isolated vertices will have zero similarity to any other vertex. Self-similarities are not calculated.

See the following paper for more details: Lada A. Adamic and Eytan Adar: Friends and neighbors on the Web. Social Networks, 25(3):211-230, 2003.

Arguments:

graph: The graph object to analyze.

res: Pointer to a matrix, the result of the calculation will be stored here. The number of its rows

is the same as the number of vertex IDs in vids, the number of columns is the number

of vertices in the graph.

vids: The vertex IDs of the vertices for which the calculation will be done.

mode: The type of neighbors to be used for the calculation in directed graphs. Possible values:

IGRAPH_OUT the outgoing edges will be considered for each node. Nodes will be

weighted according to their in-degree.

IGRAPH_IN the incoming edges will be considered for each node. Nodes will be

weighted according to their out-degree.

IGRAPH_ALL the directed graph is considered as an undirected one for the computation.

Every node is weighted according to its undirected degree.

Returns:

Error code: IGRAPH_EINVVID: invalid vertex ID.

Time complexity: $O(|V|d^2)$, |V| is the number of vertices in the graph, d is the (maximum) degree of the vertices in the graph.

Example 13.25. File examples/simple/igraph_similarity.c

Trees and forests

igraph_minimum_spanning_tree — Calculates one minimum spanning tree of a graph.

```
igraph_error_t igraph_minimum_spanning_tree(
     const igraph_t* graph, igraph_vector_int_t* res, const igraph_vector_t* wei
);
```

If the graph has more minimum spanning trees (this is always the case, except if it is a forest) this implementation returns only the same one.

Directed graphs are considered as undirected for this computation.

If the graph is not connected then its minimum spanning forest is returned. This is the set of the minimum spanning trees of each component.

Arguments:

graph: The graph object.

res: An initialized vector, the IDs of the edges that constitute a spanning tree will be returned

here. Use igraph_subgraph_edges() to extract the spanning tree as a separate

graph object.

weights: A vector containing the weights of the edges in the same order as the simple edge iterator

visits them (i.e. in increasing order of edge IDs).

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: O(|V|+|E|) for the unweighted case, $O(|E| \log |V|)$ for the weighted case. |V| is the number of vertices, |E| the number of edges in the graph.

See also:

igraph_minimum_spanning_tree_unweighted() and igraph_minimum_spanning_tree_prim() if you only need the tree as a separate graph object.

Example 13.26. File examples/simple/igraph minimum spanning tree.c

igraph_minimum_spanning_tree_unweighted — Calculates one minimum spanning tree of an unweighted graph.

If the graph has more minimum spanning trees (this is always the case, except if it is a forest) this implementation returns only the same one.

Directed graphs are considered as undirected for this computation.

If the graph is not connected then its minimum spanning forest is returned. This is the set of the minimum spanning trees of each component.

Arguments:

graph: The graph object.

mst: The minimum spanning tree, another graph object. Do not initialize this object before pass-

ing it to this function, but be sure to call <code>igraph_destroy()</code> on it if you don't need

it any more.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: O(|V|+|E|), |V| is the number of vertices, |E| the number of edges in the graph.

See also:

igraph_minimum_spanning_tree_prim() for weighted graphs, igraph_minimum_spanning_tree() if you need the IDs of the edges that constitute the spanning tree.

igraph_minimum_spanning_tree_prim — Calculates one minimum spanning tree of a weighted graph.

This function uses Prim's method for carrying out the computation, see Prim, R.C.: Shortest connection networks and some generalizations, Bell System Technical Journal, Vol. 36, 1957, 1389--1401.

If the graph has more than one minimum spanning tree, the current implementation returns always the same one.

Directed graphs are considered as undirected for this computation.

If the graph is not connected then its minimum spanning forest is returned. This is the set of the minimum spanning trees of each component.

Arguments:

graph: The graph object.

mst: The result of the computation, a graph object containing the minimum spanning tree of

the graph. Do not initialize this object before passing it to this function, but be sure to

call igraph_destroy() on it if you don't need it any more.

weights: A vector containing the weights of the edges in the same order as the simple edge iterator

visits them (i.e. in increasing order of edge IDs).

Returns:

Error code: IGRAPH_ENOMEM, not enough memory. IGRAPH_EINVAL, length of weight vector does not match number of edges.

Time complexity: $O(|E| \log |V|)$, |V| is the number of vertices, |E| the number of edges in the graph.

See also:

igraph_minimum_spanning_tree_unweighted() for unweighted graphs, igraph_minimum_spanning_tree() if you need the IDs of the edges that constitute the spanning tree.

Example 13.27. File examples/simple/igraph_minimum_spanning_tree.c

igraph_random_spanning_tree — Uniformly samples the spanning trees of a graph.

igraph_error_t igraph_random_spanning_tree(const igraph_t *graph, igraph_vector,

Performs a loop-erased random walk on the graph to uniformly sample its spanning trees. Edge directions are ignored.

Multi-graphs are supported, and edge multiplicities will affect the sampling frequency. For example, consider the 3-cycle graph 1=2-3-1, with two edges between vertices 1 and 2. Due to these parallel edges, the trees 1-2-3 and 3-1-2 will be sampled with multiplicity 2, while the tree 2-3-1 will be sampled with multiplicity 1.

Arguments:

graph: The input graph. Edge directions are ignored.

res: An initialized vector, the IDs of the edges that constitute a spanning tree will be returned

here. Use igraph_subgraph_edges() to extract the spanning tree as a separate

graph object.

vid: This parameter is relevant if the graph is not connected. If negative, a random spanning

forest of all components will be generated. Otherwise, it should be the ID of a vertex. A

random spanning tree of the component containing the vertex will be generated.

Returns:

Error code.

See also:

igraph_minimum_spanning_tree(), igraph_random_walk()

igraph_is_tree — Decides whether the graph is a tree.

igraph_error_t igraph_is_tree(const igraph_t *graph, igraph_bool_t *res, igraph_

An undirected graph is a tree if it is connected and has no cycles.

In the directed case, a possible additional requirement is that all edges are oriented away from a root (out-tree or arborescence) or all edges are oriented towards a root (in-tree or anti-arborescence). This test can be controlled using the *mode* parameter.

By convention, the null graph (i.e. the graph with no vertices) is considered not to be a tree.

Arguments:

graph: The graph object to analyze.

res: Pointer to a logical variable, the result will be stored here.

root: If not NULL, the root node will be stored here. When mode is IGRAPH_ALL or the graph

is undirected, any vertex can be the root and *root* is set to 0 (the first vertex). When *mode* is IGRAPH_OUT or IGRAPH_IN, the root is set to the vertex with zero in- or out-

degree, respectively.

mode: For a directed graph this specifies whether to test for an out-tree, an in-tree or ig-

nore edge directions. The respective possible values are: IGRAPH_OUT, IGRAPH_IN,

IGRAPH_ALL. This argument is ignored for undirected graphs.

Returns:

Error code: IGRAPH_EINVAL: invalid mode argument.

Time complexity: At most O(|V|+|E|), the number of vertices plus the number of edges in the graph.

See also:

```
igraph_is_connected()
```

Example 13.28. File examples/simple/igraph_kary_tree.c

igraph_is_forest — Decides whether the graph is a forest.

An undirected graph is a forest if it has no cycles.

In the directed case, a possible additional requirement is that edges in each tree are oriented away from the root (out-trees or arborescences) or all edges are oriented towards the root (in-trees or anti-arborescences). This test can be controlled using the *mode* parameter.

By convention, the null graph (i.e. the graph with no vertices) is considered to be a forest.

The res return value of this function is cached in the graph itself if mode is set to IGRAPH_ALL or if the graph is undirected. Calling the function multiple times with no modifications to the graph in between will return a cached value in O(1) time if the roots are not asked for.

Arguments:

graph: The graph object to analyze.

res: Pointer to a logical variable. If not NULL, then the result will be stored here.

roots: If not NULL, the root nodes will be stored here. When mode is IGRAPH_ALL or the

graph is undirected, any one vertex from each component can be the root. When mode is IGRAPH_OUT or IGRAPH_IN, all the vertices with zero in- or out-degree, respectively

are considered as root nodes.

mode: For a directed graph this specifies whether to test for an out-forest, an in-forest or ig-

nore edge directions. The respective possible values are: IGRAPH_OUT, IGRAPH_IN,

IGRAPH_ALL. This argument is ignored for undirected graphs.

Returns:

Error code: IGRAPH_EINVMODE: invalid mode argument.

Time complexity: At most O(|V|+|E|), the number of vertices plus the number of edges in the graph.

igraph_to_prufer — Converts a tree to its Prüfer sequence.

```
igraph_error_t igraph_to_prufer(const igraph_t *graph, igraph_vector_int_t* pru
```

A Prüfer sequence is a unique sequence of integers associated with a labelled tree. A tree on $n \ge 2$ vertices can be represented by a sequence of n-2 integers, each between 0 and n-1 (inclusive).

Arguments:

graph: Pointer to an initialized graph object which must be a tree on $n \ge 2$ vertices.

prufer: A pointer to the integer vector that should hold the Prüfer sequence; the vector must be

initialized and will be resized to n - 2.

Returns:

Error code:

IGRAPH_ENOMEM there is not enough memory to perform the operation.

IGRAPH_EINVAL the graph is not a tree or it is has less than vertices

See also:

igraph_from_prufer()

Transitivity or clustering coefficient

igraph_transitivity_undirected — Calculates the transitivity (clustering coefficient) of a graph.

The transitivity measures the probability that two neighbors of a vertex are connected. More precisely, this is the ratio of the triangles and connected triples in the graph, the result is a single real number. Directed graphs are considered as undirected ones and multi-edges are ignored.

Note that this measure is different from the local transitivity measure (see igraph_transitivi-ty_local_undirected()) as it calculates a single value for the whole graph.

Clustering coefficient is an alternative name for transitivity.

References:

S. Wasserman and K. Faust: Social Network Analysis: Methods and Applications. Cambridge: Cambridge University Press, 1994.

Arguments:

graph: The graph object. Edge directions and multiplicites are ignored.

res: Pointer to a real variable, the result will be stored here.

mode: Defines how to treat graphs with no connected triples. IGRAPH_TRANSITIVITY_NAN

returns NaN in this case, IGRAPH_TRANSITIVITY_ZERO returns zero.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory for temporary data.

See also:

```
igraph_transitivity_local_undirected(), igraph_transitivity_avglo-
cal_undirected().
```

Time complexity: $O(|V|*d^2)$, |V| is the number of vertices in the graph, d is the average node degree.

Example 13.29. File examples/simple/igraph_transitivity.c

igraph_transitivity_local_undirected — Calculates the local transitivity (clustering coefficient) of a graph.

The transitivity measures the probability that two neighbors of a vertex are connected. In case of the local transitivity, this probability is calculated separately for each vertex.

Note that this measure is different from the global transitivity measure (see igraph_transitiv-ity_undirected()) as it calculates a transitivity value for each vertex individually.

Clustering coefficient is an alternative name for transitivity.

References:

D. J. Watts and S. Strogatz: Collective dynamics of small-world networks. Nature 393(6684):440-442 (1998).

Arguments:

graph: The input graph. Edge directions and multiplicities are ignored.

res: Pointer to an initialized vector, the result will be stored here. It will be resized as needed.

vids: Vertex set, the vertices for which the local transitivity will be calculated.

mode: Defines how to treat vertices with degree less than two. IGRAPH_TRANSITIVITY_NAN returns NaN for these vertices, IGRAPH_TRANSITIVITY_ZERO returns zero.

Returns:

Error code.

See also:

Time complexity: $O(n*d^2)$, n is the number of vertices for which the transitivity is calculated, d is the average vertex degree.

igraph_transitivity_avglocal_undirected — Average local transitivity (clustering coefficient).

The transitivity measures the probability that two neighbors of a vertex are connected. In case of the average local transitivity, this probability is calculated for each vertex and then the average is taken. Vertices with less than two neighbors require special treatment, they will either be left out from the calculation or they will be considered as having zero transitivity, depending on the mode argument. Edge directions and edge multiplicities are ignored.

Note that this measure is different from the global transitivity measure (see igraph_transitiv-ity_undirected()) as it simply takes the average local transitivity across the whole network.

Clustering coefficient is an alternative name for transitivity.

References:

D. J. Watts and S. Strogatz: Collective dynamics of small-world networks. Nature 393(6684):440-442 (1998).

Arguments:

graph: The input graph. Edge directions and multiplicites are ignored.

res: Pointer to a real variable, the result will be stored here.

mode: Defines how to treat vertices with degree less than two. IGRAPH_TRANSITIVITY_NAN

leaves them out from averaging, IGRAPH_TRANSITIVITY_ZERO includes them with zero transitivity. The result will be NaN if the mode is IGRAPH_TRANSITIVITY_NAN

and there are no vertices with more than one neighbor.

Returns:

Error code.

See also:

```
igraph\_transitivity\_undirected(), \quad igraph\_transitivity\_local\_undirected().
```

Time complexity: $O(|V|*d^2)$, |V| is the number of vertices in the graph and d is the average degree.

igraph_transitivity_barrat — Weighted transitivity, as defined by A. Barrat.

This is a local transitivity, i.e. a vertex-level index. For a given vertex i, from all triangles in which it participates we consider the weight of the edges incident on i. The transitivity is the sum of these weights divided by twice the strength of the vertex (see igraph_strength()) and the degree of the vertex minus one. See equation (5) in Alain Barrat, Marc Barthelemy, Romualdo Pastor-Satorras, Alessandro Vespignani: The architecture of complex weighted networks, Proc. Natl. Acad. Sci. USA 101, 3747 (2004) at https://doi.org/10.1073/pnas.0400087101 for the exact formula.

Arguments:

graph: The input graph. Edge directions are ignored for directed graphs. Note that the function

does not work for non-simple graphs.

res: Pointer to an initialized vector, the result will be stored here. It will be resized as needed.

vids: The vertices for which the calculation is performed.

weights: Edge weights. If this is a null pointer, then a warning is given and igraph_transi-

tivity_local_undirected() is called.

mode: Defines how to treat vertices with zero strength. IGRAPH_TRANSITIVITY_NAN says

that the transitivity of these vertices is NaN, IGRAPH_TRANSITIVITY_ZERO says

it is zero.

Returns:

Error code.

Time complexity: $O(|V|*d^2)$, |V| is the number of vertices in the graph, d is the average node degree.

See also:

igraph_transitivity_undirected(), igraph_transitivity_local_undirected() and igraph_transitivity_avglocal_undirected() for other kinds of (non-weighted) transitivity.

Directedness conversion

igraph_to_directed — Convert an undirected graph to a directed one.

If the supplied graph is directed, this function does nothing.

Arguments:

graph: The graph object to convert.

mode: Constant, specifies the details of how exactly the conversion is done. Possible values:

IGRAPH_TO_DIRECTED_ARBI-

TRARY

The number of edges in the graph stays the same, an arbitrarily directed edge is created for each undirected edge.

Two directed edges are created for each undirected IGRAPH_TO_DIRECTED_MUTUedge, one in each direction. IGRAPH_TO_DIRECTED_RAN-Each undirected edge is converted to a randomly

IGRAPH_TO_DIRECTED_ACY-

Each undirected edge is converted to a directed edge oriented from a lower index vertex to a higher CLIC index one. If no self-loops were present, then the

oriented directed one.

result is a directed acyclic graph.

Returns:

Error code.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

igraph to undirected — Convert a directed graph to an undirected one.

```
igraph error t igraph to undirected(igraph t *graph,
                         igraph_to_undirected_t mode,
                         const igraph_attribute_combination_t *edge_comb);
```

If the supplied graph is undirected, this function does nothing.

Arguments:

The graph object to convert. graph:

Constant, specifies the details of how exactly the conversion is done. Possible valmode:

> ues: IGRAPH_TO_UNDIRECTED_EACH: the number of edges remains constant, an undirected edge is created for each directed one, this version might create graphs with multiple edges; IGRAPH_TO_UNDIRECTED_COLLAPSE: one undirected edge will be created for each pair of vertices that are connected with at least one directed edge, no multiple edges will be created. IGRAPH_TO_UNDIRECTED_MUTUAL creates an undirected edge for each pair of mutual edges in the directed graph. Nonmutual edges are lost; loop edges are kept unconditionally. This mode might create

multiple edges.

edge_comb: What to do with the edge attributes. See the igraph manual section about attribut-

es for details. NULL means that the edge attributes are lost during the conversion, except when mode is IGRAPH_TO_UNDIRECTED_EACH, in which case the edge

attributes are kept intact.

Returns:

Error code.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

Example 13.30. File examples/simple/igraph_to_undirected.c

Spectral properties

igraph_get_laplacian — Returns the Laplacian matrix of a graph.

```
igraph_error_t igraph_get_laplacian(
    const igraph_t *graph, igraph_matrix_t *res, igraph_neimode_t mode,
    igraph_laplacian_normalization_t normalization,
    const igraph_vector_t *weights
);
```

The Laplacian matrix L of a graph is defined as L_ij = - A_ij when i != j and L_ii = d_i - A_ii. Here A denotes the (possibly weighted) adjacency matrix and d_i is the degree (or strength, if weighted) of vertex i. In directed graphs, the mode parameter controls whether to use outor in-degrees. Correspondingly, the rows or columns will sum to zero. In undirected graphs, A_ii is taken to be twice the number (or total weight) of self-loops, ensuring that d_i = \sum_j A_ij. Thus, the Laplacian of an undirected graph is the same as the Laplacian of a directed one obtained by replacing each undirected edge with two reciprocal directed ones.

More compactly, L = D - A where the D is a diagonal matrix containing the degrees. The Laplacian matrix can also be normalized, with several conventional normalization methods. See igraph_laplacian_normalization_t for the methods available in igraph.

The first version of this function was written by Vincent Matossian.

Arguments:

graph: Pointer to the graph to convert.

res: Pointer to an initialized matrix object, the result is stored here. It will be resized

if needed.

mode: Controls whether to use out- or in-degrees in directed graphs. If set to

IGRAPH_ALL, edge directions will be ignored.

normalization: The normalization method to use when calculating the Laplacian matrix. See

igraph_laplacian_normalization_t for possible values.

weights: An optional vector containing non-negative edge weights, to calculate the

weighted Laplacian matrix. Set it to a null pointer to calculate the unweighted

Laplacian.

Returns:

Error code.

Time complexity: $O(|V|^2)$, |V| is the number of vertices in the graph.

Example 13.31. File examples/simple/igraph_get_laplacian.c

igraph_get_laplacian_sparse — Returns the Laplacian of a graph in a sparse matrix format.

```
igraph_error_t igraph_get_laplacian_sparse(
    const igraph_t *graph, igraph_sparsemat_t *sparseres, igraph_neimode_t mode
    igraph_laplacian_normalization_t normalization,
    const igraph_vector_t *weights
);
```

See igraph_get_laplacian() for the definition of the Laplacian matrix.

The first version of this function was written by Vincent Matossian.

Arguments:

graph: Pointer to the graph to convert.

sparseres: Pointer to an initialized sparse matrix object, the result is stored here.

mode: Controls whether to use out- or in-degrees in directed graphs. If set to

IGRAPH_ALL, edge directions will be ignored.

normalization: The normalization method to use when calculating the Laplacian matrix. See

igraph_laplacian_normalization_t for possible values.

weights: An optional vector containing non-negative edge weights, to calculate the

weighted Laplacian matrix. Set it to a null pointer to calculate the unweighted

Laplacian.

Returns:

Error code.

Time complexity: O(|E|), |E| is the number of edges in the graph.

Example 13.32. File examples/simple/igraph get laplacian sparse.c

igraph_laplacian_normalization_t — Normalization methods for a Laplacian matrix.

```
typedef enum {
    IGRAPH_LAPLACIAN_UNNORMALIZED = 0,
    IGRAPH_LAPLACIAN_SYMMETRIC = 1,
    IGRAPH_LAPLACIAN_LEFT = 2,
    IGRAPH_LAPLACIAN_RIGHT = 3
} igraph_laplacian_normalization_t;
```

Normalization methods for <code>igraph_get_laplacian()</code> and <code>igraph_get_laplacian_s-parse()</code>. In the following, A refers to the (possibly weighted) adjacency matrix and D is a diagonal matrix containing degrees (unweighted case) or strengths (weighted case). Out-, in- or total degrees are used according to the <code>mode</code> parameter.

Values:

```
IGRAPH_LAPLACIAN_UNNOR- Unnormalized Laplacian, L = D - A.
MALIZED:
```

Non-simple graphs: Multiple and loop edges

igraph_is_simple — Decides whether the input graph is a simple graph.

```
igraph_error_t igraph_is_simple(const igraph_t *graph, igraph_bool_t *res);
```

Right-stochastic normalized Laplacian, L = I - A D^-1.

A graph is a simple graph if it does not contain loop edges and multiple edges.

Arguments:

graph: The input graph.

IGRAPH_LAPLACIAN_RIGHT:

res: Pointer to a boolean constant, the result is stored here.

Returns:

Error code.

See also:

igraph_is_loop() and igraph_is_multiple() to find the loops and multiple edges, igraph_simplify() to get rid of them, or igraph_has_multiple() to decide whether there is at least one multiple edge.

Time complexity: O(|V|+|E|).

igraph_is_loop — Find the loop edges in a graph.

A loop edge is an edge from a vertex to itself.

Arguments:

graph: The input graph.

res: Pointer to an initialized boolean vector for storing the result, it will be resized as needed.

es: The edges to check, for all edges supply igraph_ess_all() here.

Returns:

Error code.

See also:

igraph_simplify() to get rid of loop edges.

Time complexity: O(e), the number of edges to check.

Example 13.33. File examples/simple/igraph_is_loop.c

igraph_is_multiple — Find the multiple edges in a graph.

An edge is a multiple edge if there is another edge with the same head and tail vertices in the graph.

Note that this function returns true only for the second or more appearances of the multiple edges.

Arguments:

graph: The input graph.

res: Pointer to a boolean vector, the result will be stored here. It will be resized as needed.

es: The edges to check. Supply igraph_ess_all() if you want to check all edges.

Returns:

Error code.

See also:

```
igraph_count_multiple(), igraph_has_multiple() and igraph_simplify().
```

Time complexity: O(e*d), e is the number of edges to check and d is the average degree (out-degree in directed graphs) of the vertices at the tail of the edges.

Example 13.34. File examples/simple/igraph_is_multiple.c

igraph_has_multiple — Check whether the graph has at least one multiple edge.

```
igraph_error_t igraph_has_multiple(const igraph_t *graph, igraph_bool_t *res);
```

An edge is a multiple edge if there is another edge with the same head and tail vertices in the graph.

The return value of this function is cached in the graph itself; calling the function multiple times with no modifications to the graph in between will return a cached value in O(1) time.

Arguments:

graph: The input graph.

res: Pointer to a boolean variable, the result will be stored here.

Returns:

Error code.

See also:

```
igraph_count_multiple(),igraph_is_multiple() and igraph_simplify().
```

Time complexity: O(e*d), e is the number of edges to check and d is the average degree (out-degree in directed graphs) of the vertices at the tail of the edges.

Example 13.35. File examples/simple/igraph_has_multiple.c

igraph_count_multiple — Count the number of appearances of the edges in a graph.

```
igraph_error_t igraph_count_multiple(const igraph_t *graph, igraph_vector_int_t
```

If the graph has no multiple edges then the result vector will be filled with ones. (An edge is a multiple edge if there is another edge with the same head and tail vertices in the graph.)

Arguments:

graph: The input graph.

res: Pointer to a vector, the result will be stored here. It will be resized as needed.

es: The edges to check. Supply igraph_ess_all() if you want to check all edges.

Returns:

Error code.

See also:

```
igraph_is_multiple() and igraph_simplify().
```

Time complexity: O(E d), E is the number of edges to check and d is the average degree (out-degree in directed graphs) of the vertices at the tail of the edges.

Mixing patterns

igraph_assortativity_nominal — Assortativity of a graph based on vertex categories.

Assuming the vertices of the input graph belong to different categories, this function calculates the assortativity coefficient of the graph. The assortativity coefficient is between minus one and one and it is one if all connections stay within categories, it is minus one, if the network is perfectly disassortative. For a randomly connected network it is (asymptotically) zero.

The unnormalized version, computed when normalized is set to false, is identical to the modularity, and is defined as follows for directed networks:

```
1/m sum ij (A ij - k^out i k^in j / m) d(i,j),
```

where m denotes the number of edges, A_{ij} is the adjacency matrix, k^{out} and k^{in} are the outand in-degrees, and d(i,j) is one if vertices i and j are in the same category and zero otherwise.

The normalized assortativity coefficient is obtained by dividing the previous expression by

```
1/m \quad sum_{ij} \quad (m - k^out_i \quad k^in_j \quad d(i,j) / m).
```

It can take any value within the interval [-1, 1].

Undirected graphs are effectively treated as directed ones with all-reciprocal edges. Thus, self-loops are taken into account twice in undirected graphs.

References:

M. E. J. Newman: Mixing patterns in networks, Phys. Rev. E 67, 026126 (2003) https://doi.org/10.1103/PhysRevE.67.026126. See section II and equation (2) for the definition of the concept.

For an educational overview of assortativity, see M. E. J. Newman, Networks: An Introduction, Oxford University Press (2010). https://doi.org/10.1093/acprof%3Aoso/9780199206650.001.0001.

Arguments:

graph: The input graph, it can be directed or undirected.

types: Integer vector giving the vertex categories. The types are represented by integers

starting at zero.

res: Pointer to a real variable, the result is stored here.

directed: Boolean, it gives whether to consider edge directions in a directed graph. It is ig-

nored for undirected graphs.

normalized: Boolean, whether to compute the usual normalized assortativity. The unnormalized

version is identical to modularity. Supply true here to compute the standard assor-

tativity.

Returns:

Error code.

Time complexity: O(|E|+t), |E| is the number of edges, t is the number of vertex types.

See also:

igraph_assortativity() for computing the assortativity based on continuous vertex values
instead of discrete categories. igraph_modularity() to compute generalized modularity.

Example 13.36. File examples/simple/igraph_assortativity_nominal.c

igraph_assortativity — Assortativity based on numeric properties of vertices.

This function calculates the assortativity coefficient of a graph based on given values x_i for each vertex i. This type of assortativity coefficient equals the Pearson correlation of the values at the two ends of the edges.

The unnormalized covariance of values, computed when *normalized* is set to false, is defined as follows in a directed graph:

```
cov(x_out, x_in) = 1/m sum_ij (A_ij - k^out_i k^in_j / m) x_i x_j,
```

where m denotes the number of edges, A_ij is the adjacency matrix, and k^out and k^in are the out- and in-degrees. x_out and x_in refer to the sets of vertex values at the start and end of the directed edges.

The normalized covariance, i.e. Pearson correlation, is obtained by dividing the previous expression by $sqrt(var(x_out))$ $sqrt(var(x_in))$, where

```
var(x_out) = 1/m sum_i k^out_i x_i^2 - (1/m sum_i k^out_i x_i^2)^2
var(x in) = 1/m sum j k^in j x j^2 - (1/m sum j k^in j x j^2)^2
```

Undirected graphs are effectively treated as directed graphs where all edges are reciprocal. Therefore, self-loops are effectively considered twice in undirected graphs.

References:

M. E. J. Newman: Mixing patterns in networks, Phys. Rev. E 67, 026126 (2003) https://doi.org/10.1103/PhysRevE.67.026126. See section III and equation (21) for the definition, and equation (26) for performing the calculation in directed graphs with the degrees as values.

M. E. J. Newman: Assortative mixing in networks, Phys. Rev. Lett. 89, 208701 (2002) http://doi.org/10.1103/PhysRevLett.89.208701. See equation (4) for performing the calculation in undirected graphs with the degrees as values.

For an educational overview of the concept of assortativity, see M. E. J. Newman, Networks: An Introduction, Oxford University Press (2010). https://doi.org/10.1093/acprof %3Aoso/9780199206650.001.0001.

Arguments:

graph: The input graph, it can be directed or undirected.

values: The vertex values, these can be arbitrary numeric values.

values_in: A second value vector to be used for the incoming edges when calculating assor-

tativity for a directed graph. Supply NULL here if you want to use the same values for outgoing and incoming edges. This argument is ignored (with a warning) if it is not a null pointer and the undirected assortativity coefficient is being calculated.

res: Pointer to a real variable, the result is stored here.

directed: Boolean, whether to consider edge directions for directed graphs. It is ignored for

undirected graphs.

normalized: Boolean, whether to compute the normalized covariance, i.e. Pearson correlation.

Supply true here to compute the standard assortativity.

Returns:

Error code.

Time complexity: O(|E|), linear in the number of edges of the graph.

See also:

igraph_assortativity_nominal() if you have discrete vertex categories instead of numeric labels, and igraph_assortativity_degree() for the special case of assortativity based on vertex degrees.

igraph_assortativity_degree — Assortativity of a graph based on vertex degree.

Assortativity based on vertex degree, please see the discussion at the documentation of <code>igraph_as-sortativity()</code> for details. This function simply calls <code>igraph_assortativity()</code> with the degrees as the vertex values and normalization enabled. In the directed case, it uses out-degrees as out-values and in-degrees as in-values.

For regular graphs, i.e. graphs in which all vertices have the same degree, computing degree correlations is not meaningful, and this function returns NaN.

Arguments:

graph: The input graph, it can be directed or undirected.

res: Pointer to a real variable, the result is stored here.

directed: Boolean, whether to consider edge directions for directed graphs. This argument is

ignored for undirected graphs. Supply true here to do the natural thing, i.e. use directed version of the measure for directed graphs and the undirected version for undirected

graphs.

Returns:

Error code.

Time complexity: O(|E|+|V|), |E| is the number of edges, |V| is the number of vertices.

See also:

igraph_assortativity() for the general function calculating assortativity for any kind of numeric vertex values.

Example 13.37. File examples/simple/igraph_assortativity_degree.c

K-Cores and K-Trusses

igraph_coreness — Finding the coreness of the vertices in a network.

The k-core of a graph is a maximal subgraph in which each vertex has at least degree k. (Degree here means the degree in the subgraph of course.). The coreness of a vertex is the highest order of a k-core containing the vertex.

This function implements the algorithm presented in Vladimir Batagelj, Matjaz Zaversnik: An O(m) Algorithm for Cores Decomposition of Networks.

Arguments:

graph: The input graph.

cores: Pointer to an initialized vector, the result of the computation will be stored here. It will be resized as needed. For each vertex it contains the highest order of a core containing the vertex.

mode: For directed graph it specifies whether to calculate in-cores, out-cores or the undirected version. It is ignored for undirected graphs. Possible values: IGRAPH_ALL undirected version, IGRAPH_IN in-cores, IGRAPH_OUT out-cores.

Returns:

Error code.

Time complexity: O(|E|), the number of edges.

igraph_trussness — Finding the "trussness" of the edges in a network.

```
igraph_error_t igraph_trussness(const igraph_t* graph, igraph_vector_int_t* tru
```

A k-truss is a subgraph in which every edge occurs in at least k-2 triangles in the subgraph. The trussness of an edge indicates the highest k-truss that the edge occurs in.

This function returns the highest k for each edge. If you are interested in a particular k-truss subgraph, you can subset the graph using to those eids which are >= k because each k-truss is a subgraph of a (k–1)-truss (thus to get all 4-trusses, take k >= 4 because the 5-trusses, 6-trusses, etc need to be included).

The current implementation of this function iteratively decrements support of each edge using O(|E|) space and $O(|E|^{1.5})$ time. The implementation does not support multigraphs; use igraph_simplify() to collapse edges before calling this function.

Arguments:

graph: The input graph. Loop edges are allowed; multigraphs are not.

truss: Pointer to initialized vector of truss values that will indicate the highest k-truss each edge

occurs in. It will be resized as needed.

Returns:

Error code.

Time complexity: It should be $O(|E|^{\Lambda}1.5)$. See Algorithm 2 in: Wang, Jia, and James Cheng. "Truss decomposition in massive networks." Proceedings of the VLDB Endowment 5.9 (2012): 812-823.

Topological sorting, directed acyclic graphs

igraph_is_dag — Checks whether a graph is a directed acyclic graph (DAG).

```
igraph_error_t igraph_is_dag(const igraph_t* graph, igraph_bool_t *res);
```

A directed acyclic graph (DAG) is a directed graph with no cycles.

The return value of this function is cached in the graph itself; calling the function multiple times with no modifications to the graph in between will return a cached value in O(1) time.

Arguments:

graph: The input graph.

res: Pointer to a boolean constant, the result is stored here.

Returns:

Error code.

Time complexity: O(|V|+|E|), where |V| and |E| are the number of vertices and edges in the original input graph.

See also:

igraph_topological_sorting() to get a possible topological sorting of a DAG.

igraph_topological_sorting — Calculate a possible topological sorting of the graph.

A topological sorting of a directed acyclic graph (DAG) is a linear ordering of its vertices where each vertex comes before all nodes to which it has edges. Every DAG has at least one topological sort, and may have many. This function returns one possible topological sort among them. If the graph contains any cycles that are not self-loops, an error is raised.

Arguments:

graph: The input graph.

res: Pointer to a vector, the result will be stored here. It will be resized if needed.

mode: Specifies how to use the direction of the edges. For IGRAPH_OUT, the sorting order en-

sures that each vertex comes before all vertices to which it has edges, so vertices with no incoming edges go first. For IGRAPH_IN, it is quite the opposite: each vertex comes before all vertices from which it receives edges. Vertices with no outgoing edges go first.

Returns:

Error code.

Time complexity: O(|V|+|E|), where |V| and |E| are the number of vertices and edges in the original input graph.

See also:

igraph_is_dag() if you are only interested in whether a given graph is a DAG or not, or igraph_feedback_arc_set() to find a set of edges whose removal makes the graph acyclic.

Example 13.38. File examples/simple/igraph_topological_sorting.c

igraph_feedback_arc_set — Feedback arc set of a graph using exact or heuristic methods.

A feedback arc set is a set of edges whose removal makes the graph acyclic. We are usually interested in *minimum* feedback arc sets, i.e. sets of edges whose total weight is minimal among all the feedback arc sets.

For undirected graphs, the problem is simple: one has to find a maximum weight spanning tree and then remove all the edges not in the spanning tree. For directed graphs, this is an NP-hard problem, and various heuristics are usually used to find an approximate solution to the problem. This function implements a few of these heuristics.

Arguments:

graph: The graph object.

result: An initialized vector, the result will be returned here.

weights: Weight vector or NULL if no weights are specified.

algo: The algorithm to use to solve the problem if the graph is directed. Possible values:

IGRAPH_FAS_EXACT_IP Finds a minimum feedback arc set using integer

programming (IP). The complexity of this algo-

rithm is exponential of course.

IGRAPH_FAS_APPROX_EADES Finds a feedback arc set using the heuristic of

Eades, Lin and Smyth (1993). This is guaranteed to be smaller than |E|/2 - |V|/6, and it is linear in the number of edges (i.e. O(|E|)). For more details, see Eades P, Lin X and Smyth WF: A fast and effective heuristic for the feedback arc set problem.

In: Proc Inf Process Lett 319-323, 1993.

Returns:

Error code: IGRAPH_EINVAL if an unknown method was specified or the weight vector is invalid.

Example 13.39. File examples/simple/igraph_feedback_arc_set.c

Example 13.40. File examples/simple/igraph_feedback_arc_set_ip.c

Time complexity: depends on algo, see the time complexities there.

Maximum cardinality search and chordal graphs

igraph_maximum_cardinality_search — Maximum
cardinality search.

This function implements the maximum cardinality search algorithm. It computes a rank alpha for each vertex, such that visiting vertices in decreasing rank order corresponds to always choosing the vertex with the most already visited neighbors as the next one to visit.

Maximum cardinality search is useful in deciding the chordality of a graph. A graph is chordal if and only if any two neighbors of a vertex which are higher in rank than it are connected to each other.

References:

Robert E Tarjan and Mihalis Yannakakis: Simple linear-time algorithms to test chordality of graphs, test acyclicity of hypergraphs, and selectively reduce acyclic hypergraphs. SIAM Journal of Computation 13, 566-579, 1984. https://doi.org/10.1137/0213035

Arguments:

graph: The input graph. Edge directions will be ignored.

alpha: Pointer to an initialized vector, the result is stored here. It will be resized, as needed.

Upon return it contains the rank of the each vertex in the range 0 to n - 1, where n

is the number of vertices.

alpham1: Pointer to an initialized vector or a NULL pointer. If not NULL, then the inverse of

alpha is stored here. In other words, the elements of alpham1 are vertex IDs in

reverse maximum cardinality search order.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in terms of the number of vertices and edges.

See also:

```
igraph_is_chordal().
```

igraph_is_chordal — Decides whether a graph is chordal.

A graph is chordal if each of its cycles of four or more nodes has a chord, i.e. an edge joining two nodes that are not adjacent in the cycle. An equivalent definition is that any chordless cycles have at most three nodes. If either <code>alpha</code> or <code>alpham1</code> is given, then the other is calculated by taking simply the inverse. If neither are given, then <code>igraph_maximum_cardinality_search()</code> is called to calculate them.

Arguments:

graph: The input graph. Edge directions will be ignored.

alpha: Either an alpha vector coming from igraph_maximum_cardinali-

ty_search() (on the same graph), or a NULL pointer.

alpham1: Either an inverse alpha vector coming from igraph_maximum_cardinali-

ty_search() (on the same graph) or a NULL pointer.

chorda1: Pointer to a boolean. If not NULL the result is stored here.

fill_in: Pointer to an initialized vector, or a NULL pointer. If not a NULL pointer, then the

fill-in, also called the chordal completion of the graph is stored here. The chordal completion is a set of edges that are needed to make the graph chordal. The vector is resized as needed. Note that the chordal completion returned by this function may not be minimal, i.e. some of the returned fill-in edges may not be needed to make the

graph chordal.

newgraph: Pointer to an uninitialized graph, or a NULL pointer. If not a null pointer, then a new

triangulated graph is created here. This essentially means adding the fill-in edges to

the original graph.

Returns:

Error code.

Time complexity: O(n).

See also:

igraph maximum cardinality search().

Matchings

igraph_is_matching — Checks whether the given matching is valid for the given graph.

This function checks a matching vector and verifies whether its length matches the number of vertices in the given graph, its values are between -1 (inclusive) and the number of vertices (exclusive), and whether there exists a corresponding edge in the graph for every matched vertex pair. For bipartite graphs, it also verifies whether the matched vertices are in different parts of the graph.

Arguments:

graph: The input graph. It can be directed but the edge directions will be ignored.

types: If the graph is bipartite and you are interested in bipartite matchings only, pass the

vertex types here. If the graph is non-bipartite, simply pass NULL.

matching: The matching itself. It must be a vector where element i contains the ID of the vertex

that vertex i is matched to, or -1 if vertex i is unmatched.

result: Pointer to a boolean variable, the result will be returned here.

See also:

igraph_is_maximal_matching() if you are also interested in whether the matching is maximal (i.e. non-extendable).

Time complexity: O(|V|+|E|) where |V| is the number of vertices and |E| is the number of edges.

Example 13.41. File examples/simple/igraph_maximum_bipartite_matching.c

igraph_is_maximal_matching — Checks whether a matching in a graph is maximal.

```
igraph_error_t igraph_is_maximal_matching(const igraph_t *graph,
```

```
const igraph_vector_bool_t *types, const igraph_
igraph_bool_t *result);
```

A matching is maximal if and only if there exists no unmatched vertex in a graph such that one of its neighbors is also unmatched.

Arguments:

graph: The input graph. It can be directed but the edge directions will be ignored.

types: If the graph is bipartite and you are interested in bipartite matchings only, pass the

vertex types here. If the graph is non-bipartite, simply pass NULL.

matching: The matching itself. It must be a vector where element i contains the ID of the vertex

that vertex i is matched to, or -1 if vertex i is unmatched.

result: Pointer to a boolean variable, the result will be returned here.

See also:

igraph_is_matching() if you are only interested in whether a matching vector is valid for a given graph.

Time complexity: O(|V|+|E|) where |V| is the number of vertices and |E| is the number of edges.

Example 13.42. File examples/simple/igraph_maximum_bipartite_matching.c

igraph_maximum_bipartite_matching — Calculates a maximum matching in a bipartite graph.

A matching in a bipartite graph is a partial assignment of vertices of the first kind to vertices of the second kind such that each vertex of the first kind is matched to at most one vertex of the second kind and vice versa, and matched vertices must be connected by an edge in the graph. The size (or cardinality) of a matching is the number of edges. A matching is a maximum matching if there exists no other matching with larger cardinality. For weighted graphs, a maximum matching is a matching whose edges have the largest possible total weight among all possible matchings.

Maximum matchings in bipartite graphs are found by the push-relabel algorithm with greedy initialization and a global relabeling after every n/2 steps where n is the number of vertices in the graph.

References: Cherkassky BV, Goldberg AV, Martin P, Setubal JC and Stolfi J: Augment or push: A computational study of bipartite matching and unit-capacity flow algorithms. ACM Journal of Experimental Algorithmics 3, 1998.

Kaya K, Langguth J, Manne F and Ucar B: Experiments on push-relabel-based maximum cardinality matching algorithms for bipartite graphs. Technical Report TR/PA/11/33 of the Centre Europeen de Recherche et de Formation Avancee en Calcul Scientifique, 2011.

Arguments:

graph: The input graph. It can be directed but the edge directions will be ignored.

types: Boolean vector giving the vertex types of the graph.

matching_size: The size of the matching (i.e. the number of matched vertex pairs will be

returned here). It may be NULL if you don't need this.

matching_weight: The weight of the matching if the edges are weighted, or the size of the

matching again if the edges are unweighted. It may be NULL if you don't

need this.

matching: The matching itself. It must be a vector where element i contains the ID of

the vertex that vertex i is matched to, or -1 if vertex i is unmatched.

weights: A null pointer (=no edge weights), or a vector giving the weights of the

edges. Note that the algorithm is stable only for integer weights.

eps: A small real number used in equality tests in the weighted bipartite match-

ing algorithm. Two real numbers are considered equal in the algorithm if their difference is smaller than eps. This is required to avoid the accumulation of numerical errors. It is advised to pass a value derived from the DBL_EPSILON constant in float.h here. If you are running the algorithm

with no weights vector, this argument is ignored.

Returns:

Error code.

Time complexity: $O(\operatorname{sqrt}(|V|)|E|)$ for unweighted graphs (according to the technical report referenced above), O(|V||E|) for weighted graphs.

Example 13.43. File examples/simple/igraph_maximum_bipartite_matching.c

Unfolding a graph into a tree

igraph_unfold_tree — Unfolding a graph into a tree, by possibly multiplicating its vertices.

A graph is converted into a tree (or forest, if it is unconnected), by performing a breadth-first search on it, and replicating vertices that were found a second, third, etc. time.

Arguments:

graph: The input graph, it can be either directed or undirected.

tree: Pointer to an uninitialized graph object, the result is stored here.

mode: For directed graphs; whether to follow paths along edge directions

 $({\tt IGRAPH_OUT}), or the opposite ({\tt IGRAPH_IN}), or ignore edge directions common and other properties of the common of the$

pletely (IGRAPH_ALL). It is ignored for undirected graphs.

roots: A numeric vector giving the root vertex, or vertices (if the graph is not connect-

ed), to start from.

vertex_index: Pointer to an initialized vector, or a null pointer. If not a null pointer, then a

mapping from the vertices in the new graph to the ones in the original is created

here.

Returns:

Error code.

Time complexity: O(n+m), linear in the number vertices and edges.

Other operations

igraph_density — Calculate the density of a graph.

The density of a graph is simply the ratio of the actual number of its edges and the largest possible number of edges it could have. The maximum number of edges depends on interpretation: are vertices allowed to have a connected to themselves? This is controlled by the *loops* parameter.

Note that density is ill-defined for graphs which have multiple edges between some pairs of vertices. Consider calling <code>igraph_simplify()</code> on such graphs.

Arguments:

graph: The input graph object.

res: Pointer to a real number, the result will be stored here.

loops: Logical constant, whether to include self-loops in the calculation. If this constant is true

then loop edges are thought to be possible in the graph (this does not necessarily mean that the graph really contains any loops). If this is false then the result is only correct if the

graph does not contain loops.

Returns:

Error code.

Time complexity: O(1).

igraph_reciprocity — Calculates the reciprocity of a directed graph.

The measure of reciprocity defines the proportion of mutual connections, in a directed graph. It is most commonly defined as the probability that the opposite counterpart of a directed edge is also included

in the graph. In adjacency matrix notation: sum(i, j, (A.*A')ij) / sum(i, j, Aij), where A.*A' is the element-wise product of matrix A and its transpose. This measure is calculated if the *mode* argument is IGRAPH RECIPROCITY DEFAULT.

Prior to igraph version 0.6, another measure was implemented, defined as the probability of mutual connection between a vertex pair if we know that there is a (possibly non-mutual) connection between them. In other words, (unordered) vertex pairs are classified into three groups: (1) disconnected, (2) non-reciprocally connected, (3) reciprocally connected. The result is the size of group (3), divided by the sum of group sizes (2)+(3). This measure is calculated if *mode* is IGRAPH_RECIPROCITY_RATIO.

Arguments:

graph: The graph object.

res: Pointer to an igraph_real_t which will contain the result.

ignore_loops: Whether to ignore loop edges.

mode: Type of reciprocity to calculate, possible values are IGRAPH_RECIPROCI-

TY_DEFAULT and IGRAPH_RECIPROCITY_RATIO, please see their de-

scription above.

Returns:

Error code: IGRAPH_EINVAL: graph has no edges IGRAPH_ENOMEM: not enough memory for temporary data.

Time complexity: O(|V|+|E|), |V| is the number of vertices, |E| is the number of edges.

Example 13.44. File examples/simple/igraph_reciprocity.c

igraph_diversity — Structural diversity index of the vertices.

This measure was defined in Nathan Eagle, Michael Macy and Rob Claxton: Network Diversity and Economic Development, Science 328, 1029--1031, 2010.

It is simply the (normalized) Shannon entropy of the incident edges' weights. $D(i)=H(i)/\log(k[i])$, and $H(i) = -sum(p[i,j] \log(p[i,j]), j=1..k[i])$, where p[i,j]=w[i,j]/sum(w[i,l], l=1..k[i]), k[i] is the (total) degree of vertex i, and w[i,j] is the weight of the edge(s) between vertex i and j. The diversity of isolated vertices will be NaN (not-a-number), while that of vertices with a single connection will be zero.

The measure works only if the graph is undirected and has no multiple edges. If the graph has multiple edges, simplify it first using <code>igraph_simplify()</code>. If the graph is directed, convert it into an undirected graph with <code>igraph_to_undirected()</code>.

Arguments:

graph: The undirected input graph.

weights: The edge weights, in the order of the edge IDs, must have appropriate length. Weights

must be non-negative.

res: An initialized vector, the results are stored here.

vids: Vertex selector that specifies the vertices which to calculate the measure.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear.

igraph_is_mutual — Check whether some edges of a directed graph are mutual.

```
igraph_error_t igraph_is_mutual(const igraph_t *graph, igraph_vector_bool_t *re
```

An (A,B) non-loop directed edge is mutual if the graph contains the (B,A) edge too. Whether directed self-loops are considered mutual is controlled by the *loops* parameter.

An undirected graph only has mutual edges, by definition.

Edge multiplicity is not considered here, e.g. if there are two (A,B) edges and one (B,A) edge, then all three are considered to be mutual.

Arguments:

graph: The input graph.

res: Pointer to an initialized vector, the result is stored here.

es: The sequence of edges to check. Supply igraph_ess_all() to check all edges.

100ps: Boolean, whether to consider directed self-loops to be mutual.

Returns:

Error code.

Time complexity: $O(n \log(d))$, n is the number of edges supplied, d is the maximum in-degree of the vertices that are targets of the supplied edges. An upper limit of the time complexity is $O(n \log(|E|))$, |E| is the number of edges in the graph.

igraph_avg_nearest_neighbor_degree — Average neighbor degree.

Calculates the average degree of the neighbors for each vertex (*knn*), and optionally, the same quantity as a function of the vertex degree (*knnk*).

For isolated vertices *knn* is set to NaN. The same is done in *knnk* for vertex degrees that don't appear in the graph.

The weighted version computes a weighted average of the neighbor degrees as

 $k_nn_u = 1/s_u sum_v w_uv k_v$

where s_u = sum_v w_uv is the sum of the incident edge weights of vertex u, i.e. its strength. The sum runs over the neighbors v of vertex u as indicated by mode. w_uv denotes the weighted adjacency matrix and k_v is the neighbors' degree, specified by neighbor_degree_mode. This is equation (6) in the reference below.

Reference:

A. Barrat, M. Barthélemy, R. Pastor-Satorras, and A. Vespignani, The architecture of complex weighted networks, Proc. Natl. Acad. Sci. USA 101, 3747 (2004). https://dx.doi.org/10.1073/pnas.0400087101

Arguments:

graph: The input graph. It may be directed.

vids: The vertices for which the calculation is performed.

mode: The type of neighbors to consider in directed graphs.

IGRAPH_OUT considers out-neighbors, IGRAPH_IN in-neighbors

and IGRAPH_ALL ignores edge directions.

neighbor_degree_mode: The type of degree to average in directed graphs. IGRAPH_OUT

averages out-degrees, IGRAPH_IN averages in-degrees and IGRAPH_ALL ignores edge directions for the degree calculation.

vids: The vertices for which the calculation is performed.

knn: Pointer to an initialized vector, the result will be stored here. It will

be resized as needed. Supply a NULL pointer here, if you only want

to calculate knnk.

knnk: Pointer to an initialized vector, the average neighbor degree as a

function of the vertex degree is stored here. The first (zeroth) element is for degree one vertices, etc. Supply a NULL pointer here if

you don't want to calculate this.

weights: Optional edge weights. Supply a null pointer here for the non-

weighted version.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

Example 13.45. File examples/simple/igraph_knn.c

igraph_get_adjacency — The adjacency matrix of a graph.

```
#define WEIGHT_OF(eid) (weights ? VECTOR(*weights)[eid] : 1)
igraph_error_t igraph_get_adjacency(
    const igraph_t *graph, igraph_matrix_t *res, igraph_get_adjacency_t type,
    const igraph_vector_t *weights, igraph_loops_t loops
);
```

The result is an adjacency matrix. Entry i, j of the matrix contains the number of edges connecting vertex i to vertex j in the unweighted case, or the total weight of edges connecting vertex i to vertex j in the weighted case.

Arguments:

graph: Pointer to the graph to convert

res: Pointer to an initialized matrix object, it will be resized if needed.

type: Constant specifying the type of the adjacency matrix to create for undirected graphs. It

is ignored for directed graphs. Possible values:

IGRAPH_GET_ADJACENCY_UP- the upper right triangle of the matrix is used.

PER

IGRAPH GET ADJACEN- the lower left triangle of the matrix is used.

CY_LOWER

IGRAPH_GET_ADJACEN- the whole matrix is used, a symmetric matrix is

CY_BOTH returned if the graph is undirected.

weights: An optional vector containing the weight of each edge in the graph. Supply a null pointer

here to make all edges have the same weight of 1.

loops: Constant specifying how loop edges should be handled. Possible values:

IGRAPH_NO_LOOPS loop edges are ignored and the diagonal of the matrix will

contain zeros only

IGRAPH_LOOPS_ONCE loop edges are counted once, i.e. a vertex with a single

unweighted loop edge will have 1 in the corresponding

diagonal entry

IGRAPH_LOOPS_TWICE loop edges are counted twice in undirected graphs, i.e. a

vertex with a single unweighted loop edge in an undirected graph will have 2 in the corresponding diagonal entry. Loop edges in directed graphs are still counted as 1. Essentially, this means that the function is counting the incident edge *stems*, which makes more sense when using

the adjacency matrix in linear algebra.

Returns:

Error code: IGRAPH_EINVAL invalid type argument.

See also:

igraph_get_adjacency_sparse() if you want a sparse matrix representation

Time complexity: O(|V||V|), |V| is the number of vertices in the graph.

igraph_get_adjacency_sparse — Returns the adjacency matrix of a graph in a sparse matrix format.

```
igraph_error_t igraph_get_adjacency_sparse(
    const igraph_t *graph, igraph_sparsemat_t *res, igraph_get_adjacency_t type
    const igraph_vector_t *weights, igraph_loops_t loops
);
```

Arguments:

graph: The input graph.

res: Pointer to an *initialized* sparse matrix. The result will be stored here. The matrix will be

resized as needed.

type: Constant specifying the type of the adjacency matrix to create for undirected graphs. It is

ignored for directed graphs. Possible values:

IGRAPH_GET_ADJACENCY_UP- the upper right triangle of the matrix is used.

PER

IGRAPH_GET_ADJACEN- the lower left triangle of the matrix is used.

CY_LOWER

IGRAPH_GET_ADJACEN- the whole matrix is used, a symmetric matrix is re-

CY_BOTH turned if the graph is undirected.

Returns:

Error code: IGRAPH_EINVAL invalid type argument.

See also:

igraph_get_adjacency(), the dense version of this function.

Time complexity: TODO.

igraph_get_stochastic — Stochastic adjacency matrix of a graph

```
igraph_error_t igraph_get_stochastic(
    const igraph_t *graph, igraph_matrix_t *res, igraph_bool_t column_wise,
    const igraph_vector_t *weights
);
```

Stochastic matrix of a graph. The stochastic matrix of a graph is its adjacency matrix, normalized row-wise or column-wise, such that the sum of each row (or column) is one.

Arguments:

graph: The input graph.

res: Pointer to an initialized matrix, the result is stored here. It will be resized as need-

ed.

column_wise: Whether to normalize column-wise.

Returns:

Error code.

Time complexity: O(|V||V|), |V| is the number of vertices in the graph.

See also:

igraph_get_stochastic_sparse(), the sparse version of this function.

igraph_get_stochastic_sparse — The stochastic adjacency matrix of a graph.

```
igraph_error_t igraph_get_stochastic_sparse(
    const igraph_t *graph, igraph_sparsemat_t *res, igraph_bool_t column_wise,
    const igraph_vector_t *weights
);
```

Stochastic matrix of a graph. The stochastic matrix of a graph is its adjacency matrix, normalized row-wise or column-wise, such that the sum of each row (or column) is one.

Arguments:

graph: The input graph.

res: Pointer to an *initialized* sparse matrix, the result is stored here. The matrix will

be resized as needed.

column_wise: Whether to normalize column-wise.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

See also:

igraph_get_stochastic(), the dense version of this function.

igraph_get_edgelist — The list of edges in a graph.

```
igraph_error_t igraph_get_edgelist(const igraph_t *graph, igraph_vector_int_t *
```

The order of the edges is given by the edge IDs.

Arguments:

graph: Pointer to the graph object

res: Pointer to an initialized vector object, it will be resized.

bycol: Logical, if true, the edges will be returned columnwise, e.g. the first edge is res[0]-

>res[|E|], the second is res[1]->res[|E|+1], etc.

Returns:

Error code.

See also:

igraph_edges() to return the result only for some edge IDs.

Time complexity: O(|E|), the number of edges in the graph.

igraph_is_acyclic — Checks whether a graph is acyclic or not.

```
igraph_error_t igraph_is_acyclic(const igraph_t *graph, igraph_bool_t *res);
```

This function checks whether a graph is acyclic or not.

Arguments:

graph: The input graph.

res: Pointer to a boolean constant, the result is stored here.

Returns:

Error code.

Time complexity: O(|V|+|E|), where |V| and |E| are the number of vertices and edges in the original input graph.

Deprecated functions

igraph_shortest_paths — Length of the shortest paths between vertices.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_distances() instead.

igraph_shortest_paths_dijkstra — Weighted shortest path lengths between vertices (deprecated).

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_distances_dijkstra() instead.

igraph_shortest_paths_bellman_ford — Weighted shortest path lengths between vertices, allowing negative weights (deprecated).

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_distances_bellman_ford() instead.

igraph_shortest_paths_johnson — Weighted shortest path lengths between vertices, using Johnson's algorithm (deprecated).

```
const igraph_vs_t to,
const igraph_vector_t *weights);
```

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_distances_johnson() instead.

igraph_get_stochastic_sparsemat — Stochastic adjacency matrix of a graph (deprecated).

This function is deprecated in favour of <code>igraph_get_stochastic_sparse()</code>, but does not work in an identical way. This function takes an <code>uninitialized igraph_sparsemat_t</code> while <code>igraph_get_stochastic_sparse()</code> takes an already initialized one.

Arguments:

graph: The input graph.

res: Pointer to an *uninitialized* sparse matrix, the result is stored here. The matrix will

be resized as needed.

column_wise: Whether to normalize column-wise. For undirected graphs this argument does not

have any effect.

Returns:

Error code.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_get_stochastic_sparse() instead.

igraph_get_sparsemat — Converts an igraph graph to a sparse matrix (deprecated).

```
igraph_error_t igraph_get_sparsemat(const igraph_t *graph, igraph_sparsemat_t *
```

If the graph is undirected, then a symmetric matrix is created.

This function is deprecated in favour of <code>igraph_get_adjacency_sparse()</code>, but does not work in an identical way. This function takes an <code>uninitialized igraph_sparsemat_t</code> while <code>igraph_get_adjacency_sparse()</code> takes an already initialized one.

Arguments:

graph: The input graph.

res: Pointer to an *uninitialized* sparse matrix. The result will be stored here.

Returns:

Error code.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_get_adjacency_sparse() instead.

igraph_laplacian — Returns the Laplacian matrix of a graph (deprecated).

```
igraph_error_t igraph_laplacian(
    const igraph_t *graph, igraph_matrix_t *res, igraph_sparsemat_t *sparseres,
    igraph_bool_t normalized, const igraph_vector_t *weights
);
```

This function produces the Laplacian matrix of a graph in either dense or sparse format. When nor-malized is set to true, the type of normalization used depends on the directnedness of the graph: symmetric normalization is used for undirected graphs and left stochastic normalization for directed graphs.

Arguments:

graph: Pointer to the graph to convert.

res: Pointer to an initialized matrix object or NULL. The dense matrix result will be

stored here.

sparseres: Pointer to an initialized sparse matrix object or NULL. The sparse matrix result will

be stored here.

mode: Controls whether to use out- or in-degrees in directed graphs. If set to

IGRAPH_ALL, edge directions will be ignored.

normalized: Boolean, whether to normalize the result.

weights: An optional vector containing non-negative edge weights, to calculate the weighted

Laplacian matrix. Set it to a null pointer to calculate the unweighted Laplacian.

Returns:

Error code.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_get_laplacian() instead.

Chapter 14. Graph cycles

Eulerian cycles and paths

These functions calculate whether an Eulerian path or cycle exists and if so, can find them.

igraph_is_eulerian — Checks whether an Eulerian path or cycle exists.

```
igraph_error_t igraph_is_eulerian(const igraph_t *graph, igraph_bool_t *has_pat
```

An Eulerian path traverses each edge of the graph precisely once. A closed Eulerian path is referred to as an Eulerian cycle.

Arguments:

graph: The graph object.

has_path: Pointer to a Boolean, will be set to true if an Eulerian path exists.

has_cycle: Pointer to a Boolean, will be set to true if an Eulerian cycle exists.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

igraph_eulerian_cycle — Finds an Eulerian cycle.

Finds an Eulerian cycle, if it exists. An Eulerian cycle is a closed path that traverses each edge precisely once.

This function uses Hierholzer's algorithm.

Arguments:

graph: The graph object.

edge_res: Pointer to an initialised vector. The indices of edges belonging to the cycle will be

stored here. May be NULL if it is not needed by the caller.

vertex_res: Pointer to an initialised vector. The indices of vertices belonging to the cycle will

be stored here. May be NULL if it is not needed by the caller.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID graph does not have an Eulerian cycle.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

igraph_eulerian_path — Finds an Eulerian path.

Finds an Eulerian path, if it exists. An Eulerian path traverses each edge precisely once.

This function uses Hierholzer's algorithm.

Arguments:

graph: The graph object.

edge_res: Pointer to an initialised vector. The indices of edges belonging to the path will be

stored here. May be NULL if it is not needed by the caller.

vertex_res: Pointer to an initialised vector. The indices of vertices belonging to the path will

be stored here. May be NULL if it is not needed by the caller.

Returns:

Error code:

IGRAPH_ENOMEM not enough memory for temporary data.

IGRAPH_EINVVID graph does not have an Eulerian path.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

Chapter 15. Graph visitors

Breadth-first search

igraph_bfs — Breadth-first search

A simple breadth-first search, with a lot of different results and the possibility to call a callback whenever a vertex is visited. It is allowed to supply null pointers as the output arguments the user is not interested in, in this case they will be ignored.

If not all vertices can be reached from the supplied root vertex, then additional root vertices will be used, in the order of their vertex IDs.

Consider using igraph_bfs_simple instead if you set most of the output arguments provided by this function to a null pointer.

Arguments:

graph: The input graph.

root: The id of the root vertex. It is ignored if the roots argument is not a null pointer.

roots: Pointer to an initialized vector, or a null pointer. If not a null pointer, then it is a

vector containing root vertices to start the BFS from. The vertices are considered in the order they appear. If a root vertex was already found while searching from

another one, then no search is conducted from it.

mode: For directed graphs, it defines which edges to follow. IGRAPH_OUT means

following the direction of the edges, IGRAPH_IN means the opposite, and IGRAPH_ALL ignores the direction of the edges. This parameter is ignored for

undirected graphs.

unreachable: Logical scalar, whether the search should visit the vertices that are unreachable

from the given root node(s). If true, then additional searches are performed until

all vertices are visited.

restricted: If not a null pointer, then it must be a pointer to a vector containing vertex IDs.

The BFS is carried out only on these vertices.

order: If not null pointer, then the vertex IDs of the graph are stored here, in the same

order as they were visited.

rank: If not a null pointer, then the rank of each vertex is stored here.

parents: If not a null pointer, then the id of the parent of each vertex is stored here. When a

vertex was not visited during the traversal, -2 will be stored as the ID of its parent.

When a vertex was visited during the traversal and it was one of the roots of the search trees, -1 will be stored as the ID of its parent.

pred: If not a null pointer, then the id of vertex that was visited before the current one

is stored here. If there is no such vertex (the current vertex is the root of a search tree), then -1 is stored as the predecessor of the vertex. If the vertex was not visited

at all, then -2 is stored for the predecessor of the vertex.

succ: If not a null pointer, then the id of the vertex that was visited after the current one

is stored here. If there is no such vertex (the current one is the last in a search tree), then -1 is stored as the successor of the vertex. If the vertex was not visited

at all, then -2 is stored for the successor of the vertex.

dist: If not a null pointer, then the distance from the root of the current search tree is

stored here for each vertex. If a vertex was not reached during the traversal, its

distance will be -1 in this vector.

callback: If not null, then it should be a pointer to a function of type igraph_bfshan-

dler_t. This function will be called, whenever a new vertex is visited.

extra: Extra argument to pass to the callback function.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

Example 15.1. File examples/simple/igraph_bfs.c

Example 15.2. File examples/simple/igraph_bfs_callback.c

igraph_bfs_simple — Breadth-first search, single-source version

```
igraph_error_t igraph_bfs_simple(
    const igraph_t *graph, igraph_integer_t root, igraph_neimode_t mode,
    igraph_vector_int_t *order, igraph_vector_int_t *layers,
    igraph_vector_int_t *parents
);
```

An alternative breadth-first search implementation to cater for the simpler use-cases when only a single breadth-first search has to be conducted from a source node and most of the output arguments from igraph_bfs are not needed. It is allowed to supply null pointers as the output arguments the user is not interested in, in this case they will be ignored.

Arguments:

graph: The input graph.

root: The id of the root vertex.

mode: For directed graphs, it defines which edges to follow. IGRAPH_OUT means following

the direction of the edges, IGRAPH IN means the opposite, and IGRAPH ALL ignores

the direction of the edges. This parameter is ignored for undirected graphs.

order: If not a null pointer, then an initialized vector must be passed here. The IDs of the

vertices visited during the traversal will be stored here, in the same order as they were

visited.

layers: If not a null pointer, then an initialized vector must be passed here. The i-th element

of the vector will contain the index into order where the vertices that are at distance i from the root are stored. In other words, if you are interested in the vertices that are at distance i from the root, you need to look in the order vector from layers[i] to

layers[i+1].

parents: If not a null pointer, then an initialized vector must be passed here. The vector will be

resized so its length is equal to the number of nodes, and it will contain the index of the parent node for each *visited* node. The values in the vector are set to -2 for vertices that

were not visited, and -1 for the root vertex.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

Example 15.3. File examples/simple/igraph_bfs_simple.c

igraph_bfshandler_t — Callback type for BFS function.

igraph_bfs() is able to call a callback function, whenever a new vertex is found, while doing the breadth-first search. This callback function must be of type igraph_bfshandler_t. It has the following arguments:

Arguments:

graph: The graph that that algorithm is working on. Of course this must not be modified.

via: The id of the vertex just found by the breadth-first search.

pred: The id of the previous vertex visited. It is -1 if there is no previous vertex, because the

current vertex is the root is a search tree.

succ: The id of the next vertex that will be visited. It is -1 if there is no next vertex, because the

current vertex is the last one in a search tree.

rank: The rank of the current vertex, it starts with zero.

dist: The distance (number of hops) of the current vertex from the root of the current search tree.

extra: The extra argument that was passed to igraph_bfs().

Returns:

IGRAPH_SUCCESS if the BFS should continue, IGRAPH_STOP if the BFS should stop and return to the caller normally. Any other value is treated as an igraph error code, terminating the search and returning to the caller with the same error code. If a BFS is is terminated prematurely, then all elements of the result vectors that were not yet calculated at the point of the termination contain NaN.

See also:

igraph_bfs()

Depth-first search

igraph_dfs — Depth-first search

A simple depth-first search, with the possibility to call a callback whenever a vertex is discovered and/ or whenever a subtree is finished. It is allowed to supply null pointers as the output arguments the user is not interested in, in this case they will be ignored.

If not all vertices can be reached from the supplied root vertex, then additional root vertices will be used, in the order of their vertex IDs.

Arguments:

graph: The input graph.

root: The id of the root vertex.

mode: For directed graphs, it defines which edges to follow. IGRAPH_OUT means

following the direction of the edges, IGRAPH_IN means the opposite, and IGRAPH_ALL ignores the direction of the edges. This parameter is ignored for

undirected graphs.

unreachable: Logical scalar, whether the search should visit the vertices that are unreachable

from the given root node(s). If true, then additional searches are performed until

all vertices are visited.

order: If not null pointer, then the vertex IDs of the graph are stored here, in the same

order as they were discovered. The tail of the vector will be padded with -1 to ensure that the length of the vector is the same as the number of vertices, even

if some vertices were not visited during the traversal.

order_out: If not a null pointer, then the vertex IDs of the graphs are stored here, in the order

of the completion of their subtree. The tail of the vector will be padded with -1 to ensure that the length of the vector is the same as the number of vertices, even

if some vertices were not visited during the traversal.

parents: If not a null pointer, then the id of the parent of each vertex is stored here. -1

will be stored for the root of the search tree; -2 will be stored for vertices that

were not visited.

dist: If not a null pointer, then the distance from the root of the current search tree is

stored here. -1 will be stored for vertices that were not visited.

in_callback: If not null, then it should be a pointer to a function of type igraph_dfshan-

dler_t. This function will be called, whenever a new vertex is discovered.

out_callback: If not null, then it should be a pointer to a function of type igraph_dfshan-

dler_t. This function will be called, whenever the subtree of a vertex is com-

pleted.

extra: Extra argument to pass to the callback function(s).

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

igraph_dfshandler_t — Callback type for the DFS function.

igraph_dfs() is able to call a callback function, whenever a new vertex is discovered, and/or whenever a subtree is completed. These callbacks must be of type igraph_dfshandler_t. They have the following arguments:

Arguments:

graph: The graph that that algorithm is working on. Of course this must not be modified.

vid: The id of the vertex just found by the depth-first search.

dist: The distance (number of hops) of the current vertex from the root of the current search tree.

extra: The extra argument that was passed to igraph_dfs().

Returns:

IGRAPH_SUCCESS if the DFS should continue, IGRAPH_STOP if the DFS should stop and return to the caller normally. Any other value is treated as an igraph error code, terminating the search and returning to the caller with the same error code. If a BFS is is terminated prematurely, then all elements of the result vectors that were not yet calculated at the point of the termination contain NaN.

See also:

```
igraph_dfs()
```

Random walks

igraph_random_walk — Performs a random walk on a graph.

Performs a random walk with a given length on a graph, from the given start vertex. Edge directions are (potentially) considered, depending on the *mode* argument.

Arguments:

graph: The input graph, it can be directed or undirected. Multiple edges are respected, so are

loop edges.

weights: A vector of non-negative edge weights. It is assumed that at least one strictly posi-

tive weight is found among the outgoing edges of each vertex. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If it is NULL,

all edges are considered to have equal weight.

vertices: An allocated vector, the result is stored here as a list of vertex IDs. It will be resized

as needed. It includes the vertex IDs of starting and ending vertices. Length of the

vertices vector: steps + 1

edges: An initialized vector, the indices of traversed edges are stored here. It will be resized

as needed. Length of the edges vector: steps

start: The start vertex for the walk.

steps: The number of steps to take. If the random walk gets stuck, then the stuck argument

specifies what happens. steps is the number of edges to traverse during the walk.

mode: How to walk along the edges in directed graphs. IGRAPH_OUT means following edge

directions, IGRAPH_IN means going opposite the edge directions, IGRAPH_ALL

means ignoring edge directions. This argument is ignored for undirected graphs.

stuck: What to do if the random walk gets stuck. IGRAPH_RANDOM_WALK_STUCK_RE-

TURN means that the function returns with a shorter walk; IGRAPH_RAN-DOM_WALK_STUCK_ERROR means that an IGRAPH_ERWSTUCK error is reported. In both cases, vertices and edges are truncated to contain the actual interrupted

walk.

Returns:

Error code: IGRAPH_ERWSTUCK if the walk got stuck.

Time complexity: O(l + d) for unweighted graphs and O(l * log(k) + d) for weighted graphs, where 1 is the length of the walk, d is the total degree of the visited nodes and k is the average degree of vertices of the given graph.

Deprecated functions

igraph_random_edge_walk — Performs a random walk on a graph and returns the traversed edges.

Performs a random walk with a given length on a graph, from the given start vertex. Edge directions are (potentially) considered, depending on the *mode* argument.

Arguments:

graph: The input graph, it can be directed or undirected. Multiple edges are respected, so are

loop edges.

weights: A vector of non-negative edge weights. It is assumed that at least one strictly posi-

tive weight is found among the outgoing edges of each vertex. Additionally, no edge weight may be NaN. If either case does not hold, an error is returned. If it is a NULL

pointer, all edges are considered to have equal weight.

edgewalk: An initialized vector; the indices of traversed edges are stored here. It will be resized

as needed.

start: The start vertex for the walk.

steps: The number of steps to take. If the random walk gets stuck, then the stuck argument

specifies what happens.

mode: How to walk along the edges in directed graphs. IGRAPH_OUT means following edge

directions, IGRAPH_IN means going opposite the edge directions, IGRAPH_ALL means ignoring edge directions. This argument is ignored for undirected graphs.

stuck: What to do if the random walk gets stuck. IGRAPH_RANDOM_WALK_STUCK_RE-

TURN means that the function returns with a shorter walk; IGRAPH_RANDOM_WALK_STUCK_ERROR means that an IGRAPH_ERWSTUCK error is reported.

In both cases, *edgewalk* is truncated to contain the actual interrupted walk.

Returns:

Error code.

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_random_walk() instead.

Chapter 16. Cliques and independent vertex sets

These functions calculate various graph properties related to cliques and independent vertex sets.

Cliques

igraph_cliques — Finds all or some cliques in a graph.

Cliques are fully connected subgraphs of a graph.

If you are only interested in the size of the largest clique in the graph, use igraph_clique_number() instead.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, http://users.aalto.fi/~pat/cliquer.html

Arguments:

graph: The input graph.

res: Pointer to a list of integer vectors, the result will be stored here. The pointer vector

will be resized if needed.

min_size: Integer specifying the minimum size of the cliques to be returned. If negative or zero,

no lower bound will be used.

max_size: Integer specifying the maximum size of the cliques to be returned. If negative or zero,

no upper bound will be used.

Returns:

Error code.

See also:

```
igraph_largest_cliques() and igraph_clique_number().
```

Time complexity: Exponential

Example 16.1. File examples/simple/igraph_cliques.c

igraph_clique_size_hist — Counts cliques of each size in the graph.

Cliques are fully connected subgraphs of a graph.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, http://users.aalto.fi/~pat/cliquer.html

Arguments:

graph: The input graph.

hist: Pointer to an initialized vector. The result will be stored here. The first element will

store the number of size-1 cliques, the second element the number of size-2 cliques,

etc. For cliques smaller than min_size, zero counts will be returned.

min_size: Integer specifying the minimum size of the cliques to be returned. If negative or zero,

no lower bound will be used.

max_size: Integer specifying the maximum size of the cliques to be returned. If negative or zero,

no upper bound will be used.

Returns:

Error code.

See also:

```
igraph_cliques() and igraph_cliques_callback()
```

Time complexity: Exponential

igraph_cliques_callback — Calls a function for each clique in the graph.

Cliques are fully connected subgraphs of a graph. This function enumerates all cliques within the given size range and calls <code>cliquehandler_fn</code> for each of them. The cliques are passed to the callback function as a pointer to an <code>igraph_vector_int_t</code>. Destroying and freeing this vector is left up to the user. Use <code>igraph_vector_int_destroy()</code> to destroy it first, then free it using <code>igraph_free()</code>.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, http://users.aalto.fi/~pat/cliquer.html

Arguments:

graph: The input graph.

min_size: Integer specifying the minimum size of the cliques to be returned. If neg-

ative or zero, no lower bound will be used.

max_size: Integer specifying the maximum size of the cliques to be returned. If neg-

ative or zero, no upper bound will be used.

cliquehandler_fn: Callback function to be called for each clique. See also

igraph_clique_handler_t.

arg: Extra argument to supply to cliquehandler_fn.

Returns:

Error code.

See also:

igraph_cliques()

Time complexity: Exponential

igraph_clique_handler_t — Type of clique handler functions.

typedef igraph_error_t igraph_clique_handler_t(const igraph_vector_int_t *cliqu

Callback type, called when a clique was found. See the details at the documentation of igraph_cliques_callback().

Arguments:

clique: The current clique. The clique is owned by the clique search routine. You do not need to

destroy or free it if you do not want to store it; however, if you want to hold on to it for a longer period of time, you need to make a copy of it on your own and store the copy itself.

arg: This extra argument was passed to igraph_cliques_callback() when it was

called.

Returns:

Error code; IGRAPH_SUCCESS to continue the search or IGRAPH_STOP to stop the search without signaling an error.

igraph_largest_cliques — Finds the largest clique(s) in a graph.

igraph_error_t igraph_largest_cliques(const igraph_t *graph, igraph_vector_int_

A clique is largest (quite intuitively) if there is no other clique in the graph which contains more vertices.

Note that this is not necessarily the same as a maximal clique, i.e. the largest cliques are always maximal but a maximal clique is not always largest.

The current implementation of this function searches for maximal cliques using igraph_maximal_cliques_callback() and drops those that are not the largest.

The implementation of this function changed between igraph 0.5 and 0.6, so the order of the cliques and the order of vertices within the cliques will almost surely be different between these two versions.

Arguments:

graph: The input graph.

res: Pointer to a list of integer vectors, the result will be stored here. The pointer vector will

be resized if needed.

Returns:

Error code.

See also:

```
igraph_cliques(), igraph_maximal_cliques()
```

Time complexity: $O(3^{(|V|/3)})$ worst case.

igraph_maximal_cliques — Finds all maximal cliques in a graph.

```
igraph_error_t igraph_maximal_cliques(
    const igraph_t *graph, igraph_vector_int_list_t *res,
    igraph_integer_t min_size, igraph_integer_t max_size
);
```

A maximal clique is a clique which can't be extended any more by adding a new vertex to it.

If you are only interested in the size of the largest clique in the graph, use igraph_clique_number() instead.

The current implementation uses a modified Bron-Kerbosch algorithm to find the maximal cliques, see: David Eppstein, Maarten Löffler, Darren Strash: Listing All Maximal Cliques in Sparse Graphs in Near-Optimal Time. Algorithms and Computation, Lecture Notes in Computer Science Volume 6506, 2010, pp 403-414.

The implementation of this function changed between igraph 0.5 and 0.6 and also between 0.6 and 0.7, so the order of the cliques and the order of vertices within the cliques will almost surely be different between these three versions.

Arguments:

graph: The input graph.

res: Pointer to a pointer vector, the result will be stored here, i.e. res will contain pointers

to igraph_vector_int_t objects which contain the indices of vertices involved in a clique. The pointer vector will be resized if needed but note that the objects in the pointer vector will not be freed. Note that vertices of a clique may be returned in

arbitrary order.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no

lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no

upper bound will be used.

Returns:

Error code.

See also:

```
igraph_maximal_independent_vertex_sets(),igraph_clique_number()
```

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

Example 16.2. File examples/simple/igraph_maximal_cliques.c

igraph_maximal_cliques_count — Count the number of maximal cliques in a graph

The current implementation uses a modified Bron-Kerbosch algorithm to find the maximal cliques, see: David Eppstein, Maarten Löffler, Darren Strash: Listing All Maximal Cliques in Sparse Graphs in Near-Optimal Time. Algorithms and Computation, Lecture Notes in Computer Science Volume 6506, 2010, pp 403-414.

Arguments:

graph: The input graph.

res: Pointer to an igraph_integer_t; the number of maximal cliques will be stored

here.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no

lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no

upper bound will be used.

Returns:

Error code.

See also:

```
igraph_maximal_cliques().
```

Time complexity: $O(d(n-d)3^{d/3})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

Example 16.3. File examples/simple/igraph_maximal_cliques.c

igraph_maximal_cliques_file — Find maximal cliques and write them to a file.

This function enumerates all maximal cliques and writes them to file.

Edge directions are ignored.

Arguments:

graph: The input graph.

outfile: Pointer to the output file, it should be writable.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no

lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no

upper bound will be used.

Returns:

Error code.

See also:

```
igraph_maximal_cliques().
```

Time complexity: $O(d(n-d)3^{d/3})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.*

igraph_maximal_cliques_subset — Maximal cliques for a subset of initial vertices

```
igraph_error_t igraph_maximal_cliques_subset(
    const igraph_t *graph, const igraph_vector_int_t *subset,
    igraph_vector_int_list_t *res, igraph_integer_t *no,
    FILE *outfile, igraph_integer_t min_size, igraph_integer_t max_size
);
```

This function enumerates all maximal cliques for a subset of initial vertices and writes them to file.

Edge directions are ignored.

Arguments:

graph: The input graph.

subset: Pointer to an igraph_vector_int_t containing the subset of initial vertices

res: Pointer to a list of integer vectors; the cliques will be stored here

no: Pointer to an igraph_integer_t; the number of maximal cliques will be stored

here.

outfile: Pointer to an output file or NULL. When not NULL, the file should be writable.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no

lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no

upper bound will be used.

Returns:

Error code.

See also:

```
igraph_maximal_cliques().
```

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

igraph_maximal_cliques_hist — Counts the number of maximal cliques of each size in a graph.

This function counts how many maximal cliques of each size are present in the graph. Size-1 maximal cliques are simply isolated vertices.

Edge directions are ignored.

Arguments:

graph: The input graph.

hist: Pointer to an initialized vector. The result will be stored here. The first element will

store the number of size-1 maximal cliques, the second element the number of size-2 maximal cliques, etc. For cliques smaller than min_size , zero counts will be re-

turned.

min_size: Integer giving the minimum size of the cliques to be returned. If negative or zero, no

lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative or zero, no upper bound will be used.

Returns:

Error code.

See also:

```
igraph_maximal_cliques().
```

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

igraph_maximal_cliques_callback — Finds maximal cliques in a graph and calls a function for each one.

This function enumerates all maximal cliques within the given size range and calls <code>cliquehan-dler_fn</code> for each of them. The cliques are passed to the callback function as a pointer to an <code>igraph_vector_int_t</code>. The vector is owned by the maximal clique search routine so users are expected to make a copy of the vector usign <code>igraph_vector_int_init_copy()</code> if they want to hold on to it.

Edge directions are ignored.

Arguments:

graph: The input graph.

cliquehandler_fn: Callback function to be called for each clique. See also

igraph_clique_handler_t.

arg: Extra argument to supply to cliquehandler_fn.

min_size: Integer giving the minimum size of the cliques to be returned. If negative

or zero, no lower bound will be used.

max_size: Integer giving the maximum size of the cliques to be returned. If negative

or zero, no upper bound will be used.

Returns:

Error code.

See also:

```
igraph_maximal_cliques().
```

Time complexity: $O(d(n-d)3^{(d/3)})$ worst case, d is the degeneracy of the graph, this is typically small for sparse graphs.

igraph_clique_number — Finds the clique number of the graph.

```
igraph_error_t igraph_clique_number(const igraph_t *graph, igraph_integer_t *no
```

The clique number of a graph is the size of the largest clique.

The current implementation of this function searches for maximal cliques using igraph_maximal_cliques_callback() and keeps track of the size of the largest clique that was found.

Arguments:

graph: The input graph.

no: The clique number will be returned to the igraph_integer_t pointed by this variable.

Returns:

Error code.

See also:

```
igraph_cliques(), igraph_largest_cliques().
```

Time complexity: $O(3^{(|V|/3)})$ worst case.

Weighted cliques

igraph_weighted_cliques — Finds all cliques in a given weight range in a vertex weighted graph.

Cliques are fully connected subgraphs of a graph. The weight of a clique is the sum of the weights of individual vertices within the clique.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, http://users.aalto.fi/~pat/cliquer.html Only positive integer vertex weights are supported.

Arguments:

graph: The input graph.

vertex_weights: A vector of vertex weights. The current implementation will truncate all

weights to their integer parts. You may pass NULL here to make each vertex

have a weight of 1.

res: Pointer to a list of integer vectors, the result will be stored here. The pointer

vector will be resized if needed.

min_weight: Integer specifying the minimum weight of the cliques to be returned. If neg-

ative or zero, no lower bound will be used.

max_weight: Integer specifying the maximum weight of the cliques to be returned. If neg-

ative or zero, no upper bound will be used.

maximal: If true, only maximal cliques will be returned

Returns:

Error code.

See also:

```
igraph_cliques(), igraph_maximal_cliques()
```

Time complexity: Exponential

igraph_largest_weighted_cliques — Finds the largest weight clique(s) in a graph.

Finds the clique(s) having the largest weight in the graph.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, http://users.aalto.fi/~pat/cliquer.html Only positive integer vertex weights are supported.

Arguments:

graph: The input graph.

vertex_weights: A vector of vertex weights. The current implementation will truncate all

weights to their integer parts. You may pass NULL here to make each vertex

have a weight of 1.

res: Pointer to a list of integer vectors, the result will be stored here. The pointer

vector will be resized if needed.

Returns:

Error code.

See also:

Time complexity: TODO

igraph_weighted_clique_number — Finds the weight of the largest weight clique in the graph.

The current implementation of this function uses version 1.21 of the Cliquer library by Sampo Niskanen and Patric R. J. Östergård, http://users.aalto.fi/~pat/cliquer.html Only positive integer vertex weights are supported.

Arguments:

graph: The input graph.

vertex_weights: A vector of vertex weights. The current implementation will truncate all

weights to their integer parts. You may pass NULL here to make each vertex

have a weight of 1.

res: The largest weight will be returned to the igraph_real_t pointed to by

this variable.

Returns:

Error code.

See also:

Time complexity: TODO

Independent vertex sets

igraph_independent_vertex_sets — Finds all independent vertex sets in a graph.

A vertex set is considered independent if there are no edges between them.

If you are interested in the size of the largest independent vertex set, use igraph_independence number() instead.

The current implementation was ported to igraph from the Very Nauty Graph Library by Keith Briggs and uses the algorithm from the paper S. Tsukiyama, M. Ide, H. Ariyoshi and I. Shirawaka. A new algorithm for generating all the maximal independent sets. SIAM J Computing, 6:505--517, 1977.

Arguments:

graph: The input graph.

res: Pointer to a list of integer vectors, the result will be stored here. The pointer vector

will be resized if needed.

min_size: Integer specifying the minimum size of the sets to be returned. If negative or zero, no

lower bound will be used.

max_size: Integer specifying the maximum size of the sets to be returned. If negative or zero,

no upper bound will be used.

Returns:

Error code.

See also:

```
igraph_largest_independent_vertex_sets(), igraph_independence_num-
ber().
```

Time complexity: TODO

Example 16.4. File examples/simple/igraph_independent_sets.c

igraph_largest_independent_vertex_sets — Finds the largest independent vertex set(s) in a graph.

An independent vertex set is largest if there is no other independent vertex set with more vertices in the graph.

The current implementation was ported to igraph from the Very Nauty Graph Library by Keith Briggs and uses the algorithm from the paper S. Tsukiyama, M. Ide, H. Ariyoshi and I. Shirawaka. A new algorithm for generating all the maximal independent sets. SIAM J Computing, 6:505--517, 1977.

Arguments:

graph: The input graph.

res: Pointer to a list of integer vectors, the result will be stored here. The pointer vector will

be resized if needed.

Returns:

Error code.

See also:

```
igraph_independent_vertex_sets(), igraph_maximal_independent_ver-
tex sets().
```

Time complexity: TODO

igraph_maximal_independent_vertex_sets — Finds all maximal independent vertex sets of a graph.

A maximal independent vertex set is an independent vertex set which can't be extended any more by adding a new vertex to it.

The algorithm used here is based on the following paper: S. Tsukiyama, M. Ide, H. Ariyoshi and I. Shirawaka. A new algorithm for generating all the maximal independent sets. SIAM J Computing, 6:505--517, 1977.

The implementation was originally written by Kevin O'Neill and modified by K M Briggs in the Very Nauty Graph Library. I simply re-wrote it to use igraph's data structures.

If you are interested in the size of the largest independent vertex set, use igraph_independence_number() instead.

Arguments:

graph: The input graph.

res: Pointer to a list of integer vectors, the result will be stored here. The pointer vector will

be resized if needed.

Returns:

Error code.

See also:

```
igraph_maximal_cliques(), igraph_independence_number()
```

Time complexity: TODO.

igraph_independence_number — Finds the independence number of the graph.

```
igraph_error_t igraph_independence_number(const igraph_t *graph, igraph_integer)
```

The independence number of a graph is the cardinality of the largest independent vertex set.

The current implementation was ported to igraph from the Very Nauty Graph Library by Keith Briggs and uses the algorithm from the paper S. Tsukiyama, M. Ide, H. Ariyoshi and I. Shirawaka. A new algorithm for generating all the maximal independent sets. SIAM J Computing, 6:505--517, 1977.

Arguments:

graph: The input graph.

no: The independence number will be returned to the <code>igraph_integer_t</code> pointed by this variable.

Returns:

Error code.

See also:

```
{\tt igraph\_independent\_vertex\_sets()}.
```

Time complexity: TODO.

Chapter 17. Graph isomorphism

The simple interface

igraph provides four set of functions to deal with graph isomorphism problems.

The igraph_isomorphic() and igraph_subisomorphic() functions make up the first set (in addition with the igraph_permute_vertices() function). These functions choose the algorithm which is best for the supplied input graph. (The choice is not very sophisticated though, see their documentation for details.)

The VF2 graph (and subgraph) isomorphism algorithm is implemented in igraph, these functions are the second set. See igraph_isomorphic_vf2() and igraph_subisomorphic_vf2() for starters.

Functions for the Bliss algorithm constitute the third set, see igraph_isomorphic_bliss().

Finally, the isomorphism classes of all directed graphs with three and four vertices and all undirected graphs with 3-6 vertices are precomputed and stored in igraph, so for these small graphs there is a separate fast path in the code that does not use more complex, generic isomorphism algorithms.

igraph_isomorphic — Are two graphs isomorphic?

In simple terms, two graphs are isomorphic if they become indistinguishable from each other once their vertex labels are removed (rendering the vertices within each graph indistiguishable). More precisely, two graphs are isomorphic if there is a one-to-one mapping from the vertices of the first one to the vertices of the second such that it transforms the edge set of the first graph into the edge set of the second. This mapping is called an *isomorphism*.

Currently, this function supports simple graphs and graphs with self-loops, but does not support multigraphs.

This function decides which graph isomorphism algorithm to be used based on the input graphs. Right now it does the following:

- 1. If one graph is directed and the other undirected then an error is triggered.
- 2. If one of the graphs has multi-edges then an error is triggered.
- 3. If the two graphs does not have the same number of vertices and edges it returns with false.
- 4. Otherwise, if the igraph_isoclass() function supports both graphs (which is true for directed graphs with 3 and 4 vertices, and undirected graphs with 3-6 vertices), an O(1) algorithm is used with precomputed data.
- 5. Otherwise Bliss is used, see igraph isomorphic bliss().

Please call the VF2 and Bliss functions directly if you need something more sophisticated, e.g. you need the isomorphic mapping.

Arguments:

graph1: The first graph.

graph2: The second graph.

iso: Pointer to a logical variable, will be set to true if the two graphs are isomorphic, and

false otherwise.

Returns:

Error code.

See also:

```
igraph_isoclass(), igraph_isoclass_subgraph(), igraph_isoclass_cre-
ate().
```

Time complexity: exponential.

igraph_subisomorphic — Decide subgraph isomorphism.

Check whether *graph2* is isomorphic to a subgraph of *graph1*. Currently this function just calls igraph_subisomorphic_vf2() for all graphs.

Currently this function does not support non-simple graphs.

Arguments:

graph1: The first input graph, may be directed or undirected. This is supposed to be the bigger

graph.

graph2: The second input graph, it must have the same directedness as graph2, or an error is

triggered. This is supposed to be the smaller graph.

iso: Pointer to a boolean, the result is stored here.

Returns:

Error code.

Time complexity: exponential.

The BLISS algorithm

Bliss is a successor of the famous NAUTY algorithm and implementation. While using the same ideas in general, with better heuristics and data structures Bliss outperforms NAUTY on most graphs.

Bliss was developed and implemented by Tommi Junttila and Petteri Kaski at Helsinki University of Technology, Finland. For more information, see the Bliss homepage at https://users.aalto.fi/~tjunttil/bliss/ and the following publication:

Tommi Junttila and Petteri Kaski: "Engineering an Efficient Canonical Labeling Tool for Large and Sparse Graphs" In ALENEX 2007, pages 135–149, 2007 https://doi.org/10.1137/1.9781611972870.13

Tommi Junttila and Petteri Kaski: "Conflict Propagation and Component Recursion for Canonical Labeling" in TAPAS 2011, pages 151–162, 2011. https://doi.org/10.1007/978-3-642-19754-3_16

Bliss works with both directed graphs and undirected graphs. It supports graphs with self-loops, but not graphs with multi-edges.

Bliss version 0.75 is included in igraph.

igraph_bliss_sh_t — Splitting heuristics for Bliss.

IGRAPH_BLISS_FL provides good performance for many graphs, and is a reasonable default choice. IGRAPH_BLISS_FSM is recommended for graphs that have some combinatorial structure, and is the default of the Bliss library's command line tool.

Values:

```
IGRAPH_BLISS_F: First non-singleton cell.
IGRAPH_BLISS_FL: First largest non-singleton cell.
IGRAPH_BLISS_FS: First smallest non-singleton cell.
IGRAPH_BLISS_FM: First maximally non-trivially connected non-singleton cell.
IGRAPH_BLISS_FLM: Largest maximally non-trivially connected non-singleton cell.
IGRAPH_BLISS_FSM: Smallest maximally non-trivially connected non-singletion cell.
```

igraph_bliss_info_t — Information about a BLISS run

```
typedef struct igraph_bliss_info_t {
   unsigned long nof_nodes;
   unsigned long nof_leaf_nodes;
   unsigned long nof_bad_nodes;
   unsigned long nof_canupdates;
   unsigned long nof_generators;
   unsigned long max_level;
   char *group_size;
} igraph_bliss_info_t;
```

Some secondary information found by the BLISS algorithm is stored here. It is useful if you wany to study the internal working of the algorithm.

Values:

```
nof_nodes: The number of nodes in the search tree.

nof_leaf_nodes: The number of leaf nodes in the search tree.
```

nof_bad_nodes: Number of bad nodes.

nof_canupdates: Number of canrep updates.

nof_generators: Number of generators of the automorphism group.

max_level: Maximum level.

group_size: The size of the automorphism group of the graph, given as a string. It should

be deallocated via igraph_free() if not needed any more.

See http://www.tcs.hut.fi/Software/bliss/index.html for details about the algorithm and these parameters

igraph_canonical_permutation — Canonical permutation using Bliss.

igraph_error_t igraph_canonical_permutation(const igraph_t *graph, const igraph_
igraph_vector_int_t *labeling, igraph_bliss_sh

This function computes the vertex permutation which transforms the graph into a canonical form, using the Bliss algorithm. Two graphs have the same canonical form if and only if they are isomorphic. Use igraph_is_same_graph() to compare two canonical forms.

Arguments:

graph: The input graph. Multiple edges between the same nodes are not supported and will

cause an incorrect result to be returned.

colors: An optional vertex color vector for the graph. Supply a null pointer is the graph is

not colored.

labeling: Pointer to a vector, the result is stored here. The permutation takes vertex 0 to the first

element of the vector, vertex 1 to the second, etc. The vector will be resized as needed.

sh: The splitting heuristics to be used in Bliss. See igraph_bliss_sh_t.

info: If not NULL then information on Bliss internals is stored here. The memory used by

this structure must to be freed when no longer needed, see igraph_bliss_in-

fo_t.

Returns:

Error code.

See also:

```
igraph_is_same_graph()
```

Time complexity: exponential, in practice it is fast for many graphs.

igraph_isomorphic_bliss — Graph isomorphism via Bliss.

This function uses the Bliss graph isomorphism algorithm, a successor of the famous NAUTY algorithm and implementation. Bliss is open source and licensed according to the GNU LGPL. See https://users.aalto.fi/~tjunttil/bliss/ for details. Currently the 0.75 version of Bliss is included in igraph.

Isomorphism testing is implemented by producing the canonical form of both graphs using igraph_canonical_permutation() and comparing them.

Arguments:

colors2:

The first input graph. Multiple edges between the same nodes are not supported and will cause an incorrect result to be returned.

graph2: The second input graph. Multiple edges between the same nodes are not supported and will cause an incorrect result to be returned.

colors1: An optional vertex color vector for the first graph. Supply a null pointer if your graph is not colored.

An optional vertex color vector for the second graph. Supply a null pointer if your graph

is not colored.

iso: Pointer to a boolean, the result is stored here.

map12: A vector or NULL pointer. If not NULL then an isomorphic mapping from graph1 to graph2 is stored here. If the input graphs are not isomorphic then this vector is cleared,

i a it will have langth zero

i.e. it will have length zero.

map21: Similar to map12, but for the mapping from graph2 to graph1.

sh: Splitting heuristics to be used for the graphs. See igraph_bliss_sh_t.

info1: If not NULL, information about the canonization of the first input graph is stored here.

Note that if the two graphs have different number of vertices or edges, then this is only partially filled. The memory used by this structure should be released when no longer

needed, see igraph_bliss_info_t for details.

info2: Same as *info1*, but for the second graph.

Returns:

Error code.

Time complexity: exponential, but in practice it is quite fast.

igraph_automorphisms — Number of automorphisms using Bliss.

The number of automorphisms of a graph is computed using Bliss. The result is returned as part of the *info* structure, in tag group_size. It is returned as a string, as it can be very high even for relatively small graphs. If the GNU MP library is used then this number is exact, otherwise a long double is used and it is only approximate. See also igraph_bliss_info_t.

Arguments:

graph: The input graph. Multiple edges between the same nodes are not supported and will cause

an incorrect result to be returned.

colors: An optional vertex color vector for the graph. Supply a null pointer is the graph is not

colored.

sh: The splitting heuristics to be used in Bliss. See igraph_bliss_sh_t.

info: The result is stored here, in particular in the group_size tag of info. The memory

used by this structure must be released when no longer needed, see igraph_blis-

s_info_t.

Returns:

Error code.

Time complexity: exponential, in practice it is fast for many graphs.

igraph_automorphism_group — Automorphism group generators using Bliss.

```
igraph_error_t igraph_automorphism_group(
    const igraph_t *graph, const igraph_vector_int_t *colors, igraph_vector_int
    igraph_bliss_sh_t sh, igraph_bliss_info_t *info);
```

The generators of the automorphism group of a graph are computed using Bliss. The generator set may not be minimal and may depend on the splitting heuristics. The generators are permutations represented using zero-based indexing.

Arguments:

graph: The input graph. Multiple edges between the same nodes are not supported and

will cause an incorrect result to be returned.

colors: An optional vertex color vector for the graph. Supply a null pointer is the graph

is not colored.

generators: Must be an initialized pointer vector. It will contain pointers to igraph_vec-

tor_int_t objects representing generators of the automorphism group.

sh: The splitting heuristics to be used in Bliss. See igraph_bliss_sh_t.

info: If not NULL then information on Bliss internals is stored here. The memory used

by this structure must to be freed when no longer needed, see igraph_blis-

s_info_t.

Returns:

Error code.

Time complexity: exponential, in practice it is fast for many graphs.

The VF2 algorithm

The VF2 algorithm can search for a subgraph in a larger graph, or check if two graphs are isomorphic. See P. Foggia, C. Sansone, M. Vento, An Improved algorithm for matching large graphs, Proc. of the 3rd IAPR-TC-15 International Workshop on Graph-based Representations, Italy, 2001.

VF2 supports both vertex and edge-colored graphs, as well as custom vertex or edge compatibility functions.

VF2 works with both directed and undirected graphs. Only simple graphs are supported. Self-loops or multi-edges must not be present in the graphs. Currently, the VF2 functions do not check that the input graph is simple: it is the responsibility of the user to pass in valid input.

igraph_isomorphic_vf2 — Isomorphism via VF2.

This function performs the VF2 algorithm via calling $igraph_get_isomorphism-s_vf2_callback()$.

Note that this function cannot be used for deciding subgraph isomorphism, use igraph_subiso-morphic_vf2() for that.

Arguments:

graph1: The first graph, may be directed or undirected.

graph2: The second graph. It must have the same directedness as graph1, otherwise

an error is reported.

vertex_color1: An optional color vector for the first graph. If color vectors are given for both

graphs, then the isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here

if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument

for explanation.

edge_color1: An optional edge color vector for the first graph. The matching edges in the

two graphs must have matching colors as well. Supply a null pointer here if

your graphs are not edge-colored.

edge_color2: The edge color vector for the second graph.

iso: Pointer to a logical constant, the result of the algorithm will be placed here.

map12: Pointer to an initialized vector or a NULL pointer. If not a NULL pointer

then the mapping from graph1 to graph2 is stored here. If the graphs are

not isomorphic then the vector is cleared (i.e. has zero elements).

map21: Pointer to an initialized vector or a NULL pointer. If not a NULL pointer

then the mapping from graph2 to graph1 is stored here. If the graphs are

not isomorphic then the vector is cleared (i.e. has zero elements).

 $node_compat_fn$: A pointer to a function of type $igraph_isocompat_t$. This function will

be called by the algorithm to determine whether two nodes are compatible.

edge_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions node_compat_fn and

edge_compat_fn.

Returns:

Error code.

See also:

Time complexity: exponential, what did you expect?

Example 17.1. File examples/simple/igraph_isomorphic_vf2.c

igraph_count_isomorphisms_vf2 — Number of isomorphisms via VF2.

This function counts the number of isomorphic mappings between two graphs. It uses the generic igraph_get_isomorphisms_vf2_callback() function.

Arguments:

graph1: The first input graph, may be directed or undirected.

graph2: The second input graph, it must have the same directedness as graph1, or

an error will be reported.

vertex_color1: An optional color vector for the first graph. If color vectors are given for both

graphs, then the isomorphism is calculated on the colored graphs; i.e. two

vertices can match only if their color also matches. Supply a null pointer here

if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument

for explanation.

edge_color1: An optional edge color vector for the first graph. The matching edges in the

two graphs must have matching colors as well. Supply a null pointer here if

your graphs are not edge-colored.

edge_color2: The edge color vector for the second graph.

count: Point to an integer, the result will be stored here.

node_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two nodes are compatible.

edge compat fn: A pointer to a function of type igraph isocompat t. This function will

be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions node_compat_fn and

edge_compat_fn.

Returns:

Error code.

Time complexity: exponential.

igraph_get_isomorphisms_vf2 — Collect all isomorphic mappings of two graphs.

This function finds all the isomorphic mappings between two simple graphs. It uses the igraph_get_isomorphisms_vf2_callback() function. Call the function with the same graph as graph1 and graph2 to get automorphisms.

Arguments:

graph1: The first input graph, may be directed or undirected.

graph2: The second input graph, it must have the same directedness as graph1, or

an error will be reported.

vertex_color1:
An optional color vector for the first graph. If color vectors are given for both

graphs, then the isomorphism is calculated on the colored graphs; i.e. two

vertices can match only if their color also matches. Supply a null pointer here

if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument

for explanation.

edge_color1: An optional edge color vector for the first graph. The matching edges in the

two graphs must have matching colors as well. Supply a null pointer here if

your graphs are not edge-colored.

edge_color2: The edge color vector for the second graph.

maps: Pointer to a list of integer vectors. On return it is empty if the input

graphs are not isomorphic. Otherwise it contains pointers to igraph_vector int t objects, each vector is an isomorphic mapping of graph2 to

graph1.

node_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two nodes are compatible.

edge_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions node_compat_fn and

edge_compat_fn.

Returns:

Error code.

Time complexity: exponential.

igraph_get_isomorphisms_vf2_callback — The generic VF2 interface

```
igraph_error_t igraph_get_isomorphisms_vf2_callback(
   const igraph_t *graph1, const igraph_t *graph2,
   const igraph_vector_int_t *vertex_color1, const igraph_vector_int_t *vertex_const igraph_vector_int_t *edge_color1, const igraph_vector_int_t *edge_col_igraph_vector_int_t *map12, igraph_vector_int_t *map21,
   igraph_isohandler_t *isohandler_fn, igraph_isocompat_t *node_compat_fn,
   igraph_isocompat_t *edge_compat_fn, void *arg
);
```

This function is an implementation of the VF2 isomorphism algorithm, see P. Foggia, C. Sansone, M. Vento, An Improved algorithm for matching large graphs, Proc. of the 3rd IAPR-TC-15 International Workshop on Graph-based Representations, Italy, 2001.

For using it you need to define a callback function of type igraph_isohandler_t. This function will be called whenever VF2 finds an isomorphism between the two graphs. The mapping between the two graphs will be also provided to this function. If the callback returns IGRAPH_SUCCESS, then the search is continued, otherwise it stops. IGRAPH_STOP as a return value can be used to indicate normal premature termination; any other return value will be treated as an igraph error code, making the caller function return the same error code as well. The callback function must not destroy the mapping vectors that are passed to it.

Arguments:

graph1: The first input graph.

graph2: The second input graph.

vertex_color1: An optional color vector for the first graph. If color vectors are given for both

graphs, then the isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null pointer here

if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument

for explanation.

edge_color1: An optional edge color vector for the first graph. The matching edges in the

two graphs must have matching colors as well. Supply a null pointer here if

your graphs are not edge-colored.

edge_color2: The edge color vector for the second graph.

map 12: Pointer to an initialized vector or NULL. If not NULL and the supplied graphs

are isomorphic then the permutation taking graph1 to graph is stored here. If not NULL and the graphs are not isomorphic then a zero-length vector is

returned.

map21: This is the same as map12, but for the permutation taking graph2 to

graph1.

isohandler_fn: The callback function to be called if an isomorphism is found. See also

igraph_isohandler_t.

node_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two nodes are compatible.

edge_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions isohandler_fn, node_com-

pat fn and edge compat fn.

Returns:

Error code.

Time complexity: exponential.

igraph_isohandler_t — Callback type, called when an isomorphism was found

See the details at the documentation of igraph get isomorphisms vf2 callback().

Arguments:

map12: The mapping from the first graph to the second.

map21: The mapping from the second graph to the first, the inverse of map12 basically.

arg: This extra argument was passed to igraph_get_isomorphisms_vf2_call-back() when it was called.

Returns:

IGRAPH_SUCCESS to continue the search, IGRAPH_STOP to terminate the search. Any other return value is interpreted as an igraph error code, which will then abort the search and return the same error code from the caller function.

igraph_isocompat_t — Callback type, called to check whether two vertices or edges are compatible

VF2 (subgraph) isomorphism functions can be restricted by defining relations on the vertices and/or edges of the graphs, and then checking whether the vertices (edges) match according to these relations.

This feature is implemented by two callbacks, one for vertices, one for edges. Every time igraph tries to match a vertex (edge) of the first (sub)graph to a vertex of the second graph, the vertex (edge) compatibility callback is called. The callback returns a logical value, giving whether the two vertices match.

Both callback functions are of type $igraph_isocompat_t$.

Arguments:

graph1: The first graph.

graph2: The second graph.

g1_num: The id of a vertex or edge in the first graph.

g2_num: The id of a vertex or edge in the second graph.

arg: Extra argument to pass to the callback functions.

Returns:

Logical scalar, whether vertex (or edge) g1_num in graph1 is compatible with vertex (or edge) g2_num in graph2.

igraph_subisomorphic_vf2 — Decide subgraph isomorphism using VF2

```
const igraph_vector_int_t *vertex_color2,
const igraph_vector_int_t *edge_color1,
const igraph_vector_int_t *edge_color2,
igraph_bool_t *iso, igraph_vector_int_t *map12,
igraph_vector_int_t *map21,
igraph_isocompat_t *node_compat_fn,
igraph_isocompat_t *edge_compat_fn,
void *arg);
```

Decides whether a subgraph of *graph1* is isomorphic to *graph2*. It uses igraph_get_subi-somorphisms_vf2_callback().

Arguments:

graph1: The first input graph, may be directed or undirected. This is supposed to be

the larger graph.

graph2: The second input graph, it must have the same directedness as graph1. This

is supposed to be the smaller graph.

vertex_color1: An optional color vector for the first graph. If color vectors are given for both

graphs, then the subgraph isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null

pointer here if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument

for explanation.

edge_color1: An optional edge color vector for the first graph. The matching edges in the

two graphs must have matching colors as well. Supply a null pointer here if

your graphs are not edge-colored.

edge_color2: The edge color vector for the second graph.

iso: Pointer to a boolean. The result of the decision problem is stored here.

map12: Pointer to a vector or NULL. If not NULL, then an isomorphic mapping from

graph1 to graph2 is stored here.

map21: Pointer to a vector of NULL. If not NULL, then an isomorphic mapping from

graph2 to graph1 is stored here.

node_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two nodes are compatible.

edge_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions node_compat_fn and

edge_compat_fn.

Returns:

Error code.

Time complexity: exponential.

igraph_count_subisomorphisms_vf2 — Number of subgraph isomorphisms using VF2

Count the number of isomorphisms between subgraphs of *graph1* and *graph2*. This function uses igraph_get_subisomorphisms_vf2_callback().

Arguments:

graph1: The first input graph, may be directed or undirected. This is supposed to be

the larger graph.

graph2: The second input graph, it must have the same directedness as graph1. This

is supposed to be the smaller graph.

vertex_color1: An optional color vector for the first graph. If color vectors are given for both

graphs, then the subgraph isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null

pointer here if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument

for explanation.

edge color1: An optional edge color vector for the first graph. The matching edges in the

two graphs must have matching colors as well. Supply a null pointer here if

your graphs are not edge-colored.

edge_color2: The edge color vector for the second graph.

count: Pointer to an integer. The number of subgraph isomorphisms is stored here.

node_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two nodes are compatible.

edge_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions node_compat_fn and

edge_compat_fn.

Returns:

Error code.

Time complexity: exponential.

igraph_get_subisomorphisms_vf2 — Return all subgraph isomorphic mappings.

igraph_error_t igraph_get_subisomorphisms_vf2(const igraph_t *graph1,

```
const igraph_t *graph2,
const igraph_vector_int_t *vertex_color1,
const igraph_vector_int_t *vertex_color2,
const igraph_vector_int_t *edge_color1,
const igraph_vector_int_t *edge_color2,
igraph_vector_int_list_t *maps,
igraph_isocompat_t *node_compat_fn,
igraph_isocompat_t *edge_compat_fn,
void *arg);
```

This function collects all isomorphic mappings of graph2 to a subgraph of graph1. It uses the igraph_get_subisomorphisms_vf2_callback() function. The graphs should be simple.

Arguments:

graph1: The first input graph, may be directed or undirected. This is supposed to be

the larger graph.

graph2: The second input graph, it must have the same directedness as graph1. This

is supposed to be the smaller graph.

vertex_color1: An optional color vector for the first graph. If color vectors are given for both

graphs, then the subgraph isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null

pointer here if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument

for explanation.

edge_color1: An optional edge color vector for the first graph. The matching edges in the

two graphs must have matching colors as well. Supply a null pointer here if

your graphs are not edge-colored.

edge_color2: The edge color vector for the second graph.

maps: Pointer to a list of integer vectors. On return it contains pointers to

igraph_vector_int_t objects, each vector is an isomorphic mapping

of graph2 to a subgraph of graph1.

node_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two nodes are compatible.

edge_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions node_compat_fn and

 $edge_compat_fn.$

Returns:

Error code.

Time complexity: exponential.

igraph_get_subisomorphisms_vf2_callback — Generic VF2 function for subgraph isomorphism problems.

```
igraph_error_t igraph_get_subisomorphisms_vf2_callback(
    const igraph_t *graph1, const igraph_t *graph2,
    const igraph_vector_int_t *vertex_color1, const igraph_vector_int_t *vertex
    const igraph_vector_int_t *edge_color1, const igraph_vector_int_t *edge_col
    igraph_vector_int_t *map12, igraph_vector_int_t *map21,
    igraph_isohandler_t *isohandler_fn, igraph_isocompat_t *node_compat_fn,
    igraph_isocompat_t *edge_compat_fn, void *arg
);
```

This function is the pair of <code>igraph_get_isomorphisms_vf2_callback()</code>, for subgraph isomorphism problems. It searches for subgraphs of <code>graph1</code> which are isomorphic to <code>graph2</code>. When it founds an isomorphic mapping it calls the supplied callback <code>isohandler_fn</code>. The mapping (and its inverse) and the additional <code>arg</code> argument are supplied to the callback.

Arguments:

graph1: The first input graph, may be directed or undirected. This is supposed to be

the larger graph.

graph2: The second input graph, it must have the same directedness as graph1. This

is supposed to be the smaller graph.

vertex_color1: An optional color vector for the first graph. If color vectors are given for both

graphs, then the subgraph isomorphism is calculated on the colored graphs; i.e. two vertices can match only if their color also matches. Supply a null

pointer here if your graphs are not colored.

vertex_color2: An optional color vector for the second graph. See the previous argument

for explanation.

edge_color1: An optional edge color vector for the first graph. The matching edges in the

two graphs must have matching colors as well. Supply a null pointer here if

your graphs are not edge-colored.

edge_color2: The edge color vector for the second graph.

map12: Pointer to a vector or NULL. If not NULL, then an isomorphic mapping from

graph1 to graph2 is stored here.

map21: Pointer to a vector of NULL. If not NULL, then an isomorphic mapping from

graph2 to graph1 is stored here.

isohandler_fn: A pointer to a function of type igraph_isohandler_t. This will be

called whenever a subgraph isomorphism is found. If the function returns IGRAPH_SUCCESS, then the search is continued. If the function returns IGRAPH_STOP, the search is terminated normally. Any other value is treat-

ed as an igraph error code.

 $node_compat_fn: \quad A \ pointer \ to \ a \ function \ of \ type \ \verb"igraph_isocompat_t". \ This \ function \ will$

be called by the algorithm to determine whether two nodes are compatible.

edge_compat_fn: A pointer to a function of type igraph_isocompat_t. This function will

be called by the algorithm to determine whether two edges are compatible.

arg: Extra argument to supply to functions isohandler_fn, node_com-

pat_fn and edge_compat_fn.

Returns:

Error code.

Time complexity: exponential.

Deprecated aliases

igraph_isomorphic_function_vf2 — The generic VF2 interface (deprecated alias).

```
igraph_error_t igraph_isomorphic_function_vf2(
    const igraph_t *graph1, const igraph_t *graph2,
    const igraph_vector_int_t *vertex_color1, const igraph_vector_int_t *vertex
    const igraph_vector_int_t *edge_color1, const igraph_vector_int_t *edge_col
    igraph_vector_int_t *map12, igraph_vector_int_t *map21,
    igraph_isohandler_t *isohandler_fn, igraph_isocompat_t *node_compat_fn,
    igraph_isocompat_t *edge_compat_fn, void *arg
);
```

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_get_isomorphisms_vf2_callback() instead.

igraph_subisomorphic_function_vf2 — Generic VF2 function for subgraph isomorphism problems (deprecated alias).

```
igraph_error_t igraph_subisomorphic_function_vf2(
    const igraph_t *graph1, const igraph_t *graph2,
    const igraph_vector_int_t *vertex_color1, const igraph_vector_int_t *vertex_const igraph_vector_int_t *edge_color1, const igraph_vector_int_t *edge_col_igraph_vector_int_t *map12, igraph_vector_int_t *map21,
    igraph_isohandler_t *isohandler_fn, igraph_isocompat_t *node_compat_fn,
    igraph_isocompat_t *edge_compat_fn, void *arg
);
```

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_get_subisomorphisms_vf2_callback() instead.

The LAD algorithm

The LAD algorithm can search for a subgraph in a larger graph, or check if two graphs are isomorphic. See Christine Solnon: AllDifferent-based Filtering for Subgraph Isomorphism. Artificial Intelligence, 174(12-13):850-864, 2010. https://doi.org/10.1016/j.artint.2010.05.002 as well as the homepage of the LAD library at http://liris.cnrs.fr/csolnon/LAD.html The implementation in igraph is based on LADv1, but it is modified to use igraph's own memory allocation and error handling.

LAD uses the concept of domains to indicate vertex compatibility when matching the pattern graph. Domains can be used to implement matching of colored vertices.

LAD works with both directed and undirected graphs. Graphs with multi-edges are not supported.

igraph_subisomorphic_lad — Check subgraph isomorphism with the LAD algorithm

Check whether pattern is isomorphic to a subgraph os target. The original LAD implementation by Christine Solnon was used as the basis of this code.

See more about LAD at http://liris.cnrs.fr/csolnon/LAD.html and in Christine Solnon: AllDifferent-based Filtering for Subgraph Isomorphism. Artificial Intelligence, 174(12-13):850-864, 2010. https://doi.org/10.1016/j.artint.2010.05.002

Arguments:

pattern: The smaller graph, it can be directed or undirected.

target: The bigger graph, it can be directed or undirected.

domains: A pointer vector, or a null pointer. If a pointer vector, then it must contain pointers

to igraph_vector_int_t objects and the length of the vector must match the number of vertices in the pattern graph. For each vertex, the IDs of the

compatible vertices in the target graph are listed.

iso: Pointer to a boolean, or a null pointer. If not a null pointer, then the boolean is set

to true if a subgraph isomorphism is found, and to false otherwise.

map: Pointer to a vector or a null pointer. If not a null pointer and a subgraph isomor-

phism is found, the matching vertices from the target graph are listed here, for each

vertex (in vertex ID order) from the pattern graph.

maps: Pointer to a list of integer vectors or a null pointer. If not a null pointer, then all

subgraph isomorphisms are stored in the vector list, in igraph_vector_int_t

objects.

induced: Boolean, whether to search for induced matching subgraphs.

time_limit: Processor time limit in seconds. Supply zero here for no limit. If the time limit is

over, then the function signals an error.

Returns:

Error code

See also:

igraph_subisomorphic_vf2() for the VF2 algorithm.

Time complexity: exponential.

Example 17.2. File examples/simple/igraph_subisomorphic_lad.c

Functions for small graphs

igraph_isoclass — Determine the isomorphism class of small graphs.

```
igraph_error_t igraph_isoclass(const igraph_t *graph, igraph_integer_t *isoclass
```

All graphs with a given number of vertices belong to a number of isomorphism classes, with every graph in a given class being isomorphic to each other.

This function gives the isomorphism class (a number) of a graph. Two graphs have the same isomorphism class if and only if they are isomorphic.

The first isomorphism class is numbered zero and it contains the edgeless graph. The last isomorphism class contains the full graph. The number of isomorphism classes for directed graphs with three vertices is 16 (between 0 and 15), for undirected graph it is only 4. For graphs with four vertices it is 218 (directed) and 11 (undirected). For 5 and 6 vertex undirected graphs, it is 34 and 156, respectively. These values can also be retrieved using <code>igraph_graph_count()</code>. For more information, see https://oeis.org/A000273 and https://oeis.org/A000088.

At the moment, 3- and 4-vertex directed graphs and 3 to 6 vertex undirected graphs are supported.

Multi-edges and self-loops are ignored by this function.

Arguments:

graph: The graph object.

isoclass: Pointer to an integer, the isomorphism class will be stored here.

Returns:

Error code.

See also:

```
igraph_isomorphic(), igraph_isoclass_subgraph(), igraph_iso-
class_create(),igraph_motifs_randesu().
```

Because of some limitations this function works only for graphs with three of four vertices.

Time complexity: O(|E|), the number of edges in the graph.

igraph_isoclass_subgraph — The isomorphism class of a subgraph of a graph.

This function identifies the isomorphism class of the subgraph induced the vertices specified in vids.

At the moment, 3- and 4-vertex directed graphs and 3 to 6 vertex undirected graphs are supported.

Multi-edges and self-loops are ignored by this function.

Arguments:

graph: The graph object.

vids: A vector containing the vertex IDs to be considered as a subgraph. Each vertex ID

should be included at most once.

isoclass: Pointer to an integer, this will be set to the isomorphism class.

Returns:

Error code.

See also:

```
igraph_isoclass(),igraph_isomorphic(),igraph_isoclass_create().
```

Time complexity: O((d+n)*n), d is the average degree in the network, and n is the number of vertices in vids.

igraph_isoclass_create — Creates a graph from the given isomorphism class.

This function creates the canonical representative graph of the given isomorphism class.

The isomorphism class is an integer between 0 and the number of unique unlabeled (i.e. non-isomorphic) graphs on the given number of vertices and give directedness. See https://oeis.org/A000273 and https://oeis.org/A000088 for the number of directed and undirected graphs on size nodes.

At the moment, 3- and 4-vertex directed graphs and 3 to 6 vertex undirected graphs are supported.

Arguments:

graph: Pointer to an uninitialized graph object.

size: The number of vertices to add to the graph.

number: The isomorphism class.

directed: Logical constant, whether to create a directed graph.

Returns:

Error code.

See also:

```
igraph_isoclass(), igraph_isoclass_subgraph(), igraph_isomorphic().
```

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the graph to create.

igraph_graph_count — The number of unlabelled graphs on the given number of vertices.

```
igraph_error_t igraph_graph_count(igraph_integer_t n, igraph_bool_t directed, id
```

Gives the number of unlabelled *simple* graphs on the specified number of vertices. The "isoclass" of a graph of this size is at most one less than this value.

This function is meant to be used in conjunction with isoclass and motif finder functions. It will only work for small *n* values for which the result is representable in an igraph_integer_t. For larger *n* values, an overflow error is raised.

Arguments:

n: The number of vertices.

directed: Boolean, whether to consider directed graphs.

count: Pointer to an integer, the result will be stored here.

Returns:

Error code.

See also:

```
igraph_isoclass(), igraph_motifs_randesu_callback().
```

Time complexity: O(1).

Utility functions

igraph permute vertices — Permute the vertices.

This function creates a new graph from the input graph by permuting its vertices according to the specified mapping. Call this function with the output of <code>igraph_canonical_permutation()</code> to create the canonical form of a graph.

Arguments:

graph: The input graph.

res: Pointer to an uninitialized graph object. The new graph is created here.

permutation: The permutation to apply. Vertex 0 is mapped to the first element of the vector,

vertex 1 to the second, etc.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in terms of the number of vertices and edges.

igraph_simplify_and_colorize — Simplify the graph and compute self-loop and edge multiplicities.

```
igraph_error_t igraph_simplify_and_colorize(
    const igraph_t *graph, igraph_t *res,
    igraph_vector_int_t *vertex_color, igraph_vector_int_t *edge_color);
```

This function creates a vertex and edge colored simple graph from the input graph. The vertex colors are computed as the number of incident self-loops to each vertex in the input graph. The edge colors are computed as the number of parallel edges in the input graph that were merged to create each edge in the simple graph.

The resulting colored simple graph is suitable for use by isomorphism checking algorithms such as VF2, which only support simple graphs, but can consider vertex and edge colors.

Arguments:

graph: The graph object, typically having self-loops or multi-edges.

res: An uninitialized graph object. The result will be stored here

vertex_color: Computed vertex colors corresponding to self-loop multiplicities.

edge_color: Computed edge colors corresponding to edge multiplicities

Returns:

Error code.

See also:

```
igraph_simplify(), igraph_isomorphic_vf2(), igraph_subisomor-
phic_vf2()
```

Deprecated functions

igraph_isomorphic_34 — Graph isomorphism for 3-4 vertices (deprecated alias).

```
igraph_error_t igraph_isomorphic_34(
     const igraph_t *graph1, const igraph_t *graph2, igraph_bool_t *iso
);
```

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph isomorphic() instead.

If you really care about performance and you *know* for sure that your input graphs are simple and have either 3 or 4 vertices for directed graphs, or 3-6 vertices for undirected graphs, you can compare their isomorphism classes obtained from <code>igraph_isoclass()</code> directly instead of calling <code>igraph_i-somorphic()</code>; this saves the cost of checking whether the graphs do not contain multiple edges or self-loops.

Arguments:

graph1: The first input graph.

graph2: The second input graph. Must have the same directedness as graph1.

iso: Pointer to a boolean, the result is stored here.

Returns:

Error code.

Time complexity: O(1).

Chapter 18. Graph coloring

igraph_vertex_coloring_greedy — Computes a vertex coloring using a greedy algorithm.

```
igraph_error_t igraph_vertex_coloring_greedy(const igraph_t *graph, igraph_vect
```

This function assigns a "color"—represented as a non-negative integer—to each vertex of the graph in such a way that neighboring vertices never have the same color. The obtained coloring is not necessarily minimal.

Vertices are colored one by one, choosing the smallest color index that differs from that of already colored neighbors. Colors are represented with non-negative integers 0, 1, 2, ...

Arguments:

graph: The input graph.

colors: Pointer to an initialized integer vector. The vertex colors will be stored here.

heuristic: The vertex ordering heuristic to use during greedy coloring. See igraph_color-

ing_greedy_t

Returns:

Error code.

Example 18.1. File examples/simple/igraph_coloring.c

igraph_coloring_greedy_t — Ordering heuristics for greedy graph coloring.

```
typedef enum {
    IGRAPH_COLORING_GREEDY_COLORED_NEIGHBORS = 0
} igraph_coloring_greedy_t;
```

Ordering heuristics for igraph_vertex_coloring_greedy().

Values:

IGRAPH_COLORING_GREEDY_COLORED_NEIGHBORS:

Choose vertex with largest number of already colored neighbors.

igraph_is_perfect — Checks if the graph is perfect.

igraph_error_t igraph_is_perfect(const igraph_t *graph, igraph_bool_t *perfect)

A perfect graph is an undirected graph in which the chromatic number of every induced subgraph equals the order of the largest clique of that subgraph. The chromatic number of a graph G is the smallest number of colors needed to color the vertices of G so that no two adjacent vertices share the same color.

Warning: This function may create the complement of the graph internally, which consumes a lot of memory. For moderately sized graphs, consider decomposing them into biconnected components and running the check separately on each component.

This implementation is based on the strong perfect graph theorem which was conjectured by Claude Berge and proved by Maria Chudnovsky, Neil Robertson, Paul Seymour, and Robin Thomas.

Arguments:

graph: The input graph. It is expected to be undirected and simple.

perfect: Pointer to an integer, the result will be stored here.

Returns:

Error code.

Time complexity: worst case exponenital, often faster in practice.

Chapter 19. Graph motifs, dyad census and triad census

This section deals with functions which find small induced subgraphs in a graph. These were first defined for subgraphs of two and three vertices by Holland and Leinhardt, and named dyad census and triad census.

igraph_dyad_census — Dyad census, as defined by Holland and Leinhardt.

Dyad census means classifying each pair of vertices of a directed graph into three categories: mutual (there is at least one edge from a to b and also from b to a); asymmetric (there is at least one edge either from a to b or from b to a, but not the other way) and null (no edges between a and b in either direction).

Holland, P.W. and Leinhardt, S. (1970). A Method for Detecting Structure in Sociometric Data. American Journal of Sociology, 70, 492-513.

Arguments:

graph: The input graph. For an undirected graph, there are no asymmetric connections.

mut: Pointer to a real, the number of mutual dyads is stored here.

asym: Pointer to a real, the number of asymmetric dyads is stored here.

null: Pointer to a real, the number of null dyads is stored here.

Returns:

Error code.

See also:

```
igraph_reciprocity(), igraph_triad_census().
```

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

igraph_triad_census — Triad census, as defined by Davis and Leinhardt.

```
igraph_error_t igraph_triad_census(const igraph_t *graph, igraph_vector_t *res)
```

Calculating the triad census means classifying every triple of vertices in a directed graph. A triple can be in one of 16 states:

- 003 A, B, C, the empty graph.
- 012 A->B, C, a graph with a single directed edge.
- 102 A<->B, C, a graph with a mutual connection between two vertices.
- 021D A<-B->C, the binary out-tree.
- 021U A->B<-C, the binary in-tree.
- 021C A->B->C, the directed line.
- 111D A<->B<-C.
- 111U A<->B->C.
- 030T A->B<-C, A->C.
- 030C A < -B < -C, A > C.
- 201 A<->B<->C.
- 120D A<-B->C, A<->C.
- 120U A->B<-C, A<->C.
- 120C A->B->C, A<->C.
- 210 A->B<->C, A<->C.
- A<->B<->C, A<->C, the complete graph.

See also Davis, J.A. and Leinhardt, S. (1972). The Structure of Positive Interpersonal Relations in Small Groups. In J. Berger (Ed.), Sociological Theories in Progress, Volume 2, 218-251. Boston: Houghton Mifflin.

This function calls <code>igraph_motifs_randesu()</code> which is an implementation of the FANMOD motif finder tool, see <code>igraph_motifs_randesu()</code> for details. Note that the order of the triads is not the same for <code>igraph_triad_census()</code> and <code>igraph_motifs_randesu()</code>.

Arguments:

graph: The input graph. A warning is given for undirected graphs, as the result is undefined for those.

res: Pointer to an initialized vector, the result is stored here in the same order as given in the list above. Note that this order is different than the one used by igraph_motifs_randesu().

Returns:

Error code.

See also:

```
igraph_motifs_randesu(), igraph_dyad_census().
```

Time complexity: TODO.

Finding triangles

igraph_adjacent_triangles — Count the number of triangles a vertex is part of.

Arguments:

graph: The input graph. Edge directions and multiplicities are ignored.

res: Initiliazed vector, the results are stored here.

vids: The vertices to perform the calculation for.

Returns:

Error mode.

See also:

```
igraph_list_triangles() to list them.
```

Time complexity: O(d^2 n), d is the average vertex degree of the queried vertices, n is their number.

igraph_list_triangles — Find all triangles in a graph.

The triangles are reported as a long list of vertex ID triplets. Use the int variant of igraph_matrix_view_from_vector() to create a matrix view into the vector where each triangle is stored in a column of the matrix (see the example).

Arguments:

graph: The input graph, edge directions are ignored. Multiple edges are ignored.

res: Pointer to an initialized integer vector, the result is stored here, in a long list of triples of vertex IDs. Each triple is a triangle in the graph. Each triangle is listed exactly once.

Returns:

Error code.

See also:

igraph_transitivity_undirected() to count the triangles, igraph_adjacen-t_triangles() to count the triangles a vertex participates in.

Time complexity: O(d^2 n), d is the average degree, n is the number of vertices.

Example 19.1. File examples/simple/igraph_list_triangles.c

Graph motifs

igraph_motifs_randesu — Count the number of motifs in a graph.

Motifs are small weakly connected induced subgraphs of a given structure in a graph. It is argued that the motif profile (i.e. the number of different motifs in the graph) is characteristic for different types of networks and network function is related to the motifs in the graph.

This function is able to find directed motifs of sizes three and four and undirected motifs of sizes three to six (i.e. the number of different subgraphs with three to six vertices in the network).

In a big network the total number of motifs can be very large, so it takes a lot of time to find all of them. In this case, a sampling method can be used. This function is capable of doing sampling via the <code>cut_prob</code> argument. This argument gives the probability that a branch of the motif search tree will not be explored. See S. Wernicke and F. Rasche: FANMOD: a tool for fast network motif detection, Bioinformatics 22(9), 1152--1153, 2006 for details. https://doi.org/10.1093/bioinformatics/btl038

Set the *cut_prob* argument to a zero vector for finding all motifs.

Directed motifs will be counted in directed graphs and undirected motifs in undirected graphs.

Arguments:

graph: The graph to find the motifs in.

hist: The result of the computation, it gives the number of motifs found for each isomor-

phism class. See <code>igraph_isoclass()</code> for help about isomorphism classes. Note that this function does *not* count isomorphism classes that are not connected and will

report NaN (more precisely IGRAPH_NAN) for them.

size: The size of the motifs to search for. For directed graphs, only 3 and 4 are implemented,

for undirected, 3 to 6. The limitation is not in the motif finding code, but the graph

isomorphism code.

cut_prob: Vector of probabilities for cutting the search tree at a given level. The first element is

the first level, etc. Supply all zeros here (of length size) to find all motifs in a graph.

Returns:

Error code.

See also:

igraph_motifs_randesu_estimate() for estimating the number of motifs in a graph, this can help to set the <code>cut_prob</code> parameter; <code>igraph_motifs_randesu_no()</code> to calculate the total number of motifs of a given size in a graph; <code>igraph_motifs_randesu_callback()</code> for calling a callback function for every motif found; <code>igraph_subisomorphic_lad()</code> for finding subgraphs on more than 4 (directed) or 6 (undirected) vertices; <code>igraph_graph_count()</code> to find the number of graph on a given number of vertices, i.e. the length of the <code>hist</code> vector.

Time complexity: TODO.

Example 19.2. File examples/simple/igraph_motifs_randesu.c

igraph_motifs_randesu_no — Count the total number of motifs in a graph.

This function counts the total number of motifs in a graph, i.e. the number of of (weakly) connected triplets or quadruplets, without assigning isomorphism classes to them.

Arguments:

graph: The graph object to study.

no: Pointer to an integer type, the result will be stored here.

size: The size of the motifs to count.

cut_prob: Vector giving the probabilities that a branch of the search tree will be cut at a given

level.

Returns:

Error code.

See also:

```
igraph_motifs_randesu(), igraph_motifs_randesu_estimate().
```

Time complexity: TODO.

igraph_motifs_randesu_estimate — Estimate the total number of motifs in a graph.

This function estimates the total number of weakly connected induced subgraphs, called motifs, of a fixed number of vertices. For example, an undirected complete graph on n vertices will have one

motif of size n, and n motifs of size n - 1. As another example, one triangle and a separate vertex will have zero motifs of size four.

This function is useful for large graphs for which it is not feasible to count all the different motifs, because there are very many of them.

The total number of motifs is estimated by taking a sample of vertices and counts all motifs in which these vertices are included. (There is also a *cut_prob* parameter which gives the probabilities to cut a branch of the search tree.)

Directed motifs will be counted in directed graphs and undirected motifs in undirected graphs.

Arguments:

graph: The graph object to study.

est: Pointer to an integer type, the result will be stored here.

size: The size of the motifs to look for.

cut_prob: Vector giving the probabilities to cut a branch of the search tree and omit counting

the motifs in that branch. It contains a probability for each level. Supply size

zeros here to count all the motifs in the sample.

sample_size: The number of vertices to use as the sample. This parameter is only used if the

parsample argument is a null pointer.

parsample: Either pointer to an initialized vector or a null pointer. If a vector then the vertex

IDs in the vector are used as a sample. If a null pointer then the <code>sample_size</code> argument is used to create a sample of vertices drawn with uniform probability.

Returns:

Error code.

See also:

```
igraph_motifs_randesu(), igraph_motifs_randesu_no().
```

Time complexity: TODO.

igraph_motifs_randesu_callback — Finds motifs in a graph and calls a function for each of them.

Similarly to <code>igraph_motifs_randesu()</code>, this function is able to find directed motifs of sizes three and four and undirected motifs of sizes three to six (i.e. the number of different subgraphs with three to six vertices in the network). However, instead of counting them, the function will call a callback function for each motif found to allow further tests or post-processing.

The *cut_prob* argument also allows sampling the motifs, just like for igraph_motifs_randesu(). Set the *cut_prob* argument to a zero vector for finding all motifs.

Arguments:

graph: The graph to find the motifs in.

size: The size of the motifs to search for. Only three and four are implemented currently.

The limitation is not in the motif finding code, but the graph isomorphism code.

cut_prob: Vector of probabilities for cutting the search tree at a given level. The first element is

the first level, etc. Supply all zeros here (of length size) to find all motifs in a graph.

callback: A pointer to a function of type igraph_motifs_handler_t. This function will

be called whenever a new motif is found.

extra: Extra argument to pass to the callback function.

Returns:

Error code.

Time complexity: TODO.

Example 19.3. File examples/simple/igraph_motifs_randesu.c

igraph_motifs_handler_t — Callback type for igraph motifs randesu callback

igraph_motifs_randesu_callback() calls a specified callback function whenever a new
motif is found during a motif search. This callback function must be of type igraph_motifs_handler_t. It has the following arguments:

Arguments:

graph: The graph that that algorithm is working on. Of course this must not be modified.

vids: The IDs of the vertices in the motif that has just been found. This vector is owned by

the motif search algorithm, so do not modify or destroy it; make a copy of it if you

need it later.

isoclass: The isomorphism class of the motif that has just been found. Use

igraph_graph_count() to find the maximum possible isoclass for graphs of a given size. See igraph_isoclass and igraph_isoclass_subgraph for

more information.

extra: The extra argument that was passed to igraph_motifs_randesu_call-

back().

Returns:

IGRAPH_SUCCESS to continue the motif search, IGRAPH_STOP to stop the motif search and return to the caller normally. Any other return value is interpreted as an igraph error code, which will terminate the search and return the same error code to the caller.

a	1	
See	0	CO

igraph_motifs_randesu_callback()

Chapter 20. Generating layouts for graph drawing

2D layout generators

Layout generator functions (or at least most of them) try to place the vertices and edges of a graph on a 2D plane or in 3D space in a way which visually pleases the human eye.

They take a graph object and a number of parameters as arguments and return an igraph_matrix_t, in which each row gives the coordinates of a vertex.

igraph_layout_random — Places the vertices uniform randomly on a plane.

igraph_error_t igraph_layout_random(const igraph_t *graph, igraph_matrix_t *res

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized as

needed.

Returns:

Error code. The current implementation always returns with success.

Time complexity: O(|V|), the number of vertices.

igraph_layout_circle — Places the vertices uniformly on a circle in arbitrary order.

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized as

needed.

der: The order of the vertices on the circle. The vertices not included here, will be placed at

(0,0). Supply igraph_vss_all() here to place vertices in the order of their vertex IDs.

Returns:

Error code.

Time complexity: O(|V|), the number of vertices.

igraph_layout_star — Generates a star-like layout.

Arguments:

graph: The input graph. Its edges are ignored by this function.

res: Pointer to an initialized matrix object. This will contain the result and will be resized

as needed.

center: The id of the vertex to put in the center. You can set it to any arbitrary value for the

special case when the input graph has no vertices; otherwise it must be between 0 and

the number of vertices minus 1.

order: A numeric vector giving the order of the vertices (including the center vertex!). If a null

pointer, then the vertices are placed in increasing vertex ID order.

Returns:

Error code.

Time complexity: O(|V|), linear in the number of vertices.

See also:

igraph_layout_circle() and other layout generators.

igraph_layout_grid — Places the vertices on a regular grid on the plane.

igraph_error_t igraph_layout_grid(const igraph_t *graph, igraph_matrix_t *res,

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized as

needed.

width: The number of vertices in a single row of the grid. When zero or negative, the width of the

grid will be the square root of the number of vertices, rounded up if needed.

Returns:

Error code. The current implementation always returns with success.

Time complexity: O(|V|), the number of vertices.

igraph_layout_graphopt — Optimizes vertex layout via the graphopt algorithm.

This is a port of the graphopt layout algorithm by Michael Schmuhl. graphopt version 0.4.1 was rewritten in C and the support for layers was removed (might be added later) and a code was a bit reorganized to avoid some unnecessary steps is the node charge (see below) is zero.

Graphopt uses physical analogies for defining attracting and repelling forces among the vertices and then the physical system is simulated until it reaches an equilibrium. (There is no simulated annealing or anything like that, so a stable fixed point is not guaranteed.)

See also http://www.schmuhl.org/graphopt/ for the original graphopt.

Arguments:

graph: The input graph.

res: Pointer to an initialized matrix, the result will be stored here and its initial

contents are used as the starting point of the simulation if the use_seed argument is true. Note that in this case the matrix should have the proper size, otherwise a warning is issued and the supplied values are ignored. If no starting positions are given (or they are invalid) then a random starting

position is used. The matrix will be resized if needed.

niter: Integer constant, the number of iterations to perform. Should be a couple of

hundred in general. If you have a large graph then you might want to only do a few iterations and then check the result. If it is not good enough you can feed it in again in the res argument. The original graphopt default is 500.

node_charge: The charge of the vertices, used to calculate electric repulsion. The original

graphopt default is 0.001.

node_mass: The mass of the vertices, used for the spring forces. The original graphopt

defaults to 30.

spring_length: The length of the springs. The original graphopt defaults to zero.

spring_constant: The spring constant, the original graphopt defaults to one.

max_sa_movement: Real constant, it gives the maximum amount of movement allowed in a

single step along a single axis. The original graphopt default is 5.

use_seed: Logical scalar, whether to use the positions in res as a starting configura-

tion. See also res above.

Returns:

Error code.

Time complexity: $O(n(|V|^2+|E|))$, n is the number of iterations, |V| is the number of vertices, |E| the number of edges. If $node_charge$ is zero then it is only O(n|E|).

igraph_layout_bipartite — Simple layout for bipartite graphs.

The layout is created by first placing the vertices in two rows, according to their types. Then the positions within the rows are optimized to minimize edge crossings, by calling igraph_layout_sugiyama().

Arguments:

graph: The input graph.

types: A boolean vector containing ones and zeros, the vertex types. Its length must match the

number of vertices in the graph.

res: Pointer to an initialized matrix, the result, the x and y coordinates are stored here.

hgap: The preferred minimum horizontal gap between vertices in the same layer (i.e. vertices

of the same type).

vgap: The distance between layers.

maxiter: Maximum number of iterations in the crossing minimization stage. 100 is a reasonable

default; if you feel that you have too many edge crossings, increase this.

Returns:

Error code.

See also:

igraph_layout_sugiyama().

The DrL layout generator

DrL is a sophisticated layout generator developed and implemented by Shawn Martin et al. As of October 2012 the original DrL homepage is unfortunately not available. You can read more about this algorithm in the following technical report: Martin, S., Brown, W.M., Klavans, R., Boyack, K.W., DrL: Distributed Recursive (Graph) Layout. SAND Reports, 2008. 2936: p. 1-10.

Only a subset of the complete DrL functionality is included in igraph, parallel runs and recursive, multi-level layouting is not supported.

The parameters of the layout are stored in an <code>igraph_layout_drl_options_t</code> structure, this can be initialized by calling the function <code>igraph_layout_drl_options_init()</code>. The fields of this structure can then be adjusted by hand if needed. The layout is calculated by an <code>igraph_layout_drl()</code> call.

igraph_layout_drl_options_t — Parameters for the DrL layout generator

```
typedef struct igraph_layout_drl_options_t {
    igraph_real_t
                    edge_cut;
   igraph_integer_t init_iterations;
    igraph real t init temperature;
    igraph_real_t
                   init_attraction;
    igraph_real_t init_damping_mult;
    igraph_integer_t liquid_iterations;
    igraph_real_t
                    liquid_temperature;
    igraph_real_t
                     liquid_attraction;
    igraph_real_t
                     liquid_damping_mult;
    igraph_integer_t expansion_iterations;
                     expansion_temperature;
    igraph_real_t
    igraph_real_t
                     expansion_attraction;
    igraph_real_t
                     expansion_damping_mult;
    igraph_integer_t cooldown_iterations;
    igraph_real_t
                     cooldown_temperature;
    igraph_real_t
                     cooldown_attraction;
    igraph_real_t
                     cooldown_damping_mult;
    igraph_integer_t crunch_iterations;
    igraph_real_t
                     crunch_temperature;
    igraph_real_t
                     crunch_attraction;
    igraph_real_t
                     crunch_damping_mult;
    igraph_integer_t simmer_iterations;
    igraph_real_t
                     simmer_temperature;
    igraph_real_t
                     simmer_attraction;
    igraph_real_t
                     simmer_damping_mult;
} igraph_layout_drl_options_t;
```

Values:

edge_cut: The edge cutting parameter. Edge cutting is done in the late

stages of the algorithm in order to achieve less dense layouts. Edges are cut if there is a lot of stress on them (a large value in the objective function sum). The edge cutting parameter is a value between 0 and 1 with 0 representing no edge cutting and 1 representing maximal edge cutting. The default value is 32/40.

init_iterations: Number of iterations, initial phase.

init_temperature: Start temperature, initial phase.

init_attraction: Attraction, initial phase.

init_damping_mult: Damping factor, initial phase.

liquid_iterations: Number of iterations in the liquid phase.

liquid_temperature: Start temperature in the liquid phase.

liquid attraction: Attraction in the liquid phase.

liquid_damping_mult: Multiplicatie damping factor, liquid phase.

Generating layouts for graph drawing

expansion_iterations: Number of iterations in the expansion phase.

expansion_temperature: Start temperature in the expansion phase.

expansion_attraction: Attraction, expansion phase.

expansion_damping_mult: Damping factor, expansion phase.

cooldown_iterations: Number of iterations in the cooldown phase.

cooldown_temperature: Start temperature in the cooldown phase.

cooldown_attraction: Attraction in the cooldown phase.

cooldown_damping_mult: Damping fact int the cooldown phase.

crunch_iterations: Number of iterations in the crunch phase.

crunch_temperature: Start temperature in the crunch phase.

crunch_attraction: Attraction in the crunch phase.

crunch_damping_mult: Damping factor in the crunch phase.

simmer_iterations: Number of iterations in the simmer phase.

simmer_temperature: Start temperature in te simmer phase.

simmer_attraction: Attraction in the simmer phase.

simmer_damping_mult: Multiplicative damping factor in the simmer phase.

igraph_layout_drl_default_t — Predefined parameter templates for the DrL layout generator

These constants can be used to initialize a set of DrL parameters. These can then be modified according to the user's needs.

Values:

IGRAPH_LAYOUT_DRL_DE- The deafult parameters.

FAULT:

IGRAPH_LAYOUT_DR- Slightly modified parameters to get a coarser layout.

L_COARSEN:

IGRAPH_LAYOUT_DR- An even coarser layout.

L_COARSEST:

IGRAPH_LAYOUT_DRL_RE- Refine an already calculated layout.

FINE:

IGRAPH_LAYOUT_DRL_FINAL: Finalize an already refined layout.

igraph_layout_drl_options_init — Initialize parameters for the DrL layout generator

This function can be used to initialize the struct holding the parameters for the DrL layout generator. There are a number of predefined templates available, it is a good idea to start from one of these by modifying some parameters.

Arguments:

options: The struct to initialize.

temp1: The template to use. Currently the following templates are sup-

plied: IGRAPH_LAYOUT_DRL_DEFAULT, IGRAPH_LAYOUT_DRL_COARSEN, IGRAPH_LAYOUT_DRL_COARSEST, IGRAPH_LAYOUT_DRL_REFINE and

IGRAPH_LAYOUT_DRL_FINAL.

Returns:

Error code.

Time complexity: O(1).

igraph_layout_drl — The DrL layout generator

This function implements the force-directed DrL layout generator. Please see more in the following technical report: Martin, S., Brown, W.M., Klavans, R., Boyack, K.W., DrL: Distributed Recursive (Graph) Layout. SAND Reports, 2008. 2936: p. 1-10.

Arguments:

graph: The input graph.

use_seed: Logical scalar, if true, then the coordinates supplied in the res argument are used as

starting points.

res: Pointer to a matrix, the result layout is stored here. It will be resized as needed.

options: The parameters to pass to the layout generator.

weights: Edge weights, pointer to a vector. If this is a null pointer then every edge will have

the same weight.

Returns:

Error code.

Time complexity: ???.

igraph_layout_drl_3d — The DrL layout generator, 3d version.

This function implements the force-directed DrL layout generator. Please see more in the technical report: Martin, S., Brown, W.M., Klavans, R., Boyack, K.W., DrL: Distributed Recursive (Graph) Layout. SAND Reports, 2008. 2936: p. 1-10.

This function uses a modified DrL generator that does the layout in three dimensions.

Arguments:

graph: The input graph.

use_seed: Logical scalar, if true, then the coordinates supplied in the res argument are used as

starting points.

res: Pointer to a matrix, the result layout is stored here. It will be resized as needed.

options: The parameters to pass to the layout generator.

weights: Edge weights, pointer to a vector. If this is a null pointer then every edge will have

the same weight.

Returns:

Error code.

Time complexity: ???.

See also:

igraph_layout_drl() for the standard 2d version.

igraph_layout_fruchterman_reingold — Places the vertices on a plane according to the Fruchterman-Reingold algorithm.

```
const igraph_vector_t *maxx,
const igraph_vector_t *miny,
const igraph_vector_t *maxy);
```

This is a force-directed layout that simulates an attractive force f_a between connected vertex pairs and a repulsive force f_r between all vertex pairs. The forces are computed as a function of the distance d between the two vertices as

```
f_a(d) = -w * d^2 and f_r(d) = 1/d,
```

where w represents the edge weight. The equilibrium distance of two connected vertices is thus $1/w^3$, assuming no other forces acting on them.

In disconnected graphs, igraph effectively inserts a weak connection of weight $n^{(-3/2)}$ between all pairs of vertices, where n is the vertex count. This ensures that components are kept near each other.

Reference:

Fruchterman, T.M.J. and Reingold, E.M.: Graph Drawing by Force-directed Placement. Software -- Practice and Experience, 21/11, 1129--1164, 1991. https://doi.org/10.1002/spe.4380211102

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be re-

sized as needed.

use_seed: Logical, if true the supplied values in the res argument are used as an initial layout,

if false a random initial layout is used.

niter: The number of iterations to do. A reasonable default value is 500.

start_temp: Start temperature. This is the maximum amount of movement allowed along one

axis, within one step, for a vertex. Currently it is decreased linearly to zero during

the iteration.

grid: Whether to use the (fast but less accurate) grid based version of the algo-

rithm. Possible values: IGRAPH_LAYOUT_GRID, IGRAPH_LAYOUT_NOGRID, IGRAPH_LAYOUT_AUTOGRID. The last one uses the grid based version only for

large graphs, currently the ones with more than 1000 vertices.

weight: Pointer to a vector containing edge weights, the attraction along the edges will be

multiplied by these. Weights must be positive. It will be ignored if it is a null-

pointer.

minx: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives

the minimum "x" coordinate for every vertex.

maxx: Same as minx, but the maximum "x" coordinates.

miny: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives

the minimum "y" coordinate for every vertex.

maxy: Same as miny, but the maximum "y" coordinates.

Returns:

Error code.

Time complexity: $O(|V|^2)$ in each iteration, |V| is the number of vertices in the graph.

igraph_layout_kamada_kawai — Places the vertices on a plane according to the Kamada-Kawai algorithm.

This is a force-directed layout. A spring is inserted between all pairs of vertices, both those which are directly connected and those that are not. The unstretched length of springs is chosen based on the undirected graph distance between the corresponding pair of vertices. Thus, in a weighted graph, increasing the weight between two vertices pushes them apart. The Young modulus of springs is inversely proportional to the graph distance, ensuring that springs between far-apart veritces will have a smaller effect on the layout.

This layout works particularly well for locally connected spatial networks such as lattices.

This layout algorithm is not suitable for large graphs. The memory requirements are of the order $O(|V|^2)$.

Reference:

Kamada, T. and Kawai, S.: An Algorithm for Drawing General Undirected Graphs. Information Processing Letters, 31/1, 7--15, 1989. https://doi.org/10.1016/0020-0190(89)90102-6

Arguments:

A graph object. graph: Pointer to an initialized matrix object. This will contain the result (x-positions in colres: umn zero and y-positions in column one) and will be resized if needed. Boolean, whether to use the values supplied in the res argument as the initial conuse_seed: figuration. If zero and there are any limits on the X or Y coordinates, then a random initial configuration is used. Otherwise the vertices are placed on a circle of radius 1 as the initial configuration. maxiter: The maximum number of iterations to perform. A reasonable default value is at least ten (or more) times the number of vertices. Stop the iteration, if the maximum delta value of the algorithm is smaller than still. It epsilon: is safe to leave it at zero, and then maxiter iterations are performed. kkconst: The Kamada-Kawai vertex attraction constant. Typical value: number of vertices. weights: Edge weights, larger values will result longer edges. Weights must be positive. Pass NULL to assume unit weights for all edges. minx: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the minimum "x" coordinate for every vertex. maxx: Same as minx, but the maximum "x" coordinates. Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the miny: minimum "y" coordinate for every vertex.

maxy: Same as miny, but the maximum "y" coordinates.

Returns:

Error code.

Time complexity: O(|V|) for each iteration, after an $O(|V|^2 \log |V|)$ initialization step. |V| is the number of vertices in the graph.

igraph_layout_gem — Layout graph according to GEM algorithm.

The GEM layout algorithm, as described in Arne Frick, Andreas Ludwig, Heiko Mehldau: A Fast Adaptive Layout Algorithm for Undirected Graphs, Proc. Graph Drawing 1994, LNCS 894, pp. 388-403, 1995.

Arguments:

graph: The input graph. Edge directions are ignored in directed graphs.

res: The result is stored here. If the use_seed argument is true (non-zero), then this

matrix is also used as the starting point of the algorithm.

use_seed: Boolean, whether to use the supplied coordinates in res as the starting point. If false

(zero), then a uniform random starting point is used.

maxiter: The maximum number of iterations to perform. Updating a single vertex counts as

an iteration. A reasonable default is 40 * n * n, where n is the number of vertices. The original paper suggests 4 * n * n, but this usually only works if the other parameters

are set up carefully.

temp_max: The maximum allowed local temperature. A reasonable default is the number of

vertices.

temp_min: The global temperature at which the algorithm terminates (even before reaching

maxiter iterations). A reasonable default is 1/10.

temp_init: Initial local temperature of all vertices. A reasonable default is the square root of the

number of vertices.

Returns:

Error code.

Time complexity: O(t * n * (n+e)), where n is the number of vertices, e is the number of edges and t is the number of time steps performed.

igraph_layout_davidson_harel — Davidson-Harel layout algorithm

This function implements the algorithm by Davidson and Harel, see Ron Davidson, David Harel: Drawing Graphs Nicely Using Simulated Annealing. ACM Transactions on Graphics 15(4), pp. 301-331, 1996.

The algorithm uses simulated annealing and a sophisticated energy function, which is unfortunately hard to parameterize for different graphs. The original publication did not disclose any parameter values, and the ones below were determined by experimentation.

The algorithm consists of two phases, an annealing phase, and a fine-tuning phase. There is no simulated annealing in the second phase.

Our implementation tries to follow the original publication, as much as possible. The only major difference is that coordinates are explicitly kept within the bounds of the rectangle of the layout.

Arguments:

C .	
graph:	The input graph, edge directions are ignored.
res:	A matrix, the result is stored here. It can be used to supply start coordinates, see use_seed.
use_seed:	Boolean, whether to use the supplied res as start coordinates.
maxiter:	The maximum number of annealing iterations. A reasonable value for smaller graphs is 10.
fineiter:	The number of fine tuning iterations. A reasonable value is $\max(10, \log 2(n))$ where n is the number of vertices.
cool_fact:	Cooling factor. A reasonable value is 0.75.
weight_node_dist:	Weight for the node-node distances component of the energy function. Reasonable value: 1.0.
weight_border:	Weight for the distance from the border component of the energy function. It can be set to zero, if vertices are allowed to sit on the border.
weight_edge_lengths:	Weight for the edge length component of the energy function, a reasonable value is the density of the graph divided by 10.
weight_edge_crossings:	Weight for the edge crossing component of the energy function, a reasonable default is 1 minus the square root of the density of the graph.
weight_node_edge_dist:	Weight for the node-edge distance component of the energy function. A reasonable value is 1 minus the density, divided by 5.

Returns:

Error code.

Time complexity: one first phase iteration has time complexity $O(n^2+m^2)$, one fine tuning iteration has time complexity O(mn). Time complexity might be smaller if some of the weights of the components of the energy function are set to zero.

igraph_layout_mds — Place the vertices on a plane using multidimensional scaling.

This layout requires a distance matrix, where the intersection of row i and column j specifies the desired distance between vertex i and vertex j. The algorithm will try to place the vertices in a space having a given number of dimensions in a way that approximates the distance relations prescribed in the distance matrix. igraph uses the classical multidimensional scaling by Torgerson; for more details, see Cox & Cox: Multidimensional Scaling (1994), Chapman and Hall, London.

If the input graph is disconnected, igraph will decompose it first into its subgraphs, lay out the subgraphs one by one using the appropriate submatrices of the distance matrix, and then merge the layouts using igraph_layout_merge_dla. Since igraph_layout_merge_dla works for 2D layouts only, you cannot run the MDS layout on disconnected graphs for more than two dimensions.

Warning: if the graph is symmetric to the exchange of two vertices (as is the case with leaves of a tree connecting to the same parent), classical multidimensional scaling may assign the same coordinates to these vertices.

Arguments:

graph: A graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized if

needed.

dist: The distance matrix. It must be symmetric and this function does not check whether the

matrix is indeed symmetric. Results are unspecified if you pass a non-symmetric matrix here. You can set this parameter to null; in this case, the shortest path lengths between

vertices will be used as distances.

dim: The number of dimensions in the embedding space. For 2D layouts, supply 2 here.

Returns:

Error code.

Added in version 0.6.

Time complexity: usually around $O(|V|^2 dim)$.

igraph_layout_lgl — Force based layout algorithm for large graphs.

Generating layouts for graph drawing

```
igraph_real_t area, igraph_real_t coolexp,
igraph_real_t repulserad, igraph_real_t cellsize,
igraph_integer_t proot);
```

This is a layout generator similar to the Large Graph Layout algorithm and program (http://lgl.source-forge.net/). But unlike LGL, this version uses a Fruchterman-Reingold style simulated annealing algorithm for placing the vertices. The speedup is achieved by placing the vertices on a grid and calculating the repulsion only for vertices which are closer to each other than a limit.

Arguments:

graph: The (initialized) graph object to place. It must be connnected; disconnected graphs

are not handled by the algorithm.

res: Pointer to an initialized matrix object to hold the result. It will be resized if needed.

maxit: The maximum number of cooling iterations to perform for each layout step. A

reasonable default is 150.

maxdelta: The maximum length of the move allowed for a vertex in a single iteration. A

reasonable default is the number of vertices.

area: This parameter gives the area of the square on which the vertices will be placed. A

reasonable default value is the number of vertices squared.

coolexp: The cooling exponent. A reasonable default value is 1.5.

repulserad: Determines the radius at which vertex-vertex repulsion cancels out attraction of

adjacent vertices. A reasonable default value is area times the number of vertices.

cellsize: The size of the grid cells, one side of the square. A reasonable default value is the

fourth root of area (or the square root of the number of vertices if area is also

left at its default value).

proot: The root vertex, this is placed first, its neighbors in the first iteration, second neigh-

bors in the second, etc. If negative then a random vertex is chosen.

Returns:

Error code.

Added in version 0.2.

Time complexity: ideally O(dia*maxit*(|V|+|E|)), |V| is the number of vertices, dia is the diameter of the graph, worst case complexity is still $O(dia*maxit*(|V|^2+|E|))$, this is the case when all vertices happen to be in the same grid cell.

Layouts for trees and acyclic graphs

igraph_layout_reingold_tilford — Reingold-Tilford layout for tree graphs.

```
const igraph_vector_int_t *roots,
const igraph_vector_int_t *rootlevel);
```

Arranges the nodes in a tree where the given node is used as the root. The tree is directed downwards and the parents are centered above its children. For the exact algorithm, see:

Reingold, E and Tilford, J: Tidier drawing of trees. IEEE Trans. Softw. Eng., SE-7(2):223--228, 1981. https://doi.org/10.1109/TSE.1981.234519

If the given graph is not a tree, a breadth-first search is executed first to obtain a possible spanning tree.

Arguments:

graph: The graph object.

res: The result, the coordinates in a matrix. The parameter should point to an initialized

matrix object and will be resized.

mode: Specifies which edges to consider when building the tree. If it is IGRAPH_OUT then

only the outgoing, if it is IGRAPH_IN then only the incoming edges of a parent are considered. If it is IGRAPH_ALL then all edges are used (this was the behavior in igraph 0.5 and before). This parameter also influences how the root vertices are

calculated, if they are not given. See the roots parameter.

roots: The index of the root vertex or root vertices. The set of roots should be specified so

that all vertices of the graph are reachable from them. Simply put, in the udirected case, one root should be given from each connected component. If roots is NULL or a pointer to an empty vector, then the roots will be selected automatically. Currently, automatic root selection prefers low eccentricity vertices in graphs with fewer than 500 vertices, and high degree vertices (acording to mode) in larger graphs. The root selection heuristic may change without notice. To ensure a consistent output, please specify the roots manually. The igraph_roots_for_tree_layout()

function gives more control over automatic root selection.

rootlevel: This argument can be useful when drawing forests which are not trees (i.e. they are

unconnected and have tree components). It specifies the level of the root vertices for every tree in the forest. It is only considered if not a null pointer and the *roots*

argument is also given (and it is not a null pointer of an empty vector).

Returns:

Error code.

Added in version 0.2.

See also:

Example 20.1. File examples/simple/igraph_layout_reingold_tilford.c

igraph_layout_reingold_tilford_circular — Circular Reingold-Tilford layout for trees.

This layout is almost the same as <code>igraph_layout_reingold_tilford()</code>, but the tree is drawn in a circular way, with the root vertex in the center.

Arguments:

graph: The graph object.

res: The result, the coordinates in a matrix. The parameter should point to an initialized

matrix object and will be resized.

mode: Specifies which edges to consider when building the tree. If it is IGRAPH_OUT then

only the outgoing, if it is IGRAPH_IN then only the incoming edges of a parent are considered. If it is IGRAPH_ALL then all edges are used (this was the behavior in igraph 0.5 and before). This parameter also influences how the root vertices are

calculated, if they are not given. See the *roots* parameter.

roots: The index of the root vertex or root vertices. The set of roots should be specified so

that all vertices of the graph are reachable from them. Simply put, in the udirected case, one root should be given from each connected component. If roots is NULL or a pointer to an empty vector, then the roots will be selected automatically. Currently, automatic root selection prefers low ecccentricity vertices in graphs with fewer than 500 vertices, and high degree vertices (acording to mode) in larger graphs. The root selection heuristic may change without notice. To ensure a consistent output, please

specify the roots manually.

rootlevel: This argument can be useful when drawing forests which are not trees (i.e. they are

unconnected and have tree components). It specifies the level of the root vertices for every tree in the forest. It is only considered if not a null pointer and the *roots*

argument is also given (and it is not a null pointer or an empty vector).

Returns:

Error code.

See also:

igraph_layout_reingold_tilford().

igraph_roots_for_tree_layout — Roots suitable for a nice tree layout.

```
igraph_error_t igraph_roots_for_tree_layout(
    const igraph_t *graph,
    igraph_neimode_t mode,
    igraph_vector_int_t *roots,
    igraph_root_choice_t heuristic);
```

This function chooses a root, or a set of roots suitable for visualizing a tree, or a tree-like graph. It is typically used with <code>igraph_layout_reingold_tilford()</code>. The principle is to select a minimal set of roots so that all other vertices will be reachable from them.

In the undirected case, one root is chosen from each connected component. In the directed case, one root is chosen from each strongly connected component that has no incoming (or outgoing) edges (depending on 'mode'). When more than one root choice is possible, vertices are prioritized based on the given heuristic.

Arguments:

graph: The graph, typically a tree, but any graph is accepted.

mode: Whether to interpret the input as undirected, a directed out-tree or in-tree.

roots: An initialized integer vector, the roots will be returned here.

heuristic: The heuristic to use for breaking ties when multiple root choices are possible.

IGRAPH_ROOT_CHOICE_DE- Choose the vertices

GREE

Choose the vertices with the highest degree (out- or in-degree in directed mode). This simple heuristic is fast even in large graphs.

IGRAPH_ROOT_CHOICE_EC-

CENTRICITY

Choose the vertices with the lowest eccentricity. This usually results in a "wide and shallow" tree layout. While this heuristic produces high-quality results, it is slow for large graphs: computing the eccentricities has quadractic complexity in the number of vertices.

Returns:

Error code.

Time complexity: depends on the heuristic.

igraph_layout_sugiyama — Sugiyama layout algorithm for layered directed acyclic graphs.

This layout algorithm is designed for directed acyclic graphs where each vertex is assigned to a layer. Layers are indexed from zero, and vertices of the same layer will be placed on the same horizontal line. The X coordinates of vertices within each layer are decided by the heuristic proposed by Sugiyama et al to minimize edge crossings.

You can also try to lay out undirected graphs, graphs containing cycles, or graphs without an a priori layered assignment with this algorithm. igraph will try to eliminate cycles and assign vertices to layers, but there is no guarantee on the quality of the layout in such cases.

The Sugiyama layout may introduce "bends" on the edges in order to obtain a visually more pleasing layout. This is achieved by adding dummy nodes to edges spanning more than one layer. The resulting layout assigns coordinates not only to the nodes of the original graph but also to the dummy nodes. The

layout algorithm will also return the extended graph with the dummy nodes. An edge in the original graph may either be mapped to a single edge in the extended graph or a *path* that starts and ends in the original source and target vertex and passes through multiple dummy vertices. In such cases, the user may also request the mapping of the edges of the extended graph back to the edges of the original graph.

For more details, see K. Sugiyama, S. Tagawa and M. Toda, "Methods for Visual Understanding of Hierarchical Systems". IEEE Transactions on Systems, Man and Cybernetics 11(2):109-125, 1981.

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and

will be resized as needed. The first |V| rows of the layout will contain the coordinates of the original graph, the remaining rows contain the positions of the dummy nodes. Therefore, you can use the result both

with graph or with extended_graph.

extended_graph: Pointer to an uninitialized graph object or NULL. The extended graph

with the added dummy nodes will be returned here. In this graph, each edge points downwards to lower layers, spans exactly one layer and the

first |V| vertices coincide with the vertices of the original graph.

extd_to_orig_eids: Pointer to a vector or NULL. If not NULL, the mapping from the edge

IDs of the extended graph back to the edge IDs of the original graph will

be stored here.

layers: The layer index for each vertex or NULL if the layers should be deter-

mined automatically by igraph.

hgap: The preferred minimum horizontal gap between vertices in the same lay-

er.

vgap: The distance between layers.

maxiter: Maximum number of iterations in the crossing minimization stage. 100

is a reasonable default; if you feel that you have too many edge crossings,

increase this.

weights: Weights of the edges. These are used only if the graph contains cycles;

igraph will tend to reverse edges with smaller weights when breaking

the cycles.

igraph_layout_umap — Layout using Uniform Manifold Approximation and Projection for Dimension Reduction.

Warning

This function is experimental and its signature is not considered final yet. We reserve the right to change the function signature without changing the major version of igraph. Use it at your own risk.

UMAP is mostly used to embed high-dimensional vectors in a low-dimensional space (most commonly by far, 2D). The algorithm is probabilistic and introduces nonlinearities, unlike e.g. PCA and similar to T-distributed Stochastic Neighbor Embedding (t-SNE). Nonlinearity helps "cluster" very similar vectors together without imposing a global geometry on the embedded space (e.g. a rigid rotation + compression in PCA).

However, UMAP uses a graph with distances associated to the edges as a key intermediate representation of the high-dimensional space, so it is also useful as a general graph layouting algorithm, hence its inclusion in igraph.

Importantly, the edge-associated distances are derived from a similarity metric between the high-dimensional vectors, often Pearson correlation:

```
corr(v1, v2) = v1 \cdot v2 / [ sqrt(v1 \cdot v1) * sqrt(v2 \cdot v2) ],
```

where . denotes the dot product. In this case, the associated distance is usually defined as:

```
d(v1, v2) = 1 - corr(v1, v2)
```

This implementation can also work with unweighted similarity graphs, in which case the distance parameter should be a null pointer and all edges beget a similarity score of 1 (a distance of 0).

While all similarity graphs are theoretically embeddable, UMAP's stochastic gradient descent approach really shines when the graph is sparse. In practice, most people feed a k-nearest neighbor (either computed exactly or approximated) similarity graph with some additional cutoff to exclude "quasi-neighbors" that lie beyond a certain distance (e.g. correlation less than 0.2).

Therefore, if you are trying to use this function to embed high-dimensional vectors, the steps are:

- 1. Compute a sparse similarity graph (either exact or approximate) from your vectors, weighted or unweighted. If unsure, compute a k-nearest neighbors graph.
- 2. If you keep the weights, convert them into distances or store them as a "weight" edge attribute and use a null pointer for the distances. If using similarity weights instead of distances, make sure they do not exceed 1.
- 3. Feed the graph (and distances, if you have them) into this function.

Note: Step 1 above involves deciding if two high-dimensional vectors "look similar" which, because of the curse of dimensionality, is in many cases a highly subjective and potentially controversial operation: thread with care and at your own risk. Two high-dimensional vectors might look similar or extremely different depending on the point of view/angle, and there are a lot of viewpoints when the dimensionality ramps up.

References:

Leland McInnes, John Healy, and James Melville. https://arxiv.org/abs/1802.03426

Arguments:

graph: Pointer to the similarity graph to find a layout for (i.e. to embed).

res: Pointer to the n by 2 matrix where the layout coordinates will be stored.

use seed: Logical, if true the supplied values in the res argument are used as an initial

layout, if false a random initial layout is used.

Generating layouts for graph drawing

distances: Pointer to a vector of edge lengths. Similarity graphs for UMAP are often

originally meant in terms of similarity weights (e.g. correlation between high-dimensional vectors) and converted into distances by crude dist = 1 - corr. That is fine here too. If this argument NULL, all lengths are assumed

to be the same.

min_dist: A fudge parameter that decides how close two unconnected vertices can be in

the embedding before feeling a repulsive force. It should be positive. Typical-

ly, 0.01 is a good number.

epochs: Number of iterations of the main stochastic gradient descent loop on the cross-

entropy. Usually, 500 epochs can be used if the graph is the graph is small

(less than 50000 edges), 50 epochs are used for larger graphs.

sampling_prob: The fraction of vertices moved at each iteration of the stochastic gradient de-

scent (epoch). At fixed number of epochs, a higher fraction makes the algorithm slower. Vice versa, a too low number will converge very slowly, possi-

bly too slowly.

Returns:

Error code.

3D layout generators

igraph_layout_random_3d — Places the vertices uniform randomly in a cube.

```
igraph_error_t igraph_layout_random_3d(const igraph_t *graph, igraph_matrix_t *
```

Vertex coordinates range from -1 to 1, and are placed in 3 columns of a matrix, with a row for each vertex.

Arguments:

graph: The graph to place.

res: Pointer to an initialized matrix object. It will be resized to hold the result.

Returns:

Error code. The current implementation always returns with success.

Added in version 0.2.

Time complexity: O(|V|), the number of vertices.

igraph_layout_sphere — Places vertices (more or less) uniformly on a sphere.

igraph_error_t igraph_layout_sphere(const igraph_t *graph, igraph_matrix_t *res

The algorithm was described in the following paper:

Distributing many points on a sphere by E.B. Saff and A.B.J. Kuijlaars, *Mathematical Intelligencer* 19.1 (1997) 5--11. https://doi.org/10.1007/BF03024331

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized as

needed.

Returns:

Error code. The current implementation always returns with success.

Added in version 0.2.

Time complexity: O(|V|), the number of vertices in the graph.

igraph_layout_grid_3d — Places the vertices on a regular grid in the 3D space.

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be resized

as needed.

width: The number of vertices in a single row of the grid. When zero or negative, the width is

determined automatically.

height: The number of vertices in a single column of the grid. When zero or negative, the height

is determined automatically.

Returns:

Error code. The current implementation always returns with success.

Time complexity: O(|V|), the number of vertices.

igraph_layout_fruchterman_reingold_3d — 3D Fruchterman-Reingold algorithm.

```
igraph_real_t start_temp,
const igraph_vector_t *weight,
const igraph_vector_t *minx,
const igraph_vector_t *maxx,
const igraph_vector_t *miny,
const igraph_vector_t *maxy,
const igraph_vector_t *minz,
const igraph_vector_t *minz,
```

This is the 3D version of the force based Fruchterman-Reingold layout. See igraph_layout_fruchterman_reingold() for the 2D version.

Arguments:

graph: Pointer to an initialized graph object.

res: Pointer to an initialized matrix object. This will contain the result and will be re-

sized as needed.

use_seed: Logical, if true the supplied values in the res argument are used as an initial layout,

if false a random initial layout is used.

niter: The number of iterations to do. A reasonable default value is 500.

start_temp: Start temperature. This is the maximum amount of movement alloved along one

axis, within one step, for a vertex. Currently it is decreased linearly to zero during

the iteration.

weight: Pointer to a vector containing edge weights, the attraction along the edges will be

multiplied by these. Weights must be positive. It will be ignored if it is a null-

pointer

minx: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives

the minimum "x" coordinate for every vertex.

maxx: Same as minx, but the maximum "x" coordinates.

miny: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives

the minimum "y" coordinate for every vertex.

maxy: Same as miny, but the maximum "y" coordinates.

minz: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives

the minimum "z" coordinate for every vertex.

maxz: Same as minz, but the maximum "z" coordinates.

Returns:

Error code.

Added in version 0.2.

Time complexity: $O(|V|^2)$ in each iteration, |V| is the number of vertices in the graph.

igraph_layout_kamada_kawai_3d — 3D version of the Kamada-Kawai layout generator.

Generating layouts for graph drawing

This is the 3D version of igraph_layout_kamada_kawai(). See the documentation of that function for more information.

This layout algorithm is not suitable for large graphs. The memory requirements are of the order $O(|V|^2)$.

Arguments:

graph: A graph object.

res: Pointer to an initialized matrix object. This will contain the result (x-, y- and z-posi-

tions in columns one through three) and will be resized if needed.

use_seed: Boolean, whether to use the values supplied in the res argument as the initial con-

figuration. If zero and there are any limits on the z, y or z coordinates, then a random initial configuration is used. Otherwise the vertices are placed uniformly on a sphere

of radius 1 as the initial configuration.

maxiter: The maximum number of iterations to perform. A reasonable default value is at least

ten (or more) times the number of vertices.

epsilon: Stop the iteration, if the maximum delta value of the algorithm is smaller than still. It

is safe to leave it at zero, and then maxiter iterations are performed.

kkconst: The Kamada-Kawai vertex attraction constant. Typical value: number of vertices.

weights: Edge weights, larger values will result longer edges. Weights must be positive. Pass

NULL to assume unit weights for all edges.

minx: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the

minimum "x" coordinate for every vertex.

maxx: Same as minx, but the maximum "x" coordinates.

miny: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the

minimum "y" coordinate for every vertex.

maxy: Same as miny, but the maximum "y" coordinates.

minz: Pointer to a vector, or a NULL pointer. If not a NULL pointer then the vector gives the

minimum "z" coordinate for every vertex.

maxz: Same as minz, but the maximum "z" coordinates.

Returns:

Error code

Time complexity: O(|V|) for each iteration, after an $O(|V|^2 \log |V|)$ initialization step. |V| is the number of vertices in the graph.

igraph_layout_umap_3d — 3D layout using UMAP.

Warning

This function is experimental and its signature is not considered final yet. We reserve the right to change the function signature without changing the major version of igraph. Use it at your own risk.

This is the 3D version of the UMAP algorithm (see igraph_layout_umap() for the 2D version).

Arguments:

graph: Pointer to the similarity graph to find a layout for (i.e. to embed).

res: Pointer to the n by 3 matrix where the layout coordinates will be stored.

use_seed: Logical, if true the supplied values in the res argument are used as an initial

layout, if false a random initial layout is used.

distances: Pointer to a vector of edge lengths. Similarity graphs for UMAP are often

originally meant in terms of similarity weights (e.g. correlation between high-dimensional vectors) and converted into distances by crude dist = 1 - corr. That is fine here too. If this argument is NULL, all lengths are assumed

to be the same.

min_dist: A fudge parameter that decides how close two unconnected vertices can be in

the embedding before feeling a repulsive force. It should be positive. Typical-

ly, 0.01 is a good number.

epochs: Number of iterations of the main stochastic gradient descent loop on the cross-

entropy. Usually, 500 epochs can be used if the graph is the graph is small

(less than 50000 edges), 50 epochs are used for larger graphs.

sampling_prob: The fraction of vertices moved at each iteration of the stochastic gradient de-

scent (epoch). At fixed number of epochs, a higher fraction makes the algorithm slower. Vice versa, a too low number will converge very slowly, possi-

bly too slowly.

Returns:

Error code.

Merging layouts

igraph_layout_merge_dla — Merges multiple layouts by using a DLA algorithm.

Generating layouts for graph drawing

```
igraph_error_t igraph_layout_merge_dla(
    const igraph_vector_ptr_t *thegraphs, const igraph_matrix_list_t *coords,
    igraph_matrix_t *res
);
```

Warning

This function is experimental and its signature is not considered final yet. We reserve the right to change the function signature without changing the major version of igraph. Use it at your own risk.

First each layout is covered by a circle. Then the layout of the largest graph is placed at the origin. Then the other layouts are placed by the DLA algorithm, larger ones first and smaller ones last.

Arguments:

thegraphs: Pointer vector containing the graph objects of which the layouts will be merged.

coords: List of matrices with the 2D layouts of the graphs in thegraphs.

res: Pointer to an initialized matrix object, the result will be stored here. It will be resized

if needed.

Returns:

Error code.

Added in version 0.2.

Time complexity: TODO.

Chapter 21. Reading and writing graphs from and to files

These functions can write a graph to a file, or read a graph from a file.

They assume that the current locale uses a decimal point and not a decimal comma. See igraph_enter_safelocale() and igraph_exit_safelocale() for more information.

Note that as **igraph** uses the traditional C streams, it is possible to read/write files from/to memory, at least on GNU operating systems supporting "non-standard" streams.

Simple edge list and similar formats

igraph_read_graph_edgelist — Reads an edge list from a file and creates a graph.

This format is simply a series of an even number of non-negative integers separated by whitespace. The integers represent vertex IDs. Placing each edge (i.e. pair of integers) on a separate line is not required, but it is recommended for readability. Edges of directed graphs are assumed to be in "from, to" order.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: Pointer to a stream, it should be readable.

n: The number of vertices in the graph. If smaller than the largest integer in the file it

will be ignored. It is thus safe to supply zero here.

directed: Logical, if true the graph is directed, if false it will be undirected.

Returns:

Error code: IGRAPH_PARSEERROR: if there is a problem reading the file, or the file is syntactically incorrect.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges. It is assumed that reading an integer requires O(1) time.

igraph_write_graph_edgelist — Writes the edge list of a graph to a file.

igraph_error_t igraph_write_graph_edgelist(const igraph_t *graph, FILE *outstre

Edges are represented as pairs of 0-based vertex indices. One edge is written per line, separated by a single space. For directed graphs edges are written in from, to order.

Arguments:

graph: The graph object to write.

outstream: Pointer to a stream, it should be writable.

Returns:

Error code: IGRAPH_EFILE if there is an error writing the file.

Time complexity: O(|E|), the number of edges in the graph. It is assumed that writing an integer to the file requires O(1) time.

igraph_read_graph_ncol — Reads an .ncol file used by LGL.

Also useful for creating graphs from "named" (and optionally weighted) edge lists.

This format is used by the Large Graph Layout program (http://lgl.sourceforge.net), and it is simply a symbolic weighted edge list. It is a simple text file with one edge per line. An edge is defined by two symbolic vertex names separated by whitespace. The vertex names themselves cannot contain whitespace. They may be followed by an optional number, the weight of the edge; the number can be negative and can be in scientific notation. If there is no weight specified to an edge it is assumed to be zero.

The resulting graph is always undirected. LGL cannot deal with files which contain multiple or loop edges, this is however not checked here, as **igraph** is happy with these.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: Pointer to a stream, it should be readable.

predefnames: Pointer to the symbolic names of the vertices in the file. If NULL is given here

then vertex IDs will be assigned to vertex names in the order of their appearance in the .ncol file. If it is not NULL and some unknown vertex names are found

in the .ncol file then new vertex ids will be assigned to them.

names: Logical value, if true the symbolic names of the vertices will be added to the

graph as a vertex attribute called "name".

weights: Whether to add the weights of the edges to the graph as an edge attribute

called "weight". IGRAPH_ADD_WEIGHTS_YES adds the weights (even if they are not present in the file, in this case they are assumed to be zero). IGRAPH_ADD_WEIGHTS_NO does not add any edge attribute. IGRAPH_ADD_WEIGHTS_IF_PRESENT adds the attribute if and only if there is at least one

explicit edge weight in the input file.

directed: Whether to create a directed graph. As this format was originally used only for

undirected graphs there is no information in the file about the directedness of the graph. Set this parameter to ${\tt IGRAPH_DIRECTED}$ or ${\tt IGRAPH_UNDIRECTED}$

to create a directed or undirected graph.

Returns:

Error code: IGRAPH_PARSEERROR: if there is a problem reading the file, or the file is syntactically incorrect.

Time complexity: O(|V|+|E|log(|V|)) if we neglect the time required by the parsing. As usual |V| is the number of vertices, while |E| is the number of edges.

See also:

igraph_read_graph_lgl(), igraph_write_graph_ncol()

igraph_write_graph_ncol — Writes the graph to a file in .ncol format.

.ncol is a format used by LGL, see igraph_read_graph_ncol() for details.

Note that having multiple or loop edges in an .ncol file breaks the LGL software but **igraph** does not check for this condition.

This format cannot represent zero-degree vertices.

Arguments:

graph: The graph to write.

outstream: The stream object to write to, it should be writable.

names: The name of a string vertex attribute, if symbolic names are to be written to the file.

Supply NULL to write vertex ids instead.

weights: The name of a numerical edge attribute, which will be written as weights to the file.

Supply NULL to skip writing edge weights.

Returns:

Error code: IGRAPH_EFILE if there is an error writing the file.

Time complexity: O(|E|), the number of edges. All file operations are expected to have time complexity O(1).

See also:

```
igraph_read_graph_ncol(), igraph_write_graph_lgl()
```

igraph_read_graph_lgl — Reads a graph from an .lgl file.

The .lgl format is used by the Large Graph Layout visualization software (http://lgl.sourceforge.net), it can describe undirected optionally weighted graphs. From the LGL manual:

The second format is the LGL file format (.lgl file suffix). This is yet another graph file format that tries to be as stingy as possible with space, yet keeping the edge file in a human readable (not binary) format. The format itself is like the following:

```
# vertex1name
vertex2name [optionalWeight]
vertex3name [optionalWeight]
```

Here, the first vertex of an edge is preceded with a pound sign '#'. Then each vertex that shares an edge with that vertex is listed one per line on subsequent lines.

LGL cannot handle loop and multiple edges or directed graphs, but in **igraph** it is not an error to have multiple and loop edges.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: A stream, it should be readable.

names: Logical value, if true the symbolic names of the vertices will be added to the graph

as a vertex attribute called "name".

weights: Whether to add the weights of the edges to the graph as an edge at-

tribute called "weight". IGRAPH_ADD_WEIGHTS_YES adds the weights (even if they are not present in the file, in this case they are assumed to be zero). IGRAPH_ADD_WEIGHTS_NO does not add any edge attribute. IGRAPH_ADD_WEIGHTS_IF_PRESENT adds the attribute if and only if there is at least one ex-

plicit edge weight in the input file.

directed: Whether to create a directed graph. As this format was originally used only for undi-

rected graphs there is no information in the file about the directedness of the graph. Set this parameter to IGRAPH_DIRECTED or IGRAPH_UNDIRECTED to create a

directed or undirected graph.

Returns:

Error code: IGRAPH_PARSEERROR: if there is a problem reading the file, or the file is syntactically incorrect.

Time complexity: O(|V|+|E|log(|V|)) if we neglect the time required by the parsing. As usual |V| is the number of vertices, while |E| is the number of edges.

See also:

```
igraph_read_graph_ncol(), igraph_write_graph_lgl()
```

Example 21.1. File examples/simple/igraph_read_graph_lgl.c

igraph_write_graph_lgl — Writes the graph to a file in .lgl format.

.lgl is a format used by LGL, see igraph_read_graph_lgl() for details.

Note that having multiple or loop edges in an .lgl file breaks the LGL software but **igraph** does not check for this condition.

Arguments:

graph: The graph to write.

outstream: The stream object to write to, it should be writable.

names: The name of a string vertex attribute, if symbolic names are to be written to the file.

Supply NULL to write vertex ids instead.

weights: The name of a numerical edge attribute, which will be written as weights to the file.

Supply NULL to skip writing edge weights.

isolates: Logical, if true isolated vertices are also written to the file. If false they will

be omitted.

Returns:

Error code: IGRAPH_EFILE if there is an error writing the file.

Time complexity: O(|E|), the number of edges if *isolates* is false, O(|V|+|E|) otherwise. All file operations are expected to have time complexity O(1).

See also:

```
igraph_read_graph_lgl(), igraph_write_graph_ncol()
```

Example 21.2. File examples/simple/igraph_write_graph_lgl.c

igraph_read_graph_dimacs_flow — Read a graph in DIMACS format.

This function reads the DIMACS file format, more specifically the version for network flow problems, see the files at ftp://dimacs.rutgers.edu/pub/netflow/general-info/

This is a line-oriented text file (ASCII) format. The first character of each line defines the type of the line. If the first character is c the line is a comment line and it is ignored. There is one problem line (p in the file, it must appear before any node and arc descriptor lines. The problem line has three fields separated by spaces: the problem type (min, max or asn), the number of vertices and number of edges in the graph. Exactly two node identification lines are expected (n), one for the source, one for the target vertex. These have two fields: the id of the vertex and the type of the vertex, either s (=source) or t (=target). Arc lines start with a and have three fields: the source vertex, the target vertex and the edge capacity.

Vertex IDs are numbered from 1.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: The file to read from.

source: Pointer to an integer, the id of the source node will be stored here. (The igraph vertex

ID, which is one less than the actual number in the file.) It is ignored if NULL.

target: Pointer to an integer, the (igraph) id of the target node will be stored here. It is ignored

if NULL.

capacity: Pointer to an initialized vector, the capacity of the edges will be stored here if not

NULL.

directed: Boolean, whether to create a directed graph.

Returns:

Error code.

Time complexity: O(|V|+|E|+c), the number of vertices plus the number of edges, plus the size of the file in characters.

See also:

```
igraph_write_graph_dimacs()
```

igraph_write_graph_dimacs_flow — Write a graph in DIMACS format.

This function writes a graph to an output stream in DIMACS format, describing a maximum flow problem. See ftp://dimacs.rutgers.edu/pub/netflow/general-info/

This file format is discussed in the documentation of igraph_read_graph_dimacs_flow(), see that for more information.

Arguments:

graph: The graph to write to the stream.

outstream: The stream.

source: Integer, the id of the source vertex for the maximum flow.

target: Integer, the id of the target vertex.

capacity: Pointer to an initialized vector containing the edge capacity values.

Returns:

Error code.

Time complexity: O(|E|), the number of edges in the graph.

See also:

igraph_read_graph_dimacs_flow()

Binary formats

igraph_read_graph_graphdb — Read a graph in the binary graph database format.

This is a binary format, used in the graph database for isomorphism testing. From the (now defunct) graph database homepage:

The graphs are stored in a compact binary format, one graph per file. The file is composed of 16 bit words, which are represented using the so-called little-endian convention, i.e. the least significant byte of the word is stored first.

Then, for each node, the file contains the list of edges coming out of the node itself. The list is represented by a word encoding its length, followed by a word for each edge, representing the destination node of the edge. Node numeration is 0-based, so the first node of the graph has index 0.

Only unlabelled graphs are implemented.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: The stream to read from.

directed: Logical scalar, whether to create a directed graph.

Returns:

Error code.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

Example 21.3. File examples/simple/igraph_read_graph_graphdb.c

GraphML format

igraph_read_graph_graphml — Reads a graph from a GraphML file.

igraph_error_t igraph_read_graph_graphml(igraph_t *graph, FILE *instream, igraph_

GraphML is an XML-based file format for representing various types of graphs. Currently only the most basic import functionality is implemented in igraph: it can read GraphML files without nested graphs and hyperedges. Attributes of the graph are loaded only if an attribute interface is attached, see igraph_set_attribute_table(). String attribute values are returned in UTF-8 encoding.

Graph attribute names are taken from the attr.name attributes of the key tags in the GraphML file. Since attr.name is not mandatory, igraph will fall back to the id attribute of the key tag if attr.name is missing.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: A stream, it should be readable.

index: If the GraphML file contains more than one graph, the one specified by this index will

be loaded. Indices start from zero, so supply zero here if your GraphML file contains

only a single graph.

Returns:

Error code: IGRAPH_PARSEERROR: if there is a problem reading the file, or the file is syntactically incorrect. IGRAPH_UNIMPLEMENTED: the GraphML functionality was disabled at compile-time

Example 21.4. File examples/simple/graphml.c

igraph_write_graph_graphml — Writes the graph to a file in GraphML format.

GraphML is an XML-based file format for representing various types of graphs. See the GraphML Primer (http://graphml.graphdrawing.org/primer/graphml-primer.html) for detailed format description.

When a numerical attribute value is NaN, it will be omitted from the file.

Arguments:

graph: The graph to write.

outstream: The stream object to write to, it should be writable.

prefixattr: Logical value, whether to put a prefix in front of the attribute names to ensure

uniqueness if the graph has vertex and edge (or graph) attributes with the same

name.

Returns:

Error code: IGRAPH_EFILE if there is an error writing the file.

Time complexity: O(|V|+|E|) otherwise. All file operations are expected to have time complexity O(1).

Example 21.5. File examples/simple/graphml.c

GML format

igraph_read_graph_gml — Read a graph in GML format.

```
igraph_error_t igraph_read_graph_gml(igraph_t *graph, FILE *instream);
```

GML is a simple textual format, see https://web.archive.org/web/20190207140002/http://www.fim.u-ni-passau.de/index.php?id=17297%26L=1 for details.

Although all syntactically correct GML can be parsed, we implement only a subset of this format. Some attributes might be ignored. Here is a list of all the differences:

- 1. Only attributes with a simple type are used: integer, real or string. If an attribute is composite, i.e. an array or a record, then it is ignored. When some values of the attribute are simple and some compound, the compositve ones are replaced with a default value (NaN for numeric, " " for string).
- 2. comment fields are not ignored. They are treated as any other field and converted to attributes.
- 3. Top level attributes except for Version and the first graph attribute are completely ignored.
- 4. There is no maximum line length or maximum keyword length.
- 5. Only the quot, amp, apos, lt and gt character entities are supported. Any other entity is passed through unchanged by the reader after issuing a warning, and is expected to be decoded by the user.
- 6. We allow inf, -inf and nan (not a number) as a real number. This is case insensitive, so nan, NaN and NAN are equivalent.

Please contact us if you cannot live with these limitations of the GML parser.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: The stream to read the GML file from.

Returns:

Error code.

Time complexity: should be proportional to the length of the file.

See also:

```
igraph_read_graph_graphml() for a more modern format,
igraph_write_graph_gml() for writing GML files.
```

Example 21.6. File examples/simple/gml.c

igraph_write_graph_gml — Write the graph to a stream in GML format.

GML is a quite general textual format, see https://web.archive.org/web/20190207140002/http://www.fim.uni-passau.de/index.php?id=17297%26L=1 for details.

The graph, vertex and edges attributes are written to the file as well, if they are numeric or string. Boolean attributes are converted to numeric, with 0 and 1 used for false and true, respectively. NaN values of numeric attributes are skipped, as NaN is not part of the GML specification and other software may not be able to read files containing them. This is consistent with <code>igraph_read_graph_gm-1()</code>, which produces NaN when an attribute value is missing. In contrast with NaN, infinite values are retained. Ensure that none of the numeric attributes values are infinite to produce a conformant GML file that can be read by other software.

As igraph is more forgiving about attribute names, it might be necessary to simplify the them before writing to the GML file. This way we'll have a syntactically correct GML file. The following simple procedure is performed on each attribute name: first the alphanumeric characters are extracted, the others are ignored. Then if the first character is not a letter then the attribute name is prefixed with "igraph". Note that this might result identical names for two attributes, igraph does not check this.

The "id" vertex attribute is treated specially. If the *id* argument is not NULL then it should be a numeric vector with the vertex IDs and the "id" vertex attribute is ignored (if there is one). If *id* is NULL and there is a numeric "id" vertex attribute, it will be used instead. If ids are not specified in either way then the regular igraph vertex IDs are used. If some of the supplied id values are invalid (non-integer or NaN), all supplied id are ignored and igraph vertex IDs are used instead.

Note that whichever way vertex IDs are specified, their uniqueness is not checked.

If the graph has edge attributes that become "source" or "target" after encoding, or the graph has an attribute that becomes "directed", they will be ignored with a warning. GML uses these attributes to specify the edge endpoints, and the graph directedness, so we cannot write them to the file. Rename them before calling this function if you want to preserve them.

Arguments:

graph: The graph to write to the stream.

outstream: The stream to write the file to.

0

options: Set of |-combinable boolean flags for writing the GML file.

All options turned off.

IGRAPH_WRITE_GML_DE-

FAULT_SW

Default options, currently equivalent to 0.

May change in future versions.

IGRAPH_WRITE_GML_ENCODE_ONLY_QUOT_SW

Do not encode any other characters than " as entities. Specifically, this option prevents the encoding of &. Useful when re-exporting a graph that was read from a GML file in which igraph could not interpret all entities, and thus passed them through without decoding.

id: Either NULL or a numeric vector with the vertex IDs. See details above.

creator: An optional string to write to the stream in the creator line. If NULL, the igraph

version with the current date and time is added. If " ", the creator line is omitted.

Otherwise, the supplied string is used verbatim.

Returns:

Error code.

Time complexity: should be proportional to the number of characters written to the file.

See also:

 $\verb|igraph_read_graph_gml(|)| for reading GML files, \verb|igraph_read_graph_graphml(|)| for a more modern format.$

Example 21.7. File examples/simple/gml.c

Pajek format

igraph_read_graph_pajek — Reads a file in Pajek format.

igraph_error_t igraph_read_graph_pajek(igraph_t *graph, FILE *instream);

Arguments:

graph: Pointer to an uninitialized graph object.

file: An already opened file handler.

Returns:

Error code.

Only a subset of the Pajek format is implemented. This is partially because this format is not very well documented, but also because **igraph** does not support some Pajek features, like multigraphs.

Starting from version 0.6.1 igraph reads bipartite (two-mode) graphs from Pajek files and add the type vertex attribute for them. Warnings are given for invalid edges, i.e. edges connecting vertices of the same type.

The list of the current limitations:

- 1. Only .net files are supported, Pajek project files (.paj) are not. These might be supported in the future if there is need for it.
- 2. Time events networks are not supported.
- 3. Hypergraphs (i.e. graphs with non-binary edges) are not supported.
- 4. Graphs with both directed and non-directed edges are not supported, are they cannot be represented in **igraph**.
- 5. Only Pajek networks are supported, permutations, hierarchies, clusters and vectors are not.
- 6. Graphs with multiple edge sets are not supported.

If there are attribute handlers installed, **igraph** also reads the vertex and edge attributes from the file. Most attributes are renamed to be more informative: color instead of c, xfact instead of x_fact, yfact instead of y_fact, labeldist instead of lr, labeldegree2 instead of lphi, framewidth instead of bw, fontsize instead of fos, rotation instead of phi, radius instead of r, diamondratio instead of q, labeldegree instead of la, vertexsize instead of size, color instead of ic, framecolor instead of bc, labelcolor instead of lc, these belong to vertices.

Edge attributes are also renamed, s to arrowsize, w to edgewidth, h1 to hook1, h2 to hook2, a1 to angle1, a2 to angle2, k1 to velocity1, k2 to velocity2, ap to arrowpos, lp to labelpos, lr to labelangle, lphi to labelangle2, la to labeldegree, fos to font-size, a to arrowtype, p to linepattern, l to label, lc to labelcolor, c to color.

In addition the following vertex attributes might be added: id if there are vertex IDs in the file, x and y or x and y and z if there are vertex coordinates in the file.

The weight edge attribute might be added if there are edge weights present.

See the pajek homepage: http://vlado.fmf.uni-lj.si/pub/networks/pajek/ for more info on Pajek and the Pajek manual: http://vlado.fmf.uni-lj.si/pub/networks/pajek/doc/pajekman.pdf for information on the Pajek file format.

Time complexity: O(|V|+|E|+|A|), |V| is the number of vertices, |E| the number of edges, |A| the number of attributes (vertex + edge) in the graph if there are attribute handlers installed.

See also:

```
igraph_write_graph_pajek() for writing Pajek files,
igraph_read_graph_graphml() for reading GraphML files.
```

Example 21.8. File examples/simple/foreign.c

igraph_write_graph_pajek — Writes a graph to a file in Pajek format.

```
igraph_error_t igraph_write_graph_pajek(const igraph_t *graph, FILE *outstream)
```

The Pajek vertex and edge parameters (like color) are determined by the attributes of the vertices and edges, of course this requires an attribute handler to be installed. The names of the corresponding

vertex and edge attributes are listed at igraph_read_graph_pajek(), e.g. the color vertex attributes determines the color (c in Pajek) parameter.

As of version 0.6.1 igraph writes bipartite graphs into Pajek files correctly, i.e. they will be also bipartite when read into Pajek. As Pajek is less flexible for bipartite graphs (the numeric IDs of the vertices must be sorted according to vertex type), igraph might need to reorder the vertices when writing a bipartite Pajek file. This effectively means that numeric vertex IDs usually change when a bipartite graph is written to a Pajek file, and then read back into igraph.

Early versions of Pajek supported only Windows-style line endings in Pajek files, but recent versions support both Windows and Unix line endings. igraph therefore uses the platform-native line endings when the input file is opened in text mode, and uses Unix-style line endings when the input file is opened in binary mode. If you are using an old version of Pajek, you are on Unix and you are having problems reading files written by igraph on a Windows machine, convert the line endings manually with a text editor or with unix2dos or iconv from the command line).

Arguments:

graph: The graph object to write.

outstream: The file to write to. It should be opened and writable. Make sure that you open the

file in binary format if you use MS Windows, otherwise end of line characters will be messed up. (igraph will be able to read back these messed up files, but Pajek won't.)

Returns:

Error code.

Time complexity: O(|V|+|E|+|A|), |V| is the number of vertices, |E| is the number of edges, |A| the number of attributes (vertex + edge) in the graph if there are attribute handlers installed.

See also:

igraph_read_graph_pajek() for reading Pajek graphs,
igraph_write_graph_graphml() for writing a graph in GraphML format, this suites
igraph graphs better.

Example 21.9. File examples/simple/igraph_write_graph_pajek.c

UCINET's DL file format

igraph_read_graph_d1 — Reads a file in the DL format of UCINET.

This is a simple textual file format used by UCINET. See http://www.analytictech.com/net-works/dataentry.htm for examples. All the forms described here are supported by igraph. Vertex names and edge weights are also supported and they are added as attributes. (If an attribute handler is attached.)

Note the specification does not mention whether the format is case sensitive or not. For igraph DL files are case sensitive, i.e. Larry and larry are not the same.

Arguments:

graph: Pointer to an uninitialized graph object.

instream: The stream to read the DL file from.

directed: Logical scalar, whether to create a directed file.

Returns:

Error code.

Time complexity: linear in terms of the number of edges and vertices, except for the matrix format, which is quadratic in the number of vertices.

Example 21.10. File examples/simple/igraph_read_graph_dl.c

Graphviz format

igraph_write_graph_dot — Write the graph to a stream in DOT format.

```
igraph_error_t igraph_write_graph_dot(const igraph_t *graph, FILE* outstream);
```

DOT is the format used by the widely known GraphViz software, see http://www.graphviz.org for details. The grammar of the DOT format can be found here: http://www.graphviz.org/doc/info/lang.html

This is only a preliminary implementation, no visualization information is written.

Arguments:

graph: The graph to write to the stream.

outstream: The stream to write the file to.

Time complexity: should be proportional to the number of characters written to the file.

See also:

igraph_write_graph_graphml() for a more modern format.

Example 21.11. File examples/simple/dot.c

Convenience functions for locale change

igraph_enter_safelocale — Temporarily set the C locale.

```
igraph_error_t igraph_enter_safelocale(igraph_safelocale_t *loc);
```

Warning

This function is experimental and its signature is not considered final yet. We reserve the right to change the function signature without changing the major version of igraph. Use it at your own risk.

igraph's foreign format readers and writers require a locale that uses a decimal point instead of a decimal comma. This is a convenience function that temporarily sets the C locale so that readers and writer would work correctly. It *must* be paired with a call to <code>igraph_exit_safelocale()</code>, otherwise a memory leak will occur.

This function tries to set the locale for the current thread only on a best-effort basis. Restricting the locale change to a single thread is not supported on all platforms. In these cases, this function falls back to using the standard setlocale() function, which affects the entire process and is not safe to use from concurrent threads.

It is generally recommended to run igraph within a thread that has been permanently set to the C locale using system-specific means. This is a convenience function for situations when this is not easily possible because the programmer is not in control of the process, such as when developing plugins/extensions. Note that processes start up in the C locale by default, thus nothing needs to be done unless the locale has been changed away from the default.

Arguments:

loc: Pointer to a variable of type igraph_safelocale_t. The current locale will be stored here, so that it can be restored using igraph_exit_safelocale().

Returns:

Error code.

Example 21.12. File examples/simple/safelocale.c

igraph_exit_safelocale — Temporarily set the C locale.

void igraph_exit_safelocale(igraph_safelocale_t *loc);

Warning

This function is experimental and its signature is not considered final yet. We reserve the right to change the function signature without changing the major version of igraph. Use it at your own risk.

Restores a locale saved by igraph_enter_safelocale() and deallocates all associated data. This function *must* be paired with a call to igraph_enter_safelocale().

Arguments:

loc: A variable of type igraph_safelocale_t, originally set by igraph_enter_safelocale().

Deprecated functions

igraph_read_graph_dimacs — Read a graph in DI-MACS format (deprecated alias).

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_read_graph_dimacs_flow() instead.

igraph_write_graph_dimacs — Write a graph in DI-MACS format (deprecated alias).

Warning

Deprecated since version 0.10.0. Please do not use this function in new code; use igraph_write_graph_dimacs_flow() instead.

Chapter 22. Maximum flows, minimum cuts and related measures

Maximum flows

igraph_maxflow — Maximum network flow between a pair of vertices.

This function implements the Goldberg-Tarjan algorithm for calculating value of the maximum flow in a directed or undirected graph. The algorithm was given in Andrew V. Goldberg, Robert E. Tarjan: A New Approach to the Maximum-Flow Problem, Journal of the ACM, 35(4), 921-940, 1988.

The input of the function is a graph, a vector of real numbers giving the capacity of the edges and two vertices of the graph, the source and the target. A flow is a function assigning positive real numbers to the edges and satisfying two requirements: (1) the flow value is less than the capacity of the edge and (2) at each vertex except the source and the target, the incoming flow (i.e. the sum of the flow on the incoming edges) is the same as the outgoing flow (i.e. the sum of the flow on the outgoing edges). The value of the flow is the incoming flow at the target vertex. The maximum flow is the flow with the maximum value.

Arguments:

graph: The input graph, either directed or undirected.

value: Pointer to a real number, the value of the maximum will be placed here, unless it

is a null pointer.

£10w: If not a null pointer, then it must be a pointer to an initialized vector. The vector

will be resized, and the flow on each edge will be placed in it, in the order of the edge IDs. For undirected graphs this argument is bit trickier, since for these the flow direction is not predetermined by the edge direction. For these graphs the elements of the flow vector can be negative, this means that the flow goes from the bigger vertex ID to the smaller one. Positive values mean that the flow goes

from the smaller vertex ID to the bigger one.

cut: A null pointer or a pointer to an initialized vector. If not a null pointer, then the

minimum cut corresponding to the maximum flow is stored here, i.e. all edge IDs

that are part of the minimum cut are stored in the vector.

partition: A null pointer or a pointer to an initialized vector. If not a null pointer, then the

first partition of the minimum cut that corresponds to the maximum flow will be placed here. The first partition is always the one that contains the source vertex.

partition2: A null pointer or a pointer to an initialized vector. If not a null pointer, then the

second partition of the minimum cut that corresponds to the maximum flow will be placed here. The second partition is always the one that contains the target vertex.

Maximum flows, minimum cuts and related measures

source: The id of the source vertex.

target: The id of the target vertex.

capacity: Vector containing the capacity of the edges. If NULL, then every edge is considered

to have capacity 1.0.

stats: Counts of the number of different operations preformed by the algorithm are stored

here.

Returns:

Error code.

Time complexity: O(|V|^3). In practice it is much faster, but i cannot prove a better lower bound for the data structure i've used. In fact, this implementation runs much faster than the hi_pr implementation discussed in B. V. Cherkassky and A. V. Goldberg: On implementing the push-relabel method for the maximum flow problem, (Algorithmica, 19:390--410, 1997) on all the graph classes i've tried.

See also:

```
igraph_mincut_value(), igraph_edge_connectivity(), igraph_ver-
tex_connectivity() for properties based on the maximum flow.
```

Example 22.1. File examples/simple/flow.c

Example 22.2. File examples/simple/flow2.c

igraph_maxflow_value — Maximum flow in a network with the push/relabel algorithm.

This function implements the Goldberg-Tarjan algorithm for calculating value of the maximum flow in a directed or undirected graph. The algorithm was given in Andrew V. Goldberg, Robert E. Tarjan: A New Approach to the Maximum-Flow Problem, Journal of the ACM, 35(4), 921-940, 1988.

The input of the function is a graph, a vector of real numbers giving the capacity of the edges and two vertices of the graph, the source and the target. A flow is a function assigning positive real numbers to the edges and satisfying two requirements: (1) the flow value is less than the capacity of the edge and (2) at each vertex except the source and the target, the incoming flow (i.e. the sum of the flow on the incoming edges) is the same as the outgoing flow (i.e. the sum of the flow on the outgoing edges). The value of the flow is the incoming flow at the target vertex. The maximum flow is the flow with the maximum value.

According to a theorem by Ford and Fulkerson (L. R. Ford Jr. and D. R. Fulkerson. Maximal flow through a network. Canadian J. Math., 8:399-404, 1956.) the maximum flow between two vertices is the same as the minimum cut between them (also called the minimum s-t cut). So igraph_st_min-cut_value() gives the same result in all cases as igraph_maxflow_value().

Note that the value of the maximum flow is the same as the minimum cut in the graph.

Arguments:

graph: The input graph, either directed or undirected.

value: Pointer to a real number, the result will be placed here.

source: The id of the source vertex.

target: The id of the target vertex.

capacity: Vector containing the capacity of the edges. If NULL, then every edge is considered

to have capacity 1.0.

stats: Counts of the number of different operations preformed by the algorithm are stored

here.

Returns:

Error code.

Time complexity: $O(|V|^3)$.

See also:

igraph_maxflow() to calculate the actual flow. igraph_mincut_value(), igraph_edge_connectivity(), igraph_vertex_connectivity() for properties based on the maximum flow.

igraph_dominator_tree — Calculates the dominator tree of a flowgraph

A flowgraph is a directed graph with a distinguished start (or root) vertex r, such that for any vertex v, there is a path from r to v. A vertex v dominates another vertex w (not equal to v), if every path from r to w contains v. Vertex v is the immediate dominator or w, v=idom(w), if v dominates w and every other dominator of w dominates v. The edges $\{(idom(w), w) | w \text{ is not } r\}$ form a directed tree, rooted at r, called the dominator tree of the graph. Vertex v dominates vertex w if and only if v is an ancestor of w in the dominator tree.

This function implements the Lengauer-Tarjan algorithm to construct the dominator tree of a directed graph. For details please see Thomas Lengauer, Robert Endre Tarjan: A fast algorithm for finding dominators in a flowgraph, ACM Transactions on Programming Languages and Systems (TOPLAS) I/1, 121--141, 1979.

Arguments:

graph: A directed graph. If it is not a flowgraph, and it contains some vertices not reachable

from the root vertex, then these vertices will be collected in the leftout vector.

root: The id of the root (or source) vertex, this will be the root of the tree.

Maximum flows, minimum cuts and related measures

dom: Pointer to an initialized vector or a null pointer. If not a null pointer, then the immediate

dominator of each vertex will be stored here. For vertices that are not reachable from

the root, -2 is stored here. For the root vertex itself, -1 is added.

domtree: Pointer to an uninitialized igraph_t, or NULL. If not a null pointer, then the dominator

tree is returned here. The graph contains the vertices that are unreachable from the root

(if any), these will be isolates.

leftout: Pointer to an initialized vector object, or NULL. If not NULL, then the IDs of the ver-

tices that are unreachable from the root vertex (and thus not part of the dominator tree)

are stored here.

mode: Constant, must be IGRAPH IN or IGRAPH OUT. If it is IGRAPH IN, then all direc-

tions are considered as opposite to the original one in the input graph.

Returns:

Error code.

Time complexity: very close to O(|E|+|V|), linear in the number of edges and vertices. More precisely, it is O(|V|+|E|alpha(|E|,|V|)), where alpha(|E|,|V|) is a functional inverse of Ackermann's function.

Example 22.3. File examples/simple/dominator_tree.c

igraph_maxflow_stats_t — A simple data type to return some statistics from the

```
typedef struct {
    igraph_integer_t nopush, norelabel, nogap, nogapnodes, nobfs;
push-relabel maximum flow solver.
```

Arguments:

nopush: The number of push operations performed.

norelabel: The number of relabel operarions performed.

nogap: The number of times the gap heuristics was used.

nogapnodes: The total number of vertices that were omitted form further calculations because

of the gap heuristics.

nobfs: The number of times the reverse BFS was run to assign good values to the height

function. This includes an initial run before the whole algorithm, so it is always

at least one.

Cuts and minimum cuts

igraph_st_mincut — Minimum cut between a source and a target vertex.

Maximum flows, minimum cuts and related measures

Finds the edge set that has the smallest total capacity among all edge sets that disconnect the source and target vertices.

The calculation is performed using maximum flow techniques, by calling igraph_maxflow().

Arguments:

graph: The input graph.

value: Pointer to a real variable, the value of the cut is stored here.

cut: Pointer to an initialized vector, the edge IDs that are included in the cut are stored

here. This argument is ignored if it is a null pointer.

partition: Pointer to an initialized vector, the vertex IDs of the vertices in the first partition of

the cut are stored here. The first partition is always the one that contains the source

vertex. This argument is ignored if it is a null pointer.

partition2: Pointer to an initialized vector, the vertex IDs of the vertices in the second partition

of the cut are stored here. The second partition is always the one that contains the

target vertex. This argument is ignored if it is a null pointer.

source: Integer, the id of the source vertex.

target: Integer, the id of the target vertex.

capacity: Vector containing the capacity of the edges. If a null pointer, then every edge is

considered to have capacity 1.0.

Returns:

Error code.

See also:

```
igraph_maxflow().
```

Time complexity: see igraph_maxflow().

igraph_st_mincut_value — The minimum s-t cut in a graph.

The minimum s-t cut in a weighted (=valued) graph is the total minimum edge weight needed to remove from the graph to eliminate all paths from a given vertex (source) to another vertex (target). Directed paths are considered in directed graphs, and undirected paths in undirected graphs.

The minimum s-t cut between two vertices is known to be same as the maximum flow between these two vertices. So this function calls <code>igraph_maxflow_value()</code> to do the calculation.

Arguments:

graph: The input graph.

value: Pointer to a real variable, the result will be stored here.

source: The id of the source vertex.

target: The id of the target vertex.

capacity: Pointer to the capacity vector, it should contain non-negative numbers and its length

should be the same the the number of edges in the graph. It can be a null pointer, then

every edge has unit capacity.

Returns:

Error code.

Time complexity: $O(|V|^3)$, see also the discussion for $igraph_maxflow_value()$, |V| is the number of vertices.

igraph_all_st_cuts — List all edge-cuts between two vertices in a directed graph

This function lists all edge-cuts between a source and a target vertex. Every cut is listed exactly once. The implemented algorithm is described in JS Provan and DR Shier: A Paradigm for listing (s,t)-cuts in graphs, Algorithmica 15, 351--372, 1996.

Arguments:

graph: The input graph, is must be directed.

cuts: An initialized list of integer vectors, the cuts are stored here. Each vector will

contain the IDs of the edges in the cut. This argument is ignored if it is a null

pointer.

partition1s: An initialized list of integer vectors, the list of vertex sets generating the actual

edge cuts are stored here. Each vector contains a set of vertex IDs. If X is such a set, then all edges going from X to the complement of X form an (s, t) edge-cut

in the graph. This argument is ignored if it is a null pointer.

source: The id of the source vertex.

target: The id of the target vertex.

Returns:

Error code.

Time complexity: O(n(|V|+|E|)), where |V| is the number of vertices, |E| is the number of edges, and n is the number of cuts.

igraph_all_st_mincuts — All minimum s-t cuts of a directed graph

This function lists all edge cuts between two vertices, in a directed graph, with minimum total capacity. Possibly, multiple cuts may have the same total capacity, although there is often only one minimum cut in weighted graphs. It is recommended to supply integer-values capacities. Otherwise, not all minimum cuts may be detected because of numerical roundoff errors. The implemented algorithm is described in JS Provan and DR Shier: A Paradigm for listing (s,t)-cuts in graphs, Algorithmica 15, 351--372, 1996.

Arguments:

graph: The input graph, it must be directed.

value: Pointer to a real number, the value of the minimum cut is stored here, unless it

is a null pointer.

cuts: An initialized pointer vector, the cuts are stored here. It is a list of pointers

to igraph_vector_int_t objects. Each vector will contain the IDs of the edges in the cut. This argument is ignored if it is a null pointer. To free all memory allocated for cuts, you need call igraph_vector_int_destroy() and then igraph_free() on each element, before destroying the pointer vector

itself.

partition1s: An initialized pointer vector, the list of vertex sets, generating the actual edge

cuts, are stored here. It is a list of pointers to igraph_vector_int_t objects. Each vector contains a set of vertex IDs. If X is such a set, then all edges going from X to the complement of X form an (s,t) edge-cut in the graph. This argument

is ignored if it is a null pointer.

source: The id of the source vertex.

target: The id of the target vertex.

capacity: Vector of edge capacities. All capacities must be strictly positive. If this is a null

pointer, then all edges are assumed to have capacity one.

Returns:

Error code.

Time complexity: O(n(|V|+|E|))+O(F), where |V| is the number of vertices, |E| is the number of edges, and n is the number of cuts; O(F) is the time complexity of the maximum flow algorithm, see $igraph_maxflow()$.

Example 22.4. File examples/simple/igraph_all_st_mincuts.c

igraph_mincut — Calculates the minimum cut in a graph.

This function calculates the minimum cut in a graph. The minimum cut is the minimum set of edges which needs to be removed to disconnect the graph. The minimum is calculated using the weights (capacity) of the edges, so the cut with the minimum total capacity is calculated.

For directed graphs an implementation based on calculating 2|V|-2 maximum flows is used. For undirected graphs we use the Stoer-Wagner algorithm, as described in M. Stoer and F. Wagner: A simple min-cut algorithm, Journal of the ACM, 44 585-591, 1997.

The first implementation of the actual cut calculation for undirected graphs was made by Gregory Benison, thanks Greg.

Arguments:

graph: The input graph.

value: Pointer to a float, the value of the cut will be stored here.

partition: Pointer to an initialized vector, the ids of the vertices in the first partition after

separating the graph will be stored here. The vector will be resized as needed. This

argument is ignored if it is a NULL pointer.

partition2: Pointer to an initialized vector the ids of the vertices in the second partition will

be stored here. The vector will be resized as needed. This argument is ignored if

it is a NULL pointer.

cut: Pointer to an initialized vector, the IDs of the edges in the cut will be stored here.

This argument is ignored if it is a NULL pointer.

capacity: A numeric vector giving the capacities of the edges. If a null pointer then all edges

have unit capacity.

Returns:

Error code.

See also:

igraph_mincut_value(), a simpler interface for calculating the value of the cut only.

Time complexity: for directed graphs it is $O(|V|^4)$, but see the remarks at $igraph_maxflow()$. For undirected graphs it is $O(|V||E|+|V|^2\log|V|)$. |V| and |E| are the number of vertices and edges respectively.

Example 22.5. File examples/simple/igraph mincut.c

igraph_mincut_value — The minimum edge cut in a graph.

The minimum edge cut in a graph is the total minimum weight of the edges needed to remove from the graph to make the graph *not* strongly connected. (If the original graph is not strongly connected then this is zero.) Note that in undirected graphs strong connectedness is the same as weak connectedness.

The minimum cut can be calculated with maximum flow techniques, although the current implementation does this only for directed graphs and a separate non-flow based implementation is used for undirected graphs. See Mechthild Stoer and Frank Wagner: A simple min-cut algorithm, Journal of the ACM 44 585--591, 1997. For directed graphs the maximum flow is calculated between a fixed vertex and all the other vertices in the graph and this is done in both directions. Then the minimum is taken to get the minimum cut.

Arguments:

graph: The input graph.

res: Pointer to a real variable, the result will be stored here.

capacity: Pointer to the capacity vector, it should contain the same number of non-negative

numbers as the number of edges in the graph. If a null pointer then all edges will have

unit capacity.

Returns:

Error code.

See also:

```
igraph_mincut(), igraph_maxflow_value(), igraph_st_mincut_value().
```

Time complexity: $O(\log(|V|)*|V|^2)$ for undirected graphs and $O(|V|^4)$ for directed graphs, but see also the discussion at the documentation of <code>igraph_maxflow_value()</code>.

igraph_gomory_hu_tree — Gomory-Hu tree of a graph.

The Gomory-Hu tree is a concise representation of the value of all the maximum flows (or minimum cuts) in a graph. The vertices of the tree correspond exactly to the vertices of the original graph in the same order. Edges of the Gomory-Hu tree are annotated by flow values. The value of the maximum flow (or minimum cut) between an arbitrary (u,v) vertex pair in the original graph is then given by the minimum flow value (i.e. edge annotation) along the shortest path between u and v in the Gomory-Hu tree.

Maximum flows, minimum cuts and related measures

This implementation uses Gusfield's algorithm to construct the Gomory-Hu tree. See the following paper for more details:

Gusfield D: Very simple methods for all pairs network flow analysis. SIAM J Comput 19(1):143-155, 1990

Arguments:

graph: The input graph.

tree: Pointer to an uninitialized graph; the result will be stored here.

flows: Pointer to an uninitialized vector; the flow values corresponding to each edge in the

Gomory-Hu tree will be returned here. You may pass a NULL pointer here if you are

not interested in the flow values.

capacity: Vector containing the capacity of the edges. If NULL, then every edge is considered

to have capacity 1.0.

Returns:

Error code.

Time complexity: $O(|V|^4)$ since it performs a max-flow calculation between vertex zero and every other vertex and max-flow is $O(|V|^3)$.

See also:

```
igraph_maxflow()
```

Connectivity

igraph_st_edge_connectivity — Edge connectivity of a pair of vertices.

The edge connectivity of two vertices (source and target) in a graph is the minimum number of edges that have to be deleted from the graph to eliminate all paths from source to target.

This function uses the maximum flow algorithm to calculate the edge connectivity.

Arguments:

graph: The input graph, it has to be directed.

res: Pointer to an integer, the result will be stored here.

source: The id of the source vertex.

target: The id of the target vertex.

Returns:

Error code.

Time complexity: $O(|V|^3)$.

See also:

igraph_maxflow_value(), igraph_edge_connectivity(), igraph_st_vertex_connectivity(),igraph_vertex_connectivity().

igraph_edge_connectivity — The minimum edge connectivity in a graph.

This is the minimum of the edge connectivity over all pairs of vertices in the graph.

The edge connectivity of a graph is the same as group adhesion as defined in Douglas R. White and Frank Harary: The cohesiveness of blocks in social networks: node connectivity and conditional density, Sociological Methodology 31:305--359, 2001.

Arguments:

graph: The input graph.

res: Pointer to an integer, the result will be stored here.

checks: Logical constant. Whether to check that the graph is connected and also the degree of

the vertices. If the graph is not (strongly) connected then the connectivity is obviously zero. Otherwise if the minimum degree is one then the edge connectivity is also one. It is a good idea to perform these checks, as they can be done quickly compared to the connectivity calculation itself. They were suggested by Peter McMahan, thanks Peter.

Returns:

Error code.

Time complexity: $O(\log(|V|)*|V|^2)$ for undirected graphs and $O(|V|^4)$ for directed graphs, but see also the discussion at the documentation of <code>igraph_maxflow_value()</code>.

See also:

igraph_st_edge_connectivity(), igraph_maxflow_value(), igraph_vertex_connectivity().

igraph_st_vertex_connectivity — The vertex connectivity of a pair of vertices.

igraph_error_t igraph_st_vertex_connectivity(const igraph_t *graph,

Maximum flows, minimum cuts and related measures

```
igraph_integer_t *res,
igraph_integer_t source,
igraph_integer_t target,
igraph_vconn_nei_t neighbors);
```

The vertex connectivity of two vertices (source and target) is the minimum number of vertices that have to be deleted to eliminate all paths from source to target. Directed paths are considered in directed graphs.

The vertex connectivity of a pair is the same as the number of different (i.e. node-independent) paths from source to target.

The current implementation uses maximum flow calculations to obtain the result.

Arguments:

graph: The input graph.

res: Pointer to an integer, the result will be stored here.

source: The id of the source vertex.

target: The id of the target vertex.

neighbors: A constant giving what to do if the two vertices are connected. Possible val-

ues: IGRAPH_VCONN_NEI_ERROR, stop with an error message, IGRAPH_V-CONN_NEGATIVE, return -1. IGRAPH_VCONN_NUMBER_OF_NODES, return the number of nodes. IGRAPH_VCONN_IGNORE, ignore the fact that the two vertices are connected and calculate the number of vertices needed to eliminate all paths except for the trivial (direct) paths between source and vertex. TODO: what about

neighbors?

Returns:

Error code.

Time complexity: $O(|V|^3)$, but see the discussion at igraph_maxflow_value().

See also:

igraph_vertex_connectivity — The vertex connectivity of a graph.

The vertex connectivity of a graph is the minimum vertex connectivity along each pairs of vertices in the graph.

The vertex connectivity of a graph is the same as group cohesion as defined in Douglas R. White and Frank Harary: The cohesiveness of blocks in social networks: node connectivity and conditional density, Sociological Methodology 31:305--359, 2001.

Arguments:

graph: The input graph.

res: Pointer to an integer, the result will be stored here.

checks: Logical constant. Whether to check that the graph is connected and also the degree of

the vertices. If the graph is not (strongly) connected then the connectivity is obviously zero. Otherwise if the minimum degree is one then the vertex connectivity is also one. It is a good idea to perform these checks, as they can be done quickly compared to the connectivity calculation itself. They were suggested by Peter McMahan, thanks Peter.

Returns:

Error code.

Time complexity: $O(|V|^5)$.

See also:

```
igraph_st_vertex_connectivity(), igraph_maxflow_value(), and
igraph_edge_connectivity().
```

Edge- and vertex-disjoint paths

igraph_edge_disjoint_paths — The maximum number of edge-disjoint paths between two vertices.

A set of paths between two vertices is called edge-disjoint if they do not share any edges. The maximum number of edge-disjoint paths are calculated by this function using maximum flow techniques. Directed paths are considered in directed graphs.

Note that the number of disjoint paths is the same as the edge connectivity of the two vertices using uniform edge weights.

Arguments:

graph: The input graph, can be directed or undirected.

res: Pointer to an integer variable, the result will be stored here.

source: The id of the source vertex.target: The id of the target vertex.

Returns:

Error code.

Time complexity: $O(|V|^3)$, but see the discussion at igraph_maxflow_value().

See also:

igraph_vertex_disjoint_paths — Maximum number of vertex-disjoint paths between two vertices.

A set of paths between two vertices is called vertex-disjoint if they share no vertices. The calculation is performed by using maximum flow techniques.

Note that the number of vertex-disjoint paths is the same as the vertex connectivity of the two vertices in most cases (if the two vertices are not connected by an edge).

Arguments:

graph: The input graph.

res: Pointer to an integer variable, the result will be stored here.

source: The id of the source vertex.

target: The id of the target vertex.

Returns:

Error code.

Time complexity: $O(|V|^3)$.

See also:

Graph adhesion and cohesion

igraph_adhesion — Graph adhesion, this is (almost) the same as edge connectivity.

This quantity is defined by White and Harary in The cohesiveness of blocks in social networks: node connectivity and conditional density, (Sociological Methodology 31:305--359, 2001) and basically it is the edge connectivity of the graph with uniform edge weights.

Arguments:

graph: The input graph, either directed or undirected.

res: Pointer to an integer, the result will be stored here.

checks: Logical constant. Whether to check that the graph is connected and also the degree of

the vertices. If the graph is not (strongly) connected then the adhesion is obviously zero. Otherwise if the minimum degree is one then the adhesion is also one. It is a good idea to perform these checks, as they can be done quickly compared to the edge connectivity

calculation itself. They were suggested by Peter McMahan, thanks Peter. *

Returns:

Error code.

Time complexity: $O(\log(|V|)*|V|^2)$ for undirected graphs and $O(|V|^4)$ for directed graphs, but see also the discussion at the documentation of igraph_maxflow_value().

See also:

```
igraph_cohesion(), igraph_maxflow_value(), igraph_edge_connectivi-
ty(),igraph_mincut_value().
```

igraph_cohesion — Graph cohesion, this is the same as vertex connectivity.

This quantity was defined by White and Harary in "The cohesiveness of blocks in social networks: node connectivity and conditional density", (Sociological Methodology 31:305--359, 2001) and it is the same as the vertex connectivity of a graph.

Arguments:

graph: The input graph.

res: Pointer to an integer variable, the result will be stored here.

checks: Logical constant. Whether to check that the graph is connected and also the degree of

the vertices. If the graph is not (strongly) connected then the cohesion is obviously zero. Otherwise if the minimum degree is one then the cohesion is also one. It is a good idea to perform these checks, as they can be done quickly compared to the vertex connectivity

calculation itself. They were suggested by Peter McMahan, thanks Peter.

Returns:

Error code.

Time complexity: $O(|V|^4)$, |V| is the number of vertices. In practice it is more like $O(|V|^4)$, see $igraph_maxflow_value()$.

See also:

igraph_vertex_connectivity(), igraph_adhesion(), igraph_maxflow_value().

Cohesive blocks

igraph_cohesive_blocks — Identifies the hierarchical cohesive block structure of a graph

Cohesive blocking is a method of determining hierarchical subsets of graph vertices based on their structural cohesion (or vertex connectivity). For a given graph G, a subset of its vertices S is said to be maximally k-cohesive if there is no superset of S with vertex connectivity greater than or equal to k. Cohesive blocking is a process through which, given a k-cohesive set of vertices, maximally l-cohesive subsets are recursively identified with l>k. Thus a hiearchy of vertex subsets is found, with the entire graph G at its root. See the following reference for details: J. Moody and D. R. White. Structural cohesion and embeddedness: A hierarchical concept of social groups. American Sociological Review, 68(1):103--127, Feb 2003.

This function implements cohesive blocking and calculates the complete cohesive block hierarchy of a graph.

Arguments:

graph: The input graph. It must be undirected and simple. See igraph_is_simple().

blocks: If not a null pointer, then it must be an initialized list of integers vectors; the co-

hesive blocks will be stored here. Each block is encoded with a vector of type

igraph_vector_int_t that contains the vertex IDs of the block.

cohesion: If not a null pointer, then it must be an initialized vector and the cohesion of the

blocks is stored here, in the same order as the blocks in the blocks pointer vector.

parent: If not a null pointer, then it must be an initialized vector and the block hierarchy is

stored here. For each block, the id (i.e. the position in the blocks pointer vector)

of its parent block is stored. For the top block in the hierarchy, -1 is stored.

block_tree: If not a null pointer, then it must be a pointer to an uninitialized graph, and the

block hierarchy is stored here as an igraph graph. The vertex IDs correspond to the

order of the blocks in the blocks vector.

Returns:

Error code.

Time complexity: TODO.

Example 22.6. File examples/simple/cohesive_blocks.c

Chapter 23. Vertex separators

igraph_is_separator — Would removing this set of vertices disconnect the graph?

Arguments:

graph: The input graph. It may be directed, but edge directions are ignored.

candidate: The candidate separator. It must not contain all vertices.

res: Pointer to a Boolean variable, the result is stored here.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number vertices and edges.

Example 23.1. File examples/simple/igraph_is_separator.c

igraph_is_minimal_separator — Decides whether a set of vertices is a minimal separator.

A set of vertices is a minimal separator, if the removal of the vertices disconnects the graph, and this is not true for any subset of the set.

This implementation first checks that the given candidate is a separator, by calling igraph_is_separator(). If it is a separator, then it checks that each subset of size n-1, where n is the size of the candidate, is not a separator.

Arguments:

graph: The input graph. It may be directed, but edge directions are ignored.

candidate: The candidate minimal separators.

res: Pointer to a boolean variable, the result is stored here.

Returns:

Error code.

Time complexity: O(n(|V|+|E|)), |V| is the number of vertices, |E| is the number of edges, n is the number vertices in the candidate separator.

Example 23.2. File examples/simple/igraph_is_minimal_separator.c

igraph_all_minimal_st_separators — List all vertex sets that are minimal (s,t) separators for some s and t.

```
igraph_error_t igraph_all_minimal_st_separators(
    const igraph_t *graph, igraph_vector_int_list_t *separators
);
```

This function lists all vertex sets that are minimal (s,t) separators for some (s,t) vertex pair.

Note that some vertex sets returned by this function may not be minimal with respect to disconnecting the graph (or increasing the number of connected components). Take for example the 5-vertex graph with edges 0-1-2-3-4-1. This function returns the vertex sets $\{1\}$, $\{2,4\}$ and $\{1,3\}$. Notice that $\{1,3\}$ is not minimal with respect to disconnecting the graph, as $\{1\}$ would be sufficient for that. However, it is minimal with respect to separating vertices 2 and 4.

See more about the implemented algorithm in Anne Berry, Jean-Paul Bordat and Olivier Cogis: Generating All the Minimal Separators of a Graph, In: Peter Widmayer, Gabriele Neyer and Stephan Eidenbenz (editors): Graph-theoretic concepts in computer science, 1665, 167--172, 1999. Springer. https://doi.org/10.1007/3-540-46784-X_17

Arguments:

graph: The input graph. It may be directed, but edge directions are ignored.

separators: An initialized pointer vector, the separators are stored here. It is a list of point-

ers to igraph_vector_int_t objects. Each vector will contain the ids of the vertices in the separator. To free all memory allocated for <code>separators</code>, you need call <code>igraph_vector_destroy()</code> and then <code>igraph_free()</code> on each element,

before destroying the pointer vector itself.

Returns:

Error code.

See also:

```
igraph_minimum_size_separators()
```

Time complexity: $O(n|V|^3)$, |V| is the number of vertices, n is the number of separators.

Example 23.3. File examples/simple/igraph_minimal_separators.c

igraph_minimum_size_separators — Find all minimum size separating vertex sets.

```
igraph_error_t igraph_minimum_size_separators(
    const igraph_t *graph, igraph_vector_int_list_t *separators
);
```

This function lists all separator vertex sets of minimum size. A vertex set is a separator if its removal disconnects the graph.

The implementation is based on the following paper: Arkady Kanevsky: Finding all minimum-size separating vertex sets in a graph, Networks 23, 533--541, 1993.

Arguments:

graph: The input graph, which must be undirected.

separators: An initialized list of integer vectors, the separators are stored here. It is a list of

pointers to igraph_vector_int_t objects. Each vector will contain the IDs of the

vertices in the separator. The separators are returned in an arbitrary order.

Returns:

Error code.

Time complexity: TODO.

Example 23.4. File examples/simple/

igraph_minimum_size_separators.c

igraph_even_tarjan_reduction — Even-Tarjan reduction of a graph.

```
igraph_error_t igraph_even_tarjan_reduction(const igraph_t *graph, igraph_t *graph_vector_t *capacity);
```

A digraph is created with twice as many vertices and edges. For each original vertex i, two vertices i' = i and i'' = i' + n are created, with a directed edge from i' to i''. For each original directed edge from i to j, two new edges are created, from i' to j'' and from i'' to j'.

This reduction is used in the paper (observation 2): Arkady Kanevsky: Finding all minimum-size separating vertex sets in a graph, Networks 23, 533--541, 1993.

The original paper where this reduction was conceived is Shimon Even and R. Endre Tarjan: Network Flow and Testing Graph Connectivity, SIAM J. Comput., 4(4), 507–518.

Arguments:

graph: A graph. Although directness is not checked, this function is commonly used only on

directed graphs.

graphbar: Pointer to a new directed graph that will contain the reduction, with twice as many

vertices and edges.

capacity: Pointer to an initialized vector or a null pointer. If not a null pointer, then it will be

filled the capacity from the reduction: the first |E| elements are 1, the remaining |E| are

equal to |V| (which is used to indicate infinity).

Returns:

Error code.

Time complexity: O(|E|+|V|).

Example 23.5. File examples/simple/even_tarjan.c

Chapter 24. Detecting community structure

Common functions related to community structure

igraph_modularity — Calculates the modularity of a graph with respect to some clusters or vertex types.

The modularity of a graph with respect to some clustering of the vertices (or assignment of vertex types) measures how strongly separated the different clusters are from each other compared to a random null model. It is defined as

```
Q = 1/(2m) sum_{ij} (A_{ij} - \# k_i k_j / (2m)) \#(c_i,c_j),
```

where m is the number of edges, A_ij is the adjacency matrix, k_i is the degree of vertex i, c_i is the cluster that vertex i belongs to (or its vertex type), #(i,j)=1 if i=j and 0 otherwise, and the sum goes over all i, j pairs of vertices. Note that in this formula, the diagonal of the adjacency matrix contains twice the number of self-loops.

The resolution parameter # allows weighting the random null model, which might be useful when finding partitions with a high modularity. Maximizing modularity with higher values of the resolution parameter typically results in more, smaller clusters when finding partitions with a high modularity. Lower values typically results in fewer, larger clusters. The original definition of modularity is retrieved when setting # = 1.

Modularity can also be calculated on directed graphs. This only requires a relatively modest change,

```
Q = 1/m sum_{ij} (A_{ij} - \# k^out_{i} k^in_{j} / m) \#(c_{i}, c_{j}),
```

where k^out_i is the out-degree of node i and k^in_j is the in-degree of node j.

Modularity on weighted graphs is also meaningful. When taking edge weights into account, A_ij equals the weight of the corresponding edge (or 0 if there is no edge), k_ij is the strength (i.e. the weighted degree) of vertex i, with similar counterparts for a directed graph, and m is the total weight of all edges.

Note that the modularity is not well-defined for graphs with no edges. igraph returns NaN for graphs with no edges; see https://github.com/igraph/israph/issues/1539 for a detailed discussion.

For the original definition of modularity, see Newman, M. E. J., and Girvan, M. (2004). Finding and evaluating community structure in networks. Physical Review E 69, 026113. https://doi.org/10.1103/PhysRevE.69.026113

For the directed definition of modularity, see Leicht, E. A., and Newman, M. E. J. (2008). Community Structure in Directed Networks. Physical Review Letters 100, 118703. https://doi.org/10.1103/PhysRevLett.100.118703

For the introduction of the resolution parameter #, see Reichardt, J., and Bornholdt, S. (2006). Statistical mechanics of community detection. Physical Review E 74, 016110. https://doi.org/10.1103/PhysRevE.74.016110

Arguments:

graph: The input graph.

membership: Numeric vector of integer values which gives the type of each vertex, i.e. the cluster

to which it belongs. It does not have to be consecutive, i.e. empty communities

are allowed.

weights: Weight vector or NULL if no weights are specified.

resolution: The resolutin parameter #. Must not be negative. Set it to 1 to use the classical

definition of modularity.

directed: Whether to use the directed or undirected version of modularity. Ignored for undi-

rected graphs.

modularity: Pointer to a real number, the result will be stored here.

Returns:

Error code.

See also:

```
igraph_modularity_matrix()
```

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges.

igraph_modularity_matrix — Calculates the modularity matrix.

This function returns the modularity matrix, which is defined as

```
B_{ij} = A_{ij} - \# k_i k_j / (2m)
```

for undirected graphs, where A_ij is the adjacency matrix, # is the resolution parameter, k_i is the degree of vertex i, and m is the number of edges in the graph. When there are no edges, or the weights add up to zero, the result is undefined.

For directed graphs the modularity matrix is changed to

```
B_{ij} = A_{ij} - \# k^out_i k^in_j / m
```

where k^out_i is the out-degree of node i and k^in_j is the in-degree of node j.

Note that self-loops in undirected graphs are multiplied by 2 in this implementation. If weights are specified, the weighted counterparts of the adjacency matrix and degrees are used.

Arguments:

graph: The input graph.

weights: Edge weights, pointer to a vector. If this is a null pointer then every edge is assumed

to have a weight of 1.

resolution: The resolution parameter #. Must not be negative. Default is 1. Lower values favor

fewer, larger communities; higher values favor more, smaller communities.

modmat: Pointer to an initialized matrix in which the modularity matrix is stored.

directed: For directed graphs: if the edges should be treated as undirected. For undirected

graphs this is ignored.

See also:

igraph_modularity()

igraph_community_optimal_modularity — Calculate the community structure with the highest modularity value

This function calculates the optimal community structure for a graph, in terms of maximal modularity score.

The calculation is done by transforming the modularity maximization into an integer programming problem, and then calling the GLPK library to solve that. Please see Ulrik Brandes et al.: On Modularity Clustering, IEEE Transactions on Knowledge and Data Engineering 20(2):172-188, 2008.

Note that modularity optimization is an NP-complete problem, and all known algorithms for it have exponential time complexity. This means that you probably don't want to run this function on larger graphs. Graphs with up to fifty vertices should be fine, graphs with a couple of hundred vertices might be possible.

Arguments:

graph: The input graph. It is always treated as undirected.

modularity: Pointer to a real number, or a null pointer. If it is not a null pointer, then a optimal

modularity value is returned here.

membership: Pointer to a vector, or a null pointer. If not a null pointer, then the membership

vector of the optimal community structure is stored here.

weights: Vector giving the weights of the edges. If it is NULL then each edge is supposed

to have the same weight.

Returns:

Error code. When GLPK is not available, IGRAPH_UNIMPLEMENTED is returned.

See also:

igraph_modularity(), igraph_community_fastgreedy() for an algorithm that finds a local optimum in a greedy way.

Time complexity: exponential in the number of vertices.

Example 24.1. File examples/simple/igraph_community_optimal_modularity.c

igraph_community_to_membership — Creates a membership vector from a community structure dendrogram.

This function creates a membership vector from a community structure dendrogram. A membership vector contains for each vertex the id of its graph component, the graph components are numbered from zero, see the same argument of <code>igraph_connected_components()</code> for an example of a membership vector.

Many community detection algorithms return with a merges matrix, igraph_community_walk-trap() and igraph_community_edge_betweenness() are two examples. The matrix contains the merge operations performed while mapping the hierarchical structure of a network. If the matrix has n-1 rows, where n is the number of vertices in the graph, then it contains the hierarchical structure of the whole network and it is called a dendrogram.

This function performs steps merge operations as prescribed by the merges matrix and returns the current state of the network.

If merges is not a complete dendrogram, it is possible to take steps steps if steps is not bigger than the number lines in merges.

Arguments:

merges: The two-column matrix containing the merge operations. See igraph_commu-

nity_walktrap() for the detailed syntax.

nodes: The number of leaf nodes in the dendrogram.

steps: Integer constant, the number of steps to take.

membership: Pointer to an initialized vector, the membership results will be stored here, if not

NULL. The vector will be resized as needed.

csize: Pointer to an initialized vector, or NULL. If not NULL then the sizes of the com-

ponents will be stored here, the vector will be resized as needed.

See also:

```
igraph_community_walktrap(), igraph_community_edge_betweenness(),
igraph_community_fastgreedy() for community structure detection algorithms.
```

Time complexity: O(|V|), the number of vertices in the graph.

igraph_reindex_membership — Makes the IDs in a membership vector contiguous.

This function reindexes component IDs in a membership vector in a way that the new IDs start from zero and go up to C-1, where C is the number of unique component IDs in the original vector. The supplied membership is expected to fall in the range 0, ..., n - 1.

Arguments:

membership: Numeric vector which gives the type of each vertex, i.e. the component to which

it belongs. The vector will be altered in-place.

new_to_old: Pointer to a vector which will contain the old component ID for each new one, or

NULL, in which case it is not returned. The vector will be resized as needed.

nb_clusters: Pointer to an integer for the number of distinct clusters. If not NULL, this will be

updated to reflect the number of distinct clusters found in membership.

Time complexity: should be O(n) for n elements.

igraph_compare_communities — Compares community structures using various metrics.

This function assesses the distance between two community structures using the variation of information (VI) metric of Meila (2003), the normalized mutual information (NMI) of Danon et al (2005), the split-join distance of van Dongen (2000), the Rand index of Rand (1971) or the adjusted Rand index of Hubert and Arabie (1985).

Some of these measures are defined based on the entropy of a discrete random variable associated with a given clustering C of vertices. Let p_i be the probability that a randomly picked vertex would be part of cluster i. Then the entropy of the clustering is

```
H(C) = - \sum_{i=1}^{n} \log p_i
```

Similarly, we can define the joint entropy of two clusterings C_1 and C_2 based on the probability p_ij that a random vertex is part of cluster i in the first clustering and cluster j in the second one:

```
H(C_1, C_2) = - \sum_{i=1}^{n} \log p_{ij}
```

The mutual information of C_1 and C_2 is then $MI(C_1, C_2) = H(C_1) + H(C_2) - H(C_1, C_2) >= 0$. A large mutual information indicates a high overlap between the two clusterings. The normalized mutual information, as computed by igraph, is

```
NMI(C_1, C_2) = 2 MI(C_1, C_2) / (H(C_1) + H(C_2)).
```

It takes its value from the interval (0, 1], with 1 achieved when the two clusterings coincide.

The variation of information is defined as $VI(C_1, C_2) = [H(C_1) - MI(C_1, C_2)] + [H(C_2) - MI(C_1, C_2)]$. Lower values of the variation of information indicate a smaller difference between the two clusterings, with VI = 0 achieved precisely when they coincide.

The Rand index is defined as the probability that the two clusterings agree about the cluster memberships of a randomly chosen vertex *pair*. All vertex pairs are considered, and the two clusterings are considered to be in agreement about the memberships of a vertex pair if either the two vertices are in the same cluster in both clusterings, or they are in different clusters in both clusterings. The Rand index is then the number of vertex pairs in agreement, divided by the total number of vertex pairs. A Rand index of zero means that the two clusterings disagree about the membership of all vertex pairs, while 1 means that the two clusterings are identical.

The adjusted Rand index is similar to the Rand index, but it takes into account that agreement between the two clusterings may also occur by chance even if the two clusterings are chosen completely randomly. The adjusted Rand index therefore subtracts the expected fraction of agreements from the value of the Rand index, and divides the result by one minus the expected fraction of agreements. The maximum value of the adjusted Rand index is still 1 (similarly to the Rand index), indicating maximum agreement, but the value may be less than zero if there is *less* agreement between the two clusterings than what would be expected by chance.

For an explanation of the split-join distance, see igraph_split_join_distance().

References:

Meil# M: Comparing clusterings by the variation of information. In: Schölkopf B, Warmuth MK (eds.). Learning Theory and Kernel Machines: 16th Annual Conference on Computational Learning Theory and 7th Kernel Workshop, COLT/Kernel 2003, Washington, DC, USA. Lecture Notes in Computer Science, vol. 2777, Springer, 2003. ISBN: 978-3-540-40720-1. https://doi.org/10.1007/978-3-540-45167-9_14

Danon L, Diaz-Guilera A, Duch J, Arenas A: Comparing community structure identification. J Stat Mech P09008, 2005. https://doi.org/10.1088/1742-5468/2005/09/P09008

van Dongen S: Performance criteria for graph clustering and Markov cluster experiments. Technical Report INS-R0012, National Research Institute for Mathematics and Computer Science in the Netherlands, Amsterdam, May 2000. https://ir.cwi.nl/pub/4461

Rand WM: Objective criteria for the evaluation of clustering methods. J Am Stat Assoc 66(336):846-850, 1971. https://doi.org/10.2307/2284239

 $Hubert\ L\ and\ Arabie\ P:\ Comparing\ partitions.\ Journal\ of\ Classification\ 2:193-218,\ 1985.\ https://doi.org/10.1007/BF01908075$

Arguments:

comm1: the membership vector of the first community structure

comm2: the membership vector of the second community structure

result: the result is stored here.

method: the comparison method to use. IGRAPH_COMMCMP_VI selects the variation of infor-

mation (VI) metric of Meila (2003), IGRAPH_COMMCMP_NMI selects the normalized mutual information measure proposed by Danon et al (2005), IGRAPH_COMMCMP_S-

PLIT_JOIN selects the split-join distance of van Dongen (2000), IGRAPH_COMMCM-P_RAND selects the unadjusted Rand index (1971) and IGRAPH_COMMCMP_ADJUST-ED_RAND selects the adjusted Rand index.

Returns:

Error code.

See also:

```
igraph_split_join_distance().
```

Time complexity: $O(n \log(n))$.

igraph_split_join_distance — Calculates the split-join distance of two community structures.

The split-join distance between partitions A and B is the sum of the projection distance of A from B and the projection distance of B from A. The projection distance is an asymmetric measure and it is defined as follows:

First, each set in partition A is evaluated against all sets in partition B. For each set in partition A, the best matching set in partition B is found and the overlap size is calculated. (Matching is quantified by the size of the overlap between the two sets). Then, the maximal overlap sizes for each set in A are summed together and subtracted from the number of elements in A.

The split-join distance will be returned in two arguments, distance12 will contain the projection distance of the first partition from the second, while distance21 will be the projection distance of the second partition from the first. This makes it easier to detect whether a partition is a subpartition of the other, since in this case, the corresponding distance will be zero.

Reference:

van Dongen S: Performance criteria for graph clustering and Markov cluster experiments. Technical Report INS-R0012, National Research Institute for Mathematics and Computer Science in the Netherlands, Amsterdam, May 2000.

Arguments:

comm1: the membership vector of the first community structure

comm2: the membership vector of the second community structure

distance 12: pointer to an igraph_integer_t, the projection distance of the first commu-

nity structure from the second one will be returned here.

distance21: pointer to an igraph_integer_t, the projection distance of the second com-

munity structure from the first one will be returned here.

Returns:

Error code.

See also:

igraph_compare_communities() with the IGRAPH_COMMCMP_SPLIT_JOIN method if you are not interested in the individual distances but only the sum of them.

Time complexity: $O(n \log(n))$.

Community structure based on statistical mechanics

igraph_community_spinglass — Community detection based on statistical mechanics.

This function implements the community structure detection algorithm proposed by Joerg Reichardt and Stefan Bornholdt. The algorithm is described in their paper: Statistical Mechanics of Community Detection, http://arxiv.org/abs/cond-mat/0603718.

From version 0.6, igraph also supports an extension to the algorithm that allows negative edge weights. This is described in V. A. Traag and Jeroen Bruggeman: Community detection in networks with positive and negative links, http://arxiv.org/abs/0811.2329.

Arguments:

graph: The input graph, it may be directed but the direction of the edges is not used

in the algorithm.

weights: The vector giving the edge weights, it may be NULL, in which case all

edges are weighted equally. The edge weights must be positive unless using the IGRAPH_SPINCOMM_IMP_NEG implementation. This condition is not

verified by the function.

modularity: Pointer to a real number, if not NULL then the modularity score of the solution

will be stored here. This is the gereralized modularity that simplifies to the one defined in M. E. J. Newman and M. Girvan, Phys. Rev. E 69, 026113

(2004), if the gamma parameter is one.

temperature: Pointer to a real number, if not NULL then the temperature at the end of the

algorithm will be stored here.

membership: Pointer to an initialized vector or NULL. If not NULL then the result of the

clustering will be stored here. For each vertex, the number of its cluster is given, with the first cluster numbered zero. The vector will be resized as

needed.

csize: Pointer to an initialized vector or NULL. If not NULL then the sizes of the

clusters will stored here in cluster number order. The vector will be resized

as needed.

spins: Integer giving the number of spins, i.e. the maximum number of clusters.

Usually it is not a program to give a high number here, the default was 25 in the original code. Even if the number of spins is high the number of clusters

in the result might be small.

parupdate: A logical constant, whether to update all spins in parallel. The default for

this argument was false in the original code. It is not implemented in the

IGRAPH_SPINCOMM_INP_NEG implementation.

starttemp: Real number, the temperature at the start. The value of this argument was

1.0 in the original code.

stoptemp: Real number, the algorithm stops at this temperature. The default was 0.01

in the original code.

coolfact: Real number, the cooling factor for the simulated annealing. The default was

0.99 in the original code.

update_rule: The type of the update rule. Possible values: IGRAPH_SPINCOMM_UP-

DATE_SIMPLE and IGRAPH_SPINCOMM_UPDATE_CONFIG. Basically this parameter defines the null model based on which the actual clustering is done. If this is $IGRAPH_SPINCOMM_UPDATE_SIMPLE$ then the random graph (i.e. G(n,p)), if it is $IGRAPH_SPINCOMM_UPDATE$ then the configuration model is used. The configuration means that the baseline for the clustering is a random graph with the same degree distribution as the input graph.

gamma: Real number. The gamma parameter of the algorithm. This defines the weight

of the missing and existing links in the quality function for the clustering. The default value in the original code was 1.0, which is equal weight to missing and existing edges. Smaller values make the existing links contibute more to the energy function which is minimized in the algorithm. Bigger values make the missing links more important. (If my understanding is correct.)

implementation: Constant, chooses between the two implementations of the spin-glass algo-

rithm that are included in igraph. IGRAPH_SPINCOMM_IMP_ORIG selects the original implementation, this is faster, IGRAPH_SPINCOMM_INP_NEG selects a new implementation by Vincent Traag that allows negative edge

weights.

gamma_minus: Real number. Parameter for the IGRAPH_SPINCOMM_IMP_NEG imple-

mentation. This specifies the balance between the importance of present and non-present negative weighted edges in a community. Smaller values of gamma_minus lead to communities with lesser negative intra-connectivity. If this argument is set to zero, the algorithm reduces to a graph coloring

algorithm, using the number of spins as the number of colors.

Returns:

Error code.

See also:

igraph_community_spinglass_single() for calculating the community of a single vertex.

Time complexity: TODO.

igraph_community_spinglass_single — Community of a single node based on statistical mechanics.

This function implements the community structure detection algorithm proposed by Joerg Reichardt and Stefan Bornholdt. It is described in their paper: Statistical Mechanics of Community Detection, http://arxiv.org/abs/cond-mat/0603718 .

This function calculates the community of a single vertex without calculating all the communities in the graph.

Arguments:

graph: The input graph, it may be directed but the direction of the edges is not used in the algorithm. weights: Pointer to a vector with the weights of the edges. Alternatively NULL can be supplied to have the same weight for every edge. vertex: The vertex ID of the vertex of which ths community is calculated. Pointer to an initialized vector, the result, the IDs of the vertices in the community community: of the input vertex will be stored here. The vector will be resized as needed. cohesion: Pointer to a real variable, if not NULL the cohesion index of the community will be stored here. Pointer to a real variable, if not NULL the adhesion index of the community will adhesion: be stored here. Pointer to an integer, if not NULL the number of edges within the community is inner_links: Pointer to an integer, if not NULL the number of edges between the community outer_links: and the rest of the graph will be stored here. The number of spins to use, this can be higher than the actual number of clusters spins: in the network, in which case some clusters will contain zero vertices. The type of the update rule. Possible values: IGRAPH SPINCOMM UPupdate rule: DATE SIMPLE and IGRAPH SPINCOMM UPDATE CONFIG. Basically this parameter defined the null model based on which the actual clustering is done. If this is ${\tt IGRAPH_SPINCOMM_UPDATE_SIMPLE}$ then the random graph (ie. G(n,p)), if it is ${\tt IGRAPH_SPINCOMM_UPDATE}$ then the configuration model is used. The configuration means that the baseline for the clustering is a random graph with the same degree distribution as the input graph.

gamma:

Real number. The gamma parameter of the algorithm. This defined the weight of the missing and existing links in the quality function for the clustering. The default value in the original code was 1.0, which is equal weight to missing and existing edges. Smaller values make the existing links contibute more to the energy function which is minimized in the algorithm. Bigger values make the missing links more important. (If my understanding is correct.)

Returns:

Error code.

See also:

igraph_community_spinglass() for the traditional version of the algorithm.

Time complexity: TODO.

Community structure based on eigenvectors of matrices

The function documented in these section implements the "leading eigenvector" method developed by Mark Newman and published in MEJ Newman: Finding community structure using the eigenvectors of matrices, Phys Rev E 74:036104 (2006).

The heart of the method is the definition of the modularity matrix, B, which is B=A-P, A being the adjacency matrix of the (undirected) network, and P contains the probability that certain edges are present according to the "configuration model" In other words, a Pij element of P is the probability that there is an edge between vertices i and j in a random network in which the degrees of all vertices are the same as in the input graph.

The leading eigenvector method works by calculating the eigenvector of the modularity matrix for the largest positive eigenvalue and then separating vertices into two community based on the sign of the corresponding element in the eigenvector. If all elements in the eigenvector are of the same sign that means that the network has no underlying community structure. Check Newman's paper to understand why this is a good method for detecting community structure.

The leading eigenvector community structure detection method is implemented in igraph_community_leading_eigenvector(). After the initial split, the following splits are done in a way to optimize modularity regarding to the original network. Note that any further refinement, for example using Kernighan-Lin, as proposed in Section V.A of Newman (2006), is not implemented here.

Example 24.2. File examples/simple/igraph_community_leading_eigenvector.c

igraph_community_leading_eigenvector — Leading eigenvector community finding (proper version).

Newman's leading eigenvector method for detecting community structure. This is the proper implementation of the recursive, divisive algorithm: each split is done by maximizing the modularity regarding the original network, see MEJ Newman: Finding community structure in networks using the eigenvectors of matrices, Phys Rev E 74:036104 (2006).

Arguments:

graph: The undirected input graph.

weights: The weights of the edges, or a null pointer for unweighted graphs.

merges: The result of the algorithm, a matrix containing the information about the

splits performed. The matrix is built in the opposite way however, it is like the result of an agglomerative algorithm. If at the end of the algorithm (after steps steps was done) there are "p" communities, then these are numbered from zero to "p-1". The first line of the matrix contains the first "merge" (which is in reality the last split) of two communities into community "p", the merge in the second line forms community "p+1", etc. The matrix should be initialized before calling and will be resized as needed. This argument is

ignored if it is NULL.

membership: The membership of the vertices after all the splits were performed will be

stored here. The vector must be initialized before calling and will be resized as needed. This argument is ignored if it is NULL. This argument can also be used to supply a starting configuration for the community finding, in the format of a membership vector. In this case the start argument must be

set to 1.

steps: The maximum number of steps to perform. It might happen that some com-

ponent (or the whole network) has no underlying community structure and no further steps can be done. If you want as many steps as possible then sup-

ply the number of vertices in the network here.

options: The options for ARPACK. Supply NULL here to use the defaults. n is always

overwritten. ncv is set to at least 4.

modularity: If not a null pointer, then it must be a pointer to a real number and the mod-

ularity score of the final division is stored here.

start: Boolean, whether to use the community structure given in the membership

argument as a starting point.

eigenvalues: Pointer to an initialized vector or a null pointer. If not a null pointer, then the

eigenvalues calculated along the community structure detection are stored

here. The non-positive eigenvalues, that do not result a split, are stored as well.

eigenvectors:

If not a null pointer, then the eigenvectors that are calculated in each step of the algorithm are stored here, in a list of vectors. Each eigenvector is stored in an igraph_vector_t object.

history:

Pointer to an initialized vector or a null pointer. If not a null pointer, then a trace of the algorithm is stored here, encoded numerically. The various operations:

IGRAPH LEVC HIST S-

TART_FULL

Start the algorithm from an initial state where each connected compo-

nent is a separate community.

IGRAPH_LEVC_HIST_S-

TART_GIVEN

Start the algorithm from a given community structure. The next value in the vector contains the initial number

of communities.

IGRAPH_LEVC_HIST_SPLIT Split a community into two commu-

> nities. The id of the splitted community is given in the next element of the history vector. The id of the first new community is the same as the id of the splitted community. The id of the second community equals to the number of communities before the split.

IGRAPH_LEVC_HIST_FAILED

Tried to split a community, but it was not worth it, as it does not result in a bigger modularity value. The id of the community is given in the next el-

ement of the vector.

callback:

A null pointer or a function of type igraph_community_leading_eigenvector_callback_t. If given, this callback function is called after each eigenvector/eigenvalue calculation. If the callback returns IGRAPH_STOP, then the community finding algorithm stops. If it returns IGRAPH_SUCCESS, the algorithm continues normally. Any other return value is considered an igraph error code and will terminete the algorithm with the same error code. See the arguments passed to the callback at the documentation of igraph_community_leading_eigenvec-

tor_callback_t.

callback_extra:

Extra argument to pass to the callback function.

Returns:

Error code.

See also:

igraph_community_walktrap() and igraph_community_spinglass() for other community structure detection methods.

Time complexity: $O(|E|+|V|^2*steps)$, |V| is the number of vertices, |E| the number of edges, "steps" the number of splits performed.

igraph_community_leading_eigenvector_call-back_t — Callback for the leading eigenvector community finding method.

```
typedef igraph_error_t igraph_community_leading_eigenvector_callback_t(
   const igraph_vector_int_t *membership,
   igraph_integer_t comm,
   igraph_real_t eigenvalue,
   const igraph_vector_t *eigenvector,
   igraph_arpack_function_t *arpack_multiplier,
   void *arpack_extra,
   void *extra);
```

The leading eigenvector community finding implementation in igraph is able to call a callback function, after each eigenvalue calculation. This callback function must be of igraph_communi-ty_leading_eigenvector_callback_t type. The following arguments are passed to the callback:

Arguments:

membership: The actual membership vector, before recording the potential change

implied by the newly found eigenvalue.

comm: The id of the community that the algorithm tried to split in the last iter-

ation. The community IDs are indexed from zero here!

eigenvalue: The eigenvalue the algorithm has just found.

eigenvector: The eigenvector corresponding to the eigenvalue the algorithm just

found.

arpack_multiplier: A function that was passed to igraph_arpack_rssolve() to

solve the last eigenproblem.

arpack_extra: The extra argument that was passed to the ARPACK solver.

extra: Extra argument that as passed to igraph_community_lead-

ing_eigenvector().

See also:

igraph_community_leading_eigenvector(), igraph_arpack_function_t, igraph_arpack_rssolve().

igraph_le_community_to_membership — Vertex membership from the leading eigenvector community structure

```
igraph_vector_int_t *csize);
```

This function creates a membership vector from the result of igraph_community_leading_eigenvector(), It takes membership and performs steps merges, according to the supplied merges matrix.

Arguments:

merges: The two-column matrix containing the merge operations. See igraph_commu-

nity_walktrap() for the detailed syntax. This is usually from the output of

the leading eigenvector community structure detection routines.

steps: The number of steps to make according to merges.

membership: Initially the starting membership vector, on output the resulting membership vector,

after performing steps merges.

csize: Optionally the sizes of the communities is stored here, if this is not a null pointer,

but an initialized vector.

Returns:

Error code.

Time complexity: O(|V|), the number of vertices.

Walktrap: Community structure based on random walks

igraph_community_walktrap — This function is the implementation of the Walktrap community

finding algorithm, see Pascal Pons, Matthieu Latapy: Computing communities in large networks using random walks, https://arxiv.org/abs/physics/0512106

Currently the original C++ implementation is used in igraph, see https://www-complexnetwork-s.lip6.fr/~latapy/PP/walktrap.html We are grateful to Matthieu Latapy and Pascal Pons for providing this source code.

In contrast to the original implementation, isolated vertices are allowed in the graph and they are assumed to have a single incident loop edge with weight 1.

Arguments:

graph: The input graph, edge directions are ignored.

weights: Numeric vector giving the weights of the edges. If it is a NULL pointer then all

edges will have equal weights. The weights are expected to be positive.

steps: Integer constant, the length of the random walks. Typically, good results are ob-

tained with values between 3-8 with 4-5 being a reasonable default.

merges: Pointer to a matrix, the merges performed by the algorithm will be stored here (if

not NULL). Each merge is a row in a two-column matrix and contains the IDs of the merged clusters. Clusters are numbered from zero and cluster numbers smaller than the number of nodes in the network belong to the individual vertices as singleton clusters. In each step a new cluster is created from two other clusters and its id will be one larger than the largest cluster id so far. This means that before the first merge we have n clusters (the number of vertices in the graph) numbered from zero to

n-1. The first merge creates cluster n, the second cluster n+1, etc.

modularity: Pointer to a vector. If not NULL then the modularity score of the current clustering

is stored here after each merge operation.

membership: Pointer to a vector. If not a NULL pointer, then the membership vector correspond-

ing to the maximal modularity score is stored here. If it is not a NULL pointer, then

neither modularity nor merges may be NULL.

Returns:

Error code.

See also:

```
igraph_community_spinglass(), igraph_community_edge_betweenness().
```

Time complexity: $O(|E||V|^2)$ in the worst case, $O(|V|^2 \log |V|)$ typically, |V| is the number of vertices, |E| is the number of edges.

Example 24.3. File examples/simple/walktrap.c

Edge betweenness based community detection

igraph_community_edge_betweenness — Community finding based on edge betweenness.

Community structure detection based on the betweenness of the edges in the network. The algorithm was invented by M. Girvan and M. Newman, see: M. Girvan and M. E. J. Newman: Community structure in social and biological networks, Proc. Nat. Acad. Sci. USA 99, 7821-7826 (2002). https://doi.org/10.1073/pnas.122653799

The idea is that the betweenness of the edges connecting two communities is typically high, as many of the shortest paths between nodes in separate communities go through them. So we gradually remove the edge with highest betweenness from the network, and recalculate edge betweenness after every removal. This way sooner or later the network splits into two components, then after a while one of these components splits again into two smaller components, and so on until all edges are removed. This is a divisive hierarchical approach, the result of which is a dendrogram.

In directed graphs, when directed is set to true, the directed version of betweenness and modularity are used, however, only splits into weakly connected components are detected.

Arguments:

graph: The input graph.

result: Pointer to an initialized vector, the result will be stored here, the IDs of the

removed edges in the order of their removal. It will be resized as needed.

It may be NULL if the edge IDs are not needed by the caller.

edge_betweenness: Pointer to an initialized vector or NULL. In the former case the edge be-

tweenness of the removed edge is stored here. The vector will be resized

as needed.

merges: Pointer to an initialized matrix or NULL. If not NULL then merges per-

formed by the algorithm are stored here. Even if this is a divisive algorithm, we can replay it backwards and note which two clusters were merged. Clusters are numbered from zero, see the *merges* argument of igraph_community_walktrap() for details. The matrix will be

resized as needed.

bridges: Pointer to an initialized vector of NULL. If not NULL then the indices into

result of all edges which caused one of the merges will be put here. This is equivalent to all edge removals which separated the network into

more components, in reverse order.

modularity: If not a null pointer, then the modularity values of the different divisions

are stored here, in the order corresponding to the merge matrix. The modularity values will take weights into account if weights is not null.

membership: If not a null pointer, then the membership vector, corresponding to the

highest modularity value, is stored here.

directed: Logical constant. Controls whether to calculate directed betweenness (i.e.

directed paths) for directed graphs, and whether to use the directed version

of modularity. It is ignored for undirected graphs.

weights: An optional vector containing edge weights. If null, the unweighted edge

betweenness scores will be calculated and used. If not null, the weighted

edge betweenness scores will be calculated and used.

Returns:

Error code.

See also:

Time complexity: $O(|V||E|^2)$, as the betweenness calculation requires O(|V||E|) and we do it |E|-1 times.

Example 24.4. File examples/simple/igraph_community_edge_betweenness.c

igraph_community_eb_get_merges — Calculating the merges, i.e. the dendrogram for an edge betweenness community structure.

This function is handy if you have a sequence of edges which are gradually removed from the network and you would like to know how the network falls apart into separate components. The edge sequence may come from the <code>igraph_community_edge_betweenness()</code> function, but this is not necessary. Note that <code>igraph_community_edge_betweenness()</code> can also calculate the dendrogram, via its <code>merges</code> argument. Merges happen when the edge removal process is run backwards and two components become connected.

Arguments:

graph: The input graph.

edges: Vector containing the edges to be removed from the network, all edges are expected

to appear exactly once in the vector.

directed: Whether to use the directed or undirected version of modularity. Will be ignored

for undirected graphs.

weights: An optional vector containing edge weights. If null, the unweighted modularity

scores will be calculated. If not null, the weighted modularity scores will be calcu-

lated. Ignored if both modularity and membership are NULL pointers.

res: Pointer to an initialized matrix, if not NULL then the dendrogram will be stored

here, in the same form as for the <code>igraph_community_walktrap()</code> function: the matrix has two columns and each line is a merge given by the IDs of the merged components. The component IDs are numbered from zero and component IDs smaller than the number of vertices in the graph belong to individual vertices. The non-trivial components containing at least two vertices are numbered from n, where n is the number of vertices in the graph. So if the first line contains a and b that means that components a and b are merged into component n, the second line

creates component n+1, etc. The matrix will be resized as needed.

bridges: Pointer to an initialized vector of NULL. If not NULL then the indices into edges

of all edges which caused one of the merges will be put here. This is equal to all edge removals which separated the network into more components, in reverse order.

modularity: If not a null pointer, then the modularity values for the different divisions, corre-

sponding to the merges matrix, will be stored here.

membership: If not a null pointer, then the membership vector for the best division (in terms of modularity) will be stored here.

Returns:

Error code.

See also:

```
igraph_community_edge_betweenness().
```

Time complexity: $O(|E|+|V|\log|V|)$, |V| is the number of vertices, |E| is the number of edges.

Community structure based on the optimization of modularity

igraph_community_fastgreedy — Finding community structure by greedy optimization of modularity.

This function implements the fast greedy modularity optimization algorithm for finding community structure, see A Clauset, MEJ Newman, C Moore: Finding community structure in very large networks, http://www.arxiv.org/abs/cond-mat/0408187 for the details.

Some improvements proposed in K Wakita, T Tsurumi: Finding community structure in mega-scale social networks, http://www.arxiv.org/abs/cs.CY/0702048v1 have also been implemented.

Arguments:

graph: The input graph. It must be a graph without multiple edges. This is checked and an

error message is given for graphs with multiple edges.

weights: Potentially a numeric vector containing edge weights. Supply a null pointer here

for unweighted graphs. The weights are expected to be non-negative.

merges: Pointer to an initialized matrix or NULL, the result of the computation is stored

here. The matrix has two columns and each merge corresponds to one merge, the IDs of the two merged components are stored. The component IDs are numbered from zero and the first n components are the individual vertices, n is the number of vertices in the graph. Component n is created in the first merge, component n +1 in the second merge, etc. The matrix will be resized as needed. If this argument

is NULL then it is ignored completely.

modularity: Pointer to an initialized vector or NULL pointer, in the former case the modularity

scores along the stages of the computation are recorded here. The vector will be

resized as needed.

membership: Pointer to a vector. If not a null pointer, then the membership vector corresponding

to the best split (in terms of modularity) is stored here.

Returns:

Error code.

See also:

igraph_community_walktrap(), igraph_community_edge_betweenness() for other community detection algorithms, igraph_community_to_membership() to convert the dendrogram to a membership vector.

Time complexity: O(|E||V|log|V|) in the worst case, $O(|E|+|V|log^2|V|)$ typically, |V| is the number of vertices, |E| is the number of edges.

Example 24.5. File examples/simple/igraph_community_fastgreedy.c

igraph_community_multilevel — Finding community structure by multi-level optimization of modularity.

This function implements the multi-level modularity optimization algorithm for finding community structure, see Blondel, V. D., Guillaume, J.-L., Lambiotte, R., & Lefebvre, E. (2008). Fast unfolding of communities in large networks. Journal of Statistical Mechanics: Theory and Experiment, 10008(10), 6. https://doi.org/10.1088/1742-5468/2008/10/P10008 for the details (preprint: http://arxiv.org/abs/0803.0476). The algorithm is sometimes known as the "Louvain" algorithm.

The algorithm is based on the modularity measure and a hierarchical approach. Initially, each vertex is assigned to a community on its own. In every step, vertices are re-assigned to communities in a local, greedy way: in a random order, each vertex is moved to the community with which it achieves the highest contribution to modularity. When no vertices can be reassigned, each community is considered a vertex on its own, and the process starts again with the merged communities. The process stops when there is only a single vertex left or when the modularity cannot be increased any more in a step.

The resolution parameter gamma allows finding communities at different resolutions. Higher values of the resolution parameter typically result in more, smaller communities. Lower values typically result in fewer, larger communities. The original definition of modularity is retrieved when setting gamma=1. Note that the returned modularity value is calculated using the indicated resolution parameter. See igraph_modularity() for more details. This function was contributed by Tom Gregorovic.

Arguments:

graph: The input graph. It must be an undirected graph.

weights: Numeric vector containing edge weights. If NULL, every edge has equal weight.

The weights are expected to be non-negative.

resolution: Resolution parameter. Must be greater than or equal to 0. Lower values favor

fewer, larger communities; higher values favor more, smaller communities. Set

it to 1 to use the classical definition of modularity.

membership: The membership vector, the result is returned here. For each vertex it gives the ID

of its community. The vector must be initialized and it will be resized accordingly.

memberships: Numeric matrix that will contain the membership vector after each level, if not

NULL. It must be initialized and it will be resized accordingly.

modularity: Numeric vector that will contain the modularity score after each level, if not

NULL. It must be initialized and it will be resized accordingly.

Returns:

Error code.

Time complexity: in average near linear on sparse graphs.

Example 24.6. File examples/simple/igraph_community_multilevel.c

igraph_community_leiden — Finding community structure using the Leiden algorithm.

This function implements the Leiden algorithm for finding community structure, see Traag, V. A., Waltman, L., & van Eck, N. J. (2019). From Louvain to Leiden: guaranteeing well-connected communities. Scientific reports, 9(1), 5233. http://dx.doi.org/10.1038/s41598-019-41695-z

It is similar to the multilevel algorithm, often called the Louvain algorithm, but it is faster and yields higher quality solutions. It can optimize both modularity and the Constant Potts Model, which does not suffer from the resolution-limit (see preprint http://arxiv.org/abs/1104.3083).

The Leiden algorithm consists of three phases: (1) local moving of nodes, (2) refinement of the partition and (3) aggregation of the network based on the refined partition, using the non-refined partition to create an initial partition for the aggregate network. In the local move procedure in the Leiden algorithm, only nodes whose neighborhood has changed are visited. Only moves that strictly improve the quality function are made. The refinement is done by restarting from a singleton partition within each cluster and gradually merging the subclusters. When aggregating, a single cluster may then be represented by several nodes (which are the subclusters identified in the refinement).

The Leiden algorithm provides several guarantees. The Leiden algorithm is typically iterated: the output of one iteration is used as the input for the next iteration. At each iteration all clusters are guaranteed to be connected and well-separated. After an iteration in which nothing has changed, all nodes and some parts are guaranteed to be locally optimally assigned. Note that even if a single iteration did not result in any change, it is still possible that a subsequent iteration might find some improvement. Each iteration explores different subsets of nodes to consider for moving from one cluster to another. Finally, asymptotically, all subsets of all clusters are guaranteed to be locally optimally assigned. For more details, please see Traag, Waltman & van Eck (2019).

The objective function being optimized is

 $1/2m \text{ sum_ij } (A_ij - \text{gamma n_i n_j})d(s_i, s_j)$

where m is the total edge weight, A_ij is the weight of edge (i, j), gamma is the so-called resolution parameter, n_i is the node weight of node i, s_i is the cluster of node i and d(x, y) = 1 if and only if x = y and 0 otherwise. By setting n_i = k_i, the degree of node i, and dividing gamma by 2m, you effectively obtain an expression for modularity. Hence, the standard modularity will be optimized when you supply the degrees as node_weights and by supplying as a resolution parameter 1.0/(2*m), with m the number of edges.

Arguments:

graph: The input graph. It must be an undirected graph.

edge_weights: Numeric vector containing edge weights. If NULL, every edge has

equal weight of 1. The weights need not be non-negative.

node_weights: Numeric vector containing node weights. If NULL, every node has

equal weight of 1.

resolution_parameter: The resolution parameter used, which is represented by gamma in

the objective function mentioned in the documentation.

beta: The randomness used in the refinement step when merging. A small

amount of randomness (beta = 0.01) typically works well.

start: Start from membership vector. If this is true, the optimization will

start from the provided membership vector. If this is false, the opti-

mization will start from a singleton partition.

n_iterations: Iterate the core Leiden algorithm for the indicated number of times.

If this is a negative number, it will continue iterating until an itera-

tion did not change the clustering.

membership: The membership vector. This is both used as the initial membership

from which optimisation starts and is updated in place. It must hence be properly initialized. When finding clusters from scratch it is typically started using a singleton clustering. This can be achieved us-

ing igraph_vector_int_init_range.

nb_clusters: The number of clusters contained in membership. If NULL, the

number of clusters will not be returned.

quality: The quality of the partition, in terms of the objective function as

included in the documentation. If NULL the quality will not be cal-

culated.

Returns:

Error code.

Time complexity: near linear on sparse graphs.

Example 24.7. File examples/simple/igraph_community_leiden.c

Fluid communities

igraph_community_fluid_communities — Community detection based on fluids interacting on the graph.

The algorithm is based on the simple idea of several fluids interacting in a non-homogeneous environment (the graph topology), expanding and contracting based on their interaction and density. This function implements the community detection method described in: Parés F, Gasulla DG, et. al. (2018) Fluid Communities: A Competitive, Scalable and Diverse Community Detection Algorithm. In: Complex Networks & Their Applications VI: Proceedings of Complex Networks 2017 (The Sixth International Conference on Complex Networks and Their Applications), Springer, vol 689, p 229.

Arguments:

graph: The input graph. The graph must be simple and connected. Empty graphs

are not supported as well as single vertex graphs. Edge directions are

ignored. Weights are not considered.

no of communities: The number of communities to be found. Must be greater than 0 and

fewer than number of vertices in the graph.

membership: The result vector mapping vertices to the communities they are assigned

to.

modularity: If not a null pointer, then it must be a pointer to a real number. The

modularity score of the detected community structure is stored here.

Returns:

Error code.

Time complexity: O(|E|)

Label propagation

igraph_community_label_propagation — Community detection based on label propagation.

This function implements the label propagation-based community detection algorithm described by Raghavan, Albert and Kumara. This version extends the original method by the ability to take edge weights into consideration and also by allowing some labels to be fixed.

Weights are taken into account as follows: when the new label of node i is determined, the algorithm iterates over all edges incident on node i and calculate the total weight of edges leading to other nodes with label 0, 1, 2, ..., k-1 (where k is the number of possible labels). The new label of node i will then be the label whose edges (among the ones incident on node i) have the highest total weight.

For directed graphs, it is important to know that labels can circulate freely only within the strongly connected components of the graph and may propagate in only one direction (or not at all) between

strongly connected components. You should treat directed edges as directed only if you are aware of the consequences.

References:

Raghavan, U.N. and Albert, R. and Kumara, S.: Near linear time algorithm to detect community structures in large-scale networks. Phys Rev E 76, 036106 (2007). https://doi.org/10.1103/Phys-RevE.76.036106

Šubelj, L.: Label propagation for clustering. Chapter in "Advances in Network Clustering and Block-modeling" edited by P. Doreian, V. Batagelj & A. Ferligoj (Wiley, New York, 2018). https://doi.org/10.1002/9781119483298.ch5 https://arxiv.org/abs/1709.05634

Arguments:

graph: The input graph. Note that the algorithm wsa originally defined for undirected

graphs. You are advised to set mode to IGRAPH_ALL if you pass a directed graph

here to treat it as undirected.

membership: The membership vector, the result is returned here. For each vertex it gives the ID

of its community (label).

mode: Whether to consider edge directions for the label propagation, and if so, which di-

rection the labels should propagate. Ignored for undirected graphs. IGRAPH_ALL means to ignore edge directions (even in directed graphs). IGRAPH_OUT means to propagate labels along the natural direction of the edges. IGRAPH_IN means to propagate labels *backwards* (i.e. from head to tail). It is advised to set this to IGRAPH_ALL unless you are specifically interested in the effect of edge direc-

tions.

weights: The weight vector, it should contain a positive weight for all the edges.

initial: The initial state. If NULL, every vertex will have a different label at the beginning.

Otherwise it must be a vector with an entry for each vertex. Non-negative values denote different labels, negative entries denote vertices without labels. Unlabeled vertices which are not reachable from any labeled ones will remain unlabeled at the end of the label propagation process, and will be labeled in an additional step to avoid returning negative values in <code>membership</code>. In undirected graphs, this happens when entire connected components are unlabeled. Then, each unlabeled component will receive its own separate label. In directed graphs, the outcome of the additional labeling should be considered undefined and may change in the

future; please do not rely on it.

fixed: Boolean vector denoting which labels are fixed. Of course this makes sense only if

you provided an initial state, otherwise this element will be ignored. Note that vertices without labels cannot be fixed. The fixed status will be ignord for these with a warning. Also note that label numbers by themselves have no meaning, and igraph may renumber labels. However, co-membership constraints will be respected: two

vertices can be fixed to be in the same or in different communities.

modularity: If not a null pointer, then it must be a pointer to a real number. The modularity score

of the detected community structure is stored here. Note that igraph will calculate the *directed* modularity if the input graph is directed, even if you set *mode* to

IGRAPH_ALL

Returns:

Error code.

Time complexity: O(m+n)

Example 24.8. File examples/simple/igraph_community_label_propagation.c

The InfoMAP algorithm

igraph_community_infomap — Find community structure that minimizes the expected description length of a random walker trajectory.

Implementation of the InfoMap community detection algorithm of Martin Rosvall and Carl T. Bergstrom.

For more details, see the visualization of the math and the map generator at https://www.mapequation.org . The original paper describing the algorithm is: M. Rosvall and C. T. Bergstrom, Maps of information flow reveal community structure in complex networks, PNAS 105, 1118 (2008) (http://dx.doi.org/10.1073/pnas.0706851105, http://arxiv.org/abs/0707.0609). A more detailed paper about the algorithm is: M. Rosvall, D. Axelsson, and C. T. Bergstrom, The map equation, Eur. Phys. J. Special Topics 178, 13 (2009). (http://dx.doi.org/10.1140/epjst/e2010-01179-1, http://arxiv.org/abs/0906.1405)

The original C++ implementation of Martin Rosvall is used, see http://www.tp.umu.se/~rosvall/downloads/infomap_undir.tgz . Intergation in igraph was done by Emmanuel Navarro (who is grateful to Martin Rosvall and Carl T. Bergstrom for providing this source code.)

Note that the graph must not contain isolated vertices.

If you want to specify a random seed (as in the original implementation) you can use $igraph_rng_seed()$.

Arguments:

graph: The input graph.
e_weights: Numeric vector giving the weights of the edges. If it is a NULL pointer then all edges will have equal weights. The weights are expected to be positive.
v_weights: Numeric vector giving the weights of the vertices. If it is a NULL pointer then all vertices will have equal weights. The weights are expected to be positive.
nb_trials: The number of attempts to partition the network (can be any integer value equal or larger than 1).
membership: Pointer to a vector. The membership vector is stored here.
codelength: Pointer to a real. If not NULL the code length of the partition is stored here.

Returns:

Error code.

See also:

```
igraph\_community\_spinglass(), \quad igraph\_community\_edge\_betweenness(), \\ igraph\_community\_walktrap().
```

Time complexity: TODO.

Chapter 25. Graphlets

Introduction

Graphlet decomposition models a weighted undirected graph via the union of potentially overlapping dense social groups. This is done by a two-step algorithm. In the first step, a candidate set of groups (a candidate basis) is created by finding cliques in the thresholded input graph. In the second step, the graph is projected onto the candidate basis, resulting in a weight coefficient for each clique in the candidate basis.

For more information on graphlet decomposition, see Hossein Azari Soufiani and Edoardo M Airoldi: "Graphlet decomposition of a weighted network", https://arxiv.org/abs/1203.2821 and http://proceedings.mlr.press/v22/azari12/azari12.pdf

igraph contains three functions for performing the graphlet decomponsition of a graph. The first is igraph_graphlets(), which performs both steps of the method and returns a list of subgraphs with their corresponding weights. The other two functions correspond to the first and second steps of the algorithm, and they are useful if the user wishes to perform them individually: igraph_graphlets_candidate_basis() and igraph_graphlets_project().

Note: The term "graphlet" is used for several unrelated concepts in the literature. If you are looking to count induced subgraphs, see <code>igraph_motifs_randesu()</code> and <code>igraph_subisomor-phic_lad()</code>.

Performing graphlet decomposition

igraph_graphlets — Calculate graphlets basis and project the graph on it

This function simply calls <code>igraph_graphlets_candidate_basis()</code> and <code>igraph_graphlets_project()</code>, and then orders the graphlets according to decreasing weights.

Arguments:

graph: The input graph, it must be a simple graph, edge directions are ignored.

weights: Weights of the edges, a vector.

cliques: An initialized list of integer vectors. The graphlet basis is stored here. Each element of

the list is an integer vector of vertex IDs, encoding a single basis subgraph.

Mu: An initialized vector, the weights of the graphlets will be stored here.

niter: Integer scalar, the number of iterations to perform for the projection step.

Returns:

Error code.

See also: igraph_graphlets_candidate_basis() and igraph_graphlet-s_project().

igraph_graphlets_candidate_basis — Calculate a candidate graphlets basis

Arguments:

graph: The input graph, it must be a simple graph, edge directions are ignored.

weights: Weights of the edges, a vector.

cliques: An initialized list of integer vectors. The graphlet basis is stored here. Each element

of the list is an integer vector of vertex IDs, encoding a single basis subgraph.

thresholds: An initialized vector, the (highest possible) weight thresholds for finding the basis

subgraphs are stored here.

Returns:

Error code.

See also: igraph_graphlets() and igraph_graphlets_project().

igraph_graphlets_project — Project a graph on a graphlets basis

Note that the graph projected does not have to be the same that was used to calculate the graphlet basis, but it is assumed that it has the same number of vertices, and the vertex IDs of the two graphs match.

Arguments:

graph: The input graph, it must be a simple graph, edge directions are ignored.

weights: Weights of the edges in the input graph, a vector.

cliques: An initialized list of integer vectors. The graphlet basis is stored here. Each element of

the list is an integer vector of vertex IDs, encoding a single basis subgraph.

Mu: An initialized vector, the weights of the graphlets will be stored here. This vector is

also used to initialize the the weight vector for the iterative algorithm, if the $\mathtt{startMu}$

argument is true (non-zero).

Graphlets

startMu: If true (non-zero), then the supplied Mu vector is used as the starting point of the iter-

ation. Otherwise a constant 1 vector is used.

niter: Integer scalar, the number of iterations to perform.

Returns:

Error code.

See also: $igraph_graphlets()$ and $igraph_graphlets_candidate_basis()$.

Chapter 26. Hierarchical random graphs

Introduction

A hierarchical random graph is an ensemble of undirected graphs with n vertices. It is defined via a binary tree with n leaf and n-1 internal vertices, where the internal vertices are labeled with probabilities. The probability that two vertices are connected in the random graph is given by the probability label at their closest common ancestor.

Please read the following two articles for more about hierarchical random graphs: A. Clauset, C. Moore, and M.E.J. Newman. Hierarchical structure and the prediction of missing links in networks. Nature 453, 98 - 101 (2008); and A. Clauset, C. Moore, and M.E.J. Newman. Structural Inference of Hierarchies in Networks. In E. M. Airoldi et al. (Eds.): ICML 2006 Ws, Lecture Notes in Computer Science 4503, 1-13. Springer-Verlag, Berlin Heidelberg (2007).

igraph contains functions for fitting HRG models to a given network (igraph_hrg_fit), for generating networks from a given HRG ensemble (igraph_hrg_game, igraph_hrg_sample), converting an igraph graph to a HRG and back (igraph_hrg_create, igraph_hrg_dendrogram), for calculating a consensus tree from a set of sampled HRGs (igraph_hrg_consensus) and for predicting missing edges in a network based on its HRG models (igraph_hrg_predict).

The igraph HRG implementation is heavily based on the code published by Aaron Clauset, at his website, http://tuvalu.santafe.edu/~aaronc/hierarchy/

Representing HRGs

igraph_hrg_t — Data structure to store a hierarchical random graph

```
typedef struct igraph_hrg_t {
    igraph_vector_int_t left;
    igraph_vector_int_t right;
    igraph_vector_t prob;
    igraph_vector_int_t vertices;
    igraph_vector_int_t edges;
} igraph_hrg_t;
```

A hierarchical random graph (HRG) can be given as a binary tree, where the internal vertices are labeled with real numbers.

Note that you don't necessarily have to know this internal representation for using the HRG functions, just pass the HRG objects created by one igraph function, to another igraph function.

It has the following members:

Values:

left:

Vector that contains the left children of the internal tree vertices. The first vertex is always the root vertex, so the first element of the vector is the left child of the root vertex. Internal vertices are denoted with negative numbers, starting from -1 and going

down, i.e. the root vertex is -1. Leaf vertices are denoted by non-negative number,

starting from zero and up.

right: Vector that contains the right children of the vertices, with the same encoding as the

left vector.

prob: The connection probabilities attached to the internal vertices, the first number belongs

to the root vertex (i.e. internal vertex -1), the second to internal vertex -2, etc.

edges: The number of edges in the subtree below the given internal vertex.

vertices: The number of vertices in the subtree below the given internal vertex, including itself.

igraph_hrg_init — Allocate memory for a HRG.

igraph_error_t igraph_hrg_init(igraph_hrg_t *hrg, igraph_integer_t n);

This function must be called before passing an igraph_hrg_t to an igraph function.

Arguments:

hrg: Pointer to the HRG data structure to initialize.

n: The number of vertices in the graph that is modeled by this HRG. It can be zero, if this is not yet known.

Returns:

Error code.

Time complexity: O(n), the number of vertices in the graph.

igraph_hrg_destroy — Deallocate memory for an HRG.

```
void igraph_hrg_destroy(igraph_hrg_t *hrg);
```

The HRG data structure can be reinitialized again with an igraph_hrg_destroy call.

Arguments:

hrg: Pointer to the HRG data structure to deallocate.

Time complexity: operating system dependent.

igraph_hrg_size — Returns the size of the HRG, the number of leaf nodes.

```
igraph_integer_t igraph_hrg_size(const igraph_hrg_t *hrg);
```

Arguments:

hrg: Pointer to the HRG.

Returns:

The number of leaf nodes in the HRG.

Time complexity: O(1).

igraph_hrg_resize — Resize a HRG.

```
igraph_error_t igraph_hrg_resize(igraph_hrg_t *hrg, igraph_integer_t newsize);
```

Arguments:

hrg: Pointer to an initialized (see igraph_hrg_init) HRG.

newsize: The new size, i.e. the number of leaf nodes.

Returns:

Error code.

Time complexity: O(n), n is the new size.

Fitting HRGs

igraph_hrg_fit — Fit a hierarchical random graph model to a network.

Arguments:

graph: The igraph graph to fit the model to. Edge directions are ignored in directed graphs.

hrg: Pointer to an initialized HRG, the result of the fitting is stored here. It can also be used

to pass a HRG to the function, that can be used as the starting point of the Markov Chain

Monte Carlo fitting, if the start argument is true.

start: Logical, whether to start the fitting from the given HRG model.

steps: Integer, the number of MCMC steps to take in the fitting procedure. If this is zero, then

the fitting stop is a convergence criteria is fulfilled.

Returns:

Error code.

Time complexity: TODO.

igraph_hrg_consensus — Calculate a consensus tree for a HRG.

The calculation can be started from the given HRG (hrg), or (if start is false), a HRG is first fitted to the given graph.

Arguments:

graph: The input graph.

parents: An initialized vector, the results are stored here. For each vertex, the id of its

parent vertex is stored, or -1, if the vertex is the root vertex in the tree. The first n vertex IDs (from 0) refer to the original vertices of the graph, the other IDs refer

to vertex groups.

weights: Numeric vector, counts the number of times a given tree split occured in the gen-

erated network samples, for each internal vertices. The order is the same as in

parents.

hrg: A hierarchical random graph. It is used as a starting point for the sampling, if the

start argument is true. It is modified along the MCMC.

start: Logical, whether to use the supplied HRG (in hrg) as a starting point for the

MCMC.

num_samples: The number of samples to generate for creating the consensus tree.

Returns:

Error code.

Time complexity: TODO.

HRG sampling

igraph_hrg_sample — Sample from a hierarchical random graph model.

```
igraph_error_t igraph_hrg_sample(const igraph_hrg_t *hrg, igraph_t *sample);
```

This function draws a single sample from a hierarchical random graph model.

Arguments:

hrg: A HRG model to sample from

sample: Pointer to an uninitialized graph; the sample is stored here.

Returns:

Error code.

Time complexity: TODO.

igraph_hrg_game — Generate a hierarchical random graph.

This function is a simple shortcut to igraph_hrg_sample. It creates a single graph from the given HRG.

Arguments:

graph: Pointer to an uninitialized graph, the new graph is created here.

hrg: The hierarchical random graph model to sample from.

Returns:

Error code.

Time complexity: TODO.

Conversion to and from igraph graphs

igraph_hrg_dendrogram — Create a dendrogram from a hierarchical random graph.

```
igraph_error_t igraph_hrg_dendrogram(
    igraph_t *graph, const igraph_hrg_t *hrg
);
```

Creates the igraph graph equivalent of an igraph_hrg_t data structure.

Arguments:

graph: Pointer to an uninitialized graph, the result is stored here.

hrg: The hierarchical random graph to convert.

Returns:

Error code.

Time complexity: O(n), the number of vertices in the graph.

igraph_hrg_create — Create a HRG from an igraph graph.

Arguments:

hrg: Pointer to an initialized igraph_hrg_t. The result is stored here.

graph: The igraph graph to convert. It must be a directed binary tree, with n-1 internal and n leaf

vertices. The root vertex must have in-degree zero.

prob: The vector of probabilities, this is used to label the internal nodes of the hierarchical ran-

dom graph.

Returns:

Error code.

Time complexity: O(n), the number of vertices in the tree.

Predicting missing edges

igraph_hrg_predict — Predict missing edges in a graph, based on HRG models.

Samples HRG models for a network, and estimated the probability that an edge was falsely observed as non-existent in the network.

Arguments:

graph: The input graph.

edges: The list of missing edges is stored here, the first two elements are the first edge,

the next two the second edge, etc.

prob: Vector of probabilies for the existence of missing edges, in the order correspond-

ing to edges.

hrg: A HRG, it is used as a starting point if start is true. It is also modified during

the MCMC sampling.

start: Logical, whether to start the MCMC from the given HRG.

num_samples: The number of samples to generate.

num_bins: Controls the resolution of the edge probabilities. Higher numbers result higher

resolution.

Returns:

Error code.

Time complexity: TODO.

Chapter 27. Embedding of graphs

Spectral embedding

igraph_adjacency_spectral_embedding — Adjacency spectral embedding

Spectral decomposition of the adjacency matrices of graphs. This function computes an n-dimensional Euclidean representation of the graph based on its adjacency matrix, A. This representation is computed via the singular value decomposition of the adjacency matrix, A=U D V^T. In the case, where the graph is a random dot product graph generated using latent position vectors in R^n for each vertex, the embedding will provide an estimate of these latent vectors.

For undirected graphs, the latent positions are calculated as $X = U^n D^{(1/2)}$ where U^n equals to the first no columns of U, and $D^{(1/2)}$ is a diagonal matrix containing the square root of the selected singular values on the diagonal.

For directed graphs, the embedding is defined as the pair $X = U^n D^(1/2)$, $Y = V^n D^(1/2)$. (For undirected graphs U=V, so it is sufficient to keep one of them.)

Arguments:

graph: The input graph, can be directed or undirected.

n: An integer scalar. This value is the embedding dimension of the spectral embedding.

Should be smaller than the number of vertices. The largest n-dimensional non-zero

singular values are used for the spectral embedding.

weights: Optional edge weights. Supply a null pointer for unweighted graphs.

which: Which eigenvalues (or singular values, for directed graphs) to use, possible values:

IGRAPH_EIGEN_LM the ones with the largest magnitude

IGRAPH_EIGEN_LA the (algebraic) largest ones

IGRAPH EIGEN SA the (algebraic) smallest ones.

For directed graphs, IGRAPH_EIGEN_LM and IGRAPH_EIGEN_LA are the same be-

cause singular values are used for the ordering instead of eigenvalues.

scaled: Whether to return X and Y (if scaled is true), or U and V.

X: Initialized matrix, the estimated latent positions are stored here.

Y: Initialized matrix or a null pointer. If not a null pointer, then the second half of the latent

positions are stored here. (For undirected graphs, this always equals X.)

D: Initialized vector or a null pointer. If not a null pointer, then the eigenvalues (for undi-

rected graphs) or the singular values (for directed graphs) are stored here.

cvec: A numeric vector, its length is the number vertices in the graph. This vector is added to

the diagonal of the adjacency matrix, before performing the SVD.

options: Options to ARPACK. See igraph_arpack_options_t for details. Supply NULL

to use the defaults. Note that the function overwrites the n (number of vertices), nev and which parameters and it always starts the calculation from a random start vector.

Returns:

Error code.

igraph_laplacian_spectral_embedding — Spectral embedding of the Laplacian of a graph

This function essentially does the same as igraph_adjacency_spectral_embedding, but works on the Laplacian of the graph, instead of the adjacency matrix.

Arguments:

graph: The input graph.

n: The number of eigenvectors (or singular vectors if the graph is directed) to use for the

embedding.

weights: Optional edge weights. Supply a null pointer for unweighted graphs.

which: Which eigenvalues (or singular values, for directed graphs) to use, possible values:

IGRAPH_EIGEN_LM the ones with the largest magnitude

IGRAPH_EIGEN_LA the (algebraic) largest ones

IGRAPH_EIGEN_SA the (algebraic) smallest ones.

For directed graphs, IGRAPH_EIGEN_LM and IGRAPH_EIGEN_LA are the same be-

cause singular values are used for the ordering instead of eigenvalues.

type: The type of the Laplacian to use. Various definitions exist for the Laplacian of a graph,

and one can choose between them with this argument. Possible values:

IGRAPH_EMBEDDING_D_A means D - A where D is the degree matrix and A is

the adjacency matrix

IGRAPH_EMBEDDING_DAD means Di times A times Di, where Di is the inverse

of the square root of the degree matrix;

IGRAPH_EMBEDDING_I_DAD means I - Di A Di, where I is the identity matrix.

scaled: Whether to return X and Y (if scaled is true), or U and V.

X: Initialized matrix, the estimated latent positions are stored here.

Y: Initialized matrix or a null pointer. If not a null pointer, then the second half of the latent

positions are stored here. (For undirected graphs, this always equals X.)

D: Initialized vector or a null pointer. If not a null pointer, then the eigenvalues (for undi-

rected graphs) or the singular values (for directed graphs) are stored here.

options: Options to ARPACK. See igraph_arpack_options_t for details. Supply NULL

to use the defaults. Note that the function overwrites the n (number of vertices), nev and which parameters and it always starts the calculation from a random start vector.

Returns:

Error code.

See also:

igraph_adjacency_spectral_embedding to embed the adjacency matrix.

igraph_dim_select — Dimensionality selection.

igraph_error_t igraph_dim_select(const igraph_vector_t *sv, igraph_integer_t *d

Dimensionality selection for singular values using profile likelihood.

The input of the function is a numeric vector which contains the measure of "importance" for each dimension.

For spectral embedding, these are the singular values of the adjacency matrix. The singular values are assumed to be generated from a Gaussian mixture distribution with two components that have different means and same variance. The dimensionality d is chosen to maximize the likelihood when the d largest singular values are assigned to one component of the mixture and the rest of the singular values assigned to the other component.

This function can also be used for the general separation problem, where we assume that the left and the right of the vector are coming from two normal distributions, with different means, and we want to know their border.

Arguments:

sv: A numeric vector, the ordered singular values.

dim: The result is stored here.

Returns:

Error code.

Time complexity: O(n), n is the number of values in sv.

See also:

 $\verb|igraph_adjacency_spectral_embedding()|.$

Chapter 28. Graph operators

Union and intersection

igraph_disjoint_union — Creates the union of two disjoint graphs.

First the vertices of the second graph will be relabeled with new vertex IDs to have two disjoint sets of vertex IDs, then the union of the two graphs will be formed. If the two graphs have |V1| and |V2| vertices and |E1| and |E2| edges respectively then the new graph will have |V1|+|V2| vertices and |E1|+|E2| edges.

Both graphs need to have the same directedness, i.e. either both directed or both undirected.

The current version of this function cannot handle graph, vertex and edge attributes, they will be lost.

Arguments:

res: Pointer to an uninitialized graph object, the result will stored here.

left: The first graph.

right: The second graph.

Returns:

Error code.

See also:

igraph_disjoint_union_many() for creating the disjoint union of more than two graphs, igraph_union() for non-disjoint union.

Time complexity: O(|V1|+|V2|+|E1|+|E2|).

Example 28.1. File examples/simple/igraph_disjoint_union.c

igraph_disjoint_union_many — The disjint union of many graphs.

First the vertices in the graphs will be relabeled with new vertex IDs to have pairwise disjoint vertex ID sets and then the union of the graphs is formed. The number of vertices and edges in the result is the total number of vertices and edges in the graphs.

All graphs need to have the same directedness, i.e. either all directed or all undirected. If the graph list has length zero, the result will be a *directed* graph with no vertices.

The current version of this function cannot handle graph, vertex and edge attributes, they will be lost.

Arguments:

res: Pointer to an uninitialized graph object, the result of the operation will be stored here.

graphs: Pointer vector, contains pointers to initialized graph objects.

Returns:

Error code.

See also:

igraph_disjoint_union() for an easier syntax if you have only two graphs, igraph_union_many() for non-disjoint union.

Time complexity: O(|V|+|E|), the number of vertices plus the number of edges in the result.

igraph_union — Calculates the union of two graphs.

The number of vertices in the result is that of the larger graph from the two arguments. The result graph contains edges which are present in at least one of the operand graphs.

Arguments:

res: Pointer to an uninitialized graph object, the result will be stored here.

left: The first graph.

right: The second graph.

edge_map1: Pointer to an initialized vector or a null pointer. If not a null pointer, it will contain

a mapping from the edges of the first argument graph (left) to the edges of the

result graph.

edge_map2: The same as edge_map1, but for the second graph, right.

Returns:

Error code.

See also:

```
igraph_union_many() for the union of many graphs, igraph_intersection() and igraph difference() for other operators.
```

Time complexity: O(|V|+|E|), |V| is the number of vertices, |E| the number of edges in the result graph.

Example 28.2. File examples/simple/igraph_union.c

igraph_union_many — Creates the union of many graphs.

```
igraph_error_t igraph_union_many(
    igraph_t *res, const igraph_vector_ptr_t *graphs,
    igraph_vector_int_list_t *edgemaps
);
```

The result graph will contain as many vertices as the largest graph among the arguments does, and an edge will be included in it if it is part of at least one operand graph.

The directedness of the operand graphs must be the same. If the graph list has length zero, the result will be a *directed* graph with no vertices.

Arguments:

res: Pointer to an uninitialized graph object, this will contain the result.

graphs: Pointer vector, contains pointers to the operands of the union operator, graph objects

of course.

edgemaps: If not a null pointer, then it must be an initialized list of integer vectors, and the map-

pings of edges from the graphs to the result graph will be stored here, in the same order as *graphs*. Each mapping is stored in a separate igraph_vector_int_t object.

Returns:

Error code.

See also:

```
igraph_union() for the union of two graphs, igraph_intersection_many(),
igraph_intersection() and igraph_difference for other operators.
```

Time complexity: O(|V|+|E|), |V| is the number of vertices in largest graph and |E| is the number of edges in the result graph.

Example 28.3. File examples/simple/igraph union.c

igraph_intersection — Collect the common edges from two graphs.

The result graph contains only edges present both in the first and the second graph. The number of vertices in the result graph is the same as the larger from the two arguments.

Arguments:

res: Pointer to an uninitialized graph object. This will contain the result of the operation.

left: The first operand, a graph object.

right: The second operand, a graph object.

edge_map1: Null pointer, or an initialized vector. If the latter, then a mapping from the edges of

the result graph, to the edges of the *left* input graph is stored here.

edge_map2: Null pointer, or an initialized vector. The same as edge_map1, but for the right

input graph.

Returns:

Error code.

See also:

igraph_intersection_many() to calculate the intersection of many graphs at once, igraph_union(), igraph_difference() for other operators.

Time complexity: O(|V|+|E|), |V| is the number of nodes, |E| is the number of edges in the smaller graph of the two. (The one containing less vertices is considered smaller.)

Example 28.4. File examples/simple/igraph_intersection.c

igraph_intersection_many — The intersection of more than two graphs.

```
igraph_error_t igraph_intersection_many(
    igraph_t *res, const igraph_vector_ptr_t *graphs,
    igraph_vector_int_list_t *edgemaps
);
```

This function calculates the intersection of the graphs stored in the *graphs* argument. Only those edges will be included in the result graph which are part of every graph in *graphs*.

The number of vertices in the result graph will be the maximum number of vertices in the argument graphs.

Arguments:

res: Pointer to an uninitialized graph object, the result of the operation will be stored here.

graphs: Pointer vector, contains pointers to graphs objects, the operands of the intersection

operator.

edgemaps: If not a null pointer, then it must be an initialized list of integer vectors, and the map-

pings of edges from the graphs to the result graph will be stored here, in the same order as graphs. Each mapping is stored in a separate igraph_vector_int_t object.

For the edges that are not in the intersection, -1 is stored.

Returns:

Error code.

See also:

```
igraph_intersection() for the intersection of two graphs, igraph_union_many(),
igraph_union() and igraph_difference() for other operators.
```

Time complexity: O(|V|+|E|), |V| is the number of vertices, |E| is the number of edges in the smallest graph (i.e. the graph having the less vertices).

Other set-like operators

igraph_difference — Calculates the difference of two graphs.

The number of vertices in the result is the number of vertices in the original graph, i.e. the left, first operand. In the results graph only edges will be included from orig which are not present in sub.

Arguments:

res: Pointer to an uninitialized graph object, the result will be stored here.

orig: The left operand of the operator, a graph object.

sub: The right operand of the operator, a graph object.

Returns:

Error code.

See also:

```
igraph_intersection() and igraph_union() for other operators.
```

Time complexity: O(|V|+|E|), |V| is the number vertices in the smaller graph, |E| is the number of edges in the result graph.

Example 28.5. File examples/simple/igraph difference.c

igraph_complementer — Creates the complementer of a graph.

The complementer graph means that all edges which are not part of the original graph will be included in the result.

Arguments:

res: Pointer to an uninitialized graph object.

graph: The original graph.

loops: Whether to add loop edges to the complementer graph.

Returns:

Error code.

See also:

```
igraph_union(), igraph_intersection() and igraph_difference().
```

Time complexity: O(|V|+|E1|+|E2|), |V| is the number of vertices in the graph, |E1| is the number of edges in the original and |E2| in the complementer graph.

Example 28.6. File examples/simple/igraph_complementer.c

igraph_compose — Calculates the composition of two graphs.

The composition of graphs contains the same number of vertices as the bigger graph of the two operands. It contains an (i,j) edge if and only if there is a k vertex, such that the first graphs contains an (i,k) edge and the second graph a (k,j) edge.

This is of course exactly the composition of two binary relations.

The two graphs must have the same directedness, otherwise the function returns with an error. Note that for undirected graphs the two relations are by definition symmetric.

Arguments:

res: Pointer to an uninitialized graph object, the result will be stored here.

g1: The firs operand, a graph object.

g2: The second operand, another graph object.

edge_map1: If not a null pointer, then it must be a pointer to an initialized vector, and a mapping

from the edges of the result graph to the edges of the first graph is stored here.

edge_map1: If not a null pointer, then it must be a pointer to an initialized vector, and a mapping

from the edges of the result graph to the edges of the second graph is stored here.

Returns:

Error code.

Time complexity: O(|V|*d1*d2), |V| is the number of vertices in the first graph, d1 and d2 the average degree in the first and second graphs.

Example 28.7. File examples/simple/igraph_compose.c

Miscellaneous operators

igraph_connect_neighborhood — Graph power: connect each vertex to its neighborhood.

This function adds new edges to the input graph. Each vertex is connected to all vertices reachable by at most *order* steps from it (unless a connection already existed). In other words, the *order* power of the graph is computed.

Note that the input graph is modified in place, no new graph is created. Call <code>igraph_copy()</code> if you want to keep the original graph as well.

For undirected graphs reachability is always symmetric: if vertex A can be reached from vertex B in at most *order* steps, then the opposite is also true. Only one undirected (A,B) edge will be added in this case.

Arguments:

graph: The input graph, this is the output graph as well.

order: Integer constant, it gives the distance within which the vertices will be connected to the

source vertex.

mode: Constant, it specifies how the neighborhood search is performed for directed graphs.

If IGRAPH_OUT then vertices reachable from the source vertex will be connected, IGRAPH_IN is the opposite. If IGRAPH_ALL then the directed graph is considered as

an undirected one.

Returns:

Error code.

See also:

 $\verb|igraph_square_lattice|() | uses this function to connect the neighborhood of the vertices.$

Time complexity: $O(|V|*d^k)$, |V| is the number of vertices in the graph, d is the average degree and k is the *order* argument.

igraph_contract_vertices — Replace multiple vertices with a single one.

This function modifies the graph by merging several vertices into one. The vertices in the modified graph correspond to groups of vertices in the input graph. No edges are removed, thus the modified graph will typically have self-loops (corresponding to in-group edges) and multi-edges (corresponding to multiple connections between two groups). Use <code>igraph_simplify()</code> to eliminate self-loops and merge multi-edges.

Arguments:

graph: The input graph. It will be modified in-place.

mapping: A vector giving the mapping. For each vertex in the original graph, it should

contain its desired ID in the result graph. In order to create "orphan vertices" that have no corresponding vertices in the original graph, ensure that the IDs are

consecutive integers starting from zero.

vertex_comb: What to do with the vertex attributes. NULL means that vertex attributes are not

kept after the contraction (not even for unaffected vertices). See the igraph manual

section about attributes for details.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number or vertices plus edges.

igraph_induced_subgraph — Creates a subgraph induced by the specified vertices.

This function collects the specified vertices and all edges between them to a new graph. As the vertex IDs in a graph always start with zero, this function very likely needs to reassign IDs to the vertices.

Arguments:

graph: The graph object.

res: The subgraph, another graph object will be stored here, do not initialize this object before

calling this function, and call igraph_destroy() on it if you don't need it any more.

vids: A vertex selector describing which vertices to keep.

imp1: This parameter selects which implementation should we use when constructing the new

graph. Basically there are two possibilities: IGRAPH_SUBGRAPH_COPY_AND_DELETE copies the existing graph and deletes the vertices that are not needed in the new graph, while IGRAPH_SUBGRAPH_CREATE_FROM_SCRATCH constructs the new graph from scratch without copying the old one. The latter is more efficient if you are extracting a relatively small subpart of a very large graph, while the former is better if you want to extract a subgraph whose size is comparable to the size of the whole graph. There is a third possibility: IGRAPH_SUBGRAPH_AUTO will select one of the two methods automatically

based on the ratio of the number of vertices in the new and the old graph.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVVID, invalid vertex ID in *vids*.

Time complexity: O(|V|+|E|), |V| and |E| are the number of vertices and edges in the original graph.

See also:

igraph_delete_vertices() to delete the specified set of vertices from a graph, the opposite
of this function.

igraph_linegraph — Create the line graph of a graph.

```
igraph_error_t igraph_linegraph(const igraph_t *graph, igraph_t *linegraph);
```

The line graph L(G) of a G undirected graph is defined as follows. L(G) has one vertex for each edge in G and two different vertices in L(G) are connected by an edge if their corresponding edges share an end point. In a multigraph, if two end points are shared, two edges are created. The vertex of a loop is counted as two end points.

The line graph L(G) of a G directed graph is slightly different, L(G) has one vertex for each edge in G and two vertices in L(G) are connected by a directed edge if the target of the first vertex's corresponding edge is the same as the source of the second vertex's corresponding edge.

Edge i in the original graph will correspond to vertex i in the line graph.

The first version of this function was contributed by Vincent Matossian, thanks.

Arguments:

graph: The input graph, may be directed or undirected.

linegraph: Pointer to an uninitialized graph object, the result is stored here.

Returns:

Error code.

Time complexity: O(|V|+|E|), the number of edges plus the number of vertices.

igraph_simplify — Removes loop and/or multiple edges from the graph.

Arguments:

graph: The graph object.

multiple: Logical, if true, multiple edges will be removed.

loops: Logical, if true, loops (self edges) will be removed.

edge_comb: What to do with the edge attributes. NULL means to discard the edge attributes after

the operation, even for edges that were unaffeccted. See the igraph manual section

about attributes for details.

Returns:

Error code: IGRAPH_ENOMEM if we are out of memory.

Time complexity: O(|V|+|E|).

Example 28.8. File examples/simple/igraph_simplify.c

igraph_subgraph_edges — Creates a subgraph with the specified edges and their endpoints.

This function collects the specified edges and their endpoints to a new graph. As the vertex IDs in a graph always start with zero, this function very likely needs to reassign IDs to the vertices.

Arguments:

graph: The graph object.

res: The subgraph, another graph object will be stored here, do *not* initialize this

object before calling this function, and call igraph_destroy() on it if

you don't need it any more.

eids: An edge selector describing which edges to keep.

delete_vertices: Whether to delete the vertices not incident on any of the specified edges as

well. If false, the number of vertices in the result graph will always be

equal to the number of vertices in the input graph.

Returns:

Error code: IGRAPH_ENOMEM, not enough memory for temporary data. IGRAPH_EINVEID, invalid edge ID in eids.

Time complexity: O(|V|+|E|), |V| and |E| are the number of vertices and edges in the original graph.

See also:

igraph_delete_edges() to delete the specified set of edges from a graph, the opposite of this function.

igraph_reverse_edges — Reverses some edges of a directed graph.

igraph_error_t igraph_reverse_edges(igraph_t *graph, const igraph_es_t eids);

This functon reverses some edges of a directed graph. The modification is done in place. All attributes, as well as the ordering of edges and vertices are preserved.

Note that is rarely necessary to reverse *all* edges, as almost all functions that handle directed graphs take a mode argument that can be set to IGRAPH_IN to effectively treat edges as reversed.

Arguments:

graph: The graph whose edges will be reversed.

es: The edges to be reversed. Pass igraph_ess_all(IGRAPH_EDGEORDER_ID) to reverse all edges.

Returns:

Error code.

Time complexity: O(1) if all edges are reversed, otherwise O(|E|) where |E| is the number of edges in the graph.

Chapter 29. Using BLAS, LAPACK and ARPACK for igraph matrices and graphs

BLAS interface in igraph

BLAS is a highly optimized library for basic linear algebra operations such as vector-vector, matrix-vector and matrix-matrix product. Please see http://www.netlib.org/blas/ for details and a reference implementation in Fortran. igraph contains some wrapper functions that can be used to call BLAS routines in a somewhat more user-friendly way. Not all BLAS routines are included in igraph, and even those which are included might not have wrappers; the extension of the set of wrapped functions will probably be driven by igraph's internal requirements. The wrapper functions usually substitute double-precision floating point arrays used by BLAS with igraph_vector_t and igraph_matrix_t instances and also remove those parameters (such as the number of rows/columns) that can be inferred from the passed arguments directly.

igraph_blas_ddot — Dot product of two vectors.

Arguments:

v1: The first vector.

v2: The second vector.

res: Pointer to a real, the result will be stored here.

Time complexity: O(n) where n is the length of the vectors.

Example 29.1. File examples/simple/blas.c

igraph_blas_dnrm2 — Euclidean norm of a vector.

```
igraph_real_t igraph_blas_dnrm2(const igraph_vector_t *v);
```

Arguments:

v: The vector.

Returns:

Real value, the norm of v.

Time complexity: O(n) where n is the length of the vector.

igraph_blas_dgemv — Matrix-vector multiplication using BLAS, vector version.

This function is a somewhat more user-friendly interface to the dgemv function in BLAS. dgemv performs the operation y = alpha*A*x + beta*y, where x and y are vectors and A is an appropriately sized matrix (symmetric or non-symmetric).

Arguments:

transpose: whether to transpose the matrix A

alpha: the constant alpha

a: the matrix A

x: the vector x

beta: the constant beta

y: the vector y (which will be modified in-place)

Time complexity: O(nk) if the matrix is of size n x k

Returns:

IGRAPH_EOVERFLOW if the matrix is too large for BLAS, IGRAPH_SUCCESS otherwise.

See also:

igraph_blas_dgemv_array if you have arrays instead of vectors.

Example 29.2. File examples/simple/blas.c

igraph_blas_dgemm — Matrix-matrix multiplication using BLAS.

```
igraph_error_t igraph_blas_dgemm(igraph_bool_t transpose_a, igraph_bool_t transpose_a, igraph_boo
```

This function is a somewhat more user-friendly interface to the dgemm function in BLAS. dgemm calculates alpha*a*b + beta*c, where a, b and c are matrices, of which a and b can be transposed.

Arguments:

transpose a: whether to transpose the matrix a

transpose_b: whether to transpose the matrix b

alpha: the constant alpha

a: the matrix a

b: the matrix b

beta: the constant beta

c: the matrix c. The result will also be stored here. If beta is zero, c will be resized

to fit the result.

Time complexity: $O(n \ m \ k)$ where matrix a is of size $n \times k$, and matrix b is of size $k \times m$.

Returns:

IGRAPH_EOVERFLOW if the matrix is too large for BLAS, IGRAPH_EINVAL if the matrices have incompatible sizes, IGRAPH_SUCCESS otherwise.

Example 29.3. File examples/simple/blas_dgemm.c

igraph_blas_dgemv_array — Matrix-vector multiplication using BLAS, array version.

This function is a somewhat more user-friendly interface to the dgemv function in BLAS. dgemv performs the operation y = alpha*A*x + beta*y, where x and y are vectors and A is an appropriately sized matrix (symmetric or non-symmetric).

Arguments:

transpose: whether to transpose the matrix A

alpha: the constant alpha

a: the matrix A

x: the vector x as a regular C array

beta: the constant beta

y: the vector y as a regular C array (which will be modified in-place)

Time complexity: O(nk) if the matrix is of size n x k

Returns:

IGRAPH_EOVERFLOW if the matrix is too large for BLAS, IGRAPH_SUCCESS otherwise.

See also:

igraph_blas_dgemv if you have vectors instead of arrays.

LAPACK interface in igraph

LAPACK is written in Fortran90 and provides routines for solving systems of simultaneous linear equations, least-squares solutions of linear systems of equations, eigenvalue problems, and singular value problems. The associated matrix factorizations (LU, Cholesky, QR, SVD, Schur, generalized Schur) are also provided, as are related computations such as reordering of the Schur factorizations and estimating condition numbers. Dense and banded matrices are handled, but not general sparse matrices. In all areas, similar functionality is provided for real and complex matrices, in both single and double precision.

igraph provides an interface to a very limited set of LAPACK functions, using the regular igraph data structures.

See more about LAPACK at http://www.netlib.org/lapack/

Matrix factorization, solving linear systems

igraph_lapack_dgetrf — LU factorization of a general M-by-N
matrix.

The factorization has the form A = P * L * U where P is a permutation matrix, L is lower triangular with unit diagonal elements (lower trapezoidal if m > n), and U is upper triangular (upper trapezoidal if m < n).

Arguments:

- a: The input/output matrix. On entry, the M-by-N matrix to be factored. On exit, the factors L and U from the factorization A = P * L * U; the unit diagonal elements of L are not stored.
- *ipiv*: An integer vector, the pivot indices are stored here, unless it is a null pointer. Row i of the matrix was interchanged with row ipiv[i].
- info: LAPACK error code. Zero on successful exit. If its value is a positive number i, it indicates that U(i,i) is exactly zero. The factorization has been completed, but the factor U is exactly singular, and division by zero will occur if it is used to solve a system of equations. If LAPACK returns an error, i.e. a negative info value, then an igraph error is generated as well.

Returns:

Error code.

Time complexity: TODO.

igraph_lapack_dgetrs — Solve general system of linear equations using LU factorization.

This function calls LAPACK to solve a system of linear equations A * X = B or A' * X = B with a general N-by-N matrix A using the LU factorization computed by igraph_lapack_dgetrf.

Arguments:

transpose: Logical scalar, whether to transpose the input matrix.

a: A matrix containing the L and U factors from the factorization $A = P^*L^*U$. L is

expected to be unitriangular, diagonal entries are those of U. If A is singular, no

warning or error wil be given and random output will be returned.

ipiv: An integer vector, the pivot indices from igraph_lapack_dgetrf() must be

given here. Row i of A was interchanged with row ipiv[i].

b: The right hand side matrix must be given here. The solution will also be placed here.

Returns:

Error code.

Time complexity: TODO.

igraph_lapack_dgesv — Solve system of linear equations with LU factorization

This function computes the solution to a real system of linear equations A * X = B, where A is an N-by-N matrix and X and B are N-by-NRHS matrices.

The LU decomposition with partial pivoting and row interchanges is used to factor A as A = P * L * U, where P is a permutation matrix, L is unit lower triangular, and U is upper triangular. The factored form of A is then used to solve the system of equations A * X = B.

Arguments:

a: Matrix. On entry the N-by-N coefficient matrix, on exit, the factors L and U from the factorization A=P*L*U; the unit diagonal elements of L are not stored.

ipiv: An integer vector or a null pointer. If not a null pointer, then the pivot indices that define the permutation matrix P, are stored here. Row i of the matrix was interchanged with row IPIV(i).

b: Matrix, on entry the right hand side matrix should be stored here. On exit, if there was no error, and the info argument is zero, then it contains the solution matrix X.

info: The LAPACK info code. If it is positive, then U(info,info) is exactly zero. In this case the factorization has been completed, but the factor U is exactly singular, so the solution could not be computed.

Returns:

Error code.

Time complexity: TODO.

Example 29.4. File examples/simple/igraph lapack dgesv.c

Eigenvalues and eigenvectors of matrices

igraph_lapack_dsyevr — Selected eigenvalues and optionally eigenvectors of a symmetric matrix

Calls the DSYEVR LAPACK function to compute selected eigenvalues and, optionally, eigenvectors of a real symmetric matrix A. Eigenvalues and eigenvectors can be selected by specifying either a range of values or a range of indices for the desired eigenvalues.

See more in the LAPACK documentation.

Arguments: Matrix, on entry it contains the symmetric input matrix. Only the leading N-by-N A: upper triangular part is used for the computation. Constant that gives which eigenvalues (and possibly the corresponding eigenvectors) which: to calculate. Possible values are IGRAPH LAPACK DSYEV ALL, all eigenvalues; IGRAPH_LAPACK_DSYEV_INTERVAL, all eigenvalues in the half-open interval (vl,vu]; IGRAPH_LAPACK_DSYEV_SELECT, the il-th through iu-th eigenvalues. v1: If which is IGRAPH_LAPACK_DSYEV_INTERVAL, then this is the lower bound of the interval to be searched for eigenvalues. See also the *vestimate* argument. If which is IGRAPH_LAPACK_DSYEV_INTERVAL, then this is the upper bound vu: of the interval to be searched for eigenvalues. See also the vestimate argument. An upper bound for the number of eigenvalues in the (vl,vu) interval, if which is vestimate: IGRAPH_LAPACK_DSYEV_INTERVAL. Memory is allocated only for the given number of eigenvalues (and eigenvectors), so this upper bound must be correct. i1: The index of the smallest eigenvalue to return, if which is IGRAPH_LA-PACK_DSYEV_SELECT. The index of the largets eigenvalue to return, if which is IGRAPH_LA-111: PACK_DSYEV_SELECT. abstol: The absolute error tolerance for the eigevalues. An approximate eigenvalue is accepted as converged when it is determined to lie in an interval [a,b] of width less than or equal to abstol + EPS * max(|a|,|b|), where EPS is the machine precision. values: An initialized vector, the eigenvalues are stored here, unless it is a null pointer. It will be resized as needed. An initialized matrix, the eigenvectors are stored in its columns, unless it is a null vectors: pointer. It will be resized as needed. An integer vector. If not a null pointer, then it will be resized to (2*max(1,M)) (M support: is a the total number of eigenvalues found). Then the support of the eigenvectors in vectors is stored here, i.e., the indices indicating the nonzero elements in vec-

tors. The i-th eigenvector is nonzero only in elements support(2*i-1) through support(2*i).

Returns:

Error code.

Time complexity: TODO.

Example 29.5. File examples/simple/igraph_lapack_dsyevr.c

igraph_lapack_dgeev — Eigenvalues and optionally eigenvectors of a non-symmetric matrix

This function calls LAPACK to compute, for an N-by-N real nonsymmetric matrix A, the eigenvalues and, optionally, the left and/or right eigenvectors.

The right eigenvector v(j) of A satisfies A * v(j) = lambda(j) * v(j) where lambda(j) is its eigenvalue. The left eigenvector u(j) of A satisfies u(j) * H * A = lambda(j) * u(j) * H * H where u(j) * H denotes the conjugate transpose of u(j).

The computed eigenvectors are normalized to have Euclidean norm equal to 1 and largest component real.

Arguments:

A: matrix. On entry it contains the N-by-N input matrix.

valuesreal: Pointer to an initialized vector, or a null pointer. If not a null pointer, then the real

parts of the eigenvalues are stored here. The vector will be resized as needed.

valuesimag: Pointer to an initialized vector, or a null pointer. If not a null pointer, then the

imaginary parts of the eigenvalues are stored here. The vector will be resized

as needed.

vectorsleft: Pointer to an initialized matrix, or a null pointer. If not a null pointer, then the

left eigenvectors are stored in the columns of the matrix. The matrix will be

resized as needed.

vectorsright: Pointer to an initialized matrix, or a null pointer. If not a null pointer, then the

right eigenvectors are stored in the columns of the matrix. The matrix will be

resized as needed.

info: This argument is used for two purposes. As an input argument it gives whether

an igraph error should be generated if the QR algorithm fails to compute all eigenvalues. If *info* is non-zero, then an error is generated, otherwise only a warning is given. On exit it contains the LAPACK error code. Zero means successful exit. A negative values means that some of the arguments had an illegal value, this always triggers an igraph error. An i positive value means that the

QR algorithm failed to compute all the eigenvalues, and no eigenvectors have been computed; element i+1:N of valuesreal and valuesimag contain eigenvalues which have converged. This case only generates an igraph error, if info was non-zero on entry.

Returns:

Error code.

Time complexity: TODO.

Example 29.6. File examples/simple/igraph_lapack_dgeev.c

igraph_lapack_dgeevx — Eigenvalues/vectors of nonsymmetric matrices, expert mode

This function calculates the eigenvalues and optionally the left and/or right eigenvectors of a nonsymmetric N-by-N real matrix.

Optionally also, it computes a balancing transformation to improve the conditioning of the eigenvalues and eigenvectors (*ilo*, *ihi*, *scale*, and *abnrm*), reciprocal condition numbers for the eigenvalues (*rconde*), and reciprocal condition numbers for the right eigenvectors (*rcondv*).

The right eigenvector v(j) of A satisfies A * v(j) = lambda(j) * v(j) where lambda(j) is its eigenvalue. The left eigenvector u(j) of A satisfies $u(j)^A H * A = lambda(j) * u(j)^A H$ where $u(j)^A H$ denotes the conjugate transpose of u(j).

The computed eigenvectors are normalized to have Euclidean norm equal to 1 and largest component real.

Balancing a matrix means permuting the rows and columns to make it more nearly upper triangular, and applying a diagonal similarity transformation D * A * D^(-1), where D is a diagonal matrix, to make its rows and columns closer in norm and the condition numbers of its eigenvalues and eigenvectors smaller. The computed reciprocal condition numbers correspond to the balanced matrix. Permuting rows and columns will not change the condition numbers (in exact arithmetic) but diagonal scaling will. For further explanation of balancing, see section 4.10.2 of the LAPACK Users' Guide.

Arguments:

balance: Scalar that indicated, whether the input matrix should be balanced. Possible val-

ues:

IGRAPH_LA- no not diagonally scale or permute.

PACK_DGEEVX_BALANCE_NONE

matrices and graphs		
	IGRAPH_LA- PACK_DGEEVX_BALANCE_PERM	perform permutations to make the matrix more nearly upper triangular. Do not diagonally scale.
	IGRAPH_LA- PACK_DGEEVX_BALANCE_S- CALE	diagonally scale the matrix, i.e. replace A by D*A*D^(-1), where D is a diagonal matrix, chosen to make the rows and columns of A more equal in norm. Do not permute.
	IGRAPH_LA- PACK_DGEEVX_BALANCE_BOTH	both diagonally scale and permute A.
A:	The input matrix, must be square.	
valuesreal:	An initialized vector, or a NULL pointer. If not a NULL pointer, then the real parts of the eigenvalues are stored here. The vector will be resized, as needed.	
valuesimag:	An initialized vector, or a NULL pointer. If not a NULL pointer, then the imaginary parts of the eigenvalues are stored here. The vector will be resized, as needed.	
vectorsleft:	An initialized matrix or a NULL pointer. If not a null pointer, then the left eigenvectors are stored here. The order corresponds to the eigenvalues and the eigenvectors are stored in a compressed form. If the j-th eigenvalue is real then column j contains the corresponding eigenvector. If the j-th and $(j+1)$ -th eigenvalues form a complex conjugate pair, then the j-th and $(j+1)$ -th columns contain their corresponding eigenvectors.	
vectorsright:	An initialized matrix or a NULL pointer. If not a null pointer, then the right eigenvectors are stored here. The format is the same, as for the <code>vectorsleft</code> argument.	
ilo:		
ihi:	ilo and ihi are integer values determined when A was balanced. The balanced $A(i,j)=0$ if $I>J$ and $J=1,,ilo-1$ or $I=ihi+1,,N$.	
scale:	Pointer to an initialized vector or a NULL pointer. If not a NULL pointer, then details of the permutations and scaling factors applied when balancing A , are stored here. If $P(j)$ is the index of the row and column interchanged with row and column j , and $D(j)$ is the scaling factor applied to row and column j , then $scale(J) = P(J)$, for $J = 1, \ldots, ilo-1$	
	scale(J) = D(J), for $J = ilo,,ihi$	
	<pre>scale(J) = P(J) for J = ihi+1,,N. The order in which the interchanges</pre>	are made is N to <i>ihi</i> +1, then 1 to <i>ilo</i> -1.
abnrm:	Pointer to a real variable, the one-norm of the balanced matrix is stored here. (The one-norm is the maximum of the sum of absolute values of elements in any column.)	
rconde:	An initialized vector or a NULL pointer. If not a null pointer, then the reciprocal condition numbers of the eigenvalues are stored here.	

rcondv: An initialized vector or a NULL pointer. If not a null pointer, then the reciprocal

condition numbers of the right eigenvectors are stored here.

info: This argument is used for two purposes. As an input argument it gives whether

an igraph error should be generated if the QR algorithm fails to compute all eigenvalues. If <code>info</code> is non-zero, then an error is generated, otherwise only a warning is given. On exit it contains the LAPACK error code. Zero means successful exit. A negative values means that some of the arguments had an illegal value, this always triggers an igraph error. An i positive value means that the QR algorithm failed to compute all the eigenvalues, and no eigenvectors have been computed; element i+1:N of <code>valuesreal</code> and <code>valuesimag</code> contain eigenvalues which have converged. This case only generated an igraph error, if

info was non-zero on entry.

Returns:

Error code.

Time complexity: TODO

Example 29.7. File examples/simple/igraph_lapack_dgeevx.c

ARPACK interface in igraph

ARPACK is a library for solving large scale eigenvalue problems. The package is designed to compute a few eigenvalues and corresponding eigenvectors of a general n by n matrix A. It is most appropriate for large sparse or structured matrices A where structured means that a matrix-vector product w <- Av requires order n rather than the usual order n^2 floating point operations. Please see http:// www.caam.rice.edu/software/ARPACK/ for details.

The eigenvalue calculation in ARPACK (in the simplest case) involves the calculation of the Av product where A is the matrix we work with and v is an arbitrary vector. A user-defined function of type igraph_arpack_function_t is expected to perform this product. If the product can be done efficiently, e.g. if the matrix is sparse, then ARPACK is usually able to calculate the eigenvalues very quickly.

In igraph, eigenvalue/eigenvector calculations usually involve the following steps:

- 1. Initialization of an igraph_arpack_options_t data structure using igraph_arpack_options_init.
- 2. Setting some options in the initialized igraph_arpack_options_t object.
- 3. Defining a function of type igraph_arpack_function_t. The input of this function is a vector, and the output should be the output matrix multiplied by the input vector.
- 4. Calling igraph_arpack_rssolve() (is the matrix is symmetric), or igraph_arpack_rnsolve().

The igraph_arpack_options_t object can be used multiple times.

If we have many eigenvalue problems to solve, then it might worth to create an <code>igraph_arpack_s-torage_t</code> object, and initialize it via <code>igraph_arpack_storage_init()</code>. This structure contains all memory needed for ARPACK (with the given upper limit regerding to the size of the eigenvalue problem). Then many problems can be solved using the same <code>igraph_arpack_storage_t</code> object, without always reallocating the required memory. The <code>igraph_arpack_storage_t</code> object needs to be destroyed by calling <code>igraph_arpack_storage_destroy()</code> on it, when it is not needed any more.

igraph does not contain all ARPACK routines, only the ones dealing with symmetric and non-symmetric eigenvalue problems using double precision real numbers.

Data structures

igraph_arpack_options_t — Options for ARPACK

```
typedef struct igraph_arpack_options_t {
    /* INPUT */
   char bmat[1];
                          /* I-standard problem, G-generalized */
                          /* Dimension of the eigenproblem */
   int n;
   char which[2];
                         /* LA, SA, LM, SM, BE */
                         /* Number of eigenvalues to be computed */
   int nev;
   igraph_real_t tol;
                         /* Stopping criterion */
                          /* Number of columns in V */
   int ncv;
   int ldv;
                         /* Leading dimension of V */
   int ishift;
                         /* 0-reverse comm., 1-exact with tridiagonal */
                         /* Maximum number of update iterations to take */
   int mxiter;
   int nb;
                         /* Block size on the recurrence, only 1 works */
   int mode;
                          /* The kind of problem to be solved (1-5)
                              1: A*x=l*x, A symmetric
                              2: A*x=1*M*x, A symm. M pos. def.
                              3: K*x = 1*M*x, K symm., M pos. semidef.
                              4: K*x = 1*KG*x, K s. pos. semidef. KG s. indef.
                              5: A*x = 1*M*x, A symm., M symm. pos. semidef. *
   int start;
                          /* 0: random, 1: use the supplied vector */
                          /* Size of temporary storage, default is fine */
   int lworkl;
   igraph_real_t sigma;
                          /* The shift for modes 3,4,5 */
   igraph_real_t sigmai; /* The imaginary part of shift for rnsolve */
    /* OUTPUT */
   int info;
                          /* What happened, see docs */
   int ierr;
                         /* What happened in the dseupd call */
   int noiter;
                         /* The number of iterations taken */
   int nconv;
   int numop;
                         /* Number of OP*x operations */
   int numopb;
                          /* Number of B*x operations if BMAT='G' */
                          /* Number of steps of re-orthogonalizations */
   int numreo;
   /* INTERNAL */
   int iparam[11];
   int ipntr[14];
} igraph_arpack_options_t;
```

This data structure contains the options of the ARPACK eigenvalue solver routines. It must be initialized by calling <code>igraph_arpack_options_init()</code> on it. Then it can be used for multiple ARPACK calls, as the ARPACK solvers do not modify it. Input options:

Values:

bmat: Character. Whether to solve a standard ('I') ot a generalized problem ('B').

n: Dimension of the eigenproblem.

which: Specifies which eigenvalues/vectors to compute. Possible values for symmetric matrices:

LA Compute nev largest (algebraic) eigenvalues.

- SA Compute nev smallest (algebraic) eigenvalues.
- LM Compute nev largest (in magnitude) eigenvalues.
- SM Compute nev smallest (in magnitude) eigenvalues.
- BE Compute nev eigenvalues, half from each end of the spectrum. When nev is odd, compute one more from the high en than from the low end.

Possible values for non-symmetric matrices:

- LM Compute nev largest (in magnitude) eigenvalues.
- SM Compute nev smallest (in magnitude) eigenvalues.
- LR Compute nev eigenvalues of largest real part.
- SR Compute nev eigenvalues of smallest real part.
- LI Compute nev eigenvalues of largest imaginary part.
- SI Compute nev eigenvalues of smallest imaginary part.

nev: The number of eigenvalues to be computed.

tol: Stopping criterion: the relative accuracy of the Ritz value is considered acceptable if its error is less than tol times its estimated value. If this is set to zero then machine precision is used.

ncv: Number of Lanczos vectors to be generated. Setting this to zero means that igraph_arpack_rssolve and igraph_arpack_rnsolve will determine a suitable value for ncv automatically.

ldv: Numberic scalar. It should be set to zero in the current igraph implementation.

ishift: Either zero or one. If zero then the shifts are provided by the user via reverse communication. If one then exact shifts with respect to the reduced tridiagonal matrix T. Please always set this to one.

mxiter: Maximum number of Arnoldi update iterations allowed.

nb: Blocksize to be used in the recurrence. Please always leave this on the default value, one.

mode: The type of the eigenproblem to be solved. Possible values if the input matrix is symmetric:

- 1. A*x=lambda*x, A is symmetric.
- 2. A*x=lambda*M*x, A is symmetric, M is symmetric positive definite.
- 3. K*x=lambda*M*x, K is symmetric, M is symmetric positive semi-definite.
- 4. K*x=lambda*KG*x, K is symmetric positive semi-definite, KG is symmetric indefinite.
- 5. A*x=lambda*M*x, A is symmetric, M is symmetric positive semi-definite. (Cayley transformed mode.)

Please note that only mode ==1 was tested and other values might not work properly. Possible values if the input matrix is not symmetric:

- 1. A*x=lambda*x.
- 2. A*x=lambda*M*x, M is symmetric positive definite.

- 3. A*x=lambda*M*x, M is symmetric semi-definite.
- 4. A*x=lambda*M*x, M is symmetric semi-definite.

Please note that only mode == 1 was tested and other values might not work properly.

start:

Whether to use the supplied starting vector (1), or use a random starting vector (0). The starting vector must be supplied in the first column of the vectors argument of the igraph_arpack_rssolve() of igraph_arpack_rnsolve() call.

Output options:

Values:

info: Error flag of ARPACK. Possible values:

- Normal exit.
- 1 Maximum number of iterations taken.
- No shifts could be applied during a cycle of the Implicitly restarted Arnoldi iteration. One possibility is to increase the size of nev relative to nev.

ARPACK can return other error flags as well, but these are converted to igraph errors, see igraph_error_type_t.

ierr: Error flag of the second ARPACK call (one eigenvalue computation usually involves

two calls to ARPACK). This is always zero, as other error codes are converted to igraph

errors.

noiter: Number of Arnoldi iterations taken.

nconv: Number of converged Ritz values. This represents the number of Ritz values that satisfy

the convergence critetion.

numop: Total number of matrix-vector multiplications.

numopb: Not used currently.

numreo: Total number of steps of re-orthogonalization.

Internal options:

Values:

lworkl: Do not modify this option.

sigma: The shift for the shift-invert mode.

sigmai: The imaginary part of the shift, for the non-symmetric or complex shift-invert mode.

iparam: Do not modify this option.

ipntr: Do not modify this option.

igraph_arpack_storage_t — Storage for ARPACK

```
typedef struct igraph_arpack_storage_t {
   int maxn, maxncv, maxldv;
```

Public members, do not modify them directly, these are considered to be read-only.

Values:

maxn: Maximum rank of matrix.

maxncv: Maximum NCV.

maxldv: Maximum LDV.

These members are considered to be private:

Values:

workl: Working memory.

workd: Working memory.

d: Memory for eigenvalues.

resid: Memory for residuals.

ax: Working memory.

select: Working memory.

di: Memory for eigenvalues, non-symmetric case only.

workev: Working memory, non-symmetric case only.

igraph_arpack_function_t — Type of the ARPACK callback function

Arguments:

to: Pointer to an igraph_real_t, the result of the matrix-vector product is expected to be

stored here.

from: Pointer to an igraph_real_t, the input matrix should be multiplied by the vector stored

here.

n: The length of the vector (which is the same as the order of the input matrix).

extra: Extra argument to the matrix-vector calculation function. This is coming from the igraph_arpack_rssolve() or igraph_arpack_rnsolve() function.

Returns:

Error code, if not zero, then the ARPACK solver considers this as an error, stops and calls the igraph error handler.

igraph_arpack_options_init — Initialize ARPACK options

```
void igraph_arpack_options_init(igraph_arpack_options_t *o);
```

Initializes ARPACK options, set them to default values. You can always pass the initialized igraph_arpack_options_t object to built-in igraph functions without any modification. The built-in igraph functions modify the options to perform their calculation, e.g. igraph_pagerank() always searches for the eigenvalue with the largest magnitude, regardless of the supplied value.

If you want to implement your own function involving eigenvalue calculation using ARPACK, however, you will likely need to set up the fields for yourself.

Arguments:

```
o: The igraph_arpack_options_t object to initialize.
```

Time complexity: O(1).

igraph_arpack_storage_init — Initialize ARPACK storage

You only need this function if you want to run multiple eigenvalue calculations using ARPACK, and want to spare the memory allocation/deallocation between each two runs. Otherwise it is safe to supply a null pointer as the storage argument of both igraph_arpack_rssolve() and igraph_arpack_rnsolve() to make memory allocated and deallocated automatically.

Don't forget to call the <code>igraph_arpack_storage_destroy()</code> function on the storage object if you don't need it any more.

Arguments:

s: The igraph_arpack_storage_t object to initialize.

maxn: The maximum order of the matrices.

maxncv: The maximum NCV parameter intended to use.

maxldv: The maximum LDV parameter intended to use.

symm: Whether symmetric or non-symmetric problems will be solved using this

igraph_arpack_storage_t. (You cannot use the same storage both with symmet-

ric and non-symmetric solvers.)

Returns:

Error code.

Time complexity: O(maxncv*(maxldv+maxn)).

igraph_arpack_storage_destroy — Deallocate ARPACK storage

void igraph_arpack_storage_destroy(igraph_arpack_storage_t *s);

Arguments:

s: The igraph_arpack_storage_t object for which the memory will be deallocated.

Time complexity: operating system dependent.

ARPACK solvers

igraph_arpack_rssolve — ARPACK solver for symmetric matrices.

This is the ARPACK solver for symmetric matrices. Please use <code>igraph_arpack_rnsolve()</code> for non-symmetric matrices.

Arguments:

fun: Pointer to an igraph_arpack_function_t object, the function that performs the

matrix-vector multiplication.

extra: An extra argument to be passed to fun.

options: An igraph_arpack_options_t object.

storage: An igraph_arpack_storage_t object, or a null pointer. In the latter case mem-

ory allocation and deallocation is performed automatically. Either this or the vectors argument must be non-null if the ARPACK iteration is started from a given starting

vector. If both are given *vectors* take precedence.

values: If not a null pointer, then it should be a pointer to an initialized vector. The eigenvalues

will be stored here. The vector will be resized as needed.

vectors: If not a null pointer, then it must be a pointer to an initialized matrix. The eigenvectors

will be stored in the columns of the matrix. The matrix will be resized as needed. Either this or the *vectors* argument must be non-null if the ARPACK iteration is started

from a given starting vector. If both are given vectors take precedence.

Returns:

Error code.

Time complexity: depends on the matrix-vector multiplication. Usually a small number of iterations is enough, so if the matrix is sparse and the matrix-vector multiplication can be done in O(n) time (the number of vertices), then the eigenvalues are found in O(n) time as well.

igraph_arpack_rnsolve — ARPACK solver for non-symmetric matrices.

Please always consider calling igraph_arpack_rssolve() if your matrix is symmetric, it is much faster. igraph_arpack_rnsolve() for non-symmetric matrices.

Note that ARPACK is not called for 2x2 matrices as an exact algebraic solution exists in these cases.

Arguments:

fun: Pointer to an igraph_arpack_function_t object, the function that performs the

matrix-vector multiplication.

extra: An extra argument to be passed to fun.

options: An igraph_arpack_options_t object.

storage: An igraph_arpack_storage_t object, or a null pointer. In the latter case mem-

ory allocation and deallocation is performed automatically.

values: If not a null pointer, then it should be a pointer to an initialized matrix. The (possibly

complex) eigenvalues will be stored here. The matrix will have two columns, the first column contains the real, the second the imaginary parts of the eigenvalues. The matrix

will be resized as needed.

vectors: If not a null pointer, then it must be a pointer to an initialized matrix. The eigenvectors

will be stored in the columns of the matrix. The matrix will be resized as needed. Note that real eigenvalues will have real eigenvectors in a single column in this matrix; however, complex eigenvalues come in conjugate pairs and the result matrix will store the eigenvector corresponding to the eigenvalue with *positive* imaginary part only. Since in this case the eigenvector is also complex, it will occupy *two* columns in the eigenvector matrix (the real and the imaginary parts, in this order). Caveat: if the eigenvalue vector returns only the eigenvalue with the *negative* imaginary part for a complex conjugate eigenvalue pair, the result vector will *still* store the eigenvector corresponding to the eigenvalue with the positive imaginary part (since this is how ARPACK works).

Returns:

Error code.

Time complexity: depends on the matrix-vector multiplication. Usually a small number of iterations is enough, so if the matrix is sparse and the matrix-vector multiplication can be done in O(n) time (the number of vertices), then the eigenvalues are found in O(n) time as well.

igraph_arpack_unpack_complex — Makes the result of the non-symmetric ARPACK solver more readable.

This function works on the output of igraph_arpack_rnsolve and brushes it up a bit: it only keeps nev eigenvalues/vectors and every eigenvector is stored in two columns of the vectors matrix.

The output of the non-symmetric ARPACK solver is somewhat hard to parse, as real eigenvectors occupy only one column in the matrix, and the complex conjugate eigenvectors are not stored at all (usually). The other problem is that the solver might return more eigenvalues than requested. The common use of this function is to call it directly after <code>igraph_arpack_rnsolve</code> with its <code>vectors</code> and <code>values</code> argument and <code>options->nev</code> as <code>nev</code>. This will add the vectors for eigenvalues with a negative imaginary part and return all vectors as 2 columns, a real and imaginary part.

Arguments:

vectors: The eigenvector matrix, as returned by igraph_arpack_rnsolve. It will be re-

sized, typically it will be larger.

values: The eigenvalue matrix, as returned by igraph_arpack_rnsolve. It will be re-

sized, typically extra, unneeded rows (=eigenvalues) will be removed.

nev: The number of eigenvalues/vectors to keep. Can be less or equal than the number orig-

inally requested from ARPACK.

Returns:

Error code.

Time complexity: linear in the number of elements in the vectors matrix.

Chapter 30. Bipartite, i.e. two-mode graphs

Bipartite networks in igraph

A bipartite network contains two kinds of vertices and connections are only possible between two vertices of different kinds. There are many natural examples, e.g. movies and actors as vertices and a movie is connected to all participating actors, etc.

igraph does not have direct support for bipartite networks, at least not at the C language level. In other words the igraph_t structure does not contain information about the vertex types. The C functions for bipartite networks usually have an additional input argument to graph, called types, a boolean vector giving the vertex types.

Most functions creating bipartite networks are able to create this extra vector, you just need to supply an initialized boolean vector to them.

Create two-mode networks

igraph_create_bipartite — Create a bipartite graph.

This is a simple wrapper function to create a bipartite graph. It does a little more than <code>igraph_create()</code>, e.g. it checks that the graph is indeed bipartite with respect to the given <code>types</code> vector. If there is an edge connecting two vertices of the same kind, then an error is reported.

Arguments:

graph: Pointer to an uninitialized graph object, the result is created here.

types: Boolean vector giving the vertex types. The length of the vector defines the number

of vertices in the graph.

edges: Vector giving the edges of the graph. The highest vertex ID in this vector must be

smaller than the length of the types vector.

directed: Boolean scalar, whether to create a directed graph.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

Example 30.1. File examples/simple/igraph_bipartite_create.c

igraph_full_bipartite — Create a full bipartite network.

A bipartite network contains two kinds of vertices and connections are only possible between two vertices of different kind. There are many natural examples, e.g. movies and actors as vertices and a movie is connected to all participating actors, etc.

igraph does not have direct support for bipartite networks, at least not at the C language level. In other words the igraph_t structure does not contain information about the vertex types. The C functions for bipartite networks usually have an additional input argument to graph, called types, a boolean vector giving the vertex types.

Most functions creating bipartite networks are able to create this extra vector, you just need to supply an initialized boolean vector to them.

Arguments:

graph: Pointer to an igraph_t object, the graph will be created here.

types: Pointer to a boolean vector. If not a null pointer, then the vertex types will be stored

here.

*n*1: Integer, the number of vertices of the first kind.

n2: Integer, the number of vertices of the second kind.

directed: Boolean, whether to create a directed graph.

mode: A constant that gives the type of connections for directed graphs. If IGRAPH_OUT,

then edges point from vertices of the first kind to vertices of the second kind; if IGRAPH_IN, then the opposite direction is realized; if IGRAPH_ALL, then mutual

edges will be created.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

See also:

igraph_full() for non-bipartite full graphs.

igraph_bipartite_game — Generate a bipartite random graph (similar to Erd#s-Rényi).

igraph_error_t igraph_bipartite_game(igraph_t *graph, igraph_vector_bool_t *typ

```
igraph_erdos_renyi_t type,
igraph_integer_t n1, igraph_integer_t n2,
igraph_real_t p, igraph_integer_t m,
igraph_bool_t directed, igraph_neimode_t mode);
```

Similarly to unipartite (one-mode) networks, we can define the G(n,p), and G(n,m) graph classes for bipartite graphs, via their generating process. In G(n,p) every possible edge between top and bottom vertices is realized with probability p, independently of the rest of the edges. In G(n,m), we uniformly choose m edges to realize.

Arguments:

graph: Pointer to an uninitialized igraph graph, the result is stored here.

types: Pointer to an initialized boolean vector, or a null pointer. If not a null pointer, then

the vertex types are stored here. Bottom vertices come first, n1 of them, then n2 top

vertices.

type: The type of the random graph, possible values:

IGRAPH_ERDOS_RENYI_GNM G(n,m) graph, m edges are selected uniformly ran-

domly in a graph with n vertices.

IGRAPH_ERDOS_RENYI_GNP G(n,p) graph, every possible edge is included in

the graph with probability p.

*n*1: The number of bottom vertices.

n2: The number of top verices.

p: The connection probability for G(n,p) graphs. It is ignored for G(n,m) graphs.

m: The number of edges for G(n,m) graphs. It is ignored for G(n,p) graphs.

directed: Boolean, whether to generate a directed graph. See also the mode argument.

mode: Specifies how to direct the edges in directed graphs. If it is IGRAPH_OUT, then di-

rected edges point from bottom vertices to top vertices. If it is IGRAPH_IN, edges point from top vertices to bottom vertices. IGRAPH_OUT and IGRAPH_IN do not generate mutual edges. If this argument is IGRAPH_ALL, then each edge direction is considered independently and mutual edges might be generated. This argument is

ignored for undirected graphs.

Returns:

Error code.

See also:

igraph_erdos_renyi_game.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

Incidence matrices

igraph_incidence — Creates a bipartite graph from an incidence matrix.

A bipartite (or two-mode) graph contains two types of vertices and edges always connect vertices of different types. An incidence matrix is an $n \times m$ matrix, n and m are the number of vertices of the two types, respectively. Nonzero elements in the matrix denote edges between the two corresponding vertices.

Note that this function can operate in two modes, depending on the *multiple* argument. If it is false, then a single edge is created for every non-zero element in the incidence matrix. If *multiple* is true, then the matrix elements are rounded up to the closest non-negative integer to get the number of edges to create between a pair of vertices.

This function does not create multiple edges if multiple is false, but might create some if it is true.

Arguments:

graph: Pointer to an uninitialized graph object.

types: Pointer to an initialized boolean vector, or a null pointer. If not a null pointer, then

the vertex types are stored here. It is resized as needed.

incidence: The incidence matrix.

directed: Gives whether to create an undirected or a directed graph.

mode: Specifies the direction of the edges in a directed graph. If IGRAPH_OUT, then edges

point from vertices of the first kind (corresponding to rows) to vertices of the second kind (corresponding to columns); if IGRAPH_IN, then the opposite direction is re-

alized; if IGRAPH_ALL, then mutual edges will be created.

multiple: How to interpret the incidence matrix elements. See details below.

Returns:

Error code.

Time complexity: O(n*m), the size of the incidence matrix.

igraph_get_incidence — Convert a bipartite graph into an incidence matrix.

Arguments:

graph: The input graph, edge directions are ignored.

types: Boolean vector containing the vertex types. All vertices in one part of the graph should

have type 0, the others type 1.

res: Pointer to an initialized matrix, the result is stored here. An element of the matrix gives

the number of edges (irrespectively of their direction) between the two corresponding vertices. The rows will correspond to vertices with type 0, the columns correspond to

vertices with type 1.

row_ids: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex IDs

(in the graph) corresponding to the rows of the result matrix are stored here.

col_ids: Pointer to an initialized vector or a null pointer. If not a null pointer, then the vertex IDs

corresponding to the columns of the result matrix are stored here.

Returns:

Error code.

Time complexity: O(n*m), n and m are number of vertices of the two different kind.

See also:

igraph_incidence() for the opposite operation.

Project two-mode graphs

igraph_bipartite_projection_size — Calculate the number of vertices and edges in the bipartite projections.

This function calculates the number of vertices and edges in the two projections of a bipartite network. This is useful if you have a big bipartite network and you want to estimate the amount of memory you would need to calculate the projections themselves.

Arguments:

graph: The input graph.

types: Boolean vector giving the vertex types of the graph.

vcount 1: Pointer to an igraph_integer_t, the number of vertices in the first projection is

stored here. May be NULL if not needed.

ecount 1: Pointer to an igraph_integer_t, the number of edges in the first projection is

stored here. May be NULL if not needed.

vcount2: Pointer to an igraph_integer_t, the number of vertices in the second projection

is stored here. May be NULL if not needed.

ecount 2: Pointer to an igraph_integer_t, the number of edges in the second projection is stored here. May be NULL if not needed.

Returns:

Error code.

See also:

igraph_bipartite_projection() to calculate the actual projection.

Time complexity: $O(|V|*d^2+|E|)$, |V| is the number of vertices, |E| is the number of edges, d is the average (total) degree of the graphs.

Example 30.2. File examples/simple/igraph_bipartite_projection.c

igraph_bipartite_projection — Create one or both projections of a bipartite (two-mode) network.

Creates one or both projections of a bipartite graph.

Arguments:

graph: The bipartite input graph. Directedness of the edges is ignored.

types: Boolean vector giving the vertex types of the graph.

proj1: Pointer to an uninitialized graph object, the first projection will be created here.

It a null pointer, then it is ignored, see also the probe1 argument.

proj2: Pointer to an uninitialized graph object, the second projection is created here,

if it is not a null pointer. See also the probe 1 argument.

multiplicity1: Pointer to a vector, or a null pointer. If not the latter, then the multiplicity of

the edges is stored here. E.g. if there is an A-C-B and also an A-D-B triple in the bipartite graph (but no more X, such that A-X-B is also in the graph), then

the multiplicity of the A-B edge in the projection will be 2.

multiplicity2: The same as multiplicity1, but for the other projection.

probe1: This argument can be used to specify the order of the projections in the re-

sulting list. When it is non-negative, then it is considered as a vertex ID and the projection containing this vertex will be the first one in the result. Setting this argument to a non-negative value implies that proj1 must be a non-null pointer. If you don't care about the ordering of the projections, pass -1 here.

Returns:

Error code.

See also:

igraph_bipartite_projection_size() to calculate the number of vertices and edges in the projections, without creating the projection graphs themselves.

Time complexity: $O(|V|*d^2+|E|)$, |V| is the number of vertices, |E| is the number of edges, d is the average (total) degree of the graphs.

Example 30.3. File examples/simple/igraph_bipartite_projection.c

Other operations on bipartite graphs

igraph_is_bipartite — Check whether a graph is bipartite.

This function checks whether a graph is bipartite. It tries to find a mapping that gives a possible division of the vertices into two classes, such that no two vertices of the same class are connected by an edge.

The existence of such a mapping is equivalent of having no circuits of odd length in the graph. A graph with loop edges cannot be bipartite.

Note that the mapping is not necessarily unique, e.g. if the graph has at least two components, then the vertices in the separate components can be mapped independently.

Arguments:

graph: The input graph.

res: Pointer to a boolean, the result is stored here.

types: Pointer to an initialized boolean vector, or a null pointer. If not a null pointer and a mapping

was found, then it is stored here. If not a null pointer, but no mapping was found, the

contents of this vector is invalid.

Returns:

Error code.

Time complexity: O(|V|+|E|), linear in the number of vertices and edges.

Chapter 31. Advanced igraph programming

Using igraph in multi-threaded programs

The igraph library is considered thread-safe if it has been compiled with thread-local storage enabled, i.e. the IGRAPH_ENABLE_TLS setting was toggled to ON and the current platform supports this feature. To check whether an igraph build is thread-safe, use the IGRAPH_THREAD_SAFE macro. When linking to external versions of igraph's dependencies, it is the responsibility of the user to check that these dependencies were also compiled to be thread-safe.

IGRAPH_THREAD_SAFE — Specifies whether igraph was built in thread-safe mode.

#define IGRAPH_THREAD_SAFE

This macro is defined to 1 if the current build of the igraph library is built in thread-safe mode, and 0 if it is not. A thread-safe igraph library attempts to use thread-local data structures instead of global ones, but note that this is not (and can not) be guaranteed for third-party libraries that igraph links to.

Thread-safe ARPACK library

Note that igraph is only thread-safe if it was built with the internal ARPACK library, i.e. the one that comes with igraph. The standard ARPACK library is not thread-safe.

Thread-safety of random number generators

The default random number generator that igraph uses is *not* guaranteed to be thread-safe. You need to set a different random number generator instance for every thread that you want to use igraph from. This is especially important if you set the seed of the random number generator to ensure reproducibility; sharing a random number generator between threads would break reproducibility as the order in which the various threads are scheduled is random, and therefore they would still receive random numbers in an unpredictable order from the shared random number generator.

Progress handlers

About progress handlers

It is often useful to report the progress of some long calculation, to allow the user to follow the computation and guess the total running time. A couple of igraph functions support this at the time of writing, hopefully more will support it in the future.

To see the progress of a computation, the user has to install a progress handler, as there is none installed by default. If an igraph function supports progress reporting, then it calls the installed progress handler periodically, and passes a percentage value to it, the percentage of computation already performed. To install a progress handler, you need to call <code>igraph_set_progress_handler()</code>. Currently there is a single pre-defined progress handler, called <code>igraph_progress_handler_stderr()</code>.

Setting up progress handlers

igraph_progress_handler_t — Type of progress handler functions

This is the type of the igraph progress handler functions. There is currently one such predefined function, igraph_progress_handler_stderr(), but the user can write and set up more sophisticated ones.

Arguments:

message: A string describing the function or algorithm that is reporting the progress. Current

igraph functions always use the name message argument if reporting from the same

function.

percent: Numeric, the percentage that was completed by the algorithm or function.

data: User-defined data. Current igraph functions that report progress pass a null pointer here.

Users can write their own progress handlers and functions with progress reporting, and

then pass some meaningfull context here.

Returns:

If the return value of the progress handler is not IGRAPH_SUCCESS, then igraph_progress() returns the error code IGRAPH_INTERRUPTED. The IGRAPH_PROGRESS() macro frees all memory and finishes the igraph function with error code IGRAPH_INTERRUPTED in this case.

igraph_set_progress_handler — Install a progress handler, or remove the current handler.

```
igraph_progress_handler_t *
igraph_set_progress_handler(igraph_progress_handler_t new_handler);
```

There is a single simple predefined progress handler: igraph_progress_handler_stderr().

Arguments:

new_handler: Pointer to a function of type igraph_progress_handler_t, the progress

handler function to install. To uninstall the current progress handler, this argument

can be a null pointer.

Returns:

Pointer to the previously installed progress handler function.

Time complexity: O(1).

igraph_progress_handler_stderr — A simple predefined progress handler.

This simple progress handler first prints *message*, and then the percentage complete value in a short message to standard error.

Arguments:

message: A string describing the function or algorithm that is reporting the progress. Current

igraph functions always use the same message argument if reporting from the same

function.

percent: Numeric, the percentage that was completed by the algorithm or function.

data: User-defined data. Current igraph functions that report progress pass a null pointer here.

Users can write their own progress handlers and functions with progress reporting, and

then pass some meaningfull context here.

Returns:

This function always returns with IGRAPH_SUCCESS.

Time complexity: O(1).

Invoking the progress handler

IGRAPH_PROGRESS — Report progress.

#define IGRAPH_PROGRESS(message, percent, data)

The standard way to report progress from an igraph function

Arguments:

message: A string, a textual message that references the calculation under progress.

percent: Numeric scalar, the percentage that is complete.

data: User-defined data, this can be used in user-defined progress handler functions, from

user-written igraph functions.

Returns:

If the progress handler returns with IGRAPH_INTERRUPTED, then this macro frees up the igraph allocated memory for temporary data and returns to the caller with IGRAPH_INTERRUPTED.

igraph_progress — Report progress

igraph_error_t igraph_progress(const char *message, igraph_real_t percent, void

Note that the usual way to report progress is the IGRAPH_PROGRESS macro, as that takes care of the return value of the progress handler.

Arguments:

message: A string describing the function or algorithm that is reporting the progress. Current

igraph functions always use the name message argument if reporting from the same

function.

percent: Numeric, the percentage that was completed by the algorithm or function.

data: User-defined data. Current igraph functions that report progress pass a null pointer here.

Users can write their own progress handlers and functions with progress reporting, and

then pass some meaningfull context here.

Returns:

If there is a progress handler installed and it does not return IGRAPH_SUCCESS, then IGRAPH_INTERRUPTED is returned.

Time complexity: O(1).

igraph_progressf — Report progress, printf-like version

igraph_error_t igraph_progressf(const char *message, igraph_real_t percent, voi
...);

This is a more flexible version of igraph_progress(), with a printf-like template string. First the template string is filled with the additional arguments and then igraph_progress() is called.

Note that there is an upper limit for the length of the message string, currently 1000 characters.

Arguments:

message: A string describing the function or algorithm that is reporting the progress. For this

function this is a template string, using the same syntax as the standard libc printf

function.

percent: Numeric, the percentage that was completed by the algorithm or function.

data: User-defined data. Current igraph functions that report progress pass a null pointer here.

Users can write their own progress handlers and functions with progress reporting, and

then pass some meaningfull context here.

. . .: Additional argument that were specified in the message argument.

Returns:

If there is a progress handler installed and it does not return IGRAPH_SUCCESS, then IGRAPH_INTERRUPTED is returned. \return

Writing progress handlers

To write a new progress handler, one needs to create a function of type igraph_progress_handler_t. The new progress handler can then be installed with the igraph_set_progress_handler() function.

One can assume that the first progress handler call from a calculation will be call with zero as the *percentage* argument, and the last call from a function will have 100 as the *percentage* argument. Note, however, that if an error happens in the middle of a computation, then the 100 percent call might be omitted.

Writing igraph functions with progress reporting

If you want to write a function that uses igraph and supports progress reporting, you need to include igraph_progress() calls in your function, usually via the IGRAPH_PROGRESS() macro.

It is good practice to always include a call to <code>igraph_progress()</code> with a zero <code>percentage</code> argument, before the computation; and another call with 100 <code>percentage</code> value after the computation is completed.

It is also good practice *not* to call <code>igraph_progress()</code> too often, as this would slow down the computation. It might not be worth to support progress reporting in functions with linear or log-linear time complexity, as these are fast, even with a large amount of data. For functions with quadratic or higher time complexity make sure that the time complexity of the progress reporting is constant or at least linear. In practice this means having at most O(n) progress checks and at most 100 <code>igraph_progress()</code> calls.

Multi-threaded programs

In multi-threaded programs, each thread has its own progress handler, if thread-local storage is supported and igraph is thread-safe. See the <code>IGRAPH_THREAD_SAFE</code> macro for checking whether an igraph build is thread-safe.

Status handlers

Status reporting

In addition to the possibility of reporting the progress of an igraph computation via igraph_progress(), it is also possible to report simple status messages from within igraph functions, without having to judge how much of the computation was performed already. For this one needs to install a status handler function.

Status handler functions must be of type igraph_status_handler_t and they can be install by a call to igraph_set_status_handler(). Currently there is a simple predefined status handler function, called igraph_status_handler_stderr(), but the user can define new ones.

igraph functions report their status via a call to the IGRAPH_STATUS() or the IGRAPH_STATUSF() macro.

Setting up status handlers

igraph_status_handler_t — The type of the igraph status handler functions

typedef igraph_error_t igraph_status_handler_t(const char *message, void *data)

Arguments:

message: The status message.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null

pointer here.

Returns:

Error code. The current calculation will abort if you return anything else than IGRAPH_SUCCESS here.

igraph_set_status_handler — Install of uninstall a status handler function.

```
igraph_status_handler_t *
igraph_set_status_handler(igraph_status_handler_t new_handler);
```

To uninstall the currently installed status handler, call this function with a null pointer.

Arguments:

new_handler: The status handler function to install.

Returns:

The previously installed status handler function.

Time complexity: O(1).

igraph_status_handler_stderr — A simple predefined status handler function.

```
igraph_error_t igraph_status_handler_stderr(const char *message, void *data);
```

A simple status handler function that writes the status message to the standard error.

Arguments:

message: The status message.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null

pointer here.

Returns:

Error code.

Time complexity: O(1).

Invoking the status handler

IGRAPH_STATUS — Report the status of an igraph function.

```
#define IGRAPH_STATUS(message, data)
```

Typically this function is called only a handful of times from an igraph function. E.g. if an algorithm has three major steps, then it is logical to call it three times, to signal the three major steps.

Arguments:

message: The status message.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null

pointer here.

Returns:

If the status handler returns with a value other than IGRAPH_SUCCESS, then the function that called this macro returns as well, with error code IGRAPH_INTERRUPTED.

IGRAPH_STATUSF — Report the status from an igraph function

```
#define IGRAPH_STATUSF(args)
```

This is the more flexible version of IGRAPH_STATUS(), having a printf-like syntax. As this macro takes variable number of arguments, they must be all supplied as a single argument, enclosed in parentheses. Then igraph_statusf() is called with the given arguments.

Arguments:

args: The arguments to pass to igraph_statusf().

Returns:

If the status handler returns with a value other than IGRAPH_SUCCESS, then the function that called this macro returns as well, with error code IGRAPH_INTERRUPTED.

igraph_status — Reports status from an igraph function.

```
igraph_error_t igraph_status(const char *message, void *data);
```

It calls the installed status handler function, if there is one. Otherwise it does nothing. Note that the standard way to report the status from an igraph function is the IGRAPH_STATUS or IGRAPH_STATUSF macro, as these take care of the termination of the calling function if the status handler returns with IGRAPH_INTERRUPTED.

Arguments:

message: The status message.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null

pointer here.

Returns:

Error code. If a status handler function was called and it did not return with IGRAPH_SUCCESS, then IGRAPH INTERRUPTED is returned by igraph status().

Time complexity: O(1).

igraph_statusf — Report status, more flexible printf-like version.

```
igraph_error_t igraph_statusf(const char *message, void *data, ...);
```

This is the more flexible version of <code>igraph_status()</code>, that has a syntax similar to the <code>printf</code> standard C library function. It substitutes the values of the additional arguments into the <code>message</code> template string and calls <code>igraph_status()</code>.

Arguments:

message: Status message template string, the syntax is the same as for the printf function.

data: Additional context, with user-defined semantics. Existing igraph functions pass a null

pointer here.

...: The additional arguments to fill the template given in the message argument.

Returns:

Error code. If a status handler function was called and it did not return with IGRAPH_SUCCESS, then IGRAPH_INTERRUPTED is returned by igraph_status().

Chapter 32. Non-graph related functions

igraph version number

igraph_version — Return the version of the igraph C library

Arguments:

version_string: Pointer to a string pointer. If not null, it is set to the igraph version string, e.g.

"0.6" or "0.5.3". This string should not be modified or deallocated.

major: If not a null pointer, then it is set to the major igraph version. E.g. for version

"0.5.3" this is 0.

minor: If not a null pointer, then it is set to the minor igraph version. E.g. for version

"0.5.3" this is 5.

subminor: If not a null pointer, then it is set to the subminor igraph version. E.g. for

version "0.5.3" this is 3.

Returns:

Error code.

Time complexity: O(1).

Example 32.1. File examples/simple/igraph_version.c

Running mean of a time series

igraph_running_mean — Calculates the running mean of a vector.

The running mean is defined by the mean of the previous binwidth values.

Arguments:

data: The vector containing the data.

res: The vector containing the result. This should be initialized before calling this function

and will be resized.

binwidth: Integer giving the width of the bin for the running mean calculation.

Returns:

Error code.

Time complexity: O(n), n is the length of the data vector.

Random sampling from very long sequences

igraph_random_sample — Generates an increasing random sequence of integers.

This function generates an increasing sequence of random integer numbers from a given interval. The algorithm is taken literally from (Vitter 1987). This method can be used for generating numbers from a *very* large interval. It is primarily created for randomly selecting some edges from the sometimes huge set of possible edges in a large graph.

Reference:

J. S. Vitter. An efficient algorithm for sequential random sampling. ACM Transactions on Mathematical Software, 13(1):58--67, 1987. https://doi.org/10.1145/23002.23003

Arguments:

res: Pointer to an initialized vector. This will hold the result. It will be resized to the proper

size.

1: The lower limit of the generation interval (inclusive). This must be less than or equal to

the upper limit, and it must be integral.

h: The upper limit of the generation interval (inclusive). This must be greater than or equal

to the lower limit, and it must be integral.

length: The number of random integers to generate.

Returns:

The error code IGRAPH_EINVAL is returned in each of the following cases: (1) The given lower limit is greater than the given upper limit, i.e. 1 > h. (2) Assuming that 1 < h and N is the sample size, the above error code is returned if N > |h - 1|, i.e. the sample size exceeds the size of the candidate pool.

Time complexity: according to (Vitter 1987), the expected running time is O(length).

Example 32.2. File examples/simple/igraph_random_sample.c

Random sampling of spatial points

igraph_sample_sphere_surface — Sample points uniformly from the surface of a sphere.

The center of the sphere is at the origin.

Arguments:

dim: The dimension of the random vectors.

n: The number of vectors to sample.

radius: Radius of the sphere, it must be positive.

positive: Whether to restrict sampling to the positive orthant.

res: Pointer to an initialized matrix, the result is stored here, each column will be a sampled

vector. The matrix is resized, as needed.

Returns:

Error code.

Time complexity: O(n*dim*g), where g is the time complexity of generating a standard normal random number.

See also:

 $igraph_sample_sphere_volume(), igraph_sample_dirichlet() for other similar samplers.$

igraph_sample_sphere_volume — Sample points uniformly from the volume of a sphere.

The center of the sphere is at the origin.

Arguments:

dim: The dimension of the random vectors.

n: The number of vectors to sample.

radius: Radius of the sphere, it must be positive.

positive: Whether to restrict sampling to the positive orthant.

res: Pointer to an initialized matrix, the result is stored here, each column will be a sampled

vector. The matrix is resized, as needed.

Returns:

Error code.

Time complexity: O(n*dim*g), where g is the time complexity of generating a standard normal random number.

See also:

igraph_sample_sphere_surface(), igraph_sample_dirichlet() for other similar samplers.

igraph_sample_dirichlet — Sample points from a Dirichlet distribution.

Arguments:

n: The number of vectors to sample.

alpha: The parameters of the Dirichlet distribution. They must be positive. The length of this

vector gives the dimension of the generated samples.

res: Pointer to an initialized matrix, the result is stored here, one sample in each column. It will

be resized, as needed.

Returns:

Error code.

Time complexity: O(n * dim * g), where dim is the dimension of the sample vectors, set by the length of alpha, and g is the time complexity of sampling from a Gamma distribution.

See also:

 $igraph_sample_sphere_surface() \ and \ igraph_sample_sphere_volume() \ for other methods to sample latent vectors.$

Convex hull of a set of points on a plane

igraph_convex_hull — Determines the convex hull of a given set of points in the 2D plane.

```
igraph_error_t igraph_convex_hull(
    const igraph_matrix_t *data, igraph_vector_int_t *resverts,
    igraph_matrix_t *rescoords
);
```

The convex hull is determined by the Graham scan algorithm. See the following reference for details:

Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. Introduction to Algorithms, Second Edition. MIT Press and McGraw-Hill, 2001. ISBN 0262032937. Pages 949-955 of section 33.3: Finding the convex hull.

Arguments:

data: vector containing the coordinates. The length of the vector must be even, since it

contains X-Y coordinate pairs.

resverts: the vector containing the result, e.g. the vector of vertex indices used as the corners

of the convex hull. Supply NULL here if you are only interested in the coordinates

of the convex hull corners.

rescoords: the matrix containing the coordinates of the selected corner vertices. Supply NULL

here if you are only interested in the vertex indices.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory

Time complexity: $O(n \log(n))$ where n is the number of vertices.

Fitting power-law distributions to empirical data

igraph_plfit_result_t — Result of fitting a power-law distribution to a vector

```
typedef struct igraph_plfit_result_t {
    igraph_bool_t continuous;
    igraph_real_t alpha;
    igraph_real_t xmin;
    igraph_real_t L;
    igraph_real_t D;
    const igraph_vector_t* data;
} igraph_plfit_result_t;
```

This data structure contains the result of <code>igraph_power_law_fit()</code>, which tries to fit a power-law distribution to a vector of numbers. The structure contains the following members:

Values:

continuous: Whether the fitted power-law distribution was continuous or discrete.

The exponent of the fitted power-law distribution. alpha:

The minimum value from which the power-law distribution was fitted. In other xmin:

words, only the values larger than xmin were used from the input vector.

The log-likelihood of the fitted parameters; in other words, the probability of ob-L:

serving the input vector given the parameters.

D: The test statistic of a Kolmogorov-Smirnov test that compares the fitted distribution

with the input vector. Smaller scores denote better fit.

The p-value of the Kolmogorov-Smirnov test; NaN if it has not been calculated yet.

Small p-values (less than 0.05) indicate that the test rejected the hypothesis that the

original data could have been drawn from the fitted power-law distribution.

data: The vector containing the original input data. May not be valid any more if the

caller already destroyed the vector.

igraph_power_law_fit — Fits a power-law distribution to a vector of numbers.

```
igraph_error_t igraph_power_law_fit(
    const igraph_vector_t* data, igraph_plfit_result_t* result,
    igraph_real_t xmin, igraph_bool_t force_continuous
);
```

This function fits a power-law distribution to a vector containing samples from a distribution (that is assumed to follow a power-law of course). In a power-law distribution, it is generally assumed that P(X=x) is proportional to x^{-alpha} , where x is a positive number and alpha is greater than 1. In many real-world cases, the power-law behaviour kicks in only above a threshold value xmin. The goal of this functions is to determine alpha if xmin is given, or to determine xmin and the corresponding value of alpha.

The function uses the maximum likelihood principle to determine alpha for a given xmin; in other words, the function will return the alpha value for which the probability of drawing the given sample is the highest. When xmin is not given in advance, the algorithm will attempt to find the optimal xmin value for which the p-value of a Kolmogorov-Smirnov test between the fitted distribution and the original sample is the largest. The function uses the method of Clauset, Shalizi and Newman to calculate the parameters of the fitted distribution. See the following reference for details:

Aaron Clauset, Cosma R .Shalizi and Mark E.J. Newman: Power-law distributions in empirical data. SIAM Review 51(4):661-703, 2009.

Arguments:

data: vector containing the samples for which a power-law distribution is to be

> fitted. Note that you have to provide the samples, not the probability density function or the cumulative distribution function. For example, if you wish to fit a power-law to the degrees of a graph, you can use the output of igraph_degree directly as an input argument to igraph_pow-

er_law_fit

result: the result of the fitting algorithm. See igraph_plfit_result_t for

> more details. Note that the p-value of the fit is not calculated by default as it is time-consuming; you need to call igraph_plfit_result calculate p value() to calculate the p-value itself

> > 652

p:

xmin: the minimum value in the sample vector where the power-law behaviour

is expected to kick in. Samples smaller than xmin will be ignored by the algorithm. Pass zero here if you want to include all the samples. If xmin is negative, the algorithm will attempt to determine its best value

automatically.

 $force_continuous$: assume that the samples in the data argument come from a continuous

distribution even if the sample vector contains integer values only (by chance). If this argument is false, igraph will assume a continuous distribution if at least one sample is non-integer and assume a discrete distrib-

ution otherwise.

Returns:

Error code: IGRAPH_ENOMEM: not enough memory IGRAPH_EINVAL: one of the arguments is invalid IGRAPH_EOVERFLOW: overflow during the fitting process IGRAPH_EUNDERFLOW: underflow during the fitting process IGRAPH_FAILURE: the underlying algorithm signaled a failure without returning a more specific error code

Time complexity: in the continuous case, $O(n \log(n))$ if xmin is given. In the discrete case, the time complexity is dominated by the complexity of the underlying L-BFGS algorithm that is used to optimize alpha. If xmin is not given, the time complexity is multiplied by the number of unique samples in the input vector (although it should be faster in practice).

Example 32.3. File examples/simple/igraph_power_law_fit.c

igraph_plfit_result_calculate_p_value — Calculates the p-value of a fitted power-law model.

The p-value is calculated by resampling the input data many times in a way that the part below the fitted x_min threshold is resampled from the input data itself, while the part above the fitted x_min threshold is drawn from the fitted power-law function. A Kolmogorov-Smirnov test is then performed for each resampled dataset and its test statistic is compared with the observed test statistic from the original dataset. The fraction of resampled datasets that have a *higher* test statistic is the returned p-value.

Note that the precision of the returned p-value depends on the number of resampling attempts. The number of resampling trials is determined by 0.25 divided by the square of the required precision. For instance, a required precision of 0.01 means that 2500 samples will be drawn.

Arguments:

model: The fitted power-law model from the igraph_power_law_fit() function

result: The calculated p-value is returned here

precision: The desired precision of the p-value. Higher values correspond to longer calculation

time. @return igraph_error_t

Comparing floats with a tolerance

igraph_cmp_epsilon — Compare two double-precision floats with a tolerance.

```
int igraph_cmp_epsilon(double a, double b, double eps);
```

Determines whether two double-precision floats are "almost equal" to each other with a given level of tolerance on the relative error.

Arguments:

a: The first float.

b: The second float.

eps: The level of tolerance on the relative error. The relative error is defined as abs(a-b) / (abs(a) + abs(b)). The two numbers are considered equal if this is less than eps.

Returns:

Zero if the two floats are nearly equal to each other within the given level of tolerance, positive number if the first float is larger, negative number if the second float is larger.

igraph_almost_equals — Compare two double-precision floats with a tolerance.

```
igraph_bool_t igraph_almost_equals(double a, double b, double eps);
```

Determines whether two double-precision floats are "almost equal" to each other with a given level of tolerance on the relative error.

Arguments:

a: The first float.

b: The second float.

eps: The level of tolerance on the relative error. The relative error is defined as abs(a-b) / (abs(a) + abs(b)). The two numbers are considered equal if this is less than eps.

Returns:

True if the two floats are nearly equal to each other within the given level of tolerance, false otherwise.

igraph_complex_almost_equals — Compare two complex numbers with a tolerance.

Determines whether two complex numbers are "almost equal" to each other with a given level of tolerance on the relative error.

Arguments:

- a: The first complex number.
- b: The second complex number.
- eps: The level of tolerance on the relative error. The relative error is defined as abs(a-b) / (abs(a) + abs(b)). The two numbers are considered equal if this is less than eps.

Returns:

True if the two complex numbers are nearly equal to each other within the given level of tolerance, false otherwise.

Chapter 33. Licenses for igraph and this manual

THE GNU GENERAL PUBLIC LICENSE

Copyright © 1989, 1991 Free Software Foundation, Inc. 51 Franklin St, Fifth Floor, Boston, MA 02110-1301 USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software--to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation's software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Library General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps: (1) copyright the software, and (2) offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author's protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors' reputations.

Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone's free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

GNU GENERAL PUBLIC LICENSE

TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

0. This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The "Program", below, refers to any such program or work, and a "work based on the Program" means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term "modification".) Each licensee is addressed as "you".

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

1. You may copy and distribute verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

- 2. You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of Section 1 above, provided that you also meet all of these conditions:
- a. You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.
- b. You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.
- c. If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you provide a warranty) and that users may redistribute the program under these conditions, and telling the user how to view a copy of this License. (Exception: if the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

- 3. You may copy and distribute the Program (or a work based on it, under Section 2) in object code or executable form under the terms of Sections 1 and 2 above provided that you also do one of the following:
- a. Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- b. Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,

c. Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.

- 4. You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.
- 5. You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Program or works based on it.
- 6. Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.
- 7. If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

8. If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may

add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.

9. The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and "any later version", you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

10. If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

NO WARRANTY

11. BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

12. IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the "copyright" line and a pointer to where the full notice is found.

<one line to give the program's name and a brief idea of what it does.>
Copyright (C) <year> <name of author>

This program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by

the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like this when it starts in an interactive mode:

Gnomovision version 69, Copyright (C) year name of author Gnomovision comes with ABSOLUTELY NO WARRANTY; for details type `show w'. This is free software, and you are welcome to redistribute it under certain conditions; type `show c' for details.

The hypothetical commands `show w' and `show c' should show the appropriate parts of the General Public License. Of course, the commands you use may be called something other than `show w' and `show c'; they could even be mouse-clicks or menu items--whatever suits your program.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a "copyright disclaimer" for the program, if necessary. Here is a sample; alter the names:

Yoyodyne, Inc., hereby disclaims all copyright interest in the program `Gnomovision' (which makes passes at compilers) written by James Hacker.

```
<signature of Ty Coon>, 1 April 1989
Ty Coon, President of Vice
```

This General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Library General Public License instead of this License.

The GNU Free Documentation License

Copyright © 2000, 2001, 2002 Free Software Foundation, Inc. 51 Franklin St, Fifth Floor, Boston, MA 02110-1301 USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document "free" in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or noncommercially. Secondarily, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of "copyleft", which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The "Document", below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as "you". You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A "Modified Version" of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A "Secondary Section" is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document's overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The "Invariant Sections" are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The "Cover Texts" are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A "Transparent" copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not "Transparent" is called "Opaque".

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The "Title Page" means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, "Title Page" means the text near the most prominent appearance of the work's title, preceding the beginning of the body of the text.

A section "Entitled XYZ" means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ

stands for a specific section name mentioned below, such as "Acknowledgements", "Dedications", "Endorsements", or "History".) To "Preserve the Title" of such a section when you modify the Document means that it remains a section "Entitled XYZ" according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section

- of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.
- B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
- C. State on the Title page the name of the publisher of the Modified Version, as the publisher.
- D. Preserve all the copyright notices of the Document.
- E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
- F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
- G. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.
- H. Include an unaltered copy of this License.
- I. Preserve the section Entitled "History", Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
- J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
- K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.
- L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
- M.Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.
- N. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.
- O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties--for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by)

any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled "History" in the various original documents, forming one section Entitled "History"; likewise combine any sections Entitled "Acknowledgements", and any sections Entitled "Dedications". You must delete all sections Entitled "Endorsements".

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an "aggregate" if the copyright resulting from the compilation is not used to limit the legal rights of the compilation's users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document's Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in

addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled "Acknowledgements", "Dedications", or "History", the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided for under this License. Any other attempt to copy, modify, sublicense or distribute the Document is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See http://www.gnu.org/copyleft/.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License "or any later version" applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation.

G.1.1 ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

Copyright (c) YEAR YOUR NAME.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. A copy of the license is included in the section entitled "GNU Free Documentation License".

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the "with...Texts." line with this:

with the Invariant Sections being LIST THEIR TITLES, with the Front-Cover Texts being LIST, and with the Back-Cover Texts being LIST.

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

Index

Α

add_edge, 25 add_edges, 26 add vertices, 26 adhesion, 556 adjacency, 208 adjacency spectral embedding, 599 adjacent triangles, 496 adilist, 211 adjlist_clear, 186 adilist destroy, 185 adjlist_get, 185 adjlist_init, 184 adilist init complementer, 185 adjlist init empty, 184 adjlist_simplify, 186 adjlist_size, 186 adjlist sort, 186 all_minimal_st_separators, 560 all_st_cuts, 548 all_st_mincuts, 549 almost_equals, 654 are_connected, 329 arpack function t, 627 arpack options init, 628 arpack options t, 624 arpack_rnsolve, 630 arpack_rssolve, 629 arpack storage destroy, 629 arpack_storage_init, 628 arpack_storage_t, 626 arpack_unpack_complex, 630 articulation_points, 371 ASSERT, 45 assortativity, 425 assortativity_degree, 426 assortativity_nominal, 423 asymmetric_preference_game, 244 atlas, 226 attribute combination, 298 attribute_combination_add, 296 attribute_combination_destroy, 297 attribute_combination_init, 296 attribute_combination_remove, 297 attribute_combination_type_t, 297 attribute_table_t, 292 attribute type t, 295 authority_score, 387 automorphisms, 473 automorphism_group, 474 average_local_efficiency, 359 average_path_length, 344 average_path_length_dijkstra, 345 avg_nearest_neighbor_degree, 437

В

barabasi aging game, 246 barabasi game, 230 betweenness, 377 betweenness_cutoff, 391 betweenness_subset, 392 bfs, 448 bfshandler t, 450 bfs_simple, 449 bibcoupling, 402 biconnected_components, 370 bipartite_game, 633 bipartite_projection, 637 bipartite_projection_size, 636 blas_ddot, 614 blas_dgemm, 615 blas_dgemv, 615 blas_dgemv_array, 616 blas_dnrm2, 614 bliss_info_t, 471 bliss_sh_t, 471 bridges, 372

C

callaway_traits_game, 241 calloc, 47 canonical_permutation, 472 cattribute_EAB, 309 cattribute EABV, 310 cattribute_EAB_set, 318 cattribute_EAB_setv, 323 cattribute_EAN, 308 cattribute_EANV, 308 cattribute_EAN_set, 317 cattribute_EAN_setv, 322 cattribute_EAS, 311 cattribute_EASV, 312 cattribute_EAS_set, 319 cattribute_EAS_setv, 324 cattribute GAB, 301 cattribute_GAB_set, 313 cattribute_GAN, 301 cattribute_GAN_set, 312 cattribute GAS, 302 cattribute_GAS_set, 314 cattribute_has_attr, 300 cattribute_list, 300 cattribute_remove_all, 326 cattribute_remove_e, 326 cattribute_remove_g, 325 cattribute remove v, 325 cattribute_VAB, 304 cattribute_VABV, 305 cattribute_VAB_set, 315 cattribute_VAB_setv, 321 cattribute_VAN, 303 cattribute_VANV, 304

cattribute_VAN_set, 315 copy, 17 cattribute_VAN_setv, 320 coreness, 427 cattribute_VAS, 306 correlated_game, 254 cattribute_VASV, 307 correlated_pair_game, 255 cattribute_VAS_set, 316 count_isomorphisms_vf2, 476 cattribute_VAS_setv, 322 count_multiple, 423 centralization, 394 count subisomorphisms vf2, 481 centralization betweenness, 396 create, 207 centralization betweenness tmax, 399 create bipartite, 632 centralization closeness, 396 centralization closeness tmax, 400 centralization_degree, 395 decompose, 369 centralization_degree_tmax, 398 decompose_destroy, 370 centralization eigenvector centrality, 397 degree, 24 centralization_eigenvector_centrality_tmax, 401 degree_sequence_game, 235 CHECK, 41 DELALL, 327 CHECK_CALLBACK, 42 DELEA, 326 circulant, 227 DELEAS, 327 cited_type_game, 249 delete_edges, 27 citing_cited_type_game, 250 delete_vertices, 27 cliques, 455 delete_vertices_idx, 28 cliques_callback, 456 DELGA, 325 clique_handler_t, 457 DELGAS, 327 clique_number, 463 **DELVA**, 326 clique_size_hist, 455 DELVAS, 327 closeness, 374 density, 435 closeness_cutoff, 388 destroy, 18 clusters, 368 deterministic_optimal_imitation, 258 cmp_epsilon, 654 de_bruijn, 226 cocitation, 402 dfs, 451 cohesion, 557 dfshandler_t, 452 cohesive_blocks, 558 diameter, 347 coloring_greedy_t, 492 diameter_dijkstra, 348 community_eb_get_merges, 580 difference, 607 community_edge_betweenness, 578 dim_select, 601 community_fastgreedy, 581 disjoint union, 603 community_fluid_communities, 584 disjoint_union_many, 603 community_infomap, 587 distances, 330 community_label_propagation, 585 distances_bellman_ford, 332 community_leading_eigenvector, 573 distances_dijkstra, 331 community_leading_eigenvector_callback_t, 576 distances_johnson, 333 community_leiden, 583 diversity, 436 community_multilevel, 582 dominator tree, 545 community_optimal_modularity, 565 dot_product_game, 253 community_spinglass, 570 dqueue_back, 157 community_spinglass_single, 572 dqueue_clear, 156 community_to_membership, 566 dqueue_destroy, 155 community_walktrap, 577 dqueue_empty, 155 compare_communities, 567 dqueue_full, 155 complementer, 607 dqueue_head, 156 complex_almost_equals, 654 dqueue_init, 154 compose, 608 dqueue_pop, 157 connected_components, 367 dqueue_pop_back, 157 connect_neighborhood, 609 dqueue_push, 157 constraint, 383 dqueue_size, 156 contract_vertices, 609 dyad_census, 494 convergence_degree, 388 convex_hull, 650

E	es_patn, 281
EAB, 310	es_range, 280
EABV, 310	es_seq, 289
EAN, 308	es_size, 286
EANV, 309	es_type, 286
	es_vector, 279
EAS, 311	es_vector_copy, 282
EASV, 312	eulerian_cycle, 446
eccentricity, 349	eulerian_path, 447
eccentricity_dijkstra, 350	even_tarjan_reduction, 561
ecount, 18	exit_safelocale, 541
ECOUNT_MAX, 29	expand_path_to_pairs, 29
edge, 19	extended_chordal_ring, 229
edges, 19	extended_enordal_ring, 227
edge_betweenness, 378	_
edge_betweenness_cutoff, 391	F
edge_betweenness_subset, 393	famous, 222
edge_connectivity, 553	FATAL, 44
edge_disjoint_paths, 555	fatal, 45
eigenvector_centrality, 385	FATALF, 44
eit_create, 286	fatalf, 45
eit_destroy, 287	fatal_handler_abort, 44
EIT_END, 288	fatal_handler_t, 44
EIT_GET, 288	feedback_arc_set, 429
EIT_NEXT, 287	FINALLY, 42
EIT_RESET, 288	FINALLY_CLEAN, 43
EIT_SIZE, 288	FINALLY_FREE, 43
empty, 16	forest_fire_game, 239
empty_attrs, 16	free, 48
enter_safelocale, 540	FROM, 20
erdos_renyi_game, 232	from_prufer, 225
ERROR, 40	full, 217
error, 40	full_bipartite, 633
ERRORF, 40	full_citation, 218
errorf, 41	full_multipartite, 218
error_handler_abort, 32	_
error_handler_ignore, 32	G
error_handler_printignore, 32	GAB, 302
error_handler_t, 31	GAN, 301
error_t, 32	GAS, 302
error_type_t, 32	generalized_petersen, 228
ess_1, 283	get_adjacency, 438
ess_all, 282	get_adjacency_sparse, 440
ess_none, 283	get_all_eids_between, 23
ess_range, 284	get_all_shortest_paths, 340
ess_seq, 289	get_all_shortest_paths_dijkstra, 341
ess_vector, 284	get_all_simple_paths, 344
establishment_game, 242	get_edgelist, 441
es_1, 279	get_edgenst, 441 get_eid, 21
es_all, 277	get_eids, 22
es_as_vector, 284	<u> </u>
es_copy, 285	get_incidence, 635
es_destroy, 285	get_isomorphisms_vf2, 477
es_incident, 278	get_isomorphisms_vf2_callback, 478
es_is_all, 285	get_k_shortest_paths, 343
es_none, 278	get_laplacian, 419
es_pairs, 280	get_laplacian_sparse, 419
es_pairs_small, 281	get_shortest_path, 335
<u>-</u>	get shortest paths, 334

get_shortest_paths_bellman_ford, 338	igraph_integer_t, 15
get_shortest_paths_dijkstra, 336	igraph_real_t, 15
get_shortest_path_bellman_ford, 339	IGRAPH_UINT_MAX, 15
get_shortest_path_dijkstra, 337	IGRAPH_UINT_MIN, 15
get_sparsemat, 444	igraph_uint_t, 15
get_stochastic, 440	incidence, 634
get_stochastic_sparse, 441	incident, 24
get_stochastic_sparsemat, 444	inclist_clear, 188
get_subisomorphisms_vf2, 482	inclist_destroy, 187
get_subisomorphisms_vf2_callback, 483	inclist_get, 188
get_widest_path, 353	inclist_init, 187
get_widest_paths, 354	inclist_size, 188
girth, 349	independence_number, 467
global_efficiency, 357	independent_vertex_sets, 465
gomory_hu_tree, 551	induced_subgraph, 610
graphlets, 589	intersection, 605
graphlets_candidate_basis, 590	intersection_many, 606
graphlets_project, 590	invalidate_cache, 29
graph_center, 351	isoclass, 487
graph_count, 489	isoclass_create, 488
grg_game, 230	isoclass_subgraph, 487
growing_random_game, 241	isocompat_t, 480
	isohandler_t, 479
Н	isomorphic, 469
harmonic_centrality, 376	isomorphic_34, 490
harmonic_centrality_cutoff, 390	isomorphic_bliss, 472
has_multiple, 422	isomorphic_function_vf2, 485
heap_clear, 159	isomorphic_vf2, 475
heap_delete_top, 160	is_acyclic, 442
heap_destroy, 159	is_bigraphical, 374
heap_empty, 159	is_bipartite, 638
heap_init, 158	is_chordal, 431
heap_init_array, 158	is_connected, 368
heap_push, 160	is_dag, 428
heap_reserve, 161	is_directed, 18
heap_size, 161	is_eulerian, 446
heap_top, 160	is_forest, 413
hrg_consensus, 595	is_graphical, 372
hrg_create, 597	is_loop, 421
hrg_dendrogram, 596	is_matching, 432
hrg_destroy, 593	is_maximal_matching, 432
hrg_fit, 594	is_minimal_separator, 559
hrg_game, 596	is_multiple, 422
hrg_init, 593	is_mutual, 437
hrg_predict, 597	is_perfect, 493
hrg_resize, 594	is_same_graph, 30
hrg_sample, 595	is_separator, 559
hrg_size, 593	is_simple, 421
hrg_t, 592	is_tree, 412
hsbm_game, 251	
hsbm_list_game, 252	K
hub_score, 386	kary_tree, 215
	kautz, 227
1	k_regular_game, 236
	0
igraph_bool_t, 15	ı
IGRAPH_INTEGER_MAX, 15	-
IGRAPH_INTEGER_MIN, 15	lapack_dgeev, 620

local_scan_0_them, 364 lapack_dgeevx, 621 lapack_dgesv, 618 local_scan_1_ecount, 363 lapack_dgetrf, 617 local_scan_1_ecount_them, 364 lapack_dgetrs, 617 local_scan_k_ecount, 363 lapack_dsyevr, 619 local_scan_k_ecount_them, 365 laplacian, 445 local_scan_neighborhood_ecount, 366 laplacian normalization t, 420 local scan subset ecount, 366 laplacian spectral embedding, 600 M largest cliques, 457 largest independent vertex sets, 466 malloc, 47 largest weighted cliques, 464 MATRIX, 94 lastcit_game, 248 matrix_add, 102 lattice, 256 matrix_add_cols, 116 layout bipartite, 505 matrix_add_constant, 102 layout_circle, 502 matrix add rows, 116 layout_davidson_harel, 512 matrix_all_almost_e, 106 layout drl, 508 matrix_all_e, 106 layout_drl_3d, 509 matrix all g, 107 layout_drl_default_t, 507 matrix all ge, 107 layout_drl_options_init, 508 matrix_all_l, 106 layout_drl_options_t, 506 matrix_all_le, 107 layout_fruchterman_reingold, 509 matrix_as_sparsemat, 149 layout_fruchterman_reingold_3d, 522 matrix_capacity, 112 layout_gem, 512 matrix_cbind, 109 layout_graphopt, 504 matrix colsum, 105 layout_grid, 503 matrix complex all almost e, 119 layout_grid_3d, 522 matrix complex create, 118 layout kamada kawai, 511 matrix complex create polar, 119 layout_kamada_kawai_3d, 523 matrix complex imag, 117 layout_lgl, 514 matrix_complex_real, 117 layout_mds, 514 matrix_complex_realimag, 118 layout_merge_dla, 525 matrix_complex_zapsmall, 119 layout_random, 502 matrix contains, 114 layout_random_3d, 521 matrix_copy, 120 layout_reingold_tilford, 515 matrix_copy_to, 97 layout_reingold_tilford_circular, 516 matrix destroy, 93 layout_sphere, 521 matrix div elements, 104 layout_star, 503 matrix_e, 120 layout_sugiyama, 518 matrix_empty, 111 layout_umap, 519 matrix e ptr, 120 layout_umap_3d, 524 matrix fill, 94 lazy_adjlist_clear, 190 matrix_get, 94 lazy_adjlist_destroy, 189 matrix_get_col, 98 lazy_adjlist_get, 189 matrix_get_ptr, 95 lazy_adjlist_init, 189 matrix_get_row, 98 lazy_adjlist_size, 190 matrix init, 93 lazy_inclist_clear, 192 matrix init copy, 93 lazy_inclist_destroy, 191 matrix isnull, 112 lazy_inclist_get, 191 matrix_is_symmetric, 113 lazy_inclist_init, 190 matrix_max, 109 lazy_inclist_size, 192 matrix maxdifference, 114 lcf, 224 matrix_min, 109 lcf_vector, 225 matrix_minmax, 110 le_community_to_membership, 576 matrix_mul_elements, 103 linegraph, 611 matrix ncol, 113 list_triangles, 496 matrix nrow, 113 local_efficiency, 358 matrix null, 94 local_scan_0, 362 matrix prod, 104

neighborhood_graphs, 361 matrix_rbind, 108 matrix_remove_col, 117 neighborhood_size, 360 matrix remove row, 116 neighbors, 23 matrix_resize, 115 matrix_resize_min, 115 0 matrix_rowsum, 105 OTHER, 21 matrix scale, 102 matrix_search, 114 P matrix select cols, 101 pagerank, 379 matrix select rows, 100 pagerank_algo_t, 378 matrix select rows cols, 101 path_length_hist, 346 matrix_set, 95 permute_vertices, 489 matrix_set_col, 99 personalized_pagerank, 380 matrix set row, 99 personalized_pagerank_vs, 381 matrix_size, 112 plfit_result_calculate_p_value, 653 matrix_sub, 103 plfit_result_t, 651 matrix sum, 104 power_law_fit, 652 matrix_swap, 98 preference_game, 243 matrix_swap_cols, 100 PROGRESS, 641 matrix_swap_rows, 100 progress, 641 matrix transpose, 105 progressf, 642 matrix_update, 97 progress_handler_stderr, 641 matrix_view, 96 progress_handler_t, 640 matrix_view_from_vector, 96 pseudo_diameter, 352 matrix_which_max, 110 psumtree_destroy, 193 matrix_which_min, 110 psumtree_get, 194 matrix_which_minmax, 111 psumtree_init, 193 matrix zapsmall, 108 psumtree_search, 194 maxdegree, 383 psumtree_size, 193 maxflow, 543 psumtree_sum, 194 maxflow_stats_t, 546 psumtree_update, 195 maxflow_value, 544 maximal_cliques, 458 R maximal_cliques_callback, 462 maximal_cliques_count, 459 radius, 351 maximal_cliques_file, 460 random edge walk, 454 maximal_cliques_hist, 461 random sample, 648 maximal_cliques_subset, 460 random_spanning_tree, 411 maximal_independent_vertex_sets, 467 random walk, 453 maximum_bipartite_matching, 433 read graph dimacs, 542 maximum_cardinality_search, 430 read_graph_dimacs_flow, 531 mincut, 550 read_graph_dl, 539 mincut_value, 551 read_graph_edgelist, 527 minimum_size_separators, 561 read_graph_gml, 535 minimum_spanning_tree, 409 read_graph_graphdb, 533 minimum_spanning_tree_prim, 410 read_graph_graphml, 534 minimum_spanning_tree_unweighted, 410 read graph 1gl, 530 modularity, 563 read graph ncol, 528 modularity_matrix, 564 read_graph_pajek, 537 moran process, 259 realize_degree_sequence, 219 motifs_handler_t, 500 realloc, 48 motifs_randesu, 497 recent_degree_aging_game, 247 motifs_randesu_callback, 499 recent_degree_game, 245 motifs_randesu_estimate, 498 reciprocity, 435 motifs_randesu_no, 498 regular tree, 216 reindex membership, 567 Ν reverse edges, 612 neighborhood, 361 rewire, 240

rewire_directed_edges, 234	similarity_dice_es, 407
rewire_edges, 233	similarity_dice_pairs, 406
ring, 214	similarity_inverse_log_weighted, 408
rngtype_glibc2, 203	similarity_jaccard, 403
rngtype_mt19937, 203	similarity_jaccard_es, 405
rngtype_pcg32, 204	similarity_jaccard_pairs, 404
rngtype_pcg64, 204	simple_interconnected_islands_game, 255
rng_bits, 198	simplify, 611
rng_default, 196	simplify_and_colorize, 490
rng_destroy, 197	sir, 264
rng_get_binom, 201	sir_destroy, 265
6 _ 6	sir_t, 264
rng_get_exp, 200	
rng_get_gamma, 201	small, 207
rng_get_geom, 202	spanner, 329
rng_get_integer, 199	sparsemat, 149
rng_get_normal, 200	sparsemat_add, 135
rng_get_pois, 202	sparsemat_add_cols, 136
rng_get_unif, 199	sparsemat_add_rows, 136
rng_get_unif01, 199	sparsemat_arpack_rnsolve, 148
rng_init, 197	sparsemat_arpack_rssolve, 147
rng_max, 198	sparsemat_as_matrix, 150
rng_name, 198	sparsemat_cholsol, 144
rng_seed, 197	sparsemat_colsums, 131
rng_set_default, 196	sparsemat_compress, 140
roots_for_tree_layout, 517	sparsemat_copy, 150
roulette_wheel_imitation, 261	sparsemat_count_nonzero, 130
running_mean, 647	sparsemat_count_nonzerotol, 130
	sparsemat_destroy, 124
S	sparsemat_diag, 151
sample_dirichlet, 650	sparsemat_droptol, 133
sample_sphere_surface, 649	sparsemat_dropzeros, 133
sample_sphere_volume, 649	sparsemat_dupl, 141
sbm_game, 251	sparsemat_entry, 132
SETEAB, 319	sparsemat_eye, 151
SETEABV, 324	sparsemat_fkeep, 132
SETEAN, 318	sparsemat_gaxpy, 136
SETEANV, 323	sparsemat_get, 127
SETEAS, 320	sparsemat_getelements, 128
SETEASV, 324	sparsemat_getelements_sorted, 128
SETGAB, 313	sparsemat_index, 125
SETGAN, 313	sparsemat_init, 121
SETGAS, 314	sparsemat_init_copy, 122
SETVAB, 316	sparsemat_init_diag, 122
SETVABV, 321	sparsemat_init_eye, 123
SETVAN, 315	sparsemat_is_cc, 127
SETVANV, 320	sparsemat_is_symmetric, 127
	sparsemat_is_triplet, 126
SETVASY 222	sparsemat_iterator_col, 139
SETVASV, 322	sparsemat_iterator_end, 138
set_attribute_table, 294	sparsemat_iterator_get, 139
set_error_handler, 39	sparsemat_iterator_idx, 140
set_progress_handler, 640	sparsemat_iterator_init, 138
set_status_handler, 644	sparsemat_iterator_next, 140
set_warning_handler, 36	sparsemat_iterator_reset, 138
shortest_paths, 442	sparsemat_iterator_row, 139
shortest_paths_bellman_ford, 443	sparsemat_lsolve, 142
shortest_paths_dijkstra, 443	sparsemat_ltsolve, 142
shortest_paths_johnson, 443	sparsemat_lu, 145
similarity dice, 405	Sparseiliat Iu. 143

sparsemat_luresol, 146 strvector_destroy, 162 sparsemat_lusol, 144 strvector_get, 163 sparsemat max, 129 strvector init, 161 sparsemat_min, 129 strvector_init_copy, 162 sparsemat_minmax, 130 strvector_merge, 166 sparsemat_multiply, 135 strvector_push_back, 164 sparsemat_ncol, 126 strvector_push_back len, 164 sparsemat_nonzero_storage, 131 strvector remove, 165 sparsemat nrow, 126 strvector remove section, 165 sparsemat numeric destroy, 147 strvector reserve, 167 sparsemat_permute, 134 strvector_resize, 167 sparsemat_print, 150 strvector_resize_min, 168 sparsemat_qr, 145 strvector_set, 163 sparsemat grresol, 146 strvector set2, 169 sparsemat_realloc, 123 strvector_set_len, 164 sparsemat_resize, 137 strvector_size, 168 sparsemat rowsums, 131 st edge connectivity, 552 sparsemat_scale, 134 st mincut, 546 sparsemat_sort, 137 st_mincut_value, 547 sparsemat_symblu, 141 st_vertex_connectivity, 553 sparsemat_symbolic_destroy, 147 subcomponent, 367 sparsemat_symbqr, 142 subgraph_edges, 612 subisomorphic, 470 sparsemat_transpose, 134 sparsemat_type, 126 subisomorphic_function_vf2, 485 sparsemat_usolve, 143 subisomorphic_lad, 486 subisomorphic_vf2, 480 sparsemat_utsolve, 143 sparsemat_view, 124 symmetric_tree, 215 sparse adjacency, 210 sparse_weighted_adjacency, 211 Т split_join_distance, 569 THREAD_SAFE, 639 square_lattice, 213 TO, 20 stack_clear, 153 topological_sorting, 428 stack_destroy, 152 to_directed, 417 stack_empty, 152 to_prufer, 413 stack_init, 151 to_undirected, 418 stack_pop, 154 transitivity_avglocal_undirected, 416 stack_push, 153 transitivity barrat, 416 stack_reserve, 152 transitivity_local_undirected, 415 stack_size, 153 transitivity_undirected, 414 stack_top, 154 tree, 256 star, 211 tree_game, 253 static_fitness_game, 237 triad_census, 494 static_power_law_game, 238 trussness, 427 STATUS, 644 turan, 219 status, 645 STATUSF, 645 U statusf, 646 unfold_tree, 434 status_handler_stderr, 644 union, 604 status_handler_t, 643 union many, 605 stochastic imitation, 262 STR, 162 strength, 384 strerror, 36 VAB, 305 strvector_add, 169 VABV, 306 strvector_append, 165 VAN, 303 strvector_capacity, 168 VANV, 304 strvector_clear, 166 VAS, 306 strvector_copy, 169 VASV, 307

vcount, 18 vector_list_get_ptr, 172 VCOUNT MAX, 28 vector_list_init, 171 VECTOR, 54 vector list insert, 177 vector_add, 60 vector_list_insert_copy, 178 vector_list_insert_new, 178 vector_add_constant, 59 vector_all_almost_e, 62 vector_list_permute, 181 vector all e, 62 vector list pop back, 177 vector_all_g, 63 vector_list_push_back, 176 vector_all_ge, 64 vector_list_push_back_copy, 176 vector all 1, 63 vector list push back new, 177 vector all le, 64 vector list remove, 179 vector_append, 57 vector_list_remove_fast, 179 vector_binsearch, 73 vector_list_replace, 173 vector binsearch2, 74 vector list reserve, 175 vector_binsearch_slice, 74 vector_list_resize, 175 vector_capacity, 70 vector_list_set, 172 vector_clear, 75 vector list size, 174 vector_colex_cmp, 66 vector_list_sort, 182 vector_colex_cmp_untyped, 67 vector_list_sort_ind, 182 vector_complex_all_almost_e, 80 vector_list_swap, 182 vector_complex_create, 79 vector_list_swap_elements, 183 vector_complex_create_polar, 80 vector_list_tail_ptr, 172 vector_complex_imag, 78 vector_max, 67 vector_complex_real, 78 vector_maxdifference, 71 vector_complex_realimag, 79 vector_min, 67 vector_complex_zapsmall, 80 vector_minmax, 68 vector_contains, 72 vector_mul, 61 vector_copy, 90 vector null, 53 vector_copy_to, 56 vector_permute, 59 vector_destroy, 52 vector_pop_back, 77 vector_difference_sorted, 83 vector_prod, 71 vector_div, 61 vector_ptr_clear, 86 vector_e, 91 vector_ptr_copy, 92 vector_empty, 69 vector_ptr_destroy, 84 vector_e_ptr, 91 vector_ptr_destroy_all, 85 vector_fill, 53 vector_ptr_e, 92 vector_floor, 62 vector_ptr_free_all, 85 vector_get, 54 vector_ptr_get, 87 vector_ptr_get_item_destructor, 90 vector_get_ptr, 55 vector_init, 51 vector_ptr_init, 83 vector_init_array, 51 vector_ptr_init_copy, 84 vector_init_copy, 52 vector_ptr_insert, 87 vector_init_range, 52 vector_ptr_permute, 89 vector_init_seq, 91 vector_ptr_pop_back, 86 vector_insert, 77 vector_ptr_push_back, 86 vector_intersect_sorted, 82 vector_ptr_resize, 88 vector_isininterval, 71 vector_ptr_set, 87 vector_ptr_set_item_destructor, 90 vector_is_any_nan, 72 VECTOR_PTR_SET_ITEM_DESTRUCTOR, 90 vector_is_nan, 72 vector_lex_cmp, 65 vector_ptr_size, 85 vector_lex_cmp_untyped, 65 vector_ptr_sort, 88 vector_list_capacity, 174 vector_ptr_sort_ind, 88 vector_list_clear, 174 vector_push_back, 76 vector_list_destroy, 171 vector_qsort_ind, 82 vector_list_discard, 180 vector_range, 53 vector_list_discard_back, 180 vector_remove, 77 vector_list_discard_fast, 181 vector_remove_section, 78

vector_reserve, 75

vector_list_empty, 173

vector_resize, 75 vector_resize_min, 76 vector reverse, 58 vector_reverse_sort, 81 vector_scale, 60 vector_search, 73 vector set, 55 vector_shuffle, 58 vector_size, 70 vector sort, 81 vector_sub, 60 vector_sum, 70 vector_swap, 57 vector_swap_elements, 58 vector_tail, 55 vector_update, 56 vector view, 56 vector_which_max, 68 vector_which_min, 68 vector_which_minmax, 69 vector_zapsmall, 64 version, 647 vertex_coloring_greedy, 492 vertex_connectivity, 554 vertex_disjoint_paths, 556 vertex_path_from_edge_path, 353 vit_create, 274 vit destroy, 275 VIT_END, 276 VIT_GET, 277 VIT_NEXT, 276 VIT_RESET, 277 VIT_SIZE, 276 vss_1, 273 vss_all, 272 vss_none, 273 vss_range, 274 vss_seq, 290 vss_vector, 274 vs_1, 268 vs_adj, 267 vs_all, 266 vs_copy, 271 vs_destroy, 271 vs_is_all, 272 vs_nonadj, 267 vs_none, 268 vs_range, 270 vs_seq, 290 vs_size, 272 vs_type, 272 vs_vector, 269 vs_vector_copy, 270 vs_vector_small, 269

WARNING, 37 warning, 37

WARNINGF, 37 warningf, 38 warning_handler_ignore, 38 warning_handler_print, 38 warning_handler_t, 36 watts_strogatz_game, 232 weighted adjacency, 209 weighted_cliques, 463 weighted_clique_number, 465 wheel, 212 widest_path_widths_dijkstra, 355 widest_path_widths_floyd_warshall, 356 write_graph_dimacs, 542 write_graph_dimacs_flow, 532 write_graph_dot, 540 write_graph_edgelist, 527 write_graph_gml, 536 write_graph_graphml, 534

write_graph_lgl, 531

write_graph_ncol, 529

write_graph_pajek, 538